


GNU Octave

A high-level interactive language for numerical computations

Edition 5 for Octave version 5.0.91

January 2019



Free Your Numbers

**John W. Eaton
David Bateman
Søren Hauberg
Rik Wehbring**

Copyright © 1996, 1997, 1999, 2000, 2001, 2002, 2005, 2006, 2007, 2011, 2013, 2015, 2016, 2017, 2018 John W. Eaton.

This is the fifth edition of the Octave documentation, and is consistent with version 5.0.91 of Octave.

Permission is granted to make and distribute verbatim copies of this manual provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this manual under the conditions for verbatim copying, provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this manual into another language, under the same conditions as for modified versions.

Portions of this document have been adapted from the `gawk`, `readline`, `gcc`, and C library manuals, published by the Free Software Foundation, Inc., 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301–1307, USA.

Table of Contents

Preface	1
Acknowledgements	1
Citing Octave in Publications	5
How You Can Contribute to Octave.....	6
Distribution	6
1 A Brief Introduction to Octave.....	7
1.1 Running Octave	7
1.2 Simple Examples	7
1.2.1 Elementary Calculations	7
1.2.2 Creating a Matrix	8
1.2.3 Matrix Arithmetic	8
1.2.4 Solving Systems of Linear Equations	8
1.2.5 Integrating Differential Equations	9
1.2.6 Producing Graphical Output	10
1.2.7 Help and Documentation	10
1.2.8 Editing What You Have Typed	10
1.3 Conventions.....	11
1.3.1 Fonts	11
1.3.2 Evaluation Notation	11
1.3.3 Printing Notation	11
1.3.4 Error Messages	12
1.3.5 Format of Descriptions	12
1.3.5.1 A Sample Function Description	12
1.3.5.2 A Sample Command Description	13
2 Getting Started	15
2.1 Invoking Octave from the Command Line.....	15
2.1.1 Command Line Options	15
2.1.2 Startup Files	18
2.2 Quitting Octave	19
2.3 Commands for Getting Help	20
2.4 Command Line Editing.....	25
2.4.1 Cursor Motion.....	25
2.4.2 Killing and Yanking	26
2.4.3 Commands for Changing Text	27
2.4.4 Letting Readline Type for You.....	27
2.4.5 Commands for Manipulating the History.....	28
2.4.6 Customizing <code>readline</code>	32
2.4.7 Customizing the Prompt	32
2.4.8 Diary and Echo Commands.....	34
2.5 How Octave Reports Errors	35

2.6 Executable Octave Programs	36
2.7 Comments in Octave Programs	37
2.7.1 Single Line Comments	37
2.7.2 Block Comments	37
2.7.3 Comments and the Help System	38
3 Data Types.....	39
3.1 Built-in Data Types	39
3.1.1 Numeric Objects	42
3.1.2 Missing Data	43
3.1.3 String Objects.....	43
3.1.4 Data Structure Objects.....	43
3.1.5 Cell Array Objects	44
3.2 User-defined Data Types	44
3.3 Object Sizes.....	44
4 Numeric Data Types	47
4.1 Matrices	48
4.1.1 Empty Matrices	51
4.2 Ranges.....	52
4.3 Single Precision Data Types.....	53
4.4 Integer Data Types.....	54
4.4.1 Integer Arithmetic.....	57
4.5 Bit Manipulations	58
4.6 Logical Values.....	60
4.7 Promotion and Demotion of Data Types	62
4.8 Predicates for Numeric Objects	62
5 Strings.....	67
5.1 Escape Sequences in String Constants	67
5.2 Character Arrays.....	68
5.3 Creating Strings.....	70
5.3.1 Concatenating Strings	70
5.3.2 Converting Numerical Data to Strings	73
5.4 Comparing Strings	76
5.5 Manipulating Strings	77
5.6 String Conversions	92
5.7 Character Class Functions	98
6 Data Containers	101
6.1 Structures.....	101
6.1.1 Basic Usage and Examples	101
6.1.2 Structure Arrays	105
6.1.3 Creating Structures	106
6.1.4 Manipulating Structures	109
6.1.5 Processing Data in Structures	113

6.2	containers.Map	114
6.3	Cell Arrays	114
6.3.1	Basic Usage of Cell Arrays.....	114
6.3.2	Creating Cell Arrays	116
6.3.3	Indexing Cell Arrays	120
6.3.4	Cell Arrays of Strings	123
6.3.5	Processing Data in Cell Arrays	124
6.4	Comma Separated Lists	124
6.4.1	Comma Separated Lists Generated from Cell Arrays	125
6.4.2	Comma Separated Lists Generated from Structure Arrays	126
7	Variables	127
7.1	Global Variables	129
7.2	Persistent Variables	130
7.3	Status of Variables	132
8	Expressions	139
8.1	Index Expressions	139
8.1.1	Advanced Indexing	141
8.2	Calling Functions	145
8.2.1	Call by Value	146
8.2.2	Recursion	147
8.2.3	Access via Handle	148
8.3	Arithmetic Operators	149
8.4	Comparison Operators	152
8.5	Boolean Expressions	153
8.5.1	Element-by-element Boolean Operators	153
8.5.2	Short-circuit Boolean Operators	155
8.6	Assignment Expressions	156
8.7	Increment Operators	159
8.8	Operator Precedence	159
9	Evaluation	161
9.1	Calling a Function by its Name	162
9.2	Evaluation in a Different Context	163
10	Statements	165
10.1	The if Statement	165
10.2	The switch Statement	167
10.2.1	Notes for the C Programmer	168
10.3	The while Statement	169
10.4	The do-until Statement	170
10.5	The for Statement	170
10.5.1	Looping Over Structure Elements	171
10.6	The break Statement	172
10.7	The continue Statement	173

10.8	The unwind_protect Statement	174
10.9	The try Statement	174
10.10	Continuation Lines.....	175

11 Functions and Scripts..... 177

11.1	Introduction to Function and Script Files.....	177
11.2	Defining Functions	177
11.3	Multiple Return Values	180
11.4	Variable-length Argument Lists	188
11.5	Ignoring Arguments	190
11.6	Variable-length Return Lists	191
11.7	Returning from a Function	192
11.8	Default Arguments.....	193
11.9	Function Files.....	193
11.9.1	Manipulating the Load Path.....	196
11.9.2	Subfunctions	199
11.9.3	Private Functions	200
11.9.4	Nested Functions.....	200
11.9.5	Overloading and Autoloading.....	202
11.9.6	Function Locking.....	204
11.9.7	Function Precedence	205
11.10	Script Files	205
11.10.1	Publish Octave Script Files.....	207
11.10.2	Publishing Markup.....	209
11.10.2.1	Using Publishing Markup in Script Files.....	209
11.10.2.2	Text Formatting	210
11.10.2.3	Sections.....	210
11.10.2.4	Preformatted Code.....	211
11.10.2.5	Preformatted Text	211
11.10.2.6	Bulleted Lists	212
11.10.2.7	Numbered Lists	212
11.10.2.8	Including File Content	212
11.10.2.9	Including Graphics	212
11.10.2.10	Including URLs	213
11.10.2.11	Mathematical Equations	213
11.10.2.12	HTML Markup	213
11.10.2.13	LaTeX Markup	214
11.11	Function Handles, Anonymous Functions, Inline Functions.....	214
11.11.1	Function Handles	214
11.11.2	Anonymous Functions.....	215
11.11.3	Inline Functions.....	216
11.12	Commands	217
11.13	Organization of Functions Distributed with Octave	218

12 Errors and Warnings	221
12.1 Handling Errors.....	221
12.1.1 Raising Errors	221
12.1.2 Catching Errors	224
12.1.3 Recovering From Errors	227
12.2 Handling Warnings.....	227
12.2.1 Issuing Warnings.....	228
12.2.2 Enabling and Disabling Warnings.....	235
13 Debugging	237
13.1 Entering Debug Mode.....	237
13.2 Leaving Debug Mode.....	238
13.3 Breakpoints.....	238
13.4 Debug Mode.....	242
13.5 Call Stack	243
13.6 Profiling	244
13.7 Profiler Example	246
14 Input and Output	251
14.1 Basic Input and Output.....	251
14.1.1 Terminal Output.....	251
14.1.1.1 Paging Screen Output	254
14.1.2 Terminal Input	256
14.1.3 Simple File I/O	258
14.1.3.1 Saving Data on Unexpected Exits.....	271
14.2 C-Style I/O Functions.....	273
14.2.1 Opening and Closing Files.....	273
14.2.2 Simple Output	275
14.2.3 Line-Oriented Input	276
14.2.4 Formatted Output	277
14.2.5 Output Conversion for Matrices	279
14.2.6 Output Conversion Syntax	279
14.2.7 Table of Output Conversions	280
14.2.8 Integer Conversions	281
14.2.9 Floating-Point Conversions	282
14.2.10 Other Output Conversions	282
14.2.11 Formatted Input	283
14.2.12 Input Conversion Syntax	284
14.2.13 Table of Input Conversions	285
14.2.14 Numeric Input Conversions	286
14.2.15 String Input Conversions	286
14.2.16 Binary I/O	286
14.2.17 Temporary Files	289
14.2.18 End of File and Errors	290
14.2.19 File Positioning	292

15 Plotting	295
15.1 Introduction to Plotting	295
15.2 High-Level Plotting	295
15.2.1 Two-Dimensional Plots	295
15.2.1.1 Axis Configuration	323
15.2.1.2 Two-dimensional Function Plotting	327
15.2.1.3 Two-dimensional Geometric Shapes	331
15.2.2 Three-Dimensional Plots	331
15.2.2.1 Aspect Ratio	358
15.2.2.2 Three-dimensional Function Plotting	359
15.2.2.3 Three-dimensional Geometric Shapes	362
15.2.3 Plot Annotations	363
15.2.4 Multiple Plots on One Page	371
15.2.5 Multiple Plot Windows	372
15.2.6 Manipulation of Plot Objects	372
15.2.7 Manipulation of Plot Windows	374
15.2.8 Use of the <code>interpreter</code> Property	378
15.2.8.1 Degree Symbol	381
15.2.9 Printing and Saving Plots	381
15.2.10 Interacting with Plots	392
15.2.11 Test Plotting Functions	392
15.3 Graphics Data Structures	393
15.3.1 Introduction to Graphics Structures	393
15.3.2 Graphics Objects	395
15.3.2.1 Creating Graphics Objects	395
15.3.2.2 Handle Functions	399
15.3.3 Graphics Object Properties	405
15.3.3.1 Root Figure Properties	405
15.3.3.2 Figure Properties	407
15.3.3.3 Axes Properties	412
15.3.3.4 Line Properties	419
15.3.3.5 Text Properties	422
15.3.3.6 Image Properties	425
15.3.3.7 Patch Properties	427
15.3.3.8 Surface Properties	430
15.3.3.9 Light Properties	434
15.3.3.10 Uimenu Properties	436
15.3.3.11 Uibuttongroup Properties	437
15.3.3.12 Uicontextmenu Properties	440
15.3.3.13 Uipanel Properties	441
15.3.3.14 Uicontrol Properties	444
15.3.3.15 Uitable Properties	446
15.3.3.16 Uitoolbar Properties	449
15.3.3.17 Uipushtool Properties	450
15.3.3.18 Uitoggletool Properties	452
15.3.4 Searching Properties	453
15.3.5 Managing Default Properties	455
15.4 Advanced Plotting	456

15.4.1	Colors	456
15.4.2	Line Styles.....	456
15.4.3	Marker Styles	457
15.4.4	Callbacks	457
15.4.5	Application-defined Data	459
15.4.6	Object Groups	460
15.4.6.1	Data Sources in Object Groups.....	465
15.4.6.2	Area Series.....	465
15.4.6.3	Bar Series	466
15.4.6.4	Contour Groups.....	467
15.4.6.5	Error Bar Series	468
15.4.6.6	Line Series	468
15.4.6.7	Quiver Group	469
15.4.6.8	Scatter Group	470
15.4.6.9	Stair Group	470
15.4.6.10	Stem Series	471
15.4.6.11	Surface Group.....	472
15.4.7	Transform Groups.....	472
15.4.8	Graphics Toolkits	473
15.4.8.1	Customizing Toolkit Behavior	473
15.4.8.2	Hardware vs. Software Rendering.....	474
16	Matrix Manipulation	475
16.1	Finding Elements and Checking Conditions.....	475
16.2	Rearranging Matrices	479
16.3	Special Utility Matrices	489
16.4	Famous Matrices	500
17	Arithmetic	509
17.1	Exponents and Logarithms.....	509
17.2	Complex Arithmetic	511
17.3	Trigonometry	512
17.4	Sums and Products	516
17.5	Utility Functions	518
17.6	Special Functions	526
17.7	Rational Approximations	538
17.8	Coordinate Transformations	539
17.9	Mathematical Constants	540
18	Linear Algebra	545
18.1	Techniques Used for Linear Algebra	545
18.2	Basic Matrix Functions	545
18.3	Matrix Factorizations	555
18.4	Functions of a Matrix	568
18.5	Specialized Solvers	570

19 Vectorization and Faster Code Execution	583
19.1 Basic Vectorization	583
19.2 Broadcasting	585
19.2.1 Broadcasting and Legacy Code	588
19.3 Function Application	588
19.4 Accumulation	593
19.5 JIT Compiler	595
19.6 Miscellaneous Techniques	596
19.7 Examples	598
20 Nonlinear Equations	599
20.1 Solvers	599
20.2 Minimizers	603
21 Diagonal and Permutation Matrices	609
21.1 Creating and Manipulating Diagonal/Permutation Matrices	609
21.1.1 Creating Diagonal Matrices	610
21.1.2 Creating Permutation Matrices	610
21.1.3 Explicit and Implicit Conversions	611
21.2 Linear Algebra with Diagonal/Permutation Matrices	612
21.2.1 Expressions Involving Diagonal Matrices	612
21.2.2 Expressions Involving Permutation Matrices	613
21.3 Functions That Are Aware of These Matrices	614
21.3.1 Diagonal Matrix Functions	614
21.3.2 Permutation Matrix Functions	614
21.4 Examples of Usage	614
21.5 Differences in Treatment of Zero Elements	615
22 Sparse Matrices	617
22.1 Creation and Manipulation of Sparse Matrices	617
22.1.1 Storage of Sparse Matrices	617
22.1.2 Creating Sparse Matrices	618
22.1.3 Finding Information about Sparse Matrices	624
22.1.4 Basic Operators and Functions on Sparse Matrices	627
22.1.4.1 Sparse Functions	628
22.1.4.2 Return Types of Operators and Functions	628
22.1.4.3 Mathematical Considerations	630
22.2 Linear Algebra on Sparse Matrices	638
22.3 Iterative Techniques Applied to Sparse Matrices	647
22.4 Real Life Example using Sparse Matrices	655
23 Numerical Integration	659
23.1 Functions of One Variable	659
23.2 Orthogonal Collocation	668
23.3 Functions of Multiple Variables	669

24 Differential Equations	677
24.1 Ordinary Differential Equations	677
24.1.1 Matlab-compatible solvers	679
24.2 Differential-Algebraic Equations	688
25 Optimization	697
25.1 Linear Programming	697
25.2 Quadratic Programming	703
25.3 Nonlinear Programming	705
25.4 Linear Least Squares	707
26 Statistics	711
26.1 Descriptive Statistics	711
26.2 Statistics on Sliding Windows of Data	719
26.3 Basic Statistical Functions	736
26.4 Correlation and Regression Analysis	739
26.5 Distributions	741
26.6 Random Number Generation	741
27 Sets	743
27.1 Set Operations	743
28 Polynomial Manipulations	747
28.1 Evaluating Polynomials	747
28.2 Finding Roots	748
28.3 Products of Polynomials	749
28.4 Derivatives / Integrals / Transforms	752
28.5 Polynomial Interpolation	753
28.6 Miscellaneous Functions	763
29 Interpolation	765
29.1 One-dimensional Interpolation	765
29.2 Multi-dimensional Interpolation	769
30 Geometry	775
30.1 Delaunay Triangulation	775
30.1.1 Plotting the Triangulation	777
30.1.2 Identifying Points in Triangulation	780
30.2 Voronoi Diagrams	782
30.3 Convex Hull	786
30.4 Interpolation on Scattered Data	788
31 Signal Processing	791

32 Image Processing	805
32.1 Loading and Saving Images	805
32.2 Displaying Images	811
32.3 Representing Images	813
32.4 Plotting on top of Images	824
32.5 Color Conversion	824
33 Audio Processing	827
33.1 Audio File Utilities	827
33.2 Audio Device Information	828
33.3 Audio Player	829
33.3.1 Playback	829
33.3.2 Properties	830
33.4 Audio Recorder	830
33.4.1 Recording	831
33.4.2 Data Retrieval	831
33.4.3 Properties	832
33.5 Audio Data Processing	832
34 Object Oriented Programming	835
34.1 Creating a Class	835
34.2 Class Methods	837
34.3 Indexing Objects	840
34.3.1 Defining Indexing And Indexed Assignment	840
34.3.2 Indexed Assignment Optimization	844
34.4 Overloading Objects	845
34.4.1 Function Overloading	845
34.4.2 Operator Overloading	846
34.4.3 Precedence of Objects	847
34.5 Inheritance and Aggregation	848
34.6 <code>classdef</code> Classes	853
34.6.1 Creating a <code>classdef</code> Class	853
34.6.2 Properties	855
34.6.3 Methods	855
34.6.4 Inheritance	857
34.6.5 Value Classes vs. Handle Classes	857
35 GUI Development	861
35.1 I/O Dialogs	861
35.2 Progress Bar	868
35.3 UI Elements	869
35.4 GUI Utility Functions	879
35.5 User-Defined Preferences	881

36 System Utilities	885
36.1 Timing Utilities	885
36.2 Filesystem Utilities	897
36.3 File Archiving Utilities	906
36.4 Networking Utilities	909
36.4.1 FTP Objects	909
36.4.2 URL Manipulation	911
36.4.3 Base64 and Binary Data Transmission	912
36.5 Controlling Subprocesses	913
36.6 Process, Group, and User IDs	921
36.7 Environment Variables	922
36.8 Current Working Directory	922
36.9 Password Database Functions	924
36.10 Group Database Functions	925
36.11 System Information	926
36.12 Hashing Functions	931
37 Packages	933
37.1 Installing and Removing Packages	933
37.2 Using Packages	937
37.3 Administrating Packages	937
37.4 Creating Packages	937
37.4.1 The DESCRIPTION File	939
37.4.2 The INDEX File	941
37.4.3 PKG_ADD and PKG_DEL Directives	942
37.4.4 Missing Components	942
Appendix A External Code Interface	945
A.1 Oct-Files	946
A.1.1 Getting Started with Oct-Files	946
A.1.2 Matrices and Arrays in Oct-Files	950
A.1.3 Character Strings in Oct-Files	953
A.1.4 Cell Arrays in Oct-Files	955
A.1.5 Structures in Oct-Files	955
A.1.6 Sparse Matrices in Oct-Files	957
A.1.6.1 Array and Sparse Class Differences	957
A.1.6.2 Creating Sparse Matrices in Oct-Files	958
A.1.6.3 Using Sparse Matrices in Oct-Files	961
A.1.7 Accessing Global Variables in Oct-Files	962
A.1.8 Calling Octave Functions from Oct-Files	962
A.1.9 Calling External Code from Oct-Files	964
A.1.10 Allocating Local Memory in Oct-Files	966
A.1.11 Input Parameter Checking in Oct-Files	966
A.1.12 Exception and Error Handling in Oct-Files	968
A.1.13 Documentation and Testing of Oct-Files	969
A.2 Mex-Files	970
A.2.1 Getting Started with Mex-Files	970

A.2.2	Working with Matrices and Arrays in Mex-Files	972
A.2.3	Character Strings in Mex-Files	974
A.2.4	Cell Arrays with Mex-Files	975
A.2.5	Structures with Mex-Files	976
A.2.6	Sparse Matrices with Mex-Files	978
A.2.7	Calling Other Functions in Mex-Files	981
A.3	Standalone Programs	982
A.4	Java Interface	986
A.4.1	Making Java Classes Available	986
A.4.2	How to use Java from within Octave	987
A.4.3	Set up the JVM	989
A.4.4	Java Interface Functions	990
Appendix B	Test and Demo Functions	997
B.1	Test Functions	997
B.2	Demonstration Functions	1005
Appendix C	Obsolete Functions	1009
Appendix D	Known Causes of Trouble	1015
D.1	Actual Bugs We Haven't Fixed Yet	1015
D.2	Reporting Bugs	1015
D.2.1	Have You Found a Bug?	1015
D.2.2	Where to Report Bugs	1016
D.2.3	How to Report Bugs	1016
D.2.4	Sending Patches for Octave	1017
D.3	How To Get Help with Octave	1018
D.4	How to Distinguish Between Octave and Matlab	1018
Appendix E	Installing Octave	1021
E.1	Build Dependencies	1021
E.1.1	Obtaining the Dependencies Automatically	1021
E.1.2	Build Tools	1021
E.1.3	External Packages	1022
E.2	Running Configure and Make	1024
E.3	Compiling Octave with 64-bit Indexing	1029
E.4	Installation Problems	1031
Appendix F	Grammar and Parser	1035
F.1	Keywords	1035
F.2	Parser	1035
Appendix G	GNU GENERAL PUBLIC LICENSE	1037
Concept Index	1049	

Function Index	1055
Operator Index	1069
Graphics Properties Index	1071

Preface

Octave was originally intended to be companion software for an undergraduate-level textbook on chemical reactor design being written by James B. Rawlings of the University of Wisconsin-Madison and John G. Ekerdt of the University of Texas.

Clearly, Octave is now much more than just another ‘courseware’ package with limited utility beyond the classroom. Although our initial goals were somewhat vague, we knew that we wanted to create something that would enable students to solve realistic problems, and that they could use for many things other than chemical reactor design problems. We find that most students pick up the basics of Octave quickly, and are using it confidently in just a few hours.

Although it was originally intended to be used to teach reactor design, it has been used in several other undergraduate and graduate courses in the Chemical Engineering Department at the University of Texas, and the math department at the University of Texas has been using it for teaching differential equations and linear algebra as well. More recently, Octave has been used as the primary computational tool for teaching Stanford’s online Machine Learning class ([ml-class.org](#)) taught by Andrew Ng. Tens of thousands of students participated in the course.

If you find Octave useful, please let us know. We are always interested to find out how Octave is being used.

Virtually everyone thinks that the name Octave has something to do with music, but it is actually the name of one of John W. Eaton’s former professors who wrote a famous textbook on chemical reaction engineering, and who was also well known for his ability to do quick ‘back of the envelope’ calculations. We hope that this software will make it possible for many people to do more ambitious computations just as easily.

Everyone is encouraged to share this software with others under the terms of the GNU General Public License (see [Appendix G \[Copying\]](#), page 1037). You are also encouraged to help make Octave more useful by writing and contributing additional functions for it, and by reporting any problems you may have.

Acknowledgements

Many people have contributed to Octave’s development. The following people have helped code parts of Octave or aided in various other ways (listed alphabetically).

Ben Abbott	Drew Abbot	NVS Abhilash
Andy Adler	Adam H. Aitkenhead	Joakim Andén
Giles Anderson	Joel Andersson	Lachlan Andrew
Pedro Angelo	Damjan Angelovski	Muthiah Annamalai
Markus Appel	Branden Archer	Willem Atsma
Marco Atzeri	Ander Aurrekoetxea	Shai Ayal
Sahil Badyal	Jeff Bai	Roger Banks
Ben Barrowes	Alexander Barth	David Bateman
Heinz Bauschke	Miguel Bazdresch	Julien Bect
Stefan Beller	Roman Belov	Markus Bergholz
Karl Berry	Atri Bhattacharya	Ethan Biery
David Billinghurst	Don Bindner	Jakub Bogusz

Moritz Borgmann
John Bradshaw
Remy Bruno
Marco Caliari
Juan Pablo Carbajal
Larrie Carr
Marco Cecchetti
Albert Chin-A-Young
Catalin Codreanu
Andre da Costa Barros
Richard Crozier
Jacob Dawid
Thomas D. Dean
Fabian Deutsch
Vivek Dogra
Carnë Draug
John W. Eaton
Paul Eggert
Garrett Euler
Francesco Faccio
Stephen Fegan
David Finkel
Jose Daniel Munoz Frias
Eduardo Gallestey
Driss Ghaddab
Michele Ginesi
Michael D. Godfrey
Glenn Golden
Brian Gough
Steffen Groot
Kyle Guinn
Kai Habel
Jaroslav Hajek
Kim Hansen
Dave Hawthorne
Piotr Held
Martin Hepperle
Yozo Hida
A. Scott Edward Hodel
Tom Holroyd
Craig Hudson
John Hunt
Alan W. Irwin
Vytautas Jančauskas
Robert Jenssen
Heikki Junes
Jarkko Kaleva

Paul Boven
Marcus Brinkmann
Clemens Buchacher
Daniel Calvelo
Jean-Francois Cardoso
David Castelow
Corbin Champion
Sunghyun Cho
J. D. Cole
Martin Costabel
Jeff Cunningham
Jorge Barros de Abreu
Philippe Defert
Christos Dimitrakakis
John Donoghue
Sergey Dudoladov
Dirk Eddelbuettel
Stephen Eglen
Edmund Grimley Evans
Gunnar Farnebäck
Ramon Garcia Fernandez
Guillaume Flandin
Brad Froehle
Walter Gautschi
Eugenio Gianniti
Nicolo Giorgetti
Dave Goel
Tomislav Goles
Alexander Graf
Etienne Grossmann
Vaibhav Gupta
Patrick Häcker
Benjamin Hall
Gene Harvey
Oliver Heimlich
Martin Helm
Jordi Gutiérrez Hermoso
Ryan Hinton
Júlio Hoffmann
David Hoover
Christopher Hulbert
Stefan Husmann
Allan Jacobs
Nicholas R. Jankowski
Cai Jianming
Matthias Jüschke
Avinoam Kalma

Richard Bovey
Max Brister
Ansgar Burchard
John C. Campbell
Joao Cardoso
Vincent Cautaerts
Clinton Chee
Carsten Clark
Jacopo Corno
Michael Creel
Martin Dalecki
Carlo de Falco
Bill Denney
Pantxo Diribarne
David M. Doolin
Pascal A. Dupuis
Pieter Eendebak
Peter Ekberg
Rolf Fabian
Massimiliano Fasi
Torsten Finke
Colin Foster
Castor Fu
Klaus Gebhardt
Hartmut Gimpel
Arun Giridhar
Michael Goffioul
Keith Goodman
Michael C. Grant
David Grundberg
Peter Gustafson
William P. Y. Hadisoeseno
Alexander Hansen
Søren Hauberg
Daniel Heiserer
Stefan Hepp
Israel Herraiz
Roman Hodek
Richard Allan Holcombe
Kurt Hornik
Cyril Humbert
Teemu Ikonen
Geoff Jacobsen
Mats Jansson
Steven G. Johnson
Atsushi Kajita
Mohamed Kamoun

Lute Kamstra	Fotios Kasolis	Thomas Kasper
Joel Keay	Mumit Khan	Paul Kienzle
Lars Kindermann	Aaron A. King	Erik Kjellson
Arno J. Klaassen	Alexander Klein	Lasse Kliemann
Geoffrey Knauth	Heine Kolلتveit	Ken Kouno
Kacper Kowalik	Endre Kozma	Daniel Kraft
Nir Krakauer	Aravindh Krishnamoorthy	Oyvind Kristiansen
Artem Krosheninnikov	Piotr Krzyzanowski	Volker Kuhlmann
Ilya Kurdyukov	Tetsuro Kurita	Ben Kurtz
Philipp Kutin	Miroslaw Kwasniak	Rafael Laboissiere
Kai Labusch	Claude Lacoursiere	Walter Landry
Bill Lash	Dirk Laurie	Maurice LeBrun
Friedrich Leisch	Michael Leitner	Johannes Leuschner
Thorsten Liebig	Torsten Lilge	Jyh-miin Lin
Timo Lindfors	Benjamin Lindner	Ross Lippert
Yu Liu	David Livings	Barbara Locsi
Sebastien Loisel	Erik de Castro Lopo	Massimo Lorenzin
Emil Lucretiu	Yi-Hong Lyu	Hoxide Ma
Colin Macdonald	James Macnicol	Jens-Uwe Mager
Stefan Mahr	Rob Mahurin	Alexander Mamonov
Ricardo Marranita	Orestes Mas	Axel Math��
Makoto Matsumoto	Tatsuro Matsuoka	Christoph Mayer
Laurent Mazet	G. D. McBain	Ronald van der Meer
J��lio Hoffmann Mendes	Ed Meyer	Thorsten Meyer
Stefan Miereis	Petr Mikulik	Mike Miller
Serviscope Minor	Stefan Monnier	Rafael Monteiro
Stephen Montgomery-Smith	Antoine Moreau	Kai P. Mueller
Amod Mulay	Armin M��ller	Hannes M��ller
Victor Munoz	Prasanna Kumar	Iain Murray
	Muralidharan	
Nicholas Musolino	Markus M��tzel	Carmen Navarrete
Todd Neal	Philip Nienhuis	Al Niessner
Felipe G. Nievinski	Rick Niles	Takuji Nishimura
Akira Noda	Kai Noda	Patrick Noffke
Victor Norton	Eric Norum	Krzesimir Nowak
Michael O'Brien	Cillian O'Driscoll	Peter O'Gorman
Thorsten Ohl	Kai T. Ohlhus	Arno Onken
Valentin Ortega-Clavero	Luis F. Ortiz	Carl Osterwisch
Janne Olavi Paanaj��rvi	Scott Pakin	Jos�� Luis Garc��a Pallero
Jason Alan Palmer	Gabriele Pannocchia	Sylvain Pelissier
Rolando Pereira	Per Persson	Primozz Peterlin
Jim Peterson	Danilo Piazzalunga	Nicholas Piper
Elias Pipping	Robert Platt	Hans Ekkehard Plesser
Sergey Plotnikov	Tom Poage	Nathan Podlich
Orion Poplawski	Ondrej Popp	Jef Poskanzer
Francesco Potort��	Konstantinos Poulios	Tejaswi D. Prakash

Jarno Rajahalme	Eduardo Ramos	Pooja Rao
James B. Rawlings	Eric S. Raymond	Balint Reczey
Joshua Redstone	Andy Register	Lukas Reichlin
Michael Reifenberger	Ernst Reissner	Jens Restemeier
Anthony Richardson	Jason Riedy	E. Joshua Rigler
Sander van Rijn	Petter Rishholm	Matthew W. Roberts
Melvin Robinson	Dmitry Roshchin	Peter Rosin
Andrew Ross	Fabio Rossi	Mark van Rossum
Joe Rothweiler	David Rörich	Kevin Ruland
Kristian Rumberg	Ryan Rusaw	Olli Saarela
Toni Saarela	Juhani Saastamoinen	Radek Salac
Mike Sander	Ben Sapp	Aleksej Saushev
Alois Schlögl	Michel D. Schmid	Julian Schnidder
Sebastian Schöps	Nicol N. Schraudolph	Sebastian Schubert
Lasse Schuirmann	Ludwig Schwardt	Thomas L. Scofield
Daniel J. Sebald	Dmitri A. Sergatskov	Vanya Sergeev
Marko Seric	Ahsan Ali Shahid	Baylis Shanks
Andriy Shinkarchuck	Robert T. Short	Joseph P. Skudlarek
John Smith	Julius Smith	Shan G. Smith
Peter L. Sondergaard	Rüdiger Sonderfeld	Joerg Specht
Quentin H. Spencer	Christoph Spiel	David Spies
Andreas Stahel	Richard Stallman	Russell Standish
Ryan Starret	Brett Stewart	Doug Stewart
Jen Stewart	Jonathan Stickel	Judd Storrs
Thomas Stuart	Bernardo Sulzbach	Ivan Sutoris
John Swensen	Daisuke Takago	Ariel Tankus
Falk Tannhäuser	Duncan Temple Lang	Matthew Tenny
Kris Thielemans	Georg Thimm	Corey Thomasson
Andrew Thornton	Olaf Till	Christophe Tournery
Thomas Treichl	Abhinav Tripathi	Karsten Trulsen
David Turner	Frederick Umminger	Utkarsh Upadhyay
José Vallet	Stefan van der Walt	Peter Van Wieren
James R. Van Zandt	Risto Vanhanen	Gregory Vanuxem
Mihas Varantsou	Ivana Varekova	Sébastien Villemot
Marco Vitetta	Daniel Wagenaar	Thomas Walter
Jun Wang	Andreas Weber	Olaf Weber
Thomas Weber	Rik Wehbring	Bob Weigel
Andreas Weingessel	Martin Weiser	Michael Weitzel
David Wells	Joachim Wiesemann	Alexander Wilms
Joe Winegarden	Georg Wiora	Sahil Yadav
Fook Fah Yap	Sean Young	Michele Zaffalon
Serhiy Zahoriya	Johannes Zarl	Michael Zeising
Federico Zenith	Claudius Zingerli	Alex Zvoleff
Richard Zweig		

Special thanks to the following people and organizations for supporting the development of Octave:

- The United States Department of Energy, through grant number DE-FG02-04ER25635.
- Ashok Krishnamurthy, David Hudak, Juan Carlos Chaves, and Stanley C. Ahalt of the Ohio Supercomputer Center.
- The National Science Foundation, through grant numbers CTS-0105360, CTS-9708497, CTS-9311420, CTS-8957123, and CNS-0540147.
- The industrial members of the Texas-Wisconsin Modeling and Control Consortium (TWMCC).
- The Paul A. Elfers Endowed Chair in Chemical Engineering at the University of Wisconsin-Madison.
- Digital Equipment Corporation, for an equipment grant as part of their External Research Program.
- Sun Microsystems, Inc., for an Academic Equipment grant.
- International Business Machines, Inc., for providing equipment as part of a grant to the University of Texas College of Engineering.
- Texaco Chemical Company, for providing funding to continue the development of this software.
- The University of Texas College of Engineering, for providing a Challenge for Excellence Research Supplement, and for providing an Academic Development Funds grant.
- The State of Texas, for providing funding through the Texas Advanced Technology Program under Grant No. 003658-078.
- Noel Bell, Senior Engineer, Texaco Chemical Company, Austin Texas.
- John A. Turner, Group Leader, Continuum Dynamics (CCS-2), Los Alamos National Laboratory, for registering the octave.org domain name.
- James B. Rawlings, Professor, University of Wisconsin-Madison, Department of Chemical and Biological Engineering.
- Richard Stallman, for writing GNU.

This project would not have been possible without the GNU software used in and to produce Octave.

Citing Octave in Publications

In view of the many contributions made by numerous developers over many years it is common courtesy to cite Octave in publications when it has been used during the course of research or the preparation of figures. The **citation** function can automatically generate a recommended citation text for Octave or any of its packages. See the help text below on how to use **citation**.

citation

citation package

Display instructions for citing GNU Octave or its packages in publications.

When called without an argument, display information on how to cite the core GNU Octave system.

When given a package name *package*, display information on citing the specific named package. Note that some packages may not yet have instructions on how to cite them.

The GNU Octave developers and its active community of package authors have invested a lot of time and effort in creating GNU Octave as it is today. Please give credit where credit is due and cite GNU Octave and its packages when you use them.

How You Can Contribute to Octave

There are a number of ways that you can contribute to help make Octave a better system. Perhaps the most important way to contribute is to write high-quality code for solving new problems, and to make your code freely available for others to use. See <https://www.gnu.org/get-involved.html> for detailed information.

If you find Octave useful, consider providing additional funding to continue its development. Even a modest amount of additional funding could make a significant difference in the amount of time that is available for development and support.

Donations supporting Octave development may be made on the web at <https://my.fsf.org/donate/working-together/octave>. These donations also help to support the Free Software Foundation

If you'd prefer to pay by check or money order, you can do so by sending a check to the FSF at the following address:

Free Software Foundation
51 Franklin Street, Suite 500
Boston, MA 02110-1335
USA

If you pay by check, please be sure to write "GNU Octave" in the memo field of your check.

If you cannot provide funding or contribute code, you can still help make Octave better and more reliable by reporting any bugs you find and by offering suggestions for ways to improve Octave. See [Appendix D \[Trouble\], page 1015](#), for tips on how to write useful bug reports.

Distribution

Octave is *free* software. This means that everyone is free to use it and free to redistribute it on certain conditions. Octave is not, however, in the public domain. It is copyrighted and there are restrictions on its distribution, but the restrictions are designed to ensure that others will have the same freedom to use and redistribute Octave that you have. The precise conditions can be found in the GNU General Public License that comes with Octave and that also appears in [Appendix G \[Copying\], page 1037](#).

To download a copy of Octave, please visit <https://www.gnu.org/download.html>.

1 A Brief Introduction to Octave

GNU Octave is a high-level language primarily intended for numerical computations. It is typically used for such problems as solving linear and nonlinear equations, numerical linear algebra, statistical analysis, and for performing other numerical experiments. It may also be used as a batch-oriented language for automated data processing.

The current version of Octave executes in a graphical user interface (GUI). The GUI hosts an Integrated Development Environment (IDE) which includes a code editor with syntax highlighting, built-in debugger, documentation browser, as well as the interpreter for the language itself. A command-line interface for Octave is also available.

GNU Octave is freely redistributable software. You may redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation. The GPL is included in this manual, see [Appendix G \[Copying\]](#), page 1037.

This manual provides comprehensive documentation on how to install, run, use, and extend GNU Octave. Additional chapters describe how to report bugs and help contribute code.

This document corresponds to Octave version 5.0.91.

1.1 Running Octave

On most systems, Octave is started with the shell command ‘`octave`’. This starts the graphical user interface. The central window in the GUI is the Octave command-line interface. In this window Octave displays an initial message and then a prompt indicating it is ready to accept input. If you have chosen the traditional command-line interface then only the command prompt appears in the same window that was running a shell. In either case, you can immediately begin typing Octave commands.

If you get into trouble, you can usually interrupt Octave by typing `Control-C` (written `C-c` for short). `C-c` gets its name from the fact that you type it by holding down `CTRL` and then pressing `c`. Doing this will normally return you to Octave’s prompt.

To exit Octave, type `quit` or `exit` at the Octave prompt.

On systems that support job control, you can suspend Octave by sending it a `SIGTSTP` signal, usually by typing `C-z`.

1.2 Simple Examples

The following chapters describe all of Octave’s features in detail, but before doing that, it might be helpful to give a sampling of some of its capabilities.

If you are new to Octave, we recommend that you try these examples to begin learning Octave by using it. Lines marked like so, ‘`octave:13>`’, are lines you type, ending each with a carriage return. Octave will respond with an answer, or by displaying a graph.

1.2.1 Elementary Calculations

Octave can easily be used for basic numerical calculations. Octave knows about arithmetic operations (`+, -, *, /`), exponentiation (`^`), natural logarithms/exponents (`log, exp`), and the trigonometric functions (`sin, cos, ...`). Moreover, Octave calculations work on real or imaginary numbers (`i,j`). In addition, some mathematical constants such as the base of

the natural logarithm (e) and the ratio of a circle's circumference to its diameter (π) are pre-defined.

For example, to verify Euler's Identity,

$$e^{i\pi} = -1$$

type the following which will evaluate to -1 within the tolerance of the calculation.

```
octave:1> exp (i*pi)
```

1.2.2 Creating a Matrix

Vectors and matrices are the basic building blocks for numerical analysis. To create a new matrix and store it in a variable so that you can refer to it later, type the command

```
octave:1> A = [ 1, 1, 2; 3, 5, 8; 13, 21, 34 ]
```

Octave will respond by printing the matrix in neatly aligned columns. Octave uses a comma or space to separate entries in a row, and a semicolon or carriage return to separate one row from the next. Ending a command with a semicolon tells Octave not to print the result of the command. For example,

```
octave:2> B = rand (3, 2);
```

will create a 3 row, 2 column matrix with each element set to a random value between zero and one.

To display the value of a variable, simply type the name of the variable at the prompt. For example, to display the value stored in the matrix B , type the command

```
octave:3> B
```

1.2.3 Matrix Arithmetic

Octave uses standard mathematical notation with the advantage over low-level languages that operators may act on scalars, vector, matrices, or N-dimensional arrays. For example, to multiply the matrix A by a scalar value, type the command

```
octave:4> 2 * A
```

To multiply the two matrices A and B , type the command

```
octave:5> A * B
```

and to form the matrix product $A^T A$, type the command

```
octave:6> A' * A
```

1.2.4 Solving Systems of Linear Equations


Systems of linear equations are ubiquitous in numerical analysis. To solve the set of linear equations $Ax = b$, use the left division operator, ' \backslash :

```
x = A \ b
```




This is conceptually equivalent to $A^{-1}b$, but avoids computing the inverse of a matrix directly.

If the coefficient matrix is singular, Octave will print a warning message and compute a minimum norm solution.

A simple example comes from chemistry and the need to obtain balanced chemical equations. Consider the burning of hydrogen and oxygen to produce water.



The equation above is not accurate. The Law of Conservation of Mass requires that the number of molecules of each type balance on the left- and right-hand sides of the equation. Writing the variable overall reaction with individual equations for hydrogen and oxygen one finds:



The solution in Octave is found in just three steps.

```
octave:1> A = [ 2, 0; 0, 2 ];
octave:2> b = [ 2; 1 ];
octave:3> x = A \ b
```

1.2.5 Integrating Differential Equations

Octave has built-in functions for solving nonlinear differential equations of the form

$$\frac{dx}{dt} = f(x, t), \quad x(t = t_0) = x_0$$

For Octave to integrate equations of this form, you must first provide a definition of the function $f(x, t)$. This is straightforward, and may be accomplished by entering the function body directly on the command line. For example, the following commands define the right-hand side function for an interesting pair of nonlinear differential equations. Note that while you are entering a function, Octave responds with a different prompt, to indicate that it is waiting for you to complete your input.

```
octave:1> function xdot = f (x, t)
>
> r = 0.25;
> k = 1.4;
> a = 1.5;
> b = 0.16;
> c = 0.9;
> d = 0.8;
>
> xdot(1) = r*x(1)*(1 - x(1)/k) - a*x(1)*x(2)/(1 + b*x(1));
> xdot(2) = c*a*x(1)*x(2)/(1 + b*x(1)) - d*x(2);
>
> endfunction
```

Given the initial condition

```
octave:2> x0 = [1; 2];
```

and the set of output times as a column vector (note that the first output time corresponds to the initial condition given above)

```
octave:3> t = linspace (0, 50, 200)';
```

it is easy to integrate the set of differential equations:

```
octave:4> x = lsode ("f", x0, t);
```

The function `lsode` uses the Livermore Solver for Ordinary Differential Equations, described in A. C. Hindmarsh, *ODEPACK, a Systematized Collection of ODE Solvers*, in: Scientific Computing, R. S. Stepleman et al. (Eds.), North-Holland, Amsterdam, 1983, pages 55–64.

1.2.6 Producing Graphical Output

To display the solution of the previous example graphically, use the command

```
octave:1> plot (t, x)
```

Octave will automatically create a separate window to display the plot.

To save a plot once it has been displayed on the screen, use the `print` command. For example,

```
print -dpdf foo.pdf
```

will create a file called `foo.pdf` that contains a rendering of the current plot in Portable Document Format. The command

```
help print
```

explains more options for the `print` command and provides a list of additional output file formats.

1.2.7 Help and Documentation

Octave has an extensive help facility. The same documentation that is available in printed form is also available from the Octave prompt, because both forms of the documentation are created from the same input file.

In order to get good help you first need to know the name of the command that you want to use. The name of this function may not always be obvious, but a good place to start is to type `help --list`. This will show you all the operators, keywords, built-in functions, and loadable functions available in the current session of Octave. An alternative is to search the documentation using the `lookfor` function (described in [Section 2.3 \[Getting Help\], page 20](#)).

Once you know the name of the function you wish to use, you can get more help on the function by simply including the name as an argument to `help`. For example,

```
help plot
```

will display the help text for the `plot` function.

The part of Octave's help facility that allows you to read the complete text of the printed manual from within Octave normally uses a separate program called Info. When you invoke Info you will be put into a menu driven program that contains the entire Octave manual. Help for using Info is provided in this manual, see [Section 2.3 \[Getting Help\], page 20](#).

1.2.8 Editing What You Have Typed

At the Octave prompt, you can recall, edit, and reissue previous commands using Emacs- or vi-style editing commands. The default keybindings use Emacs-style commands. For example, to recall the previous command, press `Control-p` (written `C-p` for short). Doing this will normally bring back the previous line of input. `C-n` will bring up the next line of

input, $C-b$ will move the cursor backward on the line, $C-f$ will move the cursor forward on the line, etc.

A complete description of the command line editing capability is given in this manual, see [Section 2.4 \[Command Line Editing\]](#), page 25.

1.3 Conventions

This section explains the notation conventions that are used in this manual. You may want to skip this section and refer back to it later.

1.3.1 Fonts

Examples of Octave code appear in this font or form: `svd (a)`. Names that represent variables or function arguments appear in this font or form: *first-number*. Commands that you type at the shell prompt appear in this font or form: ‘`octave --no-init-file`’. Commands that you type at the Octave prompt sometimes appear in this font or form: `foo --bar --baz`. Specific keys on your keyboard appear in this font or form: RET.

1.3.2 Evaluation Notation

In the examples in this manual, results from expressions that you evaluate are indicated with ‘ \Rightarrow ’. For example:

```
sqrt (2)
⇒ 1.4142
```

You can read this as “`sqrt (2)` evaluates to 1.4142”.

In some cases, matrix values that are returned by expressions are displayed like this

```
[1, 2; 3, 4] == [1, 3; 2, 4]
⇒ [ 1, 0; 0, 1 ]
```

and in other cases, they are displayed like this

```
eye (3)
⇒ 1 0 0
    0 1 0
    0 0 1
```

in order to clearly show the structure of the result.

Sometimes to help describe one expression, another expression is shown that produces identical results. The exact equivalence of expressions is indicated with ‘ \equiv ’. For example:

```
rot90 ([1, 2; 3, 4], -1)
≡
rot90 ([1, 2; 3, 4], 3)
≡
rot90 ([1, 2; 3, 4], 7)
```

1.3.3 Printing Notation

Many of the examples in this manual print text when they are evaluated. In this manual the printed text resulting from an example is indicated by ‘ \dashv ’. The value that is returned

by evaluating the expression is displayed with ‘⇒’ (1 in the next example) and follows on a separate line.

```
printf ("foo %s\n", "bar")
      ↓ foo bar
⇒ 1
```

1.3.4 Error Messages

Some examples signal errors. This normally displays an error message on your terminal. Error messages are shown on a line beginning with `error:`.

```
fieldnames ([1, 2; 3, 4])
error: fieldnames: Invalid input argument
```

1.3.5 Format of Descriptions

Functions and commands are described in this manual in a uniform format. The first line of a description contains the name of the item followed by its arguments, if any. If there are multiple ways to invoke the function then each allowable form is listed.

The description follows on succeeding lines, sometimes with examples.

1.3.5.1 A Sample Function Description

In a function description, the name of the function being described appears first. It is followed on the same line by a list of parameters. The names used for the parameters are also used in the body of the description.

After all of the calling forms have been enumerated, the next line is a concise one-sentence summary of the function.

After the summary there may be documentation on the inputs and outputs, examples of function usage, notes about the algorithm used, and references to related functions.

Here is a description of an imaginary function `foo`:

```
foo (x)
foo (x, y)
foo (x, y, ...)
```

Subtract *x* from *y*, then add any remaining arguments to the result.

The input *x* must be a numeric scalar, vector, or array.

The optional input *y* defaults to 19 if it is not supplied.

Example:

```
foo (1, [3, 5], 3, 9)
    ⇒ [ 14, 16 ]
foo (5)
    ⇒ 14
```

More generally,

```
foo (w, x, y, ...)
≡
x - w + y + ...
```

See also: bar

Any parameter whose name contains the name of a type (e.g., *integer* or *matrix*) is expected to be of that type. Parameters named *object* may be of any type. Parameters with other sorts of names (e.g., *new_file*) are discussed specifically in the description of the function. In some sections, features common to parameters of several functions are described at the beginning.

1.3.5.2 A Sample Command Description

Commands are functions that may be called without surrounding their arguments in parentheses. Command descriptions have a format similar to function descriptions. For example, here is the description for Octave's **diary** command:

```
diary
diary on
diary off
diary filename
[status, diaryfile] = diary
```

Record a list of all commands *and* the output they produce, mixed together just as they appear on the terminal.

Valid options are:

on Start recording a session in a file called **diary** in the current working directory.

off Stop recording the session in the diary file.

filename Record the session in the file named *filename*.

With no input or output arguments, **diary** toggles the current diary state.

If output arguments are requested, **diary** ignores inputs and returns the current status. The boolean *status* indicates whether recording is on or off, and *diaryfile* is the name of the file where the session is stored.

See also: `history`, `evalc`.

2 Getting Started

This chapter explains some of Octave's basic features, including how to start an Octave session, get help at the command prompt, edit the command line, and write Octave programs that can be executed as commands from your shell.

2.1 Invoking Octave from the Command Line

Normally, Octave is used interactively by running the program ‘`octave`’ without any arguments. Once started, Octave reads commands from the terminal until you tell it to exit.

You can also specify the name of a file on the command line, and Octave will read and execute the commands from the named file and then exit when it is finished.

You can further control how Octave starts by using the command-line options described in the next section, and Octave itself can remind you of the options available. Type ‘`octave --help`’ to display all available options and briefly describe their use (‘`octave -h`’ is a shorter equivalent).

2.1.1 Command Line Options

Here is a complete list of the command line options that Octave accepts.

--built-in-docstrings-file *filename*

Specify the name of the file containing documentation strings for the built-in functions of Octave. This value is normally correct and should only need to be specified in extraordinary situations.

--debug

-d Enter parser debugging mode. Using this option will cause Octave's parser to print a lot of information about the commands it reads, and is probably only useful if you are actually trying to debug the parser.

--debug-jit

Enable JIT compiler debugging and tracing.

--doc-cache-file *filename*

Specify the name of the doc cache file to use. The value of *filename* specified on the command line will override any value of `OCTAVE_DOC_CACHE_FILE` found in the environment, but not any commands in the system or user startup files that use the `doc_cache_file` function.

--echo-commands

-x Echo commands as they are executed.

--eval *code*

Evaluate *code* and exit when finished unless **--persist** is also specified.

--exec-path *path*

Specify the path to search for programs to run. The value of *path* specified on the command line will override any value of `OCTAVE_EXEC_PATH` found in the environment, but not any commands in the system or user startup files that set the built-in variable `EXEC_PATH`.

--gui Start the graphical user interface (GUI).

--help

-h Print short help message and exit.

--image-path *path*

Add path to the head of the search path for images. The value of *path* specified on the command line will override any value of `OCTAVE_IMAGE_PATH` found in the environment, but not any commands in the system or user startup files that set the built-in variable `IMAGE_PATH`.

--info-file *filename*

Specify the name of the info file to use. The value of *filename* specified on the command line will override any value of `OCTAVE_INFO_FILE` found in the environment, but not any commands in the system or user startup files that use the `info_file` function.

--info-program *program*

Specify the name of the info program to use. The value of *program* specified on the command line will override any value of `OCTAVE_INFO_PROGRAM` found in the environment, but not any commands in the system or user startup files that use the `info_program` function.

--interactive

-i Force interactive behavior. This can be useful for running Octave via a remote shell command or inside an Emacs shell buffer.

--jit-compiler

Enable the JIT compiler used for accelerating loops.

--line-editing

Force readline use for command-line editing.

--no-gui

Disable the graphical user interface (GUI) and use the command line interface (CLI) instead. This is the default behavior, but this option may be useful to override a previous --gui.

--no-history

-H Disable recording of command-line history.

--no-init-file

Don't read the initialization files `~/.octaverc` and `.octaverc`.

--no-init-path

Don't initialize the search path for function files to include default locations.

--no-line-editing

Disable command-line editing.

--no-site-file

Don't read the site-wide `octaverc` initialization files.

--no-window-system

-W Disable use of a windowing system including graphics. This forces a strictly terminal-only environment.

--norc
 -f Don't read any of the system or user initialization files at startup. This is equivalent to using both of the options --no-init-file and --no-site-file.

--path *path*
 -p *path* Add path to the head of the search path for function files. The value of *path* specified on the command line will override any value of OCTAVE_PATH found in the environment, but not any commands in the system or user startup files that set the internal load path through one of the path functions.

--persist
 Go to interactive mode after --eval or reading from a file named on the command line.

--silent
 --quiet
 -q Don't print the usual greeting and version message at startup.

--texi-macros-file *filename*
 Specify the name of the file containing Texinfo macros for use by makeinfo.

--traditional
 --braindead
 For compatibility with MATLAB, set initial values for user preferences to the following values

PS1	= ">> "
PS2	= ""
beep_on_error	= true
confirm_recursive_rmdir	= false
crash.dumps_octave_core	= false
disable_diagonal_matrix	= true
disable_permutation_matrix	= true
disable_range	= true
fixed_point_format	= true
history_timestamp_format_string	= "%-- %D %I:%M %p --%%"
print_empty_dimensions	= false
save_default_options	= "-mat-binary"
struct_levels_to_print	= 0

and disable the following warnings

```
Octave:abbreviated-property-match
Octave:data-file-in-path
Octave:function-name-clash
Octave:possible-matlab-short-circuit-operator
```

Note that this does not enable the Octave:language-extension warning, which you might want if you want to be told about writing code that works in Octave but not MATLAB (see [warning], page 228, [warning_ids], page 230).

--verbose
 -V Turn on verbose output.

`--version`

`-v` Print the program version number and exit.

`file` Execute commands from `file`. Exit when done unless `--persist` is also specified.

Octave also includes several functions which return information about the command line, including the number of arguments and all of the options.

`argv ()`

Return the command line arguments passed to Octave.

For example, if you invoked Octave using the command

```
octave --no-line-editing --silent
```

`argv` would return a cell array of strings with the elements `--no-line-editing` and `--silent`.

If you write an executable Octave script, `argv` will return the list of arguments passed to the script. See [Section 2.6 \[Executable Octave Programs\], page 36](#), for an example of how to create an executable Octave script.

`program_name ()`

Return the last component of the value returned by `program_invocation_name`.

See also: [\[program_invocation_name\]](#), page 18.

`program_invocation_name ()`

Return the name that was typed at the shell prompt to run Octave.

If executing a script from the command line (e.g., `octave foo.m`) or using an executable Octave script, the program name is set to the name of the script. See [Section 2.6 \[Executable Octave Programs\], page 36](#), for an example of how to create an executable Octave script.

See also: [\[program_name\]](#), page 18.

Here is an example of using these functions to reproduce the command line which invoked Octave.

```
printf ("%s", program_name ());
arg_list = argv ();
for i = 1:nargin
    printf (" %s", arg_list{i});
endfor
printf ("\n");
```

See [Section 6.3.3 \[Indexing Cell Arrays\], page 120](#), for an explanation of how to retrieve objects from cell arrays, and [Section 11.2 \[Defining Functions\], page 177](#), for information about the variable `nargin`.

2.1.2 Startup Files

When Octave starts, it looks for commands to execute from the files in the following list. These files may contain any valid Octave commands, including function definitions.

`octave-home/share/octave/site/m/startup/octaverc`

where `octave-home` is the directory in which Octave is installed (the default is `/usr/local`). This file is provided so that changes to the default Octave

environment can be made globally for all users at your site for all versions of Octave you have installed. Care should be taken when making changes to this file since all users of Octave at your site will be affected. The default file may be overridden by the environment variable `OCTAVE_SITE_INITFILE`.

`octave-home/share/octave/version/m/startup/octaverc`

where `octave-home` is the directory in which Octave is installed (the default is `/usr/local`), and `version` is the version number of Octave. This file is provided so that changes to the default Octave environment can be made globally for all users of a particular version of Octave. Care should be taken when making changes to this file since all users of Octave at your site will be affected. The default file may be overridden by the environment variable `OCTAVE_VERSION_INITFILE`.

`~/.octaverc`

This file is used to make personal changes to the default Octave environment.

`.octaverc`

This file can be used to make changes to the default Octave environment for a particular project. Octave searches for this file in the current directory after it reads `~/.octaverc`. Any use of the `cd` command in the `~/.octaverc` file will affect the directory where Octave searches for `.octaverc`.

If you start Octave in your home directory, commands from the file `~/.octaverc` will only be executed once.

`startup.m`

This file is used to make personal changes to the default Octave environment. It is executed for MATLAB compatibility, but `~/.octaverc` is the preferred location for configuration changes.

A message will be displayed as each of the startup files is read if you invoke Octave with the `--verbose` option but without the `--silent` option.

2.2 Quitting Octave

Shutdown is initiated with the `exit` or `quit` commands (they are equivalent). Similar to startup, Octave has a shutdown process that can be customized by user script files. During shutdown Octave will search for the script file `finish.m` in the function load path. Commands to save all workspace variables or cleanup temporary files may be placed there. Additional functions to execute on shutdown may be registered with `atexit`.

```
exit
exit (status)
quit
quit (status)
```

Exit the current Octave session.

If the optional integer value `status` is supplied, pass that value to the operating system as Octave's exit status. The default value is zero.

When exiting, Octave will attempt to run the m-file `finish.m` if it exists. User commands to save the workspace or clean up temporary files may be placed in that file. Alternatively, another m-file may be scheduled to run using `atexit`.

See also: [\[atexit\]](#), page 20.

atexit (*fcn*)
atexit (*fcn*, *flag*)

Register a function to be called when Octave exits.

For example,

```
function last_words ()
    disp ("Bye bye");
endfunction
atexit ("last_words");
```

will print the message "Bye bye" when Octave exits.

The additional argument *flag* will register or unregister *fcn* from the list of functions to be called when Octave exits. If *flag* is true, the function is registered, and if *flag* is false, it is unregistered. For example, after registering the function `last_words` above,

```
atexit ("last_words", false);
```

will remove the function from the list and Octave will not call `last_words` when it exits.

Note that `atexit` only removes the first occurrence of a function from the list, so if a function was placed in the list multiple times with `atexit`, it must also be removed from the list multiple times.

See also: [\[quit\]](#), page 19.

2.3 Commands for Getting Help

The entire text of this manual is available from the Octave prompt via the command `doc`. In addition, the documentation for individual user-written functions and variables is also available via the `help` command. This section describes the commands used for reading the manual and the documentation strings for user-supplied functions and variables. See [Section 11.9 \[Function Files\]](#), page 193, for more information about how to document the functions you write.

```
help name
help --list
help .
help
```

Display the help text for *name*.

For example, the command `help help` prints a short message describing the `help` command.

Given the single argument `--list`, list all operators, keywords, built-in functions, and loadable functions available in the current session of Octave.

Given the single argument `.`, list all operators available in the current session of Octave.

If invoked without any arguments, `help` displays instructions on how to access help from the command line.

The help command can provide information about most operators, but *name* must be enclosed by single or double quotes to prevent the Octave interpreter from acting on *name*. For example, `help "+"` displays help on the addition operator.

See also: [\[doc\]](#), page 21, [\[lookfor\]](#), page 21, [\[which\]](#), page 138, [\[info\]](#), page 22.

```
doc function_name
doc
```

Display documentation for the function *function_name* directly from an online version of the printed manual, using the GNU Info browser.

If invoked without an argument, the manual is shown from the beginning.

For example, the command `doc rand` starts the GNU Info browser at the `rand` node in the online version of the manual.

Once the GNU Info browser is running, help for using it is available using the command `C-h`.

See also: [\[help\]](#), page 20.

```
lookfor str
lookfor -all str
[fcns, help1str] = lookfor (str)
[fcns, help1str] = lookfor ("{-all", str)
```

Search for the string *str* in the documentation of all functions in the current function search path.

By default, `lookfor` looks for *str* in just the first sentence of the help string for each function found. The entire help text of each function can be searched by using the "`-all`" argument. All searches are case insensitive.

When called with no output arguments, `lookfor` prints the list of matching functions to the terminal. Otherwise, the output argument *fcns* contains the function names and *help1str* contains the first sentence from the help string of each function.

Programming Note: The ability of `lookfor` to correctly identify the first sentence of the help text is dependent on the format of the function's help. All Octave core functions are correctly formatted, but the same can not be guaranteed for external packages and user-supplied functions. Therefore, the use of the "`-all`" argument may be necessary to find related functions that are not a part of Octave.

The speed of lookup is greatly enhanced by having a cached documentation file. See `doc_cache_create` for more information.

See also: [\[help\]](#), page 20, [\[doc\]](#), page 21, [\[which\]](#), page 138, [\[path\]](#), page 197, [\[doc_cache_create\]](#), page 24.

To see what is new in the current release of Octave, use the `news` function.

```
news
news package
```

Display the current NEWS file for Octave or an installed package.

When called without an argument, display the NEWS file for Octave.

When given a package name *package*, display the current NEWS file for that package.

See also: [\[ver\]](#), page 928, [\[pkg\]](#), page 933.

info ()

Display contact information for the GNU Octave community.

warranty ()

Describe the conditions for copying and distributing Octave.

The following functions can be used to change which programs are used for displaying the documentation, and where the documentation can be found.

```
val = info_file ()
old_val = info_file (new_val)
info_file (new_val, "local")
```

Query or set the internal variable that specifies the name of the Octave info file.

The default value is *octave-home/info/octave.info*, in which *octave-home* is the root directory of the Octave installation. The default value may be overridden by the environment variable `OCTAVE_INFO_FILE`, or the command line argument `--info-file FNAME`.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [info_program], page 22, [doc], page 21, [help], page 20, [makeinfo_program], page 22.

```
val = info_program ()
old_val = info_program (new_val)
info_program (new_val, "local")
```

Query or set the internal variable that specifies the name of the info program to run.

The default value is *octave-home/libexec/octave/version/exec/arch/info* in which *octave-home* is the root directory of the Octave installation, *version* is the Octave version number, and *arch* is the system type (for example, *i686-pc-linux-gnu*). The default value may be overridden by the environment variable `OCTAVE_INFO_PROGRAM`, or the command line argument `--info-program NAME`.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [info_file], page 22, [doc], page 21, [help], page 20, [makeinfo_program], page 22.

```
val = makeinfo_program ()
old_val = makeinfo_program (new_val)
makeinfo_program (new_val, "local")
```

Query or set the internal variable that specifies the name of the program that Octave runs to format help text containing Texinfo markup commands.

The default value is `makeinfo`.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[texi_macros_file\]](#), page 23, [\[info_file\]](#), page 22, [\[info_program\]](#), page 22, [\[doc\]](#), page 21, [\[help\]](#), page 20.

```
val = texi_macros_file ()
old_val = texi_macros_file (new_val)
texi_macros_file (new_val, "local")
```

Query or set the internal variable that specifies the name of the file containing Texinfo macros that are prepended to documentation strings before they are passed to makeinfo.

The default value is *octave-home/share/octave/version/etc/macros.texi*, in which *octave-home* is the root directory of the Octave installation, and *version* is the Octave version number. The default value may be overridden by the environment variable `OCTAVE_TEXI_MACROS_FILE`, or the command line argument `--texi-macros-file FNAME`.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[makeinfo_program\]](#), page 22.

```
val = doc_cache_file ()
old_val = doc_cache_file (new_val)
doc_cache_file (new_val, "local")
```

Query or set the internal variable that specifies the name of the Octave documentation cache file.

A cache file significantly improves the performance of the `lookfor` command. The default value is *octave-home/share/octave/version/etc/doc-cache*, in which *octave-home* is the root directory of the Octave installation, and *version* is the Octave version number. The default value may be overridden by the environment variable `OCTAVE_DOC_CACHE_FILE`, or the command line argument `--doc-cache-file FNAME`.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[doc_cache_create\]](#), page 24, [\[lookfor\]](#), page 21, [\[info_program\]](#), page 22, [\[doc\]](#), page 21, [\[help\]](#), page 20, [\[makeinfo_program\]](#), page 22.

See also: [\[lookfor\]](#), page 21.

```
val = built_in_docstrings_file ()
old_val = built_in_docstrings_file (new_val)
built_in_docstrings_file (new_val, "local")
```

Query or set the internal variable that specifies the name of the file containing docstrings for built-in Octave functions.

The default value is *octave-home/share/octave/version/etc/built-in-docstrings*, in which *octave-home* is the root directory of the Octave installation,

and `version` is the Octave version number. The default value may be overridden by the environment variable `OCTAVE_BUILT_IN_DOCSTRINGS_FILE`, or the command line argument `--built-in-docstrings-file FNAME`.

Note: This variable is only used when Octave is initializing itself. Modifying it during a running session of Octave will have no effect.

```
val = suppress_verbose_help_message ()
old_val = suppress_verbose_help_message (new_val)
suppress_verbose_help_message (new_val, "local")
```

Query or set the internal variable that controls whether Octave will add additional help information to the end of the output from the `help` command and usage messages for built-in commands.

When called from inside a function with the "`local`" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

The following functions are principally used internally by Octave for generating the documentation. They are documented here for completeness and because they may occasionally be useful for users.

```
doc_cache_create (out_file, directory)
doc_cache_create (out_file)
doc_cache_create ()
```

Generate documentation cache for all functions in `directory`.

A documentation cache is generated for all functions in `directory` which may be a single string or a cell array of strings. The cache is used to speed up the function `lookfor`.

The cache is saved in the file `out_file` which defaults to the value `doc-cache` if not given.

If no directory is given (or it is the empty matrix), a cache for built-in functions, operators, and keywords is generated.

See also: [\[doc_cache_file\]](#), page 23, [\[lookfor\]](#), page 21, [\[path\]](#), page 197.

```
[text, format] = get_help_text (name)
```

Return the raw help text of function `name`.

The raw help text is returned in `text` and the format in `format`. The format is a string which is one of "`texinfo`", "`html`", or "`plain text`".

See also: [\[get_help_text_from_file\]](#), page 24.

```
[text, format] = get_help_text_from_file (fname)
```

Return the raw help text from the file `fname`.

The raw help text is returned in `text` and the format in `format`. The format is a string which is one of "`texinfo`", "`html`", or "`plain text`".

See also: [\[get_help_text\]](#), page 24.

```
text = get_first_help_sentence (name)
text = get_first_help_sentence (name, max_len)
[text, status] = get_first_help_sentence (...)

Return the first sentence of a function's help text.
```

The first sentence is defined as the text after the function declaration until either the first period (".") or the first appearance of two consecutive newlines ("\n\n"). The text is truncated to a maximum length of *max_len*, which defaults to 80. If the text must be truncated the last three characters of the text are replaced with "..." to indicate that more text was available.

The optional output argument *status* returns the status reported by `makeinfo`. If only one output argument is requested, and *status* is nonzero, a warning is displayed.

As an example, the first sentence of this help text is

```
get_first_help_sentence ("get_first_help_sentence")
→ ans = Return the first sentence of a function's help text.
```

2.4 Command Line Editing

Octave uses the GNU Readline library to provide an extensive set of command-line editing and history features. Only the most common features are described in this manual. In addition, all of the editing functions can be bound to different key strokes at the user's discretion. This manual assumes no changes from the default Emacs bindings. See the GNU Readline Library manual for more information on customizing Readline and for a complete feature list.

To insert printing characters (letters, digits, symbols, etc.), simply type the character. Octave will insert the character at the cursor and advance the cursor forward.

Many of the command-line editing functions operate using control characters. For example, the character `Control-a` moves the cursor to the beginning of the line. To type `C-a`, hold down `CTRL` and then press `a`. In the following sections, control characters such as `Control-a` are written as `C-a`.

Another set of command-line editing functions use Meta characters. To type `M-u`, hold down the `META` key and press `u`. Depending on the keyboard, the `META` key may be labeled `ALT` or even `WINDOWS`. If your terminal does not have a `META` key, you can still type Meta characters using two-character sequences starting with `ESC`. Thus, to enter `M-u`, you would type `ESC u`. The `ESC` character sequences are also allowed on terminals with real Meta keys. In the following sections, Meta characters such as `Meta-u` are written as `M-u`.

2.4.1 Cursor Motion

The following commands allow you to position the cursor.

`C-b` Move back one character.

`C-f` Move forward one character.

`BACKSPACE`

Delete the character to the left of the cursor.

`DEL` Delete the character underneath the cursor.

<i>C-d</i>	Delete the character underneath the cursor.
<i>M-f</i>	Move forward a word.
<i>M-b</i>	Move backward a word.
<i>C-a</i>	Move to the start of the line.
<i>C-e</i>	Move to the end of the line.
<i>C-l</i>	Clear the screen, reprinting the current line at the top.
<i>C-</i>	
<i>C-/</i>	Undo the last action. You can undo all the way back to an empty line.
<i>M-r</i>	Undo all changes made to this line. This is like typing the ‘undo’ command enough times to get back to the beginning.

The above table describes the most basic possible keystrokes that you need in order to do editing of the input line. On most terminals, you can also use the left and right arrow keys in place of *C-f* and *C-b* to move forward and backward.

Notice how *C-f* moves forward a character, while *M-f* moves forward a word. It is a loose convention that control keystrokes operate on characters while meta keystrokes operate on words.

The function `clc` will allow you to clear the screen from within Octave programs.

```
clc ()  
home ()
```

Clear the terminal screen and move the cursor to the upper left corner.

2.4.2 Killing and Yanking

Killing text means to delete the text from the line, but to save it away for later use, usually by *yanking* it back into the line. If the description for a command says that it ‘kills’ text, then you can be sure that you can get the text back in a different (or the same) place later.

Here is the list of commands for killing text.

<i>C-k</i>	Kill the text from the current cursor position to the end of the line.
<i>M-d</i>	Kill from the cursor to the end of the current word, or if between words, to the end of the next word.
<i>M-DEL</i>	Kill from the cursor to the start of the previous word, or if between words, to the start of the previous word.
<i>C-w</i>	Kill from the cursor to the previous whitespace. This is different than <i>M-DEL</i> because the word boundaries differ.

And, here is how to *yank* the text back into the line. Yanking means to copy the most-recently-killed text from the kill buffer.

<i>C-y</i>	Yank the most recently killed text back into the buffer at the cursor.
<i>M-y</i>	Rotate the kill-ring, and yank the new top. You can only do this if the prior command is <i>C-y</i> or <i>M-y</i> .

When you use a kill command, the text is saved in a *kill-ring*. Any number of consecutive kills save all of the killed text together, so that when you yank it back, you get it in one clean sweep. The kill ring is not line specific; the text that you killed on a previously typed line is available to be yanked back later, when you are typing another line.

2.4.3 Commands for Changing Text

The following commands can be used for entering characters that would otherwise have a special meaning (e.g., TAB, C-q, etc.), or for quickly correcting typing mistakes.

<i>C-q</i>	
<i>C-v</i>	Add the next character that you type to the line verbatim. This is how to insert things like <i>C-q</i> for example.
<i>M-TAB</i>	Insert a tab character.
<i>C-t</i>	Drag the character before the cursor forward over the character at the cursor, also moving the cursor forward. If the cursor is at the end of the line, then transpose the two characters before it.
<i>M-t</i>	Drag the word behind the cursor past the word in front of the cursor moving the cursor over that word as well.
<i>M-u</i>	Uppercase the characters following the cursor to the end of the current (or following) word, moving the cursor to the end of the word.
<i>M-l</i>	Lowercase the characters following the cursor to the end of the current (or following) word, moving the cursor to the end of the word.
<i>M-c</i>	Uppercase the character following the cursor (or the beginning of the next word if the cursor is between words), moving the cursor to the end of the word.

2.4.4 Letting Readline Type for You

The following commands allow Octave to complete command and variable names for you.

<i>TAB</i>	Attempt to do completion on the text before the cursor. Octave can complete the names of commands and variables.
<i>M-?</i>	List the possible completions of the text before the cursor.

```
val = completion_append_char ()
old_val = completion_append_char (new_val)
completion_append_char (new_val, "local")
```

Query or set the internal character variable that is appended to successful command-line completion attempts.

The default value is " " (a single space).

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

```
completion_matches (hint)
Generate possible completions given hint.
```

This function is provided for the benefit of programs like Emacs which might be controlling Octave and handling user input. The current command number is not incremented when this function is called. This is a feature, not a bug.

2.4.5 Commands for Manipulating the History

Octave normally keeps track of the commands you type so that you can recall previous commands to edit or execute them again. When you exit Octave, the most recent commands you have typed, up to the number specified by the variable `history_size`, are saved in a file. When Octave starts, it loads an initial list of commands from the file named by the variable `history_file`.

Here are the commands for simple browsing and searching the history list.

<code>LFD</code>	
<code>RET</code>	Accept the current line regardless of where the cursor is. If the line is non-empty, add it to the history list. If the line was a history line, then restore the history line to its original state.
<code>C-p</code>	Move ‘up’ through the history list.
<code>C-n</code>	Move ‘down’ through the history list.
<code>M-<</code>	Move to the first line in the history.
<code>M-></code>	Move to the end of the input history, i.e., the line you are entering!
<code>C-r</code>	Search backward starting at the current line and moving ‘up’ through the history as necessary. This is an incremental search.
<code>C-s</code>	Search forward starting at the current line and moving ‘down’ through the history as necessary.

On most terminals, you can also use the up and down arrow keys in place of `C-p` and `C-n` to move through the history list.

In addition to the keyboard commands for moving through the history list, Octave provides three functions for viewing, editing, and re-running chunks of commands from the history list.

```
history
history opt1 ...
h = history ()
h = history (opt1, ...)
```

If invoked with no arguments, `history` displays a list of commands that you have executed.

Valid options are:

<code>n</code>	
<code>-n</code>	Display only the most recent <code>n</code> lines of history.
<code>-c</code>	Clear the history list.
<code>-q</code>	Don’t number the displayed lines of history. This is useful for cutting and pasting commands using the X Window System.

-r file Read the file *file*, appending its contents to the current history list. If the name is omitted, use the default history file (normally `~/.octave_hist`).

-w file Write the current history to the file *file*. If the name is omitted, use the default history file (normally `~/.octave_hist`).

For example, to display the five most recent commands that you have typed without displaying line numbers, use the command `history -q 5`.

If invoked with a single output argument, the history will be saved to that argument as a cell string and will not be output to screen.

See also: [\[edit_history\]](#), page 29, [\[run_history\]](#), page 29.

`edit_history`

`edit_history cmd_number`

`edit_history first last`

Edit the history list using the editor named by the variable `EDITOR`.

The commands to be edited are first copied to a temporary file. When you exit the editor, Octave executes the commands that remain in the file. It is often more convenient to use `edit_history` to define functions rather than attempting to enter them directly on the command line. The block of commands is executed as soon as you exit the editor. To avoid executing any commands, simply delete all the lines from the buffer before leaving the editor.

When invoked with no arguments, edit the previously executed command; With one argument, edit the specified command *cmd_number*; With two arguments, edit the list of commands between *first* and *last*. Command number specifiers may also be negative where -1 refers to the most recently executed command. The following are equivalent and edit the most recently executed command.

`edit_history`

`edit_history -1`

When using ranges, specifying a larger number for the first command than the last command reverses the list of commands before they are placed in the buffer to be edited.

See also: [\[run_history\]](#), page 29, [\[history\]](#), page 28.

`run_history`

`run_history cmd_number`

`run_history first last`

Run commands from the history list.

When invoked with no arguments, run the previously executed command;

With one argument, run the specified command *cmd_number*;

With two arguments, run the list of commands between *first* and *last*. Command number specifiers may also be negative where -1 refers to the most recently executed command. For example, the command

`run_history`

OR

`run_history -1`

executes the most recent command again. The command

```
run_history 13 169
```

executes commands 13 through 169.

Specifying a larger number for the first command than the last command reverses the list of commands before executing them. For example:

```
disp (1)
disp (2)
run_history -1 -2
⇒
2
1
```

See also: [\[edit_history\]](#), page 29, [\[history\]](#), page 28.

Octave also allows you customize the details of when, where, and how history is saved.

```
val = history_save ()
old_val = history_save (new_val)
history_save (new_val, "local")
```

Query or set the internal variable that controls whether commands entered on the command line are saved in the history file.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[history_control\]](#), page 30, [\[history_file\]](#), page 31, [\[history_size\]](#), page 31, [\[history_timestamp_format_string\]](#), page 31.

```
val = history_control ()
old_val = history_control (new_val)
```

Query or set the internal variable that specifies how commands are saved to the history list.

The default value is an empty character string, but may be overridden by the environment variable `OCTAVE_HISTCONTROL`.

The value of `history_control` is a colon-separated list of values controlling how commands are saved on the history list. If the list of values includes `ignorespace`, lines which begin with a space character are not saved in the history list. A value of `ignoredups` causes lines matching the previous history entry to not be saved. A value of `ignoreboth` is shorthand for `ignorespace` and `ignoredups`. A value of `erasedups` causes all previous lines matching the current line to be removed from the history list before that line is saved. Any value not in the above list is ignored. If `history_control` is the empty string, all commands are saved on the history list, subject to the value of `history_save`.

See also: [\[history_file\]](#), page 31, [\[history_size\]](#), page 31, [\[history_timestamp_format_string\]](#), page 31, [\[history_save\]](#), page 30.

```
val = history_file ()  
old_val = history_file (new_val)
```

Query or set the internal variable that specifies the name of the file used to store command history.

The default value is `~/.octave_hist`, but may be overridden by the environment variable `OCTAVE_HISTFILE`.

See also: [\[history_size\]](#), page 31, [\[history_save\]](#), page 30, [\[history_timestamp_format_string\]](#), [\[history_file\]](#), page 31.

```
val = history_size ()  
old_val = history_size (new_val)
```

Query or set the internal variable that specifies how many entries to store in the history file.

The default value is 1000, but may be overridden by the environment variable `OCTAVE_HISTSIZE`.

See also: [\[history_file\]](#), page 31, [\[history_timestamp_format_string\]](#), page 31, [\[history_save\]](#), page 30.

```
val = history_timestamp_format_string ()  
old_val = history_timestamp_format_string (new_val)  
history_timestamp_format_string (new_val, "local")
```

Query or set the internal variable that specifies the format string for the comment line that is written to the history file when Octave exits.

The format string is passed to `strftime`. The default value is

```
"# Octave VERSION, %a %b %d %H:%M:%S %Y %Z <USER@HOST>"
```

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[strftime\]](#), page 887, [\[history_file\]](#), page 31, [\[history_size\]](#), page 31, [\[history_save\]](#), page 30.

```
val = EDITOR ()  
old_val = EDITOR (new_val)  
EDITOR (new_val, "local")
```

Query or set the internal variable that specifies the default text editor.

The default value is taken from the environment variable `EDITOR` when Octave starts. If the environment variable is not initialized, `EDITOR` will be set to "`emacs`".

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[edit\]](#), page 194, [\[edit_history\]](#), page 29.

2.4.6 Customizing readline

Octave uses the GNU Readline library for command-line editing and history features. Readline is very flexible and can be modified through a configuration file of commands (See the GNU Readline library for the exact command syntax). The default configuration file is normally `~/.inputrc`.

Octave provides two commands for initializing Readline and thereby changing the command line behavior.

`readline_read_init_file (file)`

Read the readline library initialization file *file*.

If *file* is omitted, read the default initialization file (normally `~/.inputrc`).

See Section “Readline Init File” in *GNU Readline Library*, for details.

See also: [readline_re_read_init_file], page 32.

`readline_re_read_init_file ()`

Re-read the last readline library initialization file that was read.

See Section “Readline Init File” in *GNU Readline Library*, for details.

See also: [readline_read_init_file], page 32.

2.4.7 Customizing the Prompt

The following variables are available for customizing the appearance of the command-line prompts. Octave allows the prompt to be customized by inserting a number of backslash-escaped special characters that are decoded as follows:

<code>\t</code>	The time.
<code>\d</code>	The date.
<code>\n</code>	Begins a new line by printing the equivalent of a carriage return followed by a line feed.
<code>\s</code>	The name of the program (usually just ‘ <code>octave</code> ’).
<code>\w</code>	The current working directory.
<code>\W</code>	The basename of the current working directory.
<code>\u</code>	The username of the current user.
<code>\h</code>	The hostname, up to the first ‘ <code>.</code> ’.
<code>\H</code>	The hostname.
<code>\#</code>	The command number of this command, counting from when Octave starts.
<code>\!</code>	The history number of this command. This differs from ‘ <code>\#</code> ’ by the number of commands in the history list when Octave starts.
<code>\\$</code>	If the effective UID is 0, a ‘ <code>#</code> ’, otherwise a ‘ <code>\$</code> ’.
<code>\nnn</code>	The character whose character code in octal is <i>nnn</i> .
<code>\\"</code>	A backslash.

```
val = PS1 ()
old_val = PS1 (new_val)
PS1 (new_val, "local")
```

Query or set the primary prompt string.

When executing interactively, Octave displays the primary prompt when it is ready to read a command.

The default value of the primary prompt string is '`'octave:\#>`'. To change it, use a command like

```
PS1 ("\\u@\\H> ")
```

which will result in the prompt '`'boris@kremvax>`' for the user '`'boris'` logged in on the host '`'kremvax.kgb.su'`. Note that two backslashes are required to enter a backslash into a double-quoted character string. See [Chapter 5 \[Strings\], page 67](#).

You can also use ANSI escape sequences if your terminal supports them. This can be useful for coloring the prompt. For example,

```
PS1 ('\\[\\033[01;31m]\\s:\\#> \\[\\033[0m\\]')
```

will give the default Octave prompt a red coloring.

When called from inside a function with the "`"local"` option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[PS2\]](#), page 33, [\[PS4\]](#), page 33.

```
val = PS2 ()
old_val = PS2 (new_val)
PS2 (new_val, "local")
```

Query or set the secondary prompt string.

The secondary prompt is printed when Octave is expecting additional input to complete a command. For example, if you are typing a `for` loop that spans several lines, Octave will print the secondary prompt at the beginning of each line after the first. The default value of the secondary prompt string is "`>`".

When called from inside a function with the "`"local"` option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[PS1\]](#), page 32, [\[PS4\]](#), page 33.

```
val = PS4 ()
old_val = PS4 (new_val)
PS4 (new_val, "local")
```

Query or set the character string used to prefix output produced when echoing commands is enabled.

The default value is "`+`". See [Section 2.4.8 \[Diary and Echo Commands\], page 34](#), for a description of echoing commands.

When called from inside a function with the "`"local"` option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[echo\]](#), page 34, [\[PS1\]](#), page 32, [\[PS2\]](#), page 33.

2.4.8 Diary and Echo Commands

Octave's diary feature allows you to keep a log of all or part of an interactive session by recording the input you type and the output that Octave produces in a separate file.

```
diary
diary on
diary off
diary filename
[status, diaryfile] = diary
```

Record a list of all commands *and* the output they produce, mixed together just as they appear on the terminal.

Valid options are:

- | | |
|-----------------|---|
| on | Start recording a session in a file called <code>diary</code> in the current working directory. |
| off | Stop recording the session in the diary file. |
| <i>filename</i> | Record the session in the file named <i>filename</i> . |

With no input or output arguments, `diary` toggles the current diary state.

If output arguments are requested, `diary` ignores inputs and returns the current status. The boolean *status* indicates whether recording is on or off, and *diaryfile* is the name of the file where the session is stored.

See also: [\[history\]](#), page 28, [\[evalc\]](#), page 161.

Sometimes it is useful to see the commands in a function or script as they are being evaluated. This can be especially helpful for debugging some kinds of problems.

```
echo
echo on
echo off
echo on all
echo off all
echo function on
echo function off
```

Control whether commands are displayed as they are executed.

Valid options are:

- | | |
|---------|---|
| on | Enable echoing of commands as they are executed in script files. |
| off | Disable echoing of commands as they are executed in script files. |
| on all | Enable echoing of commands as they are executed in script files and functions. |
| off all | Disable echoing of commands as they are executed in script files and functions. |

`function on`

Enable echoing of commands as they are executed in the named function.

```
function off
```

Disable echoing of commands as they are executed in the named function.

With no arguments, `echo` toggles the current echo state.

See also: [PS4], page 33.

2.5 How Octave Reports Errors

Octave reports two kinds of errors for invalid programs.

A *parse error* occurs if Octave cannot understand something you have typed. For example, if you misspell a keyword,

```
octave:13> function y = f (x) y = x***2; endfunction
```

Octave will respond immediately with a message like this:

```
parse error:
```

```
syntax error
```

```
>>> function y = f (x) y = x***2; endfunction
```

```
^
```

For most parse errors, Octave uses a caret (^) to mark the point on the line where it was unable to make sense of your input. In this case, Octave generated an error message because the keyword for exponentiation (`**`) was misspelled. It marked the error at the third '*' because the code leading up to this was correct but the final '*' was not understood.

Another class of error message occurs at evaluation time. These errors are called *run-time errors*, or sometimes *evaluation errors*, because they occur when your program is being *run*, or *evaluated*. For example, if after correcting the mistake in the previous function definition, you type

```
octave:13> f ()
```

Octave will respond with

```
error: 'x' undefined near line 1 column 24
error: called from:
error:    f at line 1, column 22
```

This error message has several parts, and gives quite a bit of information to help you locate the source of the error. The messages are generated from the point of the innermost error, and provide a traceback of enclosing expressions and function calls.

In the example above, the first line indicates that a variable named 'x' was found to be undefined near line 1 and column 24 of some function or expression. For errors occurring within functions, lines are counted from the beginning of the file containing the function definition. For errors occurring outside of an enclosing function, the line number indicates the input line number, which is usually displayed in the primary prompt string.

The second and third lines of the error message indicate that the error occurred within the function `f`. If the function `f` had been called from within another function, for example, `g`, the list of errors would have ended with one more line:

```
error:    g at line 1, column 17
```

These lists of function calls make it fairly easy to trace the path your program took before the error occurred, and to correct the error before trying again.

2.6 Executable Octave Programs

Once you have learned Octave, you may want to write self-contained Octave scripts, using the ‘#!’ script mechanism. You can do this on GNU systems and on many Unix systems¹.

Self-contained Octave scripts are useful when you want to write a program which users can invoke without knowing that the program is written in the Octave language. Octave scripts are also used for batch processing of data files. Once an algorithm has been developed and tested in the interactive portion of Octave, it can be committed to an executable script and used again and again on new data files.

As a trivial example of an executable Octave script, you might create a text file named `hello`, containing the following lines:

```
#! octave-interpreter-name -qf
# a sample Octave program
printf ("Hello, world!\n");
```

(where `octave-interpreter-name` should be replaced with the full path and name of your Octave binary). Note that this will only work if ‘#!’ appears at the very beginning of the file. After making the file executable (with the `chmod` command on Unix systems), you can simply type:

```
hello
```

at the shell, and the system will arrange to run Octave as if you had typed:

```
octave hello
```

The line beginning with ‘#!’ lists the full path and filename of an interpreter to be run, and an optional initial command line argument to pass to that interpreter. The operating system then runs the interpreter with the given argument and the full argument list of the executed program. The first argument in the list is the full filename of the Octave executable. The rest of the argument list will either be options to Octave, or data files, or both. The ‘-qf’ options are usually specified in stand-alone Octave programs to prevent them from printing the normal startup message, and to keep them from behaving differently depending on the contents of a particular user’s `~/.octaverc` file. See [Section 2.1 \[Invoking Octave from the Command Line\]](#), page 15.

Note that some operating systems may place a limit on the number of characters that are recognized after ‘#!’. Also, the arguments appearing in a ‘#!’ line are parsed differently by various shells/systems. The majority of them group all the arguments together in one string and pass it to the interpreter as a single argument. In this case, the following script:

```
#! octave-interpreter-name -q -f # comment
```

is equivalent to typing at the command line:

```
octave "-q -f # comment"
```

¹ The ‘#!’ mechanism works on Unix systems derived from Berkeley Unix, System V Release 4, and some System V Release 3 systems.

which will produce an error message. Unfortunately, it is not possible for Octave to determine whether it has been called from the command line or from a ‘#!’ script, so some care is needed when using the ‘#!’ mechanism.

Note that when Octave is started from an executable script, the built-in function `argv` returns a cell array containing the command line arguments passed to the executable Octave script, not the arguments passed to the Octave interpreter on the ‘#!’ line of the script. For example, the following program will reproduce the command line that was used to execute the script, not ‘-qf’.

```
#! /bin/octave -qf
printf ("%s", program_name ());
arg_list = argv ();
for i = 1:nargin
    printf (" %s", arg_list{i});
endfor
printf ("\n");
```

2.7 Comments in Octave Programs

A *comment* is some text that is included in a program for the sake of human readers, and which is NOT an executable part of the program. Comments can explain what the program does, and how it works. Nearly all programming languages have provisions for comments, because programs are typically hard to understand without them.

2.7.1 Single Line Comments

In the Octave language, a comment starts with either the sharp sign character, ‘#’, or the percent symbol ‘%’ and continues to the end of the line. Any text following the sharp sign or percent symbol is ignored by the Octave interpreter and not executed. The following example shows whole line and partial line comments.

```
function countdown
    # Count down for main rocket engines
    disp (3);
    disp (2);
    disp (1);
    disp ("Blast Off!"); # Rocket leaves pad
endfunction
```

2.7.2 Block Comments

Entire blocks of code can be commented by enclosing the code between matching ‘#{’ and ‘#}’ or ‘%{’ and ‘%}’ markers. For example,

```

function quick_countdown
    # Count down for main rocket engines
    disp (3);
    #{
        disp (2);
        disp (1);
    //}
    disp ("Blast Off!"); # Rocket leaves pad
endfunction

```

will produce a very quick countdown from '3' to "Blast Off" as the lines "disp (2);" and "disp (1);" won't be executed.

The block comment markers must appear alone as the only characters on a line (excepting whitespace) in order to be parsed correctly.

2.7.3 Comments and the Help System

The `help` command (see [Section 2.3 \[Getting Help\], page 20](#)) is able to find the first block of comments in a function and return those as a documentation string. This means that the same commands used to get help on built-in functions are available for properly formatted user-defined functions. For example, after defining the function `f` below,

```

function xdot = f (x, t)

# usage: f (x, t)
#
# This function defines the right-hand
# side functions for a set of nonlinear
# differential equations.

r = 0.25;
...
endfunction

```

the command `help f` produces the output

```

usage: f (x, t)

This function defines the right-hand
side functions for a set of nonlinear
differential equations.

```

Although it is possible to put comment lines into keyboard-composed, throw-away Octave programs, it usually isn't very useful because the purpose of a comment is to help you or another person understand the program at a later time.

The `help` parser currently only recognizes single line comments (see [Section 2.7.1 \[Single Line Comments\], page 37](#)) and not block comments for the initial help text.

3 Data Types

All versions of Octave include a number of built-in data types, including real and complex scalars and matrices, character strings, a data structure type, and an array that can contain all data types.

It is also possible to define new specialized data types by writing a small amount of C++ code. On some systems, new data types can be loaded dynamically while Octave is running, so it is not necessary to recompile all of Octave just to add a new type. See [Appendix A \[External Code Interface\], page 945](#), for more information about Octave's dynamic linking capabilities. [Section 3.2 \[User-defined Data Types\], page 44](#), describes what you must do to define a new data type for Octave.

typeinfo ()
typeinfo (*expr*)

Return the type of the expression *expr*, as a string.

If *expr* is omitted, return a cell array of strings containing all the currently installed data types.

See also: [\[class\], page 39](#), [\[isa\], page 39](#).

3.1 Built-in Data Types

The standard built-in data types are real and complex scalars and matrices, ranges, character strings, a data structure type, and cell arrays. Additional built-in data types may be added in future versions. If you need a specialized data type that is not currently provided as a built-in type, you are encouraged to write your own user-defined data type and contribute it for distribution in a future release of Octave.

The data type of a variable can be determined and changed through the use of the following functions.

classname = **class** (*obj*)
class (*s*, *id*)
class (*s*, *id*, *p*, ...)

Return the class of the object *obj*, or create a class with fields from structure *s* and name (string) *id*.

Additional arguments name a list of parent classes from which the new class is derived.

See also: [\[typeinfo\], page 39](#), [\[isa\], page 39](#).

isa (*obj*, *classname*)

Return true if *obj* is an object from the class *classname*.

classname may also be one of the following class categories:

"**float**" Floating point value comprising classes "**double**" and "**single**".

"**integer**"

Integer value comprising classes (u)int8, (u)int16, (u)int32, (u)int64.

"**numeric**"

Numeric value comprising either a floating point or integer value.

If *classname* is a cell array of string, a logical array of the same size is returned, containing true for each class to which *obj* belongs to.

See also: [class], page 39, [typeinfo], page 39.

cast (val, "type")

Convert *val* to data type *type*.

Both *val* and *type* are typically one of the following built-in classes:

```
"double"
"single"
"logical"
"char"
"int8"
"int16"
"int32"
"int64"
"uint8"
"uint16"
"uint32"
"uint64"
```

The value *val* may be modified to fit within the range of the new type.

Examples:

```
cast (-5, "uint8")
      ⇒ 0
cast (300, "int8")
      ⇒ 127
```

Programming Note: This function relies on the object *val* having a conversion method named *type*. User-defined classes may implement only a subset of the full list of types shown above. In that case, it may be necessary to call *cast* twice in order to reach the desired type. For example, the conversion to double is nearly always implemented, but the conversion to uint8 might not be. In that case, the following code will work

```
cast (cast (user_defined_val, "double"), "uint8")
```

See also: [typecast], page 40, [int8], page 54, [uint8], page 55, [int16], page 55, [uint16], page 55, [int32], page 55, [uint32], page 55, [int64], page 55, [uint64], page 55, [double], page 47, [single], page 53, [logical], page 61, [char], page 71, [class], page 39, [typeinfo], page 39.

y = typecast (x, "class")

Return a new array *y* resulting from interpreting the data of *x* in memory as data of the numeric class *class*.

Both the class of *x* and *class* must be one of the built-in numeric classes:

```
"logical"
"char"
"int8"
"int16"
"int32"
"int64"
"uint8"
"uint16"
"uint32"
"uint64"
"double"
"single"
"double complex"
"single complex"
```

the last two are only used with *class*; they indicate that a complex-valued result is requested. Complex arrays are stored in memory as consecutive pairs of real numbers. The sizes of integer types are given by their bit counts. Both logical and char are typically one byte wide; however, this is not guaranteed by C++. If your system is IEEE conformant, single and double will be 4 bytes and 8 bytes wide, respectively. "logical" is not allowed for *class*.

If the input is a row vector, the return value is a row vector, otherwise it is a column vector.

If the bit length of *x* is not divisible by that of *class*, an error occurs.

An example of the use of typecast on a little-endian machine is

```
x = uint16 ([1, 65535]);
typecast (x, "uint8")
⇒ [ 1, 0, 255, 255]
```

See also: [cast], page 40, [bitpack], page 41, [bitunpack], page 42, [swapbytes], page 41.

swapbytes (*x*)

Swap the byte order on values, converting from little endian to big endian and vice versa.

For example:

```
swapbytes (uint16 (1:4))
⇒ 256 512 768 1024
```

See also: [typecast], page 40, [cast], page 40.

y = bitpack (*x, class*)

Return a new array *y* resulting from interpreting the logical array *x* as raw bit patterns for data of the numeric class *class*.

class must be one of the built-in numeric classes:

```
"double"
"single"
"double complex"
"single complex"
"char"
"int8"
"int16"
"int32"
"int64"
"uint8"
"uint16"
"uint32"
"uint64"
```

The number of elements of *x* should be divisible by the bit length of *class*. If it is not, excess bits are discarded. Bits come in increasing order of significance, i.e., *x*(1) is bit 0, *x*(2) is bit 1, etc.

The result is a row vector if *x* is a row vector, otherwise it is a column vector.

See also: [\[bitunpack\]](#), page 42, [\[typecast\]](#), page 40.

y = bitunpack (x)

Return a logical array *y* corresponding to the raw bit patterns of *x*.

x must belong to one of the built-in numeric classes:

```
"double"
"single"
"char"
"int8"
"int16"
"int32"
"int64"
"uint8"
"uint16"
"uint32"
"uint64"
```

The result is a row vector if *x* is a row vector; otherwise, it is a column vector.

See also: [\[bitpack\]](#), page 41, [\[typecast\]](#), page 40.

3.1.1 Numeric Objects

Octave's built-in numeric objects include real, complex, and integer scalars and matrices. All built-in floating point numeric data is currently stored as double precision numbers. On systems that use the IEEE floating point format, values in the range of approximately 2.2251×10^{-308} to 1.7977×10^{308} can be stored, and the relative precision is approximately 2.2204×10^{-16} . The exact values are given by the variables `realmin`, `realmax`, and `eps`, respectively.

Matrix objects can be of any size, and can be dynamically reshaped and resized. It is easy to extract individual rows, columns, or submatrices using a variety of powerful indexing features. See [Section 8.1 \[Index Expressions\]](#), page 139.

See [Chapter 4 \[Numeric Data Types\], page 47](#), for more information.

3.1.2 Missing Data

It is possible to represent missing data explicitly in Octave using `NA` (short for “Not Available”). Missing data can only be represented when data is represented as floating point numbers. In this case missing data is represented as a special case of the representation of `NaN`.

`NA`
`NA (n)`
`NA (n, m)`
`NA (n, m, k, ...)`
`NA (... , class)`

Return a scalar, matrix, or N-dimensional array whose elements are all equal to the special constant used to designate missing values.

Note that `NA` always compares not equal to `NA` (`NA != NA`). To find `NA` values, use the `isna` function.

When called with no arguments, return a scalar with the value ‘`NA`’.

When called with a single argument, return a square matrix with the dimension specified.

When called with more than one scalar argument the first two arguments are taken as the number of rows and columns and any further arguments specify additional matrix dimensions.

The optional argument `class` specifies the return type and may be either “`double`” or “`single`”.

See also: [\[isna\], page 43](#).

`isna (x)`

Return a logical array which is true where the elements of `x` are `NA` (missing) values and false where they are not.

For example:

```
isna ([13, Inf, NA, NaN])
⇒ [ 0, 0, 1, 0 ]
```

See also: [\[isnan\], page 476](#), [\[isinf\], page 476](#), [\[isfinite\], page 477](#).

3.1.3 String Objects

A character string in Octave consists of a sequence of characters enclosed in either double-quote or single-quote marks. Internally, Octave currently stores strings as matrices of characters. All the indexing operations that work for matrix objects also work for strings.

See [Chapter 5 \[Strings\], page 67](#), for more information.

3.1.4 Data Structure Objects

Octave’s data structure type can help you to organize related objects of different types. The current implementation uses an associative array with indices limited to strings, but the syntax is more like C-style structures.

See [Section 6.1 \[Structures\], page 101](#), for more information.

3.1.5 Cell Array Objects

A Cell Array in Octave is general array that can hold any number of different data types.

See [Section 6.3 \[Cell Arrays\]](#), page 114, for more information.

3.2 User-defined Data Types

Someday I hope to expand this to include a complete description of Octave's mechanism for managing user-defined data types. Until this feature is documented here, you will have to make do by reading the code in the `ov.h`, `ops.h`, and related files from Octave's `src` directory.

3.3 Object Sizes

The following functions allow you to determine the size of a variable or expression. These functions are defined for all objects. They return `-1` when the operation doesn't make sense. For example, Octave's data structure type doesn't have rows or columns, so the `rows` and `columns` functions return `-1` for structure arguments.

`ndims (a)`

Return the number of dimensions of `a`.

For any array, the result will always be greater than or equal to 2. Trailing singleton dimensions are not counted, i.e., tailing dimensions `d` greater than 2, for which `size (a, d) = 1`.

```
ndims (ones (4, 1, 2, 1))
⇒ 3
```

See also: [\[size\]](#), page 45.

`columns (a)`

Return the number of columns of `a`. This is equivalent to `size (a, 2)`.

See also: [\[rows\]](#), page 44, [\[size\]](#), page 45, [\[length\]](#), page 45, [\[numel\]](#), page 44, [\[isscalar\]](#), page 64, [\[isvector\]](#), page 64, [\[ismatrix\]](#), page 63.

`rows (a)`

Return the number of rows of `a`. This is equivalent to `size (a, 1)`.

See also: [\[columns\]](#), page 44, [\[size\]](#), page 45, [\[length\]](#), page 45, [\[numel\]](#), page 44, [\[isscalar\]](#), page 64, [\[isvector\]](#), page 64, [\[ismatrix\]](#), page 63.

`numel (a)`

`numel (a, idx1, idx2, ...)`

Return the number of elements in the object `a`.

Optionally, if indices `idx1, idx2, ...` are supplied, return the number of elements that would result from the indexing

```
a(idx1, idx2, ...)
```

Note that the indices do not have to be scalar numbers. For example,

```
a = 1;
b = ones (2, 3);
numel (a, b)
```

will return 6, as this is the number of ways to index with *b*. Or the index could be the string ":" which represents the colon operator. For example,

```
a = ones (5, 3);
numel (a, 2, ":")
```

will return 3 as the second row has three column entries.

This method is also called when an object appears as lvalue with cs-list indexing, i.e., `object{...}` or `object(...).field`.

See also: `[size]`, page 45, `[length]`, page 45, `[ndims]`, page 44.

`length (a)`

Return the length of the object *a*.

The length is 0 for empty objects, 1 for scalars, and the number of elements for vectors. For matrix or N-dimensional objects, the length is the number of elements along the largest dimension (equivalent to `max (size (a))`).

See also: `[numel]`, page 44, `[size]`, page 45.

```
sz = size (a)
dim_sz = size (a, dim)
[rows, cols, ..., dim_N_sz] = size (...)
```

Return a row vector with the size (number of elements) of each dimension for the object *a*.

When given a second argument, *dim*, return the size of the corresponding dimension.

With a single output argument, `size` returns a row vector. When called with multiple output arguments, `size` returns the size of dimension *N* in the *N*th argument. The number of rows, dimension 1, is returned in the first argument, the number of columns, dimension 2, is returned in the second argument, etc. If there are more dimensions in *a* than there are output arguments, `size` returns the total number of elements in the remaining dimensions in the final output argument.

Example 1: single row vector output

```
size ([1, 2; 3, 4; 5, 6])
⇒ [ 3, 2 ]
```

Example 2: number of elements in 2nd dimension (columns)

```
size ([1, 2; 3, 4; 5, 6], 2)
⇒ 2
```

Example 3: number of output arguments == number of dimensions

```
[nr, nc] = size ([1, 2; 3, 4; 5, 6])
⇒ nr = 3
⇒ nc = 2
```

Example 4: number of output arguments < number of dimensions

```
[nr, remainder] = size (ones (2, 3, 4, 5))
⇒ nr = 2
⇒ remainder = 60
```

See also: `[numel]`, page 44, `[ndims]`, page 44, `[length]`, page 45, `[rows]`, page 44, `[columns]`, page 44, `[size_equal]`, page 46, `[common_size]`, page 477.

isempty (a)

Return true if *a* is an empty matrix (any one of its dimensions is zero).

See also: [isnull], page 46, [isa], page 39.

isnull (x)

Return true if *x* is a special null matrix, string, or single quoted string.

Indexed assignment with such a null value on the right-hand side should delete array elements. This function is used in place of `isempty` when overloading the indexed assignment method (`subsasgn`) for user-defined classes. `isnull` is used to distinguish between these two cases:

```
A(I) = []
```

and

```
X = [] ; A(I) = X
```

In the first assignment, the right-hand side is `[]` which is a special null value. As long as the index *I* is not empty, this code should delete elements from *A* rather than perform assignment.

In the second assignment, the right-hand side is empty (because *X* is `[]`), but it is **not** null. This code should assign the empty value to elements in *A*.

An example from Octave's built-in char class demonstrates the interpreter behavior when `isnull` is used correctly.

```
str = "Hello World";
nm = "Wally";
str(7:end) = nm           # indexed assignment
⇒ str = Hello Wally
str(7:end) = ""           # indexed deletion
⇒ str = Hello
```

See also: [isempty], page 46, [isindex], page 145.

sizeof (val)

Return the size of *val* in bytes.

See also: [whos], page 132.

size_equal (a, b, ...)

Return true if the dimensions of all arguments agree.

Trailing singleton dimensions are ignored. When called with a single argument, or no argument, `size_equal` returns true.

See also: [size], page 45, [numel], page 44, [ndims], page 44, [common_size], page 477.

squeeze (x)

Remove singleton dimensions from *x* and return the result.

Note that for compatibility with MATLAB, all objects have a minimum of two dimensions and row vectors are left unchanged.

See also: [reshape], page 482.

4 Numeric Data Types

A *numeric constant* may be a scalar, a vector, or a matrix, and it may contain complex values.

The simplest form of a numeric constant, a scalar, is a single number. Note that by default numeric constants are represented within Octave by IEEE 754 double precision (binary64) floating-point format (complex constants are stored as pairs of binary64 values). It is, however, possible to represent real integers as described in [Section 4.4 \[Integer Data Types\]](#), page 54.

If the numeric constant is a real integer, it can be defined in decimal, hexadecimal, or binary notation. Hexadecimal notation starts with ‘0x’ or ‘0X’, binary notation starts with ‘0b’ or ‘0B’, otherwise decimal notation is assumed. As a consequence, ‘0b’ is not a hexadecimal number, in fact, it is not a valid number at all.

For better readability, digits may be partitioned by the underscore separator ‘_’, which is ignored by the Octave interpreter. Here are some examples of real-valued integer constants, which all represent the same value and are internally stored as binary64:

```
42          # decimal notation
0xA         # hexadecimal notation
0b101010    # binary notation
0b10_1010   # underscore notation
round (42.1) # also binary64
```

In decimal notation, the numeric constant may be denoted as decimal fraction or even in scientific (exponential) notation. Note that this is not possible for hexadecimal or binary notation. Again, in the following example all numeric constants represent the same value:

```
.105
1.05e-1
.00105e+2
```

Unlike most programming languages, complex numeric constants are denoted as the sum of real and imaginary parts. The imaginary part is denoted by a real-valued numeric constant followed immediately by a complex value indicator (‘i’, ‘j’, ‘I’, or ‘J’ which represents $\sqrt{-1}$). No spaces are allowed between the numeric constant and the complex value indicator. Some examples of complex numeric constants that all represent the same value:

```
3 + 42i
3 + 42j
3 + 42I
3 + 42J
3.0 + 42.0i
3.0 + 0x2Ai
3.0 + 0b10_1010i
0.3e1 + 420e-1i

double (x)
Convert x to double precision type.
```

See also: [\[single\]](#), page 53.

```
complex (x)
complex (re, im)
```

Return a complex value from real arguments.

With 1 real argument *x*, return the complex result *x + 0i*.

With 2 real arguments, return the complex result *re + imi*. `complex` can often be more convenient than expressions such as *a + b*i*. For example:

```
complex ([1, 2], [3, 4])
⇒ [ 1 + 3i   2 + 4i ]
```

See also: `[real]`, page 512, `[imag]`, page 512, `[iscomplex]`, page 63, `[abs]`, page 511, `[arg]`, page 511.

4.1 Matrices

It is easy to define a matrix of values in Octave. The size of the matrix is determined automatically, so it is not necessary to explicitly state the dimensions. The expression

```
a = [1, 2; 3, 4]
```

results in the matrix

$$a = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

Elements of a matrix may be arbitrary expressions, provided that the dimensions all make sense when combining the various pieces. For example, given the above matrix, the expression

```
[ a, a ]
```

produces the matrix

```
ans =
```

```
1 2 1 2
3 4 3 4
```

but the expression

```
[ a, 1 ]
```

produces the error

```
error: number of rows must match (1 != 2) near line 13, column 6
```

(assuming that this expression was entered as the first thing on line 13, of course).

Inside the square brackets that delimit a matrix expression, Octave looks at the surrounding context to determine whether spaces and newline characters should be converted into element and row separators, or simply ignored, so an expression like

```
a = [ 1 2
      3 4 ]
```

will work. However, some possible sources of confusion remain. For example, in the expression

```
[ 1 - 1 ]
```

the ‘-’ is treated as a binary operator and the result is the scalar 0, but in the expression

```
[ 1 -1 ]
```

the ‘-’ is treated as a unary operator and the result is the vector [1, -1]. Similarly, the expression

```
[ sin (pi) ]
```

will be parsed as

```
[ sin, (pi) ]
```

and will result in an error since the `sin` function will be called with no arguments. To get around this, you must omit the space between `sin` and the opening parenthesis, or enclose the expression in a set of parentheses:

```
[ (sin (pi)) ]
```

Whitespace surrounding the single quote character (‘ ’, used as a transpose operator and for delimiting character strings) can also cause confusion. Given `a = 1`, the expression

```
[ 1 a' ]
```

results in the single quote character being treated as a transpose operator and the result is the vector [1, 1], but the expression

```
[ 1 a' ]
```

produces the error message

```
parse error:
```

```
syntax error
```

```
>>> [ 1 a' ]
```

because not doing so would cause trouble when parsing the valid expression

```
[ a 'foo' ]
```

For clarity, it is probably best to always use commas and semicolons to separate matrix elements and rows.

The maximum number of elements in a matrix is fixed when Octave is compiled. The allowable number can be queried with the function `sizemax`. Note that other factors, such as the amount of memory available on your machine, may limit the maximum size of matrices to something smaller.

`sizemax ()`

Return the largest value allowed for the size of an array.

If Octave is compiled with 64-bit indexing, the result is of class `int64`, otherwise it is of class `int32`. The maximum array size is slightly smaller than the maximum value allowable for the relevant class as reported by `intmax`.

See also: [\[intmax\]](#), [page 55](#).

When you type a matrix or the name of a variable whose value is a matrix, Octave responds by printing the matrix in with neatly aligned rows and columns. If the rows of the matrix are too large to fit on the screen, Octave splits the matrix and displays a header before each section to indicate which columns are being displayed. You can use the following variables to control the format of the output.

```
val = output_precision ()
old_val = output_precision (new_val)
output_precision (new_val, "local")
```

Query or set the internal variable that specifies the minimum number of significant figures to display for numeric output.

Note that regardless of the value set for `output_precision`, the number of digits of precision displayed is limited to 16 for double precision values and 7 for single precision values.

When called from inside a function with the "`local`" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[format\]](#), page 252, [\[fixed_point_format\]](#), page 50.

It is possible to achieve a wide range of output styles by using different values of `output_precision`. Reasonable combinations can be set using the `format` function. See [Section 14.1 \[Basic Input and Output\]](#), page 251.

```
val = split_long_rows ()
old_val = split_long_rows (new_val)
split_long_rows (new_val, "local")
```

Query or set the internal variable that controls whether rows of a matrix may be split when displayed to a terminal window.

If the rows are split, Octave will display the matrix in a series of smaller pieces, each of which can fit within the limits of your terminal width and each set of rows is labeled so that you can easily see which columns are currently being displayed. For example:

```
octave:13> rand (2,10)
ans =
```

Columns 1 through 6:

```
0.75883  0.93290  0.40064  0.43818  0.94958  0.16467
0.75697  0.51942  0.40031  0.61784  0.92309  0.40201
```

Columns 7 through 10:

```
0.90174  0.11854  0.72313  0.73326
0.44672  0.94303  0.56564  0.82150
```

When called from inside a function with the "`local`" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[format\]](#), page 252.

Octave automatically switches to scientific notation when values become very large or very small. This guarantees that you will see several significant figures for every value in a matrix. If you would prefer to see all values in a matrix printed in a fixed point format, you can set the built-in variable `fixed_point_format` to a nonzero value. But doing so is not recommended, because it can produce output that can easily be misinterpreted.

```
val = fixed_point_format ()
old_val = fixed_point_format (new_val)
fixed_point_format (new_val, "local")
```

Query or set the internal variable that controls whether Octave will use a scaled format to print matrix values.

The scaled format prints a scaling factor on the first line of output chosen such that the largest matrix element can be written with a single leading digit. For example:

```
fixed_point_format (true)
logspace (1, 7, 5)'
ans =

1.0e+07  *

0.00000
0.00003
0.00100
0.03162
1.00000
```

Notice that the first value appears to be 0 when it is actually 1. Because of the possibility for confusion you should be careful about enabling `fixed_point_format`.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[format\]](#), page 252, [\[output_precision\]](#), page 49.

4.1.1 Empty Matrices

A matrix may have one or both dimensions zero, and operations on empty matrices are handled as described by Carl de Boor in *An Empty Exercise*, SIGNUM, Volume 25, pages 2-6, 1990 and C. N. Nett and W. M. Haddad, in *A System-Theoretic Appropriate Realization of the Empty Matrix Concept*, IEEE Transactions on Automatic Control, Volume 38, Number 5, May 1993. Briefly, given a scalar s , an $m \times n$ matrix $M_{m \times n}$, and an $m \times n$ empty matrix $[]_{m \times n}$ (with either one or both dimensions equal to zero), the following are true:

$$\begin{aligned}s \cdot []_{m \times n} &= []_{m \times n} \cdot s = []_{m \times n} \\ []_{m \times n} + []_{m \times n} &= []_{m \times n} \\ []_{0 \times m} \cdot M_{m \times n} &= []_{0 \times n} \\ M_{m \times n} \cdot []_{n \times 0} &= []_{m \times 0} \\ []_{m \times 0} \cdot []_{0 \times n} &= 0_{m \times n}\end{aligned}$$

By default, dimensions of the empty matrix are printed along with the empty matrix symbol, '`[]`'. The built-in variable `print_empty_dimensions` controls this behavior.

```
val = print_empty_dimensions ()
old_val = print_empty_dimensions (new_val)
```

```
print_empty_dimensions (new_val, "local")
```

Query or set the internal variable that controls whether the dimensions of empty matrices are printed along with the empty matrix symbol, ‘[]’.

For example, the expression

```
zeros (3, 0)
```

will print

```
ans = [] (3x0)
```

When called from inside a function with the “local” option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[format\]](#), [page 252](#).

Empty matrices may also be used in assignment statements as a convenient way to delete rows or columns of matrices. See [Section 8.6 \[Assignment Expressions\]](#), [page 156](#).

When Octave parses a matrix expression, it examines the elements of the list to determine whether they are all constants. If they are, it replaces the list with a single matrix constant.

4.2 Ranges

A range is a convenient way to write a row vector with evenly spaced elements. A range expression is defined by the value of the first element in the range, an optional value for the increment between elements, and a maximum value which the elements of the range will not exceed. The base, increment, and limit are separated by colons (the ‘:’ character) and may contain any arithmetic expressions and function calls. If the increment is omitted, it is assumed to be 1. For example, the range

```
1 : 5
```

defines the set of values [1, 2, 3, 4, 5], and the range

```
1 : 3 : 5
```

defines the set of values [1, 4].

Although a range constant specifies a row vector, Octave does *not* normally convert range constants to vectors unless it is necessary to do so. This allows you to write a constant like 1 : 10000 without using 80,000 bytes of storage on a typical 32-bit workstation.

A common example of when it does become necessary to convert ranges into vectors occurs when they appear within a vector (i.e., inside square brackets). For instance, whereas

```
x = 0 : 0.1 : 1;
```

defines x to be a variable of type `range` and occupies 24 bytes of memory, the expression

```
y = [ 0 : 0.1 : 1];
```

defines y to be of type `matrix` and occupies 88 bytes of memory.

This space saving optimization may be disabled using the function `disable_range`.

```
val = disable_range ()
old_val = disable_range (new_val)
```

disable_range (*new_val*, "local")

Query or set the internal variable that controls whether ranges are stored in a special space-efficient format.

The default value is true. If this option is disabled Octave will store ranges as full matrices.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [disable_diagonal_matrix], page 609, [disable_permutation_matrix], page 609.

Note that the upper (or lower, if the increment is negative) bound on the range is not always included in the set of values, and that ranges defined by floating point values can produce surprising results because Octave uses floating point arithmetic to compute the values in the range. If it is important to include the endpoints of a range and the number of elements is known, you should use the `linspace` function instead (see Section 16.3 [Special Utility Matrices], page 489).

When adding a scalar to a range, subtracting a scalar from it (or subtracting a range from a scalar) and multiplying by scalar, Octave will attempt to avoid unpacking the range and keep the result as a range, too, if it can determine that it is safe to do so. For instance, doing

```
a = 2*(1:1e7) - 1;
```

will produce the same result as `1:2:2e7-1`, but without ever forming a vector with ten million elements.

Using zero as an increment in the colon notation, as `1:0:1` is not allowed, because a division by zero would occur in determining the number of range elements. However, ranges with zero increment (i.e., all elements equal) are useful, especially in indexing, and Octave allows them to be constructed using the built-in function `ones`. Note that because a range must be a row vector, `ones (1, 10)` produces a range, while `ones (10, 1)` does not.

When Octave parses a range expression, it examines the elements of the expression to determine whether they are all constants. If they are, it replaces the range expression with a single range constant.

4.3 Single Precision Data Types

Octave includes support for single precision data types, and most of the functions in Octave accept single precision values and return single precision answers. A single precision variable is created with the `single` function.

single (*x*)

Convert *x* to single precision type.

See also: [double], page 47.

for example:

```
sngl = single (rand (2, 2))
⇒ sngl =
  0.37569  0.92982
  0.11962  0.50876
class (sngl)
⇒ single
```

Many functions can also return single precision values directly. For example

```
ones (2, 2, "single")
zeros (2, 2, "single")
eye (2, 2, "single")
rand (2, 2, "single")
NaN (2, 2, "single")
NA (2, 2, "single")
Inf (2, 2, "single")
```

will all return single precision matrices.

4.4 Integer Data Types

Octave supports integer matrices as an alternative to using double precision. It is possible to use both signed and unsigned integers represented by 8, 16, 32, or 64 bits. It should be noted that most computations require floating point data, meaning that integers will often change type when involved in numeric computations. For this reason integers are most often used to store data, and not for calculations.

In general most integer matrices are created by casting existing matrices to integers. The following example shows how to cast a matrix into 32 bit integers.

```
float = rand (2, 2)
⇒ float = 0.37569  0.92982
          0.11962  0.50876
integer = int32 (float)
⇒ integer = 0  1
           0  1
```

As can be seen, floating point values are rounded to the nearest integer when converted.

`isinteger (x)`

Return true if *x* is an integer object (`int8`, `uint8`, `int16`, etc.).

Note that `isinteger (14)` is false because numeric constants in Octave are double precision floating point values.

See also: [\[isfloat\]](#), page 63, [\[ischar\]](#), page 68, [\[islogical\]](#), page 63, [\[isstring\]](#), page 68, [\[isnumeric\]](#), page 63, [\[isa\]](#), page 39.

`int8 (x)`

Convert *x* to 8-bit integer type.

See also: [\[uint8\]](#), page 55, [\[int16\]](#), page 55, [\[uint16\]](#), page 55, [\[int32\]](#), page 55, [\[uint32\]](#), page 55, [\[int64\]](#), page 55, [\[uint64\]](#), page 55.

uint8 (x)

Convert x to unsigned 8-bit integer type.

See also: [int8], page 54, [int16], page 55, [uint16], page 55, [int32], page 55, [uint32], page 55, [int64], page 55, [uint64], page 55.

int16 (x)

Convert x to 16-bit integer type.

See also: [int8], page 54, [uint8], page 55, [uint16], page 55, [int32], page 55, [uint32], page 55, [int64], page 55, [uint64], page 55.

uint16 (x)

Convert x to unsigned 16-bit integer type.

See also: [int8], page 54, [uint8], page 55, [int16], page 55, [int32], page 55, [uint32], page 55, [int64], page 55, [uint64], page 55.

int32 (x)

Convert x to 32-bit integer type.

See also: [int8], page 54, [uint8], page 55, [int16], page 55, [uint16], page 55, [uint32], page 55, [int64], page 55, [uint64], page 55.

uint32 (x)

Convert x to unsigned 32-bit integer type.

See also: [int8], page 54, [uint8], page 55, [int16], page 55, [uint16], page 55, [int32], page 55, [int64], page 55, [uint64], page 55.

int64 (x)

Convert x to 64-bit integer type.

See also: [int8], page 54, [uint8], page 55, [int16], page 55, [uint16], page 55, [int32], page 55, [int64], page 55, [uint64], page 55.

uint64 (x)

Convert x to unsigned 64-bit integer type.

See also: [int8], page 54, [uint8], page 55, [int16], page 55, [uint16], page 55, [int32], page 55, [uint32], page 55, [int64], page 55.

intmax ()**intmax ("type")****intmax (var)**

Return the largest integer that can be represented by a specific integer type.

The input is either a string "type" specifying an integer type, or it is an existing integer variable var.

Possible values for type are

"int8" signed 8-bit integer.

"int16" signed 16-bit integer.

"int32" signed 32-bit integer.

```
"int64"    signed 64-bit integer.
"uint8"    unsigned 8-bit integer.
"uint16"   unsigned 16-bit integer.
"uint32"   unsigned 32-bit integer.
"uint64"   unsigned 64-bit integer.
```

The default for *type* is "int32".

Example Code - query an existing variable

```
x = int8 (1);
intmax (x)
⇒ 127
```

See also: [\[intmin\]](#), page 56, [\[flintmax\]](#), page 56.

```
intmin ()
intmin ("type")
intmin (var)
```

Return the smallest integer that can be represented by a specific integer type.

The input is either a string "type" specifying an integer type, or it is an existing integer variable var.

Possible values for *type* are

```
"int8"    signed 8-bit integer.
"int16"   signed 16-bit integer.
"int32"   signed 32-bit integer.
"int64"   signed 64-bit integer.
"uint8"   unsigned 8-bit integer.
"uint16"  unsigned 16-bit integer.
"uint32"  unsigned 32-bit integer.
"uint64"  unsigned 64-bit integer.
```

The default for *type* is "int32".

Example Code - query an existing variable

```
x = int8 (1);
intmin (x)
⇒ -128
```

See also: [\[intmax\]](#), page 55, [\[flintmax\]](#), page 56.

```
flintmax ()
flintmax ("double")
flintmax ("single")
flintmax (var)
```

Return the largest integer that can be represented consecutively in a floating point value.

The input is either a string specifying a floating point type, or it is an existing floating point variable *var*.

The default type is "double", but "single" is a valid option. On IEEE 754 compatible systems, `flintmax` is 2^{53} for "double" and 2^{24} for "single".

Example Code - query an existing variable

```
x = single (1);
flintmax (x)
⇒ 16777216
```

See also: [\[intmax\]](#), page 55, [\[realmax\]](#), page 543, [\[realmin\]](#), page 543.

4.4.1 Integer Arithmetic

While many numerical computations can't be carried out in integers, Octave does support basic operations like addition and multiplication on integers. The operators `+`, `-`, `.*`, and `./` work on integers of the same type. So, it is possible to add two 32 bit integers, but not to add a 32 bit integer and a 16 bit integer.

When doing integer arithmetic one should consider the possibility of underflow and overflow. This happens when the result of the computation can't be represented using the chosen integer type. As an example it is not possible to represent the result of $10 - 20$ when using unsigned integers. Octave makes sure that the result of integer computations is the integer that is closest to the true result. So, the result of $10 - 20$ when using unsigned integers is zero.

When doing integer division Octave will round the result to the nearest integer. This is different from most programming languages, where the result is often floored to the nearest integer. So, the result of `int32 (5) ./ int32 (8)` is 1.

`idivide (x, y, op)`

Integer division with different rounding rules.

The standard behavior of integer division such as `a ./ b` is to round the result to the nearest integer. This is not always the desired behavior and `idivide` permits integer element-by-element division to be performed with different treatment for the fractional part of the division as determined by the *op* flag. *op* is a string with one of the values:

- "fix" Calculate `a ./ b` with the fractional part rounded towards zero.
- "round" Calculate `a ./ b` with the fractional part rounded towards the nearest integer.
- "floor" Calculate `a ./ b` with the fractional part rounded towards negative infinity.
- "ceil" Calculate `a ./ b` with the fractional part rounded towards positive infinity.

If *op* is not given it defaults to "fix". An example demonstrating these rounding rules is

```

idivide (int8([-3, 3]), int8(4), "fix")
⇒ 0 0
idivide (int8([-3, 3]), int8(4), "round")
⇒ -1 1
idivide (int8([-3, 3]), int8(4), "floor")
⇒ -1 0
idivide (int8([-3, 3]), int8(4), "ceil")
⇒ 0 1

```

See also: [\[ldivide\]](#), page 150, [\[rdivide\]](#), page 151.

4.5 Bit Manipulations

Octave provides a number of functions for the manipulation of numeric values on a bit by bit basis. The basic functions to set and obtain the values of individual bits are `bitset` and `bitget`.

```

C = bitset (A, n)
C = bitset (A, n, val)
    Set or reset bit(s) n of the unsigned integers in A.

```

The least significant bit is $n = 1$. $val = 0$ resets bits and $val = 1$ sets bits. If no val is specified it defaults to 1 (set bit). All inputs must be the same size or scalars.

Example 1: Set multiple bits

```

x = bitset (1, 3:5)
⇒ x =

```

```
5 9 17
```

```

dec2bin (x)
⇒
00101
01001
10001

```

Example 2: Reset and set bits

```

x = bitset ([15 14], 1, [0 1])
⇒ x =

```

```
14 15
```

See also: [\[bitand\]](#), page 59, [\[bitor\]](#), page 59, [\[bitxor\]](#), page 59, [\[bitget\]](#), page 58, [\[bitcmp\]](#), page 59, [\[bitshift\]](#), page 60, [\[intmax\]](#), page 55, [\[flintmax\]](#), page 56.

```

c = bitget (A, n)
    Return the status of bit(s) n of the unsigned integers in A.

```

The least significant bit is $n = 1$.

```

bitget (100, 8:-1:1)
⇒ 0 1 1 0 0 1 0 0

```

See also: [bitand], page 59, [bitor], page 59, [bitxor], page 59, [bitset], page 58, [bitcmp], page 59, [bitshift], page 60, [intmax], page 55, [flintmax], page 56.

The arguments to all of Octave's bitwise operations can be scalar or arrays, except for **bitcmp**, whose *k* argument must a scalar. In the case where more than one argument is an array, then all arguments must have the same shape, and the bitwise operator is applied to each of the elements of the argument individually. If at least one argument is a scalar and one an array, then the scalar argument is duplicated. Therefore

```
bitget (100, 8:-1:1)
```

is the same as

```
bitget (100 * ones (1, 8), 8:-1:1)
```

It should be noted that all values passed to the bit manipulation functions of Octave are treated as integers. Therefore, even though the example for **bitset** above passes the floating point value 10, it is treated as the bits [1, 0, 1, 0] rather than the bits of the native floating point format representation of 10.

As the maximum value that can be represented by a number is important for bit manipulation, particularly when forming masks, Octave supplies two utility functions: **flintmax** for floating point integers, and **intmax** for integer objects (**uint8**, **int64**, etc.).

Octave also includes the basic bitwise 'and', 'or', and 'exclusive or' operators.

bitand (x, y)

Return the bitwise AND of non-negative integers.

x, y must be in the range [0,intmax]

See also: [bitor], page 59, [bitxor], page 59, [bitset], page 58, [bitget], page 58, [bitcmp], page 59, [bitshift], page 60, [intmax], page 55, [flintmax], page 56.

bitor (x, y)

Return the bitwise OR of non-negative integers *x* and *y*.

See also: [bitor], page 59, [bitxor], page 59, [bitset], page 58, [bitget], page 58, [bitcmp], page 59, [bitshift], page 60, [intmax], page 55, [flintmax], page 56.

bitxor (x, y)

Return the bitwise XOR of non-negative integers *x* and *y*.

See also: [bitand], page 59, [bitor], page 59, [bitset], page 58, [bitget], page 58, [bitcmp], page 59, [bitshift], page 60, [intmax], page 55, [flintmax], page 56.

The bitwise 'not' operator is a unary operator that performs a logical negation of each of the bits of the value. For this to make sense, the mask against which the value is negated must be defined. Octave's bitwise 'not' operator is **bitcmp**.

bitcmp (A, k)

Return the *k*-bit complement of integers in *A*.

If *k* is omitted *k* = log2 (flintmax) + 1 is assumed.

```
bitcmp (7,4)
⇒ 8
dec2bin (11)
⇒ 1011
dec2bin (bitcmp (11, 6))
⇒ 110100
```

See also: [bitand], page 59, [bitor], page 59, [bitxor], page 59, [bitset], page 58, [bitget], page 58, [bitcmp], page 59, [bitshift], page 60, [flintmax], page 56.

Octave also includes the ability to left-shift and right-shift values bitwise.

```
bitshift (a, k)
bitshift (a, k, n)
```

Return a k bit shift of n -digit unsigned integers in a .

A positive k leads to a left shift; A negative value to a right shift.

If n is omitted it defaults to 64. n must be in the range [1,64].

```
bitshift (eye (3), 1)
⇒
2 0 0
0 2 0
0 0 2

bitshift (10, [-2, -1, 0, 1, 2])
⇒ 2 5 10 20 40
```

See also: [bitand], page 59, [bitor], page 59, [bitxor], page 59, [bitset], page 58, [bitget], page 58, [bitcmp], page 59, [intmax], page 55, [flintmax], page 56.

Bits that are shifted out of either end of the value are lost. Octave also uses arithmetic shifts, where the sign bit of the value is kept during a right shift. For example:

```
bitshift (-10, -1)
⇒ -5
bitshift (int8 (-1), -1)
⇒ -1
```

Note that `bitshift (int8 (-1), -1)` is `-1` since the bit representation of `-1` in the `int8` data type is `[1, 1, 1, 1, 1, 1, 1, 1]`.

4.6 Logical Values

Octave has built-in support for logical values, i.e., variables that are either `true` or `false`. When comparing two variables, the result will be a logical value whose value depends on whether or not the comparison is true.

The basic logical operations are `&`, `|`, and `!`, which correspond to “Logical And”, “Logical Or”, and “Logical Negation”. These operations all follow the usual rules of logic.

It is also possible to use logical values as part of standard numerical calculations. In this case `true` is converted to 1, and `false` to 0, both represented using double precision floating point numbers. So, the result of `true*22 - false/6` is 22.

Logical values can also be used to index matrices and cell arrays. When indexing with a logical array the result will be a vector containing the values corresponding to `true` parts of the logical array. The following example illustrates this.

```
data = [ 1, 2; 3, 4 ];
idx = (data <= 2);
data(idx)
⇒ ans = [ 1; 2 ]
```

Instead of creating the `idx` array it is possible to replace `data(idx)` with `data(data <= 2)` in the above code.

Logical values can also be constructed by casting numeric objects to logical values, or by using the `true` or `false` functions.

`logical (x)`

Convert the numeric object `x` to logical type.

Any nonzero values will be converted to true (1) while zero values will be converted to false (0). The non-numeric value NaN cannot be converted and will produce an error.

Compatibility Note: Octave accepts complex values as input, whereas MATLAB issues an error.

See also: [\[double\]](#), page 47, [\[single\]](#), page 53, [\[char\]](#), page 71.

`true (x)`

`true (n, m)`

`true (n, m, k, ...)`

Return a matrix or N-dimensional array whose elements are all logical 1.

If invoked with a single scalar integer argument, return a square matrix of the specified size.

If invoked with two or more scalar integer arguments, or a vector of integer values, return an array with given dimensions.

See also: [\[false\]](#), page 61.

`false (x)`

`false (n, m)`

`false (n, m, k, ...)`

Return a matrix or N-dimensional array whose elements are all logical 0.

If invoked with a single scalar integer argument, return a square matrix of the specified size.

If invoked with two or more scalar integer arguments, or a vector of integer values, return an array with given dimensions.

See also: [\[true\]](#), page 61.

4.7 Promotion and Demotion of Data Types

Many operators and functions can work with mixed data types. For example,

```
uint8 (1) + 1
⇒ 2
```

where the above operator works with an 8-bit integer and a double precision value and returns an 8-bit integer value. Note that the type is demoted to an 8-bit integer, rather than promoted to a double precision value as might be expected. The reason is that if Octave promoted values in expressions like the above with all numerical constants would need to be explicitly cast to the appropriate data type like

```
uint8 (1) + uint8 (1)
⇒ 2
```

which becomes difficult for the user to apply uniformly and might allow hard to find bugs to be introduced. The same applies to single precision values where a mixed operation such as

```
single (1) + 1
⇒ 2
```

returns a single precision value. The mixed operations that are valid and their returned data types are

Mixed Operation	Result
double OP single	single
double OP integer	integer
double OP char	double
double OP logical	double
single OP integer	integer
single OP char	single
single OP logical	single

The same logic applies to functions with mixed arguments such as

```
min (single (1), 0)
⇒ 0
```

where the returned value is single precision.

In the case of mixed type indexed assignments, the type is not changed. For example,

```
x = ones (2, 2);
x(1, 1) = single (2)
⇒ x = 2    1
      1    1
```

where `x` remains of the double precision type.

4.8 Predicates for Numeric Objects

Since the type of a variable may change during the execution of a program, it can be necessary to do type checking at run-time. Doing this also allows you to change the behavior of a function depending on the type of the input. As an example, this naive implementation

of **abs** returns the absolute value of the input if it is a real number, and the length of the input if it is a complex number.

```
function a = abs (x)
    if (isreal (x))
        a = sign (x) .* x;
    elseif (iscomplex (x))
        a = sqrt (real(x).^2 + imag(x).^2);
    endif
endfunction
```

The following functions are available for determining the type of a variable.

isnumeric (x)

Return true if *x* is a numeric object, i.e., an integer, real, or complex array.

Logical and character arrays are not considered to be numeric.

See also: [isinteger], page 54, [isfloat], page 63, [isreal], page 63, [iscomplex], page 63, [ischar], page 68, [islogical], page 63, [isstring], page 68, [iscell], page 116, [isstruct], page 109, [isa], page 39.

islogical (x)

isbool (x)

Return true if *x* is a logical object.

See also: [ischar], page 68, [isfloat], page 63, [isinteger], page 54, [isstring], page 68, [isnumeric], page 63, [isa], page 39.

isfloat (x)

Return true if *x* is a floating-point numeric object.

Objects of class double or single are floating-point objects.

See also: [isinteger], page 54, [ischar], page 68, [islogical], page 63, [isnumeric], page 63, [isstring], page 68, [isa], page 39.

isreal (x)

Return true if *x* is a non-complex matrix or scalar.

For compatibility with MATLAB, this includes logical and character matrices.

See also: [iscomplex], page 63, [isnumeric], page 63, [isa], page 39.

iscomplex (x)

Return true if *x* is a complex-valued numeric object.

See also: [isreal], page 63, [isnumeric], page 63, [ischar], page 68, [isfloat], page 63, [islogical], page 63, [isstring], page 68, [isa], page 39.

ismatrix (a)

Return true if *a* is a 2-D array, i.e., **size (a)** returns [M N] with non-negative M and N.

See also: [isscalar], page 64, [isvector], page 64, [iscell], page 116, [isstruct], page 109, [issparse], page 624, [isa], page 39.

isvector (x)

Return true if x is a vector.

A vector is a 2-D array where one of the dimensions is equal to 1. As a consequence a 1×1 array, or scalar, is also a vector.

See also: [isscalar], page 64, [ismatrix], page 63, [size], page 45, [rows], page 44, [columns], page 44, [length], page 45.

isrow (x)

Return true if x is a row vector, i.e., `size (x)` returns $[1 N]$ with non-negative N .

See also: [iscolumn], page 64, [isscalar], page 64, [isvector], page 64, [ismatrix], page 63.

iscolumn (x)

Return true if x is a column vector, i.e., `size (x)` returns $[N 1]$ with non-negative N .

See also: [isrow], page 64, [isscalar], page 64, [isvector], page 64, [ismatrix], page 63.

isscalar (x)

Return true if x is a scalar, i.e., `size (x)` returns $[1 1]$.

See also: [isvector], page 64, [ismatrix], page 63.

issquare (x)

Return true if x is a square matrix, i.e., `size (x)` returns $[N N]$ with non-negative N .

See also: [isscalar], page 64, [isvector], page 64, [ismatrix], page 63, [size], page 45.

issymmetric (A)**issymmetric (A, tol)****issymmetric (A, "skew")****issymmetric (A, "skew", tol)**

Return true if A is a symmetric or skew-symmetric matrix within the tolerance specified by tol .

The default tolerance is zero (uses faster code).

The type of symmetry to check may be specified with the additional input "`nonskew`" (default) for regular symmetry or "`skew`" for skew-symmetry.

Background: A matrix is symmetric if the transpose of the matrix is equal to the original matrix: $A == A.^.$. If a tolerance is given then symmetry is determined by `norm (A - A.^., Inf) / norm (A, Inf) < tol`.

A matrix is skew-symmetric if the transpose of the matrix is equal to the negative of the original matrix: $A == -A.^.$. If a tolerance is given then skew-symmetry is determined by `norm (A + A.^., Inf) / norm (A, Inf) < tol`.

See also: [ishermitian], page 64, [isdefinite], page 65.

ishermitian (A)**ishermitian (A, tol)****ishermitian (A, "skew")**

ishermitian (*A*, "skew", *tol*)

Return true if *A* is a Hermitian or skew-Hermitian matrix within the tolerance specified by *tol*.

The default tolerance is zero (uses faster code).

The type of symmetry to check may be specified with the additional input "nonskew" (default) for regular Hermitian or "skew" for skew-Hermitian.

Background: A matrix is Hermitian if the complex conjugate transpose of the matrix is equal to the original matrix: $A == A'$. If a tolerance is given then the calculation is $\text{norm}(A - A', \text{Inf}) / \text{norm}(A, \text{Inf}) < tol$.

A matrix is skew-Hermitian if the complex conjugate transpose of the matrix is equal to the negative of the original matrix: $A == -A'$. If a tolerance is given then the calculation is $\text{norm}(A + A', \text{Inf}) / \text{norm}(A, \text{Inf}) < tol$.

See also: [issymmetric], page 64, [isdefinite], page 65.

isdefinite (*A*)**isdefinite (*A*, *tol*)**

Return true if *A* is symmetric positive definite matrix within the tolerance specified by *tol*.

If *tol* is omitted, use a tolerance of $100 * \text{eps} * \text{norm}(A, \text{"fro"})$.

Background: A positive definite matrix has eigenvalues which are all greater than zero. A positive semi-definite matrix has eigenvalues which are all greater than or equal to zero. The matrix *A* is very likely to be positive semi-definite if the following two conditions hold for a suitably small tolerance *tol*.

```
isdefinite (A) => 0
isdefinite (A + 5*tol, tol) => 1
```

See also: [issymmetric], page 64, [ishermitian], page 64.

isbanded (*A*, *lower*, *upper*)

Return true if *A* is a matrix with entries confined between *lower* diagonals below the main diagonal and *upper* diagonals above the main diagonal.

lower and *upper* must be non-negative integers.

See also: [isdiag], page 65, [istril], page 65, [istriu], page 66, [bandwidth], page 546.

isdiag (*A*)

Return true if *A* is a diagonal matrix.

See also: [isbanded], page 65, [istril], page 65, [istriu], page 66, [diag], page 489, [bandwidth], page 546.

istril (*A*)

Return true if *A* is a lower triangular matrix.

A lower triangular matrix has nonzero entries only on the main diagonal and below.

See also: [istriu], page 66, [isbanded], page 65, [isdiag], page 65, [tril], page 487, [bandwidth], page 546.

istriu (A)

Return true if A is an upper triangular matrix.

An upper triangular matrix has nonzero entries only on the main diagonal and above.

See also: [isdiag], page 65, [isbanded], page 65, [istril], page 65, [triu], page 487, [bandwidth], page 546.

isprime (x)

Return a logical array which is true where the elements of x are prime numbers and false where they are not.

A prime number is conventionally defined as a positive integer greater than 1 (e.g., 2, 3, ...) which is divisible only by itself and 1. Octave extends this definition to include both negative integers and complex values. A negative integer is prime if its positive counterpart is prime. This is equivalent to `isprime (abs (x))`.

If `class (x)` is complex, then primality is tested in the domain of Gaussian integers (https://en.wikipedia.org/wiki/Gaussian_integer). Some non-complex integers are prime in the ordinary sense, but not in the domain of Gaussian integers. For example, $5 = (1+2i) * (1-2i)$ shows that 5 is not prime because it has a factor other than itself and 1. Exercise caution when testing complex and real values together in the same matrix.

Examples:

```
isprime (1:6)
⇒ 0 1 1 0 1 0
isprime ([i, 2, 3, 5])
⇒ 0 0 1 0
```

Programming Note: `isprime` is appropriate if the maximum value in x is not too large ($< 1e15$). For larger values special purpose factorization code should be used.

Compatibility Note: `matlab` does not extend the definition of prime numbers and will produce an error if given negative or complex inputs.

See also: [primes], page 525, [factor], page 524, [gcd], page 524, [lcm], page 524.

If instead of knowing properties of variables, you wish to know which variables are defined and to gather other information about the workspace itself, see [Section 7.3 \[Status of Variables\]](#), page 132.

5 Strings

A *string constant* consists of a sequence of characters enclosed in either double-quote or single-quote marks. For example, both of the following expressions

```
"parrot"
'parrot'
```

represent the string whose contents are ‘parrot’. Strings in Octave can be of any length.

Since the single-quote mark is also used for the transpose operator (see [Section 8.3 \[Arithmetic Ops\], page 149](#)) but double-quote marks have no other purpose in Octave, it is best to use double-quote marks to denote strings.

Strings can be concatenated using the notation for defining matrices. For example, the expression

```
[ "foo" , "bar" , "baz" ]
```

produces the string whose contents are ‘foobarbaz’. See [Chapter 4 \[Numeric Data Types\], page 47](#), for more information about creating matrices.

While strings can in principal store arbitrary content, most functions expect them to be UTF-8 encoded Unicode strings.

5.1 Escape Sequences in String Constants

In double-quoted strings, the backslash character is used to introduce *escape sequences* that represent other characters. For example, ‘\n’ embeds a newline character in a double-quoted string and ‘\"’ embeds a double quote character. In single-quoted strings, backslash is not a special character. Here is an example showing the difference:

```
double ("\\n")
      ⇒ 10
double ('\\n')
      ⇒ [ 92 110 ]
```

Here is a table of all the escape sequences used in Octave (within double quoted strings). They are the same as those used in the C programming language.

\\\	Represents a literal backslash, ‘\’.
\\"	Represents a literal double-quote character, “”.
\\'	Represents a literal single-quote character, ‘’.
\\0	Represents the null character, control-@, ASCII code 0.
\\a	Represents the “alert” character, control-g, ASCII code 7.
\\b	Represents a backspace, control-h, ASCII code 8.
\\f	Represents a formfeed, control-l, ASCII code 12.
\\n	Represents a newline, control-j, ASCII code 10.
\\r	Represents a carriage return, control-m, ASCII code 13.
\\t	Represents a horizontal tab, control-i, ASCII code 9.

\v	Represents a vertical tab, control-k, ASCII code 11.
\nnn	Represents the octal value <i>nnn</i> , where <i>nnn</i> are one to three digits between 0 and 7. For example, the code for the ASCII ESC (escape) character is ‘\033’.
\xhh...	Represents the hexadecimal value <i>hh</i> , where <i>hh</i> are hexadecimal digits ('0' through '9' and either 'A' through 'F' or 'a' through 'f'). Like the same construct in ANSI C, the escape sequence continues until the first non-hexadecimal digit is seen. However, using more than two hexadecimal digits produces undefined results.

In a single-quoted string there is only one escape sequence: you may insert a single quote character using two single quote characters in succession. For example,

```
'I can''t escape'  
⇒ I can't escape
```

In scripts the two different string types can be distinguished if necessary by using `is_dq_string` and `is_sq_string`.

is_dq_string (x)
Return true if *x* is a double-quoted character string.

See also: [\[is_sq_string\]](#), page 68, [\[ischar\]](#), page 68.

is_sq_string (x)
Return true if *x* is a single-quoted character string.

See also: [\[is_dq_string\]](#), page 68, [\[ischar\]](#), page 68.

5.2 Character Arrays

The string representation used by Octave is an array of characters, so internally the string "ddddddddd" is actually a row vector of length 10 containing the value 100 in all places (100 is the ASCII code of "d"). This lends itself to the obvious generalization to character matrices. Using a matrix of characters, it is possible to represent a collection of same-length strings in one variable. The convention used in Octave is that each row in a character matrix is a separate string, but letting each column represent a string is equally possible.

The easiest way to create a character matrix is to put several strings together into a matrix.

```
collection = [ "String #1"; "String #2" ];
```

This creates a 2-by-9 character matrix.

The function `ischar` can be used to test if an object is a character matrix.

ischar (x)
Return true if *x* is a character array.

See also: [\[isfloat\]](#), page 63, [\[isinteger\]](#), page 54, [\[islogical\]](#), page 63, [\[isnumeric\]](#), page 63, [\[isstring\]](#), page 68, [\[iscellstr\]](#), page 123, [\[isa\]](#), page 39.

isstring (s)
Return true if *s* is a string array.

A string array is a data type that stores strings (row vectors of characters) at each element in the array. It is distinct from character arrays which are N-dimensional arrays where each element is a single 1x1 character. It is also distinct from cell arrays of strings which store strings at each element, but use cell indexing ‘{}’ to access elements rather than string arrays which use ordinary array indexing ‘()’.

Programming Note: Octave does not yet implement string arrays so this function will always return false.

See also: [\[ischar\]](#), page 68, [\[iscellstr\]](#), page 123, [\[isfloat\]](#), page 63, [\[isinteger\]](#), page 54, [\[islogical\]](#), page 63, [\[isnumeric\]](#), page 63, [\[isa\]](#), page 39.

To test if an object is a string (i.e., a 1xN row vector of characters and not a character matrix) you can use the `ischar` function in combination with the `isrow` function as in the following example:

```
ischar (collection)
⇒ 1

ischar (collection) && isrow (collection)
⇒ 0

ischar ("my string") && isrow ("my string")
⇒ 1
```

One relevant question is, what happens when a character matrix is created from strings of different length. The answer is that Octave puts blank characters at the end of strings shorter than the longest string. It is possible to use a different character than the blank character using the `string_fill_char` function.

```
val = string_fill_char ()
old_val = string_fill_char (new_val)
string_fill_char (new_val, "local")
```

Query or set the internal variable used to pad all rows of a character matrix to the same length.

The value must be a single character and the default is " " (a single space). For example:

```
string_fill_char ("X");
[ "these"; "are"; "strings" ]
⇒ "theseXX"
"areXXXX"
"strings"
```

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

This shows a problem with character matrices. It simply isn't possible to represent strings of different lengths. The solution is to use a cell array of strings, which is described in [Section 6.3.4 \[Cell Arrays of Strings\]](#), page 123.

5.3 Creating Strings

The easiest way to create a string is, as illustrated in the introduction, to enclose a text in double-quotes or single-quotes. It is however possible to create a string without actually writing a text. The function `blanks` creates a string of a given length consisting only of blank characters (ASCII code 32).

`blanks (n)`

Return a string of n blanks.

For example:

```
blanks (10);
whos ans
⇒
Attr Name      Size          Bytes  Class
===== =====  =====
ans           1x10           10   char
```

See also: [\[repmat\]](#), page 491.

5.3.1 Concatenating Strings

Strings can be concatenated using matrix notation (see [Chapter 5 \[Strings\]](#), page 67, [Section 5.2 \[Character Arrays\]](#), page 68) which is often the most natural method. For example:

```
fullname = [fname ".txt"];
email = ["<" user "@" domain ">"];
```

In each case it is easy to see what the final string will look like. This method is also the most efficient. When using matrix concatenation the parser immediately begins joining the strings without having to process the overhead of a function call and the input validation of the associated function.

Nevertheless, there are several other functions for concatenating string objects which can be useful in specific circumstances: `char`, `strvcat`, `strcat`, and `cstrcat`. Finally, the general purpose concatenation functions can be used: see [\[cat\]](#), page 481, [\[horzcat\]](#), page 481, and [\[vertcat\]](#), page 481.

- All string concatenation functions except `cstrcat` convert numerical input into character data by taking the corresponding UTF-8 character for each element (or multi-byte sequence), as in the following example:

```
char ([98, 97, 110, 97, 110, 97])
⇒ banana
```

For conversion between locale encodings and UTF-8, see [\[unicode2native\]](#), page 97, and [\[native2unicode\]](#), page 97.

- `char` and `strvcat` concatenate vertically, while `strcat` and `cstrcat` concatenate horizontally. For example:

```
char ("an apple", "two pears")
⇒ an apple
    two pears
```

```
strcat ("oc", "tave", " is", " good", " for you")
      ⇒ octave is good for you
```

- `char` generates an empty row in the output for each empty string in the input. `strvcat`, on the other hand, eliminates empty strings.

```
char ("orange", "green", "", "red")
      ⇒ orange
          green
              red
```

```
strvcat ("orange", "green", "", "red")
      ⇒ orange
          green
              red
```

- All string concatenation functions except `cstrcat` also accept cell array data (see [Section 6.3 \[Cell Arrays\], page 114](#)). `char` and `strvcat` convert cell arrays into character arrays, while `strcat` concatenates within the cells of the cell arrays:

```
char ({"red", "green", "", "blue"})
      ⇒ red
          green
              blue
```

```
strcat ({"abc"; "ghi"}, {"def"; "jkl"})
      ⇒
      {
          [1,1] = abcdef
          [2,1] = ghijkl
      }
```

- `strcat` removes trailing white space in the arguments (except within cell arrays), while `cstrcat` leaves white space untouched. Both kinds of behavior can be useful as can be seen in the examples:

```
strcat (["dir1";"directory2"], ["/";"/"], ["file1";"file2"])
      ⇒ dir1/file1
          directory2/file2
```

```
cstrcat (["thirteen apples"; "a banana"], [" 5$";" 1$"])
      ⇒ thirteen apples 5$
          a banana      1$
```

Note that in the above example for `cstrcat`, the white space originates from the internal representation of the strings in a string array (see [Section 5.2 \[Character Arrays\], page 68](#)).

```
char (x)
char (x, ...)
char (s1, s2, ...)
```

char (cell_array)

Create a string array from one or more numeric matrices, character matrices, or cell arrays.

Arguments are concatenated vertically. The returned values are padded with blanks as needed to make each row of the string array have the same length. Empty input strings are significant and will be concatenated in the output.

For numerical input, each element is converted to the corresponding ASCII character. A range error results if an input is outside the ASCII range (0-255).

For cell arrays, each element is concatenated separately. Cell arrays converted through **char** can mostly be converted back with **cellstr**. For example:

```
char ([97, 98, 99], "", {"98", "99", 100}, "str1", ["ha", "lf"])
⇒ ["abc "
    " "
    "98 "
    "99 "
    "d "
    "str1"
    "half"]
```

See also: [\[strvcat\]](#), page 72, [\[cellstr\]](#), page 123.

strvcat (x)

strvcat (x, ...)

strvcat (s1, s2, ...)

strvcat (cell_array)

Create a character array from one or more numeric matrices, character matrices, or cell arrays.

Arguments are concatenated vertically. The returned values are padded with blanks as needed to make each row of the string array have the same length. Unlike **char**, empty strings are removed and will not appear in the output.

For numerical input, each element is converted to the corresponding ASCII character. A range error results if an input is outside the ASCII range (0-255).

For cell arrays, each element is concatenated separately. Cell arrays converted through **strvcat** can mostly be converted back with **cellstr**. For example:

```
strvcat ([97, 98, 99], "", {"98", "99", 100}, "str1", ["ha", "lf"])
⇒ ["abc "
    "98 "
    "99 "
    "d "
    "str1"
    "half"]
```

See also: [\[char\]](#), page 71, [\[strcat\]](#), page 72, [\[cstrcat\]](#), page 73.

strcat (s1, s2, ...)

Return a string containing all the arguments concatenated horizontally.

If the arguments are cell strings, **strcat** returns a cell string with the individual cells concatenated. For numerical input, each element is converted to the corresponding ASCII character. Trailing white space for any character string input is eliminated before the strings are concatenated. Note that cell string values do **not** have whitespace trimmed.

For example:

```
strcat ("|", " leading space is preserved", "|")
⇒ | leading space is preserved|
strcat ("|", "trailing space is eliminated ", "|")
⇒ |trailing space is eliminated|
strcat ("homogeneous space |", " ", "| is also eliminated")
⇒ homogeneous space || is also eliminated
s = [ "ab"; "cde" ];
strcat (s, s, s)
⇒
    "ababab"
    "cdecdecde"
s = { "ab"; "cd " };
strcat (s, s, s)
⇒
{
    [1,1] = ababab
    [2,1] = cd cd cd
}
```

See also: [\[cstrcat\]](#), page 73, [\[char\]](#), page 71, [\[strvcat\]](#), page 72.

cstrcat (*s₁*, *s₂*, ...)

Return a string containing all the arguments concatenated horizontally with trailing white space preserved.

For example:

```
cstrcat ("ab      ", "cd")
⇒ "ab      cd"
s = [ "ab"; "cde" ];
cstrcat (s, s, s)
⇒ "ab ab ab "
    "cdecdecde"
```

See also: [\[strcat\]](#), page 72, [\[char\]](#), page 71, [\[strvcat\]](#), page 72.

5.3.2 Converting Numerical Data to Strings

Apart from the string concatenation functions (see [Section 5.3.1 \[Concatenating Strings\]](#), page 70) which cast numerical data to the corresponding UTF-8 encoded characters, there are several functions that format numerical data as strings. **mat2str** and **num2str** convert real or complex matrices, while **int2str** converts integer matrices. **int2str** takes the real part of complex values and round fractional values to integer. A more flexible way to format

numerical data as strings is the `sprintf` function (see [Section 14.2.4 \[Formatted Output\]](#), page 277, [\[sprintf\]](#), page 278).

```
s = mat2str (x, n)
s = mat2str (x, n, "class")
Format real, complex, and logical matrices as strings.
```

The returned string may be used to reconstruct the original matrix by using the `eval` function.

The precision of the values is given by *n*. If *n* is a scalar then both real and imaginary parts of the matrix are printed to the same precision. Otherwise *n*(1) defines the precision of the real part and *n*(2) defines the precision of the imaginary part. The default for *n* is 15.

If the argument "class" is given then the class of *x* is included in the string in such a way that `eval` will result in the construction of a matrix of the same class.

```
mat2str ([ -1/3 + i/7; 1/3 - i/7 ], [4 2])
⇒ "[ -0.3333+0.14i; 0.3333-0.14i ]"

mat2str ([ -1/3 +i/7; 1/3 -i/7 ], [4 2])
⇒ "[ -0.3333+0i 0+0.14i; 0.3333+0i -0-0.14i ]"

mat2str (int16 ([1 -1]), "class")
⇒ "int16([1 -1])"

mat2str (logical (eye (2)))
⇒ "[true false;false true]"

isequal (x, eval (mat2str (x)))
⇒ 1
```

See also: [\[sprintf\]](#), page 278, [\[num2str\]](#), page 74, [\[int2str\]](#), page 75.

```
num2str (x)
num2str (x, precision)
num2str (x, format)
```

Convert a number (or array) to a string (or a character array).

The optional second argument may either give the number of significant digits (*precision*) to be used in the output or a format template string (*format*) as in `sprintf` (see [Section 14.2.4 \[Formatted Output\]](#), page 277). `num2str` can also process complex numbers.

Examples:

```

num2str (123.456)
⇒ 123.456

num2str (123.456, 4)
⇒ 123.5

s = num2str ([1, 1.34; 3, 3.56], "%5.1f")
⇒ s =
    1.0  1.3
    3.0  3.6
whos s
⇒ Variables in the current scope:
      Attr Name      Size      Bytes  Class
      =====  =====  =====  =====
      s          2x8           16   char
Total is 16 elements using 16 bytes

num2str (1.234 + 27.3i)
⇒ 1.234+27.3i

```

The `num2str` function is not very flexible. For better control over the results, use `sprintf` (see [Section 14.2.4 \[Formatted Output\]](#), page [277](#)).

Programming Notes:

For MATLAB compatibility, leading spaces are stripped before returning the string.

Integers larger than `flintmax` may not be displayed correctly.

For complex `x`, the format string may only contain one output conversion specification and nothing else. Otherwise, results will be unpredictable.

Any optional `format` specified by the programmer is used without modification. This is in contrast to MATLAB which tampers with the `format` based on internal heuristics.

See also: [\[sprintf\]](#), page [278](#), [\[int2str\]](#), page [75](#), [\[mat2str\]](#), page [74](#).

int2str (n)

Convert an integer (or array of integers) to a string (or a character array).

```

int2str (123)
⇒ 123

s = int2str ([1, 2, 3; 4, 5, 6])
⇒ s =
    1  2  3
    4  5  6

whos s
⇒ Variables in the current scope:
      Attr Name      Size      Bytes  Class
      =====  =====  =====  =====
      s          2x7           14   char
Total is 14 elements using 14 bytes

```

This function is not very flexible. For better control over the results, use `sprintf` (see [Section 14.2.4 \[Formatted Output\]](#), page [277](#)).

Programming Notes:

Non-integers are rounded to integers before display. Only the real part of complex numbers is displayed.

See also: [\[sprintf\]](#), page [278](#), [\[num2str\]](#), page [74](#), [\[mat2str\]](#), page [74](#).

5.4 Comparing Strings

Since a string is a character array, comparisons between strings work element by element as the following example shows:

```
GNU = "GNU's Not UNIX";
spaces = (GNU == " ")
⇒ spaces =
 0  0  0  0  0  1  0  0  0  1  0  0  0  0  0
```

To determine if two strings are identical it is necessary to use the `strcmp` function. It compares complete strings and is case sensitive. `strncmp` compares only the first N characters (with N given as a parameter). `strcmpi` and `strncmpi` are the corresponding functions for case-insensitive comparison.

`strcmp (s_1, s_2)`

Return 1 if the character strings s_1 and s_2 are the same, and 0 otherwise.

If either s_1 or s_2 is a cell array of strings, then an array of the same size is returned, containing the values described above for every member of the cell array. The other argument may also be a cell array of strings (of the same size or with only one element), char matrix or character string.

Caution: For compatibility with MATLAB, Octave's `strcmp` function returns 1 if the character strings are equal, and 0 otherwise. This is just the opposite of the corresponding C library function.

See also: [\[strcmpi\]](#), page [77](#), [\[strncmp\]](#), page [76](#), [\[strncmpi\]](#), page [77](#).

`strncmp (s_1, s_2, n)`

Return 1 if the first n characters of strings s_1 and s_2 are the same, and 0 otherwise.

```
strncmp ("abce", "abcd", 3)
⇒ 1
```

If either s_1 or s_2 is a cell array of strings, then an array of the same size is returned, containing the values described above for every member of the cell array. The other argument may also be a cell array of strings (of the same size or with only one element), char matrix or character string.

```
strncmp ("abce", {"abcd", "bca", "abc"}, 3)
⇒ [1, 0, 1]
```

Caution: For compatibility with MATLAB, Octave's `strncmp` function returns 1 if the character strings are equal, and 0 otherwise. This is just the opposite of the corresponding C library function.

See also: [\[strncmpi\]](#), page [77](#), [\[strcmp\]](#), page [76](#), [\[strcmpi\]](#), page [77](#).

strcmpi (s1, s2)

Return 1 if the character strings $s1$ and $s2$ are the same, disregarding case of alphabetic characters, and 0 otherwise.

If either $s1$ or $s2$ is a cell array of strings, then an array of the same size is returned, containing the values described above for every member of the cell array. The other argument may also be a cell array of strings (of the same size or with only one element), char matrix or character string.

Caution: For compatibility with MATLAB, Octave's strcmp function returns 1 if the character strings are equal, and 0 otherwise. This is just the opposite of the corresponding C library function.

Caution: National alphabets are not supported.

See also: [strcmp], page 76, [strncmp], page 76, [strncmpi], page 77.

strncmpi (s1, s2, n)

Return 1 if the first n character of $s1$ and $s2$ are the same, disregarding case of alphabetic characters, and 0 otherwise.

If either $s1$ or $s2$ is a cell array of strings, then an array of the same size is returned, containing the values described above for every member of the cell array. The other argument may also be a cell array of strings (of the same size or with only one element), char matrix or character string.

Caution: For compatibility with MATLAB, Octave's strncmpi function returns 1 if the character strings are equal, and 0 otherwise. This is just the opposite of the corresponding C library function.

Caution: National alphabets are not supported.

See also: [strncmp], page 76, [strcmp], page 76, [strcmpi], page 77.

5.5 Manipulating Strings

Octave supports a wide range of functions for manipulating strings. Since a string is just a matrix, simple manipulations can be accomplished using standard operators. The following example shows how to replace all blank characters with underscores.

```
quote = ...
    "First things first, but not necessarily in that order";
quote( quote == " " ) = "_"
⇒ quote =
    First_things_first,_but_not_necessarily_in_that_order
```

For more complex manipulations, such as searching, replacing, and general regular expressions, the following functions come with Octave.

deblank (s)

Remove trailing whitespace and nulls from s .

If s is a matrix, deblank trims each row to the length of longest string. If s is a cell array of strings, operate recursively on each string element.

Examples:

```
deblank ("      abc    ")
⇒ "      abc"

deblank ([" abc    "; "      def    "])
⇒ [" abc  " ; "      def"]
```

See also: [strtrim], page 78.

strtrim (s)

Remove leading and trailing whitespace from *s*.

If *s* is a matrix, *strtrim* trims each row to the length of longest string. If *s* is a cell array of strings, operate recursively on each string element.

For example:

```
strtrim ("      abc    ")
⇒ "abc"

strtrim ([" abc    "; "      def    "])
⇒ ["abc  " ; "      def"]
```

See also: [deblank], page 77.

strtrunc (s, n)

Truncate the character string *s* to length *n*.

If *s* is a character matrix, then the number of columns is adjusted.

If *s* is a cell array of strings, then the operation is performed on each cell element and the new cell array is returned.

findstr (s, t)

findstr (s, t, overlap)

This function is obsolete. Use **strfind** instead.

Return the vector of all positions in the longer of the two strings *s* and *t* where an occurrence of the shorter of the two starts.

If the optional argument *overlap* is true (default), the returned vector can include overlapping positions. For example:

```
findstr ("ababab", "a")
⇒ [1, 3, 5];
findstr ("abababa", "aba", 0)
⇒ [1, 5]
```

Caution: **findstr** is obsolete. Use **strfind** in all new code.

See also: [strfind], page 79, [strmatch], page 80, [strcmp], page 76, [strncmp], page 76, [strcmpi], page 77, [strncmpi], page 77, [find], page 477.

```
idx = strchr (str, chars)
idx = strchr (str, chars, n)
idx = strchr (str, chars, n, direction)
```

[i, j] = strchr (...)

Search for the string *str* for occurrences of characters from the set *chars*.

The return value(s), as well as the *n* and *direction* arguments behave identically as in **find**.

This will be faster than using regexp in most cases.

See also: [\[find\]](#), page 477.

index (s, t)

index (s, t, direction)

Return the position of the first occurrence of the string *t* in the string *s*, or 0 if no occurrence is found.

s may also be a string array or cell array of strings.

For example:

```
index ("Teststring", "t")
      ⇒ 4
```

If *direction* is "first", return the first element found. If *direction* is "last", return the last element found.

See also: [\[find\]](#), page 477, [\[rindex\]](#), page 79.

rindex (s, t)

Return the position of the last occurrence of the character string *t* in the character string *s*, or 0 if no occurrence is found.

s may also be a string array or cell array of strings.

For example:

```
rindex ("Teststring", "t")
      ⇒ 6
```

The **rindex** function is equivalent to **index** with *direction* set to "last".

See also: [\[find\]](#), page 477, [\[index\]](#), page 79.

idx = strfind (str, pattern)

idx = strfind (cellstr, pattern)

idx = strfind (..., "overlaps", val)

Search for *pattern* in the string *str* and return the starting index of every such occurrence in the vector *idx*.

If there is no such occurrence, or if *pattern* is longer than *str*, or if *pattern* itself is empty, then *idx* is the empty array [].

The optional argument "overlaps" determines whether the pattern can match at every position in *str* (true), or only for unique occurrences of the complete pattern (false). The default is true.

If a cell array of strings *cellstr* is specified then *idx* is a cell array of vectors, as specified above.

Examples:

```
strfind ("abababa", "aba")
⇒ [1, 3, 5]

strfind ("abababa", "aba", "overlaps", false)
⇒ [1, 5]

strfind ({"abababa", "bebebe", "ab"}, "aba")
⇒
{
    [1,1] =
        1   3   5

    [1,2] = []
    [1,3] = []
}
}
```

See also: [\[regexp\]](#), page 88, [\[regexpi\]](#), page 90, [\[find\]](#), page 477.

```
str = strjoin (cstr)
str = strjoin (cstr, delimiter)
```

Join the elements of the cell string array, *cstr*, into a single string.

If no *delimiter* is specified, the elements of *cstr* are separated by a space.

If *delimiter* is specified as a string, the cell string array is joined using the string. Escape sequences are supported.

If *delimiter* is a cell string array whose length is one less than *cstr*, then the elements of *cstr* are joined by interleaving the cell string elements of *delimiter*. Escape sequences are not supported.

```
strjoin ({'Octave','Scilab','Lush','Yorick'}, '*')
⇒ 'Octave*Scilab*Lush*Yorick'
```

See also: [\[strsplit\]](#), page 81.

```
strmatch (s, A)
strmatch (s, A, "exact")
```

This function is obsolete. Use an alternative such as `strcmp` or `strncmp` instead.

Return indices of entries of *A* which begin with the string *s*.

The second argument *A* must be a string, character matrix, or a cell array of strings.

If the third argument "*exact*" is not given, then *s* only needs to match *A* up to the length of *s*. Trailing spaces and nulls in *s* and *A* are ignored when matching.

For example:

```

strmatch ("apple", "apple juice")
⇒ 1

strmatch ("apple", ["apple "; "apple juice"; "an apple"])
⇒ [1; 2]

strmatch ("apple", ["apple "; "apple juice"; "an apple"], "exact")
⇒ [1]

```

Caution: `strmatch` is obsolete. Use `strcmp` (normal case) or `strcmp` ("exact" case) in all new code. Other replacement possibilities, depending on application, include `regexp` or `validatestring`.

See also: [\[strcmp\]](#), page 76, [\[strcmp\]](#), page 76, [\[regexp\]](#), page 88, [\[strfind\]](#), page 79, [\[validatestring\]](#), page 183.

[tok, rem] = strtok (str)
[tok, rem] = strtok (str, delim)

Find all characters in the string *str* up to, but not including, the first character which is in the string *delim*.

str may also be a cell array of strings in which case the function executes on every individual string and returns a cell array of tokens and remainders.

Leading delimiters are ignored. If *delim* is not specified, whitespace is assumed.

If *rem* is requested, it contains the remainder of the string, starting at the first delimiter.

Examples:

```

strtok ("this is the life")
⇒ "this"

[tok, rem] = strtok ("14*27+31", "+-*/")
⇒
    tok = 14
    rem = *27+31

```

See also: [\[index\]](#), page 79, [\[strsplit\]](#), page 81, [\[strchr\]](#), page 78, [\[isspace\]](#), page 99.

[cstr] = strsplit (str)
[cstr] = strsplit (str, del)
[cstr] = strsplit (..., name, value)
[cstr, matches] = strsplit (...)

Split the string *str* using the delimiters specified by *del* and return a cell string array of substrings.

If a delimiter is not specified the string is split at whitespace {" ", "\f", "\n", "\r", "\t", "\v"}. Otherwise, the delimiter, *del* must be a string or cell array of strings. By default, consecutive delimiters in the input string *s* are collapsed into one resulting in a single split.

Supported *name/value* pair arguments are:

- *collapseddelimiters* which may take the value of `true` (default) or `false`.

- `delimitertype` which may take the value of "`simple`" (default) or "`regularexpression`". A simple delimiter matches the text exactly as written. Otherwise, the syntax for regular expressions outlined in `regexp` is used.

The optional second output, `matches`, returns the delimiters which were matched in the original string.

Examples with simple delimiters:

```

strsplit ("a b c")
⇒
{
  [1,1] = a
  [1,2] = b
  [1,3] = c
}

strsplit ("a,b,c", ",")
⇒
{
  [1,1] = a
  [1,2] = b
  [1,3] = c
}

strsplit ("a foo b,bar c", {" ", ",", "foo", "bar"})
⇒
{
  [1,1] = a
  [1,2] = b
  [1,3] = c
}

strsplit ("a,,b, c", {","}, "collapsedlimiters", false)
⇒
{
  [1,1] = a
  [1,2] =
  [1,3] = b
  [1,4] =
  [1,5] = c
}

```

Examples with `regularexpression` delimiters:

```

strsplit ("a foo b,bar c", ',|\s|foo|bar', ...
          "delimitertype", "regularexpression")
⇒
{
  [1,1] = a
  [1,2] = b
}

```

```

        [1,3] = c
    }

strsplit ("a,,b, c", '[, ]', "collapsedelimters", false, ...
          "delimitertype", "regularexpression")
⇒
{
    [1,1] = a
    [1,2] =
    [1,3] = b
    [1,4] =
    [1,5] = c
}

strsplit ("a,\t,b, c", {',', '\s'}, "delimitertype", "regularexpression")
⇒
{
    [1,1] = a
    [1,2] = b
    [1,3] = c
}

strsplit ("a,\t,b, c", {',', ' ', '\t'}, "collapsedelimters", false)
⇒
{
    [1,1] = a
    [1,2] =
    [1,3] =
    [1,4] = b
    [1,5] =
    [1,6] = c
}

```

See also: [\[ostrsplit\]](#), page 83, [\[strjoin\]](#), page 80, [\[strtok\]](#), page 81, [\[regexp\]](#), page 88.

`[cstr] = ostrsplit (s, sep)`
`[cstr] = ostrsplit (s, sep, strip_empty)`

Split the string *s* using one or more separators *sep* and return a cell array of strings.

Consecutive separators and separators at boundaries result in empty strings, unless *strip_empty* is true. The default value of *strip_empty* is false.

2-D character arrays are split at separators and at the original column boundaries.

Example:

```

ostrsplit ("a,b,c", ",")
⇒
{
  [1,1] = a
  [1,2] = b
  [1,3] = c
}

ostrsplit (["a,b" ; "cde"], ",")
⇒
{
  [1,1] = a
  [1,2] = b
  [1,3] = cde
}

```

See also: [\[strsplit\]](#), page 81, [\[strtok\]](#), page 81.

```

[a, ...] = strread (str)
[a, ...] = strread (str, format)
[a, ...] = strread (str, format, format_repeat)
[a, ...] = strread (str, format, prop1, value1, ...)
[a, ...] = strread (str, format, format_repeat, prop1, value1, ...)

```

This function is obsolete. Use `textscan` instead.

Read data from a string.

The string `str` is split into words that are repeatedly matched to the specifiers in `format`. The first word is matched to the first specifier, the second to the second specifier and so forth. If there are more words than specifiers, the process is repeated until all words have been processed.

The string `format` describes how the words in `str` should be parsed. It may contain any combination of the following specifiers:

`%s` The word is parsed as a string.

`%f`

`%n` The word is parsed as a number and converted to double.

`%d`

`%u` The word is parsed as a number and converted to int32.

`/*`

`/*f`

`/*s` The word is skipped.

For `%s` and `%d`, `%f`, `%n`, `%u` and the associated `/*s`... specifiers an optional width can be specified as `%Ns`, etc. where `N` is an integer > 1 . For `%f`, format specifiers like `%N.Mf` are allowed.

`literals` In addition the format may contain literal character strings; these will be skipped during reading.

Parsed word corresponding to the first specifier are returned in the first output argument and likewise for the rest of the specifiers.

By default, *format* is "%f", meaning that numbers are read from *str*. This will do if *str* contains only numeric fields.

For example, the string

```
str = "\n
Bunny Bugs 5.5\n\
Duck Daffy -7.5e-5\n\
Penguin Tux 6"
```

can be read using

```
[a, b, c] = strread (str, "%s %s %f");
```

Optional numeric argument *format_repeat* can be used for limiting the number of items read:

- 1 (default) read all of the string until the end.
- N Read N times *nargout* items. 0 (zero) is an acceptable value for *format_repeat*.

The behavior of **strread** can be changed via property-value pairs. The following properties are recognized:

"commentstyle"

Parts of *str* are considered comments and will be skipped. *value* is the comment style and can be any of the following.

- "shell" Everything from # characters to the nearest end-of-line is skipped.
- "c" Everything between /* and */ is skipped.
- "c++" Everything from // characters to the nearest end-of-line is skipped.
- "matlab" Everything from % characters to the nearest end-of-line is skipped.
- user-supplied. Two options: (1) One string, or 1x1 cell string: Skip everything to the right of it; (2) 2x1 cell string array: Everything between the left and right strings is skipped.

"delimiter"

Any character in *value* will be used to split *str* into words (default value = any whitespace). Note that whitespace is implicitly added to the set of delimiter characters unless a "%s" format conversion specifier is supplied; see "whitespace" parameter below. The set of delimiter characters cannot be empty; if needed Octave substitutes a space as delimiter.

"emptyvalue"

Value to return for empty numeric values in non-whitespace delimited data. The default is NaN. When the data type does not support NaN (int32 for example), then default is zero.

"multipledelimsasone"

Treat a series of consecutive delimiters, without whitespace in between, as a single delimiter. Consecutive delimiter series need not be vertically "aligned".

"treatasempty"

Treat single occurrences (surrounded by delimiters or whitespace) of the string(s) in *value* as missing values.

"returnonerror"

If *value* true (1, default), ignore read errors and return normally. If false (0), return an error.

"whitespace"

Any character in *value* will be interpreted as whitespace and trimmed; the string defining whitespace must be enclosed in double quotes for proper processing of special characters like "\t". In each data field, multiple consecutive whitespace characters are collapsed into one space and leading and trailing whitespace is removed. The default value for whitespace is "\b\r\n\t" (note the space). Whitespace is always added to the set of delimiter characters unless at least one "%s" format conversion specifier is supplied; in that case only whitespace explicitly specified in "delimiter" is retained as delimiter and removed from the set of whitespace characters. If whitespace characters are to be kept as-is (in e.g., strings), specify an empty value (i.e., "") for "whitespace"; obviously, whitespace cannot be a delimiter then.

When the number of words in *str* doesn't match an exact multiple of the number of format conversion specifiers, strread's behavior depends on the last character of *str*:

last character = "\n"

Data columns are padded with empty fields or NaN so that all columns have equal length

last character is not "\n"

Data columns are not padded; strread returns columns of unequal length

See also: [textscan], page 266, [sscanf], page 284.

```
newstr = strrep (str, ptn, rep)
newstr = strrep (cellstr, ptn, rep)
newstr = strrep (..., "overlaps", val)
```

Replace all occurrences of the pattern *ptn* in the string *str* with the string *rep* and return the result.

The optional argument "overlaps" determines whether the pattern can match at every position in *str* (true), or only for unique occurrences of the complete pattern (false). The default is true.

s may also be a cell array of strings, in which case the replacement is done for each element and a cell array is returned.

Example:

```
strrep ("This is a test string", "is", "&%$")
⇒ "Th&%$ &%$ a test string"
```

See also: [\[regexprep\]](#), page 91, [\[strfind\]](#), page 79.

newstr = erase (str, ptn)

Delete all occurrences of *ptn* within *str*.

str and *ptn* can be ordinary strings, cell array of strings, or character arrays.

Examples

```
## string, single pattern
erase ("Hello World!", " World")
⇒ "Hello!"

## cellstr, single pattern
erase ({"Hello", "World!"}, "World")
⇒ {"Hello", "!"}

## string, multiple patterns
erase ("The Octave interpreter is fabulous", ...
      {"interpreter", "The"})
⇒ "Octave is fabulous"

## cellstr, multiple patterns
erase ({"The", "Octave interpreter", "is fabulous"}, ...
      {"interpreter", "The"})
⇒ {"", "Octave", "is fabulous"}
```

Programming Note: **erase** deletes the first instance of a pattern in a string when there are overlapping occurrences. For example:

```
erase ("abababa", "aba")
⇒ "b"
```

See **strrep** for processing overlaps.

See also: [\[strrep\]](#), page 86, [\[regexprep\]](#), page 91.

substr (s, offset)

substr (s, offset, len)

Return the substring of *s* which starts at character number *offset* and is *len* characters long.

Position numbering for offsets begins with 1. If *offset* is negative, extraction starts that far from the end of the string.

If *len* is omitted, the substring extends to the end of *s*. A negative value for *len* extracts to within *len* characters of the end of the string

Examples:

```

substr ("This is a test string", 6, 9)
    => "is a test"
substr ("This is a test string", -11)
    => "test string"
substr ("This is a test string", -11, -7)
    => "test"

```

This function is patterned after the equivalent function in Perl.

`[s, e, te, m, t, nm, sp] = regexp (str, pat)`
`[...] = regexp (str, pat, "opt1", ...)`

Regular expression string matching.

Search for *pat* in *str* and return the positions and substrings of any matches, or empty values if there are none.

The matched pattern *pat* can include any of the standard regex operators, including:

.	Match any character
* + ? {}	Repetition operators, representing
*	Match zero or more times
+	Match one or more times
?	Match zero or one times
{n}	Match exactly <i>n</i> times
{n,}	Match <i>n</i> or more times
{m, n}	Match between <i>m</i> and <i>n</i> times
[...] [^...]	List operators. The pattern will match any character listed between "[" and "]". If the first character is "^" then the pattern is inverted and any character except those listed between brackets will match.
	Escape sequences defined below can also be used inside list operators.
	For example, a template for a floating point number might be [-+.\\d]+.
() (?:)	Grouping operator. The first form, parentheses only, also creates a token.
	Alternation operator. Match one of a choice of regular expressions. The alternatives must be delimited by the grouping operator () above.
^ \$	Anchoring operators. Requires pattern to occur at the start (^) or end (\$) of the string.

In addition, the following escaped characters have special meaning.

\d	Match any digit
\D	Match any non-digit
\s	Match any whitespace character
\S	Match any non-whitespace character
\w	Match any word character

\W	Match any non-word character
\<	Match the beginning of a word
\>	Match the end of a word
\B	Match within a word

Implementation Note: For compatibility with MATLAB, escape sequences in *pat* (e.g., "`\n`" => newline) are expanded even when *pat* has been defined with single quotes. To disable expansion use a second backslash before the escape sequence (e.g., "`\\\n`") or use the **regexptranslate** function.

The outputs of **regexp** default to the order given below

s	The start indices of each matching substring
e	The end indices of each matching substring
te	The extents of each matched token surrounded by (...) in <i>pat</i>
m	A cell array of the text of each match
t	A cell array of the text of each token matched
nm	A structure containing the text of each matched named token, with the name being used as the fieldname. A named token is denoted by (?<name>...).
sp	A cell array of the text not returned by match, i.e., what remains if you split the string based on <i>pat</i> .

Particular output arguments, or the order of the output arguments, can be selected by additional *opt* arguments. These are strings and the correspondence between the output arguments and the optional argument are

'start'	s
'end'	e
'tokenExtents'	te
'match'	m
'tokens'	t
'names'	nm
'split'	sp

Additional arguments are summarized below.

'once' Return only the first occurrence of the pattern.

'matchcase'

Make the matching case sensitive. (default)

Alternatively, use (?-i) in the pattern.

'ignorecase'

Ignore case when matching the pattern to the string.

Alternatively, use (?i) in the pattern.

'stringanchors'
 Match the anchor characters at the beginning and end of the string.
 (default)
 Alternatively, use (?-m) in the pattern.

'lineanchors'
 Match the anchor characters at the beginning and end of the line.
 Alternatively, use (?m) in the pattern.

'dotall' The pattern . matches all characters including the newline character.
 (default)
 Alternatively, use (?s) in the pattern.

'dotexceptnewline'
 The pattern . matches all characters except the newline character.
 Alternatively, use (?-s) in the pattern.

'literalspacing'
 All characters in the pattern, including whitespace, are significant and
 are used in pattern matching. (default)
 Alternatively, use (?-x) in the pattern.

'freespacing'
 The pattern may include arbitrary whitespace and also comments begin-
 ning with the character '#'.
 Alternatively, use (?x) in the pattern.

'noemptymatch'
 Zero-length matches are not returned. (default)

'emptymatch'
 Return zero-length matches.
 `regexp ('a', 'b*', 'emptymatch')` returns [1 2] because there are
 zero or more 'b' characters at positions 1 and end-of-string.

Stack Limitation Note: Pattern searches are done with a recursive function which can overflow the program stack when there are a high number of matches. For example,

```
regexp (repmat ('a', 1, 1e5), '(a)+')
```

may lead to a segfault. As an alternative, consider constructing pattern searches that reduce the number of matches (e.g., by creatively using set complement), and then further processing the return variables (now reduced in size) with successive `regexp` searches.

See also: [\[regexpi\]](#), page 90, [\[strfind\]](#), page 79, [\[regexpexec\]](#), page 91.

```
[s, e, te, m, t, nm, sp] = regexpi (str, pat)
[...] = regexpi (str, pat, "opt1", ...)
```

Case insensitive regular expression string matching.

Search for *pat* in *str* and return the positions and substrings of any matches, or empty values if there are none. See [\[regexp\]](#), page 88, for details on the syntax of the search pattern.

See also: [\[regexp\]](#), page 88.

```
outstr = regexprep (string, pat, repstr)
outstr = regexprep (string, pat, repstr, "opt1", ...)
    Replace occurrences of pattern pat in string with repstr.
```

The pattern is a regular expression as documented for **regexp**. See [regexp], page 88.

The replacement string may contain $\$i$, which substitutes for the *i*th set of parentheses in the match string. For example,

```
regexprep ("Bill Dunn", '(\w+) (\w+)', '$2, $1')
returns "Dunn, Bill"
```

Options in addition to those of **regexp** are

'once' Replace only the first occurrence of *pat* in the result.

'warnings'

This option is present for compatibility but is ignored.

Implementation Note: For compatibility with MATLAB, escape sequences in *pat* (e.g., "`\n`" => newline) are expanded even when *pat* has been defined with single quotes. To disable expansion use a second backslash before the escape sequence (e.g., "`\\\n`") or use the **regexptranslate** function.

See also: [regexp], page 88, [regexpi], page 90, [strrep], page 86.

regexptranslate (op, s)

Translate a string for use in a regular expression.

This may include either wildcard replacement or special character escaping.

The behavior is controlled by *op* which can take the following values

"wildcard"

The wildcard characters `.`, `*`, and `?` are replaced with wildcards that are appropriate for a regular expression. For example:

```
regexptranslate ("wildcard", "*.m")
⇒ '.*\.\.m'
```

"escape" The characters `$`, `.`, `?`, `[]`, that have special meaning for regular expressions are escaped so that they are treated literally. For example:

```
regexptranslate ("escape", "12.5")
⇒ '12\.5'
```

See also: [regexp], page 88, [regexpi], page 90, [regexprep], page 91.

untabify (t)

untabify (t, tw)

untabify (t, tw, deblank)

Replace TAB characters in *t* with spaces.

The input, *t*, may be either a 2-D character array, or a cell array of character strings. The output is the same class as the input.

The tab width is specified by *tw*, and defaults to eight.

If the optional argument *deblank* is true, then the spaces will be removed from the end of the character data.

The following example reads a file and writes an untabified version of the same file with trailing spaces stripped.

```
fid = fopen ("tabbed_script.m");
text = char (fread (fid, "uchar"));
fclose (fid);
fid = fopen ("untabified_script.m", "w");
text = untabify (strsplit (text, "\n"), 8, true);
fprintf (fid, "%s\n", text{:});
fclose (fid);
```

See also: [strjust], page 96, [strsplit], page 81, [deblank], page 77.

idx = unicode_idx (str)

Return an array with the indices for each UTF-8 encoded character in *str*.

```
unicode_idx ("ääbc")
⇒ [1, 2, 2, 3, 4]
```

5.6 String Conversions

Octave supports various kinds of conversions between strings and numbers. As an example, it is possible to convert a string containing a hexadecimal number to a floating point number.

```
hex2dec ("FF")
⇒ 255
```

bin2dec (s)

Return the decimal number corresponding to the binary number represented by the string *s*.

For example:

```
bin2dec ("1110")
⇒ 14
```

Spaces are ignored during conversion and may be used to make the binary number more readable.

```
bin2dec ("1000 0001")
⇒ 129
```

If *s* is a string matrix, return a column vector with one converted number per row of *s*; Invalid rows evaluate to NaN.

If *s* is a cell array of strings, return a column vector with one converted number per cell element in *s*.

See also: [dec2bin], page 92, [base2dec], page 94, [hex2dec], page 93.

dec2bin (d, len)

Return a binary number corresponding to the non-negative integer *d*, as a string of ones and zeros.

For example:

```
dec2bin (14)
⇒ "1110"
```

If d is a matrix or cell array, return a string matrix with one row per element in d , padded with leading zeros to the width of the largest value.

The optional second argument, len , specifies the minimum number of digits in the result.

See also: [bin2dec], page 92, [dec2base], page 93, [dec2hex], page 93.

`dec2hex (d, len)`

Return the hexadecimal string corresponding to the non-negative integer d .

For example:

```
dec2hex (2748)
⇒ "ABC"
```

If d is a matrix or cell array, return a string matrix with one row per element in d , padded with leading zeros to the width of the largest value.

The optional second argument, len , specifies the minimum number of digits in the result.

See also: [hex2dec], page 93, [dec2base], page 93, [dec2bin], page 92.

`hex2dec (s)`

Return the integer corresponding to the hexadecimal number represented by the string s .

For example:

```
hex2dec ("12B")
⇒ 299
hex2dec ("12b")
⇒ 299
```

If s is a string matrix, return a column vector with one converted number per row of s ; Invalid rows evaluate to NaN.

If s is a cell array of strings, return a column vector with one converted number per cell element in s .

See also: [dec2hex], page 93, [base2dec], page 94, [bin2dec], page 92.

`dec2base (d, base)`

`dec2base (d, base, len)`

Return a string of symbols in base $base$ corresponding to the non-negative integer d .

```
dec2base (123, 3)
⇒ "11120"
```

If d is a matrix or cell array, return a string matrix with one row per element in d , padded with leading zeros to the width of the largest value.

If $base$ is a string then the characters of $base$ are used as the symbols for the digits of d . Space (' ') may not be used as a symbol.

```
dec2base (123, "aei")
⇒ "eeeia"
```

The optional third argument, len , specifies the minimum number of digits in the result.

See also: [base2dec], page 94, [dec2bin], page 92, [dec2hex], page 93.

base2dec (s, base)

Convert *s* from a string of digits in base *base* to a decimal integer (base 10).

```
base2dec ("11120", 3)
⇒ 123
```

If *s* is a string matrix, return a column vector with one value per row of *s*. If a row contains invalid symbols then the corresponding value will be NaN.

If *s* is a cell array of strings, return a column vector with one value per cell element in *s*.

If *base* is a string, the characters of *base* are used as the symbols for the digits of *s*. Space (' ') may not be used as a symbol.

```
base2dec ("yyzxx", "xyz")
⇒ 123
```

See also: [dec2base], page 93, [bin2dec], page 92, [hex2dec], page 93.

```
s = num2hex (n)
s = num2hex (n, "cell")
```

Convert a numeric array to an array of hexadecimal strings.

For example:

```
num2hex ([-1, 1, e, Inf])
⇒ "bff000000000000
    3ff000000000000
    4005bf0a8b145769
    7ff000000000000"
```

If the argument *n* is a single precision number or vector, the returned string has a length of 8. For example:

```
num2hex (single ([-1, 1, e, Inf]))
⇒ "bf800000
    3f800000
    402df854
    7f800000"
```

With the optional second argument "cell", return a cell array of strings instead of a character array.

See also: [hex2num], page 94, [hex2dec], page 93, [dec2hex], page 93.

```
n = hex2num (s)
n = hex2num (s, class)
```

Typecast a hexadecimal character array or cell array of strings to an array of numbers.

By default, the input array is interpreted as a hexadecimal number representing a double precision value. If fewer than 16 characters are given the strings are right padded with '0' characters.

Given a string matrix, **hex2num** treats each row as a separate number.

```
hex2num (["4005bf0a8b145769"; "4024000000000000"])
⇒ [2.7183; 10.000]
```

The optional second argument *class* may be used to cause the input array to be interpreted as a different value type. Possible values are

Option Characters

"int8"	2
"uint8"	2
"int16"	4
"uint16"	4
"int32"	8
"uint32"	8
"int64"	16
"uint64"	16
"char"	2
"single"	8
"double"	16

For example:

```
hex2num (["402df854"; "41200000"], "single")
⇒ [2.7183; 10.000]
```

See also: [\[num2hex\]](#), page 94, [\[hex2dec\]](#), page 93, [\[dec2hex\]](#), page 93.

str2double (s)

Convert a string to a real or complex number.

The string must be in one of the following formats where a and b are real numbers and the complex unit is 'i' or 'j':

- a + bi
- a + b*i
- a + i*b
- bi + a
- b*i + a
- i*b + a

If present, a and/or b are of the form $[+]-d[,]d[[eE][+]-]d$ where the brackets indicate optional arguments and 'd' indicates zero or more digits. The special input values Inf, NaN, and NA are also accepted.

s may be a character string, character matrix, or cell array. For character arrays the conversion is repeated for every row, and a double or complex array is returned. Empty rows in s are deleted and not returned in the numeric array. For cell arrays each character string element is processed and a double or complex array of the same dimensions as s is returned.

For unconvertible scalar or character string input str2double returns a NaN. Similarly, for character array input str2double returns a NaN for any row of s that could not be converted. For a cell array, str2double returns a NaN for any element of s for which conversion fails. Note that numeric elements in a mixed string/numeric cell array are not strings and the conversion will fail for these elements and return NaN.

`str2double` can replace `str2num`, and it avoids the security risk of using `eval` on unknown data.

See also: [\[str2num\]](#), page 96.

strjust (*s*)
strjust (*s*, *pos*)

Return the text, *s*, justified according to *pos*, which may be "left", "center", or "right".

If *pos* is omitted it defaults to "right".

Null characters are replaced by spaces. All other character data are treated as non-white space.

Example:

```
strjust (["a"; "ab"; "abc"; "abcd"])
⇒
    "    a"
    "  ab"
    " abc"
    "abcd"
```

See also: [\[deblank\]](#), page 77, [\[strrep\]](#), page 86, [\[strtrim\]](#), page 78, [\[untabify\]](#), page 91.

x = str2num (*s*)
[x, state] = str2num (*s*)

Convert the string (or character array) *s* to a number (or an array).

Examples:

```
str2num ("3.141596")
⇒ 3.141596
```

```
str2num (["1, 2, 3"; "4, 5, 6"])
⇒ 1 2 3
    4 5 6
```

The optional second output, *state*, is logically true when the conversion is successful. If the conversion fails the numeric output, *x*, is empty and *state* is false.

Caution: As `str2num` uses the `eval` function to do the conversion, `str2num` will execute any code contained in the string *s*. Use `str2double` for a safer and faster conversion.

For cell array of strings use `str2double`.

See also: [\[str2double\]](#), page 95, [\[eval\]](#), page 161.

tolower (*s*)
lower (*s*)

Return a copy of the string or cell string *s*, with each uppercase character replaced by the corresponding lowercase one; non-alphabetic characters are left unchanged.

For example:

```
tolower ("MiXeD cAsE 123")
⇒ "mixed case 123"
```

See also: [toupper], page 97.

toupper (s)
upper (s)

Return a copy of the string or cell string *s*, with each lowercase character replaced by the corresponding uppercase one; non-alphabetic characters are left unchanged.

For example:

```
toupper ("MiXeD cAsE 123")
          => "MIXED CASE 123"
```

See also: [tolower], page 96.

```
native_bytes = unicode2native (utf8_str, codepage)
native_bytes = unicode2native (utf8_str)
```

Convert UTF-8 string *utf8_str* to byte stream using *codepage*.

The character vector *utf8_str* is converted to a byte stream *native_bytes* using the code page given by *codepage*. The string *codepage* must be an identifier of a valid code page. Examples for valid code pages are "ISO-8859-1", "Shift-JIS", or "UTF-16". For a list of supported code pages, see <https://www.gnu.org/software/libiconv>. If *codepage* is omitted or empty, the system default codepage is used.

If any of the characters cannot be mapped into the codepage *codepage*, they are replaced with the appropriate substitution sequence for that codepage.

See also: [native2unicode], page 97.

```
utf8_str = native2unicode (native_bytes, codepage)
utf8_str = native2unicode (native_bytes)
```

Convert byte stream *native_bytes* to UTF-8 using *codepage*.

The numbers in the vector *native_bytes* are rounded and clipped to integers between 0 and 255. This byte stream is then mapped into the code page given by the string *codepage* and returned in the string *utf8_str*. Octave uses UTF-8 as its internal encoding. The string *codepage* must be an identifier of a valid code page. Examples for valid code pages are "ISO-8859-1", "Shift-JIS", or "UTF-16". For a list of supported code pages, see <https://www.gnu.org/software/libiconv>. If *codepage* is omitted or empty, the system default codepage is used.

If *native_bytes* is a string vector, it is returned as is.

See also: [unicode2native], page 97.

do_string_escapes (string)

Convert escape sequences in *string* to the characters they represent.

Escape sequences begin with a leading backslash ('\\') followed by 1-3 characters (.e.g., "\\n" => newline).

See also: [undo_string_escapes], page 97.

undo_string_escapes (s)

Convert special characters in strings back to their escaped forms.

For example, the expression

```
bell = "\a";
```

assigns the value of the alert character (control-g, ASCII code 7) to the string variable `bell`. If this string is printed, the system will ring the terminal bell (if it is possible). This is normally the desired outcome. However, sometimes it is useful to be able to print the original representation of the string, with the special characters replaced by their escape sequences. For example,

```
octave:13> undo_string_escapes (bell)
ans = \a
```

replaces the unprintable alert character with its printable representation.

See also: [\[do_string_escapes\]](#), page 97.

5.7 Character Class Functions

Octave also provides the following character class test functions patterned after the functions in the standard C library. They all operate on string arrays and return matrices of zeros and ones. Elements that are nonzero indicate that the condition was true for the corresponding character in the string array. For example:

```
isalpha ("!Q@WERT^Y&")
⇒ [ 0, 1, 0, 1, 1, 1, 1, 0, 1, 0 ]
```

isalnum (*s*)

Return a logical array which is true where the elements of *s* are letters or digits and false where they are not.

This is equivalent to `(isalpha (s) | isdigit (s))`.

See also: [\[isalpha\]](#), page 98, [\[isdigit\]](#), page 99, [\[ispunct\]](#), page 99, [\[isspace\]](#), page 99, [\[iscntrl\]](#), page 99.

isalpha (*s*)

Return a logical array which is true where the elements of *s* are letters and false where they are not.

This is equivalent to `(islower (s) | isupper (s))`.

See also: [\[isdigit\]](#), page 99, [\[ispunct\]](#), page 99, [\[isspace\]](#), page 99, [\[iscntrl\]](#), page 99, [\[isalnum\]](#), page 98, [\[islower\]](#), page 98, [\[isupper\]](#), page 99.

isletter (*s*)

Return a logical array which is true where the elements of *s* are letters and false where they are not.

This is an alias for the `isalpha` function.

See also: [\[isalpha\]](#), page 98, [\[isdigit\]](#), page 99, [\[ispunct\]](#), page 99, [\[isspace\]](#), page 99, [\[iscntrl\]](#), page 99, [\[isalnum\]](#), page 98.

islower (*s*)

Return a logical array which is true where the elements of *s* are lowercase letters and false where they are not.

See also: [\[isupper\]](#), page 99, [\[isalpha\]](#), page 98, [\[isletter\]](#), page 98, [\[isalnum\]](#), page 98.

isupper (s)

Return a logical array which is true where the elements of *s* are uppercase letters and false where they are not.

See also: [islower], page 98, [isalpha], page 98, [isletter], page 98, [isalnum], page 98.

isdigit (s)

Return a logical array which is true where the elements of *s* are decimal digits (0-9) and false where they are not.

See also: [isxdigit], page 99, [isalpha], page 98, [isletter], page 98, [ispunct], page 99, [isspace], page 99, [iscntrl], page 99.

isxdigit (s)

Return a logical array which is true where the elements of *s* are hexadecimal digits (0-9 and a-fA-F).

See also: [isdigit], page 99.

ispunct (s)

Return a logical array which is true where the elements of *s* are punctuation characters and false where they are not.

See also: [isalpha], page 98, [isdigit], page 99, [isspace], page 99, [iscntrl], page 99.

isspace (s)

Return a logical array which is true where the elements of *s* are whitespace characters (space, formfeed, newline, carriage return, tab, and vertical tab) and false where they are not.

See also: [iscntrl], page 99, [ispunct], page 99, [isalpha], page 98, [isdigit], page 99.

iscntrl (s)

Return a logical array which is true where the elements of *s* are control characters and false where they are not.

See also: [ispunct], page 99, [isspace], page 99, [isalpha], page 98, [isdigit], page 99.

isgraph (s)

Return a logical array which is true where the elements of *s* are printable characters (but not the space character) and false where they are not.

See also: [isprint], page 99.

isprint (s)

Return a logical array which is true where the elements of *s* are printable characters (including the space character) and false where they are not.

See also: [isgraph], page 99.

isascii (s)

Return a logical array which is true where the elements of *s* are ASCII characters (in the range 0 to 127 decimal) and false where they are not.

isstrprop (str, prop)

Test character string properties.

For example:

```
isstrprop ("abc123", "alpha")
⇒ [1, 1, 1, 0, 0, 0]
```

If *str* is a cell array, *isstrpop* is applied recursively to each element of the cell array.

Numeric arrays are converted to character strings.

The second argument *prop* must be one of

"alpha" True for characters that are alphabetic (letters).

"alnum"

"alphanum"

True for characters that are alphabetic or digits.

"lower" True for lowercase letters.

"upper" True for uppercase letters.

"digit" True for decimal digits (0-9).

"xdigit" True for hexadecimal digits (a-fA-F0-9).

"space"

"wspace" True for whitespace characters (space, formfeed, newline, carriage return, tab, vertical tab).

"punct" True for punctuation characters (printing characters except space or letter or digit).

"cntrl" True for control characters.

"graph"

"graphic"

True for printing characters except space.

"print" True for printing characters including space.

"ascii" True for characters that are in the range of ASCII encoding.

See also: [\[isalpha\]](#), page 98, [\[isalnum\]](#), page 98, [\[islower\]](#), page 98, [\[isupper\]](#), page 99, [\[isdigit\]](#), page 99, [\[isxdigit\]](#), page 99, [\[isspace\]](#), page 99, [\[ispunct\]](#), page 99, [\[iscntrl\]](#), page 99, [\[isgraph\]](#), page 99, [\[isprint\]](#), page 99, [\[isascii\]](#), page 99.

6 Data Containers

Octave includes support for three different mechanisms to contain arbitrary data types in the same variable: Structures, which are C-like, and are indexed with named fields; containers.Map objects, which store data in key/value pairs; and cell arrays, where each element of the array can have a different data type and or shape. Multiple input arguments and return values of functions are organized as another data container, the comma separated list.

6.1 Structures

Octave includes support for organizing data in structures. The current implementation uses an associative array with indices limited to strings, but the syntax is more like C-style structures.

6.1.1 Basic Usage and Examples

Here are some examples of using data structures in Octave.

Elements of structures can be of any value type. For example, the three expressions

```
x.a = 1;
x.b = [1, 2; 3, 4];
x.c = "string";
```

create a structure with three elements. The ‘.’ character separates the structure name (in the example above `x`) from the field name and indicates to Octave that this variable is a structure. To print the value of the structure you can type its name, just as for any other variable:

```
x
⇒ x =
scalar structure containing the fields:

a = 1
b =
1   2
3   4

c = string
```

Note that Octave may print the elements in any order.

Structures may be copied just like any other variable:

```
y = x
⇒ y =

scalar structure containing the fields:

a = 1
b =

1 2
3 4

c = string
```

Since structures are themselves values, structure elements may reference other structures, as well. The following statement adds the field `d` to the structure `x`. The value of field `d` is itself a data structure containing the single field `a`, which has a value of 3.

```
x.d.a = 3;
x.d
⇒ ans =

scalar structure containing the fields:

a = 3

x
⇒ x =

scalar structure containing the fields:

a = 1
b =

1 2
3 4

c = string
d =

scalar structure containing the fields:

a = 3
```

Note that when Octave prints the value of a structure that contains other structures, only a few levels are displayed. For example:

```

a.b.c.d.e = 1;
a
⇒ a =
scalar structure containing the fields:

b =
scalar structure containing the fields:

c =
scalar structure containing the fields:

d: 1x1 scalar struct

```

This prevents long and confusing output from large deeply nested structures. The number of levels to print for nested structures may be set with the function `struct_levels_to_print`, and the function `print_struct_array_contents` may be used to enable printing of the contents of structure arrays.

```

val = struct_levels_to_print ()
old_val = struct_levels_to_print (new_val)
struct_levels_to_print (new_val, "local")

```

Query or set the internal variable that specifies the number of structure levels to display.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[print_struct_array_contents\]](#), page 103.

```

val = print_struct_array_contents ()
old_val = print_struct_array_contents (new_val)
print_struct_array_contents (new_val, "local")

```

Query or set the internal variable that specifies whether to print struct array contents.

If true, values of struct array elements are printed. This variable does not affect scalar structures whose elements are always printed. In both cases, however, printing will be limited to the number of levels specified by `struct_levels_to_print`.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[struct_levels_to_print\]](#), page 103.

Functions can return structures. For example, the following function separates the real and complex parts of a matrix and stores them in two elements of the same structure variable `y`.

```
function y = f (x)
  y.re = real (x);
  y.im = imag (x);
endfunction
```

When called with a complex-valued argument, the function **f** returns the data structure containing the real and imaginary parts of the original function argument.

```
f (rand (2) + rand (2) * I)
⇒ ans =

scalar structure containing the fields:

re =
0.040239  0.242160
0.238081  0.402523

im =
0.26475  0.14828
0.18436  0.83669
```

Function return lists can include structure elements, and they may be indexed like any other variable. For example:

```
[ x.u, x.s(2:3,2:3), x.v ] = svd ([1, 2; 3, 4]);
x

⇒ x =

scalar structure containing the fields:

u =
-0.40455  -0.91451
-0.91451  0.40455

s =
0.00000  0.00000  0.00000
0.00000  5.46499  0.00000
0.00000  0.00000  0.36597

v =
-0.57605  0.81742
-0.81742  -0.57605
```

It is also possible to cycle through all the elements of a structure in a loop, using a special form of the `for` statement (see [Section 10.5.1 \[Looping Over Structure Elements\]](#), page 171).

6.1.2 Structure Arrays

A structure array is a particular instance of a structure, where each of the fields of the structure is represented by a cell array. Each of these cell arrays has the same dimensions. Conceptually, a structure array can also be seen as an array of structures with identical fields. An example of the creation of a structure array is

```
x(1).a = "string1";
x(2).a = "string2";
x(1).b = 1;
x(2).b = 2;
```

which creates a 1-by-2 structure array with two fields. Another way to create a structure array is with the `struct` function (see [Section 6.1.3 \[Creating Structures\]](#), page 106). As previously, to print the value of the structure array, you can type its name:

```
x
⇒ x =
{
    1x2 struct array containing the fields:

        a
        b
}
```

Individual elements of the structure array can be returned by indexing the variable like `x(1)`, which returns a structure with two fields:

```
x(1)
⇒ ans =
{
    a = string1
    b = 1
}
```

Furthermore, the structure array can return a comma separated list of field values (see [Section 6.4 \[Comma Separated Lists\]](#), page 124), if indexed by one of its own field names. For example:

```
x.a
⇒
ans = string1
ans = string2
```

Here is another example, using this comma separated list on the left-hand side of an assignment:

```
[x.a] = deal ("new string1", "new string2");
x(1).a
⇒ ans = new string1
x(2).a
⇒ ans = new string2
```

Just as for numerical arrays, it is possible to use vectors as indices (see [Section 8.1 \[Index Expressions\]](#), page 139):

```
x(3:4) = x(1:2);
[x([1,3]).a] = deal ("other string1", "other string2");
x.a
⇒
ans = other string1
ans = new string2
ans = other string2
ans = new string2
```

The function `size` will return the size of the structure. For the example above

```
size (x)
⇒ ans =
1 4
```

Elements can be deleted from a structure array in a similar manner to a numerical array, by assigning the elements to an empty matrix. For example

```
in = struct ("call1", {x, Inf, "last"},
            "call2", {x, Inf, "first"})
⇒ in =
{
  1x3 struct array containing the fields:

    call1
    call2
}

in(1) = [];
in.call1
⇒
ans = Inf
ans = last
```

6.1.3 Creating Structures

Besides the index operator `"."`, Octave can use dynamic naming `"(var)"` or the `struct` function to create structures. Dynamic naming uses the string value of a variable as the field name. For example:

```
a = "field2";
x.a = 1;
x.(a) = 2;
x
⇒ x =
{
  a = 1
  field2 = 2
}
```

Dynamic indexing also allows you to use arbitrary strings, not merely valid Octave identifiers (note that this does not work on MATLAB):

```
a = "long field with spaces (and funny char$)";
x.a = 1;
x.(a) = 2;
x
⇒ x =
{
  a = 1
  long field with spaces (and funny char$) = 2
}
```

The warning id `Octave:language-extension` can be enabled to warn about this usage. See [\[warning_ids\]](#), page 230.

More realistically, all of the functions that operate on strings can be used to build the correct field name before it is entered into the data structure.

```
names = ["Bill"; "Mary"; "John"];
ages = [37; 26; 31];
for i = 1:rows (names)
  database.(names(i,:)) = ages(i);
endfor
database
⇒ database =
{
  Bill = 37
  Mary = 26
  John = 31
}
```

The third way to create structures is the `struct` command. `struct` takes pairs of arguments, where the first argument in the pair is the fieldname to include in the structure and the second is a scalar or cell array, representing the values to include in the structure or structure array. For example:

```
struct ("field1", 1, "field2", 2)
⇒ ans =
{
  field1 = 1
  field2 = 2
}
```

If the values passed to **struct** are a mix of scalar and cell arrays, then the scalar arguments are expanded to create a structure array with a consistent dimension. For example:

```
s = struct ("field1", {1, "one"}, "field2", {2, "two"},
           "field3", 3);
s.field1
⇒
ans = 1
ans = one

s.field2
⇒
ans = 2
ans = two

s.field3
⇒
ans = 3
ans = 3
```

If you want to create a struct which contains a cell array as an individual field, you must wrap it in another cell array as shown in the following example:

```
struct ("field1", {{1, "one"}}, "field2", 2)
⇒ ans =
{
    field1 =
    {
        [1,1] = 1
        [1,2] = one
    }

    field2 = 2
}
```

```
s = struct ()
s = struct (field1, value1, field2, value2, ...)
s = struct (obj)
```

Create a scalar or array structure and initialize its values.

The *field1*, *field2*, ... variables are strings specifying the names of the fields and the *value1*, *value2*, ... variables can be of any type.

If the values are cell arrays, create a structure array and initialize its values. The dimensions of each cell array of values must match. Singleton cells and non-cell values are repeated so that they fill the entire array. If the cells are empty, create an empty structure array with the specified field names.

If the argument is an object, return the underlying struct.

Observe that the syntax is optimized for struct **arrays**. Consider the following examples:

```

struct ("foo", 1)
    ⇒ scalar structure containing the fields:
        foo = 1

struct ("foo", {})
    ⇒ 0x0 struct array containing the fields:
        foo

struct ("foo", { {} })
    ⇒ scalar structure containing the fields:
        foo = {}(0x0)

struct ("foo", {1, 2, 3})
    ⇒ 1x3 struct array containing the fields:
        foo

```

The first case is an ordinary scalar struct—one field, one value. The second produces an empty struct array with one field and no values, since being passed an empty cell array of struct array values. When the value is a cell array containing a single entry, this becomes a scalar struct with that single entry as the value of the field. That single entry happens to be an empty cell array.

Finally, if the value is a non-scalar cell array, then **struct** produces a **struct array**.

See also: [\[cell2struct\]](#), page 124, [\[fieldnames\]](#), page 109, [\[getfield\]](#), page 111, [\[setfield\]](#), page 110, [\[rmfield\]](#), page 111, [\[isfield\]](#), page 110, [\[orderfields\]](#), page 111, [\[isstruct\]](#), page 109, [\[structfun\]](#), page 592.

The function **isstruct** can be used to test if an object is a structure or a structure array.

```
isstruct (x)
    Return true if x is a structure or a structure array.

See also: \[ismatrix\], page 63, \[iscell\], page 116, \[isa\], page 39.
```

6.1.4 Manipulating Structures

Other functions that can manipulate the fields of a structure are given below.

```
numfields (s)
    Return the number of fields of the structure s.
```

See also: [\[fieldnames\]](#), page 109.

```

names = fieldnames (struct)
names = fieldnames (obj)
names = fieldnames (javaobj)
names = fieldnames ("javaclassname")

```

Return a cell array of strings with the names of the fields in the specified input.

When the input is a structure *struct*, the names are the elements of the structure.

When the input is an Octave object *obj*, the names are the public properties of the object.

When the input is a Java object *javaobj* or a string containing the name of a Java class *javaclassname*, the names are the public fields (data members) of the object or class.

See also: [\[numfields\]](#), page 109, [\[isfield\]](#), page 110, [\[orderfields\]](#), page 111, [\[struct\]](#), page 108, [\[methods\]](#), page 837.

```
isfield (x, "name")
isfield (x, name)
```

Return true if the *x* is a structure and it includes an element named *name*.

If *name* is a cell array of strings then a logical array of equal dimension is returned.

See also: [\[fieldnames\]](#), page 109.

```
sout = setfield (s, field, val)
sout = setfield (s, sidx1, field1, fidx1, sidx2, field2, fidx2, ..., val)
```

Return a *copy* of the structure *s* with the field member *field* set to the value *val*.

For example:

```
s = struct ();
s = setfield (s, "foo bar", 42);
```

This is equivalent to

```
s.("foo bar") = 42;
```

Note that ordinary structure syntax *s.foo bar = 42* cannot be used here, as the field name is not a valid Octave identifier because of the space character. Using arbitrary strings for field names is incompatible with MATLAB, and this usage will emit a warning if the warning ID `Octave:language-extension` is enabled. See [\[warning_ids\]](#), page 230.

With the second calling form, set a field of a structure array. The input *sidx* selects an element of the structure array, *field* specifies the field name of the selected element, and *fidx* selects which element of the field (in the case of an array or cell array). The *sidx*, *field*, and *fidx* inputs can be repeated to address nested structure array elements. The structure array index and field element index must be cell arrays while the field name must be a string.

For example:

```
s = struct ("baz", 42);
setfield (s, {1}, "foo", {1}, "bar", 54)
⇒
ans =
scalar structure containing the fields:
baz = 42
foo =
scalar structure containing the fields:
bar = 54
```

The example begins with an ordinary scalar structure to which a nested scalar structure is added. In all cases, if the structure index *sidx* is not specified it defaults to

1 (scalar structure). Thus, the example above could be written more concisely as
`setfield (s, "foo", "bar", 54)`

Finally, an example with nested structure arrays:

```
sa.foo = 1;
sa = setfield (sa, {2}, "bar", {3}, "baz", {1, 4}, 5);
sa(2).bar(3)
⇒
ans =
    scalar structure containing the fields:
    baz = 0 0 0 5
```

Here `sa` is a structure array whose field at elements 1 and 2 is in turn another structure array whose third element is a simple scalar structure. The terminal scalar structure has a field which contains a matrix value.

Note that the same result as in the above example could be achieved by:

```
sa.foo = 1;
sa(2).bar(3).baz(1,4) = 5
```

See also: [getfield], page 111, [rmfield], page 111, [orderfields], page 111, [isfield], page 110, [fieldnames], page 109, [isstruct], page 109, [struct], page 108.

```
val = getfield (s, field)
val = getfield (s, sidx1, field1, fidx1, ...)
```

Get the value of the field named `field` from a structure or nested structure `s`.

If `s` is a structure array then `sidx` selects an element of the structure array, `field` specifies the field name of the selected element, and `fidx` selects which element of the field (in the case of an array or cell array). See `setfield` for a more complete description of the syntax.

See also: [setfield], page 110, [rmfield], page 111, [orderfields], page 111, [isfield], page 110, [fieldnames], page 109, [isstruct], page 109, [struct], page 108.

```
sout = rmfield (s, "f")
sout = rmfield (s, f)
```

Return a *copy* of the structure (array) `s` with the field `f` removed.

If `f` is a cell array of strings or a character array, remove each of the named fields.

See also: [orderfields], page 111, [fieldnames], page 109, [isfield], page 110.

```
sout = orderfields (s1)
sout = orderfields (s1, s2)
sout = orderfields (s1, {cellstr})
sout = orderfields (s1, p)
[sout, p] = orderfields (...)
```

Return a *copy* of `s1` with fields arranged alphabetically, or as specified by the second input.

Given one input struct `s1`, arrange field names alphabetically.

If a second struct argument is given, arrange field names in `s1` as they appear in `s2`. The second argument may also specify the order in a cell array of strings `cellstr`. The second argument may also be a permutation vector.

The optional second output argument p is the permutation vector which converts the original name order to the new name order.

Examples:

```
s = struct ("d", 4, "b", 2, "a", 1, "c", 3);
t1 = orderfields (s)
⇒ t1 =
    scalar structure containing the fields:
      a = 1
      b = 2
      c = 3
      d = 4

t = struct ("d", {}, "c", {}, "b", {}, "a", {});
t2 = orderfields (s, t)
⇒ t2 =
    scalar structure containing the fields:
      d = 4
      c = 3
      b = 2
      a = 1

t3 = orderfields (s, [3, 2, 4, 1])
⇒ t3 =
    scalar structure containing the fields:
      a = 1
      b = 2
      c = 3
      d = 4

[t4, p] = orderfields (s, {"d", "c", "b", "a"})
⇒ t4 =
    scalar structure containing the fields:
      d = 4
      c = 3
      b = 2
      a = 1

p =
  1
  4
  2
  3
```

See also: [\[fieldnames\]](#), page 109, [\[getfield\]](#), page 111, [\[setfield\]](#), page 110, [\[rmfield\]](#), page 111, [\[isfield\]](#), page 110, [\[isstruct\]](#), page 109, [\[struct\]](#), page 108.

substruct (type, subs, ...)

Create a subscript structure for use with `subsref` or `subsasgn`.

For example:

```

idx = substruct ("()", {3, ":"})
⇒ idx =
    scalar structure containing the fields:
        type = ()
        subs =
    {
        [1,1] = 3
        [1,2] = :
    }
x = [1, 2, 3;
      4, 5, 6;
      7, 8, 9];
subsref (x, idx)
⇒ 7 8 9

```

See also: [\[subsref\]](#), page 840, [\[subsasgn\]](#), page 842.

6.1.5 Processing Data in Structures

The simplest way to process data in a structure is within a `for` loop (see [Section 10.5.1 \[Looping Over Structure Elements\]](#), page 171). A similar effect can be achieved with the `structfun` function, where a user defined function is applied to each field of the structure. See [\[structfun\]](#), page 592.

Alternatively, to process the data in a structure, the structure might be converted to another type of container before being treated.

`c = struct2cell (s)`

Create a new cell array from the objects stored in the struct object.

If f is the number of fields in the structure, the resulting cell array will have a dimension vector corresponding to `[f size(s)]`. For example:

```

s = struct ("name", {"Peter", "Hannah", "Robert"},
            "age", {23, 16, 3});
c = struct2cell (s)
⇒ c = {2x1x3 Cell Array}
c(1,1,:)(:)
⇒
{
    [1,1] = Peter
    [2,1] = Hannah
    [3,1] = Robert
}
c(2,1,:)(:)
⇒
{
    [1,1] = 23
    [2,1] = 16
    [3,1] = 3
}

```

See also: [\[cell2struct\]](#), page 124, [\[fieldnames\]](#), page 109.

6.2 containers.Map

```
m = containers.Map()
m = containers.Map(keys, vals)
m = containers.Map(keys, vals, "UniformValues", is_uniform)
m = containers.Map("KeyType", kt, "ValueType", vt)
```

Create an object of the containers.Map class that stores a list of key/value pairs.

keys is an array of *unique* keys for the map. The keys can be numeric scalars or strings. The type for numeric keys may be one of "double", "single", "int32", "uint32", "int64", or "uint64". Other numeric or logical keys will be converted to "double". A single string key may be entered as is. Multiple string keys are entered as a cell array of strings.

vals is an array of values for the map with the *same* number of elements as *keys*.

When called with no input arguments a default map is created with strings as the key type and "any" as the value type.

The "UniformValues" option specifies whether the values of the map must be strictly of the same type. If *is_uniform* is true, any values which would be added to the map are first validated to ensure they are of the correct type.

When called with "KeyType" and "ValueType" arguments, create an empty map with the specified types. The inputs *kt* and *vt* are the types for the keys and values of the map respectively. Allowed values for *kt* are "char", "double", "single", "int32", "uint32", "int64", "uint64". Allowed values for *vt* are "any", "char", "double", "single", "int32", "uint32", "int64", "uint64", "logical".

The return value *m* is an object of the containers.Map class.

See also: [\[struct\]](#), page 108.

6.3 Cell Arrays

It can be both necessary and convenient to store several variables of different size or type in one variable. A cell array is a container class able to do just that. In general cell arrays work just like *N*-dimensional arrays with the exception of the use of '{' and '}' as allocation and indexing operators.

6.3.1 Basic Usage of Cell Arrays

As an example, the following code creates a cell array containing a string and a 2-by-2 random matrix

```
c = {"a string", rand(2, 2)};
```

To access the elements of a cell array, it can be indexed with the { and } operators. Thus, the variable created in the previous example can be indexed like this:

```
c{1}
⇒ ans = a string
```

As with numerical arrays several elements of a cell array can be extracted by indexing with a vector of indexes

```
c{1:2}
⇒ ans = a string
⇒ ans =
0.593993 0.627732
0.377037 0.033643
```

The indexing operators can also be used to insert or overwrite elements of a cell array. The following code inserts the scalar 3 on the third place of the previously created cell array

```
c{3} = 3
⇒ c =
{
[1,1] = a string
[1,2] =
0.593993 0.627732
0.377037 0.033643

[1,3] = 3
}
```

Details on indexing cell arrays are explained in [Section 6.3.3 \[Indexing Cell Arrays\]](#), [page 120](#).

In general nested cell arrays are displayed hierarchically as in the previous example. In some circumstances it makes sense to reference them by their index, and this can be performed by the `celldisp` function.

```
celldisp (c)
celldisp (c, name)
```

Recursively display the contents of a cell array.

By default the values are displayed with the name of the variable `c`. However, this name can be replaced with the variable `name`. For example:

```
c = {1, 2, {31, 32}};
celldisp (c, "b")
⇒
b{1} =
1
b{2} =
2
b{3}{1} =
31
b{3}{2} =
32
```

See also: `[disp]`, [page 251](#).

To test if an object is a cell array, use the `iscell` function. For example:

```
iscell (c)
⇒ ans = 1
```

```
iscell (3)
⇒ ans = 0
```

```
iscell (x)
```

Return true if x is a cell array object.

See also: `[ismatrix]`, page 63, `[isstruct]`, page 109, `[iscellstr]`, page 123, `[isa]`, page 39.

6.3.2 Creating Cell Arrays

The introductory example (see [Section 6.3.1 \[Basic Usage of Cell Arrays\]](#), page 114) showed how to create a cell array containing currently available variables. In many situations, however, it is useful to create a cell array and then fill it with data.

The `cell` function returns a cell array of a given size, containing empty matrices. This function is similar to the `zeros` function for creating new numerical arrays. The following example creates a 2-by-2 cell array containing empty matrices

```
c = cell (2,2)
⇒ c =
{
    [1,1] = []
    [2,1] = []
    [1,2] = []
    [2,2] = []
}
```

Just like numerical arrays, cell arrays can be multi-dimensional. The `cell` function accepts any number of positive integers to describe the size of the returned cell array. It is also possible to set the size of the cell array through a vector of positive integers. In the following example two cell arrays of equal size are created, and the size of the first one is displayed

```
c1 = cell (3, 4, 5);
c2 = cell ([3, 4, 5]);
size (c1)
⇒ ans =
    3   4   5
```

As can be seen, the [\[size\]](#), page 45, function also works for cell arrays. As do other functions describing the size of an object, such as [\[length\]](#), page 45, [\[numel\]](#), page 44, [\[rows\]](#), page 44, and [\[columns\]](#), page 44.

```
cell (n)
cell (m, n)
cell (m, n, k, ...)
```

cell ([m n ...])

Create a new cell array object.

If invoked with a single scalar integer argument, return a square NxN cell array. If invoked with two or more scalar integer arguments, or a vector of integer values, return an array with the given dimensions.

See also: [cellstr], page 123, [mat2cell], page 118, [num2cell], page 117, [struct2cell], page 113.

As an alternative to creating empty cell arrays, and then filling them, it is possible to convert numerical arrays into cell arrays using the **num2cell**, **mat2cell** and **cellslices** functions.

C = num2cell (A)

C = num2cell (A, dim)

Convert the numeric matrix *A* to a cell array.

When no *dim* is specified, each element of *A* becomes a 1x1 element in the output *C*.

If *dim* is defined then individual elements of *C* contain all of the elements from *A* along the specified dimension. *dim* may also be a vector of dimensions with the same rule applied.

For example:

x = [1,2;3,4]

⇒

1	2
3	4

each element of A becomes a 1x1 element of C

num2cell (x)

⇒

```
{
    [1,1] = 1
    [2,1] = 3
    [1,2] = 2
    [2,2] = 4
}
```

all rows (dim 1) of A appear in each element of C

num2cell (x, 1)

⇒

```
{
    [1,1] =
        1
        3
    [1,2] =
        2
        4
}
```

all columns (dim 2) of A appear in each element of C

```

num2cell (x, 2)
⇒
{
  [1,1] =
    1   2
  [2,1] =
    3   4
}
## all rows and cols appear in each element of C
## (hence, only 1 output)
num2cell (x, [1, 2])
⇒
{
  [1,1] =
    1   2
    3   4
}

```

See also: [\[mat2cell\]](#), page 118.

```

C = mat2cell (A, dim1, dim2, ..., dimi, ..., dimn)
C = mat2cell (A, rowdim)
Convert the matrix A to a cell array.

```

Each dimension argument (dim_1 , dim_2 , etc.) is a vector of integers which specifies how to divide that dimension's elements amongst the new elements in the output C . The number of elements in the i -th dimension is `size (A, i)`. Because all elements in A must be partitioned, there is a requirement that `sum (di) == size (A, i)`. The size of the output cell C is `numel (dim1) x numel (dim2) x ... x numel (dimn)`.

Given a single dimensional argument, `rowdim`, the output is divided into rows as specified. All other dimensions are not divided and thus all columns (dim 2), pages (dim 3), etc. appear in each output element.

Examples

```

x = reshape (1:12, [3, 4])'
⇒
  1   2   3
  4   5   6
  7   8   9
 10  11  12

```

```
## The 4 rows (dim1) are divided in to two cell elements
## with 2 rows each.
## The 3 cols (dim2) are divided in to three cell elements
## with 1 col each.
mat2cell (x, [2,2], [1,1,1])
⇒
{
  [1,1] =
    1
    4

  [2,1] =
    7
    10

  [1,2] =
    2
    5

  [2,2] =
    8
    11

  [1,3] =
    3
    6

  [2,3] =
    9
    12
}
```

```

## The 4 rows (dim1) are divided in to two cell elements
## with a 3/1 split.
## All columns appear in each output element.
mat2cell (x, [3,1])
⇒
{
  [1,1] =
    1   2   3
    4   5   6
    7   8   9

  [2,1] =
    10  11  12
}

```

See also: [\[num2cell\]](#), page 117, [\[cell2mat\]](#), page 124.

sl = cellslices (x, lb, ub, dim)

Given an array *x*, this function produces a cell array of slices from the array determined by the index vectors *lb*, *ub*, for lower and upper bounds, respectively.

In other words, it is equivalent to the following code:

```

n = length (lb);
sl = cell (1, n);
for i = 1:length (lb)
  sl{i} = x(:,...,lb(i):ub(i),...,:);
endfor

```

The position of the index is determined by *dim*. If not specified, slicing is done along the first non-singleton dimension.

See also: [\[cell2mat\]](#), page 124, [\[cellindexmat\]](#), page 122, [\[cellfun\]](#), page 590.

6.3.3 Indexing Cell Arrays

As shown in see [Section 6.3.1 \[Basic Usage of Cell Arrays\]](#), page 114, elements can be extracted from cell arrays using the ‘{’ and ‘}’ operators. If you want to extract or access subarrays which are still cell arrays, you need to use the ‘(’ and ‘)’ operators. The following example illustrates the difference:

```

c = {"1", "2", "3"; "x", "y", "z"; "4", "5", "6"};
c{2,3}
⇒ ans = z

c(2,3)
⇒ ans =
{
  [1,1] = z
}

```

So with ‘{}’ you access elements of a cell array, while with ‘()’ you access a sub array of a cell array.

Using the ‘(’ and ‘)’ operators, indexing works for cell arrays like for multi-dimensional arrays. As an example, all the rows of the first and third column of a cell array can be set to 0 with the following command:

```
c(:, [1, 3]) = {0}
⇒ =
{
    [1,1] = 0
    [2,1] = 0
    [3,1] = 0
    [1,2] = 2
    [2,2] = y
    [3,2] = 5
    [1,3] = 0
    [2,3] = 0
    [3,3] = 0
}
```

Note, that the above can also be achieved like this:

```
c(:, [1, 3]) = 0;
```

Here, the scalar ‘0’ is automatically promoted to cell array ‘{0}’ and then assigned to the subarray of c.

To give another example for indexing cell arrays with ‘()’, you can exchange the first and the second row of a cell array as in the following command:

```
c = {1, 2, 3; 4, 5, 6};
c([1, 2], :) = c([2, 1], :)
⇒ =
{
    [1,1] = 4
    [2,1] = 1
    [1,2] = 5
    [2,2] = 2
    [1,3] = 6
    [2,3] = 3
}
```

Accessing multiple elements of a cell array with the ‘{’ and ‘}’ operators will result in a comma-separated list of all the requested elements (see [Section 6.4 \[Comma Separated Lists\], page 124](#)). Using the ‘{’ and ‘}’ operators the first two rows in the above example can be swapped back like this:

```
[c{[1,2], :}] = deal (c{[2, 1], :})
⇒ =
{
    [1,1] = 1
    [2,1] = 4
    [1,2] = 2
    [2,2] = 5
    [1,3] = 3
    [2,3] = 6
}
```

As for struct arrays and numerical arrays, the empty matrix ‘[]’ can be used to delete elements from a cell array:

```
x = {"1", "2"; "3", "4"};
x(1, :) = []
⇒ x =
{
    [1,1] = 3
    [1,2] = 4
}
```

The following example shows how to just remove the contents of cell array elements but not delete the space for them:

```
x = {"1", "2"; "3", "4"};
x(1, :) = {[[]]}
⇒ x =
{
    [1,1] = [] (0x0)
    [2,1] = 3
    [1,2] = [] (0x0)
    [2,2] = 4
}
```

The indexing operations operate on the cell array and not on the objects within the cell array. By contrast, `cellindexmat` applies matrix indexing to the objects within each cell array entry and returns the requested values.

```
y = cellindexmat (x, varargin)
Perform indexing of matrices in a cell array.
```

Given a cell array of matrices `x`, this function computes

```
Y = cell (size (X));
for i = 1:numel (X)
    Y{i} = X{i}(varargin{1}, varargin{2}, ..., varargin{N});
endfor
```

The indexing arguments may be scalar (2), arrays ([1, 3]), ranges (1:3), or the colon operator ("::"). However, the indexing keyword `end` is not available.

See also: `[cellslices]`, page 120, `[cellfun]`, page 590.

6.3.4 Cell Arrays of Strings

One common use of cell arrays is to store multiple strings in the same variable. It is also possible to store multiple strings in a character matrix by letting each row be a string. This, however, introduces the problem that all strings must be of equal length. Therefore, it is recommended to use cell arrays to store multiple strings. For cases, where the character matrix representation is required for an operation, there are several functions that convert a cell array of strings to a character array and back. `char` and `strvcat` convert cell arrays to a character array (see [Section 5.3.1 \[Concatenating Strings\], page 70](#)), while the function `cellstr` converts a character array to a cell array of strings:

```
a = ["hello"; "world"];
c = cellstr (a)
⇒ c =
{
    [1,1] = hello
    [2,1] = world
}
```

`cstr = cellstr (strmat)`

Create a new cell array object from the elements of the string array `strmat`.

Each row of `strmat` becomes an element of `cstr`. Any trailing spaces in a row are deleted before conversion.

To convert back from a `cellstr` to a character array use `char`.

See also: [\[cell\]](#), page 116, [\[char\]](#), page 71.

One further advantage of using cell arrays to store multiple strings is that most functions for string manipulations included with Octave support this representation. As an example, it is possible to compare one string with many others using the `strcmp` function. If one of the arguments to this function is a string and the other is a cell array of strings, each element of the cell array will be compared to the string argument:

```
c = {"hello", "world"};
strcmp ("hello", c)
⇒ ans =
1   0
```

The following string functions support cell arrays of strings: `char`, `strvcat`, `strcat` (see [Section 5.3.1 \[Concatenating Strings\], page 70](#)), `strcmp`, `strncmp`, `strcmpi`, `strncmpi` (see [Section 5.4 \[Comparing Strings\], page 76](#)), `str2double`, `deblank`, `strtrim`, `strtrunc`, `strfind`, `strmatch`, `regexp`, `regexpi` (see [Section 5.5 \[Manipulating Strings\], page 77](#)) and `str2double` (see [Section 5.6 \[String Conversions\], page 92](#)).

The function `iscellstr` can be used to test if an object is a cell array of strings.

`iscellstr (cell)`

Return true if every element of the cell array `cell` is a character string.

See also: [\[ischar\]](#), page 68, [\[isstring\]](#), page 68.

6.3.5 Processing Data in Cell Arrays

Data that is stored in a cell array can be processed in several ways depending on the actual data. The simplest way to process that data is to iterate through it using one or more `for` loops. The same idea can be implemented more easily through the use of the `cellfun` function that calls a user-specified function on all elements of a cell array. See [`cellfun`], page 590.

An alternative is to convert the data to a different container, such as a matrix or a data structure. Depending on the data this is possible using the `cell2mat` and `cell2struct` functions.

`m = cell2mat (c)`

Convert the cell array `c` into a matrix by concatenating all elements of `c` into a hyperrectangle.

Elements of `c` must be numeric, logical, or char matrices; or cell arrays; or structs; and `cat` must be able to concatenate them together.

See also: [`mat2cell`], page 118, [`num2cell`], page 117.

`cell2struct (cell, fields)`
`cell2struct (cell, fields, dim)`

Convert `cell` to a structure.

The number of fields in `fields` must match the number of elements in `cell` along dimension `dim`, that is `numel (fields) == size (cell, dim)`. If `dim` is omitted, a value of 1 is assumed.

```
A = cell2struct ({"Peter", "Hannah", "Robert";
                  185, 170, 168},
                  {"Name", "Height"}, 1);

A(1)
⇒
{
    Name    = Peter
    Height = 185
}
```

See also: [`struct2cell`], page 113, [`cell2mat`], page 124, [`struct`], page 108.

6.4 Comma Separated Lists

Comma separated lists¹ are the basic argument type to all Octave functions - both for input and return arguments. In the example

`max (a, b)`

'`a, b`' is a comma separated list. Comma separated lists can appear on both the right and left hand side of an assignment. For example

```
x = [1 0 1 0 0 1 1; 0 0 0 0 0 0 7];
[i, j] = find (x, 2, "last");
```

¹ Comma-separated lists are also sometimes informally referred to as *cs-lists*.

Here, ‘`x, 2, "last"`’ is a comma separated list constituting the input arguments of `find`. `find` returns a comma separated list of output arguments which is assigned element by element to the comma separated list ‘`i, j`’.

Another example of where comma separated lists are used is in the creation of a new array with `[]` (see [Section 4.1 \[Matrices\], page 48](#)) or the creation of a cell array with `{}` (see [Section 6.3.1 \[Basic Usage of Cell Arrays\], page 114](#)). In the expressions

```
a = [1, 2, 3, 4];
c = {4, 5, 6, 7};
```

both ‘`1, 2, 3, 4`’ and ‘`4, 5, 6, 7`’ are comma separated lists.

Comma separated lists cannot be directly manipulated by the user. However, both structure arrays and cell arrays can be converted into comma separated lists, and thus used in place of explicitly written comma separated lists. This feature is useful in many ways, as will be shown in the following subsections.

6.4.1 Comma Separated Lists Generated from Cell Arrays

As has been mentioned above (see [Section 6.3.3 \[Indexing Cell Arrays\], page 120](#)), elements of a cell array can be extracted into a comma separated list with the `{` and `}` operators. By surrounding this list with `[` and `]`, it can be concatenated into an array. For example:

```
a = {1, [2, 3], 4, 5, 6};
b = [a{1:4}]
⇒ b =
    1   2   3   4   5
```

Similarly, it is possible to create a new cell array containing cell elements selected with `{}`. By surrounding the list with ‘`{`’ and ‘`}`’ a new cell array will be created, as the following example illustrates:

```
a = {1, rand(2, 2), "three"};
b = { a{ [1, 3] } }
⇒ b =
{
    [1,1] =  1
    [1,2] = three
}
```

Furthermore, cell elements (accessed by `{}`) can be passed directly to a function. The list of elements from the cell array will be passed as an argument list to a given function as if it is called with the elements as individual arguments. The two calls to `printf` in the following example are identical but the latter is simpler and can handle cell arrays of an arbitrary size:

```
c = {"GNU", "Octave", "is", "Free", "Software"};
printf ("%s ", c{1}, c{2}, c{3}, c{4}, c{5});
    → GNU Octave is Free Software
printf ("%s ", c{:});
    → GNU Octave is Free Software
```

If used on the left-hand side of an assignment, a comma separated list generated with `{}` can be assigned to. An example is

```
in{1} = [10, 20, 30];
```

```

inf{2} = inf;
in{3} = "last";
in{4} = "first";
out = cell (4, 1);
[out{1:3}] = in{1 : 3};
[out{4:6}] = in{[1, 2, 4]}
⇒ out =
{
    [1,1] =
        10    20    30

    [2,1] = Inf
    [3,1] = last
    [4,1] =

        10    20    30

    [5,1] = Inf
    [6,1] = first
}

```

6.4.2 Comma Separated Lists Generated from Structure Arrays

Structure arrays can equally be used to create comma separated lists. This is done by addressing one of the fields of a structure array. For example:

```

x = ceil (randn (10, 1));
in = struct ("call1", {x, 3, "last"}, 
            "call2", {x, inf, "first"});
out = struct ("call1", cell (2, 1), "call2", cell (2, 1));
[out.call1] = find (in.call1);
[out.call2] = find (in.call2);

```

7 Variables

Variables let you give names to values and refer to them later. You have already seen variables in many of the examples. The name of a variable must be a sequence of letters, digits and underscores, but it may not begin with a digit. Octave does not enforce a limit on the length of variable names, but it is seldom useful to have variables with names longer than about 30 characters. The following are all valid variable names

```
x
x15
__foo_bar_baz__
fucnrdthsucngtagdjb
```

However, names like `__foo_bar_baz__` that begin and end with two underscores are understood to be reserved for internal use by Octave. You should not use them in code you write, except to access Octave's documented internal variables and built-in symbolic constants.

Case is significant in variable names. The symbols `a` and `A` are distinct variables.

A variable name is a valid expression by itself. It represents the variable's current value. Variables are given new values with *assignment operators* and *increment operators*. See Section 8.6 [Assignment Expressions], page 156.

There is one built-in variable with a special meaning. The `ans` variable always contains the result of the last computation, where the output wasn't assigned to any variable. The code `a = cos (pi)` will assign the value -1 to the variable `a`, but will not change the value of `ans`. However, the code `cos (pi)` will set the value of `ans` to -1.

Variables in Octave do not have fixed types, so it is possible to first store a numeric value in a variable and then to later use the same name to hold a string value in the same program. Variables may not be used before they have been given a value. Doing so results in an error.

ans [Automatic Variable]

The most recently computed result that was not explicitly assigned to a variable.

For example, after the expression

```
3^2 + 4^2
```

is evaluated, the value returned by `ans` is 25.

isvarname (name)

Return true if `name` is a valid variable name.

A valid variable name is composed of letters, digits, and underscores ("_"), and the first character must not be a digit.

See also: `[iskeyword]`, page 1035, `[exist]`, page 134, `[who]`, page 132.

```
varname = matlab.lang.makeValidName (str)
varname = matlab.lang.makeValidName (... , "ReplacementStyle", rs)
varname = matlab.lang.makeValidName (... , "Prefix", pfx)
[varname, ismodified] = matlab.lang.makeValidName (...)
```

Create valid variable name `varname` from `str`.

The input *str* must be a string or a cell array of strings. The output *varname* will be of the same type.

A valid variable name is a sequence of letters, digits, and underscores that does not begin with a digit.

The "ReplacementStyle" option specifies how invalid characters are handled. Acceptable values are

"underscore" (default)

Replace all invalid characters with an underscore ("_").

"delete" Remove any invalid character.

"hex" Replace all invalid characters with their hexadecimal representation.

Whitespace characters are always removed **prior** to the application of the "ReplacementStyle". Lowercase letters following a whitespace will be changed to uppercase.

The "Prefix" option specifies the string *pxf* to add as a prefix to the input if it begins with a digit. *pxf* must be a valid variable name itself. The default prefix is "x".

The optional output *ismodified* is a logical array indicating whether the respective element in *str* was a valid name or not.

See also: [\[iskeyword\]](#), page 1035, [\[isvarname\]](#), page 127, [\[matlab.lang.makeUniqueStrings\]](#),
page 128.

```
uniqstr = matlab.lang.makeUniqueStrings (str)
uniqstr = matlab.lang.makeUniqueStrings (str, ex)
uniqstr = matlab.lang.makeUniqueStrings (str, ex, maxlen)
[uniqstr, ismodified] = matlab.lang.makeUniqueStrings (...)
```

Construct a list of unique strings from a list of strings.

The input *str* must be a string or a cell array of strings. The output *uniqstr* will be of the same type.

The algorithm makes two strings unique by appending an underscore ("_" and a numeric count to the second string.

If *ex* is a string or a cell array of strings, *uniqstr* will contain elements that are unique between themselves and with respect to *ex*.

If *ex* is an index array or a logical array for *str* then it selects the subset of *str* that are made unique. Unselected elements are not modified.

The optional input *maxlength* specifies the maximum length of any string in *uniqstr*. If an input string cannot be made unique without exceeding *maxlength* an error is emitted.

The optional output *ismodified* is a logical array indicating whether each element in *str* was modified to make it unique.

See also: [\[unique\]](#), page 743, [\[matlab.lang.makeValidName\]](#), page 127.

```
namelengthmax ()
```

Return the MATLAB compatible maximum variable name length.

Octave is capable of storing strings up to $2^{31} - 1$ in length. However for MATLAB compatibility all variable, function, and structure field names should be shorter than the length returned by `namelengthmax`. In particular, variables stored to a MATLAB file format (*.mat) will have their names truncated to this length.

7.1 Global Variables

A variable that has been declared *global* may be accessed from within a function body without having to pass it as a formal parameter.

A variable may be declared global using a `global` declaration statement. The following statements are all global declarations.

```
global a
global a b
global c = 2
global d = 3 e f = 5
```

A global variable may only be initialized once in a `global` statement. For example, after executing the following code

```
global gvar = 1
global gvar = 2
```

the value of the global variable `gvar` is 1, not 2. Issuing a ‘`clear gvar`’ command does not change the above behavior, but ‘`clear all`’ does.

It is necessary declare a variable as global within a function body in order to access it. For example,

```
global x
function f ()
    x = 1;
endfunction
f ()
```

does *not* set the value of the global variable `x` to 1. In order to change the value of the global variable `x`, you must also declare it to be global within the function body, like this

```
function f ()
    global x;
    x = 1;
endfunction
```

Passing a global variable in a function parameter list will make a local copy and not modify the global value. For example, given the function

```
function f (x)
    x = 0
endfunction
```

and the definition of `x` as a global variable at the top level,

```
global x = 13
```

the expression

```
f (x)
```

will display the value of `x` from inside the function as 0, but the value of `x` at the top level remains unchanged, because the function works with a *copy* of its argument.

isglobal (name)

Return true if *name* is a globally visible variable.

For example:

```
global x
isglobal ("x")
⇒ 1
```

See also: [isvarname], page 127, [exist], page 134.

7.2 Persistent Variables

A variable that has been declared *persistent* within a function will retain its contents in memory between subsequent calls to the same function. The difference between persistent variables and global variables is that persistent variables are local in scope to a particular function and are not visible elsewhere.

The following example uses a persistent variable to create a function that prints the number of times it has been called.

```
function count_calls ()
    persistent calls = 0;
    printf ("'count_calls' has been called %d times\n",
           ++calls);
endfunction

for i = 1:3
    count_calls ();
endfor

⇒ 'count_calls' has been called 1 times
⇒ 'count_calls' has been called 2 times
⇒ 'count_calls' has been called 3 times
```

As the example shows, a variable may be declared persistent using a **persistent** declaration statement. The following statements are all persistent declarations.

```
persistent a
persistent a b
persistent c = 2
persistent d = 3 e f = 5
```

The behavior of persistent variables is equivalent to the behavior of static variables in C.

One restriction for persistent variables is, that neither input nor output arguments of a function can be persistent:

```
function y = foo ()
    persistent y = 0; # Not allowed!
endfunction

foo ()
⇒ error: can't make function parameter y persistent
```

Like global variables, a persistent variable may only be initialized once. For example, after executing the following code

```
persistent pvar = 1
persistent pvar = 2
```

the value of the persistent variable `pvar` is 1, not 2.

If a persistent variable is declared but not initialized to a specific value, it will contain an empty matrix. So, it is also possible to initialize a persistent variable by checking whether it is empty, as the following example illustrates.

```
function count_calls ()
    persistent calls;
    if (isempty (calls))
        calls = 0;
    endif
    printf ("'count_calls' has been called %d times\n",
           ++calls);
endfunction
```

This implementation behaves in exactly the same way as the previous implementation of `count_calls`.

The value of a persistent variable is kept in memory until it is explicitly cleared. Assuming that the implementation of `count_calls` is saved on disk, we get the following behavior.

```
for i = 1:2
    count_calls ();
endfor
→ 'count_calls' has been called 1 times
→ 'count_calls' has been called 2 times

clear
for i = 1:2
    count_calls ();
endfor
→ 'count_calls' has been called 3 times
→ 'count_calls' has been called 4 times

clear all
for i = 1:2
    count_calls ();
endfor
→ 'count_calls' has been called 1 times
→ 'count_calls' has been called 2 times

clear count_calls
for i = 1:2
    count_calls ();
endfor
```

```

└ 'count_calls' has been called 1 times
└ 'count_calls' has been called 2 times

```

That is, the persistent variable is only removed from memory when the function containing the variable is removed. Note that if the function definition is typed directly into the Octave prompt, the persistent variable will be cleared by a simple `clear` command as the entire function definition will be removed from memory. If you do not want a persistent variable to be removed from memory even if the function is cleared, you should use the `mlock` function (see [Section 11.9.6 \[Function Locking\]](#), page 204).

7.3 Status of Variables

When creating simple one-shot programs it can be very convenient to see which variables are available at the prompt. The function `who` and its siblings `whos` and `whos_line_format` will show different information about what is in memory, as the following shows.

```

str = "A random string";
who
└ Variables in the current scope:
└
└ ans  str

who
who pattern ...
who option pattern ...
C = who ("pattern", ...)
List currently defined variables matching the given patterns.

Valid pattern syntax is the same as described for the clear command. If no patterns are supplied, all variables are listed.

```

By default, only variables visible in the local scope are displayed.

The following are valid options, but may not be combined.

<code>global</code>	List variables in the global scope rather than the current scope.
<code>-regexp</code>	The patterns are considered to be regular expressions when matching the variables to display. The same pattern syntax accepted by the <code>regexp</code> function is used.
<code>-file</code>	The next argument is treated as a filename. All variables found within the specified file are listed. No patterns are accepted when reading variables from a file.

If called as a function, return a cell array of defined variable names matching the given patterns.

See also: [\[whos\]](#), page 132, [\[isglobal\]](#), page 129, [\[isvarname\]](#), page 127, [\[exist\]](#), page 134, [\[regexp\]](#), page 88.

```

whos
whos pattern ...
whos option pattern ...

```

S = whos ("pattern", ...)

Provide detailed information on currently defined variables matching the given patterns.

Options and pattern syntax are the same as for the **who** command.

Extended information about each variable is summarized in a table with the following default entries.

Attr	Attributes of the listed variable. Possible attributes are:
blank	Variable in local scope
a	Automatic variable. An automatic variable is one created by the interpreter, for example argn .
c	Variable of complex type.
f	Formal parameter (function argument).
g	Variable with global scope.
p	Persistent variable.
Name	The name of the variable.
Size	The logical size of the variable. A scalar is 1x1, a vector is 1xN or Nx1, a 2-D matrix is MxN.
Bytes	The amount of memory currently used to store the variable.
Class	The class of the variable. Examples include double, single, char, uint16, cell, and struct.

The table can be customized to display more or less information through the function **whos_line_format**.

If **whos** is called as a function, return a struct array of defined variable names matching the given patterns. Fields in the structure describing each variable are: name, size, bytes, class, global, sparse, complex, nesting, persistent.

See also: [who], page 132, [whos_line_format], page 133.

```
val = whos_line_format ()
old_val = whos_line_format (new_val)
whos_line_format (new_val, "local")
```

Query or set the format string used by the command **whos**.

A full format string is:

```
%[modifier]<command>[:width[:left-min[:balance]]];
```

The following command sequences are available:

%a	Prints attributes of variables (g=global, p=persistent, f=formal parameter, a=automatic variable).
%b	Prints number of bytes occupied by variables.
%c	Prints class names of variables.
%e	Prints elements held by variables.

- %n Prints variable names.
- %s Prints dimensions of variables.
- %t Prints type names of variables.

Every command may also have an alignment modifier:

- l Left alignment.
- r Right alignment (default).
- c Column-aligned (only applicable to command %s).

The `width` parameter is a positive integer specifying the minimum number of columns used for printing. No maximum is needed as the field will auto-expand as required.

The parameters `left-min` and `balance` are only available when the column-aligned modifier is used with the command '%s'. `balance` specifies the column number within the field width which will be aligned between entries. Numbering starts from 0 which indicates the leftmost column. `left-min` specifies the minimum field width to the left of the specified balance column.

The default format is:

```
"%a:4; %ln:6; %cs:16:6:1; %rb:12; %lc:-1;\n"
```

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [whos], page 132.

Instead of displaying which variables are in memory, it is possible to determine if a given variable is available. That way it is possible to alter the behavior of a program depending on the existence of a variable. The following example illustrates this.

```
if (! exist ("meaning", "var"))
    disp ("The program has no 'meaning'");
endif

c = exist (name)
c = exist (name, type)
```

Check for the existence of `name` as a variable, function, file, directory, or class.

The return code `c` is one of

- 1 `name` is a variable.
- 2 `name` is an absolute filename, an ordinary file in Octave's `path`, or (after appending '.m') a function file in Octave's `path`.
- 3 `name` is a '.oct' or '.mex' file in Octave's `path`.
- 5 `name` is a built-in function.
- 7 `name` is a directory.
- 8 `name` is a class. (Note: not currently implemented)

103 *name* is a function not associated with a file (entered on the command line).

0 *name* does not exist.

If the optional argument *type* is supplied, check only for symbols of the specified type. Valid types are

"var" Check only for variables.

"builtin" Check only for built-in functions.

"dir" Check only for directories.

"file" Check only for files and directories.

"class" Check only for classes. (Note: This option is accepted, but not currently implemented)

If no type is given, and there are multiple possible matches for *name*, **exist** will return a code according to the following priority list: variable, built-in function, oct-file, directory, file, class.

exist returns 2 if a regular file called *name* is present in Octave's search path. For information about other types of files not on the search path use some combination of the functions **file_in_path** and **stat** instead.

Programming Note: If *name* is implemented by a buggy .oct/.mex file, calling **exist** may cause Octave to crash. To maintain high performance, Octave trusts .oct/.mex files instead of sandboxing them.

See also: [\[file_in_loadpath\], page 198](#), [\[file_in_path\], page 904](#), [\[dir_in_loadpath\], page 199](#), [\[stat\], page 900](#).

Usually Octave will manage the memory, but sometimes it can be practical to remove variables from memory manually. This is usually needed when working with large variables that fill a substantial part of the memory. On a computer that uses the IEEE floating point format, the following program allocates a matrix that requires around 128 MB memory.

```
large_matrix = zeros (4000, 4000);
```

Since having this variable in memory might slow down other computations, it can be necessary to remove it manually from memory. The **clear** or **clearvars** functions do this.

```
clear
clear pattern ...
clear options pattern ...
```

Delete the names matching the given patterns from the symbol table.

The *pattern* may contain the following special characters:

? Match any single character.

***** Match zero or more characters.

[list] Match the list of characters specified by *list*. If the first character is ! or ^, match all characters except those specified by *list*. For example, the pattern [a-zA-Z] will match all lowercase and uppercase alphabetic characters.

For example, the command

```
clear foo b*r
```

clears the name `foo` and all names that begin with the letter ‘`b`’ and end with the letter ‘`r`’.

If `clear` is called without any arguments, all user-defined variables are cleared from the current workspace (i.e., local variables). Any global variables present will no longer be visible in the current workspace, but they will continue to exist in the global workspace. Functions are unaffected by this form of `clear`.

The following options are available in both long and short form

`all`, `-all`, `-a`

Clear all local and global user-defined variables, and all functions from the symbol table.

`-exclusive`, `-x`

Clear variables that do `not` match the following pattern.

`functions`, `-functions`, `-f`

Clear function names from the function symbol table. Persistent variables will be re-initialized to their default value unless the function has been locked in memory with `mlock`.

`global`, `-global`, `-g`

Clear global variable names.

`variables`, `-variables`, `-v`

Clear local variable names.

`classes`, `-classes`, `-c`

Clear the class structure table and all objects.

`-regexp`, `-r`

The *pattern* arguments are treated as regular expressions and any matches will be cleared.

With the exception of `-exclusive` and `-regexp`, all long options can be used without the dash as well. Note that, aside from `-exclusive`, only one other option may appear. All options must appear before any patterns.

Programming Note: The command `clear name` only clears the variable `name` when both a variable and a (shadowed) function named `name` are currently defined. For example, suppose you have defined a function `foo`, and then hidden it by performing the assignment `foo = 2`. Executing the command `clear foo` once will clear the variable definition and restore the definition of `foo` as a function. Executing `clear foo` a second time will clear the function definition.

See also: `[who]`, page 132, `[whos]`, page 132, `[exist]`, page 134, `[mlock]`, page 204.

```
clearvars
clearvars pattern ...
clearvars -regexp pattern ...
clearvars ... -except pattern ...
```

```
clearvars ... -except -regexp pattern ...
clearvars -global ...
```

Delete the variables matching the given *patterns* from memory.

The *pattern* may contain the following special characters:

- ? Match any single character.
- * Match zero or more characters.
- [*list*] Match the list of characters specified by *list*. If the first character is ! or ^, match all characters except those specified by *list*. For example, the pattern [a-zA-Z] will match all lowercase and uppercase alphabetic characters.

If the **-regexp** option is given then subsequent patterns are treated as regular expressions and any matches will be cleared.

If the **-except** option is given then subsequent patterns select variables that will **not** be cleared.

If the **-global** option is given then all patterns will be applied to global variables rather than local variables.

When called with no arguments, **clearvars** deletes all local variables.

Example Code:

Clear all variables starting with 'x' and the specific variable "foobar"

```
clearvars x* foobar
```

Clear the specific variable "foobar" and use regular expressions to clear all variables starting with 'x' or 'y'.

```
clearvars foobar -regexp ^x ^y
```

Clear all variables except for "foobar"

```
clearvars -except foobar
```

Clear all variables beginning with "foo", except for those ending in "bar"

```
clearvars foo* -except -regexp bar$
```

See also: [clear], page 135, [who], page 132, [whos], page 132, [exist], page 134.

pack ()

Consolidate workspace memory in MATLAB.

This function is provided for compatibility, but does nothing in Octave.

See also: [clear], page 135.

Information about a function or variable such as its location in the file system can also be acquired from within Octave. This is usually only useful during development of programs, and not within a program.

```
type name ...
type -q name ...
text = type ("name", ...)
```

Display the contents of *name* which may be a file, function (m-file), variable, operator, or keyword.

type normally prepends a header line describing the category of *name* such as function or variable; The **-q** option suppresses this behavior.

If no output variable is used the contents are displayed on screen. Otherwise, a cell array of strings is returned, where each element corresponds to the contents of each requested function.

which *name* . . .

Display the type of each *name*.

If *name* is defined from a function file, the full name of the file is also displayed.

See also: [help], page 20, [lookfor], page 21.

what

what *dir*

w = **what** (*dir*)

List the Octave specific files in directory *dir*.

If *dir* is not specified then the current directory is used.

If a return argument is requested, the files found are returned in the structure *w*. The structure contains the following fields:

path Full path to directory *dir*

m Cell array of m-files

mat Cell array of mat files

mex Cell array of mex files

oct Cell array of oct files

mdl Cell array of mdl files

slx Cell array of slx files

p Cell array of p-files

classes Cell array of class directories (@classname/)

packages Cell array of package directories (+pkgname/)

Compatibility Note: Octave does not support mdl, slx, and p files; nor does it support package directories. **what** will always return an empty list for these categories.

See also: [which], page 138, [ls], page 923, [exist], page 134.

8 Expressions

Expressions are the basic building block of statements in Octave. An expression evaluates to a value, which you can print, test, store in a variable, pass to a function, or assign a new value to a variable with an assignment operator.

An expression can serve as a statement on its own. Most other kinds of statements contain one or more expressions which specify data to be operated on. As in other languages, expressions in Octave include variables, array references, constants, and function calls, as well as combinations of these with various operators.

8.1 Index Expressions

An *index expression* allows you to reference or extract selected elements of a vector, a matrix (2-D), or a higher-dimensional array.

Indices may be scalars, vectors, ranges, or the special operator ‘`:`’, which selects entire rows, columns, or higher-dimensional slices.

An index expression consists of a set of parentheses enclosing M expressions separated by commas. Each individual index value, or component, is used for the respective dimension of the object that it is applied to. In other words, the first index component is used for the first dimension (rows) of the object, the second index component is used for the second dimension (columns) of the object, and so on. The number of index components M defines the dimensionality of the index expression. An index with two components would be referred to as a 2-D index because it has two dimensions.

In the simplest case, 1) all components are scalars, and 2) the dimensionality of the index expression M is equal to the dimensionality of the object it is applied to. For example:

```
A = reshape (1:8, 2, 2, 2) # Create 3-D array
A =
ans(:,:,1) =
1   3
2   4
ans(:,:,2) =
5   7
6   8
A(2, 1, 2) # second row, first column of second slice
# in third dimension: ans = 6
```

The size of the returned object in a specific dimension is equal to the number of elements in the corresponding component of the index expression. When all components are scalars, the result is a single output value. However, if any component is a vector or range then the returned values are the Cartesian product of the indices in the respective dimensions. For example:

```
A([1, 2], 1, 2) ≡ [A(1,1,2); A(2,1,2)]
⇒
ans =
5
6
```

The total number of returned values is the product of the number of elements returned for each index component. In the example above, the total is $2 \times 1 \times 1 = 2$ elements.

Notice that the size of the returned object in a given dimension is equal to the number of elements in the index expression for that dimension. In the code above, the first index component ([1, 2]) was specified as a row vector, but its shape is unimportant. The important fact is that the component specified two values, and hence the result must have a size of two in the first dimension; and because the first dimension corresponds to rows, the overall result is a column vector.

```
A(1, [2, 1, 1], 1) # result is a row vector: ans = [3, 1, 1]
A(ones (2, 2), 1, 1) # result is a column vector: ans = [1; 1; 1; 1]
```

The first line demonstrates again that the size of the output in a given dimension is equal to the number of elements in the respective indexing component. In this case, the output has three elements in the second dimension (which corresponds to columns), so the result is a row vector. The example also shows how repeating entries in the index expression can be used to replicate elements in the output. The last example further proves that the shape of the indexing component is irrelevant, it is only the number of elements ($2 \times 2 = 4$) which is important.

The above rules apply whenever the dimensionality of the index expression is greater than one ($M > 1$). However, for one-dimensional index expressions special rules apply and the shape of the output is determined by the shape of the indexing component. For example:

```
A([1, 2]) # result is a row vector: ans = [1, 2]
A([1; 2]) # result is a column vector: ans = [1; 2]
```

Note that it is permissible to use a 1-D index with a multi-dimensional object (also called linear indexing). In this case, the elements of the multi-dimensional array are taken in column-first order like Fortran. That is, the columns of the array are imagined to be stacked on top of each other to form a column vector and then the single linear index is applied to this vector.

```
A(5) # linear indexing into three-dimensional array: ans = 5
A(3:5) # result has shape of index component: ans = [3, 4, 5]
```

A colon (‘:’) may be used as an index component to select all of the elements in a specified dimension. Given the matrix,

```
A = [1, 2; 3, 4]
```

all of the following expressions are equivalent and select the first row of the matrix.

```
A(1, [1, 2]) # row 1, columns 1 and 2
A(1, 1:2) # row 1, columns in range 1-2
A(1, :) # row 1, all columns
```

When a colon is used in the special case of 1-D indexing the result is always a column vector. Creating column vectors with a colon index is a very frequently encountered code idiom and is faster and generally clearer than calling `reshape` for this case.

```
A(:)      # result is column vector: ans = [1; 2; 3; 4]
A(:)'    # result is row vector: ans = [1, 2, 3, 4]
```

In index expressions the keyword `end` automatically refers to the last entry for a particular dimension. This magic index can also be used in ranges and typically eliminates the need to call `size` or `length` to gather array bounds before indexing. For example:

```
A(1:end/2)        # first half of A => [1, 2]
A(end + 1) = 5;  # append element
A(end) = [];     # delete element
A(1:2:end)       # odd elements of A => [1, 3]
A(2:2:end)       # even elements of A => [2, 4]
A(end:-1:1)      # reversal of A => [4, 3, 2, 1]
```

8.1.1 Advanced Indexing

When it is necessary to extract subsets of entries out of an array whose indices cannot be written as a Cartesian product of components, linear indexing together with the function `sub2ind` can be used. For example:

```
A = reshape (1:8, 2, 2, 2)  # Create 3-D array
A =
ans(:,:,:,1) =
1   3
2   4
ans(:,:,:2) =
5   7
6   8
A(sub2ind (size (A), [1, 2, 1], [1, 1, 2], [1, 2, 1]))
⇒ ans = [A(1, 1, 1), A(2, 1, 2), A(1, 2, 1)]
```

An array with ‘nd’ dimensions can be indexed by an index expression which has from 1 to ‘nd’ components. For the ordinary and most common case, the number of components ‘M’ matches the number of dimensions ‘nd’. In this case the ordinary indexing rules apply and each component corresponds to the respective dimension of the array.

However, if the number of indexing components exceeds the number of dimensions ($M > nd$) then the excess components must all be singletons (1). Moreover, if $M < nd$, the behavior is equivalent to reshaping the input object so as to merge the trailing $nd - M$ dimensions into the last index dimension M. Thus, the result will have the dimensionality of the index expression, and not the original object. This is the case whenever dimensionality of the index is greater than one ($M > 1$), so that the special rules for linear indexing are not applied. This is easiest to understand with an example:

```
A = reshape (1:8, 2, 2, 2)  # Create 3-D array
A =
```

```

ans(:,:,1) =
1   3
2   4

ans(:,:,2) =
5   7
6   8

## 2-D indexing causes third dimension to be merged into second dimension.■
## Equivalent array for indexing, Atmp, is now 2x4.
Atmp = reshape (A, 2, 4)
Atmp =
1   3   5   7
2   4   6   8

A(2,1) # Reshape to 2x4 matrix, second entry of first column: ans = 2
A(2,4) # Reshape to 2x4 matrix, second entry of fourth column: ans = 8
A(:,:,1) # Reshape to 2x4 matrix, select all rows & columns, ans = Atmp

```

Note here the elegant use of the double colon to replace the call to the `reshape` function.

Another advanced use of linear indexing is to create arrays filled with a single value. This can be done by using an index of ones on a scalar value. The result is an object with the dimensions of the index expression and every element equal to the original scalar. For example, the following statements

```
a = 13;
a(ones (1, 4))
```

produce a row vector whose four elements are all equal to 13.

Similarly, by indexing a scalar with two vectors of ones it is possible to create a matrix. The following statements

```
a = 13;
a(ones (1, 2), ones (1, 3))
```

create a 2x3 matrix with all elements equal to 13. This could also have been written as

```
13(ones (2, 3))
```

It is more efficient to use indexing rather than the code construction `scalar * ones (M, N, ...)` because it avoids the unnecessary multiplication operation. Moreover, multiplication may not be defined for the object to be replicated whereas indexing an array is always defined. The following code shows how to create a 2x3 cell array from a base unit which is not itself a scalar.

```
{"Hello"}(ones (2, 3))
```

It should be noted that `ones (1, n)` (a row vector of ones) results in a range object (with zero increment). A range is stored internally as a starting value, increment, end value, and

total number of values; hence, it is more efficient for storage than a vector or matrix of ones whenever the number of elements is greater than 4. In particular, when ‘*r*’ is a row vector, the expressions

```
r(ones (1, n), :)
r(ones (n, 1), :)
```

will produce identical results, but the first one will be significantly faster, at least for ‘*r*’ and ‘*n*’ large enough. In the first case the index is held in compressed form as a range which allows Octave to choose a more efficient algorithm to handle the expression.

A general recommendation for users unfamiliar with these techniques is to use the function `repmat` for replicating smaller arrays into bigger ones, which uses such tricks.

A second use of indexing is to speed up code. Indexing is a fast operation and judicious use of it can reduce the requirement for looping over individual array elements, which is a slow operation.

Consider the following example which creates a 10-element row vector *a* containing the values $a_i = \sqrt{i}$.

```
for i = 1:10
    a(i) = sqrt (i);
endfor
```

It is quite inefficient to create a vector using a loop like this. In this case, it would have been much more efficient to use the expression

```
a = sqrt (1:10);
```

which avoids the loop entirely.

In cases where a loop cannot be avoided, or a number of values must be combined to form a larger matrix, it is generally faster to set the size of the matrix first (pre-allocate storage), and then insert elements using indexing commands. For example, given a matrix *a*,

```
[nr, nc] = size (a);
x = zeros (nr, n * nc);
for i = 1:n
    x(:,(i-1)*nc+1:i*nc) = a;
endfor
```

is considerably faster than

```
x = a;
for i = 1:n-1
    x = [x, a];
endfor
```

because Octave does not have to repeatedly resize the intermediate result.

```
ind = sub2ind (dims, i, j)
ind = sub2ind (dims, s1, s2, ..., sN)
```

Convert subscripts to linear indices.

The input *dims* is a dimension vector where each element is the size of the array in the respective dimension (see [\[size\]](#), page 45). The remaining inputs are scalars or vectors of subscripts to be converted.

The output vector *ind* contains the converted linear indices.

Background: Array elements can be specified either by a linear index which starts at 1 and runs through the number of elements in the array, or they may be specified with subscripts for the row, column, page, etc. The functions **ind2sub** and **sub2ind** interconvert between the two forms.

The linear index traverses dimension 1 (rows), then dimension 2 (columns), then dimension 3 (pages), etc. until it has numbered all of the elements. Consider the following 3-by-3 matrices:

```
[(1,1), (1,2), (1,3)]      [1, 4, 7]
[(2,1), (2,2), (2,3)] ==> [2, 5, 8]
[(3,1), (3,2), (3,3)]      [3, 6, 9]
```

The left matrix contains the subscript tuples for each matrix element. The right matrix shows the linear indices for the same matrix.

The following example shows how to convert the two-dimensional indices (2,1) and (2,3) of a 3-by-3 matrix to linear indices with a single call to **sub2ind**.

```
s1 = [2, 2];
s2 = [1, 3];
ind = sub2ind ([3, 3], s1, s2)
⇒ ind = 2 8
```

See also: [\[ind2sub\]](#), page 144, [\[size\]](#), page 45.

[s1, s2, ..., sN] = ind2sub (dims, ind)

Convert linear indices to subscripts.

The input *dims* is a dimension vector where each element is the size of the array in the respective dimension (see [\[size\]](#), page 45). The second input *ind* contains linear indices to be converted.

The outputs *s1, ..., sN* contain the converted subscripts.

Background: Array elements can be specified either by a linear index which starts at 1 and runs through the number of elements in the array, or they may be specified with subscripts for the row, column, page, etc. The functions **ind2sub** and **sub2ind** interconvert between the two forms.

The linear index traverses dimension 1 (rows), then dimension 2 (columns), then dimension 3 (pages), etc. until it has numbered all of the elements. Consider the following 3-by-3 matrices:

```
[1, 4, 7]      [(1,1), (1,2), (1,3)]
[2, 5, 8] ==> [(2,1), (2,2), (2,3)]
[3, 6, 9]      [(3,1), (3,2), (3,3)]
```

The left matrix contains the linear indices for each matrix element. The right matrix shows the subscript tuples for the same matrix.

The following example shows how to convert the two-dimensional indices (2,1) and (2,3) of a 3-by-3 matrix to linear indices with a single call to **sub2ind**.

The following example shows how to convert the linear indices 2 and 8 in a 3-by-3 matrix into subscripts.

```
ind = [2, 8];
[r, c] = ind2sub ([3, 3], ind)
⇒ r = 2 2
⇒ c = 1 3
```

If the number of output subscripts exceeds the number of dimensions, the exceeded dimensions are set to 1. On the other hand, if fewer subscripts than dimensions are provided, the exceeding dimensions are merged into the final requested dimension. For clarity, consider the following examples:

```
ind = [2, 8];
dims = [3, 3];
## same as dims = [3, 3, 1]
[r, c, s] = ind2sub (dims, ind)
⇒ r = 2 2
⇒ c = 1 3
⇒ s = 1 1
## same as dims = [9]
r = ind2sub (dims, ind)
⇒ r = 2 8
```

See also: [\[ind2sub\]](#), page 144, [\[size\]](#), page 45.

isindex (ind)
isindex (ind, n)

Return true if *ind* is a valid index.

Valid indices are either positive integers (although possibly of real data type), or logical arrays.

If present, *n* specifies the maximum extent of the dimension to be indexed. When possible the internal result is cached so that subsequent indexing using *ind* will not perform the check again.

Implementation Note: Strings are first converted to double values before the checks for valid indices are made. Unless a string contains the NULL character "\0", it will always be a valid index.

8.2 Calling Functions

A *function* is a name for a particular calculation. Because it has a name, you can ask for it by name at any point in the program. For example, the function **sqrt** computes the square root of a number.

A fixed set of functions are *built-in*, which means they are available in every Octave program. The **sqrt** function is one of these. In addition, you can define your own functions. See [Chapter 11 \[Functions and Scripts\]](#), page 177, for information about how to do this.

The way to use a function is with a *function call* expression, which consists of the function name followed by a list of *arguments* in parentheses. The arguments are expressions which give the raw materials for the calculation that the function will do. When there is more than one argument, they are separated by commas. If there are no arguments, you can

omit the parentheses, but it is a good idea to include them anyway, to clearly indicate that a function call was intended. Here are some examples:

```
sqrt (x^2 + y^2)      # One argument
ones (n, m)           # Two arguments
rand ()                # No arguments
```

Each function expects a particular number of arguments. For example, the `sqrt` function must be called with a single argument, the number to take the square root of:

`sqrt (argument)`

Some of the built-in functions take a variable number of arguments, depending on the particular usage, and their behavior is different depending on the number of arguments supplied.

Like every other expression, the function call has a value, which is computed by the function based on the arguments you give it. In this example, the value of `sqrt (argument)` is the square root of the argument. A function can also have side effects, such as assigning the values of certain variables or doing input or output operations.

Unlike most languages, functions in Octave may return multiple values. For example, the following statement

```
[u, s, v] = svd (a)
```

computes the singular value decomposition of the matrix `a` and assigns the three result matrices to `u`, `s`, and `v`.

The left side of a multiple assignment expression is itself a list of expressions, that is, a list of variable names potentially qualified by index expressions. See also [Section 8.1 \[Index Expressions\]](#), page 139, and [Section 8.6 \[Assignment Ops\]](#), page 156.

8.2.1 Call by Value

In Octave, unlike Fortran, function arguments are passed by value, which means that each argument in a function call is evaluated and assigned to a temporary location in memory before being passed to the function. There is currently no way to specify that a function parameter should be passed by reference instead of by value. This means that it is impossible to directly alter the value of a function parameter in the calling function. It can only change the local copy within the function body. For example, the function

```
function f (x, n)
  while (n-- > 0)
    disp (x);
  endwhile
endfunction
```

displays the value of the first argument `n` times. In this function, the variable `n` is used as a temporary variable without having to worry that its value might also change in the calling function. Call by value is also useful because it is always possible to pass constants for any function parameter without first having to determine that the function will not attempt to modify the parameter.

The caller may use a variable as the expression for the argument, but the called function does not know this: it only knows what value the argument had. For example, given a function called as

```
foo = "bar";
fcn (foo)
```

you should not think of the argument as being “the variable `foo`.” Instead, think of the argument as the string value, `"bar"`.

Even though Octave uses pass-by-value semantics for function arguments, values are not copied unnecessarily. For example,

```
x = rand (1000);
f (x);
```

does not actually force two 1000 by 1000 element matrices to exist *unless* the function `f` modifies the value of its argument. Then Octave must create a copy to avoid changing the value outside the scope of the function `f`, or attempting (and probably failing!) to modify the value of a constant or the value of a temporary result.

8.2.2 Recursion

With some restrictions¹, recursive function calls are allowed. A *recursive function* is one which calls itself, either directly or indirectly. For example, here is an inefficient² way to compute the factorial of a given integer:

```
function retval = fact (n)
  if (n > 0)
    retval = n * fact (n-1);
  else
    retval = 1;
  endif
endfunction
```

This function is recursive because it calls itself directly. It eventually terminates because each time it calls itself, it uses an argument that is one less than was used for the previous call. Once the argument is no longer greater than zero, it does not call itself, and the recursion ends.

The built-in variable `max_recursion_depth` specifies a limit to the recursion depth and prevents Octave from recursing infinitely. Similarly, the variable `max_stack_depth` specifies a limit to the depth of function calls, whether recursive or not. These limits help prevent stack overflow on the computer Octave is running on, so that instead of exiting with a signal, the interpreter will throw an error and return to the command prompt.

```
val = max_recursion_depth ()
old_val = max_recursion_depth (new_val)
max_recursion_depth (new_val, "local")
```

Query or set the internal limit on the number of times a function may be called recursively.

¹ Some of Octave’s functions are implemented in terms of functions that cannot be called recursively. For example, the ODE solver `lsode` is ultimately implemented in a Fortran subroutine that cannot be called recursively, so `lsode` should not be called either directly or indirectly from within the user-supplied function that `lsode` requires. Doing so will result in an error.

² It would be much better to use `prod (1:n)`, or `gamma (n+1)` instead, after first checking to ensure that the value `n` is actually a positive integer.

If the limit is exceeded, an error message is printed and control returns to the top level.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[max_stack_depth\]](#), page 148.

```
val = max_stack_depth ()
old_val = max_stack_depth (new_val)
max_stack_depth (new_val, "local")
```

Query or set the internal limit on the number of times a function may be called recursively.

If the limit is exceeded, an error message is printed and control returns to the top level.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[max_recursion_depth\]](#), page 147.

8.2.3 Access via Handle

A function may be abstracted and referenced via a function handle acquired using the special operator '@'. For example,

```
f = @plus;
f (2, 2)
⇒ 4
```

is equivalent to calling `plus (2, 2)` directly. Beyond abstraction for general programming, function handles find use in callback methods for figures and graphics by adding listeners to properties or assigning pre-existing actions, such as in the following example:

```
function mydeletefcn (h, ~, msg)
    printf (msg);
endfunction
sombrero;
set (gcf, "deletefcn", {@mydeletefcn, "Bye!\n"});
close;
```

The above will print "Bye!" to the terminal upon the closing (deleting) of the figure. There are many graphics property actions for which a callback function may be assigned, including, `buttondownfcn`, `windowscrollwheelfcn`, `createfcn`, `deletefcn`, `keypressfcn`, etc.

Note that the '@' character also plays a role in defining class functions, i.e., methods, but not as a syntactical element. Rather it begins a directory name containing methods for a class that shares the directory name sans the '@' character. See [Chapter 34 \[Object Oriented Programming\]](#), page 835.

8.3 Arithmetic Operators

The following arithmetic operators are available, and work on scalars and matrices. The element-by-element operators and functions broadcast (see [Section 19.2 \[Broadcasting\], page 585](#)).

$x + y$	Addition. If both operands are matrices, the number of rows and columns must both agree, or they must be broadcastable to the same shape.
$x .+ y$	Element-by-element addition. This operator is equivalent to <code>+</code> .
$x - y$	Subtraction. If both operands are matrices, the number of rows and columns of both must agree, or they must be broadcastable to the same shape.
$x .- y$	Element-by-element subtraction. This operator is equivalent to <code>-</code> .
$x * y$	Matrix multiplication. The number of columns of x must agree with the number of rows of y .
$x .* y$	Element-by-element multiplication. If both operands are matrices, the number of rows and columns must both agree, or they must be broadcastable to the same shape.
x / y	Right division. This is conceptually equivalent to the expression $\text{inverse}(y') * x'$ but it is computed without forming the inverse of y' . If the system is not square, or if the coefficient matrix is singular, a minimum norm solution is computed.
$x ./ y$	Element-by-element right division.
$x \setminus y$	Left division. This is conceptually equivalent to the expression $\text{inverse}(x) * y$ but it is computed without forming the inverse of x . If the system is not square, or if the coefficient matrix is singular, a minimum norm solution is computed.
$x .\setminus y$	Element-by-element left division. Each element of y is divided by each corresponding element of x .
$x ^ y$ $x ** y$	Power operator. If x and y are both scalars, this operator returns x raised to the power y . If x is a scalar and y is a square matrix, the result is computed using an eigenvalue expansion. If x is a square matrix, the result is computed by repeated multiplication if y is an integer, and by an eigenvalue expansion if y is not an integer. An error results if both x and y are matrices. The implementation of this operator needs to be improved.
$x .^ y$ $x .** y$	Element-by-element power operator. If both operands are matrices, the number of rows and columns must both agree, or they must be broadcastable to the same shape. If several complex results are possible, the one with smallest non-negative argument (angle) is taken. This rule may return a complex root even

when a real root is also possible. Use `realpow`, `realsqrt`, `cbrt`, or `nthroot` if a real result is preferred.

<code>-x</code>	Negation.
<code>+x</code>	Unary plus. This operator has no effect on the operand.
<code>x'</code>	Complex conjugate transpose. For real arguments, this operator is the same as the transpose operator. For complex arguments, this operator is equivalent to the expression <code>conj (x.'</code>)
<code>x.'</code>	Transpose.

Note that because Octave's element-by-element operators begin with a '.', there is a possible ambiguity for statements like

`1./m`

because the period could be interpreted either as part of the constant or as part of the operator. To resolve this conflict, Octave treats the expression as if you had typed

(1) `./ m`

and not

(1.) `/ m`

Although this is inconsistent with the normal behavior of Octave's lexer, which usually prefers to break the input into tokens by preferring the longest possible match at any given point, it is more useful in this case.

`ctranspose (x)`

Return the complex conjugate transpose of `x`.

This function and `x'` are equivalent.

See also: [\[transpose\]](#), page 152.

`ldivide (x, y)`

Return the element-by-element left division of `x` and `y`.

This function and `x .\ y` are equivalent.

See also: [\[rdivide\]](#), page 151, [\[mldivide\]](#), page 150, [\[times\]](#), page 151, [\[plus\]](#), page 151.

`minus (x, y)`

This function and `x - y` are equivalent.

See also: [\[plus\]](#), page 151, [\[uminus\]](#), page 152.

`mldivide (x, y)`

Return the matrix left division of `x` and `y`.

This function and `x \ y` are equivalent.

See also: [\[mrdivide\]](#), page 151, [\[ldivide\]](#), page 150, [\[rdivide\]](#), page 151.

`mpower (x, y)`

Return the matrix power operation of `x` raised to the `y` power.

This function and `x ^ y` are equivalent.

See also: [\[power\]](#), page 151, [\[mtimes\]](#), page 151, [\[plus\]](#), page 151, [\[minus\]](#), page 150.

mrdivide (x, y)

Return the matrix right division of x and y.

This function and x / y are equivalent.

See also: [mldivide], page 150, [rdivide], page 151, [plus], page 151, [minus], page 150.

mtimes (x, y)**mtimes (x1, x2, ...)**

Return the matrix multiplication product of inputs.

This function and $x * y$ are equivalent. If more arguments are given, the multiplication is applied cumulatively from left to right:

$(\dots((x_1 * x_2) * x_3) * \dots)$

See also: [times], page 151, [plus], page 151, [minus], page 150, [rdivide], page 151, [mrdivide], page 151, [mldivide], page 150, [mpower], page 150.

plus (x, y)**plus (x1, x2, ...)**

This function and $x + y$ are equivalent.

If more arguments are given, the summation is applied cumulatively from left to right:

$(\dots((x_1 + x_2) + x_3) + \dots)$

See also: [minus], page 150, [uplus], page 152.

power (x, y)

Return the element-by-element operation of x raised to the y power.

This function and $x .^ y$ are equivalent.

If several complex results are possible, returns the one with smallest non-negative argument (angle). Use **realpow**, **realsqrt**, **cbrt**, or **nthroot** if a real result is preferred.

See also: [mpower], page 150, [realpow], page 510, [realsqrt], page 510, [cbrt], page 510, [nthroot], page 510.

rdivide (x, y)

Return the element-by-element right division of x and y.

This function and $x ./ y$ are equivalent.

See also: [ldivide], page 150, [mrdivide], page 151, [times], page 151, [plus], page 151.

times (x, y)**times (x1, x2, ...)**

Return the element-by-element multiplication product of inputs.

This function and $x .* y$ are equivalent. If more arguments are given, the multiplication is applied cumulatively from left to right:

$(\dots((x_1 .* x_2) .* x_3) .* \dots)$

See also: [mtimes], page 151, [rdivide], page 151.

transpose (x)

Return the transpose of x.

This function and `x.'` are equivalent.

See also: [ctranspose], page 150.

uminus (x)

This function and `-x` are equivalent.

See also: [uplus], page 152, [minus], page 150.

uplus (x)

This function and `+x` are equivalent.

See also: [uminus], page 152, [plus], page 151, [minus], page 150.

8.4 Comparison Operators

Comparison operators compare numeric values for relationships such as equality. They are written using *relational operators*.

All of Octave's comparison operators return a value of 1 if the comparison is true, or 0 if it is false. For matrix values, they all work on an element-by-element basis. Broadcasting rules apply. See Section 19.2 [Broadcasting], page 585. For example:

```
[1, 2; 3, 4] == [1, 3; 2, 4]
⇒ 1 0
    0 1
```

According to broadcasting rules, if one operand is a scalar and the other is a matrix, the scalar is compared to each element of the matrix in turn, and the result is the same size as the matrix.

<code>x < y</code>	True if x is less than y.
<code>x <= y</code>	True if x is less than or equal to y.
<code>x == y</code>	True if x is equal to y.
<code>x >= y</code>	True if x is greater than or equal to y.
<code>x > y</code>	True if x is greater than y.
<code>x != y</code>	
<code>x ~= y</code>	True if x is not equal to y.

For complex numbers, the following ordering is defined: $z1 < z2$ if and only if

```
abs (z1) < abs (z2)
|| (abs (z1) == abs (z2) && arg (z1) < arg (z2))
```

This is consistent with the ordering used by `max`, `min` and `sort`, but is not consistent with MATLAB, which only compares the real parts.

String comparisons may also be performed with the `strcmp` function, not with the comparison operators listed above. See Chapter 5 [Strings], page 67.

eq (x, y)

Return true if the two inputs are equal.

This function is equivalent to `x == y`.

See also: [\[ne\]](#), page 153, [\[isequal\]](#), page 153, [\[le\]](#), page 153, [\[ge\]](#), page 153, [\[gt\]](#), page 153, [\[ne\]](#), page 153, [\[lt\]](#), page 153.

ge (x, y)

This function is equivalent to `x >= y`.

See also: [\[le\]](#), page 153, [\[eq\]](#), page 152, [\[gt\]](#), page 153, [\[ne\]](#), page 153, [\[lt\]](#), page 153.

gt (x, y)

This function is equivalent to `x > y`.

See also: [\[le\]](#), page 153, [\[eq\]](#), page 152, [\[ge\]](#), page 153, [\[ne\]](#), page 153, [\[lt\]](#), page 153.

isequal (x1, x2, ...)

Return true if all of `x1, x2, ...` are equal.

See also: [\[isequaln\]](#), page 153.

isequaln (x1, x2, ...)

Return true if all of `x1, x2, ...` are equal under the additional assumption that `NaN == NaN` (no comparison of `NaN` placeholders in dataset).

See also: [\[isequal\]](#), page 153.

le (x, y)

This function is equivalent to `x <= y`.

See also: [\[eq\]](#), page 152, [\[ge\]](#), page 153, [\[gt\]](#), page 153, [\[ne\]](#), page 153, [\[lt\]](#), page 153.

lt (x, y)

This function is equivalent to `x < y`.

See also: [\[le\]](#), page 153, [\[eq\]](#), page 152, [\[ge\]](#), page 153, [\[gt\]](#), page 153, [\[ne\]](#), page 153.

ne (x, y)

Return true if the two inputs are not equal.

This function is equivalent to `x != y`.

See also: [\[eq\]](#), page 152, [\[isequal\]](#), page 153, [\[le\]](#), page 153, [\[ge\]](#), page 153, [\[gt\]](#), page 153, [\[ne\]](#), page 153.

8.5 Boolean Expressions

8.5.1 Element-by-element Boolean Operators

An *element-by-element boolean expression* is a combination of comparison expressions using the boolean operators “or” (`|`), “and” (`&`), and “not” (`!`), along with parentheses to control nesting. The truth of the boolean expression is computed by combining the truth values of the corresponding elements of the component expressions. A value is considered to be false if it is zero, and true otherwise.

Element-by-element boolean expressions can be used wherever comparison expressions can be used. They can be used in `if` and `while` statements. However, a matrix value used as the condition in an `if` or `while` statement is only true if *all* of its elements are nonzero.

Like comparison operations, each element of an element-by-element boolean expression also has a numeric value (1 if true, 0 if false) that comes into play if the result of the boolean expression is stored in a variable, or used in arithmetic.

Here are descriptions of the three element-by-element boolean operators.

`boolean1 & boolean2`

Elements of the result are true if both corresponding elements of `boolean1` and `boolean2` are true.

`boolean1 | boolean2`

Elements of the result are true if either of the corresponding elements of `boolean1` or `boolean2` is true.

`! boolean`

`~ boolean` Each element of the result is true if the corresponding element of `boolean` is false.

These operators work on an element-by-element basis. For example, the expression

`[1, 0; 0, 1] & [1, 0; 2, 3]`

returns a two by two identity matrix.

For the binary operators, broadcasting rules apply. See [Section 19.2 \[Broadcasting\]](#), [page 585](#). In particular, if one of the operands is a scalar and the other a matrix, the operator is applied to the scalar and each element of the matrix.

For the binary element-by-element boolean operators, both subexpressions `boolean1` and `boolean2` are evaluated before computing the result. This can make a difference when the expressions have side effects. For example, in the expression

`a & b++`

the value of the variable `b` is incremented even if the variable `a` is zero.

This behavior is necessary for the boolean operators to work as described for matrix-valued operands.

`z = and (x, y)`

`z = and (x1, x2, ...)`

Return the logical AND of `x` and `y`.

This function is equivalent to the operator syntax `x & y`. If more than two arguments are given, the logical AND is applied cumulatively from left to right:

`(...((x1 & x2) & x3) & ...)`

See also: [\[or\]](#), [page 155](#), [\[not\]](#), [page 154](#), [\[xor\]](#), [page 475](#).

`z = not (x)`

Return the logical NOT of `x`.

This function is equivalent to the operator syntax `! x`.

See also: [\[and\]](#), [page 154](#), [\[or\]](#), [page 155](#), [\[xor\]](#), [page 475](#).

```
z = or (x, y)
z = or (x1, x2, ...)
```

Return the logical OR of x and y .

This function is equivalent to the operator syntax $x \mid y$. If more than two arguments are given, the logical OR is applied cumulatively from left to right:

```
((x1 \mid x2) \mid x3) \mid ...)
```

See also: [and], page 154, [not], page 154, [xor], page 475.

8.5.2 Short-circuit Boolean Operators

Combined with the implicit conversion to scalar values in `if` and `while` conditions, Octave's element-by-element boolean operators are often sufficient for performing most logical operations. However, it is sometimes desirable to stop evaluating a boolean expression as soon as the overall truth value can be determined. Octave's *short-circuit* boolean operators work this way.

`boolean1 && boolean2`

The expression `boolean1` is evaluated and converted to a scalar using the equivalent of the operation `all (boolean1(:))`. If it is false, the result of the overall expression is 0. If it is true, the expression `boolean2` is evaluated and converted to a scalar using the equivalent of the operation `all (boolean2(:))`. If it is true, the result of the overall expression is 1. Otherwise, the result of the overall expression is 0.

Warning: there is one exception to the rule of evaluating `all (boolean1(:))`, which is when `boolean1` is the empty matrix. The truth value of an empty matrix is always `false` so `[] && true` evaluates to `false` even though `all ([])` is `true`.

`boolean1 || boolean2`

The expression `boolean1` is evaluated and converted to a scalar using the equivalent of the operation `all (boolean1(:))`. If it is true, the result of the overall expression is 1. If it is false, the expression `boolean2` is evaluated and converted to a scalar using the equivalent of the operation `all (boolean2(:))`. If it is true, the result of the overall expression is 1. Otherwise, the result of the overall expression is 0.

Warning: the truth value of an empty matrix is always `false`, see the previous list item for details.

The fact that both operands may not be evaluated before determining the overall truth value of the expression can be important. For example, in the expression

```
a && b++
```

the value of the variable b is only incremented if the variable a is nonzero.

This can be used to write somewhat more concise code. For example, it is possible write

```
function f (a, b, c)
  if (nargin > 2 && ischar (c))
    ...
  end
```

instead of having to use two `if` statements to avoid attempting to evaluate an argument that doesn't exist. For example, without the short-circuit feature, it would be necessary to write

```
function f (a, b, c)
  if (nargin > 2)
    if (ischar (c))
      ...
    ...
```

Writing

```
function f (a, b, c)
  if (nargin > 2 & ischar (c))
  ...
  ...
```

would result in an error if `f` were called with one or two arguments because Octave would be forced to try to evaluate both of the operands for the operator '`&`'.

MATLAB has special behavior that allows the operators '`&`' and '`|`' to short-circuit when used in the truth expression for `if` and `while` statements. Octave behaves the same way for compatibility, however, the use of the '`&`' and '`|`' operators in this way is strongly discouraged and a warning will be issued. Instead, you should use the '`&&`' and '`||`' operators that always have short-circuit behavior.

Finally, the ternary operator `(?:)` is not supported in Octave. If short-circuiting is not important, it can be replaced by the `ifelse` function.

```
merge (mask, tval, fval)
ifelse (mask, tval, fval)
```

Merge elements of `true_val` and `false_val`, depending on the value of `mask`.

If `mask` is a logical scalar, the other two arguments can be arbitrary values. Otherwise, `mask` must be a logical array, and `tval`, `fval` should be arrays of matching class, or cell arrays. In the scalar mask case, `tval` is returned if `mask` is true, otherwise `fval` is returned.

In the array `mask` case, both `tval` and `fval` must be either scalars or arrays with dimensions equal to `mask`. The result is constructed as follows:

```
result(mask) = tval(mask);
result(! mask) = fval(! mask);
```

`mask` can also be arbitrary numeric type, in which case it is first converted to logical.

See also: [\[logical\]](#), page 61, [\[diff\]](#), page 476.

8.6 Assignment Expressions

An `assignment` is an expression that stores a new value into a variable. For example, the following expression assigns the value 1 to the variable `z`:

```
z = 1
```

After this expression is executed, the variable `z` has the value 1. Whatever old value `z` had before the assignment is forgotten. The '`=`' sign is called an `assignment operator`.

Assignments can store string values also. For example, the following expression would store the value "this food is good" in the variable `message`:

```

thing = "food"
predicate = "good"
message = [ "this " , thing , " is " , predicate ]

```

(This also illustrates concatenation of strings.)

Most operators (addition, concatenation, and so on) have no effect except to compute a value. If you ignore the value, you might as well not use the operator. An assignment operator is different. It does produce a value, but even if you ignore the value, the assignment still makes itself felt through the alteration of the variable. We call this a *side effect*.

The left-hand operand of an assignment need not be a variable (see Chapter 7 [Variables], page 127). It can also be an element of a matrix (see Section 8.1 [Index Expressions], page 139) or a list of return values (see Section 8.2 [Calling Functions], page 145). These are all called *lvalues*, which means they can appear on the left-hand side of an assignment operator. The right-hand operand may be any expression. It produces the new value which the assignment stores in the specified variable, matrix element, or list of return values.

It is important to note that variables do *not* have permanent types. The type of a variable is simply the type of whatever value it happens to hold at the moment. In the following program fragment, the variable `foo` has a numeric value at first, and a string value later on:

```

octave:13> foo = 1
foo = 1
octave:13> foo = "bar"
foo = bar

```

When the second assignment gives `foo` a string value, the fact that it previously had a numeric value is forgotten.

Assignment of a scalar to an indexed matrix sets all of the elements that are referenced by the indices to the scalar value. For example, if `a` is a matrix with at least two columns,

```
a(:, 2) = 5
```

sets all the elements in the second column of `a` to 5.

Assigning an empty matrix ‘`[]`’ works in most cases to allow you to delete rows or columns of matrices and vectors. See Section 4.1.1 [Empty Matrices], page 51. For example, given a 4 by 5 matrix `A`, the assignment

```
A (3, :) = []
```

deletes the third row of `A`, and the assignment

```
A (:, 1:2:5) = []
```

deletes the first, third, and fifth columns.

An assignment is an expression, so it has a value. Thus, `z = 1` as an expression has the value 1. One consequence of this is that you can write multiple assignments together:

```
x = y = z = 0
```

stores the value 0 in all three variables. It does this because the value of `z = 0`, which is 0, is stored into `y`, and then the value of `y = z = 0`, which is 0, is stored into `x`.

This is also true of assignments to lists of values, so the following is a valid expression

```
[a, b, c] = [u, s, v] = svd (a)
```

that is exactly equivalent to

```
[u, s, v] = svd (a)
a = u
b = s
c = v
```

In expressions like this, the number of values in each part of the expression need not match. For example, the expression

```
[a, b] = [u, s, v] = svd (a)
```

is equivalent to

```
[u, s, v] = svd (a)
a = u
b = s
```

The number of values on the left side of the expression can, however, not exceed the number of values on the right side. For example, the following will produce an error.

```
[a, b, c, d] = [u, s, v] = svd (a);
^- error: element number 4 undefined in return list
```

The symbol `~` may be used as a placeholder in the list of lvalues, indicating that the corresponding return value should be ignored and not stored anywhere:

```
[~, s, v] = svd (a);
```

This is cleaner and more memory efficient than using a dummy variable. The `nargout` value for the right-hand side expression is not affected. If the assignment is used as an expression, the return value is a comma-separated list with the ignored values dropped.

A very common programming pattern is to increment an existing variable with a given value, like this

```
a = a + 2;
```

This can be written in a clearer and more condensed form using the `+=` operator

```
a += 2;
```

Similar operators also exist for subtraction (`-=`), multiplication (`*=`), and division (`/=`). An expression of the form

```
expr1 op= expr2
```

is evaluated as

```
expr1 = (expr1) op (expr2)
```

where `op` can be either `+`, `-`, `*`, or `/`, as long as `expr2` is a simple expression with no side effects. If `expr2` also contains an assignment operator, then this expression is evaluated as

```
temp = expr2
expr1 = (expr1) op temp
```

where `temp` is a placeholder temporary value storing the computed result of evaluating `expr2`. So, the expression

```
a *= b+1
```

is evaluated as

```
a = a * (b+1)
```

and *not*

```
a = a * b + 1
```

You can use an assignment anywhere an expression is called for. For example, it is valid to write `x != (y = 1)` to set `y` to 1 and then test whether `x` equals 1. But this style tends to make programs hard to read. Except in a one-shot program, you should rewrite it to get rid of such nesting of assignments. This is never very hard.

8.7 Increment Operators

Increment operators increase or decrease the value of a variable by 1. The operator to increment a variable is written as ‘`++`’. It may be used to increment a variable either before or after taking its value.

For example, to pre-increment the variable `x`, you would write `++x`. This would add one to `x` and then return the new value of `x` as the result of the expression. It is exactly the same as the expression `x = x + 1`.

To post-increment a variable `x`, you would write `x++`. This adds one to the variable `x`, but returns the value that `x` had prior to incrementing it. For example, if `x` is equal to 2, the result of the expression `x++` is 2, and the new value of `x` is 3.

For matrix and vector arguments, the increment and decrement operators work on each element of the operand.

Here is a list of all the increment and decrement expressions.

- `++x` This expression increments the variable `x`. The value of the expression is the *new* value of `x`. It is equivalent to the expression `x = x + 1`.
- `--x` This expression decrements the variable `x`. The value of the expression is the *new* value of `x`. It is equivalent to the expression `x = x - 1`.
- `x++` This expression causes the variable `x` to be incremented. The value of the expression is the *old* value of `x`.
- `x--` This expression causes the variable `x` to be decremented. The value of the expression is the *old* value of `x`.

8.8 Operator Precedence

Operator precedence determines how operators are grouped, when different operators appear close by in one expression. For example, ‘`*`’ has higher precedence than ‘`+`’. Thus, the expression `a + b * c` means to multiply `b` and `c`, and then add `a` to the product (i.e., `a + (b * c)`).

You can overrule the precedence of the operators by using parentheses. You can think of the precedence rules as saying where the parentheses are assumed if you do not write parentheses yourself. In fact, it is wise to use parentheses whenever you have an unusual combination of operators, because other people who read the program may not remember what the precedence is in this case. You might forget as well, and then you too could make a mistake. Explicit parentheses will help prevent any such mistake.

When operators of equal precedence are used together, the leftmost operator groups first, except for the assignment operators, which group in the opposite order. Thus, the

expression $a - b + c$ groups as $(a - b) + c$, but the expression $a = b = c$ groups as $a = (b = c)$.

The precedence of prefix unary operators is important when another operator follows the operand. For example, $-x^2$ means $-(x^2)$, because ' $-$ ' has lower precedence than ' \wedge '.

Here is a table of the operators in Octave, in order of decreasing precedence. Unless noted, all operators group left to right.

function call and array indexing, cell array indexing, and structure element indexing

`'()' '{}' '.'`

postfix increment, and postfix decrement

`'++' '--'`

These operators group right to left.

transpose and exponentiation

`'.' '.' '^' '**' '.^' '.**'`

unary plus, unary minus, prefix increment, prefix decrement, and logical "not"

`'+' '-' '++' '--' '~~' '!'`

multiply and divide

`'*' '/' '\' '.\'' '.*' './'`

add, subtract

`'+' '-'`

colon

`'.'`

relational

`'<' '<=' '<==' '>=' '>' '!='
 '~='`

element-wise "and"

`'&'`

element-wise "or"

`'|'`

logical "and"

`'&&'`

logical "or"

`'||'`

assignment

`'=' '+=' '-=' '*=' '/=' '\=' '^=' '.*='
 './='
 '.\=' '.^=' '|=' '&='`

These operators group right to left.

9 Evaluation

Normally, you evaluate expressions simply by typing them at the Octave prompt, or by asking Octave to interpret commands that you have saved in a file.

Sometimes, you may find it necessary to evaluate an expression that has been computed and stored in a string, which is exactly what the `eval` function lets you do.

`eval (try)`

`eval (try, catch)`

Parse the string *try* and evaluate it as if it were an Octave program.

If execution fails, evaluate the optional string *catch*.

The string *try* is evaluated in the current context, so any results remain available after `eval` returns.

The following example creates the variable *A* with the approximate value of 3.1416 in the current workspace.

```
eval ("A = acos(-1);");
```

If an error occurs during the evaluation of *try* then the *catch* string is evaluated, as the following example shows:

```
eval ('error ("This is a bad example");',
      'printf ("This error occurred:\n%s\n", lasterr ());
      + This error occurred:
      This is a bad example
```

Programming Note: if you are only using `eval` as an error-capturing mechanism, rather than for the execution of arbitrary code strings, Consider using try/catch blocks or unwind_protect/unwind_protect_cleanup blocks instead. These techniques have higher performance and don't introduce the security considerations that the evaluation of arbitrary code does.

See also: [\[evalin\]](#), page 164, [\[evalc\]](#), page 161, [\[assignin\]](#), page 164, [\[feval\]](#), page 162.

The `evalc` function additionally captures any console output produced by the evaluated expression.

`s = evalc (try)`

`s = evalc (try, catch)`

Parse and evaluate the string *try* as if it were an Octave program, while capturing the output into the return variable *s*.

If execution fails, evaluate the optional string *catch*.

This function behaves like `eval`, but any output or warning messages which would normally be written to the console are captured and returned in the string *s*.

The `diary` is disabled during the execution of this function. When `system` is used, any output produced by external programs is *not* captured, unless their output is captured by the `system` function itself.

```
s = evalc ("t = 42"), t
```

```
⇒ s = t = 42
```

```
⇒ t = 42
```

See also: [eval], page 161, [diary], page 34.

9.1 Calling a Function by its Name

The `feval` function allows you to call a function from a string containing its name. This is useful when writing a function that needs to call user-supplied functions. The `feval` function takes the name of the function to call as its first argument, and the remaining arguments are given to the function.

The following example is a simple-minded function using `feval` that finds the root of a user-supplied function of one variable using Newton's method.

```
function result = newtroot (fname, x)

# usage: newtroot (fname, x)
#
#   fname : a string naming a function f(x).
#   x      : initial guess

delta = tol = sqrt (eps);
maxit = 200;
fx = feval (fname, x);
for i = 1:maxit
  if (abs (fx) < tol)
    result = x;
    return;
  else
    fx_new = feval (fname, x + delta);
    deriv = (fx_new - fx) / delta;
    x = x - fx / deriv;
    fx = fx_new;
  endif
endfor

result = x;

endfunction
```

Note that this is only meant to be an example of calling user-supplied functions and should not be taken too seriously. In addition to using a more robust algorithm, any serious code would check the number and type of all the arguments, ensure that the supplied function really was a function, etc. See [Section 4.8 \[Predicates for Numeric Objects\]](#), page 62, for a list of predicates for numeric objects, and see [Section 7.3 \[Status of Variables\]](#), page 132, for a description of the `exist` function.

`feval (name, ...)`

Evaluate the function named `name`.

Any arguments after the first are passed as inputs to the named function. For example,

```
feval ("acos", -1)
⇒ 3.1416
```

calls the function `acos` with the argument '`-1`'.

The function `feval` can also be used with function handles of any sort (see [Section 11.11.1 \[Function Handles\], page 214](#)). Historically, `feval` was the only way to call user-supplied functions in strings, but function handles are now preferred due to the cleaner syntax they offer. For example,

```
f = @exp;
feval (f, 1)
      => 2.7183
f (1)
      => 2.7183
```

are equivalent ways to call the function referred to by `f`. If it cannot be predicted beforehand whether `f` is a function handle, function name in a string, or inline function then `feval` can be used instead.

A similar function `run` exists for calling user script files, that are not necessarily on the user path

```
run script
run ("script")
Run script in the current workspace.
```

Scripts which reside in directories specified in Octave's load path, and which end with the extension "`.m`", can be run simply by typing their name. For scripts not located on the load path, use `run`.

The filename `script` can be a bare, fully qualified, or relative filename and with or without a file extension. If no extension is specified, Octave will first search for a script with the "`.m`" extension before falling back to the script name without an extension.

Implementation Note: If `script` includes a path component, then `run` first changes the working directory to the directory where `script` is found. Next, the script is executed. Finally, `run` returns to the original working directory unless `script` has specifically changed directories.

See also: [\[path\], page 197](#), [\[addpath\], page 196](#), [\[source\], page 207](#).

9.2 Evaluation in a Different Context

Before you evaluate an expression you need to substitute the values of the variables used in the expression. These are stored in the symbol table. Whenever the interpreter starts a new function it saves the current symbol table and creates a new one, initializing it with the list of function parameters and a couple of predefined variables such as `nargin`. Expressions inside the function use the new symbol table.

Sometimes you want to write a function so that when you call it, it modifies variables in your own context. This allows you to use a pass-by-name style of function, which is similar to using a pointer in programming languages such as C.

Consider how you might write `save` and `load` as m-files. For example:

```
function create_data
  x = linspace (0, 10, 10);
  y = sin (x);
  save mydata x y
endfunction
```

With `evalin`, you could write `save` as follows:

```
function save (file, name1, name2)
  f = open_save_file (file);
  save_var (f, name1, evalin ("caller", name1));
  save_var (f, name2, evalin ("caller", name2));
endfunction
```

Here, ‘caller’ is the `create_data` function and `name1` is the string “`x`”, which evaluates simply as the value of `x`.

You later want to load the values back from `mydata` in a different context:

```
function process_data
  load mydata
  ... do work ...
endfunction
```

With `assignin`, you could write `load` as follows:

```
function load (file)
  f = open_load_file (file);
  [name, val] = load_var (f);
  assignin ("caller", name, val);
  [name, val] = load_var (f);
  assignin ("caller", name, val);
endfunction
```

Here, ‘caller’ is the `process_data` function.

You can set and use variables at the command prompt using the context ‘`base`’ rather than ‘`caller`’.

These functions are rarely used in practice. One example is the `fail` (‘`code`’, ‘`pattern`’) function which evaluates ‘`code`’ in the caller’s context and checks that the error message it produces matches the given pattern. Other examples such as `save` and `load` are written in C++ where all Octave variables are in the ‘`caller`’ context and `evalin` is not needed.

evalin (*context, try*)
evalin (*context, try, catch*)

Like `eval`, except that the expressions are evaluated in the context `context`, which may be either “`caller`” or “`base`”.

See also: [\[eval\]](#), page 161, [\[assignin\]](#), page 164.

assignin (*context, varname, value*)

Assign `value` to `varname` in context `context`, which may be either “`base`” or “`caller`”.

See also: [\[evalin\]](#), page 164.

10 Statements

Statements may be a simple constant expression or a complicated list of nested loops and conditional statements.

Control statements such as `if`, `while`, and so on control the flow of execution in Octave programs. All the control statements start with special keywords such as `if` and `while`, to distinguish them from simple expressions. Many control statements contain other statements; for example, the `if` statement contains another statement which may or may not be executed.

Each control statement has a corresponding `end` statement that marks the end of the control statement. For example, the keyword `endif` marks the end of an `if` statement, and `endwhile` marks the end of a `while` statement. You can use the keyword `end` anywhere a more specific end keyword is expected, but using the more specific keywords is preferred because if you use them, Octave is able to provide better diagnostics for mismatched or missing end tokens.

The list of statements contained between keywords like `if` or `while` and the corresponding end statement is called the *body* of a control statement.

10.1 The if Statement

The `if` statement is Octave's decision-making statement. There are three basic forms of an `if` statement. In its simplest form, it looks like this:

```
if (condition)
    then-body
endif
```

condition is an expression that controls what the rest of the statement will do. The *then-body* is executed only if *condition* is true.

The condition in an `if` statement is considered true if its value is nonzero, and false if its value is zero. If the value of the conditional expression in an `if` statement is a vector or a matrix, it is considered true only if it is non-empty and *all* of the elements are nonzero. The conceptually equivalent code when *condition* is a matrix is shown below.

```
if (matrix) ≡ if (all (matrix(:)))
```

The second form of an `if` statement looks like this:

```
if (condition)
    then-body
else
    else-body
endif
```

If *condition* is true, *then-body* is executed; otherwise, *else-body* is executed.

Here is an example:

```
if (rem (x, 2) == 0)
    printf ("x is even\n");
else
    printf ("x is odd\n");
endif
```

In this example, if the expression `rem (x, 2) == 0` is true (that is, the value of `x` is divisible by 2), then the first `printf` statement is evaluated, otherwise the second `printf` statement is evaluated.

The third and most general form of the `if` statement allows multiple decisions to be combined in a single statement. It looks like this:

```
if (condition)
  then-body
elseif (condition)
  elseif-body
else
  else-body
endif
```

Any number of `elseif` clauses may appear. Each condition is tested in turn, and if one is found to be true, its corresponding `body` is executed. If none of the conditions are true and the `else` clause is present, its body is executed. Only one `else` clause may appear, and it must be the last part of the statement.

In the following example, if the first condition is true (that is, the value of `x` is divisible by 2), then the first `printf` statement is executed. If it is false, then the second condition is tested, and if it is true (that is, the value of `x` is divisible by 3), then the second `printf` statement is executed. Otherwise, the third `printf` statement is performed.

```
if (rem (x, 2) == 0)
  printf ("x is even\n");
elseif (rem (x, 3) == 0)
  printf ("x is odd and divisible by 3\n");
else
  printf ("x is odd\n");
endif
```

Note that the `elseif` keyword must not be spelled `else if`, as is allowed in Fortran. If it is, the space between the `else` and `if` will tell Octave to treat this as a new `if` statement within another `if` statement's `else` clause. For example, if you write

```
if (c1)
  body-1
else if (c2)
  body-2
endif
```

Octave will expect additional input to complete the first `if` statement. If you are using Octave interactively, it will continue to prompt you for additional input. If Octave is reading this input from a file, it may complain about missing or mismatched `end` statements, or, if you have not used the more specific `end` statements (`endif`, `endfor`, etc.), it may simply produce incorrect results, without producing any warning messages.

It is much easier to see the error if we rewrite the statements above like this,

```

if (c1)
    body-1
else
    if (c2)
        body-2
    endif

```

using the indentation to show how Octave groups the statements. See [Chapter 11 \[Functions and Scripts\]](#), page 177.

10.2 The switch Statement

It is very common to take different actions depending on the value of one variable. This is possible using the `if` statement in the following way

```

if (X == 1)
    do_something ();
elseif (X == 2)
    do_something_else ();
else
    do_something_completely_different ();
endif

```

This kind of code can however be very cumbersome to both write and maintain. To overcome this problem Octave supports the `switch` statement. Using this statement, the above example becomes

```

switch (X)
    case 1
        do_something ();
    case 2
        do_something_else ();
    otherwise
        do_something_completely_different ();
endswitch

```

This code makes the repetitive structure of the problem more explicit, making the code easier to read, and hence maintain. Also, if the variable `X` should change its name, only one line would need changing compared to one line per case when `if` statements are used.

The general form of the `switch` statement is

```

switch (expression)
    case label
        command_list
    case label
        command_list
    ...
    otherwise
        command_list
endswitch

```

where *label* can be any expression. However, duplicate *label* values are not detected, and only the *command_list* corresponding to the first match will be executed. For the **switch** statement to be meaningful at least one **case** *label command_list* clause must be present, while the **otherwise command_list** clause is optional.

If *label* is a cell array the corresponding *command_list* is executed if *any* of the elements of the cell array match *expression*. As an example, the following program will print ‘Variable is either 6 or 7’.

```
A = 7;
switch (A)
  case { 6, 7 }
    printf ("variable is either 6 or 7\n");
  otherwise
    printf ("variable is neither 6 nor 7\n");
endswitch
```

As with all other specific **end** keywords, **endswitch** may be replaced by **end**, but you can get better diagnostics if you use the specific forms.

One advantage of using the **switch** statement compared to using **if** statements is that the *labels* can be strings. If an **if** statement is used it is *not* possible to write

```
if (X == "a string") # This is NOT valid
```

since a character-to-character comparison between *X* and the string will be made instead of evaluating if the strings are equal. This special-case is handled by the **switch** statement, and it is possible to write programs that look like this

```
switch (X)
  case "a string"
    do_something
  ...
endswitch
```

10.2.1 Notes for the C Programmer

The **switch** statement is also available in the widely used C programming language. There are, however, some differences between the statement in Octave and C

- Cases are exclusive, so they don’t ‘fall through’ as do the cases in the **switch** statement of the C language.
- The *command_list* elements are not optional. Making the list optional would have meant requiring a separator between the label and the command list. Otherwise, things like

```
switch (foo)
  case (1) -2
  ...

```

would produce surprising results, as would

```
switch (foo)
  case (1)
  case (2)
    doit ();
  ...

```

particularly for C programmers. If `doit()` should be executed if `foo` is either 1 or 2, the above code should be written with a cell array like this

```
switch (foo)
  case { 1, 2 }
    doit ();
  ...

```

10.3 The while Statement

In programming, a *loop* means a part of a program that is (or at least can be) executed two or more times in succession.

The `while` statement is the simplest looping statement in Octave. It repeatedly executes a statement as long as a condition is true. As with the condition in an `if` statement, the condition in a `while` statement is considered true if its value is nonzero, and false if its value is zero. If the value of the conditional expression in a `while` statement is a vector or a matrix, it is considered true only if it is non-empty and *all* of the elements are nonzero.

Octave's `while` statement looks like this:

```
while (condition)
  body
endwhile
```

Here *body* is a statement or list of statements that we call the *body* of the loop, and *condition* is an expression that controls how long the loop keeps running.

The first thing the `while` statement does is test *condition*. If *condition* is true, it executes the statement *body*. After *body* has been executed, *condition* is tested again, and if it is still true, *body* is executed again. This process repeats until *condition* is no longer true. If *condition* is initially false, the body of the loop is never executed.

This example creates a variable `fib` that contains the first ten elements of the Fibonacci sequence.

```
fib = ones (1, 10);
i = 3;
while (i <= 10)
  fib (i) = fib (i-1) + fib (i-2);
  i++;
endwhile
```

Here the body of the loop contains two statements.

The loop works like this: first, the value of `i` is set to 3. Then, the `while` tests whether `i` is less than or equal to 10. This is the case when `i` equals 3, so the value of the `i`-th element of `fib` is set to the sum of the previous two values in the sequence. Then the `i++` increments the value of `i` and the loop repeats. The loop terminates when `i` reaches 11.

A newline is not required between the condition and the body; but using one makes the program clearer unless the body is very simple.

10.4 The do-until Statement

The **do-until** statement is similar to the **while** statement, except that it repeatedly executes a statement until a condition becomes true, and the test of the condition is at the end of the loop, so the body of the loop is always executed at least once. As with the condition in an **if** statement, the condition in a **do-until** statement is considered true if its value is nonzero, and false if its value is zero. If the value of the conditional expression in a **do-until** statement is a vector or a matrix, it is considered true only if it is non-empty and *all* of the elements are nonzero.

Octave's **do-until** statement looks like this:

```
do
  body
until (condition)
```

Here *body* is a statement or list of statements that we call the *body* of the loop, and *condition* is an expression that controls how long the loop keeps running.

This example creates a variable **fib** that contains the first ten elements of the Fibonacci sequence.

```
fib = ones (1, 10);
i = 2;
do
  i++;
  fib (i) = fib (i-1) + fib (i-2);
until (i == 10)
```

A newline is not required between the **do** keyword and the *body*; but using one makes the program clearer unless the *body* is very simple.

10.5 The for Statement

The **for** statement makes it more convenient to count iterations of a loop. The general form of the **for** statement looks like this:

```
for var = expression
  body
endfor
```

where *body* stands for any statement or list of statements, *expression* is any valid expression, and *var* may take several forms. Usually it is a simple variable name or an indexed variable. If the value of *expression* is a structure, *var* may also be a vector with two elements. See Section 10.5.1 [Looping Over Structure Elements], page 171, below.

The assignment expression in the **for** statement works a bit differently than Octave's normal assignment statement. Instead of assigning the complete result of the expression, it assigns each column of the expression to *var* in turn. If *expression* is a range, a row vector, or a scalar, the value of *var* will be a scalar each time the loop body is executed. If *var* is a column vector or a matrix, *var* will be a column vector each time the loop body is executed.

The following example shows another way to create a vector containing the first ten elements of the Fibonacci sequence, this time using the **for** statement:

```

fib = ones (1, 10);
for i = 3:10
    fib(i) = fib(i-1) + fib(i-2);
endfor

```

This code works by first evaluating the expression `3:10`, to produce a range of values from 3 to 10 inclusive. Then the variable `i` is assigned the first element of the range and the body of the loop is executed once. When the end of the loop body is reached, the next value in the range is assigned to the variable `i`, and the loop body is executed again. This process continues until there are no more elements to assign.

Within Octave it is also possible to iterate over matrices or cell arrays using the `for` statement. For example consider

```

disp ("Loop over a matrix")
for i = [1,3;2,4]
    i
endfor
disp ("Loop over a cell array")
for i = {1,"two";"three",4}
    i
endfor

```

In this case the variable `i` takes on the value of the columns of the matrix or cell matrix. So the first loop iterates twice, producing two column vectors `[1;2]`, followed by `[3;4]`, and likewise for the loop over the cell array. This can be extended to loops over multi-dimensional arrays. For example:

```

a = [1,3;2,4]; c = cat (3, a, 2*a);
for i = c
    i
endfor

```

In the above case, the multi-dimensional matrix `c` is reshaped to a two-dimensional matrix as `reshape (c, rows (c), prod (size (c)(2:end)))` and then the same behavior as a loop over a two-dimensional matrix is produced.

Although it is possible to rewrite all `for` loops as `while` loops, the Octave language has both statements because often a `for` loop is both less work to type and more natural to think of. Counting the number of iterations is very common in loops and it can be easier to think of this counting as part of looping rather than as something to do inside the loop.

10.5.1 Looping Over Structure Elements

A special form of the `for` statement allows you to loop over all the elements of a structure:

```

for [ val, key ] = expression
    body
endfor

```

In this form of the `for` statement, the value of `expression` must be a structure. If it is, `key` and `val` are set to the name of the element and the corresponding value in turn, until there are no more elements. For example:

```

x.a = 1
x.b = [1, 2; 3, 4]
x.c = "string"
for [val, key] = x
    key
    val
endfor

    -> key = a
    -> val = 1
    -> key = b
    -> val =
    ->
    ->     1 2
    ->     3 4
    ->
    -> key = c
    -> val = string

```

The elements are not accessed in any particular order. If you need to cycle through the list in a particular way, you will have to use the function `fieldnames` and sort the list yourself.

10.6 The break Statement

The `break` statement jumps out of the innermost `while`, `do-until`, or `for` loop that encloses it. The `break` statement may only be used within the body of a loop. The following example finds the smallest divisor of a given integer, and also identifies prime numbers:

```

num = 103;
div = 2;
while (div*div <= num)
    if (rem (num, div) == 0)
        break;
    endif
    div++;
endwhile
if (rem (num, div) == 0)
    printf ("Smallest divisor of %d is %d\n", num, div)
else
    printf ("%d is prime\n", num);
endif

```

When the remainder is zero in the first `while` statement, Octave immediately *breaks out* of the loop. This means that Octave proceeds immediately to the statement following the loop and continues processing. (This is very different from the `exit` statement which stops the entire Octave program.)

Here is another program equivalent to the previous one. It illustrates how the *condition* of a `while` statement could just as well be replaced with a `break` inside an `if`:

```

num = 103;
div = 2;
while (1)
    if (rem (num, div) == 0)
        printf ("Smallest divisor of %d is %d\n", num, div);
        break;
    endif
    div++;
    if (div*div > num)
        printf ("%d is prime\n", num);
        break;
    endif
endwhile

```

10.7 The continue Statement

The `continue` statement, like `break`, is used only inside `while`, `do-until`, or `for` loops. It skips over the rest of the loop body, causing the next cycle around the loop to begin immediately. Contrast this with `break`, which jumps out of the loop altogether. Here is an example:

```

# print elements of a vector of random
# integers that are even.

# first, create a row vector of 10 random
# integers with values between 0 and 100:

vec = round (rand (1, 10) * 100);

# print what we're interested in:

for x = vec
    if (rem (x, 2) != 0)
        continue;
    endif
    printf ("%d\n", x);
endfor

```

If one of the elements of `vec` is an odd number, this example skips the print statement for that element, and continues back to the first statement in the loop.

This is not a practical example of the `continue` statement, but it should give you a clear understanding of how it works. Normally, one would probably write the loop like this:

```

for x = vec
    if (rem (x, 2) == 0)
        printf ("%d\n", x);
    endif
endfor

```

10.8 The unwind_protect Statement

Octave supports a limited form of exception handling modeled after the `unwind-protect` form of Lisp.

The general form of an `unwind_protect` block looks like this:

```
unwind_protect
  body
  unwind_protect_cleanup
    cleanup
  end_unwind_protect
```

where `body` and `cleanup` are both optional and may contain any Octave expressions or commands. The statements in `cleanup` are guaranteed to be executed regardless of how control exits `body`.

This is useful to protect temporary changes to global variables from possible errors. For example, the following code will always restore the original value of the global variable `frobnosticate` even if an error occurs in the first part of the `unwind_protect` block.

```
save_frobnosticate = frobnosticate;
unwind_protect
  frobnosticate = true;
  ...
  unwind_protect_cleanup
    frobnosticate = save_frobnosticate;
  end_unwind_protect
```

Without `unwind_protect`, the value of `frobnosticate` would not be restored if an error occurs while evaluating the first part of the `unwind_protect` block because evaluation would stop at the point of the error and the statement to restore the value would not be executed.

In addition to `unwind_protect`, Octave supports another form of exception handling, the `try` block.

10.9 The try Statement

The original form of a `try` block looks like this:

```
try
  body
catch
  cleanup
end_tryCatch
```

where `body` and `cleanup` are both optional and may contain any Octave expressions or commands. The statements in `cleanup` are only executed if an error occurs in `body`.

No warnings or error messages are printed while `body` is executing. If an error does occur during the execution of `body`, `cleanup` can use the functions `lasterr` or `lasterror` to access the text of the message that would have been printed, as well as its identifier. The alternative form,

```

try
  body
catch err
  cleanup
end_tryCatch

```

will automatically store the output of `lasterror` in the structure `err`. See [Chapter 12 \[Errors and Warnings\]](#), page 221, for more information about the `lasterr` and `lasterror` functions.

10.10 Continuation Lines

In the Octave language, most statements end with a newline character and you must tell Octave to ignore the newline character in order to continue a statement from one line to the next. Lines that end with the characters `...` are joined with the following line before they are divided into tokens by Octave's parser. For example, the lines

```

x = long_variable_name ...
  + longer_variable_name ...
  - 42

```

form a single statement.

Any text between the continuation marker and the newline character is ignored. For example, the statement

```

x = long_variable_name ...      # comment one
  + longer_variable_name ...comment two
  - 42                      # last comment

```

is equivalent to the one shown above.

Inside double-quoted string constants, the character `\` has to be used as continuation marker. The `\` must appear at the end of the line just before the newline character:

```

s = "This text starts in the first line \
and is continued in the second line."

```

Input that occurs inside parentheses can be continued to the next line without having to use a continuation marker. For example, it is possible to write statements like

```

if (fine_dining_destination == on_a_boat
    || fine_dining_destination == on_a_train)
  seuss (i, will, not, eat, them, sam, i, am, i,
         will, not, eat, green, eggs, and, ham);
endif

```

without having to add to the clutter with continuation markers.

11 Functions and Scripts

Complicated Octave programs can often be simplified by defining functions. Functions can be defined directly on the command line during interactive Octave sessions, or in external files, and can be called just like built-in functions.

11.1 Introduction to Function and Script Files

There are seven different things covered in this section.

1. Typing in a function at the command prompt.
2. Storing a group of commands in a file — called a script file.
3. Storing a function in a file—called a function file.
4. Subfunctions in function files.
5. Multiple functions in one script file.
6. Private functions.
7. Nested functions.

Both function files and script files end with an extension of .m, for MATLAB compatibility. If you want more than one independent functions in a file, it must be a script file (see [Section 11.10 \[Script Files\], page 205](#)), and to use these functions you must execute the script file before you can use the functions that are in the script file.

11.2 Defining Functions

In its simplest form, the definition of a function named *name* looks like this:

```
function name
  body
endfunction
```

A valid function name is like a valid variable name: a sequence of letters, digits and underscores, not starting with a digit. Functions share the same pool of names as variables.

The function *body* consists of Octave statements. It is the most important part of the definition, because it says what the function should actually *do*.

For example, here is a function that, when executed, will ring the bell on your terminal (assuming that it is possible to do so):

```
function wakeup
  printf ("\a");
endfunction
```

The `printf` statement (see [Chapter 14 \[Input and Output\], page 251](#)) simply tells Octave to print the string "\a". The special character '\a' stands for the alert character (ASCII 7). See [Chapter 5 \[Strings\], page 67](#).

Once this function is defined, you can ask Octave to evaluate it by typing the name of the function.

Normally, you will want to pass some information to the functions you define. The syntax for passing parameters to a function in Octave is

```
function name (arg-list)
  body
endfunction
```

where *arg-list* is a comma-separated list of the function's arguments. When the function is called, the argument names are used to hold the argument values given in the call. The list of arguments may be empty, in which case this form is equivalent to the one shown above.

To print a message along with ringing the bell, you might modify the `wakeup` to look like this:

```
function wakeup (message)
  printf ("\a%s\n", message);
endfunction
```

Calling this function using a statement like this

```
wakeup ("Rise and shine!");
```

will cause Octave to ring your terminal's bell and print the message ‘Rise and shine!', followed by a newline character (the ‘\n' in the first argument to the `printf` statement).

In most cases, you will also want to get some information back from the functions you define. Here is the syntax for writing a function that returns a single value:

```
function ret-var = name (arg-list)
  body
endfunction
```

The symbol *ret-var* is the name of the variable that will hold the value to be returned by the function. This variable must be defined before the end of the function body in order for the function to return a value.

Variables used in the body of a function are local to the function. Variables named in *arg-list* and *ret-var* are also local to the function. See [Section 7.1 \[Global Variables\]](#), [page 129](#), for information about how to access global variables inside a function.

For example, here is a function that computes the average of the elements of a vector:

```
function retval = avg (v)
  retval = sum (v) / length (v);
endfunction
```

If we had written `avg` like this instead,

```
function retval = avg (v)
  if (isvector (v))
    retval = sum (v) / length (v);
  endif
endfunction
```

and then called the function with a matrix instead of a vector as the argument, Octave would have printed an error message like this:

```
error: value on right hand side of assignment is undefined
```

because the body of the `if` statement was never executed, and `retval` was never defined. To prevent obscure errors like this, it is a good idea to always make sure that the return

variables will always have values, and to produce meaningful error messages when problems are encountered. For example, `avg` could have been written like this:

```
function retval = avg (v)
    retval = 0;
    if (isvector (v))
        retval = sum (v) / length (v);
    else
        error ("avg: expecting vector argument");
    endif
endfunction
```

There is still one additional problem with this function. What if it is called without an argument? Without additional error checking, Octave will probably print an error message that won't really help you track down the source of the error. To allow you to catch errors like this, Octave provides each function with an automatic variable called `nargin`. Each time a function is called, `nargin` is automatically initialized to the number of arguments that have actually been passed to the function. For example, we might rewrite the `avg` function like this:

```
function retval = avg (v)
    retval = 0;
    if (nargin != 1)
        usage ("avg (vector)");
    endif
    if (isvector (v))
        retval = sum (v) / length (v);
    else
        error ("avg: expecting vector argument");
    endif
endfunction
```

Although Octave does not automatically report an error if you call a function with more arguments than expected, doing so probably indicates that something is wrong. Octave also does not automatically report an error if a function is called with too few arguments, but any attempt to use a variable that has not been given a value will result in an error. To avoid such problems and to provide useful messages, we check for both possibilities and issue our own error message.

nargin ()
nargin (*fcn*)

Report the number of input arguments to a function.

Called from within a function, return the number of arguments passed to the function.
At the top level, return the number of command line arguments passed to Octave.

If called with the optional argument *fcn*—a function name or handle—return the declared number of arguments that the function can accept.

If the last argument to *fcn* is `varargin` the returned value is negative. For example, the function `union` for sets is declared as

```
function [y, ia, ib] = union (a, b, varargin)
```

and

```
nargin ("union")
⇒ -3
```

Programming Note: `nargin` does not work on compiled functions (`.oct` files) such as built-in or dynamically loaded functions.

See also: [\[nargout\]](#), page 182, [\[narginchk\]](#), page 183, [\[varargin\]](#), page 188, [\[inputname\]](#), page 180.

`inputname (n)`

`inputname (n, ids_only)`

Return the name of the *n*-th argument to the calling function.

If the argument is not a simple variable name, return an empty string. As an example, a reference to a field in a structure such as `s.field` is not a simple name and will return "".

`inputname` is only useful within a function. When used at the command line it always returns an empty string.

By default, return an empty string if the *n*-th argument is not a valid variable name. If the optional argument *ids_only* is false, return the text of the argument even if it is not a valid variable name.

See also: [\[nargin\]](#), page 179, [\[nthargout\]](#), page 181.

```
val = silent_functions ()
```

```
old_val = silent_functions (new_val)
```

```
silent_functions (new_val, "local")
```

Query or set the internal variable that controls whether internal output from a function is suppressed.

If this option is disabled, Octave will display the results produced by evaluating expressions within a function body that are not terminated with a semicolon.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

11.3 Multiple Return Values

Unlike many other computer languages, Octave allows you to define functions that return more than one value. The syntax for defining functions that return multiple values is

```
function [ret-list] = name (arg-list)
    body
endfunction
```

where `name`, `arg-list`, and `body` have the same meaning as before, and `ret-list` is a comma-separated list of variable names that will hold the values returned from the function. The list of return values must have at least one element. If `ret-list` has only one element, this form of the `function` statement is equivalent to the form described in the previous section.

Here is an example of a function that returns two values, the maximum element of a vector and the index of its first occurrence in the vector.

```
function [max, idx] = vmax (v)
    idx = 1;
    max = v (idx);
    for i = 2:length (v)
        if (v (i) > max)
            max = v (i);
            idx = i;
        endif
    endfor
endfunction
```

In this particular case, the two values could have been returned as elements of a single array, but that is not always possible or convenient. The values to be returned may not have compatible dimensions, and it is often desirable to give the individual return values distinct names.

It is possible to use the **nthargout** function to obtain only some of the return values or several at once in a cell array. See [Section 3.1.5 \[Cell Array Objects\]](#), page 44.

```
nthargout (n, func, ...)
nthargout (n, ntot, func, ...)
```

Return the *n*th output argument of the function specified by the function handle or string *func*.

Any additional arguments are passed directly to *func*. The total number of arguments to call *func* with can be passed in *ntot*; by default *ntot* is *n*. The input *n* can also be a vector of indices of the output, in which case the output will be a cell array of the requested output arguments.

The intended use **nthargout** is to avoid intermediate variables. For example, when finding the indices of the maximum entry of a matrix, the following two compositions of **nthargout**

```
m = magic (5);
cell2mat (nthargout ([1, 2], @ind2sub, size (m),
                     nthargout (2, @max, m(:))))
```

$\Rightarrow 5 \quad 3$

are completely equivalent to the following lines:

```
m = magic (5);
[~, idx] = max (M(:));
[i, j] = ind2sub (size (m), idx);
[i, j]
⇒ 5   3
```

It can also be helpful to have all output arguments in a single cell in the following manner:

```
USV = nthargout ([1:3], @svd, hilb (5));
```

See also: [\[nargin\]](#), page 179, [\[nargout\]](#), page 182, [\[varargin\]](#), page 188, [\[varargout\]](#), page 188, [\[isargout\]](#), page 191.

In addition to setting `nargin` each time a function is called, Octave also automatically initializes `nargout` to the number of values that are expected to be returned. This allows you to write functions that behave differently depending on the number of values that the user of the function has requested. The implicit assignment to the built-in variable `ans` does not figure in the count of output arguments, so the value of `nargout` may be zero.

The `svd` and `lu` functions are examples of built-in functions that behave differently depending on the value of `nargout`.

It is possible to write functions that only set some return values. For example, calling the function

```
function [x, y, z] = f ()
    x = 1;
    z = 2;
endfunction
```

as

```
[a, b, c] = f ()
```

produces:

```
a = 1
```

```
b = [] (0x0)
```

```
c = 2
```

along with a warning.

```
nargout ()  
nargout (fcn)
```

Report the number of output arguments from a function.

Called from within a function, return the number of values the caller expects to receive. At the top level, `nargout` with no argument is undefined and will produce an error.

If called with the optional argument `fcn`—a function name or handle—return the number of declared output values that the function can produce.

If the final output argument is `varargout` the returned value is negative.

For example,

```
f ()
```

will cause `nargout` to return 0 inside the function `f` and

```
[s, t] = f ()
```

will cause `nargout` to return 2 inside the function `f`.

In the second usage,

```
nargout (@histc) # or nargout ("histc") using a string input
```

will return 2, because `histc` has two outputs, whereas

```
nargout (@imread)
```

will return -2, because `imread` has two outputs and the second is `varargout`.

Programming Note. `nargout` does not work for built-in functions and returns -1 for all anonymous functions.

See also: [\[nargin\]](#), page 179, [\[varargout\]](#), page 188, [\[isargout\]](#), page 191, [\[nthargout\]](#), page 181.

It is good practice at the head of a function to verify that it has been called correctly. In Octave the following idiom is seen frequently

```
if (nargin < min_#_inputs || nargin > max_#_inputs)
    print_usage ();
endif
```

which stops the function execution and prints a message about the correct way to call the function whenever the number of inputs is wrong.

For compatibility with MATLAB, `narginchk` and `nargoutchk` are available which provide similar error checking.

`narginchk` (*minargs*, *maxargs*)

Check for correct number of input arguments.

Generate an error message if the number of arguments in the calling function is outside the range *minargs* and *maxargs*. Otherwise, do nothing.

Both *minargs* and *maxargs* must be scalar numeric values. Zero, Inf, and negative values are all allowed, and *minargs* and *maxargs* may be equal.

Note that this function evaluates `nargin` on the caller.

See also: [\[nargoutchk\]](#), page 183, [\[error\]](#), page 221, [\[nargout\]](#), page 182, [\[nargin\]](#), page 179.

`nargoutchk` (*minargs*, *maxargs*)

```
msgstr = nargoutchk (minargs, maxargs, nargs)
```

```
msgstr = nargoutchk (minargs, maxargs, nargs, "string")
```

```
msgstruct = nargoutchk (minargs, maxargs, nargs, "struct")
```

Check for correct number of output arguments.

In the first form, return an error if the number of arguments is not between *minargs* and *maxargs*. Otherwise, do nothing. Note that this function evaluates the value of `nargout` on the caller so its value must have not been tampered with.

Both *minargs* and *maxargs* must be numeric scalars. Zero, Inf, and negative are all valid, and they can have the same value.

For backwards compatibility, the other forms return an appropriate error message string (or structure) if the number of outputs requested is invalid.

This is useful for checking to that the number of output arguments supplied to a function is within an acceptable range.

See also: [\[narginchk\]](#), page 183, [\[error\]](#), page 221, [\[nargout\]](#), page 182, [\[nargin\]](#), page 179.

Besides the number of arguments, inputs can be checked for various properties. `validatestring` is used for string arguments and `validateattributes` for numeric arguments.

```
validstr = validatestring (str, strarray)
validstr = validatestring (str, strarray, funcname)
validstr = validatestring (str, strarray, funcname, varname)
validstr = validatestring (..., position)
```

Verify that *str* is an element, or substring of an element, in *strarray*.

When *str* is a character string to be tested, and *strarray* is a cellstr of valid values, then *validstr* will be the validated form of *str* where validation is defined as *str* being a member or substring of *validstr*. This is useful for both verifying and expanding short options, such as "r", to their longer forms, such as "red". If *str* is a substring of *validstr*, and there are multiple matches, the shortest match will be returned if all matches are substrings of each other. Otherwise, an error will be raised because the expansion of *str* is ambiguous. All comparisons are case insensitive.

The additional inputs *funcname*, *varname*, and *position* are optional and will make any generated validation error message more specific.

Examples:

```
validatestring ("r", {"red", "green", "blue"})
⇒ "red"

validatestring ("b", {"red", "green", "blue", "black"})
⇒ error: validatestring: multiple unique matches were found for 'b':
      blue, black
```

See also: [\[strcmp\]](#), page 76, [\[strcmpi\]](#), page 77, [\[validateattributes\]](#), page 184, [\[inputParser\]](#), page 186.

```
validateattributes (A, classes, attributes)
validateattributes (A, classes, attributes, arg_idx)
validateattributes (A, classes, attributes, func_name)
validateattributes (A, classes, attributes, func_name, arg_name)
validateattributes (A, classes, attributes, func_name, arg_name,
                   arg_idx)
```

Check validity of input argument.

Confirms that the argument *A* is valid by belonging to one of *classes*, and holding all of the *attributes*. If it does not, an error is thrown, with a message formatted accordingly. The error message can be made further complete by the function name *func_name*, the argument name *arg_name*, and its position in the input *arg_idx*.

classes must be a cell array of strings (an empty cell array is allowed) with the name of classes (remember that a class name is case sensitive). In addition to the class name, the following categories names are also valid:

"float"	Floating point value comprising classes "double" and "single".
"integer"	Integer value comprising classes (u)int8, (u)int16, (u)int32, (u)int64.
"numeric"	Numeric value comprising either a floating point or integer value.
<i>attributes</i> must be a cell array with names of checks for <i>A</i> . Some of them require an additional value to be supplied right after the name (see details for each below).	
"<="	All values are less than or equal to the following value in <i>attributes</i> .

"<"	All values are less than the following value in <i>attributes</i> .
">="	All values are greater than or equal to the following value in <i>attributes</i> .
">"	All values are greater than the following value in <i>attributes</i> .
"2d"	A 2-dimensional matrix. Note that vectors and empty matrices have 2 dimensions, one of them being of length 1, or both length 0.
"3d"	Has no more than 3 dimensions. A 2-dimensional matrix is a 3-D matrix whose 3rd dimension is of length 1.
"binary"	All values are either 1 or 0.
"column"	Values are arranged in a single column.
"decreasing"	No value is NaN, and each is less than the preceding one.
"diag"	Value is a diagonal matrix.
"even"	All values are even numbers.
"finite"	All values are finite.
"increasing"	No value is NaN, and each is greater than the preceding one.
"integer"	All values are integer. This is different than using <code>isinteger</code> which only checks its an integer type. This checks that each value in <i>A</i> is an integer value, i.e., it has no decimal part.
"ncols"	Has exactly as many columns as the next value in <i>attributes</i> .
"ndims"	Has exactly as many dimensions as the next value in <i>attributes</i> .
"nondecreasing"	No value is NaN, and each is greater than or equal to the preceding one.
"nonempty"	It is not empty.
"nonincreasing"	No value is NaN, and each is less than or equal to the preceding one.
"nonnan"	No value is a NaN.
"nonnegative"	All values are non negative.
"nonsparse"	It is not a sparse matrix.
"nonzero"	No value is zero.
"nrows"	Has exactly as many rows as the next value in <i>attributes</i> .
"numel"	Has exactly as many elements as the next value in <i>attributes</i> .

```

"odd"      All values are odd numbers.

"positive"
            All values are positive.

"real"     It is a non-complex matrix.

"row"      Values are arranged in a single row.

"scalar"   It is a scalar.

"size"     Its size has length equal to the values of the next in attributes. The next
            value must be an array with the length for each dimension. To ignore the
            check for a certain dimension, the value of NaN can be used.

"square"   Is a square matrix.

"vector"   Values are arranged in a single vector (column or vector).

```

See also: `[isa]`, page 39, `[validatestring]`, page 183, `[inputParser]`, page 186.

If none of the preceding functions is sufficient there is also the class `inputParser` which can perform extremely complex input checking for functions.

`p = inputParser ()`

Create object `p` of the `inputParser` class.

This class is designed to allow easy parsing of function arguments. The class supports four types of arguments:

1. mandatory (see `addRequired`);
2. optional (see `addOptional`);
3. named (see `addParameter`);
4. switch (see `addSwitch`).

After defining the function API with these methods, the supplied arguments can be parsed with the `parse` method and the parsing results accessed with the `Results` accessor.

`inputParser.Parameters`

Return list of parameter names already defined.

`inputParser.Results`

Return structure with argument names as fieldnames and corresponding values.

`inputParser.Unmatched`

Return structure similar to `Results`, but for unmatched parameters. See the `KeepUnmatched` property.

`inputParser.UsingDefaults`

Return cell array with the names of arguments that are using default values.

`inputParser.CaseSensitive = boolean`

Set whether matching of argument names should be case sensitive. Defaults to false.

inputParser.FunctionName = name

Set function name to be used in error messages; Defaults to empty string.

inputParser.KeepUnmatched = boolean

Set whether an error should be given for non-defined arguments. Defaults to false. If set to true, the extra arguments can be accessed through **Unmatched** after the **parse** method. Note that since **Switch** and **Parameter** arguments can be mixed, it is not possible to know the unmatched type. If argument is found unmatched it is assumed to be of the **Parameter** type and it is expected to be followed by a value.

inputParser.StructExpand = boolean

Set whether a structure can be passed to the function instead of parameter/value pairs. Defaults to true.

The following example shows how to use this class:

```

function check (varargin)
    p = inputParser ();                                # create object
    p.FunctionName = "check";                         # set function name
    p.addRequired ("pack", @ischar);                  # mandatory argument
    p.addOptional ("path", pwd(), @ischar);           # optional argument

    ## create a function handle to anonymous functions for validators
    val_mat = @(x) isvector (x) && all (x <= 1) && all (x >= 0);
    p.addOptional ("mat", [0 0], val_mat);

    ## create two arguments of type "Parameter"
    val_type = @(x) any (strcmp (x, {"linear", "quadratic"}));
    p.addParameter ("type", "linear", val_type);
    val_verb = @(x) any (strcmp (x, {"low", "medium", "high"}));
    p.addParameter ("tolerance", "low", val_verb);

    ## create a switch type of argument
    p.addSwitch ("verbose");

    p.parse (varargin{:});  # Run created parser on inputs

    ## the rest of the function can access inputs by using p.Results.
    ## for example, get the tolerance input with p.Results.tolerance
endfunction

```

```

check ("mech");           # valid, use defaults for other arguments
check ();                 # error, one argument is mandatory
check (1);                # error, since ! ischar
check ("mech", "~/dev");  # valid, use defaults for other arguments

check ("mech", "~/dev", [0 1 0 0], "type", "linear"); # valid

## following is also valid. Note how the Switch argument type can
## be mixed into or before the Parameter argument type (but it
## must still appear after any Optional argument).
check ("mech", "~/dev", [0 1 0 0], "verbose", "tolerance", "high");

## following returns an error since not all optional arguments,
## 'path' and 'mat', were given before the named argument 'type'.
check ("mech", "~/dev", "type", "linear");

```

Note 1: A function can have any mixture of the four API types but they must appear in a specific order. Required arguments must be first and can be followed by any Optional arguments. Only the Parameter and Switch arguments may be mixed together and they must appear at the end.

Note 2: If both Optional and Parameter arguments are mixed in a function API then once a string Optional argument fails to validate it will be considered the end of the Optional arguments. The remaining arguments will be compared against any Parameter or Switch arguments.

See also: [\[nargin\]](#), page 179, [\[validateattributes\]](#), page 184, [\[validatestring\]](#), page 183, [\[varargin\]](#), page 188.

11.4 Variable-length Argument Lists

Sometimes the number of input arguments is not known when the function is defined. As an example think of a function that returns the smallest of all its input arguments. For example:

```
a = smallest (1, 2, 3);
b = smallest (1, 2, 3, 4);
```

In this example both `a` and `b` would be 1. One way to write the `smallest` function is

```
function val = smallest (arg1, arg2, arg3, arg4, arg5)
    body
endfunction
```

and then use the value of `nargin` to determine which of the input arguments should be considered. The problem with this approach is that it can only handle a limited number of input arguments.

If the special parameter name `varargin` appears at the end of a function parameter list it indicates that the function takes a variable number of input arguments. Using `varargin` the function looks like this

```
function val = smallest (varargin)
    body
endfunction
```

In the function body the input arguments can be accessed through the variable `varargin`. This variable is a cell array containing all the input arguments. See [Section 6.3 \[Cell Arrays\], page 114](#), for details on working with cell arrays. The `smallest` function can now be defined like this

```
function val = smallest (varargin)
    val = min ([varargin{:}]);
endfunction
```

This implementation handles any number of input arguments, but it's also a very simple solution to the problem.

A slightly more complex example of `varargin` is a function `print_arguments` that prints all input arguments. Such a function can be defined like this

```
function print_arguments (varargin)
    for i = 1:length (varargin)
        printf ("Input argument %d: ", i);
        disp (varargin{i});
    endfor
endfunction
```

This function produces output like this

```
print_arguments (1, "two", 3);
    ↘ Input argument 1: 1
    ↘ Input argument 2: two
    ↘ Input argument 3: 3

[reg, prop] = parseparams (params)
[reg, var1, ...] = parseparams (params, name1, default1, ...)
```

Return in `reg` the cell elements of `param` up to the first string element and in `prop` all remaining elements beginning with the first string element.

For example:

```
[reg, prop] = parseparams ({1, 2, "linewidth", 10})
reg =
{
    [1,1] = 1
    [1,2] = 2
}
prop =
{
    [1,1] = linewidth
    [1,2] = 10
}
```

The `parseparams` function may be used to separate regular numeric arguments from additional arguments given as property/value pairs of the `varargin` cell array.

In the second form of the call, available options are specified directly with their default values given as name-value pairs. If `params` do not form name-value pairs, or if an option occurs that does not match any of the available options, an error occurs.

When called from an m-file function, the error is prefixed with the name of the caller function.

The matching of options is case-insensitive.

See also: `[varargin]`, page 188, `[inputParser]`, page 186.

11.5 Ignoring Arguments

In the formal argument list, it is possible to use the dummy placeholder `~` instead of a name. This indicates that the corresponding argument value should be ignored and not stored to any variable.

```
function val = pick2nd (~, arg2)
    val = arg2;
endfunction
```

The value of `nargin` is not affected by using this declaration.

Return arguments can also be ignored using the same syntax. For example, the `sort` function returns both the sorted values, and an index vector for the original input which will result in a sorted output. Ignoring the second output is simple—don't request more than one output. But ignoring the first, and calculating just the second output, requires the use of the `~` placeholder.

```
x = [2, 3, 1];
[s, i] = sort (x)
⇒
s =
1   2   3

i =
3   1   2

[~, i] = sort (x)
⇒
i =
3   1   2
```

When using the `~` placeholder, commas—not whitespace—must be used to separate output arguments. Otherwise, the interpreter will view `~` as the logical not operator.

```
[~ i] = sort (x)
parse error:

invalid left hand side of assignment
```

Functions may take advantage of ignored outputs to reduce the number of calculations performed. To do so, use the `isargout` function to query whether the output argument is wanted. For example:

```

function [out1, out2] = long_function (x, y, z)
    if (isargout (1))
        ## Long calculation
        ...
        out1 = result;
    endif
    ...
endfunction

```

isargout (k)

Within a function, return a logical value indicating whether the argument *k* will be assigned to a variable on output.

If the result is false, the argument has been ignored during the function call through the use of the tilde (~) special output argument. Functions can use **isargout** to avoid performing unnecessary calculations for outputs which are unwanted.

If *k* is outside the range 1:max (nargout), the function returns false. *k* can also be an array, in which case the function works element-by-element and a logical array is returned. At the top level, **isargout** returns an error.

See also: [\[nargout\]](#), page 182, [\[varargout\]](#), page 188, [\[nthargout\]](#), page 181.

11.6 Variable-length Return Lists

It is possible to return a variable number of output arguments from a function using a syntax that's similar to the one used with the special **varargin** parameter name. To let a function return a variable number of output arguments the special output parameter name **varargout** is used. As with **varargin**, **varargout** is a cell array that will contain the requested output arguments.

As an example the following function sets the first output argument to 1, the second to 2, and so on.

```

function varargout = one_to_n ()
    for i = 1:nargout
        varargout{i} = i;
    endfor
endfunction

```

When called this function returns values like this

```

[a, b, c] = one_to_n ()
⇒ a = 1
⇒ b = 2
⇒ c = 3

```

If **varargin** (**varargout**) does not appear as the last element of the input (output) parameter list, then it is not special, and is handled the same as any other parameter name.

```

[r1, r2, ..., rn] = deal (a)
[r1, r2, ..., rn] = deal (a1, a2, ..., an)

```

Copy the input parameters into the corresponding output parameters.

If only a single input parameter is supplied, its value is copied to each of the outputs.

For example,

```
[a, b, c] = deal (x, y, z);
```

is equivalent to

```
a = x;
b = y;
c = z;
```

and

```
[a, b, c] = deal (x);
```

is equivalent to

```
a = b = c = x;
```

Programming Note: `deal` is often used with comma separated lists derived from cell arrays or structures. This is unnecessary as the interpreter can perform the same action without the overhead of a function call. For example:

```
c = {[1 2], "Three", 4};
[x, y, z] = c{:}
⇒
x =
```

```
1   2
```

```
y = Three
z = 4
```

See also: [\[cell2struct\]](#), page 124, [\[struct2cell\]](#), page 113, [\[repmat\]](#), page 491.

11.7 Returning from a Function

The body of a user-defined function can contain a `return` statement. This statement returns control to the rest of the Octave program. It looks like this:

```
return
```

Unlike the `return` statement in C, Octave's `return` statement cannot be used to return a value from a function. Instead, you must assign values to the list of return variables that are part of the `function` statement. The `return` statement simply makes it easier to exit a function from a deeply nested loop or conditional statement.

Here is an example of a function that checks to see if any elements of a vector are nonzero.

```
function retval = any_nonzero (v)
    retval = 0;
    for i = 1:length (v)
        if (v (i) != 0)
            retval = 1;
            return;
        endif
    endfor
    printf ("no nonzero elements found\n");
endfunction
```

Note that this function could not have been written using the `break` statement to exit the loop once a nonzero value is found without adding extra logic to avoid printing the message if the vector does contain a nonzero element.

`return`

When Octave encounters the keyword `return` inside a function or script, it returns control to the caller immediately. At the top level, the `return` statement is ignored. A `return` statement is assumed at the end of every function definition.

11.8 Default Arguments

Since Octave supports variable number of input arguments, it is very useful to assign default values to some input arguments. When an input argument is declared in the argument list it is possible to assign a default value to the argument like this

```
function name (arg1 = val1, ...)
    body
endfunction
```

If no value is assigned to `arg1` by the user, it will have the value `val1`.

As an example, the following function implements a variant of the classic “Hello, World” program.

```
function hello (who = "World")
    printf ("Hello, %s!\n", who);
endfunction
```

When called without an input argument the function prints the following

```
hello ();
⇒ Hello, World!
```

and when it’s called with an input argument it prints the following

```
hello ("Beautiful World of Free Software");
⇒ Hello, Beautiful World of Free Software!
```

Sometimes it is useful to explicitly tell Octave to use the default value of an input argument. This can be done writing a ‘`:`’ as the value of the input argument when calling the function.

```
hello (:);
⇒ Hello, World!
```

11.9 Function Files

Except for simple one-shot programs, it is not practical to have to define all the functions you need each time you need them. Instead, you will normally want to save them in a file so that you can easily edit them, and save them for use at a later time.

Octave does not require you to load function definitions from files before using them. You simply need to put the function definitions in a place where Octave can find them.

When Octave encounters an identifier that is undefined, it first looks for variables or functions that are already compiled and currently listed in its symbol table. If it fails to find a definition there, it searches a list of directories (the *path*) for files ending in `.m` that

have the same base name as the undefined identifier.¹ Once Octave finds a file with a name that matches, the contents of the file are read. If it defines a *single* function, it is compiled and executed. See [Section 11.10 \[Script Files\], page 205](#), for more information about how you can define more than one function in a single file.

When Octave defines a function from a function file, it saves the full name of the file it read and the time stamp on the file. If the time stamp on the file changes, Octave may reload the file. When Octave is running interactively, time stamp checking normally happens at most once each time Octave prints the prompt. Searching for new function definitions also occurs if the current working directory changes.

Checking the time stamp allows you to edit the definition of a function while Octave is running, and automatically use the new function definition without having to restart your Octave session.

To avoid degrading performance unnecessarily by checking the time stamps on functions that are not likely to change, Octave assumes that function files in the directory tree `octave-home/share/octave/version/m` will not change, so it doesn't have to check their time stamps every time the functions defined in those files are used. This is normally a very good assumption and provides a significant improvement in performance for the function files that are distributed with Octave.

If you know that your own function files will not change while you are running Octave, you can improve performance by calling `ignore_function_time_stamp ("all")`, so that Octave will ignore the time stamps for all function files. Passing "system" to this function resets the default behavior.

```
edit name
edit field value
value = edit ("get", field)
value = edit ("get", "all")
```

Edit the named function, or change editor settings.

If `edit` is called with the name of a file or function as its argument it will be opened in the text editor defined by `EDITOR`.

- If the function `name` is available in a file on your path and that file is modifiable, then it will be edited in place. If it is a system function, then it will first be copied to the directory `HOME` (see below) and then edited. If no file is found, then the m-file variant, ending with ".m", will be considered. If still no file is found, then variants with a leading "@" and then with both a leading "@" and trailing ".m" will be considered.
- If `name` is the name of a function defined in the interpreter but not in an m-file, then an m-file will be created in `HOME` to contain that function along with its current definition.
- If `name.cc` is specified, then it will search for `name.cc` in the path and try to modify it, otherwise it will create a new .cc file in the current directory. If `name` happens to be an m-file or interpreter defined function, then the text of that function will be inserted into the .cc file as a comment.

¹ The '.m' suffix was chosen for compatibility with MATLAB.

- If `name.ext` is on your path then it will be edited, otherwise the editor will be started with `name.ext` in the current directory as the filename. If `name.ext` is not modifiable, it will be copied to `HOME` before editing.

Warning: You may need to clear `name` before the new definition is available. If you are editing a `.cc` file, you will need to execute `mkoctfile name.cc` before the definition will be available.

If `edit` is called with `field` and `value` variables, the value of the control field `field` will be set to `value`.

If an output argument is requested and the first input argument is `get` then `edit` will return the value of the control field `field`. If the control field does not exist, `edit` will return a structure containing all fields and values. Thus, `edit ("get", "all")` returns a complete control structure.

The following control fields are used:

`'home'` This is the location of user local m-files. Be sure it is in your path. The default is `~/octave`.

`'author'` This is the name to put after the "`## Author:`" field of new functions. By default it guesses from the `gecos` field of the password database.

`'email'` This is the e-mail address to list after the name in the author field. By default it guesses `<$LOGNAME@$HOSTNAME>`, and if `$HOSTNAME` is not defined it uses `uname -n`. You probably want to override this. Be sure to use the format `user@host`.

`'license'`

`'gpl'` GNU General Public License (default).

`'bsd'` BSD-style license without advertising clause.

`'pd'` Public domain.

`"text"` Your own default copyright and license.

Unless you specify `'pd'`, `edit` will prepend the copyright statement with "Copyright (C) YYYY Author".

`'mode'` This value determines whether the editor should be started in async mode (editor is started in the background and Octave continues) or sync mode (Octave waits until the editor exits). Set it to `"sync"` to start the editor in sync mode. The default is `"async"` (see [system], page 913).

`'editinplace'`

Determines whether files should be edited in place, without regard to whether they are modifiable or not. The default is `false`.

```
mfilename ()  
mfilename ("fullpath")  
mfilename ("fullpathext")
```

Return the name of the currently executing file.

The base name of the currently executing script or function is returned without any extension. If called from outside an m-file, such as the command line, return the empty string.

Given the argument "fullpath", include the directory part of the filename, but not the extension.

Given the argument "fullpathext", include the directory part of the filename and the extension.

See also: [inputname], page 180, [dbstack], page 243.

```
val = ignore_function_time_stamp ()
old_val = ignore_function_time_stamp (new_val)
```

Query or set the internal variable that controls whether Octave checks the time stamp on files each time it looks up functions defined in function files.

If the internal variable is set to "system", Octave will not automatically recompile function files in subdirectories of *octave-home/lib/version* if they have changed since they were last compiled, but will recompile other function files in the search path if they change.

If set to "all", Octave will not recompile any function files unless their definitions are removed with `clear`.

If set to "none", Octave will always check time stamps on files to determine whether functions defined in function files need to be recompiled.

11.9.1 Manipulating the Load Path

When a function is called, Octave searches a list of directories for a file that contains the function declaration. This list of directories is known as the load path. By default the load path contains a list of directories distributed with Octave plus the current working directory. To see your current load path call the `path` function without any input or output arguments.

It is possible to add or remove directories to or from the load path using `addpath` and `rmpath`. As an example, the following code adds '`~/Octave`' to the load path.

```
addpath ("~/Octave")
```

After this the directory '`~/Octave`' will be searched for functions.

```
addpath (dir1, ...)
addpath (dir1, ..., option)
```

Add named directories to the function search path.

If *option* is "-begin" or 0 (the default), prepend the directory name to the current path. If *option* is "-end" or 1, append the directory name to the current path. Directories added to the path must exist.

In addition to accepting individual directory arguments, lists of directory names separated by `pathsep` are also accepted. For example:

```
addpath ("dir1:/dir2:~/dir3")
```

For each directory that is added, and that was not already in the path, `addpath` checks for the existence of a file named `PKG_ADD` (note lack of .m extension) and runs it if it exists.

See also: [path], page 197, [rmpath], page 197, [genpath], page 197, [pathdef], page 198, [savepath], page 197, [pathsep], page 198.

genpath (dir)

genpath (dir, skip, ...)

Return a path constructed from *dir* and all its subdirectories.

The path does not include package directories (beginning with '+'), old-style class directories (beginning with '@'), private directories, or any subdirectories of these types.

If additional string parameters are given, the resulting path will exclude directories with those names.

See also: [path], page 197, [addpath], page 196.

rmpath (dir1, ...)

Remove *dir1*, ... from the current function search path.

In addition to accepting individual directory arguments, lists of directory names separated by `pathsep` are also accepted. For example:

```
rmpath ("dir1:/dir2:~/dir3")
```

For each directory that is removed, `rmpath` checks for the existence of a file named `PKG_DEL` (note lack of .m extension) and runs it if it exists.

See also: [path], page 197, [addpath], page 196, [genpath], page 197, [pathdef], page 198, [savepath], page 197, [pathsep], page 198.

savepath ()

savepath (file)

status = savepath (...)

Save the unique portion of the current function search path that is not set during Octave's initialization process to *file*.

If *file* is omitted, Octave looks in the current directory for a project-specific `.octaverc` file in which to save the path information. If no such file is present then the user's configuration file `~/.octaverc` is used.

If successful, `savepath` returns 0.

The `savepath` function makes it simple to customize a user's configuration file to restore the working paths necessary for a particular instance of Octave. Assuming no filename is specified, Octave will automatically restore the saved directory paths from the appropriate `.octaverc` file when starting up. If a filename has been specified then the paths may be restored manually by calling `source file`.

See also: [path], page 197, [addpath], page 196, [rmpath], page 197, [genpath], page 197, [pathdef], page 198.

path ()

str = path ()

str = path (path1, ...)

Modify or display Octave's load path.

If *nargin* and *nargout* are zero, display the elements of Octave's load path in an easy to read format.

If *nargin* is zero and *nargout* is greater than zero, return the current load path.

If *nargin* is greater than zero, concatenate the arguments, separating them with `pathsep`. Set the internal search path to the result and return it.

No checks are made for duplicate elements.

See also: `[addpath]`, page 196, `[rmpath]`, page 197, `[genpath]`, page 197, `[pathdef]`, page 198, `[savepath]`, page 197, `[pathsep]`, page 198.

val = pathdef ()

Return the default path for Octave.

The path information is extracted from one of four sources. The possible sources, in order of preference, are:

1. `.octaverc`
2. `~/.octaverc`
3. `<OCTAVE_HOME>/.../<version>/m/startup/octaverc`
4. Octave's path prior to changes by any octaverc file.

See also: `[path]`, page 197, `[addpath]`, page 196, `[rmpath]`, page 197, `[genpath]`, page 197, `[savepath]`, page 197.

val = pathsep ()

Query the character used to separate directories in a path.

See also: `[filesep]`, page 904.

rehash ()

Reinitialize Octave's load path directory cache.

fname = file_in_loadpath (file)

fname = file_in_loadpath (file, "all")

Return the absolute name of *file* if it can be found in the list of directories specified by `path`.

If no file is found, return an empty character string.

When *file* is already an absolute name, the name is checked against the file system instead of Octave's loadpath. In this case, if *file* exists it will be returned in *fname*, otherwise an empty string is returned.

If the first argument is a cell array of strings, search each directory of the loadpath for element of the cell array and return the first that matches.

If the second optional argument "`all`" is supplied, return a cell array containing the list of all files that have the same name in the path. If no files are found, return an empty cell array.

See also: `[file_in_path]`, page 904, `[dir_in_loadpath]`, page 199, `[path]`, page 197.

restoredefaultpath ()

Restore Octave's path to its initial state at startup.

See also: [path], page 197, [addpath], page 196, [rmpath], page 197, [genpath], page 197, [pathdef], page 198, [savepath], page 197, [pathsep], page 198.

command_line_path ()

Return the command line path variable.

See also: [path], page 197, [addpath], page 196, [rmpath], page 197, [genpath], page 197, [pathdef], page 198, [savepath], page 197, [pathsep], page 198.

dirname = dir_in_loadpath (dir)**dirname = dir_in_loadpath (dir, "all")**

Return the absolute name of the loadpath element matching *dir* if it can be found in the list of directories specified by *path*.

If no match is found, return an empty character string.

The match is performed at the end of each path element. For example, if *dir* is "foo/bar", it matches the path element "/some/dir/foo/bar", but not "/some/dir/foo/bar/baz" "/some/dir/allfoo/bar". When *dir* is an absolute name, rather than just a path fragment, it is matched against the file system instead of Octave's loadpath. In this case, if *dir* exists it will be returned in *dirname*, otherwise an empty string is returned.

If the optional second argument is supplied, return a cell array containing all name matches rather than just the first.

See also: [file_in_path], page 904, [file_in_loadpath], page 198, [path], page 197.

11.9.2 Subfunctions

A function file may contain secondary functions called *subfunctions*. These secondary functions are only visible to the other functions in the same function file. For example, a file *f.m* containing

```
function f ()
  printf ("in f, calling g\n");
  g ()
endfunction
function g ()
  printf ("in g, calling h\n");
  h ()
endfunction
function h ()
  printf ("in h\n")
endfunction
```

defines a main function *f* and two subfunctions. The subfunctions *g* and *h* may only be called from the main function *f* or from the other subfunctions, but not from outside the file *f.m*.

localfunctions ()

Return a list of all local functions, i.e., subfunctions, within the current file.

The return value is a column cell array of function handles to all local functions accessible from the function from which `localfunctions` is called. Nested functions are *not* included in the list.

If the call is from the command line, an anonymous function, or a script, the return value is an empty cell array.

See also: [\[functions\]](#), page 214.

11.9.3 Private Functions

In many cases one function needs to access one or more helper functions. If the helper function is limited to the scope of a single function, then subfunctions as discussed above might be used. However, if a single helper function is used by more than one function, then this is no longer possible. In this case the helper functions might be placed in a subdirectory, called "private", of the directory in which the functions needing access to this helper function are found.

As a simple example, consider a function `func1`, that calls a helper function `func2` to do much of the work. For example:

```
function y = func1 (x)
    y = func2 (x);
endfunction
```

Then if the path to `func1` is `<directory>/func1.m`, and if `func2` is found in the directory `<directory>/private/func2.m`, then `func2` is only available for use of the functions, like `func1`, that are found in `<directory>`.

11.9.4 Nested Functions

Nested functions are similar to subfunctions in that only the main function is visible outside the file. However, they also allow for child functions to access the local variables in their parent function. This shared access mimics using a global variable to share information — but a global variable which is not visible to the rest of Octave. As a programming strategy, sharing data this way can create code which is difficult to maintain. It is recommended to use subfunctions in place of nested functions when possible.

As a simple example, consider a parent function `foo`, that calls a nested child function `bar`, with a shared variable `x`.

```
function y = foo ()
    x = 10;
    bar ();
    y = x;

    function bar ()
        x = 20;
    endfunction
endfunction

foo ()
⇒ 20
```

Notice that there is no special syntax for sharing `x`. This can lead to problems with accidental variable sharing between a parent function and its child. While normally variables are inherited, child function parameters and return values are local to the child function.

Now consider the function `foobar` that uses variables `x` and `y`. `foobar` calls a nested function `foo` which takes `x` as a parameter and returns `y`. `foo` then calls `bat` which does some computation.

```
function z = foobar ()
    x = 0;
    y = 0;
    z = foo (5);
    z += x + y;

    function y = foo (x)
        y = x + bat ();

        function z = bat ()
            z = x;
        endfunction
    endfunction
endfunction

foobar ()
⇒ 10
```

It is important to note that the `x` and `y` in `foobar` remain zero, as in `foo` they are a return value and parameter respectively. The `x` in `bat` refers to the `x` in `foo`.

Variable inheritance leads to a problem for `eval` and scripts. If a new variable is created in a parent function, it is not clear what should happen in nested child functions. For example, consider a parent function `foo` with a nested child function `bar`:

```
function y = foo (to_eval)
    bar ();
    eval (to_eval);

    function bar ()
        eval ("x = 100;");
        eval ("y = x;");
    endfunction
endfunction

foo ("x = 5;")
⇒ error: can not add variable "x" to a static workspace

foo ("y = 10;")
⇒ 10

foo ("")
⇒ 100
```

The parent function `foo` is unable to create a new variable `x`, but the child function `bar` was successful. Furthermore, even in an `eval` statement `y` in `bar` is the same `y` as in its parent function `foo`. The use of `eval` in conjunction with nested functions is best avoided.

As with subfunctions, only the first nested function in a file may be called from the outside. Inside a function the rules are more complicated. In general a nested function may call:

0. Globally visible functions
1. Any function that the nested function's parent can call
2. Sibling functions (functions that have the same parents)
3. Direct children

As a complex example consider a parent function `ex_top` with two child functions, `ex_a` and `ex_b`. In addition, `ex_a` has two more child functions, `ex_aa` and `ex_ab`. For example:

```
function ex_top ()
    ## Can call: ex_top, ex_a, and ex_b
    ## Can NOT call: ex_aa and ex_ab

    function ex_a ()
        ## Can call everything

        function ex_aa ()
            ## Can call everything
        endfunction

        function ex_ab ()
            ## Can call everything
        endfunction
    endfunction

    function ex_b ()
        ## Can call: ex_top, ex_a, and ex_b
        ## Can NOT call: ex_aa and ex_ab
    endfunction
endfunction
```

11.9.5 Overloading and Autoloading

Functions can be overloaded to work with different input arguments. For example, the operator '+' has been overloaded in Octave to work with single, double, uint8, int32, and many other arguments. The preferred way to overload functions is through classes and object oriented programming (see [Section 34.4.1 \[Function Overloading\], page 845](#)). Occasionally, however, one needs to undo user overloading and call the default function associated with a specific type. The `builtin` function exists for this purpose.

`[...] = builtin (f, ...)`

Call the base function `f` even if `f` is overloaded to another function for the given type signature.

This is normally useful when doing object-oriented programming and there is a requirement to call one of Octave's base functions rather than the overloaded one of a new class.

A trivial example which redefines the `sin` function to be the `cos` function shows how `builtin` works.

```
sin (0)
⇒ 0
function y = sin (x), y = cos (x); endfunction
sin (0)
⇒ 1
builtin ("sin", 0)
⇒ 0
```

A single dynamically linked file might define several functions. However, as Octave searches for functions based on the functions filename, Octave needs a manner in which to find each of the functions in the dynamically linked file. On operating systems that support symbolic links, it is possible to create a symbolic link to the original file for each of the functions which it contains.

However, there is at least one well known operating system that doesn't support symbolic links. Making copies of the original file for each of the functions is undesirable as it increases the amount of disk space used by Octave. Instead Octave supplies the `autoload` function, that permits the user to define in which file a certain function will be found.

```
autoload_map = autoload ()
autoload (function, file)
autoload (... , "remove")
Define function to autoload from file.
```

The second argument, `file`, should be an absolute filename or a file name in the same directory as the function or script from which the `autoload` command was run. `file` should not depend on the Octave load path.

Normally, calls to `autoload` appear in PKG_ADD script files that are evaluated when a directory is added to Octave's load path. To avoid having to hardcode directory names in `file`, if `file` is in the same directory as the PKG_ADD script then

```
autoload ("foo", "bar.oct");
```

will load the function `foo` from the file `bar.oct`. The above usage when `bar.oct` is not in the same directory, or usages such as

```
autoload ("foo", file_in_loadpath ("bar.oct"))
```

are strongly discouraged, as their behavior may be unpredictable.

With no arguments, return a structure containing the current autoload map.

If a third argument "`remove`" is given, the function is cleared and not loaded anymore during the current Octave session.

See also: [PKG_ADD], page 938.

11.9.6 Function Locking

It is sometime desirable to lock a function into memory with the `mlock` function. This is typically used for dynamically linked functions in oct-files or mex-files that contain some initialization, and it is desirable that calling `clear` does not remove this initialization.

As an example,

```
function my_function ()
    mlock ();
    ...
endfunction
```

prevents `my_function` from being removed from memory after it is called, even if `clear` is called. It is possible to determine if a function is locked into memory with the `mislocked`, and to unlock a function with `munlock`, which the following code illustrates.

```
my_function ();
mislocked ("my_function")
⇒ ans = 1
munlock ("my_function");
mislocked ("my_function")
⇒ ans = 0
```

A common use of `mlock` is to prevent persistent variables from being removed from memory, as the following example shows:

```
function count_calls ()
    mlock ();
    persistent calls = 0;
    printf ("count_calls() has been called %d times\n", ++calls);
endfunction

count_calls ();
→ count_calls() has been called 1 times

clear count_calls
count_calls ();
→ count_calls() has been called 2 times
```

`mlock` might also be used to prevent changes to an m-file, such as in an external editor, from having any effect in the current Octave session; A similar effect can be had with the `ignore_function_time_stamp` function.

`mlock ()`

Lock the current function into memory so that it can't be removed with `clear`.

See also: `[munlock]`, page 204, `[mislocked]`, page 205, `[persistent]`, page 130, `[clear]`, page 135.

`munlock ()`

`munlock (fcn)`

Unlock the named function `fcn` so that it may be removed from memory with `clear`.

If no function is named then unlock the current function.

See also: [mlock], page 204, [mislocked], page 205, [persistent], page 130, [clear], page 135.

mislocked ()
mislocked (fcn)

Return true if the named function *fcn* is locked in memory.

If no function is named then return true if the current function is locked.

See also: [mlock], page 204, [munlock], page 204, [persistent], page 130, [clear], page 135.

11.9.7 Function Precedence

Given the numerous different ways that Octave can define a function, it is possible and even likely that multiple versions of a function, might be defined within a particular scope. The precedence of which function will be used within a particular scope is given by

1. Subfunction A subfunction with the required function name in the given scope.
2. Private function A function defined within a private directory of the directory which contains the current function.
3. Class constructor A function that constructs a user class as defined in chapter Chapter 34 [Object Oriented Programming], page 835.
4. Class method An overloaded function of a class as in chapter Chapter 34 [Object Oriented Programming], page 835.
5. Command-line Function A function that has been defined on the command-line.
6. Autoload function A function that is marked as autoloaded with See [autoload], page 203.
7. A Function on the Path A function that can be found on the users load-path. There can also be Oct-file, mex-file or m-file versions of this function and the precedence between these versions are in that order.
8. Built-in function A function that is a part of core Octave such as `numel`, `size`, etc.

11.10 Script Files

A script file is a file containing (almost) any sequence of Octave commands. It is read and evaluated just as if you had typed each command at the Octave prompt, and provides a convenient way to perform a sequence of commands that do not logically belong inside a function.

Unlike a function file, a script file must *not* begin with the keyword `function`. If it does, Octave will assume that it is a function file, and that it defines a single function that should be evaluated as soon as it is defined.

A script file also differs from a function file in that the variables named in a script file are not local variables, but are in the same scope as the other variables that are visible on the command line.

Even though a script file may not begin with the `function` keyword, it is possible to define more than one function in a single script file and load (but not execute) all of them at once. To do this, the first token in the file (ignoring comments and other white space)

must be something other than `function`. If you have no other statements to evaluate, you can use a statement that has no effect, like this:

```
# Prevent Octave from thinking that this
# is a function file:

1;

# Define function one:

function one ()
...
```

To have Octave read and compile these functions into an internal form, you need to make sure that the file is in Octave's load path (accessible through the `path` function), then simply type the base name of the file that contains the commands. (Octave uses the same rules to search for script files as it does to search for function files.)

If the first token in a file (ignoring comments) is `function`, Octave will compile the function and try to execute it, printing a message warning about any non-whitespace characters that appear after the function definition.

Note that Octave does not try to look up the definition of any identifier until it needs to evaluate it. This means that Octave will compile the following statements if they appear in a script file, or are typed at the command line,

```
# not a function file:
1;
function foo ()
    do_something ();
endfunction
function do_something ()
    do_something_else ();
endfunction
```

even though the function `do_something` is not defined before it is referenced in the function `foo`. This is not an error because Octave does not need to resolve all symbols that are referenced by a function until the function is actually evaluated.

Since Octave doesn't look for definitions until they are needed, the following code will always print '`bar = 3`' whether it is typed directly on the command line, read from a script file, or is part of a function body, even if there is a function or script file called `bar.m` in Octave's path.

```
eval ("bar = 3");
bar
```

Code like this appearing within a function body could fool Octave if definitions were resolved as the function was being compiled. It would be virtually impossible to make Octave clever enough to evaluate this code in a consistent fashion. The parser would have to be able to perform the call to `eval` at compile time, and that would be impossible unless all the references in the string to be evaluated could also be resolved, and requiring that would be too restrictive (the string might come from user input, or depend on things that are not known until the function is evaluated).

Although Octave normally executes commands from script files that have the name `file.m`, you can use the function `source` to execute commands from any file.

```
source (file)
source (file, context)
```

Parse and execute the contents of *file*.

Without specifying *context*, this is equivalent to executing commands from a script file, but without requiring the file to be named `file.m` or to be on the execution path.

Instead of the current context, the script may be executed in either the context of the function that called the present function ("caller"), or the top-level context ("base").

See also: [run], page 163.

11.10.1 Publish Octave Script Files

The function `publish` provides a dynamic possibility to document your script file. Unlike static documentation, `publish` runs the script file, saves any figures and output while running the script, and presents them alongside static documentation in a desired output format. The static documentation can make use of Section 11.10.2 [Publishing Markup], page 209, to enhance and customize the output.

```
publish (file)
publish (file, output_format)
publish (file, option1, value1, ...)
publish (file, options)
output_file = publish (file, ...)
```

Generate a report from the Octave script file *file* in one of several output formats.

The generated reports interpret any Publishing Markup in comments, which is explained in detail in the GNU Octave manual. Assume the following example, using some Publishing Markup, to be the contents of the script file `pub_example.m`:

```
## Headline title
#
# Some *bold*, _italic_, or |monospaced| Text with
# a <https://www.gnu.org> link to *GNU Octave*.
##

# "Real" Octave commands to be evaluated
sombrero ()

%% MATLAB comment style ('%') is supported as well
%
% * Bulleted list item 1
% * Bulleted list item 2
%
% # Numbered list item 1
% # Numbered list item 2
```

To publish this script file, type `publish ("pub_example.m")`.

With only *file* given, a HTML report is generated in a subdirectory `html` relative to the current working directory. The Octave commands are evaluated in a separate context and any figures created while executing the script file are included in the report. All formatting syntax of *file* is treated according to the specified output format and included in the report.

Using `publish (file, output_format)` is equivalent to the function call using a structure

```
options.format = output_format;
publish (file, options)
```

which is described below. The same holds for using option/value pairs

```
options.option1 = value1;
publish (file, options)
```

The structure *options* can have the following field names. If a field name is not specified, the default value is used:

- ‘format’ — Output format of the published script file, one of ‘html’ (default), ‘doc’, ‘latex’, ‘ppt’, ‘pdf’, or ‘xml’.

The output formats ‘doc’, ‘ppt’, and ‘xml’ are not currently supported. To generate a ‘doc’ report, open a generated ‘html’ report with your office suite.

In Octave custom formats are supported by implementing all callback subfunctions in a function file named ‘`_publish_<custom format>_output__.m`’. To obtain a template for the HTML format type:

```
edit (fullfile (fileparts (which ("publish")), ...
    "private", "_publish_html_output__.m"))
```

- ‘outputDir’ — Full path of the directory where the generated report will be located. If no directory is given, the report is generated in a subdirectory `html` relative to the current working directory.
- ‘stylesheet’ — Not supported, only for MATLAB compatibility.
- ‘createThumbnail’ — Not supported, only for MATLAB compatibility.
- ‘figureSnapMethod’ — Not supported, only for MATLAB compatibility.
- ‘imageFormat’ — Desired format for any images produced while evaluating the code. The allowed image formats depend on the output format:
 - ‘html’, ‘xml’ — ‘png’ (default), any image format supported by Octave
 - ‘latex’ — ‘epsc2’ (default), any image format supported by Octave
 - ‘pdf’ — ‘jpg’ (default) or ‘bmp’, note MATLAB uses ‘bmp’ as default
 - ‘doc’ or ‘ppt’ — ‘png’ (default), ‘jpg’, ‘bmp’, or ‘tiff’
- ‘maxWidth’ and ‘maxHeight’ — Maximum width (height) of the produced images in pixels. An empty value means no restriction. Both values must be set in order for the option to work properly.
- ‘[]’ (default), integer value ≥ 0
- ‘useNewFigure’ — Use a new figure window for figures created by the evaluated code. This avoids side effects with already opened figure windows.
- ‘true’ (default) or ‘false’

- ‘`evalCode`’ — Evaluate code of the Octave source file
‘`true`’ (default) or ‘`false`’
- ‘`catchError`’ — Catch errors while evaluating code and continue
‘`true`’ (default) or ‘`false`’
- ‘`codeToEvaluate`’ — Octave commands that should be evaluated prior to publishing the script file. These Octave commands do not appear in the generated report.
- ‘`maxOutputLines`’ — Maximum number of output lines from code evaluation which are included in output.
‘`Inf`’ (default) or integer value > 0
- ‘`showCode`’ — Show the evaluated Octave commands in the generated report
‘`true`’ (default) or ‘`false`’

The option output `output_file` is a string with path and file name of the generated report.

See also: [\[grabcode\]](#), page 209.

The counterpart to `publish` is `grabcode`:

```
grabcode (url)
grabcode (filename)
code_str = grabcode (...)
```

Grab the code from a report created by the `publish` function.

The grabbed code inside the published report must be enclosed by the strings ‘##### SOURCE BEGIN #####’ and ‘##### SOURCE END #####’. The `publish` function creates this format automatically.

If no return value is requested the code is saved to a temporary file and opened in the default editor. NOTE: The temporary file must be saved under a new or the code will be lost.

If an output is requested the grabbed code will be returned as string `code_str`.

Example:

```
publish ("my_script.m");
grabcode ("html/my_script.html");
```

The example above publishes `my_script.m` to the default location `html/my_script.html`. Next, the published Octave script is grabbed to edit its content in a new temporary file.

See also: [\[publish\]](#), page 207.

11.10.2 Publishing Markup

11.10.2.1 Using Publishing Markup in Script Files

To use Publishing Markup, start by typing ‘##’ or ‘%%’ at the beginning of a new line. For MATLAB compatibility ‘%%%’ is treated the same way as ‘%%’.

The lines following ‘##’ or ‘%%’ start with one of either ‘#’ or ‘%’ followed by at least one space. These lines are interpreted as section. A section ends at the first line not starting with ‘#’ or ‘%’, or when the end of the document is reached.

A section starting in the first line of the document, followed by another start of a section that might be empty, is interpreted as a document title and introduction text.

See the example below for clarity:

```
%% Headline title
%
% Some *bold*, _italic_, or |monospaced| Text with
% a <https://www.gnu.org> link to GNU Octave.
%

# "Real" Octave commands to be evaluated
sombbrero ()

## Octave comment style supported as well
#
# * Bulleted list item 1
# * Bulleted list item 2
#
# # Numbered list item 1
# # Numbered list item 2
```

11.10.2.2 Text Formatting

Basic text formatting is supported inside sections, see the example given below:

```
##
# *bold*, _italic_, or |monospaced| Text
```

Additionally two trademark symbols are supported, just embrace the letters ‘TM’ or ‘R’.

```
##
# (TM) or (R)
```

11.10.2.3 Sections

A section is started by typing ‘##’ or ‘%%’ at the beginning of a new line. A section title can be provided by writing it, separated by a space, in the first line after ‘##’ or ‘%%’. Without a section title, the section is interpreted as a continuation of the previous section. For MATLAB compatibility ‘%%%’ is treated the same way as ‘%%’.

```

some_code ();

## Section 1
#
## Section 2

some_code ();

## 
# Still in section 2

some_code ();

%%% Section 3
%
%
```

11.10.2.4 Preformatted Code

To write preformatted code inside a section, indent the code by three spaces after ‘#’ at the beginning of each line and leave the lines above and below the code blank, except for ‘#’ at the beginning of those lines.

```

## 
# This is a syntax highlighted for-loop:
#
#   for i = 1:5
#     disp (i);
#   endfor
#
# And more usual text.
```

11.10.2.5 Preformatted Text

To write preformatted text inside a section, indent the code by two spaces after ‘#’ at the beginning of each line and leave the lines above and below the preformatted text blank, except for ‘#’ at the beginning of those lines.

```

## 
# This following text is preformatted:
#
# "To be, or not to be: that is the question:
# Whether 'tis nobler in the mind to suffer
#   The slings and arrows of outrageous fortune,
#   Or to take arms against a sea of troubles,
#   And by opposing end them? To die: to sleep;"
#
# --"Hamlet" by W. Shakespeare
```

11.10.2.6 Bulleted Lists

To create a bulleted list, type

```
##  
#  
# * Bulleted list item 1  
# * Bulleted list item 2  
#
```

to get output like

- Bulleted list item 1
- Bulleted list item 2

Notice the blank lines, except for the ‘#’ or ‘%’ before and after the bulleted list!

11.10.2.7 Numbered Lists

To create a numbered list, type

```
##  
#  
# # Numbered list item 1  
# # Numbered list item 2  
#
```

to get output like

1. Numbered list item 1
2. Numbered list item 2

Notice the blank lines, except for the ‘#’ or ‘%’ before and after the numbered list!

11.10.2.8 Including File Content

To include the content of an external file, e.g., a file called ‘my_function.m’ at the same location as the published Octave script, use the following syntax to include it with Octave syntax highlighting.

Alternatively, you can write the full or relative path to the file.

```
##  
#  
# <include>my_function.m</include>  
#  
# <include>/full/path/to/my_function.m</include>  
#  
# <include>../relative/path/to/my_function.m</include>  
#
```

11.10.2.9 Including Graphics

To include external graphics, e.g., a graphic called ‘my_graphic.png’ at the same location as the published Octave script, use the following syntax.

Alternatively, you can write the full path to the graphic.

```
##  
#  
# <<my_graphic.png>>  
#  
# <</full/path/to/my_graphic.png>>  
#  
# <<../relative/path/to/my_graphic.png>>  
#
```

11.10.2.10 Including URLs

Basically, a URL is written between an opening ‘<’ and a closing ‘>’ angle.

```
##  
# <https://www.octave.org>
```

Text that is within these angles and separated by at least one space from the URL is a displayed text for the link.

```
##  
# <https://www.octave.org GNU Octave>
```

A link starting with ‘<octave:’ followed by the name of a GNU Octave function, optionally with a displayed text, results in a link to the online GNU Octave documentations function index.

```
##  
# <octave:DISP The display function>
```

11.10.2.11 Mathematical Equations

One can insert L^AT_EX inline math, surrounded by single ‘\$’ signs, or displayed math, surrounded by double ‘\$\$’ signs, directly inside sections.

```
##  
# Some shorter inline equation $e^{ix} = \cos x + i\sin x$.  
#  
# Or more complicated formulas as displayed math:  
# $$e^x = \lim_{n\rightarrow\infty}\left(1+\frac{x}{n}\right)^n.
```

11.10.2.12 HTML Markup

If the published output is a HTML report, you can insert HTML markup, that is only visible in this kind of output.

```
##  
# <html>  
# <table style="border:1px solid black;">  
# <tr><td>1</td><td>2</td></tr>  
# <tr><td>3</td><td>3</td></tr>  
# </html>
```

11.10.2.13 LaTeX Markup

If the published output is a L^AT_EX or PDF report, you can insert L^AT_EX markup, that is only visible in this kind of output.

```
##  
# <tex>  
# Some output only visible in LaTeX or PDF reports.  
# \begin{equation}  
# e^x = \lim\limits_{n\rightarrow\infty}\left(1+\frac{x}{n}\right)^n  
# \end{equation}  
# </tex>
```

11.11 Function Handles, Anonymous Functions, Inline Functions

It can be very convenient store a function in a variable so that it can be passed to a different function. For example, a function that performs numerical minimization needs access to the function that should be minimized.

11.11.1 Function Handles

A function handle is a pointer to another function and is defined with the syntax

`@function-name`

For example,

```
f = @sin;
```

creates a function handle called `f` that refers to the function `sin`.

Function handles are used to call other functions indirectly, or to pass a function as an argument to another function like `quad` or `fsolve`. For example:

```
f = @sin;  
quad (f, 0, pi)  
⇒ 2
```

You may use `feval` to call a function using function handle, or simply write the name of the function handle followed by an argument list. If there are no arguments, you must use an empty argument list ‘()’. For example:

```
f = @sin;  
feval (f, pi/4)  
⇒ 0.70711  
f (pi/4)  
⇒ 0.70711
```

`is_function_handle (x)`

Return true if `x` is a function handle.

See also: `[isa]`, page 39, `[typeinfo]`, page 39, `[class]`, page 39, `[functions]`, page 214.

`s = functions (fcn_handle)`

Return a structure containing information about the function handle `fcn_handle`.

The structure *s* always contains these three fields:

function	The function name. For an anonymous function (no name) this will be the actual function definition.
type	Type of the function.
anonymous	The function is anonymous.
private	The function is private.
overloaded	The function overloads an existing function.
simple	The function is a built-in or m-file function.
subfunction	The function is a subfunction within an m-file.
file	The m-file that will be called to perform the function. This field is empty for anonymous and built-in functions.

In addition, some function types may return more information in additional fields.

Warning: `functions` is provided for debugging purposes only. Its behavior may change in the future and programs should not depend on any particular output format.

See also: [\[func2str\]](#), page 215, [\[str2func\]](#), page 215.

`func2str (fcn_handle)`

Return a string containing the name of the function referenced by the function handle *fcn_handle*.

See also: [\[str2func\]](#), page 215, [\[functions\]](#), page 214.

`str2func (fcn_name)`

Return a function handle constructed from the string *fcn_name*.

Previous versions of Octave accepted an optional second argument, "global", that caused str2func to ignore locally visible functions. This option is no longer supported.

See also: [\[func2str\]](#), page 215, [\[inline\]](#), page 216, [\[functions\]](#), page 214.

11.11.2 Anonymous Functions

Anonymous functions are defined using the syntax

`@(argument-list) expression`

Any variables that are not found in the argument list are inherited from the enclosing scope. Anonymous functions are useful for creating simple unnamed functions from expressions or for wrapping calls to other functions to adapt them for use by functions like `quad`. For example,

```
f = @(x) x.^2;
quad (f, 0, 10)
⇒ 333.33
```

creates a simple unnamed function from the expression `x.^2` and passes it to `quad`,

```
quad (@(x) sin (x), 0, pi)
⇒ 2
```

wraps another function, and

```
a = 1;
b = 2;
quad (@(x) betainc (x, a, b), 0, 0.4)
⇒ 0.13867
```

adapts a function with several parameters to the form required by `quad`. In this example, the values of `a` and `b` that are passed to `betainc` are inherited from the current environment.

Note that for performance reasons it is better to use handles to existing Octave functions, rather than to define anonymous functions which wrap an existing function. The integration of `sin (x)` is 5X faster if the code is written as

```
quad (@sin, 0, pi)
```

rather than using the anonymous function `@(x) sin (x)`. There are many operators which have functional equivalents that may be better choices than an anonymous function. Instead of writing

```
f = @(x, y) x + y
```

this should be coded as

```
f = @plus
```

See [Section 34.4.2 \[Operator Overloading\]](#), page 846, for a list of operators which also have a functional form.

11.11.3 Inline Functions

An inline function is created from a string containing the function body using the `inline` function. The following code defines the function $f(x) = x^2 + 2$.

```
f = inline ("x^2 + 2");
```

After this it is possible to evaluate f at any x by writing `f(x)`.

Caution: MATLAB has begun the process of deprecating inline functions. At some point in the future support will be dropped and eventually Octave will follow MATLAB and also remove inline functions. Use anonymous functions in all new code.

```
inline (str)
inline (str, arg1, ...)
inline (str, n)
```

Create an inline function from the character string `str`.

If called with a single argument, the arguments of the generated function are extracted from the function itself. The generated function arguments will then be in alphabetical order. It should be noted that `i` and `j` are ignored as arguments due to the ambiguity between their use as a variable or their use as an built-in constant. All arguments followed by a parenthesis are considered to be functions. If no arguments are found, a function taking a single argument named `x` will be created.

If the second and subsequent arguments are character strings, they are the names of the arguments of the function.

If the second argument is an integer n , the arguments are "x", "P1", ..., "PN".

Programming Note: The use of `inline` is discouraged and it may be removed from a future version of Octave. The preferred way to create functions from strings is through the use of anonymous functions (see [Section 11.11.2 \[Anonymous Functions\]](#), [page 215](#)) or `str2func`.

See also: `[argnames]`, [page 217](#), `[formula]`, [page 217](#), `[vectorize]`, [page 584](#), `[str2func]`, [page 215](#).

`argnames (fun)`

Return a cell array of character strings containing the names of the arguments of the inline function *fun*.

See also: `[inline]`, [page 216](#), `[formula]`, [page 217](#), `[vectorize]`, [page 584](#).

`formula (fun)`

Return a character string representing the inline function *fun*.

Note that `char (fun)` is equivalent to `formula (fun)`.

See also: `[char]`, [page 71](#), `[argnames]`, [page 217](#), `[inline]`, [page 216](#), `[vectorize]`, [page 584](#).

`vars = symvar (str)`

Identify the symbolic variable names in the string *str*.

Common constant names such as `i`, `j`, `pi`, `Inf` and Octave functions such as `sin` or `plot` are ignored.

Any names identified are returned in a cell array of strings. The array is empty if no variables were found.

Example:

```
symvar ("x^2 + y^2 == 4")
⇒ {
    [1,1] = x
    [2,1] = y
}
```

11.12 Commands

Commands are a special class of functions that only accept string input arguments. A command can be called as an ordinary function, but it can also be called without the parentheses. For example,

```
my_command hello world
```

is equivalent to

```
my_command ("hello", "world")
```

The general form of a command call is

```
cmdname arg1 arg2 ...
```

which translates directly to

```
cmdname ("arg1", "arg2", ...)
```

Any regular function can be used as a command if it accepts string input arguments. For example:

```
toupper lower_case_arg
⇒ ans = LOWER_CASE_ARG
```

One difficulty of commands occurs when one of the string input arguments is stored in a variable. Because Octave can't tell the difference between a variable name and an ordinary string, it is not possible to pass a variable as input to a command. In such a situation a command must be called as a function. For example:

```
strvar = "hello world";
toupper strvar
⇒ ans = STRVAR
toupper (strvar)
⇒ ans = HELLO WORLD
```

11.13 Organization of Functions Distributed with Octave

Many of Octave's standard functions are distributed as function files. They are loosely organized by topic, in subdirectories of *octave-home/lib/octave/version/m*, to make it easier to find them.

The following is a list of all the function file subdirectories, and the types of functions you will find there.

@ftp	Class functions for the FTP object.
+containers	Package for the containers classes.
audio	Functions for playing and recording sounds.
deprecated	Out-of-date functions which will eventually be removed from Octave.
elfun	Elementary functions, principally trigonometric.
general	Miscellaneous matrix manipulations, like flipud , rot90 , and triu , as well as other basic functions, like ismatrix , narginchk , etc.
geometry	Functions related to Delaunay triangulation.
gui	Functions for GUI elements like dialog, message box, etc.
help	Functions for Octave's built-in help system.
image	Image processing tools. These functions require the X Window System.
io	Input-output functions.
java	Functions related to the Java integration.
linear-algebra	Functions for linear algebra.
miscellaneous	Functions that don't really belong anywhere else.

ode	Functions to solve ordinary differential equations (ODEs).
optimization	Functions related to minimization, optimization, and root finding.
path	Functions to manage the directory path Octave uses to find functions.
pkg	Package manager for installing external packages of functions in Octave.
plot	Functions for displaying and printing two- and three-dimensional graphs.
polynomial	Functions for manipulating polynomials.
prefs	Functions implementing user-defined preferences.
set	Functions for creating and manipulating sets of unique values.
signal	Functions for signal processing applications.
sparse	Functions for handling sparse matrices.
specfun	Special functions such as <code>bessel</code> or <code>factor</code> .
special-matrix	Functions that create special matrix forms such as Hilbert or Vandermonde matrices.
startup	Octave's system-wide startup file.
statistics	Statistical functions.
strings	Miscellaneous string-handling functions.
testfun	Functions for performing unit tests on other functions.
time	Functions related to time and date processing.

12 Errors and Warnings

Octave includes several functions for printing error and warning messages. When you write functions that need to take special action when they encounter abnormal conditions, you should print the error messages using the functions described in this chapter.

Since many of Octave's functions use these functions, it is also useful to understand them, so that errors and warnings can be handled.

12.1 Handling Errors

An error is something that occurs when a program is in a state where it doesn't make sense to continue. An example is when a function is called with too few input arguments. In this situation the function should abort with an error message informing the user of the lacking input arguments.

Since an error can occur during the evaluation of a program, it is very convenient to be able to detect that an error occurred, so that the error can be fixed. This is possible with the `try` statement described in [Section 10.9 \[The try Statement\]](#), page 174.

12.1.1 Raising Errors

The most common use of errors is for checking input arguments to functions. The following example calls the `error` function if the function `f` is called without any input arguments.

```
function f (arg1)
  if (nargin == 0)
    error ("not enough input arguments");
  endif
endfunction
```

When the `error` function is called, it prints the given message and returns to the Octave prompt. This means that no code following a call to `error` will be executed.

It is also possible to assign an identification string to an error. If an error has such an ID the user can catch this error as will be described in the next section. To assign an ID to an error, simply call `error` with two string arguments, where the first is the identification string, and the second is the actual error. Note that error IDs are in the format "`NAMESPACE:ERROR-NAME`". The namespace "`Octave`" is used for Octave's own errors. Any other string is available as a namespace for user's own errors.

```
error (template, ...)
error (id, template, ...)
```

Display an error message and stop m-file execution.

Format the optional arguments under the control of the template string `template` using the same rules as the `printf` family of functions (see [Section 14.2.4 \[Formatted Output\]](#), page 277) and print the resulting message on the `stderr` stream. The message is prefixed by the character string '`error:`'.

Calling `error` also sets Octave's internal error state such that control will return to the top level without evaluating any further commands. This is useful for aborting from functions or scripts.

If the error message does not end with a newline character, Octave will print a traceback of all the function calls leading to the error. For example, given the following function definitions:

```
function f () g (); end
function g () h (); end
function h () nargin == 1 || error ("nargin != 1"); end
```

calling the function `f` will result in a list of messages that can help you to quickly find the exact location of the error:

```
f ()
error: nargin != 1
error: called from:
error:    h at line 1, column 27
error:    g at line 1, column 15
error:    f at line 1, column 15
```

If the error message ends in a newline character, Octave will print the message but will not display any traceback messages as it returns control to the top level. For example, modifying the error message in the previous example to end in a newline causes Octave to only print a single message:

```
function h () nargin == 1 || error ("nargin != 1\n"); end
f ()
error: nargin != 1
```

A null string ("") input to `error` will be ignored and the code will continue running as if the statement were a NOP. This is for compatibility with MATLAB. It also makes it possible to write code such as

```
err_msg = "";
if (CONDITION 1)
    err_msg = "CONDITION 1 found";
elseif (CONDITION2)
    err_msg = "CONDITION 2 found";
...
endif
error (err_msg);
```

which will only stop execution if an error has been found.

Implementation Note: For compatibility with MATLAB, escape sequences in *template* (e.g., "\n" => newline) are processed regardless of whether *template* has been defined with single quotes, as long as there are two or more input arguments. To disable escape sequence expansion use a second backslash before the sequence (e.g., "\\n") or use the `regexptranslate` function.

See also: [\[warning\]](#), page 228, [\[lasterror\]](#), page 224.

Since it is common to use errors when there is something wrong with the input to a function, Octave supports functions to simplify such code. When the `print_usage` function is called, it reads the help text of the function calling `print_usage`, and presents a useful error. If the help text is written in Texinfo it is possible to present an error message that only contains the function prototypes as described by the `@deftypefn` parts of the help

text. When the help text isn't written in Texinfo, the error message contains the entire help message.

Consider the following function.

```
## -*- texinfo -*-
## @deftypefn {} f (@var{arg1})
## Function help text goes here...
## @end deftypefn
function f (arg1)
  if ( nargin == 0)
    print_usage ();
  endif
endfunction
```

When it is called with no input arguments it produces the following error.

```
f ()

-| error: Invalid call to f.  Correct usage is:
-|
-|   -- f (ARG1)
-|
-|
-| Additional help for built-in functions and operators is
-| available in the online version of the manual.  Use the command
-| 'doc <topic>' to search the manual index.
-|
-| Help and information about Octave is also available on the WWW
-| at https://www.gnu.org and via the help@octave.org
-| mailing list.
```

```
print_usage ()
print_usage (name)
```

Print the usage message for the function *name*.

When called with no input arguments the `print_usage` function displays the usage message of the currently executing function.

See also: [\[help\]](#), page 20.

```
beep ()
```

Produce a beep from the speaker (or visual bell).

This function sends the alarm character "\a" to the terminal. Depending on the user's configuration this may produce an audible beep, a visual bell, or nothing at all.

See also: [\[puts\]](#), page 276, [\[fputs\]](#), page 275, [\[printf\]](#), page 277, [\[fprintf\]](#), page 277.

```
val = beep_on_error ()
old_val = beep_on_error (new_val)
beep_on_error (new_val, "local")
```

Query or set the internal variable that controls whether Octave will try to ring the terminal bell before printing an error message.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

12.1.2 Catching Errors

When an error occurs, it can be detected and handled using the `try` statement as described in [Section 10.9 \[The try Statement\]](#), page 174. As an example, the following piece of code counts the number of errors that occurs during a `for` loop.

```
number_of_errors = 0;
for n = 1:100
    try
        ...
    catch
        number_of_errors++;
    end_try_catch
endfor
```

The above example treats all errors the same. In many situations it can however be necessary to discriminate between errors, and take different actions depending on the error. The `lasterror` function returns a structure containing information about the last error that occurred. As an example, the code above could be changed to count the number of errors related to the '*' operator.

```
number_of_errors = 0;
for n = 1:100
    try
        ...
    catch
        msg = lasterror.message;
        if (strfind (msg, "operator *"))
            number_of_errors++;
        endif
    end_try_catch
endfor
```

Alternatively, the output of the `lasterror` function can be found in a variable indicated immediately after the `catch` keyword, as in the example below showing how to redirect an error as a warning:

```
try
    ...
catch err
    warning(err.identifier, err.message);
    ...
end_try_catch

lasterr = lasterror ()
lasterror (err)
lasterror ("reset")
Query or set the last error message structure.
```

When called without arguments, return a structure containing the last error message and other information related to this error. The elements of the structure are:

message	The text of the last error message
identifier	The message identifier of this error message
stack	A structure containing information on where the message occurred. This may be an empty structure if the information cannot be obtained. The fields of the structure are:
file	The name of the file where the error occurred
name	The name of function in which the error occurred
line	The line number at which the error occurred
column	An optional field with the column number at which the error occurred

The last error structure may be set by passing a scalar structure, *err*, as input. Any fields of *err* that match those above are set while any unspecified fields are initialized with default values.

If `lasterror` is called with the argument "reset", all fields are set to their default values.

See also: [\[lasterr\]](#), page 225, [\[error\]](#), page 221, [\[lastwarn\]](#), page 230.

<code>[msg, msgid] = lasterr ()</code>	
<code>lasterr (msg)</code>	Query or set the last error message.
<code>lasterr (msg, msgid)</code>	When called without input arguments, return the last error message and message identifier.
	With one argument, set the last error message to <i>msg</i> .
	With two arguments, also set the last message identifier.

See also: [\[lasterror\]](#), page 224, [\[error\]](#), page 221, [\[lastwarn\]](#), page 230.

The next example counts indexing errors. The errors are caught using the field identifier of the structure returned by the function `lasterror`.

```
number_of_errors = 0;
for n = 1:100
    try
        ...
    catch
        id = lasterror.identifier;
        if (strcmp (id, "Octave:invalid-indexing"))
            number_of_errors++;
        endif
    end_try_catch
endfor
```

The functions distributed with Octave can issue one of the following errors.

Octave:invalid-context

Indicates the error was generated by an operation that cannot be executed in the scope from which it was called. For example, the function `print_usage()` when called from the Octave prompt raises this error.

Octave:invalid-input-arg

Indicates that a function was called with invalid input arguments.

Octave:invalid-fun-call

Indicates that a function was called in an incorrect way, e.g., wrong number of input arguments.

Octave:invalid-indexing

Indicates that a data-type was indexed incorrectly, e.g., real-value index for arrays, nonexistent field of a structure.

Octave:bad-alloc

Indicates that memory couldn't be allocated.

Octave:undefined-function

Indicates a call to a function that is not defined. The function may exist but Octave is unable to find it in the search path.

When an error has been handled it is possible to raise it again. This can be useful when an error needs to be detected, but the program should still abort. This is possible using the `rethrow` function. The previous example can now be changed to count the number of errors related to the '*' operator, but still abort if another kind of error occurs.

```
number_of_errors = 0;
for n = 1:100
    try
        ...
    catch
        msg = lasterror.message;
        if (strfind (msg, "operator *"))
            number_of_errors++;
        else
            rethrow (lasterror);
        endif
    end_try_catch
endfor

rethrow (err)
```

Reissue a previous error as defined by `err`.

`err` is a structure that must contain at least the "`message`" and "`identifier`" fields. `err` can also contain a field "`stack`" that gives information on the assumed location of the error. Typically `err` is returned from `lasterror`.

See also: `[lasterror]`, page 224, `[lasterr]`, page 225, `[error]`, page 221.

```
err = errno ()
err = errno (val)
err = errno (name)
```

Query or set the system-dependent variable `errno`.

When called with no inputs, return the current value of `errno`.

When called with a numeric input `val`, set the current value of `errno` to the specified value. The previous value of `errno` is returned as `err`.

When called with a character string `name`, return the numeric value of `errno` which corresponds to the specified error code. If `name` is not a recognized error code then -1 is returned.

See also: [\[errno_list\]](#), page 227.

errno_list ()

Return a structure containing the system-dependent `errno` values.

See also: [\[errno\]](#), page 227.

12.1.3 Recovering From Errors

Octave provides several ways of recovering from errors. There are `try/catch` blocks, `unwind_protect/unwind_protect_cleanup` blocks, and finally the `onCleanup` command.

The `onCleanup` command associates an ordinary Octave variable (the trigger) with an arbitrary function (the action). Whenever the Octave variable ceases to exist—whether due to a function return, an error, or simply because the variable has been removed with `clear`—then the assigned function is executed.

The function can do anything necessary for cleanup such as closing open file handles, printing an error message, or restoring global variables to their initial values. The last example is a very convenient idiom for Octave code. For example:

```
function rand42
    old_state = rand ("state");
    restore_state = onCleanup (@() rand ("state", old_state));
    rand ("state", 42);
    ...
endfunction # rand generator state restored by onCleanup
```

obj = onCleanup (function)

Create a special object that executes a given function upon destruction.

If the object is copied to multiple variables (or cell or struct array elements) or returned from a function, `function` will be executed after clearing the last copy of the object. Note that if multiple local `onCleanup` variables are created, the order in which they are called is unspecified. For similar functionality See [Section 10.8 \[The unwind_protect Statement\]](#), page 174.

12.2 Handling Warnings

Like an error, a warning is issued when something unexpected happens. Unlike an error, a warning doesn't abort the currently running program. A simple example of a warning is

when a number is divided by zero. In this case Octave will issue a warning and assign the value `Inf` to the result.

```
a = 1/0
  -> warning: division by zero
  => a = Inf
```

12.2.1 Issuing Warnings

It is possible to issue warnings from any code using the `warning` function. In its most simple form, the `warning` function takes a string describing the warning as its input argument. As an example, the following code controls if the variable ‘`a`’ is non-negative, and if not issues a warning and sets ‘`a`’ to zero.

```
a = -1;
if (a < 0)
    warning ("'a' must be non-negative. Setting 'a' to zero.");
    a = 0;
endif
  -> 'a' must be non-negative. Setting 'a' to zero.
```

Since warnings aren’t fatal to a running program, it is not possible to catch a warning using the `try` statement or something similar. It is however possible to access the last warning as a string using the `lastwarn` function.

It is also possible to assign an identification string to a warning. If a warning has such an ID the user can enable and disable this warning as will be described in the next section. To assign an ID to a warning, simply call `warning` with two string arguments, where the first is the identification string, and the second is the actual warning. Note that warning IDs are in the format “`NAMESPACE:WARNING-NAME`”. The namespace “`Octave`” is used for Octave’s own warnings. Any other string is available as a namespace for user’s own warnings.

```
warning (template, ...)
warning (id, template, ...)
warning ("on", id)
warning ("off", id)
warning ("error", id)
warning ("query", id)
warning (state, id, "local")
warning (warning_struct)
warning_struct = warning (...)
warning (state, mode)
```

Display a warning message or control the behavior of Octave’s warning system.

The first call form uses a template `template` and optional additional arguments to display a message on the `stderr` stream. The message is formatted using the same rules as the `printf` family of functions (see [Section 14.2.4 \[Formatted Output\]](#), page [277](#)) and prefixed by the character string ‘`warning:`’. You should use this function when you want to notify the user of an unusual condition, but only when it makes sense for your program to go on. For example:

```
warning ("foo: maybe something wrong here");
```

The optional warning identifier *id* allows users to enable or disable warnings tagged by this identifier. A message identifier is a string of the form "NAMESPACE:WARNING-NAME". Octave's own warnings use the "Octave" namespace (see [warning_ids], page 230). For example:

```
warning ("MyNameSpace:check-something",
        "foo: maybe something wrong here");
```

The second call form is meant to change and/or query the state of warnings. The first input argument must be a string *state* ("on", "off", "error", or "query") followed by an optional warning identifier *id* or "all" (default).

The optional output argument *warning_struct* is a structure or structure array with fields "state" and "identifier". The *state* argument may have the following values:

"on" | "off":

Enable or disable the display of warnings identified by *id* and optionally return their previous state *stout*.

"error": Turn warnings identified by *id* into errors and optionally return their previous state *stout*.

"query": Return the current state of warnings identified by *id*.

A structure or structure array *warning_struct*, with fields "state" and "identifier", may be given as an input to achieve equivalent results. The following example shows how to temporarily disable a warning and then restore its original state:

```
loglog (-1:10);
## Disable the previous warning and save its original state
[~, id] = lastwarn ();
warnstate = warning ("off", id);
loglog (-1:10);
## Restore its original state
warning (warnstate);
```

If a final argument "local" is provided then the warning state will be set temporarily until the end of the current function. Changes to warning states that are set locally affect the current function and all functions called from the current scope. The previous warning state is restored on return from the current function. The "local" option is ignored if used in the top-level workspace.

With no input argument `warning()` is equivalent to `warning ("query", "all")` except that in the absence of an output argument, the state of warnings is displayed on `stderr`.

The level of verbosity of the warning system may also be controlled by two modes *mode*:

"backtrace":

enable/disable the display of the stack trace after the warning message

"verbose":

enable/disable the display of additional information after the warning message

In this case the *state* argument may only be "on" or "off".

Implementation Note: For compatibility with MATLAB, escape sequences in *template* (e.g., "\n" => newline) are processed regardless of whether *template* has been defined with single quotes, as long as there are two or more input arguments. To disable escape sequence expansion use a second backslash before the sequence (e.g., "\\n") or use the `regexptranslate` function.

See also: [warning_ids], page 230, [lastwarn], page 230, [error], page 221.

```
[msg, msgid] = lastwarn ()
lastwarn (msg)
lastwarn (msg, msgid)
```

Query or set the last warning message.

When called without input arguments, return the last warning message and message identifier.

With one argument, set the last warning message to *msg*.

With two arguments, also set the last message identifier.

See also: [warning], page 228, [lasterror], page 224, [lasterr], page 225.

The functions distributed with Octave can issue one of the following warnings.

Octave:abbreviated-property-match

If the `Octave:abbreviated-property-match` warning is enabled.

Octave:addpath-pkg

If the `Octave:addpath-pkg` warning is enabled, Octave will warn when a package directory (i.e., +package_name) is added to the `path`. Typically, only the parent directory which contains the package directory should be added to the load path. By default, the `Octave:addpath-pkg` warning is enabled.

Octave:array-as-logical

If the `Octave:array-as-logical` warning is enabled, Octave will warn when an array of size greater than 1x1 is used as a truth value in an if, while or until statement. By default, the `Octave:array-as-logical` warning is disabled.

Octave:array-to-scalar

If the `Octave:array-to-scalar` warning is enabled, Octave will warn when an implicit conversion from an array to a scalar value is attempted. By default, the `Octave:array-to-scalar` warning is disabled.

Octave:array-to-vector

If the `Octave:array-to-vector` warning is enabled, Octave will warn when an implicit conversion from an array to a vector value is attempted. By default, the `Octave:array-to-vector` warning is disabled.

Octave:assign-as-truth-value

If the `Octave:assign-as-truth-value` warning is enabled, a warning is issued for statements like

```
if (s = t)
```

```
...
```

since such statements are not common, and it is likely that the intent was to write

```
if (s == t)
    ...
```

instead.

There are times when it is useful to write code that contains assignments within the condition of a `while` or `if` statement. For example, statements like

```
while (c = getc ())
    ...
```

are common in C programming.

It is possible to avoid all warnings about such statements by disabling the `Octave:assign-as-truth-value` warning, but that may also let real errors like

```
if (x = 1) # intended to test (x == 1)!
    ...
```

slip by.

In such cases, it is possible suppress errors for specific statements by writing them with an extra set of parentheses. For example, writing the previous example as

```
while ((c = getc ()))
    ...
```

will prevent the warning from being printed for this statement, while allowing Octave to warn about other assignments used in conditional contexts.

By default, the `Octave:assign-as-truth-value` warning is enabled.

`Octave:associativity-change`

If the `Octave:associativity-change` warning is enabled, Octave will warn about possible changes in the meaning of some code due to changes in associativity for some operators. Associativity changes have typically been made for MATLAB compatibility. By default, the `Octave:associativity-change` warning is enabled.

`Octave:autoload-relative-file-name`

If the `Octave:autoload-relative-file-name` is enabled, Octave will warn when parsing `autoload()` function calls with relative paths to function files. This usually happens when using `autoload()` calls in PKG_ADD files, when the PKG_ADD file is not in the same directory as the .oct file referred to by the `autoload()` command. By default, the `Octave:autoload-relative-file-name` warning is enabled.

`Octave:built-in-variable-assignment`

By default, the `Octave:built-in-variable-assignment` warning is enabled.

`Octave:deprecated-function`

If the `Octave:deprecated-function` warning is enabled, a warning is issued when Octave encounters a function that is obsolete and scheduled for removal

from Octave. By default, the `Octave:deprecated-function` warning is enabled.

Octave:deprecated-keyword

If the `Octave:deprecated-keyword` warning is enabled, a warning is issued when Octave encounters a keyword that is obsolete and scheduled for removal from Octave. By default, the `Octave:deprecated-keyword` warning is enabled.

Octave:deprecated-property

If the `Octave:deprecated-property` warning is enabled, a warning is issued when Octave encounters a graphics property that is obsolete and scheduled for removal from Octave. By default, the `Octave:deprecated-property` warning is enabled.

Octave:divide-by-zero

If the `Octave:divide-by-zero` warning is enabled, a warning is issued when Octave encounters a division by zero. By default, the `Octave:divide-by-zero` warning is enabled.

Octave:eigs:UnconvergedEigenvalues

If the `Octave:eigs:UnconvergedEigenvalues` warning is enabled then the `eigs` function will issue a warning if the number of calculated eigenvalues is less than the number of requested eigenvalues. By default, the `Octave:eigs:UnconvergedEigenvalues` warning is enabled.

Octave:erase:chararray

If the `Octave:erase:chararray` warning is enabled then the `erase` function will issue a warning if the input pattern is a character array rather than a string or cell array of strings. By default, the `Octave:erase:chararray` warning is enabled.

Octave:data-file-in-path

If the `Octave:data-file-in-path` warning is enabled, a warning is issued when Octave does not find the target of a file operation such as `load` or `fopen` directly, but is able to locate the file in Octave's search `path` for files. The warning could indicate that a different file target than the programmer intended is being used. By default, the `Octave:data-file-in-path` warning is enabled.

Octave:function-name-clash

If the `Octave:function-name-clash` warning is enabled, a warning is issued when Octave finds that the name of a function defined in a function file differs from the name of the file. (If the names disagree, the name declared inside the file is ignored.) By default, the `Octave:function-name-clash` warning is enabled.

Octave:future-time-stamp

If the `Octave:future-time-stamp` warning is enabled, Octave will print a warning if it finds a function file with a time stamp that is in the future. By default, the `Octave:future-time-stamp` warning is enabled.

Octave:glyph-render

By default, the `Octave:glyph-render` warning is enabled.

Octave:imag-to-real

If the `Octave:imag-to-real` warning is enabled, a warning is printed for implicit conversions of complex numbers to real numbers. By default, the `Octave:imag-to-real` warning is disabled.

Octave:language-extension

Print warnings when using features that are unique to the Octave language and that may still be missing in MATLAB. By default, the `Octave:language-extension` warning is disabled. The `--traditional` or `--braindead` startup options for Octave may also be of use, see [Section 2.1.1 \[Command Line Options\], page 15](#).

Octave:legacy-function

If the `Octave:legacy-function` warning is enabled, a warning is issued when Octave encounters a function that MATLAB has suggested should be avoided. The function may become obsolete at some point in the future and removed, in which case the warning will change to `Octave:deprecated-function`, and the function will continue to exist for two further versions of Octave before being removed. By default, the `Octave:legacy-function` warning is enabled.

Octave:logical-conversion

By default, the `Octave:logical-conversion` warning is enabled.

Octave:lu:sparse_input

If the `Octave:lu:sparse_input` warning is enabled, Octave will warn when the `lu` function is called with a sparse input and less than four output arguments. In this case, sparsity-preserving column permutations are not performed and the result may be inaccurate. By default, the `Octave:lu:sparse_input` warning is enabled.

Octave:missing-glyph

By default, the `Octave:missing-glyph` warning is enabled.

Octave:missing-seicolon

If the `Octave:missing-seicolon` warning is enabled, Octave will warn when statements in function definitions don't end in semicolons. By default the `Octave:missing-seicolon` warning is disabled.

Octave:mixed-string-concat

If the `Octave:mixed-string-concat` warning is enabled, print a warning when concatenating a mixture of double and single quoted strings. By default, the `Octave:mixed-string-concat` warning is disabled.

Octave:nearly-singular-matrix**Octave:singular-matrix**

These warnings are emitted if a (nearly) singular matrix is inverted. By default, the `Octave:nearly-singular-matrix` and `Octave:singular-matrix` warnings are enabled.

Octave:neg-dim-as-zero

If the `Octave:neg-dim-as-zero` warning is enabled, print a warning for expressions like

`eye (-1)`

By default, the `Octave:neg-dim-as-zero` warning is disabled.

`Octave:nested-functions-coerced`

By default, the `Octave:nested-functions-coerced` warning is enabled.

`Octave:noninteger-range-as-index`

By default, the `Octave:noninteger-range-as-index` warning is enabled.

`Octave:num-to-str`

If the `Octave:num-to-str` warning is enable, a warning is printed for implicit conversions of numbers to their UTF-8 encoded character equivalents when strings are constructed using a mixture of strings and numbers in matrix notation. For example,

```
[ "f", 111, 111 ]  
⇒ "foo"
```

elicits a warning if the `Octave:num-to-str` warning is enabled. By default, the `Octave:num-to-str` warning is enabled.

`Octave:possible-matlab-short-circuit-operator`

If the `Octave:possible-matlab-short-circuit-operator` warning is enabled, Octave will warn about using the not short circuiting operators & and | inside if or while conditions. They normally never short circuit, but they do short circuit when used in a condition. By default, the `Octave:possible-matlab-short-circuit-operator` warning is enabled.

`Octave:precedence-change`

If the `Octave:precedence-change` warning is enabled, Octave will warn about possible changes in the meaning of some code due to changes in precedence for some operators. Precedence changes have typically been made for MATLAB compatibility. By default, the `Octave:precedence-change` warning is enabled.

`Octave:recursive-path-search`

By default, the `Octave:recursive-path-search` warning is enabled.

`Octave:remove-init-dir`

The path function changes the search path that Octave uses to find functions. It is possible to set the path to a value which excludes Octave's own built-in functions. If the `Octave:remove-init-dir` warning is enabled then Octave will warn when the path function has been used in a way that may render Octave unworkable. By default, the `Octave:remove-init-dir` warning is enabled.

`Octave:reload-forces-clear`

If several functions have been loaded from the same file, Octave must clear all the functions before any one of them can be reloaded. If the `Octave:reload-forces-clear` warning is enabled, Octave will warn you when this happens, and print a list of the additional functions that it is forced to clear. By default, the `Octave:reload-forces-clear` warning is enabled.

`Octave:resize-on-range-error`

If the `Octave:resize-on-range-error` warning is enabled, print a warning when a matrix is resized by an indexed assignment with indices outside the

current bounds. By default, the `Octave:resize-on-range-error` warning is disabled.

`Octave:separator-insert`

Print warning if commas or semicolons might be inserted automatically in literal matrices. By default, the `Octave:separator-insert` warning is disabled.

`Octave:shadowed-function`

By default, the `Octave:shadowed-function` warning is enabled.

`Octave:single-quote-string`

Print warning if a single quote character is used to introduce a string constant. By default, the `Octave:single-quote-string` warning is disabled.

`Octave:sqrtn:SingularMatrix`

By default, the `Octave:sqrtn:SingularMatrix` warning is enabled.

`Octave:str-to-num`

If the `Octave:str-to-num` warning is enabled, a warning is printed for implicit conversions of strings to their numeric UTF-8 encoded byte sequences. For example,

```
"abc" + 0
⇒ 97 98 99
```

elicits a warning if the `Octave:str-to-num` warning is enabled. By default, the `Octave:str-to-num` warning is disabled.

`Octave:variable-switch-label`

If the `Octave:variable-switch-label` warning is enabled, Octave will print a warning if a switch label is not a constant or constant expression. By default, the `Octave:variable-switch-label` warning is disabled.

12.2.2 Enabling and Disabling Warnings

The `warning` function also allows you to control which warnings are actually printed to the screen. If the `warning` function is called with a string argument that is either "on" or "off" all warnings will be enabled or disabled.

It is also possible to enable and disable individual warnings through their string identifications. The following code will issue a warning

```
warning ("example:non-negative-variable",
        "'a' must be non-negative. Setting 'a' to zero.");
```

while the following won't issue a warning

```
warning ("off", "example:non-negative-variable");
warning ("example:non-negative-variable",
        "'a' must be non-negative. Setting 'a' to zero.");
```


13 Debugging

Octave includes a built-in debugger to aid in the development of scripts. This can be used to interrupt the execution of an Octave script at a certain point, or when certain conditions are met. Once execution has stopped, and debug mode is entered, the symbol table at the point where execution has stopped can be examined and modified to check for errors.

The normal command-line editing and history functions are available in debug mode.

13.1 Entering Debug Mode

There are two basic means of interrupting the execution of an Octave script. These are breakpoints (see [Section 13.3 \[Breakpoints\]](#), page 238), discussed in the next section, and interruption based on some condition.

Octave supports three means to stop execution based on the values set in the functions `debug_on_interrupt`, `debug_on_warning`, and `debug_on_error`.

```
val = debug_on_interrupt ()
old_val = debug_on_interrupt (new_val)
debug_on_interrupt (new_val, "local")
```

Query or set the internal variable that controls whether Octave will try to enter debugging mode when it receives an interrupt signal (typically generated with *C-c*).

If a second interrupt signal is received before reaching the debugging mode, a normal interrupt will occur.

When called from inside a function with the "`local`" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[debug_on_error\]](#), page 237, [\[debug_on_warning\]](#), page 237.

```
val = debug_on_warning ()
old_val = debug_on_warning (new_val)
debug_on_warning (new_val, "local")
```

Query or set the internal variable that controls whether Octave will try to enter the debugger when a warning is encountered.

When called from inside a function with the "`local`" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[debug_on_error\]](#), page 237, [\[debug_on_interrupt\]](#), page 237.

```
val = debug_on_error ()
old_val = debug_on_error (new_val)
debug_on_error (new_val, "local")
```

Query or set the internal variable that controls whether Octave will try to enter the debugger when an error is encountered.

This will also inhibit printing of the normal traceback message (you will only see the top-level error message).

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[debug_on_warning\]](#), page 237, [\[debug_on_interrupt\]](#), page 237.

13.2 Leaving Debug Mode

Use either `dbcont` or `return` to leave the debug mode and continue the normal execution of the script.

`dbcont`

Leave command-line debugging mode and continue code execution normally.

See also: [\[dbstep\]](#), page 243, [\[dbquit\]](#), page 238.

To quit debug mode and return directly to the prompt without executing any additional code use `dbquit`.

`dbquit`

Quit debugging mode immediately without further code execution and return to the Octave prompt.

See also: [\[dbcont\]](#), page 238, [\[dbstep\]](#), page 243.

Finally, typing `exit` or `quit` at the debug prompt will result in Octave terminating normally.

13.3 Breakpoints

Breakpoints can be set in any m-file function by using the `dbstop` function.

```
dbstop func
dbstop func line
dbstop func line1 line2 ...
dbstop line1 ...
dbstop in func
dbstop in func at line
dbstop in func at line if "condition"
dbstop if event
dbstop if event ID
dbstop (bp_struct)
rline = dbstop ...
```

Set breakpoints for the built-in debugger.

`func` is the name of a function on the current path. When already in debug mode the `func` argument can be omitted and the current function will be used. Breakpoints at subfunctions are set with the scope operator '`>`'. For example, If `file.m` has a subfunction `func2`, then a breakpoint in `func2` can be specified by `file>func2`.

`line` is the line number at which to break. If `line` is not specified, it defaults to the first executable line in the file `func.m`. Multiple lines can be specified in a single

command; when function syntax is used, the lines may also be passed as a single vector argument (`[line1, line2, ...]`).

condition is any Octave expression that can be evaluated in the code context that exists at the breakpoint. When the breakpoint is encountered, *condition* will be evaluated, and execution will stop if *condition* is true. If *condition* cannot be evaluated, for example because it refers to an undefined variable, an error will be thrown. Expressions with side effects (such as `y++ > 1`) will alter variables, and should generally be avoided. Conditions containing quotes ('"', "'") or comment characters ('#', '%') must be enclosed in quotes. (This does not apply to conditions entered from the editor's context menu.) For example:

```
dbstop in axis at 246 if 'any (opt == "x")'
```

The form specifying event does not cause a specific breakpoint at a given function and line number. Instead it causes debug mode to be entered when certain unexpected events are encountered. Possible values are

error Stop when an error is reported. This is equivalent to specifying both `debug_on_error (true)` and `debug_on_interrupt (true)`.

caught error
Stop when an error is caught by a try-catch block (not yet implemented).

interrupt
Stop when an interrupt (*Ctrl-C*) occurs.

naninf Stop when code returns a non-finite value (not yet implemented).

warning Stop when a warning is reported. This is equivalent to specifying `debug_on_warning (true)`.

The events **error**, **caught error**, and **warning** can all be followed by a string specifying an error ID or warning ID. If that is done, only errors with the specified ID will cause execution to stop. To stop on one of a set of IDs, multiple `dbstop` commands must be issued.

Breakpoints and events can be removed using the `dbclear` command with the same syntax.

It is possible to save all breakpoints and restore them at once by issuing the commands `bp_state = dbstatus; ...; dbstop (bp_state)`.

The optional output *rline* is the real line number where the breakpoint was set. This can differ from the specified line if the line is not executable. For example, if a breakpoint attempted on a blank line then Octave will set the real breakpoint at the next executable line.

When a file is re-parsed, such as when it is modified outside the GUI, all breakpoints within the file are cleared.

See also: `[dbclear]`, page 240, `[dbstatus]`, page 240, `[dbstep]`, page 243, `[debug_on_error]`, page 237, `[debug_on_warning]`, page 237, `[debug_on_interrupt]`, page 237.

Breakpoints in class methods are also supported (e.g., `dbstop ("@class/method")`). However, breakpoints cannot be set in built-in functions (e.g., `sin`, etc.) or dynamically loaded functions (i.e., oct-files).

To set a breakpoint immediately upon entering a function use line number 1, or omit the line number entirely and just give the function name. When setting the breakpoint Octave will ignore the leading comment block, and the breakpoint will be set on the first executable statement in the function. For example:

```
dbstop ("asind", 1)
⇒ 29
```

Note that the return value of 29 means that the breakpoint was effectively set to line 29. The status of breakpoints in a function can be queried with `dbstatus`.

`dbstatus`

`dbstatus func`

`bp_list = dbstatus ...`

Report the location of active breakpoints.

When called with no input or output arguments, print the list of all functions with breakpoints and the line numbers where those breakpoints are set.

If a function name *func* is specified then only report breakpoints for the named function and its subfunctions.

The optional return argument *bp_list* is a struct array with the following fields.

<code>name</code>	The name of the function with a breakpoint. A subfunction, say <code>func2</code> within an m-file, say <code>file.m</code> , is specified as <code>file>func2</code> .
<code>file</code>	The name of the m-file where the function code is located.
<code>line</code>	The line number with the breakpoint.
<code>cond</code>	The condition that must be satisfied for the breakpoint to be active, or the empty string for unconditional breakpoints.

If `dbstop if error` is true but no explicit IDs are specified, the return value will have an empty field called "`errs`". If IDs are specified, the `errs` field will have one row per ID. If `dbstop if error` is false, there is no "`errs`" field. The "`warn`" field is set similarly by `dbstop if warning`.

See also: [\[dbstop\]](#), page 238, [\[dbclear\]](#), page 240, [\[dbwhere\]](#), page 242, [\[dblist\]](#), page 242, [\[dbstack\]](#), page 243.

Reusing the previous example, `dbstatus ("asind")` will return 29. The breakpoints listed can then be cleared with the `dbclear` function.

```
dbclear func
dbclear func line
dbclear func line1 line2 ...
dbclear line ...
dbclear all
dbclear in func
dbclear in func at line
```

```
dbclear if event
dbclear ("func")
dbclear ("func", line)
dbclear ("func", line1, line2, ...)
dbclear ("func", line1, ...)
dbclear (line, ...)
dbclear ("all")
```

Delete a breakpoint at line number *line* in the function *func*.

Arguments are

func Function name as a string variable. When already in debug mode this argument can be omitted and the current function will be used.

line Line number from which to remove a breakpoint. Multiple lines may be given as separate arguments or as a vector.

event An event such as `error`, `interrupt`, or `warning` (see [dbstop], page 238, for details).

When called without a line number specification all breakpoints in the named function are cleared.

If the requested line is not a breakpoint no action is performed.

The special keyword "all" will clear all breakpoints from all files.

See also: [dbstop], page 238, [dbstatus], page 240, [dbwhere], page 242.

A breakpoint may also be set in a subfunction. For example, if a file contains the functions

```
function y = func1 (x)
    y = func2 (x);
endfunction
function y = func2 (x)
    y = x + 1;
endfunction
```

then a breakpoint can be set at the start of the subfunction directly with

```
dbstop func1>func2
⇒ 5
```

Note that ‘>’ is the character that distinguishes subfunctions from the m-file containing them.

Another simple way of setting a breakpoint in an Octave script is the use of the `keyboard` function.

```
keyboard ()
keyboard ("prompt")
```

Stop m-file execution and enter debug mode.

When the `keyboard` function is executed, Octave prints a prompt and waits for user input. The input strings are then evaluated and the results are printed. This makes it possible to examine the values of variables within a function, and to assign new values

if necessary. To leave the prompt and return to normal execution type ‘`return`’ or ‘`dbcont`’. The `keyboard` function does not return an exit status.

If `keyboard` is invoked without arguments, a default prompt of ‘`debug>`’ is used.

See also: `[dbstop]`, page 238, `[dbcont]`, page 238, `[dbquit]`, page 238.

The `keyboard` function is placed in a script at the point where the user desires that the execution be stopped. It automatically sets the running script into the debug mode.

13.4 Debug Mode

There are three additional support functions that allow the user to find out where in the execution of a script Octave entered the debug mode, and to print the code in the script surrounding the point where Octave entered debug mode.

`dbwhere`

In debugging mode, report the current file and line number where execution is stopped.

See also: `[dbstack]`, page 243, `[dblist]`, page 242, `[dbstatus]`, page 240, `[dbcont]`, page 238, `[dbstep]`, page 243, `[dbup]`, page 244, `[dbdown]`, page 244.

```
dbtype
dbtype lineno
dbtype startl:endl
dbtype startl:end
dbtype func
dbtype func lineno
dbtype func startl:endl
dbtype func startl:end
```

Display a script file with line numbers.

When called with no arguments in debugging mode, display the script file currently being debugged.

An optional range specification can be used to list only a portion of the file. The special keyword “`end`” is a valid line number specification for the last line of the file.

When called with the name of a function, list that script file with line numbers.

See also: `[dblist]`, page 242, `[dbwhere]`, page 242, `[dbstatus]`, page 240, `[dbstop]`, page 238.

```
dblist
dblist n
```

In debugging mode, list `n` lines of the function being debugged centered around the current line to be executed.

If unspecified `n` defaults to 10 (+/- 5 lines)

See also: `[dbwhere]`, page 242, `[dbtype]`, page 242, `[dbstack]`, page 243.

You may also use `isdebugmode` to determine whether the debugger is currently active.

isdebugmode ()

Return true if in debugging mode, otherwise false.

See also: [dbwhere], page 242, [dbstack], page 243, [dbstatus], page 240.

Debug mode also allows single line stepping through a function using the command **dbstep**.

dbstep
dbstep n
dbstep in
dbstep out
dbnext ...

In debugging mode, execute the next *n* lines of code.

If *n* is omitted, execute the next single line of code. If the next line of code is itself defined in terms of an m-file remain in the existing function.

Using **dbstep in** will cause execution of the next line to step into any m-files defined on the next line.

Using **dbstep out** will cause execution to continue until the current function returns.
dbnext is an alias for **dbstep**.

See also: [dbcont], page 238, [dbquit], page 238.

When in debug mode the RETURN key will execute the last entered command. This is useful, for example, after hitting a breakpoint and entering **dbstep** once. After that, one can advance line by line through the code with only a single key stroke.

13.5 Call Stack

The function being debugged may be the leaf node of a series of function calls. After examining values in the current subroutine it may turn out that the problem occurred in earlier pieces of code. Use **dbup** and **dbdown** to move up and down through the series of function calls to locate where variables first took on the wrong values. **dbstack** shows the entire series of function calls and at what level debugging is currently taking place.

dbstack
dbstack n
dbstack -completenames
[stack, idx] = dbstack (...)

Display or return current debugging function stack information.

With optional argument *n*, omit the *n* innermost stack frames.

Although accepted, the argument **-completenames** is silently ignored. Octave always returns absolute filenames.

The arguments *n* and **-completenames** can be both specified in any order.

The optional return argument *stack* is a struct array with the following fields:

file The name of the m-file where the function code is located.

name The name of the function with a breakpoint.

line	The line number of an active breakpoint.
column	The column number of the line where the breakpoint begins.
scope	Undocumented.
context	Undocumented.

The return argument *idx* specifies which element of the *stack* struct array is currently active.

See also: [\[dbup\]](#), page 244, [\[dbdown\]](#), page 244, [\[dbwhere\]](#), page 242, [\[dblist\]](#), page 242, [\[dbstatus\]](#), page 240.

dbup
dbup *n*
 In debugging mode, move up the execution stack *n* frames.
 If *n* is omitted, move up one frame.
 See also: [\[dbstack\]](#), page 243, [\[dbdown\]](#), page 244.

dbdown
dbdown *n*
 In debugging mode, move down the execution stack *n* frames.
 If *n* is omitted, move down one frame.
 See also: [\[dbstack\]](#), page 243, [\[dbup\]](#), page 244.

13.6 Profiling

Octave supports profiling of code execution on a per-function level. If profiling is enabled, each call to a function (supporting built-ins, operators, functions in oct- and mex-files, user-defined functions in Octave code and anonymous functions) is recorded while running Octave code. After that, this data can aid in analyzing the code behavior, and is in particular helpful for finding “hot spots” in the code which use up a lot of computation time and are the best targets to spend optimization efforts on.

The main command for profiling is `profile`, which can be used to start or stop the profiler and also to query collected data afterwards. The data is returned in an Octave data structure which can then be examined or further processed by other routines or tools.

```
profile on
profile off
profile resume
profile clear
S = profile ("status")
T = profile ("info")
    Control the built-in profiler.

profile on
    Start the profiler, clearing all previously collected data if there is any.

profile off
    Stop profiling. The collected data can later be retrieved and examined
    with T = profile ("info").
```

```
profile clear
    Clear all collected profiler data.

profile resume
    Restart profiling without clearing the old data. All newly collected statistics are added to the existing ones.

S = profile ("status")
    Return a structure with information about the current status of the profiler. At the moment, the only field is ProfilerStatus which is either "on" or "off".

T = profile ("info")
    Return the collected profiling statistics in the structure T. The flat profile is returned in the field FunctionTable which is an array of structures, each entry corresponding to a function which was called and for which profiling statistics are present. In addition, the field Hierarchical contains the hierarchical call tree. Each node has an index into the FunctionTable identifying the function it corresponds to as well as data fields for number of calls and time spent at this level in the call tree.

See also: [profshow], page 245, [profexplore], page 246.
```

An easy way to get an overview over the collected data is **profshow**. This function takes the profiler data returned by **profile** as input and prints a flat profile, for instance:

Function	Attr	Time (s)	Calls
<hr/>			
>myfib	R	2.195	13529
binary <=		0.061	13529
binary -		0.050	13528
binary +		0.026	6764

This shows that most of the run time was spent executing the function ‘myfib’, and some minor proportion evaluating the listed binary operators. Furthermore, it is shown how often the function was called and the profiler also records that it is recursive.

```
profshow (data)
profshow (data, n)
profshow ()
profshow (n)
```

Display flat per-function profiler results.

Print out profiler data (execution time, number of calls) for the most critical *n* functions. The results are sorted in descending order by the total time spent in each function. If *n* is unspecified it defaults to 20.

The input *data* is the structure returned by **profile ("info")**. If unspecified, **profshow** will use the current profile dataset.

The attribute column displays ‘R’ for recursive functions, and is blank for all other function types.

See also: [profexplore], page 246, [profile], page 244.

```
profexport (dir)
profexport (dir, data)
profexport (dir, name)
profexport (dir, name, data)
```

Export profiler data as HTML.

Export the profiling data in *data* into a series of HTML files in the folder *dir*. The initial file will be *data/index.html*.

If *name* is specified, it must be a string that contains a “name” for the profile being exported. This name is included in the HTML.

The input *data* is the structure returned by `profile ("info")`. If unspecified, `profexport` will use the current profile dataset.

See also: `[profshow]`, page 245, `[profexplore]`, page 246, `[profile]`, page 244.

```
profexplore ()
profexplore (data)
```

Interactively explore hierarchical profiler output.

Assuming *data* is the structure with profile data returned by `profile ("info")`, this command opens an interactive prompt that can be used to explore the call-tree. Type `help` to get a list of possible commands. If *data* is omitted, `profile ("info")` is called and used in its place.

See also: `[profile]`, page 244, `[profshow]`, page 245.

13.7 Profiler Example

Below, we will give a short example of a profiler session. See [Section 13.6 \[Profiling\]](#), page 244, for the documentation of the profiler functions in detail. Consider the code:

```
global N A;

N = 300;
A = rand (N, N);

function xt = timesteps (steps, x0, expM)
    global N;

    if (steps == 0)
        xt = NA (N, 0);
    else
        xt = NA (N, steps);
        x1 = expM * x0;
        xt(:, 1) = x1;
        xt(:, 2 : end) = timesteps (steps - 1, x1, expM);
    endif
endfunction

function foo ()
    global N A;
```

```

initial = @(x) sin (x);
x0 = (initial (linspace (0, 2 * pi, N)))';

expA = expm (A);
xt = timesteps (100, x0, expA);
endfunction

function fib = bar (N)
if (N <= 2)
    fib = 1;
else
    fib = bar (N - 1) + bar (N - 2);
endif
endfunction

```

If we execute the two main functions, we get:

```

tic; foo; toc;
⇒ Elapsed time is 2.37338 seconds.

```

```

tic; bar (20); toc;
⇒ Elapsed time is 2.04952 seconds.

```

But this does not give much information about where this time is spent; for instance, whether the single call to `expm` is more expensive or the recursive time-stepping itself. To get a more detailed picture, we can use the profiler.

```

profile on;
foo;
profile off;

data = profile ("info");
profshow (data, 10);

```

This prints a table like:

#	Function Attr	Time (s)	Calls
7	expm	1.034	1
3	binary *	0.823	117
41	binary \	0.188	1
38	binary ^	0.126	2
43	timesteps R	0.111	101
44	NA	0.029	101
39	binary +	0.024	8
34	norm	0.011	1
40	binary -	0.004	101
33	balance	0.003	1

The entries are the individual functions which have been executed (only the 10 most important ones), together with some information for each of them. The entries like ‘binary

*' denote operators, while other entries are ordinary functions. They include both built-ins like `expm` and our own routines (for instance `timesteps`). From this profile, we can immediately deduce that `expm` uses up the largest proportion of the processing time, even though it is only called once. The second expensive operation is the matrix-vector product in the routine `timesteps`.¹

Timing, however, is not the only information available from the profile. The attribute column shows us that `timesteps` calls itself recursively. This may not be that remarkable in this example (since it's clear anyway), but could be helpful in a more complex setting. As to the question of why is there a 'binary \` in the output, we can easily shed some light on that too. Note that `data` is a structure array ([Section 6.1.2 \[Structure Arrays\], page 105](#)) which contains the field `FunctionTable`. This stores the raw data for the profile shown. The number in the first column of the table gives the index under which the shown function can be found there. Looking up `data.FunctionTable(41)` gives:

```
scalar structure containing the fields:
```

```
FunctionName = binary \
TotalTime = 0.18765
NumCalls = 1
IsRecursive = 0
Parents = 7
Children = [] (1x0)
```

Here we see the information from the table again, but have additional fields `Parents` and `Children`. Those are both arrays, which contain the indices of functions which have directly called the function in question (which is entry 7, `expm`, in this case) or been called by it (no functions). Hence, the backslash operator has been used internally by `expm`.

Now let's take a look at `bar`. For this, we start a fresh profiling session (`profile on` does this; the old data is removed before the profiler is restarted):

```
profile on;
bar (20);
profile off;

profshow (profile ("info"));
```

This gives:

¹ We only know it is the binary multiplication operator, but fortunately this operator appears only at one place in the code and thus we know which occurrence takes so much time. If there were multiple places, we would have to use the hierarchical profile to find out the exact place which uses up the time which is not covered in this example.

#	Function	Attr	Time (s)	Calls
<hr/>				
1	bar	R	2.091	13529
2	binary <=		0.062	13529
3	binary -		0.042	13528
4	binary +		0.023	6764
5	profile		0.000	1
8	false		0.000	1
6	nargin		0.000	1
7	binary !=		0.000	1
9	__profiler_enable__		0.000	1

Unsurprisingly, `bar` is also recursive. It has been called 13,529 times in the course of recursively calculating the Fibonacci number in a suboptimal way, and most of the time was spent in `bar` itself.

Finally, let's say we want to profile the execution of both `foo` and `bar` together. Since we already have the run-time data collected for `bar`, we can restart the profiler without clearing the existing data and collect the missing statistics about `foo`. This is done by:

```
profile resume;
foo;
profile off;

profshow (profile ("info"), 10);
```

As you can see in the table below, now we have both profiles mixed together.

#	Function	Attr	Time (s)	Calls
<hr/>				
1	bar	R	2.091	13529
16	expm		1.122	1
12	binary *		0.798	117
46	binary \		0.185	1
45	binary ^		0.124	2
48	timesteps	R	0.115	101
2	binary <=		0.062	13529
3	binary -		0.045	13629
4	binary +		0.041	6772
49	NA		0.036	101

14 Input and Output

Octave supports several ways of reading and writing data to or from the prompt or a file. The simplest functions for data Input and Output (I/O) are easy to use, but only provide limited control of how data is processed. For more control, a set of functions modeled after the C standard library are also provided by Octave.

14.1 Basic Input and Output

14.1.1 Terminal Output

Since Octave normally prints the value of an expression as soon as it has been evaluated, the simplest of all I/O functions is a simple expression. For example, the following expression will display the value of ‘pi’

```
pi
→ pi = 3.1416
```

This works well as long as it is acceptable to have the name of the variable (or ‘ans’) printed along with the value. To print the value of a variable without printing its name, use the function `disp`.

The `format` command offers some control over the way Octave prints values with `disp` and through the normal echoing mechanism.

```
disp (x)
str = disp (x)
Display the value of x.
```

For example:

```
disp ("The value of pi is:"), disp (pi)

→ the value of pi is:
→ 3.1416
```

Note that the output from `disp` always ends with a newline.

If an output value is requested, `disp` prints nothing and returns the formatted output in a string.

See also: [\[fdisp\]](#), page 263.

```
list_in_columns (arg, width, prefix)
```

Return a string containing the elements of `arg` listed in columns with an overall maximum width of `width` and optional prefix `prefix`.

The argument `arg` must be a cell array of character strings or a character array.

If `width` is not specified or is an empty matrix, or less than or equal to zero, the width of the terminal screen is used. Newline characters are used to break the lines in the output string. For example:

```

list_in_columns ({"abc", "def", "ghijkl", "mnop", "qrs", "tuv"}, 20)
⇒ abc      mnop
    def      qrs
    ghijkl  tuv

whos ans
⇒
Variables in the current scope:

  Attr Name      Size          Bytes  Class
  ===== ========= ====== =====
    ans        1x37           37   char

Total is 37 elements using 37 bytes

```

See also: [terminal_size], page 252.

terminal_size ()

Return a two-element row vector containing the current size of the terminal window in characters (rows and columns).

See also: [list_in_columns], page 251.

format

format options

[format, formatspacing] = format

Reset or specify the format of the output produced by `disp` and Octave's normal echoing mechanism.

This command only affects the display of numbers, but not how they are stored or computed. To change the internal representation from the default double use one of the conversion functions such as `single`, `uint8`, `int64`, etc.

By default, Octave displays 5 significant digits in a human readable form (option '`short`' paired with '`loose`' format for matrices). If `format` is invoked without any options, this default format is restored.

Valid formats for floating point numbers are listed in the following table.

`short` Fixed point format with 5 significant figures (default).

`long` Fixed point format with 16 significant figures.

As with the '`short`' format, Octave will switch to an exponential '`e`' format if it is unable to format a matrix properly using the current format.

`short e`

`long e` Exponential format. The number to be represented is split between a mantissa and an exponent (power of 10). The mantissa has 5 significant digits in the short format. In the long format, double values are displayed with 16 significant digits and single values are displayed with 8. For example, with the '`short e`' format, `pi` is displayed as `3.1416e+00`.

`short E`

`long E` Identical to '`short e`' or '`long e`' but displays an uppercase '`E`' to indicate the exponent. For example, with the '`long E`' format, `pi` is displayed as `3.141592653589793E+00`.

<code>short g</code>	
<code>long g</code>	Optimally choose between fixed point and exponential format based on the magnitude of the number. For example, with the ‘ <code>short g</code> ’ format, <code>pi . ^ [2; 4; 8; 16; 32]</code> is displayed as
	<code>ans =</code>
	9.8696
	97.409
	9488.5
	9.0032e+07
	8.1058e+15
<code>short eng</code>	
<code>long eng</code>	Identical to ‘ <code>short e</code> ’ or ‘ <code>long e</code> ’ but displays the value using an engineering format, where the exponent is divisible by 3. For example, with the ‘ <code>short eng</code> ’ format, <code>10 * pi</code> is displayed as <code>31.416e+00</code> .
<code>long G</code>	
<code>short G</code>	Identical to ‘ <code>short g</code> ’ or ‘ <code>long g</code> ’ but displays an uppercase ‘E’ to indicate the exponent.
<code>free</code>	
<code>none</code>	Print output in free format, without trying to line up columns of matrices on the decimal point. This is a raw format equivalent to the C++ code <code>std::cout << variable</code> . In general, the result is a presentation with 6 significant digits where unnecessary precision (such as trailing zeros for integers) is suppressed. Complex numbers are formatted as numeric pairs like this ‘ <code>(0.60419, 0.60709)</code> ’ instead of like this ‘ <code>0.60419 + 0.60709i</code> ’.

The following formats affect all numeric output (floating point and integer types).

<code>"+"</code>	
<code>"+" "chars"</code>	
<code>plus</code>	
<code>plus chars</code>	Print a ‘+’ symbol for matrix elements greater than zero, a ‘-’ symbol for elements less than zero, and a space for zero matrix elements. This format can be useful for examining the sparsity structure of a large matrix. For very large matrices the function <code>spy</code> which plots the sparsity pattern will be clearer.

The optional argument `chars` specifies a list of 3 characters to use for printing values greater than zero, less than zero, and equal to zero. For example, with the format “`+" "+-.`”, the matrix `[1, 0, -1; -1, 0, 1]` is displayed as

`ans =`

```
+ . -
- . +
```

bank	Print variable in a format appropriate for a currency (fixed format with two digits to the right of the decimal point). Only the real part of a variable is displayed, as the imaginary part makes no sense for a currency.
native-hex	Print the hexadecimal representation of numbers as they are stored in memory. For example, on a workstation which stores 8 byte real values in IEEE format with the least significant byte first, the value of <code>pi</code> when printed in <code>native-hex</code> format is <code>400921fb54442d18</code> .
hex	The same as <code>native-hex</code> , but always print the most significant byte first.
native-bit	Print the bit representation of numbers as stored in memory. For example, the value of <code>pi</code> is <pre>0100000000001001001000011111011 01010100010001000010110100011000</pre> (shown here in two 32 bit sections for typesetting purposes) when printed in native-bit format on a workstation which stores 8 byte real values in IEEE format with the least significant byte first.
bit	The same as <code>native-bit</code> , but always print the most significant bits first.
rat	Print a rational approximation, i.e., values are approximated as the ratio of small integers. For example, with the ‘ <code>rat</code> ’ format, <code>pi</code> is displayed as <code>355/113</code> .

The following two options affect the display of all matrices.

compact	Remove blank lines around column number labels and between matrices producing more compact output with more data per page.
loose	Insert blank lines above and below column number labels and between matrices to produce a more readable output with less data per page. (default).

If called with one or two output arguments, and no inputs, return the current format and format spacing.

See also: [\[fixed_point_format\]](#), page 50, [\[output_precision\]](#), page 49, [\[split_long_rows\]](#), page 50, [\[print_empty_dimensions\]](#), page 51, [\[rats\]](#), page 538.

14.1.1.1 Paging Screen Output

When running interactively, Octave normally sends all output directly to the Command Window. However, when using the CLI version of Octave this can create a problem because large volumes of data will stream by before you can read them. In such cases, it is better to use a paging program such as `less` or `more` which displays just one screenful at a time. With `less` (and some versions of `more`) you can also scan forward and backward, and search for specific items. The pager is enabled by the command `more on`.

Normally, no output is displayed by the pager until just before Octave is ready to print the top level prompt, or read from the standard input (for example, by using the `fscanf` or `scanf` functions). This means that there may be some delay before any output appears on your screen if you have asked Octave to perform a significant amount of work with a single

command statement. The function `fflush` may be used to force output to be sent to the pager (or any other stream) immediately.

You can select the program to run as the pager with the `PAGER` function, and configure the pager itself with the `PAGER_FLAGS` function.

```
more  
more on  
more off
```

Turn output pagination on or off.

Without an argument, `more` toggles the current state.

The current state can be determined via `page_screen_output`.

See also: [\[page_screen_output\]](#), page 255, [\[page_output_immediately\]](#), page 256, [\[PAGER\]](#), page 255, [\[PAGER_FLAGS\]](#), page 255.

```
val = PAGER ()  
old_val = PAGER (new_val)  
PAGER (new_val, "local")
```

Query or set the internal variable that specifies the program to use to display terminal output on your system.

The default value is normally "less", "more", or "pg", depending on what programs are installed on your system. See [Appendix E \[Installation\]](#), page 1021.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[PAGER_FLAGS\]](#), page 255, [\[page_output_immediately\]](#), page 256, [\[more\]](#), page 255, [\[page_screen_output\]](#), page 255.

```
val = PAGER_FLAGS ()  
old_val = PAGER_FLAGS (new_val)  
PAGER_FLAGS (new_val, "local")
```

Query or set the internal variable that specifies the options to pass to the pager.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[PAGER\]](#), page 255, [\[more\]](#), page 255, [\[page_screen_output\]](#), page 255, [\[page_output_immediately\]](#), page 256.

```
val = page_screen_output ()  
old_val = page_screen_output (new_val)  
page_screen_output (new_val, "local")
```

Query or set the internal variable that controls whether output intended for the terminal window that is longer than one page is sent through a pager.

This allows you to view one screenful at a time. Some pagers (such as `less`—see [Appendix E \[Installation\]](#), page 1021) are also capable of moving backward on the output.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [more], page 255, [page_output_immediately], page 256, [PAGER], page 255, [PAGER_FLAGS], page 255.

```
val = page_output_immediately ()
old_val = page_output_immediately (new_val)
page_output_immediately (new_val, "local")
```

Query or set the internal variable that controls whether Octave sends output to the pager as soon as it is available.

Otherwise, Octave buffers its output and waits until just before the prompt is printed to flush it to the pager.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [page_screen_output], page 255, [more], page 255, [PAGER], page 255, [PAGER_FLAGS], page 255.

fflush (*fid*)

Flush output to file descriptor *fid*.

fflush returns 0 on success and an OS dependent error value (-1 on Unix) on error.

Programming Note: Flushing is useful for ensuring that all pending output makes it to the screen before some other event occurs. For example, it is always a good idea to flush the standard output stream before calling **input**.

See also: [fopen], page 273, [fclose], page 275.

14.1.2 Terminal Input

Octave has three functions that make it easy to prompt users for input. The **input** and **menu** functions are normally used for managing an interactive dialog with a user, and the **keyboard** function is normally used for doing simple debugging.

```
ans = input (prompt)
ans = input (prompt, "s")
```

Print *prompt* and wait for user input.

For example,

```
input ("Pick a number, any number! ")
```

prints the prompt

```
Pick a number, any number!
```

and waits for the user to enter a value. The string entered by the user is evaluated as an expression, so it may be a literal constant, a variable name, or any other valid Octave code.

The number of return arguments, their size, and their class depend on the expression entered.

If you are only interested in getting a literal string value, you can call `input` with the character string "`s`" as the second argument. This tells Octave to return the string entered by the user directly, without evaluating it first.

Because there may be output waiting to be displayed by the pager, it is a good idea to always call `fflush (stdout)` before calling `input`. This will ensure that all pending output is written to the screen before your prompt.

See also: [\[yes_or_no\]](#), page 257, [\[kbhit\]](#), page 257, [\[pause\]](#), page 891, [\[menu\]](#), page 257, [\[listdlg\]](#), page 864.

```
choice = menu (title, opt1, ...)
choice = menu (title, {opt1, ...})
```

Display a menu with heading `title` and options `opt1, ...`, and wait for user input.

If the GUI is running, the menu is displayed graphically using `listdlg`. Otherwise, the title and menu options are printed on the console.

`title` is a string and the options may be input as individual strings or as a cell array of strings.

The return value `choice` is the number of the option selected by the user counting from 1. If the user aborts the dialog or makes an invalid selection then 0 is returned.

This function is useful for interactive programs. There is no limit to the number of options that may be passed in, but it may be confusing to present more than will fit easily on one screen.

See also: [\[input\]](#), page 256, [\[listdlg\]](#), page 864.

```
ans = yes_or_no ("prompt")
```

Ask the user a yes-or-no question.

Return logical true if the answer is yes or false if the answer is no.

Takes one argument, `prompt`, which is the string to display when asking the question. `prompt` should end in a space; `yes-or-no` adds the string '(yes or no)' to it. The user must confirm the answer with RET and can edit it until it has been confirmed.

See also: [\[input\]](#), page 256.

For `input`, the normal command line history and editing functions are available at the prompt.

Octave also has a function that makes it possible to get a single character from the keyboard without requiring the user to type a carriage return.

```
kbhit ()  
kbhit (1)
```

Read a single keystroke from the keyboard.

If called with an argument, don't wait for a keypress.

For example,

```
x = kbhit ();
```

will set `x` to the next character typed at the keyboard as soon as it is typed.

```
x = kbhit (1);
```

is identical to the above example, but doesn't wait for a keypress, returning the empty string if no key is available.

See also: [input], page 256, [pause], page 891.

14.1.3 Simple File I/O

The **save** and **load** commands allow data to be written to and read from disk files in various formats. The default format of files written by the **save** command can be controlled using the functions **save_default_options** and **save_precision**.

As an example the following code creates a 3-by-3 matrix and saves it to the file ‘myfile.mat’.

```
A = [ 1:3; 4:6; 7:9 ];
save myfile.mat A
```

Once one or more variables have been saved to a file, they can be read into memory using the **load** command.

```
load myfile.mat
A
└ A =
  └
  └
  └ 1   2   3
  └ 4   5   6
  └ 7   8   9

save file
save options file
save options file v1 v2 ...
save options file -struct STRUCT
save options file -struct STRUCT f1 f2 ...
save - v1 v2 ...
str = save ("-", "v1", "v2", ...)
Save the named variables v1, v2, ..., in the file file.
```

The special filename ‘-’ may be used to return the content of the variables as a string. If no variable names are listed, Octave saves all the variables in the current scope. Otherwise, full variable names or pattern syntax can be used to specify the variables to save. If the **-struct** modifier is used then the fields of the **scalar** struct are saved as if they were variables with the corresponding field names. The **-struct** option can be combined with specific field names *f1*, *f2*, ... to write only certain fields to the file.

Valid options for the **save** command are listed in the following table. Options that modify the output format override the format specified by **save_default_options**.

If **save** is invoked using the functional form

```
save ("-option1", ..., "file", "v1", ...)
```

then the *options*, *file*, and variable name arguments (*v1*, ...) must be specified as character strings.

If called with a filename of "-", write the output to stdout if nargout is 0, otherwise return the output in a character string.

- append Append to the destination instead of overwriting.
- ascii Save a matrix in a text file without a header or any other information. The matrix must be 2-D and only the real part of any complex value is written to the file. Numbers are stored in single-precision format and separated by spaces. Additional options for the -ascii format are
 - double Store numbers in double-precision format.
 - tabs Separate numbers with tabs.
- binary Save the data in Octave's binary data format.
- float-binary Save the data in Octave's binary data format but using only single precision. Use this format **only** if you know that all the values to be saved can be represented in single precision.
- hdf5 Save the data in HDF5 format. (HDF5 is a free, portable, binary format developed by the National Center for Supercomputing Applications at the University of Illinois.) This format is only available if Octave was built with a link to the HDF5 libraries.
- float-hdf5 Save the data in HDF5 format but using only single precision. Use this format **only** if you know that all the values to be saved can be represented in single precision.
- V7
- v7
- 7
- mat7-binary Save the data in MATLAB's v7 binary data format.
- V6
- v6
- 6
- mat
- mat-binary Save the data in MATLAB's v6 binary data format.
- V4
- v4
- 4
- mat4-binary Save the data in the binary format written by MATLAB version 4.
- text Save the data in Octave's text data format. (default).
- zip
- z Use the gzip algorithm to compress the file. This works on files that are compressed with gzip outside of Octave, and gzip can also be used to

convert the files for backward compatibility. This option is only available if Octave was built with a link to the zlib libraries.

The list of variables to save may use wildcard patterns containing the following special characters:

- ? Match any single character.
 - * Match zero or more characters.
 - [*list*] Match the list of characters specified by *list*. If the first character is ! or ^, match all characters except those specified by *list*. For example, the pattern [a-zA-Z] will match all lower and uppercase alphabetic characters.
- Wildcards may also be used in the field name specifications when using the **-struct** modifier (but not in the struct name itself).

Except when using the MATLAB binary data file format or the ‘-ascii’ format, saving global variables also saves the global status of the variable. If the variable is restored at a later time using ‘load’, it will be restored as a global variable.

The command

```
save -binary data a b*
```

saves the variable ‘a’ and all variables beginning with ‘b’ to the file **data** in Octave’s binary format.

See also: [\[load\]](#), page 261, [\[save_default_options\]](#), page 260, [\[save_header_format_string\]](#), [page 261](#), [\[save_precision\]](#), page 260, [\[dlmread\]](#), page 264, [\[csvread\]](#), page 265, [\[fread\]](#), [page 286](#).

There are three functions that modify the behavior of **save**.

```
val = save_default_options ()
old_val = save_default_options (new_val)
save_default_options (new_val, "local")
```

Query or set the internal variable that specifies the default options for the **save** command, and defines the default format.

The default value is “-text” (Octave’s own text-based file format). See the documentation of the **save** command for other choices.

When called from inside a function with the “local” option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[save\]](#), page 258, [\[save_header_format_string\]](#), page 261, [\[save_precision\]](#), [page 260](#).

```
val = save_precision ()
old_val = save_precision (new_val)
save_precision (new_val, "local")
```

Query or set the internal variable that specifies the number of digits to keep when saving data in text format.

The default value is 17 which is the minimum necessary for the lossless saving and restoring of IEEE-754 double values; For IEEE-754 single values the minimum value is 9. If file size is a concern, it is probably better to choose a binary format for saving data rather than to reduce the precision of the saved values.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [save_default_options], page 260.

```
val = save_header_format_string ()
old_val = save_header_format_string (new_val)
save_header_format_string (new_val, "local")
```

Query or set the internal variable that specifies the format string used for the comment line written at the beginning of text-format data files saved by Octave.

The format string is passed to `strftime` and must begin with the character '#' and contain no newline characters. If the value of `save_header_format_string` is the empty string, the header comment is omitted from text-format data files. The default value is

```
"# Created by Octave VERSION, %a %b %d %H:%M:%S %Y %Z <USER@HOST>"
```

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [strftime], page 887, [save_default_options], page 260.

```
load file
load options file
load options file v1 v2 ...
S = load ("options", "file", "v1", "v2", ...)
load file options
load file options v1 v2 ...
S = load ("file", "options", "v1", "v2", ...)
Load the named variables v1, v2, ..., from the file file.
```

If no variables are specified then all variables found in the file will be loaded. As with `save`, the list of variables to extract can be full names or use a pattern syntax. The format of the file is automatically detected but may be overridden by supplying the appropriate option.

If `load` is invoked using the functional form

```
load ("-option1", ..., "file", "v1", ...)
```

then the `options`, `file`, and variable name arguments (`v1`, ...) must be specified as character strings.

If a variable that is not marked as global is loaded from a file when a global symbol with the same name already exists, it is loaded in the global symbol table. Also, if a variable is marked as global in a file and a local symbol exists, the local symbol is moved to the global symbol table and given the value from the file.

If invoked with a single output argument, Octave returns data instead of inserting variables in the symbol table. If the data file contains only numbers (TAB- or space-delimited columns), a matrix of values is returned. Otherwise, `load` returns a structure with members corresponding to the names of the variables in the file.

The `load` command can read data stored in Octave's text and binary formats, and MATLAB's binary format. If compiled with zlib support, it can also load gzip-compressed files. It will automatically detect the type of file and do conversion from different floating point formats (currently only IEEE big and little endian, though other formats may be added in the future).

Valid options for `load` are listed in the following table.

-force	This option is accepted for backward compatibility but is ignored. Octave now overwrites variables currently in memory with those of the same name found in the file.
-ascii	Force Octave to assume the file contains columns of numbers in text format without any header or other information. Data in the file will be loaded as a single numeric matrix with the name of the variable derived from the name of the file.
-binary	Force Octave to assume the file is in Octave's binary format.
-hdf5	Force Octave to assume the file is in HDF5 format. (HDF5 is a free, portable binary format developed by the National Center for Supercomputing Applications at the University of Illinois.) Note that Octave can read HDF5 files not created by itself, but may skip some datasets in formats that it cannot support. This format is only available if Octave was built with a link to the HDF5 libraries.
-import	This option is accepted for backward compatibility but is ignored. Octave can now support multi-dimensional HDF data and automatically modifies variable names if they are invalid Octave identifiers.
-mat	
-mat-binary	
-6	
-v6	
-7	
-v7	Force Octave to assume the file is in MATLAB's version 6 or 7 binary format.
-mat4-binary	
-4	
-v4	
-v4	Force Octave to assume the file is in the binary format written by MATLAB version 4.
-text	Force Octave to assume the file is in Octave's text format.

See also: [\[save\]](#), page 258, [\[dlmwrite\]](#), page 263, [\[csvwrite\]](#), page 264, [\[fwrite\]](#), page 289.

str = fileread (filename)

Read the contents of *filename* and return it as a string.

See also: [fread], page 286, [fscanf], page 283, [importdata], page 270, [textscan], page 266, [type], page 137.

native_float_format ()

Return the native floating point format as a string.

It is possible to write data to a file in a similar way to the **disp** function for writing data to the screen. The **fdisp** works just like **disp** except its first argument is a file pointer as created by **fopen**. As an example, the following code writes to data ‘myfile.txt’.

```
fid = fopen ("myfile.txt", "w");
fdisp (fid, "3/8 is ");
fdisp (fid, 3/8);
fclose (fid);
```

See [Section 14.2.1 \[Opening and Closing Files\]](#), page 273, for details on how to use **fopen** and **fclose**.

fdisp (fid, x)

Display the value of *x* on the stream *fid*.

For example:

```
fdisp (stdout, "The value of pi is:"), fdisp (stdout, pi)

    ↗ the value of pi is:
    ↗ 3.1416
```

Note that the output from **fdisp** always ends with a newline.

See also: [disp], page 251.

Octave can also read and write matrices text files such as comma separated lists.

```
dlmwrite (file, M)
dlmwrite (file, M, delim, r, c)
dlmwrite (file, M, key, val ...)
dlmwrite (file, M, "-append", ...)
dlmwrite (fid, ...)
```

Write the numeric matrix *M* to the text file *file* using a delimiter.

file should be a filename or a writable file ID given by **fopen**.

The parameter *delim* specifies the delimiter to use to separate values on a row. If no delimiter is specified the comma character ‘,’ is used.

The value of *r* specifies the number of delimiter-only lines to add to the start of the file.

The value of *c* specifies the number of delimiters to prepend to each line of data.

If the argument “-append” is given, append to the end of *file*.

In addition, the following keyword value pairs may appear at the end of the argument list:

“append” Either “on” or “off”. See “-append” above.

"delimiter"

See *delim* above.

"newline"

The character(s) to separate each row. Three special cases exist for this option. "unix" is changed into "\n", "pc" is changed into "\r\n", and "mac" is changed into "\r". Any other value is used directly as the newline separator.

"roffset"

See *r* above.

"coffset"

See *c* above.

"precision"

The precision to use when writing the file. It can either be a format string (as used by fprintf) or a number of significant digits.

```
dlmwrite ("file.csv", reshape (1:16, 4, 4));
dlmwrite ("file.tex", a, "delimiter", "&", "newline", "\n")
```

See also: [dlmread], page 264, [csvread], page 265, [csvwrite], page 264.

```
data = dlmread (file)
data = dlmread (file, sep)
data = dlmread (file, sep, r0, c0)
data = dlmread (file, sep, range)
data = dlmread (... , "emptyvalue", EMPTYVAL)
```

Read numeric data from the text file *file* which uses the delimiter *sep* between data values.

If *sep* is not defined the separator between fields is determined from the file itself.

The optional scalar arguments *r0* and *c0* define the starting row and column of the data to be read. These values are indexed from zero, i.e., the first data row corresponds to an index of zero.

The *range* parameter specifies exactly which data elements are read. The first form of the parameter is a 4-element vector containing the upper left and lower right corners [*R0,C0,R1,C1*] where the indices are zero-based. Alternatively, a spreadsheet style form such as "A2..Q15" or "T1:AA5" can be used. The lowest alphabetical index 'A' refers to the first column. The lowest row index is 1.

file should be a filename or a file id given by fopen. In the latter case, the file is read until end of file is reached.

The "emptyvalue" option may be used to specify the value used to fill empty fields. The default is zero. Note that any non-numeric values, such as text, are also replaced by the "emptyvalue".

See also: [csvread], page 265, [textscan], page 266, [dlmwrite], page 263.

```
csvwrite (filename, x)
csvwrite (filename, x, dlm_opt1, ...)
```

Write the numeric matrix *x* to the file *filename* in comma-separated-value (CSV) format.

This function is equivalent to

```
dlmwrite (filename, x, ",", dlm_opt1, ...)
```

Any optional arguments are passed directly to `dlmwrite` (see [`dlmwrite`], page 263).

See also: [`csvread`], page 265, [`dlmwrite`], page 263, [`dlmread`], page 264.

```
x = csvread (filename)
```

```
x = csvread (filename, dlm_opt1, ...)
```

Read the comma-separated-value (CSV) file *filename* into the matrix *x*.

Note: only CSV files containing numeric data can be read.

This function is equivalent to

```
x = dlmread (filename, ", ", dlm_opt1, ...)
```

Any optional arguments are passed directly to `dlmread` (see [`dlmread`], page 264).

See also: [`dlmread`], page 264, [`textscan`], page 266, [`csvwrite`], page 264, [`dlmwrite`], page 263.

Formatted data from can be read from, or written to, text files as well.

```
[a, ...] = textread (filename)
```

```
[a, ...] = textread (filename, format)
```

```
[a, ...] = textread (filename, format, n)
```

```
[a, ...] = textread (filename, format, prop1, value1, ...)
```

```
[a, ...] = textread (filename, format, n, prop1, value1, ...)
```

This function is obsolete. Use `textscan` instead.

Read data from a text file.

The file *filename* is read and parsed according to *format*. The function behaves like `strread` except it works by parsing a file instead of a string. See the documentation of `strread` for details.

In addition to the options supported by `strread`, this function supports two more:

- **"headerlines"**: The first *value* number of lines of *filename* are skipped.
- **"endofline"**: Specify a single character or "`\r\n`". If no value is given, it will be inferred from the file. If set to "" (empty string) EOLs are ignored as delimiters.

The optional input *n* (format repeat count) specifies the number of times the format string is to be used or the number of lines to be read, whichever happens first while reading. The former is equivalent to requesting that the data output vectors should be of length *N*. Note that when reading files with format strings referring to multiple lines, *n* should rather be the number of lines to be read than the number of format string uses.

If the format string is empty (not just omitted) and the file contains only numeric data (excluding headerlines), `textread` will return a rectangular matrix with the number of columns matching the number of numeric fields on the first data line of the file. Empty fields are returned as zero values.

Examples:

```
Assume a data file like:
1 a 2 b
3 c 4 d
5 e

[a, b] = textread (f, "%f %s")
returns two columns of data, one with doubles, the other a
cellstr array:
a = [1; 2; 3; 4; 5]
b = {"a"; "b"; "c"; "d"; "e"}

[a, b] = textread (f, "%f %s", 3)
(read data into two culumns, try to use the format string
three times)
returns
a = [1; 2; 3]
b = {"a"; "b"; "c"}
```

With a data file like:

```
1
a
2
b

[a, b] = textread (f, "%f %s", 2)
returns a = 1 and b = {"a"}; i.e., the format string is used
only once because the format string refers to 2 lines of the
data file. To obtain 2x1 data output columns, specify N = 4
(number of data lines containing all requested data) rather
than 2.
```

See also: [\[textscan\]](#), page 266, [\[load\]](#), page 261, [\[dlmread\]](#), page 264, [\[fscanf\]](#), page 283, [\[strread\]](#), page 84.

```
C = textscan (fid, format)
C = textscan (fid, format, repeat)
C = textscan (fid, format, param, value, ...)
C = textscan (fid, format, repeat, param, value, ...)
C = textscan (str, ...)
[C, position, errmsg] = textscan (...)
```

Read data from a text file or string.

The string *str* or file associated with *fid* is read from and parsed according to *format*. The function is an extension of [strread](#) and [textread](#). Differences include: the ability to read from either a file or a string, additional options, and additional format specifiers.

The input is interpreted as a sequence of words, delimiters (such as whitespace), and literals. The characters that form delimiters and whitespace are determined by the

options. The format consists of format specifiers interspersed between literals. In the format, whitespace forms a delimiter between consecutive literals, but is otherwise ignored.

The output *C* is a cell array where the number of columns is determined by the number of format specifiers.

The first word of the input is matched to the first specifier of the format and placed in the first column of the output; the second is matched to the second specifier and placed in the second column and so forth. If there are more words than specifiers then the process is repeated until all words have been processed or the limit imposed by *repeat* has been met (see below).

The string *format* describes how the words in *str* should be parsed. As in *fscanf*, any (non-whitespace) text in the format that is not one of these specifiers is considered a literal. If there is a literal between two format specifiers then that same literal must appear in the input stream between the matching words.

The following specifiers are valid:

%f	
%f64	
%n	The word is parsed as a number and converted to double.
%f32	The word is parsed as a number and converted to single (float).
%d	
%d8	
%d16	
%d32	
%d64	The word is parsed as a number and converted to int8, int16, int32, or int64. If no size is specified then int32 is used.
%u	
%u8	
%u16	
%u32	
%u64	The word is parsed as a number and converted to uint8, uint16, uint32, or uint64. If no size is specified then uint32 is used.
%s	The word is parsed as a string ending at the last character before whitespace, an end-of-line, or a delimiter specified in the options.
%q	The word is parsed as a "quoted string". If the first character of the string is a double quote ("") then the string includes everything until a matching double quote—including whitespace, delimiters, and end-of-line characters. If a pair of consecutive double quotes appears in the input, it is replaced in the output by a single double quote. For examples, the input "He said ""Hello"" would return the value 'He said "Hello"'.
%c	The next character of the input is read. This includes delimiters, whitespace, and end-of-line characters.

%[...]	
%[^...]	In the first form, the word consists of the longest run consisting of only characters between the brackets. Ranges of characters can be specified by a hyphen; for example, %[0-9a-zA-Z] matches all alphanumeric characters (if the underlying character set is ASCII). Since MATLAB treats hyphens literally, this expansion only applies to alphanumeric characters. To include '-' in the set, it should appear first or last in the brackets; to include ']', it should be the first character. If the first character is '^' then the word consists of characters not listed.
%N...	For %s, %c %d, %f, %n, %u, an optional width can be specified as %Ns, etc. where N is an integer > 1. For %c, this causes exactly N characters to be read instead of a single character. For the other specifiers, it is an upper bound on the number of characters read; normal delimiters can cause fewer characters to be read. For complex numbers, this limit applies to the real and imaginary components individually. For %f and %n, format specifiers like %N.Mf are allowed, where M is an upper bound on number of characters after the decimal point to be considered; subsequent digits are skipped. For example, the specifier %8.2f would read 12.345e6 as 1.234e7.
%*...	The word specified by the remainder of the conversion specifier is skipped.
literals	In addition the format may contain literal character strings; these will be skipped during reading. If the input string does not match this literal, the processing terminates.

Parsed words corresponding to the first specifier are returned in the first output argument and likewise for the rest of the specifiers.

By default, if there is only one input argument, *format* is "%f". This means that numbers are read from the input into a single column vector. If *format* is explicitly empty ("") then *textscan* will return data in a number of columns matching the number of fields on the first data line of the input. Either of these is suitable only when the input is exclusively numeric.

For example, the string

```
str = "\n"
Bunny Bugs  5.5\n\
Duck Daffy -7.5e-5\n\
Penguin Tux  6"
```

can be read using

```
a = textscan (str, "%s %s %f");
```

The optional numeric argument *repeat* can be used for limiting the number of items read:

- 1 Read all of the string or file until the end (default).
- N Read until the first of two conditions occurs: 1) the format has been processed N times, or 2) N lines of the input have been processed. Zero (0) is an acceptable value for *repeat*. Currently, end-of-line characters

inside %q, %c, and %[...]\$ conversions do not contribute to the line count. This is incompatible with MATLAB and may change in future.

The behavior of `textscan` can be changed via property/value pairs. The following properties are recognized:

"BufSize"

This specifies the number of bytes to use for the internal buffer. A modest speed improvement may be obtained by setting this to a large value when reading a large file, especially if the input contains long strings. The default is 4096, or a value dependent on *n* if that is specified.

"CollectOutput"

A value of 1 or true instructs `textscan` to concatenate consecutive columns of the same class in the output cell array. A value of 0 or false (default) leaves output in distinct columns.

"CommentStyle"

Specify parts of the input which are considered comments and will be skipped. *value* is the comment style and can be either (1) A string or 1x1 cell string, to skip everything to the right of it; (2) A cell array of two strings, to skip everything between the first and second strings. Comments are only parsed where whitespace is accepted and do not act as delimiters.

"Delimiter"

If *value* is a string, any character in *value* will be used to split the input into words. If *value* is a cell array of strings, any string in the array will be used to split the input into words. (default value = any whitespace.)

"EmptyValue"

Value to return for empty numeric values in non-whitespace delimited data. The default is NaN. When the data type does not support NaN (int32 for example), then the default is zero.

"EndOfLine"

value can be either an empty or one character specifying the end-of-line character, or the pair "\r\n" (CRLF). In the latter case, any of "\r", "\n" or "\r\n" is counted as a (single) newline. If no value is given, "\r\n" is used.

"HeaderLines"

The first *value* number of lines of *fid* are skipped. Note that this does not refer to the first non-comment lines, but the first lines of any type.

"MultipleDelimsAsOne"

If *value* is nonzero, treat a series of consecutive delimiters, without whitespace in between, as a single delimiter. Consecutive delimiter series need not be vertically aligned. Without this option, a single delimiter before the end of the line does not cause the line to be considered to end with an empty value, but a single delimiter at the start of a line causes the line to be considered to start with an empty value.

"TreatAsEmpty"

Treat single occurrences (surrounded by delimiters or whitespace) of the string(s) in *value* as missing values.

"ReturnOnError"

If set to numerical 1 or true, return normally as soon as an error is encountered, such as trying to read a string using %f. If set to 0 or false, return an error and no data.

"Whitespace"

Any character in *value* will be interpreted as whitespace and trimmed; The default value for whitespace is " \b\r\n\t" (note the space). Unless whitespace is set to "" (empty) AND at least one "%s" format conversion specifier is supplied, a space is always part of whitespace.

When the number of words in *str* or *fid* doesn't match an exact multiple of the number of format conversion specifiers, `textscan`'s behavior depends on whether the last character of the string or file is an end-of-line as specified by the `EndOfLine` option:

last character = end-of-line

Data columns are padded with empty fields, NaN or 0 (for integer fields) so that all columns have equal length

last character is not end-of-line

Data columns are not padded; `textscan` returns columns of unequal length

The second output *position* provides the location, in characters from the beginning of the file or string, where processing stopped.

See also: [\[dlmread\]](#), page 264, [\[fscanf\]](#), page 283, [\[load\]](#), page 261, [\[strread\]](#), page 84, [\[textread\]](#), page 265.

The `importdata` function has the ability to work with a wide variety of data.

```
A = importdata (fname)
A = importdata (fname, delimiter)
A = importdata (fname, delimiter, header_rows)
[A, delimiter] = importdata (...)
[A, delimiter, header_rows] = importdata (...)
Import data from the file fname.
```

Input parameters:

- *fname* The name of the file containing data.
- *delimiter* The character separating columns of data. Use \t for tab. (Only valid for ASCII files)
- *header_rows* The number of header rows before the data begins. (Only valid for ASCII files)

Different file types are supported:

- ASCII table
Import ASCII table using the specified number of header rows and the specified delimiter.
- Image file
- MATLAB file
- Spreadsheet files (depending on external software)
- WAV file

See also: [textscan], page 266, [dlmread], page 264, [csvread], page 265, [load], page 261.

14.1.3.1 Saving Data on Unexpected Exits

If Octave for some reason exits unexpectedly it will by default save the variables available in the workspace to a file in the current directory. By default this file is named ‘octave-workspace’ and can be loaded into memory with the `load` command. While the default behavior most often is reasonable it can be changed through the following functions.

```
val = crash_dumps_octave_core ()
old_val = crash_dumps_octave_core (new_val)
crash_dumps_octave_core (new_val, "local")
```

Query or set the internal variable that controls whether Octave tries to save all current variables to the file `octave-workspace` if it crashes or receives a hangup, terminate or similar signal.

When called from inside a function with the “`local`” option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [octave_core_file_limit], page 272, [octave_core_file_name], page 272, [octave_core_file_options], page 272.

```
val = sighup_dumps_octave_core ()
old_val = sighup_dumps_octave_core (new_val)
sighup_dumps_octave_core (new_val, "local")
```

Query or set the internal variable that controls whether Octave tries to save all current variables to the file `octave-workspace` if it receives a hangup signal.

When called from inside a function with the “`local`” option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

```
val = sigquit_dumps_octave_core ()
old_val = sigquit_dumps_octave_core (new_val)
sigquit_dumps_octave_core (new_val, "local")
```

Query or set the internal variable that controls whether Octave tries to save all current variables to the file `octave-workspace` if it receives a quit signal.

When called from inside a function with the “`local`” option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

```
val = sigterm.dumps_octave_core ()
old_val = sigterm.dumps_octave_core (new_val)
sigterm.dumps_octave_core (new_val, "local")
```

Query or set the internal variable that controls whether Octave tries to save all current variables to the file `octave-workspace` if it receives a terminate signal.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

```
val = octave_core_file_options ()
old_val = octave_core_file_options (new_val)
octave_core_file_options (new_val, "local")
```

Query or set the internal variable that specifies the options used for saving the workspace data if Octave aborts.

The value of `octave_core_file_options` should follow the same format as the options for the `save` function. The default value is Octave's binary format.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[crash.dumps_octave_core\]](#), page 271, [\[octave_core_file_name\]](#), page 272, [\[octave_core_file_limit\]](#), page 272.

```
val = octave_core_file_limit ()
old_val = octave_core_file_limit (new_val)
octave_core_file_limit (new_val, "local")
```

Query or set the internal variable that specifies the maximum amount of memory that Octave will save when writing a crash dump file.

The limit is measured in kilobytes and is applied to the top-level workspace. The name of the crash dump file is specified by `octave_core_file_name`.

If `octave_core_file_options` flags specify a binary format, then `octave_core_file_limit` will be approximately the maximum size of the file. If a text file format is used, then the file could be much larger than the limit. The default value is -1 (unlimited).

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[crash.dumps_octave_core\]](#), page 271, [\[octave_core_file_name\]](#), page 272, [\[octave_core_file_options\]](#), page 272.

```
val = octave_core_file_name ()
old_val = octave_core_file_name (new_val)
octave_core_file_name (new_val, "local")
```

Query or set the internal variable that specifies the name of the file used for saving data from the top-level workspace if Octave aborts.

The default value is "`octave-workspace`"

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[crash.dumps_octave_core\]](#), page 271, [\[octave_core_file_name\]](#), page 272, [\[octave_core_file_options\]](#), page 272.

14.2 C-Style I/O Functions

Octave's C-style input and output functions provide most of the functionality of the C programming language's standard I/O library. The argument lists for some of the input functions are slightly different, however, because Octave has no way of passing arguments by reference.

In the following, *file* refers to a filename and *fid* refers to an integer file number, as returned by `fopen`.

There are three files that are always available. Although these files can be accessed using their corresponding numeric file ids, you should always use the symbolic names given in the table below, since it will make your programs easier to understand.

`stdin ()`

Return the numeric value corresponding to the standard input stream.

When Octave is used interactively, `stdin` is filtered through the command line editing functions.

See also: [\[stdout\]](#), page 273, [\[stderr\]](#), page 273.

`stdout ()`

Return the numeric value corresponding to the standard output stream.

Data written to the standard output may be filtered through the pager.

See also: [\[stdin\]](#), page 273, [\[stderr\]](#), page 273, [\[page_screen_output\]](#), page 255.

`stderr ()`

Return the numeric value corresponding to the standard error stream.

Even if paging is turned on, the standard error is not sent to the pager. It is useful for error messages and prompts.

See also: [\[stdin\]](#), page 273, [\[stdout\]](#), page 273.

14.2.1 Opening and Closing Files

When reading data from a file it must be opened for reading first, and likewise when writing to a file. The `fopen` function returns a pointer to an open file that is ready to be read or written. Once all data has been read from or written to the opened file it should be closed. The `fclose` function does this. The following code illustrates the basic pattern for writing to a file, but a very similar pattern is used when reading a file.

```
filename = "myfile.txt";
fid = fopen (filename, "w");
# Do the actual I/O here...
fclose (fid);
```

```

fid = fopen (name)
fid = fopen (name, mode)
fid = fopen (name, mode, arch)
[fid, msg] = fopen (... )
fid_list = fopen ("all")
[file, mode, arch] = fopen (fid)

```

Open a file for low-level I/O or query open files and file descriptors.

The first form of the `fopen` function opens the named file with the specified mode (read-write, read-only, etc.) and architecture interpretation (IEEE big endian, IEEE little endian, etc.), and returns an integer value that may be used to refer to the file later. If an error occurs, *fid* is set to -1 and *msg* contains the corresponding system error message. The *mode* is a one or two character string that specifies whether the file is to be opened for reading, writing, or both.

The second form of the `fopen` function returns a vector of file ids corresponding to all the currently open files, excluding the `stdin`, `stdout`, and `stderr` streams.

The third form of the `fopen` function returns information about the open file given its file id.

For example,

```
myfile = fopen ("splat.dat", "r", "ieee-le");
```

opens the file `splat.dat` for reading. If necessary, binary numeric values will be read assuming they are stored in IEEE format with the least significant bit first, and then converted to the native representation.

Opening a file that is already open simply opens it again and returns a separate file id. It is not an error to open a file several times, though writing to the same file through several different file ids may produce unexpected results.

The possible values of *mode* are

- ‘r’ (default)
Open a file for reading.
- ‘w’
Open a file for writing. The previous contents are discarded.
- ‘a’
Open or create a file for writing at the end of the file.
- ‘r+’
Open an existing file for reading and writing.
- ‘w+’
Open a file for reading or writing. The previous contents are discarded.
- ‘a+’
Open or create a file for reading or writing at the end of the file.

Append a “t” to the mode string to open the file in text mode or a “b” to open in binary mode. On Windows systems, text mode reading and writing automatically converts linefeeds to the appropriate line end character for the system (carriage-return linefeed on Windows). The default when no mode is specified is binary.

Additionally, you may append a “z” to the mode string to open a gzipped file for reading or writing. For this to be successful, you must also open the file in binary mode.

The parameter `arch` is a string specifying the default data format for the file. Valid values for `arch` are:

`"native"` or `"n"` (default)

The format of the current machine.

`"ieee-be"` or `"b"`

IEEE big endian format.

`"ieee-le"` or `"l"`

IEEE little endian format.

However, conversions are currently only supported for ‘native’, ‘ieee-be’, and ‘ieee-le’ formats.

When opening a new file that does not yet exist, permissions will be set to `0666 - umask`.

Compatibility Note: Octave opens files using buffered I/O. Small writes are accumulated until an internal buffer is filled, and then everything is written in a single operation. This is very efficient and improves performance. MATLAB, however, opens files using flushed I/O where every write operation is immediately performed. If the write operation must be performed immediately after data has been written then the write should be followed by a call to `fflush` to flush the internal buffer.

See also: `[fclose]`, page 275, `[fgets]`, page 276, `[fgetl]`, page 276, `[fscanf]`, page 283, `[fread]`, page 286, `[fputs]`, page 275, `[fdisp]`, page 263, `[fprintf]`, page 277, `[fwrite]`, page 289, `[fskip]`, page 277, `[fseek]`, page 292, `[rewind]`, page 292, `[ftell]`, page 292, `[feof]`, page 291, `[ferror]`, page 291, `[fclear]`, page 291, `[fflush]`, page 256, `[freport]`, page 291, `[umask]`, page 899.

```
fclose (fid)
fclose ("all")
status = fclose ("all")
```

Close the file specified by the file descriptor `fid`.

If successful, `fclose` returns 0, otherwise, it returns -1. The second form of the `fclose` call closes all open files except `stdin`, `stdout`, `stderr`, and any FIDs associated with gnuplot.

See also: `[fopen]`, page 273, `[fflush]`, page 256, `[freport]`, page 291.

```
is_valid_file_id (fid)
```

Return true if `fid` refers to an open file.

See also: `[freport]`, page 291, `[fopen]`, page 273.

14.2.2 Simple Output

Once a file has been opened for writing a string can be written to the file using the `fputs` function. The following example shows how to write the string ‘Free Software is needed for Free Science’ to the file ‘`free.txt`’.

```
filename = "free.txt";
fid = fopen (filename, "w");
fputs (fid, "Free Software is needed for Free Science");
fclose (fid);
```

```
fputs (fid, string)
status = fputs (fid, string)
```

Write the string *string* to the file with file descriptor *fid*.

The string is written to the file with no additional formatting. Use **fdisp** instead to automatically append a newline character appropriate for the local machine.

Return a non-negative number on success or EOF on error.

See also: [\[fdisp\]](#), page 263, [\[sprintf\]](#), page 277, [\[fwrite\]](#), page 289, [\[fopen\]](#), page 273.

A function much similar to **fputs** is available for writing data to the screen. The **puts** function works just like **fputs** except it doesn't take a file pointer as its input.

```
puts (string)
status = puts (string)
```

Write a string to the standard output with no formatting.

The string is written verbatim to the standard output. Use **disp** to automatically append a newline character appropriate for the local machine.

Return a non-negative number on success and EOF on error.

See also: [\[fputs\]](#), page 275, [\[disp\]](#), page 251.

14.2.3 Line-Oriented Input

To read from a file it must be opened for reading using **fopen**. Then a line can be read from the file using **fgetl** as the following code illustrates

```
fid = fopen ("free.txt");
txt = fgetl (fid)
    ⇒ Free Software is needed for Free Science
fclose (fid);
```

This of course assumes that the file ‘`free.txt`’ exists and contains the line ‘Free Software is needed for Free Science’.

```
str = fgetl (fid)
str = fgetl (fid, len)
```

Read characters from a file, stopping after a newline, or EOF, or *len* characters have been read.

The characters read, excluding the possible trailing newline, are returned as a string.

If *len* is omitted, **fgetl** reads until the next newline character.

If there are no more characters to read, **fgetl** returns `-1`.

To read a line and return the terminating newline see **fgets**.

See also: [\[fgets\]](#), page 276, [\[fscanf\]](#), page 283, [\[fread\]](#), page 286, [\[fopen\]](#), page 273.

```
str = fgets (fid)
str = fgets (fid, len)
```

Read characters from a file, stopping after a newline, or EOF, or *len* characters have been read.

The characters read, including the possible trailing newline, are returned as a string.

If *len* is omitted, **fgets** reads until the next newline character.

If there are no more characters to read, **fgets** returns **-1**.

To read a line and discard the terminating newline see **fgetl**.

See also: [\[fputs\]](#), page 275, [\[fgetl\]](#), page 276, [\[fscanf\]](#), page 283, [\[fread\]](#), page 286, [\[fopen\]](#), page 273.

```
nlines = fskip1 (fid)
nlines = fskip1 (fid, count)
nlines = fskip1 (fid, Inf)
```

Read and skip *count* lines from the file specified by the file descriptor *fid*.

fskip1 discards characters until an end-of-line is encountered exactly *count*-times, or until the end-of-file marker is found.

If *count* is omitted, it defaults to 1. *count* may also be **Inf**, in which case lines are skipped until the end of the file. This form is suitable for counting the number of lines in a file.

Returns the number of lines skipped (end-of-line sequences encountered).

See also: [\[fgetl\]](#), page 276, [\[fgets\]](#), page 276, [\[fscanf\]](#), page 283, [\[fopen\]](#), page 273.

14.2.4 Formatted Output

This section describes how to call **printf** and related functions.

The following functions are available for formatted output. They are modeled after the C language functions of the same name, but they interpret the format template differently in order to improve the performance of printing vector and matrix values.

Implementation Note: For compatibility with MATLAB, escape sequences in the template string (e.g., "\n" => newline) are expanded even when the template string is defined with single quotes.

```
printf (template, ...)
```

Print optional arguments under the control of the template string *template* to the stream **stdout** and return the number of characters printed.

See the Formatted Output section of the GNU Octave manual for a complete description of the syntax of the template string.

Implementation Note: For compatibility with MATLAB, escape sequences in the template string (e.g., "\n" => newline) are expanded even when the template string is defined with single quotes.

See also: [\[fprintf\]](#), page 277, [\[sprintf\]](#), page 278, [\[scanf\]](#), page 283.

```
fprintf (fid, template, ...)
fprintf (template, ...)
numbytes = fprintf (...)
```

This function is equivalent to **printf**, except that the output is written to the file descriptor *fid* instead of **stdout**.

If *fid* is omitted, the output is written to **stdout** making the function exactly equivalent to **printf**.

The optional output returns the number of bytes written to the file.

Implementation Note: For compatibility with MATLAB, escape sequences in the template string (e.g., "\n" => newline) are expanded even when the template string is defined with single quotes.

See also: [\[fputs\]](#), page 275, [\[fdisp\]](#), page 263, [\[fwrite\]](#), page 289, [\[fscanf\]](#), page 283, [\[printf\]](#), page 277, [\[sprintf\]](#), page 278, [\[fopen\]](#), page 273.

sprintf (template, ...)

This is like [printf](#), except that the output is returned as a string.

Unlike the C library function, which requires you to provide a suitably sized string as an argument, Octave's **sprintf** function returns the string, automatically sized to hold all of the items converted.

Implementation Note: For compatibility with MATLAB, escape sequences in the template string (e.g., "\n" => newline) are expanded even when the template string is defined with single quotes.

See also: [\[printf\]](#), page 277, [\[fprintf\]](#), page 277, [\[sscanf\]](#), page 284.

The **printf** function can be used to print any number of arguments. The template string argument you supply in a call provides information not only about the number of additional arguments, but also about their types and what style should be used for printing them.

Ordinary characters in the template string are simply written to the output stream as-is, while *conversion specifications* introduced by a '%' character in the template cause subsequent arguments to be formatted and written to the output stream. For example,

```
pct = 37;
filename = "foo.txt";
printf ("Processed %d%% of '%s'.\nPlease be patient.\n",
       pct, filename);
```

produces output like

```
Processed 37% of 'foo.txt'.
Please be patient.
```

This example shows the use of the '%d' conversion to specify that a scalar argument should be printed in decimal notation, the '%s' conversion to specify printing of a string argument, and the '%%' conversion to print a literal '%' character.

There are also conversions for printing an integer argument as an unsigned value in octal, decimal, or hexadecimal radix ('%o', '%u', or '%x', respectively); or as a character value ('%c').

Floating-point numbers can be printed in normal, fixed-point notation using the '%f' conversion or in exponential notation using the '%e' conversion. The '%g' conversion uses either '%e' or '%f' format, depending on what is more appropriate for the magnitude of the particular number.

You can control formatting more precisely by writing *modifiers* between the '%' and the character that indicates which conversion to apply. These slightly alter the ordinary behavior of the conversion. For example, most conversion specifications permit you to

specify a minimum field width and a flag indicating whether you want the result left- or right-justified within the field.

The specific flags and modifiers that are permitted and their interpretation vary depending on the particular conversion. They're all described in more detail in the following sections.

14.2.5 Output Conversion for Matrices

When given a matrix value, Octave's formatted output functions cycle through the format template until all the values in the matrix have been printed. For example:

```
printf ("%4.2f %10.2e %8.4g\n", hilb (3));
```

-	1.00	5.00e-01	0.3333
-	0.50	3.33e-01	0.25
-	0.33	2.50e-01	0.2

If more than one value is to be printed in a single call, the output functions do not return to the beginning of the format template when moving on from one value to the next. This can lead to confusing output if the number of elements in the matrices are not exact multiples of the number of conversions in the format template. For example:

```
printf ("%4.2f %10.2e %8.4g\n", [1, 2], [3, 4]);
```

-	1.00	2.00e+00	3
-	4.00		

If this is not what you want, use a series of calls instead of just one.

14.2.6 Output Conversion Syntax

This section provides details about the precise syntax of conversion specifications that can appear in a `printf` template string.

Characters in the template string that are not part of a conversion specification are printed as-is to the output stream.

The conversion specifications in a `printf` template string have the general form:

```
% flags width [ . precision ] type conversion
```

For example, in the conversion specifier ‘%-10.81d’, the ‘-’ is a flag, ‘10’ specifies the field width, the precision is ‘8’, the letter ‘1’ is a type modifier, and ‘d’ specifies the conversion style. (This particular type specifier says to print a numeric argument in decimal notation, with a minimum of 8 digits left-justified in a field at least 10 characters wide.)

In more detail, output conversion specifications consist of an initial ‘%’ character followed in sequence by:

- Zero or more *flag characters* that modify the normal behavior of the conversion specification.
- An optional decimal integer specifying the *minimum field width*. If the normal conversion produces fewer characters than this, the field is padded with spaces to the specified width. This is a *minimum* value; if the normal conversion produces more characters than this, the field is *not* truncated. Normally, the output is right-justified within the field.

You can also specify a field width of ‘*’. This means that the next argument in the argument list (before the actual value to be printed) is used as the field width. The value is rounded to the nearest integer. If the value is negative, this means to set the ‘-’ flag (see below) and to use the absolute value as the field width.

- An optional *precision* to specify the number of digits to be written for the numeric conversions. If the precision is specified, it consists of a period (‘.’) followed optionally by a decimal integer (which defaults to zero if omitted).

You can also specify a precision of ‘*’. This means that the next argument in the argument list (before the actual value to be printed) is used as the precision. The value must be an integer, and is ignored if it is negative.

- An optional *type modifier character*. This character is ignored by Octave’s `printf` function, but is recognized to provide compatibility with the C language `printf`.
- A character that specifies the conversion to be applied.

The exact options that are permitted and how they are interpreted vary between the different conversion specifiers. See the descriptions of the individual conversions for information about the particular options that they use.

14.2.7 Table of Output Conversions

Here is a table summarizing what all the different conversions do:

‘%d’, ‘%i’	Print an integer as a signed decimal number. See Section 14.2.8 [Integer Conversions] , page 281, for details. ‘%d’ and ‘%i’ are synonymous for output, but are different when used with <code>scanf</code> for input (see Section 14.2.13 [Table of Input Conversions] , page 285).
‘%o’	Print an integer as an unsigned octal number. See Section 14.2.8 [Integer Conversions] , page 281, for details.
‘%u’	Print an integer as an unsigned decimal number. See Section 14.2.8 [Integer Conversions] , page 281, for details.
‘%x’, ‘%X’	Print an integer as an unsigned hexadecimal number. ‘%x’ uses lowercase letters and ‘%X’ uses uppercase. See Section 14.2.8 [Integer Conversions] , page 281, for details.
‘%f’	Print a floating-point number in normal (fixed-point) notation. See Section 14.2.9 [Floating-Point Conversions] , page 282, for details.
‘%e’, ‘%E’	Print a floating-point number in exponential notation. ‘%e’ uses lowercase letters and ‘%E’ uses uppercase. See Section 14.2.9 [Floating-Point Conversions] , page 282, for details.
‘%g’, ‘%G’	Print a floating-point number in either normal (fixed-point) or exponential notation, whichever is more appropriate for its magnitude. ‘%g’ uses lowercase letters and ‘%G’ uses uppercase. See Section 14.2.9 [Floating-Point Conversions] , page 282, for details.
‘%c’	Print a single character. See Section 14.2.10 [Other Output Conversions] , page 282.
‘%s’	Print a string. See Section 14.2.10 [Other Output Conversions] , page 282.

'%%'	Print a literal '%' character. See Section 14.2.10 [Other Output Conversions] , page 282 .
------	--

If the syntax of a conversion specification is invalid, unpredictable things will happen, so don't do this. In particular, MATLAB allows a bare percentage sign '%' with no subsequent conversion character. Octave will emit an error and stop if it sees such code. When the string variable to be processed cannot be guaranteed to be free of potential format codes it is better to use the two argument form of any of the `printf` functions and set the format string to `%s`. Alternatively, for code which is not required to be backwards-compatible with MATLAB the Octave function `puts` or `disp` can be used.

```
printf (strvar);          # Unsafe if strvar contains format codes
printf ("%s", strvar);   # Safe
puts (strvar);          # Safe
```

If there aren't enough function arguments provided to supply values for all the conversion specifications in the template string, or if the arguments are not of the correct types, the results are unpredictable. If you supply more arguments than conversion specifications, the extra argument values are simply ignored; this is sometimes useful.

14.2.8 Integer Conversions

This section describes the options for the '%d', '%i', '%o', '%u', '%x', and '%X' conversion specifications. These conversions print integers in various formats.

The '%d' and '%i' conversion specifications both print an numeric argument as a signed decimal number; while '%o', '%u', and '%x' print the argument as an unsigned octal, decimal, or hexadecimal number (respectively). The '%X' conversion specification is just like '%x' except that it uses the characters 'ABCDEF' as digits instead of 'abcdef'.

The following flags are meaningful:

'-'	Left-justify the result in the field (instead of the normal right-justification).
'+'	For the signed '%d' and '%i' conversions, print a plus sign if the value is positive.
' '	For the signed '%d' and '%i' conversions, if the result doesn't start with a plus or minus sign, prefix it with a space character instead. Since the '+' flag ensures that the result includes a sign, this flag is ignored if you supply both of them.
'#'	For the '%o' conversion, this forces the leading digit to be '0', as if by increasing the precision. For '%x' or '%X', this prefixes a leading '0x' or '0X' (respectively) to the result. This doesn't do anything useful for the '%d', '%i', or '%u' conversions.
'0'	Pad the field with zeros instead of spaces. The zeros are placed after any indication of sign or base. This flag is ignored if the '-' flag is also specified, or if a precision is specified.

If a precision is supplied, it specifies the minimum number of digits to appear; leading zeros are produced if necessary. If you don't specify a precision, the number is printed with as many digits as it needs. If you convert a value of zero with an explicit precision of zero, then no characters at all are produced.

14.2.9 Floating-Point Conversions

This section discusses the conversion specifications for floating-point numbers: the ‘%f’, ‘%e’, ‘%E’, ‘%g’, and ‘%G’ conversions.

The ‘%f’ conversion prints its argument in fixed-point notation, producing output of the form $[-]ddd.ddd$, where the number of digits following the decimal point is controlled by the precision you specify.

The ‘%e’ conversion prints its argument in exponential notation, producing output of the form $[-]d.dde[+|-]dd$. Again, the number of digits following the decimal point is controlled by the precision. The exponent always contains at least two digits. The ‘%E’ conversion is similar but the exponent is marked with the letter ‘E’ instead of ‘e’.

The ‘%g’ and ‘%G’ conversions print the argument in the style of ‘%e’ or ‘%E’ (respectively) if the exponent would be less than -4 or greater than or equal to the precision; otherwise they use the ‘%f’ style. Trailing zeros are removed from the fractional portion of the result and a decimal-point character appears only if it is followed by a digit.

The following flags can be used to modify the behavior:

- ‘-’ Left-justify the result in the field. Normally the result is right-justified.
- ‘+’ Always include a plus or minus sign in the result.
- ‘ ’ If the result doesn’t start with a plus or minus sign, prefix it with a space instead. Since the ‘+’ flag ensures that the result includes a sign, this flag is ignored if you supply both of them.
- ‘#’ Specifies that the result should always include a decimal point, even if no digits follow it. For the ‘%g’ and ‘%G’ conversions, this also forces trailing zeros after the decimal point to be left in place where they would otherwise be removed.
- ‘0’ Pad the field with zeros instead of spaces; the zeros are placed after any sign. This flag is ignored if the ‘-’ flag is also specified.

The precision specifies how many digits follow the decimal-point character for the ‘%f’, ‘%e’, and ‘%E’ conversions. For these conversions, the default precision is 6. If the precision is explicitly 0, this suppresses the decimal point character entirely. For the ‘%g’ and ‘%G’ conversions, the precision specifies how many significant digits to print. Significant digits are the first digit before the decimal point, and all the digits after it. If the precision is 0 or not specified for ‘%g’ or ‘%G’, it is treated like a value of 1. If the value being printed cannot be expressed precisely in the specified number of digits, the value is rounded to the nearest number that fits.

14.2.10 Other Output Conversions

This section describes miscellaneous conversions for `printf`.

The ‘%c’ conversion prints a single character. The ‘-’ flag can be used to specify left-justification in the field, but no other flags are defined, and no precision or type modifier can be given. For example:

```
printf ("%c%c%c%c%c", "h", "e", "l", "l", "o");
prints 'hello'.
```

The ‘%s’ conversion prints a string. The corresponding argument must be a string. A precision can be specified to indicate the maximum number of characters to write; otherwise characters in the string up to but not including the terminating null character are written to the output stream. The ‘-’ flag can be used to specify left-justification in the field, but no other flags or type modifiers are defined for this conversion. For example:

```
printf ("%3s%-6s", "no", "where");
prints ' nowhere ' (note the leading and trailing spaces).
```

14.2.11 Formatted Input

Octave provides the `scanf`, `fscanf`, and `sscanf` functions to read formatted input. There are two forms of each of these functions. One can be used to extract vectors of data from a file, and the other is more ‘C-like’.

```
[val, count, errmsg] = fscanf (fid, template, size)
[v1, v2, ..., count, errmsg] = fscanf (fid, template, "C")
```

In the first form, read from `fid` according to `template`, returning the result in the matrix `val`.

The optional argument `size` specifies the amount of data to read and may be one of

`Inf` Read as much as possible, returning a column vector.

`nr` Read up to `nr` elements, returning a column vector.

`[nr, Inf]` Read as much as possible, returning a matrix with `nr` rows. If the number of elements read is not an exact multiple of `nr`, the last column is padded with zeros.

`[nr, nc]` Read up to `nr * nc` elements, returning a matrix with `nr` rows. If the number of elements read is not an exact multiple of `nr`, the last column is padded with zeros.

If `size` is omitted, a value of `Inf` is assumed.

A string is returned if `template` specifies only character conversions.

The number of items successfully read is returned in `count`.

If an error occurs, `errmsg` contains a system-dependent error message.

In the second form, read from `fid` according to `template`, with each conversion specifier in `template` corresponding to a single scalar return value. This form is more “C-like”, and also compatible with previous versions of Octave. The number of successful conversions is returned in `count`

See the Formatted Input section of the GNU Octave manual for a complete description of the syntax of the template string.

See also: `[fgets]`, page 276, `[fgetl]`, page 276, `[fread]`, page 286, `[scanf]`, page 283, `[sscanf]`, page 284, `[fopen]`, page 273.

```
[val, count, errmsg] = scanf (template, size)
[v1, v2, ..., count, errmsg] = scanf (template, "C")
```

This is equivalent to calling `fscanf` with `fid = stdin`.

It is currently not useful to call `scanf` in interactive programs.

See also: `[fscanf]`, page 283, `[sscanf]`, page 284, `[printf]`, page 277.

```
[val, count, errmsg, pos] = sscanf (string, template, size)
[v1, v2, ..., count, errmsg] = sscanf (string, template, "C")
This is like fscanf, except that the characters are taken from the string string instead
of from a stream.
```

Reaching the end of the string is treated as an end-of-file condition. In addition to the values returned by `fscanf`, the index of the next character to be read is returned in *pos*.

See also: `[fscanf]`, page 283, `[scanf]`, page 283, `[sprintf]`, page 278.

Calls to `scanf` are superficially similar to calls to `printf` in that arbitrary arguments are read under the control of a template string. While the syntax of the conversion specifications in the template is very similar to that for `printf`, the interpretation of the template is oriented more towards free-format input and simple pattern matching, rather than fixed-field formatting. For example, most `scanf` conversions skip over any amount of “white space” (including spaces, tabs, and newlines) in the input file, and there is no concept of precision for the numeric input conversions as there is for the corresponding output conversions. Ordinarily, non-whitespace characters in the template are expected to match characters in the input stream exactly.

When a *matching failure* occurs, `scanf` returns immediately, leaving the first non-matching character as the next character to be read from the stream, and `scanf` returns all the items that were successfully converted.

The formatted input functions are not used as frequently as the formatted output functions. Partly, this is because it takes some care to use them properly. Another reason is that it is difficult to recover from a matching error.

14.2.12 Input Conversion Syntax

A `scanf` template string is a string that contains ordinary multibyte characters interspersed with conversion specifications that start with ‘%’.

Any whitespace character in the template causes any number of whitespace characters in the input stream to be read and discarded. The whitespace characters that are matched need not be exactly the same whitespace characters that appear in the template string. For example, write ‘ , ’ in the template to recognize a comma with optional whitespace before and after.

Other characters in the template string that are not part of conversion specifications must match characters in the input stream exactly; if this is not the case, a matching failure occurs.

The conversion specifications in a `scanf` template string have the general form:

```
% flags width type conversion
```

In more detail, an input conversion specification consists of an initial ‘%’ character followed in sequence by:

- An optional *flag character* ‘*’, which says to ignore the text read for this specification. When `scanf` finds a conversion specification that uses this flag, it reads input as directed by the rest of the conversion specification, but it discards this input, does not return any value, and does not increment the count of successful assignments.

- An optional decimal integer that specifies the *maximum field width*. Reading of characters from the input stream stops either when this maximum is reached or when a non-matching character is found, whichever happens first. Most conversions discard initial whitespace characters, and these discarded characters don't count towards the maximum field width. Conversions that do not discard initial whitespace are explicitly documented.
- An optional type modifier character. This character is ignored by Octave's `scanf` function, but is recognized to provide compatibility with the C language `scanf`.
- A character that specifies the conversion to be applied.

The exact options that are permitted and how they are interpreted vary between the different conversion specifiers. See the descriptions of the individual conversions for information about the particular options that they allow.

14.2.13 Table of Input Conversions

Here is a table that summarizes the various conversion specifications:

'%d'	Matches an optionally signed integer written in decimal. See Section 14.2.14 [Numeric Input Conversions], page 286 .
'%i'	Matches an optionally signed integer in any of the formats that the C language defines for specifying an integer constant. See Section 14.2.14 [Numeric Input Conversions], page 286 .
'%o'	Matches an unsigned integer written in octal radix. See Section 14.2.14 [Numeric Input Conversions], page 286 .
'%u'	Matches an unsigned integer written in decimal radix. See Section 14.2.14 [Numeric Input Conversions], page 286 .
'%x', '%X'	Matches an unsigned integer written in hexadecimal radix. See Section 14.2.14 [Numeric Input Conversions], page 286 .
'%e', '%f', '%g', '%E', '%G'	Matches an optionally signed floating-point number. See Section 14.2.14 [Numeric Input Conversions], page 286 .
'%s'	Matches a string containing only non-whitespace characters. See Section 14.2.15 [String Input Conversions], page 286 .
'%c'	Matches a string of one or more characters; the number of characters read is controlled by the maximum field width given for the conversion. See Section 14.2.15 [String Input Conversions], page 286 .
'%%'	This matches a literal '%' character in the input stream. No corresponding argument is used.

If the syntax of a conversion specification is invalid, the behavior is undefined. If there aren't enough function arguments provided to supply addresses for all the conversion specifications in the template strings that perform assignments, or if the arguments are not of the correct types, the behavior is also undefined. On the other hand, extra arguments are simply ignored.

14.2.14 Numeric Input Conversions

This section describes the `scanf` conversions for reading numeric values.

The ‘%d’ conversion matches an optionally signed integer in decimal radix.

The ‘%i’ conversion matches an optionally signed integer in any of the formats that the C language defines for specifying an integer constant.

For example, any of the strings ‘10’, ‘0xa’, or ‘012’ could be read in as integers under the ‘%i’ conversion. Each of these specifies a number with decimal value 10.

The ‘%o’, ‘%u’, and ‘%x’ conversions match unsigned integers in octal, decimal, and hexadeciml radices, respectively.

The ‘%X’ conversion is identical to the ‘%x’ conversion. They both permit either uppercase or lowercase letters to be used as digits.

By default, integers are read as 32-bit quantities. With the ‘h’ modifier, 16-bit integers are used, and with the ‘l’ modifier, 64-bit integers are used.

14.2.15 String Input Conversions

This section describes the `scanf` input conversions for reading string and character values: ‘%s’ and ‘%c’.

The ‘%c’ conversion is the simplest: it matches a fixed number of characters, always. The maximum field width says how many characters to read; if you don’t specify the maximum, the default is 1. This conversion does not skip over initial whitespace characters. It reads precisely the next *n* characters, and fails if it cannot get that many.

The ‘%s’ conversion matches a string of non-whitespace characters. It skips and discards initial whitespace, but stops when it encounters more whitespace after having read something.

For example, reading the input:

```
hello, world
```

with the conversion ‘%10c’ produces “hello, wo”, but reading the same input with the conversion ‘%10s’ produces “hello,”.

14.2.16 Binary I/O

Octave can read and write binary data using the functions `fread` and `fwrite`, which are patterned after the standard C functions with the same names. They are able to automatically swap the byte order of integer data and convert among the supported floating point formats as the data are read.

```
val = fread (fid)
val = fread (fid, size)
val = fread (fid, size, precision)
val = fread (fid, size, precision, skip)
val = fread (fid, size, precision, skip, arch)
[val, count] = fread (...)
```

Read binary data from the file specified by the file descriptor *fid*.

The optional argument *size* specifies the amount of data to read and may be one of

Inf	Read as much as possible, returning a column vector.
-----	--

- nr*** Read up to *nr* elements, returning a column vector.
- [*nr*, *Inf*]** Read as much as possible, returning a matrix with *nr* rows. If the number of elements read is not an exact multiple of *nr*, the last column is padded with zeros.
- [*nr*, *nc*]** Read up to *nr * nc* elements, returning a matrix with *nr* rows. If the number of elements read is not an exact multiple of *nr*, the last column is padded with zeros.

If size is omitted, a value of **Inf** is assumed.

The optional argument *precision* is a string specifying the type of data to read and may be one of

"uint8" (default)
8-bit unsigned integer.

"int8"
"integer*1"
8-bit signed integer.

"uint16"
"ushort"
"unsigned short"
16-bit unsigned integer.

"int16"
"integer*2"
"short" 16-bit signed integer.

"uint"
"uint32"
"unsigned int"
"ulong"
"unsigned long"
32-bit unsigned integer.

"int"
"int32"
"integer*4"
"long" 32-bit signed integer.

"uint64" 64-bit unsigned integer.

"int64"
"integer*8"
64-bit signed integer.

"single"
"float"
"float32"
"real*4" 32-bit floating point number.

```

"double"
"float64"
"real*8" 64-bit floating point number.

"char"
"char*1" 8-bit single character.

"uchar"
"unsigned char"
          8-bit unsigned character.

"schar"
"signed char"
          8-bit signed character.

```

The default precision is "uint8".

The *precision* argument may also specify an optional repeat count. For example, '32*single' causes `fread` to read a block of 32 single precision floating point numbers. Reading in blocks is useful in combination with the *skip* argument.

The *precision* argument may also specify a type conversion. For example, 'int16=>int32' causes `fread` to read 16-bit integer values and return an array of 32-bit integer values. By default, `fread` returns a double precision array. The special form '*TYPE' is shorthand for 'TYPE=>TYPE'.

The conversion and repeat counts may be combined. For example, the specification '32*single=>single' causes `fread` to read blocks of single precision floating point values and return an array of single precision values instead of the default array of double precision values.

The optional argument *skip* specifies the number of bytes to skip after each element (or block of elements) is read. If it is not specified, a value of 0 is assumed. If the final block read is not complete, the final skip is omitted. For example,

```
fread (f, 10, "3*single=>single", 8)
```

will omit the final 8-byte skip because the last read will not be a complete block of 3 values.

The optional argument *arch* is a string specifying the data format for the file. Valid values are

```

"native" or "n"
          The format of the current machine.

"ieee-be" or "b"
          IEEE big endian.

"ieee-le" or "l"
          IEEE little endian.

```

If no *arch* is given the value used in the call to `fopen` which created the file descriptor is used. Otherwise, the value specified with `fread` overrides that of `fopen` and determines the data format.

The output argument *val* contains the data read from the file.

The optional return value *count* contains the number of elements read.

See also: [fwrite], page 289, [fgets], page 276, [fgetl], page 276, [fscanf], page 283, [fopen], page 273.

```
fwrite (fid, data)
fwrite (fid, data, precision)
fwrite (fid, data, precision, skip)
fwrite (fid, data, precision, skip, arch)
count = fwrite (...)
```

Write data in binary form to the file specified by the file descriptor *fid*, returning the number of values *count* successfully written to the file.

The argument *data* is a matrix of values that are to be written to the file. The values are extracted in column-major order.

The remaining arguments *precision*, *skip*, and *arch* are optional, and are interpreted as described for **fread**.

The behavior of **fwrite** is undefined if the values in *data* are too large to fit in the specified precision.

See also: [fread], page 286, [fputs], page 275, [fprintf], page 277, [fopen], page 273.

14.2.17 Temporary Files

Sometimes one needs to write data to a file that is only temporary. This is most commonly used when an external program launched from within Octave needs to access data. When Octave exits all temporary files will be deleted, so this step need not be executed manually.

```
[fid, name, msg] = mkstemp ("template")
[fid, name, msg] = mkstemp ("template", delete)
```

Return the file descriptor *fid* corresponding to a new temporary file with a unique name created from *template*.

The last six characters of *template* must be "XXXXXX" and these are replaced with a string that makes the filename unique. The file is then created with mode read/write and permissions that are system dependent (on GNU/Linux systems, the permissions will be 0600 for versions of glibc 2.0.7 and later). The file is opened in binary mode and with the **O_EXCL** flag.

If the optional argument *delete* is supplied and is true, the file will be deleted automatically when Octave exits.

If successful, *fid* is a valid file ID, *name* is the name of the file, and *msg* is an empty string. Otherwise, *fid* is -1, *name* is empty, and *msg* contains a system-dependent error message.

See also: [tempname], page 290, [tempdir], page 290, [**P_tmpdir**], page 290, [tmpfile], page 289, [fopen], page 273.

```
[fid, msg] = tmpfile ()
```

Return the file ID corresponding to a new temporary file with a unique name.

The file is opened in binary read/write ("**w+b**") mode and will be deleted automatically when it is closed or when Octave exits.

If successful, *fid* is a valid file ID and *msg* is an empty string. Otherwise, *fid* is -1 and *msg* contains a system-dependent error message.

See also: [\[tempname\]](#), page 290, [\[mkstemp\]](#), page 289, [\[tempdir\]](#), page 290, [\[P_tmpdir\]](#), page 290.

```
fname = tempname ()
fname = tempname (dir)
fname = tempname (dir, prefix)
```

Return a unique temporary filename as a string.

If *prefix* is omitted, a value of "oct-" is used.

If *dir* is also omitted, the default directory for temporary files ([P_tmpdir](#)) is used. If *dir* is provided, it must exist, otherwise the default directory for temporary files is used.

Programming Note: Because the named file is not opened by [tempname](#), it is possible, though relatively unlikely, that it will not be available by the time your program attempts to open it. If this is a concern, see [tmpfile](#).

See also: [\[mkstemp\]](#), page 289, [\[tempdir\]](#), page 290, [\[P_tmpdir\]](#), page 290, [\[tmpfile\]](#), page 289.

```
dir = tempdir ()
```

Return the name of the host system's directory for temporary files.

The directory name is taken first from the environment variable TMPDIR. If that does not exist the system default returned by [P_tmpdir](#) is used.

See also: [\[P_tmpdir\]](#), page 290, [\[tempname\]](#), page 290, [\[mkstemp\]](#), page 289, [\[tmpfile\]](#), page 289.

```
P_tmpdir ()
```

Return the name of the host system's **default** directory for temporary files.

Programming Note: The value returned by [P_tmpdir](#) is always the default location. This value may not agree with that returned from [tempdir](#) if the user has overridden the default with the TMPDIR environment variable.

See also: [\[tempdir\]](#), page 290, [\[tempname\]](#), page 290, [\[mkstemp\]](#), page 289, [\[tmpfile\]](#), page 289.

14.2.18 End of File and Errors

Once a file has been opened its status can be acquired. As an example the [feof](#) functions determines if the end of the file has been reached. This can be very useful when reading small parts of a file at a time. The following example shows how to read one line at a time from a file until the end has been reached.

```
filename = "myfile.txt";
fid = fopen (filename, "r");
while (! feof (fid) )
    text_line = fgetl (fid);
endwhile
fclose (fid);
```

Note that in some situations it is more efficient to read the entire contents of a file and then process it, than it is to read it line by line. This has the potential advantage of removing the loop in the above code.

`status = feof (fid)`

Return 1 if an end-of-file condition has been encountered for the file specified by file descriptor *fid* and 0 otherwise.

Note that **feof** will only return 1 if the end of the file has already been encountered, not if the next read operation will result in an end-of-file condition.

See also: [\[fread\]](#), page 286, [\[rewind\]](#), page 292, [\[fseek\]](#), page 292, [\[clear\]](#), page 291, [\[open\]](#), page 273.

`msg = ferror (fid)`

`[msg, err] = ferror (fid)`

`[...] = ferror (fid, "clear")`

Query the error status of the stream specified by file descriptor *fid*.

If an error condition exists then return a string *msg* describing the error. Otherwise, return an empty string "".

The second input "clear" is optional. If supplied, the error state on the stream will be cleared.

The optional second output is a numeric indication of the error status. *err* is 1 if an error condition has been encountered and 0 otherwise.

Note that **ferror** indicates if an error has already occurred, not whether the next operation will result in an error condition.

See also: [\[clear\]](#), page 291, [\[open\]](#), page 273.

`fclear (fid)`

Clear the stream state for the file specified by the file descriptor *fid*.

See also: [\[error\]](#), page 291, [\[open\]](#), page 273.

`freport ()`

Print a list of which files have been opened, and whether they are open for reading, writing, or both.

For example:

`freport ()`

+	number	mode	arch	name
+	-----	----	----	----
+	0	r	ieee-le	stdin
+	1	w	ieee-le	stdout
+	2	w	ieee-le	stderr
+	3	r	ieee-le	myfile

See also: [\[open\]](#), page 273, [\[close\]](#), page 275, [\[is_valid_file_id\]](#), page 275.

14.2.19 File Positioning

Three functions are available for setting and determining the position of the file pointer for a given file.

`pos = ftell (fid)`

Return the position of the file pointer as the number of characters from the beginning of the file specified by file descriptor *fid*.

See also: [\[fseek\]](#), page 292, [\[rewind\]](#), page 292, [\[feof\]](#), page 291, [\[fopen\]](#), page 273.

`fseek (fid, offset)`

`fseek (fid, offset, origin)`

`status = fseek (...)`

Set the file pointer to the location *offset* within the file *fid*.

The pointer is positioned *offset* characters from the *origin*, which may be one of the predefined variables `SEEK_SET` (beginning), `SEEK_CUR` (current position), or `SEEK_END` (end of file) or strings "`b0f`", "`c0f`", or "`e0f`". If *origin* is omitted, `SEEK_SET` is assumed. *offset* may be positive, negative, or zero but not all combinations of *origin* and *offset* can be realized.

`fseek` returns 0 on success and -1 on error.

See also: [\[fseek\]](#), page 277, [\[rewind\]](#), page 292, [\[ftell\]](#), page 292, [\[fopen\]](#), page 273, [\[SEEK_SET\]](#), page 292, [\[SEEK_CUR\]](#), page 292, [\[SEEK_END\]](#), page 292.

`SEEK_SET ()`

Return the numerical value to pass to `fseek` to position the file pointer relative to the beginning of the file.

See also: [\[SEEK_CUR\]](#), page 292, [\[SEEK_END\]](#), page 292, [\[fseek\]](#), page 292.

`SEEK_CUR ()`

Return the numerical value to pass to `fseek` to position the file pointer relative to the current position.

See also: [\[SEEK_SET\]](#), page 292, [\[SEEK_END\]](#), page 292, [\[fseek\]](#), page 292.

`SEEK_END ()`

Return the numerical value to pass to `fseek` to position the file pointer relative to the end of the file.

See also: [\[SEEK_SET\]](#), page 292, [\[SEEK_CUR\]](#), page 292, [\[fseek\]](#), page 292.

`rewind (fid)`

`status = rewint (fid)`

Move the file pointer to the beginning of the file specified by file descriptor *fid*.

`rewind` returns 0 for success, and -1 if an error is encountered. It is equivalent to `fseek (fid, 0, SEEK_SET)`.

See also: [\[fseek\]](#), page 292, [\[ftell\]](#), page 292, [\[fopen\]](#), page 273.

The following example stores the current file position in the variable `marker`, moves the pointer to the beginning of the file, reads four characters, and then returns to the original position.

```
marker = ftell (myfile);
rewind (myfile);
fourch = fgets (myfile, 4);
fseek (myfile, marker, SEEK_SET);
```


15 Plotting

15.1 Introduction to Plotting

Earlier versions of Octave provided plotting through the use of gnuplot. This capability is still available. But, a newer plotting capability is provided by access to OpenGL. Which plotting system is used is controlled by the `graphics_toolkit` function. See [Section 15.4.8 \[Graphics Toolkits\]](#), page 473.

The function call `graphics_toolkit ("qt")` selects the Qt/OpenGL system, `graphics_toolkit ("fltk")` selects the FLTK/OpenGL system, and `graphics_toolkit ("gnuplot")` selects the gnuplot system. The three systems may be used selectively through the use of the `graphics_toolkit` property of the `graphics` handle for each figure. This is explained in [Section 15.3 \[Graphics Data Structures\]](#), page 393. **Caution:** The OpenGL-based toolkits use single precision variables internally which limits the maximum value that can be displayed to approximately 10^{38} . If your data contains larger values you must use the gnuplot toolkit which supports values up to 10^{308} .

15.2 High-Level Plotting

Octave provides simple means to create many different types of two- and three-dimensional plots using high-level functions.

If you need more detailed control, see [Section 15.3 \[Graphics Data Structures\]](#), page 393, and [Section 15.4 \[Advanced Plotting\]](#), page 456.

15.2.1 Two-Dimensional Plots

The `plot` function allows you to create simple x-y plots with linear axes. For example,

```
x = -10:0.1:10;
plot (x, sin (x));
xlabel ("x");
ylabel ("sin (x)");
title ("Simple 2-D Plot");
```

displays a sine wave shown in [Figure 15.1](#). On most systems, this command will open a separate plot window to display the graph.




Figure 15.1: Simple Two-Dimensional Plot.

```
plot (y)
plot (x, y)
plot (x, y, fmt)
plot (..., property, value, ...)
plot (x1, y1, ..., xn, yn)
plot (hax, ...)
h = plot (...)
    Produce 2-D plots.
```

Many different combinations of arguments are possible. The simplest form is

```
plot (y)
```

where the argument is taken as the set of y coordinates and the x coordinates are taken to be the range `1:numel (y)`.

If more than one argument is given, they are interpreted as

```
plot (y, property, value, ...)
```

or

```
plot (x, y, property, value, ...)
```

or

```
plot (x, y, fmt, ...)
```

and so on. Any number of argument sets may appear. The x and y values are interpreted as follows:

- If a single data argument is supplied, it is taken as the set of y coordinates and the x coordinates are taken to be the indices of the elements, starting with 1.
- If x and y are scalars, a single point is plotted.
- `squeeze()` is applied to arguments with more than two dimensions, but no more than two singleton dimensions.

- If both arguments are vectors, the elements of y are plotted versus the elements of x .
- If x is a vector and y is a matrix, then the columns (or rows) of y are plotted versus x . (using whichever combination matches, with columns tried first.)
- If the x is a matrix and y is a vector, y is plotted versus the columns (or rows) of x . (using whichever combination matches, with columns tried first.)
- If both arguments are matrices, the columns of y are plotted versus the columns of x . In this case, both matrices must have the same number of rows and columns and no attempt is made to transpose the arguments to make the number of rows match.

Multiple property-value pairs may be specified, but they must appear in pairs. These arguments are applied to the line objects drawn by `plot`. Useful properties to modify are `"linestyle"`, `"linewidth"`, `"color"`, `"marker"`, `"markersize"`, `"markeredgecolor"`, `"markerfacecolor"`. See [Section 15.3.3.4 \[Line Properties\]](#), page 419.

The *fmt* format argument can also be used to control the plot style. It is a string composed of four optional parts: "`<linestyle><marker><color><;displayname;>`". When a marker is specified, but no linestyle, only the markers are plotted. Similarly, if a linestyle is specified, but no marker, then only lines are drawn. If both are specified then lines and markers will be plotted. If no *fmt* and no *property/value* pairs are given, then the default plot style is solid lines with no markers and the color determined by the `"colororder"` property of the current axes.

Format arguments:

`linestyle`

'-	Use solid lines (default).
'--'	Use dashed lines.
':'	Use dotted lines.
'-.'	Use dash-dotted lines.

`marker`

'+'	crosshair
'o'	circle
'*'	star
'.'	point
'x'	cross
's'	square
'd'	diamond
'^'	upward-facing triangle
'v'	downward-facing triangle
right-facing triangle	
'<'	left-facing triangle
'p'	pentagram
'h'	hexagram

`color`

```

'k'    blacK
'r'    Red
'g'    Green
'b'    Blue
'y'    Yellow
'm'    Magenta
'c'    Cyan
'w'    White

```

";*displayname*;"

Here "*displayname*" is the label to use for the plot legend.

The *fmt* argument may also be used to assign legend labels. To do so, include the desired label between semicolons after the formatting sequence described above, e.g., "+b;Key Title;". Note that the last semicolon is required and Octave will generate an error if it is left out.

Here are some plot examples:

```
plot (x, y, "or", x, y2, x, y3, "m", x, y4, "+")
```

This command will plot y with red circles, y2 with solid lines, y3 with solid magenta lines, and y4 with points displayed as '+'.

```
plot (b, "*", "markersize", 10)
```

This command will plot the data in the variable b, with points displayed as '*' and a marker size of 10.

```
t = 0:0.1:6.3;
plot (t, cos(t), "-;cos(t);", t, sin(t), "-b;sin(t);");
```

This will plot the cosine and sine functions and label them accordingly in the legend.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a vector of graphics handles to the created line objects.

To save a plot, in one of several image formats such as PostScript or PNG, use the `print` command.

See also: [\[axis\]](#), page 323, [\[box\]](#), page 367, [\[grid\]](#), page 367, [\[hold\]](#), page 375, [\[legend\]](#), page 364, [\[title\]](#), page 363, [\[xlabel\]](#), page 366, [\[ylabel\]](#), page 366, [\[xlim\]](#), page 325, [\[ylim\]](#), page 325, [\[ezplot\]](#), page 328, [\[errorbar\]](#), page 313, [\[fplot\]](#), page 327, [\[line\]](#), page 396, [\[plot3\]](#), page 348, [\[polar\]](#), page 317, [\[loglog\]](#), page 300, [\[semilogx\]](#), page 299, [\[semilogy\]](#), page 299, [\[subplot\]](#), page 371.

The `plotyy` function may be used to create a plot with two independent y axes.

```

plotyy (x1, y1, x2, y2)
plotyy (... , fun)
plotyy (... , fun1, fun2)
plotyy (hax, ...)
[ax, h1, h2] = plotyy (... )

```

Plot two sets of data with independent y-axes and a common x-axis.

The arguments $x1$ and $y1$ define the arguments for the first plot and $x2$ and $y2$ for the second.

By default the arguments are evaluated with `feval (@plot, x, y)`. However the type of plot can be modified with the *fun* argument, in which case the plots are generated by `feval (fun, x, y)`. *fun* can be a function handle, an inline function, or a string of a function name.

The function to use for each of the plots can be independently defined with *fun1* and *fun2*.

If the first argument *hax* is an axes handle, then it defines the principal axes in which to plot the $x1$ and $y1$ data.

The return value *ax* is a vector with the axes handles of the two y-axes. *h1* and *h2* are handles to the objects generated by the plot commands.

```
x = 0:0.1:2*pi;
y1 = sin (x);
y2 = exp (x - 1);
ax = plotyy (x, y1, x - 1, y2, @plot, @semilogy);
xlabel ("X");
ylabel (ax(1), "Axis 1");
ylabel (ax(2), "Axis 2");
```

See also: [\[plot\]](#), page 296.

The functions `semilogx`, `semilogy`, and `loglog` are similar to the `plot` function, but produce plots in which one or both of the axes use log scales.

```
semilogx (y)
semilogx (x, y)
semilogx (x, y, property, value, ...)
semilogx (x, y, fmt)
semilogx (hax, ...)
h = semilogx (...)
```

Produce a 2-D plot using a logarithmic scale for the x-axis.

See the documentation of `plot` for a description of the arguments that `semilogx` will accept.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the created plot.

See also: [\[plot\]](#), page 296, [\[semilogy\]](#), page 299, [\[loglog\]](#), page 300.

```
semilogy (y)
semilogy (x, y)
semilogy (x, y, property, value, ...)
semilogy (x, y, fmt)
semilogy (h, ...)
h = semilogy (...)
```

Produce a 2-D plot using a logarithmic scale for the y-axis.

See the documentation of `plot` for a description of the arguments that `semilogy` will accept.

If the first argument `hax` is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value `h` is a graphics handle to the created plot.

See also: `[plot]`, page 296, `[semilogx]`, page 299, `[loglog]`, page 300.

```
loglog (y)
loglog (x, y)
loglog (x, y, prop, value, ...)
loglog (x, y, fmt)
loglog (hax, ...)
h = loglog (...)
```

Produce a 2-D plot using logarithmic scales for both axes.

See the documentation of `plot` for a description of the arguments that `loglog` will accept.

If the first argument `hax` is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value `h` is a graphics handle to the created plot.

See also: `[plot]`, page 296, `[semilogx]`, page 299, `[semilogy]`, page 299.

The functions `bar`, `barh`, `stairs`, and `stem` are useful for displaying discrete data. For example,

```
randn ("state", 1);
hist (randn (10000, 1), 30);
xlabel ("Value");
ylabel ("Count");
title ("Histogram of 10,000 normally distributed random numbers");
```

produces the histogram of 10,000 normally distributed random numbers shown in [Figure 15.2](#). Note that, `randn ("state", 1);`, initializes the random number generator for `randn` to a known value so that the returned values are reproducible; This guarantees that the figure produced is identical to the one in this manual.




Figure 15.2: Histogram.

```

bar (y)
bar (x, y)
bar (..., w)
bar (..., style)
bar (..., prop, val, ...)
bar (hax, ...)
h = bar (..., prop, val, ...)

    Produce a bar graph from two vectors of X-Y data.

If only one argument is given, y, it is taken as a vector of Y values and the X coordinates are the range 1:numel (y).

The optional input w controls the width of the bars. A value of 1.0 will cause each bar to exactly touch any adjacent bars. The default width is 0.8.

If y is a matrix, then each column of y is taken to be a separate bar graph plotted on the same graph. By default the columns are plotted side-by-side. This behavior can be changed by the style argument which can take the following values:

"grouped" (default)
    Side-by-side bars with a gap between bars and centered over the X-coordinate.

"stacked"
    Bars are stacked so that each X value has a single bar composed of multiple segments.

"hist"
    Side-by-side bars with no gap between bars and centered over the X-coordinate.

"histc"
    Side-by-side bars with no gap between bars and left-aligned to the X-coordinate.

```

Optional property/value pairs are passed directly to the underlying patch objects. If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a vector of handles to the created "bar series" hggroups with one handle per column of the variable *y*. This series makes it possible to change a common element in one bar series object and have the change reflected in the other "bar series". For example,

```
h = bar (rand (5, 10));
set (h(1), "basevalue", 0.5);
```

changes the position on the base of all of the bar series.

The following example modifies the face and edge colors using property/value pairs.

```
bar (randn (1, 100), "facecolor", "r", "edgecolor", "b");
```

The color of the bars is taken from the figure's colormap, such that

```
bar (rand (10, 3));
colormap (summer (64));
```

will change the colors used for the bars. The color of bars can also be set manually using the "facecolor" property as shown below.

```
h = bar (rand (10, 3));
set (h(1), "facecolor", "r")
set (h(2), "facecolor", "g")
set (h(3), "facecolor", "b")
```

See also: `[barh]`, page 302, `[hist]`, page 303, `[pie]`, page 317, `[plot]`, page 296, `[patch]`, page 397.

```
barh (y)
barh (x, y)
barh (..., w)
barh (..., style)
barh (..., prop, val, ...)
barh (hax, ...)
h = barh (..., prop, val, ...)
```

Produce a horizontal bar graph from two vectors of X-Y data.

If only one argument is given, it is taken as a vector of Y values and the X coordinates are the range `1:numel (y)`.

The optional input *w* controls the width of the bars. A value of 1.0 will cause each bar to exactly touch any adjacent bars. The default width is 0.8.

If *y* is a matrix, then each column of *y* is taken to be a separate bar graph plotted on the same graph. By default the columns are plotted side-by-side. This behavior can be changed by the *style* argument which can take the following values:

"grouped" (default)

Side-by-side bars with a gap between bars and centered over the Y-coordinate.

"stacked"

Bars are stacked so that each Y value has a single bar composed of multiple segments.

"hist" Side-by-side bars with no gap between bars and centered over the Y-coordinate.

"histc" Side-by-side bars with no gap between bars and left-aligned to the Y-coordinate.

Optional property/value pairs are passed directly to the underlying patch objects.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by *gca*.

The optional return value *h* is a graphics handle to the created bar series hggroup. For a description of the use of the bar series, see [\[bar\], page 301](#).

See also: [\[bar\], page 301](#), [\[hist\], page 303](#), [\[pie\], page 317](#), [\[plot\], page 296](#), [\[patch\], page 397](#).

```
hist (y)
hist (y, nbins)
hist (y, x)
hist (y, x, norm)
hist (..., prop, val, ...)
hist (hax, ...)
[nn, xx] = hist (...)
```

Produce histogram counts or plots.

With one vector input argument, *y*, plot a histogram of the values with 10 bins. The range of the histogram bins is determined by the range of the data (difference between maximum and minimum value in *y*). Extreme values are lumped into the first and last bins. If *y* is a matrix then plot a histogram where each bin contains one bar per input column of *y*.

If the optional second argument is a scalar, *nbins*, it defines the number of bins.

If the optional second argument is a vector, *x*, it defines the centers of the bins. The width of the bins is determined from the adjacent values in the vector. The total number of bins is *numel* (*x*).

If a third argument is provided, the histogram is normalized such that the sum of the bars is equal to *norm*.

The histogram's appearance may be modified by specifying property/value pairs. For example, the face and edge color may be modified:

```
hist (randn (1, 100), 25, "facecolor", "r", "edgecolor", "b");
```

The histogram's colors also depend upon the current colormap.

```
hist (rand (10, 3));
colormap (summer ());
```

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by *gca*.

If an output is requested then no plot is made. Instead, return the values *nn* (numbers of elements) and *xx* (bin centers) such that `bar (xx, nn)` will plot the histogram.

See also: `[histc]`, page 736, `[bar]`, page 301, `[pie]`, page 317, `[rose]`, page 309.

```
stemleaf (x, caption)
stemleaf (x, caption, stem_sz)
plotstr = stemleaf (...)
```

Compute and display a stem and leaf plot of the vector *x*.

The input *x* should be a vector of integers. Any non-integer values will be converted to integer by `x = fix (x)`. By default each element of *x* will be plotted with the last digit of the element as a leaf value and the remaining digits as the stem. For example, 123 will be plotted with the stem '12' and the leaf '3'. The second argument, *caption*, should be a character array which provides a description of the data. It is included as a heading for the output.

The optional input *stem_sz* sets the width of each stem. The stem width is determined by $10^{(stem_sz + 1)}$. The default stem width is 10.

The output of `stemleaf` is composed of two parts: a "Fenced Letter Display," followed by the stem-and-leaf plot itself. The Fenced Letter Display is described in *Exploratory Data Analysis*. Briefly, the entries are as shown:

Fenced Letter Display	
#% nx -----	nx = numel (x)
M% mi md	mi median index, md median
H% hi hl	hi lower hinge index, hl,hu hinges,
1 x(1)	hs h_spreadx(1), x(nx) first
	and last data value.
----- step -----	step 1.5*h_spread
f ifl	inner fence, lower and higher
nfl	no.\ of data points within fences
F ofl	outer fence, lower and higher
nFl	no.\ of data points outside outer
	fences

The stem-and-leaf plot shows on each line the stem value followed by the string made up of the leaf digits. If the *stem_sz* is not 1 the successive leaf values are separated by ",".

With no return argument, the plot is immediately displayed. If an output argument is provided, the plot is returned as an array of strings.

The leaf digits are not sorted. If sorted leaf values are desired, use `xs = sort (x)` before calling `stemleaf (xs)`.

The stem and leaf plot and associated displays are described in: Chapter 3, *Exploratory Data Analysis* by J. W. Tukey, Addison-Wesley, 1977.

See also: `[hist]`, page 303, `[printd]`, page 305.

```
printfd (obj, filename)
out_file = printfd (...)
```

Convert any object acceptable to `disp` into the format selected by the suffix of `filename`.

If the return argument `out_file` is given, the name of the created file is returned.

This function is intended to facilitate manipulation of the output of functions such as `stemleaf`.

See also: [stemleaf], page 304.

```
stairs (y)
stairs (x, y)
stairs (..., style)
stairs (..., prop, val, ...)
stairs (hax, ...)
h = stairs (...)
[xstep, ystep] = stairs (...)
```

Produce a stairstep plot.

The arguments `x` and `y` may be vectors or matrices. If only one argument is given, it is taken as a vector of Y values and the X coordinates are taken to be the indices of the elements (`x = 1:numel (y)`).

The style to use for the plot can be defined with a line style `style` of the same format as the `plot` command.

Multiple property/value pairs may be specified, but they must appear in pairs.

If the first argument `hax` is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

If one output argument is requested, return a graphics handle to the created plot. If two output arguments are specified, the data are generated but not plotted. For example,

```
stairs (x, y);
```

and

```
[xs, ys] = stairs (x, y);
plot (xs, ys);
```

are equivalent.

See also: [bar], page 301, [hist], page 303, [plot], page 296, [stem], page 305.

```
stem (y)
stem (x, y)
stem (..., linespec)
stem (..., "filled")
stem (..., prop, val, ...)
stem (hax, ...)
h = stem (...)
Plot a 2-D stem graph.
```

If only one argument is given, it is taken as the y-values and the x-coordinates are taken from the indices of the elements.

If *y* is a matrix, then each column of the matrix is plotted as a separate stem graph. In this case *x* can either be a vector, the same length as the number of rows in *y*, or it can be a matrix of the same size as *y*.

The default color is "**b**" (blue), the default line style is **"-"**, and the default marker is "**o**". The line style can be altered by the *linespec* argument in the same manner as the *plot* command. If the "**filled**" argument is present the markers at the top of the stems will be filled in. For example,

```
x = 1:10;
y = 2*x;
stem (x, y, "r");
```

plots 10 stems with heights from 2 to 20 in red;

Optional property/value pairs may be specified to control the appearance of the plot.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by *gca*.

The optional return value *h* is a handle to a "stem series" hggroup. The single hggroup handle has all of the graphical elements comprising the plot as its children; This allows the properties of multiple graphics objects to be changed by modifying just a single property of the "stem series" hggroup.

For example,

```
x = [0:10]';
y = [sin(x), cos(x)]
h = stem (x, y);
set (h(2), "color", "g");
set (h(1), "basevalue", -1)
```

changes the color of the second "stem series" and moves the base line of the first.

Stem Series Properties

<code>linestyle</code>	The linestyle of the stem. (Default: "-")
<code>linewidth</code>	The width of the stem. (Default: 0.5)
<code>color</code>	The color of the stem, and if not separately specified, the marker. (Default: " b " [blue])
<code>marker</code>	The marker symbol to use at the top of each stem. (Default: " o ")
<code>markeredgecolor</code>	The edge color of the marker. (Default: <code>"color"</code> property)
<code>markerfacecolor</code>	The color to use for "filling" the marker. (Default: <code>"none"</code> [unfilled])
<code>markersize</code>	The size of the marker. (Default: 6)
<code>baseline</code>	The handle of the line object which implements the baseline. Use <code>set</code> with the returned handle to change graphic properties of the baseline.

`basevalue` The y-value where the baseline is drawn. (Default: 0)

See also: [stem3], page 307, [bar], page 301, [hist], page 303, [plot], page 296, [stairs], page 305.

```
stem3 (x, y, z)
stem3 (..., linespec)
stem3 (..., "filled")
stem3 (..., prop, val, ...)
stem3 (hax, ...)
h = stem3 (...)
```

Plot a 3-D stem graph.

Stems are drawn from the height z to the location in the x-y plane determined by x and y . The default color is "b" (blue), the default line style is "-", and the default marker is "o".

The line style can be altered by the `linespec` argument in the same manner as the `plot` command. If the "filled" argument is present the markers at the top of the stems will be filled in.

Optional property/value pairs may be specified to control the appearance of the plot.

If the first argument `hax` is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value `h` is a handle to the "stem series" hggroup containing the line and marker objects used for the plot. See [stem], page 305, for a description of the "stem series" object.

Example:

```
theta = 0:0.2:6;
stem3 (cos (theta), sin (theta), theta);
```

plots 31 stems with heights from 0 to 6 lying on a circle.

Implementation Note: Color definitions with RGB-triples are not valid.

See also: [stem], page 305, [bar], page 301, [hist], page 303, [plot], page 296.

```
scatter (x, y)
scatter (x, y, s)
scatter (x, y, s, c)
scatter (..., style)
scatter (..., "filled")
scatter (..., prop, val, ...)
scatter (hax, ...)
h = scatter (...)
```

Draw a 2-D scatter plot.

A marker is plotted at each point defined by the coordinates in the vectors x and y .

The size of the markers is determined by s , which can be a scalar or a vector of the same length as x and y . If s is not given, or is an empty matrix, then a default value of 36 square points is used (The marker size itself is `sqrt (s)`).

The color of the markers is determined by *c*, which can be a string defining a fixed color; a 3-element vector giving the red, green, and blue components of the color; a vector of the same length as *x* that gives a scaled index into the current colormap; or an Nx3 matrix defining the RGB color of each marker individually.

The marker to use can be changed with the *style* argument; it is a string defining a marker in the same manner as the `plot` command. If no marker is specified it defaults to "o" or circles. If the argument "filled" is given then the markers are filled.

Additional property/value pairs are passed directly to the underlying patch object.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the created scatter object.

Example:

```
x = randn (100, 1);
y = randn (100, 1);
scatter (x, y, [], sqrt (x.^2 + y.^2));
```

See also: `[scatter3]`, page 356, `[patch]`, page 397, `[plot]`, page 296.

```
plotmatrix (x, y)
plotmatrix (x)
plotmatrix (..., style)
plotmatrix (hax, ...)
[h, ax, bigax, p, pax] = plotmatrix (...)
```

Scatter plot of the columns of one matrix against another.

Given the arguments *x* and *y* that have a matching number of rows, `plotmatrix` plots a set of axes corresponding to

```
plot (x(:, i), y(:, j))
```

When called with a single argument *x* this is equivalent to

```
plotmatrix (x, x)
```

except that the diagonal of the set of axes will be replaced with the histogram `hist (x(:, i))`.

The marker to use can be changed with the *style* argument, that is a string defining a marker in the same manner as the `plot` command.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* provides handles to the individual graphics objects in the scatter plots, whereas *ax* returns the handles to the scatter plot axes objects.

bigax is a hidden axes object that surrounds the other axes, such that the commands `xlabel`, `title`, etc., will be associated with this hidden axes.

Finally, *p* returns the graphics objects associated with the histogram and *pax* the corresponding axes objects.

Example:

```
plotmatrix (randn (100, 3), "g+")
```

See also: `[scatter]`, page 307, `[plot]`, page 296.

```
pareto (y)
pareto (y, x)
pareto (hax, ...)
h = pareto (...)
```

Draw a Pareto chart.

A Pareto chart is a bar graph that arranges information in such a way that priorities for process improvement can be established; It organizes and displays information to show the relative importance of data. The chart is similar to the histogram or bar chart, except that the bars are arranged in decreasing magnitude from left to right along the x-axis.

The fundamental idea (Pareto principle) behind the use of Pareto diagrams is that the majority of an effect is due to a small subset of the causes. For quality improvement, the first few contributing causes (leftmost bars as presented on the diagram) to a problem usually account for the majority of the result. Thus, targeting these "major causes" for elimination results in the most cost-effective improvement scheme.

Typically only the magnitude data *y* is present in which case *x* is taken to be the range `1 : length (y)`. If *x* is given it may be a string array, a cell array of strings, or a numerical vector.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a 2-element vector with a graphics handle for the created bar plot and a second handle for the created line plot.

An example of the use of `pareto` is

```
Cheese = {"Cheddar", "Swiss", "Camembert", ...
          "Munster", "Stilton", "Blue"};
Sold = [105, 30, 70, 10, 15, 20];
pareto (Sold, Cheese);
```

See also: `[bar]`, page 301, `[barh]`, page 302, `[hist]`, page 303, `[pie]`, page 317, `[plot]`, page 296.

```
rose (th)
rose (th, nbins)
rose (th, bins)
rose (hax, ...)
h = rose (...)
```

`[thout rout] = rose (...)`

Plot an angular histogram.

With one vector argument, *th*, plot the histogram with 20 angular bins. If *th* is a matrix then each column of *th* produces a separate histogram.

If *nbins* is given and is a scalar, then the histogram is produced with *nbins* bins. If *bins* is a vector, then the center of each bin is defined by the values in *bins* and the number of bins is given by the number of elements in *bins*.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a vector of graphics handles to the line objects representing each histogram.

If two output arguments are requested then no plot is made and the polar vectors necessary to plot the histogram are returned instead.

Example

```
[th, r] = rose ([2*randn(1e5,1), pi + 2*randn(1e5,1)]);
polar (th, r);
```

Programming Note: When specifying bin centers with the *bins* input, the edges for bins 2 to N-1 are spaced so that *bins*(*i*) is centered between the edges. The final edge is drawn halfway between bin N and bin 1. This guarantees that all input *th* will be placed into one of the bins, but also means that for some combinations bin 1 and bin N may not be centered on the user's given values.

See also: [\[hist\]](#), page 303, [\[polar\]](#), page 317.

The **contour**, **contourf** and **contourc** functions produce two-dimensional contour plots from three-dimensional data.

```
contour (z)
contour (z, vn)
contour (x, y, z)
contour (x, y, z, vn)
contour (..., style)
contour (hax, ...)
[c, h] = contour (...)
```

Create a 2-D contour plot.

Plot level curves (contour lines) of the matrix *z*, using the contour matrix *c* computed by **contourc** from the same arguments; see the latter for their interpretation.

The appearance of contour lines can be defined with a line style *style* in the same manner as **plot**. Only line style and color are used; Any markers defined by *style* are ignored.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by **gca**.

The optional output *c* contains the contour levels in **contourc** format.

The optional return value *h* is a graphics handle to the hggroup comprising the contour lines.

Example:

```
x = 0:2;
y = x;
z = x' * y;
contour (x, y, z, 2:3)
```

See also: [\[ezcontour\]](#), page 329, [\[contourc\]](#), page 311, [\[contourf\]](#), page 311, [\[contour3\]](#), page 312, [\[clabel\]](#), page 366, [\[meshc\]](#), page 334, [\[surf\]](#), page 336, [\[caxis\]](#), page 325, [\[colormap\]](#), page 816, [\[plot\]](#), page 296.

```
contourf (z)
contourf (z, vn)
contourf (x, y, z)
contourf (x, y, z, vn)
contourf (..., style)
contourf (hax, ...)
[c, h] = contourf (...)
```

Create a 2-D contour plot with filled intervals.

Plot level curves (contour lines) of the matrix *z* and fill the region between lines with colors from the current colormap.

The level curves are taken from the contour matrix *c* computed by `contourc` for the same arguments; see the latter for their interpretation.

The appearance of contour lines can be defined with a line style *style* in the same manner as `plot`. Only line style and color are used; Any markers defined by *style* are ignored.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional output *c* contains the contour levels in `contourc` format.

The optional return value *h* is a graphics handle to the hggroup comprising the contour lines.

The following example plots filled contours of the `peaks` function.

```
[x, y, z] = peaks (50);
contourf (x, y, z, -7:9)
```

See also: `[ezcontourf]`, page 330, `[contour]`, page 310, `[contourc]`, page 311, `[contour3]`, page 312, `[clabel]`, page 366, `[meshc]`, page 334, `[surf]`, page 336, `[caxis]`, page 325, `[colormap]`, page 816, `[plot]`, page 296.

```
[c, lev] = contourc (z)
[c, lev] = contourc (z, vn)
[c, lev] = contourc (x, y, z)
[c, lev] = contourc (x, y, z, vn)
```

Compute contour lines (isolines of constant Z value).

The matrix *z* contains height values above the rectangular grid determined by *x* and *y*. If only a single input *z* is provided then *x* is taken to be `1:columns (z)` and *y* is taken to be `1:rows (z)`. The minimum data size is 2x2.

The optional input *vn* is either a scalar denoting the number of contour lines to compute or a vector containing the Z values where lines will be computed. When *vn* is a vector the number of contour lines is `numel (vn)`. However, to compute a single contour line at a given value use *vn* = `[val, val]`. If *vn* is omitted it defaults to 10.

The return value *c* is a $2 \times n$ matrix containing the contour lines in the following format

```
c = [lev1, x1, x2, ..., levn, x1, x2, ...
     len1, y1, y2, ..., lenn, y1, y2, ...]
```

in which contour line *n* has a level (height) of *levn* and length of *lenn*.

The optional return value *lev* is a vector with the Z values of the contour levels.

Example:

```
x = 0:2;
y = x;
z = x' * y;
c = contourc (x, y, z, 2:3)
⇒ c =
    2.0000   1.0000   1.0000   2.0000   2.0000   3.0000   1.5000   2.0000■
    4.0000   2.0000   2.0000   1.0000   1.0000   2.0000   2.0000   1.5000■
```

See also: [contour], page 310, [contourf], page 311, [contour3], page 312, [clabel], page 366.

```
contour3 (z)
contour3 (z, vn)
contour3 (x, y, z)
contour3 (x, y, z, vn)
contour3 (..., style)
contour3 (hax, ...)
[c, h] = contour3 (...)
```

Create a 3-D contour plot.

`contour3` plots level curves (contour lines) of the matrix `z` at a Z level corresponding to each contour. This is in contrast to `contour` which plots all of the contour lines at the same Z level and produces a 2-D plot.

The level curves are taken from the contour matrix `c` computed by `contourc` for the same arguments; see the latter for their interpretation.

The appearance of contour lines can be defined with a line style `style` in the same manner as `plot`. Only line style and color are used; Any markers defined by `style` are ignored.

If the first argument `hax` is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional output `c` are the contour levels in `contourc` format.

The optional return value `h` is a graphics handle to the hggroup comprising the contour lines.

Example:

```
contour3 (peaks (19));
colormap cool;
hold on;
surf (peaks (19), "facecolor", "none", "edgecolor", "black");
```

See also: [contour], page 310, [contourc], page 311, [contourf], page 311, [clabel], page 366, [meshc], page 334, [surfc], page 336, [caxis], page 325, [colormap], page 816, [plot], page 296.


The `errorbar`, `semilogxerr`, `semilogyerr`, and `loglogerr` functions produce plots with error bar markers. For example,

```

rand ("state", 2);
x = 0:0.1:10;
y = sin (x);
lerr = 0.1 .* rand (size (x));
uerr = 0.1 .* rand (size (x));
errorbar (x, y, lerr, uerr);
axis ([0, 10, -1.1, 1.1]);
xlabel ("x");
ylabel ("sin (x)");
title ("Errorbar plot of sin (x)");

```

produces the figure shown in [Figure 15.3](#).



[Figure 15.3: Errorbar plot.](#)

```

errorbar (y, ey)
errorbar (y, ..., fmt)
errorbar (x, y, ey)
errorbar (x, y, err, fmt)
errorbar (x, y, lerr, uerr, fmt)
errorbar (x, y, ex, ey, fmt)
errorbar (x, y, lx, ux, ly, uy, fmt)
errorbar (x1, y1, ..., fmt, xn, yn, ...)
errorbar (hax, ...)
h = errorbar (...)

```

Create a 2-D plot with errorbars.

Many different combinations of arguments are possible. The simplest form is

```
errorbar (y, ey)
```

where the first argument is taken as the set of y coordinates, the second argument ey are the errors around the y values, and the x coordinates are taken to be the indices of the elements ($1:\text{numel}(\text{y})$).

The general form of the function is

```
errorbar (x, y, err1, ..., fmt, ...)
```

After the x and y arguments there can be 1, 2, or 4 parameters specifying the error values depending on the nature of the error values and the plot format fmt .

err (scalar)

When the error is a scalar all points share the same error value. The errorbars are symmetric and are drawn from $\text{data}-\text{err}$ to $\text{data}+\text{err}$. The fmt argument determines whether err is in the x -direction, y -direction (default), or both.

err (vector or matrix)

Each data point has a particular error value. The errorbars are symmetric and are drawn from $\text{data}(n)-\text{err}(n)$ to $\text{data}(n)+\text{err}(n)$.

lerr, uerr (scalar)

The errors have a single low-side value and a single upper-side value. The errorbars are not symmetric and are drawn from $\text{data}-\text{lerr}$ to $\text{data}+\text{uerr}$.

lerr, uerr (vector or matrix)

Each data point has a low-side error and an upper-side error. The errorbars are not symmetric and are drawn from $\text{data}(n)-\text{lerr}(n)$ to $\text{data}(n)+\text{uerr}(n)$.

Any number of data sets ($x1,y1, x2,y2, \dots$) may appear as long as they are separated by a format string fmt .

If y is a matrix, x and the error parameters must also be matrices having the same dimensions. The columns of y are plotted versus the corresponding columns of x and errorbars are taken from the corresponding columns of the error parameters.

If fmt is missing, the yerrorbars ("~") plot style is assumed.

If the fmt argument is supplied then it is interpreted, as in normal plots, to specify the line style, marker, and color. In addition, fmt may include an errorbar style which **must precede** the ordinary format codes. The following errorbar styles are supported:

- '~' Set yerrorbars plot style (default).
- '>' Set xerrorbars plot style.
- '~>' Set xyerrorbars plot style.
- '#~' Set yboxes plot style.
- '#' Set xboxes plot style.
- '#~>' Set xyboxes plot style.

If the first argument hax is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a handle to the hggroup object representing the data plot and errorbars.

Note: For compatibility with MATLAB a line is drawn through all data points. However, most scientific errorbar plots are a scatter plot of points with errorbars. To accomplish this, add a marker style to the *fmt* argument such as ". ". Alternatively, remove the line by modifying the returned graphic handle with `set(h, "linestyle", "none")`.

Examples:

```
errorbar (x, y, ex, ">.r")
```

produces an xerrorbar plot of *y* versus *x* with *x* errorbars drawn from *x-ex* to *x+ex*. The marker ". " is used so no connecting line is drawn and the errorbars appear in red.

```
errorbar (x, y1, ey, "~",
          x, y2, ly, uy)
```

produces yerrorbar plots with *y1* and *y2* versus *x*. Errorbars for *y1* are drawn from *y1-ey* to *y1+ey*, errorbars for *y2* from *y2-ly* to *y2+uy*.

```
errorbar (x, y, lx, ux,
          ly, uy, "~>")
```

produces an xyerrorbar plot of *y* versus *x* in which *x* errorbars are drawn from *x-lx* to *x+ux* and *y* errorbars from *y-ly* to *y+uy*.

See also: [\[semilogxerr\]](#), page 315, [\[semilogyerr\]](#), page 316, [\[loglogerr\]](#), page 316, [\[plot\]](#), page 296.

```
semilogxerr (y, ey)
semilogxerr (y, ..., fmt)
semilogxerr (x, y, ey)
semilogxerr (x, y, err, fmt)
semilogxerr (x, y, lerr, uerr, fmt)
semilogxerr (x, y, ex, ey, fmt)
semilogxerr (x, y, lx, ux, ly, uy, fmt)
semilogxerr (x1, y1, ..., fmt, xn, yn, ...)
semilogxerr (hax, ...)
h = semilogxerr (...)
```

Produce 2-D plots using a logarithmic scale for the x-axis and errorbars at each data point.

Many different combinations of arguments are possible. The most common form is

```
semilogxerr (x, y, ey, fmt)
```

which produces a semi-logarithmic plot of *y* versus *x* with errors in the *y*-scale defined by *ey* and the plot format defined by *fmt*. See [\[errorbar\]](#), page 313, for available formats and additional information.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

See also: [\[errorbar\]](#), page 313, [\[semilogyerr\]](#), page 316, [\[loglogerr\]](#), page 316.

```

semilogyerr (y, ey)
semilogyerr (y, ..., fmt)
semilogyerr (x, y, ey)
semilogyerr (x, y, err, fmt)
semilogyerr (x, y, lerr, uerr, fmt)
semilogyerr (x, y, ex, ey, fmt)
semilogyerr (x, y, lx, ux, ly, uy, fmt)
semilogyerr (x1, y1, ..., fmt, xn, yn, ...)
semilogyerr (hax, ...)
h = semilogyerr (...)

```

Produce 2-D plots using a logarithmic scale for the y-axis and errorbars at each data point.

Many different combinations of arguments are possible. The most common form is

```
semilogyerr (x, y, ey, fmt)
```

which produces a semi-logarithmic plot of y versus x with errors in the y-scale defined by ey and the plot format defined by fmt . See [errorbar], page 313, for available formats and additional information.

If the first argument hax is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

See also: [errorbar], page 313, [semilogxerr], page 315, [loglogerr], page 316.

```

loglogerr (y, ey)
loglogerr (y, ..., fmt)
loglogerr (x, y, ey)
loglogerr (x, y, err, fmt)
loglogerr (x, y, lerr, uerr, fmt)
loglogerr (x, y, ex, ey, fmt)
loglogerr (x, y, lx, ux, ly, uy, fmt)
loglogerr (x1, y1, ..., fmt, xn, yn, ...)
loglogerr (hax, ...)
h = loglogerr (...)

```

Produce 2-D plots on a double logarithm axis with errorbars.

Many different combinations of arguments are possible. The most common form is

```
loglogerr (x, y, ey, fmt)
```

which produces a double logarithm plot of y versus x with errors in the y-scale defined by ey and the plot format defined by fmt . See [errorbar], page 313, for available formats and additional information.

If the first argument hax is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

See also: [errorbar], page 313, [semilogxerr], page 315, [semilogyerr], page 316.

Finally, the `polar` function allows you to easily plot data in polar coordinates. However, the display coordinates remain rectangular and linear. For example,

```

polar (0:0.1:10*pi, 0:0.1:10*pi);
title ("Example polar plot from 0 to 10*pi");

```

produces the spiral plot shown in Figure 15.4.




Figure 15.4: Polar plot.

```
polar (theta, rho)
polar (theta, rho, fmt)
polar (cplx)
polar (cplx, fmt)
polar (hax, ...)
h = polar (...)
```

Create a 2-D plot from polar coordinates *theta* and *rho*.

If a single complex input *cplx* is given then the real part is used for *theta* and the imaginary part is used for *rho*.

The optional argument *fmt* specifies the line format in the same way as *plot*.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by *gca*.

The optional return value *h* is a graphics handle to the created plot.

Implementation Note: The polar axis is drawn using line and text objects encapsulated in an hggroup. The hggroup properties are linked to the original axes object such that altering an appearance property, for example *fontname*, will update the polar axis. Two new properties are added to the original axes—*rtick*, *ttick*—which replace *xtick*, *ytick*. The first is a list of tick locations in the radial (*rho*) direction; The second is a list of tick locations in the angular (*theta*) direction specified in degrees, i.e., in the range 0–359.

See also: [rose], page 309, [compass], page 320, [plot], page 296.

```
pie (x)
pie (..., explode)
pie (..., labels)
```

```
pie (hax, ...);
h = pie (...);
Plot a 2-D pie chart.
```

When called with a single vector argument, produce a pie chart of the elements in *x*. The size of the *i*th slice is the percentage that the element *xi* represents of the total sum of *x*: $pct = x(i) / \text{sum}(x)$.

The optional input *explode* is a vector of the same length as *x* that, if nonzero, "explodes" the slice from the pie chart.

The optional input *labels* is a cell array of strings of the same length as *x* specifying the label for each slice.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by *gca*.

The optional return value *h* is a list of handles to the patch and text objects generating the plot.

Note: If $\text{sum}(x) \leq 1$ then the elements of *x* are interpreted as percentages directly and are not normalized by $\text{sum}(x)$. Furthermore, if the sum is less than 1 then there will be a missing slice in the pie plot to represent the missing, unspecified percentage.

See also: [pie3], page 318, [bar], page 301, [hist], page 303, [rose], page 309.

```
pie3 (x)
pie3 (... , explode)
pie3 (... , labels)
pie3 (hax, ...);
h = pie3 (...);
Plot a 3-D pie chart.
```

Called with a single vector argument, produces a 3-D pie chart of the elements in *x*. The size of the *i*th slice is the percentage that the element *xi* represents of the total sum of *x*: $pct = x(i) / \text{sum}(x)$.

The optional input *explode* is a vector of the same length as *x* that, if nonzero, "explodes" the slice from the pie chart.

The optional input *labels* is a cell array of strings of the same length as *x* specifying the label for each slice.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by *gca*.

The optional return value *h* is a list of graphics handles to the patch, surface, and text objects generating the plot.

Note: If $\text{sum}(x) \leq 1$ then the elements of *x* are interpreted as percentages directly and are not normalized by $\text{sum}(x)$. Furthermore, if the sum is less than 1 then there will be a missing slice in the pie plot to represent the missing, unspecified percentage.

See also: [pie], page 317, [bar], page 301, [hist], page 303, [rose], page 309.

```
quiver (u, v)
quiver (x, y, u, v)
quiver (... , s)
```

```
quiver (..., style)
quiver (..., "filled")
quiver (hax, ...)
h = quiver (...)
```

Plot a 2-D vector field with arrows.

Plot the (u, v) components of a vector field at the grid points defined by (x, y). If the grid is uniform then x and y can be specified as vectors and `meshgrid` is used to create the 2-D grid.

If x and y are not given they are assumed to be $(1:m, 1:n)$ where $[m, n] = \text{size}(u)$.

The optional input s is a scalar defining a scaling factor to use for the arrows of the field relative to the mesh spacing. A value of 1.0 will result in the longest vector exactly filling one grid square. A value of 0 disables all scaling. The default value is 0.9.

The style to use for the plot can be defined with a line style *style* of the same format as the `plot` command. If a marker is specified then the markers are drawn at the origin of the vectors (which are the grid points defined by x and y). When a marker is specified, the arrowhead is not drawn. If the argument "filled" is given then the markers are filled.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to a quiver object. A quiver object regroups the components of the quiver plot (body, arrow, and marker), and allows them to be changed together.

Example:

```
[x, y] = meshgrid (1:2:20);
h = quiver (x, y, sin (2*pi*x/10), sin (2*pi*y/10));
set (h, "maxheadsize", 0.33);
```

See also: [quiver3], page 319, [compass], page 320, [feather], page 320, [plot], page 296.

```
quiver3 (u, v, w)
quiver3 (x, y, z, u, v, w)
quiver3 (..., s)
quiver3 (..., style)
quiver3 (..., "filled")
quiver3 (hax, ...)
h = quiver3 (...)
```

Plot a 3-D vector field with arrows.

Plot the (u, v, w) components of a vector field at the grid points defined by (x, y, z). If the grid is uniform then x, y , and z can be specified as vectors and `meshgrid` is used to create the 3-D grid.

If x, y , and z are not given they are assumed to be $(1:m, 1:n, 1:p)$ where $[m, n] = \text{size}(u)$ and $p = \max(\text{size}(w))$.

The optional input *s* is a scalar defining a scaling factor to use for the arrows of the field relative to the mesh spacing. A value of 1.0 will result in the longest vector exactly filling one grid cube. A value of 0 disables all scaling. The default value is 0.9.

The style to use for the plot can be defined with a line style *style* of the same format as the `plot` command. If a marker is specified then the markers are drawn at the origin of the vectors (which are the grid points defined by *x*, *y*, *z*). When a marker is specified, the arrowhead is not drawn. If the argument "filled" is given then the markers are filled.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to a quiver object. A quiver object regroups the components of the quiver plot (body, arrow, and marker), and allows them to be changed together.

```
[x, y, z] = peaks (25);
surf (x, y, z);
hold on;
[u, v, w] = surfnorm (x, y, z / 10);
h = quiver3 (x, y, z, u, v, w);
set (h, "maxheadsize", 0.33);
```

See also: [\[quiver\]](#), page 318, [\[compass\]](#), page 320, [\[feather\]](#), page 320, [\[plot\]](#), page 296.

```
compass (u, v)
compass (z)
compass (..., style)
compass (hax, ...)
h = compass (...)
```

Plot the (*u*, *v*) components of a vector field emanating from the origin of a polar plot.

The arrow representing each vector has one end at the origin and the tip at [*u*(*i*), *v*(*i*)]. If a single complex argument *z* is given, then *u* = `real` (*z*) and *v* = `imag` (*z*).

The style to use for the plot can be defined with a line style *style* of the same format as the `plot` command.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a vector of graphics handles to the line objects representing the drawn vectors.

```
a = toeplitz ([1;randn(9,1)], [1,randn(1,9)]);
compass (eig (a));
```

See also: [\[polar\]](#), page 317, [\[feather\]](#), page 320, [\[quiver\]](#), page 318, [\[rose\]](#), page 309, [\[plot\]](#), page 296.

```
feather (u, v)
feather (z)
feather (..., style)
```

```
feather (hax, ...)
h = feather (...)
```

Plot the (u, v) components of a vector field emanating from equidistant points on the x-axis.

If a single complex argument z is given, then $u = \text{real} (z)$ and $v = \text{imag} (z)$.

The style to use for the plot can be defined with a line style *style* of the same format as the `plot` command.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a vector of graphics handles to the line objects representing the drawn vectors.

```
phi = [0 : 15 : 360] * pi/180;
feather (sin (phi), cos (phi));
```

See also: [\[plot\]](#), page 296, [\[quiver\]](#), page 318, [\[compass\]](#), page 320.

```
pcolor (x, y, c)
pcolor (c)
pcolor (hax, ...)
h = pcolor (...)
```

Produce a 2-D density plot.

A `pcolor` plot draws rectangles with colors from the matrix *c* over the two-dimensional region represented by the matrices *x* and *y*. *x* and *y* are the coordinates of the mesh's vertices and are typically the output of `meshgrid`. If *x* and *y* are vectors, then a typical vertex is $(x(j), y(i), c(i,j))$. Thus, columns of *c* correspond to different *x* values and rows of *c* correspond to different *y* values.

The values in *c* are scaled to span the range of the current colormap. Limits may be placed on the color axis by the command `caxis`, or by setting the `clim` property of the parent axis.

The face color of each cell of the mesh is determined by interpolating the values of *c* for each of the cell's vertices; Contrast this with `imagesc` which renders one cell for each element of *c*.

`shading` modifies an attribute determining the manner by which the face color of each cell is interpolated from the values of *c*, and the visibility of the cells' edges. By default the attribute is "faceted", which renders a single color for each cell's face with the edge visible.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the created surface object.

See also: [\[caxis\]](#), page 325, [\[shading\]](#), page 356, [\[meshgrid\]](#), page 347, [\[contour\]](#), page 310, [\[imagesc\]](#), page 812.

```
area (y)
area (x, y)
area (... , lvl)
```

```
area (... , prop , val , ... )
area (hax , ... )
h = area (... )
```

Area plot of the columns of y.

This plot shows the contributions of each column value to the row sum. It is functionally similar to `plot (x, cumsum (y, 2))`, except that the area under the curve is shaded.

If the x argument is omitted it defaults to `1:rows (y)`. A value *lvl* can be defined that determines where the base level of the shading under the curve should be defined. The default level is 0.

Additional property/value pairs are passed directly to the underlying patch object.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the hggroup object comprising the area patch objects. The "BaseValue" property of the hggroup can be used to adjust the level where shading begins.

Example: Verify identity $\sin^2 + \cos^2 = 1$

```
t = linspace (0, 2*pi, 100)';
y = [sin(t).^2, cos(t).^2];
area (t, y);
legend ("sin^2", "cos^2", "location", "NorthEastOutside");
```

See also: [\[plot\]](#), page 296, [\[patch\]](#), page 397.

```
fill (x, y, c)
fill (x1, y1, c1, x2, y2, c2)
fill (... , prop , val)
fill (hax , ... )
h = fill (... )
```

Create one or more filled 2-D polygons.

The inputs x and y are the coordinates of the polygon vertices. If the inputs are matrices then the rows represent different vertices and each column produces a different polygon. `fill` will close any open polygons before plotting.

The input c determines the color of the polygon. The simplest form is a single color specification such as a `plot` format or an RGB-triple. In this case the polygon(s) will have one unique color. If c is a vector or matrix then the color data is first scaled using `caxis` and then indexed into the current colormap. A row vector will color each polygon (a column from matrices x and y) with a single computed color. A matrix c of the same size as x and y will compute the color of each vertex and then interpolate the face color between the vertices.

Multiple property/value pairs for the underlying patch object may be specified, but they must appear in pairs.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a vector of graphics handles to the created patch objects.

Example: red square

```
vertices = [0 0
            1 0
            1 1
            0 1];
fill (vertices(:,1), vertices(:,2), "r");
axis([-0.5 1.5, -0.5 1.5])
axis equal
```

See also: [\[patch\]](#), page 397, [\[caxis\]](#), page 325, [\[colormap\]](#), page 816.

```
comet (y)
comet (x, y)
comet (x, y, p)
comet (hax, ...)
```

Produce a simple comet style animation along the trajectory provided by the input coordinate vectors (x, y).

If x is not specified it defaults to the indices of y .

The speed of the comet may be controlled by p , which represents the time each point is displayed before moving to the next one. The default for p is 0.1 seconds.

If the first argument hax is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

See also: [\[comet3\]](#), page 323.

```
comet3 (z)
comet3 (x, y, z)
comet3 (x, y, z, p)
comet3 (hax, ...)
```

Produce a simple comet style animation along the trajectory provided by the input coordinate vectors (x, y, z).

If only z is specified then x, y default to the indices of z .

The speed of the comet may be controlled by p , which represents the time each point is displayed before moving to the next one. The default for p is 0.1 seconds.

If the first argument hax is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

See also: [\[comet\]](#), page 323.

15.2.1.1 Axis Configuration

The axis function may be used to change the axis limits of an existing plot and various other axis properties, such as the aspect ratio and the appearance of tic marks. By default, high level plotting functions such as `plot` reset axes properties. Any customization of properties, for example by calling `axis`, `xlim`, etc., should happen after the plot is done or, alternatively, after calling the [\[hold function\]](#), page 375.

```
axis ()
axis ([x_lo x_hi])
```

```
axis ([x_lo x_hi y_lo y_hi])
axis ([x_lo x_hi y_lo y_hi z_lo z_hi])
axis ([x_lo x_hi y_lo y_hi z_lo z_hi c_lo c_hi])
axis (option)
axis (option1, option2, ...)
axis (hax, ...)
limits = axis ()
```

Set axis limits and appearance.

The argument *limits* should be a 2-, 4-, 6-, or 8-element vector. The first and second elements specify the lower and upper limits for the x-axis. The third and fourth specify the limits for the y-axis, the fifth and sixth specify the limits for the z-axis, and the seventh and eighth specify the limits for the color axis. The special values -Inf and Inf may be used to indicate that the limit should be automatically computed based on the data in the axes.

Without any arguments, **axis** turns autoscaling on.

With one output argument, **limits = axis** returns the current axis limits.

The vector argument specifying limits is optional, and additional string arguments may be used to specify various axis properties.

The following options control the aspect ratio of the axes.

- "square"** Force a square axis aspect ratio.
- "equal"** Force x-axis unit distance to equal y-axis (and z-axis) unit distance.
- "normal"** Restore default aspect ratio.

The following options control the way axis limits are interpreted.

- "auto"**
- "auto[xyz]"** Set the specified axes to have nice limits around the data or all if no axes are specified.
- "manual"** Fix the current axes limits.
- "tight"** Fix axes to the limits of the data.
- "image"** Equivalent to **"tight"** and **"equal"**.
- "vis3d"** Set aspect ratio modes to **"manual"** for rotation without stretching.

The following options affect the appearance of tick marks.

- "tic[xyz]"** Turn tick marks on for all axes, or turn them on for the specified axes and off for the remainder.
- "label[xyz]"** Turn tick labels on for all axes, or turn them on for the specified axes and off for the remainder.
- "nolabel"** Turn tick labels off for all axes.

Note: If there are no tick marks for an axes then there can be no labels.

The following options affect the direction of increasing values on the axes.

"xy" Default y-axis, larger values are near the top.

"ij" Reverse y-axis, smaller values are near the top.

The following options affects the visibility of the axes.

"on" Make the axes visible.

"off" Hide the axes.

If the first argument *hax* is an axes handle, then operate on this axes rather than the current axes returned by `gca`.

Example 1: set X/Y limits and force a square aspect ratio

```
axis ([1, 2, 3, 4], "square");
```

Example 2: enable tick marks on all axes, enable tick mark labels only on the y-axis

```
axis ("tic", "labely");
```

See also: [\[xlim\]](#), page 325, [\[ylim\]](#), page 325, [\[zlim\]](#), page 325, [\[caxis\]](#), page 325, [\[daspect\]](#), page 358, [\[pbaspect\]](#), page 358, [\[box\]](#), page 367, [\[grid\]](#), page 367.

Similarly the axis limits of the colormap can be changed with the `caxis` function.

```
caxis ([cmin cmax])
caxis ("auto")
caxis ("manual")
caxis (hax, ...)
limits = caxis ()
```

Query or set color axis limits for plots.

The *limits* argument should be a 2-element vector specifying the lower and upper limits to assign to the first and last value in the colormap. Data values outside this range are clamped to the first and last colormap entries.

If the "auto" option is given then automatic colormap limits are applied. The automatic algorithm sets *cmin* to the minimum data value and *cmax* to the maximum data value. If "manual" is specified then the "climmode" property is set to "manual" and the numeric values in the "clim" property are used for limits.

If the first argument *hax* is an axes handle, then operate on this axes rather than the current axes returned by `gca`.

Called without arguments the current color axis limits are returned.

Programming Note: The color axis affects the display of image, patch, and surface graphics objects, but **only** if the "cdata" property has indexed data and the "cdatamapping" property is set to "scaled". Graphic objects with true color `cdata`, or "direct" `cdatamapping` are not affected.

See also: [\[colormap\]](#), page 816, [\[axis\]](#), page 323.

The `xlim`, `ylim`, and `zlim` functions may be used to get or set individual axis limits. Each has the same form.

```
xlimits = xlim ()
xmode = xlim ("mode")
xlim ([x_lo x_hi])
xlim ("auto")
xlim ("manual")
xlim (hax, ...)
```

Query or set the limits of the x-axis for the current plot.

Called without arguments `xlim` returns the x-axis limits of the current plot.

With the input query "mode", return the current x-limit calculation mode which is either "auto" or "manual".

If passed a 2-element vector [x_lo x_hi], the limits of the x-axis are set to these values and the mode is set to "manual". The special values -Inf and Inf can be used to indicate that either the lower axis limit or upper axis limit should be automatically calculated.

The current plotting mode can be changed by using either "auto" or "manual" as the argument.

If the first argument `hax` is an axes handle, then operate on this axes rather than the current axes returned by `gca`.

Programming Note: The `xlim` function operates by modifying the "xlim" and "xlimmode" properties of an axes object. These properties can be directly inspected and altered with `get/set`.

See also: `[ylim]`, page 325, `[zlim]`, page 325, `[axis]`, page 323, `[set]`, page 402, `[get]`, page 402, `[gca]`, page 401.

The `xticks`, `yticks`, `zticks`, `rticks`, and `thetaticks` functions may be used to get or set the tick mark locations and modes on the respective axis. Each has the same form, although mode options are not currently available for `rticks`, and `thetaticks`.

```
tickval = xticks
mode = xticks ("mode")
xticks (tickval)
xticks ("auto")
xticks ("manual")
... = xticks (hax, ...)
```

Query or set the tick values on the x-axis of the current axis.

When called without an argument, return the current tick locations as specified in the "xtick" axes property. These locations can be changed by calling `xticks` with a vector of tick values. Note: ascending order is not required.

When called with argument "mode", `xticks` returns the current value of the axes property "xtickmode". This property can be changed by calling `xticks` with either "auto" (algorithm determines tick positions) or "manual" (tick values remain fixed regardless of axes resizing or rotation). Note: Specifying xtick values will also set the property "xtickmode" to "manual".

If the first argument `hax` is an axes handle, then operate on this axis rather than the current axes returned by `gca`.

Requesting a return value when calling `xticks` to set a property value will result in an error.

See also: [xticklabels], page 327, [yticks], page 326, [zticks], page 326, [rticks], page 326, [thetaticks], page 326, [get], page 402, [set], page 402.

The `xticklabels`, `yticklabels`, and `zticklabels` functions may be used to get or set the label assigned to each tick location and the labeling mode on the respective axis. Each has the same form.

```
tickval = xticklabels
mode = xticklabels ("mode")
xticklabels (tickval)
xticklabels ("auto")
xticklabels ("manual")
... = xticklabels (hax, ...)
```

Query or set the tick labels on the x-axis of the current axis.

When called without an argument, return a cell array of strings of the current tick labels as specified in the "xticklabel" axes property. These labels can be changed by calling `xticklabels` with a cell array of strings. Note: a vector of numbers will be mapped to a cell array of strings. If fewer labels are specified than the current number of ticks, blank labels will be appended to the array.

When called with argument "mode", `xticklabels` returns the current value of the axes property "xticklabelmode". This property can be changed by calling `xticklabels` with either "auto" (algorithm determines tick labels) or "manual" (tick labels remain fixed). Note: Specifying `xticklabel` values will also set the "xticklabelmode" and "xticks" properties to "manual".

If the first argument `hax` is an axes handle, then operate on this axis rather than the current axes returned by `gca`.

Requesting a return value when calling `xticklabels` to set a property value will result in an error.

See also: [xticks], page 326, [yticklabels], page 327, [zticklabels], page 327, [get], page 402, [set], page 402.

15.2.1.2 Two-dimensional Function Plotting

Octave can plot a function from a function handle, inline function, or string defining the function without the user needing to explicitly create the data to be plotted. The function `fplot` also generates two-dimensional plots with linear axes using a function name and limits for the range of the x-coordinate instead of the x and y data. For example,

```
fplot (@sin, [-10, 10], 201);
```

produces a plot that is equivalent to the one above, but also includes a legend displaying the name of the plotted function.

```
fplot (fn)
fplot (fn, limits)
fplot (..., tol)
fplot (..., n)
```

```
fplot (..., fmt)
fplot (..., property, value, ...)
fplot (hax, ...)
[x, y] = fplot (...)
```

Plot a function *fn* within the range defined by *limits*.

fn is a function handle, inline function, or string containing the name of the function to evaluate.

The limits of the plot are of the form [*xlo*, *xhi*] or [*xlo*, *xhi*, *ylo*, *yhi*]. If no limits are specified the default is [-5, 5].

The next three arguments are all optional and any number of them may be given in any order.

tol is the relative tolerance to use for the plot and defaults to 2e-3 (.2%).

n is the minimum number of points to use. When *n* is specified, the maximum stepsize will be (*xhi* - *xlo*) / *n*. More than *n* points may still be used in order to meet the relative tolerance requirement.

The *fmt* argument specifies the linestyle to be used by the plot command.

Multiple property-value pairs may also be specified, but they must appear in pairs. These arguments are applied to the line objects drawn by `plot`.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

With no output arguments, the results are immediately plotted. With two output arguments, the 2-D plot data is returned. The data can subsequently be plotted manually with `plot (x, y)`.

Example:

```
fplot (@cos, [0, 2*pi])
fplot ("[cos(x), sin(x)]", [0, 2*pi])
```

Programming Notes:

`fplot` works best with continuous functions. Functions with discontinuities are unlikely to plot well. This restriction may be removed in the future.

`fplot` performance is better when the function accepts and returns a vector argument. Consider this when writing user-defined functions and use element-by-element operators such as `.*`, `./`, etc. See the function `vectorize` for potentially converting inline or anonymous functions to vectorized versions.

See also: [\[ezplot\]](#), page 328, [\[plot\]](#), page 296, [\[vectorize\]](#), page 584.

Other functions that can create two-dimensional plots directly from a function include `ezplot`, `ezcontour`, `ezcontourf` and `ezpolar`.

```
ezplot (f)
ezplot (f2v)
ezplot (fx, )
ezplot (..., dom)
ezplot (..., n)
ezplot (hax, ...)
```

`h = ezplot (...)`

Plot the 2-D curve defined by the function f .

The function f may be a string, inline function, or function handle and can have either one or two variables. If f has one variable, then the function is plotted over the domain $-2\pi < x < 2\pi$ with 500 points.

If $f2v$ is a function of two variables then the implicit function $f(x, y) = 0$ is calculated over the meshed domain $-2\pi \leq x \mid y \leq 2\pi$ with 60 points in each dimension.

For example:

```
ezplot (@(x, y) x.^2 - y.^2 - 1)
```

If two functions are passed as inputs then the parametric function

```
x = fx (t)
y = fy (t)
```

is plotted over the domain $-2\pi \leq t \leq 2\pi$ with 500 points.

If dom is a two element vector, it represents the minimum and maximum values of both x and y , or t for a parametric plot. If dom is a four element vector, then the minimum and maximum values are $[xmin xmax ymin ymax]$.

n is a scalar defining the number of points to use in plotting the function.

If the first argument hax is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value h is a vector of graphics handles to the created line objects.

See also: [\[plot\]](#), page 296, [\[ezplot3\]](#), page 359, [\[ezpolar\]](#), page 330, [\[ezcontour\]](#), page 329, [\[ezcontourf\]](#), page 330, [\[ezmesh\]](#), page 359, [\[ezmeshc\]](#), page 360, [\[ezsurf\]](#), page 361, [\[ezsurfc\]](#), page 361.

`ezcontour (f)`

`ezcontour (..., dom)`

`ezcontour (..., n)`

`ezcontour (hax, ...)`

`h = ezcontour (...)`

Plot the contour lines of a function.

f is a string, inline function, or function handle with two arguments defining the function. By default the plot is over the meshed domain $-2\pi \leq x \mid y \leq 2\pi$ with 60 points in each dimension.

If dom is a two element vector, it represents the minimum and maximum values of both x and y . If dom is a four element vector, then the minimum and maximum values are $[xmin xmax ymin ymax]$.

n is a scalar defining the number of points to use in each dimension.

If the first argument hax is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value h is a graphics handle to the created plot.

Example:

```
f = @(x,y) sqrt (abs (x .* y)) ./ (1 + x.^2 + y.^2);
ezcontour (f, [-3, 3]);
```

See also: [contour], page 310, [ezcontourf], page 330, [ezplot], page 328, [ezmeshc], page 360, [ezsurf], page 361.

```
ezcontourf (f)
ezcontourf (... , dom)
ezcontourf (... , n)
ezcontourf (hax, ...)
h = ezcontourf (...)
```

Plot the filled contour lines of a function.

f is a string, inline function, or function handle with two arguments defining the function. By default the plot is over the meshed domain $-2\pi \leq x \mid y \leq 2\pi$ with 60 points in each dimension.

If *dom* is a two element vector, it represents the minimum and maximum values of both *x* and *y*. If *dom* is a four element vector, then the minimum and maximum values are [*xmin xmax ymin ymax*].

n is a scalar defining the number of points to use in each dimension.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by *gca*.

The optional return value *h* is a graphics handle to the created plot.

Example:

```
f = @(x,y) sqrt (abs (x .* y)) ./ (1 + x.^2 + y.^2);
ezcontourf (f, [-3, 3]);
```

See also: [contourf], page 311, [ezcontour], page 329, [ezplot], page 328, [ezmeshc], page 360, [ezsurf], page 361.

```
ezpolar (f)
ezpolar (... , dom)
ezpolar (... , n)
ezpolar (hax, ...)
h = ezpolar (...)
```

Plot a 2-D function in polar coordinates.

The function *f* is a string, inline function, or function handle with a single argument. The expected form of the function is *rho* = *f(theta)*. By default the plot is over the domain $0 \leq \theta \leq 2\pi$ with 500 points.

If *dom* is a two element vector, it represents the minimum and maximum values of *theta*.

n is a scalar defining the number of points to use in plotting the function.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by *gca*.

The optional return value *h* is a graphics handle to the created plot.

Example:

```
ezpolar (@(t) sin (5/4 * t), [0, 8*pi]);
```

See also: [polar], page 317, [ezplot], page 328.

15.2.1.3 Two-dimensional Geometric Shapes

```
rectangle ()
rectangle (..., "Position", pos)
rectangle (..., "Curvature", curv)
rectangle (..., "EdgeColor", ec)
rectangle (..., "FaceColor", fc)
rectangle (hax, ...)
h = rectangle (...)
```

Draw a rectangular patch defined by *pos* and *curv*.

The variable *pos*(1:2) defines the lower left-hand corner of the patch and *pos*(3:4) defines its width and height. By default, the value of *pos* is [0, 0, 1, 1].

The variable *curv* defines the curvature of the sides of the rectangle and may be a scalar or two-element vector with values between 0 and 1. A value of 0 represents no curvature of the side, whereas a value of 1 means that the side is entirely curved into the arc of a circle. If *curv* is a two-element vector, then the first element is the curvature along the x-axis of the patch and the second along y-axis.

If *curv* is a scalar, it represents the curvature of the shorter of the two sides of the rectangle and the curvature of the other side is defined by

```
min (pos(1:2)) / max (pos(1:2)) * curv
```

Additional property/value pairs are passed to the underlying patch command.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by *gca*.

The optional return value *h* is a graphics handle to the created rectangle object.

See also: [\[patch\]](#), page 397, [\[line\]](#), page 396, [\[cylinder\]](#), page 362, [\[ellipsoid\]](#), page 363, [\[sphere\]](#), page 362.

15.2.2 Three-Dimensional Plots

The function *mesh* produces mesh surface plots. For example,

```
tx = ty = linspace (-8, 8, 41)';
[xx, yy] = meshgrid (tx, ty);
r = sqrt (xx .^ 2 + yy .^ 2) + eps;
tz = sin (r) ./ r;
mesh (tx, ty, tz);
xlabel ("tx");
ylabel ("ty");
zlabel ("tz");
title ("3-D Sombrero plot");
```

produces the familiar “sombrero” plot shown in [Figure 15.5](#). Note the use of the function *meshgrid* to create matrices of X and Y coordinates to use for plotting the Z data. The *ndgrid* function is similar to *meshgrid*, but works for N-dimensional matrices.




Figure 15.5: Mesh plot.

The `meshc` function is similar to `mesh`, but also produces a plot of contours for the surface.

The `plot3` function displays arbitrary three-dimensional data, without requiring it to form a surface. For example,

```
t = 0:0.1:10*pi;
r = linspace (0, 1, numel (t));
z = linspace (0, 1, numel (t));
plot3 (r.*sin (t), r.*cos (t), z);
xlabel ("r.*sin (t)");
ylabel ("r.*cos (t)");
zlabel ("z");
title ("plot3 display of 3-D helix");
```

displays the spiral in three dimensions shown in Figure 15.6.




Figure 15.6: Three-dimensional spiral.

Finally, the `view` function changes the viewpoint for three-dimensional plots.

```
mesh (x, y, z)
mesh (z)
mesh (... , c)
mesh (... , prop, val, ...)
mesh (hax, ...)
h = mesh (...)
```

Plot a 3-D wireframe mesh.

The wireframe mesh is plotted using rectangles. The vertices of the rectangles [x, y] are typically the output of `meshgrid`. over a 2-D rectangular region in the x - y plane. z determines the height above the plane of each vertex. If only a single z matrix is given, then it is plotted over the meshgrid $x = 1:\text{columns}(z)$, $y = 1:\text{rows}(z)$. Thus, columns of z correspond to different x values and rows of z correspond to different y values.

The color of the mesh is computed by linearly scaling the z values to fit the range of the current colormap. Use `caxis` and/or change the colormap to control the appearance. Optionally, the color of the mesh can be specified independently of z by supplying a color matrix, c .

Any property/value pairs are passed directly to the underlying surface object.

If the first argument hax is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value h is a graphics handle to the created surface object.

See also: [\[ezmesh\]](#), page 359, [\[meshc\]](#), page 334, [\[meshz\]](#), page 334, [\[trimesh\]](#), page 777, [\[contour\]](#), page 310, [\[surf\]](#), page 335, [\[surface\]](#), page 398, [\[meshgrid\]](#), page 347, [\[hidden\]](#), page 335, [\[shading\]](#), page 356, [\[colormap\]](#), page 816, [\[caxis\]](#), page 325.

```
meshc (x, y, z)
meshc (z)
meshc (..., c)
meshc (..., prop, val, ...)
meshc (hax, ...)
h = meshc (...)
```

Plot a 3-D wireframe mesh with underlying contour lines.

The wireframe mesh is plotted using rectangles. The vertices of the rectangles [*x*, *y*] are typically the output of **meshgrid**. over a 2-D rectangular region in the *x*-*y* plane. *z* determines the height above the plane of each vertex. If only a single *z* matrix is given, then it is plotted over the meshgrid *x* = 1:**columns** (*z*), *y* = 1:**rows** (*z*). Thus, columns of *z* correspond to different *x* values and rows of *z* correspond to different *y* values.

The color of the mesh is computed by linearly scaling the *z* values to fit the range of the current colormap. Use **caxis** and/or change the colormap to control the appearance.

Optionally the color of the mesh can be specified independently of *z* by supplying a color matrix, *c*.

Any property/value pairs are passed directly to the underlying surface object.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by **gca**.

The optional return value *h* is a 2-element vector with a graphics handle to the created surface object and to the created contour plot.

See also: [ezmeshc], page 360, [mesh], page 333, [meshz], page 334, [contour], page 310, [surf], page 336, [surface], page 398, [meshgrid], page 347, [hidden], page 335, [shading], page 356, [colormap], page 816, [caxis], page 325.

```
meshz (x, y, z)
meshz (z)
meshz (..., c)
meshz (..., prop, val, ...)
meshz (hax, ...)
h = meshz (...)
```

Plot a 3-D wireframe mesh with a surrounding curtain.

The wireframe mesh is plotted using rectangles. The vertices of the rectangles [*x*, *y*] are typically the output of **meshgrid**. over a 2-D rectangular region in the *x*-*y* plane. *z* determines the height above the plane of each vertex. If only a single *z* matrix is given, then it is plotted over the meshgrid *x* = 1:**columns** (*z*), *y* = 1:**rows** (*z*). Thus, columns of *z* correspond to different *x* values and rows of *z* correspond to different *y* values.

The color of the mesh is computed by linearly scaling the *z* values to fit the range of the current colormap. Use **caxis** and/or change the colormap to control the appearance.

Optionally the color of the mesh can be specified independently of *z* by supplying a color matrix, *c*.

Any property/value pairs are passed directly to the underlying surface object.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the created surface object.

See also: `[mesh]`, page 333, `[meshc]`, page 334, `[contour]`, page 310, `[surf]`, page 335, `[surface]`, page 398, `[waterfall]`, page 357, `[meshgrid]`, page 347, `[hidden]`, page 335, `[shading]`, page 356, `[colormap]`, page 816, `[caxis]`, page 325.

`hidden`

`hidden on`

`hidden off`

`mode = hidden (...)`

Control mesh hidden line removal.

When called with no argument the hidden line removal state is toggled.

When called with one of the modes "on" or "off" the state is set accordingly.

The optional output argument *mode* is the current state.

Hidden Line Removal determines what graphic objects behind a mesh plot are visible. The default is for the mesh to be opaque and lines behind the mesh are not visible. If hidden line removal is turned off then objects behind the mesh can be seen through the faces (openings) of the mesh, although the mesh grid lines are still opaque.

See also: `[mesh]`, page 333, `[meshc]`, page 334, `[meshz]`, page 334, `[ezmesh]`, page 359, `[ezmeshc]`, page 360, `[trimesh]`, page 777, `[waterfall]`, page 357.

`surf (x, y, z)`

`surf (z)`

`surf (... , c)`

`surf (... , prop, val, ...)`

`surf (hax, ...)`

`h = surf (...)`

Plot a 3-D surface mesh.

The surface mesh is plotted using shaded rectangles. The vertices of the rectangles $[x, y]$ are typically the output of `meshgrid`. over a 2-D rectangular region in the x-y plane. *z* determines the height above the plane of each vertex. If only a single *z* matrix is given, then it is plotted over the meshgrid `x = 1:columns (z)`, `y = 1:rows (z)`. Thus, columns of *z* correspond to different *x* values and rows of *z* correspond to different *y* values.

The color of the surface is computed by linearly scaling the *z* values to fit the range of the current colormap. Use `caxis` and/or change the colormap to control the appearance.

Optionally, the color of the surface can be specified independently of *z* by supplying a color matrix, *c*.

Any property/value pairs are passed directly to the underlying surface object.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the created surface object.

Note: The exact appearance of the surface can be controlled with the **shading** command or by using **set** to control surface object properties.

See also: [ezsurf], page 361, [surf], page 335, [surfl], page 336, [surfnorm], page 337, [trisurf], page 778, [contour], page 310, [mesh], page 333, [surface], page 398, [meshgrid], page 347, [hidden], page 335, [shading], page 356, [colormap], page 816, [caxis], page 325.

```
surf (x, y, z)
surf (z)
surf (..., c)
surf (..., prop, val, ...)
surf (hax, ...)
h = surf (...)
```

Plot a 3-D surface mesh with underlying contour lines.

The surface mesh is plotted using shaded rectangles. The vertices of the rectangles [x, y] are typically the output of **meshgrid**, over a 2-D rectangular region in the x-y plane. z determines the height above the plane of each vertex. If only a single z matrix is given, then it is plotted over the meshgrid `x = 1:columns(z), y = 1:rows(z)`. Thus, columns of z correspond to different x values and rows of z correspond to different y values.

The color of the surface is computed by linearly scaling the z values to fit the range of the current colormap. Use **caxis** and/or change the colormap to control the appearance.

Optionally, the color of the surface can be specified independently of z by supplying a color matrix, c.

Any property/value pairs are passed directly to the underlying surface object.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by **gca**.

The optional return value *h* is a graphics handle to the created surface object.

Note: The exact appearance of the surface can be controlled with the **shading** command or by using **set** to control surface object properties.

See also: [ezsurf], page 361, [surf], page 335, [surfl], page 336, [surfnorm], page 337, [trisurf], page 778, [contour], page 310, [mesh], page 333, [surface], page 398, [meshgrid], page 347, [hidden], page 335, [shading], page 356, [colormap], page 816, [caxis], page 325.

```
surfl (z)
surfl (x, y, z)
surfl (..., ls)
surfl (x, y, z, ls, P)
surfl (..., "cdata")
surfl (..., "light")
surfl (hax, ...)
h = surfl (...)
```

Plot a 3-D surface using shading based on various lighting models.

The surface mesh is plotted using shaded rectangles. The vertices of the rectangles $[x, y]$ are typically the output of `meshgrid`. over a 2-D rectangular region in the x-y plane. z determines the height above the plane of each vertex. If only a single z matrix is given, then it is plotted over the meshgrid $x = 1:\text{columns}(z)$, $y = 1:\text{rows}(z)$. Thus, columns of z correspond to different x values and rows of z correspond to different y values.

The default lighting mode "cdata", changes the cdata property of the surface object to give the impression of a lighted surface.

The alternate mode "light" creates a light object to illuminate the surface.

The light source location may be specified using `lsrc` which can be a 2-element vector [azimuth, elevation] in degrees, or a 3-element vector [lx, ly, lz]. The default value is rotated 45 degrees counterclockwise to the current view.

The material properties of the surface can be specified using a 4-element vector $P = [AM \ D \ SP \ EXP]$ which defaults to $p = [0.55 \ 0.6 \ 0.4 \ 10]$.

- "AM" strength of ambient light
- "D" strength of diffuse reflection
- "SP" strength of specular reflection
- "EXP" specular exponent

If the first argument `hax` is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value `h` is a graphics handle to the created surface object.

Example:

```
colormap (bone (64));
surf (peaks);
shading interp;
```

See also: [\[diffuse\]](#), page 345, [\[specular\]](#), page 345, [\[surf\]](#), page 335, [\[shading\]](#), page 356, [\[colormap\]](#), page 816, [\[caxis\]](#), page 325.

```
surfnorm (x, y, z)
surfnorm (z)
surfnorm (..., prop, val, ...)
surfnorm (hax, ...)
[nx, ny, nz] = surfnorm (...)
```

Find the vectors normal to a meshgridded surface.

If x and y are vectors, then a typical vertex is $(x(j), y(i), z(i,j))$. Thus, columns of z correspond to different x values and rows of z correspond to different y values. If only a single input z is given then x is taken to be $1:\text{columns}(z)$ and y is $1:\text{rows}(z)$.

If no return arguments are requested, a surface plot with the normal vectors to the surface is plotted.

Any property/value input pairs are assigned to the surface object.

If the first argument `hax` is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

If output arguments are requested then the components of the normal vectors are returned in *nx*, *ny*, and *nz* and no plot is made. The normal vectors are unnormalized (magnitude != 1). To normalize, use

```
len = sqrt (nx.^2 + ny.^2 + nz.^2);
nx ./= len; ny ./= len; nz ./= len;
```

An example of the use of **surfnorm** is

```
surfnorm (peaks (25));
```

Algorithm: The normal vectors are calculated by taking the cross product of the diagonals of each of the quadrilateral faces in the meshgrid to find the normal vectors at the center of each face. Next, for each meshgrid point the four nearest normal vectors are averaged to obtain the final normal to the surface at the meshgrid point.

For surface objects, the "VertexNormals" property contains equivalent information, except possibly near the boundary of the surface where different interpolation schemes may yield slightly different values.

See also: [isonormals], page 340, [quiver3], page 319, [surf], page 335, [meshgrid], page 347.

```
fv = isosurface (v, isoval)
fv = isosurface (v)
fv = isosurface (x, y, z, v, isoval)
fv = isosurface (x, y, z, v)
fvc = isosurface (..., col)
fv = isosurface (..., "noshare")
fv = isosurface (..., "verbose")
[f, v] = isosurface (...)
[f, v, c] = isosurface (...)
isosurface (...)
```

Calculate isosurface of 3-D volume data.

An isosurface connects points with the same value and is analogous to a contour plot, but in three dimensions.

The input argument *v* is a three-dimensional array that contains data sampled over a volume.

The input *isoval* is a scalar that specifies the value for the isosurface. If *isoval* is omitted or empty, a "good" value for an isosurface is determined from *v*.

When called with a single output argument **isosurface** returns a structure array *fv* that contains the fields *faces* and *vertices* computed at the points *[x, y, z] = meshgrid (1:1, 1:m, 1:n)* where *[l, m, n] = size (v)*. The output *fv* can be used directly as input to the **patch** function.

If called with additional input arguments *x*, *y*, and *z* that are three-dimensional arrays with the same size as *v* or vectors with lengths corresponding to the dimensions of *v*, then the volume data is taken at the specified points. If *x*, *y*, or *z* are empty, the grid corresponds to the indices (1:n) in the respective direction (see [meshgrid], page 347).

The optional input argument *col*, which is a three-dimensional array of the same size as *v*, specifies coloring of the isosurface. The color data is interpolated, as necessary,

to match *isoval*. The output structure array, in this case, has the additional field *facevertexexdata*.

If given the string input argument "*noshare*", vertices may be returned multiple times for different faces. The default behavior is to eliminate vertices shared by adjacent faces with *unique* which may be time consuming.

The string input argument "*verbose*" is supported for MATLAB compatibility, but has no effect.

Any string arguments must be passed after the other arguments.

If called with two or three output arguments, return the information about the faces *f*, vertices *v*, and color data *c* as separate arrays instead of a single structure array.

If called with no output argument, the isosurface geometry is directly plotted with the *patch* command and a light object is added to the axes if not yet present.

For example,

```
[x, y, z] = meshgrid (1:5, 1:5, 1:5);
v = rand (5, 5, 5);
isosurface (x, y, z, v, .5);
```

will directly draw a random isosurface geometry in a graphics window.

An example of an isosurface geometry with different additional coloring:

```
N = 15;      # Increase number of vertices in each direction
iso = .4;    # Change isovalue to .1 to display a sphere
lin = linspace (0, 2, N);
[x, y, z] = meshgrid (lin, lin, lin);
v = abs ((x-.5).^2 + (y-.5).^2 + (z-.5).^2);
figure ();

subplot (2,2,1); view (-38, 20);
[f, vert] = isosurface (x, y, z, v, iso);
p = patch ("Faces", f, "Vertices", vert, "EdgeColor", "none");
pbaspect ([1 1 1]);
isonormals (x, y, z, v, p)
set (p, "FaceColor", "green", "FaceLighting", "gouraud");
light ("Position", [1 1 5]);

subplot (2,2,2); view (-38, 20);
p = patch ("Faces", f, "Vertices", vert, "EdgeColor", "blue");
pbaspect ([1 1 1]);
isonormals (x, y, z, v, p)
set (p, "FaceColor", "none", "EdgeLighting", "gouraud");
light ("Position", [1 1 5]);

subplot (2,2,3); view (-38, 20);
[f, vert, c] = isosurface (x, y, z, v, iso, y);
p = patch ("Faces", f, "Vertices", vert, "FaceVertexCData", c, ...
           "FaceColor", "interp", "EdgeColor", "none");
pbaspect ([1 1 1]);
isonormals (x, y, z, v, p)
set (p, "FaceLighting", "gouraud");
light ("Position", [1 1 5]);

subplot (2,2,4); view (-38, 20);
p = patch ("Faces", f, "Vertices", vert, "FaceVertexCData", c, ...
```

```

    "FaceColor", "interp", "EdgeColor", "blue");
pbaspect ([1 1 1]);
isonormals (x, y, z, v, p)
set (p, "FaceLighting", "gouraud");
light ("Position", [1 1 5]);

```

See also: [\[isonormals\]](#), page 340, [\[isocolors\]](#), page 341, [\[isocaps\]](#), page 340, [\[smooth3\]](#), page 342, [\[reducevolume\]](#), page 343, [\[reducepatch\]](#), page 343, [\[patch\]](#), page 397.

```

vn = isonormals (val, vert)
vn = isonormals (val, hp)
vn = isonormals (x, y, z, val, vert)
vn = isonormals (x, y, z, val, hp)
vn = isonormals (..., "negate")
isonormals (val, hp)
isonormals (x, y, z, val, hp)
isonormals (..., "negate")

```

Calculate normals to an isosurface.

The vertex normals *vn* are calculated from the gradient of the 3-dimensional array *val* (size: *l*mx*n*) containing the data for an isosurface geometry. The normals point towards smaller values in *val*.

If called with one output argument *vn*, and the second input argument *vert* holds the vertices of an isosurface, then the normals *vn* are calculated at the vertices *vert* on a grid given by `[x, y, z] = meshgrid (1:1, 1:m, 1:n)`. The output argument *vn* has the same size as *vert* and can be used to set the "VertexNormals" property of the corresponding patch.

If called with additional input arguments *x*, *y*, and *z*, which are 3-dimensional arrays with the same size as *val*, then the volume data is taken at these points. Instead of the vertex data *vert*, a patch handle *hp* can be passed to the function.

If the last input argument is the string "negate", compute the reverse vector normals of an isosurface geometry (i.e., pointed towards larger values in *val*).

If no output argument is given, the property "VertexNormals" of the patch associated with the patch handle *hp* is changed directly.

See also: [\[isosurface\]](#), page 338, [\[isocolors\]](#), page 341, [\[smooth3\]](#), page 342.

```

fvc = isocaps (v, isoval)
fvc = isocaps (v)
fvc = isocaps (x, y, z, v, isoval)
fvc = isocaps (x, y, z, v)
fvc = isocaps (..., which_caps)
fvc = isocaps (..., which_plane)
fvc = isocaps (..., "verbose")
[faces, vertices, fvcd] = isocaps (...)
isocaps (...)

```

Create end-caps for isosurfaces of 3-D data.

This function places caps at the open ends of isosurfaces.

The input argument *v* is a three-dimensional array that contains data sampled over a volume.

The input *isoval* is a scalar that specifies the value for the isosurface. If *isoval* is omitted or empty, a "good" value for an isosurface is determined from *v*.

When called with a single output argument, *isocaps* returns a structure array *fvc* with the fields: *faces*, *vertices*, and *facevertexcdata*. The results are computed at the points $[x, y, z] = \text{meshgrid}(1:1, 1:m, 1:n)$ where $[1, m, n] = \text{size}(v)$. The output *fvc* can be used directly as input to the *patch* function.

If called with additional input arguments *x*, *y*, and *z* that are three-dimensional arrays with the same size as *v* or vectors with lengths corresponding to the dimensions of *v*, then the volume data is taken at the specified points. If *x*, *y*, or *z* are empty, the grid corresponds to the indices $(1:n)$ in the respective direction (see [*meshgrid*], page 347).

The optional parameter *which_caps* can have one of the following string values which defines how the data will be enclosed:

- "**a**", "**a**" (default)
for end-caps that enclose the data above *isoval*.
- "**b**", "**b**"
for end-caps that enclose the data below *isoval*.

The optional parameter *which_plane* can have one of the following string values to define which end-cap should be drawn:

- "**all**" (default)
for all of the end-caps.
- "**xmin**" for end-caps at the lower x-plane of the data.
- "**xmax**" for end-caps at the upper x-plane of the data.
- "**ymin**" for end-caps at the lower y-plane of the data.
- "**ymax**" for end-caps at the upper y-plane of the data.
- "**zmin**" for end-caps at the lower z-plane of the data.
- "**zmax**" for end-caps at the upper z-plane of the data.

The string input argument "**verbose**" is supported for MATLAB compatibility, but has no effect.

If called with two or three output arguments, the data for faces *faces*, vertices *vertices*, and the color data *facevertexcdata* are returned in separate arrays instead of a single structure.

If called with no output argument, the end-caps are drawn directly in the current figure with the *patch* command.

See also: [*isosurface*], page 338, [*isonormals*], page 340, [*patch*], page 397.

```
cdat = isocolors (c, v)
cdat = isocolors (x, y, z, c, v)
cdat = isocolors (x, y, z, r, g, b, v)
cdat = isocolors (r, g, b, v)
```

```
cdat = isocolors (..., p)
isocolors (...)
```

Compute isosurface colors.

If called with one output argument, and the first input argument *c* is a three-dimensional array that contains indexed color values, and the second input argument *v* are the vertices of an isosurface geometry, then return a matrix *cdat* with color data information for the geometry at computed points $[x, y, z] = \text{meshgrid} (1:1, 1:m, 1:n)$. The output argument *cdat* can be used to manually set the "FaceVertexCData" property of an isosurface patch object.

If called with additional input arguments *x*, *y* and *z* which are three-dimensional arrays of the same size as *c* then the color data is taken at those specified points.

Instead of indexed color data *c*, *isocolors* can also be called with RGB values *r*, *g*, *b*. If input arguments *x*, *y*, *z* are not given then *meshgrid* computed values are used.

Optionally, a patch handle *p* can be given as the last input argument to all function call variations and the vertex data will be extracted from the isosurface patch object. Finally, if no output argument is given then the colors of the patch given by the patch handle *p* are changed.

See also: [isosurface], page 338, [isonormals], page 340.

```
smoothed_data = smooth3 (data)
smoothed_data = smooth3 (data, method)
smoothed_data = smooth3 (data, method, sz)
smoothed_data = smooth3 (data, method, sz, std_dev)
```

Smooth values of 3-dimensional matrix *data*.

This function can be used, for example, to reduce the impact of noise in *data* before calculating isosurfaces.

data must be a non-singleton 3-dimensional matrix. The smoothed data from this matrix is returned in *smoothed_data* which is of the same size as *data*.

The option input *method* determines which convolution kernel is used for the smoothing process. Possible choices:

"box", "b" (default)

to use a convolution kernel with sharp edges.

"gaussian", "g"

to use a convolution kernel that is represented by a non-correlated trivariate normal distribution function.

sz is either a vector of 3 elements representing the size of the convolution kernel in x-, y- and z-direction or a scalar, in which case the same size is used in all three dimensions. The default value is 3.

When *method* is "gaussian", *std_dev* defines the standard deviation of the trivariate normal distribution function. *std_dev* is either a vector of 3 elements representing the standard deviation of the Gaussian convolution kernel in x-, y- and z-directions or a scalar, in which case the same value is used in all three dimensions. The default value is 0.65.

See also: [isosurface], page 338, [isonormals], page 340, [patch], page 397.

```
[nx, ny, nz, nv] = reducevolume (v, r)
[nx, ny, nz, nv] = reducevolume (x, y, z, v, r)
nv = reducevolume (...)
```

Reduce the volume of the dataset in *v* according to the values in *r*.

v is a matrix that is non-singleton in the first 3 dimensions.

r can either be a vector of 3 elements representing the reduction factors in the x-, y-, and z-directions or a scalar, in which case the same reduction factor is used in all three dimensions.

`reducevolume` reduces the number of elements of *v* by taking only every *r*-th element in the respective dimension.

Optionally, *x*, *y*, and *z* can be supplied to represent the set of coordinates of *v*. They can either be matrices of the same size as *v* or vectors with sizes according to the dimensions of *v*, in which case they are expanded to matrices (see [meshgrid], page 347).

If `reducevolume` is called with two arguments then *x*, *y*, and *z* are assumed to match the respective indices of *v*.

The reduced matrix is returned in *nv*.

Optionally, the reduced set of coordinates are returned in *nx*, *ny*, and *nz*, respectively.

Examples:

```
v = reshape (1:6*8*4, [6 8 4]);
nv = reducevolume (v, [4 3 2]);
v = reshape (1:6*8*4, [6 8 4]);
x = 1:3:24; y = -14:5:11; z = linspace (16, 18, 4);
[nx, ny, nz, nv] = reducevolume (x, y, z, v, [4 3 2]);
```

See also: [\[isosurface\]](#), page 338, [\[isornormals\]](#), page 340.

```
reduced_fv = reducepatch (fv)
reduced_fv = reducepatch (faces, vertices)
reduced_fv = reducepatch (patch_handle)
reducepatch (patch_handle)
reduced_fv = reducepatch (... , reduction_factor)
reduced_fv = reducepatch (... , "fast")
reduced_fv = reducepatch (... , "verbose")
[reduced_faces, reduced_vertices] = reducepatch (...)
```

Reduce the number of faces and vertices in a patch object while retaining the overall shape of the patch.

The input patch can be represented by a structure *fv* with the fields *faces* and *vertices*, by two matrices *faces* and *vertices* (see, e.g., the result of `isosurface`), or by a handle to a patch object *patch_handle* (see [\[patch\]](#), page 397).

The number of faces and vertices in the patch is reduced by iteratively collapsing the shortest edge of the patch to its midpoint (as discussed, e.g., here: <https://libigl.github.io/libigl/tutorial/tutorial.html#meshdecimation>).

Currently, only patches consisting of triangles are supported. The resulting patch also consists only of triangles.

If `reducepatch` is called with a handle to a valid patch `patch_handle`, and without any output arguments, then the given patch is updated immediately.

If the `reduction_factor` is omitted, the resulting structure `reduced_fv` includes approximately 50% of the faces of the original patch. If `reduction_factor` is a fraction between 0 (excluded) and 1 (excluded), a patch with approximately the corresponding fraction of faces is determined. If `reduction_factor` is an integer greater than or equal to 1, the resulting patch has approximately `reduction_factor` faces. Depending on the geometry of the patch, the resulting number of faces can differ from the given value of `reduction_factor`. This is especially true when many shared vertices are detected.

For the reduction, it is necessary that vertices of touching faces are shared. Shared vertices are detected automatically. This detection can be skipped by passing the optional string argument "`fast`".

With the optional string arguments "`verbose`", additional status messages are printed to the command window.

Any string input arguments must be passed after all other arguments.

If called with one output argument, the reduced faces and vertices are returned in a structure `reduced_fv` with the fields `faces` and `vertices` (see the one output option of `isosurface`).

If called with two output arguments, the reduced faces and vertices are returned in two separate matrices `reduced_faces` and `reduced_vertices`.

See also: [\[isosurface\]](#), page 338, [\[isornormals\]](#), page 340, [\[reducevolume\]](#), page 343, [\[patch\]](#), page 397.

```
shrinkfaces (p, sf)
nfv = shrinkfaces (p, sf)
nfv = shrinkfaces (fv, sf)
nfv = shrinkfaces (f, v, sf)
[nf, nv] = shrinkfaces (...)
```

Reduce the size of faces in a patch by the shrink factor `sf`.

The patch object can be specified by a graphics handle (`p`), a patch structure (`fv`) with the fields "`faces`" and "`vertices`", or as two separate matrices (`f`, `v`) of faces and vertices.

The shrink factor `sf` is a positive number specifying the percentage of the original area the new face will occupy. If no factor is given the default is 0.3 (a reduction to 30% of the original size). A factor greater than 1.0 will result in the expansion of faces.

Given a patch handle as the first input argument and no output parameters, perform the shrinking of the patch faces in place and redraw the patch.

If called with one output argument, return a structure with fields "`faces`", "`vertices`", and "`facevertexcdata`" containing the data after shrinking. This structure can be used directly as an input argument to the `patch` function.

Caution: Performing the shrink operation on faces which are not convex can lead to undesirable results.

Example: a triangulated 3/4 circle and the corresponding shrunken version.

```
[phi r] = meshgrid (linspace (0, 1.5*pi, 16), linspace (1, 2, 4));
tri = delaunay (phi(:), r(:));
v = [r(:).*sin(phi(:)) r(:).*cos(phi(:))];
clf ()
p = patch ("Faces", tri, "Vertices", v, "FaceColor", "none");
fv = shrinkfaces (p);
patch (fv)
axis equal
grid on
```

See also: [\[patch\]](#), page 397.

diffuse (sx, sy, sz, lv)

Calculate the diffuse reflection strength of a surface defined by the normal vector elements *sx*, *sy*, *sz*.

The light source location vector *lv* can be given as a 2-element vector [azimuth, elevation] in degrees or as a 3-element vector [x, y, z].

See also: [\[specular\]](#), page 345, [\[surfl\]](#), page 336.

specular (sx, sy, sz, lv, vv)

specular (sx, sy, sz, lv, vv, se)

Calculate the specular reflection strength of a surface defined by the normal vector elements *sx*, *sy*, *sz* using Phong's approximation.

The light source location and viewer location vectors are specified using parameters *lv* and *vv* respectively. The location vectors can given as 2-element vectors [azimuth, elevation] in degrees or as 3-element vectors [x, y, z].

An optional sixth argument specifies the specular exponent (spread) *se*. If not given, *se* defaults to 10.

See also: [\[diffuse\]](#), page 345, [\[surfl\]](#), page 336.

lighting (type)

lighting (hax, type)

Set the lighting of patch or surface graphic objects.

Valid arguments for *type* are

"flat" Draw objects with faceted lighting effects.

"gouraud"

Draw objects with linear interpolation of the lighting effects between the vertices.

"none" Draw objects without light and shadow effects.

If the first argument *hax* is an axes handle, then change the lighting effects of objects in this axes, rather than the current axes returned by [gca](#).

The lighting effects are only visible if at least one light object is present and visible in the same axes.

See also: [\[light\]](#), page 398, [\[fill\]](#), page 322, [\[mesh\]](#), page 333, [\[patch\]](#), page 397, [\[pcolor\]](#), page 321, [\[surf\]](#), page 335, [\[surface\]](#), page 398, [\[shading\]](#), page 356.

```

material shiny
material dull
material metal
material default
material ([as, ds, ss])
material ([as, ds, ss, se])
material ([as, ds, ss, se, scr])
material (hlist, ...)
mtypes = material ()
refl_props = material (mtype_string)

```

Set reflectance properties for the lighting of surfaces and patches.

This function changes the ambient, diffuse, and specular strengths, as well as the specular exponent and specular color reflectance, of all **patch** and **surface** objects in the current axes. This can be used to simulate, to some extent, the reflectance properties of certain materials when used with **light**.

When called with a string, the aforementioned properties are set according to the values in the following table:

<i>mtype</i>	ambient-strength	diffuse-strength	specular-strength	specular-exponent	specular-color-reflectance
"shiny"	0.3	0.6	0.9	20	1.0
"dull"	0.3	0.8	0.0	10	1.0
"metal"	0.3	0.3	1.0	25	0.5
"default"	"default"	"default"	"default"	"default"	"default"

When called with a vector of three elements, the ambient, diffuse, and specular strengths of all **patch** and **surface** objects in the current axes are updated. An optional fourth vector element updates the specular exponent, and an optional fifth vector element updates the specular color reflectance.

A list of graphic handles can also be passed as the first argument. In this case, the properties of these handles and all child **patch** and **surface** objects will be updated.

Additionally, **material** can be called with a single output argument. If called without input arguments, a column cell vector **mtypes** with the strings for all available materials is returned. If the one input argument **mtype_string** is the name of a material, a 1x5 cell vector **refl_props** with the reflectance properties of that material is returned. In both cases, no graphic properties are changed.

See also: [\[light\]](#), page 398, [\[fill\]](#), page 322, [\[mesh\]](#), page 333, [\[patch\]](#), page 397, [\[pcolor\]](#), page 321, [\[surf\]](#), page 335, [\[surface\]](#), page 398.

```

camlight
camlight right
camlight left
camlight headlight
camlight (az, el)
camlight (... , style)

```

```
camlight (hl, ...)
camlight (hax, ...)
h = camlight (...)
```

Add a light object to a figure using a simple interface.

When called with no arguments, a light object is added to the current plot and is placed slightly above and to the right of the camera's current position: this is equivalent to `camlight right`. The commands `camlight left` and `camlight headlight` behave similarly with the placement being either left of the camera position or centered on the camera position.

For more control, the light position can be specified by an azimuthal rotation `az` and an elevation angle `el`, both in degrees, relative to the current properties of the camera.

The optional string `style` specifies whether the light is a local point source ("`local`", the default) or placed at infinite distance ("`infinite`").

If the first argument `hl` is a handle to a light object, then act on this light object rather than creating a new object.

If the first argument `hax` is an axes handle, then create a new light object in this axes, rather than the current axes returned by `gca`.

The optional return value `h` is a graphics handle to the light object. This can be used to move or further change properties of the light object.

Examples:

Add a light object to a plot

```
sphere (36);
camlight
```

Position the light source exactly

```
camlight (45, 30);
```

Here the light is first pitched upwards (see [\[camup\], page 353](#)) from the camera position (see [\[campos\], page 350](#)) by 30 degrees. It is then yawed by 45 degrees to the right. Both rotations are centered around the camera target (see [\[camtarget\], page 352](#)).

Return a handle to further manipulate the light object

```
clf
sphere (36);
hl = camlight ("left");
set (hl, "color", "r");
```

See also: [\[light\], page 398](#).

```
[xx, yy] = meshgrid (x, y)
[xx, yy, zz] = meshgrid (x, y, z)
[xx, yy] = meshgrid (x)
[xx, yy, zz] = meshgrid (x)
```

Given vectors of `x` and `y` coordinates, return matrices `xx` and `yy` corresponding to a full 2-D grid.

The rows of `xx` are copies of `x`, and the columns of `yy` are copies of `y`. If `y` is omitted, then it is assumed to be the same as `x`.

If the optional *z* input is given, or *zz* is requested, then the output will be a full 3-D grid. If *z* is omitted and *zz* is requested, it is assumed to be the same as *y*.

meshgrid is most frequently used to produce input for a 2-D or 3-D function that will be plotted. The following example creates a surface plot of the “sombrero” function.

```
f = @(x,y) sin (sqrt (x.^2 + y.^2)) ./ sqrt (x.^2 + y.^2);
range = linspace (-8, 8, 41);
[X, Y] = meshgrid (range, range);
Z = f (X, Y);
surf (X, Y, Z);
```

Programming Note: **meshgrid** is restricted to 2-D or 3-D grid generation. The **ndgrid** function will generate 1-D through N-D grids. However, the functions are not completely equivalent. If *x* is a vector of length *M* and *y* is a vector of length *N*, then **meshgrid** will produce an output grid which is *NxM*. **ndgrid** will produce an output which is *MxN* (transpose) for the same input. Some core functions expect **meshgrid** input and others expect **ndgrid** input. Check the documentation for the function in question to determine the proper input format.

See also: [\[ndgrid\]](#), page 348, [\[mesh\]](#), page 333, [\[contour\]](#), page 310, [\[surf\]](#), page 335.

```
[y1, y2, ..., yn] = ndgrid (x1, x2, ..., xn)
[y1, y2, ..., yn] = ndgrid (x)
```

Given *n* vectors *x₁*, ..., *x_n*, **ndgrid** returns *n* arrays of dimension *n*.

The elements of the *i*-th output argument contains the elements of the vector *x_i* repeated over all dimensions different from the *i*-th dimension. Calling **ndgrid** with only one input argument *x* is equivalent to calling **ndgrid** with all *n* input arguments equal to *x*:

```
[y1, y2, ..., yn] = ndgrid (x, ..., x)
```

Programming Note: **ndgrid** is very similar to the function **meshgrid** except that the first two dimensions are transposed in comparison to **meshgrid**. Some core functions expect **meshgrid** input and others expect **ndgrid** input. Check the documentation for the function in question to determine the proper input format.

See also: [\[meshgrid\]](#), page 347.

```
plot3 (x, y, z)
plot3 (x, y, z, prop, value, ...)
plot3 (x, y, z, fmt)
plot3 (x, cplx)
plot3 (cplx)
plot3 (hax, ...)
h = plot3 (...)
```

Produce 3-D plots.

Many different combinations of arguments are possible. The simplest form is

```
plot3 (x, y, z)
```

in which the arguments are taken to be the vertices of the points to be plotted in three dimensions. If all arguments are vectors of the same length, then a single continuous line is drawn. If all arguments are matrices, then each column of is treated as a

separate line. No attempt is made to transpose the arguments to make the number of rows match.

If only two arguments are given, as

```
plot3 (x, cplx)
```

the real and imaginary parts of the second argument are used as the *y* and *z* coordinates, respectively.

If only one argument is given, as

```
plot3 (cplx)
```

the real and imaginary parts of the argument are used as the *y* and *z* values, and they are plotted versus their index.

Arguments may also be given in groups of three as

```
plot3 (x1, y1, z1, x2, y2, z2, ...)
```

in which each set of three arguments is treated as a separate line or set of lines in three dimensions.

To plot multiple one- or two-argument groups, separate each group with an empty format string, as

```
plot3 (x1, c1, "", c2, "", ...)
```

Multiple property-value pairs may be specified which will affect the line objects drawn by *plot3*. If the *fmt* argument is supplied it will format the line objects in the same manner as *plot*.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by *gca*.

The optional return value *h* is a graphics handle to the created plot.

Example:

```
z = [0:0.05:5];
plot3 (cos (2*pi*z), sin (2*pi*z), z, ";helix;");
plot3 (z, exp (2i*pi*z), ";complex sinusoid;");
```

See also: [\[ezplot3\]](#), page 359, [\[plot\]](#), page 296.

```
view (azimuth, elevation)
view ([azimuth elevation])
view ([x y z])
view (2)
view (3)
view (hax, ...)
[azimuth, elevation] = view ()
```

Query or set the viewpoint for the current axes.

The parameters *azimuth* and *elevation* can be given as two arguments or as 2-element vector. The viewpoint can also be specified with Cartesian coordinates *x*, *y*, and *z*.

The call *view* (2) sets the viewpoint to *azimuth* = 0 and *elevation* = 90, which is the default for 2-D graphs.

The call *view* (3) sets the viewpoint to *azimuth* = -37.5 and *elevation* = 30, which is the default for 3-D graphs.

If the first argument *hax* is an axes handle, then operate on this axes rather than the current axes returned by `gca`.

If no inputs are given, return the current *azimuth* and *elevation*.

```
camlookat ()
camlookat (h)
camlookat (handle_list)
camlookat (hax)
```

Move the camera and adjust its properties to look at objects.

When the input is a handle *h*, the camera is set to point toward the center of the bounding box of *h*. The camera's position is adjusted so the bounding box approximately fills the field of view.

This command fixes the camera's viewing direction (`camtarget()` – `campos()`), camera up vector (see [camup], page 353) and viewing angle (see [camva], page 354). The camera target (see [camtarget], page 352) and camera position (see [campos], page 350) are changed.

If the argument is a list *handle_list*, then a single bounding box for all the objects is computed and the camera is then adjusted as above.

If the argument is an axis object *hax*, then the children of the axis are used as *handle_list*. When called with no inputs, it uses the current axis (see [gca], page 401).

See also: [camorbit], page 351, [camzoom], page 354, [camroll], page 352.

```
P = campos ()
campos ([x y z])
mode = campos ("mode")
campos (mode)
campos (ax, ...)
```

Set or get the camera position.

The default camera position is determined automatically based on the scene. For example, to get the camera position:

```
hf = figure();
peaks()
p = campos ()
⇒ p =
-27.394 -35.701 64.079
```

We can then move the camera further up the z-axis:

```
campos (p + [0 0 10])
campos ()
⇒ ans =
-27.394 -35.701 74.079
```

Having made that change, the camera position *mode* is now manual:

```
campos ("mode")
⇒ manual
```

We can set it back to automatic:

```
campos ("auto")
campos ()
⇒ ans =
-27.394 -35.701 64.079
close (hf)
```

By default, these commands affect the current axis; alternatively, an axis can be specified by the optional argument *ax*.

See also: [camup], page 353, [camtarget], page 352, [camva], page 354.

```
camorbit (theta, phi)
camorbit (theta, phi, coorsys)
camorbit (theta, phi, coorsys, dir)
camorbit (theta, phi, "data")
camorbit (theta, phi, "data", "z")
camorbit (theta, phi, "data", "x")
camorbit (theta, phi, "data", "y")
camorbit (theta, phi, "data", [x y z])
camorbit (theta, phi, "camera")
camorbit (hax, ...)
```

Rotate the camera up/down and left/right around its target.

Move the camera *phi* degrees up and *theta* degrees to the right, as if it were in an orbit around its target. Example:

```
sphere ()
camorbit (30, 20)
```

These rotations are centered around the camera target (see [camtarget], page 352). First the camera position is pitched up or down by rotating it *phi* degrees around an axis orthogonal to both the viewing direction (specifically `camtarget()` - `campos()`) and the camera “up vector” (see [camup], page 353). Example:

```
camorbit (0, 20)
```

The second rotation depends on the coordinate system *coorsys* and direction *dir* inputs. The default for *coorsys* is “*data*”. In this case, the camera is yawed left or right by rotating it *theta* degrees around an axis specified by *dir*. The default for *dir* is “*z*”, corresponding to the vector [0, 0, 1]. Example:

```
camorbit (30, 0)
```

When *coorsys* is set to “*camera*”, the camera is moved left or right by rotating it around an axis parallel to the camera up vector (see [camup], page 353). The input *dir* should not be specified in this case. Example:

```
camorbit (30, 0, "camera")
```

(Note: the rotation by *phi* is unaffected by “*camera*”.)

The `camorbit` command modifies two camera properties: see [campos], page 350, and see [camup], page 353.

By default, this command affects the current axis; alternatively, an axis can be specified by the optional argument *hax*.

See also: [camzoom], page 354, [camroll], page 352, [camlookat], page 350.

```
camroll (theta)
camroll (ax, theta)
```

Roll the camera.

Roll the camera clockwise by *theta* degrees. For example, the following command will roll the camera by 30 degrees clockwise (to the right); this will cause the scene to appear to roll by 30 degrees to the left:

```
peaks ()
camroll (30)
```

Roll the camera back:

```
camroll (-30)
```

The following command restores the default camera roll:

```
camup ("auto")
```

By default, these commands affect the current axis; alternatively, an axis can be specified by the optional argument *ax*.

See also: [camzoom], page 354, [camorbit], page 351, [camlookat], page 350, [camup], page 353.

```
T = camtarget ()
camtarget ([x y z])
mode = camtarget ("mode")
camtarget (mode)
camtarget (ax, ...)
```

Set or get where the camera is pointed.

The camera target is a point in space where the camera is pointing. Usually, it is determined automatically based on the scene:

```
hf = figure();
sphere (36)
v = camtarget ()
⇒ v =
    0    0    0
```

We can turn the camera to point at a new target:

```
camtarget ([1 1 1])
camtarget ()
⇒ 1 1 1
```

Having done so, the camera target *mode* is manual:

```
camtarget ("mode")
⇒ manual
```

This means, for example, adding new objects to the scene will not retarget the camera:

```
hold on;
peaks ()
camtarget ()
⇒ 1 1 1
```

We can reset it to be automatic:

```
camtarget ("auto")
camtarget ()
⇒ 0 0 0.76426
close (hf)
```

By default, these commands affect the current axis; alternatively, an axis can be specified by the optional argument *ax*.

See also: [\[campos\]](#), page 350, [\[camup\]](#), page 353, [\[camva\]](#), page 354.

```
up = camup ()
camup ([x y z])
mode = camup ("mode")
camup (mode)
camup (ax, ...)
```

Set or get the camera up vector.

By default, the camera is oriented so that “up” corresponds to the positive z-axis:

```
hf = figure ();
sphere (36)
v = camup ()
⇒ v =
0 0 1
```

Specifying a new “up vector” rolls the camera and sets the mode to manual:

```
camup ([1 1 0])
camup ()
⇒ 1 1 0
camup ("mode")
⇒ manual
```

Modifying the up vector does not modify the camera target (see [\[camtarget\]](#), page 352). Thus, the camera up vector might not be orthogonal to the direction of the camera’s view:

```
camup ([1 2 3])
dot (camup (), camtarget () - campos ())
⇒ 6...
```

A consequence is that “pulling back” on the up vector does not pitch the camera view (as that would require changing the target). Setting the up vector is thus typically used only to roll the camera. A more intuitive command for this purpose is see [\[camroll\]](#), page 352.

Finally, we can reset the up vector to automatic mode:

```
camup ("auto")
camup ()
⇒ 0 0 1
close (hf)
```

By default, these commands affect the current axis; alternatively, an axis can be specified by the optional argument *ax*.

See also: [\[campos\]](#), page 350, [\[camtarget\]](#), page 352, [\[camva\]](#), page 354.

```
a = camva ()
camva (a)
mode = camva ("mode")
camva (mode)
camva (ax, ...)
```

Set or get the camera viewing angle.

The camera has a viewing angle which determines how much can be seen. By default this is:

```
hf = figure();
sphere (36)
a = camva ()
⇒ a = 10.340
```

To get a wider-angle view, we could double the viewing angle. This will also set the mode to manual:

```
camva (2*a)
camva ("mode")
⇒ manual
```

We can set it back to automatic:

```
camva ("auto")
camva ("mode")
⇒ auto
camva ()
⇒ ans = 10.340
close (hf)
```

By default, these commands affect the current axis; alternatively, an axis can be specified by the optional argument *ax*.

See also: [campos], page 350, [camtarget], page 352, [camup], page 353.

```
camzoom (zf)
camzoom (ax, zf)
```

Zoom the camera in or out.

A value of *zf* larger than 1 “zooms in” such that the scene appears magnified:

```
hf = figure ();
sphere (36)
camzoom (1.2)
```

A value smaller than 1 “zooms out” so the camera can see more of the scene:

```
camzoom (0.5)
```

Technically speaking, zooming affects the “viewing angle”. The following command resets to the default zoom:

```
camva ("auto")
close (hf)
```

By default, these commands affect the current axis; alternatively, an axis can be specified by the optional argument *ax*.

See also: [camroll], page 352, [camorbit], page 351, [camlookat], page 350, [camva], page 354.

```
slice (x, y, z, v, sx, sy, sz)
slice (x, y, z, v, xi, yi, zi)
slice (v, sx, sy, sz)
slice (v, xi, yi, zi)
slice (..., method)
slice (hax, ...)
h = slice (...)
```

Plot slices of 3-D data/scalar fields.

Each element of the 3-dimensional array *v* represents a scalar value at a location given by the parameters *x*, *y*, and *z*. The parameters *x*, *x*, and *z* are either 3-dimensional arrays of the same size as the array *v* in the "meshgrid" format or vectors. The parameters *xi*, etc. respect a similar format to *x*, etc., and they represent the points at which the array *vi* is interpolated using interp3. The vectors *sx*, *sy*, and *sz* contain points of orthogonal slices of the respective axes.

If *x*, *y*, *z* are omitted, they are assumed to be *x* = 1:size (*v*, 2), *y* = 1:size (*v*, 1) and *z* = 1:size (*v*, 3).

method is one of:

- "nearest" Return the nearest neighbor.
- "linear" Linear interpolation from nearest neighbors.
- "cubic" Cubic interpolation from four nearest neighbors (not implemented yet).
- "spline" Cubic spline interpolation—smooth first and second derivatives throughout the curve.

The default method is "linear".

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by gca.

The optional return value *h* is a graphics handle to the created surface object.

Examples:

```
[x, y, z] = meshgrid (linspace (-8, 8, 32));
v = sin (sqrt (x.^2 + y.^2 + z.^2)) ./ (sqrt (x.^2 + y.^2 + z.^2));
slice (x, y, z, v, [], 0, []);

[xi, yi] = meshgrid (linspace (-7, 7));
zi = xi + yi;
slice (x, y, z, v, xi, yi, zi);
```

See also: [interp3], page 770, [surface], page 398, [pcolor], page 321.

```
ribbon (y)
ribbon (x, y)
ribbon (x, y, width)
```

```
ribbon (hax, ...)
h = ribbon (...)
```

Draw a ribbon plot for the columns of *y* vs. *x*.

If *x* is omitted, a vector containing the row numbers is assumed (`1:rows (Y)`). Alternatively, *x* can also be a vector with same number of elements as rows of *y* in which case the same *x* is used for each column of *y*.

The optional parameter *width* specifies the width of a single ribbon (default is 0.75).

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a vector of graphics handles to the surface objects representing each ribbon.

See also: [\[surface\]](#), page 398, [\[waterfall\]](#), page 357.

```
shading (type)
shading (hax, type)
```

Set the shading of patch or surface graphic objects.

Valid arguments for *type* are

"flat" Single colored patches with invisible edges.

"faceted"

Single colored patches with black edges.

"interp" Colors between patch vertices are interpolated and the patch edges are invisible.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

See also: [\[fill\]](#), page 322, [\[mesh\]](#), page 333, [\[patch\]](#), page 397, [\[pcolor\]](#), page 321, [\[surf\]](#), page 335, [\[surface\]](#), page 398, [\[hidden\]](#), page 335, [\[lighting\]](#), page 345.

```
scatter3 (x, y, z)
scatter3 (x, y, z, s)
scatter3 (x, y, z, s, c)
scatter3 (..., style)
scatter3 (..., "filled")
scatter3 (..., prop, val)
scatter3 (hax, ...)
h = scatter3 (...)
```

Draw a 3-D scatter plot.

A marker is plotted at each point defined by the coordinates in the vectors *x*, *y*, and *z*.

The size of the markers is determined by *s*, which can be a scalar or a vector of the same length as *x*, *y*, and *z*. If *s* is not given, or is an empty matrix, then a default value of 8 points is used.

The color of the markers is determined by *c*, which can be a string defining a fixed color; a 3-element vector giving the red, green, and blue components of the color; a

vector of the same length as *x* that gives a scaled index into the current colormap; or an Nx3 matrix defining the RGB color of each marker individually.

The marker to use can be changed with the *style* argument, that is a string defining a marker in the same manner as the `plot` command. If no marker is specified it defaults to "o" or circles. If the argument "filled" is given then the markers are filled.

Additional property/value pairs are passed directly to the underlying patch object.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the hggroup object representing the points.

```
[x, y, z] = peaks (20);
scatter3 (x(:), y(:), z(:), [], z(:));
```

See also: [\[scatter\]](#), page 307, [\[patch\]](#), page 397, [\[plot\]](#), page 296.

```
waterfall (x, y, z)
waterfall (z)
waterfall (..., c)
waterfall (..., prop, val, ...)
waterfall (hax, ...)
h = waterfall (...)
```

Plot a 3-D waterfall plot.

A waterfall plot is similar to a `meshz` plot except only mesh lines for the rows of *z* (*x*-values) are shown.

The wireframe mesh is plotted using rectangles. The vertices of the rectangles [*x*, *y*] are typically the output of `meshgrid`. over a 2-D rectangular region in the *x*-*y* plane. *z* determines the height above the plane of each vertex. If only a single *z* matrix is given, then it is plotted over the meshgrid `x = 1:columns (z)`, `y = 1:rows (z)`. Thus, columns of *z* correspond to different *x* values and rows of *z* correspond to different *y* values.

The color of the mesh is computed by linearly scaling the *z* values to fit the range of the current colormap. Use `caxis` and/or change the colormap to control the appearance.

Optionally the color of the mesh can be specified independently of *z* by supplying a color matrix, *c*.

Any property/value pairs are passed directly to the underlying surface object.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the created surface object.

See also: [\[meshz\]](#), page 334, [\[mesh\]](#), page 333, [\[meshc\]](#), page 334, [\[contour\]](#), page 310, [\[surf\]](#), page 335, [\[surface\]](#), page 398, [\[ribbon\]](#), page 355, [\[meshgrid\]](#), page 347, [\[hidden\]](#), page 335, [\[shading\]](#), page 356, [\[colormap\]](#), page 816, [\[caxis\]](#), page 325.

15.2.2.1 Aspect Ratio

For three-dimensional plots the aspect ratio can be set for data with `daspect` and for the plot box with `pbaspect`. See [Section 15.2.1.1 \[Axis Configuration\]](#), page 323, for controlling the x-, y-, and z-limits for plotting.

```
data_aspect_ratio = daspect ()
daspect (data_aspect_ratio)
daspect (mode)
data_aspect_ratio_mode = daspect ("mode")
daspect (hax, ...)
```

Query or set the data aspect ratio of the current axes.

The aspect ratio is a normalized 3-element vector representing the span of the x, y, and z-axis limits.

`daspect (mode)`

Set the data aspect ratio mode of the current axes. `mode` is either "auto" or "manual".

`daspect ("mode")`

Return the data aspect ratio mode of the current axes.

`daspect (hax, ...)`

Operate on the axes in handle `hax` instead of the current axes.

See also: [\[axis\]](#), page 323, [\[pbaspect\]](#), page 358, [\[xlim\]](#), page 325, [\[ylim\]](#), page 325, [\[zlim\]](#), page 325.

```
plot_box_aspect_ratio = pbaspect ( )
pbaspect (plot_box_aspect_ratio)
pbaspect (mode)
plot_box_aspect_ratio_mode = pbaspect ("mode")
pbaspect (hax, ...)
```

Query or set the plot box aspect ratio of the current axes.

The aspect ratio is a normalized 3-element vector representing the rendered lengths of the x, y, and z axes.

`pbaspect(mode)`

Set the plot box aspect ratio mode of the current axes. `mode` is either "auto" or "manual".

`pbaspect ("mode")`

Return the plot box aspect ratio mode of the current axes.

`pbaspect (hax, ...)`

Operate on the axes in handle `hax` instead of the current axes.

See also: [\[axis\]](#), page 323, [\[daspect\]](#), page 358, [\[xlim\]](#), page 325, [\[ylim\]](#), page 325, [\[zlim\]](#), page 325.

15.2.2.2 Three-dimensional Function Plotting

```
ezplot3 (fx, fy, fz)
ezplot3 (..., dom)
ezplot3 (..., n)
ezplot3 (..., "animate")
ezplot3 (hax, ...)
h = ezplot3 (...)
```

Plot a parametrically defined curve in three dimensions.

fx, , and *fz* are strings, inline functions, or function handles with one argument defining the function. By default the plot is over the domain $0 \leq t \leq 2\pi$ with 500 points.

If *dom* is a two element vector, it represents the minimum and maximum values of *t*. *n* is a scalar defining the number of points to use in plotting the function.

If the "animate" option is given then the plotting is animated in the style of **comet3**.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by **gca**.

The optional return value *h* is a graphics handle to the created plot.

```
fx = @(t) cos (t);
fy = @(t) sin (t);
fz = @(t) t;
ezplot3 (fx, fy, fz, [0, 10*pi], 100);
```

See also: [\[plot3\]](#), page 348, [\[comet3\]](#), page 323, [\[ezplot\]](#), page 328, [\[ezmesh\]](#), page 359, [\[ezsurf\]](#), page 361.

```
ezmesh (f)
ezmesh (fx, fy, fz)
ezmesh (..., dom)
ezmesh (..., n)
ezmesh (..., "circ")
ezmesh (hax, ...)
h = ezmesh (...)
```

Plot the mesh defined by a function.

f is a string, inline function, or function handle with two arguments defining the function. By default the plot is over the meshed domain $-2\pi \leq x | y \leq 2\pi$ with 60 points in each dimension.

If three functions are passed, then plot the parametrically defined function [*fx(s, t)*, *fy(s, t)*, *fz(s, t)*].

If *dom* is a two element vector, it represents the minimum and maximum values of both *x* and *y*. If *dom* is a four element vector, then the minimum and maximum values are [*xmin xmax ymin ymax*].

n is a scalar defining the number of points to use in each dimension.

If the argument "circ" is given, then the function is plotted over a disk centered on the middle of the domain *dom*.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the created surface object.

Example 1: 2-argument function

```
f = @(x,y) sqrt (abs (x .* y)) ./ (1 + x.^2 + y.^2);
ezmesh (f, [-3, 3]);
```

Example 2: parametrically defined function

```
fx = @(s,t) cos (s) .* cos (t);
fy = @(s,t) sin (s) .* cos (t);
fz = @(s,t) sin (t);
ezmesh (fx, fy, fz, [-pi, pi, -pi/2, pi/2], 20);
```

See also: [\[mesh\]](#), page 333, [\[ezmeshc\]](#), page 360, [\[ezplot\]](#), page 328, [\[ezsurf\]](#), page 361, [\[ezsurfc\]](#), page 361, [\[hidden\]](#), page 335.

```
ezmeshc (f)
ezmeshc (fx, fy, fz)
ezmeshc (..., dom)
ezmeshc (..., n)
ezmeshc (..., "circ")
ezmeshc (hax, ...)
h = ezmeshc (...)
```

Plot the mesh and contour lines defined by a function.

f is a string, inline function, or function handle with two arguments defining the function. By default the plot is over the meshed domain $-2\pi \leq x \mid y \leq 2\pi$ with 60 points in each dimension.

If three functions are passed, then plot the parametrically defined function [*fx(s, t)*, *fy(s, t)*, *fz(s, t)*].

If *dom* is a two element vector, it represents the minimum and maximum values of both *x* and *y*. If *dom* is a four element vector, then the minimum and maximum values are [*xmin xmax ymin ymax*].

n is a scalar defining the number of points to use in each dimension.

If the argument "*circ*" is given, then the function is plotted over a disk centered on the middle of the domain *dom*.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a 2-element vector with a graphics handle for the created mesh plot and a second handle for the created contour plot.

Example: 2-argument function

```
f = @(x,y) sqrt (abs (x .* y)) ./ (1 + x.^2 + y.^2);
ezmeshc (f, [-3, 3]);
```

See also: [\[meshc\]](#), page 334, [\[ezmesh\]](#), page 359, [\[ezplot\]](#), page 328, [\[ezsurf\]](#), page 361, [\[ezsurfc\]](#), page 361, [\[hidden\]](#), page 335.

```
ezsurf (f)
ezsurf (fx, fy, fz)
ezsurf (..., dom)
ezsurf (..., n)
ezsurf (..., "circ")
ezsurf (hax, ...)
ezsurf (...)
```

Plot the surface defined by a function.

f is a string, inline function, or function handle with two arguments defining the function. By default the plot is over the meshed domain $-2\pi \leq x \mid y \leq 2\pi$ with 60 points in each dimension.

If three functions are passed, then plot the parametrically defined function [*fx(s, t)*, *fy(s, t)*, *fz(s, t)*].

If *dom* is a two element vector, it represents the minimum and maximum values of both *x* and *y*. If *dom* is a four element vector, then the minimum and maximum values are [*xmin xmax ymin ymax*].

n is a scalar defining the number of points to use in each dimension.

If the argument "circ" is given, then the function is plotted over a disk centered on the middle of the domain *dom*.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by *gca*.

The optional return value *h* is a graphics handle to the created surface object.

Example 1: 2-argument function

```
f = @(x,y) sqrt (abs (x .* y)) ./ (1 + x.^2 + y.^2);
ezsurf (f, [-3, 3]);
```

Example 2: parametrically defined function

```
fx = @(s,t) cos (s) .* cos (t);
fy = @(s,t) sin (s) .* cos (t);
fz = @(s,t) sin (t);
ezsurf (fx, fy, fz, [-pi, pi, -pi/2, pi/2], 20);
```

See also: [\[surf\]](#), page 335, [\[ezsurfc\]](#), page 361, [\[ezplot\]](#), page 328, [\[ezmesh\]](#), page 359, [\[ezmeshc\]](#), page 360, [\[shading\]](#), page 356.

```
ezsurfc (f)
ezsurfc (fx, fy, fz)
ezsurfc (..., dom)
ezsurfc (..., n)
ezsurfc (..., "circ")
ezsurfc (hax, ...)
ezsurfc (...)
```

Plot the surface and contour lines defined by a function.

f is a string, inline function, or function handle with two arguments defining the function. By default the plot is over the meshed domain $-2\pi \leq x \mid y \leq 2\pi$ with 60 points in each dimension.

If three functions are passed, then plot the parametrically defined function [$fx(s, t)$, $fy(s, t)$, $fz(s, t)$].

If dom is a two element vector, it represents the minimum and maximum values of both x and y . If dom is a four element vector, then the minimum and maximum values are [$xmin$ $xmax$ $ymin$ $ymax$].

n is a scalar defining the number of points to use in each dimension.

If the argument "circ" is given, then the function is plotted over a disk centered on the middle of the domain dom .

If the first argument hax is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value h is a 2-element vector with a graphics handle for the created surface plot and a second handle for the created contour plot.

Example:

```
f = @(x,y) sqrt (abs (x .* y)) ./ (1 + x.^2 + y.^2);
ezsurf (f, [-3, 3]);
```

See also: [\[surf\]](#), page 336, [\[ezsurf\]](#), page 361, [\[ezplot\]](#), page 328, [\[ezmesh\]](#), page 359, [\[ezmeshc\]](#), page 360, [\[shading\]](#), page 356.

15.2.2.3 Three-dimensional Geometric Shapes

```
cylinder
cylinder (r)
cylinder (r, n)
cylinder (hax, ...)
[x, y, z] = cylinder (...)
```

Plot a 3-D unit cylinder.

The optional input r is a vector specifying the radius along the unit z-axis. The default is [1 1] indicating radius 1 at $Z == 0$ and at $Z == 1$.

The optional input n determines the number of faces around the circumference of the cylinder. The default value is 20.

If the first argument hax is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

If outputs are requested `cylinder` returns three matrices in `meshgrid` format, such that `surf (x, y, z)` generates a unit cylinder.

Example:

```
[x, y, z] = cylinder (10:-1:0, 50);
surf (x, y, z);
title ("a cone");
```

See also: [\[ellipsoid\]](#), page 363, [\[rectangle\]](#), page 331, [\[sphere\]](#), page 362.

```
sphere ()
sphere (n)
sphere (hax, ...)
```

```
[x, y, z] = sphere (...)
```

Plot a 3-D unit sphere.

The optional input *n* determines the number of faces around the circumference of the sphere. The default value is 20.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

If outputs are requested `sphere` returns three matrices in `meshgrid` format such that `surf (x, y, z)` generates a unit sphere.

Example:

```
[x, y, z] = sphere (40);
surf (3*x, 3*y, 3*z);
axis equal;
title ("sphere of radius 3");
```

See also: [cylinder], page 362, [ellipsoid], page 363, [rectangle], page 331.

```
ellipsoid (xc, yc, zc, xr, yr, zr, n)
```

```
ellipsoid (... , n)
```

```
ellipsoid (hax, ...)
```

```
[x, y, z] = ellipsoid (...)
```

Plot a 3-D ellipsoid.

The inputs *xc*, *yc*, *zc* specify the center of the ellipsoid. The inputs *xr*, *yr*, *zr* specify the semi-major axis lengths.

The optional input *n* determines the number of faces around the circumference of the cylinder. The default value is 20.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

If outputs are requested `ellipsoid` returns three matrices in `meshgrid` format, such that `surf (x, y, z)` generates the ellipsoid.

See also: [cylinder], page 362, [rectangle], page 331, [sphere], page 362.

15.2.3 Plot Annotations

You can add titles, axis labels, legends, and arbitrary text to an existing plot. For example:

```
x = -10:0.1:10;
plot (x, sin (x));
title ("sin(x) for x = -10:0.1:10");
xlabel ("x");
ylabel ("sin (x)");
text (pi, 0.7, "arbitrary text");
legend ("sin (x)");
```

The functions `grid` and `box` may also be used to add grid and border lines to the plot. By default, the grid is off and the border lines are on.

Finally, arrows, text and rectangular or elliptic boxes can be added to highlight parts of a plot using the `annotation` function. Those objects are drawn in an invisible axes, on top of every other axes.

```
title (string)
title (string, prop, val, ...)
title (hax, ...)
h = title (...)
```

Specify the string used as a title for the current axis.

An optional list of *property/value* pairs can be used to change the appearance of the created title text object.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the created text object.

See also: `[xlabel]`, page 366, `[ylabel]`, page 366, `[zlabel]`, page 366, `[text]`, page 366.

```
legend ()
legend (str1, str2, ...)
legend (charmat)
legend ({cellstr})
legend (... , "location", pos)
legend (... , "orientation", orient)
legend (hax, ...)
legend (hobjs, ...)
legend (hax, hobjs, ...)
legend ("option")
legend (... , {cellstr}, property, value, ...)
[hleg, hleg_obj, hplot, labels] = legend (...)
```

Display a legend for the current axes using the specified strings as labels.

Legend entries may be specified as individual character string arguments, a character array, or a cell array of character strings. When label names might be confused with options to `legend`, the labels should be protected by specifying them as a cell array of strings.

If the first argument *hax* is an axes handle, then add a legend to this axes, rather than the current axes returned by `gca`.

Legend labels are associated with the axes' children; The first label is assigned to the first object that was plotted in the axes, the second label to the next object plotted, etc. To label specific data objects, without labeling all objects, provide their graphic handles in the input *hobjs*.

The optional parameter *pos* specifies the location of the legend as follows:

pos	location of the legend
north	center top
south	center bottom
east	right center
west	left center
northeast	right top (default)
northwest	left top
southeast	right bottom

southwest left bottom

outside can be appended to any location string

which will place the legend outside the axes

The optional parameter *orient* determines if the legend elements are placed vertically or horizontally. The allowed values are "vertical" (default) or "horizontal".

The following customizations are available using *option*:

"show"	Show legend on the plot
"hide"	Hide legend on the plot
"toggle"	Toggle between "hide" and "show"
"boxon"	Show a box around legend (default)
"boxoff"	Hide the box around legend
"right"	Place label text to the right of the keys (default)
"left"	Place label text to the left of the keys
"off"	Delete the legend object

The `legend` function creates a graphics object which has various properties that can be manipulated with `get/set`. Alternatively, properties can be set directly when calling `legend` by including *property/value* pairs. If using this calling form, the labels must be specified as a cell array of strings.

The optional output values are

<i>hleg</i>	The graphics handle of the legend object.
<i>hleg_obj</i>	Graphics handles to the text, patch, and line objects which form the legend.
<i>hplot</i>	Graphics handles to the plot objects which were used in making the legend.
<i>labels</i>	A cell array of strings of the labels in the legend.

Implementation Note: The legend label text is either provided in the call to `legend` or is taken from the `DisplayName` property of the graphics objects. Only data objects, such as line, patch, and surface, have this property whereas axes, figures, etc. do not so they are never present in a legend. If no labels or `DisplayName` properties are available, then the label text is simply "data1", "data2", ..., "dataN". No more than 20 data labels will be automatically generated. To label more, call `legend` explicitly and provide all labels.

The legend `FontSize` property is initially set to 90% of the axes `FontSize` to which it is attached. Use `set` to override this if necessary.

A legend is implemented as an additional axes object with the `tag` property set to "legend". Properties of the legend object may be manipulated directly by using `set`.

```
text (x, y, string)
text (x, y, z, string)
text (... , prop, val, ...)
h = text (...)
```

Create a text object with text *string* at position *x*, *y*, (*z*) on the current axes.

Multiple locations can be specified if *x*, *y*, (*z*) are vectors. Multiple strings can be specified with a character matrix or a cell array of strings.

Optional property/value pairs may be used to control the appearance of the text.

The optional return value *h* is a vector of graphics handles to the created text objects.

Programming Note: The full list of properties is documented at [Section 15.3.3.5 \[Text Properties\]](#), page 422.

See also: [\[gtext\]](#), page 392, [\[title\]](#), page 363, [\[xlabel\]](#), page 366, [\[ylabel\]](#), page 366, [\[zlabel\]](#), page 366.

See [Section 15.3.3.5 \[Text Properties\]](#), page 422, for the properties that you can set.

```
xlabel (string)
xlabel (string, property, val, ...)
xlabel (hax, ...)
h = xlabel (...)
```

Specify the string used to label the x-axis of the current axis.

An optional list of *property/value* pairs can be used to change the properties of the created text label.

If the first argument *hax* is an axes handle, then operate on this axes rather than the current axes returned by [gca](#).

The optional return value *h* is a graphics handle to the created text object.

See also: [\[ylabel\]](#), page 366, [\[zlabel\]](#), page 366, [\[datetick\]](#), page 896, [\[title\]](#), page 363, [\[text\]](#), page 366.

```
clabel (c, h)
clabel (c, h, v)
clabel (c, h, "manual")
clabel (c)
clabel (... , prop, val, ...)
h = clabel (...)
```

Add labels to the contours of a contour plot.

The contour levels are specified by the contour matrix *c* which is returned by [contour](#), [contourc](#), [contourf](#), and [contour3](#). Contour labels are rotated to match the local line orientation and centered on the line. The position of labels along the contour line is chosen randomly.

If the argument *h* is a handle to a contour group object, then label this plot rather than the one in the current axes returned by [gca](#).

By default, all contours are labeled. However, the contours to label can be specified by the vector *v*. If the "manual" argument is given then the contours to label can be selected with the mouse.

Additional property/value pairs that are valid properties of text objects can be given and are passed to the underlying text objects. Moreover, the contour group property "LabelSpacing" is available which determines the spacing between labels on a contour to be specified. The default is 144 points, or 2 inches.

The optional return value *h* is a vector of graphics handles to the text objects representing each label. The "userdata" property of the text objects contains the numerical value of the contour label.

An example of the use of `clabel` is

```
[c, h] = contour (peaks (), -4 : 6);
clabel (c, h, -4:2:6, "fontsize", 12);
```

See also: [contour], page 310, [contourf], page 311, [contour3], page 312, [meshc], page 334, [surf], page 336, [text], page 366.

`box`
`box on`
`box off`
`box (hax, ...)`

Control display of the axes border.

The argument may be either "on" or "off". If it is omitted, the current box state is toggled.

If the first argument *hax* is an axes handle, then operate on this axes rather than the current axes returned by `gca`.

See also: [axis], page 323, [grid], page 367.

`grid`
`grid on`
`grid off`
`grid minor`
`grid minor on`
`grid minor off`
`grid (hax, ...)`

Control the display of plot grid lines.

The function state input may be either "on" or "off". If it is omitted, the current grid state is toggled.

When the first argument is "minor" all subsequent commands modify the minor grid rather than the major grid.

If the first argument *hax* is an axes handle, then operate on this axes rather than the current axes returned by `gca`.

To control the grid lines for an individual axes use the `set` function. For example:

```
set (gca, "ygrid", "on");
```

See also: [axis], page 323, [box], page 367.

`colorbar`
`colorbar (..., loc)`

```
colorbar (delete_option)
colorbar (hcb, ...)
colorbar (hax, ...)
colorbar (... , "peer", hax, ...)
colorbar (... , "location", loc, ...)
colorbar (... , prop, val, ...)
h = colorbar (...)
```

Add a colorbar to the current axes.

A colorbar displays the current colormap along with numerical rulings so that the color scale can be interpreted.

The optional input *loc* determines the location of the colorbar. If present, it must be the last argument to `colorbar`. Valid values for *loc* are

"EastOutside"

Place the colorbar outside the plot to the right. This is the default.

"East"

Place the colorbar inside the plot to the right.

"WestOutside"

Place the colorbar outside the plot to the left.

"West"

Place the colorbar inside the plot to the left.

"NorthOutside"

Place the colorbar above the plot.

"North"

Place the colorbar at the top of the plot.

"SouthOutside"

Place the colorbar under the plot.

"South"

Place the colorbar at the bottom of the plot.

To remove a colorbar from a plot use any one of the following keywords for the *delete_option*: "off", "delete", "hide".

If the first argument *hax* is an axes handle, then the colorbar is added to this axes, rather than the current axes returned by `gca`. Alternatively, If the argument "peer" is given, then the following argument is treated as the axes handle in which to add the colorbar. The "peer" calling syntax may be removed in the future and is not recommended.

If the first argument *hcb* is a handle to a colorbar object, then operate on this colorbar directly.

Additional property/value pairs are passed directly to the underlying axes object.

The optional return value *h* is a graphics handle to the created colorbar object.

Implementation Note: A colorbar is created as an additional axes object with the "tag" property set to "colorbar". The created object has the extra property "location" which controls the positioning of the colorbar.

See also: [\[colormap\]](#), page 816.

```

annotation (type)
annotation ("line", x, y)
annotation ("arrow", x, y)
annotation ("doublearrow", x, y)
annotation ("textarrow", x, y)
annotation ("textbox", pos)
annotation ("rectangle", pos)
annotation ("ellipse", pos)
annotation (..., prop, val)
annotation (hf, ...)
h = annotation (...)

```

Draw annotations to emphasize parts of a figure.

You may build a default annotation by specifying only the *type* of the annotation.

Otherwise you can select the type of annotation and then set its position using either *x* and *y* coordinates for line-based annotations or a position vector *pos* for others. In either case, coordinates are interpreted using the "units" property of the annotation object. The default is "normalized", which means the lower left hand corner of the figure has coordinates '[0 0]' and the upper right hand corner '[1 1]'.

If the first argument *hf* is a figure handle, then plot into this figure, rather than the current figure returned by *gcf*.

Further arguments can be provided in the form of *prop/val* pairs to customize the annotation appearance.

The optional return value *h* is a graphics handle to the created annotation object. This can be used with the *set* function to customize an existing annotation object.

All annotation objects share two properties:

- "units": the units in which coordinates are interpreted.
Its value may be one of "centimeters" | "characters" | "inches" | "{normalized}" | "pixels" | "points".
- "position": a four-element vector [x0 y0 width height].
The vector specifies the coordinates (x0,y0) of the origin of the annotation object, its width, and its height. The width and height may be negative, depending on the orientation of the object.

Valid annotation types and their specific properties are described below:

- "line"** Constructs a line. *x* and *y* must be two-element vectors specifying the *x* and *y* coordinates of the two ends of the line.
The line can be customized using "linewidth", "linestyle", and "color" properties the same way as for **line** objects.
- "arrow"** Construct an arrow. The second point in vectors *x* and *y* specifies the arrowhead coordinates.
Besides line properties, the arrowhead can be customized using "headlength", "headwidth", and "headstyle" properties. Supported values for "headstyle" property are: ["diamond" | "ellipse" | "plain" | "rectangle" | "vback1" | "{vback2}" | "vback3"]

"doublearrow"

Construct a double arrow. Vectors *x* and *y* specify the arrowhead coordinates.

The line and the arrowhead can be customized as for arrow annotations, but some property names are duplicated: "head1length"/"head2length", "head1width"/"head2width", etc. The index 1 marks the properties of the arrowhead at the first point in *x* and *y* coordinates.

"textarrow"

Construct an arrow with a text label at the opposite end from the arrowhead.

Use the "string" property to change the text string. The line and the arrowhead can be customized as for arrow annotations, and the text can be customized using the same properties as `text` graphics objects. Note, however, that some text property names are prefixed with "text" to distinguish them from arrow properties: "textbackgroundcolor", "textcolor", "textedgecolor", "textlinewidth", "textmargin", "textrotation".

"textbox"

Construct a box with text inside. *pos* specifies the "position" property of the annotation.

Use the "string" property to change the text string. You may use "backgroundcolor", "edgecolor", "linestyle", and "linewidth" properties to customize the box background color and edge appearance. A limited set of `text` objects properties are also available; Besides "font..." properties, you may also use "horizontalalignment" and "verticalalignment" to position the text inside the box.

Finally, the "fitboxtotext" property controls the actual extent of the box. If "on" (the default) the box limits are fitted to the text extent.

"rectangle"

Construct a rectangle. *pos* specifies the "position" property of the annotation.

You may use "facecolor", "color", "linestyle", and "linewidth" properties to customize the rectangle background color and edge appearance.

"ellipse"

Construct an ellipse. *pos* specifies the "position" property of the annotation.

See "`rectangle`" annotations for customization.

See also: `[xlabel]`, page 366, `[ylabel]`, page 366, `[zlabel]`, page 366, `[title]`, page 363, `[text]`, page 366, `[gtext]`, page 392, `[legend]`, page 364, `[colorbar]`, page 367.

15.2.4 Multiple Plots on One Page

Octave can display more than one plot in a single figure. The simplest way to do this is to use the `subplot` function to divide the plot area into a series of subplot windows that are indexed by an integer. For example,

```
subplot (2, 1, 1)
fplot (@sin, [-10, 10]);
subplot (2, 1, 2)
fplot (@cos, [-10, 10]);
```

creates a figure with two separate axes, one displaying a sine wave and the other a cosine wave. The first call to `subplot` divides the figure into two plotting areas (two rows and one column) and makes the first plot area active. The grid of plot areas created by `subplot` is numbered in row-major order (left to right, top to bottom). After plotting a sine wave, the next call to `subplot` activates the second subplot area, but does not re-partition the figure.

```
subplot (rows, cols, index)
subplot (rows, cols, index, hax)
subplot (rcn)
subplot (hax)
subplot (..., "align")
subplot (..., "replace")
subplot (..., "position", pos)
subplot (..., prop, val, ...)
hax = subplot (...)
```

Set up a plot grid with *rows* by *cols* subwindows and set the current axes for plotting (`gca`) to the location given by *index*.

If an axes handle *hax* is provided after the `(rows, cols, index)` arguments, the corresponding axes is turned into a subplot.

If only one numeric argument is supplied, then it must be a three digit value specifying the number of rows in digit 1, the number of columns in digit 2, and the plot index in digit 3.

The plot index runs row-wise; First, all columns in a row are numbered and then the next row is filled.

For example, a plot with 2x3 grid will have plot indices running as follows:

1	2	3
4	5	6

index may also be a vector. In this case, the new axes will enclose the grid locations specified. The first demo illustrates this:

```
demo ("subplot", 1)
```

The index of the subplot to make active may also be specified by its axes handle, *hax*, returned from a previous `subplot` command.

If the option "align" is given then the plot boxes of the subwindows will align, but this may leave no room for axes tick marks or labels.

If the option "replace" is given then the subplot axes will be reset, rather than just switching the current axes for plotting to the requested subplot.

The "position" property can be used to exactly position the subplot axes within the current figure. The option *pos* is a 4-element vector [x, y, width, height] that determines the location and size of the axes. The values in *pos* are normalized in the range [0,1].

Any property/value pairs are passed directly to the underlying axes object.

If the output *hax* is requested, subplot returns the axes handle for the subplot. This is useful for modifying the properties of a subplot using `set`.

See also: `[axes]`, page 396, `[plot]`, page 296, `[gca]`, page 401, `[set]`, page 402.

15.2.5 Multiple Plot Windows

You can open multiple plot windows using the `figure` function. For example,

```
figure (1);
fplot (@sin, [-10, 10]);
figure (2);
fplot (@cos, [-10, 10]);
```

creates two figures, with the first displaying a sine wave and the second a cosine wave. Figure numbers must be positive integers.

```
figure
figure n
figure (n)
figure (... , "property", value, ...)
h = figure (...)
```

Create a new figure window for plotting.

If no arguments are specified, a new figure with the next available number is created.

If called with an integer *n*, and no such numbered figure exists, then a new figure with the specified number is created. If the figure already exists then it is made visible and becomes the current figure for plotting.

Multiple property-value pairs may be specified for the figure object, but they must appear in pairs.

The optional return value *h* is a graphics handle to the created figure object.

Programming Note: The full list of properties is documented at [Section 15.3.3.2 \[Figure Properties\]](#), page 407.

See also: `[axes]`, page 396, `[gcf]`, page 400, `[shg]`, page 377, `[clf]`, page 376, `[close]`, page 377.

15.2.6 Manipulation of Plot Objects

```
pan
pan on
```

```
pan off
pan xon
pan yon
pan (hfig, option)
```

Control the interactive panning mode of a figure in the GUI.

Given the option "on" or "off", set the interactive pan mode on or off.

With no arguments, toggle the current pan mode on or off.

Given the option "xon" or "yon", enable pan mode for the x or y axis only.

If the first argument *hfig* is a figure, then operate on the given figure rather than the current figure as returned by *gcf*.

See also: [rotate3d], page 373, [zoom], page 373.

```
rotate (h, direction, alpha)
rotate (... , origin)
```

Rotate the plot object *h* through *alpha* degrees around the line with direction *direction* and origin *origin*.

The default value of *origin* is the center of the axes object that is the parent of *h*.

If *h* is a vector of handles, they must all have the same parent axes object.

Graphics objects that may be rotated are lines, surfaces, patches, and images.

```
rotate3d
rotate3d on
rotate3d off
rotate3d (hfig, option)
```

Control the interactive 3-D rotation mode of a figure in the GUI.

Given the option "on" or "off", set the interactive rotate mode on or off.

With no arguments, toggle the current rotate mode on or off.

If the first argument *hfig* is a figure, then operate on the given figure rather than the current figure as returned by *gcf*.

See also: [pan], page 372, [zoom], page 373.

```
zoom
zoom (factor)
zoom on
zoom off
zoom xon
zoom yon
zoom out
zoom reset
zoom (hfig, option)
```

Zoom the current axes object or control the interactive zoom mode of a figure in the GUI.

Given a numeric argument greater than zero, zoom by the given factor. If the zoom factor is greater than one, zoom in on the plot. If the factor is less than one, zoom

out. If the zoom factor is a two- or three-element vector, then the elements specify the zoom factors for the x, y, and z axes respectively.

Given the option "on" or "off", set the interactive zoom mode on or off.

With no arguments, toggle the current zoom mode on or off.

Given the option "xon" or "yon", enable zoom mode for the x or y-axis only.

Given the option "out", zoom to the initial zoom setting.

Given the option "reset", store the current zoom setting so that `zoom out` will return to this zoom level.

If the first argument `hfig` is a figure, then operate on the given figure rather than the current figure as returned by `gcf`.

See also: [\[pan\]](#), page 372, [\[rotate3d\]](#), page 373.

15.2.7 Manipulation of Plot Windows

By default, Octave refreshes the plot window when a prompt is printed, or when waiting for input. The `drawnow` function is used to cause a plot window to be updated.

```
drawnow ()
drawnow ("expose")
drawnow (term, file, debug_file)
```

Update figure windows and their children.

The event queue is flushed and any callbacks generated are executed.

With the optional argument "expose", only graphic objects are updated and no other events or callbacks are processed.

The third calling form of `drawnow` is for debugging and is undocumented.

See also: [\[refresh\]](#), page 374.

Only figures that are modified will be updated. The `refresh` function can also be used to cause an update of the current figure, even if it is not modified.

```
refresh ()
refresh (h)
```

Refresh a figure, forcing it to be redrawn.

When called without an argument the current figure is redrawn. Otherwise, the figure with graphic handle `h` is redrawn.

See also: [\[drawnow\]](#), page 374.

Normally, high-level plot functions like `plot` or `mesh` call `newplot` to determine whether the state of the target axes should be initialized (the default) or if subsequent plots should be drawn on top of previous ones. To have two plots drawn over one another, use the `hold` function or manually change the axes [\[nextplot\]](#), page 415, property. For example,

```
hold on;
x = -10:0.1:10;
plot (x, sin (x));
plot (x, cos (x));
hold off;
```

displays sine and cosine waves on the same axes. If the hold state is off, consecutive plotting commands like this will only display the last plot.

```
newplot ()
newplot (hfig)
newplot (hax)
hax = newplot (...)
```

Prepare graphics engine to produce a new plot.

This function is called at the beginning of all high-level plotting functions. It is not normally required in user programs. `newplot` queries the "NextPlot" field of the current figure and axes to determine what to do.

Figure NextPlot

"new"

Action

Create a new figure and make it the current figure.

"add" (default)

Add new graphic objects to the current figure.

"replacechildren"

Delete child objects whose HandleVisibility is set to "on". Set NextPlot property to "add". This typically clears a figure, but leaves in place hidden objects such as menubars. This is equivalent to `clf`.

"replace"

Delete all child objects of the figure and reset all figure properties to their defaults. However, the following four properties are not reset: Position, Units, PaperPosition, PaperUnits. This is equivalent to `clf reset`.

Axes NextPlot

"add"

Action

Add new graphic objects to the current axes. This is equivalent to `hold on`.

"replacechildren"

Delete child objects whose HandleVisibility is set to "on", but leave axes properties unmodified. This typically clears a plot, but preserves special settings such as log scaling for axes. This is equivalent to `cla`.

"replace" (default)

Delete all child objects of the axes and reset all axes properties to their defaults. However, the following properties are not reset: Position, Units. This is equivalent to `cla reset`.

If the optional input `hfig` or `hax` is given then prepare the specified figure or axes rather than the current figure and axes.

The optional return value `hax` is a graphics handle to the created axes object (not figure).

Caution: Calling `newplot` may change the current figure and current axes.

```
hold
hold on
```

```
hold off
hold (hax, ...)
```

Toggle or set the "hold" state of the plotting engine which determines whether new graphic objects are added to the plot or replace the existing objects.

hold on Retain plot data and settings so that subsequent plot commands are displayed on a single graph. Line color and line style are advanced for each new plot added.

hold all (deprecated)

Equivalent to **hold on**.

hold off Restore default graphics settings which clear the graph and reset axes properties before each new plot command. (default).

hold Toggle the current hold state.

When given the additional argument *hax*, the hold state is modified for this axes rather than the current axes returned by **gca**.

To query the current hold state use the **ishold** function.

See also: [\[ishold\]](#), page 376, [\[cla\]](#), page 377, [\[clf\]](#), page 376, [\[newplot\]](#), page 375.

```
ishold
ishold (hax)
ishold (hfig)
```

Return true if the next plot will be added to the current plot, or false if the plot device will be cleared before drawing the next plot.

If the first argument is an axes handle *hax* or figure handle *hfig* then operate on this plot rather than the current one.

See also: [\[hold\]](#), page 375, [\[newplot\]](#), page 375.

To clear the current figure, call the **clf** function. To clear the current axis, call the **cla** function. To bring the current figure to the top of the window stack, call the **shg** function. To delete a graphics object, call **delete** on its index. To close the figure window, call the **close** function.

```
clf
clf reset
clf (hfig)
clf (hfig, "reset")
h = clf (...)
```

Clear the current figure window.

clf operates by deleting child graphics objects with visible handles (**HandleVisibility** = "on").

If the optional argument "**reset**" is specified, delete all child objects including those with hidden handles and reset all figure properties to their defaults. However, the following properties are not reset: Position, Units, PaperPosition, PaperUnits.

If the first argument *hfig* is a figure handle, then operate on this figure rather than the current figure returned by **gcf**.

The optional return value *h* is the graphics handle of the figure window that was cleared.

See also: [\[cla\]](#), page 377, [\[close\]](#), page 377, [\[delete\]](#), page 377, [\[reset\]](#), page 456.

```
cla
cla reset
cla (hax)
cla (hax, "reset")
```

Clear the current or specified (*hax*) axes object.

`cla` operates by deleting child graphic objects with visible handles (`HandleVisibility = "on"`). This typically clears the axes of any visual objects, but leaves in place axes limits, tick marks and labels, camera view, etc. In addition, the automatic coloring and styling of lines is reset by changing the axes properties `ColorOrderIndex`, `LineStyleOrderIndex` to 1.

If the optional argument "`reset`" is specified, delete all child objects, including those with hidden handles, and reset all axes properties to their defaults. However, the following properties are not reset: `Position`, `Units`.

If the first argument *hax* is an axes handle, then operate on this axes rather than the current axes returned by `gca`.

See also: [\[clf\]](#), page 376, [\[delete\]](#), page 377, [\[reset\]](#), page 456.

`shg`

Show the graph window.

This function makes the current figure visible, and places it on top of all other plot windows.

Programming Note: `shg` is equivalent to `figure (gcf)` assuming that a current figure exists.

See also: [\[figure\]](#), page 372, [\[drawnow\]](#), page 374, [\[gcf\]](#), page 400.

```
delete (file)
delete (file1, file2, ...)
delete (handle)
```

Delete the named file or graphics handle.

file may contain globbing patterns such as '`*`'. Multiple files to be deleted may be specified in the same function call.

handle may be a scalar or vector of graphic handles to delete.

Programming Note: Deleting graphics objects is the proper way to remove features from a plot without clearing the entire figure.

See also: [\[clf\]](#), page 376, [\[cla\]](#), page 377, [\[unlink\]](#), page 898, [\[rmdir\]](#), page 899.

```
close
close (h)
close figname
close all
close all hidden
```

close all force

Close figure window(s).

When called with no arguments, close the current figure. This is equivalent to **close (gcf)**. If the input *h* is a graphic handle, or vector of graphics handles, then close each figure in *h*. The figure to close may also be specified by name *figname* which is matched against the "Name" property of all figures.

If the argument "all" is given then all figures with visible handles (HandleVisibility = "on") are closed.

If the additional argument "hidden" is given then all figures, including hidden ones, are closed.

If the additional argument "force" is given then figures are closed even when "closerequestfcn" has been altered to prevent closing the window.

Implementation Note: **close** operates by making the handle *h* the current figure, and then calling the function specified by the "closerequestfcn" property of the figure. By default, the function **closereq** is used. It is possible that the function invoked will delay or abort removing the figure. To remove a figure without executing any callback functions use **delete**. When writing a callback function to close a window do not use **close** to avoid recursion.

See also: [closereq], page 378, [delete], page 377.

closereq ()

Close the current figure and delete all graphics objects associated with it.

By default, the "closerequestfcn" property of a new plot figure points to this function.

See also: [close], page 377, [delete], page 377.

15.2.8 Use of the interpreter Property

All text objects—such as titles, labels, legends, and text—include the property "interpreter" that determines the manner in which special control sequences in the text are rendered.

The interpreter property can take three values: "none", "tex", "latex". If the interpreter is set to "none" then no special rendering occurs—the displayed text is a verbatim copy of the specified text. Currently, the "latex" interpreter is not implemented for on-screen display and is equivalent to "none". Note that Octave does not parse or validate the text strings when in "latex" mode—it is the responsibility of the programmer to generate valid strings which may include wrapping sections that should appear in Math mode with '\$' characters.

The "tex" option implements a subset of TeX functionality when rendering text. This allows the insertion of special glyphs such as Greek characters or mathematical symbols. Special characters are inserted by using a backslash (\) character followed by a code, as shown in [Table 15.1](#).

Besides special glyphs, the formatting of the text can be changed within the string by using the codes

\bf

Bold font

\it	Italic font
\sl	Oblique Font
\rm	Normal font

These codes may be used in conjunction with the { and } characters to limit the change to a part of the string. For example,

```
xlabel ('{\bf H} = a {\bf V}')
```

where the character 'a' will not appear in bold font. Note that to avoid having Octave interpret the backslash character in the strings, the strings themselves should be in single quotes.

It is also possible to change the fontname and size within the text

\fontname{fontname}	Specify the font to use
\fontsize{size}	Specify the size of the font to use

The color of the text may also be changed inline using either a string (e.g., "red") or numerically with a Red-Green-Blue (RGB) specification (e.g., [1 0 0], also red).

\color{color}	Specify the color as a string
\color[rgb]{R G B}	Specify the color numerically

Finally, superscripting and subscripting can be controlled with the '^' and '_' characters. If the '^' or '_' is followed by a { character, then all of the block surrounded by the { } pair is superscripted or subscripted. Without the { } pair, only the character immediately following the '^' or '_' is changed.

Greek Lowercase Letters

Code	Sym	Code	Sym	Code	Sym
\alpha	α	\beta	β	\gamma	γ
\delta	δ	\epsilon	ϵ	\zeta	ζ
\eta	η	\theta	θ	\vartheta	ϑ
\iota	ι	\kappa	κ	\lambda	λ
\mu	μ	\nu	ν	\xi	ξ
\o	\o	\pi	π	\varpi	ϖ
\rho	ρ	\sigma	σ	\varsigma	ς
\tau	τ	\upsilon	υ	\phi	ϕ
\chi	χ	\psi	ψ	\omega	ω

Greek Uppercase Letters

Code	Sym	Code	Sym	Code	Sym
\Gamma	Γ	\Delta	Δ	\Theta	Θ
\Lambda	Λ	\Xi	Ξ	\Pi	Π
\Sigma	Σ	\Upsilon	Υ	\Phi	Φ
\Psi	Ψ	\Omega	Ω		

Misc Symbols Type Ord

Code	Sym	Code	Sym	Code	Sym
\aleph	\aleph	\wp	\wp	\Re	\Re
\Im	\Im	\partial	∂	\infty	∞
\prime	\prime	\nabla	∇	\surd	\surd
\angle	\angle	\forall	\forall	\exists	\exists
\neg	\neg	\clubsuit	\clubsuit	\diamondsuit	\diamondsuit
\heartsuit	\heartsuit	\spadesuit	\spadesuit		

“Large” Operators

Code	Sym	Code	Sym	Code	Sym
\int	\int				

Binary operators

Code	Sym	Code	Sym	Code	Sym
\pm	\pm	\cdot	\cdot	\times	\times
\ast	\ast	\circ	\circ	\bullet	\bullet
\div	\div	\cap	\cap	\cup	\cup
\vee	\vee	\wedge	\wedge	\oplus	\oplus
\otimes	\otimes	\oslash	\oslash		

Table 15.1: Available special characters in TeX mode

Relations

Code	Sym	Code	Sym	Code	Sym
\leq	\leq	\subset	\subset	\subseteq	\subseteq
\in	\in	\geq	\geq	\supset	\supset
\supseteqq	\supseteq	\ni	\ni	\mid	\mid
\equiv	\equiv	\sim	\sim	\approx	\approx
\cong	\cong	\propto	\propto	\perp	\perp

Arrows

Code	Sym	Code	Sym	Code	Sym
\leftarrow	\leftarrow	\Leftarrow	\Leftarrow	\rightarrow	\rightarrow
\rightarrow	\rightarrow	\Leftrightarrow	\Leftrightarrow	\uparrow	\uparrow
\downarrow	\downarrow				

Openings and Closings

Code	Sym	Code	Sym	Code	Sym
\lfloor	\lfloor	\langle	\langle	\lceil	\lceil
\rfloor	\rfloor	\rangle	\rangle	\rceil	\rceil

Alternate Names

Code	Sym	Code	Sym	Code	Sym
\neq	\neq				

Other (not in Appendix F Tables)

Code	Sym	Code	Sym	Code	Sym
\ldots	\dots	\emptyset	\emptyset	\copyright	\circledC
\deg	\circ				

Table 15.1: Available special characters in TeX mode (cont.)

15.2.8.1 Degree Symbol

Conformance to both TeX and MATLAB with respect to the \circ symbol is impossible. While TeX translates this symbol to Unicode 2218 (U+2218), MATLAB maps this to Unicode 00B0 (U+00B0) instead. Octave has chosen to follow the TeX specification, but has added the additional symbol \deg which maps to the degree symbol (U+00B0).

15.2.9 Printing and Saving Plots

The print command allows you to send plots to your printer and to save plots in a variety of formats. For example,

```
print -dpssc
```

prints the current figure to a color PostScript printer. And,

```
print -deps foo.eps
```

saves the current figure to an encapsulated PostScript file called `foo.eps`.

The current graphic toolkits produce very similar graphic displays, but differ in their capability to display unusual text and in their ability to print such text. In general, the "`tex`" interpreter (default) is the best all-around performer for both on-screen display and printing. However, for the reproduction of complicated text formulas the "`latex`" interpreter is preferred. The "`latex`" interpreter will not display symbols on-screen, but the printed output will be correct. When printing, use one of the `standalone` options which provide full access to L^AT_EX commands.

A complete example showing the capabilities of text printing using the `-dpdflatexstandalone` option is:

```
x = 0:0.01:3;
hf = figure ();
plot (x, erf (x));
hold on;
plot (x, x, "r");
axis ([0, 3, 0, 1]);
text (0.65, 0.6175, ...
      ['$\displaystyle\leftarrow x = \frac{2}{\sqrt{\pi}}$',...
       '\int_0^x e^{-t^2} dt = 0.6175'],
      "interpreter", "latex");
xlabel ("x");
ylabel ("erf (x)");
title ("erf (x) with text annotation");
print (hf, "plot15_7.pdf", "-dpdflatexstandalone");
system ("pdflatex plot15_7");
open plot15_7.pdf
```

The result of this example can be seen in [Figure 15.7](#)




Figure 15.7: Example of inclusion of text with use of `-dpdflatexstandalone`

```

print ()
print (options)
print (filename, options)
print (hfig, ...)
rgb = print ("-RGBImage", ...)

```

Format a figure for printing and either save it to a file, send it to a printer, or return an RGB image.

filename defines the name of the output file. If the *filename* has no suffix then one is inferred from the specified device and appended to the *filename*. When neither a *filename* nor the "`-RGBImage`" option is present, the output is sent to the printer. The various options and *filename* arguments may be given in any order, except for the figure handle argument *hfig* which must be first if it is present.

Example: Print to a file using PDF and JPEG formats.

```

figure (1);
clf ();
surf (peaks);
print figure1.pdf    # The extension specifies the format
print -djpeg figure1 # Will produce "figure1.jpg" file

```

If the first argument is a handle *hfig* to a figure object then it specifies the figure to print. By default, the current figure returned by `gcf` is printed.

For outputs to paged formats, for example, PostScript and PDF, the page size is specified by the figure's `papersize` property together with the `paperunits` property. The location and size of the plot on the page are specified by the figure's `paperposition` property. The orientation of the page is specified by the figure's `paperorientation` property.

For non-page formats—for example, image formats like JPEG—the width and height of the output are specified by the figure’s `paperposition(3:4)` property values.

The `print` command supports many *options*:

-fh Specify the handle, *h*, of the figure to be printed.

Example: Print figure 1.

```
figure (1);
clf ();
surf (peaks);
figure (2);
print -f1 figure1.pdf
## Equivalent functional form:
print (1, "figure1.pdf")
```

-Pprinter

Set the *printer* name to which the plot is sent if no *filename* is specified.

Example: Print to printer named PS_printer using PostScript format.

```
clf ();
surf (peaks);
print -dpswrite -PPS_printer
```

-RGBImage

Return an M-by-N-by-3 RGB image of the figure. The size of the image depends on the formatting options. This is similar to taking a screen capture of the plot, but formatting options may be changed such as the resolution or monochrome/color.

Example: Get the pixels of a figure image.

```
clf ();
surf (peaks);
rgb = print ("-RGBImage");
```

-opengl

-painters

Specifies whether the `opengl` (pixel-based) or `painters` (vector-based) renderer is used. This is equivalent to changing the figure’s “`Renderer`” property. When the figure `RendererMode` property is “`auto`” Octave will use the “`opengl`” renderer for raster formats (e.g., JPEG) and “`painters`” for vector formats (e.g., PDF). Those options are only supported for the “`qt`” graphics toolkit.

-svgconvert

For OpenGL-based graphic toolkits, this enables a different backend toolchain with enhanced characteristics. The toolchain adds support for printing arbitrary characters and fonts in PDF outputs; it avoids some anti-aliasing artifacts in the rendering of patch and surface objects (particularly for 2-D scenes); and it supports transparency of line, patch, and surface objects.

This option only affects PDF outputs, unless it is combined with **-painters** option, in which case raster outputs are also affected.

Caution: **-svgconvert** may lead to inaccurate rendering of image objects.

-portrait

-landscape

Specify the orientation of the plot for printed output. For non-printed output the aspect ratio of the output corresponds to the plot area defined by the "paperposition" property in the orientation specified. This option is equivalent to changing the figure's "paperorientation" property.

-fillpage

-bestfit When using a page-based format (PDF, PostScript, printer) ignore the "paperposition" property and have the plot occupy the entire page. The option **-fillpage** will stretch the plot to occupy the page with 0.25 inch margins all around. The option **-bestfit** will expand the plot to take up as much room as possible on the page **without** distorting the original aspect ratio of the plot.

-color

-mono Color or monochrome output.

-solid

-dashed Force all lines to be solid or dashed, respectively.

-noui

Don't print uicontrol objects such as pushbuttons which may overlay the plot. This is the default behavior and it is not possible to include uicontrol objects in the output without using an external screen capture tool.

-rNUM

Resolution of bitmaps in dots per inch (DPI). For both metafiles and SVG the default is the screen resolution; for other formats the default is 150 DPI. To specify screen resolution, use "**-r0**".

Example: high resolution raster output.

```
clf ();
surf (peaks (), "facelighting", "gouraud");
light ();
print ("-r600", "lit_peaks.png");
```

-Sxsize,ysize

Plot size in pixels for raster formats including PNG, JPEG, PNG, and (unusually (SVG)). For all vector formats, including PDF, PS, and EPS, the plot size is specified in points. This option is equivalent to changing the width and height of the output by setting the figure property **paperposition(3:4)**. When using the command form of the print function you must quote the **xsize,ysize** option to prevent the Octave interpreter from recognizing the embedded comma (','). For example, by writing "**-S640,480**".

-tight

-loose

Force a tight or loose bounding box for EPS files. The default is tight.

- preview** Add a preview to EPS files. Supported formats are:
 - interchange**
Provide an interchange preview.
 - metafile**
Provide a metafile preview.
 - pict**
Provide a pict preview.
 - tiff**
Provide a TIFF preview.
- append** Append PostScript or PDF output to an existing file of the same type.
- F**
-F:size
-F:size Use *fontname* and/or *fontsize* for all text. *fontname* is ignored for some devices: dxf, fig, hpgl, etc.
- ddevice** The available output format is specified by the option *device*, and is one of the following (devices marked with a "*" are only available with the Gnuplot toolkit):

Vector Formats

- pdf**
pdfcrop Portable Document Format. The **pdfcrop** device removes the default surrounding page.

The OpenGL-based graphics toolkits have limited support for text. Limitations include using only ASCII characters (e.g., no Greek letters) and support for just three base PostScript fonts: Helvetica (the default), Times, or Courier. Any other font will be replaced by Helvetica.

For an enhanced output with complete text support and basic transparency, use the **-svgconvert** option.

- ps**
ps2
psc
psc2 PostScript (level 1 and 2, mono and color). The OpenGL-based graphics toolkits always generate PostScript level 3.0 and have limited support for text.

- eps**
eps2
epsc
epsc2 Encapsulated PostScript (level 1 and 2, mono and color). The OpenGL-based toolkits always generate PostScript level 3.0 and have limited support for text.

`pslatex`
`epslatex`
`pdflatex`
`pslatexstandalone`
`epslatexstandalone`
`pdflatexstandalone`

Generate a \LaTeX file `filename.tex` for the text portions of a plot and a file `filename.(ps|eps|pdf)` for the remaining graphics. The graphics file suffix `.ps|eps|pdf` is determined by the specified device type. The \LaTeX file produced by the ‘`standalone`’ option can be processed directly by \LaTeX . The file generated without the ‘`standalone`’ option is intended to be included from another \LaTeX document. In either case, the \LaTeX file contains an `\includegraphics` command so that the generated graphics file is automatically included when the \LaTeX file is processed. The text that is written to the \LaTeX file contains the strings **exactly** as they were specified in the plot. If any special characters of the \TeX mode interpreter were used, the file must be edited before \LaTeX processing. Specifically, the special characters must be enclosed with dollar signs (`$... $`), and other characters that are recognized by \LaTeX may also need editing (e.g., braces). The ‘`pdflatex`’ device, and any of the ‘`standalone`’ formats, are not available with the Gnuplot toolkit.

`epscairo*`
`pdfcairo*`
`epscairolatex*`
`pdfcairolatex*`
`epscairolatexstandalone*`
`pdfcairolatexstandalone*`

Generate output with Cairo renderer. The devices ‘`epscairo`’ and ‘`pdfcairo`’ are synonymous with the ‘`epsc`’ device. The \LaTeX variants generate a \LaTeX file, `filename.tex`, for the text portions of a plot, and an image file, `filename.(eps|pdf)`, for the graph portion of the plot. The ‘`standalone`’ variants behave as described for ‘`epslatexstandalone`’ above.

<code>svg</code>	Scalable Vector Graphics
<code>canvas*</code>	Javascript-based drawing on an HTML5 canvas viewable in a web browser.
<code>cdr*</code>	
<code>corel*</code>	CorelDraw
<code>cgm*</code>	Computer Graphics Metafile, Version 1, ANSI X3.122-1986
<code>dxf</code>	AutoCAD

emf	
meta	Microsoft Enhanced Metafile
fig	XFig. For the Gnuplot graphics toolkit, the additional options <code>-textspecial</code> or <code>-textnormal</code> can be used to control whether the special flag should be set for the text in the figure. (default is <code>-textnormal</code>)
hpgl	HP plotter language
ill	
aifm	Adobe Illustrator (obsolete for Gnuplot versions > 4.2)
latex*	
eepic*	L <small>A</small> T <small>E</small> X picture environment and extended picture environment.
mf*	Metafont
tikz	
tikzstandalone*	Generate a L <small>A</small> T <small>E</small> X file using PGF/TikZ format. The OpenGL-based toolkits create a PGF file while Gnuplot creates a TikZ file. The ‘tikzstandalone’ device produces a L <small>A</small> T <small>E</small> X document which includes the TikZ file.

Raster Formats

png	Portable Network Graphics
jpg	
jpeg	JPEG image
tif	
tiff	
tiffn	TIFF image with LZW compression (tif, tiff) or uncompressed (tiffn).
gif	GIF image
pbm	PBMplus
dumb*	ASCII art

If the device is omitted, it is inferred from the file extension, or if there is no filename then it is sent to the printer as PostScript.

-dghostscript_device

Additional devices are supported by Ghostscript. Some examples are:

ljet2p	HP LaserJet IIP
pcx24b	24-bit color PCX file format
ppm	Portable Pixel Map file format

For a complete list of available formats and devices type `system ("gs -h")`.

When Ghostscript output is sent to a printer the size is determined by the figure's "papersize" property. When the output is sent to a file the size is determined by the plot box defined by the figure's "paperposition" property.

-Gghostscript_command

Specify the command for calling Ghostscript. For Unix the default is "gs" and for Windows it is "gswin32c".

-TextAlphaBits=n

-GraphicsAlphaBits=n

Octave is able to produce output for various printers, bitmaps, and vector formats by using Ghostscript. For bitmap and printer output anti-aliasing is applied using Ghostscript's TextAlphaBits and GraphicsAlphaBits options. The default number of bits are 4 and 1 respectively. Allowed values for N are 1, 2, or 4.

See also: [saveas], page 389, [hgsave], page 390, [getframe], page 815, [orient], page 389, [figure], page 372.

saveas (*h, filename*)

saveas (*h, filename, fmt*)

Save graphic object *h* to the file *filename* in graphic format *fmt*.

All device formats accepted by `print` may be used. Common formats are:

ps PostScript

eps Encapsulated PostScript

pdf Portable Document Format

jpg JPEG Image

png Portable Network Graphics image

emf Enhanced MetaFile

tif TIFF Image, compressed

If *fmt* is omitted it is extracted from the extension of *filename*. The default format when there is no extension is "pdf".

```
clf ();
surf (peaks);
saveas (1, "figure1.png");
```

See also: [print], page 383, [hgsave], page 390, [orient], page 389.

orient (*orientation*)

orient (*hfig, orientation*)

orientation = orient ()

orientation = orient (*hfig*)

Query or set the print orientation for figure *hfig*.

Valid values for `orientation` are "portrait", "landscape", and "tall".

The "landscape" option changes the orientation so the plot width is larger than the plot height. The "paperposition" is also modified so that the plot fills the page, while leaving a 0.25 inch border.

The "tall" option sets the orientation to "portrait" and fills the page with the plot, while leaving a 0.25 inch border.

The "portrait" option (default) changes the orientation so the plot height is larger than the plot width. It also restores the default "paperposition" property.

When called with no arguments, return the current print orientation.

If the argument `hfig` is omitted, then operate on the current figure returned by `gcf`.

See also: [\[print\]](#), page 383, [\[saveas\]](#), page 389.

`print` and `saveas` are used when work on a plot has finished and the output must be in a publication-ready format. During intermediate stages it is often better to save the graphics object and all of its associated information so that changes—to colors, axis limits, marker styles, etc.—can be made easily from within Octave. The `hgsave/hgload` commands can be used to save and re-create a graphics object.

```
hgsave (filename)
hgsave (h, filename)
hgsave (h, filename, fmt)
```

Save the graphics handle(s) `h` to the file `filename` in the format `fmt`.

If unspecified, `h` is the current figure as returned by `gcf`.

When `filename` does not have an extension the default filename extension `.ofig` will be appended.

If present, `fmt` must be one of the following:

- `-binary`, `-float-binary`
- `-hdf5`, `-float-hdf5`
- `-V7`, `-v7`, `-7`, `-mat7-binary`
- `-V6`, `-v6`, `-6`, `-mat6-binary`
- `-text`
- `-zip`, `-z`

The default format is `-binary` to minimize storage.

Programming Note: When producing graphics for final publication use `print` or `saveas`. When it is important to be able to continue to edit a figure as an Octave object, use `hgsave/hgload`.

See also: [\[hgload\]](#), page 390, [\[hdl2struct\]](#), page 404, [\[savefig\]](#), page 391, [\[saveas\]](#), page 389, [\[print\]](#), page 383.

```
h = hgload (filename)
[h, old_prop] = hgload (filename, prop_struct)
```

Load the graphics objects in `filename` into a vector of graphics handles `h`.

If *filename* has no extension, Octave will try to find the file with and without the default extension `.ofig`.

If provided, the elements of structure `prop_struct` will be used to override the properties of top-level objects stored in *filename*, and the saved values from *filename* will be stored in `old_prop`. `old_prop` is a cell array matching the size of *h*; each cell contains a structure of the existing property names and values before being overridden.

See also: [\[openfig\]](#), page 391, [\[hgsave\]](#), page 390, [\[struct2hdl\]](#), page 404.

```
openfig
openfig (filename)
openfig (... , copies)
openfig (... , visibility)
h = openfig (...)
```

Read saved figure window(s) from *filename* and return graphics handle(s) *h*.

By default, *filename* is `"Untitled.fig"`. If a full path is not specified, the file opened will be the first one encountered in the load path. If *filename* is not found and does not have an extension, a search will take place for the first file in the load path with extension `".fig"` or `".ofig"`, in that order.

copies is an optional input indicating whether a new figure should be created ("new") or whether an existing figure may be reused ("reuse"). An existing figure may be reused if the `"FileName"` property matches the specified input *filename*. When a figure is reused it becomes the active figure and is shown on top of other figures. If the figure was offscreen, it is re-positioned to be onscreen. The default value for *copies* is "new".

visibility is an optional input indicating whether to show the figure ("visible") or not ("invisible"). When *visibility* is specified as an input to `openfig` it overrides the visibility setting stored in *filename*.

See also: [\[open\]](#), page 914, [\[hgload\]](#), page 390, [\[savefig\]](#), page 391, [\[struct2hdl\]](#), page 404.

```
savefig ()
savefig (h)
savefig (filename)
savefig (h, filename)
savefig (h, filename, "compact")
```

Save figure windows specified by graphics handle(s) *h* to file *filename*.

If unspecified, *h* is the current figure returned by `gcf`.

If unspecified, *filename* is set to `"Untitled.fig"`. If *filename* does not have an extension then the default extension `".fig"` will be added.

If the optional third input "compact" is present then the data will be compressed to save more space.

See also: [\[hgsave\]](#), page 390, [\[hdl2struct\]](#), page 404, [\[openfig\]](#), page 391.

15.2.10 Interacting with Plots

The user can select points on a plot with the `ginput` function or select the position at which to place text on the plot with the `gtext` function using the mouse.

```
[x, y, buttons] = ginput (n)
[x, y, buttons] = ginput ()
```

Return the position and type of mouse button clicks and/or key strokes in the current figure window.

If *n* is defined, then capture *n* events before returning. When *n* is not defined `ginput` will loop until the return key RET is pressed.

The return values *x*, *y* are the coordinates where the mouse was clicked in the units of the current axes. The return value *button* is 1, 2, or 3 for the left, middle, or right button. If a key is pressed the ASCII value is returned in *button*.

Implementation Note: `ginput` is intended for 2-D plots. For 3-D plots see the `currentpoint` property of the current axes which can be transformed with knowledge of the current `view` into data units.

See also: [\[gtext\]](#), page 392, [\[waitforbuttonpress\]](#), page 392.

```
waitForButtonPress ()
b = waitForButtonPress ()
```

Wait for mouse click or key press over the current figure window.

The return value of *b* is 0 if a mouse button was pressed or 1 if a key was pressed.

See also: [\[waitfor\]](#), page 881, [\[ginput\]](#), page 392, [\[kbhit\]](#), page 257.

```
gtext (s)
gtext ({s1, s2, ...})
gtext ({s1; s2; ...})
gtext (..., prop, val, ...)
h = gtext (...)
```

Place text on the current figure using the mouse.

The text is defined by the string *s*. If *s* is a cell string organized as a row vector then each string of the cell array is written to a separate line. If *s* is organized as a column vector then one string element of the cell array is placed for every mouse click.

Optional property/value pairs are passed directly to the underlying text objects.

The optional return value *h* is a graphics handle to the created text object(s).

See also: [\[ginput\]](#), page 392, [\[text\]](#), page 366.

More sophisticated user interaction mechanisms can be obtained using the `ui*` family of functions, see [Section 35.3 \[UI Elements\]](#), page 869.

15.2.11 Test Plotting Functions

The functions `sombrero` and `peaks` provide a way to check that plotting is working. Typing either `sombrero` or `peaks` at the Octave prompt should display a three-dimensional plot.

```
sombrero ()
sombrero (n)
z = sombrero (...)
[x, y, z] = sombrero (...)
```

Plot the familiar 3-D sombrero function.

The function plotted is

$$z = \frac{\sin(\sqrt{(x^2 + y^2)})}{\sqrt{(x^2 + y^2)}}$$

Called without a return argument, `sombrero` plots the surface of the above function over the meshgrid [-8,8] using `surf`.

If n is a scalar the plot is made with n grid lines. The default value for n is 41.

When called with output arguments, return the data for the function evaluated over the meshgrid. This can subsequently be plotted with `surf (x, y, z)`.

See also: [\[peaks\]](#), page 393, [\[meshgrid\]](#), page 347, [\[mesh\]](#), page 333, [\[surf\]](#), page 335.

```
peaks ()
peaks (n)
peaks (x, y)
z = peaks (...)
[x, y, z] = peaks (...)
```

Plot a function with lots of local maxima and minima.

The function has the form

$$f(x, y) = 3(1 - x)^2 e^{(-x^2 - (y+1)^2)} - 10 \left(\frac{x}{5} - x^3 - y^5 \right) - \frac{1}{3} e^{(-(x+1)^2 - y^2)}$$

Called without a return argument, `peaks` plots the surface of the above function using `surf`.

If n is a scalar, `peaks` plots the value of the above function on an n -by- n mesh over the range [-3,3]. The default value for n is 49.

If x and y are specified then the function value is calculated over the specified grid of vertices.

When called with output arguments, return the data for the function evaluated over the meshgrid. This can subsequently be plotted with `surf (x, y, z)`.

See also: [\[sombrero\]](#), page 392, [\[meshgrid\]](#), page 347, [\[mesh\]](#), page 333, [\[surf\]](#), page 335.

15.3 Graphics Data Structures

15.3.1 Introduction to Graphics Structures

The graphics functions use pointers, which are of class `graphics_handle`, in order to address the data structures which control visual display. A graphics handle may point to any one of a number of different base object types and these objects are the graphics data structures

themselves. The primitive graphic object types are: `figure`, `axes`, `line`, `text`, `patch`, `surface`, `text`, `image`, and `light`.

Each of these objects has a function by the same name, and, each of these functions returns a graphics handle pointing to an object of the corresponding type. In addition there are several functions which operate on properties of the graphics objects and which also return handles: the functions `plot` and `plot3` return a handle pointing to an object of type `line`, the function `subplot` returns a handle pointing to an object of type `axes`, the function `fill` returns a handle pointing to an object of type `patch`, the functions `area`, `bar`, `barh`, `contour`, `contourf`, `contour3`, `surf`, `mesh`, `surf3`, `meshc`, `errorbar`, `quiver`, `quiver3`, `scatter`, `scatter3`, `stair`, `stem`, `stem3` each return a handle to a complex data structure as documented in [Data Sources], page 465.

The graphics objects are arranged in a hierarchy:

1. The root object is returned by `groot` (historically, equivalent to the handle 0). In other words, `get (groot)` returns the properties of the root object.
2. Below the root are `figure` objects.
3. Below the `figure` objects are `axes` or `hggroup` objects.
4. Below the `axes` objects are `line`, `text`, `patch`, `surface`, `image`, and `light` objects.

Graphics handles may be distinguished from function handles (see Section 11.11.1 [Function Handles], page 214) by means of the function `ishghandle`. `ishghandle` returns true if its argument is a handle of a graphics object. In addition, a figure or axes object may be tested using `isfigure` or `isaxes` respectively. To test for a specific type of graphics handle, such as a patch or line object, use `isgraphics`. The more specific test functions return true only if the argument is both a graphics handle and of the correct type (figure, axes, specified object type).

The `get` and `set` commands are used to obtain and set the values of properties of graphics objects. In addition, the `get` command may be used to obtain property names.

For example, the property "type" of the graphics object pointed to by the graphics handle `h` may be displayed by:

```
get (h, "type")
```

The properties and their current values may be obtained in the form of a structure using `s = get (h)`, where `h` is the handle of a graphics object. If only the names of the properties and the allowed values (for radio properties only) are wanted, one may use `set (h)`.

Thus, for example:

```
h = figure ();
get (h, "type")
⇒ ans = figure
set (h)
⇒
    alphamap:
    beingdeleted: [ {off} | on ]
    busyaction: [ cancel | {queue} ]
    buttondownfcn:
    clipping: [ off | {on} ]
    closerequestfcn:
```

```

color:
colormap:
createfcn:
currentaxes:
deletefcn:
dockcontrols: [ {off} | on ]
filename:
graphicssmoothing: [ off | {on} ]
handlevisibility: [ callback | off | {on} ]
...

```

The uses of `get` and `set` are further explained in [\[get\]](#), page 402, [\[set\]](#), page 402.

`res = isprop (obj, "prop")`

Return true if *prop* is a property of the object *obj*.

obj may also be an array of objects in which case *res* will be a logical array indicating whether each handle has the property *prop*.

For plotting, *obj* is a handle to a graphics object. Otherwise, *obj* should be an instance of a class.

See also: [\[get\]](#), page 402, [\[set\]](#), page 402, [\[ismethod\]](#), page 837, [\[isobject\]](#), page 836.

15.3.2 Graphics Objects

The hierarchy of graphics objects was explained above. See [Section 15.3.1 \[Introduction to Graphics Structures\]](#), page 393. Here the specific objects are described, and the properties contained in these objects are discussed. Keep in mind that graphics objects are always referenced by *handle*.

root figure The top level of the hierarchy and the parent of all figure objects. Use `groot` to obtain the handle of the root graphics object.

figure A figure window.

axes A set of axes. This object is a child of a figure object and may be a parent of line, text, image, patch, surface, or light objects.

line A line in two or three dimensions.

text Text annotations.

image A bitmap image.

patch A filled polygon, currently limited to two dimensions.

surface A three-dimensional surface.

light A light object used for lighting effects on patches and surfaces.

15.3.2.1 Creating Graphics Objects

You can create any graphics object primitive by calling the function of the same name as the object; In other words, `figure`, `axes`, `line`, `text`, `image`, `patch`, `surface`, and `light` functions. These fundamental graphic objects automatically become children of the current

axes object as if `hold on` was in place. Separately, axes will automatically become children of the current figure object and figures will become children of the root object.

If this auto-joining feature is not desired then it is important to call `newplot` first to prepare a new figure and axes for plotting. Alternatively, the easier way is to call a high-level graphics routine which will both create the plot and then populate it with low-level graphics objects. Instead of calling `line`, use `plot`. Or use `surf` instead of `surface`. Or use `fill` instead of `patch`.

```
axes ()
axes (property, value, ...)
axes (hpar, property, value, ...)
axes (hax)
h = axes (...)
```

Create a Cartesian axes object and return a handle to it, or set the current axes to `hax`.

Called without any arguments, or with *property/value* pairs, construct a new axes. The optional argument *hpar* is a graphics handle specifying the parent for the new axes and may be a figure, uipanel, or uitab.

Called with a single axes handle argument `hax`, the function makes `hax` the current axes (as returned by `gca`). It also makes the figure which contains `hax` the current figure (as returned by `gcf`). Finally, it restacks the parent object's `children` property so that the axes `hax` appears before all other axes handles in the list. This causes `hax` to be displayed on top of any other axes objects (Z-order stacking). In addition it restacks any legend or colorbar objects associated with `hax` so that they are also visible.

Programming Note: The full list of properties is documented at [Section 15.3.3.3 \[Axes Properties\]](#), page 412.

See also: `[gca]`, page 401, `[set]`, page 402, `[get]`, page 402.

```
line ()
line (x, y)
line (x, y, z)
line ("xdata", x, "ydata", y)
line ("xdata", x, "ydata", y, "zdata", z)
line (..., property, value)
line (hax, ...)
h = line (...)
```

Create a line object from *x* and *y* (and possibly *z*) and insert it in the current axes.

In the standard calling form the data *x*, *y*, and *z* may be scalars, vectors, or matrices. In the case of matrix inputs, `line` will attempt to orient scalars and vectors so the results can be plotted. This requires that one of the dimensions of the vector match either the number of rows or the number of columns of the matrix.

In the low-level calling form (50% higher performance) where the data is specified by name (`line ("xdata", x, ...)`) the data must be vectors. If no data is specified (`line ()`) then `x == y = [0, 1]`.

Multiple property-value pairs may be specified for the line object, but they must appear in pairs.

If called with only *property/value* pairs then any unspecified properties use their default values as specified on the root object.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle (or vector of handles) to the line objects created.

Programming Note: The full list of properties is documented at [Section 15.3.3.4 \[Line Properties\], page 419](#).

The function `line` differs from `plot` in that line objects are inserted in to the current axes without first clearing the plot.

See also: [\[image\]](#), page 812, [\[patch\]](#), page 397, [\[rectangle\]](#), page 331, [\[surface\]](#), page 398, [\[text\]](#), page 366.

```
patch ()
patch (x, y, c)
patch (x, y, z, c)
patch ("Faces", faces, "Vertices", verts, ...)
patch (..., prop, val, ...)
patch (..., propstruct, ...)
patch (hax, ...)
h = patch (...)
```

Create patch object in the current axes with vertices at locations (*x*, *y*) and of color *c*.

If the vertices are matrices of size MxN then each polygon patch has M vertices and a total of N polygons will be created. If some polygons do not have M vertices use NaN to represent "no vertex". If the *z* input is present then 3-D patches will be created.

The color argument *c* can take many forms. To create polygons which all share a single color use a string value (e.g., "r" for red), a scalar value which is scaled by `caxis` and indexed into the current colormap, or a 3-element RGB vector with the precise TrueColor.

If *c* is a vector of length N then the *i*th polygon will have a color determined by scaling entry *c*(*i*) according to `caxis` and then indexing into the current colormap. More complicated coloring situations require directly manipulating patch property/value pairs.

Instead of specifying polygons by matrices *x* and *y*, it is possible to present a unique list of vertices and then a list of polygon faces created from those vertices. In this case the "Vertices" matrix will be an Nx2 (2-D patch) or Nx3 (3-D patch). The MxN "Faces" matrix describes M polygons having N vertices—each row describes a single polygon and each column entry is an index into the "Vertices" matrix to identify a vertex. The patch object can be created by directly passing the property/value pairs "Vertices"/verts, "Faces"/faces as inputs.

Instead of using property/value pairs, any property can be set by passing a structure `propstruct` with the respective field names.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the created patch object.

Programming Note: The full list of properties is documented at [Section 15.3.3.7 \[Patch Properties\], page 427](#). Useful patch properties include: "cdata", "edgecolor", "facecolor", "faces", and "facevertexcdata".

See also: `[fill]`, page 322, `[get]`, page 402, `[set]`, page 402.

```
surface (x, y, z, c)
surface (x, y, z)
surface (z, c)
surface (z)
surface (..., prop, val, ...)
surface (hax, ...)
h = surface (...)
```

Create a surface graphic object given matrices *x* and *y* from `meshgrid` and a matrix of values *z* corresponding to the *x* and *y* coordinates of the surface.

If *x* and *y* are vectors, then a typical vertex is $(x(j), y(i), z(i,j))$. Thus, columns of *z* correspond to different *x* values and rows of *z* correspond to different *y* values. If only a single input *z* is given then *x* is taken to be `1:columns(z)` and *y* is `1:rows(z)`.

Any property/value input pairs are assigned to the surface object.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the created surface object.

Programming Note: The full list of properties is documented at [Section 15.3.3.8 \[Surface Properties\], page 430](#).

See also: `[surf]`, page 335, `[mesh]`, page 333, `[patch]`, page 397, `[line]`, page 396.

```
light ()
light (..., "prop", val, ...)
light (hax, ...)
h = light (...)
```

Create a light object in the current axes or for axes *hax*.

When a light object is present in an axes object, and the properties "EdgeLighting" or "FaceLighting" of a patch or `surface` object are set to a value other than "none", these objects are drawn with lighting effects. Supported values for Lighting properties are "none" (no lighting effects), "flat" (faceted look of the objects), and "gouraud" (linear interpolation of the lighting effects between the vertices). If the lighting mode is set to "flat", the "FaceNormals" property is used for lighting. For "gouraud", the "VertexNormals" property is used.

Up to eight light objects are supported per axes. (Implementation dependent)

Lighting is only supported for OpenGL graphic toolkits (i.e., "fltk" and "qt").

A light object has the following properties which alter the appearance of the plot.

"Color": The color of the light can be passed as an RGB-vector (e.g., [1 0 0] for red) or as a string (e.g., "r" for red). The default color is white ([1 1 1]).

"Position": The direction from which the light emanates as a 1x3-vector. The default direction is [1 0 1].

"Style": This string defines whether the light emanates from a light source at infinite distance ("infinite") or from a local point source ("local"). The default is "infinite".

If the first argument *hax* is an axes handle, then add the light object to this axes, rather than the current axes returned by `gca`.

The optional return value *h* is a graphics handle to the created light object.

Programming Note: The full list of properties is documented at [Section 15.3.3.9 \[Light Properties\]](#), page 434.

See also: [\[lighting\]](#), page 345, [\[material\]](#), page 346, [\[patch\]](#), page 397, [\[surface\]](#), page 398.

15.3.2.2 Handle Functions

To determine whether a variable is a graphics object index, or an index to an axes or figure, use the functions `ishghandle`, `isgraphics`, `isaxes`, and `isfigure`.

`ishghandle (h)`

Return true if *h* is a graphics handle and false otherwise.

h may also be a matrix of handles in which case a logical array is returned that is true where the elements of *h* are graphics handles and false where they are not.

See also: [\[isgraphics\]](#), page 399, [\[isaxes\]](#), page 400, [\[isfigure\]](#), page 400, [\[ishandle\]](#), page 399.

`isgraphics (h)`

`isgraphics (h, type)`

Return true if *h* is a graphics handle (of type *type*) and false otherwise.

When no *type* is specified the function is equivalent to `ishghandle`.

See also: [\[ishghandle\]](#), page 399, [\[ishandle\]](#), page 399, [\[isaxes\]](#), page 400, [\[isfigure\]](#), page 400.

`ishandle (h)`

Return true if *h* is a handle to a graphics or Java object and false otherwise.

h may also be a matrix of handles in which case a logical array is returned that is true where the elements of *h* are handles to graphics or Java objects and false where they are not.

Programming Note: It is often more useful to test for a specific object type. To determine if a handle belongs to a graphics object use `ishghandle` or `isgraphics`. To determine if a handle belongs to a Java object use `isjava`.

See also: [\[ishghandle\]](#), page 399, [\[isgraphics\]](#), page 399, [\[isjava\]](#), page 991.

isaxes (h)

Return true if *h* is an axes graphics handle and false otherwise.

If *h* is a matrix then return a logical array which is true where the elements of *h* are axes graphics handles and false where they are not.

See also: [isfigure], page 400, [ishghandle], page 399, [isgraphics], page 399.

isfigure (h)

Return true if *h* is a figure graphics handle and false otherwise.

If *h* is a matrix then return a logical array which is true where the elements of *h* are figure graphics handles and false where they are not.

See also: [isaxes], page 400, [ishghandle], page 399, [isgraphics], page 399.

The function `gcf` returns an index to the current figure object, or creates one if none exists. Similarly, `gca` returns the current axes object, or creates one (and its parent figure object) if none exists.

h = groot ()

Return a handle to the root graphics object.

The root graphics object is the ultimate parent of all graphics objects.

In addition, the root object contains information about the graphics system as a whole such as the `ScreenSize`. Use `get (groot)` to find out what information is available.

Defaults for the graphic system as a whole are specified by setting properties of the root graphics object that begin with "Default". For example, to set the default font for all text objects to FreeSans use

```
set (groot, "DefaultTextFontName", "FreeSans")
```

Default properties can be deleted by using `set` with the special property value of "`remove`". To undo the default font assignment above use

```
set (groot, "DefaultTextFontName", "remove")
```

Programming Note: The root graphics object is identified by the special handle value of 0. At some point this unique value may change, but code can be made resistant to future changes by using `groot` which is guaranteed to always return the root graphics object.

See also: [gcf], page 400, [gca], page 401, [get], page 402, [set], page 402.

h = gcf ()

Return a handle to the current figure.

The current figure is the default target for graphics output. If multiple figures exist, `gcf` returns the last created figure or the last figure that was clicked on with the mouse.

If a current figure does not exist, create one and return its handle. The handle may then be used to examine or set properties of the figure. For example,

```
fplot (@sin, [-10, 10]);
fig = gcf ();
set (fig, "numbertitle", "off", "name", "sin plot")
```

plots a sine wave, finds the handle of the current figure, and then renames the figure window to describe the contents.

Note: To find the current figure without creating a new one if it does not exist, query the "CurrentFigure" property of the root graphics object.

```
get (groot, "currentfigure");
```

See also: [gca], page 401, [gco], page 401, [gcbf], page 459, [gcbo], page 458, [get], page 402, [set], page 402.

h = gca ()

Return a handle to the current axes object.

The current axes is the default target for graphics output. In the case of a figure with multiple axes, **gca** returns the last created axes or the last axes that was clicked on with the mouse.

If no current axes object exists, create one and return its handle. The handle may then be used to examine or set properties of the axes. For example,

```
ax = gca ();
set (ax, "position", [0.5, 0.5, 0.5, 0.5]);
```

creates an empty axes object and then changes its location and size in the figure window.

Note: To find the current axes without creating a new axes object if it does not exist, query the "CurrentAxes" property of a figure.

```
get (gcf, "currentaxes");
```

See also: [gcf], page 400, [gco], page 401, [gcbf], page 459, [gcbo], page 458, [get], page 402, [set], page 402.

h = gco ()

h = gco (fig)

Return a handle to the current object of the current figure, or a handle to the current object of the figure with handle *fig*.

The current object of a figure is the object that was last clicked on. It is stored in the "CurrentObject" property of the target figure.

If the last mouse click did not occur on any child object of the figure, then the current object is the figure itself.

If no mouse click occurred in the target figure, this function returns an empty matrix.

Programming Note: The value returned by this function is not necessarily the same as the one returned by **gcbo** during callback execution. An executing callback can be interrupted by another callback and the current object may be changed.

See also: [gcbo], page 458, [gca], page 401, [gcf], page 400, [gcbf], page 459, [get], page 402, [set], page 402.

The **get** and **set** functions may be used to examine and set properties for graphics objects. For example,

```
get (groot)
⇒ ans =
{
  type = root
  currentfigure = [] (0x0)
  children = [] (0x0)
  visible = on
  ...
}
```

returns a structure containing all the properties of the root figure. As with all functions in Octave, the structure is returned by value, so modifying it will not modify the internal root figure plot object. To do that, you must use the `set` function. Also, note that in this case, the `currentfigure` property is empty, which indicates that there is no current figure window.

The `get` function may also be used to find the value of a single property. For example,

```
get (gca (), "xlim")
⇒ [ 0 1 ]
```

returns the range of the x-axis for the current axes object in the current figure.

To set graphics object properties, use the `set` function. For example,

```
set (gca (), "xlim", [-10, 10]);
```

sets the range of the x-axis for the current axes object in the current figure to ‘[-10, 10]’.

Default property values can also be queried if the `set` function is called without a value argument. When only one argument is given (a graphic handle) then a structure with defaults for all properties of the given object type is returned. For example,

```
set (gca ())
```

returns a structure containing the default property values for axes objects. If `set` is called with two arguments (a graphic handle and a property name) then only the defaults for the requested property are returned.

```
val = get (h)
val = get (h, p)
```

Return the value of the named property *p* from the graphics handle *h*.

If *p* is omitted, return the complete property list for *h*.

If *h* is a vector, return a cell array including the property values or lists respectively.

See also: [\[set\]](#), page 402.

```
set (h, property, value, ...)
set (h, properties, values)
set (h, pv)
value_list = set (h, property)
all_value_list = set (h)
```

Set named property values for the graphics handle (or vector of graphics handles) *h*.

There are three ways to give the property names and values:

- as a comma separated list of *property, value* pairs

Here, each *property* is a string containing the property name, each *value* is a value of the appropriate type for the property.

- as a cell array of strings *properties* containing property names and a cell array *values* containing property values.

In this case, the number of columns of *values* must match the number of elements in *properties*. The first column of *values* contains values for the first entry in *properties*, etc. The number of rows of *values* must be 1 or match the number of elements of *h*. In the first case, each handle in *h* will be assigned the same values. In the latter case, the first handle in *h* will be assigned the values from the first row of *values* and so on.

- as a structure array *pv*

Here, the field names of *pv* represent the property names, and the field values give the property values. In contrast to the previous case, all elements of *pv* will be set in all handles in *h* independent of the dimensions of *pv*.

`set` is also used to query the list of values a named property will take. `clist = set(h, "property")` will return the list of possible values for "property" in the cell list *clist*. If no output variable is used then the list is formatted and printed to the screen.

If no property is specified (`slist = set(h)`) then a structure *slist* is returned where the fieldnames are the properties of the object *h* and the fields are the list of possible values for each property. If no output variable is used then the list is formatted and printed to the screen.

For example,

```
hf = figure();
set (hf, "paperorientation")
⇒ [ landscape | {portrait} ]
```

shows the paperorientation property can take two values with the default being "portrait".

See also: [\[get\]](#), page 402.

```
parent = ancestor (h, type)
```

```
parent = ancestor (h, type, "toplevel")
```

Return the first ancestor of handle object *h* whose type matches *type*, where *type* is a character string.

If *type* is a cell array of strings, return the first parent whose type matches any of the given type strings.

If the handle object *h* itself is of type *type*, return *h*.

If "toplevel" is given as a third argument, return the highest parent in the object hierarchy that matches the condition, instead of the first (nearest) one.

See also: [\[findobj\]](#), page 453, [\[findall\]](#), page 454, [\[allchild\]](#), page 404.

`h = allchild (handles)`

Find all children, including hidden children, of a graphics object.

This function is similar to `get (h, "children")`, but also returns hidden objects (`HandleVisibility = "off"`).

If `handles` is a scalar, `h` will be a vector. Otherwise, `h` will be a cell matrix of the same size as `handles` and each cell will contain a vector of handles.

See also: [\[findall\]](#), page 454, [\[findobj\]](#), page 453, [\[get\]](#), page 402, [\[set\]](#), page 402.

`findfigs ()`

Find all visible figures that are currently off the screen and move them onto the screen.

See also: [\[allchild\]](#), page 404, [\[figure\]](#), page 372, [\[get\]](#), page 402, [\[set\]](#), page 402.

Figures can be printed or saved in many graphics formats with `print` and `saveas`. Occasionally, however, it may be useful to save the original Octave handle graphic directly so that further modifications can be made such as modifying a title or legend.

This can be accomplished with the following functions by

```
fig_struct = hdl2struct (gcf);
save myplot.fig -struct fig_struct;
...
fig_struct = load ("myplot.fig");
struct2hdl (fig_struct);
```

`s = hdl2struct (h)`

Return a structure, `s`, whose fields describe the properties of the object, and its children, associated with the handle, `h`.

The fields of the structure `s` are `"type"`, `"handle"`, `"properties"`, `"children"`, and `"special"`.

See also: [\[struct2hdl\]](#), page 404, [\[hgsave\]](#), page 390, [\[findobj\]](#), page 453.

`h = struct2hdl (s)`**`h = struct2hdl (s, p)`****`h = struct2hdl (s, p, hilev)`**

Construct a graphics handle object `h` from the structure `s`.

The structure must contain the fields `"handle"`, `"type"`, `"children"`, `"properties"`, and `"special"`.

If the handle of an existing figure or axes is specified, `p`, the new object will be created as a child of that object. If no parent handle is provided then a new figure and the necessary children will be constructed using the default values from the root figure.

A third boolean argument `hilev` can be passed to specify whether the function should preserve listeners/callbacks, e.g., for legends or hggroups. The default is false.

See also: [\[hdl2struct\]](#), page 404, [\[hgload\]](#), page 390, [\[findobj\]](#), page 453.

`hnew = copyobj (horig)`**`hnew = copyobj (horig, hparent)`**

Construct a copy of the graphic objects associated with the handles `horig` and return new handles `hnew` to the new objects.

If a parent handle *hparent* (root, figure, axes, or hggroup) is specified, the copied object will be created as a child of *hparent*.

If *horig* is a vector of handles, and *hparent* is a scalar, then each handle in the vector *hnew* has its "Parent" property set to *hparent*. Conversely, if *horig* is a scalar and *hparent* a vector, then each parent object will receive a copy of *horig*. If *horig* and *hparent* are both vectors with the same number of elements then *hnew(i)* will have parent *hparent(i)*.

See also: [\[struct2hdl\]](#), page 404, [\[hdl2struct\]](#), page 404, [\[findobj\]](#), page 453.

15.3.3 Graphics Object Properties

In this Section the graphics object properties are discussed in detail, starting with the root figure properties and continuing through the objects hierarchy. The documentation about a specific graphics object can be displayed using `doc` function, e.g., `doc ("axes properties")` will show [Section 15.3.3.3 \[Axes Properties\]](#), page 412.

The allowed values for radio (string) properties can be retrieved programmatically or displayed using the one or two arguments call to `set` function. See [\[set\]](#), page 402.

In the following documentation, default values are enclosed in { }.

15.3.3.1 Root Figure Properties

The root figure properties are:

`beingdeleted: {"off"} | "on"`
 `beingdeleted` is unused.

`busyaction: "cancel" | {"queue"}`
 `busyaction` is unused.

`buttondownfcn: string | function handle, def. [] (0x0)`
 `buttondownfcn` is unused.

`callbackobject` (read-only): graphics handle, def. [] (0x0)
 Graphics handle of the current object whose callback is executing.

`children` (read-only): vector of graphics handles, def. [] (0x1)
 Graphics handles of the root's children.

`clipping: "off" | {"on"}`
 `clipping` is unused.

`commandwindowsize` (read-only): def. [0 0]
`createfcn: string | function handle, def. [] (0x0)`
 `createfcn` is unused.

`currentfigure`: graphics handle, def. [] (0x0)
 Graphics handle of the current figure.

`deletefcn: string | function handle, def. [] (0x0)`
 `deletefcn` is unused.

`fixedwidthfontname: string, def. "Courier"`
`handlevisibility: "callback" | "off" | {"on"}`
 `handlevisibility` is unused.

hittest: "off" | {"on"}
 hittest is unused.

interruptible: "off" | {"on"}
 interruptible is unused.

monitorpositions (read-only):
 monitorpositions is unused.

parent: graphics handle, def. [] (0x0)
 Root figure has no parent graphics object. **parent** is always empty.

pickableparts: "all" | "none" | {"visible"}
 pickableparts is unused.

pointerlocation: two-element vector, def. [0 0]
 pointerlocation is unused.

pointerwindow (read-only): graphics handle, def. 0
 pointerwindow is unused.

screendepth (read-only): double

screenpixelsperinch (read-only): double

screensize (read-only): four-element vector

selected: {"off"} | "on"
 selected is unused.

selectionhighlight: "off" | {"on"}
 selectionhighlight is unused.

showhiddenhandles: {"off"} | "on"
 If **showhiddenhandles** is "on", all graphics objects handles are visible in their parents' children list, regardless of the value of their **handlevisibility** property.

tag: string, def. ""
 A user-defined string to label the graphics object.

type (read-only): string
 Class name of the graphics object. **type** is always "root"

uicontextmenu: graphics handle, def. [] (0x0)
 uicontextmenu is unused.

units: "centimeters" | "characters" | "inches" | "normalized" | {"pixels"} | "points"

userdata: Any Octave data, def. [] (0x0)
 User-defined data to associate with the graphics object.

visible: "off" | {"on"}
 visible is unused.

15.3.3.2 Figure Properties

The `figure` properties are:

`alphamap`: def. 64-by-1 double

Transparency is not yet implemented for figure objects. `alphamap` is unused.

`beingdeleted`: {"off"} | "on"

`busyaction`: "cancel" | {"queue"}

Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its `interruptible` property set to "off". The `busyaction` property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

`buttondownfcn`: string | function handle, def. [] (0x0)

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`children` (read-only): vector of graphics handles, def. [] (0x1)

Graphics handles of the figure's children.

`clipping`: "off" | {"on"}

`clipping` is unused.

`closerequestfcn`: string | function handle, def. "closereq"

Function that is executed when a figure is deleted. See [\[closereq function\], page 378](#).

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`color`: colorspec, def. [1 1 1]

Color of the figure background. See [Section 15.4.1 \[colorspec\], page 456](#).

`colormap`: N-by-3 matrix, def. 64-by-3 double

A matrix containing the RGB color map for the current axes.

`createfcn`: string | function handle, def. [] (0x0)

Callback function executed immediately after figure has been created. Function is set by using default property on root object, e.g., `set (groot, "defaultfigurecreatefcn", 'disp ("figure created!")')`.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`currentaxes`: graphics handle, def. [] (0x0)

Handle to the graphics object of the current axes.

`currentcharacter` (read-only): def. ""

`currentcharacter` is unused.

`currentobject` (read-only): graphics handle, def. [] (0x0)

`currentpoint` (read-only): two-element vector, def. [0; 0]

A 1-by-2 matrix which holds the coordinates of the point over which the mouse pointer was when a mouse event occurred. The X and Y coordinates are in

units defined by the figure's `units` property and their origin is the lower left corner of the plotting area.

Events which set `currentpoint` are

A mouse button was pressed
always

A mouse button was released
only if the figure's callback `windowbuttonupfcn` is defined

The pointer was moved while pressing the mouse button (drag)
only if the figure's callback `windowbuttonmotionfcn` is defined

`deletefcn`: string | function handle, def. [] (0x0)

Callback function executed immediately before figure is deleted.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`dockcontrols`: {"off"} | "on"

`dockcontrols` is unused.

`filename`: string, def. "

The filename used when saving the plot figure.

`graphicssmoothing`: "off" | {"on"}

Use smoothing techniques to reduce the appearance of jagged lines.

`handlevisibility`: "callback" | "off" | {"on"}

If `handlevisibility` is "off", the figure's handle is not visible in its parent's "children" property.

`hittest`: "off" | {"on"}

Specify whether figure processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [\[pickableparts property\], page 410](#).

`integerhandle`: "off" | {"on"}

Assign the next lowest unused integer as the Figure number.

`interruptible`: "off" | {"on"}

Specify whether this object's callback functions may be interrupted by other callbacks. By default `interruptible` is "on" and callbacks that make use of `drawnow`, `figure`, `waitfor`, `getframe` or `pause` functions are eventually interrupted. See [Section 15.4.4 \[Callbacks section\], page 457](#).

`inverthardcopy`: "off" | {"on"}

Replace the figure and axes background color with white when printing.

`keypressfcn`: string | function handle, def. [] (0x0)

Callback function executed when a keystroke event happens while the figure has focus. The actual key that was pressed can be retrieved using the second

argument 'evt' of the function. For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

keyreleasefcn: string | function handle, def. [] (0x0)

With **keypressfcn**, the keyboard callback functions. These callback functions are called when a key is pressed/released respectively. The functions are called with two input arguments. The first argument holds the handle of the calling figure. The second argument holds an event structure which has the following members:

Character:

The ASCII value of the key

Key: Lowercase value of the key

Modifier:

A cell array containing strings representing the modifiers pressed with the key.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

menubar: {"figure"} | "none"

Control the display of the figure menu bar at the top of the figure.

name: string, def. ""

Name to be displayed in the figure title bar. The name is displayed to the right of any title determined by the **numbertitle** property.

nextplot: {"add"} | "new" | "replace" | "replacechildren"

nextplot is used by high level plotting functions to decide what to do with axes already present in the figure. See [\[newplot function\], page 375](#).

number (read-only):

numbertitle: "off" | {"on"}

Display "Figure" followed by the numerical figure handle value in the figure title bar.

outerposition: four-element vector, def. [-1 -1 -1 -1]

Specify the position and size of the figure including the top menubar and the bottom status bar. The four elements of the vector are the coordinates of the lower left corner and width and height of the figure. See [\[units property\], page 411](#).

paperorientation: "landscape" | {"portrait"}

The value for the **papersize**, and **paperposition** properties depends upon **paperorientation**. The horizontal and vertical values for **papersize** and **paperposition** reverse order when **paperorientation** is switched between "portrait" and "landscape".

paperposition: four-element vector, def. [0.25000 2.50000 8.00000 6.00000]

Vector [**left** **bottom** **width** **height**] defining the position and size of the figure (in **paperunits** units) on the printed page. The position [**left** **bottom**] defines the lower left corner of the figure on the page, and the size is defined by [**width**

height]. For output formats not implicitly rendered on paper, **width** and **height** define the size of the image and the position information is ignored. Setting **paperposition** also forces the **paperpositionmode** property to be set to "manual".

paperpositionmode: "auto" | {"manual"}

If **paperpositionmode** is set to "auto", the **paperposition** property is automatically computed: the printed figure will have the same size as the on-screen figure and will be centered on the output page. Setting the **paperpositionmode** to "auto" does not modify the value of the **paperposition** property.

papersize: two-element vector, def. [8.5000 11.0000]

Vector [**width** **height**] defining the size of the paper for printing. Setting the **papersize** property to a value, not associated with one of the defined **papertypes** and consistent with the setting for **paperorientation**, forces the **papertype** property to the value "<custom>". If **papersize** is set to a value associated with a supported **papertype** and consistent with the **paperorientation**, the **papertype** value is modified to the associated value.

papertype: "<custom>" | "a" | "a0" | "a1" | "a2" | "a3" | "a4" | "a5" | "arch-a" | "arch-b" | "arch-c" | "arch-d" | "arch-e" | "b" | "b0" | "b1" | "b2" | "b3" | "b4" | "b5" | "c" | "d" | "e" | "tabloid" | "uslegal" | {"usletter"}

Name of the paper used for printed output. Setting **papertype** also changes **papersize**, while maintaining consistency with the **paperorientation** property.

paperunits: "centimeters" | {"inches"} | "normalized" | "points"

The unit used to compute the **paperposition** property. For **paperunits** set to "pixels", the conversion between physical units (ex: "inches") and "pixels" is dependent on the **screenpixelsperinch** property of the root object.

parent: graphics handle, def. 0

Handle of the parent graphics object.

pickableparts (read-only): "all" | "none" | {"visible"}

pickableparts is unused.

pointer: {"arrow"} | "botl" | "botr" | "bottom" | "circle" | "cross" | "crosshair" | "custom" | "fleur" | "fullcrosshair" | "hand" | "ibeam" | "left" | "right" | "top" | "topl" | "topr" | "watch"
pointer is unused.

pointershapecdata: def. 16-by-16 double

pointershapecdata is unused.

pointershapehotspot: def. [0 0]

pointershapehotspot is unused.

position: four-element vector, def. [300 200 560 420]

Specify the position and size of the figure canvas. The four elements of the vector are the coordinates of the lower left corner and width and height of the figure. See [\[units property\], page 411](#).

renderer: {"opengl"} | "painters"
 Renderering engine used for printing when **renderermode** is "manual". Setting **renderer** also forces the **renderermode** property to be set to "manual".

renderermode: {"auto"} | "manual"
 Control whether the renderering engine used for printing is chosen automatically or specified by the **renderer** property. See [print function], page 383.

resize: "off" | {"on"}
 Control whether the figure can be resized by dragging the window borders and corners using a mouse. When **resize** is "off" mouse interactions are disabled but the figure can still be resized by changing its "position" property.

resizefcn: string | function handle, def. [] (0x0)
 resizefcn is deprecated. Use **sizechangedfcn** instead.

selected: {"off"} | "on"
selectionhighlight: "off" | {"on"}
selectiontype: "alt" | "extend" | {"normal"} | "open"
sizechangedfcn: string | function handle, def. [] (0x0)
 Callback triggered when the figure window size is changed.
 For information on how to write graphics listener functions see Section 15.4.4 [Callbacks section], page 457.

tag: string, def. ""
 A user-defined string to label the graphics object.

toolbar: {"auto"} | "figure" | "none"
 Control the display of the toolbar (along the bottom of the menubar) and the status bar. When set to "auto", the display is based on the value of the **menubar** property.

type (read-only): string
 Class name of the graphics object. **type** is always "figure"

uicontextmenu: graphics handle, def. [] (0x0)
 Graphics handle of the uicontextmenu object that is currently associated to this figure object.

units: "centimeters" | "characters" | "inches" | "normalized" | {"pixels"} | "points"
 The unit used to compute the **position** and **outerposition** properties.

userdata: Any Octave data, def. [] (0x0)
 User-defined data to associate with the graphics object.

visible: "off" | {"on"}
 If **visible** is "off", the figure is not rendered on screen.

windowbuttondownfcn: string | function handle, def. [] (0x0)
 See [windowbuttonupfcn property], page 411.

windowbuttonmotionfcn: string | function handle, def. [] (0x0)
 See [windowbuttonupfcn property], page 411.

`windowbuttonupfcn`: string | function handle, def. [] (0x0)

With `windowbuttondownfcn` and `windowbuttonmotionfcn`, the mouse callback functions. These callback functions are called when a mouse button is pressed, dragged, or released respectively. When these callback functions are executed, the `currentpoint` property holds the current coordinates of the cursor.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`windowkeypressfcn`: string | function handle, def. [] (0x0)

Function that is executed when a key is pressed and the figure has focus.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`windowkeyreleasefcn`: string | function handle, def. [] (0x0)

Function that is executed when a key is released and the figure has focus.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`windowscrollwheelfcn`: string | function handle, def. [] (0x0)

`windowscrollwheelfcn` is unused.

`windowstyle`: "docked" | "modal" | {"normal"}

The window style of a figure. One of the following values:

`normal` Set the window style as non modal.

`modal` Set the window as modal so that it will stay on top of all normal figures.

`docked` Setting the window style as docked currently does not dock the window.

Changing modes of a visible figure may cause the figure to close and reopen.

15.3.3 Axes Properties

The `axes` properties are:

`activepositionproperty`: {"outerposition"} | "position"

Specify which of "position" or "outerposition" properties takes precedence when axes annotations extent changes. See [\[position property\], page 416](#), and [\[outerposition property\], page 416](#).

`alim`: def. [0 1]

Transparency is not yet implemented for axes objects. `alim` is unused.

`alimmode`: {"auto"} | "manual"

`ambientlightcolor`: def. [1 1 1]

`ambientlightcolor` is unused.

`beingdeleted`: {"off"} | "on"

`box`: {"off"} | "on"

Control whether the axes has a surrounding box.

boxstyle: {"back"} | "full"
 For 3-D axes, control whether the "full" box is drawn or only the 3 "back" axes

busyaction: "cancel" | {"queue"}
 Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its **interruptible** property set to "off". The **busyaction** property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

buttondownfcn: string | function handle, def. [] (0x0)
 For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

cameraposition: three-element vector, def. [0.50000 0.50000 9.16025]
camerapositionmode: {"auto"} | "manual"
cameratarget: three-element vector, def. [0.50000 0.50000 0.50000]
cameratargetmode: {"auto"} | "manual"
cameraupvector: three-element vector, def. [-0 1 0]
cameraupvectormode: {"auto"} | "manual"
cameraviewangle: scalar, def. 6.6086
cameraviewanglemode: {"auto"} | "manual"
children (read-only): vector of graphics handles, def. [] (0x1)
 Graphics handles of the axes's children.

clim: two-element vector, def. [0 1]
 Define the limits for the color axis of image children. Setting **clim** also forces the **climmode** property to be set to "manual". See [\[pcolor function\], page 321](#).

climmode: {"auto"} | "manual"
clipping: "off" | {"on"}
 clipping is unused.

clippingstyle: {"3dbox"} | "rectangle"
 clippingstyle is unused.

color: colorspec, def. [1 1 1]
 Color of the axes background. See [Section 15.4.1 \[colorspec\], page 456](#).

colormap: def. 64-by-3 double
colororder: N-by-3 RGB matrix, def. 7-by-3 double
 RGB values used by plot function for automatic line coloring.

colororderindex: def. 1
 colororderindex is unused.

createfcn: string | function handle, def. [] (0x0)
 Callback function executed immediately after axes has been created.
 Function is set by using default property on root object, e.g., `set (groot, "defaultaxescreatefcn", 'disp ("axes created!")')`.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

currentpoint: 2-by-3 matrix, def. 2-by-3 double

Matrix [`xf`, `yf`, `zf`; `xb`, `yb`, `zb`] which holds the coordinates (in axes data units) of the point over which the mouse pointer was when the mouse button was pressed. If a mouse callback function is defined, `currentpoint` holds the pointer coordinates at the time the mouse button was pressed. For 3-D plots, the first row of the returned matrix specifies the point nearest to the current camera position and the second row the furthest point. The two points forms a line which is perpendicular to the screen.

dataaspectratio: three-element vector, def. [1 1 1]

Specify the relative height and width of the data displayed in the axes. Setting `dataaspectratio` to [1, 2] causes the length of one unit as displayed on the x-axis to be the same as the length of 2 units on the y-axis. See [\[daspect function\], page 358](#). Setting `dataaspectratio` also forces the `dataaspectratio mode` property to be set to "manual".

dataaspectratio mode: {"auto"} | "manual"

deletefcn: string | function handle, def. [] (0x0)

Callback function executed immediately before axes is deleted.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

fontangle: "italic" | {"normal"}

Control whether the font is italic or normal.

fontname: string, def. "*"

Name of font used for text rendering. When setting this property, the text rendering engine will search for a matching font in your system. If none is found then text is rendered using a default sans serif font (same as the default "*" value).

fontsize: scalar, def. 10

Size of the font used for text rendering. See [\[fontunits property\], page 414](#).

fontsMOOTHING: "off" | {"on"}

`fontsMOOTHING` is unused.

fontunits: "centimeters" | "inches" | "normalized" | "pixels" | {"points"}

Units used to interpret the "fontsize" property.

fontweight: "bold" | {"normal"}

Control the variant of the base font used for text rendering.

gridalpha: def. 0.15000

Transparency is not yet implemented for axes objects. `gridalpha` is unused.

gridalphemode: {"auto"} | "manual"

`gridalphemode` is unused.

gridcolor: def. [0.15000 0.15000 0.15000]

`gridcolor` is unused.

gridcolormode: {"auto"} | "manual"
gridcolormode is unused.

gridlinestyle: {"-"} | {"--"} | {"-."} | {":"} | "none"
handlevisibility: "callback" | "off" | {"on"}
If handlevisibility is "off", the axes's handle is not visible in its parent's "children" property.

hittest: "off" | {"on"}
Specify whether axes processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [pickableparts property], page 416.

interruptible: "off" | {"on"}
Specify whether this object's callback functions may be interrupted by other callbacks. By default interruptible is "on" and callbacks that make use of drawnow, figure, waitfor, getframe or pause functions are eventually interrupted. See Section 15.4.4 [Callbacks section], page 457.

labelfontsizemultiplier: def. 1.1000
Ratio between the x/y/zlabel fontsize and the tick label fontsize

layer: {"bottom"} | "top"
Control whether the axes is drawn below child graphics objects (ticks, labels, etc. covered by plotted objects) or above.

linestyleorder: def. {"-"}
linestyleorder is unused.

linestyleorderindex: def. 1
linestyleorderindex is unused.

linewidth: def. 0.50000
Width of the main axes lines

minorgridalpha: def. 0.25000
Transparency is not yet implemented for axes objects. minorgridalpha is unused.

minorgridalphemode: {"auto"} | "manual"
minorgridalphemode is unused.

minorgridcolor: def. [0.10000 0.10000 0.10000]
minorgridcolor is unused.

minorgridcolormode: {"auto"} | "manual"
minorgridcolormode is unused.

minorgridlinestyle: {"-"} | {"--"} | {"-."} | {":"} | "none"
mousewheelzoom: scalar in the range (0, 1), def. 0.50000
Fraction of axes limits to zoom for each wheel movement.

nextplot: "add" | {"replace"} | "replacechildren"

nextplot is used by high level plotting functions to decide what to do with graphics objects already present in the axes. See [[newplot function](#)], page 375. The state of **nextplot** is typically controlled using the **hold** function. See [[hold function](#)], page 375.

outerposition: four-element vector, def. [0 0 1 1]

Specify the position of the plot including titles, axes, and legend. The four elements of the vector are the coordinates of the lower left corner and width and height of the plot, in units normalized to the width and height of the plot window. For example, [0.2, 0.3, 0.4, 0.5] sets the lower left corner of the axes at (0.2,0.3) and the width and height to be 0.4 and 0.5 respectively. See [[position property](#)], page 416.

parent: graphics handle

Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}

Specify whether axes will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the axes or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "hittest" property will determine how they are processed. See [[hittest property](#)], page 415.

plotboxaspectratio: def. [1 1 1]

See [[pbaspect function](#)], page 358. Setting **plotboxaspectratio** also forces the **plotboxaspectratio mode** property to be set to "manual".

plotboxaspectratio mode: {"auto"} | "manual"

position: four-element vector, def. [0.13000 0.11000 0.77500 0.81500]

Specify the position of the plot excluding titles, axes, and legend. The four elements of the vector are the coordinates of the lower left corner and width and height of the plot, in units normalized to the width and height of the plot window. For example, [0.2, 0.3, 0.4, 0.5] sets the lower left corner of the axes at (0.2,0.3) and the width and height to be 0.4 and 0.5 respectively. See [[outerposition property](#)], page 416.

projection: {"orthographic"} | "perspective"

projection is unused.

selected: {"off"} | "on"

selectionhighlight: "off" | {"on"}

sortmethod: "childorder" | {"depth"}

sortmethod is unused.

tag: string, def. ""

A user-defined string to label the graphics object.

tickdir: {"in"} | "out"

Control whether axes tick marks project "in" to the plot box or "out". Setting **tickdir** also forces the **tickdir mode** property to be set to "manual".

tickdirmode: {"auto"} | "manual"
ticklabelinterpreter: "latex" | "none" | {"tex"}
 Control the way x/y/zticklabel properties are interpreted. See
 [Use of the interpreter property], page 378.

ticklength: two-element vector, def. [0.010000 0.025000]
 Two-element vector [2Dlen 3Dlen] specifying the length of the tickmarks relative to the longest visible axis.

tightinset (read-only): four-element vector
 Size of the [left bottom right top] margins around the axes that enclose labels and title annotations.

title: graphics handle
 Graphics handle of the title text object.

titlefontsizemultiplier: positive scalar, def. 1.1000
 Ratio between the title fontsize and the tick label fontsize

titlefontweight: {"bold"} | "normal"
 Control variant of base font used for the axes title.

type (read-only): string
 Class name of the graphics object. type is always "axes"

uicontextmenu: graphics handle, def. [] (0x0)
 Graphics handle of the uicontextmenu object that is currently associated to this axes object.

units: "centimeters" | "characters" | "inches" | {"normalized"} | "pixels" | "points"
 Units used to interpret the "position", "outerposition", and "tightinset" properties.

userdata: Any Octave data, def. [] (0x0)
 User-defined data to associate with the graphics object.

view: two-element vector, def. [0 90]
 Two-element vector [azimuth elevation] specifying the viewpoint for three-dimensional plots

visible: "off" | {"on"}
 If visible is "off", the axes is not rendered on screen.

xaxislocation: {"bottom"} | "origin" | "top"
 Control the x axis location.

xcolor: {colorspec} | "none", def. [0.15000 0.15000 0.15000]
 Color of the x-axis. See Section 15.4.1 [colorspec], page 456. Setting xcolor also forces the xcolormode property to be set to "manual".

xcolormode: {"auto"} | "manual"
xdir: {"normal"} | "reverse"
 Direction of the x axis: "normal" is left to right.

xgrid: {"off"} | "on"
Control whether major x grid lines are displayed.

xlabel: graphics handle
Graphics handle of the x label text object.

xlim: two-element vector, def. [0 1]
Two-element vector [xmin xmax] specifying the limits for the x-axis. Setting **xlim** also forces the **xlimmode** property to be set to "manual". See [xlim function], page 325.

xlimmode: {"auto"} | "manual"

xminorgrid: {"off"} | "on"
Control whether minor x grid lines are displayed.

xminortick: {"off"} | "on"

xscale: {"linear"} | "log"

xtick: vector
Position of x tick marks. Setting **xtick** also forces the **xtickmode** property to be set to "manual".

xticklabel: string | cell array of strings, def. 1-by-6 cell
Labels of x tick marks. Setting **xticklabel** also forces the **xticklabelmode** property to be set to "manual".

xticklabelmode: {"auto"} | "manual"

xticklabelrotation: def. 0
xticklabelrotation is unused.

xtickmode: {"auto"} | "manual"

yaxislocation: {"left"} | "origin" | "right"
Control the y-axis location.

ycolor: {colorspec} | "none", def. [0.15000 0.15000 0.15000]
Color of the y-axis. See Section 15.4.1 [colorspec], page 456.

ycolormode: {"auto"} | "manual"

ydir: {"normal"} | "reverse"
Direction of the y-axis: "normal" is bottom to top.

ygrid: {"off"} | "on"
Control whether major y grid lines are displayed.

ylabel: graphics handle
Graphics handle of the y label text object.

ylim: two-element vector, def. [0 1]
Two-element vector [ymin ymax] specifying the limits for the y-axis. Setting **ylim** also forces the **ylimmode** property to be set to "manual". See [ylim function], page 325.

ylimmode: {"auto"} | "manual"

yminorgrid: {"off"} | "on"
Control whether minor y grid lines are displayed.

```

yminortick: {"off"} | "on"
yscale: {"linear"} | "log"
ytick: vector
    Position of y tick marks. Setting ytick also forces the ytickmode property to be set to "manual".

yticklabel: string | cell array of strings, def. 1-by-6 cell
    Labels of y tick marks. Setting yticklabel also forces the yticklabelmode property to be set to "manual".

yticklabelmode: {"auto"} | "manual"
yticklabelrotation: def. 0
    yticklabelrotation is unused.

ytickmode: {"auto"} | "manual"
zcolor: {colorspec} | "none", def. [0.15000 0.15000 0.15000]
    Color of the z-axis. See Section 15.4.1 \[colorspec\], page 456.

zcolormode: {"auto"} | "manual"
zdir: {"normal"} | "reverse"
zgrid: {"off"} | "on"
    Control whether major z grid lines are displayed.

zlabel: graphics handle
    Graphics handle of the z label text object.

zlim: two-element vector, def. [0 1]
    Two-element vector [zmin zmax] specifying the limits for the z-axis. Setting zlim also forces the zlimmode property to be set to "manual". See \[zlim function\], page 325.

zlimmode: {"auto"} | "manual"
zminorgrid: {"off"} | "on"
    Control whether minor z grid lines are displayed.

zminortick: {"off"} | "on"
zscale: {"linear"} | "log"
ztick: vector
    Position of z tick marks. Setting ztick also forces the ztickmode property to be set to "manual".

zticklabel: string | cell array of strings, def. 1-by-6 cell
    Labels of z tick marks. Setting zticklabel also forces the zticklabelmode property to be set to "manual".

zticklabelmode: {"auto"} | "manual"
zticklabelrotation: def. 0
    zticklabelrotation is unused.

ztickmode: {"auto"} | "manual"

```

15.3.3.4 Line Properties

The line properties are:

beingdeleted: {"off"} | "on"
busyaction: "cancel" | {"queue"}

Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its **interruptible** property set to "off". The **busyaction** property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

buttondownfcn: string | function handle, def. [] (0x0)

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

children (read-only): vector of graphics handles, def. [] (0x1)
children is unused.

clipping: "off" | {"on"}

If **clipping** is "on", the line is clipped in its parent axes limits.

color: colorspec, def. [0 0 0]

Color of the line object. See [Section 15.4.1 \[colorspec\], page 456](#).

createfcn: string | function handle, def. [] (0x0)

Callback function executed immediately after line has been created. Function is set by using default property on root object, e.g., `set (groot, "defaultlinecreatefcn", 'disp ("line created!")')`.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

deletefcn: string | function handle, def. [] (0x0)

Callback function executed immediately before line is deleted.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

displayname: string | cell array of strings, def. ""

Text for the legend entry corresponding to this line.

handlevisibility: "callback" | "off" | {"on"}

If **handlevisibility** is "off", the line's handle is not visible in its parent's "children" property.

hittest: "off" | {"on"}

Specify whether line processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [\[pickableparts property\], page 421](#).

interruptible: "off" | {"on"}

Specify whether this object's callback functions may be interrupted by other callbacks. By default **interruptible** is "on" and callbacks that make use of

`drawnow`, `figure`, `waitfor`, `getframe` or `pause` functions are eventually interrupted. See [Section 15.4.4 \[Callbacks section\]](#), page [457](#).

linejoin: "chamfer" | "miter" | {"round"}
Control the shape of the junction of line segments. This property currently only affects the printed output.

linestyle: {"-"} | {"--"} | {"-."} | {"::"} | {"none"}
See [Section 15.4.2 \[Line Styles\]](#), page [456](#).

linewidth: def. 0.50000
Width of the line object measured in points.

marker: "*" | "+" | "." | "<" | ">" | "^" | "d" | "diamond" | "h" | "hexagram" | {"none"} | "o" | "p" | "pentagram" | "s" | "square" | "v" | "x"
Shape of the marker for each data point. See [Section 15.4.3 \[Marker Styles\]](#), page [457](#).

markeredgecolor: {"auto"} | {"none"}
Color of the edge of the markers. When set to "auto", the marker edges have the same color as the line. If set to "none", no marker edges are displayed. This property can also be set to any color. See [Section 15.4.1 \[colorspec\]](#), page [456](#).

markerfacecolor: "auto" | {"none"}
Color of the face of the markers. When set to "auto", the marker faces have the same color as the line. If set to "none", the marker faces are not displayed. This property can also be set to any color. See [Section 15.4.1 \[colorspec\]](#), page [456](#).

markersize: scalar, def. 6
Size of the markers measured in points.

parent: graphics handle
Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}
Specify whether line will accept mouse clicks. By default, `pickableparts` is "visible" and only visible parts of the line or its children may react to mouse clicks. When `pickableparts` is "all" both visible and invisible parts (or children) may react to mouse clicks. When `pickableparts` is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "`hittest`" property will determine how they are processed. See [\[hittest property\]](#), page [420](#).

selected: {"off"} | {"on"}

selectionhighlight: "off" | {"on"}

tag: string, def. ""
A user-defined string to label the graphics object.

type (read-only): string
Class name of the graphics object. `type` is always "line"

uicontextmenu: graphics handle, def. [] (0x0)
Graphics handle of the uicontextmenu object that is currently associated to this line object.

userdata: Any Octave data, def. [] (0x0)
 User-defined data to associate with the graphics object.

visible: "off" | {"on"}
 If **visible** is "off", the line is not rendered on screen.

xdata: vector, def. [0 1]
 Vector of x data to be plotted.

xdatasource: string, def. ""
 Name of a vector in the current base workspace to use as x data.

ydata: vector, def. [0 1]
 Vector of y data to be plotted.

ydatasource: string, def. ""
 Name of a vector in the current base workspace to use as y data.

zdata: vector, def. [] (0x0)
 Vector of z data to be plotted.

zdatasource: string, def. ""
 Name of a vector in the current base workspace to use as z data.

15.3.3.5 Text Properties

The **text** properties are:

backgroundcolor: colorspec, def. "none"
 Color of the background area. See [Section 15.4.1 \[colorspec\]](#), page 456.

beingdeleted: {"off"} | "on"

busyaction: "cancel" | {"queue"}
 Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its **interruptible** property set to "off". The **busyaction** property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\]](#), page 457.

buttondownfcn: string | function handle, def. [] (0x0)
 For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\]](#), page 457.

children (read-only): vector of graphics handles, def. [] (0x1)
 children is unused.

clipping: "off" | {"on"}
 If **clipping** is "on", the text is clipped in its parent axes limits.

color: colorspec, def. [0 0 0]
 Color of the text. See [Section 15.4.1 \[colorspec\]](#), page 456.

createfcn: string | function handle, def. [] (0x0)
 Callback function executed immediately after text has been created.
 Function is set by using default property on root object, e.g., **set** (**groot**, "defaulttextcreatefcn", 'disp ("text created!")').

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

deletefcn: string | function handle, def. [] (0x0)

Callback function executed immediately before text is deleted.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

edgecolor: colorspec, def. "none"

Color of the outline of the background area. See [Section 15.4.1 \[colorspec\], page 456](#).

editing: {"off"} | "on"

editing is unused.

extent (read-only): four-element vector

Vector [x0 y0 width height] indicating the size and location of the text string.

fontangle: "italic" | {"normal"} | "oblique"

Control whether the font is italic or normal.

fontname: string, def. "*"

Name of font used for text rendering. When setting this property, the text rendering engine will search for a matching font in your system. If none is found then text is rendered using a default sans serif font (same as the default "*" value).

fontsize: scalar, def. 10

Size of the font used for text rendering. See [\[fontunits property\], page 423](#).

fontunits: "centimeters" | "inches" | "normalized" | "pixels" | {"points"}

Units used to interpret the "fontsize" property.

fontweight: "bold" | "demi" | "light" | {"normal"}

Control the variant of the base font used for text rendering.

handlevisibility: "callback" | "off" | {"on"}

If **handlevisibility** is "off", the text's handle is not visible in its parent's "children" property.

hittest: "off" | {"on"}

Specify whether text processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [\[pickableparts property\], page 424](#).

horizontalalignment: "center" | {"left"} | "right"

interpreter: "latex" | "none" | {"tex"}

Control the way the "string" property is interpreted. See [\[Use of the interpreter property\], page 378](#).

interruptible: "off" | {"on"}

Specify whether this object's callback functions may be interrupted by other callbacks. By default **interruptible** is "on" and callbacks that make use of **drawnow**, **figure**, **waitfor**, **getframe** or **pause** functions are eventually interrupted. See [Section 15.4.4 \[Callbacks section\], page 457](#).

linestyle: {"-"} | {"--"} | {"-."} | {":"} | {"none"}

Style of the outline. See [Section 15.4.2 \[Line Styles\], page 456](#).

linewidth: scalar, def. 0.50000

Width of the outline.

margin: scalar, def. 3

Margins between the borders of the background area and the texts. The value is currently interpreted as pixels, regardless of the "**fontunits**" property.

parent: graphics handle

Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}

Specify whether text will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the text or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "**hittest**" property will determine how they are processed. See [\[hittest property\], page 423](#).

position: four-element vector, def. [0 0 0]

Vector [X0 Y0 Z0] where X0, Y0 and Z0 indicate the position of the text anchor as defined by **verticalalignment** and **horizontalalignment**.

rotation: scalar, def. 0

The angle of rotation for the displayed text, measured in degrees.

selected: {"off"} | "on"

selectionhighlight: "off" | {"on"}

string: string, def. ""

The text object string content.

tag: string, def. ""

A user-defined string to label the graphics object.

type (read-only): string

Class name of the graphics object. **type** is always "text"

uicontextmenu: graphics handle, def. [] (0x0)

Graphics handle of the uicontextmenu object that is currently associated to this text object.

units: "centimeters" | {"data"} | "inches" | "normalized" | "pixels" | "points"

userdata: Any Octave data, def. [] (0x0)

User-defined data to associate with the graphics object.

`verticalalignment: "baseline" | "bottom" | "cap" | {"middle"} | "top"`

`visible: "off" | {"on"}`

If `visible` is "off", the text is not rendered on screen.

15.3.3.6 Image Properties

The `image` properties are:

`alphadata: scalar | matrix, def. 1`

Transparency is not yet implemented for image objects. `alphadata` is unused.

`alphadatamapping: "direct" | {"none"} | "scaled"`

Transparency is not yet implemented for image objects. `alphadatamapping` is unused.

`beingdeleted: {"off"} | "on"`

`busyaction: "cancel" | {"queue"}`

Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its `interruptible` property set to "off". The `busyaction` property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

`buttondownfcn: string | function handle, def. [] (0x0)`

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`cdata: matrix, def. 64-by-64 double`

`cdatamapping: {"direct"} | "scaled"`

`children (read-only): vector of graphics handles, def. [] (0x1)`

`children` is unused.

`clipping: "off" | {"on"}`

If `clipping` is "on", the image is clipped in its parent axes limits.

`createfcn: string | function handle, def. [] (0x0)`

Callback function executed immediately after image has been created. Function is set by using default property on root object, e.g., `set (groot, "defaultimagecreatefcn", 'disp ("image created!")')`.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`deletefcn: string | function handle, def. [] (0x0)`

Callback function executed immediately before image is deleted.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`handlevisibility: "callback" | "off" | {"on"}`

If `handlevisibility` is "off", the image's handle is not visible in its parent's "children" property.

hittest: "off" | {"on"}

Specify whether image processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [pickableparts property], page 426.

interruptible: "off" | {"on"}

Specify whether this object's callback functions may be interrupted by other callbacks. By default **interruptible** is "on" and callbacks that make use of **drawnow**, **figure**, **waitfor**, **getframe** or **pause** functions are eventually interrupted. See Section 15.4.4 [Callbacks section], page 457.

parent: graphics handle

Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}

Specify whether image will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the image or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "hittest" property will determine how they are processed. See [hittest property], page 425.

selected: {"off"} | "on"

selectionhighlight: "off" | {"on"}

tag: string, def. ""

A user-defined string to label the graphics object.

type (read-only): string

Class name of the graphics object. **type** is always "image"

uicontextmenu: graphics handle, def. [] (0x0)

Graphics handle of the uicontextmenu object that is currently associated to this image object.

userdata: Any Octave data, def. [] (0x0)

User-defined data to associate with the graphics object.

visible: "off" | {"on"}

If **visible** is "off", the image is not rendered on screen.

xdata: two-element vector, def. [1 64]

Two-element vector [**xfirst** **xlast**] specifying the x coordinates of the centers of the first and last columns of the image.

Setting **xdata** to the empty matrix ([])) will restore the default value of [1 **columns**(**image**)].

ydata: two-element vector, def. [1 64]

Two-element vector [**yfirst** **ylast**] specifying the y coordinates of the centers of the first and last rows of the image.

Setting `ydata` to the empty matrix (`[]`) will restore the default value of `[1 rows(image)]`.

15.3.3.7 Patch Properties

The patch properties are:

`alphadatamapping`: "direct" | "none" | {"scaled"}

Transparency is not yet implemented for patch objects. `alphadatamapping` is unused.

`ambientstrength`: scalar, def. 0.30000

Strength of the ambient light. Value between 0.0 and 1.0

`backfacelighting`: "lit" | {"reverselit"} | "unlit"

"lit": The normals are used as is for lighting. "reverselit": The normals are always oriented towards the point of view. "unlit": Faces with normals pointing away from the point of view are unlit.

`beingdeleted`: {"off"} | "on"

`busyaction`: "cancel" | {"queue"}

Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its `interruptible` property set to "off". The `busyaction` property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

`buttondownfcn`: string | function handle, def. [] (0x0)

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`cdata`: scalar | matrix, def. [] (0x0)

Data defining the patch object color. Patch color can be defined for faces or for vertices.

If `cdata` is a scalar index into the current colormap or a RGB triplet, it defines the color of all faces.

If `cdata` is an N-by-1 vector of indices or an N-by-3 (RGB) matrix, it defines the color of each one of the N faces.

If `cdata` is an N-by-M or an N-by-M-by-3 (RGB) matrix, it defines the color at each vertex.

`cdatamapping`: "direct" | {"scaled"}

`children` (read-only): vector of graphics handles, def. [] (0x1)

`children` is unused.

`clipping`: "off" | {"on"}

If `clipping` is "on", the patch is clipped in its parent axes limits.

`createfcn`: string | function handle, def. [] (0x0)

Callback function executed immediately after patch has been created. Function is set by using default property on root object, e.g., `set (groot, "defaultpatchcreatefcn", 'disp ("patch created!")')`.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

deletefcn: string | function handle, def. [] (0x0)

Callback function executed immediately before patch is deleted.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

diffusestrength: scalar, def. 0.60000

Strength of the diffuse reflex. Value between 0.0 (no diffuse reflex) and 1.0 (full diffuse reflex).

displayname: def. ""

Text of the legend entry corresponding to this patch.

edgealpha: scalar | matrix, def. 1

Transparency is not yet implemented for patch objects. **edgealpha** is unused.

edgecolor: def. [0 0 0]

edgelighting: "flat" | "gouraud" | {"none"} | "phong"

When set to a value other than "none", the edges of the object are drawn with light and shadow effects. Supported values are "none" (no lighting effects), "flat" (faceted look) and "gouraud" (linear interpolation of the lighting effects between the vertices). "phong" is deprecated and has the same effect as "gouraud".

facealpha: scalar | "flat" | "interp", def. 1

Transparency level of the faces of the patch object. Only double values are supported at present where a value of 0 means complete transparency and a value of 1 means solid faces without transparency. Setting the property to "flat" or "interp" causes the faces to not being rendered. Additionally, the faces are not sorted from back to front which might lead to unexpected results when rendering layered transparent faces.

facecolor: {colors} | "none" | "flat" | "interp", def. [0 0 0]

facelighting: {"flat"} | "gouraud" | "none" | "phong"

When set to a value other than "none", the faces of the object are drawn with light and shadow effects. Supported values are "none" (no lighting effects), "flat" (faceted look) and "gouraud" (linear interpolation of the lighting effects between the vertices). "phong" is deprecated and has the same effect as "gouraud".

facenormals: def. [] (0x0)

Face normals are used for lighting the edges or faces if the **edgelighting** or **facelighting** properties are set to "flat". Setting **facenormals** also forces the **facenormalsmode** property to be set to "manual"

facenormalsmode: {"auto"} | "manual"

If this property is set to "auto", **facenormals** are automatically calculated if the **edgelighting** or **facelighting** property are set to "flat" and at least one light object is present and visible in the same axes.

faces: vector | matrix, def. [1 2 3]

facevertexalphadata: scalar | matrix, def. [] (0x0)
Transparency is not yet implemented for patch objects. **facevertexalphadata** is unused.

facevertexcdata: def. [] (0x0)

handlevisibility: "callback" | "off" | {"on"}
If **handlevisibility** is "off", the patch's handle is not visible in its parent's "children" property.

hittest: "off" | {"on"}
Specify whether patch processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [pickableparts property], page 429.

interruptible: "off" | {"on"}
Specify whether this object's callback functions may be interrupted by other callbacks. By default **interruptible** is "on" and callbacks that make use of **drawnow**, **figure**, **waitfor**, **getframe** or **pause** functions are eventually interrupted. See Section 15.4.4 [Callbacks section], page 457.

linestyle: {"-"} | {"--"} | {"-."} | {":"} | {"none"}

linewidth: def. 0.50000

marker: "*" | "+" | "." | "<" | ">" | "^" | "d" | "diamond" | "h" | "hexagram" | {"none"} | "o" | "p" | "pentagram" | "s" | "square" | "v" | "x"
See [line marker property], page 421.

markeredgecolor: {"auto"} | "flat" | "none"
See [line markeredgecolor property], page 421.

markerfacecolor: "auto" | "flat" | {"none"}
See [line markerfacecolor property], page 421.

markersize: scalar, def. 6
See [line markersize property], page 421.

parent: graphics handle
Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}
Specify whether patch will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the patch or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "hittest" property will determine how they are processed. See [hittest property], page 429.

selected: {"off"} | "on"

selectionhighlight: "off" | {"on"}

specularcolorreflectance: scalar, def. 1
Reflectance for specular color. Value between 0.0 (color of underlying face) and 1.0 (color of light source).

specularexponent: scalar, def. 10
Exponent for the specular reflex. The lower the value, the more the reflex is spread out.

specularstrength: scalar, def. 0.90000
Strength of the specular reflex. Value between 0.0 (no specular reflex) and 1.0 (full specular reflex).

tag: string, def. ""
A user-defined string to label the graphics object.

type (read-only): string
Class name of the graphics object. **type** is always "patch"

uicontextmenu: graphics handle, def. [] (0x0)
Graphics handle of the uicontextmenu object that is currently associated to this patch object.

userdata: Any Octave data, def. [] (0x0)
User-defined data to associate with the graphics object.

vertexnormals: def. [] (0x0)
Vertex normals are used for lighting the edges or faces if the **edgelighting** or **facelighting** properties are set to "gouraud". Setting **vertexnormals** also forces the **vertexnormalsmode** property to be set to "manual"

vertexnormalsmode: {"auto"} | "manual"
If this property is set to "auto", **vertexnormals** are automatically calculated if the **edgelighting** or **facelighting** property are set to "gouraud" and at least one **light** object is present and visible in the same axes.

vertices: vector | matrix, def. 3-by-2 double

visible: "off" | {"on"}
If **visible** is "off", the patch is not rendered on screen.

xdata: vector | matrix, def. [0; 1; 0]

ydata: vector | matrix, def. [1; 1; 0]

zdata: vector | matrix, def. [] (0x0)

15.3.3.8 Surface Properties

The **surface** properties are:

alphadata: scalar | matrix, def. 1
Transparency is not yet implemented for surface objects. **alphadata** is unused.

alphadatamapping: "direct" | "none" | {"scaled"}
Transparency is not yet implemented for surface objects. **alphadatamapping** is unused.

ambientstrength: scalar, def. 0.30000
 Strength of the ambient light. Value between 0.0 and 1.0

backfacelighting: "lit" | {"reverselit"} | "unlit"
 "lit": The normals are used as is for lighting. "reverselit": The normals are always oriented towards the point of view. "unlit": Faces with normals pointing away from the point of view are unlit.

beingdeleted: {"off"} | "on"

busyaction: "cancel" | {"queue"}
 Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its **interruptible** property set to "off". The **busyaction** property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

buttondownfcn: string | function handle, def. [] (0x0)
 For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

cdata: matrix, def. 3-by-3 double

cdatamapping: "direct" | {"scaled"}

cdatasource: def. ""

children (read-only): vector of graphics handles, def. [] (0x1)
 children is unused.

clipping: "off" | {"on"}
 If **clipping** is "on", the surface is clipped in its parent axes limits.

createfcn: string | function handle, def. [] (0x0)
 Callback function executed immediately after surface has been created. Function is set by using default property on root object, e.g., `set (groot, "defaultsurfacecreatefcn", 'disp ("surface created!")')`.
 For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

deletefcn: string | function handle, def. [] (0x0)
 Callback function executed immediately before surface is deleted.
 For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

diffusestrength: scalar, def. 0.60000
 Strength of the diffuse reflex. Value between 0.0 (no diffuse reflex) and 1.0 (full diffuse reflex).

displayname: def. ""
 Text for the legend entry corresponding to this surface.

edgealpha: scalar, def. 1
 Transparency is not yet implemented for surface objects. **edgealpha** is unused.

`edgecolor`: def. [0 0 0]

`edgelighting`: "flat" | "gouraud" | {"none"} | "phong"

When set to a value other than "none", the edges of the object are drawn with light and shadow effects. Supported values are "none" (no lighting effects), "flat" (faceted look) and "gouraud" (linear interpolation of the lighting effects between the vertices). "phong" is deprecated and has the same effect as "gouraud".

`facealpha`: scalar | "flat" | "interp" | "texturemap", def. 1

Transparency level of the faces of the surface object. Only double values are supported at present where a value of 0 means complete transparency and a value of 1 means solid faces without transparency. Setting the property to "flat", "interp" or "texturemap" causes the faces to not being rendered. Additionally, the faces are not sorted from back to front which might lead to unexpected results when rendering layered transparent faces.

`facecolor`: {"flat"} | "interp" | "none" | "texturemap"

`facelighting`: {"flat"} | "gouraud" | "none" | "phong"

When set to a value other than "none", the faces of the object are drawn with light and shadow effects. Supported values are "none" (no lighting effects), "flat" (faceted look) and "gouraud" (linear interpolation of the lighting effects between the vertices). "phong" is deprecated and has the same effect as "gouraud".

`facenormals`: def. [] (0x0)

Face normals are used for lighting the edges or faces if the `edgelighting` or `facelighting` properties are set to "flat". Setting `facenormals` also forces the `facenormalsmode` property to be set to "manual"

`facenormalsmode`: {"auto"} | "manual"

If this property is set to "auto", `facenormals` are automatically calculated if the `edgelighting` or `facelighting` property are set to "flat" and at least one light object is present and visible in the same axes.

`handlevisibility`: "callback" | "off" | {"on"}

If `handlevisibility` is "off", the surface's handle is not visible in its parent's "children" property.

`hittest`: "off" | {"on"}

Specify whether surface processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [pickableparts property], page 433.

`interruptible`: "off" | {"on"}

Specify whether this object's callback functions may be interrupted by other callbacks. By default `interruptible` is "on" and callbacks that make use of `drawnow`, `figure`, `waitfor`, `getframe` or `pause` functions are eventually interrupted. See Section 15.4.4 [Callbacks section], page 457.

linestyle: `{"-"} | {"--"} | {"-."} | {":"} | {"none"}`
See [Section 15.4.2 \[Line Styles\]](#), page 456.

linewidth: def. 0.50000
See [\[line linewidth property\]](#), page 421.

marker: `"*" | "+" | "." | "<" | ">" | "^" | "d" | "diamond" | "h" | "hexagram" | {"none"} | "o" | "p" | "pentagram" | "s" | "square" | "v" | "x"`
See [Section 15.4.3 \[Marker Styles\]](#), page 457.

markeredgecolor: `{"auto"}` | `"flat"` | `"none"`
See [\[line markeredgecolor property\]](#), page 421.

markerfacecolor: `"auto"` | `"flat"` | `{"none"}`
See [\[line markerfacecolor property\]](#), page 421.

markersize: scalar, def. 6
See [\[line markersize property\]](#), page 421.

meshstyle: `{"both"}` | `"column"` | `"row"`

parent: graphics handle
Handle of the parent graphics object.

pickableparts: `"all" | "none" | {"visible"}`
Specify whether surface will accept mouse clicks. By default, `pickableparts` is `"visible"` and only visible parts of the surface or its children may react to mouse clicks. When `pickableparts` is `"all"` both visible and invisible parts (or children) may react to mouse clicks. When `pickableparts` is `"none"` mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the `"hittest"` property will determine how they are processed. See [\[hittest property\]](#), page 432.

selected: `{"off"}` | `"on"`

selectionhighlight: `"off"` | `{"on"}`

specularcolorreflectance: scalar, def. 1
Reflectance for specular color. Value between 0.0 (color of underlying face) and 1.0 (color of light source).

specularexponent: scalar, def. 10
Exponent for the specular reflex. The lower the value, the more the reflex is spread out.

specularstrength: scalar, def. 0.90000
Strength of the specular reflex. Value between 0.0 (no specular reflex) and 1.0 (full specular reflex).

tag: string, def. `""`
A user-defined string to label the graphics object.

type (read-only): string
Class name of the graphics object. `type` is always `"surface"`

uicontextmenu: graphics handle, def. `[] (0x0)`
Graphics handle of the uicontextmenu object that is currently associated to this surface object.

userdata: Any Octave data, def. [] (0x0)
 User-defined data to associate with the graphics object.

vertexnormals: def. [] (0x0)
 Vertex normals are used for lighting the edges or faces if the **edgelighting** or **facelighting** properties are set to "gouraud". Setting **vertexnormals** also forces the **vertexnormalsmode** property to be set to "manual"

vertexnormalsmode: {"auto"} | "manual"
 If this property is set to "auto", **vertexnormals** are automatically calculated if the **edgelighting** or **facelighting** property are set to "gouraud" and at least one **light** object is present and visible in the same axes.

visible: "off" | {"on"}
 If **visible** is "off", the surface is not rendered on screen.

xdata: matrix, def. [1 2 3]
xdatasource: def. ""
ydata: matrix, def. [1; 2; 3]
ydatasource: def. ""
zdata: matrix, def. 3-by-3 double
zdatasource: def. ""

15.3.3.9 Light Properties

The **light** properties are:

beingdeleted: {"off"} | "on"
busyaction: "cancel" | {"queue"}
 Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its **interruptible** property set to "off". The **busyaction** property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

buttondownfcn: string | function handle, def. [] (0x0)
 For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

children (read-only): vector of graphics handles, def. [] (0x1)
 children is unused.

clipping: "off" | {"on"}
 If **clipping** is "on", the light is clipped in its parent axes limits.

color: colorspec, def. [1 1 1]
 Color of the light source. See [Section 15.4.1 \[colorspec\], page 456](#).

createfcn: string | function handle, def. [] (0x0)
 Callback function executed immediately after light has been created.
 Function is set by using default property on root object, e.g., **set** (**groot**, "defaultlightcreatefcn", 'disp ("light created!")').

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

deletefcn: string | function handle, def. [] (0x0)

Callback function executed immediately before light is deleted.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

handlevisibility: "callback" | "off" | {"on"}

If **handlevisibility** is "off", the light's handle is not visible in its parent's "children" property.

hitTest: "off" | {"on"}

Specify whether light processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [\[pickableparts property\], page 435](#).

interruptible: "off" | {"on"}

Specify whether this object's callback functions may be interrupted by other callbacks. By default **interruptible** is "on" and callbacks that make use of **drawnow**, **figure**, **waitfor**, **getframe** or **pause** functions are eventually interrupted. See [Section 15.4.4 \[Callbacks section\], page 457](#).

parent: graphics handle

Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}

Specify whether light will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the light or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "hitTest" property will determine how they are processed. See [\[hitTest property\], page 435](#).

position: def. [1 0 1]

Position of the light source.

selected: {"off"} | "on"

selectionHighlight: "off" | {"on"}

style: {"infinite"} | "local"

This string defines whether the light emanates from a light source at infinite distance ("infinite") or from a local point source ("local").

tag: string, def. ""

A user-defined string to label the graphics object.

type (read-only): string

Class name of the graphics object. **type** is always "light"

uicontextmenu: graphics handle, def. [] (0x0)
 Graphics handle of the uicontextmenu object that is currently associated to this light object.

userdata: Any Octave data, def. [] (0x0)
 User-defined data to associate with the graphics object.

visible: "off" | {"on"}
 If **visible** is "off", the light is not rendered on screen.

15.3.3.10 Uimenu Properties

The **uimenu** properties are:

accelerator: def. ""

beingdeleted: {"off"} | "on"

busyaction: "cancel" | {"queue"}
 Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its **interruptible** property set to "off". The **busyaction** property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

buttondownfcn: string | function handle, def. [] (0x0)
 buttondownfcn is unused.

callback: def. [] (0x0)

checked: {"off"} | "on"

children (read-only): vector of graphics handles, def. [] (0x1)
 Graphics handles of the **uimenu**'s children.

clipping: "off" | {"on"}
 If **clipping** is "on", the **uimenu** is clipped in its parent axes limits.

createfcn: string | function handle, def. [] (0x0)
 Callback function executed immediately after **uimenu** has been created.
 Function is set by using default property on root object, e.g., **set** (**groot**, "defaultuimenucreatefcn", 'disp ("uimenu created!")').
 For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

deletefcn: string | function handle, def. [] (0x0)
 Callback function executed immediately before **uimenu** is deleted.
 For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

enable: "off" | {"on"}

foregroundcolor: def. [0 0 0]

handlevisibility: "callback" | "off" | {"on"}
 If **handlevisibility** is "off", the **uimenu**'s handle is not visible in its parent's "children" property.

hittest: "off" | {"on"}

Specify whether uimenu processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [pickableparts property], page 437.

interruptible: "off" | {"on"}

Specify whether this object's callback functions may be interrupted by other callbacks. By default **interruptible** is "on" and callbacks that make use of **drawnow**, **figure**, **waitfor**, **getframe** or **pause** functions are eventually interrupted. See Section 15.4.4 [Callbacks section], page 457.

label: def. ""

parent: graphics handle

Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}

Specify whether uimenu will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the uimenu or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "hittest" property will determine how they are processed. See [hittest property], page 436.

position: def. 4

selected: {"off"} | "on"

selectionhighlight: "off" | {"on"}

separator: {"off"} | "on"

tag: string, def. ""

A user-defined string to label the graphics object.

type (read-only): string

Class name of the graphics object. **type** is always "uimenu"

uicontextmenu: graphics handle, def. [] (0x0)

Graphics handle of the uicontextmenu object that is currently associated to this uimenu object.

userdata: Any Octave data, def. [] (0x0)

User-defined data to associate with the graphics object.

visible: "off" | {"on"}

If **visible** is "off", the uimenu is not rendered on screen.

15.3.3.11 Uibuttongroup Properties

The **uibuttongroup** properties are:

backgroundcolor: def. [0.94000 0.94000 0.94000]
beingdeleted: {"off"} | "on"
bordertype: "beveledin" | "beveledout" | {"etchedin"} | "etchedout" | "line" | "none"
borderwidth: def. 1
busyaction: "cancel" | {"queue"}

Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its `interruptible` property set to "off". The `busyaction` property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

buttondownfcn: string | function handle, def. [] (0x0)
For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

children (read-only): vector of graphics handles, def. [] (0x1)
Graphics handles of the uibuttongroup's children.

clipping: "off" | {"on"}
If `clipping` is "on", the uibuttongroup is clipped in its parent axes limits.

createfcn: string | function handle, def. [] (0x0)
Callback function executed immediately after uibuttongroup has been created. Function is set by using default property on root object, e.g., `set (groot, "defaultuibuttongroupcreatefcn", 'disp ("uibuttongroup created!")')`.
For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

deletefcn: string | function handle, def. [] (0x0)
Callback function executed immediately before uibuttongroup is deleted.
For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

fontangle: "italic" | {"normal"} | "oblique"
Control whether the font is italic or normal.

fontname: string, def. "*"
Name of font used for text rendering. When setting this property, the text rendering engine will search for a matching font in your system. If none is found then text is rendered using a default sans serif font (same as the default "*" value).

fontsize: scalar, def. 10
Size of the font used for text rendering. See [\[fontunits property\], page 438](#).

fontunits: "centimeters" | "inches" | "normalized" | "pixels" | {"points"}
Units used to interpret the "fontsize" property.

fontweight: "bold" | "demi" | "light" | {"normal"}
Control the variant of the base font used for text rendering.

foregroundcolor: def. [0 0 0]
handlevisibility: "callback" | "off" | {"on"}
If **handlevisibility** is "off", the uibuttongroup's handle is not visible in its parent's "children" property.

highlightcolor: def. [1 1 1]
hittest: "off" | {"on"}
Specify whether uibuttongroup processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [\[pickableparts property\], page 439](#).

interruptible: "off" | {"on"}
Specify whether this object's callback functions may be interrupted by other callbacks. By default **interruptible** is "on" and callbacks that make use of **drawnow**, **figure**, **waitfor**, **getframe** or **pause** functions are eventually interrupted. See [Section 15.4.4 \[Callbacks section\], page 457](#).

parent: graphics handle
Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}
Specify whether uibuttongroup will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the uibuttongroup or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "hittest" property will determine how they are processed. See [\[hittest property\], page 439](#).

position: def. [0 0 1 1]
resizefcn: string | function handle, def. [] (0x0)
resizefcn is deprecated. Use **sizechangedfcn** instead.

selected: {"off"} | "on"
selectedobject: def. [] (0x0)
selectionchangedfcn: def. [] (0x0)
selectionhighlight: "off" | {"on"}
shadowcolor: def. [0.70000 0.70000 0.70000]
sizechangedfcn: string | function handle, def. [] (0x0)
Callback triggered when the uibuttongroup size is changed.
For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

tag: string, def. ""
A user-defined string to label the graphics object.

```

title: def. ""
titleposition: "centerbottom" | "centertop" | "leftbottom" | {"lefttop"} |
"rightbottom" | "righttop"
type (read-only): string
    Class name of the graphics object. type is always "uibuttongroup"
uicontextmenu: graphics handle, def. [] (0x0)
    Graphics handle of the uicontextmenu object that is currently associated to
    this uibuttongroup object.
units: "centimeters" | "characters" | "inches" | {"normalized"} | "pixels" |
"points"
userdata: Any Octave data, def. [] (0x0)
    User-defined data to associate with the graphics object.
visible: "off" | {"on"}
    If visible is "off", the uibuttongroup is not rendered on screen.

```

15.3.3.12 Uicontextmenu Properties

The uicontextmenu properties are:

```

beingdeleted: {"off"} | "on"
busyaction: "cancel" | {"queue"}
    Define how Octave handles the execution of this object's callback properties
    when it is unable to interrupt another object's executing callback. This is only
    relevant when the currently executing callback object has its interruptible
    property set to "off". The busyaction property of the interrupting callback
    object indicates whether the interrupting callback is queued ("queue" (default))
    or discarded ("cancel"). See Section 15.4.4 \[Callbacks section\], page 457.
buttondownfcn: string | function handle, def. [] (0x0)
    buttondownfcn is unused.
callback: def. [] (0x0)
children (read-only): vector of graphics handles, def. [] (0x1)
    Graphics handles of the uicontextmenu's children.
clipping: "off" | {"on"}
    If clipping is "on", the uicontextmenu is clipped in its parent axes limits.
createfcn: string | function handle, def. [] (0x0)
    Callback function executed immediately after uicontextmenu has been
    created. Function is set by using default property on root object, e.g., set
    (groot, "defaultuicontextmenucreatefcn", 'disp ("uicontextmenu
created!")').
    For information on how to write graphics listener functions see Section 15.4.4
\[Callbacks section\], page 457.
deletefcn: string | function handle, def. [] (0x0)
    Callback function executed immediately before uicontextmenu is deleted.
    For information on how to write graphics listener functions see Section 15.4.4
\[Callbacks section\], page 457.

```

handlevisibility: "callback" | "off" | {"on"}
 If `handlevisibility` is "off", the `uicontextmenu`'s handle is not visible in its parent's "children" property.

hitTest: "off" | {"on"}
 Specify whether `uicontextmenu` processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the `uicontextmenu`, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [pickableparts property], page 441.

interruptible: "off" | {"on"}
 Specify whether this object's callback functions may be interrupted by other callbacks. By default `interruptible` is "on" and callbacks that make use of `drawnow`, `figure`, `waitfor`, `getframe` or `pause` functions are eventually interrupted. See Section 15.4.4 [Callbacks section], page 457.

parent: graphics handle
 Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}
 Specify whether `uicontextmenu` will accept mouse clicks. By default, `pickableparts` is "visible" and only visible parts of the `uicontextmenu` or its children may react to mouse clicks. When `pickableparts` is "all" both visible and invisible parts (or children) may react to mouse clicks. When `pickableparts` is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "hitTest" property will determine how they are processed. See [hitTest property], page 441.

position: def. [0 0]

selected: {"off"} | "on"

selectionhighlight: "off" | {"on"}

tag: string, def. ""
 A user-defined string to label the graphics object.

type (read-only): string
 Class name of the graphics object. `type` is always "uicontextmenu"

uicontextmenu: graphics handle, def. [] (0x0)
 Graphics handle of the `uicontextmenu` object that is currently associated to this `uicontextmenu` object.

userdata: Any Octave data, def. [] (0x0)
 User-defined data to associate with the graphics object.

visible: "off" | {"on"}
 If `visible` is "off", the `uicontextmenu` is not rendered on screen.

15.3.3.13 Uipanel Properties

The `uipanel` properties are:

backgroundcolor: def. [0.94000 0.94000 0.94000]
beingdeleted: {"off"} | "on"
bordertype: "beveledin" | "beveledout" | {"etchedin"} | "etchedout" | "line" | "none"
borderwidth: def. 1
busyaction: "cancel" | {"queue"}

Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its **interruptible** property set to "off". The **busyaction** property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

buttondownfcn: string | function handle, def. [] (0x0)

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

children (read-only): vector of graphics handles, def. [] (0x1)

Graphics handles of the uipanel's children.

clipping: "off" | {"on"}

If **clipping** is "on", the uipanel is clipped in its parent axes limits.

createfcn: string | function handle, def. [] (0x0)

Callback function executed immediately after uipanel has been created. Function is set by using default property on root object, e.g., `set (groot, "defaultuipanelcreatefcn", 'disp ("uipanel created!")')`.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

deletefcn: string | function handle, def. [] (0x0)

Callback function executed immediately before uipanel is deleted.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

fontangle: "italic" | {"normal"} | "oblique"

Control whether the font is italic or normal.

fontname: string, def. "*"

Name of font used for text rendering. When setting this property, the text rendering engine will search for a matching font in your system. If none is found then text is rendered using a default sans serif font (same as the default "*" value).

fontsize: scalar, def. 10

Size of the font used for text rendering. See [\[fontunits property\], page 442](#).

fontunits: "centimeters" | "inches" | "normalized" | "pixels" | {"points"}

Units used to interpret the "fontsize" property.

fontweight: "bold" | "demi" | "light" | {"normal"}

Control the variant of the base font used for text rendering.

foregroundcolor: def. [0 0 0]

handlevisibility: "callback" | "off" | {"on"}
If **handlevisibility** is "off", the uipanel's handle is not visible in its parent's "children" property.

highlightcolor: def. [1 1 1]

hittest: "off" | {"on"}
Specify whether uipanel processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [pickableparts property], page 443.

interruptible: "off" | {"on"}
Specify whether this object's callback functions may be interrupted by other callbacks. By default **interruptible** is "on" and callbacks that make use of **drawnow**, **figure**, **waitfor**, **getframe** or **pause** functions are eventually interrupted. See Section 15.4.4 [Callbacks section], page 457.

parent: graphics handle
Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}
Specify whether uipanel will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the uipanel or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "**hittest**" property will determine how they are processed. See [hittest property], page 443.

position: def. [0 0 1 1]

resizefcn: string | function handle, def. [] (0x0)
resizefcn is deprecated. Use **sizechangedfcn** instead.

selected: {"off"} | "on"

selectionhighlight: "off" | {"on"}

shadowcolor: def. [0.70000 0.70000 0.70000]

sizechangedfcn: string | function handle, def. [] (0x0)
Callback triggered when the uipanel size is changed.
For information on how to write graphics listener functions see Section 15.4.4 [Callbacks section], page 457.

tag: string, def. ""
A user-defined string to label the graphics object.

title: def. ""

titleposition: "centerbottom" | "centertop" | "leftbottom" | {"lefttop"} | "rightbottom" | "righttop"

type (read-only): string
Class name of the graphics object. **type** is always "uipanel"

uicontextmenu: graphics handle, def. [] (0x0)
 Graphics handle of the uicontextmenu object that is currently associated to this uipanel object.

units: "centimeters" | "characters" | "inches" | {"normalized"} | "pixels" | "points"

userdata: Any Octave data, def. [] (0x0)
 User-defined data to associate with the graphics object.

visible: "off" | {"on"}
 If **visible** is "off", the uipanel is not rendered on screen.

15.3.3.14 Uicontrol Properties

The **uicontrol** properties are:

backgroundcolor: def. [0.94000 0.94000 0.94000]
beingdeleted: {"off"} | "on"
busyaction: "cancel" | {"queue"}
 Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its **interruptible** property set to "off". The **busyaction** property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

buttondownfcn: string | function handle, def. [] (0x0)
 For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

callback: def. [] (0x0)
cdata: def. [] (0x0)
children (read-only): vector of graphics handles, def. [] (0x1)
 Graphics handles of the uicontrol's children.

clipping: "off" | {"on"}
 If **clipping** is "on", the uicontrol is clipped in its parent axes limits.

createfcn: string | function handle, def. [] (0x0)
 Callback function executed immediately after uicontrol has been created. Function is set by using default property on root object, e.g., `set (groot, "defaultuicontrolcreatefcn", 'disp ("uicontrol created!")')`.
 For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

deletefcn: string | function handle, def. [] (0x0)
 Callback function executed immediately before uicontrol is deleted.
 For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

enable: "inactive" | "off" | {"on"}
extent (read-only): four-element vector
fontangle: "italic" | {"normal"} | "oblique"
Control whether the font is italic or normal.

fontname: string, def. "*"
Name of font used for text rendering. When setting this property, the text rendering engine will search for a matching font in your system. If none is found then text is rendered using a default sans serif font (same as the default "*" value).

fontsize: scalar, def. 10
Size of the font used for text rendering. See [fontunits property], page 445.

fontunits: "centimeters" | "inches" | "normalized" | "pixels" | {"points"}
Units used to interpret the "fontsize" property.

fontweight: "bold" | "demi" | "light" | {"normal"}
Control the variant of the base font used for text rendering.

foregroundcolor: def. [0 0 0]

handlevisibility: "callback" | "off" | {"on"}
If **handlevisibility** is "off", the uicontrol's handle is not visible in its parent's "children" property.

hittest: "off" | {"on"}
Specify whether uicontrol processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [pickableparts property], page 445.

horizontalalignment: {"center"} | "left" | "right"

interruptible: "off" | {"on"}
Specify whether this object's callback functions may be interrupted by other callbacks. By default **interruptible** is "on" and callbacks that make use of **drawnow**, **figure**, **waitfor**, **getframe** or **pause** functions are eventually interrupted. See Section 15.4.4 [Callbacks section], page 457.

keypressfcn: def. [] (0x0)

listboxtop: def. 1

max: def. 1

min: def. 0

parent: graphics handle
Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}
Specify whether uicontrol will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the uicontrol or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse

clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "hittest" property will determine how they are processed. See [hittest property], page 445.

```
position: def. [0 0 80 30]
selected: {"off"} | "on"
selectionhighlight: "off" | {"on"}
sliderstep: def. [0.010000 0.100000]
string: def. ""
style: "checkbox" | "edit" | "frame" | "listbox" | "popupmenu" | {"pushbutton"}
| "radiobutton" | "slider" | "text" | "togglebutton"
tag: string, def. ""
    A user-defined string to label the graphics object.

tooltipstring: def. ""
type (read-only): string
    Class name of the graphics object. type is always "uicontrol"

uicontextmenu: graphics handle, def. [] (0x0)
    Graphics handle of the uicontextmenu object that is currently associated to
    this uicontrol object.

units: "centimeters" | "characters" | "inches" | "normalized" | {"pixels"} |
"points"
userdata: Any Octave data, def. [] (0x0)
    User-defined data to associate with the graphics object.

value: def. 0
verticalalignment: "bottom" | {"middle"} | "top"
visible: "off" | {"on"}
    If visible is "off", the uicontrol is not rendered on screen.
```

15.3.3.15 Uitable Properties

The uitable properties are:

```
backgroundcolor: def. 2-by-3 double
beingdeleted: {"off"} | "on"
busyaction: "cancel" | {"queue"}
    Define how Octave handles the execution of this object's callback properties
    when it is unable to interrupt another object's executing callback. This is only
    relevant when the currently executing callback object has its interruptible
    property set to "off". The busyaction property of the interrupting callback
    object indicates whether the interrupting callback is queued ("queue" (default))
    or discarded ("cancel"). See Section 15.4.4 [Callbacks section], page 457.
```

buttondownfcn: string | function handle, def. [] (0x0)

For information on how to write graphics listener functions see Section 15.4.4
[Callbacks section], page 457.

celleditcallback: def. [] (0x0)

cellselectioncallback: def. [] (0x0)

children (read-only): vector of graphics handles, def. [] (0x1)
Graphics handles of the uitable's children.

clipping: "off" | {"on"}
If **clipping** is "on", the uitable is clipped in its parent axes limits.

columneditable: def. [] (0x0)

columnformat: def. {}(0x0)

columnname: def. "numbered"

columnwidth: def. "auto"

createfcn: string | function handle, def. [] (0x0)
Callback function executed immediately after uitable has been created.
Function is set by using default property on root object, e.g., `set (groot, "defaultuitablecreatefcn", 'disp ("uitable created!")')`.
For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

data: def. [] (0x0)

deletefcn: string | function handle, def. [] (0x0)
Callback function executed immediately before uitable is deleted.
For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

enable: "off" | {"on"}

extent (read-only): four-element vector

fontangle: "italic" | {"normal"} | "oblique"
Control whether the font is italic or normal.

fontname: string, def. "*"
Name of font used for text rendering. When setting this property, the text rendering engine will search for a matching font in your system. If none is found then text is rendered using a default sans serif font (same as the default "*" value).

fontsize: scalar, def. 10
Size of the font used for text rendering. See [\[fontunits property\], page 447](#).

fontunits: "centimeters" | "inches" | "normalized" | "pixels" | {"points"}
Units used to interpret the "fontsize" property.

fontweight: "bold" | "demi" | "light" | {"normal"}
Control the variant of the base font used for text rendering.

foregroundcolor: def. [0 0 0]

handlevisibility: "callback" | "off" | {"on"}
If **handlevisibility** is "off", the uitable's handle is not visible in its parent's "children" property.

hittest: "off" | {"on"}
Specify whether uitable processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating

the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [pickableparts property], page 448.

interruptible: "off" | {"on"}

Specify whether this object's callback functions may be interrupted by other callbacks. By default **interruptible** is "on" and callbacks that make use of **drawnow**, **figure**, **waitfor**, **getframe** or **pause** functions are eventually interrupted. See Section 15.4.4 [Callbacks section], page 457.

keypressfcn: def. [] (0x0)

keyreleasefcn: def. [] (0x0)

parent: graphics handle

Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}

Specify whether uitable will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the uitable or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "hittest" property will determine how they are processed. See [hittest property], page 447.

position: def. [20 20 300 300]

rearrangeablecolumns: {"off"} | "on"

rownome: def. "numbered"

rowstriping: "off" | {"on"}

selected: {"off"} | "on"

selectionhighlight: "off" | {"on"}

tag: string, def. ""

A user-defined string to label the graphics object.

tooltipstring: def. ""

type (read-only): string

Class name of the graphics object. **type** is always "uitable"

uicontextmenu: graphics handle, def. [] (0x0)

Graphics handle of the uicontextmenu object that is currently associated to this uitable object.

units: "centimeters" | "characters" | "inches" | "normalized" | {"pixels"} | "points"

userdata: Any Octave data, def. [] (0x0)

User-defined data to associate with the graphics object.

visible: "off" | {"on"}

If **visible** is "off", the uitable is not rendered on screen.

15.3.3.16 Uitoolbar Properties

The `uitoolbar` properties are:

`beingdeleted: {"off"} | "on"`

`busyaction: "cancel" | {"queue"}`

Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its `interruptible` property set to `"off"`. The `busyaction` property of the interrupting callback object indicates whether the interrupting callback is queued (`"queue"` (default)) or discarded (`"cancel"`). See [Section 15.4.4 \[Callbacks section\], page 457](#).

`buttondownfcn: string | function handle, def. [] (0x0)`

`buttondownfcn` is unused.

`children (read-only): vector of graphics handles, def. [] (0x1)`

Graphics handles of the `uitoolbar`'s children.

`clipping: "off" | {"on"}`

If `clipping` is `"on"`, the `uitoolbar` is clipped in its parent axes limits.

`createfcn: string | function handle, def. [] (0x0)`

Callback function executed immediately after `uitoolbar` has been created. Function is set by using default property on root object, e.g., `set (groot, "defaultuitoolbarcreatefcn", 'disp ("uitoolbar created!")')`.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`deletefcn: string | function handle, def. [] (0x0)`

Callback function executed immediately before `uitoolbar` is deleted.

For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

`handlevisibility: "callback" | "off" | {"on"}`

If `handlevisibility` is `"off"`, the `uitoolbar`'s handle is not visible in its parent's `"children"` property.

`hittest: "off" | {"on"}`

Specify whether `uitoolbar` processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the `"buttondownfcn"`, showing the `uicontextmenu`, and eventually becoming the root `"currentobject"`. This property is only relevant when the object can accept mouse clicks which is determined by the `"pickableparts"` property. See [\[pickableparts property\], page 450](#).

`interruptible: "off" | {"on"}`

Specify whether this object's callback functions may be interrupted by other callbacks. By default `interruptible` is `"on"` and callbacks that make use of `drawnow`, `figure`, `waitfor`, `getframe` or `pause` functions are eventually interrupted. See [Section 15.4.4 \[Callbacks section\], page 457](#).

parent: graphics handle
 Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}
 Specify whether uitoolbar will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the uitoolbar or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "**hittest**" property will determine how they are processed. See [hittest property], page 449.

selected: {"off"} | "on"
selectionhighlight: "off" | {"on"}
tag: string, def. ""
 A user-defined string to label the graphics object.

type (read-only): string
 Class name of the graphics object. **type** is always "uitoolbar"

uicontextmenu: graphics handle, def. [] (0x0)
 Graphics handle of the uicontextmenu object that is currently associated to this uitoolbar object.

userdata: Any Octave data, def. [] (0x0)
 User-defined data to associate with the graphics object.

visible: "off" | {"on"}
 If **visible** is "off", the uitoolbar is not rendered on screen.

15.3.3.17 Uipushtool Properties

The **uipushtool** properties are:

beingdeleted: {"off"} | "on"
busyaction: "cancel" | {"queue"}
 Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its **interruptible** property set to "off". The **busyaction** property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See Section 15.4.4 [Callbacks section], page 457.

buttondownfcn: string | function handle, def. [] (0x0)
 buttondownfcn is unused.

cdata: def. [] (0x0)
children (read-only): vector of graphics handles, def. [] (0x1)
 Graphics handles of the uipushtool's children.

clickedcallback: def. [] (0x0)
clipping: "off" | {"on"}
 If **clipping** is "on", the uipushtool is clipped in its parent axes limits.

createfcn: string | function handle, def. [] (0x0)
Callback function executed immediately after uipushtool has been created. Function is set by using default property on root object, e.g., `set (groot, "defaultuipushtoolcreatefcn", 'disp ("uipushtool created!")')`.
For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

deletefcn: string | function handle, def. [] (0x0)
Callback function executed immediately before uipushtool is deleted.
For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

enable: "off" | {"on"}
handlevisibility: "callback" | "off" | {"on"}
If `handlevisibility` is "off", the uipushtool's handle is not visible in its parent's "children" property.

hitTest: "off" | {"on"}
Specify whether uipushtool processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [\[pickableparts property\], page 451](#).

interruptible: "off" | {"on"}
Specify whether this object's callback functions may be interrupted by other callbacks. By default `interruptible` is "on" and callbacks that make use of `drawnow`, `figure`, `waitFor`, `getframe` or `pause` functions are eventually interrupted. See [Section 15.4.4 \[Callbacks section\], page 457](#).

parent: graphics handle
Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}
Specify whether uipushtool will accept mouse clicks. By default, `pickableparts` is "visible" and only visible parts of the uipushtool or its children may react to mouse clicks. When `pickableparts` is "all" both visible and invisible parts (or children) may react to mouse clicks. When `pickableparts` is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "hitTest" property will determine how they are processed. See [\[hitTest property\], page 451](#).

selected: {"off"} | "on"
selectionhighlight: "off" | {"on"}
separator: {"off"} | "on"
tag: string, def. ""
A user-defined string to label the graphics object.

tooltipstring: def. ""
type (read-only): string
 Class name of the graphics object. **type** is always "uipushtool"
uicontextmenu: graphics handle, def. [] (0x0)
 Graphics handle of the uicontextmenu object that is currently associated to this uipushtool object.
userdata: Any Octave data, def. [] (0x0)
 User-defined data to associate with the graphics object.
visible: "off" | {"on"}
 If **visible** is "off", the uipushtool is not rendered on screen.

15.3.3.18 Uitoggletool Properties

The **uitoggletool** properties are:

beingdeleted: {"off"} | "on"
busyaction: "cancel" | {"queue"}
 Define how Octave handles the execution of this object's callback properties when it is unable to interrupt another object's executing callback. This is only relevant when the currently executing callback object has its **interruptible** property set to "off". The **busyaction** property of the interrupting callback object indicates whether the interrupting callback is queued ("queue" (default)) or discarded ("cancel"). See [Section 15.4.4 \[Callbacks section\], page 457](#).

buttondownfcn: string | function handle, def. [] (0x0)
 buttondownfcn is unused.

cdata: def. [] (0x0)

children (read-only): vector of graphics handles, def. [] (0x1)
 Graphics handles of the uitoggletool's children.

clickedcallback: def. [] (0x0)

clipping: "off" | {"on"}
 If **clipping** is "on", the uitoggletool is clipped in its parent axes limits.

createfcn: string | function handle, def. [] (0x0)
 Callback function executed immediately after uitoggletool has been created. Function is set by using default property on root object, e.g., `set (groot, "defaultuitoggletoolcreatefcn", 'disp ("uitoggletool created!")')`. For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

deletefcn: string | function handle, def. [] (0x0)
 Callback function executed immediately before uitoggletool is deleted. For information on how to write graphics listener functions see [Section 15.4.4 \[Callbacks section\], page 457](#).

enable: "off" | {"on"}
handlevisibility: "callback" | "off" | {"on"}
 If **handlevisibility** is "off", the uitoggletool's handle is not visible in its parent's "children" property.

hittest: "off" | {"on"}

Specify whether uitoggletool processes mouse events or passes them to ancestors of the object. When enabled, the object may respond to mouse clicks by evaluating the "buttondownfcn", showing the uicontextmenu, and eventually becoming the root "currentobject". This property is only relevant when the object can accept mouse clicks which is determined by the "pickableparts" property. See [pickableparts property], page 453.

interruptible: "off" | {"on"}

Specify whether this object's callback functions may be interrupted by other callbacks. By default **interruptible** is "on" and callbacks that make use of **drawnow**, **figure**, **waitfor**, **getframe** or **pause** functions are eventually interrupted. See Section 15.4.4 [Callbacks section], page 457.

offcallback: def. [] (0x0)

oncallback: def. [] (0x0)

parent: graphics handle

Handle of the parent graphics object.

pickableparts: "all" | "none" | {"visible"}

Specify whether uitoggletool will accept mouse clicks. By default, **pickableparts** is "visible" and only visible parts of the uitoggletool or its children may react to mouse clicks. When **pickableparts** is "all" both visible and invisible parts (or children) may react to mouse clicks. When **pickableparts** is "none" mouse clicks on the object are ignored and transmitted to any objects underneath this one. When an object is configured to accept mouse clicks the "hittest" property will determine how they are processed. See [hittest property], page 452.

selected: {"off"} | "on"

selectionhighlight: "off" | {"on"}

separator: {"off"} | "on"

state: {"off"} | "on"

tag: string, def. ""

A user-defined string to label the graphics object.

tooltipstring: def. ""

type (read-only): string

Class name of the graphics object. **type** is always "uitoggletool"

uicontextmenu: graphics handle, def. [] (0x0)

Graphics handle of the uicontextmenu object that is currently associated to this uitoggletool object.

userdata: Any Octave data, def. [] (0x0)

User-defined data to associate with the graphics object.

visible: "off" | {"on"}

If **visible** is "off", the uitoggletool is not rendered on screen.

15.3.4 Searching Properties

h = **findobj** ()

```

h = findobj (prop_name, prop_value, ...)
h = findobj (prop_name, prop_value, "-logical_op", prop_name,
             prop_value)
h = findobj ("-property", prop_name)
h = findobj ("-regexp", prop_name, pattern)
h = findobj (hlist, ...)
h = findobj (hlist, "flat", ...)
h = findobj (hlist, "-depth", d, ...)

```

Find graphics objects with specified properties.

When called without arguments, return all graphic objects beginning with the root object (0) and including all of its descendants.

The simplest form for narrowing the results is

```
findobj (prop_name, prop_value)
```

which returns the handles of all objects which have a property named *prop_name* that has the value *prop_value*. If multiple property/value pairs are specified then only objects meeting all of the conditions (equivalent to `-and`) are returned.

The search can be limited to a particular set of objects and their descendants, by passing a handle or set of handles *hlist* as the first argument.

The depth of the object hierarchy to search can be limited with the `"-depth"` argument. An example of searching through only three generations of children is:

```
findobj (hlist, "-depth", 3, prop_name, prop_value)
```

Specifying a depth *d* of 0 limits the search to the set of objects passed in *hlist*. A depth of 0 is also equivalent to the `"flat"` argument. The default depth value is `Inf` which includes all descendants.

A specified logical operator may be used between *prop_name*, *prop_value* pairs. The supported logical operators are: `"-and"`, `"-or"`, `"-xor"`, `"-not"`. Example code to locate all figure and axes objects is

```
findobj ("type", "figure", "-or", "type", "axes")
```

Objects may also be matched by comparing a regular expression to the property values, where property values that match `regexp` (*prop_value*, *pattern*) are returned.

Finally, objects which have a property name can be found with the `"-property"` option. For example, code to locate objects with a `"meshstyle"` property is

```
findobj ("-property", "meshstyle")
```

Implementation Note: The search only includes objects with visible handles (`HandleVisibility = "on"`). See [\[findall\]](#), page 454, to search for all objects including hidden ones.

See also: [\[findall\]](#), page 454, [\[allchild\]](#), page 404, [\[get\]](#), page 402, [\[set\]](#), page 402.

```

h = findall ()
h = findall (prop_name, prop_value, ...)
h = findall (prop_name, prop_value, "-logical_op", prop_name,
             prop_value)
h = findall ("-property", prop_name)
h = findall ("-regexp", prop_name, pattern)

```

```
h = findall (hlist, ...)
h = findall (hlist, "flat", ...)
h = findall (hlist, "-depth", d, ...)
```

Find graphics object, including hidden ones, with specified properties.

The return value *h* is a list of handles to the found graphic objects.

`findall` performs the same search as `findobj`, but it includes hidden objects (HandleVisibility = "off"). For full documentation, see [`findobj`], page 453.

See also: [`findobj`], page 453, [`allchild`], page 404, [`get`], page 402, [`set`], page 402.

15.3.5 Managing Default Properties

Object properties have two classes of default values, *factory defaults* (the initial values) and *user-defined defaults*, which may override the factory defaults.

Although default values may be set for any object, they are set in parent objects and apply to child objects, of the specified object type. For example, setting the default `color` property of `line` objects to "green", for the `root` object, will result in all `line` objects inheriting the `color` "green" as the default value.

```
set (groot, "defaultlinecolor", "green");
```

sets the default line color for all objects. The rule for constructing the property name to set a default value is

```
default + object-type + property-name
```

This rule can lead to some strange looking names, for example `defaultlinelinenwidth`" specifies the default `linenwidth` property for `line` objects.

The example above used the root figure object so the default property value will apply to all `line` objects. However, default values are hierarchical, so defaults set in a figure objects override those set in the root figure object. Likewise, defaults set in axes objects override those set in figure or root figure objects. For example,

```
subplot (2, 1, 1);
set (groot, "defaultlinecolor", "red");
set (1, "defaultlinecolor", "green");
set (gca (), "defaultlinecolor", "blue");
line (1:10, rand (1, 10));
subplot (2, 1, 2);
line (1:10, rand (1, 10));
figure (2)
line (1:10, rand (1, 10));
```

produces two figures. The line in first subplot window of the first figure is blue because it inherits its color from its parent axes object. The line in the second subplot window of the first figure is green because it inherits its color from its parent figure object. The line in the second figure window is red because it inherits its color from the global root figure parent object.

To remove a user-defined default setting, set the `default` property to the value "remove". For example,

```
set (gca (), "defaultlinecolor", "remove");
```

removes the user-defined default line color setting from the current axes object. To quickly remove all user-defined defaults use the `reset` function.

By default, high level plotting functions such as `plot` reset and redefine axes properties independently from the defaults. An example of such property is the axes `box` property: it is set on by high level 2-D graphics functions regardless of the property "defaultaxesbox". Use the `hold` function to prevent this behavior:

```
set (groot, "defaultaxesbox", "off");
subplot (2, 1, 1);
plot (1:10)
title ("Box is on anyway")
subplot (2, 1, 2);
hold on
plot (1:10)
title ("Box is off")
```

`reset (h)`

Reset the properties of the graphic object *h* to their default values.

For figures, the properties "position", "units", "windowstyle", and "paperunits" are not affected. For axes, the properties "position" and "units" are not affected.

The input *h* may also be a vector of graphic handles in which case each individual object will be reset.

See also: [\[cla\]](#), page 377, [\[clf\]](#), page 376, [\[newplot\]](#), page 375.

Getting the "default" property of an object returns a list of user-defined defaults set for the object. For example,

```
get (gca (), "default");
```

returns a list of user-defined default values for the current axes object.

Factory default values are stored in the root figure object. The command

```
get (groot, "factory");
```

returns a list of factory defaults.

15.4 Advanced Plotting

15.4.1 Colors

Colors may be specified as RGB triplets with values ranging from zero to one, or by name. Recognized color names include "blue", "black", "cyan", "green", "magenta", "red", "white", and "yellow".

15.4.2 Line Styles

Line styles are specified by the following properties:

`linestyle`

May be one of

"--"	Solid line. [default]
------	-----------------------

"--"	Dashed line.
Dotted line.	
"-. "	A dash-dot line.
"none"	No line. Points will still be marked using the current Marker Style.

linewidth

A number specifying the width of the line. The default is 1. A value of 2 is twice as wide as the default, etc.

15.4.3 Marker Styles

Marker styles are specified by the following properties:

marker A character indicating a plot marker to be placed at each data point, or "none", meaning no markers should be displayed.

markeredgecolor

The color of the edge around the marker, or "auto", meaning that the edge color is the same as the face color. See [Section 15.4.1 \[Colors\]](#), page 456.

markerfacecolor

The color of the marker, or "none" to indicate that the marker should not be filled. See [Section 15.4.1 \[Colors\]](#), page 456.

markersize

A number specifying the size of the marker. The default is 1. A value of 2 is twice as large as the default, etc.

The `colstyle` function will parse a plot-style specification and will return the color, line, and marker values that would result.

[style, color, marker, msg] = colstyle (style)

Parse the line specification `style` and return the line style, color, and markers given.

In the case of an error, the string `msg` will return the text of the error.

15.4.4 Callbacks

Callback functions can be associated with graphics objects and triggered after certain events occur. The basic structure of all callback function is

```
function mycallback (hsrc, evt)
  ...
endfunction
```

where `hsrc` is a handle to the source of the callback, and `evt` gives some event specific data.

The function can be provided as a function handle to a plain Octave function, as an anonymous function, or as a string representing an Octave command. The latter syntax is not recommended since syntax errors will only occur when the string is evaluated. See [Section 11.11 \[Function Handles section\]](#), page 214.

This can then be associated with an object either at the object's creation, or later with the `set` function. For example,

```
plot (x, "DeleteFcn", @(h, e) disp ("Window Deleted"))
```

where at the moment that the plot is deleted, the message "Window Deleted" will be displayed.

Additional user arguments can be passed to callback functions, and will be passed after the two default arguments. For example:

```
plot (x, "DeleteFcn", {@mycallback, "1"})
...
function mycallback (h, evt, arg1)
    fprintf ("Closing plot %d\n", arg1);
endfunction
```

Caution: The second argument in callback functions—`evt`—is only partially implemented in the Qt graphics toolkit:

- Mouse click events: `evt` is a class `double` value: 1 for left, 2 for middle, and 3 for right click.
- Key press events: `evt` is a structure with fields `Key` (string), `Character` (string), and `Modifier` (cell array of strings).
- Other events: `evt` is a class `double` empty matrix.

The basic callback functions that are available for all graphics objects are

- `CreateFcn`: called at the moment of the objects creation. It is not called if the object is altered in any way, and so it only makes sense to define this callback in the function call that defines the object. Callbacks that are added to `CreateFcn` later with the `set` function will never be executed.
- `DeleteFcn`: called at the moment an object is deleted.
- `ButtonDownFcn`: called if a mouse button is pressed while the pointer is over this object. Note, that the gnuplot interface does not implement this callback.

By default callback functions are queued (they are executed one after the other in the event queue) unless the `drawnow`, `figure`, `waitfor`, `getframe`, or `pause` functions are used. If an executing callback invokes one of those functions, it causes Octave to flush the event queue, which results in the executing callback being interrupted.

It is possible to specify that an object's callbacks should not be interrupted by setting the object's `interruptible` property to "off". In this case, Octave decides what to do based on the `busyaction` property of the `interrupting` callback object:

<code>queue</code> (the default)	The interrupting callback is executed after the executing callback has returned.
<code>cancel</code>	The interrupting callback is discarded.

The `interruptible` property has no effect when the interrupting callback is a `deletefcn`, or a figure `resizefcn` or `closerequestfcn`. Those callbacks always interrupt the executing callback.

The handle to the object that holds the callback being executed can be obtained with the `gcbo` function. The handle to the ancestor figure of this object may be obtained using the `gcbf` function.

```
h = gcbo ()
[h, fig] = gcbo ()
```

Return a handle to the object whose callback is currently executing.

If no callback is executing, this function returns the empty matrix. This handle is obtained from the root object property "CallbackObject".

When called with a second output argument, return the handle of the figure containing the object whose callback is currently executing. If no callback is executing the second output is also set to the empty matrix.

See also: [gcbf], page 459, [gco], page 401, [gca], page 401, [gcf], page 400, [get], page 402, [set], page 402.

```
fig = gcbf ()
```

Return a handle to the figure containing the object whose callback is currently executing.

If no callback is executing, this function returns the empty matrix. The handle returned by this function is the same as the second output argument of gcbo.

See also: [gcbo], page 458, [gcf], page 400, [gco], page 401, [gca], page 401, [get], page 402, [set], page 402.

Callbacks can equally be added to properties with the `addlistener` function described below.

15.4.5 Application-defined Data

Octave has a provision for attaching application-defined data to a graphics handle. The data can be anything which is meaningful to the application, and will be completely ignored by Octave.

```
setappdata (h, name, value)
setappdata (h, name1, value1, name2, value3, ...)
setappdata (h, {name1, name2, ...}, {value1, value2, ...})
```

Set the application data *name* to *value* for the graphics object with handle *h*.

h may also be a vector of graphics handles. If the application data with the specified *name* does not exist, it is created.

Multiple *name*/*value* pairs can be specified. Alternatively, a cell array of *names* and a corresponding cell array of *values* can be specified.

See also: [getappdata], page 459, [isappdata], page 460, [rmappdata], page 460, [guidata], page 879, [get], page 402, [set], page 402, [getpref], page 881, [setpref], page 882.

```
value = getappdata (h, name)
appdata = getappdata (h)
```

Return the *value* of the application data *name* for the graphics object with handle *h*.

h may also be a vector of graphics handles. If no second argument *name* is given then `getappdata` returns a structure, *appdata*, whose fields correspond to the appdata properties.

See also: [setappdata], page 459, [isappdata], page 460, [rmappdata], page 460, [guidata], page 879, [get], page 402, [set], page 402, [getpref], page 881, [setpref], page 882.

```
rmappdata (h, name)
rmappdata (h, name1, name2, ...)
```

Delete the application data *name* from the graphics object with handle *h*.

h may also be a vector of graphics handles. Multiple application data names may be supplied to delete several properties at once.

See also: [setappdata], page 459, [getappdata], page 459, [isappdata], page 460.

```
valid = isappdata (h, name)
```

Return true if the named application data, *name*, exists for the graphics object with handle *h*.

h may also be a vector of graphics handles.

See also: [getappdata], page 459, [setappdata], page 459, [rmappdata], page 460, [guidata], page 879, [get], page 402, [set], page 402, [getpref], page 881, [setpref], page 882.

15.4.6 Object Groups

A number of Octave high level plot functions return groups of other graphics objects or they return graphics objects that have their properties linked in such a way that changes to one of the properties results in changes in the others. A graphic object that groups other objects is an `hggroup`

```
hggroup ()
hggroup (hax)
hggroup (..., property, value, ...)
h = hggroup (...)
```

Create handle graphics group object with axes parent *hax*.

If no parent is specified, the group is created in the current axes.

Multiple property/value pairs may be specified for the `hggroup`, but they must appear in pairs.

The optional return value *h* is a graphics handle to the created `hggroup` object.

Programming Note: An `hggroup` is a way to group base graphics objects such as line objects or patch objects into a single unit which can react appropriately. For example, the individual lines of a contour plot are collected into a single `hggroup` so that they can be made visible/invisible with a single command, `set (hg_handle, "visible", "off")`.

See also: [addproperty], page 461, [addlistener], page 462.

For example a simple use of a `hggroup` might be

```

x = 0:0.1:10;
hg = hggroup ();
plot (x, sin (x), "color", [1, 0, 0], "parent", hg);
hold on
plot (x, cos (x), "color", [0, 1, 0], "parent", hg);
set (hg, "visible", "off");

```

which groups the two plots into a single object and controls their visibility directly. The default properties of an `hggroup` are the same as the set of common properties for the other graphics objects. Additional properties can be added with the `addproperty` function.

**addproperty (*name*, *h*, *type*)
addproperty (*name*, *h*, *type*, *arg*, ...)**

Create a new property named *name* in graphics object *h*.

type determines the type of the property to create. *args* usually contains the default value of the property, but additional arguments might be given, depending on the type of the property.

The supported property types are:

string	A string property. <i>arg</i> contains the default string value.
any	An un-typed property. This kind of property can hold any octave value. <i>args</i> contains the default value.
radio	A string property with a limited set of accepted values. The first argument must be a string with all accepted values separated by a vertical bar (' '). The default value can be marked by enclosing it with a '{' '}' pair. The default value may also be given as an optional second string argument.
boolean	A boolean property. This property type is equivalent to a radio property with "on off" as accepted values. <i>arg</i> contains the default property value.
double	A scalar double property. <i>arg</i> contains the default value.
handle	A handle property. This kind of property holds the handle of a graphics object. <i>arg</i> contains the default handle value. When no default value is given, the property is initialized to the empty matrix.
data	A data (matrix) property. <i>arg</i> contains the default data value. When no default value is given, the data is initialized to the empty matrix.
color	A color property. <i>arg</i> contains the default color value. When no default color is given, the property is set to black. An optional second string argument may be given to specify an additional set of accepted string values (like a radio property).

type may also be the concatenation of a core object type and a valid property name for that object type. The property created then has the same characteristics as the referenced property (type, possible values, hidden state...). This allows one to clone an existing property into the graphics object *h*.

Examples:

```
addproperty ("my_property", gcf, "string", "a string value");
addproperty ("my_radio", gcf, "radio", "val_1|val_2|{val_3}");
addproperty ("my_style", gcf, "linelinenstyle", "--");
```

See also: [\[addlistener\]](#), page 462, [\[hggroup\]](#), page 460.

Once a property is added to an `hggroup`, it is not linked to any other property of either the children of the group, or any other graphics object. Add so to control the way in which this newly added property is used, the `addlistener` function is used to define a callback function that is executed when the property is altered.

`addlistener (h, prop, fcn)`

Register `fcn` as listener for the property `prop` of the graphics object `h`.

Property listeners are executed (in order of registration) when the property is set. The new value is already available when the listeners are executed.

`prop` must be a string naming a valid property in `h`.

`fcn` can be a function handle, a string or a cell array whose first element is a function handle. If `fcn` is a function handle, the corresponding function should accept at least 2 arguments, that will be set to the object handle and the empty matrix respectively. If `fcn` is a string, it must be any valid octave expression. If `fcn` is a cell array, the first element must be a function handle with the same signature as described above. The next elements of the cell array are passed as additional arguments to the function.

Example:

```
function my_listener (h, dummy, p1)
    fprintf ("my_listener called with p1=%s\n", p1);
endfunction

addlistener (gcf, "position", {@my_listener, "my string"})
```

See also: [\[dellistener\]](#), page 462, [\[addproperty\]](#), page 461, [\[hggroup\]](#), page 460.

`dellistener (h, prop, fcn)`

Remove the registration of `fcn` as a listener for the property `prop` of the graphics object `h`.

The function `fcn` must be the same variable (not just the same value), as was passed to the original call to `addlistener`.

If `fcn` is not defined then all listener functions of `prop` are removed.

Example:

```
function my_listener (h, dummy, p1)
    fprintf ("my_listener called with p1=%s\n", p1);
endfunction

c = {@my_listener, "my string"};
addlistener (gcf, "position", c);
dellistener (gcf, "position", c);
```

See also: [\[addlistener\]](#), page 462.

An example of the use of these two functions might be

```
x = 0:0.1:10;
hg = hggroup ();
h = plot (x, sin (x), "color", [1, 0, 0], "parent", hg);
addproperty ("linestyle", hg, "linelinenstyle", get (h, "linestyle"));
addlistener (hg, "linestyle", @update_props);
hold on
plot (x, cos (x), "color", [0, 1, 0], "parent", hg);

function update_props (h, d)
    set (get (h, "children"), "linestyle", get (h, "linestyle"));
endfunction
```

that adds a `linestyle` property to the `hggroup` and propagating any changes its value to the children of the group. The `linkprop` function can be used to simplify the above to be

```
x = 0:0.1:10;
hg = hggroup ();
h1 = plot (x, sin (x), "color", [1, 0, 0], "parent", hg);
addproperty ("linestyle", hg, "linelinenstyle", get (h, "linestyle"));
hold on
h2 = plot (x, cos (x), "color", [0, 1, 0], "parent", hg);
hlink = linkprop ([hg, h1, h2], "color");

hlink = linkprop (h, "prop")
hlink = linkprop (h, {"prop1", "prop2", ...})
```

Link graphic object properties, such that a change in one is propagated to the others.

The input `h` is a vector of graphic handles to link.

`prop` may be a string when linking a single property, or a cell array of strings for multiple properties. During the linking process all properties in `prop` will initially be set to the values that exist on the first object in the list `h`.

The function returns `hlink` which is a special object describing the link. As long as the reference `hlink` exists, the link between graphic objects will be active. This means that `hlink` must be preserved in a workspace variable, a global variable, or otherwise stored using a function such as `setappdata` or `guidata`. To unlink properties, execute `clear hlink`.

An example of the use of `linkprop` is

```
x = 0:0.1:10;
subplot (1,2,1);
h1 = plot (x, sin (x));
subplot (1,2,2);
h2 = plot (x, cos (x));
hlink = linkprop ([h1, h2], {"color", "linestyle"});
set (h1, "color", "green");
set (h2, "linestyle", "--");
```

See also: [\[linkaxes\]](#), page 464, [\[addlistener\]](#), page 462.

linkaxes (hax)

linkaxes (hax, optstr)

Link the axis limits of 2-D plots such that a change in one is propagated to the others.

The axes handles to be linked are passed as the first argument *hax*.

The optional second argument is a string which defines which axis limits will be linked.

The possible values for *optstr* are:

"x" Link x-axes

"y" Link y-axes

"xy" (default)

Link both axes

"off" Turn off linking

If unspecified the default is to link both X and Y axes.

When linking, the limits from the first axes in *hax* are applied to the other axes in the list. Subsequent changes to any one of the axes will be propagated to the others.

See also: [\[linkprop\]](#), page 463, [\[addproperty\]](#), page 461.

These capabilities are used in a number of basic graphics objects. The **hggroup** objects created by the functions of Octave contain one or more graphics object and are used to:

- group together multiple graphics objects,
- create linked properties between different graphics objects, and
- to hide the nominal user data, from the actual data of the objects.

For example the **stem** function creates a stem series where each **hggroup** of the stem series contains two line objects representing the body and head of the stem. The **ydata** property of the **hggroup** of the stem series represents the head of the stem, whereas the body of the stem is between the baseline and this value. For example

```

h = stem (1:4)
get (h, "xdata")
⇒ [ 1   2   3   4] '
get (get (h, "children")(1), "xdata")
⇒ [ 1   1 NaN   2   2 NaN   3   3 NaN   4   4 NaN] '

```

shows the difference between the **xdata** of the **hggroup** of a stem series object and the underlying line.

The basic properties of such group objects is that they consist of one or more linked **hggroup**, and that changes in certain properties of these groups are propagated to other members of the group. Whereas, certain properties of the members of the group only apply to the current member.

In addition the members of the group can also be linked to other graphics objects through callback functions. For example the baseline of the **bar** or **stem** functions is a line object, whose length and position are automatically adjusted, based on changes to the corresponding **hggroup** elements.

15.4.6.1 Data Sources in Object Groups

All of the group objects contain data source parameters. There are string parameters that contain an expression that is evaluated to update the relevant data property of the group when the `refreshdata` function is called.

```
refreshdata ()
refreshdata (h)
refreshdata (h, workspace)
```

Evaluate any ‘datasource’ properties of the current figure and update the plot if the corresponding data has changed.

If the first argument *h* is a list of graphic handles, then operate on these objects rather than the current figure returned by `gcf`.

The optional second argument *workspace* can take the following values:

- “base” Evaluate the datasource properties in the base workspace. (default).
- “caller” Evaluate the datasource properties in the workspace of the function that called `refreshdata`.

An example of the use of `refreshdata` is:

```
x = 0:0.1:10;
y = sin (x);
plot (x, y, "ydatasource", "y");
for i = 1 : 100
    pause (0.1);
    y = sin (x + 0.1*i);
    refreshdata ();
endfor
```

15.4.6.2 Area Series

Area series objects are created by the `area` function. Each of the `hggroup` elements contains a single patch object. The properties of the area series are

`basevalue`

The value where the base of the area plot is drawn.

`linewidth`

`linestyle`

The line width and style of the edge of the patch objects making up the areas.

See [Section 15.4.2 \[Line Styles\], page 456](#).

`edgecolor`

`facecolor`

The line and fill color of the patch objects making up the areas. See [Section 15.4.1 \[Colors\], page 456](#).

`xdata`

`ydata` The x and y coordinates of the original columns of the data passed to `area` prior to the cumulative summation used in the `area` function.

```
xdatasource
ydatasource
    Data source variables.
```

15.4.6.3 Bar Series

Bar series objects are created by the `bar` or `barm` functions. Each `hggroup` element contains a single patch object. The properties of the bar series are

```
showbaseline
baseline
basevalue
```

The property `showbaseline` flags whether the baseline of the bar series is displayed (default is "on"). The handle of the graphics object representing the baseline is given by the `baseline` property and the y-value of the baseline by the `basevalue` property.

Changes to any of these properties are propagated to the other members of the bar series and to the baseline itself. Equally, changes in the properties of the base line itself are propagated to the members of the corresponding bar series.

```
barwidth
barlayout
horizontal
```

The property `barwidth` is the width of the bar corresponding to the `width` variable passed to `bar` or `barm`. Whether the bar series is "grouped" or "stacked" is determined by the `barlayout` property and whether the bars are horizontal or vertical by the `horizontal` property.

Changes to any of these property are propagated to the other members of the bar series.

```
linewidth
linestyle
```

The line width and style of the edge of the patch objects making up the bars. See [Section 15.4.2 \[Line Styles\], page 456](#).

```
edgecolor
facecolor
```

The line and fill color of the patch objects making up the bars. See [Section 15.4.1 \[Colors\], page 456](#).

`xdata` The nominal x positions of the bars. Changes in this property are propagated to the other members of the bar series.

`ydata` The y value of the bars in the `hggroup`.

```
xdatasource
ydatasource
```

Data source variables.

15.4.6.4 Contour Groups

Contour group objects are created by the `contour`, `contourf`, and `contour3` functions. They are also one of the handles returned by the `surf` and `mesh` functions. The properties of the contour group are

`contourmatrix`

A read only property that contains the data return by `contourc` used to create the contours of the plot.

`fill`

A radio property that can have the values "on" or "off" that flags whether the contours to plot are to be filled.

`zlevelmode`

`zlevel`

The radio property `zlevelmode` can have the values "none", "auto", or "manual". When its value is "none" there is no z component to the plotted contours. When its value is "auto" the z value of the plotted contours is at the same value as the contour itself. If the value is "manual", then the z value at which to plot the contour is determined by the `zlevel` property.

`levellistmode`

`levellist`

`levelstepmode`

`levelstep`

If `levellistmode` is "manual", then the levels at which to plot the contours is determined by `levellist`. If `levellistmode` is set to "auto", then the distance between contours is determined by `levelstep`. If both `levellistmode` and `levelstepmode` are set to "auto", then there are assumed to be 10 equal spaced contours.

`textlistmode`

`textlist`

`textstepmode`

`textstep`

If `textlistmode` is "manual", then the labeled contours is determined by `textlist`. If `textlistmode` is set to "auto", then the distance between labeled contours is determined by `textstep`. If both `textlistmode` and `textstepmode` are set to "auto", then there are assumed to be 10 equal spaced labeled contours.

`showtext` Flag whether the contour labels are shown or not.

`labelspacing`

The distance between labels on a single contour in points.

`linewidth`

`linestyle`

`linecolor`

The properties of the contour lines. The properties `linewidth` and `linestyle` are similar to the corresponding properties for lines. The property `linecolor` is a color property (see [Section 15.4.1 \[Colors\]](#), page 456), that can also have the values of "none" or "auto". If `linecolor` is "none", then no contour line is drawn. If `linecolor` is "auto" then the line color is determined by the colormap.

```
xdata
ydata
zdata      The original x, y, and z data of the contour lines.

xdatasource
ydatasource
zdatasource
Data source variables.
```

15.4.6.5 Error Bar Series

Error bar series are created by the `errorbar` function. Each `hggroup` element contains two line objects representing the data and the errorbars separately. The properties of the error bar series are

```
color      The RGB color or color name of the line objects of the error bars. See
          Section 15.4.1 [Colors], page 456.

linewidth
linestyle
The line width and style of the line objects of the error bars. See Section 15.4.2
[Line Styles], page 456.

marker
markeredgecolor
markerfacecolor
markersize
The line and fill color of the markers on the error bars. See Section 15.4.1
[Colors], page 456.

xdata
ydata
ldata
udata
xldata
xudata      The original x, y, l, u, xl, xu data of the error bars.

xdatasource
ydatasource
ldatasource
udatasource
xldatasource
xudatasource
Data source variables.
```

15.4.6.6 Line Series

Line series objects are created by the `plot` and `plot3` functions and are of the type `line`. The properties of the line series with the ability to add data sources.

```
color      The RGB color or color name of the line objects. See Section 15.4.1 [Colors],
          page 456.
```

linewidth
linestyle

The line width and style of the line objects. See [Section 15.4.2 \[Line Styles\]](#), page 456.

marker

markeredgecolor
markerfacecolor
markersize

The line and fill color of the markers. See [Section 15.4.1 \[Colors\]](#), page 456.

xdata

ydata

zdata The original x, y and z data.

xdatasource
ydatasource
zdatasource

Data source variables.

15.4.6.7 Quiver Group

Quiver series objects are created by the `quiver` or `quiver3` functions. Each `hggroup` element of the series contains three line objects as children representing the body and head of the arrow, together with a marker as the point of origin of the arrows. The properties of the quiver series are

autoscale

autoscalefactor

Flag whether the length of the arrows is scaled or defined directly from the `u`, `v` and `w` data. If the arrow length is flagged as being scaled by the `autoscale` property, then the length of the autoscaled arrow is controlled by the `autoscalefactor`.

maxheadsize

This property controls the size of the head of the arrows in the quiver series. The default value is 0.2.

showarrowhead

Flag whether the arrow heads are displayed in the quiver plot.

color

The RGB color or color name of the line objects of the quiver. See [Section 15.4.1 \[Colors\]](#), page 456.

linewidth

linestyle

The line width and style of the line objects of the quiver. See [Section 15.4.2 \[Line Styles\]](#), page 456.

marker

markerfacecolor
markersize

The line and fill color of the marker objects at the origin of the arrows. See [Section 15.4.1 \[Colors\]](#), page 456.

```

xdata
ydata
zdata      The origins of the values of the vector field.

udata
vdata
wdata      The values of the vector field to plot.

xdatasource
ydatasource
zdatasource
udatasource
vdatasource
wdatasource

```

Data source variables.

15.4.6.8 Scatter Group

Scatter series objects are created by the `scatter` or `scatter3` functions. A single `hggroup` element contains as many children as there are points in the scatter plot, with each child representing one of the points. The properties of the stem series are

```

linewidth
      The line width of the line objects of the points. See Section 15.4.2 \[Line Styles\],  

page 456.

marker
markeredgecolor
markerfacecolor
      The line and fill color of the markers of the points. See Section 15.4.1 \[Colors\],  

page 456.

xdata
ydata
zdata      The original x, y and z data of the stems.

cdata      The color data for the points of the plot. Each point can have a separate color,  

          or a unique color can be specified.

sizedata  The size data for the points of the plot. Each point can its own size or a unique  

          size can be specified.

xdatasource
ydatasource
zdatasource
cdatasource
sizedatasource

```

Data source variables.

15.4.6.9 Stair Group

Stair series objects are created by the `stair` function. Each `hggroup` element of the series contains a single line object as a child representing the stair. The properties of the stair series are

color The RGB color or color name of the line objects of the stairs. See [Section 15.4.1 \[Colors\]](#), page 456.

linewidth

linestyle The line width and style of the line objects of the stairs. See [Section 15.4.2 \[Line Styles\]](#), page 456.

marker

markeredgecolor

markerfacecolor

markersize The line and fill color of the markers on the stairs. See [Section 15.4.1 \[Colors\]](#), page 456.

xdata

ydata The original x and y data of the stairs.

xdatasource

ydatasource Data source variables.

15.4.6.10 Stem Series

Stem series objects are created by the `stem` or `stem3` functions. Each `hggroup` element contains a single line object as a child representing the stems. The properties of the stem series are

showbaseline

baseline

basevalue The property `showbaseline` flags whether the baseline of the stem series is displayed (default is "on"). The handle of the graphics object representing the baseline is given by the `baseline` property and the y-value (or z-value for `stem3`) of the baseline by the `basevalue` property.

Changes to any of these property are propagated to the other members of the stem series and to the baseline itself. Equally changes in the properties of the base line itself are propagated to the members of the corresponding stem series.

color The RGB color or color name of the line objects of the stems. See [Section 15.4.1 \[Colors\]](#), page 456.

linewidth

linestyle The line width and style of the line objects of the stems. See [Section 15.4.2 \[Line Styles\]](#), page 456.

marker

markeredgecolor

markerfacecolor

markersize The line and fill color of the markers on the stems. See [Section 15.4.1 \[Colors\]](#), page 456.

```

xdata
ydata
zdata      The original x, y and z data of the stems.

xdatasource
ydatasource
zdatasource
Data source variables.

```

15.4.6.11 Surface Group

Surface group objects are created by the `surf` or `mesh` functions, but are equally one of the handles returned by the `surfc` or `meshc` functions. The surface group is of the type `surface`.

The properties of the surface group are

```

edgecolor
facecolor

```

The RGB color or color name of the edges or faces of the surface. See [Section 15.4.1 \[Colors\], page 456](#).

```

linewidth
linestyle

```

The line width and style of the lines on the surface. See [Section 15.4.2 \[Line Styles\], page 456](#).

```

marker
markeredgecolor
markerfacecolor
markersize

```

The line and fill color of the markers on the surface. See [Section 15.4.1 \[Colors\], page 456](#).

```

xdata
ydata
zdata
cdata      The original x, y, z and c data.

xdatasource
ydatasource
zdatasource
cdatasource

```

Data source variables.

15.4.7 Transform Groups

```

h = hgtransform ()
h = hgtransform (property, value, ...)
h = hgtransform (hax, ...)

```

Create a graphics transform object.

FIXME: Need to write documentation. FIXME: Add 'makehgtransform' to seealso list when it is implemented.

See also: [\[hggroup\]](#), page 460.

15.4.8 Graphics Toolkits

```
name = graphics_toolkit ()  
name = graphics_toolkit (hlist)  
graphics_toolkit (name)  
graphics_toolkit (hlist, name)
```

Query or set the default graphics toolkit which is assigned to new figures.

With no inputs, return the current default graphics toolkit. If the input is a list of figure graphic handles, *hlist*, then return the name of the graphics toolkit in use for each figure.

When called with a single input *name* set the default graphics toolkit to *name*. If the toolkit is not already loaded, it is initialized by calling the function `__init_name__`. If the first input is a list of figure handles, *hlist*, then the graphics toolkit is set to *name* for these figures only.

See also: [\[available_graphics_toolkits\]](#), page 473.

```
available_graphics_toolkits ()  
Return a cell array of registered graphics toolkits.
```

See also: [\[graphics_toolkit\]](#), page 473, [\[register_graphics_toolkit\]](#), page 473.

```
loaded_graphics_toolkits ()  
Return a cell array of the currently loaded graphics toolkits.
```

See also: [\[available_graphics_toolkits\]](#), page 473.

```
register_graphics_toolkit (toolkit)  
List toolkit as an available graphics toolkit.
```

See also: [\[available_graphics_toolkits\]](#), page 473.

15.4.8.1 Customizing Toolkit Behavior

The specific behavior of the backend toolkit may be modified using the following utility functions. Note: Not all functions apply to every graphics toolkit.

```
[prog, args] = gnuplot_binary ()  
[old_prog, old_args] = gnuplot_binary (new_prog)  
[old_prog, old_args] = gnuplot_binary (new_prog, arg1, ...)  
Query or set the name of the program invoked by the plot command when the graphics toolkit is set to "gnuplot".
```

Additional arguments to pass to the external plotting program may also be given. The default value is "gnuplot" with no additional arguments. See [Appendix E \[Installation\]](#), page 1021.

See also: [\[graphics_toolkit\]](#), page 473.

15.4.8.2 Hardware vs. Software Rendering

On Windows platforms, Octave uses software rendering for the OpenGL graphics toolkits ("qt" and "fltk") by default. This is done to avoid rendering and printing issues due to imperfect OpenGL driver implementations for diverse graphic cards from different vendors. As a down-side, software rendering might be considerably slower than hardware accelerated rendering. To permanently switch back to hardware accelerated rendering with your graphic card drivers, rename the following file while Octave is closed:

```
octave-home\bin\opengl32.dll
```

where *octave-home* is the directory in which Octave is installed (the default is C:\Octave\Octave-*version*).

16 Matrix Manipulation

There are a number of functions available for checking to see if the elements of a matrix meet some condition, and for rearranging the elements of a matrix. For example, Octave can easily tell you if all the elements of a matrix are finite, or are less than some specified value. Octave can also rotate the elements, extract the upper- or lower-triangular parts, or sort the columns of a matrix.

16.1 Finding Elements and Checking Conditions

The functions `any` and `all` are useful for determining whether any or all of the elements of a matrix satisfy some condition. The `find` function is also useful in determining which elements of a matrix meet a specified condition.

`any (x)`
`any (x, dim)`

For a vector argument, return true (logical 1) if any element of the vector is nonzero.

For a matrix argument, return a row vector of logical ones and zeros with each element indicating whether any of the elements of the corresponding column of the matrix are nonzero. For example:

```
any (eye (2, 4))
⇒ [ 1, 1, 0, 0 ]
```

If the optional argument `dim` is supplied, work along dimension `dim`. For example:

```
any (eye (2, 4), 2)
⇒ [ 1; 1 ]
```

See also: [\[all\]](#), page 475.

`all (x)`
`all (x, dim)`

For a vector argument, return true (logical 1) if all elements of the vector are nonzero.

For a matrix argument, return a row vector of logical ones and zeros with each element indicating whether all of the elements of the corresponding column of the matrix are nonzero. For example:

```
all ([2, 3; 1, 0])
⇒ [ 1, 0 ]
```

If the optional argument `dim` is supplied, work along dimension `dim`.

See also: [\[any\]](#), page 475.

Since the comparison operators (see [Section 8.4 \[Comparison Ops\]](#), page 152) return matrices of ones and zeros, it is easy to test a matrix for many things, not just whether the elements are nonzero. For example,

```
all (all (rand (5) < 0.9))
⇒ 0
```

tests a random 5 by 5 matrix to see if all of its elements are less than 0.9.

Note that in conditional contexts (like the test clause of `if` and `while` statements) Octave treats the test as if you had typed `all (all (condition))`.

```
z = xor (x, y)
z = xor (x1, x2, ...)
```

Return the exclusive or of *x* and *y*.

For boolean expressions *x* and *y*, **xor** (*x*, *y*) is true if and only if one of *x* or *y* is true. Otherwise, if *x* and *y* are both true or both false, **xor** returns false.

The truth table for the xor operation is

<i>x</i>	<i>y</i>	<i>z</i>
-	-	-
0	0	0
1	0	1
0	1	1
1	1	0

If more than two arguments are given the xor operation is applied cumulatively from left to right:

```
(...((x1 XOR x2) XOR x3) XOR ...)
```

See also: [and], page 154, [or], page 155, [not], page 154.

```
diff (x)
diff (x, k)
diff (x, k, dim)
```

If *x* is a vector of length *n*, **diff** (*x*) is the vector of first differences $x_2 - x_1, \dots, x_n - x_{n-1}$.

If *x* is a matrix, **diff** (*x*) is the matrix of column differences along the first non-singleton dimension.

The second argument is optional. If supplied, **diff** (*x*, *k*), where *k* is a non-negative integer, returns the *k*-th differences. It is possible that *k* is larger than the first non-singleton dimension of the matrix. In this case, **diff** continues to take the differences along the next non-singleton dimension.

The dimension along which to take the difference can be explicitly stated with the optional variable *dim*. In this case the *k*-th order differences are calculated along this dimension. In the case where *k* exceeds **size** (*x*, *dim*) an empty matrix is returned.

See also: [sort], page 485, [merge], page 156.

```
isinf (x)
```

Return a logical array which is true where the elements of *x* are infinite and false where they are not.

For example:

```
isinf ([13, Inf, NA, NaN])
⇒ [ 0, 1, 0, 0 ]
```

See also: [isfinite], page 477, [isnan], page 476, [isna], page 43.

```
isnan (x)
```

Return a logical array which is true where the elements of *x* are NaN values and false where they are not.

NA values are also considered NaN values. For example:

```
isnan ([13, Inf, NA, NaN])
⇒ [ 0, 0, 1, 1 ]
```

See also: [isna], page 43, [isinf], page 476, [isfinite], page 477.

isfinite (x)

Return a logical array which is true where the elements of *x* are finite values and false where they are not.

For example:

```
isfinite ([13, Inf, NA, NaN])
⇒ [ 1, 0, 0, 0 ]
```

See also: [isinf], page 476, [isnan], page 476, [isna], page 43.

[err, yi, ...] = common_size (xi, ...)

Determine if all input arguments are either scalar or of common size.

If true, *err* is zero, and *yi* is a matrix of the common size with all entries equal to *xi* if this is a scalar or *xi* otherwise. If the inputs cannot be brought to a common size, *err* is 1, and *yi* is *xi*. For example:

```
[err, a, b] = common_size ([1 2; 3 4], 5)
⇒ err = 0
⇒ a = [ 1, 2; 3, 4 ]
⇒ b = [ 5, 5; 5, 5 ]
```

This is useful for implementing functions where arguments can either be scalars or of common size.

See also: [size], page 45, [size_equal], page 46, [numel], page 44, [ndims], page 44.

```
idx = find (x)
idx = find (x, n)
idx = find (x, n, direction)
[i, j] = find (...)
[i, j, v] = find (...)
```

Return a vector of indices of nonzero elements of a matrix, as a row if *x* is a row vector or as a column otherwise.

To obtain a single index for each matrix element, Octave pretends that the columns of a matrix form one long vector (like Fortran arrays are stored). For example:

```
find (eye (2))
⇒ [ 1; 4 ]
```

If two inputs are given, *n* indicates the maximum number of elements to find from the beginning of the matrix or vector.

If three inputs are given, *direction* should be one of "first" or "last", requesting only the first or last *n* indices, respectively. However, the indices are always returned in ascending order.

If two outputs are requested, `find` returns the row and column indices of nonzero elements of a matrix. For example:

```
[i, j] = find (2 * eye (2))
⇒ i = [ 1; 2 ]
⇒ j = [ 1; 2 ]
```

If three outputs are requested, `find` also returns a vector containing the nonzero values. For example:

```
[i, j, v] = find (3 * eye (2))
⇒ i = [ 1; 2 ]
⇒ j = [ 1; 2 ]
⇒ v = [ 3; 3 ]
```

Note that this function is particularly useful for sparse matrices, as it extracts the nonzero elements as vectors, which can then be used to create the original matrix. For example:

```
sz = size (a);
[i, j, v] = find (a);
b = sparse (i, j, v, sz(1), sz(2));
```

See also: [\[nonzeros\]](#), page 624.

```
idx = lookup (table, y)
idx = lookup (table, y, opt)
```

Lookup values in a **sorted** table.

This function is usually used as a prelude to interpolation.

If `table` is increasing, of length `N` and `idx = lookup (table, y)`, then `table(idx(i)) <= y(i) < table(idx(i+1))` for all `y(i)` within the table. If `y(i) < table(1)` then `idx(i)` is 0. If `y(i) >= table(end)` or `isnan (y(i))` then `idx(i)` is `N`.

If the table is decreasing, then the tests are reversed. For non-strictly monotonic tables, empty intervals are always skipped. The result is undefined if `table` is not monotonic, or if `table` contains a NaN.

The complexity of the lookup is $O(M \log(N))$ where M is the size of `y`. In the special case when `y` is also sorted, the complexity is $O(\min(M \log(N), M+N))$.

`table` and `y` can also be cell arrays of strings (or `y` can be a single string). In this case, string lookup is performed using lexicographical comparison.

If `opts` is specified, it must be a string with letters indicating additional options.

- `m` Match. `table(idx(i)) == y(i)` if `y(i)` occurs in `table`; otherwise, `idx(i)` is zero.
- `b` Boolean. `idx(i)` is a logical 1 or 0, indicating whether `y(i)` is contained in `table` or not.
- `l` Left. For numeric lookups the leftmost subinterval shall be extended to minus infinity (i.e., all indices at least 1).
- `r` Right. For numeric lookups the rightmost subinterval shall be extended to infinity (i.e., all indices at most `N-1`).

Note: If `table` is not sorted the results from `lookup` will be unpredictable.

If you wish to check if a variable exists at all, instead of properties its elements may have, consult [Section 7.3 \[Status of Variables\]](#), page 132.

16.2 Rearranging Matrices

`fliplr (x)`

Flip array left to right.

Return a copy of x with the order of the columns reversed. In other words, x is flipped left-to-right about a vertical axis. For example:

```
fliplr ([1, 2; 3, 4])
⇒ 2 1
    4 3
```

See also: [\[flipud\]](#), page 479, [\[flip\]](#), page 479, [\[rot90\]](#), page 480, [\[rotdim\]](#), page 480.

`flipud (x)`

Flip array upside down.

Return a copy of x with the order of the rows reversed. In other words, x is flipped upside-down about a horizontal axis. For example:

```
flipud ([1, 2; 3, 4])
⇒ 3 4
    1 2
```

See also: [\[fliplr\]](#), page 479, [\[flip\]](#), page 479, [\[rot90\]](#), page 480, [\[rotdim\]](#), page 480.

`flip (x)`

`flip (x, dim)`

Return a copy of array x flipped across dimension dim .

If dim is unspecified it defaults to the first non-singleton dimension.

Examples:

```
## row vector
flip ([1 2 3 4])
⇒ 4 3 2 1

## column vector
flip ([1; 2; 3; 4])
⇒ 4
    3
    2
    1

## 2-D matrix along dimension 1
flip ([1 2; 3 4])
⇒ 3 4
    1 2

## 2-D matrix along dimension 2
```

```
flip ([1 2; 3 4], 2)
⇒ 2 1
4 3
```

See also: [fliplr], page 479, [flipud], page 479, [rot90], page 480, [rotdim], page 480, [permute], page 482, [transpose], page 152.

rot90 (A)

rot90 (A, k)

Rotate array by 90 degree increments.

Return a copy of *A* with the elements rotated counterclockwise in 90-degree increments.

The second argument is optional, and specifies how many 90-degree rotations are to be applied (the default value is 1). Negative values of *k* rotate the matrix in a clockwise direction. For example,

```
rot90 ([1, 2; 3, 4], -1)
⇒ 3 1
4 2
```

rotates the given matrix clockwise by 90 degrees. The following are all equivalent statements:

```
rot90 ([1, 2; 3, 4], -1)
rot90 ([1, 2; 3, 4], 3)
rot90 ([1, 2; 3, 4], 7)
```

The rotation is always performed on the plane of the first two dimensions, i.e., rows and columns. To perform a rotation on any other plane, use **rotdim**.

See also: [rotdim], page 480, [fliplr], page 479, [flipud], page 479, [flip], page 479.

rotdim (x)

rotdim (x, n)

rotdim (x, n, plane)

Return a copy of *x* with the elements rotated counterclockwise in 90-degree increments.

The second argument *n* is optional, and specifies how many 90-degree rotations are to be applied (the default value is 1). Negative values of *n* rotate the matrix in a clockwise direction.

The third argument is also optional and defines the plane of the rotation. If present, *plane* is a two element vector containing two different valid dimensions of the matrix. When *plane* is not given the first two non-singleton dimensions are used.

For example,

```
rotdim ([1, 2; 3, 4], -1, [1, 2])
⇒ 3 1
4 2
```

rotates the given matrix clockwise by 90 degrees. The following are all equivalent statements:

```
rotdim ([1, 2; 3, 4], -1, [1, 2])
rotdim ([1, 2; 3, 4], 3, [1, 2])
rotdim ([1, 2; 3, 4], 7, [1, 2])
```

See also: [rot90], page 480, [fliplr], page 479, [flipud], page 479, [flip], page 479.

cat (dim, array1, array2, ..., arrayN)

Return the concatenation of N-D array objects, *array1*, *array2*, ..., *arrayN* along dimension *dim*.

```
A = ones (2, 2);
B = zeros (2, 2);
cat (2, A, B)
⇒ 1 1 0 0
    1 1 0 0
```

Alternatively, we can concatenate *A* and *B* along the second dimension in the following way:

[*A*, *B*]

dim can be larger than the dimensions of the N-D array objects and the result will thus have *dim* dimensions as the following example shows:

```
cat (4, ones (2, 2), zeros (2, 2))
⇒ ans(:,:,1,1) =
    1 1
    1 1

ans(:,:,1,2) =
    0 0
    0 0
```

See also: [horzcat], page 481, [vertcat], page 481.

horzcat (array1, array2, ..., arrayN)

Return the horizontal concatenation of N-D array objects, *array1*, *array2*, ..., *arrayN* along dimension 2.

Arrays may also be concatenated horizontally using the syntax for creating new matrices. For example:

```
hcat = [ array1, array2, ... ]
```

See also: [cat], page 481, [vertcat], page 481.

vertcat (array1, array2, ..., arrayN)

Return the vertical concatenation of N-D array objects, *array1*, *array2*, ..., *arrayN* along dimension 1.

Arrays may also be concatenated vertically using the syntax for creating new matrices. For example:

```
vcat = [ array1; array2; ... ]
```

See also: [cat], page 481, [horzcat], page 481.

permute (A, perm)

Return the generalized transpose for an N-D array object *A*.

The permutation vector *perm* must contain the elements `1:ndims (A)` (in any order, but each element must appear only once). The *N*th dimension of *A* gets remapped to dimension `PERM(N)`. For example:

```
x = zeros ([2, 3, 5, 7]);
size (x)
⇒ 2 3 5 7

size (permute (x, [2, 1, 3, 4]))
⇒ 3 2 5 7

size (permute (x, [1, 3, 4, 2]))
⇒ 2 5 7 3

## The identity permutation
size (permute (x, [1, 2, 3, 4]))
⇒ 2 3 5 7
```

See also: [\[ipermute\]](#), page 482.

ipermute (A, iperm)

The inverse of the `permute` function.

The expression

```
ipermute (permute (A, perm), perm)
```

returns the original array *A*.

See also: [\[permute\]](#), page 482.

```
reshape (A, m, n, ...)
reshape (A, [m n ...])
reshape (A, ..., [], ...)
reshape (A, size)
```

Return a matrix with the specified dimensions (*m, n, ...*) whose elements are taken from the matrix *A*.

The elements of the matrix are accessed in column-major order (like Fortran arrays are stored).

The following code demonstrates reshaping a 1x4 row vector into a 2x2 square matrix.

```
reshape ([1, 2, 3, 4], 2, 2)
⇒ 1 3
    2 4
```

Note that the total number of elements in the original matrix (`prod (size (A))`) must match the total number of elements in the new matrix (`prod ([m n ...])`).

A single dimension of the return matrix may be left unspecified and Octave will determine its size automatically. An empty matrix (`[]`) is used to flag the unspecified dimension.

See also: [resize], page 483, [vec], page 488, [postpad], page 489, [cat], page 481, [squeeze], page 46.

```
resize (x, m)
resize (x, m, n, ...)
resize (x, [m n ...])
```

Resize *x* cutting off elements as necessary.

In the result, element with certain indices is equal to the corresponding element of *x* if the indices are within the bounds of *x*; otherwise, the element is set to zero.

In other words, the statement

```
y = resize (x, dv)
```

is equivalent to the following code:

```
y = zeros (dv, class (x));
sz = min (dv, size (x));
for i = 1:length (sz)
    idx{i} = 1:sz(i);
endfor
y(idx{:}) = x(idx{:});
```

but is performed more efficiently.

If only *m* is supplied, and it is a scalar, the dimension of the result is *m*-by-*m*. If *m*, *n*, ... are all scalars, then the dimensions of the result are *m*-by-*n*-by-.... If given a vector as input, then the dimensions of the result are given by the elements of that vector.

An object can be resized to more dimensions than it has; in such case the missing dimensions are assumed to be 1. Resizing an object to fewer dimensions is not possible.

See also: [reshape], page 482, [postpad], page 489, [prepad], page 488, [cat], page 481.

```
y = circshift (x, n)
y = circshift (x, n, dim)
```

Circularly shift the values of the array *x*.

n must be a vector of integers no longer than the number of dimensions in *x*. The values of *n* can be either positive or negative, which determines the direction in which the values of *x* are shifted. If an element of *n* is zero, then the corresponding dimension of *x* will not be shifted. If *n* is a scalar and no *dim* is specified then the shift is applied to the first non-singular dimension.

If a scalar *dim* is given then operate along the specified dimension. In this case *n* must be a scalar as well.

Examples:

```
x = [1, 2, 3;
      4, 5, 6;
      7, 8, 9];
## positive shift on rows (1st non-singular dim)
circshift (x, 1)
⇒
```

```

    7   8   9
    1   2   3
    4   5   6
## negative shift on rows (1st non-singular dim)
circshift (x, -2)
⇒
    7   8   9
    1   2   3
    4   5   6
## no shift of rows, shift columns by 1 (2nd dimension)
circshift (x, [0,1])
⇒
    3   1   2
    6   4   5
    9   7   8
## shift columns (2nd dimension)
circshift (x, 1, 2)
⇒
    3   1   2
    6   4   5
    9   7   8

```

See also: [\[permute\]](#), page 482, [\[ipermute\]](#), page 482, [\[shiftdim\]](#), page 484.

shift (x, b)

shift (x, b, dim)

If *x* is a vector, perform a circular shift of length *b* of the elements of *x*.

If *x* is a matrix, do the same for each column of *x*.

If the optional *dim* argument is given, operate along this dimension.

y = shiftdim (x, n)

[y, ns] = shiftdim (x)

Shift the dimensions of *x* by *n*, where *n* must be an integer scalar.

When *n* is positive, the dimensions of *x* are shifted to the left, with the leading dimensions circulated to the end. If *n* is negative, then the dimensions of *x* are shifted to the right, with *n* leading singleton dimensions added.

Called with a single argument, **shiftdim**, removes the leading singleton dimensions, returning the number of dimensions removed in the second output argument *ns*.

For example:

```

x = ones (1, 2, 3);
size (shiftdim (x, -1))
    ⇒ 1 1 2 3
size (shiftdim (x, 1))
    ⇒ 2 3
[b, ns] = shiftdim (x)
⇒ b =
    1 1 1
    1 1 1
⇒ ns = 1

```

See also: [reshape], page 482, [permute], page 482, [ipermute], page 482, [circshift], page 483, [squeeze], page 46.

```

[s, i] = sort (x)
[s, i] = sort (x, dim)
[s, i] = sort (x, mode)
[s, i] = sort (x, dim, mode)

```

Return a copy of *x* with the elements arranged in increasing order.

For matrices, **sort** orders the elements within columns

For example:

```

sort ([1, 2; 2, 3; 3, 1])
    ⇒ 1 1
    2 2
    3 3

```

If the optional argument *dim* is given, then the matrix is sorted along the dimension defined by *dim*. The optional argument *mode* defines the order in which the values will be sorted. Valid values of *mode* are "ascend" or "descend".

The **sort** function may also be used to produce a matrix containing the original row indices of the elements in the sorted matrix. For example:

```

[s, i] = sort ([1, 2; 2, 3; 3, 1])
    ⇒ s = 1 1
        2 2
        3 3
    ⇒ i = 1 3
        2 1
        3 2

```

For equal elements, the indices are such that equal elements are listed in the order in which they appeared in the original list.

Sorting of complex entries is done first by magnitude (**abs (z)**) and for any ties by phase angle (**angle (z)**). For example:

```

sort ([1+i; 1; 1-i])
    ⇒ 1 + 0i
        1 - 1i
        1 + 1i

```

NaN values are treated as being greater than any other value and are sorted to the end of the list.

The **sort** function may also be used to sort strings and cell arrays of strings, in which case ASCII dictionary order (uppercase 'A' precedes lowercase 'a') of the strings is used.

The algorithm used in **sort** is optimized for the sorting of partially ordered lists.

See also: [\[sortrows\]](#), page 486, [\[issorted\]](#), page 486.

```
[s, i] = sortrows (A)
[s, i] = sortrows (A, c)
```

Sort the rows of the matrix *A* according to the order of the columns specified in *c*.

By default (*c* omitted, or a particular column unspecified in *c*) an ascending sort order is used. However, if elements of *c* are negative then the corresponding column is sorted in descending order. If the elements of *A* are strings then a lexicographical sort is used.

Example: sort by column 2 in descending order, then 3 in ascending order

```
x = [ 7, 1, 4;
      8, 3, 5;
      9, 3, 6 ];
sortrows (x, [-2, 3])
⇒ 8 3 5
  9 3 6
  7 1 4
```

See also: [\[sort\]](#), page 485.

```
issorted (a)
issorted (a, mode)
issorted (a, "rows", mode)
```

Return true if the array is sorted according to *mode*, which may be either "ascending", "descending", or "either".

By default, *mode* is "ascending". NaNs are treated in the same manner as **sort**.

If the optional argument "rows" is supplied, check whether the array is sorted by rows as output by the function **sortrows** (with no options).

This function does not support sparse matrices.

See also: [\[sort\]](#), page 485, [\[sortrows\]](#), page 486.

```
nth_element (x, n)
nth_element (x, n, dim)
```

Select the *n*-th smallest element of a vector, using the ordering defined by **sort**.

The result is equivalent to **sort(x)(n)**.

n can also be a contiguous range, either ascending **1:u** or descending **u:-1:1**, in which case a range of elements is returned.

If *x* is an array, **nth_element** operates along the dimension defined by *dim*, or the first non-singleton dimension if *dim* is not given.

Programming Note: `nth_element` encapsulates the C++ standard library algorithms `nth_element` and `partial_sort`. On average, the complexity of the operation is $O(M^* \log(K))$, where $M = \text{size}(x, \text{dim})$ and $K = \text{length}(n)$. This function is intended for cases where the ratio K/M is small; otherwise, it may be better to use `sort`.

See also: [sort], page 485, [min], page 520, [max], page 519.

```
A_LO = tril (A)
A_LO = tril (A, k)
A_LO = tril (A, k, pack)
```

Return a new matrix formed by extracting the lower triangular part of the matrix A , and setting all other elements to zero.

The optional second argument specifies how many diagonals above or below the main diagonal should also be set to zero. The default value of k is zero which includes the main diagonal as part of the result. If the value of k is a nonzero integer then the selection of elements starts at an offset of k diagonals above the main diagonal for positive k or below the main diagonal for negative k . The absolute value of k may not be greater than the number of subdiagonals or superdiagonals.

Example 1 : exclude main diagonal

```
tril (ones (3), -1)
⇒ 0 0 0
    1 0 0
    1 1 0
```

Example 2 : include first superdiagonal

```
tril (ones (3), 1)
⇒ 1 1 0
    1 1 1
    1 1 1
```

If the optional third argument "pack" is given then the extracted elements are not inserted into a matrix, but instead stacked column-wise one above another, and returned as a column vector.

See also: [triu], page 487, [istril], page 65, [diag], page 489.

```
A_UP = triu (A)
A_UP = triu (A, k)
A_UP = triu (A, k, pack)
```

Return a new matrix formed by extracting the upper triangular part of the matrix A , and setting all other elements to zero.

The optional second argument specifies how many diagonals above or below the main diagonal should also be set to zero. The default value of k is zero which includes the main diagonal as part of the result. If the value of k is a nonzero integer then the selection of elements starts at an offset of k diagonals above the main diagonal for positive k or below the main diagonal for negative k . The absolute value of k may not be greater than the number of subdiagonals or superdiagonals.

Example 1 : exclude main diagonal

```
triu (ones (3), 1)
⇒ 0 1 1
    0 0 1
    0 0 0
```

Example 2 : include first subdiagonal

```
triu (ones (3), -1)
⇒ 1 1 1
    1 1 1
    0 1 1
```

If the optional third argument "pack" is given then the extracted elements are not inserted into a matrix, but instead stacked column-wise one above another, and returned as a column vector.

See also: [tril], page 487, [istriu], page 66, [diag], page 489.

```
v = vec (x)
v = vec (x, dim)
```

Return the vector obtained by stacking the columns of the matrix *x* one above the other.

Without *dim* this is equivalent to *x*(:).

If *dim* is supplied, the dimensions of *v* are set to *dim* with all elements along the last dimension. This is equivalent to `shiftdim (x(:), 1-dim)`.

See also: [vech], page 488, [resize], page 483, [cat], page 481.

`vech (x)`

Return the vector obtained by eliminating all superdiagonal elements of the square matrix *x* and stacking the result one column above the other.

This has uses in matrix calculus where the underlying matrix is symmetric and it would be pointless to keep values above the main diagonal.

See also: [vec], page 488.

```
prepad (x, l)
prepad (x, l, c)
prepad (x, l, c, dim)
```

Prepend the scalar value *c* to the vector *x* until it is of length *l*. If *c* is not given, a value of 0 is used.

If `length (x) > l`, elements from the beginning of *x* are removed until a vector of length *l* is obtained.

If *x* is a matrix, elements are prepended or removed from each row.

If the optional argument *dim* is given, operate along this dimension.

If *dim* is larger than the dimensions of *x*, the result will have *dim* dimensions.

See also: [postpad], page 489, [cat], page 481, [resize], page 483.

```
postpad (x, l)
postpad (x, l, c)
postpad (x, l, c, dim)
```

Append the scalar value c to the vector x until it is of length l . If c is not given, a value of 0 is used.

If $\text{length}(x) > l$, elements from the end of x are removed until a vector of length l is obtained.

If x is a matrix, elements are appended or removed from each row.

If the optional argument dim is given, operate along this dimension.

If dim is larger than the dimensions of x , the result will have dim dimensions.

See also: [prepad], page 488, [cat], page 481, [resize], page 483.

```
M = diag (v)
M = diag (v, k)
M = diag (v, m, n)
v = diag (M)
v = diag (M, k)
```

Return a diagonal matrix with vector v on diagonal k .

The second argument is optional. If it is positive, the vector is placed on the k -th superdiagonal. If it is negative, it is placed on the $-k$ -th subdiagonal. The default value of k is 0, and the vector is placed on the main diagonal. For example:

```
diag ([1, 2, 3], 1)
      ⇒  0  1  0  0
          0  0  2  0
          0  0  0  3
          0  0  0  0
```

The 3-input form returns a diagonal matrix with vector v on the main diagonal and the resulting matrix being of size m rows \times n columns.

Given a matrix argument, instead of a vector, `diag` extracts the k -th diagonal of the matrix.

```
blkdiag (A, B, C, ...)
```

Build a block diagonal matrix from A, B, C, \dots

All arguments must be numeric and either two-dimensional matrices or scalars. If any argument is of type sparse, the output will also be sparse.

See also: [diag], page 489, [horzcat], page 481, [vertcat], page 481, [sparse], page 622.

16.3 Special Utility Matrices

```
eye (n)
eye (m, n)
eye ([m n])
eye (..., class)
```

Return an identity matrix.

If invoked with a single scalar argument n , return a square $N \times N$ identity matrix. If supplied two scalar arguments (m, n) , `eye` takes them to be the number of rows and columns. If given a vector with two elements, `eye` uses the values of the elements as the number of rows and columns, respectively. For example:

```
eye (3)
⇒ 1 0 0
  0 1 0
  0 0 1
```

The following expressions all produce the same result:

```
eye (2)
≡
eye (2, 2)
≡
eye (size ([1, 2; 3, 4]))
```

The optional argument *class*, allows `eye` to return an array of the specified type, like

```
val = zeros (n,m, "uint8")
```

Calling `eye` with no arguments is equivalent to calling it with an argument of 1. Any negative dimensions are treated as zero. These odd definitions are for compatibility with MATLAB.

See also: [\[speye\]](#), page 620, [\[ones\]](#), page 490, [\[zeros\]](#), page 490.

```
ones (n)
ones (m, n)
ones (m, n, k, ...)
ones ([m n ...])
ones (..., class)
```

Return a matrix or N-dimensional array whose elements are all 1.

If invoked with a single scalar integer argument n , return a square $N \times N$ matrix.

If invoked with two or more scalar integer arguments, or a vector of integer values, return an array with the given dimensions.

To create a constant matrix whose values are all the same use an expression such as

```
val_matrix = val * ones (m, n)
```

The optional argument *class* specifies the class of the return array and defaults to double. For example:

```
val = ones (m,n, "uint8")
```

See also: [\[zeros\]](#), page 490.

```
zeros (n)
zeros (m, n)
zeros (m, n, k, ...)
zeros ([m n ...])
zeros (..., class)
```

Return a matrix or N-dimensional array whose elements are all 0.

If invoked with a single scalar integer argument, return a square $N \times N$ matrix.

If invoked with two or more scalar integer arguments, or a vector of integer values, return an array with the given dimensions.

The optional argument *class* specifies the class of the return array and defaults to double. For example:

```
val = zeros (m,n, "uint8")
```

See also: [ones], page 490.

```
repmat (A, m)
repmat (A, m, n)
repmat (A, m, n, p ...)
repmat (A, [m n])
repmat (A, [m n p ...])
```

Repeat matrix or N-D array.

Form a block matrix of size *m* by *n*, with a copy of matrix *A* as each element.

If *n* is not specified, form an *m* by *m* block matrix. For copying along more than two dimensions, specify the number of times to copy across each dimension *m*, *n*, *p*, ..., in a vector in the second argument.

See also: [bsxfun], page 586, [kron], page 569, [repelems], page 491.

```
repelems (x, r)
```

Construct a vector of repeated elements from *x*.

r is a 2x*N* integer matrix specifying which elements to repeat and how often to repeat each element. Entries in the first row, *r*(1,*j*), select an element to repeat. The corresponding entry in the second row, *r*(2,*j*), specifies the repeat count. If *x* is a matrix then the columns of *x* are imagined to be stacked on top of each other for purposes of the selection index. A row vector is always returned.

Conceptually the result is calculated as follows:

```
y = [];
for i = 1:columns (r)
    y = [y, x(r(1,i)*ones(1, r(2,i)))] ;
endfor
```

See also: [repmat], page 491, [cat], page 481.

```
xxx = repelem (x, R)
xxx = repelem (x, R_1, ..., R_n)
```

Construct an array of repeated elements from *x* and repeat instructions *R*_1,

x must be a scalar, vector, or N-dimensional array.

A repeat instruction *R*_j must either be a scalar or a vector. If the instruction is a scalar then each component of *x* in dimension *j* is repeated *R*_j times. If the instruction is a vector then it must have the same number of elements as the corresponding dimension *j* of *x*. In this case, the *k*th component of dimension *j* is repeated *R*_j(*k*) times.

If *x* is a scalar or vector then **repelem** may be called with just a single repeat instruction *R* and **repelem** will return a vector with the same orientation as the input.

If x is a matrix then at least two R_j s must be specified.

Note: Using `repelem` with a vector x and a vector for R_j is equivalent to Run Length Decoding.

Examples:

```
A = [1 2 3 4 5];
B = [2 1 0 1 2];
repelem (A, B)
⇒ 1 1 2 4 5 5

A = magic (3)
⇒ A =
  8 1 6
  3 5 7
  4 9 2
B1 = [1 2 3];
B2 = 2;
repelem (A, B1, B2)
⇒ 8 8 1 1 6 6
  3 3 5 5 7 7
  3 3 5 5 7 7
  4 4 9 9 2 2
  4 4 9 9 2 2
  4 4 9 9 2 2
```

More R_j may be specified than the number of dimensions of x . Any excess R_j must be scalars (because x 's size in those dimensions is only 1), and x will be replicated in those dimensions accordingly.

```
A = [1 2 3 4 5];
B1 = 2;
B2 = [2 1 3 0 2];
B3 = 3;
repelem (A, B1, B2, B3)
⇒ ans(:,:,1) =
  1 1 2 3 3 3 5 5
  1 1 2 3 3 3 5 5

ans(:,:,2) =
  1 1 2 3 3 3 5 5
  1 1 2 3 3 3 5 5

ans(:,:,3) =
  1 1 2 3 3 3 5 5
  1 1 2 3 3 3 5 5
```

R_j must be specified in order. A placeholder of 1 may be used for dimensions which do not need replication.

```

repelem ([ -1, 0; 0, 1], 1, 2, 1, 2)
⇒ ans(:,:,1,1) =
-1 -1 0 0
0 0 1 1

ans(:,:,1,2) =
-1 -1 0 0
0 0 1 1

```

If fewer R_{-j} are given than the number of dimensions in x , `repelem` will assume R_{-j} is 1 for those dimensions.

```

A = cat (3, [-1 0; 0 1], [-1 0; 0 1])
⇒ ans(:,:,1) =
-1 0
0 1

ans(:,:,2) =
-1 0
0 1

repelem (A, 2, 3)
⇒ ans(:,:,1) =
-1 -1 -1 0 0 0
-1 -1 -1 0 0 0
0 0 0 1 1 1
0 0 0 1 1 1

ans(:,:,2) =
-1 -1 -1 0 0 0
-1 -1 -1 0 0 0
0 0 0 1 1 1
0 0 0 1 1 1

```

`repelem` preserves the class of x , and works with strings, cell arrays, NA, and NAN inputs. If any R_{-j} is 0 the output will be an empty array.

```

repelem ("Octave", 2, 3)
⇒ 000ccctttaaavvveee
000ccctttaaavvveee

repelem ([1 2 3; 1 2 3], 2, 0)
⇒ [] (4x0)

```

See also: [\[cat\]](#), page 481, [\[kron\]](#), page 569, [\[repmat\]](#), page 491.

The functions `linspace` and `logspace` make it very easy to create vectors with evenly or logarithmically spaced elements. See [Section 4.2 \[Ranges\]](#), page 52.

`linspace (start, end)`
`linspace (start, end, n)`

Return a row vector with n linearly spaced elements between $start$ and end .

If the number of elements is greater than one, then the endpoints *start* and *end* are always included in the range. If *start* is greater than *end*, the elements are stored in decreasing order. If the number of points is not specified, a value of 100 is used.

The `linspace` function returns a row vector when both *start* and *end* are scalars. If one, or both, inputs are vectors, then `linspace` transforms them to column vectors and returns a matrix where each row is an independent sequence between `start(row_n), end(row_n)`.

For compatibility with MATLAB, return the second argument (*end*) when only a single value ($n = 1$) is requested.

See also: `[colon]`, page 844, `[logspace]`, page 494.

```
logspace (a, b)
logspace (a, b, n)
logspace (a, pi, n)
```

Return a row vector with *n* elements logarithmically spaced from 10^a to 10^b .

If *n* is unspecified it defaults to 50.

If *b* is equal to π , the points are between 10^a and π , *not* 10^a and 10^π , in order to be compatible with the corresponding MATLAB function.

Also for compatibility with MATLAB, return the right-hand side of the range (10^b) when just a single value is requested.

See also: `[linspace]`, page 493.

```
rand (n)
rand (m, n, ...)
rand ([m n ...])
v = rand ("state")
rand ("state", v)
rand ("state", "reset")
v = rand ("seed")
rand ("seed", v)
rand ("seed", "reset")
rand (... , "single")
rand (... , "double")
```

Return a matrix with random elements uniformly distributed on the interval (0, 1).

The arguments are handled the same as the arguments for `eye`.

You can query the state of the random number generator using the form

```
v = rand ("state")
```

This returns a column vector *v* of length 625. Later, you can restore the random number generator to the state *v* using the form

```
rand ("state", v)
```

You may also initialize the state vector from an arbitrary vector of length ≤ 625 for *v*. This new state will be a hash based on the value of *v*, not *v* itself.

By default, the generator is initialized from `/dev/urandom` if it is available, otherwise from CPU time, wall clock time, and the current fraction of a second. Note that this

differs from MATLAB, which always initializes the state to the same state at startup. To obtain behavior comparable to MATLAB, initialize with a deterministic state vector in Octave's startup files (see [Section 2.1.2 \[Startup Files\], page 18](#)).

To compute the pseudo-random sequence, `rand` uses the Mersenne Twister with a period of $2^{19937} - 1$ (See M. Matsumoto and T. Nishimura, *Mersenne Twister: A 623-dimensionally equidistributed uniform pseudorandom number generator*, ACM Trans. on Modeling and Computer Simulation Vol. 8, No. 1, pp. 3–30, January 1998, <http://www.math.sci.hiroshima-u.ac.jp/~m-mat/MT/emt.html>). Do **not** use for cryptography without securely hashing several returned values together, otherwise the generator state can be learned after reading 624 consecutive values.

Older versions of Octave used a different random number generator. The new generator is used by default as it is significantly faster than the old generator, and produces random numbers with a significantly longer cycle time. However, in some circumstances it might be desirable to obtain the same random sequences as produced by the old generators. To do this the keyword "`seed`" is used to specify that the old generators should be used, as in

```
rand ("seed", val)
```

which sets the seed of the generator to `val`. The seed of the generator can be queried with

```
s = rand ("seed")
```

However, it should be noted that querying the seed will not cause `rand` to use the old generators, only setting the seed will. To cause `rand` to once again use the new generators, the keyword "`state`" should be used to reset the state of the `rand`.

The state or seed of the generator can be reset to a new random value using the "`reset`" keyword.

The class of the value returned can be controlled by a trailing "`double`" or "`single`" argument. These are the only valid classes.

See also: [\[randn\], page 496](#), [\[rande\], page 496](#), [\[randg\], page 498](#), [\[randp\], page 497](#).

```
randi (imax)
randi (imax, n)
randi (imax, m, n, ...)
randi ([imin imax], ...)
randi (..., "class")
```

Return random integers in the range `1:imax`.

Additional arguments determine the shape of the return matrix. When no arguments are specified a single random integer is returned. If one argument `n` is specified then a square matrix ($n \times n$) is returned. Two or more arguments will return a multi-dimensional matrix ($m \times n \times \dots$).

The integer range may optionally be described by a two element matrix with a lower and upper bound in which case the returned integers will be on the interval $[imin, imax]$.

The optional argument `class` will return a matrix of the requested type. The default is "`double`".

The following example returns 150 integers in the range 1–10.

```
ri = randi (10, 150, 1)
```

Implementation Note: `randi` relies internally on `rand` which uses class "double" to represent numbers. This limits the maximum integer (`imax`) and range (`imax - imin`) to the value returned by the `flintmax` function. For IEEE floating point numbers this value is $2^{53} - 1$.

See also: `[rand]`, page 494.

```
randn (n)
randn (m, n, ...)
randn ([m n ...])
v = randn ("state")
randn ("state", v)
randn ("state", "reset")
v = randn ("seed")
randn ("seed", v)
randn ("seed", "reset")
randn (..., "single")
randn (..., "double")
```

Return a matrix with normally distributed random elements having zero mean and variance one.

The arguments are handled the same as the arguments for `rand`.

By default, `randn` uses the Marsaglia and Tsang “Ziggurat technique” to transform from a uniform to a normal distribution.

The class of the value returned can be controlled by a trailing "double" or "single" argument. These are the only valid classes.

Reference: G. Marsaglia and W.W. Tsang, *Ziggurat Method for Generating Random Variables*, J. Statistical Software, vol 5, 2000, <https://www.jstatsoft.org/v05/i08/>

See also: `[rand]`, page 494, `[rande]`, page 496, `[randg]`, page 498, `[randp]`, page 497.

```
rande (n)
rande (m, n, ...)
rande ([m n ...])
v = rande ("state")
rande ("state", v)
rande ("state", "reset")
v = rande ("seed")
rande ("seed", v)
rande ("seed", "reset")
rande (..., "single")
rande (..., "double")
```

Return a matrix with exponentially distributed random elements.

The arguments are handled the same as the arguments for `rand`.

By default, `rande` uses the Marsaglia and Tsang “Ziggurat technique” to transform from a uniform to an exponential distribution.

The class of the value returned can be controlled by a trailing “`double`” or “`single`” argument. These are the only valid classes.

Reference: G. Marsaglia and W.W. Tsang, *Ziggurat Method for Generating Random Variables*, J. Statistical Software, vol 5, 2000, <https://www.jstatsoft.org/v05/i08/>

See also: `[rand]`, page 494, `[randn]`, page 496, `[randg]`, page 498, `[randp]`, page 497.

```
randp (l, n)
randp (l, m, n, ...)
randp (l, [m n ...])
v = randp ("state")
randp ("state", v)
randp ("state", "reset")
v = randp ("seed")
randp ("seed", v)
randp ("seed", "reset")
randp (..., "single")
randp (..., "double")
```

Return a matrix with Poisson distributed random elements with mean value parameter given by the first argument, l .

The arguments are handled the same as the arguments for `rand`, except for the argument l .

Five different algorithms are used depending on the range of l and whether or not l is a scalar or a matrix.

For scalar $l \leq 12$, use direct method.

W.H. Press, et al., *Numerical Recipes in C*, Cambridge University Press, 1992.

For scalar $l > 12$, use rejection method.[1]

W.H. Press, et al., *Numerical Recipes in C*, Cambridge University Press, 1992.

For matrix $l \leq 10$, use inversion method.[2]

E. Stadlober, et al., WinRand source code, available via FTP.

For matrix $l > 10$, use patchwork rejection method.

E. Stadlober, et al., WinRand source code, available via FTP, or H. Zechner, *Efficient sampling from continuous and discrete unimodal distributions*, Doctoral Dissertation, 156pp., Technical University Graz, Austria, 1994.

For $l > 1e8$, use normal approximation.

L. Montanet, et al., *Review of Particle Properties*, Physical Review D 50 p1284, 1994.

The class of the value returned can be controlled by a trailing “`double`” or “`single`” argument. These are the only valid classes.

See also: [rand], page 494, [randn], page 496, [rande], page 496, [randg], page 498.

```
randg (a, n)
randg (a, m, n, ...)
randg (a, [m n ...])
v = randg ("state")
randg ("state", v)
randg ("state", "reset")
v = randg ("seed")
randg ("seed", v)
randg ("seed", "reset")
randg (..., "single")
randg (..., "double")
```

Return a matrix with `gamma (a, 1)` distributed random elements.

The arguments are handled the same as the arguments for `rand`, except for the argument `a`.

This can be used to generate many distributions:

```
gamma (a, b) for a > -1, b > 0
    r = b * randg (a)

beta (a, b) for a > -1, b > -1
    r1 = randg (a, 1)
    r = r1 / (r1 + randg (b, 1))

Erlang (a, n)
    r = a * randg (n)

chisq (df) for df > 0
    r = 2 * randg (df / 2)

t (df) for 0 < df < inf (use randn if df is infinite)
    r = randn () / sqrt (2 * randg (df / 2) / df)

F (n1, n2) for 0 < n1, 0 < n2
    ## r1 equals 1 if n1 is infinite
    r1 = 2 * randg (n1 / 2) / n1
    ## r2 equals 1 if n2 is infinite
    r2 = 2 * randg (n2 / 2) / n2
    r = r1 / r2

negative binomial (n, p) for n > 0, 0 < p <= 1
    r = randp ((1 - p) / p * randg (n))

non-central chisq (df, L), for df >= 0 and L > 0
    (use chisq if L = 0)
        r = randp (L / 2)
        r(r > 0) = 2 * randg (r(r > 0))
        r(df > 0) += 2 * randg (df(df > 0)/2)

Dirichlet (a1, ... ak)
    r = (randg (a1), ..., randg (ak))
    r = r / sum (r)
```

The class of the value returned can be controlled by a trailing "double" or "single" argument. These are the only valid classes.

See also: [rand], page 494, [randn], page 496, [rande], page 496, [randp], page 497.

The generators operate in the new or old style together, it is not possible to mix the two. Initializing any generator with "state" or "seed" causes the others to switch to the same style for future calls.

The state of each generator is independent and calls to different generators can be interleaved without affecting the final result. For example,

```
rand ("state", [11, 22, 33]);
randn ("state", [44, 55, 66]);
u = rand (100, 1);
n = randn (100, 1);
```

and

```
rand ("state", [11, 22, 33]);
randn ("state", [44, 55, 66]);
u = zeros (100, 1);
n = zeros (100, 1);
for i = 1:100
    u(i) = rand ();
    n(i) = randn ();
end
```

produce equivalent results. When the generators are initialized in the old style with "seed" only `rand` and `randn` are independent, because the old `rande`, `randg` and `randp` generators make calls to `rand` and `randn`.

The generators are initialized with random states at start-up, so that the sequences of random numbers are not the same each time you run Octave.¹ If you really do need to reproduce a sequence of numbers exactly, you can set the state or seed to a specific value.

If invoked without arguments, `rand` and `randn` return a single element of a random sequence.

The original `rand` and `randn` functions use Fortran code from RANLIB, a library of Fortran routines for random number generation, compiled by Barry W. Brown and James Lovato of the Department of Biomathematics at The University of Texas, M.D. Anderson Cancer Center, Houston, TX 77030.

randperm (n)
randperm (n, m)

Return a row vector containing a random permutation of `1:n`.

If `m` is supplied, return `m` unique entries, sampled without replacement from `1:n`.

The complexity is $O(n)$ in memory and $O(m)$ in time, unless $m < n/5$, in which case $O(m)$ memory is used as well. The randomization is performed using `rand()`. All permutations are equally likely.

See also: [perms], page 737.

¹ The old versions of `rand` and `randn` obtain their initial seeds from the system clock.

16.4 Famous Matrices

The following functions return famous matrix forms.

```
gallery (name)
```

```
gallery (name, args)
```

Create interesting matrices for testing.

```
c = gallery ("cauchy", x)
```

```
c = gallery ("cauchy", x, y)
```

Create a Cauchy matrix.

```
c = gallery ("chebspec", n)
```

```
c = gallery ("chebspec", n, k)
```

Create a Chebyshev spectral differentiation matrix.

```
c = gallery ("chebvand", p)
```

```
c = gallery ("chebvand", m, p)
```

Create a Vandermonde-like matrix for the Chebyshev polynomials.

```
a = gallery ("chow", n)
```

```
a = gallery ("chow", n, alpha)
```

```
a = gallery ("chow", n, alpha, delta)
```

Create a Chow matrix – a singular Toeplitz lower Hessenberg matrix.

```
c = gallery ("circul", v)
```

Create a circulant matrix.

```
a = gallery ("clement", n)
```

```
a = gallery ("clement", n, k)
```

Create a tridiagonal matrix with zero diagonal entries.

```
c = gallery ("compar", a)
```

```
c = gallery ("compar", a, k)
```

Create a comparison matrix.

```
a = gallery ("condex", n)
```

```
a = gallery ("condex", n, k)
```

```
a = gallery ("condex", n, k, theta)
```

Create a ‘counterexample’ matrix to a condition estimator.

```
a = gallery ("cycol", [m n])
```

```
a = gallery ("cycol", n)
```

```
a = gallery (..., k)
```

Create a matrix whose columns repeat cyclically.

```
[c, d, e] = gallery ("dorr", n)
```

```
[c, d, e] = gallery ("dorr", n, theta)
```

```
a = gallery ("dorr", ...)
```

Create a diagonally dominant, ill-conditioned, tridiagonal matrix.

```

a = gallery ("dramadah", n)
a = gallery ("dramadah", n, k)
Create a (0, 1) matrix whose inverse has large integer entries.

a = gallery ("fiedler", c)
Create a symmetric Fiedler matrix.

a = gallery ("forsythe", n)
a = gallery ("forsythe", n, alpha)
a = gallery ("forsythe", n, alpha, lambda)
Create a Forsythe matrix (a perturbed Jordan block).

f = gallery ("frank", n)
f = gallery ("frank", n, k)
Create a Frank matrix (ill-conditioned eigenvalues).

c = gallery ("gcdmat", n)
Create a greatest common divisor matrix.

c is an n-by-n matrix whose values correspond to the greatest common divisor of its
coordinate values, i.e., c(i,j) correspond gcd (i, j).

a = gallery ("gearmat", n)
a = gallery ("gearmat", n, i)
a = gallery ("gearmat", n, i, j)
Create a Gear matrix.

g = gallery ("grcar", n)
g = gallery ("grcar", n, k)
Create a Toeplitz matrix with sensitive eigenvalues.

a = gallery ("hanowa", n)
a = gallery ("hanowa", n, d)
Create a matrix whose eigenvalues lie on a vertical line in the complex plane.

v = gallery ("house", x)
[v, beta] = gallery ("house", x)
Create a householder matrix.

a = gallery ("integerdata", imax, [M N ...], j)
a = gallery ("integerdata", imax, M, N, ..., j)
a = gallery ("integerdata", [imin, imax], [M N ...], j)
a = gallery ("integerdata", [imin, imax], M, N, ..., j)
a = gallery ("integerdata", ..., "class")
Create a matrix with random integers in the range [1, imax]. If imin is given then
the integers are in the range [imin, imax].
```

The second input is a matrix of dimensions describing the size of the output. The dimensions can also be input as comma-separated arguments.

The input *j* is an integer index in the range [0, 2^32-1]. The values of the output matrix are always exactly the same (reproducibility) for a given size input and *j* index.

The final optional argument determines the class of the resulting matrix. Possible values for *class*: "uint8", "uint16", "uint32", "int8", "int16", int32", "single", "double". The default is "double".

```
a = gallery ("invhess", x)
a = gallery ("invhess", x, y)
    Create the inverse of an upper Hessenberg matrix.

a = gallery ("invol", n)
    Create an involutory matrix.

a = gallery ("ipjfact", n)
a = gallery ("ipjfact", n, k)
    Create a Hankel matrix with factorial elements.

a = gallery ("jordbloc", n)
a = gallery ("jordbloc", n, lambda)
    Create a Jordan block.

u = gallery ("kahan", n)
u = gallery ("kahan", n, theta)
u = gallery ("kahan", n, theta, pert)
    Create a Kahan matrix (upper trapezoidal).

a = gallery ("kms", n)
a = gallery ("kms", n, rho)
    Create a Kac-Murdock-Szegő Toeplitz matrix.

b = gallery ("krylov", a)
b = gallery ("krylov", a, x)
b = gallery ("krylov", a, x, j)
    Create a Krylov matrix.

a = gallery ("lauchli", n)
a = gallery ("lauchli", n, mu)
    Create a Lauchli matrix (rectangular).

a = gallery ("lehmer", n)
    Create a Lehmer matrix (symmetric positive definite).

t = gallery ("lesp", n)
    Create a tridiagonal matrix with real, sensitive eigenvalues.

a = gallery ("lotkin", n)
    Create a Lotkin matrix.

a = gallery ("minij", n)
    Create a symmetric positive definite matrix MIN(i,j).

a = gallery ("moler", n)
a = gallery ("moler", n, alpha)
    Create a Moler matrix (symmetric positive definite).
```

[a, t] = gallery ("neumann", n)

Create a singular matrix from the discrete Neumann problem (sparse).

a = gallery ("normaldata", [M N . . .], j)

a = gallery ("normaldata", M, N, . . . , j)

a = gallery ("normaldata", . . . , "class")

Create a matrix with random samples from the standard normal distribution (mean = 0, std = 1).

The first input is a matrix of dimensions describing the size of the output. The dimensions can also be input as comma-separated arguments.

The input *j* is an integer index in the range [0, 2^32-1]. The values of the output matrix are always exactly the same (reproducibility) for a given size input and *j* index.

The final optional argument determines the class of the resulting matrix. Possible values for *class*: "single", "double". The default is "double".

q = gallery ("orthog", n)

q = gallery ("orthog", n, k)

Create orthogonal and nearly orthogonal matrices.

a = gallery ("parter", n)

Create a Parter matrix (a Toeplitz matrix with singular values near pi).

p = gallery ("pei", n)

p = gallery ("pei", n, alpha)

Create a Pei matrix.

a = gallery ("Poisson", n)

Create a block tridiagonal matrix from Poisson's equation (sparse).

a = gallery ("prolate", n)

a = gallery ("prolate", n, w)

Create a prolate matrix (symmetric, ill-conditioned Toeplitz matrix).

h = gallery ("randhess", x)

Create a random, orthogonal upper Hessenberg matrix.

a = gallery ("rando", n)

a = gallery ("rando", n, k)

Create a random matrix with elements -1, 0 or 1.

a = gallery ("randsvd", n)

a = gallery ("randsvd", n, kappa)

a = gallery ("randsvd", n, kappa, mode)

a = gallery ("randsvd", n, kappa, mode, k1)

a = gallery ("randsvd", n, kappa, mode, k1, ku)

Create a random matrix with pre-assigned singular values.

a = gallery ("redheff", n)

Create a zero and ones matrix of Redheffer associated with the Riemann hypothesis.

```

a = gallery ("riemann", n)
Create a matrix associated with the Riemann hypothesis.

a = gallery ("ris", n)
Create a symmetric Hankel matrix.

a = gallery ("smoke", n)
a = gallery ("smoke", n, k)
Create a complex matrix, with a 'smoke ring' pseudospectrum.

t = gallery ("toeppd", n)
t = gallery ("toeppd", n, m)
t = gallery ("toeppd", n, m, w)
t = gallery ("toeppd", n, m, w, theta)
Create a symmetric positive definite Toeplitz matrix.

p = gallery ("toeppen", n)
p = gallery ("toeppen", n, a)
p = gallery ("toeppen", n, a, b)
p = gallery ("toeppen", n, a, b, c)
p = gallery ("toeppen", n, a, b, c, d)
p = gallery ("toeppen", n, a, b, c, d, e)
Create a pentadiagonal Toeplitz matrix (sparse).

a = gallery ("tridiag", x, y, z)
a = gallery ("tridiag", n)
a = gallery ("tridiag", n, c, d, e)
Create a tridiagonal matrix (sparse).

t = gallery ("triw", n)
t = gallery ("triw", n, alpha)
t = gallery ("triw", n, alpha, k)
Create an upper triangular matrix discussed by Kahan, Golub, and Wilkinson.

a = gallery ("uniformdata", [M N . . .], j)
a = gallery ("uniformdata", M, N, . . ., j)
a = gallery ("uniformdata", . . ., "class")
Create a matrix with random samples from the standard uniform distribution (range [0,1]).

The first input is a matrix of dimensions describing the size of the output. The dimensions can also be input as comma-separated arguments.

The input j is an integer index in the range [0, 2^32-1]. The values of the output matrix are always exactly the same (reproducibility) for a given size input and j index.

The final optional argument determines the class of the resulting matrix. Possible values for class: "single", "double". The default is "double".

a = gallery ("wathen", nx, ny)
a = gallery ("wathen", nx, ny, k)
Create the Wathen matrix.

```

[a, b] = gallery ("wilk", n)

Create various specific matrices devised/discussed by Wilkinson.

hadamard (n)

Construct a Hadamard matrix (H_n) of size n -by- n .

The size n must be of the form $2^k * p$ in which p is one of 1, 12, 20 or 28. The returned matrix is normalized, meaning $H_n(:, 1) == 1$ and $H_n(1, :) == 1$.

Some of the properties of Hadamard matrices are:

- **kron** (H_m , H_n) is a Hadamard matrix of size m -by- n .
- $H_n * H_n' = n * \text{eye}(n)$.
- The rows of H_n are orthogonal.
- $\det(A) \leq \text{abs}(\det(H_n))$ for all A with $\text{abs}(A(i, j)) \leq 1$.
- Multiplying any row or column by -1 and the matrix will remain a Hadamard matrix.

See also: [compan], page 748, [hankel], page 505, [toeplitz], page 507.

hankel (c)

hankel (c, r)

Return the Hankel matrix constructed from the first column c , and (optionally) the last row r .

If the last element of c is not the same as the first element of r , the last element of c is used. If the second argument is omitted, it is assumed to be a vector of zeros with the same size as c .

A Hankel matrix formed from an m -vector c , and an n -vector r , has the elements

$$H(i, j) = \begin{cases} c_{i+j-1}, & i + j - 1 \leq m; \\ r_{i+j-m}, & \text{otherwise.} \end{cases}$$

See also: [hadamard], page 505, [toeplitz], page 507.

hilb (n)

Return the Hilbert matrix of order n .

The i, j element of a Hilbert matrix is defined as

$$H(i, j) = \frac{1}{(i + j - 1)}$$

Hilbert matrices are close to being singular which make them difficult to invert with numerical routines. Comparing the condition number of a random matrix 5x5 matrix with that of a Hilbert matrix of order 5 reveals just how difficult the problem is.

```
cond (rand (5))
    => 14.392
cond (hilb (5))
    => 4.7661e+05
```

See also: [invhilb], page 506.

invhilb (n)

Return the inverse of the Hilbert matrix of order n .

This can be computed exactly using

$$\begin{aligned} A_{ij} &= -1^{i+j}(i+j-1) \binom{n+i-1}{n-j} \binom{n+j-1}{n-i} \binom{i+j-2}{i-2}^2 \\ &= \frac{p(i)p(j)}{(i+j-1)} \end{aligned}$$

where

$$p(k) = -1^k \binom{k+n-1}{k-1} \binom{n}{k}$$

The validity of this formula can easily be checked by expanding the binomial coefficients in both formulas as factorials. It can be derived more directly via the theory of Cauchy matrices. See J. W. Demmel, *Applied Numerical Linear Algebra*, p. 92.

Compare this with the numerical calculation of **inverse (hilb (n))**, which suffers from the ill-conditioning of the Hilbert matrix, and the finite precision of your computer's floating point arithmetic.

See also: [hilb], page 505.

magic (n)

Create an n -by- n magic square.

A magic square is an arrangement of the integers $1:n^2$ such that the row sums, column sums, and diagonal sums are all equal to the same value.

Note: n must be a scalar greater than or equal to 3. If you supply n less than 3, **magic** returns either a nonmagic square, or else the degenerate magic squares 1 and [].

pascal (n)**pascal (n, t)**

Return the Pascal matrix of order n if $t = 0$.

The default value of t is 0.

When $t = 1$, return the pseudo-lower triangular Cholesky factor of the Pascal matrix (The sign of some columns may be negative). This matrix is its own inverse, that is **pascal (n, 1) ^ 2 == eye (n)**.

If $t = -1$, return the true Cholesky factor with strictly positive values on the diagonal.

If $t = 2$, return a transposed and permuted version of **pascal (n, 1)**, which is the cube root of the identity matrix. That is, **pascal (n, 2) ^ 3 == eye (n)**.

See also: [chol], page 555.

rosser ()

Return the Rosser matrix.

This is a difficult test case used to evaluate eigenvalue algorithms.

See also: [wilkinson], page 507, [eig], page 547.

toeplitz (c)
toeplitz (c, r)

Return the Toeplitz matrix constructed from the first column c , and optionally the first row r .

If the second argument is omitted, the first row is taken to be the same as the first column. If the first element of r is not the same as the first element of c , the first element of c is used.

A Toeplitz, or diagonal-constant, matrix has the same value along each diagonal. Although it need not be square, it often is. An MxN Toeplitz matrix has the form:

$$\begin{bmatrix} c_1 & r_2 & r_3 & \cdots & r_n \\ c_2 & c_1 & r_2 & \cdots & r_{n-1} \\ c_3 & c_2 & c_1 & \cdots & r_{n-2} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ c_m & c_{m-1} & c_{m-2} & \cdots & cm - n + 1 \end{bmatrix}$$

See also: [hankel], page 505.

vander (c)
vander (c, n)

Return the Vandermonde matrix whose next to last column is c .

If n is specified, it determines the number of columns; otherwise, n is taken to be equal to the length of c .

A Vandermonde matrix has the form:

$$\begin{bmatrix} c_1^{n-1} & \cdots & c_1^2 & c_1 & 1 \\ c_2^{n-1} & \cdots & c_2^2 & c_2 & 1 \\ \vdots & \ddots & \vdots & \vdots & \vdots \\ c_n^{n-1} & \cdots & c_n^2 & c_n & 1 \end{bmatrix}$$

See also: [polyfit], page 753.

wilkinson (n)

Return the Wilkinson matrix of order n .

Wilkinson matrices are symmetric and tridiagonal with pairs of nearly, but not exactly, equal eigenvalues. They are useful in testing the behavior and performance of eigenvalue solvers.

See also: [rosser], page 506, [eig], page 547.

17 Arithmetic

Unless otherwise noted, all of the functions described in this chapter will work for real and complex scalar, vector, or matrix arguments. Functions described as *mapping functions* apply the given operation individually to each element when given a matrix argument. For example:

```
sin ([1, 2; 3, 4])
⇒ 0.84147   0.90930
    0.14112 -0.75680
```

17.1 Exponents and Logarithms

`exp (x)`

Compute e^x for each element of x .

To compute the matrix exponential, see [Chapter 18 \[Linear Algebra\]](#), page 545.

See also: [\[log\]](#), page 509.

`expm1 (x)`

Compute $e^x - 1$ accurately in the neighborhood of zero.

See also: [\[exp\]](#), page 509.

`log (x)`

Compute the natural logarithm, $\ln(x)$, for each element of x .

To compute the matrix logarithm, see [Chapter 18 \[Linear Algebra\]](#), page 545.

See also: [\[exp\]](#), page 509, [\[log1p\]](#), page 509, [\[log2\]](#), page 509, [\[log10\]](#), page 509, [\[logspace\]](#), page 494.

`reallog (x)`

Return the real-valued natural logarithm of each element of x .

If any element results in a complex return value `reallog` aborts and issues an error.

See also: [\[log\]](#), page 509, [\[realpow\]](#), page 510, [\[realsqrt\]](#), page 510.

`log1p (x)`

Compute $\ln(1 + x)$ accurately in the neighborhood of zero.

See also: [\[log\]](#), page 509, [\[exp\]](#), page 509, [\[expm1\]](#), page 509.

`log10 (x)`

Compute the base-10 logarithm of each element of x .

See also: [\[log\]](#), page 509, [\[log2\]](#), page 509, [\[logspace\]](#), page 494, [\[exp\]](#), page 509.

`log2 (x)`

`[f, e] = log2 (x)`

Compute the base-2 logarithm of each element of x .

If called with two output arguments, split x into binary mantissa and exponent so that $\frac{1}{2} \leq |f| < 1$ and e is an integer. If $x = 0$, $f = e = 0$.

See also: [\[pow2\]](#), page 510, [\[log\]](#), page 509, [\[log10\]](#), page 509, [\[exp\]](#), page 509.

pow2 (x)
pow2 (f, e)

With one input argument, compute 2^x for each element of x.

With two input arguments, return $f \cdot 2^e$.

See also: [log2], page 509, [nextpow2], page 510, [power], page 151.

n = nextpow2 (x)

Compute the exponent for the smallest power of two larger than the input.

For each element in the input array x, return the first integer n such that $2^n \geq |x|$.

See also: [pow2], page 510, [log2], page 509.

realpow (x, y)

Compute the real-valued, element-by-element power operator.

This is equivalent to $x.^y$, except that **realpow** reports an error if any return value is complex.

See also: [power], page 151, [reallog], page 509, [realsqrt], page 510.

sqrt (x)

Compute the square root of each element of x.

If x is negative, a complex result is returned.

To compute the matrix square root, see Chapter 18 [Linear Algebra], page 545.

See also: [realsqrt], page 510, [nthroot], page 510.

realsqrt (x)

Return the real-valued square root of each element of x.

If any element results in a complex return value **realsqrt** aborts and issues an error.

See also: [sqrt], page 510, [realpow], page 510, [reallog], page 509.

cbrt (x)

Compute the real cube root of each element of x.

Unlike $x^{(1/3)}$, the result will be negative if x is negative.

See also: [nthroot], page 510.

nthroot (x, n)

Compute the real (non-complex) n-th root of x.

x must have all real entries and n must be a scalar. If n is an even integer and x has negative entries then **nthroot** aborts and issues an error.

Example:

```
nthroot (-1, 3)
⇒ -1
(-1) ^ (1 / 3)
⇒ 0.50000 - 0.86603i
```

See also: [realsqrt], page 510, [sqrt], page 510, [cbrt], page 510.

17.2 Complex Arithmetic

In the descriptions of the following functions, z is the complex number $x + iy$, where i is defined as $\sqrt{-1}$.

abs (z)

Compute the magnitude of z .

The magnitude is defined as $|z| = \sqrt{x^2 + y^2}$.

For example:

```
abs (3 + 4i)
      ⇒ 5
```

See also: [arg], page 511.

arg (z)

angle (z)

Compute the argument, i.e., angle of z .

This is defined as, $\theta = \text{atan2}(y, x)$, in radians.

For example:

```
arg (3 + 4i)
      ⇒ 0.92730
```

See also: [abs], page 511.

conj (z)

Return the complex conjugate of z .

The complex conjugate is defined as $\bar{z} = x - iy$.

See also: [real], page 512, [imag], page 512.

cplxpairs (z)

cplxpairs (z, tol)

cplxpairs (z, tol, dim)

Sort the numbers z into complex conjugate pairs ordered by increasing real part.

The negative imaginary complex numbers are placed first within each pair. All real numbers (those with `abs (imag (z)) / abs (z) < tol`) are placed after the complex pairs.

`tol` is a weighting factor in the range $[0, 1]$ which determines the tolerance of the matching. The default value is `100 * eps` and the resulting tolerance for a given complex pair is `tol * abs (z(i))`.

By default the complex pairs are sorted along the first non-singleton dimension of z . If `dim` is specified, then the complex pairs are sorted along this dimension.

Signal an error if some complex numbers could not be paired. Signal an error if all complex numbers are not exact conjugates (to within `tol`). Note that there is no defined order for pairs with identical real parts but differing imaginary parts.

```
cplxpairs (exp (2i*pi*[0:4]'/5)) == exp (2i*pi*[3; 2; 4; 1; 0]/5)
```

imag (z)

Return the imaginary part of z as a real number.

See also: [real], page 512, [conj], page 511.

real (z)

Return the real part of z .

See also: [imag], page 512, [conj], page 511.

17.3 Trigonometry

Octave provides the following trigonometric functions where angles are specified in radians. To convert from degrees to radians multiply by $\pi/180$ or use the `deg2rad` function. For example, `sin (30 * pi/180)` returns the sine of 30 degrees. As an alternative, Octave provides a number of trigonometric functions which work directly on an argument specified in degrees. These functions are named after the base trigonometric function with a ‘d’ suffix. As an example, `sin` expects an angle in radians while `sind` expects an angle in degrees.

Octave uses the C library trigonometric functions. It is expected that these functions are defined by the ISO/IEC 9899 Standard. This Standard is available at: <http://www.open-std.org/jtc1/sc22/wg14/www/docs/n1124.pdf>. Section F.9.1 deals with the trigonometric functions. The behavior of most of the functions is relatively straightforward. However, there are some exceptions to the standard behavior. Many of the exceptions involve the behavior for -0. The most complex case is `atan2`. Octave exactly implements the behavior given in the Standard. Including `atan2(±0, −0)` returns $\pm\pi$.

It should be noted that MATLAB uses different definitions which apparently do not distinguish -0.

rad = deg2rad (deg)

Convert degrees to radians.

The input deg must be a scalar, vector, or N-dimensional array of double or single floating point values. deg may be complex in which case the real and imaginary components are converted separately.

The output rad is the same size and shape as deg with degrees converted to radians using the conversion constant $\pi/180$.

Example:

```
deg2rad ([0, 90, 180, 270, 360])
⇒ 0.00000 1.57080 3.14159 4.71239 6.28319
```

See also: [rad2deg], page 512.

deg = rad2deg (rad)

Convert radians to degrees.

The input rad must be a scalar, vector, or N-dimensional array of double or single floating point values. rad may be complex in which case the real and imaginary components are converted separately.

The output deg is the same size and shape as rad with radians converted to degrees using the conversion constant $180/\pi$.

Example:

```
rad2deg ([0, pi/2, pi, 3/2*pi, 2*pi])
⇒ 0    90   180   270   360
```

See also: [deg2rad], page 512.

sin (x)

Compute the sine for each element of x in radians.

See also: [asin], page 513, [sind], page 515, [sinh], page 514.

cos (x)

Compute the cosine for each element of x in radians.

See also: [acos], page 513, [cosd], page 515, [cosh], page 514.

tan (z)

Compute the tangent for each element of x in radians.

See also: [atan], page 513, [tand], page 515, [tanh], page 514.

sec (x)

Compute the secant for each element of x in radians.

See also: [asec], page 513, [secd], page 515, [sech], page 514.

csc (x)

Compute the cosecant for each element of x in radians.

See also: [acsc], page 514, [cscd], page 515, [csch], page 514.

cot (x)

Compute the cotangent for each element of x in radians.

See also: [acot], page 514, [cotd], page 516, [coth], page 514.

asin (x)

Compute the inverse sine in radians for each element of x .

See also: [sin], page 513, [asind], page 516.

acos (x)

Compute the inverse cosine in radians for each element of x .

See also: [cos], page 513, [acosd], page 516.

atan (x)

Compute the inverse tangent in radians for each element of x .

See also: [tan], page 513, [atand], page 516.

asec (x)

Compute the inverse secant in radians for each element of x .

See also: [sec], page 513, [secd], page 516.

acsc (x)

Compute the inverse cosecant in radians for each element of x.

See also: [csc], page 513, [acscl], page 516.

acot (x)

Compute the inverse cotangent in radians for each element of x.

See also: [cot], page 513, [acotd], page 516.

sinh (x)

Compute the hyperbolic sine for each element of x.

See also: [asinh], page 514, [cosh], page 514, [tanh], page 514.

cosh (x)

Compute the hyperbolic cosine for each element of x.

See also: [acosh], page 514, [sinh], page 514, [tanh], page 514.

tanh (x)

Compute hyperbolic tangent for each element of x.

See also: [atanh], page 514, [sinh], page 514, [cosh], page 514.

sech (x)

Compute the hyperbolic secant of each element of x.

See also: [asech], page 514.

csch (x)

Compute the hyperbolic cosecant of each element of x.

See also: [acsch], page 515.

coth (x)

Compute the hyperbolic cotangent of each element of x.

See also: [acoth], page 515.

asinh (x)

Compute the inverse hyperbolic sine for each element of x.

See also: [sinh], page 514.

acosh (x)

Compute the inverse hyperbolic cosine for each element of x.

See also: [cosh], page 514.

atanh (x)

Compute the inverse hyperbolic tangent for each element of x.

See also: [tanh], page 514.

asech (x)

Compute the inverse hyperbolic secant of each element of x.

See also: [sech], page 514.

acsch (x)

Compute the inverse hyperbolic cosecant of each element of x.

See also: [csch], page 514.

acoth (x)

Compute the inverse hyperbolic cotangent of each element of x.

See also: [coth], page 514.

atan2 (y, x)

Compute atan (y / x) for corresponding elements of y and x.

y and x must match in size and orientation. The signs of elements of y and x are used to determine the quadrants of each resulting value.

This function is equivalent to `arg (complex (x, y))`.

See also: [tan], page 513, [tand], page 515, [tanh], page 514, [atanh], page 514.

Octave provides the following trigonometric functions where angles are specified in degrees. These functions produce true zeros at the appropriate intervals rather than the small round-off error that occurs when using radians. For example:

```
cosd (90)
⇒ 0
cos (pi/2)
⇒ 6.1230e-17
```

sind (x)

Compute the sine for each element of x in degrees.

Returns zero for elements where $x/180$ is an integer.

See also: [asind], page 516, [sin], page 513.

cosd (x)

Compute the cosine for each element of x in degrees.

Returns zero for elements where $(x-90)/180$ is an integer.

See also: [acosd], page 516, [cos], page 513.

tand (x)

Compute the tangent for each element of x in degrees.

Returns zero for elements where $x/180$ is an integer and Inf for elements where $(x-90)/180$ is an integer.

See also: [atand], page 516, [tan], page 513.

secd (x)

Compute the secant for each element of x in degrees.

See also: [asecd], page 516, [sec], page 513.

cscd (x)

Compute the cosecant for each element of x in degrees.

See also: [acsqd], page 516, [csc], page 513.

cotd (x)

Compute the cotangent for each element of x in degrees.

See also: [acotd], page 516, [cot], page 513.

asind (x)

Compute the inverse sine in degrees for each element of x .

See also: [sind], page 515, [asin], page 513.

acosd (x)

Compute the inverse cosine in degrees for each element of x .

See also: [cosd], page 515, [acos], page 513.

atand (x)

Compute the inverse tangent in degrees for each element of x .

See also: [tand], page 515, [atan], page 513.

atan2d (y, x)

Compute atan (y / x) in degrees for corresponding elements from y and x .

See also: [tand], page 515, [atan2], page 515.

asecd (x)

Compute the inverse secant in degrees for each element of x .

See also: [secd], page 515, [asec], page 513.

acsqd (x)

Compute the inverse cosecant in degrees for each element of x .

See also: [csqd], page 515, [acsc], page 514.

acotd (x)

Compute the inverse cotangent in degrees for each element of x .

See also: [cotd], page 516, [acot], page 514.

17.4 Sums and Products

sum (x)**sum (x, dim)****sum (..., "native")****sum (..., "double")****sum (..., "extra")**

Sum of elements along dimension dim .

If dim is omitted, it defaults to the first non-singleton dimension.

The optional "type" input determines the class of the variable used for calculations. By default, operations on floating point inputs (double or single) are performed in their native data type, while operations on integer, logical, and character data types are performed using doubles. If the argument "native" is given, then the operation is performed in the same type as the original argument.

For example:

```
sum ([true, true])
    => 2
sum ([true, true], "native")
    => true
```

If "double" is given the sum is performed in double precision even for single precision inputs.

For double precision inputs, the "extra" option will use a more accurate algorithm than straightforward summation. For single precision inputs, "extra" is the same as "double". For all other data type "extra" has no effect.

See also: [\[cumsum\]](#), page 517, [\[sumsq\]](#), page 518, [\[prod\]](#), page 517.

```
prod (x)
prod (x, dim)
prod (... , "native")
prod (... , "double")
```

Product of elements along dimension *dim*.

If *dim* is omitted, it defaults to the first non-singleton dimension.

The optional "type" input determines the class of the variable used for calculations. If the argument "native" is given, then the operation is performed in the same type as the original argument, rather than the default double type.

For example:

```
prod ([true, true])
    => 1
prod ([true, true], "native")
    => true
```

On the contrary, if "double" is given, the operation is performed in double precision even for single precision inputs.

See also: [\[cumprod\]](#), page 518, [\[sum\]](#), page 516.

```
cumsum (x)
cumsum (x, dim)
cumsum (... , "native")
cumsum (... , "double")
```

Cumulative sum of elements along dimension *dim*.

If *dim* is omitted, it defaults to the first non-singleton dimension. For example:

```
cumsum ([1, 2; 3, 4; 5, 6])
    =>
      1   2
      4   6
      9  12
```

See [sum](#) for an explanation of the optional parameters "native" and "double".

See also: [\[sum\]](#), page 516, [\[cumprod\]](#), page 518.

cumprod (x)
cumprod (x, dim)

Cumulative product of elements along dimension *dim*.

If *dim* is omitted, it defaults to the first non-singleton dimension. For example:

```
cumprod ([1, 2; 3, 4; 5, 6])
⇒ 1   2
    3   8
    15  48
```

See also: [prod], page 517, [cumsum], page 517.

sumsq (x)
sumsq (x, dim)

Sum of squares of elements along dimension *dim*.

If *dim* is omitted, it defaults to the first non-singleton dimension.

This function is conceptually equivalent to computing

```
sum (x .* conj (x), dim)
```

but it uses less memory and avoids calling `conj` if *x* is real.

See also: [sum], page 516, [prod], page 517.

17.5 Utility Functions

ceil (x)

Return the smallest integer not less than *x*.

This is equivalent to rounding towards positive infinity.

If *x* is complex, return `ceil (real (x)) + ceil (imag (x)) * I`.

```
ceil([-2.7, 2.7])
⇒ -2      3
```

See also: [floor], page 518, [round], page 519, [fix], page 518.

fix (x)

Truncate fractional portion of *x* and return the integer portion.

This is equivalent to rounding towards zero. If *x* is complex, return `fix (real (x)) + fix (imag (x)) * I`.

```
fix([-2.7, 2.7])
⇒ -2      2
```

See also: [ceil], page 518, [floor], page 518, [round], page 519.

floor (x)

Return the largest integer not greater than *x*.

This is equivalent to rounding towards negative infinity. If *x* is complex, return `floor (real (x)) + floor (imag (x)) * I`.

```
floor([-2.7, 2.7])
⇒ -3      2
```

See also: [ceil], page 518, [round], page 519, [fix], page 518.

round (x)

Return the integer nearest to x.

If x is complex, return `round (real (x)) + round (imag (x)) * I`. If there are two nearest integers, return the one further away from zero.

```
round([-2.7, 2.7])
⇒ -3    3
```

See also: [ceil], page 518, [floor], page 518, [fix], page 518, [roundb], page 519.

roundb (x)

Return the integer nearest to x. If there are two nearest integers, return the even one (banker's rounding).

If x is complex, return `roundb (real (x)) + roundb (imag (x)) * I`.

See also: [round], page 519.

max (x)

```
max(x, [], dim)
```

```
[w, iw] = max(x)
```

```
max(x, y)
```

Find maximum values in the array x.

For a vector argument, return the maximum value. For a matrix argument, return a row vector with the maximum value of each column. For a multi-dimensional array, `max` operates along the first non-singleton dimension.

If the optional third argument `dim` is present then operate along this dimension. In this case the second argument is ignored and should be set to the empty matrix.

For two matrices (or a matrix and a scalar), return the pairwise maximum.

Thus,

```
max(max(x))
```

returns the largest element of the 2-D matrix x, and

```
max(2:5, pi)
⇒ 3.1416 3.1416 4.0000 5.0000
```

compares each element of the range 2:5 with pi, and returns a row vector of the maximum values.

For complex arguments, the magnitude of the elements are used for comparison. If the magnitudes are identical, then the results are ordered by phase angle in the range (-pi, pi]. Hence,

```
max([-1 i 1 -i])
⇒ -1
```

because all entries have magnitude 1, but -1 has the largest phase angle with value pi.

If called with one input and two output arguments, `max` also returns the first index of the maximum value(s). Thus,

```
[x, ix] = max([1, 3, 5, 2, 5])
⇒ x = 5
ix = 3
```

See also: `[min]`, page 520, `[cummax]`, page 520, `[cummin]`, page 521.

```
min (x)
min (x, [], dim)
[w, iw] = min (x)
min (x, y)
```

Find minimum values in the array *x*.

For a vector argument, return the minimum value. For a matrix argument, return a row vector with the minimum value of each column. For a multi-dimensional array, `min` operates along the first non-singleton dimension.

If the optional third argument *dim* is present then operate along this dimension. In this case the second argument is ignored and should be set to the empty matrix.

For two matrices (or a matrix and a scalar), return the pairwise minimum.

Thus,

```
min (min (x))
```

returns the smallest element of the 2-D matrix *x*, and

```
min (2:5, pi)
⇒ 2.0000 3.0000 3.1416 3.1416
```

compares each element of the range `2:5` with `pi`, and returns a row vector of the minimum values.

For complex arguments, the magnitude of the elements are used for comparison. If the magnitudes are identical, then the results are ordered by phase angle in the range $(-\pi, \pi]$. Hence,

```
min ([-1 i 1 -i])
⇒ -i
```

because all entries have magnitude 1, but `-i` has the smallest phase angle with value $-\pi/2$.

If called with one input and two output arguments, `min` also returns the first index of the minimum value(s). Thus,

```
[x, ix] = min ([1, 3, 0, 2, 0])
⇒ x = 0
ix = 3
```

See also: `[max]`, page 519, `[cummin]`, page 521, `[cummax]`, page 520.

```
cummax (x)
cummax (x, dim)
[w, iw] = cummax (...)
```

Return the cumulative maximum values along dimension *dim*.

If *dim* is unspecified it defaults to column-wise operation. For example:

```
cummax ([1 3 2 6 4 5])
⇒ 1 3 3 6 6 6
```

If called with two output arguments the index of the maximum value is also returned.

```
[w, iw] = cummax ([1 3 2 6 4 5])
⇒
w = 1 3 3 6 6 6
iw = 1 2 2 4 4 4
```

See also: [cummin], page 521, [max], page 519, [min], page 520.

```
cummin (x)
cummin (x, dim)
[w, iw] = cummin (x)
```

Return the cumulative minimum values along dimension *dim*.

If *dim* is unspecified it defaults to column-wise operation. For example:

```
cummin ([5 4 6 2 3 1])
⇒ 5 4 4 2 2 1
```

If called with two output arguments the index of the minimum value is also returned.

```
[w, iw] = cummin ([5 4 6 2 3 1])
⇒
w = 5 4 4 2 2 1
iw = 1 2 2 4 4 6
```

See also: [cummax], page 520, [min], page 520, [max], page 519.

```
hypot (x, y)
hypot (x, y, z, ...)
```

Compute the element-by-element square root of the sum of the squares of *x* and *y*.

This is equivalent to `sqrt (x.^2 + y.^2)`, but is calculated in a manner that avoids overflows for large values of *x* or *y*.

`hypot` can also be called with more than 2 arguments; in this case, the arguments are accumulated from left to right:

```
hypot (hypot (x, y), z)
hypot (hypot (hypot (x, y), z), w), etc.
```

```
dx = gradient (m)
[dx, dy, dz, ...] = gradient (m)
[...] = gradient (m, s)
[...] = gradient (m, x, y, z, ...)
[...] = gradient (f, x0)
[...] = gradient (f, x0, s)
[...] = gradient (f, x0, x, y, ...)
```

Calculate the gradient of sampled data or a function.

If *m* is a vector, calculate the one-dimensional gradient of *m*. If *m* is a matrix the gradient is calculated for each dimension.

`[dx, dy] = gradient (m)` calculates the one-dimensional gradient for *x* and *y* direction if *m* is a matrix. Additional return arguments can be used for multi-dimensional matrices.

A constant spacing between two points can be provided by the *s* parameter. If *s* is a scalar, it is assumed to be the spacing for all dimensions. Otherwise, separate values of the spacing can be supplied by the *x*, ... arguments. Scalar values specify an equidistant spacing. Vector values for the *x*, ... arguments specify the coordinate for that dimension. The length must match their respective dimension of *m*.

At boundary points a linear extrapolation is applied. Interior points are calculated with the first approximation of the numerical gradient

$$y'(i) = 1/(x(i+1)-x(i-1)) * (y(i-1)-y(i+1)).$$

If the first argument *f* is a function handle, the gradient of the function at the points in *x0* is approximated using central difference. For example, `gradient (@cos, 0)` approximates the gradient of the cosine function in the point *x0* = 0. As with sampled data, the spacing values between the points from which the gradient is estimated can be set via the *s* or *dx*, *dy*, ... arguments. By default a spacing of 1 is used.

See also: `[diff]`, page 476, `[del2]`, page 523.

`dot (x, y, dim)`

Compute the dot product of two vectors.

If *x* and *y* are matrices, calculate the dot products along the first non-singleton dimension.

If the optional argument *dim* is given, calculate the dot products along this dimension.

This is equivalent to `sum (conj (X) .* Y, dim)`, but avoids forming a temporary array and is faster. When *X* and *Y* are column vectors, the result is equivalent to *X'* * *Y*.

See also: `[cross]`, page 522, `[divergence]`, page 522.

`cross (x, y)`

`cross (x, y, dim)`

Compute the vector cross product of two 3-dimensional vectors *x* and *y*.

If *x* and *y* are matrices, the cross product is applied along the first dimension with three elements.

The optional argument *dim* forces the cross product to be calculated along the specified dimension.

Example Code:

```
cross ([1, 1, 0], [0, 1, 1])
⇒
1   -1   1
```

See also: `[dot]`, page 522, `[curl]`, page 523, `[divergence]`, page 522.

`div = divergence (x, y, z, fx, fy, fz)`

`div = divergence (fx, fy, fz)`

`div = divergence (x, y, fx, fy)`

`div = divergence (fx, fy)`

Calculate divergence of a vector field given by the arrays *fx*, , and *fz* or *fx*, *fy* respectively.

$$\operatorname{div} F(x, y, z) = \partial_x F + \partial_y F + \partial_z F$$

The coordinates of the vector field can be given by the arguments x , y , z or x , y respectively.

See also: [curl], page 523, [gradient], page 521, [del2], page 523, [dot], page 522.

```
[cx, cy, cz, v] = curl (x, y, z, fx, fy, fz)
[cz, v] = curl (x, y, fx, fy)
[...] = curl (fx, fy, fz)
[...] = curl (fx, fy)
v = curl (...)
```

Calculate curl of vector field given by the arrays fx , fy , and fz or fx , fy respectively.

$$\operatorname{curl} F(x, y, z) = \left(\frac{\partial d}{\partial y} F_z - \frac{\partial d}{\partial z} F_y, \frac{\partial d}{\partial z} F_x - \frac{\partial d}{\partial x} F_z, \frac{\partial d}{\partial x} F_y - \frac{\partial d}{\partial y} F_x \right)$$

The coordinates of the vector field can be given by the arguments x , y , z or x , y respectively. v calculates the scalar component of the angular velocity vector in direction of the z -axis for two-dimensional input. For three-dimensional input the scalar rotation is calculated at each grid point in direction of the vector field at that point.

See also: [divergence], page 522, [gradient], page 521, [del2], page 523, [cross], page 522.

```
L = del2 (M)
L = del2 (M, h)
L = del2 (M, dx, dy, ...)
```

Calculate the discrete Laplace operator (∇^2).

For a 2-dimensional matrix M this is defined as

$$L = \frac{1}{4} \left(\frac{d^2}{dx^2} M(x, y) + \frac{d^2}{dy^2} M(x, y) \right)$$

For N -dimensional arrays the sum in parentheses is expanded to include second derivatives over the additional higher dimensions.

The spacing between evaluation points may be defined by h , which is a scalar defining the equidistant spacing in all dimensions. Alternatively, the spacing in each dimension may be defined separately by dx , dy , etc. A scalar spacing argument defines equidistant spacing, whereas a vector argument can be used to specify variable spacing. The length of the spacing vectors must match the respective dimension of M . The default spacing value is 1.

Dimensions with fewer than 3 data points are skipped. Boundary points are calculated from the linear extrapolation of interior points.

Example: Second derivative of $2*x^3$

```
f = @(x) 2*x.^3;
dd = @(x) 12*x;
x = 1:6;
L = 4*del2 (f(x));
assert (L, dd (x));
```

See also: [gradient], page 521, [diff], page 476.

factorial (n)

Return the factorial of n where n is a real non-negative integer.

If n is a scalar, this is equivalent to `prod (1:n)`. For vector or matrix arguments, return the factorial of each element in the array.

For non-integers see the generalized factorial function `gamma`. Note that the factorial function grows large quite quickly, and even with double precision values overflow will occur if $n > 171$. For such cases consider `gammaln`.

See also: [prod], page 517, [gamma], page 535, [gammaln], page 537.

```
pf = factor (q)
[pf, n] = factor (q)
```

Return the prime factorization of q .

The prime factorization is defined as `prod (pf) == q` where every element of pf is a prime number. If $q == 1$, return 1. The output pf is of the same numeric class as the input.

With two output arguments, return the unique prime factors pf and their multiplicities. That is, `prod (pf .^ n) == q`.

Implementation Note: The input q must be less than `flintmax` ($9.0072e+15$) in order to factor correctly.

See also: [gcd], page 524, [lcm], page 524, [isprime], page 66, [primes], page 525.

```
g = gcd (a1, a2, ...)
[g, v1, ...] = gcd (a1, a2, ...)
```

Compute the greatest common divisor of $a1, a2, \dots$.

If more than one argument is given then all arguments must be the same size or scalar. In this case the greatest common divisor is calculated for each element individually. All elements must be ordinary or Gaussian (complex) integers. Note that for Gaussian integers, the gcd is only unique up to a phase factor (multiplication by 1, -1, i, or -i), so an arbitrary greatest common divisor among the four possible is returned.

Optional return arguments $v1, \dots$, contain integer vectors such that,

$$g = v_1 a_1 + v_2 a_2 + \dots$$

Example code:

```
gcd ([15, 9], [20, 18])
⇒ 5 9
```

See also: [lcm], page 524, [factor], page 524, [isprime], page 66.

```
lcm (x, y)
lcm (x, y, ...)
```

Compute the least common multiple of x and y , or of the list of all arguments.

All elements must be numeric and of the same size or scalar.

See also: [factor], page 524, [gcd], page 524, [isprime], page 66.

rem (x, y)

Return the remainder of the division x / y .

The remainder is computed using the expression

$$x - y \cdot \text{fix}(x ./ y)$$

An error message is printed if the dimensions of the arguments do not agree, or if either argument is complex.

Programming Notes: When calculating with floating point numbers (double, single), values within a few eps of an integer will be rounded to that integer before computation for compatibility with MATLAB. Any floating point integers greater than **flintmax** (2^{53} for double) will not compute correctly. For larger integer values convert the input to **uint64** before calling this function.

By convention,

```
rem (x, 0) = NaN  if x is a floating point variable
rem (x, 0) = 0    if x is an integer variable
rem (x, y)  returns a value with the signbit from x
```

For the opposite conventions see the **mod** function. In general, **rem** is best when computing the remainder after division of two *positive* numbers. For negative numbers, or when the values are periodic, **mod** is a better choice.

See also: [mod], page 525.

mod (x, y)

Compute the modulo of x and y.

Conceptually this is given by

$$x - y \cdot \text{floor}(x ./ y)$$

and is written such that the correct modulus is returned for integer types. This function handles negative values correctly. That is, **mod (-1, 3)** is 2, not -1, as **rem (-1, 3)** returns.

An error results if the dimensions of the arguments do not agree, or if either of the arguments is complex.

Programming Notes: When calculating with floating point numbers (double, single), values within a few eps of an integer will be rounded to that integer before computation for compatibility with MATLAB. Any floating point integers greater than **flintmax** (2^{53} for double) will not compute correctly. For larger integer values convert the input to **uint64** before calling this function.

By convention,

```
mod (x, 0) = x
mod (x, y)      returns a value with the signbit from y
```

For the opposite conventions see the **rem** function. In general, **mod** is a better choice than **rem** when any of the inputs are negative numbers or when the values are periodic.

See also: [rem], page 525.

p = primes (n)

Return all primes up to n.

The output data class (double, single, uint32, etc.) is the same as the input class of n . The algorithm used is the Sieve of Eratosthenes.

Note: If you need a specific number of primes you can use the fact that the distance from one prime to the next is, on average, proportional to the logarithm of the prime. Integrating, one finds that there are about k primes less than $k \log(5k)$.

See also [list_primes](#) if you need a specific number n of primes.

See also: [\[list_primes\]](#), page 526, [\[isprime\]](#), page 66.

list_primes ()

list_primes (n)

List the first n primes.

If n is unspecified, the first 25 primes are listed.

See also: [\[primes\]](#), page 525, [\[isprime\]](#), page 66.

sign (x)

Compute the *signum* function.

This is defined as

$$\text{sign}(x) = \begin{cases} 1, & x > 0; \\ 0, & x = 0; \\ -1, & x < 0. \end{cases}$$

For complex arguments, **sign** returns $x ./ \text{abs} (x)$.

Note that **sign** (-0.0) is 0. Although IEEE 754 floating point allows zero to be signed, 0.0 and -0.0 compare equal. If you must test whether zero is signed, use the **signbit** function.

See also: [\[signbit\]](#), page 526.

signbit (x)

Return logical true if the value of x has its sign bit set and false otherwise.

This behavior is consistent with the other logical functions. See [Section 4.6 \[Logical Values\]](#), page 60. The behavior differs from the C language function which returns nonzero if the sign bit is set.

This is not the same as $x < 0.0$, because IEEE 754 floating point allows zero to be signed. The comparison $-0.0 < 0.0$ is false, but **signbit** (-0.0) will return a nonzero value.

See also: [\[sign\]](#), page 526.

17.6 Special Functions

[a, ierr] = airy (k, z, opt)

Compute Airy functions of the first and second kind, and their derivatives.

K	Function	Scale factor (if "opt" is supplied)
0	Ai (Z)	$\exp((2/3) * Z * \sqrt(Z))$
1	dAi(Z)/dZ	$\exp((2/3) * Z * \sqrt(Z))$
2	Bi (Z)	$\exp(-\text{abs}(\text{real}((2/3) * Z * \sqrt(Z))))$
3	dBi(Z)/dZ	$\exp(-\text{abs}(\text{real}((2/3) * Z * \sqrt(Z))))$

The function call `airy (z)` is equivalent to `airy (0, z)`.

The result is the same size as `z`.

If requested, `ierr` contains the following status information and is the same size as the result.

0. Normal return.
1. Input error, return `NaN`.
2. Overflow, return `Inf`.
3. Loss of significance by argument reduction results in less than half of machine accuracy.
4. Loss of significance by argument reduction, output may be inaccurate.
5. Error—no computation, algorithm termination condition not met, return `NaN`.

```
J = besselj (alpha, x)
J = besselj (alpha, x, opt)
[J, ierr] = besselj (...)
```

Compute Bessel functions of the first kind.

The order of the Bessel function `alpha` must be real. The points for evaluation `x` may be complex.

If the optional argument `opt` is 1 or `true`, the result `J` is multiplied by `exp (-abs (imag (x)))`.

If `alpha` is a scalar, the result is the same size as `x`. If `x` is a scalar, the result is the same size as `alpha`. If `alpha` is a row vector and `x` is a column vector, the result is a matrix with `length (x)` rows and `length (alpha)` columns. Otherwise, `alpha` and `x` must conform and the result will be the same size.

If requested, `ierr` contains the following status information and is the same size as the result.

0. Normal return.
1. Input error, return `NaN`.
2. Overflow, return `Inf`.
3. Loss of significance by argument reduction results in less than half of machine accuracy.
4. Loss of significance by argument reduction, output may be inaccurate.
5. Error—no computation, algorithm termination condition not met, return `NaN`.

See also: `[bessely]`, page 527, `[besseli]`, page 528, `[besselk]`, page 529, `[besselh]`, page 529.

```
Y = bessely (alpha, x)
Y = bessely (alpha, x, opt)
[Y, ierr] = bessely (...)
```

Compute Bessel functions of the second kind.

The order of the Bessel function `alpha` must be real. The points for evaluation `x` may be complex.

If the optional argument *opt* is 1 or true, the result *Y* is multiplied by $\exp(-\text{abs}(\text{imag}(x)))$.

If *alpha* is a scalar, the result is the same size as *x*. If *x* is a scalar, the result is the same size as *alpha*. If *alpha* is a row vector and *x* is a column vector, the result is a matrix with `length(x)` rows and `length(alpha)` columns. Otherwise, *alpha* and *x* must conform and the result will be the same size.

If requested, *ierr* contains the following status information and is the same size as the result.

0. Normal return.
1. Input error, return NaN.
2. Overflow, return Inf.
3. Loss of significance by argument reduction results in less than half of machine accuracy.
4. Complete loss of significance by argument reduction, return NaN.
5. Error—no computation, algorithm termination condition not met, return NaN.

See also: [\[besselj\]](#), page 527, [\[besseli\]](#), page 528, [\[besselk\]](#), page 529, [\[besselh\]](#), page 529.

```
I = besseli (alpha, x)
I = besseli (alpha, x, opt)
[I, ierr] = besseli (...)
```

Compute modified Bessel functions of the first kind.

The order of the Bessel function *alpha* must be real. The points for evaluation *x* may be complex.

If the optional argument *opt* is 1 or true, the result *I* is multiplied by $\exp(-\text{abs}(\text{real}(x)))$.

If *alpha* is a scalar, the result is the same size as *x*. If *x* is a scalar, the result is the same size as *alpha*. If *alpha* is a row vector and *x* is a column vector, the result is a matrix with `length(x)` rows and `length(alpha)` columns. Otherwise, *alpha* and *x* must conform and the result will be the same size.

If requested, *ierr* contains the following status information and is the same size as the result.

0. Normal return.
1. Input error, return NaN.
2. Overflow, return Inf.
3. Loss of significance by argument reduction results in less than half of machine accuracy.
4. Complete loss of significance by argument reduction, return NaN.
5. Error—no computation, algorithm termination condition not met, return NaN.

See also: [\[besselk\]](#), page 529, [\[besselj\]](#), page 527, [\[bessely\]](#), page 527, [\[besselh\]](#), page 529.

```
K = besselk (alpha, x)
K = besselk (alpha, x, opt)
[K, ierr] = besselk (...)
```

Compute modified Bessel functions of the second kind.

The order of the Bessel function *alpha* must be real. The points for evaluation *x* may be complex.

If the optional argument *opt* is 1 or true, the result *K* is multiplied by `exp (x)`.

If *alpha* is a scalar, the result is the same size as *x*. If *x* is a scalar, the result is the same size as *alpha*. If *alpha* is a row vector and *x* is a column vector, the result is a matrix with `length (x)` rows and `length (alpha)` columns. Otherwise, *alpha* and *x* must conform and the result will be the same size.

If requested, *ierr* contains the following status information and is the same size as the result.

0. Normal return.
1. Input error, return NaN.
2. Overflow, return Inf.
3. Loss of significance by argument reduction results in less than half of machine accuracy.
4. Complete loss of significance by argument reduction, return NaN.
5. Error—no computation, algorithm termination condition not met, return NaN.

See also: [besseli], page 528, [besselj], page 527, [bessely], page 527, [besselh], page 529.

```
H = besselh (alpha, x)
H = besselh (alpha, k, x)
H = besselh (alpha, k, x, opt)
[H, ierr] = besselh (...)
```

Compute Bessel functions of the third kind (Hankel functions).

The order of the Bessel function *alpha* must be real. The kind of Hankel function is specified by *k* and may be either first (*k* = 1) or second (*k* = 2). The default is Hankel functions of the first kind. The points for evaluation *x* may be complex.

If the optional argument *opt* is 1 or true, the result is multiplied by `exp (-I*x)` for *k* = 1 or `exp (I*x)` for *k* = 2.

If *alpha* is a scalar, the result is the same size as *x*. If *x* is a scalar, the result is the same size as *alpha*. If *alpha* is a row vector and *x* is a column vector, the result is a matrix with `length (x)` rows and `length (alpha)` columns. Otherwise, *alpha* and *x* must conform and the result will be the same size.

If requested, *ierr* contains the following status information and is the same size as the result.

0. Normal return.
1. Input error, return NaN.
2. Overflow, return Inf.

3. Loss of significance by argument reduction results in less than half of machine accuracy.
4. Complete loss of significance by argument reduction, return NaN.
5. Error—no computation, algorithm termination condition not met, return NaN.

See also: [besselj], page 527, [bessely], page 527, [besseli], page 528, [besselk], page 529.

beta (a, b)

Compute the Beta function for real inputs a and b .

The Beta function definition is

$$B(a, b) = \frac{\Gamma(a)\Gamma(b)}{\Gamma(a+b)}.$$

The Beta function can grow quite large and it is often more useful to work with the logarithm of the output rather than the function directly. See [betaln], page 531, for computing the logarithm of the Beta function in an efficient manner.

See also: [betaln], page 531, [betainc], page 530, [betaincinv], page 530.

betainc (x, a, b)

betainc (x, a, b, tail)

Compute the incomplete beta function.

This is defined as

$$I_x(a, b) = \frac{1}{B(a, b)} \int_0^x t^{a-1} (1-t)^{b-1} dt$$

with real x in the range $[0,1]$. The inputs a and b must be real and strictly positive (> 0). If one of the inputs is not a scalar then the other inputs must be scalar or of compatible dimensions.

By default, $tail$ is "lower" and the incomplete beta function integrated from 0 to x is computed. If $tail$ is "upper" then the complementary function integrated from x to 1 is calculated. The two choices are related by

$$\text{betainc}(x, a, b, \text{"upper"}) = 1 - \text{betainc}(x, a, b, \text{"lower"}).$$

betainc uses a more sophisticated algorithm than subtraction to get numerically accurate results when the "lower" value is small.

Reference: A. Cuyt, V. Brevik Petersen, B. Verdonk, H. Waadeland, W.B. Jones, *Handbook of Continued Fractions for Special Functions*, ch. 18.

See also: [beta], page 530, [betaincinv], page 530, [betaln], page 531.

betaincinv (y, a, b)

betaincinv (y, a, b, "lower")

betaincinv (y, a, b, "upper")

Compute the inverse of the normalized incomplete beta function.

The normalized incomplete beta function is defined as

$$I_x(a, b) = \frac{1}{B(a, b)} \int_0^x t^{a-1} (1-t)^{b-1} dt$$

If two inputs are scalar, then **betaincinv** (*y, a, b*) is returned for each of the other inputs.

If two or more inputs are not scalar, the sizes of them must agree, and **betaincinv** is applied element-by-element.

The variable *y* must be in the interval [0,1], while *a* and *b* must be real and strictly positive.

By default, *tail* is "lower" and the inverse of the incomplete beta function integrated from 0 to *x* is computed. If *tail* is "upper" then the complementary function integrated from *x* to 1 is inverted.

The function is computed by standard Newton's method, by solving

$$y - I_x(a, b) = 0$$

See also: [\[betainc\]](#), page 530, [\[beta\]](#), page 530, [\[betaln\]](#), page 531.

betaln (a, b)

Compute the natural logarithm of the Beta function for real inputs *a* and *b*.

betaln is defined as

$$\text{betaln}(a, b) = \ln(B(a, b)) \equiv \ln\left(\frac{\Gamma(a)\Gamma(b)}{\Gamma(a+b)}\right).$$

and is calculated in a way to reduce the occurrence of underflow.

The Beta function can grow quite large and it is often more useful to work with the logarithm of the output rather than the function directly.

See also: [\[beta\]](#), page 530, [\[betainc\]](#), page 530, [\[betaincinv\]](#), page 530, [\[gammaln\]](#), page 537.

bincoeff (n, k)

Return the binomial coefficient of *n* and *k*.

The binomial coefficient is defined as

$$\binom{n}{k} = \frac{n(n-1)(n-2)\cdots(n-k+1)}{k!}$$

For example:

```
bincoeff (5, 2)
          => 10
```

In most cases, the **nchoosek** function is faster for small scalar integer arguments. It also warns about loss of precision for big arguments.

See also: [\[nchoosek\]](#), page 737.

commutation_matrix (m, n)

Return the commutation matrix $K_{m,n}$ which is the unique $mn \times mn$ matrix such that $K_{m,n} \cdot \text{vec}(A) = \text{vec}(A^T)$ for all $m \times n$ matrices A .

If only one argument m is given, $K_{m,m}$ is returned.

See Magnus and Neudecker (1988), *Matrix Differential Calculus with Applications in Statistics and Econometrics*.

cosint (x)

Compute the cosine integral function:

$$\text{Ci}(x) = - \int_x^\infty \frac{\cos(t)}{t} dt$$

An equivalent definition is

$$\text{Ci}(x) = \gamma + \log(x) + \int_0^x \frac{\cos(t) - 1}{t} dt$$

Reference:

M. Abramowitz and I.A. Stegun, *Handbook of Mathematical Functions* 1964.

See also: [sinint], page 538, [expint], page 534, [cos], page 513.

duplication_matrix (n)

Return the duplication matrix D_n which is the unique $n^2 \times n(n+1)/2$ matrix such that $D_n * \text{vech}(A) = \text{vec}(A)$ for all symmetric $n \times n$ matrices A .

See Magnus and Neudecker (1988), *Matrix Differential Calculus with Applications in Statistics and Econometrics*.

dawson (z)

Compute the Dawson (scaled imaginary error) function.

The Dawson function is defined as

$$\frac{\sqrt{\pi}}{2} e^{-z^2} \text{erfi}(z) \equiv -i \frac{\sqrt{\pi}}{2} e^{-z^2} \text{erf}(iz)$$

See also: [erfc], page 534, [erf], page 533, [erfcx], page 534, [erfi], page 534, [erfinv], page 534, [erfcinv], page 534.

[sn, cn, dn, err] = ellipj (u, m)
[sn, cn, dn, err] = ellipj (u, m, tol)

Compute the Jacobi elliptic functions sn , cn , and dn of complex argument u and real parameter m .

If m is a scalar, the results are the same size as u . If u is a scalar, the results are the same size as m . If u is a column vector and m is a row vector, the results are matrices with `length (u)` rows and `length (m)` columns. Otherwise, u and m must conform in size and the results will be the same size as the inputs.

The value of u may be complex. The value of m must be $0 \leq m \leq 1$.

The optional input *tol* is currently ignored (MATLAB uses this to allow faster, less accurate approximation).

If requested, *err* contains the following status information and is the same size as the result.

0. Normal return.
1. Error—no computation, algorithm termination condition not met, return NaN.

Reference: Milton Abramowitz and Irene A Stegun, *Handbook of Mathematical Functions*, Chapter 16 (Sections 16.4, 16.13, and 16.15), Dover, 1965.

See also: [ellipke], page 533.

```
k = ellipke (m)
k = ellipke (m, tol)
[k, e] = ellipke (...)
```

Compute complete elliptic integrals of the first $K(m)$ and second $E(m)$ kind.

m must be a scalar or real array with $-Inf \leq m \leq 1$.

The optional input *tol* controls the stopping tolerance of the algorithm and defaults to *eps* (class (*m*)). The tolerance can be increased to compute a faster, less accurate approximation.

When called with one output only elliptic integrals of the first kind are returned.

Mathematical Note:

Elliptic integrals of the first kind are defined as

$$K(m) = \int_0^1 \frac{dt}{\sqrt{(1-t^2)(1-mt^2)}}$$

Elliptic integrals of the second kind are defined as

$$E(m) = \int_0^1 \frac{\sqrt{1-mt^2}}{\sqrt{1-t^2}} dt$$

Reference: Milton Abramowitz and Irene A. Stegun, *Handbook of Mathematical Functions*, Chapter 17, Dover, 1965.

See also: [ellipj], page 532.

erf (z)

Compute the error function.

The error function is defined as

$$\operatorname{erf}(z) = \frac{2}{\sqrt{\pi}} \int_0^z e^{-t^2} dt$$

See also: [erfc], page 534, [erfcx], page 534, [erfi], page 534, [dawson], page 532, [erfinv], page 534, [erfcinv], page 534.

erfc (z)

Compute the complementary error function.

The complementary error function is defined as $1 - \text{erf}(z)$.

See also: [erfcinv], page 534, [erfcx], page 534, [erfi], page 534, [dawson], page 532, [erf], page 533, [erfinv], page 534.

erfcx (z)

Compute the scaled complementary error function.

The scaled complementary error function is defined as

$$e^{z^2} \text{erfc}(z) \equiv e^{z^2} (1 - \text{erf}(z))$$

See also: [erfc], page 534, [erf], page 533, [erfi], page 534, [dawson], page 532, [erfinv], page 534, [erfcinv], page 534.

erfi (z)

Compute the imaginary error function.

The imaginary error function is defined as

$$-i\text{erf}(iz)$$

See also: [erfc], page 534, [erf], page 533, [erfcx], page 534, [dawson], page 532, [erfinv], page 534, [erfcinv], page 534.

erfinv (x)

Compute the inverse error function.

The inverse error function is defined such that

$$\text{erf}(y) == x$$

See also: [erf], page 533, [erfc], page 534, [erfcx], page 534, [erfi], page 534, [dawson], page 532, [erfcinv], page 534.

erfcinv (x)

Compute the inverse complementary error function.

The inverse complementary error function is defined such that

$$\text{erfc}(y) == x$$

See also: [erfc], page 534, [erf], page 533, [erfcx], page 534, [erfi], page 534, [dawson], page 532, [erfinv], page 534.

expint (x)

Compute the exponential integral.

The exponential integral is defined as:

$$E_1(x) = \int_x^\infty \frac{e^{-t}}{t} dt$$

Note: For compatibility, this function uses the MATLAB definition of the exponential integral. Most other sources refer to this particular value as $E_1(x)$, and the exponential integral as

$$\text{Ei}(x) = - \int_{-x}^{\infty} \frac{e^{-t}}{t} dt.$$

The two definitions are related, for positive real values of x , by $E_1(-x) = -\text{Ei}(x) - i\pi$.

References:

M. Abramowitz and I.A. Stegun, *Handbook of Mathematical Functions*, 1964.

N. Bleistein and R.A. Handelsman, *Asymptotic expansions of integrals*, 1986.

See also: [cosint], page 532, [sinint], page 538, [exp], page 509.

gamma (z)

Compute the Gamma function.

The Gamma function is defined as

$$\Gamma(z) = \int_0^{\infty} t^{z-1} e^{-t} dt.$$

Programming Note: The gamma function can grow quite large even for small input values. In many cases it may be preferable to use the natural logarithm of the gamma function (`gammaln`) in calculations to minimize loss of precision. The final result is then `exp(result_using_gammaln)`.

See also: [gammainc], page 535, [gammaln], page 537, [factorial], page 524.

gammainc (x, a)

gammainc (x, a, tail)

Compute the normalized incomplete gamma function.

This is defined as

$$\gamma(x, a) = \frac{1}{\Gamma(a)} \int_0^x t^{a-1} e^{-t} dt$$

with the limiting value of 1 as x approaches infinity. The standard notation is $P(a, x)$, e.g., Abramowitz and Stegun (6.5.1).

If a is scalar, then `gammainc (x, a)` is returned for each element of x and vice versa.

If neither x nor a is scalar then the sizes of x and a must agree, and `gammainc` is applied element-by-element. The elements of a must be non-negative.

By default, `tail` is "lower" and the incomplete gamma function integrated from 0 to x is computed. If `tail` is "upper" then the complementary function integrated from x to infinity is calculated.

If `tail` is "scaledlower", then the lower incomplete gamma function is multiplied by $\Gamma(a + 1) \exp(x)x^{-a}$. If `tail` is "scaledupper", then the upper incomplete gamma function is multiplied by the same quantity.

References:

M. Abramowitz and I.A. Stegun, *Handbook of mathematical functions*, Dover publications, Inc., 1972.

W. Gautschi, *A computational procedure for incomplete gamma functions*, ACM Trans. Math Software, pp. 466–481, Vol 5, No. 4, 2012.

W. H. Press, S. A. Teukolsky, W. T. Vetterling, and B. P. Flannery, *Numerical Recipes in Fortran 77*, ch. 6.2, Vol 1, 1992.

See also: [\[gamma\]](#), page 535, [\[gammaincinv\]](#), page 536, [\[gammaln\]](#), page 537.

gammaincinv (*y, a*)

gammaincinv (*y, a, tail*)

Compute the inverse of the normalized incomplete gamma function.

The normalized incomplete gamma function is defined as

$$\gamma(x, a) = \frac{1}{\Gamma(a)} \int_0^x t^{a-1} e^{-t} dt$$

and **gammaincinv** (**gammainc** (*x, a*), *a*) = *x* for each non-negative value of *x*. If *a* is scalar then **gammaincinv** (*y, a*) is returned for each element of *y* and vice versa.

If neither *y* nor *a* is scalar then the sizes of *y* and *a* must agree, and **gammaincinv** is applied element-by-element. The variable *y* must be in the interval [0, 1] while *a* must be real and positive.

By default, *tail* is "lower" and the inverse of the incomplete gamma function integrated from 0 to *x* is computed. If *tail* is "upper", then the complementary function integrated from *x* to infinity is inverted.

The function is computed with Newton's method by solving

$$y - \gamma(x, a) = 0$$

Reference: A. Gil, J. Segura, and N. M. Temme, *Efficient and accurate algorithms for the computation and inversion of the incomplete gamma function ratios*, SIAM J. Sci. Computing, pp. A2965–A2981, Vol 34, 2012.

See also: [\[gammainc\]](#), page 535, [\[gamma\]](#), page 535, [\[gammaln\]](#), page 537.

l = legendre (*n, x*)

l = legendre (*n, x, normalization*)

Compute the associated Legendre function of degree *n* and order *m* = 0 . . . *n*.

The value *n* must be a real non-negative integer.

x is a vector with real-valued elements in the range [-1, 1].

The optional argument *normalization* may be one of "unnorm", "sch", or "norm". The default if no normalization is given is "unnorm".

When the optional argument *normalization* is "unnorm", compute the associated Legendre function of degree *n* and order *m* and return all values for *m* = 0 . . . *n*. The return value has one dimension more than *x*.

The associated Legendre function of degree *n* and order *m*:

$$P_n^m(x) = (-1)^m (1 - x^2)^{m/2} \frac{d^m}{dx^m} P_n(x)$$

with Legendre polynomial of degree n :

$$P(x) = \frac{1}{2^n n!} \left(\frac{d^n}{dx^n} (x^2 - 1)^n \right)$$

`legendre (3, [-1.0, -0.9, -0.8])` returns the matrix:

x	-1.0	-0.9	-0.8
<hr/>			
m=0	-1.00000	-0.47250	-0.08000
m=1	0.00000	-1.99420	-1.98000
m=2	0.00000	-2.56500	-4.32000
m=3	0.00000	-1.24229	-3.24000

When the optional argument *normalization* is "sch", compute the Schmidt semi-normalized associated Legendre function. The Schmidt semi-normalized associated Legendre function is related to the unnormalized Legendre functions by the following:

For Legendre functions of degree n and order 0:

$$SP_n^0(x) = P_n^0(x)$$

For Legendre functions of degree n and order m :

$$SP_n^m(x) = P_n^m(x)(-1)^m \left(\frac{2(n-m)!}{(n+m)!} \right)^{0.5}$$

When the optional argument *normalization* is "norm", compute the fully normalized associated Legendre function. The fully normalized associated Legendre function is related to the unnormalized associated Legendre functions by the following:

For Legendre functions of degree n and order m

$$NP_n^m(x) = P_n^m(x)(-1)^m \left(\frac{(n+0.5)(n-m)!}{(n+m)!} \right)^{0.5}$$

`gammaln (x)`
`lgamma (x)`

Return the natural logarithm of the gamma function of x .

See also: [\[gamma\]](#), page 535, [\[gammainc\]](#), page 535.

`psi (z)`
`psi (k, z)`

Compute the psi (polygamma) function.

The polygamma functions are the k th derivative of the logarithm of the gamma function. If unspecified, k defaults to zero. A value of zero computes the digamma function, a value of 1, the trigamma function, and so on.

The digamma function is defined:

$$\Psi(z) = \frac{d(\log(\Gamma(z)))}{dx}$$

When computing the digamma function (when k equals zero), z can have any value real or complex value. However, for polygamma functions (k higher than 0), z must be real and non-negative.

See also: [gamma], page 535, [gammainc], page 535, [gammaln], page 537.

sinint (x)

Compute the sine integral function:

$$\text{Si}(x) = \int_0^x \frac{\sin(t)}{t} dt$$

Reference: M. Abramowitz and I.A. Stegun, *Handbook of Mathematical Functions*, 1964.

See also: [cosint], page 532, [expint], page 534, [sin], page 513.

17.7 Rational Approximations

```
s = rat (x)
s = rat (x, tol)
[n, d] = rat (...)
```

Find a rational approximation of x to within the tolerance defined by tol .

If unspecified, the default tolerance is $1e-6 * \text{norm} (x(:), 1)$.

When called with one output argument, return a string containing a continued fraction expansion (multiple terms).

When called with two output arguments, return numeric matrices for the numerator and denominator of a fractional representation of x such that $x = n ./ d$.

For example:

```
s = rat (pi)
⇒ s = 3 + 1/(7 + 1/16)
```

```
[n, d] = rat (pi)
⇒ n = 355
⇒ d = 113
```

```
n / d - pi
⇒ 0.00000026676
```

Programming Note: With one output **rat** produces a string which is a continued fraction expansion. To produce a string which is a simple fraction (one numerator, one denominator) use **rats**.

See also: [rats], page 538, [format], page 252.

rats (x, len)

Convert x into a rational approximation represented as a string.

The string can be converted back into a matrix as follows:

```
r = rats (hilb (4));
x = str2num (r)
```

The optional second argument defines the maximum length of the string representing the elements of x . By default len is 9.

If the length of the smallest possible rational approximation exceeds len , an asterisk (*) padded with spaces will be returned instead.

See also: [\[format\]](#), page 252, [\[rat\]](#), page 538.

17.8 Coordinate Transformations

```
[theta, r] = cart2pol (x, y)
[theta, r, z] = cart2pol (x, y, z)
[theta, r] = cart2pol (C)
[theta, r, z] = cart2pol (C)
P = cart2pol (...)
```

Transform Cartesian coordinates to polar or cylindrical coordinates.

The inputs x , y (, and z) must be the same shape, or scalar. If called with a single matrix argument then each row of C represents the Cartesian coordinate $(x, y$ (, z)).

θ describes the angle relative to the positive x-axis.

r is the distance to the z-axis $(0, 0, z)$.

If only a single return argument is requested then return a matrix P where each row represents one polar/(cylindrical) coordinate $(\theta, \phi$ (, z)).

See also: [\[pol2cart\]](#), page 539, [\[cart2sph\]](#), page 539, [\[sph2cart\]](#), page 540.

```
[x, y] = pol2cart (theta, r)
[x, y, z] = pol2cart (theta, r, z)
[x, y] = pol2cart (P)
[x, y, z] = pol2cart (P)
C = pol2cart (...)
```

Transform polar or cylindrical coordinates to Cartesian coordinates.

The inputs θ , r , (and z) must be the same shape, or scalar. If called with a single matrix argument then each row of P represents the polar/(cylindrical) coordinate $(\theta, r$ (, z)).

θ describes the angle relative to the positive x-axis.

r is the distance to the z-axis $(0, 0, z)$.

If only a single return argument is requested then return a matrix C where each row represents one Cartesian coordinate $(x, y$ (, z)).

See also: [\[cart2pol\]](#), page 539, [\[sph2cart\]](#), page 540, [\[cart2sph\]](#), page 539.

```
[theta, phi, r] = cart2sph (x, y, z)
[theta, phi, r] = cart2sph (C)
S = cart2sph (...)
```

Transform Cartesian coordinates to spherical coordinates.

The inputs x , y , and z must be the same shape, or scalar. If called with a single matrix argument then each row of C represents the Cartesian coordinate (x, y, z) .

θ describes the angle relative to the positive x-axis.

phi is the angle relative to the xy-plane.

r is the distance to the origin (0, 0, 0).

If only a single return argument is requested then return a matrix *S* where each row represents one spherical coordinate (*theta*, *phi*, *r*).

See also: [sph2cart], page 540, [cart2pol], page 539, [pol2cart], page 539.

```
[x, y, z] = sph2cart (theta, phi, r)
[x, y, z] = sph2cart (S)
C = sph2cart (...)
```

Transform spherical coordinates to Cartesian coordinates.

The inputs *theta*, *phi*, and *r* must be the same shape, or scalar. If called with a single matrix argument then each row of *S* represents the spherical coordinate (*theta*, *phi*, *r*).

theta describes the angle relative to the positive x-axis.

phi is the angle relative to the xy-plane.

r is the distance to the origin (0, 0, 0).

If only a single return argument is requested then return a matrix *C* where each row represents one Cartesian coordinate (x, y, z).

See also: [cart2sph], page 539, [pol2cart], page 539, [cart2pol], page 539.

17.9 Mathematical Constants

```
e
e (n)
e (n, m)
e (n, m, k, ...)
e (..., class)
```

Return a scalar, matrix, or N-dimensional array whose elements are all equal to the base of natural logarithms.

The constant *e* satisfies the equation $\log(e) = 1$.

When called with no arguments, return a scalar with the value *e*.

When called with a single argument, return a square matrix with the dimension specified.

When called with more than one scalar argument the first two arguments are taken as the number of rows and columns and any further arguments specify additional matrix dimensions.

The optional argument *class* specifies the return type and may be either "double" or "single".

See also: [log], page 509, [exp], page 509, [pi], page 540, [I], page 541.

```
pi
pi (n)
pi (n, m)
pi (n, m, k, ...)
```

pi (..., class)

Return a scalar, matrix, or N-dimensional array whose elements are all equal to the ratio of the circumference of a circle to its diameter(π).

Internally, **pi** is computed as '`4.0 * atan (1.0)`'.

When called with no arguments, return a scalar with the value of π .

When called with a single argument, return a square matrix with the dimension specified.

When called with more than one scalar argument the first two arguments are taken as the number of rows and columns and any further arguments specify additional matrix dimensions.

The optional argument *class* specifies the return type and may be either "**double**" or "**single**".

See also: [e], page 540, [I], page 541.

I**I (n)****I (n, m)****I (n, m, k, ...)****I (..., class)**

Return a scalar, matrix, or N-dimensional array whose elements are all equal to the pure imaginary unit, defined as $\sqrt{-1}$.

I, and its equivalents **i**, **j**, and **J**, are functions so any of the names may be reused for other purposes (such as **i** for a counter variable).

When called with no arguments, return a scalar with the value *i*.

When called with a single argument, return a square matrix with the dimension specified.

When called with more than one scalar argument the first two arguments are taken as the number of rows and columns and any further arguments specify additional matrix dimensions.

The optional argument *class* specifies the return type and may be either "**double**" or "**single**".

See also: [e], page 540, [pi], page 540, [log], page 509, [exp], page 509.

Inf**Inf (n)****Inf (n, m)****Inf (n, m, k, ...)****Inf (..., class)**

Return a scalar, matrix or N-dimensional array whose elements are all equal to the IEEE representation for positive infinity.

Infinity is produced when results are too large to be represented using the IEEE floating point format for numbers. Two common examples which produce infinity are division by zero and overflow.

[1/0 e^800]

⇒ Inf Inf

When called with no arguments, return a scalar with the value ‘`Inf`’.

When called with a single argument, return a square matrix with the dimension specified.

When called with more than one scalar argument the first two arguments are taken as the number of rows and columns and any further arguments specify additional matrix dimensions.

The optional argument `class` specifies the return type and may be either “`double`” or “`single`”.

See also: [\[isinf\]](#), page 476, [\[NaN\]](#), page 542.

`NaN`

`NaN (n)`

`NaN (n, m)`

`NaN (n, m, k, ...)`

`NaN (... , class)`

Return a scalar, matrix, or N-dimensional array whose elements are all equal to the IEEE symbol `NaN` (Not a Number).

`NaN` is the result of operations which do not produce a well defined numerical result. Common operations which produce a `NaN` are arithmetic with infinity ($\infty - \infty$), zero divided by zero (0/0), and any operation involving another `NaN` value (5 + `NaN`).

Note that `NaN` always compares not equal to `NaN` (`NaN != NaN`). This behavior is specified by the IEEE standard for floating point arithmetic. To find `NaN` values, use the `isnan` function.

When called with no arguments, return a scalar with the value ‘`NaN`’.

When called with a single argument, return a square matrix with the dimension specified.

When called with more than one scalar argument the first two arguments are taken as the number of rows and columns and any further arguments specify additional matrix dimensions.

The optional argument `class` specifies the return type and may be either “`double`” or “`single`”.

See also: [\[isnan\]](#), page 476, [\[Inf\]](#), page 541.

`eps`

`eps (x)`

`eps (n, m)`

`eps (n, m, k, ...)`

`eps (... , class)`

Return a scalar, matrix or N-dimensional array whose elements are all `eps`, the machine precision.

More precisely, `eps` is the relative spacing between any two adjacent numbers in the machine’s floating point system. This number is obviously system dependent. On machines that support IEEE floating point arithmetic, `eps` is approximately 2.2204×10^{-16} for double precision and 1.1921×10^{-7} for single precision.

When called with no arguments, return a scalar with the value `eps` (1.0).

Given a single argument `x`, return the distance between `x` and the next largest value.

When called with more than one argument the first two arguments are taken as the number of rows and columns and any further arguments specify additional matrix dimensions. The optional argument `class` specifies the return type and may be either "double" or "single".

See also: `[realmax]`, page 543, `[realmin]`, page 543, `[intmax]`, page 55, `[flintmax]`, page 56.

```
realmax
realmax (n)
realmax (n, m)
realmax (n, m, k, ...)
realmax (..., class)
```

Return a scalar, matrix, or N-dimensional array whose elements are all equal to the largest floating point number that is representable.

The actual value is system dependent. On machines that support IEEE floating point arithmetic, `realmax` is approximately 1.7977×10^{308} for double precision and 3.4028×10^{38} for single precision.

When called with no arguments, return a scalar with the value `realmax ("double")`.

When called with a single argument, return a square matrix with the dimension specified.

When called with more than one scalar argument the first two arguments are taken as the number of rows and columns and any further arguments specify additional matrix dimensions.

The optional argument `class` specifies the return type and may be either "double" or "single".

See also: `[realmin]`, page 543, `[intmax]`, page 55, `[flintmax]`, page 56, `[eps]`, page 542.

```
realmin
realmin (n)
realmin (n, m)
realmin (n, m, k, ...)
realmin (..., class)
```

Return a scalar, matrix, or N-dimensional array whose elements are all equal to the smallest normalized floating point number that is representable.

The actual value is system dependent. On machines that support IEEE floating point arithmetic, `realmin` is approximately 2.2251×10^{-308} for double precision and 1.1755×10^{-38} for single precision.

When called with no arguments, return a scalar with the value `realmin ("double")`.

When called with a single argument, return a square matrix with the dimension specified.

When called with more than one scalar argument the first two arguments are taken as the number of rows and columns and any further arguments specify additional matrix dimensions.

The optional argument *class* specifies the return type and may be either "double" or "single".

See also: [\[realmax\]](#), page 543, [\[intmin\]](#), page 56, [\[eps\]](#), page 542.

18 Linear Algebra

This chapter documents the linear algebra functions provided in Octave. Reference material for many of these functions may be found in Golub and Van Loan, *Matrix Computations*, 2nd Ed., Johns Hopkins, 1989, and in the *LAPACK Users' Guide*, SIAM, 1992. The *LAPACK Users' Guide* is available at: <http://www.netlib.org/lapack/lug/>

A common text for engineering courses is G. Strang, *Linear Algebra and Its Applications*, 4th Edition. It has become a widespread reference for linear algebra. An alternative is P. Lax *Linear Algebra and Its Applications*, and also is a good choice. It claims to be suitable for high school students with substantial mathematical interests as well as first-year undergraduates.

18.1 Techniques Used for Linear Algebra

Octave includes a polymorphic solver that selects an appropriate matrix factorization depending on the properties of the matrix itself. Generally, the cost of determining the matrix type is small relative to the cost of factorizing the matrix itself. In any case the matrix type is cached once it is calculated so that it is not re-determined each time it is used in a linear equation.

The selection tree for how the linear equation is solved or a matrix inverse is formed is given by:

1. If the matrix is upper or lower triangular sparse use a forward or backward substitution using the LAPACK xTRTRS function, and goto 4.
2. If the matrix is square, Hermitian with a real positive diagonal, attempt Cholesky factorization using the LAPACK xPOTRF function.
3. If the Cholesky factorization failed or the matrix is not Hermitian with a real positive diagonal, and the matrix is square, factorize using the LAPACK xGETRF function.
4. If the matrix is not square, or any of the previous solvers flags a singular or near singular matrix, find a least squares solution using the LAPACK xGELSD function.

The user can force the type of the matrix with the `matrix_type` function. This overcomes the cost of discovering the type of the matrix. However, it should be noted that identifying the type of the matrix incorrectly will lead to unpredictable results, and so `matrix_type` should be used with care.

It should be noted that the test for whether a matrix is a candidate for Cholesky factorization, performed above, and by the `matrix_type` function, does not make certain that the matrix is Hermitian. However, the attempt to factorize the matrix will quickly detect a non-Hermitian matrix.

18.2 Basic Matrix Functions

```
AA = balance (A)
AA = balance (A, opt)
[DD, AA] = balance (A, opt)
[D, P, AA] = balance (A, opt)
```

[CC, DD, AA, BB] = balance (A, B, opt)

Balance the matrix A to reduce numerical errors in future calculations.

Compute $AA = DD \setminus A * DD$ in which AA is a matrix whose row and column norms are roughly equal in magnitude, and $DD = P * D$, in which P is a permutation matrix and D is a diagonal matrix of powers of two. This allows the equilibration to be computed without round-off. Results of eigenvalue calculation are typically improved by balancing first.

If two output values are requested, **balance** returns the diagonal D and the permutation P separately as vectors. In this case, $DD = eye(n)(:,P) * diag(D)$, where n is the matrix size.

If four output values are requested, compute $AA = CC*A*DD$ and $BB = CC*B*DD$, in which AA and BB have nonzero elements of approximately the same magnitude and CC and DD are permuted diagonal matrices as in DD for the algebraic eigenvalue problem.

The eigenvalue balancing option *opt* may be one of:

"noperm", "S"

Scale only; do not permute.

"noscal", "P"

Permute only; do not scale.

Algebraic eigenvalue balancing uses standard LAPACK routines.

Generalized eigenvalue problem balancing uses Ward's algorithm (SIAM Journal on Scientific and Statistical Computing, 1981).

bw = bandwidth (A, type)
[lower, upper] = bandwidth (A)

Compute the bandwidth of A .

The *type* argument is the string "**lower**" for the lower bandwidth and "**upper**" for the upper bandwidth. If no *type* is specified return both the lower and upper bandwidth of A .

The lower/upper bandwidth of a matrix is the number of subdiagonals/superdiagonals with nonzero entries.

See also: [\[isbanded\]](#), page 65, [\[isdiag\]](#), page 65, [\[istril\]](#), page 65, [\[istriu\]](#), page 66.

cond (A)
cond (A, p)

Compute the p -norm condition number of a matrix with respect to inversion.

cond (A) is defined as $\| A \|_p * \| A^{-1} \|_p$.

By default, $p = 2$ is used which implies a (relatively slow) singular value decomposition. Other possible selections are $p = 1$, **Inf**, "**fro**" which are generally faster. See **norm** for a full discussion of possible p values.

The condition number of a matrix quantifies the sensitivity of the matrix inversion operation when small changes are made to matrix elements. Ideally the condition number will be close to 1. When the number is large this indicates small changes (such as underflow or round-off error) will produce large changes in the resulting

output. In such cases the solution results from numerical computing are not likely to be accurate.

See also: [condeig], page 640, [rcond], page 554, [condeig], page 547, [norm], page 551, [svd], page 565.

```
c = condeig (a)
[v, lambda, c] = condeig (a)
```

Compute condition numbers of a matrix with respect to eigenvalues.

The condition numbers are the reciprocals of the cosines of the angles between the left and right eigenvectors; Large values indicate that the matrix has multiple distinct eigenvalues.

The input *a* must be a square numeric matrix.

The outputs are:

- *c* is a vector of condition numbers for the eigenvalues of *a*.
- *v* is the matrix of right eigenvectors of *a*. The result is equivalent to calling [*v*, *lambda*] = eig (*a*).
- *lambda* is the diagonal matrix of eigenvalues of *a*. The result is equivalent to calling [*v*, *lambda*] = eig (*a*).

Example

```
a = [1, 2; 3, 4];
c = condeig (a)
⇒ c =
    1.0150
    1.0150
```

See also: [eig], page 547, [cond], page 546, [balance], page 545.

```
det (A)
[d, rcond] = det (A)
```

Compute the determinant of *A*.

Return an estimate of the reciprocal condition number if requested.

Programming Notes: Routines from LAPACK are used for full matrices and code from UMFPACK is used for sparse matrices.

The determinant should not be used to check a matrix for singularity. For that, use any of the condition number functions: cond, condest, rcond.

See also: [cond], page 546, [condeig], page 640, [rcond], page 554.

```
lambda = eig (A)
lambda = eig (A, B)
[V, lambda] = eig (A)
[V, lambda] = eig (A, B)
[V, lambda, W] = eig (A)
[V, lambda, W] = eig (A, B)
[...] = eig (A, balanceOption)
[...] = eig (A, B, algorithm)
```

`[...] = eig (... , eigvalOption)`

Compute the eigenvalues (*lambda*) and optionally the right eigenvectors (*V*) and the left eigenvectors (*W*) of a matrix or pair of matrices.

The flag *balanceOption* can be one of:

"**balance**" (default)

Preliminary balancing is on.

"**nobalance**"

Disables preliminary balancing.

The flag *eigvalOption* can be one of:

"**matrix**" Return the eigenvalues in a diagonal matrix. (default if 2 or 3 outputs are requested)

"**vector**" Return the eigenvalues in a column vector. (default if only 1 output is requested, e.g., *lambda* = `eig (A)`)

The flag *algorithm* can be one of:

"**chol**" Use the Cholesky factorization of *B*. (default if *A* is symmetric (Hermitian) and *B* is symmetric (Hermitian) positive definite)

"**qz**" Use the QZ algorithm. (used whenever *A* or *B* are not symmetric)

both are symmetric

no flag

at least one is not symmetric

"**chol**"

"**qz**"

chol

"**chol**"

"**qz**"

qz

"**qz**"

The eigenvalues returned by `eig` are not ordered.

See also: [eigs], page 644, [svd], page 565.

`G = givens (x, y)`

`[c, s] = givens (x, y)`

Compute the Givens rotation matrix *G*.

The Givens matrix is a 2×2 orthogonal matrix

$$G = \begin{bmatrix} c & s \\ -s' & c \end{bmatrix}$$

such that

$$G \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} * \\ 0 \end{bmatrix}$$

with *x* and *y* scalars.

If two output arguments are requested, return the factors *c* and *s* rather than the Givens rotation matrix.

For example:

```
givens (1, 1)
⇒ 0.70711  0.70711
    -0.70711  0.70711
```

Note: The Givens matrix represents a counterclockwise rotation of a 2-D plane and can be used to introduce zeros into a matrix prior to complete factorization.

See also: [planerot], page 549, [qr], page 559.

```
S = gsvd (A, B)
[U, V, X, C, S] = gsvd (A, B)
[U, V, X, C, S] = gsvd (A, B, 0)
```

Compute the generalized singular value decomposition of (A, B) .

The generalized singular value decomposition is defined by the following relations:

$$\begin{aligned} A &= UCX^\dagger \\ B &= VSX^\dagger \\ C^\dagger C + S^\dagger S &= \text{eye}(columns(A)) \end{aligned}$$

The function `gsvd` normally returns just the vector of generalized singular values

$$\sqrt{\frac{\text{diag}(C^\dagger C)}{\text{diag}(S^\dagger S)}}$$

If asked for five return values, it also computes U , V , X , and C .

If the optional third input is present, `gsvd` constructs the "economy-sized" decomposition where the number of columns of U , V and the number of rows of C , S is less than or equal to the number of columns of A . This option is not yet implemented.

Programming Note: the code is a wrapper to the corresponding LAPACK `dggsvd` and `zggsvd` routines.

See also: [\[svd\]](#), page 565.

`[G, y] = planerot (x)`

Compute the Givens rotation matrix for the two-element column vector x .

The Givens matrix is a 2×2 orthogonal matrix

$$G = \begin{bmatrix} c & s \\ -s' & c \end{bmatrix}$$

such that

$$G \begin{bmatrix} x(1) \\ x(2) \end{bmatrix} = \begin{bmatrix} * \\ 0 \end{bmatrix}$$

Note: The Givens matrix represents a counterclockwise rotation of a 2-D plane and can be used to introduce zeros into a matrix prior to complete factorization.

See also: [\[givens\]](#), page 548, [\[qr\]](#), page 559.

```
x = inv (A)
[x, rcond] = inv (A)
```

Compute the inverse of the square matrix A .

Return an estimate of the reciprocal condition number if requested, otherwise warn of an ill-conditioned matrix if the reciprocal condition number is small.

In general it is best to avoid calculating the inverse of a matrix directly. For example, it is both faster and more accurate to solve systems of equations ($A^*x = b$) with $y = A \setminus b$, rather than $y = \text{inv}(A) * b$.

If called with a sparse matrix, then in general x will be a full matrix requiring significantly more storage. Avoid forming the inverse of a sparse matrix if possible.

See also: [\[ldivide\]](#), page 150, [\[rdivide\]](#), page 151, [\[pinv\]](#), page 553.

```
x = linsolve (A, b)
x = linsolve (A, b, opts)
[x, R] = linsolve (...)

Solve the linear system A*x = b.
```

With no options, this function is equivalent to the left division operator (`x = A \ b`) or the matrix-left-divide function (`x = mldivide (A, b)`).

Octave ordinarily examines the properties of the matrix `A` and chooses a solver that best matches the matrix. By passing a structure `opts` to `linsolve` you can inform Octave directly about the matrix `A`. In this case Octave will skip the matrix examination and proceed directly to solving the linear system.

Warning: If the matrix `A` does not have the properties listed in the `opts` structure then the result will not be accurate AND no warning will be given. When in doubt, let Octave examine the matrix and choose the appropriate solver as this step takes little time and the result is cached so that it is only done once per linear system.

Possible `opts` fields (set value to true/false):

LT	<code>A</code> is lower triangular
UT	<code>A</code> is upper triangular
UHESS	<code>A</code> is upper Hessenberg (currently makes no difference)
SYM	<code>A</code> is symmetric or complex Hermitian (currently makes no difference)
POSDEF	<code>A</code> is positive definite
RECT	<code>A</code> is general rectangular (currently makes no difference)
TRANSA	Solve <code>A'*x = b</code> if true rather than <code>A*x = b</code>

The optional second output `R` is the inverse condition number of `A` (zero if matrix is singular).

See also: [\[mldivide\]](#), page 150, [\[matrix_type\]](#), page 550, [\[rcond\]](#), page 554.

```
type = matrix_type (A)
type = matrix_type (A, "nocompute")
A = matrix_type (A, type)
A = matrix_type (A, "upper", perm)
A = matrix_type (A, "lower", perm)
A = matrix_type (A, "banded", nl, nu)
```

Identify the matrix type or mark a matrix as a particular type.

This allows more rapid solutions of linear equations involving `A` to be performed.

Called with a single argument, `matrix_type` returns the type of the matrix and caches it for future use.

Called with more than one argument, `matrix_type` allows the type of the matrix to be defined.

If the option "`nocompute`" is given, the function will not attempt to guess the type if it is still unknown. This is useful for debugging purposes.

The possible matrix types depend on whether the matrix is full or sparse, and can be one of the following

- "unknown"**
Remove any previously cached matrix type, and mark type as unknown.
- "full"** Mark the matrix as full.
- "positive definite"**
Probable full positive definite matrix.
- "diagonal"**
Diagonal matrix. (Sparse matrices only)
- "permuted diagonal"**
Permuted Diagonal matrix. The permutation does not need to be specifically indicated, as the structure of the matrix explicitly gives this. (Sparse matrices only)
- "upper"** Upper triangular. If the optional third argument *perm* is given, the matrix is assumed to be a permuted upper triangular with the permutations defined by the vector *perm*.
- "lower"** Lower triangular. If the optional third argument *perm* is given, the matrix is assumed to be a permuted lower triangular with the permutations defined by the vector *perm*.
- "banded"**
- "banded positive definite"**
Banded matrix with the band size of *nl* below the diagonal and *nu* above it. If *nl* and *nu* are 1, then the matrix is tridiagonal and treated with specialized code. In addition the matrix can be marked as probably a positive definite. (Sparse matrices only)
- "singular"**
The matrix is assumed to be singular and will be treated with a minimum norm solution.

Note that the matrix type will be discovered automatically on the first attempt to solve a linear equation involving *A*. Therefore `matrix_type` is only useful to give Octave hints of the matrix type. Incorrectly defining the matrix type will result in incorrect results from solutions of linear equations; it is entirely **the responsibility of the user** to correctly identify the matrix type.

Also, the test for positive definiteness is a low-cost test for a Hermitian matrix with a real positive diagonal. This does not guarantee that the matrix is positive definite, but only that it is a probable candidate. When such a matrix is factorized, a Cholesky factorization is first attempted, and if that fails the matrix is then treated with an LU factorization. Once the matrix has been factorized, `matrix_type` will return the correct classification of the matrix.

```
norm (A)
norm (A, p)
```

norm (A, p, opt)

Compute the p-norm of the matrix A.

If the second argument is not given, p = 2 is used.

If A is a matrix (or sparse matrix):

p = 1 1-norm, the largest column sum of the absolute values of A.

p = 2 Largest singular value of A.

p = Inf or "inf"

Infinity norm, the largest row sum of the absolute values of A.

p = "fro"

Frobenius norm of A, sqrt (sum (diag (A' * A))).

other p, p > 1

maximum norm (A*x, p) such that norm (x, p) == 1

If A is a vector or a scalar:

p = Inf or "inf"

max (abs (A)).

p = -Inf min (abs (A)).

p = "fro"

Frobenius norm of A, sqrt (sumsq (abs (A))).

p = 0 Hamming norm—the number of nonzero elements.

other p, p > 1

p-norm of A, (sum (abs (A) .^ p)) ^ (1/p).

other p, p < 1

the p-pseudonorm defined as above.

If opt is the value "rows", treat each row as a vector and compute its norm. The result is returned as a column vector. Similarly, if opt is "columns" or "cols" then compute the norms of each column and return a row vector.

See also: [normest], page 639, [normest1], page 639, [vecnorm], page 554, [cond], page 546, [svd], page 565.

null (A)**null (A, tol)**

Return an orthonormal basis of the null space of A.

The dimension of the null space is taken as the number of singular values of A not greater than tol. If the argument tol is missing, it is computed as

max (size (A)) * max (svd (A)) * eps

See also: [orth], page 552.

orth (A)**orth (A, tol)**

Return an orthonormal basis of the range space of A.

The dimension of the range space is taken as the number of singular values of A greater than tol . If the argument tol is missing, it is computed as

```
max (size (A)) * max (svd (A)) * eps
```

See also: [null], page 552.

[y , h] = **mgorth** (x , v)

Orthogonalize a given column vector x with respect to a set of orthonormal vectors comprising the columns of v using the modified Gram-Schmidt method.

On exit, y is a unit vector such that:

```
norm (y) = 1
v' * y = 0
x = [v, y]*h'
```

pinv (x)

pinv (x , tol)

Return the Moore-Penrose pseudoinverse of x .

Singular values less than tol are ignored.

If the second argument is omitted, it is taken to be

```
tol = max ([rows(x), columns(x)]) * norm (x) * eps
```

See also: [inv], page 549, [ldivide], page 150.

rank (A)

rank (A , tol)

Compute the rank of matrix A , using the singular value decomposition.

The rank is taken to be the number of singular values of A that are greater than the specified tolerance tol . If the second argument is omitted, it is taken to be

```
tol = max (size (A)) * sigma(1) * eps;
```

where **eps** is machine precision and **sigma(1)** is the largest singular value of A .

The rank of a matrix is the number of linearly independent rows or columns and equals the dimension of the row and column space. The function **orth** may be used to compute an orthonormal basis of the column space.

For testing if a system $A*x = b$ of linear equations is solvable, one can use

```
rank (A) == rank ([A b])
```

In this case, $x = A \setminus b$ finds a particular solution x . The general solution is x plus the null space of matrix A . The function **null** may be used to compute a basis of the null space.

Example:

```
A = [1 2 3
      4 5 6
      7 8 9];
rank (A)
⇒ 2
```

In this example, the number of linearly independent rows is only 2 because the final row is a linear combination of the first two rows:

```
A(3,:) == -A(1,:)+2*A(2,:)
```

See also: [null], page 552, [orth], page 552, [sprank], page 642, [svd], page 565, [eps], page 542.

c = rcond (A)

Compute the 1-norm estimate of the reciprocal condition number as returned by LAPACK.

If the matrix is well-conditioned then *c* will be near 1 and if the matrix is poorly conditioned it will be close to 0.

The matrix *A* must not be sparse. If the matrix is sparse then **condest (A)** or **rcond (full (A))** should be used instead.

See also: [cond], page 546, [condest], page 640.

trace (A)

Compute the trace of *A*, the sum of the elements along the main diagonal.

The implementation is straightforward: **sum (diag (A))**.

See also: [eig], page 547.

rref (A)

rref (A, tol)

[r, k] = rref (...)

Return the reduced row echelon form of *A*.

tol defaults to **eps * max (size (A)) * norm (A, inf)**.

The optional return argument *k* contains the vector of "bound variables", which are those columns on which elimination has been performed.

n = vecnorm (A)

n = vecnorm (A, p)

n = vecnorm (A, p, dim)

Return the *p*-norm of the elements of *A* along dimension *dim*.

The *p*-norm of a vector is defined as

$$\|A\|_p = \left[\sum_{i=1}^N |A_i|^p \right]^{1/p}$$

If *p* is omitted it defaults to 2 (Euclidean norm). *p* can be **Inf** (absolute value of largest element).

If *dim* is omitted the first non-singleton dimension is used.

See also: [norm], page 551.

18.3 Matrix Factorizations

```
R = chol (A)
[R, p] = chol (A)
[R, p, Q] = chol (A)
[R, p, Q] = chol (A, "vector")
[L, ...] = chol (..., "lower")
[R, ...] = chol (..., "upper")
```

Compute the upper Cholesky factor, R , of the real symmetric or complex Hermitian positive definite matrix A .

The upper Cholesky factor R is computed by using the upper triangular part of matrix A and is defined by $R^T R = A$.

Calling `chol` using the optional "`upper`" flag has the same behavior. In contrast, using the optional "`lower`" flag, `chol` returns the lower triangular factorization, computed by using the lower triangular part of matrix A , such that $LL^T = A$.

Called with one output argument `chol` fails if matrix A is not positive definite. Note that if matrix A is not real symmetric or complex Hermitian then the lower triangular part is considered to be the (complex conjugate) transpose of the upper triangular part, or vice versa, given the "`lower`" flag.

Called with two or more output arguments p flags whether the matrix A was positive definite and `chol` does not fail. A zero value of p indicates that matrix A is positive definite and R gives the factorization. Otherwise, p will have a positive value.

If called with three output arguments matrix A must be sparse and a sparsity preserving row/column permutation is applied to matrix A prior to the factorization. That is R is the factorization of $A(Q, Q)$ such that $R^T R = Q^T A Q$.

The sparsity preserving permutation is generally returned as a matrix. However, given the optional flag "`vector`", Q will be returned as a vector such that $R^T R = A(Q, Q)$. In general the lower triangular factorization is significantly faster for sparse matrices.

See also: [\[hess\]](#), page 557, [\[lu\]](#), page 557, [\[qr\]](#), page 559, [\[qz\]](#), page 562, [\[schur\]](#), page 564, [\[svd\]](#), page 565, [\[ichol\]](#), page 652, [\[cholinv\]](#), page 555, [\[chol2inv\]](#), page 555, [\[cholupdate\]](#), page 556, [\[cholinsert\]](#), page 556, [\[choldelete\]](#), page 556, [\[cholshift\]](#), page 556.

`cholinv (A)`

Compute the inverse of the symmetric positive definite matrix A using the Cholesky factorization.

See also: [\[chol\]](#), page 555, [\[chol2inv\]](#), page 555, [\[inv\]](#), page 549.

`chol2inv (U)`

Invert a symmetric, positive definite square matrix from its Cholesky decomposition, U .

Note that U should be an upper-triangular matrix with positive diagonal elements. `chol2inv (U)` provides `inv (U'*U)` but it is much faster than using `inv`.

See also: [\[chol\]](#), page 555, [\[cholinv\]](#), page 555, [\[inv\]](#), page 549.

[R1, info] = cholupdate (R, u, op)

Update or downdate a Cholesky factorization.

Given an upper triangular matrix R and a column vector u , attempt to determine another upper triangular matrix $R1$ such that

- $R1' * R1 = R' * R + u * u'$ if op is "+"
- $R1' * R1 = R' * R - u * u'$ if op is "-"

If op is "-", $info$ is set to

- 0 if the downdate was successful,
- 1 if $R' * R - u * u'$ is not positive definite,
- 2 if R is singular.

If $info$ is not present, an error message is printed in cases 1 and 2.

See also: [chol], page 555, [cholinsert], page 556, [choldelete], page 556, [cholshift], page 556.

R1 = cholinsert (R, j, u)

[R1, info] = cholinsert (R, j, u)

Update a Cholesky factorization given a row or column to insert in the original factored matrix.

Given a Cholesky factorization of a real symmetric or complex Hermitian positive definite matrix $A = R' * R$, R upper triangular, return the Cholesky factorization of $A1$, where $A1(p,p) = A$, $A1(:,j) = A1(j,:)' = u$ and $p = [1:j-1,j+1:n+1]$. $u(j)$ should be positive.

On return, $info$ is set to

- 0 if the insertion was successful,
- 1 if $A1$ is not positive definite,
- 2 if R is singular.

If $info$ is not present, an error message is printed in cases 1 and 2.

See also: [chol], page 555, [cholupdate], page 556, [choldelete], page 556, [cholshift], page 556.

R1 = choldelete (R, j)

Update a Cholesky factorization given a row or column to delete from the original factored matrix.

Given a Cholesky factorization of a real symmetric or complex Hermitian positive definite matrix $A = R' * R$, R upper triangular, return the Cholesky factorization of $A(p,p)$, where $p = [1:j-1,j+1:n+1]$.

See also: [chol], page 555, [cholupdate], page 556, [cholinsert], page 556, [cholshift], page 556.

R1 = cholshift (R, i, j)

Update a Cholesky factorization given a range of columns to shift in the original factored matrix.

Given a Cholesky factorization of a real symmetric or complex Hermitian positive definite matrix $A = R'*R$, R upper triangular, return the Cholesky factorization of $A(p,p)$, where p is the permutation

$p = [1:i-1, \text{shift}(i:j, 1), j+1:n]$ if $i < j$

or

$p = [1:j-1, \text{shift}(j:i,-1), i+1:n]$ if $j < i$.

See also: [\[chol\]](#), page 555, [\[cholupdate\]](#), page 556, [\[cholinsert\]](#), page 556, [\[choldelete\]](#), page 556.

`H = hess (A)`

`[P, H] = hess (A)`

Compute the Hessenberg decomposition of the matrix A .

The Hessenberg decomposition is

$$A = PHP^T$$

where P is a square unitary matrix ($P^T P = I$), and H is upper Hessenberg ($H_{i,j} = 0, \forall i > j + 1$).

The Hessenberg decomposition is usually used as the first step in an eigenvalue computation, but has other applications as well (see Golub, Nash, and Van Loan, IEEE Transactions on Automatic Control, 1979).

See also: [\[eig\]](#), page 547, [\[chol\]](#), page 555, [\[lu\]](#), page 557, [\[qr\]](#), page 559, [\[qz\]](#), page 562, [\[schur\]](#), page 564, [\[svd\]](#), page 565.

`[L, U] = lu (A)`

`[L, U, P] = lu (A)`

`[L, U, P, Q] = lu (S)`

`[L, U, P, Q, R] = lu (S)`

`[...] = lu (S, thres)`

`y = lu (...)`

`[...] = lu (... , "vector")`

Compute the LU decomposition of A .

If A is full then subroutines from LAPACK are used, and if A is sparse then UMFPACK is used.

The result is returned in a permuted form, according to the optional return value P . For example, given the matrix $a = [1, 2; 3, 4]$,

`[l, u, p] = lu (a)`

returns

```

l =
1.00000 0.00000
0.33333 1.00000

u =
3.00000 4.00000
0.00000 0.66667

p =
0 1
1 0

```

The matrix is not required to be square.

When called with two or three output arguments and a sparse input matrix, `lu` does not attempt to perform sparsity preserving column permutations. Called with a fourth output argument, the sparsity preserving column transformation Q is returned, such that $P * A * Q = L * U$. This is the **preferred** way to call `lu` with sparse input matrices.

Called with a fifth output argument and a sparse input matrix, `lu` attempts to use a scaling factor R on the input matrix such that $P * (R \setminus A) * Q = L * U$. This typically leads to a sparser and more stable factorization.

An additional input argument `thres`, that defines the pivoting threshold can be given. `thres` can be a scalar, in which case it defines the UMFPACK pivoting tolerance for both symmetric and unsymmetric cases. If `thres` is a 2-element vector, then the first element defines the pivoting tolerance for the unsymmetric UMFPACK pivoting strategy and the second for the symmetric strategy. By default, the values defined by `spparms` are used ([0.1, 0.001]).

Given the string argument "vector", `lu` returns the values of P and Q as vector values, such that for full matrix, $A(P, :) = L * U$, and $R(P, :) * A(:, Q) = L * U$.

With two output arguments, returns the permuted forms of the upper and lower triangular matrices, such that $A = L * U$. With one output argument `y`, then the matrix returned by the LAPACK routines is returned. If the input matrix is sparse then the matrix L is embedded into U to give a return value similar to the full case. For both full and sparse matrices, `lu` loses the permutation information.

See also: [\[luupdate\]](#), page 558, [\[ilu\]](#), page 653, [\[chol\]](#), page 555, [\[hess\]](#), page 557, [\[qr\]](#), page 559, [\[qz\]](#), page 562, [\[schur\]](#), page 564, [\[svd\]](#), page 565.

```
[L, U] = luupdate (L, U, x, y)
[L, U, P] = luupdate (L, U, P, x, y)
```

Given an LU factorization of a real or complex matrix $A = L * U$, L lower unit trapezoidal and U upper trapezoidal, return the LU factorization of $A + x * y.'$, where x and y are column vectors (rank-1 update) or matrices with equal number of columns (rank-k update).

Optionally, row-pivoted updating can be used by supplying a row permutation (pivoting) matrix P ; in that case, an updated permutation matrix is returned. Note that if L , U , P is a pivoted LU factorization as obtained by `lu`:

`[L, U, P] = lu (A);`

then a factorization of $A + x * y.$ ' can be obtained either as

`[L1, U1] = lu (L, U, P*x, y)`

or

`[L1, U1, P1] = lu (L, U, P, x, y)`

The first form uses the unpivoted algorithm, which is faster, but less stable. The second form uses a slower pivoted algorithm, which is more stable.

The matrix case is done as a sequence of rank-1 updates; thus, for large enough k , it will be both faster and more accurate to recompute the factorization from scratch.

See also: [\[lu\]](#), page 557, [\[cholupdate\]](#), page 556, [\[qrupdate\]](#), page 561.

```
[Q, R] = qr (A)
[Q, R, P] = qr (A) # non-sparse A
X = qr (A) # non-sparse A
R = qr (A) # sparse A
[C, R] = qr (A, B)
[...] = qr (... , 0)
[...] = qr (... , "vector")
[...] = qr (... , "matrix")
```

Compute the QR factorization of A , using standard LAPACK subroutines.

The QR factorization is $QR = A$ where Q is an orthogonal matrix and R is upper triangular.

For example, given the matrix $A = [1, 2; 3, 4]$,

`[Q, R] = qr (A)`

returns

`Q =`

```
-0.31623 -0.94868
-0.94868 0.31623
```

`R =`

```
-3.16228 -4.42719
0.00000 -0.63246
```

which multiplied together return the original matrix

```
Q * R
⇒
1.0000 2.0000
3.0000 4.0000
```

If just a single return value is requested then it is either R , if A is sparse, or X , such that $R = \text{triu}(X)$ if A is full. (Note: unlike most commands, the single return value is not the first return value when multiple values are requested.)

If the matrix A is full, and a third output P is requested, then `qr` calculates the permuted QR factorization $QR = AP$ where Q is an orthogonal matrix, R is upper triangular, and P is a permutation matrix.

The permuted QR factorization has the additional property that the diagonal entries of R are ordered by decreasing magnitude. In other words, `abs(diag(R))` will be ordered from largest to smallest.

For example, given the matrix $A = [1, 2; 3, 4]$,

```
[Q, R, P] = qr (A)
```

returns

```
Q =
```

```
-0.44721 -0.89443
-0.89443 0.44721
```

```
R =
```

```
-4.47214 -3.13050
0.00000 0.44721
```

```
P =
```

```
0 1
1 0
```

If the input matrix A is sparse then the sparse QR factorization is computed using CSPARSE. Because the matrix Q is, in general, a full matrix, it is recommended to request only one return value R . In that case, the computation avoids the construction of Q and returns R such that $R = \text{chol}(A' * A)$.

If an additional matrix B is supplied and two return values are requested, then `qr` returns C , where $C = Q' * B$. This allows the least squares approximation of $A \setminus B$ to be calculated as

```
[C, R] = qr (A, B)
x = R \ C
```

If the final argument is the string "vector" then P is a permutation vector (of the columns of A) instead of a permutation matrix. In this case, the defining relationship is

```
Q * R = A(:, P)
```

The default, however, is to return a permutation matrix and this may be explicitly specified by using a final argument of "matrix".

If the final argument is the scalar 0 an "economy" factorization is returned. When the original matrix A has size $M \times N$ and $M > N$ then the "economy" factorization will calculate just N rows in R and N columns in Q and omit the zeros in R . If M

$\leq N$ there is no difference between the economy and standard factorizations. When calculating an "economy" factorization the output P is always a vector rather than a matrix.

Background: The QR factorization has applications in the solution of least squares problems

$$\min_x \|Ax - b\|_2$$

for overdetermined systems of equations (i.e., A is a tall, thin matrix).

The permuted QR factorization $[Q, R, P] = qr(A)$ allows the construction of an orthogonal basis of $\text{span}(A)$.

See also: [chol], page 555, [hess], page 557, [lu], page 557, [qz], page 562, [schur], page 564, [svd], page 565, [qrupdate], page 561, [qrinsert], page 561, [qrdelete], page 562, [qrshift], page 562.

[Q1, R1] = qrupdate (Q, R, u, v)

Update a QR factorization given update vectors or matrices.

Given a QR factorization of a real or complex matrix $A = Q^*R$, Q unitary and R upper trapezoidal, return the QR factorization of $A + u^*v'$, where u and v are column vectors (rank-1 update) or matrices with equal number of columns (rank-k update). Notice that the latter case is done as a sequence of rank-1 updates; thus, for k large enough, it will be both faster and more accurate to recompute the factorization from scratch.

The QR factorization supplied may be either full (Q is square) or economized (R is square).

See also: [qr], page 559, [qrinsert], page 561, [qrdelete], page 562, [qrshift], page 562.

[Q1, R1] = qrinsert (Q, R, j, x, orient)

Update a QR factorization given a row or column to insert in the original factored matrix.

Given a QR factorization of a real or complex matrix $A = Q^*R$, Q unitary and R upper trapezoidal, return the QR factorization of $[A(:,1:j-1) \times A(:,j:n)]$, where u is a column vector to be inserted into A (if *orient* is "col"), or the QR factorization of $[A(1:j-1,:);x;A(:,j:n)]$, where x is a row vector to be inserted into A (if *orient* is "row").

The default value of *orient* is "col". If *orient* is "col", u may be a matrix and j an index vector resulting in the QR factorization of a matrix B such that $B(:,j)$ gives u and $B(:,j) = []$ gives A . Notice that the latter case is done as a sequence of k insertions; thus, for k large enough, it will be both faster and more accurate to recompute the factorization from scratch.

If *orient* is "col", the QR factorization supplied may be either full (Q is square) or economized (R is square).

If *orient* is "row", full factorization is needed.

See also: [qr], page 559, [qrupdate], page 561, [qrdelete], page 562, [qrshift], page 562.

[Q1, R1] = qrdelete (Q, R, j, orient)

Update a QR factorization given a row or column to delete from the original factored matrix.

Given a QR factorization of a real or complex matrix $A = Q^*R$, Q unitary and R upper trapezoidal, return the QR factorization of $[A(:,1:j-1), U, A(:,j:n)]$, where u is a column vector to be inserted into A (if *orient* is "col"), or the QR factorization of $[A(1:j-1,:); X; A(:,j:n)]$, where x is a row *orient* is "row"). The default value of *orient* is "col".

If *orient* is "col", j may be an index vector resulting in the QR factorization of a matrix B such that $A(:,j) = []$ gives B . Notice that the latter case is done as a sequence of k deletions; thus, for k large enough, it will be both faster and more accurate to recompute the factorization from scratch.

If *orient* is "col", the QR factorization supplied may be either full (Q is square) or economized (R is square).

If *orient* is "row", full factorization is needed.

See also: [qr], page 559, [qrupdate], page 561, [qrinsert], page 561, [qrshift], page 562.

[Q1, R1] = qrshift (Q, R, i, j)

Update a QR factorization given a range of columns to shift in the original factored matrix.

Given a QR factorization of a real or complex matrix $A = Q^*R$, Q unitary and R upper trapezoidal, return the QR factorization of $A(:,p)$, where p is the permutation $p = [1:i-1, \text{shift}(i:j, 1), j+1:n]$ if $i < j$

or

$p = [1:j-1, \text{shift}(j:i,-1), i+1:n]$ if $j < i$.

See also: [qr], page 559, [qrupdate], page 561, [qrinsert], page 561, [qrdelete], page 562.

lambda = qz (A, B)

[AA, BB, Q, Z, V, W, lambda] = qz (A, B)

[AA, BB, Z] = qz (A, B, opt)

[AA, BB, Z, lambda] = qz (A, B, opt)

Compute the QZ decomposition of a generalized eigenvalue problem.

The generalized eigenvalue problem is defined as

$$Ax = \lambda Bx$$

There are three calling forms of the function:

1. **lambda = qz (A, B)**

Compute the generalized eigenvalues λ .

2. **[AA, BB, Q, Z, V, W, lambda] = qz (A, B)**

Compute QZ decomposition, generalized eigenvectors, and generalized eigenvalues.

$$AV = BV\text{diag}(\lambda)$$

$$W^T A = \text{diag}(\lambda) W^T B$$

$$AA = Q^T AZ, BB = Q^T BZ$$

with Q and Z orthogonal (unitary for complex case).

3. **[AA, BB, Z {, lambda}] = qz (A, B, opt)**

As in form 2 above, but allows ordering of generalized eigenpairs for, e.g., solution of discrete time algebraic Riccati equations. Form 3 is not available for complex matrices, and does not compute the generalized eigenvectors V , W , nor the orthogonal matrix Q .

opt for ordering eigenvalues of the GEP pencil. The leading block of the revised pencil contains all eigenvalues that satisfy:

"N"	unordered (default)
"S"	small: leading block has all $ \lambda < 1$
"B"	big: leading block has all $ \lambda \geq 1$
"-"	negative real part: leading block has all eigenvalues in the open left half-plane
non-negative real part: leading block has all eigenvalues in the closed right half-plane	

Note: **qz** performs permutation balancing, but not scaling (see [balance], page 545), which may be lead to less accurate results than **eig**. The order of output arguments was selected for compatibility with MATLAB.

See also: [eig], page 547, [ordeig], page 565, [balance], page 545, [lu], page 557, [chol], page 555, [hess], page 557, [qr], page 559, [qzhess], page 563, [schur], page 564, [svd], page 565.

[aa, bb, q, z] = qzhess (A, B)

Compute the Hessenberg-triangular decomposition of the matrix pencil (A, B) , returning $aa = q * A * z$, $bb = q * B * z$, with q and z orthogonal.

For example:

```
[aa, bb, q, z] = qzhess ([1, 2; 3, 4], [5, 6; 7, 8])
⇒ aa =
   -3.02244  -4.41741
    0.92998   0.69749
⇒ bb =
   -8.60233  -9.99730
    0.00000  -0.23250
⇒ q =
   -0.58124  -0.81373
   -0.81373   0.58124
⇒ z =
Diagonal Matrix
   1   0
   0   1
```

The Hessenberg-triangular decomposition is the first step in Moler and Stewart's QZ decomposition algorithm.

Algorithm taken from Golub and Van Loan, *Matrix Computations*, 2nd edition.

See also: [lu], page 557, [chol], page 555, [hess], page 557, [qr], page 559, [qz], page 562, [schur], page 564, [svd], page 565.

```
S = schur (A)
S = schur (A, "real")
S = schur (A, "complex")
S = schur (A, opt)
[U, S] = schur (...)
```

Compute the Schur decomposition of A .

The Schur decomposition is defined as

$$S = U^T A U$$

where U is a unitary matrix ($U^T U$ is identity) and S is upper triangular. The eigenvalues of A (and S) are the diagonal elements of S . If the matrix A is real, then the real Schur decomposition is computed, in which the matrix U is orthogonal and S is block upper triangular with blocks of size at most 2×2 along the diagonal. The diagonal elements of S (or the eigenvalues of the 2×2 blocks, when appropriate) are the eigenvalues of A and S .

The default for real matrices is a real Schur decomposition. A complex decomposition may be forced by passing the flag "**complex**".

The eigenvalues are optionally ordered along the diagonal according to the value of opt . $opt = "a"$ indicates that all eigenvalues with negative real parts should be moved to the leading block of S (used in **are**), $opt = "d"$ indicates that all eigenvalues with magnitude less than one should be moved to the leading block of S (used in **dare**), and $opt = "u"$, the default, indicates that no ordering of eigenvalues should occur. The leading k columns of U always span the A -invariant subspace corresponding to the k leading eigenvalues of S .

The Schur decomposition is used to compute eigenvalues of a square matrix, and has applications in the solution of algebraic Riccati equations in control (see **are** and **dare**).

See also: [rsf2csf], page 564, [ordschur], page 565, [ordeig], page 565, [lu], page 557, [chol], page 555, [hess], page 557, [qr], page 559, [qz], page 562, [svd], page 565.

```
[U, T] = rsf2csf (UR, TR)
```

Convert a real, upper quasi-triangular Schur form TR to a complex, upper triangular Schur form T .

Note that the following relations hold:

$UR \cdot TR \cdot UR^T = UTU^\dagger$ and $U^\dagger U$ is the identity matrix I .

Note also that U and T are not unique.

See also: [schur], page 564.

[UR, SR] = ordschur (U, S, select)

Reorders the real Schur factorization (U, S) obtained with the **schur** function, so that selected eigenvalues appear in the upper left diagonal blocks of the quasi triangular Schur matrix.

The logical vector *select* specifies the selected eigenvalues as they appear along *S*'s diagonal.

For example, given the matrix $A = [1, 2; 3, 4]$, and its Schur decomposition

[U, S] = schur (A)

which returns

$U =$

```
-0.82456 -0.56577
0.56577 -0.82456
```

$S =$

```
-0.37228 -1.00000
0.00000 5.37228
```

It is possible to reorder the decomposition so that the positive eigenvalue is in the upper left corner, by doing:

[U, S] = ordschur (U, S, [0,1])

See also: [\[schur\]](#), page 564, [\[ordeig\]](#), page 565.

lambda = ordeig (A)
lambda = ordeig (A, B)

Return the eigenvalues of quasi-triangular matrices in their order of appearance in the matrix *A*.

The quasi-triangular matrix *A* is usually the result of a Schur factorization. If called with a second input *B* then the generalized eigenvalues of the pair *A, B* are returned in the order of appearance of the matrix $A - \lambda B$. The pair *A, B* is usually the result of a QZ decomposition.

See also: [\[ordschur\]](#), page 565, [\[eig\]](#), page 547, [\[schur\]](#), page 564, [\[qz\]](#), page 562.

angle = subspace (A, B)

Determine the largest principal angle between two subspaces spanned by the columns of matrices *A* and *B*.

s = svd (A)
[U, S, V] = svd (A)
[U, S, V] = svd (A, "econ")
[U, S, V] = svd (A, 0)

Compute the singular value decomposition of *A*.

The singular value decomposition is defined by the relation

$$A = USV^\dagger$$

The function `svd` normally returns only the vector of singular values. When called with three return values, it computes U , S , and V . For example,

```
svd (hilb (3))
```

returns

```
ans =
```

```
1.4083189
0.1223271
0.0026873
```

and

```
[u, s, v] = svd (hilb (3))
```

returns

```
u =
```

```
-0.82704  0.54745  0.12766
-0.45986 -0.52829 -0.71375
-0.32330 -0.64901  0.68867
```

```
s =
```

```
1.40832  0.00000  0.00000
0.00000  0.12233  0.00000
0.00000  0.00000  0.00269
```

```
v =
```

```
-0.82704  0.54745  0.12766
-0.45986 -0.52829 -0.71375
-0.32330 -0.64901  0.68867
```

When given a second argument that is not 0, `svd` returns an economy-sized decomposition, eliminating the unnecessary rows or columns of U or V .

If the second argument is exactly 0, then the choice of decomposition is based on the matrix A . If A has more rows than columns then an economy-sized decomposition is returned, otherwise a regular decomposition is calculated.

Algorithm Notes: When calculating the full decomposition (left and right singular matrices in addition to singular values) there is a choice of two routines in LAPACK. The default routine used by Octave is `gesdd` which is 5X faster than the alternative `gesvd`, but may use more memory and may be less accurate for some matrices. See the documentation for `svd_driver` for more information.

See also: [\[svd_driver\]](#), page 567, [\[svds\]](#), page 646, [\[eig\]](#), page 547, [\[lu\]](#), page 557, [\[chol\]](#), page 555, [\[hess\]](#), page 557, [\[qr\]](#), page 559, [\[qz\]](#), page 562.

```
val = svd_driver ()
old_val = svd_driver (new_val)
svd_driver (new_val, "local")
```

Query or set the underlying LAPACK driver used by `svd`.

Currently recognized values are "`gesdd`" and "`gesvd`". The default is "`gesvd`".

When called from inside a function with the "`local`" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

Algorithm Notes: The LAPACK library provides two routines for calculating the full singular value decomposition (left and right singular matrices as well as singular values). When calculating just the singular values the following discussion is not relevant.

The newer `gesdd` routine is based on a Divide-and-Conquer algorithm that is 5X faster than the alternative `gesvd`, which is based on QR factorization. However, the new algorithm can use significantly more memory. For an MxN input matrix the memory usage is of order $O(\min(M,N)^2)$, whereas the alternative is of order $O(\max(M,N))$.

Beyond speed and memory issues, there have been instances where some input matrices were not accurately decomposed by `gesdd`. See currently active bug <https://savannah.gnu.org/bugs/?55564>. Until these accuracy issues are resolved in a new version of the LAPACK library, the default driver in Octave has been set to "`gesvd`".

See also: [\[svd\]](#), page 565.

`[housv, beta, zer] = housh (x, j, z)`

Compute Householder reflection vector `housv` to reflect `x` to be the `j`-th column of identity, i.e.,

```
(I - beta*housv*housv')x = norm (x)*e(j) if x(j) < 0,
(I - beta*housv*housv')x = -norm (x)*e(j) if x(j) >= 0
```

Inputs

`x` vector

`j` index into vector

`z` threshold for zero (usually should be the number 0)

Outputs (see Golub and Van Loan):

`beta` If `beta = 0`, then no reflection need be applied (`zer` set to 0)

`housv` householder vector

`[u, h, nu] = krylov (A, V, k, eps1, pflg)`

Construct an orthogonal basis `u` of a block Krylov subspace.

The block Krylov subspace has the following form:

```
[v a*v a^2*v ... a^(k+1)*v]
```

The construction is made with Householder reflections to guard against loss of orthogonality.

If V is a vector, then h contains the Hessenberg matrix such that $a*u == u*h+rk*ek'$, in which $rk = a*u(:,k)-u*h(:,k)$, and ek' is the vector $[0, 0, \dots, 1]$ of length k . Otherwise, h is meaningless.

If V is a vector and k is greater than $\text{length}(A) - 1$, then h contains the Hessenberg matrix such that $a*u == u*h$.

The value of nu is the dimension of the span of the Krylov subspace (based on eps1).

If b is a vector and k is greater than $m-1$, then h contains the Hessenberg decomposition of A .

The optional parameter eps1 is the threshold for zero. The default value is $1e-12$.

If the optional parameter $pflg$ is nonzero, row pivoting is used to improve numerical behavior. The default value is 0.

Reference: A. Hodel, P. Misra, *Partial Pivoting in the Computation of Krylov Subspaces of Large Sparse Systems*, Proceedings of the 42nd IEEE Conference on Decision and Control, December 2003.

18.4 Functions of a Matrix

`expm (A)`

Return the exponential of a matrix.

The matrix exponential is defined as the infinite Taylor series

$$\exp(A) = I + A + \frac{A^2}{2!} + \frac{A^3}{3!} + \dots$$

However, the Taylor series is *not* the way to compute the matrix exponential; see Moler and Van Loan, *Nineteen Dubious Ways to Compute the Exponential of a Matrix*, SIAM Review, 1978. This routine uses Ward's diagonal Padé approximation method with three step preconditioning (SIAM Journal on Numerical Analysis, 1977). Diagonal Padé approximations are rational polynomials of matrices $D_q(A)^{-1}N_q(A)$ whose Taylor series matches the first $2q + 1$ terms of the Taylor series above; direct evaluation of the Taylor series (with the same preconditioning steps) may be desirable in lieu of the Padé approximation when $D_q(A)$ is ill-conditioned.

See also: [\[logm\]](#), page 568, [\[sqrtm\]](#), page 569.

```
s = logm (A)
s = logm (A, opt_iters)
[s, iters] = logm (...)
```

Compute the matrix logarithm of the square matrix A .

The implementation utilizes a Padé approximant and the identity

```
logm (A) = 2^k * logm (A^(1 / 2^k))
```

The optional input opt_iters is the maximum number of square roots to compute and defaults to 100.

The optional output $iters$ is the number of square roots actually computed.

See also: [\[expm\]](#), page 568, [\[sqrtm\]](#), page 569.

```
s = sqrtm (A)
[s, error_estimate] = sqrtm (A)
```

Compute the matrix square root of the square matrix A .

Ref: N.J. Higham. *A New sqrtm for MATLAB*. Numerical Analysis Report No. 336, Manchester Centre for Computational Mathematics, Manchester, England, January 1999.

See also: [expm], page 568, [logm], page 568.

```
kron (A, B)
kron (A1, A2, ...)
```

Form the Kronecker product of two or more matrices.

This is defined block by block as

```
x = [ a(i,j)*b ]
```

For example:

```
kron (1:4, ones (3, 1))
      ⇒  1   2   3   4
          1   2   3   4
          1   2   3   4
```

If there are more than two input arguments A_1, A_2, \dots, A_n the Kronecker product is computed as

```
kron (kron (A1, A2), ..., An)
```

Since the Kronecker product is associative, this is well-defined.

```
blkmm (A, B)
```

Compute products of matrix blocks.

The blocks are given as 2-dimensional subarrays of the arrays A, B . The size of A must have the form $[m, k, \dots]$ and size of B must be $[k, n, \dots]$. The result is then of size $[m, n, \dots]$ and is computed as follows:

```
for i = 1:prod (size (A)(3:end))
    C(:,:,i) = A(:,:,i) * B(:,:,i)
endfor
```

```
X = sylvester (A, B, C)
```

Solve the Sylvester equation.

The Sylvester equation is defined as:

$$AX + XB = C$$

The solution is computed using standard LAPACK subroutines.

For example:

```
sylvester ([1, 2; 3, 4], [5, 6; 7, 8], [9, 10; 11, 12])
      ⇒ [ 0.50000, 0.66667; 0.66667, 0.50000 ]
```

18.5 Specialized Solvers

```
x = bicg (A, b)
x = bicg (A, b, tol)
x = bicg (A, b, tol, maxit)
x = bicg (A, b, tol, maxit, M)
x = bicg (A, b, tol, maxit, M1, M2)
x = bicg (A, b, tol, maxit, M, [], x0)
x = bicg (A, b, tol, maxit, M1, M2, x0)
x = bicg (A, b, tol, maxit, M, [], x0, ...)
x = bicg (A, b, tol, maxit, M1, M2, x0, ...)
[x, flag, relres, iter, resvec] = bicg (A, b, ...)
```

Solve the linear system of equations $A * x = b$ by means of the Bi-Conjugate Gradient iterative method.

The input arguments are:

- A is the matrix of the linear system and it must be square. A can be passed as a matrix, function handle, or inline function `Afun` such that `Afun (x, "notransp") = A * x` and `Afun (x, "transp") = A' * x`. Additional parameters to `Afun` may be passed after `x0`.
- b is the right-hand side vector. It must be a column vector with the same number of rows as A .
- tol is the required relative tolerance for the residual error, $b - A * x$. The iteration stops if `norm (b - A * x) ≤ tol * norm (b)`. If tol is omitted or empty, then a tolerance of 1e-6 is used.
- $maxit$ is the maximum allowed number of iterations; if $maxit$ is omitted or empty then a value of 20 is used.
- $M1, M2$ are the preconditioners. The preconditioner M is given as $M = M1 * M2$. Both $M1$ and $M2$ can be passed as a matrix or as a function handle or inline function `g` such that `g (x, "notransp") = M1 \ x` or `g (x, "notransp") = M2 \ x` and `g (x, "transp") = M1' \ x` or `g (x, "transp") = M2' \ x`. If $M1$ is omitted or empty, then preconditioning is not applied. The preconditioned system is theoretically equivalent to applying the `bicg` method to the linear system `inv (M1) * A * inv (M2) * y = inv (M1) * b` and `inv (M2') * A' * inv (M1') * z = inv (M2') * b` and then setting `x = inv (M2) * y`.
- $x0$ is the initial guess. If $x0$ is omitted or empty then the function sets $x0$ to a zero vector by default.

Any arguments which follow `x0` are treated as parameters, and passed in an appropriate manner to any of the functions (`Afun` or `Mfun`) or that have been given to `bicg`.

The output parameters are:

- x is the computed approximation to the solution of $A * x = b$. If the algorithm did not converge, then x is the iteration which has the minimum residual.
- $flag$ indicates the exit status:
 - 0: The algorithm converged to within the prescribed tolerance.

- 1: The algorithm did not converge and it reached the maximum number of iterations.
- 2: The preconditioner matrix is singular.
- 3: The algorithm stagnated, i.e., the absolute value of the difference between the current iteration x and the previous is less than $\text{eps} * \text{norm}(x, 2)$.
- 4: The algorithm could not continue because intermediate values became too small or too large for reliable computation.
- relres is the ratio of the final residual to its initial value, measured in the Euclidean norm.
- iter is the iteration which x is computed.
- resvec is a vector containing the residual at each iteration. The total number of iterations performed is given by $\text{length}(\text{resvec}) - 1$.

Consider a trivial problem with a tridiagonal matrix

```
n = 20;
A = toeplitz (sparse ([1, 1], [1, 2], [2, 1] * n ^ 2, 1, n)) + ...
    toeplitz (sparse (1, 2, -1, 1, n) * n / 2, ...
               sparse (1, 2, 1, 1, n) * n / 2);
b = A * ones (n, 1);
restart = 5;
[M1, M2] = ilu (A); # in this tridiag case, it corresponds to lu (A)
M = M1 * M2;
Afun = @(x, string) strcmp (string, "notransp") * (A * x) + ...
        strcmp (string, "transp") * (A' * x);
Mfun = @(x, string) strcmp (string, "notransp") * (M \ x) + ...
        strcmp (string, "transp") * (M' \ x);
M1fun = @(x, string) strcmp (string, "notransp") * (M1 \ x) + ...
        strcmp (string, "transp") * (M1' \ x);
M2fun = @(x, string) strcmp (string, "notransp") * (M2 \ x) + ...
        strcmp (string, "transp") * (M2' \ x);
```

EXAMPLE 1: simplest usage of `bicg`

```
x = bicg (A, b)
```

EXAMPLE 2: `bicg` with a function that computes $A*x$ and $A'*x$

```
x = bicg (Afun, b, [], n)
```

EXAMPLE 3: `bicg` with a preconditioner matrix M

```
x = bicg (A, b, 1e-6, n, M)
```

EXAMPLE 4: `bicg` with a function as preconditioner

```
x = bicg (Afun, b, 1e-6, n, Mfun)
```

EXAMPLE 5: `bicg` with preconditioner matrices $M1$ and $M2$

```
x = bicg (A, b, 1e-6, n, M1, M2)
```

EXAMPLE 6: `bicg` with functions as preconditioners

```
x = bicg (Afun, b, 1e-6, n, M1fun, M2fun)
```

EXAMPLE 7: `bicg` with as input a function requiring an argument

```
function y = Ap (A, x, string, z)
## compute A^z * x or (A^z)' * x
y = x;
if (strcmp (string, "notransp"))
    for i = 1:z
        y = A * y;
    endfor
elseif (strcmp (string, "transp"))
    for i = 1:z
        y = A' * y;
    endfor
endif
endfunction

Apfun = @(x, string, p) Ap (A, x, string, p);
x = bicg (Apfun, b, [], [], [], [], [], 2);
```

References:

1. Y. Saad, *Iterative Methods for Sparse Linear Systems*, Second edition, 2003, SIAM.

See also: [\[bicgstab\]](#), page 572, [\[cgs\]](#), page 574, [\[gmres\]](#), page 576, [\[pcg\]](#), page 647, [\[qmr\]](#), page 579, [\[tfqmr\]](#), page 580.

```
x = bicgstab (A, b, tol, maxit, M1, M2, x0, ...)
x = bicgstab (A, b, tol, maxit, M, [], x0, ...)
[x, flag, relres, iter, resvec] = bicgstab (A, b, ...)
```

Solve $A x = b$ using the stabilized Bi-conjugate gradient iterative method.

The input parameters are:

- A is the matrix of the linear system and it must be square. A can be passed as a matrix, function handle, or inline function `Afun` such that `Afun(x) = A * x`. Additional parameters to `Afun` are passed after `x0`.
- b is the right hand side vector. It must be a column vector with the same number of rows as A .
- tol is the required relative tolerance for the residual error, $b - A * x$. The iteration stops if $\text{norm}(b - A * x) \leq tol * \text{norm}(b)$. If tol is omitted or empty, then a tolerance of 1e-6 is used.
- $maxit$ the maximum number of outer iterations, if not given or set to [] the default value `min(20, numel(b))` is used.
- $M1, M2$ are the preconditioners. The preconditioner M is given as $M = M1 * M2$. Both $M1$ and $M2$ can be passed as a matrix or as a function handle or inline function `g` such that `g(x) = M1 \ x` or `g(x) = M2 \ x`. The technique used is the right preconditioning, i.e., it is solved $A * \text{inv}(M) * y = b$ and then $x = \text{inv}(M) * y$.
- $x0$ the initial guess, if not given or set to [] the default value `zeros(size(b))` is used.

The arguments which follow $x0$ are treated as parameters, and passed in a proper way to any of the functions (A or M) which are passed to `bicgstab`.

The output parameters are:

- x is the approximation computed. If the method doesn't converge then it is the iterated with the minimum residual.
- $flag$ indicates the exit status:
 - 0: iteration converged to the within the chosen tolerance
 - 1: the maximum number of iterations was reached before convergence
 - 2: the preconditioner matrix is singular
 - 3: the algorithm reached stagnation
 - 4: the algorithm can't continue due to a division by zero
- $relres$ is the relative residual obtained with as $(A*x - b) / \text{norm}(b)$.
- $iter$ is the (possibly half) iteration which x is computed. If it is an half iteration then it is $iter + 0.5$
- $resvec$ is a vector containing the residual of each half and total iteration (There are also the half iterations since x is computed in two steps at each iteration). Doing $(\text{length}(resvec) - 1) / 2$ is possible to see the total number of (total) iterations performed.

Let us consider a trivial problem with a tridiagonal matrix

```
n = 20;
A = toeplitz (sparse ([1, 1], [1, 2], [2, 1] * n ^ 2, 1, n)) + ...
     toeplitz (sparse (1, 2, -1, 1, n) * n / 2, ...
               sparse (1, 2, 1, 1, n) * n / 2);
b = A * ones (n, 1);
restart = 5;
[M1, M2] = ilu (A); # in this tridiag case, it corresponds to lu (A)
M = M1 * M2;
Afun = @(x) A * x;
Mfun = @(x) M \ x;
M1fun = @(x) M1 \ x;
M2fun = @(x) M2 \ x;
```

EXAMPLE 1: simplest usage of `bicgstab`

```
x = bicgstab (A, b, [], n)
```

EXAMPLE 2: `bicgstab` with a function which computes $A * x$

```
x = bicgstab (Afun, b, [], n)
```

EXAMPLE 3: `bicgstab` with a preconditioner matrix M

```
x = bicgstab (A, b, [], 1e-06, n, M)
```

EXAMPLE 4: `bicgstab` with a function as preconditioner

```
x = bicgstab (Afun, b, 1e-6, n, Mfun)
```

EXAMPLE 5: `bicgstab` with preconditioner matrices $M1$ and $M2$

```
x = bicgstab (A, b, [], 1e-6, n, M1, M2)
```

EXAMPLE 6: `bicgstab` with functions as preconditioners

```
x = bicgstab (Afun, b, 1e-6, n, M1fun, M2fun)
```

EXAMPLE 7: `bicgstab` with as input a function requiring an argument

```
function y = Ap (A, x, z) # compute A^z * x
    y = x;
    for i = 1:z
        y = A * y;
    endfor
endfunction
```

```
Apfun = @(x, string, p) Ap (A, x, string, p);
x = bicgstab (Apfun, b, [], [], [], [], 2);
```

EXAMPLE 8: explicit example to show that `bicgstab` uses a right preconditioner

```
[M1, M2] = ilu (A + 0.1 * eye (n)); # factorization of A perturbed
M = M1 * M2;
```

```
## reference solution computed by bicgstab after one iteration
[x_ref, f1] = bicgstab (A, b, [], 1, M)
```

```
## right preconditioning
[y, f1] = bicgstab (A / M, b, [], 1)
x = M \ y # compare x and x_ref
```

References:

1. Y. Saad, *Iterative Methods for Sparse Linear Systems*, Second edition, 2003, SIAM

See also: `[bicg]`, page 570, `[cgs]`, page 574, `[gmres]`, page 576, `[pcg]`, page 647, `[qmr]`, page 579, `[tfqmr]`, page 580.

```
x = cgs (A, b, tol, maxit, M1, M2, x0, ...)
x = cgs (A, b, tol, maxit, M, [], x0, ...)
[x, flag, relres, iter, resvec] = cgs (A, b, ...)
```

Solve $A x = b$, where A is a square matrix, using the Conjugate Gradients Squared method.

The input arguments are:

- A is the matrix of the linear system and it must be square. A can be passed as a matrix, function handle, or inline function `Afun` such that `Afun(x) = A * x`. Additional parameters to `Afun` are passed after `x0`.
- b is the right hand side vector. It must be a column vector with same number of rows of A .
- tol is the relative tolerance, if not given or set to [] the default value 1e-6 is used.
- $maxit$ the maximum number of outer iterations, if not given or set to [] the default value `min (20, numel (b))` is used.
- $M1, M2$ are the preconditioners. The preconditioner matrix is given as $M = M1 * M2$. Both $M1$ and $M2$ can be passed as a matrix or as a function handle or

inline function g such that $g(x) = M1 \setminus x$ or $g(x) = M2 \setminus x$. If $M1$ is empty or not passed then no preconditioners are applied. The technique used is the right preconditioning, i.e., it is solved $A \cdot \text{inv}(M) \cdot y = b$ and then $x = \text{inv}(M) \cdot y$.

- $x0$ the initial guess, if not given or set to `[]` the default value `zeros (size (b))` is used.

The arguments which follow $x0$ are treated as parameters, and passed in a proper way to any of the functions (A or P) which are passed to `cgs`.

The output parameters are:

- x is the approximation computed. If the method doesn't converge then it is the iterated with the minimum residual.
- $flag$ indicates the exit status:
 - 0: iteration converged to the within the chosen tolerance
 - 1: the maximum number of iterations was reached before convergence
 - 2: the preconditioner matrix is singular
 - 3: the algorithm reached stagnation
 - 4: the algorithm can't continue due to a division by zero
- $relres$ is the relative residual obtained with as $(A \cdot x - b) / \text{norm}(b)$.
- $iter$ is the iteration which x is computed.
- $resvec$ is a vector containing the residual at each iteration. Doing `length(resvec) - 1` is possible to see the total number of iterations performed.

Let us consider a trivial problem with a tridiagonal matrix

```
n = 20;
A = toeplitz (sparse ([1, 1], [1, 2], [2, 1] * n ^ 2, 1, n)) + ...
    toeplitz (sparse (1, 2, -1, 1, n) * n / 2, ...
    sparse (1, 2, 1, 1, n) * n / 2);
b = A * ones (n, 1);
restart = 5;
[M1, M2] = ilu (A); # in this tridiag case it corresponds to chol (A)'■
M = M1 * M2;
Afun = @(x) A * x;
Mfun = @(x) M \ x;
M1fun = @(x) M1 \ x;
M2fun = @(x) M2 \ x;
```

EXAMPLE 1: simplest usage of `cgs`

```
x = cgs (A, b, [], n)
```

EXAMPLE 2: `cgs` with a function which computes $A \cdot x$

```
x = cgs (Afun, b, [], n)
```

EXAMPLE 3: `cgs` with a preconditioner matrix M

```
x = cgs (A, b, [], 1e-06, n, M)
```

EXAMPLE 4: `cgs` with a function as preconditioner

```
x = cgs (Afun, b, 1e-6, n, Mfun)
```

EXAMPLE 5: `cgs` with preconditioner matrices $M1$ and $M2$

```
x = cgs (A, b, [], 1e-6, n, M1, M2)
```

EXAMPLE 6: `cgs` with functions as preconditioners

```
x = cgs (Afun, b, 1e-6, n, M1fun, M2fun)
```

EXAMPLE 7: `cgs` with as input a function requiring an argument

```
function y = Ap (A, x, z) # compute A^z * x
    y = x;
    for i = 1:z
        y = A * y;
    endfor
endfunction
Apfun = @(x, string, p) Ap (A, x, string, p);
x = cgs (Apfun, b, [], [], [], [], [], 2);
```

EXAMPLE 8: explicit example to show that `cgs` uses a right preconditioner

```
[M1, M2] = ilu (A + 0.3 * eye (n)); # factorization of A perturbed
M = M1 * M2;

## reference solution computed by cgs after one iteration
[x_ref, f1] = cgs (A, b, [], 1, M)

## right preconditioning
[y, f1] = cgs (A / M, b, [], 1)
x = M \ y # compare x and x_ref
```

References:

1. Y. Saad, *Iterative Methods for Sparse Linear Systems*, Second edition, 2003, SIAM

See also: `[pcg]`, page 647, `[bicgstab]`, page 572, `[bicg]`, page 570, `[gmres]`, page 576, `[qmr]`, page 579, `[tfqmr]`, page 580.

```
x = gmres (A, b, restart, tol, maxit, M1, M2, x0, ...)
x = gmres (A, b, restart, tol, maxit, M, [], x0, ...)
[x, flag, relres, iter, resvec] = gmres (A, b, ...)
```

Solve $A x = b$ using the Preconditioned GMRES iterative method with restart, a.k.a. PGMRES(restart).

The input arguments are:

- A is the matrix of the linear system and it must be square. A can be passed as a matrix, function handle, or inline function `Afun` such that `Afun(x) = A * x`. Additional parameters to `Afun` are passed after `x0`.
- b is the right hand side vector. It must be a column vector with the same numbers of rows as A .
- `restart` is the number of iterations before that the method restarts. If it is `[]` or $N = \text{numel}(b)$, then the restart is not applied.

- tol is the required relative tolerance for the preconditioned residual error, $\text{inv}(M) * (b - a * x)$. The iteration stops if $\text{norm}(\text{inv}(M) * (b - a * x)) \leq tol * \text{norm}(\text{inv}(M) * b)$. If tol is omitted or empty, then a tolerance of 1e-6 is used.
- $maxit$ is the maximum number of outer iterations, if not given or set to $[]$, then the default value $\min(10, N / \text{restart})$ is used. Note that, if restart is empty, then $maxit$ is the maximum number of iterations. If restart and $maxit$ are not empty, then the maximum number of iterations is $\text{restart} * maxit$. If both restart and $maxit$ are empty, then the maximum number of iterations is set to $\min(10, N)$.
- $M1, M2$ are the preconditioners. The preconditioner M is given as $M = M1 * M2$. Both $M1$ and $M2$ can be passed as a matrix, function handle, or inline function g such that $g(x) = M1 \setminus x$ or $g(x) = M2 \setminus x$. If $M1$ is $[]$ or not given, then the preconditioner is not applied. The technique used is the left-preconditioning, i.e., it is solved $\text{inv}(M) * A * x = \text{inv}(M) * b$ instead of $A * x = b$.
- $x0$ is the initial guess, if not given or set to $[]$, then the default value $\text{zeros}(\text{size}(b))$ is used.

The arguments which follow $x0$ are treated as parameters, and passed in a proper way to any of the functions (A or M or $M1$ or $M2$) which are passed to `gmres`.

The outputs are:

- x the computed approximation. If the method does not converge, then it is the iterated with minimum residual.
- $flag$ indicates the exit status:
 - 0 : iteration converged to within the specified tolerance
 - 1 : maximum number of iterations exceeded
 - 2 : the preconditioner matrix is singular
 - 3 : algorithm reached stagnation (the relative difference between two consecutive iterations is less than eps)
- $relres$ is the value of the relative preconditioned residual of the approximation x .
- $iter$ is a vector containing the number of outer iterations and inner iterations performed to compute x . That is:
 - $iter(1)$: number of outer iterations, i.e., how many times the method restarted. (if restart is empty or N , then it is 1, if not $1 \leq iter(1) \leq maxit$).
 - $iter(2)$: the number of iterations performed before the restart, i.e., the method restarts when $iter(2) = \text{restart}$. If restart is empty or N , then $1 \leq iter(2) \leq maxit$.

To be more clear, the approximation x is computed at the iteration $(\text{iter}(1) - 1) * \text{restart} + \text{iter}(2)$. Since the output x corresponds to the minimal preconditioned residual solution, the total number of iterations that the method performed is given by $\text{length(resvec)} - 1$.

- $resvec$ is a vector containing the preconditioned relative residual at each iteration, including the 0-th iteration $\text{norm}(A * x0 - b)$.

Let us consider a trivial problem with a tridiagonal matrix

```

n = 20;
A = toeplitz (sparse ([1, 1], [1, 2], [2, 1] * n ^ 2, 1, n)) + ...
    toeplitz (sparse (1, 2, -1, 1, n) * n / 2, ...
    sparse (1, 2, 1, 1, n) * n / 2);
b = A * ones (n, 1);
restart = 5;
[M1, M2] = ilu (A); # in this tridiag case, it corresponds to lu (A)
M = M1 * M2;
Afun = @(x) A * x;
Mfun = @(x) M \ x;
M1fun = @(x) M1 \ x;
M2fun = @(x) M2 \ x;

```

EXAMPLE 1: simplest usage of `gmres`

```
x = gmres (A, b, [], [], n)
```

EXAMPLE 2: `gmres` with a function which computes $A * x$

```
x = gmres (Afun, b, [], [], n)
```

EXAMPLE 3: usage of `gmres` with the restart

```
x = gmres (A, b, restart);
```

EXAMPLE 4: `gmres` with a preconditioner matrix M with and without restart

```
x = gmres (A, b, [], 1e-06, n, M)
x = gmres (A, b, restart, 1e-06, n, M)
```

EXAMPLE 5: `gmres` with a function as preconditioner

```
x = gmres (Afun, b, [], 1e-6, n, Mfun)
```

EXAMPLE 6: `gmres` with preconditioner matrices $M1$ and $M2$

```
x = gmres (A, b, [], 1e-6, n, M1, M2)
```

EXAMPLE 7: `gmres` with functions as preconditioners

```
x = gmres (Afun, b, 1e-6, n, M1fun, M2fun)
```

EXAMPLE 8: `gmres` with as input a function requiring an argument

```

function y = Ap (A, x, p) # compute  $A^p * x$ 
    y = x;
    for i = 1:p
        y = A * y;
    endfor
endfunction
Apfun = @(x, p) Ap (A, x, p);
x = gmres (Apfun, b, [], [], [], [], [], 2);

```

EXAMPLE 9: explicit example to show that `gmres` uses a left preconditioner

```
[M1, M2] = ilu (A + 0.1 * eye (n)); # factorization of A perturbed
M = M1 * M2;

## reference solution computed by gmres after two iterations
[x_ref, f1] = gmres (A, b, [], [], 1, M)

## left preconditioning
[x, f1] = gmres (M \ A, M \ b, [], [], 1)
x # compare x and x_ref
```

References:

1. Y. Saad, *Iterative Methods for Sparse Linear Systems*, Second edition, 2003, SIAM

See also: [bicg], page 570, [bicgstab], page 572, [cgs], page 574, [pcg], page 647, [pcr], page 650, [qmr], page 579, [tfqmr], page 580.

```
x = qmr (A, b, rtol, maxit, M1, M2, x0)
x = qmr (A, b, rtol, maxit, P)
[x, flag, relres, iter, resvec] = qmr (A, b, ...)
Solve A x = b using the Quasi-Minimal Residual iterative method (without lookahead).
– rtol is the relative tolerance, if not given or set to [] the default value 1e-6 is used.
– maxit the maximum number of outer iterations, if not given or set to [] the default value min (20, numel (b)) is used.
– x0 the initial guess, if not given or set to [] the default value zeros (size (b)) is used.
```

A can be passed as a matrix or as a function handle or inline function *f* such that *f*(*x*, "notransp") = *A***x* and *f*(*x*, "transp") = *A'***x*.

The preconditioner *P* is given as *P* = *M1* * *M2*. Both *M1* and *M2* can be passed as a matrix or as a function handle or inline function *g* such that *g*(*x*, "notransp") = *M1* \ *x* or *g*(*x*, "notransp") = *M2* \ *x* and *g*(*x*, "transp") = *M1'* \ *x* or *g*(*x*, "transp") = *M2'* \ *x*.

If called with more than one output parameter

- *flag* indicates the exit status:
 - 0: iteration converged to the within the chosen tolerance
 - 1: the maximum number of iterations was reached before convergence
 - 3: the algorithm reached stagnation

(the value 2 is unused but skipped for compatibility).

- *relres* is the final value of the relative residual.
- *iter* is the number of iterations performed.
- *resvec* is a vector containing the residual norms at each iteration.

References:

1. R. Freund and N. Nachtigal, *QMR: a quasi-minimal residual method for non-Hermitian linear systems*, Numerische Mathematik, 1991, 60, pp. 315-339.
2. R. Barrett, M. Berry, T. Chan, J. Demmel, J. Donato, J. Dongarra, V. Eijkhour, R. Pozo, C. Romine, and H. van der Vorst, *Templates for the solution of linear systems: Building blocks for iterative methods*, SIAM, 2nd ed., 1994.

See also: [bicg], page 570, [bicgstab], page 572, [cgs], page 574, [gmres], page 576, [pcg], page 647.

```
x = tfqmr (A, b, tol, maxit, M1, M2, x0, ...)
x = tfqmr (A, b, tol, maxit, M, [], x0, ...)
[x, flag, relres, iter, resvec] = tfqmr (A, b, ...)
```

Solve $A x = b$ using the Transpose-Tree qmr method, based on the cgs.

The input parameters are:

- A is the matrix of the linear system and it must be square. A can be passed as a matrix, function handle, or inline function `Afun` such that `Afun(x) = A * x`. Additional parameters to `Afun` are passed after `x0`.
- b is the right hand side vector. It must be a column vector with the same number of rows as A .
- `tol` is the relative tolerance, if not given or set to [] the default value 1e-6 is used.
- `maxit` the maximum number of outer iterations, if not given or set to [] the default value `min (20, numel (b))` is used. To be compatible, since the method as different behaviors in the iteration number is odd or even, is considered as iteration in `tfqmr` the entire odd-even cycle. That is, to make an entire iteration, the algorithm performs two sub-iterations: the odd one and the even one.
- $M1, M2$ are the preconditioners. The preconditioner M is given as $M = M1 * M2$. Both $M1$ and $M2$ can be passed as a matrix or as a function handle or inline function `g` such that `g(x) = M1 \ x` or `g(x) = M2 \ x`. The technique used is the right-preconditioning, i.e., it is solved `A*inv(M)*y = b` and then `x = inv(M)*y`, instead of `A x = b`.
- `x0` the initial guess, if not given or set to [] the default value `zeros (size (b))` is used.

The arguments which follow `x0` are treated as parameters, and passed in a proper way to any of the functions (A or M) which are passed to `tfqmr`.

The output parameters are:

- x is the approximation computed. If the method doesn't converge then it is the iterated with the minimum residual.
- `flag` indicates the exit status:
 - 0: iteration converged to the within the chosen tolerance
 - 1: the maximum number of iterations was reached before convergence
 - 2: the preconditioner matrix is singular
 - 3: the algorithm reached stagnation

- 4: the algorithm can't continue due to a division by zero
- *relres* is the relative residual obtained as $(A*x - b) / \text{norm}(b)$.
- *iter* is the iteration which *x* is computed.
- *resvec* is a vector containing the residual at each iteration (including $\text{norm}(b - A*x_0)$). Doing *length(resvec) - 1* is possible to see the total number of iterations performed.

Let us consider a trivial problem with a tridiagonal matrix

```
n = 20;
A = toeplitz (sparse ([1, 1], [1, 2], [2, 1] * n ^ 2, 1, n)) + ...
    toeplitz (sparse (1, 2, -1, 1, n) * n / 2, ...
    sparse (1, 2, 1, 1, n) * n / 2);
b = A * ones (n, 1);
restart = 5;
[M1, M2] = ilu (A); # in this tridiag case it corresponds to chol (A)';
M = M1 * M2;
Afun = @(x) A * x;
Mfun = @(x) M \ x;
M1fun = @(x) M1 \ x;
M2fun = @(x) M2 \ x;
```

EXAMPLE 1: simplest usage of *tfqmr*

```
x = tfqmr (A, b, [], n)
```

EXAMPLE 2: *tfqmr* with a function which computes $A * x$

```
x = tfqmr (Afun, b, [], n)
```

EXAMPLE 3: *tfqmr* with a preconditioner matrix *M*

```
x = tfqmr (A, b, [], 1e-06, n, M)
```

EXAMPLE 4: *tfqmr* with a function as preconditioner

```
x = tfqmr (Afun, b, 1e-6, n, Mfun)
```

EXAMPLE 5: *tfqmr* with preconditioner matrices *M1* and *M2*

```
x = tfqmr (A, b, [], 1e-6, n, M1, M2)
```

EXAMPLE 6: *tfmqr* with functions as preconditioners

```
x = tfmqr (Afun, b, 1e-6, n, M1fun, M2fun)
```

EXAMPLE 7: *tfqmr* with as input a function requiring an argument

```
function y = Ap (A, x, z) # compute  $A^z * x$ 
    y = x;
    for i = 1:z
        y = A * y;
    endfor
endfunction
Apfun = @(x, string, p) Ap (A, x, string, p);
x = tfqmr (Apfun, b, [], [], [], [], [], 2);
```

EXAMPLE 8: explicit example to show that `tfqmr` uses a right preconditioner

```
[M1, M2] = ilu (A + 0.3 * eye (n)); # factorization of A perturbed
M = M1 * M2;

## reference solution computed by tfqmr after one iteration
[x_ref, f1] = tfqmr (A, b, [], 1, M)

## right preconditioning
[y, f1] = tfqmr (A / M, b, [], 1)
x = M \ y # compare x and x_ref
```

References:

1. Y. Saad, *Iterative Methods for Sparse Linear Systems*, Second edition, 2003, SIAM

See also: [\[bicg\]](#), page 570, [\[bicgstab\]](#), page 572, [\[cgs\]](#), page 574, [\[gmres\]](#), page 576, [\[pcg\]](#), page 647, [\[qmr\]](#), page 579, [\[pcr\]](#), page 650.

19 Vectorization and Faster Code Execution

Vectorization is a programming technique that uses vector operations instead of element-by-element loop-based operations. Besides frequently producing more succinct Octave code, vectorization also allows for better optimization in the subsequent implementation. The optimizations may occur either in Octave's own Fortran, C, or C++ internal implementation, or even at a lower level depending on the compiler and external numerical libraries used to build Octave. The ultimate goal is to make use of your hardware's vector instructions if possible or to perform other optimizations in software.

Vectorization is not a concept unique to Octave, but it is particularly important because Octave is a matrix-oriented language. Vectorized Octave code will see a dramatic speed up (10X–100X) in most cases.

This chapter discusses vectorization and other techniques for writing faster code.

19.1 Basic Vectorization

To a very good first approximation, the goal in vectorization is to write code that avoids loops and uses whole-array operations. As a trivial example, consider

```
for i = 1:n
    for j = 1:m
        c(i,j) = a(i,j) + b(i,j);
    endfor
endfor
```

compared to the much simpler

```
c = a + b;
```

This isn't merely easier to write; it is also internally much easier to optimize. Octave delegates this operation to an underlying implementation which, among other optimizations, may use special vector hardware instructions or could conceivably even perform the additions in parallel. In general, if the code is vectorized, the underlying implementation has more freedom about the assumptions it can make in order to achieve faster execution.

This is especially important for loops with "cheap" bodies. Often it suffices to vectorize just the innermost loop to get acceptable performance. A general rule of thumb is that the "order" of the vectorized body should be greater or equal to the "order" of the enclosing loop.

As a less trivial example, instead of

```
for i = 1:n-1
    a(i) = b(i+1) - b(i);
endfor
```

write

```
a = b(2:n) - b(1:n-1);
```

This shows an important general concept about using arrays for indexing instead of looping over an index variable. See [Section 8.1 \[Index Expressions\]](#), page 139. Also use

boolean indexing generously. If a condition needs to be tested, this condition can also be written as a boolean index. For instance, instead of

```
for i = 1:n
    if (a(i) > 5)
        a(i) == 20
    endif
endfor
```

write

```
a(a>5) == 20;
```

which exploits the fact that `a > 5` produces a boolean index.

Use elementwise vector operators whenever possible to avoid looping (operators like `.*` and `.^`). See [Section 8.3 \[Arithmetic Ops\], page 149](#). For simple inline functions, the `vectorize` function can do this automatically.

`vectorize (fun)`

Create a vectorized version of the inline function `fun` by replacing all occurrences of `*`, `/`, etc., with `.*`, `./`, etc.

This may be useful, for example, when using inline functions with numerical integration or optimization where a vector-valued function is expected.

```
fcn = vectorize (inline ("x^2 - 1"))
⇒ fcn = f(x) = x.^2 - 1
quadv (fcn, 0, 3)
⇒ 6
```

See also: [\[inline\], page 216](#), [\[formula\], page 217](#), [\[argnames\], page 217](#).

Also exploit broadcasting in these elementwise operators both to avoid looping and unnecessary intermediate memory allocations. See [Section 19.2 \[Broadcasting\], page 585](#).

Use built-in and library functions if possible. Built-in and compiled functions are very fast. Even with an m-file library function, chances are good that it is already optimized, or will be optimized more in a future release.

For instance, even better than

```
a = b(2:n) - b(1:n-1);
```

is

```
a = diff (b);
```

Most Octave functions are written with vector and array arguments in mind. If you find yourself writing a loop with a very simple operation, chances are that such a function already exists. The following functions occur frequently in vectorized code:

- Index manipulation
 - `find`
 - `sub2ind`
 - `ind2sub`
 - `sort`
 - `unique`

- lookup
- ifelse / merge
- Repetition
 - repmat
 - repelems
- Vectorized arithmetic
 - sum
 - prod
 - cumsum
 - cumprod
 - sumsq
 - diff
 - dot
 - cummax
 - cummin
- Shape of higher dimensional arrays
 - reshape
 - resize
 - permute
 - squeeze
 - deal

19.2 Broadcasting

Broadcasting refers to how Octave binary operators and functions behave when their matrix or array operands or arguments differ in size. Since version 3.6.0, Octave now automatically broadcasts vectors, matrices, and arrays when using elementwise binary operators and functions. Broadly speaking, smaller arrays are “broadcast” across the larger one, until they have a compatible shape. The rule is that corresponding array dimensions must either

1. be equal, or
2. one of them must be 1.

In case all dimensions are equal, no broadcasting occurs and ordinary element-by-element arithmetic takes place. For arrays of higher dimensions, if the number of dimensions isn’t the same, then missing trailing dimensions are treated as 1. When one of the dimensions is 1, the array with that singleton dimension gets copied along that dimension until it matches the dimension of the other array. For example, consider

```
x = [1 2 3;
      4 5 6;
      7 8 9];

y = [10 20 30];

x + y
```

Without broadcasting, $x + y$ would be an error because the dimensions do not agree. However, with broadcasting it is as if the following operation were performed:

```
x = [1 2 3
      4 5 6
      7 8 9];

y = [10 20 30
      10 20 30
      10 20 30];

x + y
⇒   11   22   33
     14   25   36
     17   28   39
```

That is, the smaller array of size [1 3] gets copied along the singleton dimension (the number of rows) until it is [3 3]. No actual copying takes place, however. The internal implementation reuses elements along the necessary dimension in order to achieve the desired effect without copying in memory.

Both arrays can be broadcast across each other, for example, all pairwise differences of the elements of a vector with itself:

```
y - y'
⇒   0   10   20
    -10   0   10
    -20  -10   0
```

Here the vectors of size [1 3] and [3 1] both get broadcast into matrices of size [3 3] before ordinary matrix subtraction takes place.

A special case of broadcasting that may be familiar is when all dimensions of the array being broadcast are 1, i.e., the array is a scalar. Thus for example, operations like $x - 42$ and `max(x, 2)` are basic examples of broadcasting.

For a higher-dimensional example, suppose `img` is an RGB image of size [$m n 3$] and we wish to multiply each color by a different scalar. The following code accomplishes this with broadcasting,

```
img .*= permute ([0.8, 0.9, 1.2], [1, 3, 2]);
```

Note the usage of `permute` to match the dimensions of the [0.8, 0.9, 1.2] vector with `img`.

For functions that are not written with broadcasting semantics, `bsxfun` can be useful for coercing them to broadcast.

`bsxfun (f, A, B)`

Apply a binary function f element-by-element to two array arguments A and B , expanding singleton dimensions in either input argument as necessary.

f is a function handle, inline function, or string containing the name of the function to evaluate. The function f must be capable of accepting two column-vector arguments of equal length, or one column vector argument and a scalar.

The dimensions of A and B must be equal or singleton. The singleton dimensions of the arrays will be expanded to the same dimensionality as the other array.

See also: [arrayfun], page 588, [cellfun], page 590.

Broadcasting is only applied if either of the two broadcasting conditions hold. As usual, however, broadcasting does not apply when two dimensions differ and neither is 1:

```
x = [1 2 3
      4 5 6];
y = [10 20
      30 40];
x + y
```

This will produce an error about nonconformant arguments.

Besides common arithmetic operations, several functions of two arguments also broadcast. The full list of functions and operators that broadcast is

plus	+	.	+							
minus	-	.	-							
times	.	*								
rdivide	.	/								
ldivide	.	\								
power	.	^	.	**						
lt	<									
le	<=									
eq	==									
gt	>									
ge	>=									
ne	!=	~	=							
and	&									
or										
atan2										
hypot										
max										
min										
mod										
rem										
xor										
+ =	- =	.+ =	.- =	.*=	./=	.\ =	.^ =	.** =	& =	=

Beware of resorting to broadcasting if a simpler operation will suffice. For matrices a and b , consider the following:

```
c = sum (permute (a, [1, 3, 2]) .* permute (b, [3, 2, 1]), 3);
```

This operation broadcasts the two matrices with permuted dimensions across each other during elementwise multiplication in order to obtain a larger 3-D array, and this array is then summed along the third dimension. A moment of thought will prove that this operation is simply the much faster ordinary matrix multiplication, $c = a*b;$.

A note on terminology: “broadcasting” is the term popularized by the Numpy numerical environment in the Python programming language. In other programming languages and environments, broadcasting may also be known as *binary singleton expansion* (BSX, in MATLAB, and the origin of the name of the `bsxfun` function), *recycling* (R programming language), *single-instruction multiple data* (SIMD), or *replication*.

19.2.1 Broadcasting and Legacy Code

The new broadcasting semantics almost never affect code that worked in previous versions of Octave. Consequently, all code inherited from MATLAB that worked in previous versions of Octave should still work without change in Octave. The only exception is code such as

```
try
  c = a.*b;
catch
  c = a.*a;
end_tryCatch
```

that may have relied on matrices of different size producing an error. Because such operation is now valid Octave syntax, this will no longer produce an error. Instead, the following code should be used:

```
if (isequal (size (a), size (b)))
  c = a .* b;
else
  c = a .* a;
endif
```

19.3 Function Application

As a general rule, functions should already be written with matrix arguments in mind and should consider whole matrix operations in a vectorized manner. Sometimes, writing functions in this way appears difficult or impossible for various reasons. For those situations, Octave provides facilities for applying a function to each element of an array, cell, or struct.

```
arrayfun (func, A)
x = arrayfun (func, A)
x = arrayfun (func, A, b, ...)
[x, y, ...] = arrayfun (func, A, ...)
arrayfun (..., "UniformOutput", val)
arrayfun (..., "ErrorHandler", errfunc)
```

Execute a function on each element of an array.

This is useful for functions that do not accept array arguments. If the function does accept array arguments it is better to call the function directly.

The first input argument `func` can be a string, a function handle, an inline function, or an anonymous function. The input argument `A` can be a logic array, a numeric array, a string array, a structure array, or a cell array. By a call of the function `arrayfun` all elements of `A` are passed on to the named function `func` individually.

The named function can also take more than two input arguments, with the input arguments given as third input argument `b`, fourth input argument `c`, ... If given

more than one array input argument then all input arguments must have the same sizes, for example:

```
arrayfun (@atan2, [1, 0], [0, 1])
⇒ [ 1.57080 0.00000 ]
```

If the parameter *val* after a further string input argument "UniformOutput" is set **true** (the default), then the named function *func* must return a single element which then will be concatenated into the return value and is of type matrix. Otherwise, if that parameter is set to **false**, then the outputs are concatenated in a cell array. For example:

```
arrayfun (@(x,y) x:y, "abc", "def", "UniformOutput", false)
⇒
{
    [1,1] = abcd
    [1,2] = bcde
    [1,3] = cdef
}
```

If more than one output arguments are given then the named function must return the number of return values that also are expected, for example:

```
[A, B, C] = arrayfun (@find, [10; 0], "UniformOutput", false)
⇒
A =
{
    [1,1] = 1
    [2,1] = [](0x0)
}
B =
{
    [1,1] = 1
    [2,1] = [](0x0)
}
C =
{
    [1,1] = 10
    [2,1] = [](0x0)
}
```

If the parameter *errfunc* after a further string input argument "ErrorHandler" is another string, a function handle, an inline function, or an anonymous function, then *errfunc* defines a function to call in the case that *func* generates an error. The definition of the function must be of the form

```
function [...] = errfunc (s, ...)
```

where there is an additional input argument to *errfunc* relative to *func*, given by *s*. This is a structure with the elements "identifier", "message", and "index" giving, respectively, the error identifier, the error message, and the index of the array elements that caused the error. The size of the output argument of *errfunc* must have

the same size as the output argument of *func*, otherwise a real error is thrown. For example:

```
function y = ferr (s, x), y = "MyString"; endfunction
arrayfun (@str2num, [1234],
          "UniformOutput", false, "ErrorHandler", @ferr)
⇒
{
  [1,1] = MyString
}
```

See also: [\[spfun\]](#), page 590, [\[cellfun\]](#), page 590, [\[structfun\]](#), page 592.

y = spfun (f, S)

Compute $f(S)$ for the nonzero values of S .

This results in a sparse matrix with the same structure as S . The function f can be passed as a string, a function handle, or an inline function.

See also: [\[arrayfun\]](#), page 588, [\[cellfun\]](#), page 590, [\[structfun\]](#), page 592.

```
cellfun (name, C)
cellfun ("size", C, k)
cellfun ("isclass", C, class)
cellfun (func, C)
cellfun (func, C, D)
[a, ...] = cellfun (...)
cellfun (..., "ErrorHandler", errfunc)
cellfun (..., "UniformOutput", val)
```

Evaluate the function named *name* on the elements of the cell array *C*.

Elements in *C* are passed on to the named function individually. The function *name* can be one of the functions

isempty Return 1 for empty elements.

islogical Return 1 for logical elements.

isnumeric Return 1 for numeric elements.

isreal Return 1 for real elements.

length Return a vector of the lengths of cell elements.

ndims Return the number of dimensions of each element.

numel Return the number of elements contained within each cell element. The number is the product of the dimensions of the object at each cell element.

prodofsize Return the size along the *k*-th dimension.

isclass Return 1 for elements of *class*.

Additionally, `cellfun` accepts an arbitrary function *func* in the form of an inline function, function handle, or the name of a function (in a character string). The function can take one or more arguments, with the inputs arguments given by *C*, *D*, etc. Equally the function can return one or more output arguments. For example:

```
cellfun ("atan2", {1, 0}, {0, 1})
⇒ [ 1.57080 0.00000 ]
```

The number of output arguments of `cellfun` matches the number of output arguments of the function. The outputs of the function will be collected into the output arguments of `cellfun` like this:

```
function [a, b] = twoouts (x)
    a = x;
    b = x*x;
endfunction
[aa, bb] = cellfun (@twoouts, {1, 2, 3})
⇒
    aa =
        1 2 3
    bb =
        1 4 9
```

Note that per default the output argument(s) are arrays of the same size as the input arguments. Input arguments that are singleton (1x1) cells will be automatically expanded to the size of the other arguments.

If the parameter "`UniformOutput`" is set to true (the default), then the function must return scalars which will be concatenated into the return array(s). If "`UniformOutput`" is false, the outputs are concatenated into a cell array (or cell arrays). For example:

```
cellfun ("tolower", {"Foo", "Bar", "FooBar"},
    "UniformOutput", false)
⇒ {"foo", "bar", "foobar"}
```

Given the parameter "`ErrorHandler`", then `errfunc` defines a function to call in case *func* generates an error. The form of the function is

```
function [...] = errfunc (s, ...)
```

where there is an additional input argument to `errfunc` relative to *func*, given by *s*. This is a structure with the elements "`identifier`", "`message`", and "`index`" giving respectively the error identifier, the error message, and the index into the input arguments of the element that caused the error. For example:

```
function y = foo (s, x), y = NaN; endfunction
cellfun ("factorial", {-1,2}, "ErrorHandler", @foo)
⇒ [NaN 2]
```

Use `cellfun` intelligently. The `cellfun` function is a useful tool for avoiding loops. It is often used with anonymous function handles; however, calling an anonymous function involves an overhead quite comparable to the overhead of an m-file function. Passing a handle to a built-in function is faster, because the interpreter is not involved in the internal loop. For example:

```
a = {...}
v = cellfun (@(x) det (x), a); # compute determinants
v = cellfun (@det, a); # faster
```

See also: [\[arrayfun\]](#), page 588, [\[structfun\]](#), page 592, [\[spfun\]](#), page 590.

```
structfun (func, S)
[A, ...] = structfun (...)

structfun (... , "ErrorHandler", errfunc)
structfun (... , "UniformOutput", val)
```

Evaluate the function named *name* on the fields of the structure *S*. The fields of *S* are passed to the function *func* individually.

structfun accepts an arbitrary function *func* in the form of an inline function, function handle, or the name of a function (in a character string). In the case of a character string argument, the function must accept a single argument named *x*, and it must return a string value. If the function returns more than one argument, they are returned as separate output variables.

If the parameter "UniformOutput" is set to true (the default), then the function must return a single element which will be concatenated into the return value. If "UniformOutput" is false, the outputs are placed into a structure with the same fieldnames as the input structure.

```
s.name1 = "John Smith";
s.name2 = "Jill Jones";
structfun (@(x) regexp (x, '(\w+)$', "matches"){1}, s,
           "UniformOutput", false)
⇒ scalar structure containing the fields:
    name1 = Smith
    name2 = Jones
```

Given the parameter "ErrorHandler", *errfunc* defines a function to call in case *func* generates an error. The form of the function is

```
function [...] = errfunc (se, ...)
```

where there is an additional input argument to *errfunc* relative to *func*, given by *se*. This is a structure with the elements "identifier", "message" and "index", giving respectively the error identifier, the error message, and the index into the input arguments of the element that caused the error. For an example on how to use an error handler, see [\[cellfun\]](#), page 590.

See also: [\[cellfun\]](#), page 590, [\[arrayfun\]](#), page 588, [\[spfun\]](#), page 590.

Consistent with earlier advice, seek to use Octave built-in functions whenever possible for the best performance. This advice applies especially to the four functions above. For example, when adding two arrays together element-by-element one could use a handle to the built-in addition function `@plus` or define an anonymous function `@(x,y) x + y`. But, the anonymous function is 60% slower than the first method. See [Section 34.4.2 \[Operator Overloading\]](#), page 846, for a list of basic functions which might be used in place of anonymous ones.

19.4 Accumulation

Whenever it's possible to categorize according to indices the elements of an array when performing a computation, accumulation functions can be useful.

```
accumarray (subs, vals, sz, func, fillval, issparse)
accumarray (subs, vals, ...)
```

Create an array by accumulating the elements of a vector into the positions defined by their subscripts.

The subscripts are defined by the rows of the matrix *subs* and the values by *vals*. Each row of *subs* corresponds to one of the values in *vals*. If *vals* is a scalar, it will be used for each of the row of *subs*. If *subs* is a cell array of vectors, all vectors must be of the same length, and the subscripts in the *k*th vector must correspond to the *k*th dimension of the result.

The size of the matrix will be determined by the subscripts themselves. However, if *sz* is defined it determines the matrix size. The length of *sz* must correspond to the number of columns in *subs*. An exception is if *subs* has only one column, in which case *sz* may be the dimensions of a vector and the subscripts of *subs* are taken as the indices into it.

The default action of **accumarray** is to sum the elements with the same subscripts. This behavior can be modified by defining the *func* function. This should be a function or function handle that accepts a column vector and returns a scalar. The result of the function should not depend on the order of the subscripts.

The elements of the returned array that have no subscripts associated with them are set to zero. Defining *fillval* to some other value allows these values to be defined. This behavior changes, however, for certain values of *func*. If *func* is `@min` (respectively, `@max`) then the result will be filled with the minimum (respectively, maximum) integer if *vals* is of integral type, logical false (respectively, logical true) if *vals* is of logical type, zero if *fillval* is zero and all values are non-positive (respectively, non-negative), and NaN otherwise.

By default **accumarray** returns a full matrix. If *issparse* is logically true, then a sparse matrix is returned instead.

The following **accumarray** example constructs a frequency table that in the first column counts how many occurrences each number in the second column has, taken from the vector *x*. Note the usage of **unique** for assigning to all repeated elements of *x* the same index (see [**unique**], page 743).

```
x = [91, 92, 90, 92, 90, 89, 91, 89, 90, 100, 100, 100];
[u, ~, j] = unique (x);
[accumarray(j', 1), u']
⇒ 2    89
    3    90
    2    91
    2    92
    3    100
```

Another example, where the result is a multi-dimensional 3-D array and the default value (zero) appears in the output:

```
accumarray ([1, 1, 1;
            2, 1, 2;
            2, 3, 2;
            2, 1, 2;
            2, 3, 2], 101:105)
⇒ ans(:,:,1) = [101, 0, 0; 0, 0, 0]
⇒ ans(:,:,2) = [0, 0, 0; 206, 0, 208]
```

The sparse option can be used as an alternative to the `sparse` constructor (see [[sparse](#)], [page 622](#)). Thus

```
sparse (i, j, sv)
```

can be written with `accumarray` as

```
accumarray ([i, j], sv', [], [], 0, true)
```

For repeated indices, `sparse` adds the corresponding value. To take the minimum instead, use `min` as an accumulator function:

```
accumarray ([i, j], sv', [], @min, 0, true)
```

The complexity of `accumarray` in general for the non-sparse case is generally $O(M+N)$, where N is the number of subscripts and M is the maximum subscript (linearized in multi-dimensional case). If `func` is one of `@sum` (default), `@max`, `@min` or `@(x) {x}`, an optimized code path is used. Note that for general reduction function the interpreter overhead can play a major part and it may be more efficient to do multiple `accumarray` calls and compute the results in a vectorized manner.

See also: [[accumdim](#)], [page 594](#), [[unique](#)], [page 743](#), [[sparse](#)], [page 622](#).

`accumdim (subs, vals, dim, n, func, fillval)`

Create an array by accumulating the slices of an array into the positions defined by their subscripts along a specified dimension.

The subscripts are defined by the index vector `subs`. The dimension is specified by `dim`. If not given, it defaults to the first non-singleton dimension. The length of `subs` must be equal to `size (vals, dim)`.

The extent of the result matrix in the working dimension will be determined by the subscripts themselves. However, if `n` is defined it determines this extent.

The default action of `accumdim` is to sum the subarrays with the same subscripts. This behavior can be modified by defining the `func` function. This should be a function or function handle that accepts an array and a dimension, and reduces the array along this dimension. As a special exception, the built-in `min` and `max` functions can be used directly, and `accumdim` accounts for the middle empty argument that is used in their calling.

The slices of the returned array that have no subscripts associated with them are set to zero. Defining `fillval` to some other value allows these values to be defined.

An example of the use of `accumdim` is:

```

accumdim ([1, 2, 1, 2, 1], [ 7, -10,    4;
                           -5, -12,    8;
                           -12,   2,    8;
                           -10,   9,   -3;
                           -5,  -3,  -13])
⇒ [-10,-11,-1;-15,-3,5]

```

See also: [\[accumarray\]](#), page 593.

19.5 JIT Compiler

Vectorization is the preferred technique for eliminating loops and speeding up code. Nevertheless, it is not always possible to replace every loop. In such situations it may be worth trying Octave's **experimental** Just-In-Time (JIT) compiler.

A JIT compiler works by analyzing the body of a loop, translating the Octave statements into another language, compiling the new code segment into an executable, and then running the executable and collecting any results. The process is not simple and there is a significant amount of work to perform for each step. It can still make sense, however, if the number of loop iterations is large. Because Octave is an interpreted language every time through a loop Octave must parse the statements in the loop body before executing them. With a JIT compiler this is done just once when the body is translated to another language.

The JIT compiler is a very new feature in Octave and not all valid Octave statements can currently be accelerated. However, if no other technique is available it may be worth benchmarking the code with JIT enabled. The function `jit_enable` is used to turn compilation on or off. The function `jit_startcnt` sets the threshold for acceleration. Loops with iteration counts above `jit_startcnt` will be accelerated. The functions `jit_failcnt` and `debug_jit` are not likely to be of use to anyone not working directly on the implementation of the JIT compiler.

```

val = jit_enable ()
old_val = jit_enable (new_val)
jit_enable (new_val, "local")

```

Query or set the internal variable that enables Octave's JIT compiler.

When called from inside a function with the "`local`" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[jit_startcnt\]](#), page 595, [\[debug_jit\]](#), page 596.

```

val = jit_startcnt ()
old_val = jit_startcnt (new_val)
jit_startcnt (new_val, "local")

```

Query or set the internal variable that determines whether JIT compilation will take place for a specific loop.

Because compilation is a costly operation it does not make sense to employ JIT when the loop count is low. By default only loops with greater than 1000 iterations will be accelerated.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[jit_enable\]](#), page 595, [\[jit_failcnt\]](#), page 596, [\[debug_jit\]](#), page 596.

```
val = jit_failcnt ()
old_val = jit_failcnt (new_val)
jit_failcnt (new_val, "local")
```

Query or set the internal variable that counts the number of JIT fail exceptions for Octave's JIT compiler.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[jit_enable\]](#), page 595, [\[jit_startcnt\]](#), page 595, [\[debug_jit\]](#), page 596.

```
val = debug_jit ()
old_val = debug_jit (new_val)
debug_jit (new_val, "local")
```

Query or set the internal variable that determines whether debugging/tracing is enabled for Octave's JIT compiler.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[jit_enable\]](#), page 595, [\[jit_startcnt\]](#), page 595.

19.6 Miscellaneous Techniques

Here are some other ways of improving the execution speed of Octave programs.

- Avoid computing costly intermediate results multiple times. Octave currently does not eliminate common subexpressions. Also, certain internal computation results are cached for variables. For instance, if a matrix variable is used multiple times as an index, checking the indices (and internal conversion to integers) is only done once.
- Be aware of lazy copies (copy-on-write). When a copy of an object is created, the data is not immediately copied, but rather shared. The actual copying is postponed until the copied data needs to be modified. For example:

```
a = zeros (1000); # create a 1000x1000 matrix
b = a; # no copying done here
b(1) = 1; # copying done here
```

Lazy copying applies to whole Octave objects such as matrices, cells, struct, and also individual cell or struct elements (not array elements).

Additionally, index expressions also use lazy copying when Octave can determine that the indexed portion is contiguous in memory. For example:

```
a = zeros (1000); # create a 1000x1000 matrix
b = a(:,10:100); # no copying done here
b = a(10:100,:); # copying done here
```

This applies to arrays (matrices), cell arrays, and structs indexed using ‘()’. Index expressions generating comma-separated lists can also benefit from shallow copying in some cases. In particular, when *a* is a struct array, expressions like *{a.x}*, *{a(:,2).x}* will use lazy copying, so that data can be shared between a struct array and a cell array.

Most indexing expressions do not live longer than their parent objects. In rare cases, however, a lazily copied slice outlasts its parent, in which case it becomes orphaned, still occupying unnecessarily more memory than needed. To provide a remedy working in most real cases, Octave checks for orphaned lazy slices at certain situations, when a value is stored into a "permanent" location, such as a named variable or cell or struct element, and possibly economizes them. For example:

```
a = zeros (1000); # create a 1000x1000 matrix
b = a(:,10:100); # lazy slice
a = []; # the original "a" array is still allocated
c{1} = b; # b is reallocated at this point
```

- Avoid deep recursion. Function calls to m-file functions carry a relatively significant overhead, so rewriting a recursion as a loop often helps. Also, note that the maximum level of recursion is limited.
- Avoid resizing matrices unnecessarily. When building a single result matrix from a series of calculations, set the size of the result matrix first, then insert values into it. Write

```
result = zeros (big_n, big_m)
for i = over:and_over
    ridx = ...
    cidx = ...
    result(ridx, cidx) = new_value ();
endfor
```

instead of

```
result = [];
for i = ever:and_ever
    result = [ result, new_value() ];
endfor
```

Sometimes the number of items can not be computed in advance, and stack-like operations are needed. When elements are being repeatedly inserted or removed from the end of an array, Octave detects it as stack usage and attempts to use a smarter memory management strategy by pre-allocating the array in bigger chunks. This strategy is also applied to cell and struct arrays.

```
a = [];
while (condition)
    ...
    a(end+1) = value; # "push" operation
    ...
    a(end) = [] ; # "pop" operation
    ...
endwhile
```

- Avoid calling `eval` or `feval` excessively. Parsing input or looking up the name of a function in the symbol table are relatively expensive operations.

If you are using `eval` merely as an exception handling mechanism, and not because you need to execute some arbitrary text, use the `try` statement instead. See [Section 10.9 \[The try Statement\]](#), page 174.

- Use `ignore_function_time_stamp` when appropriate. If you are calling lots of functions, and none of them will need to change during your run, set the variable `ignore_function_time_stamp` to "all". This will stop Octave from checking the time stamp of a function file to see if it has been updated while the program is being run.

19.7 Examples

The following are examples of vectorization questions asked by actual users of Octave and their solutions.

- For a vector `A`, the following loop

```
n = length (A);
B = zeros (n, 2);
for i = 1:length (A)
    ## this will be two columns, the first is the difference and
    ## the second the mean of the two elements used for the diff.
    B(i,:) = [A(i+1)-A(i), (A(i+1) + A(i))/2];
endfor
```

can be turned into the following one-liner:

```
B = [diff(A) (:), 0.5*(A(1:end-1)+A(2:end)) (:)]
```

Note the usage of colon indexing to flatten an intermediate result into a column vector. This is a common vectorization trick.

20 Nonlinear Equations

20.1 Solvers

Octave can solve sets of nonlinear equations of the form

$$f(x) = 0$$

using the function `fsolve`, which is based on the MINPACK subroutine `hybrd`. This is an iterative technique so a starting point must be provided. This also has the consequence that convergence is not guaranteed even if a solution exists.

```
fsolve (fcn, x0)
fsolve (fcn, x0, options)
[x, fval, info, output, fjac] = fsolve (...)
```

Solve a system of nonlinear equations defined by the function `fcn`.

`fcn` should accept a vector (array) defining the unknown variables, and return a vector of left-hand sides of the equations. Right-hand sides are defined to be zeros. In other words, this function attempts to determine a vector `x` such that `fcn (x)` gives (approximately) all zeros.

`x0` is an initial guess for the solution. The shape of `x0` is preserved in all calls to `fcn`, but otherwise is treated as a column vector.

`options` is a structure specifying additional parameters which control the algorithm. Currently, `fsolve` recognizes these options: "AutoScaling", "ComplexEqn", "FinDiffType", "FunValCheck", "Jacobian", "MaxFunEvals", "MaxIter", "OutputFcn", "TolFun", "TolX", "TypicalX", and "Updating".

If "AutoScaling" is "on", the variables will be automatically scaled according to the column norms of the (estimated) Jacobian. As a result, "TolFun" becomes scaling-independent. By default, this option is "off" because it may sometimes deliver unexpected (though mathematically correct) results.

If "ComplexEqn" is "on", `fsolve` will attempt to solve complex equations in complex variables, assuming that the equations possess a complex derivative (i.e., are holomorphic). If this is not what you want, you should unpack the real and imaginary parts of the system to get a real system.

If "Jacobian" is "on", it specifies that `fcn`—when called with 2 output arguments—also returns the Jacobian matrix of right-hand sides at the requested point.

"MaxFunEvals" proscribes the maximum number of function evaluations before optimization is halted. The default value is `100 * number_of_variables`, i.e., `100 * length (x0)`. The value must be a positive integer.

If "Updating" is "on", the function will attempt to use Broyden updates to update the Jacobian, in order to reduce the number of Jacobian calculations. If your user function always calculates the Jacobian (regardless of number of output arguments) then this option provides no advantage and should be disabled.

"TolX" specifies the termination tolerance in the unknown variables, while "TolFun" is a tolerance for equations. Default is `1e-6` for both "TolX" and "TolFun".

For a description of the other options, see `optimset`. To initialize an options structure with default values for `fsolve` use `options = optimset ("fsolve")`.

The first output `x` is the solution while the second output `fval` contains the value of the function `fcn` evaluated at `x` (ideally a vector of all zeros).

The third output `info` reports whether the algorithm succeeded and may take one of the following values:

- 1 Converged to a solution point. Relative residual error is less than specified by `TolFun`.
- 2 Last relative step size was less than `TolX`.
- 3 Last relative decrease in residual was less than `TolFun`.
- 0 Iteration limit (either `MaxIter` or `MaxFunEvals`) exceeded.
- 1 Stopped by `OutputFcn`.
- 3 The trust region radius became excessively small.

`output` is a structure containing runtime information about the `fsolve` algorithm. Fields in the structure are:

iterations
Number of iterations through loop.

successful
Number of successful iterations.

funcCount
Number of function evaluations.

The final output `fjac` contains the value of the Jacobian evaluated at `x`.

Note: If you only have a single nonlinear equation of one variable, using `fzero` is usually a much better idea.

Note about user-supplied Jacobians: As an inherent property of the algorithm, a Jacobian is always requested for a solution vector whose residual vector is already known, and it is the last accepted successful step. Often this will be one of the last two calls, but not always. If the savings by reusing intermediate results from residual calculation in Jacobian calculation are significant, the best strategy is to employ `OutputFcn`: After a vector is evaluated for residuals, if `OutputFcn` is called with that vector, then the intermediate results should be saved for future Jacobian evaluation, and should be kept until a Jacobian evaluation is requested or until `OutputFcn` is called with a different vector, in which case they should be dropped in favor of this most recent vector. A short example how this can be achieved follows:

```
function [fval, fjac] = user_func (x, optimvalues, state)
persistent sav = [], sav0 = [];
if (nargin == 1)
    ## evaluation call
    if (nargout == 1)
        sav0.x = x; # mark saved vector
        ## calculate fval, save results to sav0.
```

```

elseif (nargout == 2)
    ## calculate fjac using sav.
    endif
else
    ## outputfcn call.
    if (all (x == sav0.x))
        sav = sav0;
    endif
    ## maybe output iteration status, etc.
endif
endfunction

## ...

fsolve (@user_func, x0, optimset ("OutputFcn", @user_func, ...))

```

See also: [fzero], page 602, [optimset], page 709.

The following is a complete example. To solve the set of equations

$$\begin{aligned} -2x^2 + 3xy + 4\sin(y) - 6 &= 0 \\ 3x^2 - 2xy^2 + 3\cos(x) + 4 &= 0 \end{aligned}$$

you first need to write a function to compute the value of the given function. For example:

```

function y = f (x)
y = zeros (2, 1);
y(1) = -2*x(1)^2 + 3*x(1)*x(2) + 4*sin(x(2)) - 6;
y(2) = 3*x(1)^2 - 2*x(1)*x(2)^2 + 3*cos(x(1)) + 4;
endfunction

```

Then, call `fsolve` with a specified initial condition to find the roots of the system of equations. For example, given the function `f` defined above,

```
[x, fval, info] = fsolve (@f, [1; 2])
```

results in the solution

```

x =
0.57983
2.54621

fval =
-5.7184e-10
5.5460e-10

info = 1

```

A value of `info = 1` indicates that the solution has converged.

When no Jacobian is supplied (as in the example above) it is approximated numerically. This requires more function evaluations, and hence is less efficient. In the example above we could compute the Jacobian analytically as

$$\begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} \end{bmatrix} = \begin{bmatrix} 3x_2 - 4x_1 & 4\cos(x_2) + 3x_1 \\ -2x_2^2 - 3\sin(x_1) + 6x_1 & -4x_1x_2 \end{bmatrix}$$

and compute it with the following Octave function

```
function [y, jac] = f (x)
    y = zeros (2, 1);
    y(1) = -2*x(1)^2 + 3*x(1)*x(2) + 4*sin(x(2)) - 6;
    y(2) = 3*x(1)^2 - 2*x(1)*x(2)^2 + 3*cos(x(1)) + 4;
    if (nargout == 2)
        jac = zeros (2, 2);
        jac(1,1) = 3*x(2) - 4*x(1);
        jac(1,2) = 4*cos(x(2)) + 3*x(1);
        jac(2,1) = -2*x(2)^2 - 3*sin(x(1)) + 6*x(1);
        jac(2,2) = -4*x(1)*x(2);
    endif
endfunction
```

The Jacobian can then be used with the following call to `fsolve`:

```
[x, fval, info] = fsolve (@f, [1; 2], optimset ("jacobian", "on"));
```

which gives the same solution as before.

```
fzero (fun, x0)
fzero (fun, x0, options)
[x, fval, info, output] = fzero (...)
```

Find a zero of a univariate function.

`fun` is a function handle, inline function, or string containing the name of the function to evaluate.

`x0` should be a two-element vector specifying two points which bracket a zero. In other words, there must be a change in sign of the function between `x0(1)` and `x0(2)`. More mathematically, the following must hold

```
sign (fun(x0(1))) * sign (fun(x0(2))) <= 0
```

If `x0` is a single scalar then several nearby and distant values are probed in an attempt to obtain a valid bracketing. If this is not successful, the function fails.

`options` is a structure specifying additional options. Currently, `fzero` recognizes these options: "FunValCheck", "MaxFunEvals", "MaxIter", "OutputFcn", and "TolX".

"`MaxFunEvals`" proscribes the maximum number of function evaluations before the search is halted. The default value is `Inf`. The value must be a positive integer.

"`MaxIter`" proscribes the maximum number of algorithm iterations before the search is halted. The default value is `Inf`. The value must be a positive integer.

"`TolX`" specifies the termination tolerance for the solution `x`. The default value is `eps`.

For a description of the other options, see [\[optimset\], page 709](#). To initialize an options structure with default values for `fzero` use `options = optimset ("fzero")`.

On exit, the function returns `x`, the approximate zero point, and `fval`, the function evaluated at `x`.

The third output `info` reports whether the algorithm succeeded and may take one of the following values:

- 1 The algorithm converged to a solution.
- 0 Maximum number of iterations or function evaluations has been reached.
- -1 The algorithm has been terminated by a user `OutputFcn`.
- -5 The algorithm may have converged to a singular point.

`output` is a structure containing runtime information about the `fzero` algorithm. Fields in the structure are:

- `iterations` Number of iterations through loop.
- `funcCount` Number of function evaluations.
- `algorithm` The string "`bisection, interpolation`".
- `bracketx` A two-element vector with the final bracketing of the zero along the x-axis.
- `brackety` A two-element vector with the final bracketing of the zero along the y-axis.

See also: [\[optimset\], page 709](#), [\[fsolve\], page 599](#).

20.2 Minimizers

Often it is useful to find the minimum value of a function rather than just the zeroes where it crosses the x-axis. `fminbnd` is designed for the simpler, but very common, case of a univariate function where the interval to search is bounded. For unbounded minimization of a function with potentially many variables use `fminunc` or `fminsearch`. The two functions use different internal algorithms and some knowledge of the objective function is required. For functions which can be differentiated, `fminunc` is appropriate. For functions with discontinuities, or for which a gradient search would fail, use `fminsearch`. See [Chapter 25 \[Optimization\], page 697](#), for minimization with the presence of constraint functions. Note that searches can be made for maxima by simply inverting the objective function ($F_{max} = -F_{min}$).

```
x = fminbnd (fun, a, b)
x = fminbnd (fun, a, b, options)
[x, fval, info, output] = fminbnd (...)
```

Find a minimum point of a univariate function.

`fun` must be a function handle or name.

The starting interval is specified by `a` (left boundary) and `b` (right boundary). The endpoints must be finite.

`options` is a structure specifying additional parameters which control the algorithm. Currently, `fminbnd` recognizes these options: "Display", "FunValCheck", "MaxFunEvals", "MaxIter", "OutputFcn", "TolX".

"**MaxFunEvals**" proscribes the maximum number of function evaluations before optimization is halted. The default value is 500. The value must be a positive integer.

"**MaxIter**" proscribes the maximum number of algorithm iterations before optimization is halted. The default value is 500. The value must be a positive integer.

"**TolX**" specifies the termination tolerance for the solution x . The default is $1e-4$.

For a description of the other options, see [\[optimset\], page 709](#). To initialize an options structure with default values for **fminbnd** use `options = optimset ("fminbnd")`.

On exit, the function returns x , the approximate minimum point, and $fval$, the function evaluated x .

The third output *info* reports whether the algorithm succeeded and may take one of the following values:

- 1 The algorithm converged to a solution.
- 0 Iteration limit (either **MaxIter** or **MaxFunEvals**) exceeded.
- -1 The algorithm was terminated by a user **OutputFcn**.

Programming Notes: The search for a minimum is restricted to be in the finite interval bound by a and b . If you have only one initial point to begin searching from then you will need to use an unconstrained minimization algorithm such as **fminunc** or **fminsearch**. **fminbnd** internally uses a Golden Section search strategy.

See also: [\[fzero\], page 602](#), [\[fminunc\], page 604](#), [\[fminsearch\], page 605](#), [\[optimset\], page 709](#).

```
fminunc (fcn, x0)
fminunc (fcn, x0, options)
[x, fval, info, output, grad, hess] = fminunc (fcn, ...)
Solve an unconstrained optimization problem defined by the function fcn.
fcn should accept a vector (array) defining the unknown variables, and return the objective function value, optionally with gradient. fminunc attempts to determine a vector x such that fcn (x) is a local minimum.
x0 determines a starting guess. The shape of x0 is preserved in all calls to fcn, but otherwise is treated as a column vector.
```

options is a structure specifying additional parameters which control the algorithm. Currently, **fminunc** recognizes these options: "**AutoScaling**", "**FinDiffType**", "**FunValCheck**", "**GradObj**", "**MaxFunEvals**", "**MaxIter**", "**OutputFcn**", "**TolFun**", "**TolX**", "**TypicalX**".

If "**AutoScaling**" is "on", the variables will be automatically scaled according to the column norms of the (estimated) Jacobian. As a result, "**TolFun**" becomes scaling-independent. By default, this option is "off" because it may sometimes deliver unexpected (though mathematically correct) results.

If "**GradObj**" is "on", it specifies that *fcn*—when called with two output arguments—also returns the Jacobian matrix of partial first derivatives at the requested point.

"**MaxFunEvals**" proscribes the maximum number of function evaluations before optimization is halted. The default value is `100 * number_of_variables`, i.e., `100 * length (x0)`. The value must be a positive integer.

"**MaxIter**" proscribes the maximum number of algorithm iterations before optimization is halted. The default value is 400. The value must be a positive integer.

"**TolX**" specifies the termination tolerance for the unknown variables *x*, while "**TolFun**" is a tolerance for the objective function value *fval*. The default is **1e-6** for both options.

For a description of the other options, see [optimset](#).

On return, *x* is the location of the minimum and *fval* contains the value of the objective function at *x*.

info may be one of the following values:

- 1 Converged to a solution point. Relative gradient error is less than specified by **TolFun**.
- 2 Last relative step size was less than **TolX**.
- 3 Last relative change in function value was less than **TolFun**.
- 0 Iteration limit exceeded—either maximum number of algorithm iterations **MaxIter** or maximum number of function evaluations **MaxFunEvals**.
- 1 Algorithm terminated by **OutputFcn**.
- 3 The trust region radius became excessively small.

Optionally, **fminunc** can return a structure with convergence statistics (*output*), the output gradient (*grad*) at the solution *x*, and approximate Hessian (*hess*) at the solution *x*.

Application Notes: If the objective function is a single nonlinear equation of one variable then using **fminbnd** is usually a better choice.

The algorithm used by **fminunc** is a gradient search which depends on the objective function being differentiable. If the function has discontinuities it may be better to use a derivative-free algorithm such as **fminsearch**.

See also: [\[fminbnd\]](#), page 603, [\[fminsearch\]](#), page 605, [\[optimset\]](#), page 709.

```
x = fminsearch (fun, x0)
x = fminsearch (fun, x0, options)
x = fminsearch (fun, x0, options, fun_arg1, fun_arg2, ...)
x = fminsearch (problem)
[x, fval, exitflag, output] = fminsearch (...)
```

Find a value of *x* which minimizes the multi-variable function *fun*.

The search begins at the point *x0* and iterates using the Nelder & Mead Simplex algorithm (a derivative-free method). This algorithm is better-suited to functions which have discontinuities or for which a gradient-based search such as **fminunc** fails.

Options for the search are provided in the parameter *options* using the function [optimset](#). Currently, **fminsearch** accepts the options: "**Display**", "**FunValCheck**", "**MaxFunEvals**", "**MaxIter**", "**OutputFcn**", "**TolFun**", "**TolX**".

"**MaxFunEvals**" proscribes the maximum number of function evaluations before optimization is halted. The default value is **200 * number_of_variables**, i.e., **200 * length (x0)**. The value must be a positive integer.

"**MaxIter**" proscribes the maximum number of algorithm iterations before optimization is halted. The default value is `200 * number_of_variables`, i.e., `200 * length (x0)`. The value must be a positive integer.

For a description of the other options, see `optimset`. To initialize an options structure with default values for `fminsearch` use `options = optimset ("fminsearch")`.

Additional inputs for the function *fun* can be passed as the fourth and higher arguments. To pass function arguments while using the default *options* values, use `[]` for *options*.

`fminsearch` may also be called with a single structure argument with the following fields:

objective

The objective function.

x0

The initial point.

solver

Must be set to "fminsearch".

options

A structure returned from `optimset` or an empty matrix to indicate that defaults should be used.

The field `options` is optional. All others are required.

On exit, the function returns *x*, the minimum point, and *fval*, the function value at the minimum.

The third output `exitflag` reports whether the algorithm succeeded and may take one of the following values:

- 1 if the algorithm converged (size of the simplex is smaller than `TolX AND` the step in function value between iterations is smaller than `TolFun`).
- 0 if the maximum number of iterations or the maximum number of function evaluations are exceeded.
- 1 if the iteration is stopped by the "OutputFcn".

The fourth output is a structure *output* containing runtime about the algorithm. Fields in the structure are `funcCount` containing the number of function calls to *fun*, `iterations` containing the number of iteration steps, `algorithm` with the name of the search algorithm (always: "Nelder-Mead simplex direct search"), and `message` with the exit message.

Example:

```
fminsearch (@(x) (x(1)-5).^2+(x(2)-8).^4, [0;0])
```

Note: If you need to find the minimum of a single variable function it is probably better to use `fminbnd`.

See also: `[fminbnd]`, page 603, `[fminunc]`, page 604, `[optimset]`, page 709.

The function `humps` is a useful function for testing zero and extrema finding functions.

```
y = humps (x)
[x, y] = humps (x)
```

Evaluate a function with multiple minima, maxima, and zero crossings.

The output *y* is the evaluation of the rational function:

$$y = -\frac{1200x^4 - 2880x^3 + 2036x^2 - 340x - 88}{200x^4 - 480x^3 + 406x^2 - 138x - 17}$$

x may be a scalar, vector or array. If *x* is omitted, the default range [0:0.05:1] is used.

When called with two output arguments, [x, y], x will contain the input values, and y will contain the output from **humps**.

Programming Notes: **humps** has two local maxima located near *x* = 0.300 and 0.893, a local minimum near *x* = 0.637, and zeros near *x* = -0.132 and 1.300. **humps** is a useful function for testing algorithms which find zeros or local minima and maxima.

Try **demo humps** to see a plot of the **humps** function.

See also: [\[fzero\]](#), page 602, [\[fminbnd\]](#), page 603, [\[fminunc\]](#), page 604, [\[fminsearch\]](#), page 605.

21 Diagonal and Permutation Matrices

21.1 Creating and Manipulating Diagonal/Permutation Matrices

A diagonal matrix is defined as a matrix that has zero entries outside the main diagonal; that is, $D_{ij} = 0$ if $i \neq j$. Most often, square diagonal matrices are considered; however, the definition can equally be applied to non-square matrices, in which case we usually speak of a rectangular diagonal matrix.

A permutation matrix is defined as a square matrix that has a single element equal to unity in each row and each column; all other elements are zero. That is, there exists a permutation (vector) p such that $P_{ij} = 1$ if $j = p_i$ and $P_{ij} = 0$ otherwise.

Octave provides special treatment of real and complex rectangular diagonal matrices, as well as permutation matrices. They are stored as special objects, using efficient storage and algorithms, facilitating writing both readable and efficient matrix algebra expressions in the Octave language. The special treatment may be disabled by using the functions `disable_diagonal_matrix` and `disable_permutation_matrix`.

```
val = disable_diagonal_matrix ()
old_val = disable_diagonal_matrix (new_val)
disable_diagonal_matrix (new_val, "local")
```

Query or set the internal variable that controls whether diagonal matrices are stored in a special space-efficient format.

The default value is true. If this option is disabled Octave will store diagonal matrices as full matrices.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[disable_range\]](#), page 52, [\[disable_permutation_matrix\]](#), page 609.

```
val = disable_permutation_matrix ()
old_val = disable_permutation_matrix (new_val)
disable_permutation_matrix (new_val, "local")
```

Query or set the internal variable that controls whether permutation matrices are stored in a special space-efficient format.

The default value is true. If this option is disabled Octave will store permutation matrices as full matrices.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[disable_range\]](#), page 52, [\[disable_diagonal_matrix\]](#), page 609.

The space savings are significant as demonstrated by the following code.

```
x = diag (rand (10, 1));
xf = full (x);
sizeof (x)
⇒ 80
sizeof (xf)
⇒ 800
```

21.1.1 Creating Diagonal Matrices

The most common and easiest way to create a diagonal matrix is using the built-in function `diag`. The expression `diag (v)`, with `v` a vector, will create a square diagonal matrix with elements on the main diagonal given by the elements of `v`, and size equal to the length of `v`. `diag (v, m, n)` can be used to construct a rectangular diagonal matrix. The result of these expressions will be a special diagonal matrix object, rather than a general matrix object.

Diagonal matrix with unit elements can be created using `eye`. Some other built-in functions can also return diagonal matrices. Examples include `balance` or `inv`.

Example:

```
diag (1:4)
⇒
Diagonal Matrix

1   0   0   0
0   2   0   0
0   0   3   0
0   0   0   4

diag (1:3,5,3)

⇒
Diagonal Matrix

1   0   0
0   2   0
0   0   3
0   0   0
0   0   0
```

21.1.2 Creating Permutation Matrices

For creating permutation matrices, Octave does not introduce a new function, but rather overrides an existing syntax: permutation matrices can be conveniently created by indexing an identity matrix by permutation vectors. That is, if `q` is a permutation vector of length `n`, the expression

```
P = eye (n) (:, q);
```

will create a permutation matrix - a special matrix object.

```
eye (n) (q, :)
```

will also work (and create a row permutation matrix), as well as

```
eye (n) (q1, q2).
```

For example:

```
eye (4) ([1,3,2,4], :)
```

\Rightarrow

Permutation Matrix

1	0	0	0
0	0	1	0
0	1	0	0
0	0	0	1

```
eye (4) (:, [1,3,2,4])
```

\Rightarrow

Permutation Matrix

1	0	0	0
0	0	1	0
0	1	0	0
0	0	0	1

Mathematically, an identity matrix is both diagonal and permutation matrix. In Octave, `eye (n)` returns a diagonal matrix, because a matrix can only have one class. You can convert this diagonal matrix to a permutation matrix by indexing it by an identity permutation, as shown below. This is a special property of the identity matrix; indexing other diagonal matrices generally produces a full matrix.

```
eye (3)
```

\Rightarrow

Diagonal Matrix

1	0	0
0	1	0
0	0	1

```
eye(3)(1:3,:)
```

\Rightarrow

Permutation Matrix

1	0	0
0	1	0
0	0	1

Some other built-in functions can also return permutation matrices. Examples include `inv` or `lu`.

21.1.3 Explicit and Implicit Conversions

The diagonal and permutation matrices are special objects in their own right. A number of operations and built-in functions are defined for these matrices to use special, more efficient

code than would be used for a full matrix in the same place. Examples are given in further sections.

To facilitate smooth mixing with full matrices, backward compatibility, and compatibility with MATLAB, the diagonal and permutation matrices should allow any operation that works on full matrices, and will either treat it specially, or implicitly convert themselves to full matrices.

Instances include matrix indexing, except for extracting a single element or a leading submatrix, indexed assignment, or applying most mapper functions, such as `exp`.

An explicit conversion to a full matrix can be requested using the built-in function `full`. It should also be noted that the diagonal and permutation matrix objects will cache the result of the conversion after it is first requested (explicitly or implicitly), so that subsequent conversions will be very cheap.

21.2 Linear Algebra with Diagonal/Permutation Matrices

As has been already said, diagonal and permutation matrices make it possible to use efficient algorithms while preserving natural linear algebra syntax. This section describes in detail the operations that are treated specially when performed on these special matrix objects.

21.2.1 Expressions Involving Diagonal Matrices

Assume D is a diagonal matrix. If M is a full matrix, then $D*M$ will scale the rows of M . That means, if $S = D*M$, then for each pair of indices i,j it holds

$$S_{ij} = D_{ii}M_{ij}$$

Similarly, $M*D$ will do a column scaling.

The matrix D may also be rectangular, m -by- n where $m \neq n$. If $m < n$, then the expression $D*M$ is equivalent to

$$D(:,1:m) * M(1:m,:),$$

i.e., trailing $n-m$ rows of M are ignored. If $m > n$, then $D*M$ is equivalent to

$$[D(1:n,:) * M; zeros(m-n, columns(M))],$$

i.e., null rows are appended to the result. The situation for right-multiplication $M*D$ is analogous.

The expressions $D \setminus M$ and M / D perform inverse scaling. They are equivalent to solving a diagonal (or rectangular diagonal) in a least-squares minimum-norm sense. In exact arithmetic, this is equivalent to multiplying by a pseudoinverse. The pseudoinverse of a rectangular diagonal matrix is again a rectangular diagonal matrix with swapped dimensions, where each nonzero diagonal element is replaced by its reciprocal. The matrix division algorithms do, in fact, use division rather than multiplication by reciprocals for better numerical accuracy; otherwise, they honor the above definition. Note that a diagonal matrix is never truncated due to ill-conditioning; otherwise, it would not be of much use for scaling. This is typically consistent with linear algebra needs. A full matrix that only happens to be diagonal (and is thus not a special object) is of course treated normally.

Multiplication and division by diagonal matrices work efficiently also when combined with sparse matrices, i.e., $D*S$, where D is a diagonal matrix and S is a sparse matrix scales

the rows of the sparse matrix and returns a sparse matrix. The expressions $S*D$, $D\backslash S$, S/D work analogically.

If $D1$ and $D2$ are both diagonal matrices, then the expressions

```
D1 + D2
D1 - D2
D1 * D2
D1 / D2
D1 \ D2
```

again produce diagonal matrices, provided that normal dimension matching rules are obeyed. The relations used are same as described above.

Also, a diagonal matrix D can be multiplied or divided by a scalar, or raised to a scalar power if it is square, producing diagonal matrix result in all cases.

A diagonal matrix can also be transposed or conjugate-transposed, giving the expected result. Extracting a leading submatrix of a diagonal matrix, i.e., $D(1:m, 1:n)$, will produce a diagonal matrix, other indexing expressions will implicitly convert to full matrix.

Adding a diagonal matrix to a full matrix only operates on the diagonal elements. Thus,

```
A = A + eps * eye (n)
```

is an efficient method of augmenting the diagonal of a matrix. Subtraction works analogically.

When involved in expressions with other element-by-element operators, $.*$, $./$, $.\backslash$ or $.^$, an implicit conversion to full matrix will take place. This is not always strictly necessary but chosen to facilitate better consistency with MATLAB.

21.2.2 Expressions Involving Permutation Matrices

If P is a permutation matrix and M a matrix, the expression $P*M$ will permute the rows of M . Similarly, $M*P$ will yield a column permutation. Matrix division $P\backslash M$ and M/P can be used to do inverse permutation.

The previously described syntax for creating permutation matrices can actually help an user to understand the connection between a permutation matrix and a permuting vector. Namely, the following holds, where $I = eye (n)$ is an identity matrix:

$$I(p,:) * M = (I*M)(p,:) = M(p,:)$$

Similarly,

$$M * I(:,p) = (M*I)(:,p) = M(:,p)$$

The expressions $I(p,:)$ and $I(:,p)$ are permutation matrices.

A permutation matrix can be transposed (or conjugate-transposed, which is the same, because a permutation matrix is never complex), inverting the permutation, or equivalently, turning a row-permutation matrix into a column-permutation one. For permutation matrices, transpose is equivalent to inversion, thus $P\backslash M$ is equivalent to $P'*M$. Transpose of a permutation matrix (or inverse) is a constant-time operation, flipping only a flag internally, and thus the choice between the two above equivalent expressions for inverse permuting is completely up to the user's taste.

Multiplication and division by permutation matrices works efficiently also when combined with sparse matrices, i.e., $P*S$, where P is a permutation matrix and S is a sparse

matrix permutes the rows of the sparse matrix and returns a sparse matrix. The expressions $S*P$, $P\backslash S$, S/P work analogically.

Two permutation matrices can be multiplied or divided (if their sizes match), performing a composition of permutations. Also a permutation matrix can be indexed by a permutation vector (or two vectors), giving again a permutation matrix. Any other operations do not generally yield a permutation matrix and will thus trigger the implicit conversion.

21.3 Functions That Are Aware of These Matrices

This section lists the built-in functions that are aware of diagonal and permutation matrices on input, or can return them as output. Passed to other functions, these matrices will in general trigger an implicit conversion. (Of course, user-defined dynamically linked functions may also work with diagonal or permutation matrices).

21.3.1 Diagonal Matrix Functions

inv and *pinv* can be applied to a diagonal matrix, yielding again a diagonal matrix. *det* will use an efficient straightforward calculation when given a diagonal matrix, as well as *cond*. The following mapper functions can be applied to a diagonal matrix without converting it to a full one: *abs*, *real*, *imag*, *conj*, *sqrt*. A diagonal matrix can also be returned from the *balance* and *svd* functions. The *sparse* function will convert a diagonal matrix efficiently to a sparse matrix.

21.3.2 Permutation Matrix Functions

inv and *pinv* will invert a permutation matrix, preserving its specialness. *det* can be applied to a permutation matrix, efficiently calculating the sign of the permutation (which is equal to the determinant).

A permutation matrix can also be returned from the built-in functions *lu* and *qr*, if a pivoted factorization is requested.

The *sparse* function will convert a permutation matrix efficiently to a sparse matrix. The *find* function will also work efficiently with a permutation matrix, making it possible to conveniently obtain the permutation indices.

21.4 Examples of Usage

The following can be used to solve a linear system $A*x = b$ using the pivoted LU factorization:

```
[L, U, P] = lu (A); ## now L*U = P*A
x = U \ (L \ P) * b;
```

This is one way to normalize columns of a matrix X to unit norm:

```
s = norm (X, "columns");
X ./= diag (s);
```

The same can also be accomplished with broadcasting (see [Section 19.2 \[Broadcasting\]](#), page 585):

```
s = norm (X, "columns");
X ./= s;
```

The following expression is a way to efficiently calculate the sign of a permutation, given by a permutation vector p . It will also work in earlier versions of Octave, but slowly.

```
det (eye (length (p))(p, :))
```

Finally, here's how to solve a linear system $A*x = b$ with Tikhonov regularization (ridge regression) using SVD (a skeleton only):

```
m = rows (A); n = columns (A);
[U, S, V] = svd (A);
## determine the regularization factor alpha
## alpha = ...
## transform to orthogonal basis
b = U'*b;
## Use the standard formula, replacing A with S.
## S is diagonal, so the following will be very fast and accurate.
x = (S'*S + alpha^2 * eye (n)) \ (S' * b);
## transform to solution basis
x = V*x;
```

21.5 Differences in Treatment of Zero Elements

Making diagonal and permutation matrices special matrix objects in their own right and the consequent usage of smarter algorithms for certain operations implies, as a side effect, small differences in treating zeros. The contents of this section apply also to sparse matrices, discussed in the following chapter. (see [Chapter 22 \[Sparse Matrices\], page 617](#))

The IEEE floating point standard defines the result of the expressions $0*Inf$ and $0*NaN$ as NaN . This is widely agreed to be a good compromise. Numerical software dealing with structured and sparse matrices (including Octave) however, almost always makes a distinction between a "numerical zero" and an "assumed zero". A "numerical zero" is a zero value occurring in a place where any floating-point value could occur. It is normally stored somewhere in memory as an explicit value. An "assumed zero", on the contrary, is a zero matrix element implied by the matrix structure (diagonal, triangular) or a sparsity pattern; its value is usually not stored explicitly anywhere, but is implied by the underlying data structure.

The primary distinction is that an assumed zero, when multiplied by any number, or divided by any nonzero number, yields **always** a zero, even when, e.g., multiplied by **Inf** or divided by **NaN**. The reason for this behavior is that the numerical multiplication is not actually performed anywhere by the underlying algorithm; the result is just assumed to be zero. Equivalently, one can say that the part of the computation involving assumed zeros is performed symbolically, not numerically.

This behavior not only facilitates the most straightforward and efficient implementation of algorithms, but also preserves certain useful invariants, like:

- scalar * diagonal matrix is a diagonal matrix
- sparse matrix / scalar preserves the sparsity pattern
- permutation matrix * matrix is equivalent to permuting rows

all of these natural mathematical truths would be invalidated by treating assumed zeros as numerical ones.

Note that MATLAB does not strictly follow this principle and converts assumed zeros to numerical zeros in certain cases, while not doing so in other cases. As of today, there are no intentions to mimic such behavior in Octave.

Examples of effects of assumed zeros vs. numerical zeros:

```
Inf * eye (3)
⇒
Inf      0      0
 0      Inf      0
 0      0      Inf

Inf * speye (3)
⇒
Compressed Column Sparse (rows = 3, cols = 3, nnz = 3 [33%])

(1, 1) -> Inf
(2, 2) -> Inf
(3, 3) -> Inf

Inf * full (eye (3))
⇒
Inf      NaN     NaN
NaN      Inf     NaN
NaN      NaN     Inf

diag (1:3) * [NaN; 1; 1]
⇒
NaN
2
3

sparse (1:3,1:3,1:3) * [NaN; 1; 1]
⇒
NaN
2
3
[1,0,0;0,2,0;0,0,3] * [NaN; 1; 1]
⇒
NaN
NaN
NaN
```

22 Sparse Matrices

22.1 Creation and Manipulation of Sparse Matrices

The size of mathematical problems that can be treated at any particular time is generally limited by the available computing resources. Both, the speed of the computer and its available memory place limitation on the problem size.

There are many classes of mathematical problems which give rise to matrices, where a large number of the elements are zero. In this case it makes sense to have a special matrix type to handle this class of problems where only the nonzero elements of the matrix are stored. Not only does this reduce the amount of memory to store the matrix, but it also means that operations on this type of matrix can take advantage of the a priori knowledge of the positions of the nonzero elements to accelerate their calculations.

A matrix type that stores only the nonzero elements is generally called sparse. It is the purpose of this document to discuss the basics of the storage and creation of sparse matrices and the fundamental operations on them.

22.1.1 Storage of Sparse Matrices

It is not strictly speaking necessary for the user to understand how sparse matrices are stored. However, such an understanding will help to get an understanding of the size of sparse matrices. Understanding the storage technique is also necessary for those users wishing to create their own oct-files.

There are many different means of storing sparse matrix data. What all of the methods have in common is that they attempt to reduce the complexity and storage given a priori knowledge of the particular class of problems that will be solved. A good summary of the available techniques for storing sparse matrix is given by Saad¹. With full matrices, knowledge of the point of an element of the matrix within the matrix is implied by its position in the computers memory. However, this is not the case for sparse matrices, and so the positions of the nonzero elements of the matrix must equally be stored.

An obvious way to do this is by storing the elements of the matrix as triplets, with two elements being their position in the array (rows and column) and the third being the data itself. This is conceptually easy to grasp, but requires more storage than is strictly needed.

The storage technique used within Octave is the compressed column format. It is similar to the Yale format.² In this format the position of each element in a row and the data are stored as previously. However, if we assume that all elements in the same column are stored adjacent in the computers memory, then we only need to store information on the number of nonzero elements in each column, rather than their positions. Thus assuming that the matrix has more nonzero elements than there are columns in the matrix, we win in terms of the amount of memory used.

In fact, the column index contains one more element than the number of columns, with the first element always being zero. The advantage of this is a simplification in the code, in

¹ Y. Saad "SPARSKIT: A basic toolkit for sparse matrix computation", 1994, <https://www-users.cs.umn.edu/~saad/software/SPARSKIT/paper.ps>

² https://en.wikipedia.org/wiki/Sparse_matrix#Yale_format

that there is no special case for the first or last columns. A short example, demonstrating this in C is.

```
for (j = 0; j < nc; j++)
    for (i = cidx(j); i < cidx(j+1); i++)
        printf ("nonzero element (%i,%i) is %d\n",
            ridx(i), j, data(i));
```

A clear understanding might be had by considering an example of how the above applies to an example matrix. Consider the matrix

1	2	0	0
0	0	0	3
0	0	0	4

The nonzero elements of this matrix are

(1, 1)	⇒ 1
(1, 2)	⇒ 2
(2, 4)	⇒ 3
(3, 4)	⇒ 4

This will be stored as three vectors *cidx*, *ridx* and *data*, representing the column indexing, row indexing and data respectively. The contents of these three vectors for the above matrix will be

```
cidx = [0, 1, 2, 2, 4]
ridx = [0, 0, 1, 2]
data = [1, 2, 3, 4]
```

Note that this is the representation of these elements with the first row and column assumed to start at zero, while in Octave itself the row and column indexing starts at one. Thus the number of elements in the *i*-th column is given by *cidx* (*i* + 1) - *cidx* (*i*).

Although Octave uses a compressed column format, it should be noted that compressed row formats are equally possible. However, in the context of mixed operations between mixed sparse and dense matrices, it makes sense that the elements of the sparse matrices are in the same order as the dense matrices. Octave stores dense matrices in column major ordering, and so sparse matrices are equally stored in this manner.

A further constraint on the sparse matrix storage used by Octave is that all elements in the rows are stored in increasing order of their row index, which makes certain operations faster. However, it imposes the need to sort the elements on the creation of sparse matrices. Having disordered elements is potentially an advantage in that it makes operations such as concatenating two sparse matrices together easier and faster, however it adds complexity and speed problems elsewhere.

22.1.2 Creating Sparse Matrices

There are several means to create sparse matrix.

Returned from a function

There are many functions that directly return sparse matrices. These include *speye*, *sprand*, *diag*, etc.

Constructed from matrices or vectors

The function `sparse` allows a sparse matrix to be constructed from three vectors representing the row, column and data. Alternatively, the function `spconvert` uses a three column matrix format to allow easy importation of data from elsewhere.

Created and then filled

The function `sparse` or `spalloc` can be used to create an empty matrix that is then filled by the user

From a user binary program

The user can directly create the sparse matrix within an oct-file.

There are several basic functions to return specific sparse matrices. For example the sparse identity matrix, is a matrix that is often needed. It therefore has its own function to create it as `speye (n)` or `speye (r, c)`, which creates an n -by- n or r -by- c sparse identity matrix.

Another typical sparse matrix that is often needed is a random distribution of random elements. The functions `sprand` and `sprandn` perform this for uniform and normal random distributions of elements. They have exactly the same calling convention, where `sprand (r, c, d)`, creates an r -by- c sparse matrix with a density of filled elements of d .

Other functions of interest that directly create sparse matrices, are `diag` or its generalization `spdiags`, that can take the definition of the diagonals of the matrix and create the sparse matrix that corresponds to this. For example,

```
s = diag (sparse (randn (1,n)), -1);
```

creates a sparse $(n+1)$ -by- $(n+1)$ sparse matrix with a single diagonal defined.

```
B = spdiags (A)
[B, d] = spdiags (A)
B = spdiags (A, d)
A = spdiags (v, d, A)
A = spdiags (v, d, m, n)
```

A generalization of the function `diag`.

Called with a single input argument, the nonzero diagonals d of A are extracted.

With two arguments the diagonals to extract are given by the vector d .

The other two forms of `spdiags` modify the input matrix by replacing the diagonals. They use the columns of v to replace the diagonals represented by the vector d . If the sparse matrix A is defined then the diagonals of this matrix are replaced. Otherwise a matrix of m by n is created with the diagonals given by the columns of v .

Negative values of d represent diagonals below the main diagonal, and positive values of d diagonals above the main diagonal.

For example:

```
spdiags (reshape (1:12, 4, 3), [-1 0 1], 5, 4)
⇒ 5 10 0 0
    1 6 11 0
    0 2 7 12
    0 0 3 8
    0 0 0 4
```

See also: [\[diag\]](#), page 489.

s = speye (m, n)

s = speye (m)

s = speye (sz)

Return a sparse identity matrix of size $m \times n$.

The implementation is significantly more efficient than **sparse (eye (m))** as the full matrix is not constructed.

Called with a single argument a square matrix of size m -by- m is created. If called with a single vector argument sz , this argument is taken to be the size of the matrix to create.

See also: [\[sparse\]](#), page 622, [\[spdiags\]](#), page 619, [\[eye\]](#), page 489.

r = spones (S)

Replace the nonzero entries of S with ones.

This creates a sparse matrix with the same structure as S .

See also: [\[sparse\]](#), page 622, [\[sprand\]](#), page 620, [\[sprandn\]](#), page 620, [\[sprandsym\]](#), page 621, [\[spfun\]](#), page 590, [\[spy\]](#), page 626.

sprand (m, n, d)

sprand (m, n, d, rc)

sprand (s)

Generate a sparse matrix with uniformly distributed random values.

The size of the matrix is $m \times n$ with a density of values d . d must be between 0 and 1. Values will be uniformly distributed on the interval $(0, 1)$.

If called with a single matrix argument, a sparse matrix is generated with random values wherever the matrix s is nonzero.

If called with a scalar fourth argument rc , a random sparse matrix with reciprocal condition number rc is generated. If rc is a vector, then it specifies the first singular values of the generated matrix (`length (rc) <= min (m, n)`).

See also: [\[sprandn\]](#), page 620, [\[sprandsym\]](#), page 621, [\[rand\]](#), page 494.

sprandn (m, n, d)

sprandn (m, n, d, rc)

sprandn (s)

Generate a sparse matrix with normally distributed random values.

The size of the matrix is $m \times n$ with a density of values d . d must be between 0 and 1. Values will be normally distributed with a mean of 0 and a variance of 1.

If called with a single matrix argument, a sparse matrix is generated with random values wherever the matrix s is nonzero.

If called with a scalar fourth argument rc , a random sparse matrix with reciprocal condition number rc is generated. If rc is a vector, then it specifies the first singular values of the generated matrix (`length (rc) <= min (m, n)`).

See also: [\[sprand\]](#), page 620, [\[sprandsym\]](#), page 621, [\[randn\]](#), page 496.

```
sprandsym (n, d)
sprandsym (s)
```

Generate a symmetric random sparse matrix.

The size of the matrix will be $n \times n$, with a density of values given by d . d must be between 0 and 1 inclusive. Values will be normally distributed with a mean of zero and a variance of 1.

If called with a single matrix argument, a random sparse matrix is generated wherever the matrix s is nonzero in its lower triangular part.

See also: [sprand], page 620, [sprandn], page 620, [spones], page 620, [sparse], page 622.

The recommended way for the user to create a sparse matrix, is to create two vectors containing the row and column index of the data and a third vector of the same size containing the data to be stored. For example,

```
ri = ci = d = [];
for j = 1:c
    ri = [ri; randperm(r,n)'];
    ci = [ci; j*ones(n,1)];
    d = [d; rand(n,1)];
endfor
s = sparse (ri, ci, d, r, c);
```

creates an r -by- c sparse matrix with a random distribution of n ($< r$) elements per column. The elements of the vectors do not need to be sorted in any particular order as Octave will sort them prior to storing the data. However, pre-sorting the data will make the creation of the sparse matrix faster.

The function *spconvert* takes a three or four column real matrix. The first two columns represent the row and column index respectively and the third and four columns, the real and imaginary parts of the sparse matrix. The matrix can contain zero elements and the elements can be sorted in any order. Adding zero elements is a convenient way to define the size of the sparse matrix. For example:

```
s = spconvert ([1 2 3 4; 1 3 4 4; 1 2 3 0]');
⇒ Compressed Column Sparse (rows=4, cols=4, nnz=3)
(1 , 1) -> 1
(2 , 3) -> 2
(3 , 4) -> 3
```

An example of creating and filling a matrix might be

```
k = 5;
nz = r * k;
s = spalloc (r, c, nz)
for j = 1:c
    idx = randperm (r);
    s (:, j) = [zeros(r - k, 1); ...
                 rand(k, 1)] (idx);
endfor
```

It should be noted, that due to the way that the Octave assignment functions are written that the assignment will reallocate the memory used by the sparse matrix at each iteration of the above loop. Therefore the `spalloc` function ignores the `nz` argument and does not pre-assign the memory for the matrix. Therefore, it is vitally important that code using the above structure should be vectorized as much as possible to minimize the number of assignments and reduce the number of memory allocations.

`FM = full (SM)`

Return a full storage matrix from a sparse, diagonal, or permutation matrix, or a range.

See also: [\[sparse\]](#), page 622, [\[issparse\]](#), page 624.

`s = spalloc (m, n, nz)`

Create an m -by- n sparse matrix with pre-allocated space for at most nz nonzero elements.

This is useful for building a matrix incrementally by a sequence of indexed assignments. Subsequent indexed assignments after `spalloc` will reuse the pre-allocated memory, provided they are of one of the simple forms

- $s(I:J) = x$
- $s(:, I:J) = x$
- $s(K:L, I:J) = x$

and that the following conditions are met:

- the assignment does not decrease $\text{nnz}(S)$.
- after the assignment, $\text{nnz}(S)$ does not exceed nz .
- no index is out of bounds.

Partial movement of data may still occur, but in general the assignment will be more memory and time efficient under these circumstances. In particular, it is possible to efficiently build a pre-allocated sparse matrix from a contiguous block of columns.

The amount of pre-allocated memory for a given matrix may be queried using the function `nzmax`.

See also: [\[nzmax\]](#), page 624, [\[sparse\]](#), page 622.

```
s = sparse (a)
s = sparse (i, j, sv, m, n)
s = sparse (i, j, sv)
s = sparse (m, n)
s = sparse (i, j, s, m, n, "unique")
s = sparse (i, j, sv, m, n, nzmax)
```

Create a sparse matrix from a full matrix, or row, column, value triplets.

If `a` is a full matrix, convert it to a sparse matrix representation, removing all zero values in the process.

Given the integer index vectors `i` and `j`, and a 1-by-`nnz` vector of real or complex values `sv`, construct the sparse matrix $S(i(k), j(k)) = sv(k)$ with overall dimensions `m` and `n`. If any of `sv`, `i` or `j` are scalars, they are expanded to have a common size.

If m or n are not specified their values are derived from the maximum index in the vectors i and j as given by $m = \max(i)$, $n = \max(j)$.

Note: if multiple values are specified with the same i, j indices, the corresponding value in s will be the sum of the values at the repeated location. See `accumarray` for an example of how to produce different behavior, such as taking the minimum instead.

If the option "unique" is given, and more than one value is specified at the same i, j indices, then the last specified value will be used.

`sparse(m, n)` will create an empty mxn sparse matrix and is equivalent to `sparse([], [], [], m, n)`

The argument `nzmax` is ignored but accepted for compatibility with MATLAB.

Example 1 (sum at repeated indices):

```
i = [1 1 2]; j = [1 1 2]; sv = [3 4 5];
sparse (i, j, sv, 3, 4)
⇒
Compressed Column Sparse (rows = 3, cols = 4, nnz = 2 [17%])

(1, 1) -> 7
(2, 2) -> 5
```

Example 2 ("unique" option):

```
i = [1 1 2]; j = [1 1 2]; sv = [3 4 5];
sparse (i, j, sv, 3, 4, "unique")
⇒
Compressed Column Sparse (rows = 3, cols = 4, nnz = 2 [17%])

(1, 1) -> 4
(2, 2) -> 5
```

See also: [\[full\]](#), page 622, [\[accumarray\]](#), page 593, [\[spalloc\]](#), page 622, [\[spdiags\]](#), page 619, [\[speye\]](#), page 620, [\[spones\]](#), page 620, [\[sprand\]](#), page 620, [\[sprandn\]](#), page 620, [\[sprandsym\]](#), page 621, [\[spconvert\]](#), page 623, [\[spfun\]](#), page 590.

x = spconvert (m)

Convert a simple sparse matrix format easily generated by other programs into Octave's internal sparse format.

The input m is either a 3 or 4 column real matrix, containing the row, column, real, and imaginary parts of the elements of the sparse matrix. An element with a zero real and imaginary part can be used to force a particular matrix size.

See also: [\[sparse\]](#), page 622.

The above problem of memory reallocation can be avoided in oct-files. However, the construction of a sparse matrix from an oct-file is more complex than can be discussed here. See [Appendix A \[External Code Interface\]](#), page 945, for a full description of the techniques involved.

22.1.3 Finding Information about Sparse Matrices

There are a number of functions that allow information concerning sparse matrices to be obtained. The most basic of these is `issparse` that identifies whether a particular Octave object is in fact a sparse matrix.

Another very basic function is `nnz` that returns the number of nonzero entries there are in a sparse matrix, while the function `nzmax` returns the amount of storage allocated to the sparse matrix. Note that Octave tends to crop unused memory at the first opportunity for sparse objects. There are some cases of user created sparse objects where the value returned by `nzmax` will not be the same as `nnz`, but in general they will give the same result. The function `spstats` returns some basic statistics on the columns of a sparse matrix including the number of elements, the mean and the variance of each column.

`issparse (x)`

Return true if `x` is a sparse matrix.

See also: [\[ismatrix\]](#), page 63.

`n = nnz (a)`

Return the number of nonzero elements in `a`.

See also: [\[nzmax\]](#), page 624, [\[nonzeros\]](#), page 624, [\[find\]](#), page 477.

`v = nonzeros (A)`

Return a column vector of the nonzero values of the matrix `A`.

See also: [\[find\]](#), page 477, [\[nnz\]](#), page 624.

`n = nzmax (SM)`

Return the amount of storage allocated to the sparse matrix `SM`.

Note that Octave tends to crop unused memory at the first opportunity for sparse objects. Thus, in general the value of `nzmax` will be the same as `nnz` except for some cases of user-created sparse objects.

See also: [\[nnz\]](#), page 624, [\[salloc\]](#), page 622, [\[sparse\]](#), page 622.

`[count, mean, var] = spstats (S)`

`[count, mean, var] = spstats (S, j)`

Return the stats for the nonzero elements of the sparse matrix `S`.

`count` is the number of nonzeros in each column, `mean` is the mean of the nonzeros in each column, and `var` is the variance of the nonzeros in each column.

Called with two input arguments, if `S` is the data and `j` is the bin number for the data, compute the stats for each bin. In this case, bins can contain data values of zero, whereas with `spstats (S)` the zeros may disappear.

When solving linear equations involving sparse matrices Octave determines the means to solve the equation based on the type of the matrix (see [Section 22.2 \[Sparse Linear Algebra\]](#), page 638). Octave probes the matrix type when the `div (/)` or `ldiv (\)` operator is first used with the matrix and then caches the type. However the `matrix_type` function can be used

to determine the type of the sparse matrix prior to use of the `div` or `ldiv` operators. For example,

```
a = tril (sprandn (1024, 1024, 0.02), -1) ...
+ speye (1024);
matrix_type (a);
ans = Lower
```

shows that Octave correctly determines the matrix type for lower triangular matrices. `matrix_type` can also be used to force the type of a matrix to be a particular type. For example:

```
a = matrix_type (tril (sprandn (1024, ...
1024, 0.02), -1) + speye (1024), "Lower");
```

This allows the cost of determining the matrix type to be avoided. However, incorrectly defining the matrix type will result in incorrect results from solutions of linear equations, and so it is entirely the responsibility of the user to correctly identify the matrix type.

There are several graphical means of finding out information about sparse matrices. The first is the `spy` command, which displays the structure of the nonzero elements of the matrix. See [Figure 22.1](#), for an example of the use of `spy`. More advanced graphical information can be obtained with the `treeplot`, `etreeplot` and `gplot` commands.





Figure 22.1: Structure of simple sparse matrix.

One use of sparse matrices is in graph theory, where the interconnections between nodes are represented as an adjacency matrix. That is, if the i -th node in a graph is connected to the j -th node. Then the ij -th node (and in the case of undirected graphs the ji -th node) of the sparse adjacency matrix is nonzero. If each node is then associated with a set of coordinates, then the `gplot` command can be used to graphically display the interconnections between nodes.

As a trivial example of the use of `gplot` consider the example,

```
A = sparse ([2,6,1,3,2,4,3,5,4,6,1,5],  
           [1,1,2,2,3,3,4,4,5,5,6,6],1,6,6);  
xy = [0,4,8,6,4,2;5,0,5,7,5,7]';  
gplot (A,xy)
```

which creates an adjacency matrix A where node 1 is connected to nodes 2 and 6, node 2 with nodes 1 and 3, etc. The coordinates of the nodes are given in the n-by-2 matrix xy . See [Figure 22.2](#).



[Figure 22.2](#): Simple use of the `gplot` command.

The dependencies between the nodes of a Cholesky factorization can be calculated in linear time without explicitly needing to calculate the Cholesky factorization by the `etree` command. This command returns the elimination tree of the matrix and can be displayed graphically by the command `treeplot (etree (A))` if A is symmetric or `treeplot (etree (A+A'))` otherwise.

```
spy (x)  
spy (... , markersize)  
spy (... , line_spec)
```

Plot the sparsity pattern of the sparse matrix x .

If the argument *markersize* is given as a scalar value, it is used to determine the point size in the plot.

If the string *line_spec* is given it is passed to `plot` and determines the appearance of the plot.

See also: [\[plot\]](#), page 296, [\[gplot\]](#), page 627.

```
p = etree (S)  
p = etree (S, typ)  
[p, q] = etree (S, typ)
```

Return the elimination tree for the matrix S .

By default S is assumed to be symmetric and the symmetric elimination tree is returned. The argument typ controls whether a symmetric or column elimination tree is returned. Valid values of typ are "sym" or "col", for symmetric or column elimination tree respectively.

Called with a second argument, `etree` also returns the postorder permutations on the tree.

`etreeplot (A)`

`etreeplot (A, node_style, edge_style)`

Plot the elimination tree of the matrix A or $A+A'$ if A is not symmetric.

The optional parameters `node_style` and `edge_style` define the output style.

See also: [\[treeplot\]](#), page 627, [\[gplot\]](#), page 627.

`gplot (A, xy)`

`gplot (A, xy, line_style)`

`[x, y] = gplot (A, xy)`

Plot a graph defined by A and xy in the graph theory sense.

A is the adjacency matrix of the array to be plotted and xy is an n -by-2 matrix containing the coordinates of the nodes of the graph.

The optional parameter `line_style` defines the output style for the plot. Called with no output arguments the graph is plotted directly. Otherwise, return the coordinates of the plot in x and y .

See also: [\[treeplot\]](#), page 627, [\[etreeplot\]](#), page 627, [\[spy\]](#), page 626.

`treeplot (tree)`

`treeplot (tree, node_style, edge_style)`

Produce a graph of tree or forest.

The first argument is vector of predecessors.

The optional parameters `node_style` and `edge_style` define the output plot style.

The complexity of the algorithm is $O(n)$ in terms of time and memory requirements.

See also: [\[etreeplot\]](#), page 627, [\[gplot\]](#), page 627.

`treelayout (tree)`

`treelayout (tree, permutation)`

treelayout lays out a tree or a forest.

The first argument `tree` is a vector of predecessors.

The parameter `permutation` is an optional postorder permutation.

The complexity of the algorithm is $O(n)$ in terms of time and memory requirements.

See also: [\[etreeplot\]](#), page 627, [\[gplot\]](#), page 627, [\[treeplot\]](#), page 627.

22.1.4 Basic Operators and Functions on Sparse Matrices

22.1.4.1 Sparse Functions

Many Octave functions have been overloaded to work with either sparse or full matrices. There is no difference in calling convention when using an overloaded function with a sparse matrix, however, there is also no access to potentially sparse-specific features. At any time the sparse matrix specific version of a function can be used by explicitly calling its function name.

The table below lists all of the sparse functions of Octave. Note that the names of the specific sparse forms of the functions are typically the same as the general versions with a *sp* prefix. In the table below, and in the rest of this article, the specific sparse versions of functions are used.

Generate sparse matrices:

spalloc, spdiags, speye, sprand, sprandn, sprandsym

Sparse matrix conversion:

full, sparse, spconvert

Manipulate sparse matrices

issparse, nnz, nonzeros, nzmax, spfun, spones, spy

Graph Theory:

etree, etreeplot, gplot, treeplot

Sparse matrix reordering:

amd, ccolamd, colamd, colperm, csymamd, dmperm, symamd, randperm, symrcm

Linear algebra:

conde, eigs, matrix_type, normest, normest1, sprank, spaugment, svds

Iterative techniques:

ichol, ilu, pcg, pcr

Miscellaneous:

spparms, symbfact, spstats

In addition all of the standard Octave mapper functions (i.e., basic math functions that take a single argument) such as *abs*, etc. can accept sparse matrices. The reader is referred to the documentation supplied with these functions within Octave itself for further details.

22.1.4.2 Return Types of Operators and Functions

The two basic reasons to use sparse matrices are to reduce the memory usage and to not have to do calculations on zero elements. The two are closely related in that the computation time on a sparse matrix operator or function is roughly linear with the number of nonzero elements.

Therefore, there is a certain density of nonzero elements of a matrix where it no longer makes sense to store it as a sparse matrix, but rather as a full matrix. For this reason operators and functions that have a high probability of returning a full matrix will always return one. For example adding a scalar constant to a sparse matrix will almost always make it a full matrix, and so the example,

```

speye (3) + 0
⇒ 1 0 0
 0 1 0
 0 0 1

```

returns a full matrix as can be seen.

Additionally, if `sparse_auto_mutate` is true, all sparse functions test the amount of memory occupied by the sparse matrix to see if the amount of storage used is larger than the amount used by the full equivalent. Therefore `speye (2) * 1` will return a full matrix as the memory used is smaller for the full version than the sparse version.

As all of the mixed operators and functions between full and sparse matrices exist, in general this does not cause any problems. However, one area where it does cause a problem is where a sparse matrix is promoted to a full matrix, where subsequent operations would resparsify the matrix. Such cases are rare, but can be artificially created, for example `(fliplr (speye (3)) + speye (3)) - speye (3)` gives a full matrix when it should give a sparse one. In general, where such cases occur, they impose only a small memory penalty.

There is however one known case where this behavior of Octave's sparse matrices will cause a problem. That is in the handling of the `diag` function. Whether `diag` returns a sparse or full matrix depending on the type of its input arguments. So

```
a = diag (sparse ([1,2,3]), -1);
```

should return a sparse matrix. To ensure this actually happens, the `sparse` function, and other functions based on it like `speye`, always returns a sparse matrix, even if the memory used will be larger than its full representation.

```

val = sparse_auto_mutate ()
old_val = sparse_auto_mutate (new_val)
sparse_auto_mutate (new_val, "local")

```

Query or set the internal variable that controls whether Octave will automatically mutate sparse matrices to full matrices to save memory.

For example:

```

s = speye (3);
sparse_auto_mutate (false);
s(:, 1) = 1;
typeinfo (s)
⇒ sparse matrix
sparse_auto_mutate (true);
s(1, :) = 1;
typeinfo (s)
⇒ matrix

```

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

Note that the `sparse_auto_mutate` option is incompatible with MATLAB, and so it is off by default.

22.1.4.3 Mathematical Considerations

The attempt has been made to make sparse matrices behave in exactly the same manner as their full counterparts. However, there are certain differences and especially differences with other products sparse implementations.

First, the "./" and ".^" operators must be used with care. Consider what the examples

```
s = speye (4);
a1 = s .^ 2;
a2 = s .^ s;
a3 = s .^ -2;
a4 = s ./ 2;
a5 = 2 ./ s;
a6 = s ./ s;
```

will give. The first example of *s* raised to the power of 2 causes no problems. However *s* raised element-wise to itself involves a large number of terms $0.^0$ which is 1. There $s.^s$ is a full matrix.

Likewise $s.^{-2}$ involves terms like $0.^{-2}$ which is infinity, and so $s.^{-2}$ is equally a full matrix.

For the "./" operator $s./2$ has no problems, but $2./s$ involves a large number of infinity terms as well and is equally a full matrix. The case of $s./s$ involves terms like $0./0$ which is a NaN and so this is equally a full matrix with the zero elements of *s* filled with NaN values.

The above behavior is consistent with full matrices, but is not consistent with sparse implementations in other products.

A particular problem of sparse matrices comes about due to the fact that as the zeros are not stored, the sign-bit of these zeros is equally not stored. In certain cases the sign-bit of zero is important. For example:

```
a = 0 ./ [-1, 1; 1, -1];
b = 1 ./ a
⇒ -Inf      Inf
      Inf      -Inf
c = 1 ./ sparse (a)
⇒  Inf      Inf
      Inf      Inf
```

To correct this behavior would mean that zero elements with a negative sign-bit would need to be stored in the matrix to ensure that their sign-bit was respected. This is not done at this time, for reasons of efficiency, and so the user is warned that calculations where the sign-bit of zero is important must not be done using sparse matrices.

In general any function or operator used on a sparse matrix will result in a sparse matrix with the same or a larger number of nonzero elements than the original matrix. This is particularly true for the important case of sparse matrix factorizations. The usual way to address this is to reorder the matrix, such that its factorization is sparser than the factorization of the original matrix. That is the factorization of $L * U = P * S * Q$ has sparser terms *L* and *U* than the equivalent factorization $L * U = S$.

Several functions are available to reorder depending on the type of the matrix to be factorized. If the matrix is symmetric positive-definite, then *symamd* or *csymamd* should be used. Otherwise *amd*, *colamd* or *ccolamd* should be used. For completeness the reordering functions *colperm* and *randperm* are also available.

See [Figure 22.3](#), for an example of the structure of a simple positive definite matrix.




Figure 22.3: Structure of simple sparse matrix.

The standard Cholesky factorization of this matrix can be obtained by the same command that would be used for a full matrix. This can be visualized with the command `r = chol (A); spy (r);`. See [Figure 22.4](#). The original matrix had 598 nonzero terms, while this Cholesky factorization has 10200, with only half of the symmetric matrix being stored. This is a significant level of fill in, and although not an issue for such a small test case, can represent a large overhead in working with other sparse matrices.

The appropriate sparsity preserving permutation of the original matrix is given by *symamd* and the factorization using this reordering can be visualized using the command `q = symamd (A); r = chol (A(q,q)); spy (r)`. This gives 399 nonzero terms which is a significant improvement.

The Cholesky factorization itself can be used to determine the appropriate sparsity preserving reordering of the matrix during the factorization. In that case this might be obtained with three return arguments as `[r, p, q] = chol (A); spy (r)`.




Figure 22.4: Structure of the unpermuted Cholesky factorization of the above matrix.




Figure 22.5: Structure of the permuted Cholesky factorization of the above matrix.

In the case of an asymmetric matrix, the appropriate sparsity preserving permutation is `colamd` and the factorization using this reordering can be visualized using the command `q = colamd (A); [l, u, p] = lu (A(:,q)); spy (l+u)`.

Finally, Octave implicitly reorders the matrix when using the `div (/)` and `ldiv (\)` operators, and so no the user does not need to explicitly reorder the matrix to maximize performance.

```
p = amd (S)
p = amd (S, opts)
```

Return the approximate minimum degree permutation of a matrix.

This is a permutation such that the Cholesky factorization of $S(p, p)$ tends to be sparser than the Cholesky factorization of S itself. `amd` is typically faster than `symamd` but serves a similar purpose.

The optional parameter `opts` is a structure that controls the behavior of `amd`. The fields of the structure are

`opts.dense` Determines what `amd` considers to be a dense row or column of the input matrix. Rows or columns with more than `max (16, (dense * sqrt (n)))` entries, where n is the order of the matrix S , are ignored by `amd` during the calculation of the permutation. The value of `dense` must be a positive scalar and the default value is 10.0

`opts.aggressive`

If this value is a nonzero scalar, then `amd` performs aggressive absorption. The default is not to perform aggressive absorption.

The author of the code itself is Timothy A. Davis (see <http://faculty.cse.tamu.edu/davis/suitesparse.html>).

See also: `[symamd]`, page 636, `[colamd]`, page 634.

```
p = ccolamd (S)
p = ccolamd (S, knobs)
p = ccolamd (S, knobs, cmember)
[p, stats] = ccolamd (...)
```

Constrained column approximate minimum degree permutation.

`p = ccolamd (S)` returns the column approximate minimum degree permutation vector for the sparse matrix S . For a non-symmetric matrix S , $S(:, p)$ tends to have sparser LU factors than S . `chol (S(:, p)' * S(:, p))` also tends to be sparser than `chol (S' * S)`. `p = ccolamd (S, 1)` optimizes the ordering for `lu (S(:, p))`. The ordering is followed by a column elimination tree post-ordering.

`knobs` is an optional 1-element to 5-element input vector, with a default value of [0 10 10 1 0] if not present or empty. Entries not present are set to their defaults.

`knobs(1)` if nonzero, the ordering is optimized for `lu (S(:, p))`. It will be a poor ordering for `chol (S(:, p)' * S(:, p))`. This is the most important knob for `ccolamd`.

`knobs(2)` if S is m-by-n, rows with more than `max (16, knobs(2) * sqrt (n))` entries are ignored.

`knobs(3)` columns with more than `max (16, knobs(3) * sqrt (min (m, n)))` entries are ignored and ordered last in the output permutation (subject to the `cmember` constraints).

`knobs(4)` if nonzero, aggressive absorption is performed.

`knobs(5)` if nonzero, statistics and knobs are printed.

cmember is an optional vector of length n . It defines the constraints on the column ordering. If $\text{cmember}(j) = c$, then column j is in constraint set c (c must be in the range 1 to n). In the output permutation p , all columns in set 1 appear first, followed by all columns in set 2, and so on. $\text{cmember} = \text{ones}(1, n)$ if not present or empty. $\text{ccolamd}(S, [], 1 : n)$ returns $1 : n$

$p = \text{ccolamd}(S)$ is about the same as $p = \text{colamd}(S)$. *knobs* and its default values differ. *colamd* always does aggressive absorption, and it finds an ordering suitable for both $\text{lu}(S(:, p))$ and $\text{chol}(S(:, p)' * S(:, p))$; it cannot optimize its ordering for $\text{lu}(S(:, p))$ to the extent that $\text{ccolamd}(S, 1)$ can.

stats is an optional 20-element output vector that provides data about the ordering and the validity of the input matrix S . Ordering statistics are in $\text{stats}(1 : 3)$. $\text{stats}(1)$ and $\text{stats}(2)$ are the number of dense or empty rows and columns ignored by CCOLAMD and $\text{stats}(3)$ is the number of garbage collections performed on the internal data structure used by CCOLAMD (roughly of size $2.2 * \text{nnz}(S) + 4 * m + 7 * n$ integers).

$\text{stats}(4 : 7)$ provide information if CCOLAMD was able to continue. The matrix is OK if $\text{stats}(4)$ is zero, or 1 if invalid. $\text{stats}(5)$ is the rightmost column index that is unsorted or contains duplicate entries, or zero if no such column exists. $\text{stats}(6)$ is the last seen duplicate or out-of-order row index in the column index given by $\text{stats}(5)$, or zero if no such row index exists. $\text{stats}(7)$ is the number of duplicate or out-of-order row indices. $\text{stats}(8 : 20)$ is always zero in the current version of CCOLAMD (reserved for future use).

The authors of the code itself are S. Larimore, T. Davis and S. Rajamanickam in collaboration with J. Bilbert and E. Ng. Supported by the National Science Foundation (DMS-9504974, DMS-9803599, CCR-0203270), and a grant from Sandia National Lab. See <http://faculty.cse.tamu.edu/davis/suitesparse.html> for *ccolamd*, *csymamd*, *amd*, *colamd*, *symamd*, and other related orderings.

See also: [\[colamd\]](#), page 634, [\[csymamd\]](#), page 635.

```
p = colamd(S)
p = colamd(S, knobs)
[p, stats] = colamd(S)
[p, stats] = colamd(S, knobs)
```

Compute the column approximate minimum degree permutation.

$p = \text{colamd}(S)$ returns the column approximate minimum degree permutation vector for the sparse matrix S . For a non-symmetric matrix S , $S(:, p)$ tends to have sparser LU factors than S . The Cholesky factorization of $S(:, p)' * S(:, p)$ also tends to be sparser than that of $S' * S$.

knobs is an optional one- to three-element input vector. If S is m -by- n , then rows with more than $\max(16, \text{knobs}(1) * \sqrt{n})$ entries are ignored. Columns with more than $\max(16, \text{knobs}(2) * \sqrt{\min(m, n)})$ entries are removed prior to ordering, and ordered last in the output permutation p . Only completely dense rows or columns are removed if $\text{knobs}(1)$ and $\text{knobs}(2)$ are < 0 , respectively. If $\text{knobs}(3)$ is nonzero, *stats* and *knobs* are printed. The default is $\text{knobs} = [10 10 0]$. Note that *knobs* differs from earlier versions of *colamd*.

`stats` is an optional 20-element output vector that provides data about the ordering and the validity of the input matrix S . Ordering statistics are in `stats(1:3)`. `stats(1)` and `stats(2)` are the number of dense or empty rows and columns ignored by COLAMD and `stats(3)` is the number of garbage collections performed on the internal data structure used by COLAMD (roughly of size $2.2 * \text{nnz}(S) + 4 * m + 7 * n$ integers).

Octave built-in functions are intended to generate valid sparse matrices, with no duplicate entries, with ascending row indices of the nonzeros in each column, with a non-negative number of entries in each column (!) and so on. If a matrix is invalid, then COLAMD may or may not be able to continue. If there are duplicate entries (a row index appears two or more times in the same column) or if the row indices in a column are out of order, then COLAMD can correct these errors by ignoring the duplicate entries and sorting each column of its internal copy of the matrix S (the input matrix S is not repaired, however). If a matrix is invalid in other ways then COLAMD cannot continue, an error message is printed, and no output arguments (p or `stats`) are returned. COLAMD is thus a simple way to check a sparse matrix to see if it's valid.

`stats(4:7)` provide information if COLAMD was able to continue. The matrix is OK if `stats(4)` is zero, or 1 if invalid. `stats(5)` is the rightmost column index that is unsorted or contains duplicate entries, or zero if no such column exists. `stats(6)` is the last seen duplicate or out-of-order row index in the column index given by `stats(5)`, or zero if no such row index exists. `stats(7)` is the number of duplicate or out-of-order row indices. `stats(8:20)` is always zero in the current version of COLAMD (reserved for future use).

The ordering is followed by a column elimination tree post-ordering.

The authors of the code itself are Stefan I. Larimore and Timothy A. Davis. The algorithm was developed in collaboration with John Gilbert, Xerox PARC, and Esmond Ng, Oak Ridge National Laboratory. (see <http://faculty.cse.tamu.edu/davis/suitesparse.html>)

See also: `[colperm]`, page 635, `[symamd]`, page 636, `[ccolamd]`, page 633.

`p = colperm (s)`

Return the column permutations such that the columns of $s (:, p)$ are ordered in terms of increasing number of nonzero elements.

If s is symmetric, then p is chosen such that $s (p, p)$ orders the rows and columns with increasing number of nonzeros elements.

`p = csymamd (S)`

`p = csymamd (S, knobs)`

`p = csymamd (S, knobs, cmember)`

`[p, stats] = csymamd (...)`

For a symmetric positive definite matrix S , return the permutation vector p such that $S(p, p)$ tends to have a sparser Cholesky factor than S .

Sometimes `csymamd` works well for symmetric indefinite matrices too. The matrix S is assumed to be symmetric; only the strictly lower triangular part is referenced. S must be square. The ordering is followed by an elimination tree post-ordering.

knobs is an optional 1-element to 3-element input vector, with a default value of [10 1 0]. Entries not present are set to their defaults.

knobs(1) If S is n-by-n, then rows and columns with more than $\max(16, \text{knobs}(1)*\sqrt{n})$ entries are ignored, and ordered last in the output permutation (subject to the *cmember* constraints).

knobs(2) If nonzero, aggressive absorption is performed.

knobs(3) If nonzero, statistics and knobs are printed.

cmember is an optional vector of length n. It defines the constraints on the ordering. If *cmember*(j) = S, then row/column j is in constraint set c (c must be in the range 1 to n). In the output permutation p , rows/columns in set 1 appear first, followed by all rows/columns in set 2, and so on. *cmember* = ones (1,n) if not present or empty. *csymamd* ($S, [], 1:n$) returns 1:n.

$p = \text{csymamd } (S)$ is about the same as $p = \text{symamd } (S)$. *knobs* and its default values differ.

stats(4:7) provide information if CCOLAMD was able to continue. The matrix is OK if *stats*(4) is zero, or 1 if invalid. *stats*(5) is the rightmost column index that is unsorted or contains duplicate entries, or zero if no such column exists. *stats*(6) is the last seen duplicate or out-of-order row index in the column index given by *stats*(5), or zero if no such row index exists. *stats*(7) is the number of duplicate or out-of-order row indices. *stats*(8:20) is always zero in the current version of CCOLAMD (reserved for future use).

The authors of the code itself are S. Larimore, T. Davis and S. Rajamanickam in collaboration with J. Bilbert and E. Ng. Supported by the National Science Foundation (DMS-9504974, DMS-9803599, CCR-0203270), and a grant from Sandia National Lab. See <http://faculty.cse.tamu.edu/davis/suitesparse.html> for *ccolamd*, *colamd*, *csymamd*, *amd*, *colamd*, *symamd*, and other related orderings.

See also: [\[symamd\]](#), page 636, [\[ccolamd\]](#), page 633.

```
p = dmperm (S)
[p, q, r, S] = dmperm (S)
```

Perform a Dulmage-Mendelsohn permutation of the sparse matrix S .

With a single output argument *dmperm* performs the row permutations p such that $S(p, :)$ has no zero elements on the diagonal.

Called with two or more output arguments, returns the row and column permutations, such that $S(p, q)$ is in block triangular form. The values of r and S define the boundaries of the blocks. If S is square then $r == S$.

The method used is described in: A. Pothen & C.-J. Fan. *Computing the Block Triangular Form of a Sparse Matrix*. ACM Trans. Math. Software, 16(4):303-324, 1990.

See also: [\[colamd\]](#), page 634, [\[ccolamd\]](#), page 633.

```
p = symamd (S)
p = symamd (S, knobs)
[p, stats] = symamd (S)
```

[*p*, *stats*] = symamd (*S*, *knobs*)

For a symmetric positive definite matrix *S*, returns the permutation vector *p* such that $S(p, p)$ tends to have a sparser Cholesky factor than *S*.

Sometimes **symamd** works well for symmetric indefinite matrices too. The matrix *S* is assumed to be symmetric; only the strictly lower triangular part is referenced. *S* must be square.

knobs is an optional one- to two-element input vector. If *S* is n-by-n, then rows and columns with more than $\max(16, \text{knobs}(1) * \sqrt{n})$ entries are removed prior to ordering, and ordered last in the output permutation *p*. No rows/columns are removed if *knobs*(1) < 0. If *knobs*(2) is nonzero, *stats* and *knobs* are printed. The default is *knobs* = [10 0]. Note that *knobs* differs from earlier versions of **symamd**.

stats is an optional 20-element output vector that provides data about the ordering and the validity of the input matrix *S*. Ordering statistics are in *stats*(1:3). *stats*(1) = *stats*(2) is the number of dense or empty rows and columns ignored by SYMAMD and *stats*(3) is the number of garbage collections performed on the internal data structure used by SYMAMD (roughly of size $8.4 * \text{nnz}(\text{tril}(S, -1)) + 9 * n$ integers).

Octave built-in functions are intended to generate valid sparse matrices, with no duplicate entries, with ascending row indices of the nonzeros in each column, with a non-negative number of entries in each column (!) and so on. If a matrix is invalid, then SYMAMD may or may not be able to continue. If there are duplicate entries (a row index appears two or more times in the same column) or if the row indices in a column are out of order, then SYMAMD can correct these errors by ignoring the duplicate entries and sorting each column of its internal copy of the matrix *S* (the input matrix *S* is not repaired, however). If a matrix is invalid in other ways then SYMAMD cannot continue, an error message is printed, and no output arguments (*p* or *stats*) are returned. SYMAMD is thus a simple way to check a sparse matrix to see if it's valid.

stats(4:7) provide information if SYMAMD was able to continue. The matrix is OK if *stats*(4) is zero, or 1 if invalid. *stats*(5) is the rightmost column index that is unsorted or contains duplicate entries, or zero if no such column exists. *stats*(6) is the last seen duplicate or out-of-order row index in the column index given by *stats*(5), or zero if no such row index exists. *stats*(7) is the number of duplicate or out-of-order row indices. *stats*(8:20) is always zero in the current version of SYMAMD (reserved for future use).

The ordering is followed by a column elimination tree post-ordering.

The authors of the code itself are Stefan I. Larimore and Timothy A. Davis. The algorithm was developed in collaboration with John Gilbert, Xerox PARC, and Esmond Ng, Oak Ridge National Laboratory. (see <http://faculty.cse.tamu.edu/davis/suitesparse.html>)

See also: **[colperm]**, page 635, **[colamd]**, page 634.

***p* = symrcm (*S*)**

Return the symmetric reverse Cuthill-McKee permutation of *S*.

p is a permutation vector such that $S(p, p)$ tends to have its diagonal elements closer to the diagonal than S . This is a good preordering for LU or Cholesky factorization of matrices that come from “long, skinny” problems. It works for both symmetric and asymmetric S .

The algorithm represents a heuristic approach to the NP-complete bandwidth minimization problem. The implementation is based in the descriptions found in

E. Cuthill, J. McKee. *Reducing the Bandwidth of Sparse Symmetric Matrices*. Proceedings of the 24th ACM National Conference, 157–172 1969, Brandon Press, New Jersey.

A. George, J.W.H. Liu. *Computer Solution of Large Sparse Positive Definite Systems*, Prentice Hall Series in Computational Mathematics, ISBN 0-13-165274-5, 1981.

See also: [colperm], page 635, [colamd], page 634, [symamd], page 636.

22.2 Linear Algebra on Sparse Matrices

Octave includes a polymorphic solver for sparse matrices, where the exact solver used to factorize the matrix, depends on the properties of the sparse matrix itself. Generally, the cost of determining the matrix type is small relative to the cost of factorizing the matrix itself, but in any case the matrix type is cached once it is calculated, so that it is not re-determined each time it is used in a linear equation.

The selection tree for how the linear equation is solve is

1. If the matrix is diagonal, solve directly and goto 8
2. If the matrix is a permuted diagonal, solve directly taking into account the permutations. Goto 8
3. If the matrix is square, banded and if the band density is less than that given by `spparms ("banden")` continue, else goto 4.
 - a. If the matrix is tridiagonal and the right-hand side is not sparse continue, else goto 3b.
 1. If the matrix is Hermitian, with a positive real diagonal, attempt Cholesky factorization using LAPACK `xPTSV`.
 2. If the above failed or the matrix is not Hermitian with a positive real diagonal use Gaussian elimination with pivoting using LAPACK `xGTSV`, and goto 8.
 - b. If the matrix is Hermitian with a positive real diagonal, attempt Cholesky factorization using LAPACK `xPBTRF`.
 - c. if the above failed or the matrix is not Hermitian with a positive real diagonal use Gaussian elimination with pivoting using LAPACK `xGBTRF`, and goto 8.
4. If the matrix is upper or lower triangular perform a sparse forward or backward substitution, and goto 8
5. If the matrix is an upper triangular matrix with column permutations or lower triangular matrix with row permutations, perform a sparse forward or backward substitution, and goto 8
6. If the matrix is square, Hermitian with a real positive diagonal, attempt sparse Cholesky factorization using CHOLMOD.

7. If the sparse Cholesky factorization failed or the matrix is not Hermitian with a real positive diagonal, and the matrix is square, factorize, solve, and perform one refinement iteration using UMFPACK.
8. If the matrix is not square, or any of the previous solvers flags a singular or near singular matrix, find a minimum norm solution using CXSPARSE³.

The band density is defined as the number of nonzero values in the band divided by the total number of values in the full band. The banded matrix solvers can be entirely disabled by using `spparms` to set `bandden` to 1 (i.e., `spparms ("bandden", 1)`).

The QR solver factorizes the problem with a Dulmage-Mendelsohn decomposition, to separate the problem into blocks that can be treated as over-determined, multiple well determined blocks, and a final over-determined block. For matrices with blocks of strongly connected nodes this is a big win as LU decomposition can be used for many blocks. It also significantly improves the chance of finding a solution to over-determined problems rather than just returning a vector of NaN's.

All of the solvers above, can calculate an estimate of the condition number. This can be used to detect numerical stability problems in the solution and force a minimum norm solution to be used. However, for narrow banded, triangular or diagonal matrices, the cost of calculating the condition number is significant, and can in fact exceed the cost of factoring the matrix. Therefore the condition number is not calculated in these cases, and Octave relies on simpler techniques to detect singular matrices or the underlying LAPACK code in the case of banded matrices.

The user can force the type of the matrix with the `matrix_type` function. This overcomes the cost of discovering the type of the matrix. However, it should be noted that identifying the type of the matrix incorrectly will lead to unpredictable results, and so `matrix_type` should be used with care.

```
nest = normest (A)
nest = normest (A, tol)
[nest, iter] = normest (...)
```

Estimate the 2-norm of the matrix *A* using a power series analysis.

This is typically used for large matrices, where the cost of calculating `norm (A)` is prohibitive and an approximation to the 2-norm is acceptable.

tol is the tolerance to which the 2-norm is calculated. By default *tol* is 1e-6.

The optional output *iter* returns the number of iterations that were required for `normest` to converge.

See also: `[normest1]`, page 639, `[norm]`, page 551, `[cond]`, page 546, `[condestd]`, page 640.

```
nest = normest1 (A)
nest = normest1 (A, t)
nest = normest1 (A, t, x0)
nest = normest1 (Afun, t, x0, p1, p2, ...)
[nest, v] = normest1 (A, ...)
```

³ The CHOLMOD, UMFPACK and CXSPARSE packages were written by Tim Davis and are available at <http://faculty.cse.tamu.edu/davis/suitesparse.html>

```
[nest, v, w] = normest1 (A, ...)
[nest, v, w, iter] = normest1 (A, ...)
```

Estimate the 1-norm of the matrix A using a block algorithm.

`normest1` is best for large sparse matrices where only an estimate of the norm is required. For small to medium sized matrices, consider using `norm (A, 1)`. In addition, `normest1` can be used for the estimate of the 1-norm of a linear operator A when matrix-vector products $A * x$ and $A' * x$ can be cheaply computed. In this case, instead of the matrix A , a function `Afun (flag, x)` is used; it must return:

- the dimension n of A , if `flag` is "dim"
- true if A is a real operator, if `flag` is "real"
- the result $A * x$, if `flag` is "notransp"
- the result $A' * x$, if `flag` is "transp"

A typical case is A defined by $b \wedge m$, in which the result $A * x$ can be computed without even forming explicitly $b \wedge m$ by:

```
y = x;
for i = 1:m
    y = b * y;
endfor
```

The parameters $p1, p2, \dots$ are arguments of `Afun (flag, x, p1, p2, \dots)`.

The default value for t is 2. The algorithm requires matrix-matrix products with sizes $n \times n$ and $n \times t$.

The initial matrix $x0$ should have columns of unit 1-norm. The default initial matrix $x0$ has the first column `ones (n, 1) / n` and, if $t > 1$, the remaining columns with random elements $-1 / n, 1 / n$, divided by n .

On output, `nest` is the desired estimate, `v` and `w` are vectors such that $w = A * v$, with $\text{norm} (w, 1) = c * \text{norm} (v, 1)$. `iter` contains in `iter(1)` the number of iterations (the maximum is hardcoded to 5) and in `iter(2)` the total number of products $A * x$ or $A' * x$ performed by the algorithm.

Algorithm Note: `normest1` uses random numbers during evaluation. Therefore, if consistent results are required, the "state" of the random generator should be fixed before invoking `normest1`.

Reference: N. J. Higham and F. Tisseur, *A block algorithm for matrix 1-norm estimation, with application to 1-norm pseudospectra*, SIAM J. Matrix Anal. Appl., pp. 1185–1201, Vol 21, No. 4, 2000.

See also: [\[normest\]](#), page 639, [\[norm\]](#), page 551, [\[cond\]](#), page 546, [\[condestd\]](#), page 640.

```
cest = condest (A)
cest = condest (A, t)
cest = condest (A, solvefun, t, p1, p2, ...)
cest = condest (Afcn, solvefun, t, p1, p2, ...)
[cest, v] = condest (...)
```

Estimate the 1-norm condition number of a square matrix A using t test vectors and a randomized 1-norm estimator.

The optional input t specifies the number of test vectors (default 5).

If the matrix is not explicit, e.g., when estimating the condition number of A given an LU factorization, `condest` uses the following functions:

- `Afcn` which must return
 - the dimension n of a , if `flag` is "dim"
 - true if a is a real operator, if `flag` is "real"
 - the result $a * x$, if `flag` is "notransp"
 - the result $a' * x$, if `flag` is "transp"
- `solvefcn` which must return
 - the dimension n of a , if `flag` is "dim"
 - true if a is a real operator, if `flag` is "real"
 - the result $a \backslash x$, if `flag` is "notransp"
 - the result $a' \backslash x$, if `flag` is "transp"

The parameters $p1, p2, \dots$ are arguments of `Afcn(flag, x, p1, p2, ...)` and `solvefcn(flag, x, p1, p2, ...)`.

The principal output is the 1-norm condition number estimate `cest`.

The optional second output is an approximate null vector when `cest` is large; it satisfies the equation `norm(A*v, 1) == norm(A, 1) * norm(v, 1) / est`.

Algorithm Note: `condest` uses a randomized algorithm to approximate the 1-norms. Therefore, if consistent results are required, the "state" of the random generator should be fixed before invoking `condest`.

References:

- N.J. Higham and F. Tisseur, *A Block Algorithm for Matrix 1-Norm Estimation, with an Application to 1-Norm Pseudospectra*. SIMAX vol 21, no 4, pp 1185-1201. <https://dx.doi.org/10.1137/S0895479899356080>
- N.J. Higham and F. Tisseur, *A Block Algorithm for Matrix 1-Norm Estimation, with an Application to 1-Norm Pseudospectra*. <https://citeseer.ist.psu.edu/223007.html>

See also: [cond], page 546, [norm], page 551, [normest1], page 639, [normest], page 639.

```
spparms()
vals = spparms()
[keys, vals] = spparms()
val = spparms(key)
spparms(vals)
spparms("default")
spparms("tight")
spparms(key, val)
```

Query or set the parameters used by the sparse solvers and factorization functions.

The first four calls above get information about the current settings, while the others change the current settings. The parameters are stored as pairs of keys and values, where the values are all floats and the keys are one of the following strings:

'spumoni' Printing level of debugging information of the solvers (default 0)

'ths_rel'	Included for compatibility. Not used. (default 1)
'ths_abs'	Included for compatibility. Not used. (default 1)
'exact_d'	Included for compatibility. Not used. (default 0)
'supernd'	Included for compatibility. Not used. (default 3)
'rreduce'	Included for compatibility. Not used. (default 3)
'wh_frac'	Included for compatibility. Not used. (default 0.5)
'autommd'	Flag whether the LU/QR and the '\' and '/' operators will automatically use the sparsity preserving mmd functions (default 1)
'autoamd'	Flag whether the LU and the '\' and '/' operators will automatically use the sparsity preserving amd functions (default 1)
'piv_tol'	The pivot tolerance of the UMFPACK solvers (default 0.1)
'sym_tol'	The pivot tolerance of the UMFPACK symmetric solvers (default 0.001)
'bandden'	The density of nonzero elements in a banded matrix before it is treated by the LAPACK banded solvers (default 0.5)
'umfpack'	Flag whether the UMFPACK or mmd solvers are used for the LU, '\' and '/' operations (default 1)

The value of individual keys can be set with `spparms (key, val)`. The default values can be restored with the special keyword "default". The special keyword "tight" can be used to set the mmd solvers to attempt a sparser solution at the potential cost of longer running time.

See also: [chol], page 555, [colamd], page 634, [lu], page 557, [qr], page 559, [symamd], page 636.

`p = sprank (S)`

Calculate the structural rank of the sparse matrix S .

Note that only the structure of the matrix is used in this calculation based on a Dulmage-Mendelsohn permutation to block triangular form. As such the numerical rank of the matrix S is bounded by `sprank (S) >= rank (S)`. Ignoring floating point errors `sprank (S) == rank (S)`.

See also: [dmperm], page 636.

```
[count, h, parent, post, R] = symbfact (S)
[...] = symbfact (S, typ)
[...] = symbfact (S, typ, mode)
```

Perform a symbolic factorization analysis of the sparse matrix S .

The input variables are

S S is a real or complex sparse matrix.

typ Is the type of the factorization and can be one of

"sym" (default)

Factorize S . Assumes S is symmetric and uses the upper triangular portion of the matrix.

	"col"	Factorize $S' * S$.
	"row"	Factorize $S * S'$.
	"lo"	Factorize S' . Assumes S is symmetric and uses the lower triangular portion of the matrix.
<i>mode</i>		When <i>mode</i> is unspecified return the Cholesky factorization for R . If <i>mode</i> is "lower" or "L" then return the conjugate transpose R' which is a lower triangular factor. The conjugate transpose version is faster and uses less memory, but still returns the same values for all other outputs: <i>count</i> , <i>h</i> , <i>parent</i> , and <i>post</i> .

The output variables are:

<i>count</i>	The row counts of the Cholesky factorization as determined by <i>typ</i> . The computational difficulty of performing the true factorization using chol is <code>sum(count.^2)</code> .
<i>h</i>	The height of the elimination tree.
<i>parent</i>	The elimination tree itself.
<i>post</i>	A sparse boolean matrix whose structure is that of the Cholesky factorization as determined by <i>typ</i> .

See also: [\[chol\]](#), page 555, [\[etree\]](#), page 626, [\[treelayout\]](#), page 627.

For non square matrices, the user can also utilize the **spaugment** function to find a least squares solution to a linear equation.

s = spaugment (A, c)

Create the augmented matrix of A .

This is given by

```
[c * eye(m, m), A;
   A', zeros(n, n)]
```

This is related to the least squares solution of $A \setminus b$, by

```
s * [r / c; x] = [b, zeros(n, columns(b))]
```

where r is the residual error

```
r = b - A * x
```

As the matrix s is symmetric indefinite it can be factorized with **lu**, and the minimum norm solution can therefore be found without the need for a **qr** factorization. As the residual error will be **zeros (m, m)** for underdetermined problems, and example can be

```
m = 11; n = 10; mn = max (m, n);
A = spdiags ([ones(mn,1), 10*ones(mn,1), -ones(mn,1)],
              [-1, 0, 1], m, n);
x0 = A \ ones (m,1);
s = spaugment (A);
[L, U, P, Q] = lu (s);
x1 = Q * (U \ (L \ (P * [ones(m,1); zeros(n,1)])));
x1 = x1(end - n + 1 : end);
```

To find the solution of an overdetermined problem needs an estimate of the residual error r and so it is more complex to formulate a minimum norm solution using the `sпаugment` function.

In general the left division operator is more stable and faster than using the `sпаugment` function.

See also: [\[mldivide\]](#), page 150.

Finally, the function `eigs` can be used to calculate a limited number of eigenvalues and eigenvectors based on a selection criteria and likewise for `svds` which calculates a limited number of singular values and vectors.

```
d = eigs (A)
d = eigs (A, k)
d = eigs (A, k, sigma)
d = eigs (A, k, sigma, opts)
d = eigs (A, B)
d = eigs (A, B, k)
d = eigs (A, B, k, sigma)
d = eigs (A, B, k, sigma, opts)
d = eigs (af, n)
d = eigs (af, n, B)
d = eigs (af, n, k)
d = eigs (af, n, B, k)
d = eigs (af, n, k, sigma)
d = eigs (af, n, B, k, sigma)
d = eigs (af, n, k, sigma, opts)
d = eigs (af, n, B, k, sigma, opts)
[V, d] = eigs (A, ...)
[V, d] = eigs (af, n, ...)
[V, d, flag] = eigs (A, ...)
[V, d, flag] = eigs (af, n, ...)
```

Calculate a limited number of eigenvalues and eigenvectors of A , based on a selection criteria.

The number of eigenvalues and eigenvectors to calculate is given by k and defaults to 6.

By default, `eigs` solve the equation $A\nu = \lambda\nu$, where λ is a scalar representing one of the eigenvalues, and ν is the corresponding eigenvector. If given the positive definite matrix B then `eigs` solves the general eigenvalue equation $A\nu = \lambda B\nu$.

The argument `sigma` determines which eigenvalues are returned. `sigma` can be either a scalar or a string. When `sigma` is a scalar, the k eigenvalues closest to `sigma` are returned. If `sigma` is a string, it must have one of the following values.

- "`lm`" Largest Magnitude (default).
- "`sm`" Smallest Magnitude.
- "`la`" Largest Algebraic (valid only for real symmetric problems).
- "`sa`" Smallest Algebraic (valid only for real symmetric problems).

"be"	Both Ends, with one more from the high-end if k is odd (valid only for real symmetric problems).
"lr"	Largest Real part (valid only for complex or unsymmetric problems).
"sr"	Smallest Real part (valid only for complex or unsymmetric problems).
"li"	Largest Imaginary part (valid only for complex or unsymmetric problems).
"si"	Smallest Imaginary part (valid only for complex or unsymmetric problems).

If $opts$ is given, it is a structure defining possible options that `eigs` should use. The fields of the $opts$ structure are:

<code>issym</code>	If af is given, then flags whether the function af defines a symmetric problem. It is ignored if A is given. The default is false.
<code>isreal</code>	If af is given, then flags whether the function af defines a real problem. It is ignored if A is given. The default is true.
<code>tol</code>	Defines the required convergence tolerance, calculated as $tol * \text{norm}(A)$. The default is <code>eps</code> .
<code>maxit</code>	The maximum number of iterations. The default is 300.
<code>p</code>	The number of Lanczos basis vectors to use. More vectors will result in faster convergence, but a greater use of memory. The optimal value of p is problem dependent and should be in the range $k + 1$ to n . The default value is $2 * k$.
<code>v0</code>	The starting vector for the algorithm. An initial vector close to the final vector will speed up convergence. The default is for ARPACK to randomly generate a starting vector. If specified, $v0$ must be an n -by-1 vector where $n = \text{rows}(A)$
<code>disp</code>	The level of diagnostic printout (0 1 2). If <code>disp</code> is 0 then diagnostics are disabled. The default value is 0.
<code>cholB</code>	Flag if <code>chol(B)</code> is passed rather than B . The default is false.
<code>permB</code>	The permutation vector of the Cholesky factorization of B if <code>cholB</code> is true. It is obtained by $[R, \sim, permB] = chol(B, "vector")$. The default is <code>1:n</code> .

It is also possible to represent A by a function denoted af . af must be followed by a scalar argument n defining the length of the vector argument accepted by af . af can be a function handle, an inline function, or a string. When af is a string it holds the name of the function to use.

af is a function of the form $y = af(x)$ where the required return value of af is determined by the value of $sigma$. The four possible forms are

$A * x$	if $sigma$ is not given or is a string other than "sm".
$A \setminus x$	if $sigma$ is 0 or "sm".

`(A - sigma * I) \ x`

for the standard eigenvalue problem, where `I` is the identity matrix of the same size as `A`.

`(A - sigma * B) \ x`

for the general eigenvalue problem.

The return arguments of `eigs` depend on the number of return arguments requested. With a single return argument, a vector `d` of length `k` is returned containing the `k` eigenvalues that have been found. With two return arguments, `V` is a `n`-by-`k` matrix whose columns are the `k` eigenvectors corresponding to the returned eigenvalues. The eigenvalues themselves are returned in `d` in the form of a `n`-by-`k` matrix, where the elements on the diagonal are the eigenvalues.

Given a third return argument `flag`, `eigs` returns the status of the convergence. If `flag` is 0 then all eigenvalues have converged. Any other value indicates a failure to converge.

This function is based on the ARPACK package, written by R. Lehoucq, K. Maschhoff, D. Sorensen, and C. Yang. For more information see <http://www.caam.rice.edu/software/ARPACK/>.

See also: `[eig]`, page 547, `[svds]`, page 646.

```
s = svds (A)
s = svds (A, k)
s = svds (A, k, sigma)
s = svds (A, k, sigma, opts)
[u, s, v] = svds (...)
[u, s, v, flag] = svds (...)
```

Find a few singular values of the matrix `A`.

The singular values are calculated using

```
[m, n] = size (A);
s = eigs ([sparse(m, m), A;
           A', sparse(n, n)])
```

The eigenvalues returned by `eigs` correspond to the singular values of `A`. The number of singular values to calculate is given by `k` and defaults to 6.

The argument `sigma` specifies which singular values to find. When `sigma` is the string '`'L'`', the default, the largest singular values of `A` are found. Otherwise, `sigma` must be a real scalar and the singular values closest to `sigma` are found. As a corollary, `sigma = 0` finds the smallest singular values. Note that for relatively small values of `sigma`, there is a chance that the requested number of singular values will not be found. In that case `sigma` should be increased.

`opts` is a structure defining options that `svds` will pass to `eigs`. The possible fields of this structure are documented in `eigs`. By default, `svds` sets the following three fields:

`tol` The required convergence tolerance for the singular values. The default value is `1e-10`. `eigs` is passed `tol / sqrt(2)`.

`maxit` The maximum number of iterations. The default is 300.

disp The level of diagnostic printout (0|1|2). If **disp** is 0 then diagnostics are disabled. The default value is 0.

If more than one output is requested then **svds** will return an approximation of the singular value decomposition of A

$$A_{\text{approx}} = u * s * v'$$

where A_{approx} is a matrix of size A but only rank k .

flag returns 0 if the algorithm has successfully converged, and 1 otherwise. The test for convergence is

$$\text{norm}(A * v - u * s, 1) \leq tol * \text{norm}(A, 1)$$

svds is best for finding only a few singular values from a large sparse matrix. Otherwise, **svd (full (A))** will likely be more efficient.

See also: [\[svd\]](#), page 565, [\[eigs\]](#), page 644.

22.3 Iterative Techniques Applied to Sparse Matrices

The left division `\` and right division `/` operators, discussed in the previous section, use direct solvers to resolve a linear equation of the form $x = A \backslash b$ or $x = b / A$. Octave also includes a number of functions to solve sparse linear equations using iterative techniques.

```
x = pcg (A, b, tol, maxit, m1, m2, x0, ...)
x = pcg (A, b, tol, maxit, M, [], x0, ...)
[x, flag, relres, iter, resvec, eigest] = pcg (A, b, ...)
```

Solve the linear system of equations $A * x = b$ by means of the Preconditioned Conjugate Gradient iterative method.

The input arguments are:

- A is the matrix of the linear system and it must be square. A can be passed as a matrix, function handle, or inline function **Afun** such that **Afun(x) = A * x**. Additional parameters to **Afun** may be passed after **x0**.

A has to be Hermitian and Positive Definite (HPD). If **pcg** detects A not to be positive definite, a warning is printed and the **flag** output is set.

- b is the right-hand side vector.
- tol is the required relative tolerance for the residual error, $b - A * x$. The iteration stops if $\text{norm}(b - A * x) \leq tol * \text{norm}(b)$. If **tol** is omitted or empty, then a tolerance of 1e-6 is used.
- $maxit$ is the maximum allowed number of iterations; if **maxit** is omitted or empty then a value of 20 is used.
- m is a HPD preconditioning matrix. For any decomposition $m = p1 * p2$ such that $\text{inv}(p1) * A * \text{inv}(p2)$ is HPD, the conjugate gradient method is formally applied to the linear system $\text{inv}(p1) * A * \text{inv}(p2) * y = \text{inv}(p1) * b$, with $x = \text{inv}(p2) * y$ (split preconditioning). In practice, at each iteration of the conjugate gradient method a linear system with matrix m is solved with **mldivide**. If a particular factorization $m = m1 * m2$ is available (for instance, an incomplete Cholesky factorization of A), the two matrices $m1$ and $m2$ can be passed and the relative linear systems are solved with the **mldivide** operator. Note that a proper

choice of the preconditioner may dramatically improve the overall performance of the method. Instead of matrices $m1$ and $m2$, the user may pass two functions which return the results of applying the inverse of $m1$ and $m2$ to a vector. If $m1$ is omitted or empty `[]`, then no preconditioning is applied. If no factorization of m is available, $m2$ can be omitted or left `[]`, and the input variable $m1$ can be used to pass the preconditioner m .

- $x0$ is the initial guess. If $x0$ is omitted or empty then the function sets $x0$ to a zero vector by default.

The arguments which follow $x0$ are treated as parameters, and passed in an appropriate manner to any of the functions (A or $m1$ or $m2$) that have been given to `pcg`. See the examples below for further details.

The output arguments are:

- x is the computed approximation to the solution of $A * x = b$. If the algorithm did not converge, then x is the iteration which has the minimum residual.
- $flag$ reports on the convergence:
 - 0: The algorithm converged to within the prescribed tolerance.
 - 1: The algorithm did not converge and it reached the maximum number of iterations.
 - 2: The preconditioner matrix is singular.
 - 3: The algorithm stagnated, i.e., the absolute value of the difference between the current iteration x and the previous is less than `eps * norm(x, 2)`.
 - 4: The algorithm detects that the input (preconditioned) matrix is not HPD.
- $relres$ is the ratio of the final residual to its initial value, measured in the Euclidean norm.
- $iter$ indicates the iteration of x which it was computed. Since the output x corresponds to the minimal residual solution, the total number of iterations that the method performed is given by `length(resvec) - 1`.
- $resvec$ describes the convergence history of the method. `resvec(i, 1)` is the Euclidean norm of the residual, and `resvec(i, 2)` is the preconditioned residual norm, after the $(i-1)$ -th iteration, $i = 1, 2, \dots, iter+1$. The preconditioned residual norm is defined as $r' * (m \setminus r)$ where $r = b - A * x$, see also the description of m . If `eigest` is not required, only `resvec(:, 1)` is returned.
- `eigest` returns the estimate for the smallest `eigest(1)` and largest `eigest(2)` eigenvalues of the preconditioned matrix $P = m \setminus A$. In particular, if no preconditioning is used, the estimates for the extreme eigenvalues of A are returned. `eigest(1)` is an overestimate and `eigest(2)` is an underestimate, so that `eigest(2) / eigest(1)` is a lower bound for `cond(P, 2)`, which nevertheless in the limit should theoretically be equal to the actual value of the condition number.

Let us consider a trivial problem with a tridiagonal matrix

```

n = 10;
A = toeplitz (sparse ([1, 1], [1, 2], [2, 1], 1, n));
b = A * ones (n, 1);
M1 = ichol (A); # in this tridiagonal case it corresponds to chol (A)'■
M2 = M1';
M = M1 * M2;
Afun = @(x) A * x;
Mfun = @(x) M \ x;
M1fun = @(x) M1 \ x;
M2fun = @(x) M2 \ x;

```

EXAMPLE 1: Simplest use of pcg

```
x = pcg (A, b)
```

EXAMPLE 2: pcg with a function which computes $A * x$

```
x = pcg (Afun, b)
```

EXAMPLE 3: pcg with a preconditioner matrix M

```
x = pcg (A, b, 1e-06, 100, M)
```

EXAMPLE 4: pcg with a function as preconditioner

```
x = pcg (Afun, b, 1e-6, 100, Mfun)
```

EXAMPLE 5: pcg with preconditioner matrices $M1$ and $M2$

```
x = pcg (A, b, 1e-6, 100, M1, M2)
```

EXAMPLE 6: pcg with functions as preconditioners

```
x = pcg (Afun, b, 1e-6, 100, M1fun, M2fun)
```

EXAMPLE 7: pcg with as input a function requiring an argument

```

function y = Ap (A, x, p) # compute  $A^p * x$ 
y = x;
for i = 1:p
    y = A * y;
endfor
endfunction
Apfun = @(x, p) Ap (A, x, p);
x = pcg (Apfun, b, [], [], [], [], [], 2);

```

EXAMPLE 8: explicit example to show that pcg uses a split preconditioner

```

M1 = ichol (A + 0.1 * eye (n)); # factorization of A perturbed
M2 = M1';
M = M1 * M2;

## reference solution computed by pcg after two iterations
[x_ref, f1] = pcg (A, b, [], 2, M)

## split preconditioning
[y, f1] = pcg ((M1 \ A) / M2, M1 \ b, [], 2)
x = M2 \ y # compare x and x_ref

```

References:

1. C.T. Kelley, *Iterative Methods for Linear and Nonlinear Equations*, SIAM, 1995.
(the base PCG algorithm)
2. Y. Saad, *Iterative Methods for Sparse Linear Systems*, PWS 1996. (condition number estimate from PCG) Revised version of this book is available online at
<https://www-users.cs.umn.edu/~saad/books.html>

See also: [sparse], page 622, [pcr], page 650, [gmres], page 576, [bicg], page 570, [bicgstab], page 572, [cgs], page 574.

```
x = pcr (A, b, tol, maxit, m, x0, ...)
[x, flag, relres, iter, resvec] = pcr (...)
```

Solve the linear system of equations $A * x = b$ by means of the Preconditioned Conjugate Residuals iterative method.

The input arguments are

- A can be either a square (preferably sparse) matrix or a function handle, inline function or string containing the name of a function which computes $A * x$. In principle A should be symmetric and non-singular; if `pcr` finds A to be numerically singular, you will get a warning message and the `flag` output parameter will be set.
- b is the right hand side vector.
- tol is the required relative tolerance for the residual error, $b - A * x$. The iteration stops if $\text{norm}(b - A * x) \leq tol * \text{norm}(b - A * x_0)$. If tol is empty or is omitted, the function sets $tol = 1e-6$ by default.
- $maxit$ is the maximum allowable number of iterations; if [] is supplied for $maxit$, or `pcr` has less arguments, a default value equal to 20 is used.
- m is the (left) preconditioning matrix, so that the iteration is (theoretically) equivalent to solving by `pcr P * x = m \ b`, with $P = m \ A$. Note that a proper choice of the preconditioner may dramatically improve the overall performance of the method. Instead of matrix m , the user may pass a function which returns the results of applying the inverse of m to a vector (usually this is the preferred way of using the preconditioner). If [] is supplied for m , or m is omitted, no preconditioning is applied.
- x_0 is the initial guess. If x_0 is empty or omitted, the function sets x_0 to a zero vector by default.

The arguments which follow x_0 are treated as parameters, and passed in a proper way to any of the functions (A or m) which are passed to `pcr`. See the examples below for further details.

The output arguments are

- x is the computed approximation to the solution of $A * x = b$.
- $flag$ reports on the convergence. $flag = 0$ means the solution converged and the tolerance criterion given by tol is satisfied. $flag = 1$ means that the $maxit$ limit for the iteration count was reached. $flag = 3$ reports a `pcr` breakdown, see [1] for details.

- *relres* is the ratio of the final residual to its initial value, measured in the Euclidean norm.
- *iter* is the actual number of iterations performed.
- *resvec* describes the convergence history of the method, so that *resvec* (*i*) contains the Euclidean norms of the residual after the (*i*-1)-th iteration, *i* = 1, 2, ..., *iter*+1.

Let us consider a trivial problem with a diagonal matrix (we exploit the sparsity of *A*)

```
n = 10;
A = sparse (diag (1:n));
b = rand (N, 1);
```

EXAMPLE 1: Simplest use of pcr

```
x = pcr (A, b)
```

EXAMPLE 2: pcr with a function which computes *A* * *x*.

```
function y = apply_a (x)
    y = [1:10]' .* x;
endfunction
```

```
x = pcr ("apply_a", b)
```

EXAMPLE 3: Preconditioned iteration, with full diagnostics. The preconditioner (quite strange, because even the original matrix *A* is trivial) is defined as a function

```
function y = apply_m (x)
    k = floor (length (x) - 2);
    y = x;
    y(1:k) = x(1:k) ./ [1:k]';
endfunction

[x, flag, relres, iter, resvec] = ...
    pcr (A, b, [], [], "apply_m")
semilogy ([1:iter+1], resvec);
```

EXAMPLE 4: Finally, a preconditioner which depends on a parameter *k*.

```
function y = apply_m (x, varargin)
    k = varargin{1};
    y = x;
    y(1:k) = x(1:k) ./ [1:k]';
endfunction

[x, flag, relres, iter, resvec] = ...
    pcr (A, b, [], [], "apply_m", [], 3)
```

References:

[1] W. Hackbusch, *Iterative Solution of Large Sparse Systems of Equations*, section 9.5.4; Springer, 1994

See also: [\[sparse\]](#), page 622, [\[pcg\]](#), page 647.

The speed with which an iterative solver converges to a solution can be accelerated with the use of a pre-conditioning matrix M . In this case the linear equation $M^{-1} * x = M^{-1} * A \setminus b$ is solved instead. Typical pre-conditioning matrices are partial factorizations of the original matrix.

`L = ichol (A)`

`L = ichol (A, opts)`

Compute the incomplete Cholesky factorization of the sparse square matrix A .

By default, `ichol` uses only the lower triangle of A and produces a lower triangular factor L such that $L*L'$ approximates A .

The factor given by this routine may be useful as a preconditioner for a system of linear equations being solved by iterative methods such as PCG (Preconditioned Conjugate Gradient).

The factorization may be modified by passing options in a structure `opts`. The option name is a field of the structure and the setting is the value of field. Names and specifiers are case sensitive.

`type` Type of factorization.

`"nofill"` (default)

Incomplete Cholesky factorization with no fill-in (IC(0)).

`"ict"`

Incomplete Cholesky factorization with threshold dropping (ICT).

`diagcomp` A non-negative scalar `alpha` for incomplete Cholesky factorization of $A + alpha * diag (diag (A))$ instead of A . This can be useful when A is not positive definite. The default value is 0.

`droptol` A non-negative scalar specifying the drop tolerance for factorization if performing ICT. The default value is 0 which produces the complete Cholesky factorization.

Non-diagonal entries of L are set to 0 unless

`abs (L(i,j)) >= droptol * norm (A(j:end, j), 1).`

`michol` Modified incomplete Cholesky factorization:

`"off"` (default)

Row and column sums are not necessarily preserved.

`"on"`

The diagonal of L is modified so that row (and column) sums are preserved even when elements have been dropped during the factorization. The relationship preserved is: $A * e = L * L' * e$, where e is a vector of ones.

`shape`

`"lower"` (default)

Use only the lower triangle of A and return a lower triangular factor L such that $L*L'$ approximates A .

`"upper"`

Use only the upper triangle of A and return an upper triangular factor U such that $U'*U$ approximates A .

EXAMPLES

The following problem demonstrates how to factorize a sample symmetric positive definite matrix with the full Cholesky decomposition and with the incomplete one.

```
A = [ 0.37, -0.05, -0.05, -0.07;
      -0.05,  0.116,  0.0,   -0.05;
      -0.05,  0.0,    0.116, -0.05;
      -0.07, -0.05, -0.05,  0.202];
A = sparse (A);
nnz (tril (A))
ans = 9
L = chol (A, "lower");
nnz (L)
ans = 10
norm (A - L * L', "fro") / norm (A, "fro")
ans = 1.1993e-16
opts.type = "nofill";
L = ichol (A, opts);
nnz (L)
ans = 9
norm (A - L * L', "fro") / norm (A, "fro")
ans = 0.019736
```

Another example for decomposition is a finite difference matrix used to solve a boundary value problem on the unit square.

```
nx = 400; ny = 200;
hx = 1 / (nx + 1); hy = 1 / (ny + 1);
Dxx = spdiags ([ones(nx, 1), -2*ones(nx, 1), ones(nx, 1)],
                [-1 0 1], nx, nx) / (hx ^ 2);
Dyy = spdiags ([ones(ny, 1), -2*ones(ny, 1), ones(ny, 1)],
                [-1 0 1], ny, ny) / (hy ^ 2);
A = -kron (Dxx, speye (ny)) - kron (speye (nx), Dyy);
nnz (tril (A))
ans = 239400
opts.type = "nofill";
L = ichol (A, opts);
nnz (tril (A))
ans = 239400
norm (A - L * L', "fro") / norm (A, "fro")
ans = 0.062327
```

References for implemented algorithms:

- [1] Y. Saad. "Preconditioning Techniques." *Iterative Methods for Sparse Linear Systems*, PWS Publishing Company, 1996.
- [2] M. Jones, P. Plassmann: *An Improved Incomplete Cholesky Factorization*, 1992.

See also: [\[chol\]](#), page 555, [\[ilu\]](#), page 653, [\[pcg\]](#), page 647.

[ilu \(A\)](#)
[ilu \(A, opts\)](#)

```
[L, U] = ilu (...)  
[L, U, P] = ilu (...)
```

Compute the incomplete LU factorization of the sparse square matrix A .

`ilu` returns a unit lower triangular matrix L , an upper triangular matrix U , and optionally a permutation matrix P , such that $L*U$ approximates $P*A$.

The factors given by this routine may be useful as preconditioners for a system of linear equations being solved by iterative methods such as BICG (BiConjugate Gradients) or GMRES (Generalized Minimum Residual Method).

The factorization may be modified by passing options in a structure `opts`. The option name is a field of the structure and the setting is the value of field. Names and specifiers are case sensitive.

`type` Type of factorization.

"nofill" (default)

ILU factorization with no fill-in (ILU(0)).

Additional supported options: `milu`.

"crout" Crout version of ILU factorization (ILUC).

Additional supported options: `milu`, `droptol`.

"ilutp" ILU factorization with threshold and pivoting.

Additional supported options: `milu`, `droptol`, `udiag`, `thresh`.

`droptol` A non-negative scalar specifying the drop tolerance for factorization. The default value is 0 which produces the complete LU factorization.

Non-diagonal entries of U are set to 0 unless

`abs (U(i,j)) >= droptol * norm (A(:,j))`.

Non-diagonal entries of L are set to 0 unless

`abs (L(i,j)) >= droptol * norm (A(:,j))/U(j,j)`.

`milu` Modified incomplete LU factorization:

"row" Row-sum modified incomplete LU factorization. The factorization preserves row sums: $A * e = L * U * e$, where e is a vector of ones.

"col" Column-sum modified incomplete LU factorization. The factorization preserves column sums: $e' * A = e' * L * U$.

"off" (default)

Row and column sums are not necessarily preserved.

`udiag` If true, any zeros on the diagonal of the upper triangular factor are replaced by the local drop tolerance `droptol * norm (A(:,j))/U(j,j)`. The default is false.

`thresh` Pivot threshold for factorization. It can range between 0 (diagonal pivoting) and 1 (default), where the maximum magnitude entry in the column is chosen to be the pivot.

If `ilu` is called with just one output, the returned matrix is $L + U - \text{speye}(\text{size}(A))$, where L is unit lower triangular and U is upper triangular.

With two outputs, `ilu` returns a unit lower triangular matrix L and an upper triangular matrix U . For `opts.type == "ilutp"`, one of the factors is permuted based on the value of `opts.milu`. When `opts.milu == "row"`, U is a column permuted upper triangular factor. Otherwise, L is a row-permuted unit lower triangular factor.

If there are three named outputs and `opts.milu != "row"`, P is returned such that L and U are incomplete factors of $P*A$. When `opts.milu == "row"`, P is returned such that L and U are incomplete factors of $A*P$.

EXAMPLES

```
A = gallery ("neumann", 1600) + speye (1600);
opts.type = "nofill";
nnz (A)
ans = 7840

nnz (lu (A))
ans = 126478

nnz (ilu (A, opts))
ans = 7840
```

This shows that A has 7,840 nonzeros, the complete LU factorization has 126,478 nonzeros, and the incomplete LU factorization, with 0 level of fill-in, has 7,840 nonzeros, the same amount as A . Taken from: <http://www.mathworks.com/help/matlab/ref/ilu.html>

```
A = gallery ("wathen", 10, 10);
b = sum (A, 2);
tol = 1e-8;
maxit = 50;
opts.type = "crout";
opts.droptol = 1e-4;
[L, U] = ilu (A, opts);
x = bicg (A, b, tol, maxit, L, U);
norm (A * x - b, inf)
```

This example uses ILU as preconditioner for a random FEM-Matrix, which has a large condition number. Without L and U BICG would not converge.

See also: [\[lu\]](#), page 557, [\[ichol\]](#), page 652, [\[bicg\]](#), page 570, [\[gmres\]](#), page 576.

22.4 Real Life Example using Sparse Matrices

A common application for sparse matrices is in the solution of Finite Element Models. Finite element models allow numerical solution of partial differential equations that do not have closed form solutions, typically because of the complex shape of the domain.

In order to motivate this application, we consider the boundary value Laplace equation. This system can model scalar potential fields, such as heat or electrical potential. Given a medium Ω with boundary $\partial\Omega$. At all points on the $\partial\Omega$ the boundary conditions are

known, and we wish to calculate the potential in Ω . Boundary conditions may specify the potential (Dirichlet boundary condition), its normal derivative across the boundary (Neumann boundary condition), or a weighted sum of the potential and its derivative (Cauchy boundary condition).

In a thermal model, we want to calculate the temperature in Ω and know the boundary temperature (Dirichlet condition) or heat flux (from which we can calculate the Neumann condition by dividing by the thermal conductivity at the boundary). Similarly, in an electrical model, we want to calculate the voltage in Ω and know the boundary voltage (Dirichlet) or current (Neumann condition after diving by the electrical conductivity). In an electrical model, it is common for much of the boundary to be electrically isolated; this is a Neumann boundary condition with the current equal to zero.

The simplest finite element models will divide Ω into simplexes (triangles in 2D, pyramids in 3D). We take as a 3-D example a cylindrical liquid filled tank with a small non-conductive ball from the EIDORS project⁴. This model is designed to reflect an application of electrical impedance tomography, where current patterns are applied to such a tank in order to image the internal conductivity distribution. In order to describe the FEM geometry, we have a matrix of vertices `nodes` and simplices `elems`.

The following example creates a simple rectangular 2-D electrically conductive medium with 10 V and 20 V imposed on opposite sides (Dirichlet boundary conditions). All other edges are electrically isolated.

```

node_y = [1;1.2;1.5;1.8;2]*ones(1,11);
node_x = ones(5,1)*[1,1.05,1.1,1.2, ...
                   1.3,1.5,1.7,1.8,1.9,1.95,2];
nodes = [node_x(:), node_y(:)];

[h,w] = size (node_x);
elems = [];
for idx = 1:w-1
    widx = (idx-1)*h;
    elems = [elems; ...
              widx+[(1:h-1);(2:h);h+(1:h-1)]'; ...
              widx+[(2:h);h+(2:h);h+(1:h-1)]'];
endfor

E = size (elems,1); # No. of simplices
N = size (nodes,1); # No. of vertices
D = size (elems,2); # dimensions+1

```

This creates a N-by-2 matrix `nodes` and a E-by-3 matrix `elems` with values, which define finite element triangles:

⁴ EIDORS - Electrical Impedance Tomography and Diffuse optical Tomography Reconstruction Software
<http://eidors3d.sourceforge.net>

```

nodes(1:7,:)
 1.00 1.00 1.00 1.00 1.00 1.05 1.05 ...
 1.00 1.20 1.50 1.80 2.00 1.00 1.20 ...

elems(1:7,:)
 1   2   3   4   2   3   4 ...
 2   3   4   5   7   8   9 ...
 6   7   8   9   6   7   8 ...

```

Using a first order FEM, we approximate the electrical conductivity distribution in Ω as constant on each simplex (represented by the vector `conductivity`). Based on the finite element geometry, we first calculate a system (or stiffness) matrix for each simplex (represented as 3-by-3 elements on the diagonal of the element-wise system matrix `SE`). Based on `SE` and a N-by-DE connectivity matrix `C`, representing the connections between simplices and vertices, the global connectivity matrix `S` is calculated.

```

## Element conductivity
conductivity = [1*ones(1,16), ...
                2*ones(1,48), 1*ones(1,16)];

## Connectivity matrix
C = sparse ((1:D*E), reshape (elems', ...
                                D*E, 1), 1, D*E, N);

## Calculate system matrix
Siidx = floor ([0:D*E-1]'/D) * D * ...
              ones(1,D) + ones(D*E,1)*(1:D) ;
Sjidx = [1:D*E]*ones (1,D);
Sdata = zeros (D*E,D);
dfact = factorial (D-1);
for j = 1:E
    a = inv ([ones(D,1), ...
              nodes(elems(j,:), :)]);
    const = conductivity(j) * 2 / ...
            dfact / abs (det (a));
    Sdata(D*(j-1)+(1:D),:) = const * ...
        a(2:D,:)' * a(2:D,:);
endfor
## Element-wise system matrix
SE = sparse(Siidx,Sjidx,Sdata);
## Global system matrix
S = C'* SE *C;

```

The system matrix acts like the conductivity S in Ohm's law $SV = I$. Based on the Dirichlet and Neumann boundary conditions, we are able to solve for the voltages at each vertex V .

```

## Dirichlet boundary conditions
D_nodes = [1:5, 51:55];
D_value = [10*ones(1,5), 20*ones(1,5)];

```

```

V = zeros (N,1);
V(D_nodes) = D_value;
idx = 1:N; # vertices without Dirichlet
# boundary condns
idx(D_nodes) = [];

## Neumann boundary conditions. Note that
## N_value must be normalized by the
## boundary length and element conductivity
N_nodes = [];
N_value = [];

Q = zeros (N,1);
Q(N_nodes) = N_value;

V(idx) = S(idx,idx) \ ( Q(idx) - ...
    S(idx,D_nodes) * V(D_nodes));


```

Finally, in order to display the solution, we show each solved voltage value in the z-axis for each simplex vertex. See [Figure 22.6](#).

```

elemx = elems(:,[1,2,3,1])';
xelems = reshape (nodes(elemx, 1), 4, E);
yelems = reshape (nodes(elemx, 2), 4, E);
velems = reshape (V(elemx), 4, E);
plot3 (xelems, yelems, velems,"k");
print "grid.eps";

```



[Figure 22.6](#): Example finite element model the showing triangular elements. The height of each vertex corresponds to the solution value.

23 Numerical Integration

Octave comes with several built-in functions for computing the integral of a function numerically (termed quadrature). These functions all solve 1-dimensional integration problems.

23.1 Functions of One Variable

Octave supports five different adaptive quadrature algorithms for computing the integral

$$\int_a^b f(x)dx$$

of a function f over the interval from a to b . These are

<code>quad</code>	Numerical integration based on Gaussian quadrature.
<code>quadv</code>	Numerical integration using an adaptive vectorized Simpson's rule.
<code>quadl</code>	Numerical integration using an adaptive Lobatto rule.
<code>quadgk</code>	Numerical integration using an adaptive Gauss-Konrod rule.
<code>quadcc</code>	Numerical integration using adaptive Clenshaw-Curtis rules.
	In addition, the following functions are also provided:
<code>integral</code>	A compatibility wrapper function that will choose between <code>quadv</code> and <code>quadgk</code> depending on the integrand and options chosen.
<code>trapz</code> , <code>cumtrapz</code>	Numerical integration of data using the trapezoidal method.

The best quadrature algorithm to use depends on the integrand. If you have empirical data, rather than a function, the choice is `trapz` or `cumtrapz`. If you are uncertain about the characteristics of the integrand, `quadcc` will be the most robust as it can handle discontinuities, singularities, oscillatory functions, and infinite intervals. When the integrand is smooth `quadgk` may be the fastest of the algorithms.

Function	Characteristics
<code>quad</code>	Low accuracy with nonsmooth integrands
<code>quadv</code>	Medium accuracy with smooth integrands
<code>quadl</code>	Medium accuracy with smooth integrands. Slower than <code>quadgk</code> .
<code>quadgk</code>	Medium accuracy ($1e-6 - 1e-9$) with smooth integrands. Handles oscillatory functions and infinite bounds
<code>quadcc</code>	Low to High accuracy with nonsmooth/smooth integrands Handles oscillatory functions, singularities, and infinite bounds

Here is an example of using `quad` to integrate the function

$$f(x) = x \sin(1/x) \sqrt{|1-x|}$$

from $x = 0$ to $x = 3$.

This is a fairly difficult integration (plot the function over the range of integration to see why).

The first step is to define the function:

```
function y = f (x)
    y = x .* sin (1./x) .* sqrt (abs (1 - x));
endfunction
```

Note the use of the ‘dot’ forms of the operators. This is not necessary for the `quad` integrator, but is required by the other integrators. In any case, it makes it much easier to generate a set of points for plotting because it is possible to call the function with a vector argument to produce a vector result.

The second step is to call `quad` with the limits of integration:

```
[q, ier, nfun, err] = quad ("f", 0, 3)
⇒ 1.9819
⇒ 1
⇒ 5061
⇒ 1.1522e-07
```

Although `quad` returns a nonzero value for `ier`, the result is reasonably accurate (to see why, examine what happens to the result if you move the lower bound to 0.1, then 0.01, then 0.001, etc.).

The function “`f`” can be the string name of a function, a function handle, or an inline function. These options make it quite easy to do integration without having to fully define a function in an m-file. For example:

```
# Verify integral (x^3) = x^4/4
f = inline ("x.^3");
quadgk (f, 0, 1)
⇒ 0.25000

# Verify gamma function = (n-1)! for n = 4
f = @(x) x.^3 .* exp (-x);
quadcc (f, 0, Inf)
⇒ 6.0000

q = quad (f, a, b)
q = quad (f, a, b, tol)
q = quad (f, a, b, tol, sing)
[q, ier, nfun, err] = quad (...)
```

Numerically evaluate the integral of `f` from `a` to `b` using Fortran routines from QUADPACK.

`f` is a function handle, inline function, or a string containing the name of the function to evaluate. The function must have the form `y = f (x)` where `y` and `x` are scalars.

`a` and `b` are the lower and upper limits of integration. Either or both may be infinite.

The optional argument `tol` is a vector that specifies the desired accuracy of the result. The first element of the vector is the desired absolute tolerance, and the second element is the desired relative tolerance. To choose a relative test only, set the absolute tolerance to zero. To choose an absolute test only, set the relative tolerance to zero. Both tolerances default to `sqrt (eps)` or approximately 1.5e-8.

The optional argument *sing* is a vector of values at which the integrand is known to be singular.

The result of the integration is returned in *q*.

ier contains an integer error code (0 indicates a successful integration).

nfun indicates the number of function evaluations that were made.

err contains an estimate of the error in the solution.

The function **quad_options** can set other optional parameters for **quad**.

Note: because **quad** is written in Fortran it cannot be called recursively. This prevents its use in integrating over more than one variable by routines **dblquad** and **triplequad**.

See also: [quad_options], page 661, [quadv], page 661, [quadl], page 662, [quadgk], page 663, [quadcc], page 664, [trapz], page 667, [dblquad], page 669, [triplequad], page 670.

```
quad_options ()
val = quad_options (opt)
quad_options (opt, val)
```

Query or set options for the function **quad**.

When called with no arguments, the names of all available options and their current values are displayed.

Given one argument, return the value of the option *opt*.

When called with two arguments, **quad_options** sets the option *opt* to value *val*.

Options include

"absolute tolerance"

Absolute tolerance; may be zero for pure relative error test.

"relative tolerance"

Non-negative relative tolerance. If the absolute tolerance is zero, the relative tolerance must be greater than or equal to **max** (50**eps*, 0.5e-28).

"single precision absolute tolerance"

Absolute tolerance for single precision; may be zero for pure relative error test.

"single precision relative tolerance"

Non-negative relative tolerance for single precision. If the absolute tolerance is zero, the relative tolerance must be greater than or equal to **max** (50**eps*, 0.5e-28).

```
q = quadv (f, a, b)
q = quadv (f, a, b, tol)
q = quadv (f, a, b, tol, trace)
q = quadv (f, a, b, tol, trace, p1, p2, ...)
[q, nfun] = quadv (...)
```

Numerically evaluate the integral of *f* from *a* to *b* using an adaptive Simpson's rule.

f is a function handle, inline function, or string containing the name of the function to evaluate. *quadv* is a vectorized version of *quad* and the function defined by *f* must accept a scalar or vector as input and return a scalar, vector, or array as output.

a and *b* are the lower and upper limits of integration. Both limits must be finite.

The optional argument *tol* defines the absolute tolerance used to stop the adaptation procedure. The default value is 1e-6.

The algorithm used by *quadv* involves recursively subdividing the integration interval and applying Simpson's rule on each subinterval. If *trace* is true then after computing each of these partial integrals display: (1) the total number of function evaluations, (2) the left end of the subinterval, (3) the length of the subinterval, (4) the approximation of the integral over the subinterval.

Additional arguments *p1*, etc., are passed directly to the function *f*. To use default values for *tol* and *trace*, one may pass empty matrices ([]).

The result of the integration is returned in *q*.

The optional output *nfun* indicates the total number of function evaluations performed.

Note: *quadv* is written in Octave's scripting language and can be used recursively in *dblquad* and *triplequad*, unlike the *quad* function.

See also: [\[quad\]](#), page 660, [\[quadl\]](#), page 662, [\[quadgk\]](#), page 663, [\[quadcc\]](#), page 664, [\[trapz\]](#), page 667, [\[dblquad\]](#), page 669, [\[triplequad\]](#), page 670, [\[integral\]](#), page 665, [\[integral2\]](#), page 672, [\[integral3\]](#), page 673.

```
q = quadl (f, a, b)
q = quadl (f, a, b, tol)
q = quadl (f, a, b, tol, trace)
q = quadl (f, a, b, tol, trace, p1, p2, ...)
[q, nfun] = quadl (...)
```

Numerically evaluate the integral of *f* from *a* to *b* using an adaptive Lobatto rule.

f is a function handle, inline function, or string containing the name of the function to evaluate. The function *f* must be vectorized and return a vector of output values when given a vector of input values.

a and *b* are the lower and upper limits of integration. Both limits must be finite.

The optional argument *tol* defines the absolute tolerance with which to perform the integration. The default value is 1e-6.

The algorithm used by *quadl* involves recursively subdividing the integration interval. If *trace* is defined then for each subinterval display: (1) the total number of function evaluations, (2) the left end of the subinterval, (3) the length of the subinterval, (4) the approximation of the integral over the subinterval.

Additional arguments *p1*, etc., are passed directly to the function *f*. To use default values for *tol* and *trace*, one may pass empty matrices ([]).

The result of the integration is returned in *q*.

The optional output *nfun* indicates the total number of function evaluations performed.

Reference: W. Gander and W. Gautschi, *Adaptive Quadrature - Revisited*, BIT Vol. 40, No. 1, March 2000, pp. 84–101. <https://www.inf.ethz.ch/personal/gander/>

See also: [quad], page 660, [quadv], page 661, [quadgk], page 663, [quadcc], page 664, [trapz], page 667, [dblquad], page 669, [triplequad], page 670, [integral], page 665, [integral2], page 672, [integral3], page 673.

```
q = quadgk (f, a, b)
q = quadgk (f, a, b, abstol)
q = quadgk (f, a, b, abstol, trace)
q = quadgk (f, a, b, prop, val, ...)
[q, err] = quadgk (...)
```

Numerically evaluate the integral of f from a to b using adaptive Gauss-Kronrod quadrature.

f is a function handle, inline function, or string containing the name of the function to evaluate. The function f must be vectorized and return a vector of output values when given a vector of input values.

a and b are the lower and upper limits of integration. Either or both limits may be infinite or contain weak end singularities. Variable transformation will be used to treat any infinite intervals and weaken the singularities. For example:

```
quadgk (@(x) 1 ./ (sqrt (x) .* (x + 1)), 0, Inf)
```

Note that the formulation of the integrand uses the element-by-element operator $.$ / and all user functions to `quadgk` should do the same.

The optional argument tol defines the absolute tolerance used to stop the integration procedure. The default value is 1e-10 (1e-5 for single).

The algorithm used by `quadgk` involves subdividing the integration interval and evaluating each subinterval. If $trace$ is true then after computing each of these partial integrals display: (1) the number of subintervals at this step, (2) the current estimate of the error err , (3) the current estimate for the integral q .

Alternatively, properties of `quadgk` can be passed to the function as pairs "prop", `val`. Valid properties are

AbsTol Define the absolute error tolerance for the quadrature. The default absolute tolerance is 1e-10 (1e-5 for single).

RelTol Define the relative error tolerance for the quadrature. The default relative tolerance is 1e-6 (1e-4 for single).

MaxIntervalCount

`quadgk` initially subdivides the interval on which to perform the quadrature into 10 intervals. Subintervals that have an unacceptable error are subdivided and re-evaluated. If the number of subintervals exceeds 650 subintervals at any point then a poor convergence is signaled and the current estimate of the integral is returned. The property "MaxIntervalCount" can be used to alter the number of subintervals that can exist before exiting.

WayPoints

Discontinuities in the first derivative of the function to integrate can be flagged with the "WayPoints" property. This forces the ends of a subinterval to fall on the breakpoints of the function and can result in significantly improved estimation of the error in the integral, faster computation, or both. For example,

```
quadgk (@(x) abs (1 - x.^2), 0, 2, "Waypoints", 1)
```

signals the breakpoint in the integrand at $x = 1$.

Trace If logically true `quadgk` prints information on the convergence of the quadrature at each iteration.

If any of a , b , or $waypoints$ is complex then the quadrature is treated as a contour integral along a piecewise continuous path defined by the above. In this case the integral is assumed to have no edge singularities. For example,

```
quadgk (@(z) log (z), 1+1i, 1+1i, "WayPoints",
[1-1i, -1,-1i, -1+1i])
```

integrates $\log(z)$ along the square defined by $[1+1i, 1-1i, -1-1i, -1+1i]$.

The result of the integration is returned in q .

err is an approximate bound on the error in the integral $\text{abs}(q - I)$, where I is the exact value of the integral.

Reference: L.F. Shampine, "Vectorized adaptive quadrature in MATLAB", Journal of Computational and Applied Mathematics, pp. 131–140, Vol 211, Issue 2, Feb 2008.

See also: [\[quad\]](#), page 660, [\[quadv\]](#), page 661, [\[quadl\]](#), page 662, [\[quadcc\]](#), page 664, [\[trapz\]](#), page 667, [\[dblquad\]](#), page 669, [\[triplequad\]](#), page 670, [\[integral\]](#), page 665, [\[integral2\]](#), page 672, [\[integral3\]](#), page 673.

```
q = quadcc (f, a, b)
q = quadcc (f, a, b, tol)
q = quadcc (f, a, b, tol, sing)
[q, err, nr_points] = quadcc (...)
```

Numerically evaluate the integral of f from a to b using doubly-adaptive Clenshaw-Curtis quadrature.

f is a function handle, inline function, or string containing the name of the function to evaluate. The function f must be vectorized and must return a vector of output values if given a vector of input values. For example,

```
f = @(x) x .* sin (1./x) .* sqrt (abs (1 - x));
```

which uses the element-by-element “dot” form for all operators.

a and b are the lower and upper limits of integration. Either or both limits may be infinite. `quadcc` handles an infinite limit by substituting the variable of integration with $x = \tan(\pi/2*u)$.

The optional argument tol is a 1- or 2-element vector that specifies the desired accuracy of the result. The first element of the vector is the desired absolute tolerance, and the second element is the desired relative tolerance. To choose a relative test only, set the absolute tolerance to zero. To choose an absolute test only, set the relative

tolerance to zero. The default absolute tolerance is 1e-10 (1e-5 for single), and the default relative tolerance is 1e-6 (1e-4 for single).

The optional argument *sing* contains a list of points where the integrand has known singularities, or discontinuities in any of its derivatives, inside the integration interval. For the example above, which has a discontinuity at $x=1$, the call to **quadcc** would be as follows

```
int = quadcc (f, a, b, [], [ 1 ]);
```

The result of the integration is returned in *q*.

err is an estimate of the absolute integration error.

nr_points is the number of points at which the integrand was evaluated.

If the adaptive integration did not converge, the value of *err* will be larger than the requested tolerance. Therefore, it is recommended to verify this value for difficult integrands.

quadcc is capable of dealing with non-numeric values of the integrand such as **NaN** or **Inf**. If the integral diverges, and **quadcc** detects this, then a warning is issued and **Inf** or **-Inf** is returned.

Note: **quadcc** is a general purpose quadrature algorithm and, as such, may be less efficient for a smooth or otherwise well-behaved integrand than other methods such as **quadgk**.

The algorithm uses Clenshaw-Curtis quadrature rules of increasing degree in each interval and bisects the interval if either the function does not appear to be smooth or a rule of maximum degree has been reached. The error estimate is computed from the L2-norm of the difference between two successive interpolations of the integrand over the nodes of the respective quadrature rules.

Implementation Note: For Octave versions ≤ 4.2 , **quadcc** accepted a single tolerance argument which specified the relative tolerance. For versions 4.4 and 5, **quadcc** will issue a warning when called with a single tolerance argument indicating that the meaning of this input has changed from relative tolerance to absolute tolerance. The warning ID for this message is "Octave:quadcc:RelTol-conversion". The warning may be disabled with **warning ("off", "Octave:quadcc:RelTol-conversion")**.

Reference: P. Gonnet, *Increasing the Reliability of Adaptive Quadrature Using Explicit Interpolants*, ACM Transactions on Mathematical Software, Vol. 37, Issue 3, Article No. 3, 2010.

See also: **[quad]**, page 660, **[quadv]**, page 661, **[quadl]**, page 662, **[quadgk]**, page 663, **[trapz]**, page 667, **[dblquad]**, page 669, **[triplequad]**, page 670.

```
q = integral (f, a, b)
q = integral (f, a, b, prop, val, ...)
```

Numerically evaluate the integral of *f* from *a* to *b* using adaptive quadrature.

integral is a wrapper for **quadcc** (general scalar integrands), **quadgk** (integrals with specified integration paths), and **quadv** (array-valued integrands) that is intended to provide MATLAB compatibility. More control of the numerical integration may be achievable by calling the various quadrature functions directly.

f is a function handle, inline function, or string containing the name of the function to evaluate. The function *f* must be vectorized and return a vector of output values when given a vector of input values.

a and *b* are the lower and upper limits of integration. Either or both limits may be infinite or contain weak end singularities. If either or both limits are complex, **integral** will perform a straight line path integral. Alternatively, a complex domain path can be specified using the "Waypoints" option (see below).

Additional optional parameters can be specified using "property", value pairs. Valid properties are:

Waypoints

Specifies points to be used in defining subintervals of the quadrature algorithm, or if *a*, *b*, or *waypoints* are complex then the quadrature is calculated as a contour integral along a piecewise continuous path. For more detail see **quadgk**.

ArrayValued

integral expects *f* to return a scalar value unless *arrayvalued* is specified as true. This option will cause **integral** to perform the integration over the entire array and return *q* with the same dimensions as returned by *f*. For more detail see **quadv**.

AbsTol Define the absolute error tolerance for the quadrature. The default absolute tolerance is 1e-10 (1e-5 for single).

RelTol Define the relative error tolerance for the quadrature. The default relative tolerance is 1e-6 (1e-4 for single).

Adaptive quadrature is used to minimize the estimate of error until the following is satisfied:

$$\text{error} \leq \max(AbsTol, RelTol \cdot |q|)$$

Known MATLAB incompatibilities:

1. If tolerances are left unspecified, and any integration limits or waypoints are of type **single**, then Octave's integral functions automatically reduce the default absolute and relative error tolerances as specified above. If tighter tolerances are desired they must be specified. MATLAB leaves the tighter tolerances appropriate for **double** inputs in place regardless of the class of the integration limits.
2. As a consequence of using **quadcc**, **quadgk**, and **quadv**, certain option combinations are not supported. Currently, "ArrayValued" cannot be combined with "RelTol" or "Waypoints".

See also: [\[integral2\]](#), page 672, [\[integral3\]](#), page 673, [\[quad\]](#), page 660, [\[quadgk\]](#), page 663, [\[quadv\]](#), page 661, [\[quadl\]](#), page 662, [\[quadcc\]](#), page 664, [\[trapz\]](#), page 667, [\[dblquad\]](#), page 669, [\[triplequad\]](#), page 670.

Sometimes one does not have the function, but only the raw (x, y) points from which to perform an integration. This can occur when collecting data in an experiment. The **trapz**

function can integrate these values as shown in the following example where "data" has been collected on the cosine function over the range [0, pi/2).

```
x = 0:0.1:pi/2; # Uniformly spaced points
y = cos (x);
trapz (x, y)
⇒ 0.99666
```

The answer is reasonably close to the exact value of 1. Ordinary quadrature is sensitive to the characteristics of the integrand. Empirical integration depends not just on the integrand, but also on the particular points chosen to represent the function. Repeating the example above with the sine function over the range [0, pi/2) produces far inferior results.

```
x = 0:0.1:pi/2; # Uniformly spaced points
y = sin (x);
trapz (x, y)
⇒ 0.92849
```

However, a slightly different choice of data points can change the result significantly. The same integration, with the same number of points, but spaced differently produces a more accurate answer.

```
x = linspace (0, pi/2, 16); # Uniformly spaced, but including endpoint
y = sin (x);
trapz (x, y)
⇒ 0.99909
```

In general there may be no way of knowing the best distribution of points ahead of time. Or the points may come from an experiment where there is no freedom to select the best distribution. In any case, one must remain aware of this issue when using `trapz`.

```
q = trapz (y)
q = trapz (x, y)
q = trapz (..., dim)
```

Numerically evaluate the integral of points *y* using the trapezoidal method.

`trapz (y)` computes the integral of *y* along the first non-singleton dimension. When the argument *x* is omitted an equally spaced *x* vector with unit spacing (1) is assumed. `trapz (x, y)` evaluates the integral with respect to the spacing in *x* and the values in *y*. This is useful if the points in *y* have been sampled unevenly.

If the optional *dim* argument is given, operate along this dimension.

Application Note: If *x* is not specified then unit spacing will be used. To scale the integral to the correct value you must multiply by the actual spacing value (*deltaX*). As an example, the integral of x^3 over the range [0, 1] is $x^4/4$ or 0.25. The following code uses `trapz` to calculate the integral in three different ways.

```

x = 0:0.1:1;
y = x.^3;
## No scaling
q = trapz (y)
⇒ q = 2.5250
## Approximation to integral by scaling
q * 0.1
⇒ 0.25250
## Same result by specifying x
trapz (x, y)
⇒ 0.25250

```

See also: [\[cumtrapz\]](#), page 668.

```

q = cumtrapz (y)
q = cumtrapz (x, y)
q = cumtrapz (... , dim)

```

Cumulative numerical integration of points y using the trapezoidal method.

`cumtrapz (y)` computes the cumulative integral of y along the first non-singleton dimension. Where `trapz` reports only the overall integral sum, `cumtrapz` reports the current partial sum value at each point of y .

When the argument x is omitted an equally spaced x vector with unit spacing (1) is assumed. `cumtrapz (x, y)` evaluates the integral with respect to the spacing in x and the values in y . This is useful if the points in y have been sampled unevenly.

If the optional dim argument is given, operate along this dimension.

Application Note: If x is not specified then unit spacing will be used. To scale the integral to the correct value you must multiply by the actual spacing value ($\text{delta}X$).

See also: [\[trapz\]](#), page 667, [\[cumsum\]](#), page 517.

23.2 Orthogonal Collocation

```
[r, amat, bmat, q] = colloc (n, "left", "right")
```

Compute derivative and integral weight matrices for orthogonal collocation.

Reference: J. Villadsen, M. L. Michelsen, *Solution of Differential Equation Models by Polynomial Approximation*.

Here is an example of using `colloc` to generate weight matrices for solving the second order differential equation $u' - \alpha u'' = 0$ with the boundary conditions $u(0) = 0$ and $u(1) = 1$.

First, we can generate the weight matrices for n points (including the endpoints of the interval), and incorporate the boundary conditions in the right hand side (for a specific value of α).

```

n = 7;
alpha = 0.1;
[r, a, b] = colloc (n-2, "left", "right");
at = a(2:n-1,2:n-1);
bt = b(2:n-1,2:n-1);
rhs = alpha * b(2:n-1,n) - a(2:n-1,n);

```

Then the solution at the roots r is

```
u = [ 0; (at - alpha * bt) \ rhs; 1]
⇒ [ 0.00; 0.004; 0.01 0.00; 0.12; 0.62; 1.00 ]
```

23.3 Functions of Multiple Variables

Octave includes several functions for computing the integral of functions of multiple variables. This procedure can generally be performed by creating a function that integrates f with respect to x , and then integrates that function with respect to y . This procedure can be performed manually using the following example which integrates the function:

$$f(x, y) = \sin(\pi xy)\sqrt{xy}$$

for x and y between 0 and 1.

Using `quadgk` in the example below, a double integration can be performed. (Note that any of the 1-D quadrature functions can be used in this fashion except for `quad` since it is written in Fortran and cannot be called recursively.)

```
function q = g(y)
q = ones (size (y));
for i = 1:length (y)
    f = @(x) sin (pi*x.*y(i)) .* sqrt (x.*y(i));
    q(i) = quadgk (f, 0, 1);
endfor
endfunction

I = quadgk ("g", 0, 1)
⇒ 0.30022
```

The algorithm above is implemented in the function `dblquad` for integrals over two variables. The 3-D equivalent of this process is implemented in `triplequad` for integrals over three variables. As an example, the result above can be replicated with a call to `dblquad` as shown below.

```
I = dblquad (@(x, y) sin (pi*x.*y) .* sqrt (x.*y), 0, 1, 0, 1)
⇒ 0.30022

dblquad (f, xa, xb, ya, yb)
dblquad (f, xa, xb, ya, yb, tol)
dblquad (f, xa, xb, ya, yb, tol, quadf)
dblquad (f, xa, xb, ya, yb, tol, quadf, ...)
```

Numerically evaluate the double integral of f .

f is a function handle, inline function, or string containing the name of the function to evaluate. The function f must have the form $z = f(x, y)$ where x is a vector and y is a scalar. It should return a vector of the same length and orientation as x .

xa , ya and xb , yb are the lower and upper limits of integration for x and y respectively. The underlying integrator determines whether infinite bounds are accepted.

The optional argument tol defines the absolute tolerance used to integrate each sub-integral. The default value is $1e-6$.

The optional argument *quadf* specifies which underlying integrator function to use. Any choice but *quad* is available and the default is *quadcc*.

Additional arguments, are passed directly to *f*. To use the default value for *tol* or *quadf* one may pass ':' or an empty matrix ([]).

See also: [\[integral2\]](#), page 672, [\[integral3\]](#), page 673, [\[triplequad\]](#), page 670, [\[quad\]](#), page 660, [\[quadv\]](#), page 661, [\[quadl\]](#), page 662, [\[quadgk\]](#), page 663, [\[quadcc\]](#), page 664, [\[trapz\]](#), page 667.

```
triplequad (f, xa, xb, ya, yb, za, zb)
triplequad (f, xa, xb, ya, yb, za, zb, tol)
triplequad (f, xa, xb, ya, yb, za, zb, tol, quadf)
triplequad (f, xa, xb, ya, yb, za, zb, tol, quadf, ...)
```

Numerically evaluate the triple integral of *f*.

f is a function handle, inline function, or string containing the name of the function to evaluate. The function *f* must have the form $w = f(x, y, z)$ where either *x* or *y* is a vector and the remaining inputs are scalars. It should return a vector of the same length and orientation as *x* or *y*.

xa, *ya*, *za* and *xb*, *yb*, *zb* are the lower and upper limits of integration for *x*, *y*, and *z* respectively. The underlying integrator determines whether infinite bounds are accepted.

The optional argument *tol* defines the absolute tolerance used to integrate each sub-integral. The default value is 1e-6.

The optional argument *quadf* specifies which underlying integrator function to use. Any choice but *quad* is available and the default is *quadcc*.

Additional arguments, are passed directly to *f*. To use the default value for *tol* or *quadf* one may pass ':' or an empty matrix ([]).

See also: [\[integral3\]](#), page 673, [\[integral2\]](#), page 672, [\[dblquad\]](#), page 669, [\[quad\]](#), page 660, [\[quadv\]](#), page 661, [\[quadl\]](#), page 662, [\[quadgk\]](#), page 663, [\[quadcc\]](#), page 664, [\[trapz\]](#), page 667.

The recursive algorithm for quadrature presented above is referred to as "iterated". A separate 2-D integration method is implemented in the function *quad2d*. This function performs a "tiled" integration by subdividing the integration domain into rectangular regions and performing separate integrations over those domains. The domains are further subdivided in areas requiring refinement to reach the desired numerical accuracy. For certain functions this method can be faster than the 2-D iteration used in the other functions above.

```
q = quad2d (f, xa, xb, ya, yb)
q = quad2d (f, xa, xb, ya, yb, prop, val, ...)
[q, err, iter] = quad2d (...)
```

Numerically evaluate the two-dimensional integral of *f* using adaptive quadrature over the two-dimensional domain defined by *xa*, *xb*, *ya*, *yb* using tiled integration. Additionally, *ya* and *yb* may be scalar functions of *x*, allowing for the integration over non-rectangular domains.

f is a function handle, inline function, or string containing the name of the function to evaluate. The function f must be of the form $z = f(x, y)$ where x is a vector and y is a scalar. It should return a vector of the same length and orientation as x .

Additional optional parameters can be specified using "property", value pairs. Valid properties are:

AbsTol Define the absolute error tolerance for the quadrature. The default value is 1e-10 (1e-5 for single).

RelTol Define the relative error tolerance for the quadrature. The default value is 1e-6 (1e-4 for single).

MaxFunEvals

The maximum number of function calls to the vectorized function f . The default value is 5000.

Singular Enable/disable transforms to weaken singularities on the edge of the integration domain. The default value is *true*.

Vectorized

Option to disable vectorized integration, forcing Octave to use only scalar inputs when calling the integrand. The default value is *false*.

FailurePlot

If `quad2d` fails to converge to the desired error tolerance before `MaxFunEvals` is reached, a plot of the areas that still need refinement is created. The default value is *false*.

Adaptive quadrature is used to minimize the estimate of error until the following is satisfied:

$$\text{error} \leq \max(AbsTol, RelTol \cdot |q|)$$

The optional output `err` is an approximate bound on the error in the integral `abs(q - I)`, where I is the exact value of the integral. The optional output `iter` is the number of vectorized function calls to the function f that were used.

Example 1 : integrate a rectangular region in x-y plane

```
f = @(x,y) 2*ones (size (x));
q = quad2d (f, 0, 1, 0, 1)
⇒ q = 2
```

The result is a volume, which for this constant-value integrand, is just $\text{Length} * \text{Width} * \text{Height}$.

Example 2 : integrate a triangular region in x-y plane

```
f = @(x,y) 2*ones (size (x));
ymax = @(x) 1 - x;
q = quad2d (f, 0, 1, 0, ymax)
⇒ q = 1
```

The result is a volume, which for this constant-value integrand, is the Triangle Area $\times \text{Height}$ or $1/2 * \text{Base} * \text{Width} * \text{Height}$.

Programming Notes: If there are singularities within the integration region it is best to split the integral and place the singularities on the boundary.

Known MATLAB incompatibility: If tolerances are left unspecified, and any integration limits are of type `single`, then Octave's integral functions automatically reduce the default absolute and relative error tolerances as specified above. If tighter tolerances are desired they must be specified. MATLAB leaves the tighter tolerances appropriate for `double` inputs in place regardless of the class of the integration limits.

Reference: L.F. Shampine, MATLAB program for quadrature in 2D, Applied Mathematics and Computation, pp. 266–274, Vol 1, 2008.

See also: [\[integral2\]](#), page 672, [\[dblquad\]](#), page 669, [\[integral\]](#), page 665, [\[quad\]](#), page 660, [\[quadgk\]](#), page 663, [\[quadv\]](#), page 661, [\[quadl\]](#), page 662, [\[quadcc\]](#), page 664, [\[trapz\]](#), page 667, [\[integral3\]](#), page 673, [\[triplequad\]](#), page 670.

Finally, the functions `integral2` and `integral3` are provided as general 2-D and 3-D integration functions. They will auto-select between iterated and tiled integration methods and, unlike `dblquad` and `triplequad`, will work with non-rectangular integration domains.

```
q = integral2 (f, xa, xb, ya, yb)
q = integral2 (f, xa, xb, ya, yb, prop, val, ...)
[q, err] = integral2 (...)
```

Numerically evaluate the two-dimensional integral of f using adaptive quadrature over the two-dimensional domain defined by xa , xb , ya , yb (scalars may be finite or infinite). Additionally, ya and yb may be scalar functions of x , allowing for integration over non-rectangular domains.

f is a function handle, inline function, or string containing the name of the function to evaluate. The function f must be of the form $z = f(x, y)$ where x is a vector and y is a scalar. It should return a vector of the same length and orientation as x .

Additional optional parameters can be specified using "property", value pairs. Valid properties are:

<code>AbsTol</code>	Define the absolute error tolerance for the quadrature. The default value is 1e-10 (1e-5 for <code>single</code>).
<code>RelTol</code>	Define the relative error tolerance for the quadrature. The default value is 1e-6 (1e-4 for <code>single</code>).
<code>Method</code>	Specify the two-dimensional integration method to be used, with valid options being "auto" (default), "tiled", or "iterated". When using "auto", Octave will choose the "tiled" method unless any of the integration limits are infinite.

`Vectorized`

Enable or disable vectorized integration. A value of `false` forces Octave to use only scalar inputs when calling the integrand, which enables integrands $f(x, y)$ that have not been vectorized and only accept x and y as scalars to be used. The default value is `true`.

Adaptive quadrature is used to minimize the estimate of error until the following is satisfied:

$$\text{error} \leq \max(AbsTol, RelTol \cdot |q|)$$

`err` is an approximate bound on the error in the integral `abs (q - I)`, where I is the exact value of the integral.

Example 1 : integrate a rectangular region in x-y plane

```
f = @(x,y) 2*ones (size (x));
q = integral2 (f, 0, 1, 0, 1)
⇒ q = 2
```

The result is a volume, which for this constant-value integrand, is just $\text{Length} * \text{Width} * \text{Height}$.

Example 2 : integrate a triangular region in x-y plane

```
f = @(x,y) 2*ones (size (x));
ymax = @(x) 1 - x;
q = integral2 (f, 0, 1, 0, ymax)
⇒ q = 1
```

The result is a volume, which for this constant-value integrand, is the Triangle Area \times Height or $1/2 * \text{Base} * \text{Width} * \text{Height}$.

Programming Notes: If there are singularities within the integration region it is best to split the integral and place the singularities on the boundary.

Known MATLAB incompatibility: If tolerances are left unspecified, and any integration limits are of type `single`, then Octave's integral functions automatically reduce the default absolute and relative error tolerances as specified above. If tighter tolerances are desired they must be specified. MATLAB leaves the tighter tolerances appropriate for `double` inputs in place regardless of the class of the integration limits.

Reference: L.F. Shampine, *MATLAB program for quadrature in 2D*, Applied Mathematics and Computation, pp. 266–274, Vol 1, 2008.

See also: [\[quad2d\]](#), page 670, [\[dblquad\]](#), page 669, [\[integral\]](#), page 665, [\[quad\]](#), page 660, [\[quadgk\]](#), page 663, [\[quadv\]](#), page 661, [\[quadl\]](#), page 662, [\[quadcc\]](#), page 664, [\[trapz\]](#), page 667, [\[integral3\]](#), page 673, [\[triplequad\]](#), page 670.

```
q = integral3 (f, xa, xb, ya, yb, za, zb)
q = integral3 (f, xa, xb, ya, yb, za, zb, prop, val, ...)
```

Numerically evaluate the three-dimensional integral of f using adaptive quadrature over the three-dimensional domain defined by xa , xb , ya , yb , za , zb (scalars may be finite or infinite). Additionally, ya and yb may be scalar functions of x and za , and zb maybe be scalar functions of x and y , allowing for integration over non-rectangular domains.

f is a function handle, inline function, or string containing the name of the function to evaluate. The function f must be of the form $z = f(x, y)$ where x is a vector and y is a scalar. It should return a vector of the same length and orientation as x .

Additional optional parameters can be specified using "property", value pairs. Valid properties are:

`AbsTol` Define the absolute error tolerance for the quadrature. The default value is 1e-10 (1e-5 for single).

`RelTol` Define the relative error tolerance for the quadrature. The default value is 1e-6 (1e-4 for single).

Method Specify the two-dimensional integration method to be used, with valid options being "auto" (default), "tiled", or "iterated". When using "auto", Octave will choose the "tiled" method unless any of the integration limits are infinite.

Vectorized

Enable or disable vectorized integration. A value of `false` forces Octave to use only scalar inputs when calling the integrand, which enables integrands $f(x, y)$ that have not been vectorized and only accept x and y as scalars to be used. The default value is `true`.

Adaptive quadrature is used to minimize the estimate of error until the following is satisfied:

$$\text{error} \leq \max(\text{AbsTol}, \text{RelTol} \cdot |q|)$$

`err` is an approximate bound on the error in the integral `abs (q - I)`, where I is the exact value of the integral.

Example 1 : integrate over a rectangular volume

```
f = @(x,y,z) ones (size (x));
q = integral3 (f, 0, 1, 0, 1, 0, 1)
⇒ q = 1.00000
```

For this constant-value integrand, the result is a volume which is just $\text{Length} * \text{Width} * \text{Height}$.

Example 2 : integrate over a spherical volume

```
f = @(x,y) ones (size (x));
ymax = @(x) sqrt (1 - x.^2);
zmax = @(x,y) sqrt (1 - x.^2 - y.^2);
q = integral3 (f, 0, 1, 0, ymax, 0, zmax)
⇒ q = 0.52360
```

For this constant-value integrand, the result is a volume which is 1/8th of a unit sphere or $1/8 * 4/3 * \pi$.

Programming Notes: If there are singularities within the integration region it is best to split the integral and place the singularities on the boundary.

Known MATLAB incompatibility: If tolerances are left unspecified, and any integration limits are of type `single`, then Octave's integral functions automatically reduce the default absolute and relative error tolerances as specified above. If tighter tolerances are desired they must be specified. MATLAB leaves the tighter tolerances appropriate for `double` inputs in place regardless of the class of the integration limits.

Reference: L.F. Shampine, MATLAB program for quadrature in 2D, Applied Mathematics and Computation, pp. 266–274, Vol 1, 2008.

See also: [\[triplequad\]](#), page 670, [\[integral\]](#), page 665, [\[quad\]](#), page 660, [\[quadgk\]](#), page 663, [\[quadv\]](#), page 661, [\[quadl\]](#), page 662, [\[quadcc\]](#), page 664, [\[trapz\]](#), page 667, [\[integral2\]](#), page 672, [\[quad2d\]](#), page 670, [\[dblquad\]](#), page 669.

The above integrations can be fairly slow, and that problem increases exponentially with the dimensionality of the integral. Another possible solution for 2-D integration is to use

Orthogonal Collocation as described in the previous section (see [Section 23.2 \[Orthogonal Collocation\]](#), page 668). The integral of a function $f(x, y)$ for x and y between 0 and 1 can be approximated using n points by

$$\int_0^1 \int_0^1 f(x, y) dx dy \approx \sum_{i=1}^n \sum_{j=1}^n q_i q_j f(r_i, r_j),$$

where q and r is as returned by `colloc (n)`. The generalization to more than two variables is straight forward. The following code computes the studied integral using $n = 8$ points.

```
f = @(x,y) sin (pi*x*y') .* sqrt (x*y');
n = 8;
[t, ~, ~, q] = colloc (n);
I = q'*f(t,t)*q;
⇒ 0.30022
```

It should be noted that the number of points determines the quality of the approximation. If the integration needs to be performed between a and b , instead of 0 and 1, then a change of variables is needed.

24 Differential Equations

Octave has built-in functions for solving ordinary differential equations, and differential-algebraic equations. All solvers are based on reliable ODE routines written in Fortran.

24.1 Ordinary Differential Equations

The function `lsode` can be used to solve ODEs of the form

$$\frac{dx}{dt} = f(x, t)$$

using Hindmarsh's ODE solver LSODE.

```
[x, istate, msg] = lsode (fcn, x_0, t)
[x, istate, msg] = lsode (fcn, x_0, t, t_crit)
```

Ordinary Differential Equation (ODE) solver.

The set of differential equations to solve is

$$\frac{dx}{dt} = f(x, t)$$

with

$$x(t_0) = x_0$$

The solution is returned in the matrix `x`, with each row corresponding to an element of the vector `t`. The first element of `t` should be t_0 and should correspond to the initial state of the system `x_0`, so that the first row of the output is `x_0`.

The first argument, `fcn`, is a string, inline, or function handle that names the function `f` to call to compute the vector of right hand sides for the set of equations. The function must have the form

```
xdot = f (x, t)
```

in which `xdot` and `x` are vectors and `t` is a scalar.

If `fcn` is a two-element string array or a two-element cell array of strings, inline functions, or function handles, the first element names the function `f` described above, and the second element names a function to compute the Jacobian of `f`. The Jacobian function must have the form

```
jac = j (x, t)
```

in which `jac` is the matrix of partial derivatives

$$J = \frac{\partial f_i}{\partial x_j} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} & \dots & \frac{\partial f_1}{\partial x_N} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} & \dots & \frac{\partial f_2}{\partial x_N} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f_3}{\partial x_1} & \frac{\partial f_3}{\partial x_2} & \dots & \frac{\partial f_3}{\partial x_N} \end{bmatrix}$$

The second argument specifies the initial state of the system `x_0`. The third argument is a vector, `t`, specifying the time values for which a solution is sought.

The fourth argument is optional, and may be used to specify a set of times that the ODE solver should not integrate past. It is useful for avoiding difficulties with singularities and points where there is a discontinuity in the derivative.

After a successful computation, the value of *istate* will be 2 (consistent with the Fortran version of LSODE).

If the computation is not successful, *istate* will be something other than 2 and *msg* will contain additional information.

You can use the function `lsode_options` to set optional parameters for `lsode`.

See also: [\[daspk\]](#), page 688, [\[dassl\]](#), page 691, [\[dasrt\]](#), page 694.

```
lsode_options ()
val = lsode_options (opt)
lsode_options (opt, val)
```

Query or set options for the function `lsode`.

When called with no arguments, the names of all available options and their current values are displayed.

Given one argument, return the value of the option *opt*.

When called with two arguments, `lsode_options` sets the option *opt* to value *val*.

Options include

"absolute tolerance"

Absolute tolerance. May be either vector or scalar. If a vector, it must match the dimension of the state vector.

"relative tolerance"

Relative tolerance parameter. Unlike the absolute tolerance, this parameter may only be a scalar.

The local error test applied at each integration step is

```
abs (local error in x(i)) <= ...
      rtol * abs (y(i)) + atol(i)
```

"integration method"

A string specifying the method of integration to use to solve the ODE system. Valid values are

"adams"

"non-stiff"

No Jacobian used (even if it is available).

"bdf"

"stiff" Use stiff backward differentiation formula (BDF) method. If a function to compute the Jacobian is not supplied, `lsode` will compute a finite difference approximation of the Jacobian matrix.

"initial step size"

The step size to be attempted on the first step (default is determined automatically).

"maximum order"

Restrict the maximum order of the solution method. If using the Adams method, this option must be between 1 and 12. Otherwise, it must be between 1 and 5, inclusive.

"maximum step size"

Setting the maximum stepsize will avoid passing over very large regions (default is not specified).

"minimum step size"

The minimum absolute step size allowed (default is 0).

"step limit"

Maximum number of steps allowed (default is 100000).

Here is an example of solving a set of three differential equations using `lsode`. Given the function

```
## oregonator differential equation
function xdot = f (x, t)

    xdot = zeros (3,1);

    xdot(1) = 77.27 * (x(2) - x(1)*x(2) + x(1) ...
        - 8.375e-06*x(1)^2);
    xdot(2) = (x(3) - x(1)*x(2) - x(2)) / 77.27;
    xdot(3) = 0.161*(x(1) - x(3));

endfunction
```

and the initial condition `x0 = [4; 1.1; 4]`, the set of equations can be integrated using the command

```
t = linspace (0, 500, 1000);

y = lsode ("f", x0, t);
```

If you try this, you will see that the value of the result changes dramatically between $t = 0$ and 5, and again around $t = 305$. A more efficient set of output points might be

```
t = [0, logspace(-1, log10(303), 150), ...
    logspace(log10(304), log10(500), 150)];
```

See Alan C. Hindmarsh, *ODEPACK, A Systematized Collection of ODE Solvers*, in Scientific Computing, R. S. Stepleman, editor, (1983) for more information about the inner workings of `lsode`.

An m-file for the differential equation used above is included with the Octave distribution in the examples directory under the name `oregonator.m`.

24.1.1 Matlab-compatible solvers

Octave also provides a set of solvers for initial value problems for Ordinary Differential Equations that have a MATLAB-compatible interface. The options for this class of methods are set using the functions.

- `odeset`

- **odeget**

Currently implemented solvers are:

- Runge-Kutta methods

- **ode23** integrates a system of non-stiff ordinary differential equations (ODEs) or index-1 differential-algebraic equations (DAEs). It uses the third-order Bogacki-Shampine method and adapts the local step size in order to satisfy a user-specified tolerance. The solver requires three function evaluations per integration step.
- **ode45** integrates a system of non-stiff ODEs (or index-1 DAEs) using the high-order, variable-step Dormand-Prince method. It requires six function evaluations per integration step, but may take larger steps on smooth problems than **ode23**: potentially offering improved efficiency at smaller tolerances.

- Linear multistep methods

- **ode15s** integrates a system of stiff ODEs (or index-1 DAEs) using a variable step, variable order method based on Backward Difference Formulas (BDF).
- **ode15i** integrates a system of fully-implicit ODEs (or index-1 DAEs) using the same variable step, variable order method as **ode15s**. The function **decic** can be used to compute consistent initial conditions.

```
[t, y] = ode45 (fun, trange, init)
[t, y] = ode45 (fun, trange, init, ode_opt)
[t, y, te, ye, ie] = ode45 (...)
solution = ode45 (...)
ode45 (...)
```

Solve a set of non-stiff Ordinary Differential Equations (non-stiff ODEs) with the well known explicit Dormand-Prince method of order 4.

fun is a function handle, inline function, or string containing the name of the function that defines the ODE: $y' = f(t, y)$. The function must accept two inputs where the first is time *t* and the second is a column vector of unknowns *y*.

trange specifies the time interval over which the ODE will be evaluated. Typically, it is a two-element vector specifying the initial and final times ($[t_{\text{init}}, t_{\text{final}}]$). If there are more than two elements then the solution will also be evaluated at these intermediate time instances.

By default, **ode45** uses an adaptive timestep with the **integrate_adaptive** algorithm. The tolerance for the timestep computation may be changed by using the options "RelTol" and "AbsTol".

init contains the initial value for the unknowns. If it is a row vector then the solution *y* will be a matrix in which each column is the solution for the corresponding initial value in *init*.

The optional fourth argument *ode_opt* specifies non-default options to the ODE solver. It is a structure generated by **odeset**.

The function typically returns two outputs. Variable *t* is a column vector and contains the times where the solution was found. The output *y* is a matrix in which each column refers to a different unknown of the problem and each row corresponds to a time in *t*.

The output can also be returned as a structure *solution* which has a field *x* containing a row vector of times where the solution was evaluated and a field *y* containing the solution matrix such that each column corresponds to a time in *x*. Use **fieldnames** (*solution*) to see the other fields and additional information returned. If no output arguments are requested, and no **OutputFcn** is specified in *ode_opt*, then the **OutputFcn** is set to **odeplot** and the results of the solver are plotted immediately. If using the "Events" option then three additional outputs may be returned. *te* holds the time when an Event function returned a zero. *ye* holds the value of the solution at time *te*. *ie* contains an index indicating which Event function was triggered in the case of multiple Event functions.

Example: Solve the Van der Pol equation

```
fvdp = @(t,y) [y(2); (1 - y(1)^2) * y(2) - y(1)];
[t,y] = ode45 (fvdp, [0, 20], [2, 0]);
```

See also: [\[odeset\]](#), page 685, [\[odeget\]](#), page 687, [\[ode23\]](#), page 681, [\[ode15s\]](#), page 682.

```
[t, y] = ode23 (fun, trange, init)
[t, y] = ode23 (fun, trange, init, ode_opt)
[t, y, te, ye, ie] = ode23 (...)
solution = ode23 (...)
ode23 (...)
```

Solve a set of non-stiff Ordinary Differential Equations (non-stiff ODEs) with the well known explicit Bogacki-Shampine method of order 3.

fun is a function handle, inline function, or string containing the name of the function that defines the ODE: $y' = f(t, y)$. The function must accept two inputs where the first is time *t* and the second is a column vector of unknowns *y*.

trange specifies the time interval over which the ODE will be evaluated. Typically, it is a two-element vector specifying the initial and final times ([*tinit*, *tfinal*]). If there are more than two elements then the solution will also be evaluated at these intermediate time instances.

By default, **ode23** uses an adaptive timestep with the **integrate_adaptive** algorithm. The tolerance for the timestep computation may be changed by using the options "RelTol" and "AbsTol".

init contains the initial value for the unknowns. If it is a row vector then the solution *y* will be a matrix in which each column is the solution for the corresponding initial value in *init*.

The optional fourth argument *ode_opt* specifies non-default options to the ODE solver. It is a structure generated by [odeset](#).

The function typically returns two outputs. Variable *t* is a column vector and contains the times where the solution was found. The output *y* is a matrix in which each column refers to a different unknown of the problem and each row corresponds to a time in *t*.

The output can also be returned as a structure *solution* which has a field *x* containing a row vector of times where the solution was evaluated and a field *y* containing the solution matrix such that each column corresponds to a time in *x*. Use **fieldnames** (*solution*) to see the other fields and additional information returned.

If no output arguments are requested, and no `OutputFcn` is specified in `ode_opt`, then the `OutputFcn` is set to `odeplot` and the results of the solver are plotted immediately. If using the "Events" option then three additional outputs may be returned. `te` holds the time when an Event function returned a zero. `ye` holds the value of the solution at time `te`. `ie` contains an index indicating which Event function was triggered in the case of multiple Event functions.

Example: Solve the Van der Pol equation

```
fvpd = @(t,y) [y(2); (1 - y(1)^2) * y(2) - y(1)];
[t,y] = ode23 (fvpd, [0, 20], [2, 0]);
```

Reference: For the definition of this method see https://en.wikipedia.org/wiki/List_of_Runge-Kutta_methods.

See also: `[odeset]`, page 685, `[odeget]`, page 687, `[ode45]`, page 680, `[ode15s]`, page 682.

```
[t, y] = ode15s (fun, trange, y0)
[t, y] = ode15s (fun, trange, y0, ode_opt)
[t, y, te, ye, ie] = ode15s (...)

solution = ode15s (...)

ode15s (...)
```

Solve a set of stiff Ordinary Differential Equations (ODEs) or stiff semi-explicit index 1 Differential Algebraic Equations (DAEs).

`ode15s` uses a variable step, variable order BDF (Backward Differentiation Formula) method that ranges from order 1 to 5.

`fun` is a function handle, inline function, or string containing the name of the function that defines the ODE: $y' = f(t, y)$. The function must accept two inputs where the first is time t and the second is a column vector of unknowns y .

`trange` specifies the time interval over which the ODE will be evaluated. Typically, it is a two-element vector specifying the initial and final times (`[tinit, tfinal]`). If there are more than two elements then the solution will also be evaluated at these intermediate time instances.

`init` contains the initial value for the unknowns. If it is a row vector then the solution `y` will be a matrix in which each column is the solution for the corresponding initial value in `init`.

The optional fourth argument `ode_opt` specifies non-default options to the ODE solver. It is a structure generated by `odeset`.

The function typically returns two outputs. Variable `t` is a column vector and contains the times where the solution was found. The output `y` is a matrix in which each column refers to a different unknown of the problem and each row corresponds to a time in `t`.

The output can also be returned as a structure `solution` which has a field `x` containing a row vector of times where the solution was evaluated and a field `y` containing the solution matrix such that each column corresponds to a time in `x`. Use `fieldnames (solution)` to see the other fields and additional information returned.

If no output arguments are requested, and no `OutputFcn` is specified in `ode_opt`, then the `OutputFcn` is set to `odeplot` and the results of the solver are plotted immediately.

If using the "Events" option then three additional outputs may be returned. *te* holds the time when an Event function returned a zero. *ye* holds the value of the solution at time *te*. *ie* contains an index indicating which Event function was triggered in the case of multiple Event functions.

Example: Solve Robertson's equations:

```
function r = robertson_dae (t, y)
    r = [ -0.04*y(1) + 1e4*y(2)*y(3)
          0.04*y(1) - 1e4*y(2)*y(3) - 3e7*y(2)^2
          y(1) + y(2) + y(3) - 1 ];
endfunction
opt = odeset ("Mass", [1 0 0; 0 1 0; 0 0 0], "MStateDependence", "none");
[t,y] = ode15s (@robertson_dae, [0, 1e3], [1; 0; 0], opt);
```

See also: [\[decic\]](#), page 684, [\[odeset\]](#), page 685, [\[odeget\]](#), page 687, [\[ode23\]](#), page 681, [\[ode45\]](#), page 680.

```
[t, y] = ode15i (fun, trange, y0, yp0)
[t, y] = ode15i (fun, trange, y0, yp0, ode_opt)
[t, y, te, ye, ie] = ode15i (...)
```

solution = *ode15i* (...)

ode15i (...)

Solve a set of fully-implicit Ordinary Differential Equations (ODEs) or index 1 Differential Algebraic Equations (DAEs).

ode15i uses a variable step, variable order BDF (Backward Differentiation Formula) method that ranges from order 1 to 5.

fun is a function handle, inline function, or string containing the name of the function that defines the ODE: $0 = f(t, y, yp)$. The function must accept three inputs where the first is time *t*, the second is the function value *y* (a column vector), and the third is the derivative value *yp* (a column vector).

trange specifies the time interval over which the ODE will be evaluated. Typically, it is a two-element vector specifying the initial and final times (*[tinit, tfinal]*). If there are more than two elements then the solution will also be evaluated at these intermediate time instances.

y0 and *yp0* contain the initial values for the unknowns *y* and *yp*. If they are row vectors then the solution *y* will be a matrix in which each column is the solution for the corresponding initial value in *y0* and *yp0*.

y0 and *yp0* must be consistent initial conditions, meaning that $f(t, y_0, y_{p0}) = 0$ is satisfied. The function *decic* may be used to compute consistent initial conditions given initial guesses.

The optional fifth argument *ode_opt* specifies non-default options to the ODE solver. It is a structure generated by *odeset*.

The function typically returns two outputs. Variable *t* is a column vector and contains the times where the solution was found. The output *y* is a matrix in which each column refers to a different unknown of the problem and each row corresponds to a time in *t*.

The output can also be returned as a structure *solution* which has a field *x* containing a row vector of times where the solution was evaluated and a field *y* con-

taining the solution matrix such that each column corresponds to a time in `x`. Use `fieldnames (solution)` to see the other fields and additional information returned. If no output arguments are requested, and no `OutputFcn` is specified in `ode_opt`, then the `OutputFcn` is set to `odeplot` and the results of the solver are plotted immediately. If using the "Events" option then three additional outputs may be returned. `te` holds the time when an Event function returned a zero. `ye` holds the value of the solution at time `te`. `ie` contains an index indicating which Event function was triggered in the case of multiple Event functions.

Example: Solve Robertson's equations:

```
function r = robertson_dae (t, y, yp)
  r = [ -(yp(1) + 0.04*y(1) - 1e4*y(2)*y(3))
        -(yp(2) - 0.04*y(1) + 1e4*y(2)*y(3) + 3e7*y(2)^2)
        y(1) + y(2) + y(3) - 1 ];
endfunction
[t,y] = ode15i (@robertson_dae, [0, 1e3], [1; 0; 0], [-1e-4; 1e-4; 0]);
```

See also: `[decic]`, page 684, `[odeset]`, page 685, `[odeget]`, page 687.

```
[y0_new, yp0_new] = decic (fun, t0, y0, fixed_y0, fixed_yp0)
[y0_new, yp0_new] = decic (fun, t0, y0, fixed_y0, yp0, fixed_yp0,
                           options)
[y0_new, yp0_new, resnorm] = decic (...)
```

Compute consistent implicit ODE initial conditions `y0_new` and `yp0_new` given initial guesses `y0` and `yp0`.

A maximum of `length (y0)` components between `fixed_y0` and `fixed_yp0` may be chosen as fixed values.

`fun` is a function handle. The function must accept three inputs where the first is time `t`, the second is a column vector of unknowns `y`, and the third is a column vector of unknowns `yp`.

`t0` is the initial time such that `fun(t0, y0_new, yp0_new) = 0`, specified as a scalar.

`y0` is a vector used as the initial guess for `y`.

`fixed_y0` is a vector which specifies the components of `y0` to hold fixed. Choose a maximum of `length (y0)` components between `fixed_y0` and `fixed_yp0` as fixed values. Set `fixed_y0(i)` component to 1 if you want to fix the value of `y0(i)`. Set `fixed_y0(i)` component to 0 if you want to allow the value of `y0(i)` to change.

`yp0` is a vector used as the initial guess for `yp`.

`fixed_yp0` is a vector which specifies the components of `yp0` to hold fixed. Choose a maximum of `length (yp0)` components between `fixed_y0` and `fixed_yp0` as fixed values. Set `fixed_yp0(i)` component to 1 if you want to fix the value of `yp0(i)`. Set `fixed_yp0(i)` component to 0 if you want to allow the value of `yp0(i)` to change.

The optional seventh argument `options` is a structure array. Use `odeset` to generate this structure. The relevant options are `RelTol` and `AbsTol` which specify the error thresholds used to compute the initial conditions.

The function typically returns two outputs. Variable `y0_new` is a column vector and contains the consistent initial value of `y`. The output `yp0_new` is a column vector and contains the consistent initial value of `yp`.

The optional third output `resnorm` is the norm of the vector of residuals. If `resnorm` is small, `decic` has successfully computed the initial conditions. If the value of `resnorm` is large, use `RelTol` and `AbsTol` to adjust it.

Example: Compute initial conditions for Robertson's equations:

```
function r = robertson_dae (t, y, yp)
    r = [ -(yp(1) + 0.04*y(1) - 1e4*y(2)*y(3))
          -(yp(2) - 0.04*y(1) + 1e4*y(2)*y(3) + 3e7*y(2)^2)
          y(1) + y(2) + y(3) - 1 ];
endfunction
[y0_new,yp0_new] = decic (@robertson_dae, 0, [1; 0; 0], [1; 1; 0],
[-1e-4; 1; 0], [0; 0; 0]);
```

See also: [\[ode15i\]](#), page 683, [\[odeset\]](#), page 685.

```
odeselect = odeset ()
odeselect = odeset ("field1", value1, "field2", value2, ...)
odeselect = odeset (oldstruct, "field1", value1, "field2", value2, ...)
odeselect = odeset (oldstruct, newstruct)
odeset ()
```

Create or modify an ODE options structure.

When called with no input argument and one output argument, return a new ODE options structure that contains all possible fields initialized to their default values. If no output argument is requested, display a list of the common ODE solver options along with their default value.

If called with name-value input argument pairs "field1", "value1", "field2", "value2", ... return a new ODE options structure with all the most common option fields initialized, and set the values of the fields "field1", "field2", ... to the values `value1`, `value2`,

If called with an input structure `oldstruct` then overwrite the values of the options "field1", "field2", ... with new values `value1`, `value2`, ... and return the modified structure.

When called with two input ODE options structures `oldstruct` and `newstruct` overwrite all values from the structure `oldstruct` with new values from the structure `newstruct`. Empty values in `newstruct` will not overwrite values in `oldstruct`.

The most commonly used ODE options, which are always assigned a value by `odeset`, are the following:

`AbsTol`: positive scalar | vector, def. `1e-6`

Absolute error tolerance.

`BDF`: {"off"} | "on"

Use BDF formulas in implicit multistep methods. Note: This option is not yet implemented.

`Events`: function_handle

Event function. An event function must have the form `[value, isterminal, direction] = my_events_f (t, y)`

`InitialSlope`: vector

Consistent initial slope vector for DAE solvers.

InitialStep: positive scalar
 Initial time step size.

Jacobian: matrix | function_handle
 Jacobian matrix, specified as a constant matrix or a function of time and state.

JConstant: {"off"} | "on"
 Specify whether the Jacobian is a constant matrix or depends on the state.

JPattern: sparse matrix
 If the Jacobian matrix is sparse and non-constant but maintains a constant sparsity pattern, specify the sparsity pattern.

Mass: matrix | function_handle
 Mass matrix, specified as a constant matrix or a function of time and state.

MassSingular: {"maybe"} | "yes" | "on"
 Specify whether the mass matrix is singular.

MaxOrder: {5} | 4 | 3 | 2 | 1
 Maximum order of formula.

MaxStep: positive scalar
 Maximum time step value.

MStateDependence: {"weak"} | "none" | "strong"
 Specify whether the mass matrix depends on the state or only on time.

MvPattern: sparse matrix
 If the mass matrix is sparse and non-constant but maintains a constant sparsity pattern, specify the sparsity pattern. *Note:* This option is not yet implemented.

NonNegative: scalar | vector
 Specify elements of the state vector that are expected to remain non-negative during the simulation.

NormControl: {"off"} | "on"
 Control error relative to the 2-norm of the solution, rather than its absolute value.

OutputFcn: function_handle
 Function to monitor the state during the simulation. For the form of the function to use see `odeplot`.

OutputSel: scalar | vector
 Indices of elements of the state vector to be passed to the output monitoring function.

Refine: positive scalar
 Specify whether output should be returned only at the end of each time step or also at intermediate time instances. The value should be a scalar

indicating the number of equally spaced time points to use within each timestep at which to return output. *Note:* This option is not yet implemented.

RelTol: positive scalar

Relative error tolerance.

Stats: {"off"} | "on"

Print solver statistics after simulation.

Vectorized: {"off"} | "on"

Specify whether `odefun` can be passed multiple values of the state at once.

Field names that are not in the above list are also accepted and added to the result structure.

See also: [odeget], page 687.

`val = odeget (ode_opt, field)`

`val = odeget (ode_opt, field, default)`

Query the value of the property `field` in the ODE options structure `ode_opt`.

If called with two input arguments and the first input argument `ode_opt` is an ODE option structure and the second input argument `field` is a string specifying an option name, then return the option value `val` corresponding to `field` from `ode_opt`.

If called with an optional third input argument, and `field` is not set in the structure `ode_opt`, then return the default value `default` instead.

See also: [odeset], page 685.

`stop_solve = odeplot (t, y, flag)`

Open a new figure window and plot the solution of an ode problem at each time step during the integration.

The types and values of the input parameters `t` and `y` depend on the input `flag` that is of type string. Valid values of `flag` are:

"init" The input `t` must be a column vector of length 2 with the first and last time step (`[tffirst tlast]`). The input `y` contains the initial conditions for the ode problem (`y0`).

"" The input `t` must be a scalar double specifying the time for which the solution in input `y` was calculated.

"done" The inputs should be empty, but are ignored if they are present.

`odeplot` always returns false, i.e., don't stop the ode solver.

Example: solve an anonymous implementation of the "Van der Pol" equation and display the results while solving.

```
fvdp = @(t,y) [y(2); (1 - y(1)^2) * y(2) - y(1)];
```

```
opt = odeset ("OutputFcn", @odeplot, "RelTol", 1e-6);
sol = ode45 (fvdp, [0 20], [2 0], opt);
```

Background Information: This function is called by an ode solver function if it was specified in the "OutputFcn" property of an options structure created with `odeset`. The ode solver will initially call the function with the syntax `odeplot ([tfirst, tlast], y0, "init")`. The function initializes internal variables, creates a new figure window, and sets the x limits of the plot. Subsequently, at each time step during the integration the ode solver calls `odeplot (t, y, [])`. At the end of the solution the ode solver calls `odeplot ([] , [] , "done")` so that `odeplot` can perform any clean-up actions required.

See also: [\[odeset\]](#), page 685, [\[odeget\]](#), page 687, [\[ode23\]](#), page 681, [\[ode45\]](#), page 680.

24.2 Differential-Algebraic Equations

The function `daspk` can be used to solve DAEs of the form

$$0 = f(\dot{x}, x, t), \quad x(t=0) = x_0, \dot{x}(t=0) = \dot{x}_0$$

where $\dot{x} = \frac{dx}{dt}$ is the derivative of x . The equation is solved using Petzold's DAE solver DASPK.

[x, xdot, istate, msg] = daspk (fcn, x_0, xdot_0, t, t_crit)

Solve a set of differential-algebraic equations.

`daspk` solves the set of equations

$$0 = f(x, \dot{x}, t)$$

with

$$x(t_0) = x_0, \dot{x}(t_0) = \dot{x}_0$$

The solution is returned in the matrices `x` and `xdot`, with each row in the result matrices corresponding to one of the elements in the vector `t`. The first element of `t` should be t_0 and correspond to the initial state of the system `x_0` and its derivative `xdot_0`, so that the first row of the output `x` is `x_0` and the first row of the output `xdot` is `xdot_0`.

The first argument, `fcn`, is a string, inline, or function handle that names the function `f` to call to compute the vector of residuals for the set of equations. It must have the form

```
res = f (x, xdot, t)
```

in which `x`, `xdot`, and `res` are vectors, and `t` is a scalar.

If `fcn` is a two-element string array or a two-element cell array of strings, inline functions, or function handles, the first element names the function `f` described above, and the second element names a function to compute the modified Jacobian

$$J = \frac{\partial f}{\partial x} + c \frac{\partial f}{\partial \dot{x}}$$

The modified Jacobian function must have the form

```
jac = j (x, xdot, t, c)
```

The second and third arguments to `daspk` specify the initial condition of the states and their derivatives, and the fourth argument specifies a vector of output times at which the solution is desired, including the time corresponding to the initial condition.

The set of initial states and derivatives are not strictly required to be consistent. If they are not consistent, you must use the `daspk_options` function to provide additional information so that `daspk` can compute a consistent starting point.

The fifth argument is optional, and may be used to specify a set of times that the DAE solver should not integrate past. It is useful for avoiding difficulties with singularities and points where there is a discontinuity in the derivative.

After a successful computation, the value of `istate` will be greater than zero (consistent with the Fortran version of DASPK).

If the computation is not successful, the value of `istate` will be less than zero and `msg` will contain additional information.

You can use the function `daspk_options` to set optional parameters for `daspk`.

See also: [dassl], page 691.

```
daspk_options ()
val = daspk_options (opt)
daspk_options (opt, val)
```

Query or set options for the function `daspk`.

When called with no arguments, the names of all available options and their current values are displayed.

Given one argument, return the value of the option `opt`.

When called with two arguments, `daspk_options` sets the option `opt` to value `val`.

Options include

"absolute tolerance"

Absolute tolerance. May be either vector or scalar. If a vector, it must match the dimension of the state vector, and the relative tolerance must also be a vector of the same length.

"relative tolerance"

Relative tolerance. May be either vector or scalar. If a vector, it must match the dimension of the state vector, and the absolute tolerance must also be a vector of the same length.

The local error test applied at each integration step is

```
abs (local error in x(i))
<= rtol(i) * abs (Y(i)) + atol(i)
```

"compute consistent initial condition"

Denoting the differential variables in the state vector by '`Y_d`' and the algebraic variables by '`Y_a`', `ddaspk` can solve one of two initialization problems:

1. Given `Y_d`, calculate `Y_a` and `Y'_d`
2. Given `Y'`, calculate `Y`.

In either case, initial values for the given components are input, and initial guesses for the unknown components must also be provided as input. Set this option to 1 to solve the first problem, or 2 to solve the second (the default is 0, so you must provide a set of initial conditions that are consistent).

If this option is set to a nonzero value, you must also set the "**algebraic variables**" option to declare which variables in the problem are algebraic.

"use initial condition heuristics"

Set to a nonzero value to use the initial condition heuristics options described below.

"initial condition heuristics"

A vector of the following parameters that can be used to control the initial condition calculation.

MXNIT Maximum number of Newton iterations (default is 5).

MXNJ Maximum number of Jacobian evaluations (default is 6).

MXNH Maximum number of values of the artificial stepsize parameter to be tried if the "**compute consistent initial condition**" option has been set to 1 (default is 5).

Note that the maximum total number of Newton iterations allowed is $\text{MXNIT} \times \text{MXNJ} \times \text{MXNH}$ if the "**compute consistent initial condition**" option has been set to 1 and $\text{MXNIT} \times \text{MXNJ}$ if it is set to 2.

LSOFF Set to a nonzero value to disable the linesearch algorithm (default is 0).

STPTOL Minimum scaled step in linesearch algorithm (default is $\text{eps}^{(2/3)}$).

EPINIT Swing factor in the Newton iteration convergence test. The test is applied to the residual vector, premultiplied by the approximate Jacobian. For convergence, the weighted RMS norm of this vector (scaled by the error weights) must be less than $\text{EPINIT} \times \text{EPCON}$, where $\text{EPCON} = 0.33$ is the analogous test constant used in the time steps. The default is $\text{EPINIT} = 0.01$.

"print initial condition info"

Set this option to a nonzero value to display detailed information about the initial condition calculation (default is 0).

"exclude algebraic variables from error test"

Set to a nonzero value to exclude algebraic variables from the error test. You must also set the "**algebraic variables**" option to declare which variables in the problem are algebraic (default is 0).

"algebraic variables"

A vector of the same length as the state vector. A nonzero element indicates that the corresponding element of the state vector is an algebraic variable (i.e., its derivative does not appear explicitly in the equation set).

This option is required by the "compute consistent initial condition" and "exclude algebraic variables from error test" options.

"enforce inequality constraints"

Set to one of the following values to enforce the inequality constraints specified by the "inequality constraint types" option (default is 0).

1. To have constraint checking only in the initial condition calculation.
2. To enforce constraint checking during the integration.
3. To enforce both options 1 and 2.

"inequality constraint types"

A vector of the same length as the state specifying the type of inequality constraint. Each element of the vector corresponds to an element of the state and should be assigned one of the following codes

- | | |
|----|--------------------------------|
| -2 | Less than zero. |
| -1 | Less than or equal to zero. |
| 0 | Not constrained. |
| 1 | Greater than or equal to zero. |
| 2 | Greater than zero. |

This option only has an effect if the "enforce inequality constraints" option is nonzero.

"initial step size"

Differential-algebraic problems may occasionally suffer from severe scaling difficulties on the first step. If you know a great deal about the scaling of your problem, you can help to alleviate this problem by specifying an initial stepsize (default is computed automatically).

"maximum order"

Restrict the maximum order of the solution method. This option must be between 1 and 5, inclusive (default is 5).

"maximum step size"

Setting the maximum stepsize will avoid passing over very large regions (default is not specified).

Octave also includes DASSL, an earlier version of DASPK, and DASRT, which can be used to solve DAEs with constraints (stopping conditions).

```
[x, xdot, istate, msg] = dassl (fcn, x_0, xdot_0, t, t_crit)
```

Solve a set of differential-algebraic equations.

dassl solves the set of equations

$$0 = f(x, \dot{x}, t)$$

with

$$x(t_0) = x_0, \dot{x}(t_0) = \dot{x}_0$$

The solution is returned in the matrices *x* and *xdot*, with each row in the result matrices corresponding to one of the elements in the vector *t*. The first element of *t* should be *t*₀ and correspond to the initial state of the system *x*₀ and its derivative *xdot*₀, so that the first row of the output *x* is *x*₀ and the first row of the output *xdot* is *xdot*₀.

The first argument, *fcn*, is a string, inline, or function handle that names the function *f* to call to compute the vector of residuals for the set of equations. It must have the form

```
res = f (x, xdot, t)
```

in which *x*, *xdot*, and *res* are vectors, and *t* is a scalar.

If *fcn* is a two-element string array or a two-element cell array of strings, inline functions, or function handles, the first element names the function *f* described above, and the second element names a function to compute the modified Jacobian

$$J = \frac{\partial f}{\partial x} + c \frac{\partial f}{\partial \dot{x}}$$

The modified Jacobian function must have the form

```
jac = j (x, xdot, t, c)
```

The second and third arguments to **dassl** specify the initial condition of the states and their derivatives, and the fourth argument specifies a vector of output times at which the solution is desired, including the time corresponding to the initial condition.

The set of initial states and derivatives are not strictly required to be consistent. In practice, however, DASSL is not very good at determining a consistent set for you, so it is best if you ensure that the initial values result in the function evaluating to zero.

The fifth argument is optional, and may be used to specify a set of times that the DAE solver should not integrate past. It is useful for avoiding difficulties with singularities and points where there is a discontinuity in the derivative.

After a successful computation, the value of *istate* will be greater than zero (consistent with the Fortran version of DASSL).

If the computation is not successful, the value of *istate* will be less than zero and *msg* will contain additional information.

You can use the function **dassl_options** to set optional parameters for **dassl**.

See also: [\[daspk\]](#), page 688, [\[dasrt\]](#), page 694, [\[lsode\]](#), page 677.

```
dassl_options ()
val = dassl_options (opt)
dassl_options (opt, val)
```

Query or set options for the function `dassl`.

When called with no arguments, the names of all available options and their current values are displayed.

Given one argument, return the value of the option `opt`.

When called with two arguments, `dassl_options` sets the option `opt` to value `val`.

Options include

`"absolute tolerance"`

Absolute tolerance. May be either vector or scalar. If a vector, it must match the dimension of the state vector, and the relative tolerance must also be a vector of the same length.

`"relative tolerance"`

Relative tolerance. May be either vector or scalar. If a vector, it must match the dimension of the state vector, and the absolute tolerance must also be a vector of the same length.

The local error test applied at each integration step is

```
abs (local error in x(i))
<= rtol(i) * abs (Y(i)) + atol(i)
```

`"compute consistent initial condition"`

If nonzero, `dassl` will attempt to compute a consistent set of initial conditions. This is generally not reliable, so it is best to provide a consistent set and leave this option set to zero.

`"enforce nonnegativity constraints"`

If you know that the solutions to your equations will always be non-negative, it may help to set this parameter to a nonzero value. However, it is probably best to try leaving this option set to zero first, and only setting it to a nonzero value if that doesn't work very well.

`"initial step size"`

Differential-algebraic problems may occasionally suffer from severe scaling difficulties on the first step. If you know a great deal about the scaling of your problem, you can help to alleviate this problem by specifying an initial stepsize.

`"maximum order"`

Restrict the maximum order of the solution method. This option must be between 1 and 5, inclusive.

`"maximum step size"`

Setting the maximum stepsize will avoid passing over very large regions (default is not specified).

`"step limit"`

Maximum number of integration steps to attempt on a single call to the underlying Fortran code.

```
[x, xdot, t_out, istat, msg] = dasrt (fcn, g, x_0, xdot_0, t)
... = dasrt (fcn, g, x_0, xdot_0, t, t_crit)
... = dasrt (fcn, x_0, xdot_0, t)
... = dasrt (fcn, x_0, xdot_0, t, t_crit)
Solve a set of differential-algebraic equations.
dasrt solves the set of equations
```

$$0 = f(x, \dot{x}, t)$$

with

$$x(t_0) = x_0, \dot{x}(t_0) = \dot{x}_0$$

with functional stopping criteria (root solving).

The solution is returned in the matrices *x* and *xdot*, with each row in the result matrices corresponding to one of the elements in the vector *t_out*. The first element of *t* should be *t*₀ and correspond to the initial state of the system *x_0* and its derivative *xdot_0*, so that the first row of the output *x* is *x_0* and the first row of the output *xdot* is *xdot_0*.

The vector *t* provides an upper limit on the length of the integration. If the stopping condition is met, the vector *t_out* will be shorter than *t*, and the final element of *t_out* will be the point at which the stopping condition was met, and may not correspond to any element of the vector *t*.

The first argument, *fcn*, is a string, inline, or function handle that names the function *f* to call to compute the vector of residuals for the set of equations. It must have the form

```
res = f (x, xdot, t)
```

in which *x*, *xdot*, and *res* are vectors, and *t* is a scalar.

If *fcn* is a two-element string array or a two-element cell array of strings, inline functions, or function handles, the first element names the function *f* described above, and the second element names a function to compute the modified Jacobian

$$J = \frac{\partial f}{\partial x} + c \frac{\partial f}{\partial \dot{x}}$$

The modified Jacobian function must have the form

```
jac = j (x, xdot, t, c)
```

The optional second argument names a function that defines the constraint functions whose roots are desired during the integration. This function must have the form

```
g_out = g (x, t)
```

and return a vector of the constraint function values. If the value of any of the constraint functions changes sign, DASRT will attempt to stop the integration at the point of the sign change.

If the name of the constraint function is omitted, **dasrt** solves the same problem as **daspk** or **dassl**.

Note that because of numerical errors in the constraint functions due to round-off and integration error, DASRT may return false roots, or return the same root at two or more nearly equal values of T . If such false roots are suspected, the user should consider smaller error tolerances or higher precision in the evaluation of the constraint functions.

If a root of some constraint function defines the end of the problem, the input to DASRT should nevertheless allow integration to a point slightly past that root, so that DASRT can locate the root by interpolation.

The third and fourth arguments to `dasrt` specify the initial condition of the states and their derivatives, and the fourth argument specifies a vector of output times at which the solution is desired, including the time corresponding to the initial condition.

The set of initial states and derivatives are not strictly required to be consistent. In practice, however, DASSL is not very good at determining a consistent set for you, so it is best if you ensure that the initial values result in the function evaluating to zero.

The sixth argument is optional, and may be used to specify a set of times that the DAE solver should not integrate past. It is useful for avoiding difficulties with singularities and points where there is a discontinuity in the derivative.

After a successful computation, the value of `istate` will be greater than zero (consistent with the Fortran version of DASSL).

If the computation is not successful, the value of `istate` will be less than zero and `msg` will contain additional information.

You can use the function `dasrt_options` to set optional parameters for `dasrt`.

See also: [\[dasrt_options\]](#), page 695, [\[daspk\]](#), page 688, [\[dasrt\]](#), page 694, [\[lsode\]](#), page 677.

```
dasrt_options ()  
val = dasrt_options (opt)  
dasrt_options (opt, val)
```

Query or set options for the function `dasrt`.

When called with no arguments, the names of all available options and their current values are displayed.

Given one argument, return the value of the option `opt`.

When called with two arguments, `dasrt_options` sets the option `opt` to value `val`.

Options include

"absolute tolerance"

Absolute tolerance. May be either vector or scalar. If a vector, it must match the dimension of the state vector, and the relative tolerance must also be a vector of the same length.

"relative tolerance"

Relative tolerance. May be either vector or scalar. If a vector, it must match the dimension of the state vector, and the absolute tolerance must also be a vector of the same length.

The local error test applied at each integration step is

```
abs (local error in x(i)) <= ...
rtol(i) * abs (Y(i)) + atol(i)
```

"initial step size"

Differential-algebraic problems may occasionally suffer from severe scaling difficulties on the first step. If you know a great deal about the scaling of your problem, you can help to alleviate this problem by specifying an initial stepsize.

"maximum order"

Restrict the maximum order of the solution method. This option must be between 1 and 5, inclusive.

"maximum step size"

Setting the maximum stepsize will avoid passing over very large regions.

"step limit"

Maximum number of integration steps to attempt on a single call to the underlying Fortran code.

See K. E. Brenan, et al., *Numerical Solution of Initial-Value Problems in Differential-Algebraic Equations*, North-Holland (1989) for more information about the implementation of DASSL.

25 Optimization

Octave comes with support for solving various kinds of optimization problems. Specifically Octave can solve problems in Linear Programming, Quadratic Programming, Nonlinear Programming, and Linear Least Squares Minimization.

25.1 Linear Programming

Octave can solve Linear Programming problems using the `glpk` function. That is, Octave can solve

$$\min_x c^T x$$

subject to the linear constraints $Ax = b$ where $x \geq 0$.

The `glpk` function also supports variations of this problem.

```
[xopt, fmin, errnum, extra] = glpk (c, A, b, lb, ub, ctype, vartype,
    sense, param)
```

Solve a linear program using the GNU GLPK library.

Given three arguments, `glpk` solves the following standard LP:

$$\min_x C^T x$$

subject to

$$Ax = b \quad x \geq 0$$

but may also solve problems of the form

$$[\min_x | \max_x] C^T x$$

subject to

$$Ax [= | \leq | \geq] b \quad LB \leq x \leq UB$$

Input arguments:

- `c` A column array containing the objective function coefficients.
- `A` A matrix containing the constraints coefficients.
- `b` A column array containing the right-hand side value for each constraint in the constraint matrix.
- `lb` An array containing the lower bound on each of the variables. If `lb` is not supplied, the default lower bound for the variables is zero.
- `ub` An array containing the upper bound on each of the variables. If `ub` is not supplied, the default upper bound is assumed to be infinite.
- `ctype` An array of characters containing the sense of each constraint in the constraint matrix. Each element of the array may be one of the following values
- `"F"` A free (unbounded) constraint (the constraint is ignored).

	"U"	An inequality constraint with an upper bound ($A(i,:)*x \leq b(i)$).
	"S"	An equality constraint ($A(i,:)*x = b(i)$).
	"L"	An inequality with a lower bound ($A(i,:)*x \geq b(i)$).
	"D"	An inequality constraint with both upper and lower bounds ($A(i,:)*x \geq -b(i)$) and ($A(i,:)*x \leq b(i)$).
vartype		A column array containing the types of the variables.
	"C"	A continuous variable.
	"I"	An integer variable.
sense		If <i>sense</i> is 1, the problem is a minimization. If <i>sense</i> is -1, the problem is a maximization. The default value is 1.
param		A structure containing the following parameters used to define the behavior of solver. Missing elements in the structure take on default values, so you only need to set the elements that you wish to change from the default.
		Integer parameters:
	msglev (default: 1)	Level of messages output by solver routines:
	0 (GLP_MSG_OFF)	No output.
	1 (GLP_MSG_ERR)	Error and warning messages only.
	2 (GLP_MSG_ON)	Normal output.
	3 (GLP_MSG_ALL)	Full output (includes informational messages).
scale (default: 16)		Scaling option. The values can be combined with the bitwise OR operator and may be the following:
	1 (GLP_SF_GM)	Geometric mean scaling.
	16 (GLP_SF_EQ)	Equilibration scaling.
	32 (GLP_SF_2N)	Round scale factors to power of two.
	64 (GLP_SF_SKIP)	Skip if problem is well scaled.

Alternatively, a value of 128 (`GLP_SF_AUTO`) may be also specified, in which case the routine chooses the scaling options automatically.

```
dual (default: 1)
    Simplex method option:
        1 (GLP_PRIMAL)
            Use two-phase primal simplex.
        2 (GLP_DUALP)
            Use two-phase dual simplex, and if it fails, switch
            to the primal simplex.
        3 (GLP_DUAL)
            Use two-phase dual simplex.

price (default: 34)
    Pricing option (for both primal and dual simplex):
        17 (GLP_PT_STD)
            Textbook pricing.
        34 (GLP_PT_PSE)
            Steepest edge pricing.

itlim (default: intmax)
    Simplex iterations limit. It is decreased by one each time
    when one simplex iteration has been performed, and reaching
    zero value signals the solver to stop the search.

outfrq (default: 200)
    Output frequency, in iterations. This parameter specifies how
    frequently the solver sends information about the solution to
    the standard output.

branch (default: 4)
    Branching technique option (for MIP only):
        1 (GLP_BR_FFW)
            First fractional variable.
        2 (GLP_BR_LFV)
            Last fractional variable.
        3 (GLP_BR_MFV)
            Most fractional variable.
        4 (GLP_BR_DTH)
            Heuristic by Driebeck and Tomlin.
        5 (GLP_BR_PCH)
            Hybrid pseudocost heuristic.
```

```

btrack (default: 4)
    Backtracking technique option (for MIP only):
      1 (GLP_BT_DFS)
          Depth first search.
      2 (GLP_BT_BFS)
          Breadth first search.
      3 (GLP_BT_BLB)
          Best local bound.
      4 (GLP_BT_BPH)
          Best projection heuristic.

presol (default: 1)
    If this flag is set, the simplex solver uses the built-in LP
    presolver. Otherwise the LP presolver is not used.

lpsolver (default: 1)
    Select which solver to use. If the problem is a MIP problem
    this flag will be ignored.
      1           Revised simplex method.
      2           Interior point method.

rtest (default: 34)
    Ratio test technique:
      17 (GLP_RT_STD)
          Standard ("textbook").
      34 (GLP_RT_HAR)
          Harris' two-pass ratio test.

tmlim (default: intmax)
    Searching time limit, in milliseconds.

outdly (default: 0)
    Output delay, in seconds. This parameter specifies how long
    the solver should delay sending information about the solution
    to the standard output.

save (default: 0)
    If this parameter is nonzero, save a copy of the problem in
    CPLEX LP format to the file "outpb.lp". There is currently
    no way to change the name of the output file.

Real parameters:

tolbnd (default: 1e-7)
    Relative tolerance used to check if the current basic solution
    is primal feasible. It is not recommended that you change
    this parameter unless you have a detailed understanding of
    its purpose.

```

toldj (default: 1e-7)

Absolute tolerance used to check if the current basic solution is dual feasible. It is not recommended that you change this parameter unless you have a detailed understanding of its purpose.

tolpiv (default: 1e-10)

Relative tolerance used to choose eligible pivotal elements of the simplex table. It is not recommended that you change this parameter unless you have a detailed understanding of its purpose.

objl1 (default: -DBL_MAX)

Lower limit of the objective function. If the objective function reaches this limit and continues decreasing, the solver stops the search. This parameter is used in the dual simplex method only.

obju1 (default: +DBL_MAX)

Upper limit of the objective function. If the objective function reaches this limit and continues increasing, the solver stops the search. This parameter is used in the dual simplex only.

tolint (default: 1e-5)

Relative tolerance used to check if the current basic solution is integer feasible. It is not recommended that you change this parameter unless you have a detailed understanding of its purpose.

tolobj (default: 1e-7)

Relative tolerance used to check if the value of the objective function is not better than in the best known integer feasible solution. It is not recommended that you change this parameter unless you have a detailed understanding of its purpose.

Output values:

xopt The optimizer (the value of the decision variables at the optimum).

fopt The optimum value of the objective function.

errnum Error code.

0 No error.

1 (GLP_EBADB)
 Invalid basis.

2 (GLP_ESING)
 Singular matrix.

3 (GLP_ECOND)
 Ill-conditioned matrix.

4 (GLP_EBOUND)	Invalid bounds.
5 (GLP_EFAIL)	Solver failed.
6 (GLP_EOBJLL)	Objective function lower limit reached.
7 (GLP_EOBJUL)	Objective function upper limit reached.
8 (GLP_EITLIM)	Iterations limit exhausted.
9 (GLP_ETMLIM)	Time limit exhausted.
10 (GLP_ENOPFS)	No primal feasible solution.
11 (GLP_ENODFS)	No dual feasible solution.
12 (GLP_EROOT)	Root LP optimum not provided.
13 (GLP_ESTOP)	Search terminated by application.
14 (GLP_EMIPGAP)	Relative MIP gap tolerance reached.
15 (GLP_ENOFEAS)	No primal/dual feasible solution.
16 (GLP_ENOCVG)	No convergence.
17 (GLP_EINSTAB)	Numerical instability.
18 (GLP_EDATA)	Invalid data.
19 (GLP_ERANGE)	Result out of range.
extra	A data structure containing the following fields:
lambda	Dual variables.
redcosts	Reduced Costs.
time	Time (in seconds) used for solving LP/MIP problem.
status	Status of the optimization.
1 (GLP_UNDEF)	Solution status is undefined.

- 2 (GLP_FEAS)
Solution is feasible.
- 3 (GLP_INFEAS)
Solution is infeasible.
- 4 (GLP_NOFEAS)
Problem has no feasible solution.
- 5 (GLP_OPT)
Solution is optimal.
- 6 (GLP_UNBND)
Problem has no unbounded solution.

Example:

```
c = [10, 6, 4]';
A = [ 1, 1, 1;
      10, 4, 5;
      2, 2, 6];
b = [100, 600, 300]';
lb = [0, 0, 0]';
ub = [];
ctype = "UUU";
vartype = "CCC";
s = -1;

param.msglev = 1;
param.itlim = 100;

[xmin, fmin, status, extra] = ...
    glpk (c, A, b, lb, ub, ctype, vartype, s, param);
```

25.2 Quadratic Programming

Octave can also solve Quadratic Programming problems, this is

$$\min_x \frac{1}{2} x^T H x + q^T x$$

subject to

$$Ax = b \quad l \leq x \leq u \quad A_{lb} \leq A_{in}x \leq A_{ub}$$

```
[x, obj, info, lambda] = qp (x0, H)
[x, obj, info, lambda] = qp (x0, H, q)
[x, obj, info, lambda] = qp (x0, H, q, A, b)
[x, obj, info, lambda] = qp (x0, H, q, A, b, lb, ub)
[x, obj, info, lambda] = qp (x0, H, q, A, b, lb, ub, A_lb, A_in, A_ub)
[x, obj, info, lambda] = qp (... , options)
```

Solve a quadratic program (QP).

Solve the quadratic program defined by

$$\min_x \frac{1}{2} x^T H x + x^T q$$

subject to

$$Ax = b \quad lb \leq x \leq ub \quad A_{lb} \leq A_{in}x \leq A_{ub}$$

using a null-space active-set method.

Any bound (A , b , lb , ub , A_{in} , A_{lb} , A_{ub}) may be set to the empty matrix (`[]`) if not present. The constraints A and A_{in} are matrices with each row representing a single constraint. The other bounds are scalars or vectors depending on the number of constraints. The algorithm is faster if the initial guess is feasible.

options An optional structure containing the following parameter(s) used to define the behavior of the solver. Missing elements in the structure take on default values, so you only need to set the elements that you wish to change from the default.

MaxIter (default: 200)

Maximum number of iterations.

info Structure containing run-time information about the algorithm. The following fields are defined:

solveiter

The number of iterations required to find the solution.

info An integer indicating the status of the solution.

0 The problem is feasible and convex. Global solution found.

1 The problem is not convex. Local solution found.

2 The problem is not convex and unbounded.

3 Maximum number of iterations reached.

6 The problem is infeasible.

```
x = pqpnnonneg (c, d)
x = pqpnnonneg (c, d, x0)
x = pqpnnonneg (c, d, x0, options)
[x, minval] = pqpnnonneg (...)

[x, minval, exitflag] = pqpnnonneg (...)

[x, minval, exitflag, output] = pqpnnonneg (...)

[x, minval, exitflag, output, lambda] = pqpnnonneg (...)

Minimize 1/2*x'*c*x + d'*x subject to x >= 0.
```

c and d must be real matrices, and c must be symmetric and positive definite.

$x0$ is an optional initial guess for the solution x .

options is an options structure to change the behavior of the algorithm (see [optimset], page 709). `pqpnnonneg` recognizes one option: "MaxIter".

Outputs:

x The solution matrix

<i>minval</i>	The minimum attained model value, $1/2*xmin'*c*xmin + d'*xmin$
<i>exitflag</i>	An indicator of convergence. 0 indicates that the iteration count was exceeded, and therefore convergence was not reached; >0 indicates that the algorithm converged. (The algorithm is stable and will converge given enough iterations.)
<i>output</i>	A structure with two fields: <ul style="list-style-type: none"> • "algorithm": The algorithm used ("nlnls") • "iterations": The number of iterations taken.
<i>lambda</i>	Undocumented output

See also: [lsqnonneg], page 708, [qp], page 703, [optimset], page 709.

25.3 Nonlinear Programming

Octave can also perform general nonlinear minimization using a successive quadratic programming solver.

```
[x, obj, info, iter, nf, lambda] = sqp (x0, phi)
[...] = sqp (x0, phi, g)
[...] = sqp (x0, phi, g, h)
[...] = sqp (x0, phi, g, h, lb, ub)
[...] = sqp (x0, phi, g, h, lb, ub, maxiter)
[...] = sqp (x0, phi, g, h, lb, ub, maxiter, tol)
```

Minimize an objective function using sequential quadratic programming (SQP).

Solve the nonlinear program

$$\min_x \phi(x)$$

subject to

$$g(x) = 0 \quad h(x) \geq 0 \quad lb \leq x \leq ub$$

using a sequential quadratic programming method.

The first argument is the initial guess for the vector *x0*.

The second argument is a function handle pointing to the objective function *phi*. The objective function must accept one vector argument and return a scalar.

The second argument may also be a 2- or 3-element cell array of function handles. The first element should point to the objective function, the second should point to a function that computes the gradient of the objective function, and the third should point to a function that computes the Hessian of the objective function. If the gradient function is not supplied, the gradient is computed by finite differences. If the Hessian function is not supplied, a BFGS update formula is used to approximate the Hessian.

When supplied, the gradient function *phi{2}* must accept one vector argument and return a vector. When supplied, the Hessian function *phi{3}* must accept one vector argument and return a matrix.

The third and fourth arguments *g* and *h* are function handles pointing to functions that compute the equality constraints and the inequality constraints, respectively. If

the problem does not have equality (or inequality) constraints, then use an empty matrix (`[]`) for `g` (or `h`). When supplied, these equality and inequality constraint functions must accept one vector argument and return a vector.

The third and fourth arguments may also be 2-element cell arrays of function handles. The first element should point to the constraint function and the second should point to a function that computes the gradient of the constraint function:

$$\left(\frac{\partial f(x)}{\partial x_1}, \frac{\partial f(x)}{\partial x_2}, \dots, \frac{\partial f(x)}{\partial x_N} \right)^T$$

The fifth and sixth arguments, `lb` and `ub`, contain lower and upper bounds on `x`. These must be consistent with the equality and inequality constraints `g` and `h`. If the arguments are vectors then `x(i)` is bound by `lb(i)` and `ub(i)`. A bound can also be a scalar in which case all elements of `x` will share the same bound. If only one bound (`lb`, `ub`) is specified then the other will default to `(-realmax, +realmax)`.

The seventh argument `maxiter` specifies the maximum number of iterations. The default value is 100.

The eighth argument `tol` specifies the tolerance for the stopping criteria. The default value is `sqrts (eps)`.

The value returned in `info` may be one of the following:

- 101 The algorithm terminated normally. All constraints meet the specified tolerance.
- 102 The BFGS update failed.
- 103 The maximum number of iterations was reached.
- 104 The stepsize has become too small, i.e., Δx , is less than `tol * norm (x)`.

An example of calling `sqp`:

```
function r = g (x)
  r = [ sumsq(x)-10;
        x(2)*x(3)-5*x(4)*x(5);
        x(1)^3+x(2)^3+1 ];
endfunction

function obj = phi (x)
  obj = exp (prod (x)) - 0.5*(x(1)^3+x(2)^3+1)^2;
endfunction

x0 = [-1.8; 1.7; 1.9; -0.8; -0.8];

[x, obj, info, iter, nf, lambda] = sqp (x0, @phi, @g, [])

x =
-1.71714
 1.59571
```

```

1.82725
-0.76364
-0.76364

obj = 0.053950
info = 101
iter = 8
nf = 10
lambda =

-0.0401627
0.0379578
-0.0052227

```

See also: [qp], page 703.

25.4 Linear Least Squares

Octave also supports linear least squares minimization. That is, Octave can find the parameter b such that the model $y = xb$ fits data (x, y) as well as possible, assuming zero-mean Gaussian noise. If the noise is assumed to be isotropic the problem can be solved using the ‘\’ or ‘/’ operators, or the `ols` function. In the general case where the noise is assumed to be anisotropic the `gls` is needed.

`[beta, sigma, r] = ols (y, x)`

Ordinary least squares (OLS) estimation.

OLS applies to the multivariate model $y = xb + e$ where y is a $t \times p$ matrix, x is a $t \times k$ matrix, b is a $k \times p$ matrix, and e is a $t \times p$ matrix.

Each row of y is a p -variate observation in which each column represents a variable. Likewise, the rows of x represent k -variate observations or possibly designed values. Furthermore, the collection of observations x must be of adequate rank, k , otherwise b cannot be uniquely estimated.

The observation errors, e , are assumed to originate from an underlying p -variate distribution with zero mean and p -by- p covariance matrix S , both constant conditioned on x . Furthermore, the matrix S is constant with respect to each observation such that $\bar{e} = 0$ and $\text{cov}(\text{vec}(e)) = \text{kron}(s, I)$. (For cases that don’t meet this criteria, such as autocorrelated errors, see generalized least squares, `gls`, for more efficient estimations.)

The return values `beta`, `sigma`, and `r` are defined as follows.

`beta` The OLS estimator for matrix b . `beta` is calculated directly via $(x^T x)^{-1} x^T y$ if the matrix $x^T x$ is of full rank. Otherwise, `beta = pinv (x) * y` where `pinv (x)` denotes the pseudoinverse of x .

`sigma` The OLS estimator for the matrix s ,

$$\text{sigma} = (y - x * \text{beta})' * (y - x * \text{beta}) / (\text{t-rank}(x))$$

`r` The matrix of OLS residuals, `r = y - x * beta`.

See also: [gls], page 708, [pinv], page 553.

```
[beta, v, r] = gls (y, x, o)
```

Generalized least squares (GLS) model.

Perform a generalized least squares estimation for the multivariate model $y = x b + e$ where y is a $t \times p$ matrix, x is a $t \times k$ matrix, b is a $k \times p$ matrix and e is a $t \times p$ matrix.

Each row of y is a p -variate observation in which each column represents a variable. Likewise, the rows of x represent k -variate observations or possibly designed values. Furthermore, the collection of observations x must be of adequate rank, k , otherwise b cannot be uniquely estimated.

The observation errors, e , are assumed to originate from an underlying p -variate distribution with zero mean but possibly heteroscedastic observations. That is, in general, $\bar{e} = 0$ and $\text{cov}(\text{vec}(e)) = s^2 o$ in which s is a scalar and o is a $tp \times tp$ matrix.

The return values beta , v , and r are defined as follows.

beta The GLS estimator for matrix b .

v The GLS estimator for scalar s^2 .

r The matrix of GLS residuals, $r = y - x * \text{beta}$.

See also: [\[ols\]](#), page 707.

```
x = lsqnonneg (c, d)
x = lsqnonneg (c, d, x0)
x = lsqnonneg (c, d, x0, options)
[x, resnorm] = lsqnonneg (...)

[x, resnorm, residual] = lsqnonneg (...)

[x, resnorm, residual, exitflag] = lsqnonneg (...)

[x, resnorm, residual, exitflag, output] = lsqnonneg (...)

[x, resnorm, residual, exitflag, output, lambda] = lsqnonneg (...)
```

Minimize $\text{norm} (c*x - d)$ subject to $x \geq 0$.

c and d must be real matrices.

$x0$ is an optional initial guess for the solution x .

$options$ is an options structure to change the behavior of the algorithm (see [\[optimset\]](#), page 709). `lsqnonneg` recognizes these options: "MaxIter", "TolX".

Outputs:

resnorm The squared 2-norm of the residual: `norm (c*x-d)^2`

residual The residual: $d - c*x$

exitflag An indicator of convergence. 0 indicates that the iteration count was exceeded, and therefore convergence was not reached; >0 indicates that the algorithm converged. (The algorithm is stable and will converge given enough iterations.)

output A structure with two fields:

- "algorithm": The algorithm used ("nnls")
- "iterations": The number of iterations taken.

lambda Undocumented output

See also: [pqpnnonneg], page 704, [lscov], page 709, [optimset], page 709.

```
x = lscov (A, b)
x = lscov (A, b, V)
x = lscov (A, b, V, alg)
[x, stdx, mse, S] = lscov (...)
```

Compute a generalized linear least squares fit.

Estimate *x* under the model $b = Ax + w$, where the noise *w* is assumed to follow a normal distribution with covariance matrix σ^2V .

If the size of the coefficient matrix *A* is *n*-by-*p*, the size of the vector/array of constant terms *b* must be *n*-by-*k*.

The optional input argument *V* may be an *n*-element vector of positive weights (inverse variances), or an *n*-by-*n* symmetric positive semi-definite matrix representing the covariance of *b*. If *V* is not supplied, the ordinary least squares solution is returned.

The *alg* input argument, a guidance on solution method to use, is currently ignored. Besides the least-squares estimate matrix *x* (*p*-by-*k*), the function also returns *stdx* (*p*-by-*k*), the error standard deviation of estimated *x*; *mse* (*k*-by-1), the estimated data error covariance scale factors (σ^2); and *S* (*p*-by-*p*, or *p*-by-*p*-by-*k* if *k* > 1), the error covariance of *x*.

Reference: Golub and Van Loan (1996), *Matrix Computations (3rd Ed.)*, Johns Hopkins, Section 5.6.3

See also: [ols], page 707, [gls], page 708, [lsqpnnonneg], page 708.

```
optimset ()
options = optimset ()
options = optimset (par, val, ...)
options = optimset (old, par, val, ...)
options = optimset (old, new)
```

Create options structure for optimization functions.

When called without any input or output arguments, *optimset* prints a list of all valid optimization parameters.

When called with one output and no inputs, return an options structure with all valid option parameters initialized to [].

When called with a list of parameter/value pairs, return an options structure with only the named parameters initialized.

When the first input is an existing options structure *old*, the values are updated from either the *par/val* list or from the options structure *new*.

Valid parameters are:

AutoScaling

ComplexEqn

Display Request verbose display of results from optimizations. Values are:

"off" [default]

No display.

"iter"	Display intermediate results for every loop iteration.
"final"	Display the result of the final loop iteration.
"notify"	Display the result of the final loop iteration if the function has failed to converge.

FinDiffType

FunValCheck

When enabled, display an error if the objective function returns an invalid value (a complex number, NaN, or Inf). Must be set to "on" or "off" [default]. Note: the functions `fzero` and `fminbnd` correctly handle Inf values and only complex values or NaN will cause an error in this case.

GradObj

When set to "on", the function to be minimized must return a second argument which is the gradient, or first derivative, of the function at the point x . If set to "off" [default], the gradient is computed via finite differences.

Jacobian

When set to "on", the function to be minimized must return a second argument which is the Jacobian, or first derivative, of the function at the point x . If set to "off" [default], the Jacobian is computed via finite differences.

MaxFunEvals

Maximum number of function evaluations before optimization stops. Must be a positive integer.

MaxIter

Maximum number of algorithm iterations before optimization stops. Must be a positive integer.

OutputFcn

A user-defined function executed once per algorithm iteration.

TolFun

Termination criterion for the function output. If the difference in the calculated objective function between one algorithm iteration and the next is less than `TolFun` the optimization stops. Must be a positive scalar.

TolX

Termination criterion for the function input. If the difference in x , the current search point, between one algorithm iteration and the next is less than `TolX` the optimization stops. Must be a positive scalar.

TypicalX

Updating

See also: [\[optimget\]](#), page 710.

`optimget (options, parname)`

`optimget (options, parname, default)`

Return the specific option `parname` from the optimization options structure `options` created by `optimset`.

If `parname` is not defined then return `default` if supplied, otherwise return an empty matrix.

See also: [\[optimset\]](#), page 709.

26 Statistics

Octave has support for various statistical methods. The emphasis is on basic descriptive statistics, but the Octave Forge statistics package includes probability distributions, statistical tests, random number generation, and much more.

The functions that analyze data all assume that multi-dimensional data is arranged in a matrix where each row is an observation, and each column is a variable. Thus, the matrix defined by

```
a = [ 0.9, 0.7;
      0.1, 0.1;
      0.5, 0.4 ];
```

contains three observations from a two-dimensional distribution. While this is the default data arrangement, most functions support different arrangements.

It should be noted that the statistics functions don't test for data containing NaN, NA, or Inf. These values need to be detected and dealt with explicitly. See [[isnan](#)], page 476, [[isna](#)], page 43, [[isinf](#)], page 476, [[isfinite](#)], page 477.

26.1 Descriptive Statistics

One principal goal of descriptive statistics is to represent the essence of a large data set concisely. Octave provides the mean, median, and mode functions which all summarize a data set with just a single number corresponding to the central tendency of the data.

```
mean (x)
mean (x, dim)
mean (x, opt)
mean (x, dim, opt)
mean (... , outtype)
```

Compute the mean of the elements of the vector *x*.

The mean is defined as

$$\text{mean}(x) = \bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

where *N* is the number of elements of *x*.

If *x* is a matrix, compute the mean for each column and return them in a row vector.

If the optional argument *dim* is given, operate along this dimension.

The optional argument *opt* selects the type of mean to compute. The following options are recognized:

- "a" Compute the (ordinary) arithmetic mean. [default]
- "g" Compute the geometric mean.
- "h" Compute the harmonic mean.

The optional argument *outtype* selects the data type of the output value. The following options are recognized:

- "default"** Output will be of class double unless *x* is of class single, in which case the output will also be single.
- "double"** Output will be of class double.
- "native"** Output will be the same class as *x* unless *x* is of class logical in which case it returns of class double.

Both *dim* and *opt* are optional. If both are supplied, either may appear first.

See also: [median], page 712, [mode], page 712.

median (*x*)

median (*x*, *dim*)

Compute the median value of the elements of the vector *x*.

When the elements of *x* are sorted, say *s* = **sort** (*x*), the median is defined as

$$\text{median}(x) = \begin{cases} s(\lceil N/2 \rceil), & N \text{ odd}; \\ (s(N/2) + s(N/2 + 1))/2, & N \text{ even}. \end{cases}$$

where *N* is the number of elements of *x*.

If *x* is of a discrete type such as integer or logical, then the case of even *N* rounds up (or toward **true**).

If *x* is a matrix, compute the median value for each column and return them in a row vector.

If the optional *dim* argument is given, operate along this dimension.

See also: [mean], page 711, [mode], page 712.

mode (*x*)

mode (*x*, *dim*)

[*m*, *f*, *c*] = **mode** (...)

Compute the most frequently occurring value in a dataset (mode).

mode determines the frequency of values along the first non-singleton dimension and returns the value with the highest frequency. If two, or more, values have the same frequency **mode** returns the smallest.

If the optional argument *dim* is given, operate along this dimension.

The return variable *f* is the number of occurrences of the mode in the dataset.

The cell array *c* contains all of the elements with the maximum frequency.

See also: [mean], page 711, [median], page 712.

Using just one number, such as the mean, to represent an entire data set may not give an accurate picture of the data. One way to characterize the fit is to measure the dispersion of the data. Octave provides several functions for measuring dispersion.

```
[s, l] = bounds (x)
[s, l] = bounds (x, dim)
[s, l] = bounds (..., "nanflag")
```

Return the smallest and largest values of the input data x .

If x is a vector, the bounds are calculated over the elements of x . If x is a matrix, the bounds are calculated for each column. For a multi-dimensional array, the bounds are calculated over the first non-singleton dimension.

If the optional argument dim is given, operate along this dimension.

The optional argument "nanflag" defaults to "omitnan" which does not include NaN values in the result. If the argument "includenan" is given, and there is a NaN present, then the result for both smallest (s) and largest (l) elements will be NaN.

The bounds are a quickly computed measure of the dispersion of a data set, but are less accurate than `iqr` if there are outlying data points.

See also: [range], page 713, [iqr], page 713, [mad], page 713, [std], page 714.

```
range (x)
range (x, dim)
```

Return the range, i.e., the difference between the maximum and the minimum of the input data.

If x is a vector, the range is calculated over the elements of x . If x is a matrix, the range is calculated over each column of x .

If the optional argument dim is given, operate along this dimension.

The range is a quickly computed measure of the dispersion of a data set, but is less accurate than `iqr` if there are outlying data points.

See also: [bounds], page 712, [iqr], page 713, [mad], page 713, [std], page 714.

```
iqr (x)
iqr (x, dim)
```

Return the interquartile range, i.e., the difference between the upper and lower quartile of the input data.

If x is a matrix, do the above for first non-singleton dimension of x .

If the optional argument dim is given, operate along this dimension.

As a measure of dispersion, the interquartile range is less affected by outliers than either `range` or `std`.

See also: [bounds], page 712, [mad], page 713, [range], page 713, [std], page 714.

```
mad (x)
mad (x, opt)
mad (x, opt, dim)
```

Compute the mean or median absolute deviation of the elements of x .

The mean absolute deviation is defined as

```
mad = mean (abs (x - mean (x)))
```

The median absolute deviation is defined as

```
mad = median (abs (x - median (x)))
```

If x is a matrix, compute `mad` for each column and return results in a row vector. For a multi-dimensional array, the calculation is done over the first non-singleton dimension.

The optional argument opt determines whether mean or median absolute deviation is calculated. The default is 0 which corresponds to mean absolute deviation; A value of 1 corresponds to median absolute deviation.

If the optional argument dim is given, operate along this dimension.

As a measure of dispersion, `mad` is less affected by outliers than `std`.

See also: `[bounds]`, page 712, `[range]`, page 713, `[iqr]`, page 713, `[std]`, page 714, `[mean]`, page 711, `[median]`, page 712.

`meansq (x)`

`meansq (x, dim)`

Compute the mean square of the elements of the vector x .

The mean square is defined as

$$\text{meansq}(x) = \frac{\sum_{i=1}^N x_i^2}{N}$$

where N is the number of elements of x .

If x is a matrix, return a row vector containing the mean square of each column.

If the optional argument dim is given, operate along this dimension.

See also: `[var]`, page 715, `[std]`, page 714, `[moment]`, page 716.

`std (x)`

`std (x, opt)`

`std (x, opt, dim)`

Compute the standard deviation of the elements of the vector x .

The standard deviation is defined as

$$\text{std}(x) = \sigma = \sqrt{\frac{\sum_{i=1}^N (x_i - \bar{x})^2}{N - 1}}$$

where \bar{x} is the mean value of x and N is the number of elements of x .

If x is a matrix, compute the standard deviation for each column and return them in a row vector.

The argument opt determines the type of normalization to use. Valid values are

- 0: normalize with $N - 1$, provides the square root of the best unbiased estimator of the variance [default]
- 1: normalize with N , this provides the square root of the second moment around the mean

If the optional argument dim is given, operate along this dimension.

See also: `[var]`, page 715, `[bounds]`, page 712, `[mad]`, page 713, `[range]`, page 713, `[iqr]`, page 713, `[mean]`, page 711, `[median]`, page 712.

In addition to knowing the size of a dispersion it is useful to know the shape of the data set. For example, are data points massed to the left or right of the mean? Octave provides several common measures to describe the shape of the data set. Octave can also calculate moments allowing arbitrary shape measures to be developed.

```
var (x)
var (x, opt)
var (x, opt, dim)
```

Compute the variance of the elements of the vector x .

The variance is defined as

$$\text{var}(x) = \sigma^2 = \frac{\sum_{i=1}^N (x_i - \bar{x})^2}{N - 1}$$

where \bar{x} is the mean value of x and N is the number of elements of x .

If x is a matrix, compute the variance for each column and return them in a row vector.

The argument opt determines the type of normalization to use. Valid values are

- 0: normalize with $N - 1$, provides the best unbiased estimator of the variance [default]
 - 1: normalizes with N , this provides the second moment around the mean
- If N is equal to 1 the value of opt is ignored and normalization by N is used.

If the optional argument dim is given, operate along this dimension.

See also: [\[cov\]](#), page 739, [\[std\]](#), page 714, [\[skewness\]](#), page 715, [\[kurtosis\]](#), page 716, [\[moment\]](#), page 716.

```
skewness (x)
skewness (x, flag)
skewness (x, flag, dim)
```

Compute the sample skewness of the elements of x .

The sample skewness is defined as

$$\text{skewness}(x) = \frac{\frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^3}{\sigma^3},$$

where N is the length of x , \bar{x} its mean and σ its (uncorrected) standard deviation.

The optional argument $flag$ controls which normalization is used. If $flag$ is equal to 1 (default value, used when $flag$ is omitted or empty), return the sample skewness as defined above. If $flag$ is equal to 0, return the adjusted skewness coefficient instead:

$$\text{skewness}(x) = \frac{\sqrt{N(N - 1)}}{N - 2} \times \frac{\frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^3}{\sigma^3}$$

The adjusted skewness coefficient is obtained by replacing the sample second and third central moments by their bias-corrected versions.

If x is a matrix, or more generally a multi-dimensional array, return the skewness along the first non-singleton dimension. If the optional dim argument is given, operate along this dimension.

See also: [\[var\]](#), page 715, [\[kurtosis\]](#), page 716, [\[moment\]](#), page 716.

kurtosis (x)
kurtosis (x, flag)
kurtosis (x, flag, dim)

Compute the sample kurtosis of the elements of *x*.

The sample kurtosis is defined as

$$\kappa_1 = \frac{\frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^4}{\sigma^4},$$

where *N* is the length of *x*, \bar{x} its mean, and σ its (uncorrected) standard deviation.

The optional argument *flag* controls which normalization is used. If *flag* is equal to 1 (default value, used when *flag* is omitted or empty), return the sample kurtosis as defined above. If *flag* is equal to 0, return the "bias-corrected" kurtosis coefficient instead:

$$\kappa_0 = 3 + \frac{N-1}{(N-2)(N-3)} ((N+1) \kappa_1 - 3(N-1))$$

The bias-corrected kurtosis coefficient is obtained by replacing the sample second and fourth central moments by their unbiased versions. It is an unbiased estimate of the population kurtosis for normal populations.

If *x* is a matrix, or more generally a multi-dimensional array, return the kurtosis along the first non-singleton dimension. If the optional *dim* argument is given, operate along this dimension.

See also: [var], page 715, [skewness], page 715, [moment], page 716.

moment (x, p)
moment (x, p, type)
moment (x, p, dim)
moment (x, p, type, dim)
moment (x, p, dim, type)

Compute the *p*-th central moment of the vector *x*.

The *p*-th central moment of *x* is defined as:

$$\frac{\sum_{i=1}^N (x_i - \bar{x})^p}{N}$$

where \bar{x} is the mean value of *x* and *N* is the number of elements of *x*.

If *x* is a matrix, return the row vector containing the *p*-th central moment of each column.

If the optional argument *dim* is given, operate along this dimension.

The optional string *type* specifies the type of moment to be computed. Valid options are:

- "c" Central Moment (default).
- "a"
- "ac" Absolute Central Moment. The moment about the mean ignoring sign defined as

$$\frac{\sum_{i=1}^N |x_i - \bar{x}|^p}{N}$$

"**r**" Raw Moment. The moment about zero defined as

$$\text{moment}(x) = \frac{\sum_{i=1}^N x_i^p}{N}$$

"**ar**" Absolute Raw Moment. The moment about zero ignoring sign defined as

$$\frac{\sum_{i=1}^N |x_i|^p}{N}$$

If both *type* and *dim* are given they may appear in any order.

See also: [var], page 715, [skewness], page 715, [kurtosis], page 716.

```
q = quantile (x)
q = quantile (x, p)
q = quantile (x, p, dim)
q = quantile (x, p, dim, method)
```

For a sample, *x*, calculate the quantiles, *q*, corresponding to the cumulative probability values in *p*. All non-numeric values (NaNs) of *x* are ignored.

If *x* is a matrix, compute the quantiles for each column and return them in a matrix, such that the *i*-th row of *q* contains the *p(i)*th quantiles of each column of *x*.

If *p* is unspecified, return the quantiles for [0.00 0.25 0.50 0.75 1.00]. The optional argument *dim* determines the dimension along which the quantiles are calculated. If *dim* is omitted it defaults to the first non-singleton dimension.

The methods available to calculate sample quantiles are the nine methods used by R (<https://www.r-project.org/>). The default value is *method* = 5.

Discontinuous sample quantile methods 1, 2, and 3

1. Method 1: Inverse of empirical distribution function.
2. Method 2: Similar to method 1 but with averaging at discontinuities.
3. Method 3: SAS definition: nearest even order statistic.

Continuous sample quantile methods 4 through 9, where *p(k)* is the linear interpolation function respecting each method's representative cdf.

4. Method 4: $p(k) = k/N$. That is, linear interpolation of the empirical cdf, where *N* is the length of *P*.
5. Method 5: $p(k) = (k - 0.5)/N$. That is, a piecewise linear function where the knots are the values midway through the steps of the empirical cdf.
6. Method 6: $p(k) = k/(N + 1)$.
7. Method 7: $p(k) = (k - 1)/(N - 1)$.
8. Method 8: $p(k) = (k - 1/3)/(N + 1/3)$. The resulting quantile estimates are approximately median-unbiased regardless of the distribution of *x*.
9. Method 9: $p(k) = (k - 3/8)/(N + 1/4)$. The resulting quantile estimates are approximately unbiased for the expected order statistics if *x* is normally distributed.

Hyndman and Fan (1996) recommend method 8. Maxima, S, and R (versions prior to 2.0.0) use 7 as their default. Minitab and SPSS use method 6. MATLAB uses method 5.

References:

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) The New S Language. Wadsworth & Brooks/Cole.
- Hyndman, R. J. and Fan, Y. (1996) Sample quantiles in statistical packages, American Statistician, 50, 361–365.
- R: A Language and Environment for Statistical Computing; <https://cran.r-project.org/doc/manuals/fullrefman.pdf>.

Examples:

```
x = randi (1000, [10, 1]); # Create empirical data in range 1-1000
q = quantile (x, [0, 1]); # Return minimum, maximum of distribution
q = quantile (x, [0.25 0.5 0.75]); # Return quartiles of distribution
```

See also: [\[prctile\]](#), page 718.

```
q = prctile (x)
q = prctile (x, p)
q = prctile (x, p, dim)
```

For a sample x , compute the quantiles, q , corresponding to the cumulative probability values, p , in percent.

If x is a matrix, compute the percentiles for each column and return them in a matrix, such that the i -th row of q contains the $p(i)$ th percentiles of each column of x .

If p is unspecified, return the quantiles for [0 25 50 75 100].

The optional argument dim determines the dimension along which the percentiles are calculated. If dim is omitted it defaults to the first non-singleton dimension.

Programming Note: All non-numeric values (NaNs) of x are ignored.

See also: [\[quantile\]](#), page 717.

A summary view of a data set can be generated quickly with the **statistics** function.

```
statistics (x)
statistics (x, dim)
```

Return a vector with the minimum, first quartile, median, third quartile, maximum, mean, standard deviation, skewness, and kurtosis of the elements of the vector x .

If x is a matrix, calculate statistics over the first non-singleton dimension.

If the optional argument dim is given, operate along this dimension.

See also: [\[min\]](#), page 520, [\[max\]](#), page 519, [\[median\]](#), page 712, [\[mean\]](#), page 711, [\[std\]](#), page 714, [\[skewness\]](#), page 715, [\[kurtosis\]](#), page 716.

26.2 Statistics on Sliding Windows of Data

It is often useful to calculate descriptive statistics over a subsection (i.e., window) of a full dataset. Octave provides the function `movfun` which will call an arbitrary function handle with windows of data and accumulate the results. Many of the most commonly desired functions, such as the moving average over a window of data (`movmean`), are already provided.

```
y = movfun (fcn, x, wlen)
y = movfun (fcn, x, [nb, na])
y = movfun (..., "property", value)
```

Apply function *fcn* to a moving window of length *wlen* on data *x*.

If *wlen* is a scalar, the function *fcn* is applied to a moving window of length *wlen*. When *wlen* is an odd number the window is symmetric and includes $(wlen - 1) / 2$ elements on either side of the central element. For example, when calculating the output at index 5 with a window length of 3, `movfun` uses data elements [4, 5, 6]. If *wlen* is an even number, the window is asymmetric and has *wlen/2* elements to the left of the central element and *wlen/2 - 1* elements to the right of the central element. For example, when calculating the output at index 5 with a window length of 4, `movfun` uses data elements [3, 4, 5, 6].

If *wlen* is an array with two elements `[nb, na]`, the function is applied to a moving window `-nb:na`. This window includes *nb* number of elements *before* the current element and *na* number of elements *after* the current element. The current element is always included. For example, given *wlen* = [3, 0], the data used to calculate index 5 is [2, 3, 4, 5].

During calculations the data input *x* is reshaped into a 2-dimensional *wlen*-by-*N* matrix and *fcn* is called on this new matrix. Therefore, *fcn* must accept an array input argument and apply the computation along dimension 1, i.e., down the columns of the array.

When applied to an array (possibly multi-dimensional) with *n* columns, *fcn* may return a result in either of two formats: Format 1) an array of size 1-by-*n*-by-*dim3*-by-...-by-*dimN*. This is the typical output format from Octave core functions. Type `demo ("movfun", 5)` for an example of this use case. Format 2) a row vector of length *n * numel_higher_dims* where *numel_higher_dims* is `prod (size (x) (3:end))`. The output of *fcn* for the *i*-th input column must be found in the output at indices `i:n:(n*numel_higher_dims)`. This format is useful when concatenating functions into arrays, or when using `nthargout`. Type `demo ("movfun", 6)` for an example of this case.

The calculation can be controlled by specifying *property/value* pairs. Valid properties are

"dim" Operate along the dimension specified, rather than the default of the first non-singleton dimension.

"Endpoints"

This property controls how results are calculated at the boundaries (endpoints) of the window. Possible values are:

"shrink" (default)

The window is truncated at the beginning and end of the array to exclude elements for which there is no source data. For example, with a window of length 3, $y(1) = fcn(x(1:2))$, and $y(end) = fcn(x(end-1:end))$.

"discard"

Any y values that use a window extending beyond the original data array are deleted. For example, with a 10-element data vector and a window of length 3, the output will contain only 8 elements. The first element would require calculating the function over indices [0, 1, 2] and is therefore discarded. The last element would require calculating the function over indices [9, 10, 11] and is therefore discarded.

"fill"

Any window elements outside the data array are replaced by NaN. For example, with a window of length 3, $y(1) = fcn([NaN, x(1:2)])$, and $y(end) = fcn([x(end-1:end), NaN])$. This option usually results in y having NaN values at the boundaries, although it is influenced by how fcn handles NaN, and also by the property "nancond".

user_value

Any window elements outside the data array are replaced by the specified value *user_value* which must be a numeric scalar. For example, with a window of length 3, $y(1) = fcn([user_value, x(1:2)])$, and $y(end) = fcn([x(end-1:end), user_value])$. A common choice for *user_value* is 0.

"same"

Any window elements outside the data array are replaced by the value of x at the boundary. For example, with a window of length 3, $y(1) = fcn([x(1), x(1:2)])$, and $y(end) = fcn([x(end-1:end), x(end)])$.

"periodic"

The window is wrapped so that any missing data elements are taken from the other side of the data. For example, with a window of length 3, $y(1) = fcn([x(end), x(1:2)])$, and $y(end) = fcn([x(end-1:end), x(1)])$.

Note that for some of these choices, the window size at the boundaries is not the same as for the central part, and fcn must work in these cases.

"nancond"

Controls whether NaN and NA values should be included (value: "includenan"), or excluded (value: "omitnan"), from the data passed to fcn . The default is "includenan". Caution: The "omitnan" option is not yet implemented.

"outdim"

A row vector that selects which dimensions of the calculation will appear in the output y . This is only useful when fcn returns an N-dimensional array in Format 1. The default is to return all output dimensions.

Programming Note: The property "outdim" can be used to save memory when the output of *fcn* has many dimensions, or when a wrapper to the base function that selects the desired outputs is too costly. When memory is not an issue, the easiest way to select output dimensions is to first calculate the complete result with `movfun` and then filter that result with indexing. If code complexity is not an issue then a wrapper can be created using anonymous functions. For example, if `basefcn` is a function returning a K -dimensional row output, and only dimension D is desired, then the following wrapper could be used.

```
fcn = @(x) basefcn (x)(:,size(x,2) * (D-1) + (1:size(x,2)));
y = movfun (@fcn, ...);
```

See also: [movslice], page 721, [prepad], page 488, [postpad], page 489, [permute], page 482, [reshape], page 482.

`slicidx = movslice (N, wlen)`

`[slicidx, C, Cpre, Cpost, win] = movslice (...)`

Generate indices to slice a vector of length N in to windows of length $wlen$.

FIXME: Document inputs N, wlen

FIXME: Document outputs slicidx, C, Cpre, Cpost, win.

See also: [movfun], page 719.

`y = movmad (x, wlen)`

`y = movmad (x, [na, nb])`

`y = movmad (..., dim)`

`y = movmad (..., "nancond")`

`y = movmad (..., property, value)`

Calculate the moving mean absolute deviation over a sliding window of length $wlen$ on data x .

If $wlen$ is a scalar, the function `mad` is applied to a moving window of length $wlen$. When $wlen$ is an odd number the window is symmetric and includes $(wlen - 1) / 2$ elements on either side of the central element. For example, when calculating the output at index 5 with a window length of 3, `movmad` uses data elements [4, 5, 6]. If $wlen$ is an even number, the window is asymmetric and has $wlen/2$ elements to the left of the central element and $wlen/2 - 1$ elements to the right of the central element. For example, when calculating the output at index 5 with a window length of 4, `movmad` uses data elements [3, 4, 5, 6].

If $wlen$ is an array with two elements $[nb, na]$, the function is applied to a moving window $-nb:na$. This window includes nb number of elements *before* the current element and na number of elements *after* the current element. The current element is always included. For example, given $wlen = [3, 0]$, the data used to calculate index 5 is [2, 3, 4, 5].

If the optional argument *dim* is given, operate along this dimension.

The optional string argument "nancond" controls whether NaN and NA values should be included ("includenan"), or excluded ("omitnan"), from the data passed to `mad`. The default is "includenan". Caution: the "omitnan" option is not yet implemented.

The calculation can be controlled by specifying *property*/*value* pairs. Valid properties are

"Endpoints"

This property controls how results are calculated at the boundaries (endpoints) of the window. Possible values are:

"shrink" (default)

The window is truncated at the beginning and end of the array to exclude elements for which there is no source data. For example, with a window of length 3, $y(1) = \text{mad}(x(1:2))$, and $y(\text{end}) = \text{mad}(x(\text{end}-1:\text{end}))$.

"discard"

Any y values that use a window extending beyond the original data array are deleted. For example, with a 10-element data vector and a window of length 3, the output will contain only 8 elements. The first element would require calculating the function over indices [0, 1, 2] and is therefore discarded. The last element would require calculating the function over indices [9, 10, 11] and is therefore discarded.

"fill"

Any window elements outside the data array are replaced by NaN. For example, with a window of length 3, $y(1) = \text{mad}([\text{NaN}, x(1:2)])$, and $y(\text{end}) = \text{mad}([x(\text{end}-1:\text{end}), \text{NaN}])$. This option usually results in y having NaN values at the boundaries, although it is influenced by how mad handles NaN, and also by the property "nancond".

user_value

Any window elements outside the data array are replaced by the specified value *user_value* which must be a numeric scalar. For example, with a window of length 3, $y(1) = \text{mad}([\text{user_value}, x(1:2)])$, and $y(\text{end}) = \text{mad}([x(\text{end}-1:\text{end}), \text{user_value}])$. A common choice for *user_value* is 0.

"same"

Any window elements outside the data array are replaced by the value of x at the boundary. For example, with a window of length 3, $y(1) = \text{mad}([x(1), x(1:2)])$, and $y(\text{end}) = \text{mad}([x(\text{end}-1:\text{end}), x(\text{end})])$.

"periodic"

The window is wrapped so that any missing data elements are taken from the other side of the data. For example, with a window of length 3, $y(1) = \text{mad}([x(\text{end}), x(1:2)])$, and $y(\text{end}) = \text{mad}([x(\text{end}-1:\text{end}), x(1)])$.

"SamplePoints"

Caution: This option is not yet implemented.

Programming Note: This function is a wrapper which calls `movfun`. For additional options and documentation, See [\[movfun\]](#), page 719.

See also: [movfun], page 719, [movslice], page 721, [movmax], page 723, [movmean], page 724, [movmedian], page 726, [movmin], page 727, [movprod], page 729, [movstd], page 731, [movsum], page 732, [movvar], page 734.

```
y = movmax (x, wlen)
y = movmax (x, [na, nb])
y = movmax (..., dim)
y = movmax (..., "nancond")
y = movmax (..., property, value)
```

Calculate the moving maximum over a sliding window of length *wlen* on data *x*.

If *wlen* is a scalar, the function `max` is applied to a moving window of length *wlen*. When *wlen* is an odd number the window is symmetric and includes $(wlen - 1) / 2$ elements on either side of the central element. For example, when calculating the output at index 5 with a window length of 3, `movmax` uses data elements [4, 5, 6]. If *wlen* is an even number, the window is asymmetric and has *wlen*/2 elements to the left of the central element and *wlen*/2 - 1 elements to the right of the central element. For example, when calculating the output at index 5 with a window length of 4, `movmax` uses data elements [3, 4, 5, 6].

If *wlen* is an array with two elements `[nb, na]`, the function is applied to a moving window `-nb:na`. This window includes *nb* number of elements *before* the current element and *na* number of elements *after* the current element. The current element is always included. For example, given *wlen* = [3, 0], the data used to calculate index 5 is [2, 3, 4, 5].

If the optional argument *dim* is given, operate along this dimension.

The optional string argument "nancond" controls whether `NaN` and `NA` values should be included ("includenan"), or excluded ("omitnan"), from the data passed to `max`. The default is "includenan". Caution: the "omitnan" option is not yet implemented.

The calculation can be controlled by specifying *property*/*value* pairs. Valid properties are

"Endpoints"

This property controls how results are calculated at the boundaries (endpoints) of the window. Possible values are:

"shrink" (default)

The window is truncated at the beginning and end of the array to exclude elements for which there is no source data. For example, with a window of length 3, $y(1) = \max(x(1:2))$, and $y(end) = \max(x(end-1:end))$.

"discard"

Any *y* values that use a window extending beyond the original data array are deleted. For example, with a 10-element data vector and a window of length 3, the output will contain only 8 elements. The first element would require calculating the function over indices [0, 1, 2] and is therefore discarded. The last element would require calculating the function over indices [9, 10, 11] and is therefore discarded.

"fill"	Any window elements outside the data array are replaced by NaN. For example, with a window of length 3, $y(1) = \max([\text{NaN}, x(1:2)])$, and $y(\text{end}) = \max([x(\text{end}-1:\text{end}), \text{NaN}])$. This option usually results in y having NaN values at the boundaries, although it is influenced by how max handles NaN, and also by the property "nancond".
user_value	Any window elements outside the data array are replaced by the specified value user_value which must be a numeric scalar. For example, with a window of length 3, $y(1) = \max([\text{user_value}, x(1:2)])$, and $y(\text{end}) = \max([x(\text{end}-1:\text{end}), \text{user_value}])$. A common choice for user_value is 0.
"same"	Any window elements outside the data array are replaced by the value of x at the boundary. For example, with a window of length 3, $y(1) = \max([x(1), x(1:2)])$, and $y(\text{end}) = \max([x(\text{end}-1:\text{end}), x(\text{end})])$.
"periodic"	The window is wrapped so that any missing data elements are taken from the other side of the data. For example, with a window of length 3, $y(1) = \max([x(\text{end}), x(1:2)])$, and $y(\text{end}) = \max([x(\text{end}-1:\text{end}), x(1)])$.

"SamplePoints"

Caution: This option is not yet implemented.

Programming Note: This function is a wrapper which calls `movfun`. For additional options and documentation, See [movfun], page 719.

See also: [movfun], page 719, [movslice], page 721, [movmad], page 721, [movmean], page 724, [movmedian], page 726, [movmin], page 727, [movprod], page 729, [movstd], page 731, [movsum], page 732, [movvar], page 734.

```
y = movmean (x, wlen)
y = movmean (x, [na, nb])
y = movmean (..., dim)
y = movmean (..., "nancond")
y = movmean (..., property, value)
```

Calculate the moving average over a sliding window of length wlen on data x.

If wlen is a scalar, the function `mean` is applied to a moving window of length wlen. When wlen is an odd number the window is symmetric and includes $(wlen - 1) / 2$ elements on either side of the central element. For example, when calculating the output at index 5 with a window length of 3, `movmean` uses data elements [4, 5, 6]. If wlen is an even number, the window is asymmetric and has $wlen/2$ elements to the left of the central element and $wlen/2 - 1$ elements to the right of the central element. For example, when calculating the output at index 5 with a window length of 4, `movmean` uses data elements [3, 4, 5, 6].

If wlen is an array with two elements [nb, na], the function is applied to a moving window -nb:na. This window includes nb number of elements before the current

element and *na* number of elements *after* the current element. The current element is always included. For example, given `wlen = [3, 0]`, the data used to calculate index 5 is `[2, 3, 4, 5]`.

If the optional argument `dim` is given, operate along this dimension.

The optional string argument "`nancond`" controls whether `Nan` and `NA` values should be included ("`includenan`"), or excluded ("`omitnan`"), from the data passed to `mean`. The default is "`includenan`". Caution: the "`omitnan`" option is not yet implemented. The calculation can be controlled by specifying *property/value* pairs. Valid properties are

"Endpoints"

This property controls how results are calculated at the boundaries (endpoints) of the window. Possible values are:

"shrink" (default)

The window is truncated at the beginning and end of the array to exclude elements for which there is no source data. For example, with a window of length 3, `y(1) = mean (x(1:2))`, and `y(end) = mean (x(end-1:end))`.

"discard"

Any *y* values that use a window extending beyond the original data array are deleted. For example, with a 10-element data vector and a window of length 3, the output will contain only 8 elements. The first element would require calculating the function over indices `[0, 1, 2]` and is therefore discarded. The last element would require calculating the function over indices `[9, 10, 11]` and is therefore discarded.

"fill"

Any window elements outside the data array are replaced by `Nan`. For example, with a window of length 3, `y(1) = mean ([NaN, x(1:2)])`, and `y(end) = mean ([x(end-1:end), NaN])`. This option usually results in *y* having `Nan` values at the boundaries, although it is influenced by how `mean` handles `Nan`, and also by the property "`nancond`".

user_value

Any window elements outside the data array are replaced by the specified value `user_value` which must be a numeric scalar. For example, with a window of length 3, `y(1) = mean ([user_value, x(1:2)])`, and `y(end) = mean ([x(end-1:end), user_value])`. A common choice for `user_value` is 0.

"same"

Any window elements outside the data array are replaced by the value of *x* at the boundary. For example, with a window of length 3, `y(1) = mean ([x(1), x(1:2)])`, and `y(end) = mean ([x(end-1:end), x(end)])`.

"periodic"

The window is wrapped so that any missing data elements are taken from the other side of the data. For example, with a

window of length 3, $y(1) = \text{mean}([x(\text{end}), x(1:2)])$, and $y(\text{end}) = \text{mean}([x(\text{end}-1:\text{end}), x(1)])$.

"SamplePoints"

Caution: This option is not yet implemented.

Programming Note: This function is a wrapper which calls `movfun`. For additional options and documentation, See [movfun], page 719.

See also: [movfun], page 719, [movslice], page 721, [movmad], page 721, [movmax], page 723, [movmedian], page 726, [movmin], page 727, [movprod], page 729, [movstd], page 731, [movsum], page 732, [movvar], page 734.

```
y = movmedian (x, wlen)
y = movmedian (x, [na, nb])
y = movmedian (..., dim)
y = movmedian (..., "nancond")
y = movmedian (..., property, value)
```

Calculate the moving median over a sliding window of length *wlen* on data *x*.

If *wlen* is a scalar, the function `movmedian` is applied to a moving window of length *wlen*. When *wlen* is an odd number the window is symmetric and includes $(wlen - 1) / 2$ elements on either side of the central element. For example, when calculating the output at index 5 with a window length of 3, `movmedian` uses data elements [4, 5, 6]. If *wlen* is an even number, the window is asymmetric and has *wlen*/2 elements to the left of the central element and *wlen*/2 - 1 elements to the right of the central element. For example, when calculating the output at index 5 with a window length of 4, `movmedian` uses data elements [3, 4, 5, 6].

If *wlen* is an array with two elements [*nb*, *na*], the function is applied to a moving window *-nb:na*. This window includes *nb* number of elements *before* the current element and *na* number of elements *after* the current element. The current element is always included. For example, given *wlen* = [3, 0], the data used to calculate index 5 is [2, 3, 4, 5].

If the optional argument *dim* is given, operate along this dimension.

The optional string argument "nancond" controls whether NaN and NA values should be included ("includenan"), or excluded ("omitnan"), from the data passed to `movmedian`. The default is "includenan". Caution: the "omitnan" option is not yet implemented.

The calculation can be controlled by specifying *property/value* pairs. Valid properties are

"Endpoints"

This property controls how results are calculated at the boundaries (endpoints) of the window. Possible values are:

"shrink" (default)

The window is truncated at the beginning and end of the array to exclude elements for which there is no source data.

For example, with a window of length 3, $y(1) = \text{movmedian}(x(1:2))$, and $y(\text{end}) = \text{movmedian}(x(\text{end}-1:\text{end}))$.

"discard"

Any *y* values that use a window extending beyond the original data array are deleted. For example, with a 10-element data vector and a window of length 3, the output will contain only 8 elements. The first element would require calculating the function over indices [0, 1, 2] and is therefore discarded. The last element would require calculating the function over indices [9, 10, 11] and is therefore discarded.

"fill"

Any window elements outside the data array are replaced by NaN. For example, with a window of length 3, *y*(1) = `movmedian ([NaN, x(1:2)])`, and *y*(end) = `movmedian ([x(end-1:end), NaN])`. This option usually results in *y* having NaN values at the boundaries, although it is influenced by how `movmedian` handles NaN, and also by the property "nancond".

user_value

Any window elements outside the data array are replaced by the specified value *user_value* which must be a numeric scalar. For example, with a window of length 3, *y*(1) = `movmedian ([user_value, x(1:2)])`, and *y*(end) = `movmedian ([x(end-1:end), user_value])`. A common choice for *user_value* is 0.

"same"

Any window elements outside the data array are replaced by the value of *x* at the boundary. For example, with a window of length 3, *y*(1) = `movmedian ([x(1), x(1:2)])`, and *y*(end) = `movmedian ([x(end-1:end), x(end)])`.

"periodic"

The window is wrapped so that any missing data elements are taken from the other side of the data. For example, with a window of length 3, *y*(1) = `movmedian ([x(end), x(1:2)])`, and *y*(end) = `movmedian ([x(end-1:end), x(1)])`.

"SamplePoints"

Caution: This option is not yet implemented.

Programming Note: This function is a wrapper which calls `movfun`. For additional options and documentation, See [\[movfun\]](#), page 719.

See also: [\[movfun\]](#), page 719, [\[movslice\]](#), page 721, [\[movmad\]](#), page 721, [\[movmax\]](#), page 723, [\[movmean\]](#), page 724, [\[movmin\]](#), page 727, [\[movprod\]](#), page 729, [\[movstd\]](#), page 731, [\[movsum\]](#), page 732, [\[movvar\]](#), page 734.

```
y = movmin (x, wlen)
y = movmin (x, [na, nb])
y = movmin (..., dim)
y = movmin (..., "nancond")
y = movmin (..., property, value)
```

Calculate the moving minimum over a sliding window of length *wlen* on data *x*.

If `wlen` is a scalar, the function `min` is applied to a moving window of length `wlen`. When `wlen` is an odd number the window is symmetric and includes $(wlen - 1) / 2$ elements on either side of the central element. For example, when calculating the output at index 5 with a window length of 3, `movmin` uses data elements [4, 5, 6]. If `wlen` is an even number, the window is asymmetric and has `wlen/2` elements to the left of the central element and `wlen/2 - 1` elements to the right of the central element. For example, when calculating the output at index 5 with a window length of 4, `movmin` uses data elements [3, 4, 5, 6].

If `wlen` is an array with two elements [`nb`, `na`], the function is applied to a moving window `-nb:na`. This window includes `nb` number of elements *before* the current element and `na` number of elements *after* the current element. The current element is always included. For example, given `wlen = [3, 0]`, the data used to calculate index 5 is [2, 3, 4, 5].

If the optional argument `dim` is given, operate along this dimension.

The optional string argument "`nancond`" controls whether `NaN` and `NA` values should be included ("`includenan`"), or excluded ("`omitnan`"), from the data passed to `min`. The default is "`includenan`". Caution: the "`omitnan`" option is not yet implemented. The calculation can be controlled by specifying *property/value* pairs. Valid properties are

"Endpoints"

This property controls how results are calculated at the boundaries (endpoints) of the window. Possible values are:

"shrink" (default)

The window is truncated at the beginning and end of the array to exclude elements for which there is no source data. For example, with a window of length 3, `y(1) = min (x(1:2))`, and `y(end) = min (x(end-1:end))`.

"discard"

Any `y` values that use a window extending beyond the original data array are deleted. For example, with a 10-element data vector and a window of length 3, the output will contain only 8 elements. The first element would require calculating the function over indices [0, 1, 2] and is therefore discarded. The last element would require calculating the function over indices [9, 10, 11] and is therefore discarded.

"fill"

Any window elements outside the data array are replaced by `NaN`. For example, with a window of length 3, `y(1) = min ([NaN, x(1:2)])`, and `y(end) = min ([x(end-1:end), NaN])`. This option usually results in `y` having `NaN` values at the boundaries, although it is influenced by how `min` handles `NaN`, and also by the property "`nancond`".

`user_value`

Any window elements outside the data array are replaced by the specified value `user_value` which must be a numeric scalar. For example, with a window of length 3, `y(1)`

$= \min ([\text{user_value}, \text{x}(1:2)]), \quad \text{and} \quad \text{y}(\text{end}) = \min ([\text{x}(\text{end}-1:\text{end}), \text{user_value}]).$ A common choice for `user_value` is 0.

"`same`" Any window elements outside the data array are replaced by the value of `x` at the boundary. For example, with a window of length 3, $\text{y}(1) = \min ([\text{x}(1), \text{x}(1:2)])$, and $\text{y}(\text{end}) = \min ([\text{x}(\text{end}-1:\text{end}), \text{x}(\text{end})]).$

"`periodic`"

The window is wrapped so that any missing data elements are taken from the other side of the data. For example, with a window of length 3, $\text{y}(1) = \min ([\text{x}(\text{end}), \text{x}(1:2)])$, and $\text{y}(\text{end}) = \min ([\text{x}(\text{end}-1:\text{end}), \text{x}(1)])$.

"`SamplePoints`"

Caution: This option is not yet implemented.

Programming Note: This function is a wrapper which calls `movfun`. For additional options and documentation, See [\[movfun\]](#), page 719.

See also: [\[movfun\]](#), page 719, [\[movslice\]](#), page 721, [\[movmad\]](#), page 721, [\[movmax\]](#), page 723, [\[movmean\]](#), page 724, [\[movmedian\]](#), page 726, [\[movprod\]](#), page 729, [\[movstd\]](#), page 731, [\[movsum\]](#), page 732, [\[movvar\]](#), page 734.

```
y = movprod (x, wlen)
y = movprod (x, [na, nb])
y = movprod (..., dim)
y = movprod (..., "nancond")
y = movprod (..., property, value)
```

Calculate the moving product over a sliding window of length `wlen` on data `x`.

If `wlen` is a scalar, the function `movprod` is applied to a moving window of length `wlen`. When `wlen` is an odd number the window is symmetric and includes $(wlen - 1) / 2$ elements on either side of the central element. For example, when calculating the output at index 5 with a window length of 3, `movprod` uses data elements [4, 5, 6]. If `wlen` is an even number, the window is asymmetric and has $wlen/2$ elements to the left of the central element and $wlen/2 - 1$ elements to the right of the central element. For example, when calculating the output at index 5 with a window length of 4, `movprod` uses data elements [3, 4, 5, 6].

If `wlen` is an array with two elements `[nb, na]`, the function is applied to a moving window $-nb:na$. This window includes `nb` number of elements *before* the current element and `na` number of elements *after* the current element. The current element is always included. For example, given `wlen = [3, 0]`, the data used to calculate index 5 is [2, 3, 4, 5].

If the optional argument `dim` is given, operate along this dimension.

The optional string argument "`nancond`" controls whether `NaN` and `NA` values should be included ("`includenan`"), or excluded ("`omitnan`"), from the data passed to `movprod`. The default is "`includenan`". Caution: the "`omitnan`" option is not yet implemented.

The calculation can be controlled by specifying *property*/*value* pairs. Valid properties are

"Endpoints"

This property controls how results are calculated at the boundaries (endpoints) of the window. Possible values are:

"shrink" (default)

The window is truncated at the beginning and end of the array to exclude elements for which there is no source data.

For example, with a window of length 3, $y(1) = \text{movprod}(x(1:2))$, and $y(\text{end}) = \text{movprod}(x(\text{end}-1:\text{end}))$.

"discard"

Any y values that use a window extending beyond the original data array are deleted. For example, with a 10-element data vector and a window of length 3, the output will contain only 8 elements. The first element would require calculating the function over indices [0, 1, 2] and is therefore discarded. The last element would require calculating the function over indices [9, 10, 11] and is therefore discarded.

"fill"

Any window elements outside the data array are replaced by NaN. For example, with a window of length 3, $y(1) = \text{movprod}([\text{NaN}, x(1:2)])$, and $y(\text{end}) = \text{movprod}([x(\text{end}-1:\text{end}), \text{NaN}])$. This option usually results in y having NaN values at the boundaries, although it is influenced by how movprod handles NaN, and also by the property "nancond".

user_value

Any window elements outside the data array are replaced by the specified value *user_value* which must be a numeric scalar. For example, with a window of length 3, $y(1) = \text{movprod}([\text{user_value}, x(1:2)])$, and $y(\text{end}) = \text{movprod}([x(\text{end}-1:\text{end}), \text{user_value}])$. A common choice for *user_value* is 0.

"same"

Any window elements outside the data array are replaced by the value of x at the boundary. For example, with a window of length 3, $y(1) = \text{movprod}([x(1), x(1:2)])$, and $y(\text{end}) = \text{movprod}([x(\text{end}-1:\text{end}), x(\text{end})])$.

"periodic"

The window is wrapped so that any missing data elements are taken from the other side of the data. For example, with a window of length 3, $y(1) = \text{movprod}([x(\text{end}), x(1:2)])$, and $y(\text{end}) = \text{movprod}([x(\text{end}-1:\text{end}), x(1)])$.

"SamplePoints"

Caution: This option is not yet implemented.

Programming Note: This function is a wrapper which calls `movfun`. For additional options and documentation, See [\[movfun\]](#), page 719.

See also: [movfun], page 719, [movslice], page 721, [movmad], page 721, [movmax], page 723, [movmean], page 724, [movmedian], page 726, [movmin], page 727, [movstd], page 731, [movsum], page 732, [movvar], page 734.

```
y = movstd (x, wlen)
y = movstd (x, [na, nb])
y = movstd (..., dim)
y = movstd (..., "nancond")
y = movstd (..., property, value)
```

Calculate the moving standard deviation over a sliding window of length *wlen* on data *x*.

If *wlen* is a scalar, the function **movstd** is applied to a moving window of length *wlen*. When *wlen* is an odd number the window is symmetric and includes $(wlen - 1) / 2$ elements on either side of the central element. For example, when calculating the output at index 5 with a window length of 3, **movstd** uses data elements [4, 5, 6]. If *wlen* is an even number, the window is asymmetric and has *wlen*/2 elements to the left of the central element and *wlen*/2 - 1 elements to the right of the central element. For example, when calculating the output at index 5 with a window length of 4, **movstd** uses data elements [3, 4, 5, 6].

If *wlen* is an array with two elements [*nb*, *na*], the function is applied to a moving window *-nb:na*. This window includes *nb* number of elements *before* the current element and *na* number of elements *after* the current element. The current element is always included. For example, given *wlen* = [3, 0], the data used to calculate index 5 is [2, 3, 4, 5].

If the optional argument *dim* is given, operate along this dimension.

The optional string argument "nancond" controls whether NaN and NA values should be included ("includenan"), or excluded ("omitnan"), from the data passed to **movstd**. The default is "includenan". Caution: the "omitnan" option is not yet implemented.

The calculation can be controlled by specifying *property/value* pairs. Valid properties are

"Endpoints"

This property controls how results are calculated at the boundaries (endpoints) of the window. Possible values are:

"shrink" (default)

The window is truncated at the beginning and end of the array to exclude elements for which there is no source data.

For example, with a window of length 3, *y*(1) = **movstd** (*x*(1:2)), and *y*(*end*) = **movstd** (*x*(*end*-1:*end*)).

"discard"

Any *y* values that use a window extending beyond the original data array are deleted. For example, with a 10-element data vector and a window of length 3, the output will contain only 8 elements. The first element would require calculating the function over indices [0, 1, 2] and is therefore discarded.

The last element would require calculating the function over indices [9, 10, 11] and is therefore discarded.

"fill" Any window elements outside the data array are replaced by NaN. For example, with a window of length 3, $y(1) = \text{movstd}([\text{NaN}, x(1:2)])$, and $y(\text{end}) = \text{movstd}([x(\text{end}-1:\text{end}), \text{NaN}])$. This option usually results in y having NaN values at the boundaries, although it is influenced by how `movstd` handles NaN, and also by the property "nancond".

user_value Any window elements outside the data array are replaced by the specified value `user_value` which must be a numeric scalar. For example, with a window of length 3, $y(1) = \text{movstd}([\text{user_value}, x(1:2)])$, and $y(\text{end}) = \text{movstd}([x(\text{end}-1:\text{end}), \text{user_value}])$. A common choice for `user_value` is 0.

"same" Any window elements outside the data array are replaced by the value of x at the boundary. For example, with a window of length 3, $y(1) = \text{movstd}([x(1), x(1:2)])$, and $y(\text{end}) = \text{movstd}([x(\text{end}-1:\text{end}), x(\text{end})])$.

"periodic" The window is wrapped so that any missing data elements are taken from the other side of the data. For example, with a window of length 3, $y(1) = \text{movstd}([x(\text{end}), x(1:2)])$, and $y(\text{end}) = \text{movstd}([x(\text{end}-1:\text{end}), x(1)])$.

"SamplePoints"

Caution: This option is not yet implemented.

Programming Note: This function is a wrapper which calls `movfun`. For additional options and documentation, See [\[movfun\]](#), page 719.

See also: [\[movfun\]](#), page 719, [\[movslice\]](#), page 721, [\[movmad\]](#), page 721, [\[movmax\]](#), page 723, [\[movmean\]](#), page 724, [\[movmedian\]](#), page 726, [\[movmin\]](#), page 727, [\[movprod\]](#), page 729, [\[movsum\]](#), page 732, [\[movvar\]](#), page 734.

```
y = movsum (x, wlen)
y = movsum (x, [na, nb])
y = movsum (..., dim)
y = movsum (..., "nancond")
y = movsum (..., property, value)
```

Calculate the moving sum over a sliding window of length `wlen` on data `x`.

If `wlen` is a scalar, the function `movsum` is applied to a moving window of length `wlen`. When `wlen` is an odd number the window is symmetric and includes $(wlen - 1) / 2$ elements on either side of the central element. For example, when calculating the output at index 5 with a window length of 3, `movsum` uses data elements [4, 5, 6]. If `wlen` is an even number, the window is asymmetric and has `wlen/2` elements to the left of the central element and `wlen/2 - 1` elements to the right of the central

element. For example, when calculating the output at index 5 with a window length of 4, `movsum` uses data elements [3, 4, 5, 6].

If `wlen` is an array with two elements [`nb`, `na`], the function is applied to a moving window `-nb:na`. This window includes `nb` number of elements *before* the current element and `na` number of elements *after* the current element. The current element is always included. For example, given `wlen = [3, 0]`, the data used to calculate index 5 is [2, 3, 4, 5].

If the optional argument `dim` is given, operate along this dimension.

The optional string argument "nancond" controls whether NaN and NA values should be included ("includenan"), or excluded ("omitnan"), from the data passed to `movsum`. The default is "includenan". Caution: the "omitnan" option is not yet implemented.

The calculation can be controlled by specifying *property/value* pairs. Valid properties are

"Endpoints"

This property controls how results are calculated at the boundaries (endpoints) of the window. Possible values are:

"shrink" (default)

The window is truncated at the beginning and end of the array to exclude elements for which there is no source data.

For example, with a window of length 3, `y(1) = movsum (x(1:2))`, and `y(end) = movsum (x(end-1:end))`.

"discard"

Any `y` values that use a window extending beyond the original data array are deleted. For example, with a 10-element data vector and a window of length 3, the output will contain only 8 elements. The first element would require calculating the function over indices [0, 1, 2] and is therefore discarded. The last element would require calculating the function over indices [9, 10, 11] and is therefore discarded.

"fill"

Any window elements outside the data array are replaced by NaN. For example, with a window of length 3, `y(1) = movsum ([NaN, x(1:2)])`, and `y(end) = movsum ([x(end-1:end), NaN])`. This option usually results in `y` having NaN values at the boundaries, although it is influenced by how `movsum` handles NaN, and also by the property "nancond".

`user_value`

Any window elements outside the data array are replaced by the specified value `user_value` which must be a numeric scalar. For example, with a window of length 3, `y(1) = movsum ([user_value, x(1:2)])`, and `y(end) = movsum ([x(end-1:end), user_value])`. A common choice for `user_value` is 0.

"same"

Any window elements outside the data array are replaced by the value of `x` at the boundary. For example, with a window

of length 3, $y(1) = \text{movsum}([x(1), x(1:2)])$, and $y(\text{end}) = \text{movsum}([x(\text{end}-1:\text{end}), x(\text{end})])$.

"periodic"

The window is wrapped so that any missing data elements are taken from the other side of the data. For example, with a window of length 3, $y(1) = \text{movsum}([x(\text{end}), x(1:2)])$, and $y(\text{end}) = \text{movsum}([x(\text{end}-1:\text{end}), x(1)])$.

"SamplePoints"

Caution: This option is not yet implemented.

Programming Note: This function is a wrapper which calls `movfun`. For additional options and documentation, See [movfun], page 719.

See also: [movfun], page 719, [movslice], page 721, [movmad], page 721, [movmax], page 723, [movmean], page 724, [movmedian], page 726, [movmin], page 727, [movprod], page 729, [movstd], page 731, [movvar], page 734.

```
y = movvar (x, wlen)
y = movvar (x, [na, nb])
y = movvar (..., dim)
y = movvar (..., "nancond")
y = movvar (..., property, value)
```

Calculate the moving variance over a sliding window of length `wlen` on data `x`.

If `wlen` is a scalar, the function `var` is applied to a moving window of length `wlen`. When `wlen` is an odd number the window is symmetric and includes $(wlen - 1) / 2$ elements on either side of the central element. For example, when calculating the output at index 5 with a window length of 3, `movvar` uses data elements [4, 5, 6]. If `wlen` is an even number, the window is asymmetric and has `wlen/2` elements to the left of the central element and `wlen/2 - 1` elements to the right of the central element. For example, when calculating the output at index 5 with a window length of 4, `movvar` uses data elements [3, 4, 5, 6].

If `wlen` is an array with two elements `[nb, na]`, the function is applied to a moving window `-nb:na`. This window includes `nb` number of elements *before* the current element and `na` number of elements *after* the current element. The current element is always included. For example, given `wlen = [3, 0]`, the data used to calculate index 5 is [2, 3, 4, 5].

If the optional argument `dim` is given, operate along this dimension.

The optional string argument "nancond" controls whether NaN and NA values should be included ("includenan"), or excluded ("omitnan"), from the data passed to `var`. The default is "includenan". Caution: the "omitnan" option is not yet implemented. The calculation can be controlled by specifying `property/value` pairs. Valid properties are

"Endpoints"

This property controls how results are calculated at the boundaries (endpoints) of the window. Possible values are:

"shrink" (default)

The window is truncated at the beginning and end of the array to exclude elements for which there is no source data. For example, with a window of length 3, $y(1) = \text{var}(x(1:2))$, and $y(\text{end}) = \text{var}(x(\text{end}-1:\text{end}))$.

"discard"

Any y values that use a window extending beyond the original data array are deleted. For example, with a 10-element data vector and a window of length 3, the output will contain only 8 elements. The first element would require calculating the function over indices [0, 1, 2] and is therefore discarded. The last element would require calculating the function over indices [9, 10, 11] and is therefore discarded.

"fill"

Any window elements outside the data array are replaced by NaN. For example, with a window of length 3, $y(1) = \text{var}([\text{NaN}, x(1:2)])$, and $y(\text{end}) = \text{var}([x(\text{end}-1:\text{end}), \text{NaN}])$. This option usually results in y having NaN values at the boundaries, although it is influenced by how var handles NaN, and also by the property "nancond".

user_value

Any window elements outside the data array are replaced by the specified value *user_value* which must be a numeric scalar. For example, with a window of length 3, $y(1) = \text{var}([\text{user_value}, x(1:2)])$, and $y(\text{end}) = \text{var}([x(\text{end}-1:\text{end}), \text{user_value}])$. A common choice for *user_value* is 0.

"same"

Any window elements outside the data array are replaced by the value of x at the boundary. For example, with a window of length 3, $y(1) = \text{var}([x(1), x(1:2)])$, and $y(\text{end}) = \text{var}([x(\text{end}-1:\text{end}), x(\text{end})])$.

"periodic"

The window is wrapped so that any missing data elements are taken from the other side of the data. For example, with a window of length 3, $y(1) = \text{var}([x(\text{end}), x(1:2)])$, and $y(\text{end}) = \text{var}([x(\text{end}-1:\text{end}), x(1)])$.

"SamplePoints"

Caution: This option is not yet implemented.

Programming Note: This function is a wrapper which calls `movfun`. For additional options and documentation, See [\[movfun\]](#), page 719.

See also: [\[movfun\]](#), page 719, [\[movslice\]](#), page 721, [\[movmad\]](#), page 721, [\[movmax\]](#), page 723, [\[movmean\]](#), page 724, [\[movmedian\]](#), page 726, [\[movmin\]](#), page 727, [\[movprod\]](#), page 729, [\[movstd\]](#), page 731, [\[movsum\]](#), page 732.

26.3 Basic Statistical Functions

Octave supports various helpful statistical functions. Many are useful as initial steps to prepare a data set for further analysis. Others provide different measures from those of the basic descriptive statistics.

center (x)

center (x, dim)

Center data by subtracting its mean.

If x is a vector, subtract its mean.

If x is a matrix, do the above for each column.

If the optional argument dim is given, operate along this dimension.

Programming Note: `center` has obvious application for normalizing statistical data. It is also useful for improving the precision of general numerical calculations. Whenever there is a large value that is common to a batch of data, the mean can be subtracted off, the calculation performed, and then the mean added back to obtain the final answer.

See also: [zscore], page 736.

`z = zsore (x)`

`z = zsore (x, opt)`

`z = zsore (x, opt, dim)`

`[z, mu, sigma] = zsore (...)`

Compute the Z score of x .

If x is a vector, subtract its mean and divide by its standard deviation. If the standard deviation is zero, divide by 1 instead.

The optional parameter opt determines the normalization to use when computing the standard deviation and has the same definition as the corresponding parameter for `std`.

If x is a matrix, calculate along the first non-singleton dimension. If the third optional argument dim is given, operate along this dimension.

The optional outputs mu and $sigma$ contain the mean and standard deviation.

See also: [mean], page 711, [std], page 714, [center], page 736.

`n = histc (x, edges)`

`n = histc (x, edges, dim)`

`[n, idx] = histc (...)`

Compute histogram counts.

When x is a vector, the function counts the number of elements of x that fall in the histogram bins defined by `edges`. This must be a vector of monotonically increasing values that define the edges of the histogram bins. $n(k)$ contains the number of elements in x for which $\text{edges}(k) \leq x < \text{edges}(k + 1)$. The final element of n contains the number of elements of x exactly equal to the last element of `edges`.

When x is an N -dimensional array, the computation is carried out along dimension dim . If not specified dim defaults to the first non-singleton dimension.

When a second output argument is requested an index matrix is also returned. The *idx* matrix has the same size as *x*. Each element of *idx* contains the index of the histogram bin in which the corresponding element of *x* was counted.

See also: [\[hist\]](#), page 303.

[unique](#) function documented at [\[unique\]](#), page 743, is often useful for statistics.

```
c = nchoosek (n, k)
c = nchoosek (set, k)
```

Compute the binomial coefficient of *n* or list all possible combinations of a set of items.

If *n* is a scalar then calculate the binomial coefficient of *n* and *k* which is defined as

$$\binom{n}{k} = \frac{n(n-1)(n-2)\cdots(n-k+1)}{k!} = \frac{n!}{k!(n-k)!}$$

This is the number of combinations of *n* items taken in groups of size *k*.

If the first argument is a vector, *set*, then generate all combinations of the elements of *set*, taken *k* at a time, with one row per combination. The result *c* has *k* columns and *nchoosek (length (set), k)* rows.

For example:

How many ways can three items be grouped into pairs?

```
nchoosek (3, 2)
⇒ 3
```

What are the possible pairs?

```
nchoosek (1:3, 2)
⇒ 1 2
    1 3
    2 3
```

Programming Note: When calculating the binomial coefficient **nchoosek** works only for non-negative, integer arguments. Use **bincoeff** for non-integer and negative scalar arguments, or for computing many binomial coefficients at once with vector inputs for *n* or *k*.

See also: [\[bincoeff\]](#), page 531, [\[perms\]](#), page 737.

perms (v)

Generate all permutations of vector *v* with one row per permutation.

Results are returned in inverse lexicographic order. The result has size **factorial (n) * n**, where *n* is the length of *v*. Any repetitions are included in the output. To generate just the unique permutations use **unique (perms (v), "rows") (end:-1:1,:)**.

Example

```
perms ([1, 2, 3])
⇒
 3 2 1
 3 1 2
 2 3 1
 2 1 3
 1 3 2
 1 2 3
```

Programming Note: The maximum length of *v* should be less than or equal to 10 to limit memory consumption.

See also: [permute], page 482, [randperm], page 499, [nchoosek], page 737.

```
ranks (x)
ranks (x, dim)
ranks (x, dim, rtype)
```

Return the ranks (in the sense of order statistics) of *x* along the first non-singleton dimension adjusted for ties.

If the optional *dim* argument is given, operate along this dimension.

The optional parameter *rtype* determines how ties are handled. All examples below assume an input of [1, 2, 2, 4].

```
0 or "fractional" (default) for fractional ranking (1, 2.5,
2.5, 4);
1 or "competition" for competition ranking (1, 2, 2, 4);
2 or "modified" for modified competition ranking (1, 3, 3, 4);
3 or "ordinal" for ordinal ranking (1, 2, 3, 4);
4 or "dense" for dense ranking (1, 2, 2, 3).
```

See also: [spearman], page 740, [kendall], page 740.

```
run_count (x, n)
run_count (x, n, dim)
```

Count the upward runs along the first non-singleton dimension of *x* of length 1, 2, ..., *n*-1 and greater than or equal to *n*.

If the optional argument *dim* is given then operate along this dimension.

See also: [runlength], page 738.

```
count = runlength (x)
[count, value] = runlength (x)
```

Find the lengths of all sequences of common values.

count is a vector with the lengths of each repeated value.

The optional output *value* contains the value that was repeated in the sequence.

```
runlength ([2, 2, 0, 4, 4, 4, 0, 1, 1, 1, 1])
⇒ 2 1 3 1 4
```

See also: [run_count], page 738.

26.4 Correlation and Regression Analysis

`cov (x)`
`cov (x, opt)`
`cov (x, y)`
`cov (x, y, opt)`

Compute the covariance matrix.

If each row of x and y is an observation, and each column is a variable, then the (i, j) -th entry of `cov (x, y)` is the covariance between the i -th variable in x and the j -th variable in y .

$$\sigma_{ij} = \frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})$$

where \bar{x} and \bar{y} are the mean values of x and y .

If called with one argument, compute `cov (x, x)`, the covariance between the columns of x .

The argument *opt* determines the type of normalization to use. Valid values are

- 0: normalize with $N-1$, provides the best unbiased estimator of the covariance [default]
- 1: normalize with N , this provides the second moment around the mean

Compatibility Note:: Octave always treats rows of x and y as multivariate random variables. For two inputs, however, MATLAB treats x and y as two univariate distributions regardless of their shapes, and will calculate `cov ([x(:), y(:)])` whenever the number of elements in x and y are equal. This will result in a 2×2 matrix. Code relying on MATLAB's definition will need to be changed when running in Octave.

See also: [\[corr\]](#), page 739.

`corr (x)`
`corr (x, y)`

Compute matrix of correlation coefficients.

If each row of x and y is an observation and each column is a variable, then the (i, j) -th entry of `corr (x, y)` is the correlation between the i -th variable in x and the j -th variable in y .

$$\text{corr}(x, y) = \frac{\text{cov}(x, y)}{\text{std}(x) \text{std}(y)}$$

If called with one argument, compute `corr (x, x)`, the correlation between the columns of x .

See also: [\[cov\]](#), page 739.

`r = corrcoef (x)`
`r = corrcoef (x, y)`
`r = corrcoef (..., param, value, ...)`
`[r, p] = corrcoef (...)`
`[r, p, lci, hci] = corrcoef (...)`

Compute a matrix of correlation coefficients.

x is an array where each column contains a variable and each row is an observation. If a second input *y* (of the same size as *x*) is given then calculate the correlation coefficients between *x* and *y*.

param, value are optional pairs of parameters and values which modify the calculation. Valid options are:

- "alpha" Confidence level used for the bounds of the confidence interval, *lci* and *hci*. Default is 0.05, i.e., 95% confidence interval.
- "rows" Determine processing of NaN values. Acceptable values are "all", "complete", and "pairwise". Default is "all". With "complete", only the rows without NaN values are considered. With "pairwise", the selection of NaN-free rows is made for each pair of variables.

Output *r* is a matrix of Pearson's product moment correlation coefficients for each pair of variables.

Output *p* is a matrix of pair-wise p-values testing for the null hypothesis of a correlation coefficient of zero.

Outputs *lci* and *hci* are matrices containing, respectively, the lower and higher bounds of the 95% confidence interval of each correlation coefficient.

See also: [corr], page 739, [cov], page 739.

spearman (x)

spearman (x, y)

Compute Spearman's rank correlation coefficient ρ .

For two data vectors *x* and *y*, Spearman's ρ is the correlation coefficient of the ranks of *x* and *y*.

If *x* and *y* are drawn from independent distributions, ρ has zero mean and variance $1/(N - 1)$, where *N* is the length of the *x* and *y* vectors, and is asymptotically normally distributed.

spearman (x) is equivalent to **spearman (x, x)**.

See also: [ranks], page 738, [kendall], page 740.

kendall (x)

kendall (x, y)

Compute Kendall's τ .

For two data vectors *x*, *y* of common length *N*, Kendall's τ is the correlation of the signs of all rank differences of *x* and *y*; i.e., if both *x* and *y* have distinct entries, then

$$\tau = \frac{1}{N(N - 1)} \sum_{i,j} \text{sign}(q_i - q_j) \text{sign}(r_i - r_j)$$

in which the q_i and r_i are the ranks of *x* and *y*, respectively.

If *x* and *y* are drawn from independent distributions, Kendall's τ is asymptotically normal with mean 0 and variance $\frac{2(2N+5)}{9N(N-1)}$.

kendall (x) is equivalent to **kendall (x, x)**.

See also: [ranks], page 738, [spearman], page 740.

26.5 Distributions

Octave has functions for computing the Probability Density Function (PDF), the Cumulative Distribution function (CDF), and the quantile (the inverse of the CDF) for arbitrary user-defined distributions (discrete) and for experimental data (empirical).

The following table summarizes the supported distributions (in alphabetical order).

Distribution	PDF	CDF	Quantile
Univariate Discrete	discrete_pdf	discrete_cdf	discrete_inv
Empirical	empirical_pdf	empirical_cdf	empirical_inv

`discrete_pdf (x, v, p)`

For each element of x , compute the probability density function (PDF) at x of a univariate discrete distribution which assumes the values in v with probabilities p .

`discrete_cdf (x, v, p)`

For each element of x , compute the cumulative distribution function (CDF) at x of a univariate discrete distribution which assumes the values in v with probabilities p .

`discrete_inv (x, v, p)`

For each element of x , compute the quantile (the inverse of the CDF) at x of the univariate distribution which assumes the values in v with probabilities p .

`empirical_pdf (x, data)`

For each element of x , compute the probability density function (PDF) at x of the empirical distribution obtained from the univariate sample $data$.

`empirical_cdf (x, data)`

For each element of x , compute the cumulative distribution function (CDF) at x of the empirical distribution obtained from the univariate sample $data$.

`empirical_inv (x, data)`

For each element of x , compute the quantile (the inverse of the CDF) at x of the empirical distribution obtained from the univariate sample $data$.

26.6 Random Number Generation

Octave can generate random numbers from a large number of distributions. The random number generators are based on the random number generators described in [Section 16.3 \[Special Utility Matrices\], page 489](#).

The following table summarizes the available random number generators (in alphabetical order).

Distribution	Function
Univariate Discrete Distribution	discrete_rnd
Empirical Distribution	empirical_rnd
Exponential Distribution	rande
Gamma Distribution	randg
Poisson Distribution	randp
Standard Normal Distribution	randn
Uniform Distribution	rand
Uniform Distribution (integers)	randi

```
discrete_rnd (v, p)
discrete_rnd (v, p, r)
discrete_rnd (v, p, r, c, ...)
discrete_rnd (v, p, [sz])
```

Return a matrix of random samples from the univariate distribution which assumes the values in *v* with probabilities *p*.

When called with a single size argument, return a square matrix with the dimension specified. When called with more than one scalar argument the first two arguments are taken as the number of rows and columns and any further arguments specify additional matrix dimensions. The size may also be specified with a vector of dimensions *sz*.

If no size arguments are given then the result matrix is the common size of *v* and *p*.

```
empirical_rnd (data)
empirical_rnd (data, r)
empirical_rnd (data, r, c, ...)
empirical_rnd (data, [sz])
```

Return a matrix of random samples from the empirical distribution obtained from the univariate sample *data*.

When called with a single size argument, return a square matrix with the dimension specified. When called with more than one scalar argument the first two arguments are taken as the number of rows and columns and any further arguments specify additional matrix dimensions. The size may also be specified with a vector of dimensions *sz*.

If no size arguments are given then the result matrix is a random ordering of the sample *data*.

27 Sets

Octave has a number of functions for managing sets of data. A set is defined as a collection of unique elements and is typically represented by a vector of numbers sorted in ascending order. Any vector or matrix can be converted to a set by removing duplicates through the use of the `unique` function. However, it isn't necessary to explicitly create a set as all of the functions which operate on sets will convert their input to a set before proceeding.

```
unique (x)
unique (x, "rows")
[y, i, j] = unique (...)
[y, i, j] = unique (... , "first")
[y, i, j] = unique (... , "last")
```

Return the unique elements of x sorted in ascending order.

If the input x is a column vector then return a column vector; Otherwise, return a row vector. x may also be a cell array of strings.

If the optional argument "`rows`" is given then return the unique rows of x sorted in ascending order. The input must be a 2-D matrix to use this option.

If requested, return index vectors i and j such that $y = x(i)$ and $x = y(j)$.

Additionally, if i is a requested output then one of "`first`" or "`last`" may be given as an input. If "`last`" is specified, return the highest possible indices in i , otherwise, if "`first`" is specified, return the lowest. The default is "`last`".

See also: [\[union\]](#), page 744, [\[intersect\]](#), page 743, [\[setdiff\]](#), page 744, [\[setxor\]](#), page 744, [\[ismember\]](#), page 744.

27.1 Set Operations

Octave supports several basic set operations. Octave can compute the union, intersection, and difference of two sets. Octave also supports the *Exclusive Or* set operation.

The functions for set operations all work in the same way by accepting two input sets and returning a third set. As an example, assume that a and b contains two sets, then

```
union (a, b)
```

computes the union of the two sets.

Finally, determining whether elements belong to a set can be done with the `ismember` function. Because sets are ordered this operation is very efficient and is of order $O(\log_2(n))$ which is preferable to the `find` function which is of order $O(n)$.

```
c = intersect (a, b)
c = intersect (a, b, "rows")
[c, ia, ib] = intersect (...)
```

Return the unique elements common to both a and b sorted in ascending order.

If a and b are both row vectors then return a row vector; Otherwise, return a column vector. The inputs may also be cell arrays of strings.

If the optional input "`rows`" is given then return the common rows of a and b . The inputs must be 2-D matrices to use this option.

If requested, return index vectors *ia* and *ib* such that $c = a(ia)$ and $c = b(ib)$.

See also: [unique], page 743, [union], page 744, [setdiff], page 744, [setxor], page 744, [ismember], page 744.

```
c = union (a, b)
c = union (a, b, "rows")
[c, ia, ib] = union (...)
```

Return the unique elements that are in either *a* or *b* sorted in ascending order.

If *a* and *b* are both row vectors then return a row vector; Otherwise, return a column vector. The inputs may also be cell arrays of strings.

If the optional input "rows" is given then return rows that are in either *a* or *b*. The inputs must be 2-D matrices to use this option.

The optional outputs *ia* and *ib* are index vectors such that $a(ia)$ and $b(ib)$ are disjoint sets whose union is *c*.

See also: [unique], page 743, [intersect], page 743, [setdiff], page 744, [setxor], page 744, [ismember], page 744.

```
c = setdiff (a, b)
c = setdiff (a, b, "rows")
[c, ia] = setdiff (...)
```

Return the unique elements in *a* that are not in *b* sorted in ascending order.

If *a* is a row vector return a row vector; Otherwise, return a column vector. The inputs may also be cell arrays of strings.

If the optional input "rows" is given then return the rows in *a* that are not in *b*. The inputs must be 2-D matrices to use this option.

If requested, return the index vector *ia* such that $c = a(ia)$.

See also: [unique], page 743, [union], page 744, [intersect], page 743, [setxor], page 744, [ismember], page 744.

```
c = setxor (a, b)
c = setxor (a, b, "rows")
[c, ia, ib] = setxor (...)
```

Return the unique elements exclusive to sets *a* or *b* sorted in ascending order.

If *a* and *b* are both row vectors then return a row vector; Otherwise, return a column vector. The inputs may also be cell arrays of strings.

If the optional input "rows" is given then return the rows exclusive to sets *a* and *b*. The inputs must be 2-D matrices to use this option.

If requested, return index vectors *ia* and *ib* such that $a(ia)$ and $b(ib)$ are disjoint sets whose union is *c*.

See also: [unique], page 743, [union], page 744, [intersect], page 743, [setdiff], page 744, [ismember], page 744.

```
tf = ismember (a, s)
tf = ismember (a, s, "rows")
```

[tf, s_idx] = ismember (...)

Return a logical matrix *tf* with the same shape as *a* which is true (1) if the element in *a* is found in *s* and false (0) if it is not.

If a second output argument is requested then the index into *s* of each matching element is also returned.

```
a = [3, 10, 1];
s = [0:9];
[tf, s_idx] = ismember (a, s)
⇒ tf = [1, 0, 1]
⇒ s_idx = [4, 0, 2]
```

The inputs *a* and *s* may also be cell arrays.

```
a = {"abc"};
s = {"abc", "def"};
[tf, s_idx] = ismember (a, s)
⇒ tf = [1, 0]
⇒ s_idx = [1, 0]
```

If the optional third argument "rows" is given then compare rows in *a* with rows in *s*. The inputs must be 2-D matrices with the same number of columns to use this option.

```
a = [1:3; 5:7; 4:6];
s = [0:2; 1:3; 2:4; 3:5; 4:6];
[tf, s_idx] = ismember (a, s, "rows")
⇒ tf = logical ([1; 0; 1])
⇒ s_idx = [2; 0; 5];
```

See also: [\[lookup\]](#), page 478, [\[unique\]](#), page 743, [\[union\]](#), page 744, [\[intersect\]](#), page 743, [\[setdiff\]](#), page 744, [\[setxor\]](#), page 744.

powerset (a)

powerset (a, "rows")

Compute the powerset (all subsets) of the set *a*.

The set *a* must be a numerical matrix or a cell array of strings. The output will always be a cell array of either vectors or strings.

With the optional argument "rows", each row of the set *a* is considered one element of the set. The input must be a 2-D numeric matrix to use this argument.

See also: [\[unique\]](#), page 743, [\[union\]](#), page 744, [\[intersect\]](#), page 743, [\[setdiff\]](#), page 744, [\[setxor\]](#), page 744, [\[ismember\]](#), page 744.

28 Polynomial Manipulations

In Octave, a polynomial is represented by its coefficients (arranged in descending order). For example, a vector c of length $N + 1$ corresponds to the following polynomial of order N

$$p(x) = c_1x^N + \dots + c_Nx + c_{N+1}.$$

28.1 Evaluating Polynomials

The value of a polynomial represented by the vector c can be evaluated at the point x very easily, as the following example shows:

```
N = length (c) - 1;
val = dot (x.^ (N:-1:0), c);
```

While the above example shows how easy it is to compute the value of a polynomial, it isn't the most stable algorithm. With larger polynomials you should use more elegant algorithms, such as Horner's Method, which is exactly what the Octave function `polyval` does.

In the case where x is a square matrix, the polynomial given by c is still well-defined. As when x is a scalar the obvious implementation is easily expressed in Octave, but also in this case more elegant algorithms perform better. The `polyvalm` function provides such an algorithm.

```
y = polyval (p, x)
y = polyval (p, x, [], mu)
[y, dy] = polyval (p, x, s)
[y, dy] = polyval (p, x, s, mu)
```

Evaluate the polynomial p at the specified values of x .

If x is a vector or matrix, the polynomial is evaluated for each of the elements of x .

When mu is present, evaluate the polynomial for $(x-mu(1))/mu(2)$.

In addition to evaluating the polynomial, the second output represents the prediction interval, $y +/- dy$, which contains at least 50% of the future predictions. To calculate the prediction interval, the structured variable s , originating from `polyfit`, must be supplied.

See also: [\[polyvalm\]](#), page 747, [\[polyaffine\]](#), page 753, [\[polyfit\]](#), page 753, [\[roots\]](#), page 748, [\[poly\]](#), page 763.

polyvalm (c , x)

Evaluate a polynomial in the matrix sense.

`polyvalm` (c , x) will evaluate the polynomial in the matrix sense, i.e., matrix multiplication is used instead of element by element multiplication as used in `polyval`.

The argument x must be a square matrix.

See also: [\[polyval\]](#), page 747, [\[roots\]](#), page 748, [\[poly\]](#), page 763.

28.2 Finding Roots

Octave can find the roots of a given polynomial. This is done by computing the companion matrix of the polynomial (see the `compan` function for a definition), and then finding its eigenvalues.

`roots (c)`

Compute the roots of the polynomial c .

For a vector c with N components, return the roots of the polynomial

$$c_1x^{N-1} + \cdots + c_{N-1}x + c_N.$$

As an example, the following code finds the roots of the quadratic polynomial

$$p(x) = x^2 - 5.$$

```
c = [1, 0, -5];
roots (c)
⇒ 2.2361
⇒ -2.2361
```

Note that the true result is $\pm\sqrt{5}$ which is roughly ± 2.2361 .

See also: [\[poly\]](#), page 763, [\[compan\]](#), page 748, [\[fzero\]](#), page 602.

`z = polyeig (C0, C1, ..., Cl)`

`[v, z] = polyeig (C0, C1, ..., Cl)`

Solve the polynomial eigenvalue problem of degree l .

Given an $n \times n$ matrix polynomial

$$C(s) = C_0 + C_1 s + \dots + C_l s^l$$

`polyeig` solves the eigenvalue problem

$$(C_0 + C_1 z + \dots + C_l z^l) v = 0.$$

Note that the eigenvalues z are the zeros of the matrix polynomial. z is a row vector with $n \times l$ elements. v is a matrix ($n \times n \times l$) with columns that correspond to the eigenvectors.

See also: [\[eig\]](#), page 547, [\[eigs\]](#), page 644, [\[compan\]](#), page 748.

`compan (c)`

Compute the companion matrix corresponding to polynomial coefficient vector c .

The companion matrix is

$$A = \begin{bmatrix} -c_2/c_1 & -c_3/c_1 & \cdots & -c_N/c_1 & -c_{N+1}/c_1 \\ 1 & 0 & \cdots & 0 & 0 \\ 0 & 1 & \cdots & 0 & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \cdots & 1 & 0 \end{bmatrix}.$$

The eigenvalues of the companion matrix are equal to the roots of the polynomial.

See also: [\[roots\]](#), page 748, [\[poly\]](#), page 763, [\[eig\]](#), page 547.

```
[multp, idxp] = mpoles (p)
[multp, idxp] = mpoles (p, tol)
[multp, idxp] = mpoles (p, tol, reorder)
```

Identify unique poles in *p* and their associated multiplicity.

The output is ordered from largest pole to smallest pole.

If the relative difference of two poles is less than *tol* then they are considered to be multiples. The default value for *tol* is 0.001.

If the optional parameter *reorder* is zero, poles are not sorted.

The output *multp* is a vector specifying the multiplicity of the poles. *multp(n)* refers to the multiplicity of the *N*th pole *p(idxp(n))*.

For example:

```
p = [2 3 1 1 2];
[m, n] = mpoles (p)
⇒ m = [1; 1; 2; 1; 2]
⇒ n = [2; 5; 1; 4; 3]
⇒ p(n) = [3, 2, 2, 1, 1]
```

See also: [residue], page 751, [poly], page 763, [roots], page 748, [conv], page 749, [deconv], page 750.

28.3 Products of Polynomials

```
conv (a, b)
conv (a, b, shape)
```

Convolve two vectors *a* and *b*.

When *a* and *b* are the coefficient vectors of two polynomials, the convolution represents the coefficient vector of the product polynomial.

The size of the result is determined by the optional *shape* argument which takes the following values

shape = "full"

Return the full convolution. (default) The result is a vector with length equal to *length(a)* + *length(b)* - 1.

shape = "same"

Return the central part of the convolution with the same size as *a*.

shape = "valid"

Return only the parts which do not include zero-padded edges. The size of the result is *max(size(a) - size(b) + 1, 0)*.

See also: [deconv], page 750, [conv2], page 750, [convn], page 749, [fftconv], page 794.

```
C = convn (A, B)
```

```
C = convn (A, B, shape)
```

Return the n-D convolution of *A* and *B*.

The size of the result is determined by the optional *shape* argument which takes the following values

```
shape = "full"
    Return the full convolution. (default)

shape = "same"
    Return central part of the convolution with the same size as A. The
    central part of the convolution begins at the indices floor ([size(B)/2]
    + 1).

shape = "valid"
    Return only the parts which do not include zero-padded edges. The size
    of the result is max (size (A) - size (B) + 1, 0).
```

See also: [conv2], page 750, [conv], page 749.

```
b = deconv (y, a)
[b, r] = deconv (y, a)
    Deconvolve two vectors (polynomial division).
[b, r] = deconv (y, a) solves for b and r such that y = conv (a, b) + r.
If y and a are polynomial coefficient vectors, b will contain the coefficients of the
polynomial quotient and r will be a remainder polynomial of lowest order.
```

See also: [conv], page 749, [residue], page 751.

```
conv2 (A, B)
conv2 (v1, v2, m)
conv2 (..., shape)
    Return the 2-D convolution of A and B.
```

The size of the result is determined by the optional *shape* argument which takes the following values

```
shape = "full"
    Return the full convolution. (default)

shape = "same"
    Return the central part of the convolution with the same size as A. The
    central part of the convolution begins at the indices floor ([size(B)/2]
    + 1).

shape = "valid"
    Return only the parts which do not include zero-padded edges. The size
    of the result is max (size (A) - size (B) + 1, 0).
```

When the third argument is a matrix, return the convolution of the matrix *m* by the vector *v1* in the column direction and by the vector *v2* in the row direction.

See also: [conv], page 749, [convn], page 749.

```
q = polygcd (b, a)
q = polygcd (b, a, tol)
    Find the greatest common divisor of two polynomials.
```

This is equivalent to the polynomial found by multiplying together all the common roots. Together with deconv, you can reduce a ratio of two polynomials.

The tolerance *tol* defaults to `sqrt (eps)`.

Caution: This is a numerically unstable algorithm and should not be used on large polynomials.

Example code:

```
polygcd (poly (1:8), poly (3:12)) - poly (3:8)
⇒ [ 0, 0, 0, 0, 0, 0, 0 ]
deconv (poly (1:8), polygcd (poly (1:8), poly (3:12))) - poly (1:2)
⇒ [ 0, 0, 0 ]
```

See also: `[poly]`, page 763, `[roots]`, page 748, `[conv]`, page 749, `[deconv]`, page 750, `[residue]`, page 751.

```
[r, p, k, e] = residue (b, a)
[b, a] = residue (r, p, k)
[b, a] = residue (r, p, k, e)
```

The first calling form computes the partial fraction expansion for the quotient of the polynomials, *b* and *a*.

The quotient is defined as

$$\frac{B(s)}{A(s)} = \sum_{m=1}^M \frac{r_m}{(s - p_m)_m^e} + \sum_{i=1}^N k_i s^{N-i}.$$

where *M* is the number of poles (the length of the *r*, *p*, and *e*), the *k* vector is a polynomial of order *N* – 1 representing the direct contribution, and the *e* vector specifies the multiplicity of the *m*-th residue's pole.

For example,

```
b = [1, 1, 1];
a = [1, -5, 8, -4];
[r, p, k, e] = residue (b, a)
⇒ r = [-2; 7; 3]
⇒ p = [2; 2; 1]
⇒ k = [] (0x0)
⇒ e = [1; 2; 1]
```

which represents the following partial fraction expansion

$$\frac{s^2 + s + 1}{s^3 - 5s^2 + 8s - 4} = \frac{-2}{s - 2} + \frac{7}{(s - 2)^2} + \frac{3}{s - 1}$$

The second calling form performs the inverse operation and computes the reconstituted quotient of polynomials, *b*(*s*)/*a*(*s*), from the partial fraction expansion; represented by the residues, poles, and a direct polynomial specified by *r*, *p* and *k*, and the pole multiplicity *e*.

If the multiplicity, *e*, is not explicitly specified the multiplicity is determined by the function `mpoles`.

For example:

```
r = [-2; 7; 3];
p = [2; 2; 1];
k = [1, 0];
[b, a] = residue (r, p, k)
⇒ b = [1, -5, 9, -3, 1]
⇒ a = [1, -5, 8, -4]
```

where `mpoles` is used to determine `e = [1; 2; 1]`

Alternatively the multiplicity may be defined explicitly, for example,

```
r = [7; 3; -2];
p = [2; 1; 2];
k = [1, 0];
e = [2; 1; 1];
[b, a] = residue (r, p, k, e)
⇒ b = [1, -5, 9, -3, 1]
⇒ a = [1, -5, 8, -4]
```

which represents the following partial fraction expansion

$$\frac{-2}{s-2} + \frac{7}{(s-2)^2} + \frac{3}{s-1} + s = \frac{s^4 - 5s^3 + 9s^2 - 3s + 1}{s^3 - 5s^2 + 8s - 4}$$

See also: `[mpoles]`, page 749, `[poly]`, page 763, `[roots]`, page 748, `[conv]`, page 749, `[deconv]`, page 750.

28.4 Derivatives / Integrals / Transforms

Octave comes with functions for computing the derivative and the integral of a polynomial. The functions `polyder` and `polyint` both return new polynomials describing the result. As an example we'll compute the definite integral of $p(x) = x^2 + 1$ from 0 to 3.

```
c = [1, 0, 1];
integral = polyint (c);
area = polyval (integral, 3) - polyval (integral, 0)
⇒ 12

polyder (p)
[k] = polyder (a, b)
[q, d] = polyder (b, a)
```

Return the coefficients of the derivative of the polynomial whose coefficients are given by the vector `p`.

If a pair of polynomials is given, return the derivative of the product `a * b`.

If two inputs and two outputs are given, return the derivative of the polynomial quotient b/a . The quotient numerator is in `q` and the denominator in `d`.

See also: `[polyint]`, page 753, `[polyval]`, page 747, `[polyreduce]`, page 763.

polyint (p)
polyint (p, k)

Return the coefficients of the integral of the polynomial whose coefficients are represented by the vector *p*.

The variable *k* is the constant of integration, which by default is set to zero.

See also: [polyder], page 752, [polyval], page 747.

polyaffine (f, mu)

Return the coefficients of the polynomial vector *f* after an affine transformation.

If *f* is the vector representing the polynomial $f(x)$, then $g = \text{polyaffine}(f, mu)$ is the vector representing:

$$g(x) = f((x - mu(1)) / mu(2))$$

See also: [polyval], page 747, [polyfit], page 753.

28.5 Polynomial Interpolation

Octave comes with good support for various kinds of interpolation, most of which are described in Chapter 29 [Interpolation], page 765. One simple alternative to the functions described in the aforementioned chapter, is to fit a single polynomial, or a piecewise polynomial (spline) to some given data points. To avoid a highly fluctuating polynomial, one most often wants to fit a low-order polynomial to data. This usually means that it is necessary to fit the polynomial in a least-squares sense, which just is what the **polyfit** function does.

***p* = polyfit (*x*, *y*, *n*)**
[*p*, *s*] = polyfit (*x*, *y*, *n*)
[*p*, *s*, *mu*] = polyfit (*x*, *y*, *n*)

Return the coefficients of a polynomial $p(x)$ of degree *n* that minimizes the least-squares-error of the fit to the points [*x*, *y*].

If *n* is a logical vector, it is used as a mask to selectively force the corresponding polynomial coefficients to be used or ignored.

The polynomial coefficients are returned in a row vector.

The optional output *s* is a structure containing the following fields:

‘R’	Triangular factor R from the QR decomposition.
‘X’	The Vandermonde matrix used to compute the polynomial coefficients.
‘C’	The unscaled covariance matrix, formally equal to the inverse of x^*x , but computed in a way minimizing roundoff error propagation.
‘df’	The degrees of freedom.
‘normr’	The norm of the residuals.
‘yf’	The values of the polynomial for each value of x.

The second output may be used by **polyval** to calculate the statistical error limits of the predicted values. In particular, the standard deviation of *p* coefficients is given by

```
sqrt (diag (s.C)/s.df)*s.normr.
```

When the third output, *mu*, is present the coefficients, *p*, are associated with a polynomial in

```
xhat = (x - mu(1)) / mu(2)
```

where *mu*(1) = mean (*x*), and *mu*(2) = std (*x*).

This linear transformation of *x* improves the numerical stability of the fit.

See also: [polyval], page 747, [polyaffine], page 753, [roots], page 748, [vander], page 507, [zscore], page 736.

In situations where a single polynomial isn't good enough, a solution is to use several polynomials pieced together. The function **splinefit** fits a piecewise polynomial (spline) to a set of data.

```
pp = splinefit (x, y, breaks)
pp = splinefit (x, y, p)
pp = splinefit (..., "periodic", periodic)
pp = splinefit (..., "robust", robust)
pp = splinefit (..., "beta", beta)
pp = splinefit (..., "order", order)
pp = splinefit (..., "constraints", constraints)
```

Fit a piecewise cubic spline with breaks (knots) *breaks* to the noisy data, *x* and *y*.

x is a vector, and *y* is a vector or N-D array. If *y* is an N-D array, then *x*(*j*) is matched to *y*(:,:, :,*j*).

p is a positive integer defining the number of intervals along *x*, and *p*+1 is the number of breaks. The number of points in each interval differ by no more than 1.

The optional property *periodic* is a logical value which specifies whether a periodic boundary condition is applied to the spline. The length of the period is **max** (*breaks*) - **min** (*breaks*). The default value is **false**.

The optional property *robust* is a logical value which specifies if robust fitting is to be applied to reduce the influence of outlying data points. Three iterations of weighted least squares are performed. Weights are computed from previous residuals. The sensitivity of outlier identification is controlled by the property *beta*. The value of *beta* is restricted to the range, $0 < \text{beta} < 1$. The default value is *beta* = $1/2$. Values close to 0 give all data equal weighting. Increasing values of *beta* reduce the influence of outlying data. Values close to unity may cause instability or rank deficiency.

The fitted spline is returned as a piecewise polynomial, *pp*, and may be evaluated using **ppval**.

The splines are constructed of polynomials with degree *order*. The default is a cubic, *order*=3. A spline with *P* pieces has *P*+*order* degrees of freedom. With periodic boundary conditions the degrees of freedom are reduced to *P*.

The optional property, *constraints*, is a structure specifying linear constraints on the fit. The structure has three fields, "xc", "yc", and "cc".

"xc" Vector of the x-locations of the constraints.

"yc" Constraining values at the locations *xc*. The default is an array of zeros.

"cc" Coefficients (matrix). The default is an array of ones. The number of rows is limited to the order of the piecewise polynomials, *order*.

Constraints are linear combinations of derivatives of order 0 to *order*-1 according to

$$cc(1,j) \cdot y(xc(j)) + cc(2,j) \cdot y'(xc(j)) + \dots = yc(:,\dots,:,j).$$

See also: [interp1], page 765, [unmkpp], page 761, [ppval], page 762, [spline], page 768, [pchip], page 800, [ppder], page 762, [ppint], page 762, [ppjumps], page 762.

The number of *breaks* (or knots) used to construct the piecewise polynomial is a significant factor in suppressing the noise present in the input data, *x* and *y*. This is demonstrated by the example below.

```

x = 2 * pi * rand (1, 200);
y = sin (x) + sin (2 * x) + 0.2 * randn (size (x));
## Uniform breaks
breaks = linspace (0, 2 * pi, 41); % 41 breaks, 40 pieces
pp1 = splinefit (x, y, breaks);
## Breaks interpolated from data
pp2 = splinefit (x, y, 10); % 11 breaks, 10 pieces
## Plot
xx = linspace (0, 2 * pi, 400);
y1 = ppval (pp1, xx);
y2 = ppval (pp2, xx);
plot (x, y, ".", xx, [y1; y2])
axis tight
ylim auto
legend ({"data", "41 breaks, 40 pieces", "11 breaks, 10 pieces"})

```

The result of which can be seen in Figure 28.1.




Figure 28.1: Comparison of fitting a piecewise polynomial with 41 breaks to one with 11 breaks. The fit with the large number of breaks exhibits a fast ripple that is not present in the underlying function.

The piecewise polynomial fit, provided by `splinefit`, has continuous derivatives up to the *order*-1. For example, a cubic fit has continuous first and second derivatives. This is demonstrated by the code

```
## Data (200 points)
x = 2 * pi * rand (1, 200);
y = sin (x) + sin (2 * x) + 0.1 * randn (size (x));
## Piecewise constant
pp1 = splinefit (x, y, 8, "order", 0);
## Piecewise linear
pp2 = splinefit (x, y, 8, "order", 1);
## Piecewise quadratic
pp3 = splinefit (x, y, 8, "order", 2);
## Piecewise cubic
pp4 = splinefit (x, y, 8, "order", 3);
## Piecewise quartic
pp5 = splinefit (x, y, 8, "order", 4);
## Plot
xx = linspace (0, 2 * pi, 400);
y1 = ppval (pp1, xx);
y2 = ppval (pp2, xx);
y3 = ppval (pp3, xx);
y4 = ppval (pp4, xx);
y5 = ppval (pp5, xx);
plot (x, y, ".", xx, [y1; y2; y3; y4; y5])
axis tight
```

```
ylim auto
legend ({"data", "order 0", "order 1", "order 2", "order 3", "order 4"})
```

The result of which can be seen in Figure 28.2.




Figure 28.2: Comparison of a piecewise constant, linear, quadratic, cubic, and quartic polynomials with 8 breaks to noisy data. The higher order solutions more accurately represent the underlying function, but come with the expense of computational complexity.

When the underlying function to provide a fit to is periodic, `splinefit` is able to apply the boundary conditions needed to manifest a periodic fit. This is demonstrated by the code below.

```
## Data (100 points)
x = 2 * pi * [0, (rand (1, 98)), 1];
y = sin (x) - cos (2 * x) + 0.2 * randn (size (x));
## No constraints
pp1 = splinefit (x, y, 10, "order", 5);
## Periodic boundaries
pp2 = splinefit (x, y, 10, "order", 5, "periodic", true);
## Plot
xx = linspace (0, 2 * pi, 400);
y1 = ppval (pp1, xx);
y2 = ppval (pp2, xx);
plot (x, y, ".", xx, [y1; y2])
axis tight
ylim auto
legend ({"data", "no constraints", "periodic"})
```

The result of which can be seen in Figure 28.3.




Figure 28.3: Comparison of piecewise polynomial fits to a noisy periodic function with, and without, periodic boundary conditions.

More complex constraints may be added as well. For example, the code below illustrates a periodic fit with values that have been clamped at the endpoints, and a second periodic fit which is hinged at the endpoints.

```

## Data (200 points)
x = 2 * pi * rand (1, 200);
y = sin (2 * x) + 0.1 * randn (size (x));
## Breaks
breaks = linspace (0, 2 * pi, 10);
## Clamped endpoints, y = y' = 0
xc = [0, 0, 2*pi, 2*pi];
cc = [(eye (2)), (eye (2))];
con = struct ("xc", xc, "cc", cc);
pp1 = splinefit (x, y, breaks, "constraints", con);
## Hinged periodic endpoints, y = 0
con = struct ("xc", 0);
pp2 = splinefit (x, y, breaks, "constraints", con, "periodic", true);
## Plot
xx = linspace (0, 2 * pi, 400);
y1 = ppval (pp1, xx);
y2 = ppval (pp2, xx);
plot (x, y, ".", xx, [y1; y2])
axis tight
ylim auto
legend ({"data", "clamped", "hinged periodic"})

```

The result of which can be seen in [Figure 28.4](#).




Figure 28.4: Comparison of two periodic piecewise cubic fits to a noisy periodic signal. One fit has its endpoints clamped and the second has its endpoints hinged.

The `splinefit` function also provides the convenience of a *robust* fitting, where the effect of outlying data is reduced. In the example below, three different fits are provided. Two with differing levels of outlier suppression and a third illustrating the non-robust solution.

```

## Data
x = linspace (0, 2*pi, 200);
y = sin (x) + sin (2 * x) + 0.05 * randn (size (x));
## Add outliers
x = [x, linspace(0,2*pi,60)];
y = [y, -ones(1,60)];
## Fit splines with hinged conditions
con = struct ("xc", [0, 2*pi]);
## Robust fitting, beta = 0.25
pp1 = splinefit (x, y, 8, "constraints", con, "beta", 0.25);
## Robust fitting, beta = 0.75
pp2 = splinefit (x, y, 8, "constraints", con, "beta", 0.75);
## No robust fitting
pp3 = splinefit (x, y, 8, "constraints", con);
## Plot
xx = linspace (0, 2*pi, 400);
y1 = ppval (pp1, xx);
y2 = ppval (pp2, xx);
y3 = ppval (pp3, xx);
plot (x, y, ".", xx, [y1; y2; y3])
legend ({"data with outliers","robust, beta = 0.25", ...
"robust, beta = 0.75", "no robust fitting"})
axis tight

```

```
ylim auto
```

The result of which can be seen in Figure 28.5.




Figure 28.5: Comparison of two different levels of robust fitting ($\beta = 0.25$ and 0.75) to noisy data combined with outlying data. A conventional fit, without robust fitting ($\beta = 0$) is also included.

A very specific form of polynomial interpretation is the Padé approximant. For control systems, a continuous-time delay can be modeled very simply with the approximant.

```
[num, den] = padecoef (T)
[num, den] = padecoef (T, N)
```

Compute the N th-order Padé approximant of the continuous-time delay T in transfer function form.

The Padé approximant of e^{-sT} is defined by the following equation

$$e^{-sT} \approx \frac{P_n(s)}{Q_n(s)}$$

where both $P_n(s)$ and $Q_n(s)$ are N^{th} -order rational functions defined by the following expressions

$$P_n(s) = \sum_{k=0}^N \frac{(2N-k)!N!}{(2N)!k!(N-k)!} (-sT)^k$$

$$Q_n(s) = P_n(-s)$$

The inputs T and N must be non-negative numeric scalars. If N is unspecified it defaults to 1.

The output row vectors num and den contain the numerator and denominator coefficients in descending powers of s . Both are N th-order polynomials.

For example:

```
t = 0.1;
n = 4;
[num, den] = padecoef (t, n)
⇒ num =
1.0000e-04 -2.0000e-02 1.8000e+00 -8.4000e+01 1.6800e+03

⇒ den =
1.0000e-04 2.0000e-02 1.8000e+00 8.4000e+01 1.6800e+03
```

The function, `ppval`, evaluates the piecewise polynomials, created by `mkpp` or other means, and `unmkpp` returns detailed information about the piecewise polynomial.

The following example shows how to combine two linear functions and a quadratic into one function. Each of these functions is expressed on adjoined intervals.

```
x = [-2, -1, 1, 2];
p = [ 0, 1, 0;
      1, -2, 1;
      0, -1, 1 ];
pp = mkpp (x, p);
xi = linspace (-2, 2, 50);
yi = ppval (pp, xi);
plot (xi, yi);

pp = mkpp (breaks, coefs)
pp = mkpp (breaks, coefs, d)
```

Construct a piecewise polynomial (`pp`) structure from sample points `breaks` and coefficients `coefs`.

`breaks` must be a vector of strictly increasing values. The number of intervals is given by `ni = length (breaks) - 1`.

When m is the polynomial order `coefs` must be of size: ni -by- $(m + 1)$.

The i -th row of `coefs`, `coefs(i , :)`, contains the coefficients for the polynomial over the i -th interval, ordered from highest (m) to lowest (0) degree.

`coefs` may also be a multi-dimensional array, specifying a vector-valued or array-valued polynomial. In that case the polynomial order m is defined by the length of the last dimension of `coefs`. The size of first dimension(s) are given by the scalar or vector `d`. If `d` is not given it is set to 1. In this case `p(r , i , :)` contains the coefficients for the r -th polynomial defined on interval i . In any case `coefs` is reshaped to a 2-D matrix of size `[ni*prod(d) m]`.

Programming Note: `ppval` evaluates polynomials at `xi - breaks(i)`, i.e., it subtracts the lower endpoint of the current interval from `xi`. This must be taken into account when creating piecewise polynomials objects with `mkpp`.

See also: [\[unmkpp\]](#), page 761, [\[ppval\]](#), page 762, [\[spline\]](#), page 768, [\[pchip\]](#), page 800, [\[ppder\]](#), page 762, [\[ppint\]](#), page 762, [\[ppjumps\]](#), page 762.

`[x, p, n, k, d] = unmkpp (pp)`

Extract the components of a piecewise polynomial structure `pp`.

This function is the inverse of `mkpp`: it extracts the inputs to `mkpp` needed to create the piecewise polynomial structure `PP`. The code below makes this relation explicit:

```
[breaks, coefs, numinter, order, dim] = unmkpp (pp);
pp2 = mkpp (breaks, coefs, dim);
```

The piecewise polynomial structure `pp2` obtained in this way, is identical to the original `pp`. The same can be obtained by directly accessing the fields of the structure `pp`.

The components are:

<code>x</code>	Sample points or breaks.
<code>p</code>	Polynomial coefficients for points in sample interval. <code>p(i, :)</code> contains the coefficients for the polynomial over interval <code>i</code> ordered from highest to lowest degree. If <code>d > 1</code> , then <code>p</code> is a matrix of size <code>[n*prod(d) m]</code> , where the <code>i + (1:d)</code> rows are the coefficients of all the <code>d</code> polynomials in the interval <code>i</code> .
<code>n</code>	Number of polynomial pieces or intervals, <code>n = length (x) - 1</code> .
<code>k</code>	Order of the polynomial plus 1.
<code>d</code>	Number of polynomials defined for each interval.

See also: [\[mkpp\]](#), page 761, [\[ppval\]](#), page 762, [\[spline\]](#), page 768, [\[pchip\]](#), page 800.

`yi = ppval (pp, xi)`

Evaluate the piecewise polynomial structure `pp` at the points `xi`.

If `pp` describes a scalar polynomial function, the result is an array of the same shape as `xi`. Otherwise, the size of the result is `[pp.dim, length(xi)]` if `xi` is a vector, or `[pp.dim, size(xi)]` if it is a multi-dimensional array.

See also: [\[mkpp\]](#), page 761, [\[unmkpp\]](#), page 761, [\[spline\]](#), page 768, [\[pchip\]](#), page 800.

`ppd = ppder (pp)`

`ppd = ppder (pp, m)`

Compute the piecewise `m`-th derivative of a piecewise polynomial struct `pp`.

If `m` is omitted the first derivative is calculated.

See also: [\[mkpp\]](#), page 761, [\[ppval\]](#), page 762, [\[ppint\]](#), page 762.

`ppi = ppint (pp)`

`ppi = ppint (pp, c)`

Compute the integral of the piecewise polynomial struct `pp`.

`c`, if given, is the constant of integration.

See also: [\[mkpp\]](#), page 761, [\[ppval\]](#), page 762, [\[ppder\]](#), page 762.

`jumps = ppjumps (pp)`

Evaluate the boundary jumps of a piecewise polynomial.

If there are `n` intervals, and the dimensionality of `pp` is `d`, the resulting array has dimensions `[d, n-1]`.

See also: [\[mkpp\]](#), page 761.

28.6 Miscellaneous Functions

poly (A)

poly (x)

If A is a square N -by- N matrix, **poly (A)** is the row vector of the coefficients of $\det(z * \text{eye}(N) - A)$, the characteristic polynomial of A .

For example, the following code finds the eigenvalues of A which are the roots of **poly (A)**.

```
roots (poly (eye (3)))
⇒ 1.00001 + 0.00001i
    1.00001 - 0.00001i
    0.99999 + 0.00000i
```

In fact, all three eigenvalues are exactly 1 which emphasizes that for numerical performance the **eig** function should be used to compute eigenvalues.

If x is a vector, **poly (x)** is a vector of the coefficients of the polynomial whose roots are the elements of x . That is, if c is a polynomial, then the elements of $d = \text{roots}(\text{poly}(c))$ are contained in c . The vectors c and d are not identical, however, due to sorting and numerical errors.

See also: [\[roots\]](#), page 748, [\[eig\]](#), page 547.

polyout (c)

polyout (c, x)

str = polyout (...)

Display a formatted version of the polynomial c .

The formatted polynomial

$$c(x) = c_1 x^n + \dots + c_n x + c_{n+1}$$

is returned as a string or written to the screen if **nargout** is zero.

The second argument x specifies the variable name to use for each term and defaults to the string "s".

See also: [\[polyreduce\]](#), page 763.

polyreduce (c)

Reduce a polynomial coefficient vector to a minimum number of terms by stripping off any leading zeros.

See also: [\[polyout\]](#), page 763.

29 Interpolation

29.1 One-dimensional Interpolation

Octave supports several methods for one-dimensional interpolation, most of which are described in this section. Section 28.5 [Polynomial Interpolation], page 753, and Section 30.4 [Interpolation on Scattered Data], page 788, describe additional methods.

```
yi = interp1 (x, y, xi)
yi = interp1 (y, xi)
yi = interp1 (..., method)
yi = interp1 (..., extrap)
yi = interp1 (..., "left")
yi = interp1 (..., "right")
pp = interp1 (..., "pp")
```

One-dimensional interpolation.

Interpolate input data to determine the value of *yi* at the points *xi*. If not specified, *x* is taken to be the indices of *y* (`1:length (y)`). If *y* is a matrix or an N-dimensional array, the interpolation is performed on each column of *y*.

The interpolation *method* is one of:

"nearest"

Return the nearest neighbor.

"previous"

Return the previous neighbor.

"next"

Return the next neighbor.

"linear" (default)

Linear interpolation from nearest neighbors.

"pchip"

Piecewise cubic Hermite interpolating polynomial—shape-preserving interpolation with smooth first derivative.

"cubic"

Cubic interpolation (same as "`pchip`").

"spline"

Cubic spline interpolation—smooth first and second derivatives throughout the curve.

Adding `**` to the start of any method above forces `interp1` to assume that *x* is uniformly spaced, and only *x(1)* and *x(2)* are referenced. This is usually faster, and is never slower. The default method is "`linear`".

If *extrap* is the string "`extrap`", then extrapolate values beyond the endpoints using the current *method*. If *extrap* is a number, then replace values beyond the endpoints with that number. When unspecified, *extrap* defaults to `NA`.

If the string argument "`pp`" is specified, then *xi* should not be supplied and `interp1` returns a piecewise polynomial object. This object can later be used with `ppval` to evaluate the interpolation. There is an equivalence, such that `ppval (interp1 (x, y, method, "pp"), xi) == interp1 (x, y, xi, method, "extrap")`.

Duplicate points in x specify a discontinuous interpolant. There may be at most 2 consecutive points with the same value. If x is increasing, the default discontinuous interpolant is right-continuous. If x is decreasing, the default discontinuous interpolant is left-continuous. The continuity condition of the interpolant may be specified by using the options "left" or "right" to select a left-continuous or right-continuous interpolant, respectively. Discontinuous interpolation is only allowed for "nearest" and "linear" methods; in all other cases, the x -values must be unique.

An example of the use of `interp1` is

```
xf = [0:0.05:10];
yf = sin (2*pi*xf/5);
xp = [0:10];
yp = sin (2*pi*xp/5);
lin = interp1 (xp, yp, xf);
near = interp1 (xp, yp, xf, "nearest");
pch = interp1 (xp, yp, xf, "pchip");
spl = interp1 (xp, yp, xf, "spline");
plot (xf,yf,"r-", xf,near,"g-", xf,lin,"b-", xf,pch,"c-", xf,spl,"m",
      xp,yp,"r*");
legend ("original", "nearest", "linear", "pchip", "spline");
```

See also: [\[pchip\]](#), page 800, [\[spline\]](#), page 768, [\[interpft\]](#), page 767, [\[interp2\]](#), page 769, [\[interp3\]](#), page 770, [\[interpn\]](#), page 771.

There are some important differences between the various interpolation methods. The "spline" method enforces that both the first and second derivatives of the interpolated values have a continuous derivative, whereas the other methods do not. This means that the results of the "spline" method are generally smoother. If the function to be interpolated is in fact smooth, then "spline" will give excellent results. However, if the function to be evaluated is in some manner discontinuous, then "pchip" interpolation might give better results.

This can be demonstrated by the code

```
t = -2:2;
dt = 1;
ti = -2:0.025:2;
dti = 0.025;
y = sign (t);
ys = interp1 (t,y,ti,"spline");
yp = interp1 (t,y,ti,"pchip");
ddys = diff (diff (ys)./dti) ./ dti;
ddyp = diff (diff (yp)./dti) ./ dti;
figure (1);
plot (ti,ys,"r-", ti,yp,"g-");
legend ("spline", "pchip", 4);
figure (2);
plot (ti,ddys,"r+", ti,ddyp,"g*");
legend ("spline", "pchip");
```

The result of which can be seen in [Figure 29.1](#) and [Figure 29.2](#).




Figure 29.1: Comparison of "pchip" and "spline" interpolation methods for a step function




Figure 29.2: Comparison of the second derivative of the "pchip" and "spline" interpolation methods for a step function

Fourier interpolation, is a resampling technique where a signal is converted to the frequency domain, padded with zeros and then reconverted to the time domain.

```
interpft (x, n)
interpft (x, n, dim)
    Fourier interpolation.
```

If x is a vector then x is resampled with n points. The data in x is assumed to be equispaced. If x is a matrix or an N-dimensional array, the interpolation is performed on each column of x .

If dim is specified, then interpolate along the dimension dim .

`interpft` assumes that the interpolated function is periodic, and so assumptions are made about the endpoints of the interpolation.

See also: [interp1], page 765.

There are two significant limitations on Fourier interpolation. First, the function signal is assumed to be periodic, and so non-periodic signals will be poorly represented at the edges. Second, both the signal and its interpolation are required to be sampled at equispaced points. An example of the use of `interpft` is

```
t = 0 : 0.3 : pi; dt = t(2)-t(1);
n = length (t); k = 100;
ti = t(1) + [0 : k-1]*dt*n/k;
y = sin (4*t + 0.3) .* cos (3*t - 0.1);
yp = sin (4*ti + 0.3) .* cos (3*ti - 0.1);
plot (ti, yp, "g", ti, interp1 (t, y, ti, "spline"), "b", ...
      ti, interpft (y, k), "c", t, y, "r+");
legend ("sin(4t+0.3)cos(3t-0.1)", "spline", "interpft", "data");
```

which demonstrates the poor behavior of Fourier interpolation for non-periodic functions, as can be seen in Figure 29.3.




Figure 29.3: Comparison of `interp1` and `interpft` for non-periodic data

In addition, the support functions `spline` and `lookup` that underlie the `interp1` function can be called directly.

```
pp = spline (x, y)
yi = spline (x, y, xi)
```

Return the cubic spline interpolant of points x and y .

When called with two arguments, return the piecewise polynomial pp that may be used with `ppval` to evaluate the polynomial at specific points.

When called with a third input argument, `spline` evaluates the spline at the points xi . The third calling form `spline (x, y, xi)` is equivalent to `ppval (spline (x, y), xi)`.

The variable x must be a vector of length n .

y can be either a vector or array. If y is a vector it must have a length of either n or $n + 2$. If the length of y is n , then the "not-a-knot" end condition is used. If the length of y is $n + 2$, then the first and last values of the vector y are the values of the first derivative of the cubic spline at the endpoints.

If y is an array, then the size of y must have the form

$$[s_1, s_2, \dots, s_k, n]$$

or

$$[s_1, s_2, \dots, s_k, n + 2].$$

The array is reshaped internally to a matrix where the leading dimension is given by

$$s_1 s_2 \cdots s_k$$

and each row of this matrix is then treated separately. Note that this is exactly the opposite of `interp1` but is done for MATLAB compatibility.

See also: `[pchip]`, page 800, `[ppval]`, page 762, `[mkpp]`, page 761, `[unmkpp]`, page 761.

29.2 Multi-dimensional Interpolation

There are three multi-dimensional interpolation functions in Octave, with similar capabilities. Methods using Delaunay tessellation are described in [Section 30.4 \[Interpolation on Scattered Data\]](#), page 788.

```
zi = interp2 (x, y, z, xi, yi)
zi = interp2 (z, xi, yi)
zi = interp2 (z, n)
zi = interp2 (z)
zi = interp2 (..., method)
zi = interp2 (..., method, extrap)
```

Two-dimensional interpolation.

Interpolate reference data x , y , z to determine zi at the coordinates xi , yi . The reference data x , y can be matrices, as returned by `meshgrid`, in which case the sizes of x , y , and z must be equal. If x , y are vectors describing a grid then `length (x) == columns (z)` and `length (y) == rows (z)`. In either case the input data must be strictly monotonic.

If called without x , y , and just a single reference data matrix z , the 2-D region $x = 1:\text{columns}(z)$, $y = 1:\text{rows}(z)$ is assumed. This saves memory if the grid is regular and the distance between points is not important.

If called with a single reference data matrix z and a refinement value n , then perform interpolation over a grid where each original interval has been recursively subdivided n times. This results in 2^{n-1} additional points for every interval in the original grid. If n is omitted a value of 1 is used. As an example, the interval $[0,1]$ with $n=2$ results in a refined interval with points at $[0, 1/4, 1/2, 3/4, 1]$.

The interpolation *method* is one of:

- "nearest"
Return the nearest neighbor.
- "linear" (default)
Linear interpolation from nearest neighbors.
- "pchip"
Piecewise cubic Hermite interpolating polynomial—shape-preserving interpolation with smooth first derivative.
- "cubic"
Cubic interpolation (same as "pchip").
- "spline"
Cubic spline interpolation—smooth first and second derivatives throughout the curve.

extrap is a scalar number. It replaces values beyond the endpoints with *extrap*. Note that if *extrap* is used, *method* must be specified as well. If *extrap* is omitted and the *method* is "spline", then the extrapolated values of the "spline" are used. Otherwise the default *extrap* value for any other *method* is "NA".

See also: [interp1], page 765, [interp3], page 770, [interpn], page 771, [meshgrid], page 347.

```
vi = interp3 (x, y, z, v, xi, yi, zi)
vi = interp3 (v, xi, yi, zi)
vi = interp3 (v, n)
vi = interp3 (v)
vi = interp3 (..., method)
vi = interp3 (..., method, extrapval)
```

Three-dimensional interpolation.

Interpolate reference data x , y , z , v to determine vi at the coordinates xi , yi , zi . The reference data x , y , z can be matrices, as returned by `meshgrid`, in which case the sizes of x , y , z , and v must be equal. If x , y , z are vectors describing a cubic grid then `length(x) == columns(v)`, `length(y) == rows(v)`, and `length(z) == size(v, 3)`. In either case the input data must be strictly monotonic.

If called without x , y , z , and just a single reference data matrix v , the 3-D region $x = 1:\text{columns}(v)$, $y = 1:\text{rows}(v)$, $z = 1:\text{size}(v, 3)$ is assumed. This saves memory if the grid is regular and the distance between points is not important.

If called with a single reference data matrix v and a refinement value n , then perform interpolation over a 3-D grid where each original interval has been recursively subdivided n times. This results in 2^{n-1} additional points for every interval in the

original grid. If n is omitted a value of 1 is used. As an example, the interval [0,1] with $n==2$ results in a refined interval with points at [0, 1/4, 1/2, 3/4, 1].

The interpolation *method* is one of:

- "nearest" Return the nearest neighbor.
- "linear" (default) Linear interpolation from nearest neighbors.
- "cubic" Piecewise cubic Hermite interpolating polynomial—shape-preserving interpolation with smooth first derivative (not implemented yet).
- "spline" Cubic spline interpolation—smooth first and second derivatives throughout the curve.

extrapval is a scalar number. It replaces values beyond the endpoints with *extrapval*. Note that if *extrapval* is used, *method* must be specified as well. If *extrapval* is omitted and the *method* is "spline", then the extrapolated values of the "spline" are used. Otherwise the default *extrapval* value for any other *method* is "NA".

See also: [\[interp1\]](#), page 765, [\[interp2\]](#), page 769, [\[interpn\]](#), page 771, [\[meshgrid\]](#), page 347.

```
vi = interpn (x1, x2, ..., v, y1, y2, ...)
vi = interpn (v, y1, y2, ...)
vi = interpn (v, m)
vi = interpn (v)
vi = interpn (..., method)
vi = interpn (..., method, extrapval)
```

Perform n -dimensional interpolation, where n is at least two.

Each element of the n -dimensional array *v* represents a value at a location given by the parameters *x₁*, *x₂*, ..., *x_n*. The parameters *x₁*, *x₂*, ..., *x_n* are either n -dimensional arrays of the same size as the array *v* in the "ndgrid" format or vectors. The parameters *y₁*, etc. respect a similar format to *x₁*, etc., and they represent the points at which the array *vi* is interpolated.

If *x₁*, ..., *x_n* are omitted, they are assumed to be *x₁ = 1 : size (v, 1)*, etc. If *m* is specified, then the interpolation adds a point half way between each of the interpolation points. This process is performed *m* times. If only *v* is specified, then *m* is assumed to be 1.

The interpolation *method* is one of:

- "nearest" Return the nearest neighbor.
- "linear" (default) Linear interpolation from nearest neighbors.
- "pchip" Piecewise cubic Hermite interpolating polynomial—shape-preserving interpolation with smooth first derivative (not implemented yet).
- "cubic" Cubic interpolation (same as "pchip" [not implemented yet]).

"spline" Cubic spline interpolation—smooth first and second derivatives throughout the curve.

The default method is "linear".

`extrapval` is a scalar number. It replaces values beyond the endpoints with `extrapval`. Note that if `extrapval` is used, `method` must be specified as well. If `extrapval` is omitted and the `method` is "spline", then the extrapolated values of the "spline" are used. Otherwise the default `extrapval` value for any other `method` is "NA".

See also: [interp1], page 765, [interp2], page 769, [interp3], page 770, [spline], page 768, [ndgrid], page 348.

A significant difference between `interp3` and the other two multi-dimensional interpolation functions is the fashion in which the dimensions are treated. For `interp2` and `interp3`, the y-axis is considered to be the columns of the matrix, whereas the x-axis corresponds to the rows of the array. As Octave indexes arrays in column major order, the first dimension of any array is the columns, and so `interp3` effectively reverses the 'x' and 'y' dimensions. Consider the example,

```
x = y = z = -1:1;
f = @(x,y,z) x.^2 - y - z.^2;
[xx, yy, zz] = meshgrid (x, y, z);
v = f (xx,yy,zz);
xi = yi = zi = -1:0.1:1;
[xxi, yyi, zzi] = meshgrid (xi, yi, zi);
vi = interp3 (x, y, z, v, xxi, yyi, zzi, "spline");
[xxi, yyi, zzi] = ndgrid (xi, yi, zi);
vi2 = interpn (x, y, z, v, xxi, yyi, zzi, "spline");
mesh (zi, yi, squeeze (vi2(1,:,:)));
```

where `vi` and `vi2` are identical. The reversal of the dimensions is treated in the `meshgrid` and `ndgrid` functions respectively. The result of this code can be seen in [Figure 29.4](#).




Figure 29.4: Demonstration of the use of `interp2`

30 Geometry

Much of the geometry code in Octave is based on the Qhull library¹. Some of the documentation for Qhull, particularly for the options that can be passed to `delaunay`, `voronoi` and `convhull`, etc., is relevant to Octave users.

30.1 Delaunay Triangulation

The Delaunay triangulation is constructed from a set of circum-circles. These circum-circles are chosen so that there are at least three of the points in the set to triangulation on the circumference of the circum-circle. None of the points in the set of points falls within any of the circum-circles.

In general there are only three points on the circumference of any circum-circle. However, in some cases, and in particular for the case of a regular grid, 4 or more points can be on a single circum-circle. In this case the Delaunay triangulation is not unique.

```
tri = delaunay (x, y)
tetr = delaunay (x, y, z)
tri = delaunay (x)
tri = delaunay (... , options)
```

Compute the Delaunay triangulation for a 2-D or 3-D set of points.

For 2-D sets, the return value `tri` is a set of triangles which satisfies the Delaunay circum-circle criterion, i.e., no data point from $[x, y]$ is within the circum-circle of the defining triangle. The set of triangles `tri` is a matrix of size $[n, 3]$. Each row defines a triangle and the three columns are the three vertices of the triangle. The value of `tri(i,j)` is an index into `x` and `y` for the location of the j -th vertex of the i -th triangle.

For 3-D sets, the return value `tetr` is a set of tetrahedrons which satisfies the Delaunay circum-circle criterion, i.e., no data point from $[x, y, z]$ is within the circum-circle of the defining tetrahedron. The set of tetrahedrons is a matrix of size $[n, 4]$. Each row defines a tetrahedron and the four columns are the four vertices of the tetrahedron. The value of `tetr(i,j)` is an index into `x`, `y`, `z` for the location of the j -th vertex of the i -th tetrahedron.

The input `x` may also be a matrix with two or three columns where the first column contains x-data, the second y-data, and the optional third column contains z-data.

An optional final argument, which must be a string or cell array of strings, contains options passed to the underlying qhull command. See the documentation for the Qhull library for details <http://www.qhull.org/html/qh-quick.htm#options>. The default options are {"Qt", "Qbb", "Qc", "Qz"}.

If `options` is not present or [] then the default arguments are used. Otherwise, `options` replaces the default argument list. To append user options to the defaults it is necessary to repeat the default arguments in `options`. Use a null string to pass no arguments.

¹ Barber, C.B., Dobkin, D.P., and Huhdanpaa, H.T., *The Quickhull Algorithm for Convex Hulls*, ACM Trans. on Mathematical Software, 22(4):469–483, Dec 1996, <http://www.qhull.org>

```

x = rand (1, 10);
y = rand (1, 10);
tri = delaunay (x, y);
triplot (tri, x, y);
hold on;
plot (x, y, "r*");
axis ([0,1,0,1]);

```

See also: [\[delaunayn\]](#), page 776, [\[convhull\]](#), page 786, [\[voronoi\]](#), page 782, [\[triplot\]](#), page 777, [\[trimesh\]](#), page 777, [\[tetramesh\]](#), page 778, [\[trisurf\]](#), page 778.

For 3-D inputs `delaunay` returns a set of tetrahedra that satisfy the Delaunay circum-circle criteria. Similarly, `delaunayn` returns the N-dimensional simplex satisfying the Delaunay circum-circle criteria. The N-dimensional extension of a triangulation is called a tessellation.

`T = delaunay (pts)`

`T = delaunay (pts, options)`

Compute the Delaunay triangulation for an N-dimensional set of points.

The Delaunay triangulation is a tessellation of the convex hull of a set of points such that no N-sphere defined by the N-triangles contains any other points from the set.

The input matrix `pts` of size [n, dim] contains n points in a space of dimension dim. The return matrix `T` has size [m, dim+1]. Each row of `T` contains a set of indices back into the original set of points `pts` which describes a simplex of dimension dim. For example, a 2-D simplex is a triangle and 3-D simplex is a tetrahedron.

An optional second argument, which must be a string or cell array of strings, contains options passed to the underlying `qhull` command. See the documentation for the `Qhull` library for details <http://www.qhull.org/html/qh-quick.htm#options>. The default options depend on the dimension of the input:

- 2-D and 3-D: `options = {"Qt", "Qbb", "Qc", "Qz"}`
- 4-D and higher: `options = {"Qt", "Qbb", "Qc", "Qx"}`

If `options` is not present or [] then the default arguments are used. Otherwise, `options` replaces the default argument list. To append user options to the defaults it is necessary to repeat the default arguments in `options`. Use a null string to pass no arguments.

See also: [\[delaunay\]](#), page 775, [\[convhulln\]](#), page 787, [\[voronoin\]](#), page 783, [\[trimesh\]](#), page 777, [\[tetramesh\]](#), page 778.

An example of a Delaunay triangulation of a set of points is

```

rand ("state", 2);
x = rand (10, 1);
y = rand (10, 1);
T = delaunay (x, y);
X = [ x(T(:,1)); x(T(:,2)); x(T(:,3)); x(T(:,1)) ];
Y = [ y(T(:,1)); y(T(:,2)); y(T(:,3)); y(T(:,1)) ];
axis ([0, 1, 0, 1]);
plot (X, Y, "b", x, y, "r*");

```

The result of which can be seen in Figure 30.1.




Figure 30.1: Delaunay triangulation of a random set of points

30.1.1 Plotting the Triangulation

Octave has the functions `triplot`, `trimesh`, and `trisurf` to plot the Delaunay triangulation of a 2-dimensional set of points. `tetramesh` will plot the triangulation of a 3-dimensional set of points.

```
triplot (tri, x, y)
triplot (tri, x, y, linespec)
h = triplot (...)
```

Plot a 2-D triangular mesh.

tri is typically the output of a Delaunay triangulation over the grid of *x*, *y*. Every row of *tri* represents one triangle and contains three indices into $[x, y]$ which are the vertices of the triangles in the *x*-*y* plane.

The linestyle to use for the plot can be defined with the argument *linespec* of the same format as the `plot` command.

The optional return value *h* is a graphics handle to the created patch object.

See also: `[plot]`, page 296, `[trimesh]`, page 777, `[trisurf]`, page 778, `[delaunay]`, page 775.

```
trimesh (tri, x, y, z, c)
trimesh (tri, x, y, z)
trimesh (tri, x, y)
trimesh (..., prop, val, ...)
h = trimesh (...)
```

Plot a 3-D triangular wireframe mesh.

In contrast to `mesh`, which plots a mesh using rectangles, `trimesh` plots the mesh using triangles.

tri is typically the output of a Delaunay triangulation over the grid of *x*, *y*. Every row of *tri* represents one triangle and contains three indices into $[x, y]$ which are the vertices of the triangles in the *x*-*y* plane. *z* determines the height above the plane of each vertex. If no *z* input is given then the triangles are plotted as a 2-D figure.

The color of the trimesh is computed by linearly scaling the *z* values to fit the range of the current colormap. Use **caxis** and/or change the colormap to control the appearance.

Optionally, the color of the mesh can be specified independently of *z* by supplying *c*, which is a vector for colormap data, or a matrix with three columns for RGB data. The number of colors specified in *c* must either equal the number of vertices in *z* or the number of triangles in *tri*.

Any property/value pairs are passed directly to the underlying patch object.

The optional return value *h* is a graphics handle to the created patch object.

See also: [\[mesh\]](#), page 333, [\[tetramesh\]](#), page 778, [\[triplot\]](#), page 777, [\[trisurf\]](#), page 778, [\[delaunay\]](#), page 775, [\[patch\]](#), page 397, [\[hidden\]](#), page 335.

```
trisurf (tri, x, y, z, c)
trisurf (tri, x, y, z)
trisurf (..., prop, val, ...)
h = trisurf (...)
```

Plot a 3-D triangular surface.

In contrast to **surf**, which plots a surface mesh using rectangles, **trisurf** plots the mesh using triangles.

tri is typically the output of a Delaunay triangulation over the grid of *x*, *y*. Every row of *tri* represents one triangle and contains three indices into $[x, y]$ which are the vertices of the triangles in the *x*-*y* plane. *z* determines the height above the plane of each vertex.

The color of the trisurf is computed by linearly scaling the *z* values to fit the range of the current colormap. Use **caxis** and/or change the colormap to control the appearance.

Optionally, the color of the mesh can be specified independently of *z* by supplying *c*, which is a vector for colormap data, or a matrix with three columns for RGB data. The number of colors specified in *c* must either equal the number of vertices in *z* or the number of triangles in *tri*. When specifying the color at each vertex the triangle will be colored according to the color of the first vertex only (see patch documentation and the "FaceColor" property when set to "**flat**").

Any property/value pairs are passed directly to the underlying patch object.

The optional return value *h* is a graphics handle to the created patch object.

See also: [\[surf\]](#), page 335, [\[triplot\]](#), page 777, [\[trimesh\]](#), page 777, [\[delaunay\]](#), page 775, [\[patch\]](#), page 397, [\[shading\]](#), page 356.

```
tetramesh (T, X)
tetramesh (T, X, C)
tetramesh (..., property, val, ...)
```

```
h = tetramesh (...)
```

Display the tetrahedrons defined in the m-by-4 matrix T as 3-D patches.

T is typically the output of a Delaunay triangulation of a 3-D set of points. Every row of T contains four indices into the n-by-3 matrix X of the vertices of a tetrahedron. Every row in X represents one point in 3-D space.

The vector C specifies the color of each tetrahedron as an index into the current colormap. The default value is 1:m where m is the number of tetrahedrons; the indices are scaled to map to the full range of the colormap. If there are more tetrahedrons than colors in the colormap then the values in C are cyclically repeated.

Calling `tetramesh(..., "property", "value", ...)` passes all property/value pairs directly to the patch function as additional arguments.

The optional return value h is a vector of patch handles where each handle represents one tetrahedron in the order given by T . A typical use case for h is to turn the respective patch "visible" property "on" or "off".

Type `demo tetramesh` to see examples on using `tetramesh`.

See also: `[trimesh]`, page 777, `[delaunay]`, page 775, `[delaunayn]`, page 776, `[patch]`, page 397.

The difference between `triplot`, and `trimesh` or `trisurf`, is that the former only plots the 2-dimensional triangulation itself, whereas the second two plot the value of a function $f(x, y)$. An example of the use of the `triplot` function is

```
rand ("state", 2)
x = rand (20, 1);
y = rand (20, 1);
tri = delaunay (x, y);
triplot (tri, x, y);
```

which plots the Delaunay triangulation of a set of random points in 2-dimensions. The output of the above can be seen in [Figure 30.2](#).




Figure 30.2: Delaunay triangulation of a random set of points

30.1.2 Identifying Points in Triangulation

It is often necessary to identify whether a particular point in the N-dimensional space is within the Delaunay tessellation of a set of points in this N-dimensional space, and if so which N-simplex contains the point and which point in the tessellation is closest to the desired point. The functions `tsearch` and `dsearch` perform this function in a triangulation, and `tsearchn` and `dsearchn` in an N-dimensional tessellation.

To identify whether a particular point represented by a vector p falls within one of the simplices of an N-simplex, we can write the Cartesian coordinates of the point in a parametric form with respect to the N-simplex. This parametric form is called the Barycentric Coordinates of the point. If the points defining the N-simplex are given by $N + 1$ vectors $t(i, :)$, then the Barycentric coordinates defining the point p are given by

$$p = \text{beta} * t$$

where beta contains $N + 1$ values that together as a vector represent the Barycentric coordinates of the point p . To ensure a unique solution for the values of beta an additional criteria of

$$\text{sum}(\text{beta}) == 1$$

is imposed, and we can therefore write the above as

$$p - t(end, :) = \text{beta}(1:end-1) * (t(1:end-1, :) - \text{ones}(N, 1) * t(end, :))$$

Solving for beta we can then write

$$\begin{aligned} \text{beta}(1:end-1) &= (p - t(end, :)) / \\ &\quad (t(1:end-1, :) - \text{ones}(N, 1) * t(end, :)) \\ \text{beta}(end) &= \text{sum}(\text{beta}(1:end-1)) \end{aligned}$$

which gives the formula for the conversion of the Cartesian coordinates of the point p to the Barycentric coordinates beta . An important property of the Barycentric coordinates is that for all points in the N-simplex

$$0 \leq \text{beta}(i) \leq 1$$

Therefore, the test in `tsearch` and `tsearchn` essentially only needs to express each point in terms of the Barycentric coordinates of each of the simplices of the N-simplex and test the values of beta . This is exactly the implementation used in `tsearchn`. `tsearch` is optimized for 2-dimensions and the Barycentric coordinates are not explicitly formed.

`idx = tsearch (x, y, t, xi, yi)`

Search for the enclosing Delaunay convex hull.

For $t = \text{delaunay} (x, y)$, finds the index in t containing the points (xi, yi) . For points outside the convex hull, idx is NaN.

See also: [\[delaunay\]](#), page 775, [\[delaunayn\]](#), page 776.

`idx = tsearchn (x, t, xi)`

`[idx, p] = tsearchn (x, t, xi)`

Search for the enclosing Delaunay convex hull.

For $t = \text{delaunayn} (x)$, finds the index in t containing the points xi . For points outside the convex hull, idx is NaN.

If requested `tsearchn` also returns the Barycentric coordinates p of the enclosing triangles.

See also: [\[delaunay\]](#), page 775, [\[delaunayn\]](#), page 776.

An example of the use of `tsearch` can be seen with the simple triangulation

```
x = [-1; -1; 1; 1];
y = [-1; 1; -1; 1];
tri = [1, 2, 3; 2, 3, 4];
```

consisting of two triangles defined by tri . We can then identify which triangle a point falls in like

```
tsearch (x, y, tri, -0.5, -0.5)
⇒ 1
tsearch (x, y, tri, 0.5, 0.5)
⇒ 2
```

and we can confirm that a point doesn't lie within one of the triangles like

```
tsearch (x, y, tri, 2, 2)
⇒ NaN
```

The `dsearch` and `dsearchn` find the closest point in a tessellation to the desired point. The desired point does not necessarily have to be in the tessellation, and even if it is the returned point of the tessellation does not have to be one of the vertexes of the N-simplex within which the desired point is found.

`idx = dsearch (x, y, tri, xi, yi)`

`idx = dsearch (x, y, tri, xi, yi, s)`

Return the index idx of the closest point in x , y to the elements $[xi(:), yi(:)]$.

The variable *s* is accepted for compatibility but is ignored.

See also: [dsearchn], page 782, [tsearch], page 781.

```
idx = dsearchn (x, tri, xi)
idx = dsearchn (x, tri, xi, outval)
idx = dsearchn (x, xi)
[idx, d] = dsearchn (...)
```

Return the index *idx* of the closest point in *x* to the elements *xi*.

If *outval* is supplied, then the values of *xi* that are not contained within one of the simplices *tri* are set to *outval*. Generally, *tri* is returned from delaunayn (*x*).

See also: [dsearch], page 781, [tsearch], page 781.

An example of the use of **dsearch**, using the above values of *x*, *y* and *tri* is

```
dsearch (x, y, tri, -2, -2)
⇒ 1
```

If you wish the points that are outside the tessellation to be flagged, then **dsearchn** can be used as

```
dsearchn ([x, y], tri, [-2, -2], NaN)
⇒ NaN
dsearchn ([x, y], tri, [-0.5, -0.5], NaN)
⇒ 1
```

where the point outside the tessellation are then flagged with **NaN**.

30.2 Voronoi Diagrams

A Voronoi diagram or Voronoi tessellation of a set of points *s* in an *N*-dimensional space, is the tessellation of the *N*-dimensional space such that all points in *v(p)*, a partitions of the tessellation where *p* is a member of *s*, are closer to *p* than any other point in *s*. The Voronoi diagram is related to the Delaunay triangulation of a set of points, in that the vertexes of the Voronoi tessellation are the centers of the circum-circles of the simplices of the Delaunay tessellation.

```
voronoi (x, y)
voronoi (x, y, options)
voronoi (..., "linespec")
voronoi (hax, ...)
h = voronoi (...)
[vx, vy] = voronoi (...)
```

Plot the Voronoi diagram of points (*x*, *y*).

The Voronoi facets with points at infinity are not drawn.

The *options* argument, which must be a string or cell array of strings, contains options passed to the underlying qhull command. See the documentation for the Qhull library for details <http://www.qhull.org/html/qh-quick.htm#options>.

If "linespec" is given it is used to set the color and line style of the plot.

If an axes graphics handle *hax* is supplied then the Voronoi diagram is drawn on the specified axes rather than in a new figure.

If a single output argument is requested then the Voronoi diagram will be plotted and a graphics handle *h* to the plot is returned.

`[vx, vy] = voronoi (...)` returns the Voronoi vertices instead of plotting the diagram.

```
x = rand (10, 1);
y = rand (size (x));
h = convhull (x, y);
[vx, vy] = voronoi (x, y);
plot (vx, vy, "-b", x, y, "o", x(h), y(h), "-g");
legend ("", "points", "hull");
```

See also: [voronoin], page 783, [delaunay], page 775, [convhull], page 786.

`[C, F] = voronoin (pts)`
`[C, F] = voronoin (pts, options)`

Compute N-dimensional Voronoi facets.

The input matrix *pts* of size $[n, \text{dim}]$ contains *n* points in a space of dimension *dim*.

C contains the points of the Voronoi facets. The list *F* contains, for each facet, the indices of the Voronoi points.

An optional second argument, which must be a string or cell array of strings, contains options passed to the underlying qhull command. See the documentation for the Qhull library for details <http://www.qhull.org/html/qh-quick.htm#options>.

The default options depend on the dimension of the input:

- 2-D and 3-D: *options* = {"Qbb"}
- 4-D and higher: *options* = {"Qbb", "Qx"}

If *options* is not present or [] then the default arguments are used. Otherwise, *options* replaces the default argument list. To append user options to the defaults it is necessary to repeat the default arguments in *options*. Use a null string to pass no arguments.

See also: [voronoi], page 782, [convhulln], page 787, [delaunay], page 776.

An example of the use of `voronoi` is

```
rand ("state",9);
x = rand (10,1);
y = rand (10,1);
tri = delaunay (x, y);
[vx, vy] = voronoi (x, y, tri);
triplot (tri, x, y, "b");
hold on;
plot (vx, vy, "r");
```

The result of which can be seen in [Figure 30.3](#). Note that the circum-circle of one of the triangles has been added to this figure, to make the relationship between the Delaunay tessellation and the Voronoi diagram clearer.




Figure 30.3: Delaunay triangulation (blue lines) and Voronoi diagram (red lines) of a random set of points

Additional information about the size of the facets of a Voronoi diagram, and which points of a set of points is in a polygon can be had with the `polyarea` and `inpolygon` functions respectively.

```
polyarea (x, y)
polyarea (x, y, dim)
```

Determine area of a polygon by triangle method.

The variables `x` and `y` define the vertex pairs, and must therefore have the same shape. They can be either vectors or arrays. If they are arrays then the columns of `x` and `y` are treated separately and an area returned for each.

If the optional `dim` argument is given, then `polyarea` works along this dimension of the arrays `x` and `y`.

An example of the use of `polyarea` might be

```
rand ("state", 2);
x = rand (10, 1);
y = rand (10, 1);
[c, f] = voronoin ([x, y]);
af = zeros (size (f));
for i = 1 : length (f)
    af(i) = polyarea (c (f {i, :}, 1), c (f {i, :}, 2));
endfor
```

Facets of the Voronoi diagram with a vertex at infinity have infinity area. A simplified version of `polyarea` for rectangles is available with `rectint`

```
area = rectint (a, b)
```

Compute area or volume of intersection of rectangles or N-D boxes.

Compute the area of intersection of rectangles in *a* and rectangles in *b*. N-dimensional boxes are supported in which case the volume, or hypervolume is computed according to the number of dimensions.

2-dimensional rectangles are defined as [xpos ypos width height] where xpos and ypos are the position of the bottom left corner. Higher dimensions are supported where the coordinates for the minimum value of each dimension follow the length of the box in that dimension, e.g., [xpos ypos zpos kpos ... width height depth k_length ...].

Each row of *a* and *b* define a rectangle, and if both define multiple rectangles, then the output, *area*, is a matrix where the *i*-th row corresponds to the *i*-th row of *a* and the *j*-th column corresponds to the *j*-th row of *b*.

See also: [\[polyarea\]](#), page 784.

```
in = inpolygon (x, y, xv, yv)
[in, on] = inpolygon (x, y, xv, yv)
```

For a polygon defined by vertex points (*xv*, *yv*), return true if the points (*x*, *y*) are inside (or on the boundary) of the polygon; Otherwise, return false.

The input variables *x* and *y*, must have the same dimension.

The optional output *on* returns true if the points are exactly on the polygon edge, and false otherwise.

See also: [\[delaunay\]](#), page 775.

An example of the use of *inpolygon* might be

```
randn ("state", 2);
x = randn (100, 1);
y = randn (100, 1);
vx = cos (pi * [-1 : 0.1: 1]);
vy = sin (pi * [-1 : 0.1 : 1]);
in = inpolygon (x, y, vx, vy);
plot (vx, vy, x(in), y(in), "r+", x(!in), y(!in), "bo");
axis([-2, 2, -2, 2]);
```

The result of which can be seen in [Figure 30.4](#).




Figure 30.4: Demonstration of the `inpolygon` function to determine the points inside a polygon

30.3 Convex Hull

The convex hull of a set of points is the minimum convex envelope containing all of the points. Octave has the functions `convhull` and `convhulln` to calculate the convex hull of 2-dimensional and N-dimensional sets of points.

```
H = convhull (x, y)
H = convhull (x, y, z)
H = convhull (x)
H = convhull (..., options)
[H, V] = convhull (...)

Compute the convex hull of a 2-D or 3-D set of points.
```

The hull H is a linear index vector into the original set of points that specifies which points form the enclosing hull. For 2-D inputs only, the output is ordered in a counterclockwise manner around the hull.

The input x may also be a matrix with two or three columns where the first column contains x-data, the second y-data, and the optional third column contains z-data.

An optional final argument, which must be a string or cell array of strings, contains options passed to the underlying qhull command. See the documentation for the Qhull library for details <http://www.qhull.org/html/qh-quick.htm#options>. The default option is `{"Qt"}`.

If $options$ is not present or `[]` then the default arguments are used. Otherwise, $options$ replaces the default argument list. To append user options to the defaults it is necessary to repeat the default arguments in $options$. Use a null string to pass no arguments.

If the second output V is requested the volume of the enclosing convex hull is calculated.

See also: [convhulln], page 787, [delaunay], page 775, [voronoi], page 782.

```
h = convhulln (pts)
h = convhulln (pts, options)
[h, v] = convhulln (...)
```

Compute the convex hull of the set of points *pts*.

pts is a matrix of size [*n*, *dim*] containing *n* points in a space of dimension *dim*.

The hull *h* is an index vector into the set of points and specifies which points form the enclosing hull.

An optional second argument, which must be a string or cell array of strings, contains options passed to the underlying qhull command. See the documentation for the Qhull library for details <http://www.qhull.org/html/qh-quick.htm#options>. The default options depend on the dimension of the input:

- 2D, 3D, 4D: *options* = {"Qt"}
- 5D and higher: *options* = {"Qt", "Qx"}

If *options* is not present or [] then the default arguments are used. Otherwise, *options* replaces the default argument list. To append user options to the defaults it is necessary to repeat the default arguments in *options*. Use a null string to pass no arguments.

If the second output *v* is requested the volume of the enclosing convex hull is calculated.

See also: [convhull], page 786, [delaunayn], page 776, [voronoin], page 783.

An example of the use of **convhull** is

```
x = -3:0.05:3;
y = abs (sin (x));
k = convhull (x, y);
plot (x(k), y(k), "r-", x, y, "b+");
axis ([-3.05, 3.05, -0.05, 1.05]);
```

The output of the above can be seen in **Figure 30.5**.




Figure 30.5: The convex hull of a simple set of points

30.4 Interpolation on Scattered Data

An important use of the Delaunay tessellation is that it can be used to interpolate from scattered data to an arbitrary set of points. To do this the N-simplex of the known set of points is calculated with `delaunay` or `delaunayn`. Then the simplices in to which the desired points are found are identified. Finally the vertices of the simplices are used to interpolate to the desired points. The functions that perform this interpolation are `griddata`, `griddata3` and `griddatan`.

```
zi = griddata (x, y, z, xi, yi)
zi = griddata (x, y, z, xi, yi, method)
[xi, yi, zi] = griddata (...)
```

Generate a regular mesh from irregular data using interpolation.

The function is defined by $z = f(x, y)$. Inputs x , y , z are vectors of the same length or x , y are vectors and z is matrix.

The interpolation points are all (xi, yi) . If xi , yi are vectors then they are made into a 2-D mesh.

The interpolation method can be "nearest", "cubic" or "linear". If method is omitted it defaults to "linear".

See also: [\[griddata3\]](#), page 788, [\[griddatan\]](#), page 789, [\[delaunay\]](#), page 775.

```
vi = griddata3 (x, y, z, v, xi, yi, zi)
vi = griddata3 (x, y, z, v, xi, yi, zi, method)
vi = griddata3 (x, y, z, v, xi, yi, zi, method, options)
```

Generate a regular mesh from irregular data using interpolation.

The function is defined by $v = f(x, y, z)$. The interpolation points are specified by xi , yi , zi .

The interpolation method can be "nearest" or "linear". If method is omitted it defaults to "linear".

The optional argument *options* is passed directly to Qhull when computing the Delaunay triangulation used for interpolation. See `delaunayn` for information on the defaults and how to pass different values.

See also: `[griddata]`, page 788, `[griddatan]`, page 789, `[delaunayn]`, page 776.

```
yi = griddatan (x, y, xi)
yi = griddatan (x, y, xi, method)
yi = griddatan (x, y, xi, method, options)
```

Generate a regular mesh from irregular data using interpolation.

The function is defined by $y = f(x)$. The interpolation points are all xi .

The interpolation method can be "nearest" or "linear". If method is omitted it defaults to "linear".

The optional argument *options* is passed directly to Qhull when computing the Delaunay triangulation used for interpolation. See `delaunayn` for information on the defaults and how to pass different values.

See also: `[griddata]`, page 788, `[griddata3]`, page 788, `[delaunayn]`, page 776.

An example of the use of the `griddata` function is

```
rand ("state", 1);
x = 2*rand (1000,1) - 1;
y = 2*rand (size (x)) - 1;
z = sin (2*(x.^2+y.^2));
[xx,yy] = meshgrid (linspace (-1,1,32));
zz = griddata (x, y, z, xx, yy);
mesh (xx, yy, zz);
```

that interpolates from a random scattering of points, to a uniform grid. The output of the above can be seen in [Figure 30.6](#).




Figure 30.6: Interpolation from a scattered data to a regular grid

31 Signal Processing

This chapter describes the signal processing and fast Fourier transform functions available in Octave. Fast Fourier transforms are computed with the FFTW or FFTPACK libraries depending on how Octave is built.

fft (x)
fft (x, n)
fft (x, n, dim)

Compute the discrete Fourier transform of *x* using a Fast Fourier Transform (FFT) algorithm.

The FFT is calculated along the first non-singleton dimension of the array. Thus if *x* is a matrix, **fft (x)** computes the FFT for each column of *x*.

If called with two arguments, *n* is expected to be an integer specifying the number of elements of *x* to use, or an empty matrix to specify that its value should be ignored. If *n* is larger than the dimension along which the FFT is calculated, then *x* is resized and padded with zeros. Otherwise, if *n* is smaller than the dimension along which the FFT is calculated, then *x* is truncated.

If called with three arguments, *dim* is an integer specifying the dimension of the matrix along which the FFT is performed.

See also: [\[ifft\]](#), page 791, [\[fft2\]](#), page 791, [\[fftn\]](#), page 792, [\[fftw\]](#), page 792.

ifft (x)
ifft (x, n)
ifft (x, n, dim)

Compute the inverse discrete Fourier transform of *x* using a Fast Fourier Transform (FFT) algorithm.

The inverse FFT is calculated along the first non-singleton dimension of the array. Thus if *x* is a matrix, **fft (x)** computes the inverse FFT for each column of *x*.

If called with two arguments, *n* is expected to be an integer specifying the number of elements of *x* to use, or an empty matrix to specify that its value should be ignored. If *n* is larger than the dimension along which the inverse FFT is calculated, then *x* is resized and padded with zeros. Otherwise, if *n* is smaller than the dimension along which the inverse FFT is calculated, then *x* is truncated.

If called with three arguments, *dim* is an integer specifying the dimension of the matrix along which the inverse FFT is performed.

See also: [\[fft\]](#), page 791, [\[ifft2\]](#), page 792, [\[fftn\]](#), page 792, [\[fftw\]](#), page 792.

fft2 (A)
fft2 (A, m, n)

Compute the two-dimensional discrete Fourier transform of *A* using a Fast Fourier Transform (FFT) algorithm.

The optional arguments *m* and *n* may be used to specify the number of rows and columns of *A* to use. If either of these is larger than the size of *A*, *A* is resized and padded with zeros.

If A is a multi-dimensional matrix, each two-dimensional sub-matrix of A is treated separately.

See also: [\[ifft2\]](#), page 792, [\[fft\]](#), page 791, [\[fftn\]](#), page 792, [\[fftw\]](#), page 792.

ifft2 (A)
ifft2 (A, m, n)

Compute the inverse two-dimensional discrete Fourier transform of A using a Fast Fourier Transform (FFT) algorithm.

The optional arguments m and n may be used specify the number of rows and columns of A to use. If either of these is larger than the size of A , A is resized and padded with zeros.

If A is a multi-dimensional matrix, each two-dimensional sub-matrix of A is treated separately.

See also: [\[fft2\]](#), page 791, [\[ifft\]](#), page 791, [\[fftn\]](#), page 792, [\[fftw\]](#), page 792.

fftn (A)
fftn (A, size)

Compute the N-dimensional discrete Fourier transform of A using a Fast Fourier Transform (FFT) algorithm.

The optional vector argument size may be used specify the dimensions of the array to be used. If an element of size is smaller than the corresponding dimension of A , then the dimension of A is truncated prior to performing the FFT. Otherwise, if an element of size is larger than the corresponding dimension then A is resized and padded with zeros.

See also: [\[fftn\]](#), page 792, [\[fft\]](#), page 791, [\[fft2\]](#), page 791, [\[fftw\]](#), page 792.

iifftn (A)
iifftn (A, size)

Compute the inverse N-dimensional discrete Fourier transform of A using a Fast Fourier Transform (FFT) algorithm.

The optional vector argument size may be used specify the dimensions of the array to be used. If an element of size is smaller than the corresponding dimension of A , then the dimension of A is truncated prior to performing the inverse FFT. Otherwise, if an element of size is larger than the corresponding dimension then A is resized and padded with zeros.

See also: [\[fftn\]](#), page 792, [\[ifft\]](#), page 791, [\[ifft2\]](#), page 792, [\[fftw\]](#), page 792.

Octave uses the FFTW libraries to perform FFT computations. When Octave starts up and initializes the FFTW libraries, they read a system wide file (on a Unix system, it is typically `/etc/fftw/wisdom`) that contains information useful to speed up FFT computations. This information is called the *wisdom*. The system-wide file allows wisdom to be shared between all applications using the FFTW libraries.

Use the `fftw` function to generate and save wisdom. Using the utilities provided together with the FFTW libraries (`fftw-wisdom` on Unix systems), you can even add wisdom generated by Octave to the system-wide wisdom file.

```
method = fftw ("planner")
fftw ("planner", method)
wisdom = fftw ("dwisdom")
fftw ("dwisdom", wisdom)
fftw ("threads", nthreads)
nthreads = fftw ("threads")
```

Manage FFTW wisdom data.

Wisdom data can be used to significantly accelerate the calculation of the FFTs, but implies an initial cost in its calculation. When the FFTW libraries are initialized, they read a system wide wisdom file (typically in `/etc/fftw/wisdom`), allowing wisdom to be shared between applications other than Octave. Alternatively, the `fftw` function can be used to import wisdom. For example,

```
wisdom = fftw ("dwisdom")
```

will save the existing wisdom used by Octave to the string `wisdom`. This string can then be saved to a file and restored using the `save` and `load` commands respectively. This existing wisdom can be re-imported as follows

```
fftw ("dwisdom", wisdom)
```

If `wisdom` is an empty string, then the wisdom used is cleared.

During the calculation of Fourier transforms further wisdom is generated. The fashion in which this wisdom is generated is also controlled by the `fftw` function. There are five different manners in which the wisdom can be treated:

"estimate"

Specifies that no run-time measurement of the optimal means of calculating a particular is performed, and a simple heuristic is used to pick a (probably sub-optimal) plan. The advantage of this method is that there is little or no overhead in the generation of the plan, which is appropriate for a Fourier transform that will be calculated once.

"measure"

In this case a range of algorithms to perform the transform is considered and the best is selected based on their execution time.

"patient"

Similar to **"measure"**, but a wider range of algorithms is considered.

"exhaustive"

Like **"measure"**, but all possible algorithms that may be used to treat the transform are considered.

"hybrid"

As run-time measurement of the algorithm can be expensive, this is a compromise where **"measure"** is used for transforms up to the size of 8192 and beyond that the **"estimate"** method is used.

The default method is **"estimate"**. The current method can be queried with

```
method = fftw ("planner")
```

or set by using

```
fftw ("planner", method)
```

Note that calculated wisdom will be lost when restarting Octave. However, the wisdom data can be reloaded if it is saved to a file as described above. Saved wisdom files should not be used on different platforms since they will not be efficient and the point of calculating the wisdom is lost.

The number of threads used for computing the plans and executing the transforms can be set with

```
fftw ("threads", NTHREADS)
```

Note that octave must be compiled with multi-threaded FFTW support for this feature. The number of processors available to the current process is used per default.

See also: [fft], page 791, [ifft], page 791, [fft2], page 791, [ifft2], page 792, [fftn], page 792, [ifftn], page 792.

```
fftconv (x, y)
```

```
fftconv (x, y, n)
```

Convolve two vectors using the FFT for computation.

`c = fftconv (x, y)` returns a vector of length equal to `length (x) + length (y) - 1`. If `x` and `y` are the coefficient vectors of two polynomials, the returned value is the coefficient vector of the product polynomial.

The computation uses the FFT by calling the function `fftfilt`. If the optional argument `n` is specified, an `N`-point FFT is used.

See also: [deconv], page 750, [conv], page 749, [conv2], page 750.

```
fftfilt (b, x)
```

```
fftfilt (b, x, n)
```

Filter `x` with the FIR filter `b` using the FFT.

If `x` is a matrix, filter each column of the matrix.

Given the optional third argument, `n`, `fftfilt` uses the overlap-add method to filter `x` with `b` using an `N`-point FFT. The FFT size must be an even power of 2 and must be greater than or equal to the length of `b`. If the specified `n` does not meet these criteria, it is automatically adjusted to the nearest value that does.

See also: [filter], page 794, [filter2], page 795.

```
y = filter (b, a, x)
```

```
[y, sf] = filter (b, a, x, si)
```

```
[y, sf] = filter (b, a, x, [], dim)
```

```
[y, sf] = filter (b, a, x, si, dim)
```

Apply a 1-D digital filter to the data `x`.

`filter` returns the solution to the following linear, time-invariant difference equation:

$$\sum_{k=0}^N a_{k+1} y_{n-k} = \sum_{k=0}^M b_{k+1} x_{n-k}, \quad 1 \leq n \leq P$$

where $a \in \mathbb{R}^{N-1}$, $b \in \mathbb{R}^{M-1}$, and $x \in \mathbb{R}^P$. The result is calculated over the first non-singleton dimension of `x` or over `dim` if supplied.

An equivalent form of the equation is:

$$y_n = - \sum_{k=1}^N c_{k+1} y_{n-k} + \sum_{k=0}^M d_{k+1} x_{n-k}, \quad 1 \leq n \leq P$$

where $c = a/a_1$ and $d = b/a_1$.

If the fourth argument *si* is provided, it is taken as the initial state of the system and the final state is returned as *sf*. The state vector is a column vector whose length is equal to the length of the longest coefficient vector minus one. If *si* is not supplied, the initial state vector is set to all zeros.

In terms of the Z Transform, *y* is the result of passing the discrete-time signal *x* through a system characterized by the following rational system function:

$$H(z) = \frac{\sum_{k=0}^M d_{k+1} z^{-k}}{1 + \sum_{k=1}^N c_{k+1} z^{-k}}$$

See also: [filter2], page 795, [fftfilt], page 794, [freqz], page 795.

```
y = filter2 (b, x)
y = filter2 (b, x, shape)
```

Apply the 2-D FIR filter *b* to *x*.

If the argument *shape* is specified, return an array of the desired shape. Possible values are:

- "full" pad *x* with zeros on all sides before filtering.
- "same" unpadded *x* (default)
- "valid" trim *x* after filtering so edge effects are not included.

Note this is just a variation on convolution, with the parameters reversed and *b* rotated 180 degrees.

See also: [conv2], page 750.

```
[h, w] = freqz (b, a, n, "whole")
[h, w] = freqz (b)
[h, w] = freqz (b, a)
[h, w] = freqz (b, a, n)
h = freqz (b, a, w)
[h, w] = freqz (..., Fs)
freqz (...)
```

Return the complex frequency response *h* of the rational IIR filter whose numerator and denominator coefficients are *b* and *a*, respectively.

The response is evaluated at *n* angular frequencies between 0 and 2π .

The output value *w* is a vector of the frequencies.

If *a* is omitted, the denominator is assumed to be 1 (this corresponds to a simple FIR filter).

If *n* is omitted, a value of 512 is assumed. For fastest computation, *n* should factor into a small number of small primes.

If the fourth argument, "whole", is omitted the response is evaluated at frequencies between 0 and π .

freqz (b, a, w)

Evaluate the response at the specific frequencies in the vector *w*. The values for *w* are measured in radians.

[...] = freqz (... , Fs)

Return frequencies in Hz instead of radians assuming a sampling rate *Fs*. If you are evaluating the response at specific frequencies *w*, those frequencies should be requested in Hz rather than radians.

freqz (...)

Plot the magnitude and phase response of *h* rather than returning them.

See also: [\[freqz-plot\]](#), page 796.

freqz_plot (w, h)

freqz_plot (w, h, freq_norm)

Plot the magnitude and phase response of *h*.

If the optional *freq_norm* argument is true, the frequency vector *w* is in units of normalized radians. If *freq_norm* is false, or not given, then *w* is measured in Hertz.

See also: [\[freqz\]](#), page 795.

sinc (x)

Compute the sinc function.

Return $\sin(\pi x)/(\pi x)$.

b = unwrap (x)

b = unwrap (x, tol)

b = unwrap (x, tol, dim)

Unwrap radian phases by adding or subtracting multiples of 2π as appropriate to remove jumps greater than *tol*.

tol defaults to π .

Unwrap will work along the dimension *dim*. If *dim* is unspecified it defaults to the first non-singleton dimension.

[a, b] = arch_fit (y, x, p, iter, gamma, a0, b0)

Fit an ARCH regression model to the time series *y* using the scoring algorithm in Engle's original ARCH paper.

The model is

$$\begin{aligned} y(t) &= b(1) * x(t,1) + \dots + b(k) * x(t,k) + e(t), \\ h(t) &= a(1) + a(2) * e(t-1)^2 + \dots + a(p+1) * e(t-p)^2 \end{aligned}$$

in which $e(t)$ is $N(0, h(t))$, given a time-series vector y up to time $t - 1$ and a matrix of (ordinary) regressors x up to t . The order of the regression of the residual variance is specified by p .

If invoked as `arch_fit (y, k, p)` with a positive integer k , fit an ARCH(k, p) process, i.e., do the above with the t -th row of x given by

`[1, y(t-1), ..., y(t-k)]`

Optionally, one can specify the number of iterations $iter$, the updating factor $gamma$, and initial values $a0$ and $b0$ for the scoring algorithm.

arch_rnd (a, b, t)

Simulate an ARCH sequence of length t with AR coefficients b and CH coefficients a .

The result $y(t)$ follows the model

$$y(t) = b(1) + b(2) * y(t-1) + \dots + b(lb) * y(t-lb+1) + e(t),$$

where $e(t)$, given y up to time $t - 1$, is $N(0, h(t))$, with

$$h(t) = a(1) + a(2) * e(t-1)^2 + \dots + a(la) * e(t-la+1)^2$$

[pval, lm] = arch_test (y, x, p)

For a linear regression model

$$y = x * b + e$$

perform a Lagrange Multiplier (LM) test of the null hypothesis of no conditional heteroscedascity against the alternative of CH(p).

I.e., the model is

$$y(t) = b(1) * x(t,1) + \dots + b(k) * x(t,k) + e(t),$$

given y up to $t - 1$ and x up to t , $e(t)$ is $N(0, h(t))$ with

$$h(t) = v + a(1) * e(t-1)^2 + \dots + a(p) * e(t-p)^2,$$

and the null is $a(1) == \dots == a(p) == 0$.

If the second argument is a scalar integer, k , perform the same test in a linear autoregression model of order k , i.e., with

`[1, y(t-1), ..., y(t-k)]`

as the t -th row of x .

Under the null, LM approximately has a chisquare distribution with p degrees of freedom and $pval$ is the p -value (1 minus the CDF of this distribution at LM) of the test.

If no output argument is given, the p -value is displayed.

arma_rnd (a, b, v, t, n)

Return a simulation of the ARMA model.

The ARMA model is defined by

$$\begin{aligned} x(n) = & a(1) * x(n-1) + \dots + a(k) * x(n-k) \\ & + e(n) + b(1) * e(n-1) + \dots + b(l) * e(n-l) \end{aligned}$$

in which k is the length of vector a , l is the length of vector b and e is Gaussian white noise with variance v . The function returns a vector of length t .

The optional parameter n gives the number of dummy $x(i)$ used for initialization, i.e., a sequence of length $t+n$ is generated and $x(n+1:t+n)$ is returned. If n is omitted, $n = 100$ is used.

autoreg_matrix (y, k)

Given a time series (vector) y , return a matrix with ones in the first column and the first k lagged values of y in the other columns.

In other words, for $t > k$, $[1, y(t-1), \dots, y(t-k)]$ is the t -th row of the result.

The resulting matrix may be used as a regressor matrix in autoregressions.

bartlett (m)

Return the filter coefficients of a Bartlett (triangular) window of length m .

For a definition of the Bartlett window see, e.g., A.V. Oppenheim & R. W. Schafer, *Discrete-Time Signal Processing*.

blackman (m)

blackman (m, "periodic")

blackman (m, "symmetric")

Return the filter coefficients of a Blackman window of length m .

If the optional argument "**periodic**" is given, the periodic form of the window is returned. This is equivalent to the window of length $m+1$ with the last coefficient removed. The optional argument "**symmetric**" is equivalent to not specifying a second argument.

For a definition of the Blackman window, see, e.g., A.V. Oppenheim & R. W. Schafer, *Discrete-Time Signal Processing*.

detrend (x, p)

If x is a vector, **detrend (x, p)** removes the best fit of a polynomial of order p from the data x .

If x is a matrix, **detrend (x, p)** does the same for each column in x .

The second argument p is optional. If it is not specified, a value of 1 is assumed. This corresponds to removing a linear trend.

The order of the polynomial can also be given as a string, in which case p must be either "**constant**" (corresponds to $p=0$) or "**linear**" (corresponds to $p=1$).

See also: [\[polyfit\]](#), page 753.

[d, dd] = diffpara (x, a, b)

Return the estimator d for the differencing parameter of an integrated time series.

The frequencies from $[2 * pi * a/t, 2 * pi * b/T]$ are used for the estimation. If b is omitted, the interval $[2 * pi/T, 2 * pi * a/T]$ is used. If both b and a are omitted then $a = 0.5 * sqrt(T)$ and $b = 1.5 * sqrt(T)$ is used, where T is the sample size. If x is a matrix, the differencing parameter of each column is estimated.

The estimators for all frequencies in the intervals described above is returned in dd .

The value of d is simply the mean of dd .

Reference: P.J. Brockwell & R.A. Davis. *Time Series: Theory and Methods*. Springer 1987.

durbinlevinson (c, oldphi, oldv)

Perform one step of the Durbin-Levinson algorithm.

The vector *c* specifies the autocovariances [*gamma_0*, ..., *gamma_t*] from lag 0 to *t*, *oldphi* specifies the coefficients based on *c(t-1)* and *oldv* specifies the corresponding error.

If *oldphi* and *oldv* are omitted, all steps from 1 to *t* of the algorithm are performed.

fftshift (x)**fftshift (x, dim)**

Perform a shift of the vector *x*, for use with the **fft** and **ifft** functions, in order to move the frequency 0 to the center of the vector or matrix.

If *x* is a vector of *N* elements corresponding to *N* time samples spaced by *dt*, then **fftshift (fft (x))** corresponds to frequencies

$$f = [-(ceil((N-1)/2):-1:1), 0, (1:floor((N-1)/2))] * df$$

where *df* = $1/(N * dt)$.

If *x* is a matrix, the same holds for rows and columns. If *x* is an array, then the same holds along each dimension.

The optional *dim* argument can be used to limit the dimension along which the permutation occurs.

See also: [\[ifftshift\]](#), page 799.

ifftshift (x)**ifftshift (x, dim)**

Undo the action of the **fftshift** function.

For even length *x*, **fftshift** is its own inverse, but odd lengths differ slightly.

See also: [\[fftshift\]](#), page 799.

fractdiff (x, d)

Compute the fractional differences $(1 - L)^d x$ where *L* denotes the lag-operator and *d* is greater than -1.

hamming (m)**hamming (m, "periodic")****hamming (m, "symmetric")**

Return the filter coefficients of a Hamming window of length *m*.

If the optional argument "periodic" is given, the periodic form of the window is returned. This is equivalent to the window of length *m*+1 with the last coefficient removed. The optional argument "symmetric" is equivalent to not specifying a second argument.

For a definition of the Hamming window see, e.g., A.V. Oppenheim & R. W. Schafer, *Discrete-Time Signal Processing*.

hanning (m)**hanning (m, "periodic")****hanning (m, "symmetric")**

Return the filter coefficients of a Hanning window of length *m*.

If the optional argument "periodic" is given, the periodic form of the window is returned. This is equivalent to the window of length $m+1$ with the last coefficient removed. The optional argument "symmetric" is equivalent to not specifying a second argument.

For a definition of the Hanning window see, e.g., A.V. Oppenheim & R. W. Schafer, *Discrete-Time Signal Processing*.

hurst (x)

Estimate the Hurst parameter of sample x via the rescaled range statistic.

If x is a matrix, the parameter is estimated for every column.

pp = pchip (x, y)

yi = pchip (x, y, xi)

Return the Piecewise Cubic Hermite Interpolating Polynomial (pchip) of points x and y .

If called with two arguments, return the piecewise polynomial pp that may be used with **ppval** to evaluate the polynomial at specific points.

When called with a third input argument, **pchip** evaluates the pchip polynomial at the points xi . The third calling form is equivalent to **ppval (pchip (x, y), xi)**.

The variable x must be a strictly monotonic vector (either increasing or decreasing) of length n .

y can be either a vector or array. If y is a vector then it must be the same length n as x . If y is an array then the size of y must have the form

$$[s_1, s_2, \dots, s_k, n]$$

The array is reshaped internally to a matrix where the leading dimension is given by

$$s_1 s_2 \cdots s_k$$

and each row of this matrix is then treated separately. Note that this is exactly opposite to **interp1** but is done for MATLAB compatibility.

See also: [\[spline\]](#), page 768, [\[ppval\]](#), page 762, [\[mkpp\]](#), page 761, [\[unmkpp\]](#), page 761.

[Pxx, w] = periodogram (x)
[Pxx, w] = periodogram (x, win)
[Pxx, w] = periodogram (x, win, nfft)
[Pxx, f] = periodogram (x, win, nfft, Fs)
[Pxx, f] = periodogram (..., "range")
periodogram (...)

Return the periodogram (Power Spectral Density) of x .

The possible inputs are:

x

data vector. If x is real-valued a one-sided spectrum is estimated. If x is complex-valued, or "range" specifies "twosided", the full spectrum is estimated.

<i>win</i>	window weight data. If window is empty or unspecified a default rectangular window is used. Otherwise, the window is applied to the signal (<i>x</i> .* <i>win</i>) before computing the periodogram. The window data must be a vector of the same length as <i>x</i> .
<i>nfft</i>	number of frequency bins. The default is 256 or the next higher power of 2 greater than the length of <i>x</i> (<code>max (256, 2.^nextpow2 (length (x)))</code>). If <i>nfft</i> is greater than the length of the input then <i>x</i> will be zero-padded to the length of <i>nfft</i> .
<i>Fs</i>	sampling rate. The default is 1.
<i>range</i>	range of spectrum. "onesided" computes spectrum from [0:nfft/2+1]. "twosided" computes spectrum from [0:nfft-1].

The optional second output *w* are the normalized angular frequencies. For a one-sided calculation *w* is in the range [0, pi] if *nfft* is even and [0, pi) if *nfft* is odd. Similarly, for a two-sided calculation *w* is in the range [0, 2*pi] or [0, 2*pi) depending on *nfft*. If a sampling frequency is specified, *Fs*, then the output frequencies *f* will be in the range [0, *Fs*/2] or [0, *Fs*/2) for one-sided calculations. For two-sided calculations the range will be [0, *Fs*).

When called with no outputs the periodogram is immediately plotted in the current figure window.

See also: [\[fft\]](#), page 791.

sinetone (freq, rate, sec, ampl)

Return a sinetone of frequency *freq* with a length of *sec* seconds at sampling rate *rate* and with amplitude *ampl*.

The arguments *freq* and *ampl* may be vectors of common size.

The defaults are *rate* = 8000, *sec* = 1, and *ampl* = 64.

See also: [\[sinewave\]](#), page 801.

sinewave (m, n, d)

Return an *m*-element vector with *i*-th element given by `sin (2 * pi * (i+d-1) / n)`.

The default value for *d* is 0 and the default value for *n* is *m*.

See also: [\[sinetone\]](#), page 801.

spectral_adf (c)

spectral_adf (c, win)

spectral_adf (c, win, b)

Return the spectral density estimator given a vector of autocovariances *c*, window name *win*, and bandwidth, *b*.

The window name, e.g., "triangle" or "rectangle" is used to search for a function called *win_lw*.

If *win* is omitted, the triangle window is used.

If *b* is omitted, $1 / \sqrt{\text{length}(\text{x})}$ is used.

See also: [\[spectral_xdf\]](#), page 802.

```
spectral_xdf (x)
spectral_xdf (x, win)
spectral_xdf (x, win, b)
```

Return the spectral density estimator given a data vector *x*, window name *win*, and bandwidth, *b*.

The window name, e.g., "triangle" or "rectangle" is used to search for a function called *win_sw*.

If *win* is omitted, the triangle window is used.

If *b* is omitted, $1 / \sqrt{\text{length}(\text{x})}$ is used.

See also: [spectral_adf], page 801.

```
spencer (x)
```

Return Spencer's 15 point moving average of each column of *x*.

```
y = stft (x)
y = stft (x, win_size)
y = stft (x, win_size, inc)
y = stft (x, win_size, inc, num_coef)
y = stft (x, win_size, inc, num_coef, win_type)
[y, c] = stft (...)
```

Compute the short-time Fourier transform of the vector *x* with *num_coef* coefficients by applying a window of *win_size* data points and an increment of *inc* points.

Before computing the Fourier transform, one of the following windows is applied:

```
"hanning"
    win_type = 1
"hamming"
    win_type = 2
"rectangle"
    win_type = 3
```

The window names can be passed as strings or by the *win_type* number.

The following defaults are used for unspecified arguments: *win_size* = 80, *inc* = 24, *num_coef* = 64, and *win_type* = 1.

y = *stft* (*x*, ...) returns the absolute values of the Fourier coefficients according to the *num_coef* positive frequencies.

[*y*, *c*] = *stft* (*x*, ...) returns the entire STFT-matrix *y* and a 3-element vector *c* containing the window size, increment, and window type, which is needed by the *synthesis* function.

See also: [synthesis], page 802.

```
x = synthesis (y, c)
```

Compute a signal from its short-time Fourier transform *y* and a 3-element vector *c* specifying window size, increment, and window type.

The values *y* and *c* can be derived by

```
[y, c] = stft (x, ...)
```

See also: [stft], page 802.

[*a*, *v*] = **yulewalker** (*c*)

Fit an AR (*p*)-model with Yule-Walker estimates given a vector *c* of autocovariances [*gamma_0*, ..., *gamma_p*].

Returns the AR coefficients, *a*, and the variance of white noise, *v*.

32 Image Processing

Since an image is basically a matrix, Octave is a very powerful environment for processing and analyzing images. To illustrate how easy it is to do image processing in Octave, the following example will load an image, smooth it by a 5-by-5 averaging filter, and compute the gradient of the smoothed image.

```
I = imread ("myimage.jpg");
S = conv2 (I, ones (5, 5) / 25, "same");
[Dx, Dy] = gradient (S);
```

In this example `S` contains the smoothed image, and `Dx` and `Dy` contains the partial spatial derivatives of the image.

32.1 Loading and Saving Images

The first step in most image processing tasks is to load an image into Octave which is done with the `imread` function. The `imwrite` function is the corresponding function for writing images to the disk.

In summary, most image processing code will follow the structure of this code

```
I = imread ("my_input_image.img");
J = process_my_image (I);
imwrite (J, "my_output_image.img");

[img, map, alpha] = imread (filename)
[...] = imread (url)
[...] = imread (... , ext)
[...] = imread (... , idx)
[...] = imread (... , param1, value1, ...)
```

Read images from various file formats.

Read an image as a matrix from the file `filename` or from the online resource `url`. If neither is given, but `ext` was specified, look for a file with the extension `ext`.

The size and class of the output depends on the format of the image. A color image is returned as an `MxNx3` matrix. Grayscale and black-and-white images are of size `MxN`. Multipage images will have an additional 4th dimension.

The bit depth of the image determines the class of the output: `"uint8"`, `"uint16"`, or `"single"` for grayscale and color, and `"logical"` for black-and-white. Note that indexed images always return the indexes for a colormap, independent of whether `map` is a requested output. To obtain the actual RGB image, use `ind2rgb`. When more than one indexed image is being read, `map` is obtained from the first. In some rare cases this may be incorrect and `imfinfo` can be used to obtain the colormap of each image.

See the Octave manual for more information in representing images.

Some file formats, such as TIFF and GIF, are able to store multiple images in a single file. `idx` can be a scalar or vector specifying the index of the images to read. By default, Octave will read only the first page.

Depending on the file format, it is possible to configure the reading of images with *parameter, value* pairs. The following options are supported:

"Frames" or "Index"

This is an alternative method to specify *idx*. When specifying it in this way, its value can also be the string "*all*".

"Info" This option exists for MATLAB compatibility, but has no effect. For maximum performance when reading multiple images from a single file, use the "*Index*" option.

"PixelRegion"

Controls the image region that is read. The value must be a cell array with two arrays of 3 elements {[*rows*], [*cols*]}. The elements in the array are the start, increment, and end pixel to be read. If the increment value is omitted it defaults to 1. For example, the following are all equivalent:

```
imread (filename, "PixelRegion", {[200 600], [300 700]});  
imread (filename, "PixelRegion", {[200 1 600], [300 1 700]});  
imread (filename)(200:600, 300:700);
```

See also: [\[imwrite\]](#), page 806, [\[imfinfo\]](#), page 808, [\[imformats\]](#), page 810.

```
imwrite (img, filename)  
imwrite (img, filename, ext)  
imwrite (img, map, filename)  
imwrite (... , param1, val1, ... )
```

Write images in various file formats.

The image *img* can be a binary, grayscale, RGB, or multi-dimensional image. The size and class of *img* should be the same as what should be expected when reading it with **imread**: the 3rd and 4th dimensions reserved for color space, and multiple pages respectively. If it's an indexed image, the colormap *map* must also be specified.

If *ext* is not supplied, the file extension of *filename* is used to determine the format. The actual supported formats are dependent on options made during the build of Octave. Use **imformats** to check the support of the different image formats.

Depending on the file format, it is possible to configure the writing of images with *param, val* pairs. The following options are supported:

'Alpha' Alpha (transparency) channel for the image. This must be a matrix with same class, and number of rows and columns of *img*. In case of a multipage image, the size of the 4th dimension must also match and the third dimension must be a singleton. By default, image will be completely opaque.

'Compression'

Compression to use one the image. Can be one of the following: "none" (default), "bzip", "fax3", "fax4", "jpeg", "lzw", "rle", or "deflate". Note that not all compression types are available for all image formats in which it defaults to your Magick library.

'DelayTime'

For formats that accept animations (such as GIF), controls for how long a frame is displayed until it moves to the next one. The value must be scalar (which will applied to all frames in *img*), or a vector of length equal to the number of frames in *im*. The value is in seconds, must be between 0 and 655.35, and defaults to 0.5.

'DisposalMethod'

For formats that accept animations (such as GIF), controls what happens to a frame before drawing the next one. Its value can be one of the following strings: "doNotSpecify" (default); "leaveInPlace"; "restoreBG"; and "restorePrevious", or a cell array of those string with length equal to the number of frames in *img*.

'LoopCount'

For formats that accept animations (such as GIF), controls how many times the sequence is repeated. A value of Inf means an infinite loop (default), a value of 0 or 1 that the sequence is played only once (loops zero times), while a value of 2 or above loops that number of times (looping twice means it plays the complete sequence 3 times). This option is ignored when there is only a single image at the end of writing the file.

'Quality'

Set the quality of the compression. The value should be an integer between 0 and 100, with larger values indicating higher visual quality and lower compression. Defaults to 75.

'WriteMode'

Some file formats, such as TIFF and GIF, are able to store multiple images in a single file. This option specifies if *img* should be appended to the file (if it exists) or if a new file should be created for it (possibly overwriting an existing file). The value should be the string "Overwrite" (default), or "Append".

Despite this option, the most efficient method of writing a multipage image is to pass a 4 dimensional *img* to `imwrite`, the same matrix that could be expected when using `imread` with the option "Index" set to "all".

See also: [\[imread\]](#), page 805, [\[imfinfo\]](#), page 808, [\[imformats\]](#), page 810.

```
val = IMAGE_PATH ()
old_val = IMAGE_PATH (new_val)
IMAGE_PATH (new_val, "local")
```

Query or set the internal variable that specifies a colon separated list of directories in which to search for image files.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[EXEC_PATH\]](#), page 916, [\[OCTAVE_HOME\]](#), page 927, [\[OCTAVE_EXEC_HOME\]](#), page 927.

It is possible to get information about an image file on disk, without actually reading it into Octave. This is done using the `imfinfo` function which provides read access to many of the parameters stored in the header of the image file.

```
info = imfinfo (filename)
info = imfinfo (url)
info = imfinfo (... , ext)
```

Read image information from a file.

`imfinfo` returns a structure containing information about the image stored in the file `filename`. If there is no file `filename`, and `ext` was specified, it will look for a file named `filename` and extension `ext`, i.e., a file named `filename.ext`.

The output structure `info` contains the following fields:

`'Filename'`

The full name of the image file.

`'FileModDate'`

Date of last modification to the file.

`'FileSize'`

Number of bytes of the image on disk

`'Format'` Image format (e.g., "jpeg").

`'Height'` Image height in pixels.

`'Width'` Image Width in pixels.

`'BitDepth'`

Number of bits per channel per pixel.

`'ColorType'`

Image type. Value is "grayscale", "indexed", "truecolor", "CMYK", or "undefined".

`'XResolution'`

X resolution of the image.

`'YResolution'`

Y resolution of the image.

`'ResolutionUnit'`

Units of image resolution. Value is "Inch", "Centimeter", or "undefined".

`'DelayTime'`

Time in 1/100ths of a second (0 to 65535) which must expire before displaying the next image in an animated sequence.

`'LoopCount'`

Number of iterations to loop an animation.

`'ByteOrder'`

Endian option for formats that support it. Value is "little-endian", "big-endian", or "undefined".

'Gamma'	Gamma level of the image. The same color image displayed on two different workstations may look different due to differences in the display monitor.
'Quality'	JPEG/MIFF/PNG compression level. Value is an integer in the range [0 100].
'DisposalMethod'	Only valid for GIF images, control how successive frames are rendered (how the preceding frame is disposed of) when creating a GIF animation. Values can be "doNotSpecify", "leaveInPlace", "restoreBG", or "restorePrevious". For non-GIF files, value is an empty string.
'Chromaticities'	Value is a 1x8 Matrix with the x,y chromaticity values for white, red, green, and blue points, in that order.
'Comment'	Image comment.
'Compression'	Compression type. Value can be "none", "bzip", "fax3", "fax4", "jpeg", "lzw", "rle", "deflate", "lzma", "jpeg2000", "jbig2", "jbig2", or "undefined".
'Colormap'	Colormap for each image.
'Orientation'	The orientation of the image with respect to the rows and columns. Value is an integer between 1 and 8 as defined in the TIFF 6 specifications, and for MATLAB compatibility.
'Software'	Name and version of the software or firmware of the camera or image input device used to generate the image.
'Make'	The manufacturer of the recording equipment. This is the manufacture of the DSC, scanner, video digitizer or other equipment that generated the image.
'Model'	The model name or model number of the recording equipment as mentioned on the field "Make".
'DateTime'	The date and time of image creation as defined by the Exif standard, i.e., it is the date and time the file was changed.
'ImageDescription'	The title of the image as defined by the Exif standard.
'Artist'	Name of the camera owner, photographer or image creator.
'Copyright'	Copyright notice of the person or organization claiming rights to the image.

`'DigitalCamera'`

A struct with information retrieved from the Exif tag.

`'GPSInfo'` A struct with geotagging information retrieved from the Exif tag.

See also: [imread], page 805, [imwrite], page 806, [imshow], page 811, [imformats], page 810.

By default, Octave's image IO functions (`imread`, `imwrite`, and `imfinfo`) use the `GraphicsMagick` library for their operations. This means a vast number of image formats is supported but considering the large amount of image formats in science and its commonly closed nature, it is impossible to have a library capable of reading them all. Because of this, the function `imformats` keeps a configurable list of available formats, their extensions, and what functions should the image IO functions use. This allows one to expand Octave's image IO capabilities by creating functions aimed at acting on specific file formats.

While it would be possible to call the extra functions directly, properly configuring Octave with `imformats` allows one to keep a consistent code that is abstracted from file formats.

It is important to note that a file format is not actually defined by its file extension and that `GraphicsMagick` is capable to read and write more file formats than the ones listed by `imformats`. What this means is that even with an incorrect or missing extension the image may still be read correctly, and that even unlisted formats are not necessarily unsupported.

```
imformats ()
formats = imformats (ext)
formats = imformats (format)
formats = imformats ("add", format)
formats = imformats ("remove", ext)
formats = imformats ("update", ext, format)
formats = imformats ("factory")
```

Manage supported image formats.

`formats` is a structure with information about each supported file format, or from a specific format `ext`, the value displayed on the field `ext`. It contains the following fields:

ext	The name of the file format. This may match the file extension but Octave will automatically detect the file format.
-----	--

description	A long description of the file format.
isa	A function handle to confirm if a file is of the specified format.
write	A function handle to write if a file is of the specified format.
read	A function handle to open files the specified format.
info	A function handle to obtain image information of the specified format.
alpha	Logical value if format supports alpha channel (transparency or matte).
multipage	Logical value if format supports multipage (multiple images per file).

It is possible to change the way Octave manages file formats with the options "`add`", "`remove`", and "`update`", and supplying a structure `format` with the required fields. The option "`factory`" resets the configuration to the default.

This can be used by Octave packages to extend the image reading capabilities Octave, through use of the `PKG_ADD` and `PKG_DEL` commands.

See also: [\[imfinfo\]](#), page 808, [\[imread\]](#), page 805, [\[imwrite\]](#), page 806.

32.2 Displaying Images

A natural part of image processing is visualization of an image. The most basic function for this is the `imshow` function that shows the image given in the first input argument.

```
imshow (im)
imshow (im, limits)
imshow (im, map)
imshow (rgb, ...)
imshow (filename)
imshow (..., string_param1, value1, ...)
h = imshow (...)
```

Display the image `im`, where `im` can be a 2-dimensional (grayscale image) or a 3-dimensional (RGB image) matrix.

If `limits` is a 2-element vector `[low, high]`, the image is shown using a display range between `low` and `high`. If an empty matrix is passed for `limits`, the display range is computed as the range between the minimal and the maximal value in the image.

If `map` is a valid color map, the image will be shown as an indexed image using the supplied color map.

If a filename is given instead of an image, the file will be read and shown.

If given, the parameter `string_param1` has value `value1`. `string_param1` can be any of the following:

`"displayrange"`

`value1` is the display range as described above.

`"colormap"`

`value1` is the colormap to use when displaying an indexed image.

`"xdata"`

If `value1` is a 2-element vector, it must contain horizontal image limits in the form `[xfirst, xlast]`, where `xfirst` and `xlast` are the abscissa of the centers of the corner pixels. Otherwise `value1` must be a vector and only the first and last elements will be used for `xfirst` and `xlast` respectively.

`"ydata"`

If `value1` is a 2-element vector, it must contain vertical image limits in the form `[yfirst, ylast]`, where `yfirst` and `ylast` are the ordinates of the center of the corner pixels. Otherwise `value1` must be a vector and only the first and last elements will be used for `yfirst` and `ylast` respectively.

The optional return value `h` is a graphics handle to the image.

See also: [\[image\]](#), page 812, [\[imagesc\]](#), page 812, [\[colormap\]](#), page 816, [\[gray2ind\]](#), page 814, [\[rgb2ind\]](#), page 814.

```
image (img)
image (x, y, img)
image (... , "prop", val, ...)
image ("prop1", val1, ...)
h = image (...)
```

Display a matrix as an indexed color image.

The elements of *img* are indices into the current colormap.

x and *y* are optional 2-element vectors, [*min*, *max*], which specify the coordinates of the centers of the corner pixels. If a range is specified as [*max*, *min*] then the image will be reversed along that axis. For convenience, *x* and *y* may be specified as N-element vectors matching the length of the data in *img*. However, only the first and last elements will be used to determine the axis limits.

Multiple property/value pairs may be specified for the image object, but they must appear in pairs.

The optional return value *h* is a graphics handle to the image.

Implementation Note: The origin (0, 0) for images is located in the upper left. For ordinary plots, the origin is located in the lower left. Octave handles this inversion by plotting the data normally, and then reversing the direction of the y-axis by setting the *ydir* property to "reverse". This has implications whenever an image and an ordinary plot need to be overlaid. The recommended solution is to display the image and then plot the reversed ydata using, for example, *flipud* (*ydata*).

Calling Forms: The *image* function can be called in two forms: High-Level and Low-Level. When invoked with normal options, the High-Level form is used which first calls *newplot* to prepare the graphic figure and axes. When the only inputs to *image* are property/value pairs the Low-Level form is used which creates a new instance of an image object and inserts it in the current axes.

Graphic Properties: The full list of properties is documented at [Section 15.3.3.6 \[Image Properties\]](#), page 425.

See also: [\[imshow\]](#), page 811, [\[imagesc\]](#), page 812, [\[colormap\]](#), page 816.

```
imagesc (img)
imagesc (x, y, img)
imagesc (... , climits)
imagesc (... , "prop", val, ...)
imagesc ("prop1", val1, ...)
imagesc (hax, ...)
h = imagesc (...)
```

Display a scaled version of the matrix *img* as a color image.

The colormap is scaled so that the entries of the matrix occupy the entire colormap. If *climits* = [*lo*, *hi*] is given, then that range is set to the "clim" of the current axes.

x and *y* are optional 2-element vectors, [*min*, *max*], which specify the coordinates of the centers of the corner pixels. If a range is specified as [*max*, *min*] then the image will be reversed along that axis. For convenience, *x* and *y* may be specified

as N-element vectors matching the length of the data in *img*. However, only the first and last elements will be used to determine the image limits.

The optional return value *h* is a graphics handle to the image.

Calling Forms: The `imagesc` function can be called in two forms: High-Level and Low-Level. When invoked with normal options, the High-Level form is used which first calls `newplot` to prepare the graphic figure and axes. When the only inputs to `image` are property/value pairs the Low-Level form is used which creates a new instance of an image object and inserts it in the current axes.

See also: `[image]`, page 812, `[imshow]`, page 811, `[caxis]`, page 325.

32.3 Representing Images

In general Octave supports four different kinds of images, grayscale images, RGB images, binary images, and indexed images. A grayscale image is represented with an M-by-N matrix in which each element corresponds to the intensity of a pixel. An RGB image is represented with an M-by-N-by-3 array where each 3-vector corresponds to the red, green, and blue intensities of each pixel.

The actual meaning of the value of a pixel in a grayscale or RGB image depends on the class of the matrix. If the matrix is of class `double` pixel intensities are between 0 and 1, if it is of class `uint8` intensities are between 0 and 255, and if it is of class `uint16` intensities are between 0 and 65535.

A binary image is an M-by-N matrix of class `logical`. A pixel in a binary image is black if it is `false` and white if it is `true`.

An indexed image consists of an M-by-N matrix of integers and a C-by-3 color map. Each integer corresponds to an index in the color map, and each row in the color map corresponds to an RGB color. The color map must be of class `double` with values between 0 and 1.

The following convenience functions are available for conversion between image formats.

```
im2double (img)
im2double (img, "indexed")
```

Convert image to double precision.

The conversion of *img* to double precision, is dependent on the type of input image. The following input classes are supported:

`'uint8, uint16, and int16'`

The range of values from the class is scaled to the interval [0 1].

`'logical'` True and false values are assigned a value of 0 and 1 respectively.

`'single'` Values are cast to double.

`'double'` Returns the same image.

If *img* is an indexed image, then the second argument should be the string "`indexed`". If so, then *img* must either be of floating point class, or unsigned integer class and it will simply be cast to double. If it is an integer class, a +1 offset is applied.

See also: `[double]`, page 47.

```
img = gray2ind (I)
img = gray2ind (I, n)
img = gray2ind (BW)
img = gray2ind (BW, n)
[img, map] = gray2ind (...)
```

Convert a grayscale or binary intensity image to an indexed image.

The indexed image will consist of n different intensity values. If not given n defaults to 64 for grayscale images or 2 for binary black and white images.

The output img is of class `uint8` if n is less than or equal to 256; Otherwise the return class is `uint16`.

See also: [\[ind2gray\]](#), page 814, [\[rgb2ind\]](#), page 814.

```
I = ind2gray (x, map)
```

Convert a color indexed image to a grayscale intensity image.

The image x must be an indexed image which will be converted using the colormap map . If map does not contain enough colors for the image, pixels in x outside the range are mapped to the last color in the map before conversion to grayscale.

The output I is of the same class as the input x and may be one of `uint8`, `uint16`, `single`, or `double`.

Implementation Note: There are several ways of converting colors to grayscale intensities. This functions uses the luminance value obtained from `rgb2gray` which is $I = 0.299*R + 0.587*G + 0.114*B$. Other possibilities include the value component from `rgb2hsv` or using a single color channel from `ind2rgb`.

See also: [\[gray2ind\]](#), page 814, [\[ind2rgb\]](#), page 814.

```
[x, map] = rgb2ind (rgb)
```

```
[x, map] = rgb2ind (R, G, B)
```

Convert an image in red-green-blue (RGB) color space to an indexed image.

The input image rgb can be specified as a single matrix of size $M \times N \times 3$, or as three separate variables, R , G , and B , its three color channels, red, green, and blue.

It outputs an indexed image x and a colormap map to interpret an image exactly the same as the input. No dithering or other form of color quantization is performed. The output class of the indexed image x can be `uint8`, `uint16` or `double`, whichever is required to specify the number of unique colors in the image (which will be equal to the number of rows in map) in order.

Multi-dimensional indexed images (of size $M \times N \times 3 \times K$) are also supported, both via a single input (rgb) or its three color channels as separate variables.

See also: [\[ind2rgb\]](#), page 814, [\[rgb2hsv\]](#), page 824, [\[rgb2gray\]](#), page 825.

```
rgb = ind2rgb (x, map)
```

```
[R, G, B] = ind2rgb (x, map)
```

Convert an indexed image to red, green, and blue color components.

The image x must be an indexed image which will be converted using the colormap map . If map does not contain enough colors for the image, pixels in x outside the range are mapped to the last color in the map.

The output may be a single RGB image (MxNx3 matrix where M and N are the original image x dimensions, one for each of the red, green and blue channels). Alternatively, the individual red, green, and blue color matrices of size MxN may be returned.

Multi-dimensional indexed images (of size MxNx1xK) are also supported.

See also: [rgb2ind], page 814, [ind2gray], page 814, [hsv2rgb], page 824.

Octave also provides tools to produce and work with movie frame structures. Those structures encapsulate the image data ("cdata" field) together with the corresponding colormap ("colormap" field).

```
frame = getframe ()
frame = getframe (hax)
frame = getframe (hfig)
frame = getframe (..., rect)
```

Capture a figure or axes as a movie frame structure.

Without an argument, capture the current axes excluding ticklabels, title, and x/y/zlabels. The returned structure *frame* has a field *cdata*, which contains the actual image data in the form of an NxMx3 (RGB) uint8 matrix, and a field *colormap* which is provided for MATLAB compatibility but is always empty.

If the first argument *hax* is an axes handle, then capture this axes, rather than the current axes returned by *gca*.

If the first argument *hfig* is a figure handle then the entire corresponding figure canvas is captured.

Finally, if a second argument *rect* is provided it must be a four-element vector ([left bottom width height]) defining the region inside the figure to be captured. Regardless of the figure "units" property, *rect* must be defined in **pixels**.

See also: [im2frame], page 816, [frame2im], page 816, [movie], page 815.

```
movie (mov)
movie (mov, n)
movie (mov, n, fps)
movie (h, ...)
```

Play a movie defined by an array of frame structures.

The movie *mov* must be a struct array of frames with fields "cdata" and "colormap", as returned by the *getframe* function. By default all images are displayed once, at 12 fps, in the current axes.

The optional argument *n* is a scalar or vector of integers that controls the number of times the movie is displayed and which particular frames are shown:

First element:

- n*(1) > 0 Play the movie *n*(1) times.
- n*(1) < 0 Play the movie **abs** (*n*(1)) times alternatively in forward and backward order.

Other elements (if any):

Indices of the frames in *mov* that will be displayed.

If the first argument is a handle to a figure or axes *h*, the movie is played in that figure or axes instead of the current axes.

See also: [getframe], page 815, [im2frame], page 816, [frame2im], page 816.

[*x*, *map*] = frame2im (*frame*)

Convert movie frame to indexed image.

A movie frame is simply a struct with the fields "cdata" and "colormap".

Support for N-dimensional images or movies is given when *frame* is a struct array. In such cases, *x* will be a MxNx1xK or MxNx3xK for indexed and RGB movies respectively, with each frame concatenated along the 4th dimension.

See also: [im2frame], page 816, [getframe], page 815.

im2frame (*rgb*)

im2frame (*x*, *map*)

Convert image to movie frame.

A movie frame is simply a struct with the fields "cdata" and "colormap".

Support for N-dimensional images is given when each image projection, matrix sizes of MxN and MxNx3 for RGB images, is concatenated along the fourth dimension. In such cases, the returned value is a struct array.

See also: [frame2im], page 816.

The **colormap** function is used to change the colormap of the current axes or figure.

```
cmap = colormap ()
cmap = colormap (map)
cmap = colormap ("default")
cmap = colormap (map_name)
cmap = colormap (hax, ...)
colormap map_name
```

Query or set the current colormap.

With no input arguments, **colormap** returns the current color map.

colormap (map) sets the current colormap to *map*. The colormap should be an *n* row by 3 column matrix. The columns contain red, green, and blue intensities respectively. All entries must be between 0 and 1 inclusive. The new colormap is returned.

colormap ("default") restores the default colormap (the **viridis** map with 64 entries). The default colormap is returned.

The map may also be specified by a string, *map_name*, which is the name of a function that returns a colormap.

If the first argument *hax* is an axes handle, then the colormap for those axes is queried or set.

For convenience, it is also possible to use this function with the command form, **colormap map_name**.

The list of built-in colormaps is:

Map	Description
viridis	default
jet	colormap traversing blue, cyan, green, yellow, red.
cubehelix	colormap traversing black, blue, green, red, white with increasing intensity.
hsv	cyclic colormap traversing Hue, Saturation, Value space.
rainbow	colormap traversing red, yellow, blue, green, violet.
hot	colormap traversing black, red, orange, yellow, white.
cool	colormap traversing cyan, purple, magenta.
spring	colormap traversing magenta to yellow.
summer	colormap traversing green to yellow.
autumn	colormap traversing red, orange, yellow.
winter	colormap traversing blue to green.
gray	colormap traversing black to white in shades of gray.
bone	colormap traversing black, gray-blue, white.
copper	colormap traversing black to light copper.
pink	colormap traversing black, gray-pink, white.
ocean	colormap traversing black, dark-blue, white.
colorcube	equally spaced colors in RGB color space.
flag	cyclic 4-color map of red, white, blue, black.
lines	cyclic colormap with colors from axes "ColorOrder" property.
prism	cyclic 6-color map of red, orange, yellow, green, blue, violet.
white	all white colormap (no colors).

See also: [viridis], page 821, [jet], page 820, [cubehelix], page 819, [hsv], page 820, [rainbow], page 821, [hot], page 819, [cool], page 818, [spring], page 821, [summer], page 821, [autumn], page 818, [winter], page 822, [gray], page 819, [bone], page 818, [copper], page 819, [pink], page 820, [ocean], page 820, [colorcube], page 818, [flag], page 819, [lines], page 820, [prism], page 821, [white], page 821.

iscolormap (cmap)

Return true if *cmap* is a colormap.

A colormap is a real matrix, of class single or double, with 3 columns. Each row represents a single color. The 3 columns contain red, green, and blue intensities respectively.

All values in a colormap should be in the [0 1] range but this is not enforced. Each function must decide what to do for values outside this range.

See also: [colormap], page 816, [rgbplot], page 817.

The following functions return predefined colormaps, the same that can be requested by name using the **colormap** function.

rgbplot (cmap) **rgbplot (cmap, style)**

`h = rgbplot (...)`

Plot the components of a colormap.

Two different *styles* are available for displaying the *cmap*:

profile (default)

Plot the RGB line profile of the colormap for each of the channels (red, green and blue) with the plot lines colored appropriately. Each line represents the intensity of an RGB component across the colormap.

composite Draw the colormap across the X-axis so that the actual index colors are visible rather than the individual color components.

The optional return value *h* is a graphics handle to the created plot.

Run `demo rgbplot` to see an example of `rgbplot` and each style option.

See also: [colormap], page 816.

`map = autumn ()`

`map = autumn (n)`

Create color colormap. This colormap ranges from red through orange to yellow.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

`map = bone ()`

`map = bone (n)`

Create color colormap. This colormap varies from black to white with gray-blue shades.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

`map = colordcube ()`

`map = colordcube (n)`

Create color colormap. This colormap is composed of as many equally spaced colors (not grays) in the RGB color space as possible.

If there are not a perfect number *n* of regularly spaced colors then the remaining entries in the colormap are gradients of pure red, green, blue, and gray.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

`map = cool ()`

`map = cool (n)`

Create color colormap. The colormap varies from cyan to magenta.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = copper ()  
map = copper (n)
```

Create color colormap. This colormap varies from black to a light copper tone.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = cubehelix ()  
map = cubehelix (n)
```

Create cubehelix colormap.

This colormap varies from black to white going through blue, green, and red tones while maintaining a monotonically increasing perception of intensity. This is achieved by traversing a color cube from black to white through a helix, hence the name cubehelix, while taking into account the perceived brightness of each channel according to the NTSC specifications from 1953.

```
rgbplot (cubehelix (256))
```

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

Reference: Green, D. A., 2011, "A colour scheme for the display of astronomical intensity images", Bulletin of the Astronomical Society of India, 39, 289.

See also: [colormap], page 816.

```
map = flag ()  
map = flag (n)
```

Create color colormap. This colormap cycles through red, white, blue, and black with each index change.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = gray ()  
map = gray (n)
```

Create gray colormap. This colormap varies from black to white with shades of gray.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = hot ()  
map = hot (n)
```

Create color colormap. This colormap ranges from black through dark red, red, orange, yellow, to white.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

hsv (n)

Create color colormap. This colormap begins with red, changes through yellow, green, cyan, blue, and magenta, before returning to red.

It is useful for displaying periodic functions. The map is obtained by linearly varying the hue through all possible values while keeping constant maximum saturation and value. The equivalent code is `hsv2rgb ([(0:N-1)'/N, ones(N,2)])`.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = jet ()  
map = jet (n)
```

Create color colormap. This colormap ranges from dark blue through blue, cyan, green, yellow, red, to dark red.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = lines ()  
map = lines (n)
```

Create color colormap. This colormap is composed of the list of colors in the current axes "ColorOrder" property. The default is blue, orange, yellow, purple, green, light blue, and dark red.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = ocean ()  
map = ocean (n)
```

Create color colormap. This colormap varies from black to white with shades of blue.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = pink ()  
map = pink (n)
```

Create color colormap. This colormap varies from black to white with shades of gray-pink.

This colormap gives a sepia tone when used on grayscale images.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = prism ()  
map = prism (n)
```

Create color colormap. This colormap cycles through red, orange, yellow, green, blue and violet with each index change.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = rainbow ()  
map = rainbow (n)
```

Create color colormap. This colormap ranges from red through orange, yellow, green, blue, to violet.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = spring ()  
map = spring (n)
```

Create color colormap. This colormap varies from magenta to yellow.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = summer ()  
map = summer (n)
```

Create color colormap. This colormap varies from green to yellow.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = viridis ()  
map = viridis (n)
```

Create color colormap. This colormap ranges from dark purplish-blue through blue, green, to yellow.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = white ()  
map = white (n)
```

Create color colormap. This colormap is completely white.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
map = winter ()  
map = winter (n)
```

Create color colormap. This colormap varies from blue to green.

The argument *n* must be a scalar. If unspecified, the length of the current colormap, or 64, is used.

See also: [colormap], page 816.

```
cmap = contrast (x)  
cmap = contrast (x, n)
```

Return a gray colormap that maximizes the contrast in an image.

The returned colormap will have *n* rows. If *n* is not defined then the size of the current colormap is used.

See also: [colormap], page 816, [brighten], page 822.

The following three functions modify the existing colormap rather than replace it.

```
map_out = brighten (beta)  
map_out = brighten (map, beta)  
map_out = brighten (h, beta)  
brighten (...)
```

Brighten or darken a colormap.

The argument *beta* must be a scalar between -1 and 1, where a negative value darkens and a positive value brightens the colormap.

If the *map* argument is omitted, the function is applied to the current colormap.

The first argument can also be a valid graphics handle *h*, in which case `brighten` is applied to the colormap associated with this handle.

If no output is specified then the result is written to the current colormap.

See also: [colormap], page 816, [contrast], page 822.

```
spinmap ()  
spinmap (t)  
spinmap (t, inc)  
spinmap ("inf")
```

Cycle the colormap for *t* seconds with a color increment of *inc*.

Both parameters are optional. The default cycle time is 5 seconds and the default increment is 2. If the option "inf" is given then cycle continuously until *Control-C* is pressed.

When rotating, the original color 1 becomes color 2, color 2 becomes color 3, etc. A positive or negative increment is allowed and a higher value of *inc* will cause faster cycling through the colormap.

See also: [colormap], page 816.

```
whitebg ()  
whitebg (color)  
whitebg ("none")
```

```
whitebg (hfig)
whitebg (hfig, color)
whitebg (hfig, "none")
```

Invert the colors in the current color scheme.

The root properties are also inverted such that all subsequent plots will use the new color scheme.

If the optional argument *color* is present then the background color is set to *color* rather than inverted. *color* may be a string representing one of the eight known colors or an RGB triplet. The special string argument "*none*" restores the plot to the factory default colors.

If the first argument *hfig* is a figure handle or list of figure handles, then operate on these figures rather than the current figure returned by *gcf*. The root properties will not be changed unless 0 is in the list of figures.

Programming Note: *whitebg* operates by changing the color properties of the children of the specified figures. Only objects with a single color are affected. For example, a patch with a single "FaceColor" will be changed, but a patch with shading ("interp") will not be modified. For inversion, the new color is simply the inversion in RGB space: *cnew* = [1-R 1-G 1-B]. When a color is specified, the axes and figure are set to the new color, and the color of child objects are then adjusted to have some contrast (visibility) against the new background.

See also: [\[reset\]](#), page 456, [\[get\]](#), page 402, [\[set\]](#), page 402.

The following functions can be used to manipulate colormaps.

```
[Y, newmap] = cmunique (X, map)
[Y, newmap] = cmunique (RGB)
[Y, newmap] = cmunique (I)
```

Convert an input image *X* to an output indexed image *Y* which uses the smallest colormap possible *newmap*.

When the input is an indexed image (*X* with colormap *map*) the output is a colormap *newmap* from which any repeated rows have been eliminated. The output image, *Y*, is the original input image with the indices adjusted to match the new, possibly smaller, colormap.

When the input is an RGB image (an MxNx3 array), the output colormap will contain one entry for every unique color in the original image. In the worst case the new map could have as many rows as the number of pixels in the original image.

When the input is a grayscale image *I*, the output colormap will contain one entry for every unique intensity value in the original image. In the worst case the new map could have as many rows as the number of pixels in the original image.

Implementation Details:

newmap is always an Mx3 matrix, even if the input image is an intensity grayscale image *I* (all three RGB planes are assigned the same value).

The output image is of class uint8 if the size of the new colormap is less than or equal to 256. Otherwise, the output image is of class double.

See also: [\[rgb2ind\]](#), page 814, [\[gray2ind\]](#), page 814.

```
[Y, newmap] = cmpermute (X, map)
[Y, newmap] = cmpermute (X, map, index)
Reorder colors in a colormap.
```

When called with only two arguments, `cmpermute` randomly rearranges the colormap *map* and returns a new colormap *newmap*. It also returns the indexed image *Y* which is the equivalent of the original input image *X* when displayed using *newmap*.

When called with an optional third argument the order of colors in the new colormap is defined by *index*.

Caution: *index* should not have repeated elements or the function will fail.

32.4 Plotting on top of Images

If gnuplot is being used to display images it is possible to plot on top of images. Since an image is a matrix it is indexed by row and column values. The plotting system is, however, based on the traditional (x, y) system. To minimize the difference between the two systems Octave places the origin of the coordinate system in the point corresponding to the pixel at $(1, 1)$. So, to plot points given by row and column values on top of an image, one should simply call `plot` with the column values as the first argument and the row values as the second. As an example the following code generates an image with random intensities between 0 and 1, and shows the image with red circles over pixels with an intensity above 0.99.

```
I = rand (100, 100);
[row, col] = find (I > 0.99);
hold ("on");
imshow (I);
plot (col, row, "ro");
hold ("off");
```

32.5 Color Conversion

Octave supports conversion from the RGB color system to the HSV color system and vice versa. It is also possible to convert from a color RGB image to a grayscale image.

```
hsv_map = rgb2hsv (rgb_map)
hsv_img = rgb2hsv (rgb_img)
```

Transform a colormap or image from RGB to HSV color space.

A color in the RGB space consists of red, green, and blue intensities.

A color in HSV space is represented by hue, saturation and value (brightness) levels in a cylindrical coordinate system. Hue is the azimuth and describes the dominant color. Saturation is the radial distance and gives the amount of hue mixed into the color. Value is the height and is the amount of light in the color.

Output class and size will be the same as input.

See also: [\[hsv2rgb\]](#), page 824, [\[rgb2ind\]](#), page 814, [\[rgb2gray\]](#), page 825.

```
rgb_map = hsv2rgb (hsv_map)
rgb_img = hsv2rgb (hsv_img)
```

Transform a colormap or image from HSV to RGB color space.

A color in HSV space is represented by hue, saturation and value (brightness) levels in a cylindrical coordinate system. Hue is the azimuth and describes the dominant color. Saturation is the radial distance and gives the amount of hue mixed into the color. Value is the height and is the amount of light in the color.

The input can be both a colormap or RGB image. In the case of floating point input, values are expected to be on the [0 1] range. In the case of hue (azimuth), since the value corresponds to an angle, $\text{mod}(h, 1)$ is used.

```
>> hsv2rgb ([0.5 1 1])
⇒ ans = 0 1 1

>> hsv2rgb ([2.5 1 1])
⇒ ans = 0 1 1

>> hsv2rgb ([3.5 1 1])
⇒ ans = 0 1 1
```

Output class and size will be the same as input.

See also: [rgb2hsv], page 824, [ind2rgb], page 814.

I = `rgb2gray (rgb_img)`
gray_map = `rgb2gray (rgb_map)`

Transform an image or colormap from red-green-blue (RGB) color space to a grayscale intensity image.

The input may be of class uint8, int8, uint16, int16, single, or double. The output is of the same class as the input.

Implementation Note: The grayscale intensity is calculated as

$$I = 0.298936*R + 0.587043*G + 0.114021*B$$

which corresponds to the luminance channel when RGB is translated to YIQ as documented in <https://en.wikipedia.org/wiki/YIQ>.

See also: [rgb2hsv], page 824, [rgb2ind], page 814.

33 Audio Processing

33.1 Audio File Utilities

The following functions allow you to read, write and retrieve information about audio files. Various formats are supported including wav, flac and ogg vorbis.

`info = audioinfo (filename)`

Return information about an audio file specified by *filename*.

The output *info* is a structure containing the following fields:

`'Filename'`

Name of the audio file.

`'CompressionMethod'`

Audio compression method. Unused, only present for compatibility with MATLAB.

`'NumChannels'`

Number of audio channels.

`'SampleRate'`

Sample rate of the audio, in Hertz.

`'TotalSamples'`

Number of samples in the file.

`'Duration'`

Duration of the audio, in seconds.

`'BitsPerSample'`

Number of bits per sample.

`'BitRate'` Audio bit rate. Unused, only present for compatibility with MATLAB.

`'Title'` "Title" audio metadata value as a string, or empty if not present.

`'Artist'` "Artist" audio metadata value as a string, or empty if not present.

`'Comment'` "Comment" audio metadata value as a string, or empty if not present.

See also: [audioread], page 827, [audiowrite], page 828.

`[y, fs] = audioread (filename)`

`[y, fs] = audioread (filename, samples)`

`[y, fs] = audioread (filename, datatype)`

`[y, fs] = audioread (filename, samples, datatype)`

Read the audio file *filename* and return the audio data *y* and sampling rate *fs*.

The audio data is stored as matrix with rows corresponding to audio frames and columns corresponding to channels.

The optional two-element vector argument *samples* specifies starting and ending frames.

The optional argument *datatype* specifies the datatype to return. If it is "native", then the type of data depends on how the data is stored in the audio file.

See also: [audiowrite], page 828, [audioformats], page 828, [audioinfo], page 827.

```
audiowrite (filename, y, fs)
audiowrite (filename, y, fs, name, value, ...)
```

Write audio data from the matrix *y* to *filename* at sampling rate *fs* with the file format determined by the file extension.

Additional name/value argument pairs may be used to specify the following options:

'BitsPerSample'

Number of bits per sample. Valid values are 8, 16, 24, and 32. Default is 16.

'BitRate' Valid argument name, but ignored. Left for compatibility with MATLAB.

'Quality' Quality setting for the Ogg Vorbis compressor. Values can range between 0 and 100 with 100 being the highest quality setting. Default is 75.

'Title' Title for the audio file.

'Artist' Artist name.

'Comment' Comment.

See also: [audioread], page 827, [audioformats], page 828, [audioinfo], page 827.

```
audioformats ()
audioformats (format)
```

Display information about all supported audio formats.

If the optional argument *format* is given, then display only formats with names that start with *format*.

See also: [audioread], page 827, [audiowrite], page 828.

33.2 Audio Device Information

```
devinfo = audiodevinfo ()
devs = audiodevinfo (io)
name = audiodevinfo (io, id)
id = audiodevinfo (io, name)
id = audiodevinfo (io, rate, bits, chans)
supports = audiodevinfo (io, id, rate, bits, chans)
```

Return a structure describing the available audio input and output devices.

The *devinfo* structure has two fields "input" and "output". The value of each field is a structure array with fields "Name", "DriverVersion" and "ID" describing an audio device.

If the optional argument *io* is 1, return information about input devices only. If it is 0, return information about output devices only. If *io* is the only argument supplied, return the number of input or output devices available.

If the optional argument *id* is provided, return information about the corresponding device.

If the optional argument *name* is provided, return the id of the named device.

Given a sampling rate, bits per sample, and number of channels for an input or output device, return the ID of the first device that supports playback or recording using the specified parameters.

If also given a device ID, return true if the device supports playback or recording using those parameters.

33.3 Audio Player

The following methods are used to create and use audioplayer objects. These objects can be used to play back audio data stored in Octave matrices and arrays. The audioplayer object supports playback from various devices available to the system, blocking and non-blocking playback, convenient pausing and resuming and much more.

```
player = audioplayer (y, fs)
player = audioplayer (y, fs, nbits)
player = audioplayer (y, fs, nbites, id)
player = audioplayer (recorder)
player = audioplayer (recorder, id)
```

Create an audioplayer object that will play back data *y* at sample rate *fs*.

The optional arguments *nbites*, and *id* specify the bit depth and player device id, respectively. Device IDs may be found using the audiodevinfo function. Given an audioplayer object, use the data from the object to initialize the player.

The signal *y* can be a vector or a two-dimensional array.

The following example will create an audioplayer object that will play back one second of white noise at 44100 sample rate using 8 bits per sample.

```
y = 0.25 * randn (2, 44100);
player = audioplayer (y, 44100, 8);
play (player);
```

33.3.1 Playback

The following methods are used to control player playback.

```
play (player)
play (player, start)
play (player, limits)
```

Play audio stored in the audioplayer object *player* without blocking.

Given optional argument *start*, begin playing at *start* samples in the recording. Given a two-element vector *limits*, begin and end playing at the number of samples specified by the elements of the vector.

```
playblocking (player)
playblocking (player, start)
playblocking (player, limits)
```

Play audio stored in the audioplayer object *player* with blocking.

Given optional argument *start*, begin playing at *start* samples in the recording. Given a two-element vector *limits*, begin and end playing at the number of samples specified by the elements of the vector.

pause (player)

Pause the audioplayer *player*.

resume (player)

Resume playback for the paused audioplayer object *player*.

stop (player)

Stop the playback for the audioplayer *player* and reset the relevant variables to their starting values.

isplaying (player)

Return true if the audioplayer object *player* is currently playing back audio and false otherwise.

33.3.2 Properties

The remaining couple of methods are used to get and set various properties of the audioplayer object.

value = get (player, name)

values = get (player)

Return the value of the property identified by *name*.

If *name* is a cell array return the values of the properties identified by the elements of the cell array. Given only the player object, return a scalar structure with values of all properties of *player*. The field names of the structure correspond to property names.

set (player, name, value)

set (player, properties)

properties = set (player)

Set the value of property specified by *name* to a given *value*.

If *name* and *value* are cell arrays, set each property to the corresponding value. Given a structure of *properties* with fields corresponding to property names, set the value of those properties to the field values. Given only the audioplayer object, return a structure of settable properties.

33.4 Audio Recorder

The following methods are used to create and use audiorecorder objects. These objects can be used to record audio data from various devices available to the system. You can use convenient methods to retrieve that data or audioplayer objects created from that data. Methods for blocking and non-blocking recording, pausing and resuming recording and much more is available.

recorder = audiorecorder ()

recorder = audiorecorder (fs, nbts, channels)

recorder = audiorecorder (fs, nbts, channels, id)

Create an audiorecorder object recording 8 bit mono audio at 8000 Hz sample rate.

The optional arguments *fs*, *nbts*, *channels*, and *id* specify the sample rate, bit depth, number of channels and recording device id, respectively. Device IDs may be found using the audiodevinfo function.

33.4.1 Recording

The following methods control the recording process.

record (recorder)

record (recorder, length)

Record audio without blocking using the audiorecorder object *recorder* until stopped or paused by the *stop* or *pause* method.

Given the optional argument *length*, record for *length* seconds.

recordblocking (recorder, length)

Record audio with blocking (synchronous I/O).

The length of the recording in seconds (*length*) must be specified.

pause (recorder)

Pause recording with audiorecorder object *recorder*.

resume (recorder)

Resume recording with the paused audiorecorder object *recorder*.

stop (recorder)

Stop the audiorecorder object *recorder* and clean up any audio streams.

isrecording (recorder)

Return true if the audiorecorder object *recorder* is currently recording audio and false otherwise.

33.4.2 Data Retrieval

The following methods allow you to retrieve recorded audio data in various ways.

data = getaudiodata (recorder)

data = getaudiodata (recorder, datatype)

Return recorder audio data as a matrix with values between -1.0 and 1.0 and with as many columns as there are channels in the recorder.

Given the optional argument *datatype*, convert the recorded data to the specified type, which may be one of "double", "single", "int16", "int8" or "uint8".

player = getplayer (recorder)

Return an audioplayer object with data recorded by the audiorecorder object *recorder*.

player = play (recorder)

player = play (recorder, start)

player = play (recorder, [start, end])

Play the audio recorded in *recorder* and return a corresponding audioplayer object.

If the optional argument *start* is provided, begin playing *start* seconds in to the recording.

If the optional argument *end* is provided, stop playing at *end* seconds in the recording.

33.4.3 Properties

The remaining two methods allow you to read or alter the properties of audiorecorder objects.

```
value = get (recorder, name)
values = get (recorder)
```

Return the *value* of the property identified by *name*.

If *name* is a cell array, return the values of the properties corresponding to the elements of the cell array. Given only the recorder object, return a scalar structure with values of all properties of *recorder*. The field names of the structure correspond to property names.

```
set (recorder, name, value)
set (recorder, properties)
properties = set (recorder)
```

Set the value of property specified by *name* to a given *value*.

If *name* and *value* are cell arrays of the same size, set each property to a corresponding value. Given a structure with fields corresponding to property names, set the value of those properties to the corresponding field values. Given only the recorder object, return a structure of settable properties.

33.5 Audio Data Processing

Octave provides a few functions for dealing with audio data. An audio ‘sample’ is a single output value from an A/D converter, i.e., a small integer number (usually 8 or 16 bits), and audio data is just a series of such samples. It can be characterized by three parameters: the sampling rate (measured in samples per second or Hz, e.g., 8000 or 44100), the number of bits per sample (e.g., 8 or 16), and the number of channels (1 for mono, 2 for stereo, etc.).

There are many different formats for representing such data. Currently, only the two most popular, *linear encoding* and *mu-law encoding*, are supported by Octave. There is an excellent FAQ on audio formats by Guido van Rossum guido@cwi.nl which can be found at any FAQ ftp site, in particular in the directory `/pub/usenet/news.answers/audio-fmts` of the archive site rtfm.mit.edu.

Octave simply treats audio data as vectors of samples (non-mono data are not supported yet). It is assumed that audio files using linear encoding have one of the extensions `lin` or `raw`, and that files holding data in mu-law encoding end in `au`, `mu`, or `snd`.

```
y = lin2mu (x, n)
```

Convert audio data from linear to mu-law.

Mu-law values use 8-bit unsigned integers. Linear values use *n*-bit signed integers or floating point values in the range $-1 \leq x \leq 1$ if *n* is 0.

If *n* is not specified it defaults to 0, 8, or 16 depending on the range of values in *x*.

See also: [\[mu2lin\]](#), page 832.

```
y = mu2lin (x, n)
```

Convert audio data from mu-law to linear.

Mu-law values are 8-bit unsigned integers. Linear values use n -bit signed integers or floating point values in the range $-1 \leq y \leq 1$ if n is 0.

If n is not specified it defaults to 0.

See also: [\[lin2mu\]](#), page 832.

record (sec)

record (sec, fs)

Record sec seconds of audio from the system's default audio input at a sampling rate of 8000 samples per second.

If the optional argument fs is given, it specifies the sampling rate for recording.

For more control over audio recording, use the **audiorecorder** class.

See also: [\[sound\]](#), page 833, [\[soundsc\]](#), page 833.

sound (y)

sound (y, fs)

sound (y, fs, nbits)

Play audio data y at sample rate fs to the default audio device.

The audio signal y can be a vector or a two-column array, representing mono or stereo audio, respectively.

If fs is not given, a default sample rate of 8000 samples per second is used.

The optional argument $nbits$ specifies the bit depth to play to the audio device and defaults to 8 bits.

For more control over audio playback, use the **audioplayer** class.

See also: [\[soundsc\]](#), page 833, [\[record\]](#), page 833.

soundsc (y)

soundsc (y, fs)

soundsc (y, fs, nbits)

soundsc (..., [ymin, ymax])

Scale the audio data y and play it at sample rate fs to the default audio device.

The audio signal y can be a vector or a two-column array, representing mono or stereo audio, respectively.

If fs is not given, a default sample rate of 8000 samples per second is used.

The optional argument $nbits$ specifies the bit depth to play to the audio device and defaults to 8 bits.

By default, y is automatically normalized to the range $[-1, 1]$. If the range $[ymin, ymax]$ is given, then elements of y that fall within the range $ymin \leq y \leq ymax$ are scaled to the range $[-1, 1]$ instead.

For more control over audio playback, use the **audioplayer** class.

See also: [\[sound\]](#), page 833, [\[record\]](#), page 833.

34 Object Oriented Programming

Octave has the ability to create user-defined classes—including the capabilities of operator and function overloading. Classes can protect internal properties so that they may not be altered accidentally which facilitates data encapsulation. In addition, rules can be created to address the issue of class precedence in mixed class operations.

This chapter discusses the means of constructing a user class, how to query and set the properties of a class, and how to overload operators and functions. Throughout this chapter real code examples are given using a class designed for polynomials.

34.1 Creating a Class

This chapter illustrates user-defined classes and object oriented programming through a custom class designed for polynomials. This class was chosen for its simplicity which does not distract unnecessarily from the discussion of the programming features of Octave. Even so, a bit of background on the goals of the polynomial class is necessary before the syntax and techniques of Octave object oriented programming are introduced.

The polynomial class is used to represent polynomials of the form

$$a_0 + a_1x + a_2x^2 + \dots a_nx^n$$

where a_0 , a_1 , etc. are elements of \mathbb{R} . Thus the polynomial can be represented by a vector

```
a = [a0, a1, a2, ..., an];
```

This is a sufficient specification to begin writing the constructor for the polynomial class. All object oriented classes in Octave must be located in a directory that is the name of the class prepended with the ‘@’ symbol. For example, the polynomial class will have all of its methods defined in the @polynomial directory.

The constructor for the class must be the name of the class itself; in this example the constructor resides in the file @polynomial/polynomial.m. Ideally, even when the constructor is called with no arguments it should return a valid object. A constructor for the polynomial class might look like

```
## -*- texinfo -*-
## @deftypefn {} {} polynomial ()
## @deftypefnx {} {} polynomial (@var{a})
## Create a polynomial object representing the polynomial
##
## @example
## a0 + a1 * x + a2 * x^2 + @dots{} + an * x^n
## @end example
##
## @noindent
## from a vector of coefficients [a0 a1 a2 @dots{} an].
## @end deftypefn

function p = polynomial (a)
```

```

if (nargin > 1)
    print_usage ();
endif

if (nargin == 0)
    p.poly = 0;
    p = class (p, "polynomial");
else
    if (isa (a, "polynomial"))
        p = a;
    elseif (isreal (a) && isvector (a))
        p.poly = a(:).' ; # force row vector
        p = class (p, "polynomial");
    else
        error ("@polynomial: A must be a real vector");
    endif
endif

endfunction

```

Note that the return value of the constructor must be the output of the `class` function. The first argument to the `class` function is a structure and the second is the name of the class itself. An example of calling the `class` constructor to create an instance is

```
p = polynomial ([1, 0, 1]);
```

Methods are defined by m-files in the class directory and can have embedded documentation the same as any other m-file. The help for the constructor can be obtained by using the constructor name alone, that is, for the polynomial constructor `help polynomial` will return the help string. Help can be restricted to a particular class by using the class directory name followed by the method. For example, `help @polynomial/polynomial` is another way of displaying the help string for the polynomial constructor. This second means is the only way to obtain help for the overloaded methods and functions of a class.

The same specification mechanism can be used wherever Octave expects a function name. For example `type @polynomial/disp` will print the code of the `disp` method of the `polynomial` class to the screen, and `dbstop @polynomial/disp` will set a breakpoint at the first executable line of the `disp` method of the `polynomial` class.

To check whether a variable belongs to a user class, the `isobject` and `isa` functions can be used. For example:

```

p = polynomial ([1, 0, 1]);
isobject (p)
⇒ 1
isa (p, "polynomial")
⇒ 1

isobject (x)

```

Return true if `x` is a class object.

See also: `[class]`, page 39, `[typeinfo]`, page 39, `[isa]`, page 39, `[ismethod]`, page 837, `[isprop]`, page 395.

The available methods of a class can be displayed with the `methods` function.

```
methods (obj)
methods ("classname")
mtds = methods (...)
```

List the names of the public methods for the object *obj* or the named class *classname*.

obj may be an Octave class object or a Java object. *classname* may be the name of an Octave class or a Java class.

When called with no output arguments, `methods` prints the list of method names to the screen. Otherwise, the output argument *mtds* contains the list in a cell array of strings.

See also: [\[fieldnames\]](#), page 109.

To inquire whether a particular method exists for a user class, the `ismethod` function can be used.

```
ismethod (obj, method)
ismethod (clsname, method)
```

Return true if the string *method* is a valid method of the object *obj* or of the class *clsname*.

See also: [\[isprop\]](#), page 395, [\[isobject\]](#), page 836.

For a polynomial class it makes sense to have a method to compute its roots.

```
function r = roots (p)
    r = roots (fliplr (p.poly));
endfunction
```

We can check for the existence of the `roots`-method by calling:

```
p = polynomial ([1, 0, 1]);
ismethod (p, "roots")
⇒ 1
```

34.2 Class Methods

There are a number of basic class methods that can (and should) be defined to allow the contents of the classes to be queried and set. The most basic of these is the `disp` method. The `disp` method is used by Octave whenever a class should be displayed on the screen. Usually this is the result of an Octave expression that doesn't end with a semicolon. If this method is not defined, then Octave won't print anything when displaying the contents of a class which can be confusing.

An example of a `disp` method for the polynomial class might be

```
function disp (p)

    a = p.poly;
    first = true;
    for i = 1 : length (a);
        if (a(i) != 0)
```

```

if (first)
    first = false;
elseif (a(i) > 0 || isnan (a(i)))
    printf (" +");
endif
if (a(i) < 0)
    printf (" -");
endif
if (i == 1)
    printf (" %.5g", abs (a(i)));
elseif (abs (a(i)) != 1)
    printf (" %.5g *", abs (a(i)));
endif
if (i > 1)
    printf (" X");
endif
if (i > 2)
    printf (" ^ %d", i - 1);
endif
endif
endfor

if (first)
    printf (" 0");
endif
printf ("\n");

endfunction

```

To be consistent with the Octave graphic handle classes, a class should also define the `get` and `set` methods. The `get` method accepts one or two arguments. The first argument is an object of the appropriate class. If no second argument is given then the method should return a structure with all the properties of the class. If the optional second argument is given it should be a property name and the specified property should be retrieved.

```

function val = get (p, prop)

if (nargin < 1 || nargin > 2)
    print_usage ();
endif

if (nargin == 1)
    val.poly = p.poly;
else
    if (! ischar (prop))
        error ("@polynomial/get: PROPERTY must be a string");
    endif

```

```

switch (prop)
  case "poly"
    val = p.poly;
  otherwise
    error ('@polynomial/get: invalid PROPERTY "%s"', prop);
  endswitch
endif

endfunction

```

Similarly, the first argument to the `set` method should be an object and any additional arguments should be property/value pairs.

```

function pout = set (p, varargin)

if (numel (varargin) < 2 || rem (numel (varargin), 2) != 0)
  error ("@polynomial/set: expecting PROPERTY/VALUE pairs");
endif

pout = p;
while (numel (varargin) > 1)
  prop = varargin{1};
  val = varargin{2};
  varargin(1:2) = [];
  if (! ischar (prop) || ! strcmp (prop, "poly"))
    error ("@polynomial/set: invalid PROPERTY for polynomial class");
  elseif (! (isreal (val) && isvector (val)))
    error ("@polynomial/set: VALUE must be a real vector");
  endif

  pout.poly = val(:).' ; # force row vector
endwhile

endfunction

```

Note that Octave does not implement pass by reference; Therefore, to modify an object requires an assignment statement using the return value from the `set` method.

```
p = set (p, "poly", [1, 0, 0, 0, 1]);
```

The `set` method makes use of the `subsasgn` method of the class, and therefore this method must also be defined. The `subsasgn` method is discussed more thoroughly in the next section (see [Section 34.3 \[Indexing Objects\]](#), page 840).

Finally, user classes can be considered to be a special type of a structure, and they can be saved to a file in the same manner as a structure. For example:

```

p = polynomial ([1, 0, 1]);
save userclass.mat p
clear p
load userclass.mat

```

All of the file formats supported by `save` and `load` are supported. In certain circumstances a user class might contain a field that it doesn't make sense to save, or a field that needs to be initialized before it is saved. This can be done with the `saveobj` method of the class.

`b = saveobj (a)`

Method of a class to manipulate an object prior to saving it to a file.

The function `saveobj` is called when the object `a` is saved using the `save` function. An example of the use of `saveobj` might be to remove fields of the object that don't make sense to be saved or it might be used to ensure that certain fields of the object are initialized before the object is saved. For example:

```
function b = saveobj (a)
    b = a;
    if (isempty (b.field))
        b.field = initfield (b);
    endif
endfunction
```

See also: [\[loadobj\]](#), page 840, [\[class\]](#), page 39.

`saveobj` is called just prior to saving the class to a file. Similarly, the `loadobj` method is called just after a class is loaded from a file, and can be used to ensure that any removed fields are reinserted into the user object.

`b = loadobj (a)`

Method of a class to manipulate an object after loading it from a file.

The function `loadobj` is called when the object `a` is loaded using the `load` function. An example of the use of `loadobj` might be to add fields to an object that don't make sense to be saved. For example:

```
function b = loadobj (a)
    b = a;
    b.addmissingfield = addfield (b);
endfunction
```

See also: [\[saveobj\]](#), page 840, [\[class\]](#), page 39.

34.3 Indexing Objects

34.3.1 Defining Indexing And Indexed Assignment

Objects can be indexed with parentheses or braces, either like `obj(idx)` or like `obj{idx}`, or even like `obj(idx).field`. However, it is up to the programmer to decide what this indexing actually means. In the case of the polynomial class `p(n)` might mean either the coefficient of the n -th power of the polynomial, or it might be the evaluation of the polynomial at n . The meaning of this subscripted referencing is determined by the `subsref` method.

`subsref (val, idx)`

Perform the subscripted element selection operation on `val` according to the subscript specified by `idx`.

The subscript *idx* must be a structure array with fields ‘type’ and ‘subs’. Valid values for ‘type’ are “()”, “{}”, and “.”. The ‘subs’ field may be either “:” or a cell array of index values.

The following example shows how to extract the first two columns of a matrix

```
val = magic (3)
⇒ val = [ 8   1   6
          3   5   7
          4   9   2 ]
idx.type = "()";
idx.subs = {":", 1:2};
subsref (val, idx)
⇒ [ 8   1
      3   5
      4   9 ]
```

Note that this is the same as writing `val(:, 1:2)`.

If *idx* is an empty structure array with fields ‘type’ and ‘subs’, return *val*.

See also: [\[subsasgn\]](#), page 842, [\[substruct\]](#), page 112.

For example, this class uses the convention that indexing with “()” evaluates the polynomial and indexing with “{}” returns the *n*-th coefficient (of the *n*-th power). The code for the `subsref` method looks like

```
function r = subsref (p, s)

if (isempty (s))
    error ("@polynomial/subsref: missing index");
endif

switch (s(1).type)

case "()"
    idx = s(1).subs;
    if (numel (idx) != 1)
        error ("@polynomial/subsref: need exactly one index");
    endif
    r = polyval (fliplr (p.poly), idx{1});

case "{}"
    idx = s(1).subs;
    if (numel (idx) != 1)
        error ("@polynomial/subsref: need exactly one index");
    endif

    if (isnumeric (idx{1}))
        r = p.poly(idx{1}+1);
    else
        r = p.poly(idx{1});
```

```

    endif

  case "."
    fld = s.subs;
    if (! strcmp (fld, "poly"))
      error ('@polynomial/subsref: invalid property "%s"', fld);
    endif
    r = p.poly;

  otherwise
    error ("@polynomial/subsref: invalid subscript type");

  endswitch

  if (numel (s) > 1)
    r = subsref (r, s(2:end));
  endif

endfunction

```

The equivalent functionality for subscripted assignments uses the `subsasgn` method.

`subsasgn (val, idx, rhs)`

Perform the subscripted assignment operation according to the subscript specified by `idx`.

The subscript `idx` must be a structure array with fields ‘`type`’ and ‘`subs`’. Valid values for ‘`type`’ are “`()`”, “`{}`”, and “`.`”. The ‘`subs`’ field may be either “`:`” or a cell array of index values.

The following example shows how to set the two first columns of a 3-by-3 matrix to zero.

```

val = magic (3);
idx.type = "()";
idx.subs = {":", 1:2};
subsasgn (val, idx, 0)
⇒ [ 0   0   6
    0   0   7
    0   0   2 ]

```

Note that this is the same as writing `val(:, 1:2) = 0`.

If `idx` is an empty structure array with fields ‘`type`’ and ‘`subs`’, return `rhs`.

See also: [\[subsref\]](#), page 840, [\[substruct\]](#), page 112, [\[optimize_subsasgn_calls\]](#), page 842.

```

val = optimize_subsasgn_calls ()
old_val = optimize_subsasgn_calls (new_val)
optimize_subsasgn_calls (new_val, "local")

```

Query or set the internal flag for `subsasgn` method call optimizations.

If true, Octave will attempt to eliminate the redundant copying when calling the `subsasgn` method of a user-defined class.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[subsasgn\]](#), page 842.

Note that the `subsref` and `subsasgn` methods always receive the whole index chain, while they usually handle only the first element. It is the responsibility of these methods to handle the rest of the chain (if needed), usually by forwarding it again to `subsref` or `subsasgn`.

If you wish to use the `end` keyword in subscripted expressions of an object, then there must be an `end` method defined. For example, the `end` method for the polynomial class might look like

```
function r = end (obj, index_pos, num_indices)

if (num_indices != 1)
    error ("polynomial object may only have one index");
endif

r = length (obj.poly) - 1;

endfunction
```

which is a fairly generic `end` method that has a behavior similar to the `end` keyword for Octave Array classes. An example using the polynomial class is then

```
p = polynomial ([1,2,3,4]);
p{end-1}
⇒ 3
```

Objects can also be used themselves as the index in a subscripted expression and this is controlled by the `subsindex` function.

```
idx = subsindex (obj)
Convert an object to an index vector.
```

When `obj` is a class object defined with a class constructor, then `subsindex` is the overloading method that allows the conversion of this class object to a valid indexing vector. It is important to note that `subsindex` must return a zero-based real integer vector of the class "double". For example, if the class constructor were

```
function obj = myclass (a)
    obj = class (struct ("a", a), "myclass");
endfunction
```

then the `subsindex` function

```
function idx = subsindex (obj)
    idx = double (obj.a) - 1.0;
endfunction
```

could be used as follows

```
a = myclass (1:4);
b = 1:10;
b(a)
⇒ 1 2 3 4
```

See also: [class], page 39, [subsref], page 840, [subsasgn], page 842.

Finally, objects can be used like ranges by providing a `colon` method.

```
r = colon (base, limit)
r = colon (base, increment, limit)
```

Return the result of the colon expression corresponding to `base`, `limit`, and optionally, `increment`.

This function is equivalent to the operator syntax `base : limit` or `base : increment : limit`.

See also: [linspace], page 493.

34.3.2 Indexed Assignment Optimization

Octave's ubiquitous lazily-copied pass-by-value semantics implies a problem for performance of user-defined `subsasgn` methods. Imagine the following call to `subsasgn`

```
ss = substruct ("()", {1});
x = subsasgn (x, ss, 1);
```

where the corresponding method looking like this:

```
function x = subsasgn (x, ss, val)
...
x.myfield (ss.subs{1}) = val;
endfunction
```

The problem is that on entry to the `subsasgn` method, `x` is still referenced from the caller's scope, which means that the method will first need to unshare (copy) `x` and `x.myfield` before performing the assignment. Upon completing the call, unless an error occurs, the result is immediately assigned to `x` in the caller's scope, so that the previous value of `x.myfield` is forgotten. Hence, the Octave language implies a copy of N elements (N being the size of `x.myfield`), where modifying just a single element would actually suffice. In other words, a constant-time operation is degraded to linear-time one. This may be a real problem for user classes that intrinsically store large arrays.

To partially solve the problem Octave uses a special optimization for user-defined `subsasgn` methods coded as m-files. When the method gets called as a result of the built-in assignment syntax (not a direct `subsasgn` call as shown above), i.e., `x(1) = 1`, **AND** if the `subsasgn` method is declared with identical input and output arguments, as in the example above, then Octave will ignore the copy of `x` inside the caller's scope; therefore, any changes made to `x` during the method execution will directly affect the caller's copy as well. This allows, for instance, defining a polynomial class where modifying a single element takes constant time.

It is important to understand the implications that this optimization brings. Since no extra copy of `x` will exist in the caller's scope, it is *solely* the callee's responsibility to not

leave `x` in an invalid state if an error occurs during the execution. Also, if the method partially changes `x` and then errors out, the changes *will* affect `x` in the caller's scope. Deleting or completely replacing `x` inside `subsasgn` will not do anything, however, only indexed assignments matter.

Since this optimization may change the way code works (especially if badly written), a built-in variable `optimize_subsasgn_calls` is provided to control it. It is on by default. Another way to avoid the optimization is to declare `subsasgn` methods with different output and input arguments like this:

```
function y = subsasgn (x, ss, val)
    ...
endfunction
```

34.4 Overloading Objects

34.4.1 Function Overloading

Any Octave function can be overloaded, and this allows an object-specific version of a function to be called as needed. A pertinent example for the polynomial class might be to overload the `polyval` function.

```
function [y, dy] = polyval (p, varargin)

    if (nargout > 1)
        [y, dy] = polyval (fliplr (p.poly), varargin{:});
    else
        y = polyval (fliplr (p.poly), varargin{:});
    endif

endfunction
```

This function just hands off the work to the normal Octave `polyval` function. Another interesting example of an overloaded function for the polynomial class is the `plot` function.

```
function h = plot (p, varargin)

    n = 128;
    rmax = max (abs (roots (p.poly)));
    x = [0 : (n - 1)] / (n - 1) * 2.2 * rmax - 1.1 * rmax;
    if (nargout > 0)
        h = plot (x, polyval (p, x), varargin{:});
    else
        plot (x, polyval (p, x), varargin{:});
    endif

endfunction
```

which allows polynomials to be plotted in the domain near the region of the roots of the polynomial.

Functions that are of particular interest for overloading are the class conversion functions such as `double`. Overloading these functions allows the `cast` function to work with a user

class. It can also aid in the use of a class object with methods and functions from other classes since the object can be transformed to the requisite input form for the new function. An example `double` function for the polynomial class might look like

```
function a = double (p)
  a = p.poly;
endfunction
```

34.4.2 Operator Overloading

The following table shows, for each built-in numerical operation, the corresponding function name to use when providing an overloaded method for a user class.

Operation	Method	Description
<code>a + b</code>	<code>plus (a, b)</code>	Binary addition
<code>a - b</code>	<code>minus (a, b)</code>	Binary subtraction
<code>+a</code>	<code>uplus (a)</code>	Unary addition
<code>-a</code>	<code>uminus (a)</code>	Unary subtraction
<code>a .* b</code>	<code>times (a, b)</code>	Element-wise multiplication
<code>a * b</code>	<code>mtimes (a, b)</code>	Matrix multiplication
<code>a ./ b</code>	<code>rdivide (a, b)</code>	Element-wise right division
<code>a / b</code>	<code>mrdivide (a, b)</code>	Matrix right division
<code>a .\ b</code>	<code>ldivide (a, b)</code>	Element-wise left division
<code>a \ b</code>	<code>mldivide (a, b)</code>	Matrix left division
<code>a .^ b</code>	<code>power (a, b)</code>	Element-wise power
<code>a ^ b</code>	<code>mpower (a, b)</code>	Matrix power
<code>a < b</code>	<code>lt (a, b)</code>	Less than
<code>a <= b</code>	<code>le (a, b)</code>	Less than or equal to
<code>a > b</code>	<code>gt (a, b)</code>	Greater than
<code>a >= b</code>	<code>ge (a, b)</code>	Greater than or equal to
<code>a == b</code>	<code>eq (a, b)</code>	Equal to
<code>a != b</code>	<code>ne (a, b)</code>	Not equal to
<code>a & b</code>	<code>and (a, b)</code>	Logical and
<code>a b</code>	<code>or (a, b)</code>	Logical or
<code>!a</code>	<code>not (a)</code>	Logical not
<code>a'</code>	<code>ctranspose (a)</code>	Complex conjugate transpose
<code>a.'</code>	<code>transpose (a)</code>	Transpose
<code>a:b</code>	<code>colon (a, b)</code>	Two element range
<code>a:b:c</code>	<code>colon (a, b, c)</code>	Three element range
<code>[a, b]</code>	<code>horzcat (a, b)</code>	Horizontal concatenation
<code>[a; b]</code>	<code>vertcat (a, b)</code>	Vertical concatenation
<code>a(s₁, ..., s_n)</code>	<code>subsref (a, s)</code>	Subscripted reference
<code>a(s₁, ..., s_n) = b</code>	<code>subsasgn (a, s, b)</code>	Subscripted assignment
<code>b(a)</code>	<code>subsindex (a)</code>	Convert object to index
<code>disp</code>	<code>disp (a)</code>	Object display

Table 34.1: Available overloaded operators and their corresponding class method

An example `mtimes` method for the polynomial class might look like

```

function p = mtimes (a, b)
  p = polynomial (conv (double (a), double (b)));
endfunction

```

34.4.3 Precedence of Objects

Many functions and operators take two or more arguments and the situation can easily arise where these functions are called with objects of different classes. It is therefore necessary to determine the precedence of which method from which class to call when there are mixed objects given to a function or operator. To do this the **superiorto** and **inferiorto** functions can be used

superiorto (*class_name*, ...)

When called from a class constructor, mark the object currently constructed as having a higher precedence than *class_name*.

More than one such class can be specified in a single call. This function may *only* be called from a class constructor.

See also: [\[inferiorto\]](#), page 847.

inferiorto (*class_name*, ...)

When called from a class constructor, mark the object currently constructed as having a lower precedence than *class_name*.

More than one such class can be specified in a single call. This function may *only* be called from a class constructor.

See also: [\[superiorto\]](#), page 847.

With the polynomial class, consider the case

```
2 * polynomial ([1, 0, 1]);
```

that mixes an object of the class "double" with an object of the class "polynomial". In this case the return type should be "polynomial" and so the **superiorto** function is used in the class constructor. In particular the polynomial class constructor would be modified to

```

## -*- texinfo -*-
## @deftypefn {} {} polynomial ()
## @deftypefnx {} {} polynomial (@var{a})
## Create a polynomial object representing the polynomial
##
## @example
## a0 + a1 * x + a2 * x^2 + @dots{} + an * x^n
## @end example
##
## @noindent
## from a vector of coefficients [a0 a1 a2 @dots{} an].
## @end deftypefn

function p = polynomial (a)

```

```

if (nargin > 1)
    print_usage ();
endif

if (nargin == 0)
    p.poly = [0];
    p = class (p, "polynomial");
else
    if (strcmp (class (a), "polynomial"))
        p = a;
    elseif (isreal (a) && isvector (a))
        p.poly = a(:).' ; # force row vector
        p = class (p, "polynomial");
    else
        error ("@polynomial: A must be a real vector");
    endif
endif

superiorto ("double");

endfunction

```

Note that user classes *always* have higher precedence than built-in Octave types. Thus, marking the polynomial class higher than the "**double**" class is not actually necessary.

When confronted with two objects of equal precedence, Octave will use the method of the object that appears first in the list of arguments.

34.5 Inheritance and Aggregation

Using classes to build new classes is supported by Octave through the use of both inheritance and aggregation.

Class inheritance is provided by Octave using the **class** function in the class constructor. As in the case of the polynomial class, the Octave programmer will create a structure that contains the data fields required by the class, and then call the **class** function to indicate that an object is to be created from the structure. Creating a child of an existing object is done by creating an object of the parent class and providing that object as the third argument of the **class** function.

This is most easily demonstrated by example. Suppose the programmer needs a FIR filter, i.e., a filter with a numerator polynomial but a denominator of 1. In traditional Octave programming this would be performed as follows.

```

>> x = [some data vector];
>> n = [some coefficient vector];
>> y = filter (n, 1, x);

```

The equivalent behavior can be implemented as a class **@FIRfilter**. The constructor for this class is the file **FIRfilter.m** in the class directory **@FIRfilter**.

```

## -- texinfo --
## @deftypefn {} {} FIRfilter ()

```

```

## @deftypefnx {} {} FIRfilter (@var{p})
## Create a FIR filter with polynomial @var{p} as coefficient vector.
## @end deftypefn

function f = FIRfilter (p)

    if (nargin > 1)
        print_usage ();
    endif

    if (nargin == 0)
        p = @polynomial ([1]);
    elseif (! isa (p, "polynomial"))
        error ("@FIRfilter: P must be a polynomial object");
    endif

    f.polynomial = [];
    f = class (f, "FIRfilter", p);

endfunction

```

As before, the leading comments provide documentation for the class constructor. This constructor is very similar to the polynomial class constructor, except that a polynomial object is passed as the third argument to the `class` function, telling Octave that the `FIRfilter` class will be derived from the polynomial class. The FIR filter class itself does not have any data fields, but it must provide a struct to the `class` function. Given that the `@polynomial` constructor will add an element named *polynomial* to the object struct, the `@FIRfilter` just initializes a struct with a dummy field *polynomial* which will later be overwritten.

Note that the sample code always provides for the case in which no arguments are supplied. This is important because Octave will call a constructor with no arguments when loading objects from saved files in order to determine the inheritance structure.

A class may be a child of more than one class (see [\[class\], page 39](#)), and inheritance may be nested. There is no limitation to the number of parents or the level of nesting other than memory or other physical issues.

For the `FIRfilter` class, more control about the object display is desired. Therefore, the `display` method rather than the `disp` method is overloaded (see [Section 34.2 \[Class Methods\], page 837](#)). A simple example might be

```

function display (f)
    printf ("%s.polynomial", inputname (1));
    disp (f.polynomial);
endfunction

```

Note that the `FIRfilter`'s `display` method relies on the `disp` method from the `polynomial` class to actually display the filter coefficients. Furthermore, note that in the `display` method it makes sense to start the method with the line `printf ("%s =", inputname (1))` to be consistent with the rest of Octave which prints the variable name to be displayed followed by the value. In general it is not recommended to overload the `display` function.

display (obj)

Display the contents of the object *obj* prepended by its name.

The Octave interpreter calls the **display** function whenever it needs to present a class on-screen. Typically, this would be a statement which does not end in a semicolon to suppress output. For example:

```
myclass (...)
```

Or:

```
myobj = myclass (...)
```

In general, user-defined classes should overload the **disp** method to avoid the default output:

```
myobj = myclass (...)  
⇒ myobj =
```

```
<class myclass>
```

When overloading the **display** method instead, one has to take care of properly displaying the object's name. This can be done by using the **inputname** function.

See also: [\[disp\]](#), page 251, [\[class\]](#), page 39, [\[subsref\]](#), page 840, [\[subsasgn\]](#), page 842.

Once a constructor and display method exist, it is possible to create an instance of the class. It is also possible to check the class type and examine the underlying structure.

```
octave:1> f = FIRfilter (polynomial ([1 1 1]/3))  
f.polynomial = 0.33333 + 0.33333 * X + 0.33333 * X ^ 2  
octave:2> class (f)  
ans = FIRfilter  
octave:3> isa (f, "FIRfilter")  
ans = 1  
octave:4> isa (f, "polynomial")  
ans = 1  
octave:5> struct (f)  
ans =  
  
scalar structure containing the fields:  
  
polynomial = 0.33333 + 0.33333 * X + 0.33333 * X ^ 2
```

The only thing remaining to make this class usable is a method for processing data. But before that, it is usually desirable to also have a way of changing the data stored in a class. Since the fields in the underlying struct are private by default, it is necessary to provide a mechanism to access the fields. The **subsref** method may be used for both tasks.

```
function r = subsref (f, x)  
  
switch (x.type)  
  
case "()"  
n = f.polynomial;  
r = filter (n.poly, 1, x.subs{1});
```

```

case "."
    fld = x.subs;
    if (!strcmp (fld, "polynomial"))
        error ('@FIRfilter/subsref: invalid property "%s"', fld);
    endif
    r = f.polynomial;

otherwise
    error ("@FIRfilter/subsref: invalid subscript type for FIR filter");■

endswitch

endfunction

```

The "()" case allows us to filter data using the polynomial provided to the constructor.

```

octave:2> f = FIRfilter (polynomial ([1 1 1]/3));
octave:3> x = ones (5,1);
octave:4> y = f(x)
y =

```

$$\begin{aligned} &0.33333 \\ &0.66667 \\ &1.00000 \\ &1.00000 \\ &1.00000 \end{aligned}$$

The " ." case allows us to view the contents of the polynomial field.

```

octave:1> f = FIRfilter (polynomial ([1 1 1]/3));
octave:2> f.polynomial
ans = 0.33333 + 0.33333 * X + 0.33333 * X ^ 2

```

In order to change the contents of the object a `subsasgn` method is needed. For example, the following code makes the polynomial field publicly writable

```

function fout = subsasgn (f, index, val)

switch (index.type)
  case "."
    fld = index.subs;
    if (! strcmp (fld, "polynomial"))
      error ('@FIRfilter/subsasgn: invalid property "%s"', fld);
    endif
    fout = f;
    fout.polynomial = val;

  otherwise
    error ("@FIRfilter/subsasgn: Invalid index type")
  endswitch

endfunction

```

so that

```

octave:1> f = FIRfilter ();
octave:2> f.polynomial = polynomial ([1 2 3])
f.polynomial = 1 + 2 * X + 3 * X ^ 2

```

Defining the FIRfilter class as a child of the polynomial class implies that a FIRfilter object may be used any place that a polynomial object may be used. This is not a normal use of a filter. It may be a more sensible design approach to use aggregation rather than inheritance. In this case, the polynomial is simply a field in the class structure. A class constructor for the aggregation case might be

```

## -*- texinfo -*-
## @deftypefn {} {} FIRfilter ()
## @deftypefnx {} {} FIRfilter (@var{p})
## Create a FIR filter with polynomial @var{p} as coefficient vector.
## @end deftypefn

function f = FIRfilter (p)

if (nargin > 1)
  print_usage ();
endif

if (nargin == 0)
  f.polynomial = @polynomial ([1]);
else
  if (! isa (p, "polynomial"))
    error ("@FIRfilter: P must be a polynomial object");
  endif

  f.polynomial = p;
endif

```

```
f = class (f, "FIRfilter");

endfunction
```

For this example only the constructor needs changing, and all other class methods stay the same.

34.6 classdef Classes

Since version 4.0, Octave has limited support for `classdef` classes. In contrast to the aforementioned classes, called *old style classes* in this section, `classdef` classes can be defined within a single m-file. Other innovations of `classdef` classes are:

- **access rights** for properties and methods,
- **static methods**, i.e., methods that are independent of an object, and
- the distinction between **value and handle classes**.

Several features have to be added in future versions of Octave to be fully compatible to MATLAB. An overview of what is missing can be found at <https://wiki.octave.org/Classdef>.

34.6.1 Creating a classdef Class

A very basic `classdef` value class (see [Section 34.6.5 \[Value Classes vs. Handle Classes\], page 857](#)) is defined by:

```
classdef some_class
  properties
  endproperties

  methods
  endmethods
endclassdef
```

In contrast to old style classes, the `properties-endproperties` block as well as the `methods-endmethods` block can be used to define properties and methods of the class. Because both blocks are empty, they can be omitted in this particular case.

For simplicity, a more advanced implementation of a `classdef` class is shown using the `polynomial` example again (see [Section 34.1 \[Creating a Class\], page 835](#)):

```
classdef polynomial2
  properties
    poly = 0;
  endproperties

  methods
    function p = polynomial2 (a)
      if (nargin > 1)
        print_usage ();
      endif
    endfunction
```

```

if (nargin == 1)
  if (isa (a, "polynomial2"))
    p.poly = a.poly;
  elseif (isreal (a) && isvector (a))
    p.poly = a(:).' ; # force row vector
  else
    error ("polynomial2: A must be a real vector");
  endif
endif
endfunction

function disp (p)
  a = p.poly;
  first = true;
  for i = 1 : length (a);
    if (a(i) != 0)
      if (first)
        first = false;
      elseif (a(i) > 0 || isnan (a(i)))
        printf (" +");
      endif
      if (a(i) < 0)
        printf (" -");
      endif
      if (i == 1)
        printf (" %.5g", abs (a(i)));
      elseif (abs (a(i)) != 1)
        printf (" %.5g *", abs (a(i)));
      endif
      if (i > 1)
        printf (" X");
      endif
      if (i > 2)
        printf (" ^ %d", i - 1);
      endif
    endif
  endfor

  if (first)
    printf (" 0");
  endif
  printf ("\n");
endfunction
endmethods
endclassdef

```

An object of class `polynomial2` is created by calling the class constructor:

```
>> p = polynomial2 ([1, 0, 1])
⇒ p =
1 + x ^ 2
```

34.6.2 Properties

All class properties must be defined within **properties** blocks. The definition of a default value for a property is optional and can be omitted. The default initial value for each class property is `[]`.

A **properties** block can have additional attributes to specify access rights or to define constants:

```
classdef some_class
    properties (Access = mode)
        prop1
    endproperties

    properties (SetAccess = mode, GetAccess = mode)
        prop2
    endproperties

    properties (Constant = true)
        prop3 = pi ()
    endproperties

    properties
        prop4 = 1337
    endproperties
endclassdef
```

where *mode* can be one of:

public	The properties can be accessed from everywhere.
private	The properties can only be accessed from class methods. Subclasses of that class cannot access them.
protected	The properties can only be accessed from class methods and from subclasses of that class.

When creating an object of `some_class`, *prop1* has the default value `[]` and reading from and writing to *prop1* is defined by a single *mode*. For *prop2* the read and write access can be set differently. Finally, *prop3* is a constant property which can only be initialized once within the **properties** block.

By default, in the example *prop4*, properties are not constant and have public read and write access.

34.6.3 Methods

All class methods must be defined within **methods** blocks. An exception to this rule is described at the end of this subsection. Those **methods** blocks can have additional attributes

specifying the access rights or whether the methods are static, i.e., methods that can be called without creating an object of that class.

```
classdef some_class
methods
    function obj = some_class ()
        disp ("New instance created.");
    endfunction

    function disp (obj)
        disp ("Here is some_class.");
    endfunction
endmethods

methods (Access = mode)
    function r = func (obj, r)
        r = 2 * r;
    endfunction
endmethods

methods (Static = true)
    function c = circumference (radius)
        c = 2 * pi () .* radius;
    endfunction
endmethods
endclassdef
```

The constructor of the class is declared in the `methods` block and must have the same name as the class and exactly one output argument which is an object of its class.

It is also possible to overload built-in or inherited methods, like the `disp` function in the example above to tell Octave how objects of `some_class` should be displayed (see [Section 34.2 \[Class Methods\], page 837](#)).

In general, the first argument in a method definition is always the object that it is called from. Class methods can either be called by passing the object as the first argument to that method or by calling the object followed by a dot (".") and the method's name with subsequent arguments:

```
>> obj = some_class ();
New instance created.
>> disp (obj); # both are
>> obj.disp (); # equal
```

In `some_class`, the method `func` is defined within a `methods` block setting the `Access` attribute to `mode`, which is one of:

public The methods can be accessed from everywhere.

private The methods can only be accessed from other class methods. Subclasses of that class cannot access them.

protected

The methods can only be accessed from other class methods and from subclasses of that class.

The default access for methods is **public**.

Finally, the method **circumference** is defined in a static **methods** block and can be used without creating an object of **some_class**. This is useful for methods, that do not depend on any class properties. The class name and the name of the static method, separated by a dot ("."), call this static method. In contrast to non-static methods, the object is not passed as first argument even if called using an object of **some_class**.

```
>> some_class.circumference (3)
⇒ ans = 18.850
>> obj = some_class ();
New instance created.
>> obj.circumference (3)
⇒ ans = 18.850
```

Additionally, class methods can be defined as functions in a folder of the same name as the class prepended with the '@' symbol (see [Section 34.1 \[Creating a Class\]](#), page 835). The main **classdef** file has to be stored in this class folder as well.

34.6.4 Inheritance

Classes can inherit from other classes. In this case all properties and methods of the superclass are inherited to the subclass, considering their access rights. Use this syntax to inherit from **superclass**:

```
classdef subclass < superclass
...
endclassdef
```

34.6.5 Value Classes vs. Handle Classes

There are two intrinsically different types of **classdef** classes, whose major difference is the behavior regarding variable assignment. The first type are **value classes**:

```
classdef value_class
    properties
        prop1
    endproperties

    methods
        function obj = set_prop1 (obj, val)
            obj.prop1 = val;
        endfunction
    endmethods
endclassdef
```

Assigning an object of that class to another variable essentially creates a new object:

```
>> a = value_class ();
>> a.prop1 = 1;
>> b = a;
>> b.prop1 = 2;
>> b.prop1
⇒ ans = 2
>> a.prop1
⇒ ans = 1
```

But that also means that you might have to assign the output of a method that changes properties back to the object manually:

```
>> a = value_class ();
>> a.prop1 = 1;
>> a.set_prop1 (3);
⇒ ans =

<object value_class>

>> ans.prop1
⇒ ans = 3
>> a.prop1
⇒ ans = 1
```

The second type are **handle classes**. Those classes have to be derived from the abstract **handle** class:

```
classdef handle_class < handle
    properties
        prop1
    endproperties

    methods
        function set_prop1 (obj, val)
            obj.prop1 = val;
        endfunction
    endmethods
endclassdef
```

In the following example, the variables **a** and **b** refer to the very same object of class **handle_class**:

```
>> a = handle_class ();
>> a.prop1 = 1;
>> b = a;
>> b.prop1 = 2;
>> b.prop1
⇒ ans = 2
>> a.prop1
⇒ ans = 2
```

Object properties that are modified by a method of an handle class are changed persistently:

```
>> a.set_prop1 (3);
>> a.prop1
⇒ ans = 3
```


35 GUI Development

Octave is principally a batch or command-line language. However, it does offer some features for constructing graphical interfaces that interact with users.

The GUI elements available are I/O dialogs, a progress bar, and UI elements for plot windows. For example, rather than hardcoding a filename for output results a script can open a dialog box and allow the user to choose a file. Similarly, if a calculation is expected to take a long time a script can display a progress bar. The various UI elements can be used to fully customize the plot window with menubars, toolbars, context menus, pushbuttons, sliders, etc.

Several utility functions make it possible to store private data for use with a GUI which will not pollute the user's variable space.

Finally, a program written in Octave might want to have long term storage of preferences or state variables. This can be done with user-defined preferences.

35.1 I/O Dialogs

Simple dialog menus are available for choosing directories or files. They return a string variable which can then be used with any command requiring a filename.

```
dirname = uigetdir ()
dirname = uigetdir (init_path)
dirname = uigetdir (init_path, dialog_name)
```

Open a GUI dialog for selecting a directory.

If *init_path* is not given the current working directory is used.

dialog_name may be used to customize the dialog title.

See also: [\[uigetfile\]](#), page 861, [\[uiputfile\]](#), page 862.

```
[fname, fpath, fltidx] = uigetfile ()
[...] = uigetfile (flt)
[...] = uigetfile (flt, dialog_name)
[...] = uigetfile (flt, dialog_name, default_file)
[...] = uigetfile (...,"Position", [px py])
[...] = uigetfile (...,"MultiSelect", mode)
```

Open a GUI dialog for selecting a file and return the filename *fname*, the path to this file *fpath*, and the filter index *fltidx*.

flt contains a (list of) file filter string(s) in one of the following formats:

`"/path/to/filename.ext"`

If a filename is given then the file extension is extracted and used as filter. In addition, the path is selected as current path and the filename is selected as default file. Example: `uigetfile ("myfun.m")`

A single file extension `"*.ext"`

Example: `uigetfile ("*.ext")`

A 2-column cell array

containing a file extension in the first column and a brief description in the second column. Example: `uigetfile ({".ext", "My Description";".xyz", "XYZ-Format"})`

The filter string can also contain a semicolon separated list of filter extensions. Example: `uigetfile ({"*.gif;*.png;*.jpg", "Supported Picture Formats"})`

A directory name or path name

If the folder name of path name contains a trailing file separator, the contents of that folder will be displayed. If no trailing file separator is present the parent directory is listed. The substring to the right of the rightmost file separator (if any) will be interpreted as a file or directory name and if that file or directory exists it will be highlighted. If the path name or directory name is wholly or partly nonexistent, the current working directory will be displayed. No filter will be active.

`dialog_name` can be used to customize the dialog title.

If `default_file` is given then it will be selected in the GUI dialog. If, in addition, a path is given it is also used as current path.

The screen position of the GUI dialog can be set using the "Position" key and a 2-element vector containing the pixel coordinates. Two or more files can be selected when setting the "MultiSelect" key to "on". In that case `fname` is a cell array containing the files.

See also: [\[uiputfile\]](#), page 862, [\[uigetdir\]](#), page 861.

```
[fname, fpath, fltidx] = uiputfile ()
[fname, fpath, fltidx] = uiputfile (flt)
[fname, fpath, fltidx] = uiputfile (flt, dialog_name)
[fname, fpath, fltidx] = uiputfile (flt, dialog_name, default_file)
```

Open a GUI dialog for selecting a file.

`flt` contains a (list of) file filter string(s) in one of the following formats:

`"/path/to/filename.ext"`

If a filename is given the file extension is extracted and used as filter. In addition the path is selected as current path and the filename is selected as default file. Example: `uiputfile ("myfun.m")`

`"*.ext"` A single file extension. Example: `uiputfile ("*.ext")`

`{"*.ext", "My Description"}`

A 2-column cell array containing the file extension in the 1st column and a brief description in the 2nd column. Example: `uiputfile ({"*.ext", "My Description"; ".xyz", "XYZ-Format"})`

The filter string can also contain a semicolon separated list of filter extensions. Example: `uiputfile ({"*.gif;*.png;*.jpg", "Supported Picture Formats"})`

`dialog_name` can be used to customize the dialog title. If `default_file` is given it is preselected in the GUI dialog. If, in addition, a path is given it is also used as current path.

See also: [\[uigetfile\]](#), page 861, [\[uigetdir\]](#), page 861.

Additionally, there are dialog boxes for displaying help messages, warnings, or errors, and for getting text input from the user.

```
errordlg ()  
errordlg (msg)  
errordlg (msg, title)  
errordlg (msg, title, opt)  
h = errordlg (...)
```

Display an error dialog box with error message *msg* and caption *title*.

The default error message is "This is the default error string." and the default caption is "Error Dialog".

The error message may have multiple lines separated by newline characters ("\\n"), or it may be a cellstr array with one element for each line.

The third optional argument *opt* controls the behavior of the dialog. See [msgbox](#) for details.

The return value *h* is a handle to the figure object used for building the dialog.

Examples:

```
errordlg ("Some fancy error occurred.");  
errordlg ("Some fancy error\nwith two lines.");  
errordlg ({"Some fancy error", "with two lines."});  
errordlg ("Some fancy error occurred.", "Fancy caption");
```

See also: [\[helpdlg\]](#), page 863, [\[warndlg\]](#), page 867, [\[msgbox\]](#), page 865, [\[inputdlg\]](#), page 864, [\[listdlg\]](#), page 864, [\[questdlg\]](#), page 866.

```
helpdlg ()  
helpdlg (msg)  
helpdlg (msg, title)  
h = helpdlg (...)
```

Display a help dialog box with help message *msg* and caption *title*.

The default help message is "This is the default help string." and the default caption is "Help Dialog".

The help message may have multiple lines separated by newline characters ("\\n"), or it may be a cellstr array with one element for each line.

The return value *h* is a handle to the figure object used for building the dialog.

Examples:

```
helpdlg ("Some helpful text for the user.");  
helpdlg ("Some helpful text\nwith two lines.");  
helpdlg ({"Some helpful text", "with two lines."});  
helpdlg ("Some helpful text for the user.", "Fancy caption");
```

See also: [\[errordlg\]](#), page 863, [\[warndlg\]](#), page 867, [\[msgbox\]](#), page 865, [\[inputdlg\]](#), page 864, [\[listdlg\]](#), page 864, [\[questdlg\]](#), page 866.

```
cstr = inputdlg (prompt)
cstr = inputdlg (prompt, title)
cstr = inputdlg (prompt, title, rowscols)
cstr = inputdlg (prompt, title, rowscols, defaults)
cstr = inputdlg (prompt, title, rowscols, defaults, options)
```

Return user input from a multi-textfield dialog box in a cell array of strings, or an empty cell array if the dialog is closed by the Cancel button.

Inputs:

- prompt* A cell array with strings labeling each text field. This input is required.
- title* String to use for the caption of the dialog. The default is "Input Dialog".
- rowscols* Specifies the size of the text fields and can take three forms:
 1. a scalar value which defines the number of rows used for each text field.
 2. a vector which defines the individual number of rows used for each text field.
 3. a matrix which defines the individual number of rows and columns used for each text field. In the matrix each row describes a single text field. The first column specifies the number of input rows to use and the second column specifies the text field width.
- defaults* A list of default values to place in each text field. It must be a cell array of strings with the same size as *prompt*.
- options* Not supported, only for MATLAB compatibility.

Example:

```
prompt = {"Width", "Height", "Depth"};
defaults = {"1.10", "2.20", "3.30"};
rowscols = [1,10; 2,20; 3,30];
dims = inputdlg (prompt, "Enter Box Dimensions", ...
                  rowscols, defaults);
```

See also: [\[errordlg\]](#), page 863, [\[helpdlg\]](#), page 863, [\[listdlg\]](#), page 864, [\[msgbox\]](#), page 865, [\[questdlg\]](#), page 866, [\[warndlg\]](#), page 867.

[sel, ok] = listdlg (key, value, ...)

Return user inputs from a list dialog box in a vector of selection indices (*sel*) and a flag indicating how the user closed the dialog box (*ok*).

The indices in *sel* are 1-based.

The value of *ok* is 1 if the user closed the box with the OK button, otherwise it is 0 and *sel* is empty.

Input arguments are specified in form of *key, value* pairs. The "ListString" argument pair **must** be specified.

Valid *key* and *value* pairs are:

"ListString"

a cell array of strings specifying the items to list in the dialog.

"SelectionMode"
can be either **"Single"** (only one item may be selected at a time) or
"Multiple" (default).

"ListSize"
a two-element vector [*width*, *height*] specifying the size of the list field
in pixels. The default is [160, 300].

"InitialValue"
a vector containing 1-based indices of elements which will be pre-selected
when the list dialog is first displayed. The default is 1 (first item).

"Name" a string to be used as the dialog caption. Default is "".

"PromptString"
a cell array of strings to be displayed above the list of items. Default is {}.

"OKString"
a string used to label the OK button. Default is "OK".

"CancelString"
a string used to label the Cancel button. Default is "Cancel".

Example:

```
my_options = {"An item", "another", "yet another"};
[sel, ok] = listdlg ("ListString", my_options,
                     "SelectionMode", "Multiple");
if (ok == 1)
    disp ("You selected:");
    for i = 1:numel (sel)
        disp (sprintf ("\t%s", my_options{sel(i)}));
    endfor
else
    disp ("You cancelled.");
endif
```

See also: [\[menu\]](#), page 257, [\[errordlg\]](#), page 863, [\[helpdlg\]](#), page 863, [\[inputdlg\]](#),
page 864, [\[msgbox\]](#), page 865, [\[questdlg\]](#), page 866, [\[warndlg\]](#), page 867.

```
h = msgbox (msg)
h = msgbox (msg, title)
h = msgbox (msg, title, icon)
h = msgbox (msg, title, "custom", cdata)
h = msgbox (msg, title, "custom", cdata, colormap)
h = msgbox (... , opt)
```

Display *msg* using a message dialog box.

The message may have multiple lines separated by newline characters ("\\n"), or it
may be a cellstr array with one element for each line.

The optional input *title* (character string) can be used to decorate the dialog caption.

The optional argument `icon` selects a dialog icon. It can be one of "`none`" (default), "`error`", "`help`", "`warn`", or "`custom`". The latter must be followed by an image array `cdata`, and for indexed images the associated colormap.

The final optional argument `opt` controls the behavior of the dialog. If `opt` is a string, it may be one of

`"non-modal"` (default)

The dialog is normal.

`"modal"`

The dialog is displayed "`modal`" which means it prevents users from interacting with any other GUI element until the dialog has been closed.

`"replace"`

If any dialogs already exist with the same title, the most recent is reused and all others are closed. The resulting dialog is set "`non-modal`".

If `opt` is a structure, it must contain fields "`WindowStyle`" and "`Interpreter`":

`"WindowStyle"`

The value must be "`non-modal`", "`modal`", or "`replace`". See above.

`"Interpreter"`

Controls the "`interpreter`" property of the text object used for displaying the message. The value must be "`tex`" (default), "`none`", or "`latex`".

The return value `h` is a handle to the figure object used for building the dialog.

Examples:

```
msgbox ("Some message for the user.");
msgbox ("Some message\nwith two lines.");
msgbox ({"Some message", "with two lines."});
msgbox ("Some message for the user.", "Fancy caption");
```

```
## A message dialog box with error icon
msgbox ("Some message for the user.", "Fancy caption", "error");
```

See also: [\[errordlg\]](#), page 863, [\[helpdlg\]](#), page 863, [\[inputdlg\]](#), page 864, [\[listdlg\]](#), page 864, [\[questdlg\]](#), page 866, [\[warndlg\]](#), page 867.

```
btn = questdlg (msg)
btn = questdlg (msg, title)
btn = questdlg (msg, title, default)
btn = questdlg (msg, title, btn1, btn2, default)
btn = questdlg (msg, title, btn1, btn2, btn3, default)
```

Display `msg` using a question dialog box and return the caption of the activated button.

The message may have multiple lines separated by newline characters ("`\n`"), or it may be a cellstr array with one element for each line.

The optional `title` (character string) can be used to specify the dialog caption. It defaults to "`Question Dialog`".

The dialog may contain two or three buttons which will all close the dialog.

The string *default* identifies the default button, which is activated by pressing the ENTER key. It must match one of the strings given in *btn1*, *btn2*, or *btn3*.

If only *msg* and *title* are specified, three buttons with the default captions "Yes", "No", and "Cancel" are used.

If only two button captions, *btn1* and *btn2*, are specified the dialog will have only these two buttons.

Examples:

```
btn = questdlg ("Close Octave?", "Some fancy title", ...
                "Yes", "No", "No");
if (strcmp (btn, "Yes"))
    exit ();
endif
```

See also: [errordlg], page 863, [helpdlg], page 863, [inputdlg], page 864, [listdlg], page 864, [msgbox], page 865, [warndlg], page 867.

```
warndlg ()
warndlg (msg)
warndlg (msg, title)
warndlg (msg, title, opt)
h = warndlg (...)
```

Display a warning dialog box with warning message *msg* and caption *title*.

The default warning message is "This is the default warning string." and the default caption is "Warning Dialog".

The warning message may have multiple lines separated by newline characters ("\n"), or it may be a cellstr array with one element for each line.

The third optional argument *opt* controls the behavior of the dialog. See `msgbox` for details.

The return value *h* is a handle to the figure object used for building the dialog.

Examples:

```
warndlg ("Some warning text for the user.");
warndlg ("Some warning text\nwith two lines.");
warndlg ({"Some warning text", "with two lines."});
warndlg ("Some warning text for the user.", "Fancy caption");
```

See also: [errordlg], page 863, [helpdlg], page 863, [msgbox], page 865, [inputdlg], page 864, [listdlg], page 864, [questdlg], page 866.

For creating new dialog types, there is a dialog function.

```
h = dialog ()
h = dialog ("property", value, ...)
```

Create an empty modal dialog window to which other uicontrols can be added.

The dialog box is a figure object with properties as recommended for a dialog box.

The default properties differing from a figure are:

```
buttondownfcn
    if isempty (allchild(gcbf)), close (gcbf), endif
```

```

colormap []
color      defaultuicontrolbackgroundcolor
dockcontrols
    off
handlevisibility
    callback
integerhandle
    off
inverthardcopy
    off
menubar   none
numbertitle
    off
paperpositionmode
    auto
resize     off
windowstyle
    modal

```

Multiple property-value pairs may be specified for the dialog object, but they must appear in pairs.

The return value *h* is a graphics handle to the created figure.

Example:

```

## create an empty dialog window titled "Dialog Example"
h = dialog ("name", "Dialog Example");

## create a button (default style)
b = uicontrol (h, "string", "OK",
                "position", [10 10 150 40],
                "callback", "delete (gcf)");

## wait for dialog to resume or close
uiwait (h);

```

See also: [errordlg], page 863, [msgbox], page 865, [questdlg], page 866, [warndlg], page 867, [figure], page 372, [uiwait], page 880.

35.2 Progress Bar

```

h = waitbar (frac)
h = waitbar (frac, msg)
h = waitbar (... , "createcancelbtn", fcn, ...)
h = waitbar (... , prop, val, ...)

```

```
waitbar (frac)
waitbar (frac, h)
waitbar (frac, h, msg)
```

Return a handle *h* to a new progress indicator ("waitbar") object.

The waitbar is filled to fraction *frac* which must be in the range [0, 1].

The optional message *msg* is centered and displayed above the waitbar.

A cancel button can be added to the bottom of the waitbar using the "createcancelbtn" property of waitbar figures. The action to be executed when the user presses the button is specified using a string or function handle *fcn*.

The appearance of the waitbar figure window can be configured by passing *prop/val* pairs to the function.

When called with a single input the current waitbar, if it exists, is updated to the new value *frac*. If there are multiple outstanding waitbars they can be updated individually by passing the handle *h* of the specific waitbar to modify.

See also: [\[delete\]](#), page 377.

35.3 UI Elements

The ui* series of functions work best with the qt graphics toolkit, although some functionality is available with the fltk toolkit. There is no support for the gnuplot toolkit.

```
uipanel (property, value, ...)
uipanel (parent, property, value, ...)
 = uipanel (...)
```

Create a uipanel object.

uipanels are used as containers to group other uicontrol objects.

If *parent* is omitted then a uipanel for the current figure is created. If no figure is available, a new figure is created first.

If *parent* is given then a uipanel relative to *parent* is created.

Any provided property value pairs will override the default values of the created uipanel object.

Uipanel properties are documented at [Section 15.3.3.13 \[Uipanel Properties\]](#), page 441.

The optional return value *hui* is a graphics handle to the created uipanel object.

Examples:

```
## create figure and panel on it
f = figure;
p = uipanel ("title", "Panel Title", "position", [.25 .25 .5 .5]);

## add two buttons to the panel
b1 = uicontrol ("parent", p, "string", "A Button", ...
                "position", [18 10 150 36]);
b2 = uicontrol ("parent", p, "string", "Another Button", ...
                "position", [18 60 150 36]);
```

See also: [\[figure\]](#), page 372, [\[uicontrol\]](#), page 870.

```
hui = uibuttongroup (property, value, ...)  

hui = uibuttongroup (parent, property, value, ...)  

uibuttongroup (h)
```

Create a uibuttongroup object and return a handle to it.

A uibuttongroup is used to group uicontrol objects.

If *parent* is omitted then a uibuttongroup for the current figure is created. If no figure is available, a new figure is created first.

If *parent* is given then a uibuttongroup relative to *parent* is created.

Any provided property value pairs will override the default values of the created uibuttongroup object.

Properties of uibuttongroup objects are documented at [Section 15.3.3.11 \[Uibuttongroup Properties\]](#), page 437.

Examples:

```
## Create figure and panel on it
f = figure;
## Create a button group
gp = uibuttongroup (f, "Position", [ 0 0.5 1 1])
## Create a buttons in the group
b1 = uiicontrol (gp, "style", "radiobutton", ...
                 "string", "Choice 1", ...
                 "Position", [ 10 150 100 50 ]);
b2 = uiicontrol (gp, "style", "radiobutton", ...
                 "string", "Choice 2", ...
                 "Position", [ 10 50 100 30 ]);
## Create a button not in the group
b3 = uiicontrol (f, "style", "radiobutton", ...
                 "string", "Not in the group", ...
                 "Position", [ 10 50 100 50 ]);
```

When called with a single argument *h* which is a handle to an existing uibuttongroup object, switch the focus to the specified uibuttongroup. This functionality is not currently implemented.

See also: [\[figure\]](#), page 372, [\[uipanel\]](#), page 869.

```
hui = uiicontrol (property, value, ...)  

hui = uiicontrol (parent, property, value, ...)  

uiicontrol (h)
```

Create a uiicontrol object and return a handle to it.

A uiicontrol object is used to create simple interactive controls such as push buttons, checkboxes, edit and list controls.

If *parent* is omitted then a uiicontrol for the current figure is created. If no figure is available, a new figure is created first.

If *parent* is given then a uiicontrol relative to *parent* is created.

Any provided property value pairs will override the default values of the created uiicontrol object.

Properties of uicontrol objects are documented at [Section 15.3.3.14 \[Uicontrol Properties\]](#), page 444.

The type of uicontrol created is specified by the *style* property. If no style property is provided, a push button will be created.

Valid styles for uicontrol are:

"checkbox"	Create a checkbox control that allows user on/off selection.
"edit"	Create an edit control that allows user input of single or multiple lines of text.
"listbox"	Create a listbox control that displays a list of items and allows user selection of single or multiple items.
"popupmenu"	Create a popupmenu control that displays a list of options that can be selected when the user clicks on the control.
"pushbutton"	Create a push button control that allows user to press to cause an action.
"radiobutton"	Create a radio button control intended to be used for mutually exclusive input in a group of radiobutton controls.
"slider"	Create a slider control that allows user selection from a range of values by sliding knob on the control.
"text"	Create a static text control to display single or multiple lines of text.
"togglebutton"	Create a toggle button control that appears like a push button but allows the user to select between two states.

Examples:

```
## Create figure and panel on it
f = figure;
## Create a button (default style)
b1 = uicontrol (f, "string", "A Button", ...
                "position", [10 10 150 40]);
## Create an edit control
e1 = uicontrol (f, "style", "edit", "string", "editable text", ...
                "position", [10 60 300 40]);
## Create a checkbox
c1 = uicontrol (f, "style", "checkbox", "string", "a checkbox", ...
                "position", [10 120 150 40]);
```

When called with a single argument *h* which is a handle to an existing uicontrol object, switch the keyboard focus to the specified uicontrol. As a result, the uicontrol object will receive keyboard events that can be processed using the "keypressfcn" callback.

See also: [\[figure\]](#), page 372, [\[uipanel\]](#), page 869.

```
hui = uitable (property, value, ...)  

hui = uitable (parent, property, value, ...)
```

Create a uitable object and return a handle to it.

A uitable object is used to show tables of data in a figure window.

If *parent* is omitted then a uitable for the current figure is created. If no figure is available, a new figure is created first.

If *parent* is given then a uitable relative to *parent* is created.

Any provided property value pairs will override the default values of the created uitable object.

Properties of uitable objects are documented at [Section 15.3.3.15 \[Uitable Properties\]](#), [page 446](#).

Examples:

```
## Create figure and place a table on it
f = figure ();
m = magic (8);
t = uitable (f, "Data", m, "ColumnWidth", { 40 });

## Create a table with labeled rows and columns
f = figure ();
d = reshape (1:9, [3, 3]);
row_names = { "Row1", "Row2", "Row3" };
col_names = { "Col1", "Col2", "Col3" };
t = uitable (f, "Data", d, ...
             "RowName", row_names, "ColumnName", col_names);

p = get (t, "Position");
e = get (t, "Extent");
p(3:4) = e(3:4);
set (t, "Position", p);

## Long demo with callbacks
function uitable_demo ()
    f = figure ("Name", "uitable Demo", "Menu", "none", ...
                "Position", [10 10 1000 680]);

    ## A basic example
    d = { "char" , "A string";
          "double" , 12.3456789;
          "complex", 1+2i;
          "bool"   , true;
          "single" , single (12.3456789);
          "int8"   , int8 (-128);
          "uint8"  , uint8 (128);
          "int16"  , int16 (-32768);
          "uint16" , uint16 (32768);
```

```
"int32" , int32 (-2147483648);
"uint32" , uint32 (2147483648);
"int64" , int64 (-2147483649);
"uint64" , uint64 (2147843649)};
```

```
popup_options = {"A", "B", "C", "D", "E"};
```

```
columnformat_options = { [] , "char", "pop-up", "numeric", ...
                         "short", "short e", "short eng", ...
                         "short g", "long", "long e", ...
                         "long eng", "long g", "bank", "+", ...
                         "rat", "logical"};
```

```
columnformat_values = columnformat_options;
columnformat_values{1} = "";
columnformat_values{3} = popup_options;
```

```
default_data = repmat (d(:,2), 1, columns (columnformat_options));
b_add = uicontrol (f, "Position", [285 630 600 50], ...
                   "UserData", [rows(d), 1], ...
                   "Style", "pushbutton", ...
                   "String", "Set data at selected point to selected datatype");■
```

```
l_type_table = uicontrol (f, "Position", [ 0 603 120 25 ], ...
                           "String", "Datatype Table:", ...
                           "Style", "text");
t_type_table = uitable (f, "Position", [ 0 530 1000 70 ], ...
                        "Data", transpose (d(:, 2)), ...
                        "ColumnName", transpose (d(:, 1)), ...
                        "RowName", "Value", ...
                        "CellSelectionCallback", ...
                        @(x, y) set (b_add, "UserData", y.Indices ));
```

```
l_point_table = uicontrol (f, "Position", [ 0 640 60 25 ], ...
                            "String", "Point:", ...
                            "Style", "text");
t_point_table = uitable (f, "Position", [ 80 630 160 42 ], ...
                        "RowName", [], ...
                        "ColumnName", {"x", "y"}, ...
                        "Data", [ 1, 1 ], ...
                        "ColumnEditable", true);
```

```
l_editable_table = uicontrol (f, "Position", [ 0 502 200 25 ], ...
                               "Style", "text", ...
                               "String", "Set Data Columns Editable:");
t_editable_table = ...
uitable (f, "Position", [ 0 434 1000 65 ], ...
         "Data", repmat (false, 1, columns (default_data)), ...■
```

```

    "ColumnEditable", true);

l_format_table = uicontrol (f, "Position", [ 0 406 200 25 ], ...
    "Style", "text", ...
    "String", "Set Data Column Format:");
t_format_table = ...
    uitable (f, "Position", [ 0 338 1000 65 ], ...
        "Data", columnformat_options, ...
        "ColumnEditable", true, ...
        "ColumnFormat", arrayfun (@(x) {columnformat_options}, ...
            1:columns (columnformat_options)));■

l_data_table = uicontrol (f, "Style", "text", ...
    "String", "Data:", ...
    "Position", [ 0 310 60 25 ]);
t_data_table = uitable (f, "Position", [ 0 15 1000 290 ], ...
    "Data", default_data, ...
    "ColumnFormat", columnformat_values);

set (t_format_table, ...
    "CellEditCallback", ...
    @(x, y) update_column_format (y.NewData, y.Indices, ...
        t_data_table, popup_options));
set (t_point_table, "CellEditCallback", ...
    @(x, y) validate_point_table (x, y, t_data_table));
set (t_editable_table, "CellEditCallback", ...
    @(x,y) set (t_data_table, ...
        "ColumnEditable", get (t_editable_table, "Data")));■
set (b_add, ...
    "Callback", @(x, y) update_data (b_add, t_point_table, ...
        t_type_table, t_data_table));■
set (t_data_table, "CellSelectionCallback", ...
    @(x, y) update_point_table (y.Indices, t_point_table));
endfunction

function validate_point_table (h, dat, t_data_table)
if (! (dat.NewData > 0 && ...
    dat.NewData < size (get (t_data_table, "Data")), dat.Indices(1, 1)) + 1))■

    d = get (h, "Data");
    d(dat.Indices) = 1;
    set (h, "Data", d);
endif
endfunction

function update_column_format (format, indices, t_data_table, ...
    popup_options)

```

```

cf = get (t_data_table, "ColumnFormat");
if (strcmp (format, "[]"))
    format = "";
elseif (strcmp (format, "pop-up"))
    format = popup_options;
endif
cf{indices(1,2)} = format;
set (t_data_table, "ColumnFormat", cf);
endfunction

function update_point_table (indices, t_point_table)
if (isempty (indices))
    indices = [1, 1];
endif
set (t_point_table, "Data", indices(1,:));
endfunction

function update_data (b_add, t_point_table, t_type_table, ...
                     t_data_table)
indices = get (b_add, "UserData");
if (isempty (indices))
    indices = [1, 1];
endif
d = get (t_data_table, "Data");
t_type_table_data = get (t_type_table, "Data");
p = get (t_point_table, "Data");
d(p(1,2), p(1,1)) = t_type_table_data(indices(1,2));
set (t_data_table, "Data", d);
endfunction

```

See also: [\[figure\]](#), page 372, [\[uicontrol\]](#), page 870.

hui = uimenu (*property*, *value*, ...)

hui = uimenu (*h*, *property*, *value*, ...)

Create a uimenu object and return a handle to it.

If *h* is omitted then a top-level menu for the current figure is created. If *h* is given then a submenu relative to *h* is created.

uimenu objects have the following specific properties:

"accelerator"

A string containing the key combination together with CTRL to execute this menu entry (e.g., "x" for CTRL+x).

"callback"

Is the function called when this menu entry is executed. It can be either a function string (e.g., "myfun"), a function handle (e.g., @myfun) or a cell array containing the function handle and arguments for the callback function (e.g., {@myfun, arg1, arg2}).

"checked"

Can be set "on" or "off". Sets a mark at this menu entry.

"enable" Can be set "on" or "off". If disabled the menu entry cannot be selected and it is grayed out.

"foregroundcolor"

A color value setting the text color for this menu entry.

"label" A string containing the label for this menu entry. A "&"-symbol can be used to mark the "accelerator" character (e.g., "E&xit")

"position"

An scalar value containing the relative menu position. The entry with the lowest value is at the first position starting from left or top.

"separator"

Can be set "on" or "off". If enabled it draws a separator line above the current position. It is ignored for top level entries.

The full list of properties is documented at [Section 15.3.3.10 \[Uimenu Properties\], page 436](#).

Examples:

```
f = uimenu ("label", "&File", "accelerator", "f");
e = uimenu ("label", "&Edit", "accelerator", "e");
uimenu (f, "label", "Close", "accelerator", "q", ...
    "callback", "close (gcf)");
uimenu (e, "label", "Toggle &Grid", "accelerator", "g", ...
    "callback", "grid (gca)");
```

See also: [\[figure\]](#), [page 372](#).

```
hui = uicontextmenu (property, value, ...)
hui = uicontextmenu (h, property, value, ...)
```

Create a uicontextmenu object and return a handle to it.

If *h* is omitted then a uicontextmenu for the current figure is created. If no figure is available, a new figure is created first.

If *h* is given then a uicontextmenu relative to *h* is created.

Any provided property value pairs will override the default values of the created uicontextmenu object.

Uicontextmenu properties are documented at [Section 15.3.3.12 \[Uicontextmenu Properties\], page 440](#).

Examples:

```

## create figure and uicontextmenu
f = figure ();
c = uicontextmenu (f);

## create menus in the context menu
m1 = uimenu ("parent", c, "label", "Menu item 1", ...
             "callback", "disp('menu item 1')");
m2 = uimenu ("parent", c, "label", "Menu item 2", ...
             "callback", "disp('menu item 2')");

## set the context menu for the figure
set (f, "uicontextmenu", c);

```

See also: [figure], page 372, [uimenu], page 875.

uitoolbar (*property, value, ...*)
uitoolbar (*parent, property, value, ...*)
hui = **uitoolbar** (...)

Create a uitoolbar object. A uitoolbar displays uitoggletool and uipushtool buttons.

If *parent* is omitted then a uitoolbar for the current figure is created. If no figure is available, a new figure is created first.

If *parent* is given then a uitoolbar relative to *parent* is created.

Any provided property value pairs will override the default values of the created uitoolbar object.

Uitoolbar properties are documented at [Section 15.3.3.16 \[Uitoolbar Properties\], page 449](#).

The optional return value *hui* is a graphics handle to the created uitoolbar object.

Examples:

```

% create figure without a default toolbar
f = figure ("toolbar", "none");
% create empty toolbar
t = uitoolbar (f);

```

See also: [figure], page 372, [uitoggletool], page 878, [uipushtool], page 877.

uipushtool (*property, value, ...*)
uipushtool (*parent, property, value, ...*)
hui = **uipushtool** (...)

Create a uipushtool object.

uipushtools are buttons that appear on a figure toolbar. The button is created with a border that is shown when the user hovers over the button. An image can be set using the cdata property.

If *parent* is omitted then a uipushtool for the current figure is created. If no figure is available, a new figure is created first. If a figure is available, but does not contain a uitoolbar, a uitoolbar will be created.

If *parent* is given then a uipushtool is created on the *parent* uitoolbar.

Any provided property value pairs will override the default values of the created uipushtool object.

Uipushtool properties are documented at [Section 15.3.3.17 \[Uipushtool Properties\], page 450](#).

The optional return value *hui* is a graphics handle to the created uipushtool object.

Examples:

```
% create figure without a default toolbar
f = figure ("toolbar", "none");
% create empty toolbar
t = uitablebar (f);
% create a 19x19x3 black square
img=zeros(19,19,3);
% add pushtool button to toolbar
b = uipushtool (t, "cdata", img);
```

See also: [\[figure\]](#), page 372, [\[uitablebar\]](#), page 877, [\[uitoggletool\]](#), page 878.

uitoggletool (*property*, *value*, ...)
uitoggletool (*parent*, *property*, *value*, ...)
hui = **uitoggletool** (...)

Create a uitoggletool object.

uitoggletool are togglebuttons that appear on a figure toolbar. The button is created with a border that is shown when the user hovers over the button. An image can be set using the cdata property.

If *parent* is omitted then a uitoggletool for the current figure is created. If no figure is available, a new figure is created first. If a figure is available, but does not contain a uitablebar, a uitablebar will be created.

If *parent* is given then a uitoggletool is created on the *parent* uitablebar.

Any provided property value pairs will override the default values of the created uitoggletool object.

Uitoggletool properties are documented at [Section 15.3.3.18 \[Uitoggletool Properties\], page 452](#).

The optional return value *hui* is a graphics handle to the created uitoggletool object.

Examples:

```
% create figure without a default toolbar
f = figure ("toolbar", "none");
% create empty toolbar
t = uitablebar (f);
% create a 19x19x3 black square
img=zeros(19,19,3);
% add uitoggletool button to toolbar
b = uitoggletool (t, "cdata", img);
```

See also: [\[figure\]](#), page 372, [\[uitablebar\]](#), page 877, [\[uipushtool\]](#), page 877.

35.4 GUI Utility Functions

These functions do not implement a GUI element but are useful when developing programs that do. The functions `uiwait`, `uiresume`, and `waitfor` are only available with the `qt` or `fltk` toolkits.

```
data = guidata (h)
guidata (h, data)
```

Query or set user-custom GUI data.

The GUI data is stored in the figure handle `h`. If `h` is not a figure handle then it's parent figure will be used for storage.

`data` must be a single object which means it is usually preferable for it to be a data container such as a cell array or struct so that additional data items can be added easily.

See also: [\[getappdata\]](#), page 459, [\[setappdata\]](#), page 459, [\[get\]](#), page 402, [\[set\]](#), page 402, [\[getpref\]](#), page 881, [\[setpref\]](#), page 882.

```
hdata = guihandles (h)
hdata = guihandles
```

Return a structure of object handles for the figure associated with handle `h`.

If no handle is specified the current figure returned by `gcf` is used.

The fieldname for each entry of `hdata` is taken from the "`tag`" property of the graphic object. If the tag is empty then the handle is not returned. If there are multiple graphic objects with the same tag then the entry in `hdata` will be a vector of handles. `guihandles` includes all possible handles, including those for which "`HandleVisibility`" is "`off`".

See also: [\[guidata\]](#), page 879, [\[findobj\]](#), page 453, [\[findall\]](#), page 454, [\[allchild\]](#), page 404.

```
have_window_system ()
```

Return true if a window system is available (X11, Windows, or Apple OS X) and false otherwise.

See also: [\[isguirunning\]](#), page 879.

```
isguirunning ()
```

Return true if Octave is running in GUI mode and false otherwise.

See also: [\[have_window_system\]](#), page 879.

```
movegui
movegui (h)
movegui (pos)
movegui (h, pos)
movegui (h, event)
movegui (h, event, pos)
```

Move a figure specified by figure handle `h` to a position on the screen defined by `pos`.

`h` is a figure handle, or a handle to a graphics object. In the latter case, its parent figure will be used. If unspecified, `h` will be set to the handle of the relevant figure

if a callback is being executed (`gcbf`), otherwise it will be set to the handle of the current figure (`gcf`).

`pos` is either a two-value numeric vector or a string. If `pos` is numeric then it must be of the form `[h, v]` specifying the horizontal and vertical offsets of the figure with respect to the screen. A positive value indicates the offset between the left (or bottom for the vertical component) of the screen, and the left (or bottom) of the figure. A negative value indicates the offset between the right (or top) of the screen and the right (or top) of the figure.

Possible values for `pos` as a string are

<code>north</code>	Top center of the screen.
<code>south</code>	Bottom center of the screen.
<code>east</code>	Right center of the screen.
<code>west</code>	Left center of the screen.
<code>northeast</code>	Top right of the screen.
<code>northwest</code>	Top left of the screen.
<code>southeast</code>	Bottom right of the screen.
<code>southwest</code>	Bottom left of the screen.
<code>center</code>	Center of the screen.
<code>onscreen</code> (<code>default</code>)	The figure will be minimally moved to be entirely visible on the screen, with a 30 pixel extra padding from the sides of the screen. This is the default value if none is provided.

`event` contains event data that will be ignored. This construct facilitates a call to `movegui` from a callback.

`openvar (name)`

Open the variable `name` in the graphical Variable Editor.

`uiwait`

`uiwait (h)`

`uiwait (h, timeout)`

Suspend program execution until the figure with handle `h` is deleted or `uiresume` is called.

When no figure handle is specified this function uses the current figure. If the figure handle is invalid or there is no current figure, this functions returns immediately.

When specified, `timeout` defines the number of seconds to wait for the figure deletion or the `uiresume` call. The timeout value must be at least 1. If a smaller value is

specified, a warning is issued and a timeout value of 1 is used instead. If a non-integer value is specified, it is truncated towards 0. If *timeout* is not specified, the program execution is suspended indefinitely.

See also: [\[uiresume\]](#), page 881, [\[waitfor\]](#), page 881.

uiresume (h)

Resume program execution suspended with [uiwait](#).

The handle *h* must be the same as the one specified in [uiwait](#). If the handle is invalid or there is no [uiwait](#) call pending for the figure with handle *h*, this function does nothing.

See also: [\[uiwait\]](#), page 880.

```
waitfor (h)
waitfor (h, prop)
waitfor (h, prop, value)
waitfor (..., "timeout", timeout)
```

Suspend the execution of the current program until a condition is satisfied on the graphics handle *h*.

While the program is suspended graphics events are still processed normally, allowing callbacks to modify the state of graphics objects. This function is reentrant and can be called from a callback, while another [waitfor](#) call is pending at the top-level.

In the first form, program execution is suspended until the graphics object *h* is destroyed. If the graphics handle is invalid or if *h* is the root figure handle and no property *prop* was provided, the function returns immediately.

In the second form, execution is suspended until the graphics object is destroyed or the property named *prop* is modified. If the graphics handle is invalid or the property does not exist, the function returns immediately.

In the third form, execution is suspended until the graphics object is destroyed or the property named *prop* is set to *value*. The function [isequal](#) is used to compare property values. If the graphics handle is invalid, the property does not exist or the property is already set to *value*, the function returns immediately.

An optional timeout can be specified using the property "timeout". This timeout value is the number of seconds to wait for the condition to be true. *timeout* must be at least 1. If a smaller value is specified, a warning is issued and a value of 1 is used instead. If the timeout value is not an integer, it is truncated towards 0.

To define a condition on a property named "timeout", use the string '\timeout' instead.

In all cases, typing CTRL-C stops program execution immediately.

See also: [\[waitforbuttonpress\]](#), page 392, [\[isequal\]](#), page 153.

35.5 User-Defined Preferences

```
val = getpref ("group", "pref")
val = getpref ("group", "pref", default)
```

```
{val1, val2, ...} = getpref ("group", {"pref1", "pref2", ...})
prefstruct = getpref ("group")
prefstruct = getpref ()
```

Return the preference value corresponding to the named preference *pref* in the preference group *group*.

The named preference group must be a string.

If *pref* does not exist in *group* and *default* is specified, create the preference with value *default* and return *default*.

The preference *pref* may be a string or cell array of strings. If it is a cell array of strings then a cell array of preferences is returned.

The corresponding default value *default* may be any Octave value, .e.g., double, struct, cell array, object, etc. Or, if *pref* is a cell array of strings then *default* must be a cell array of values with the same size as *pref*.

If neither *pref* nor *default* are specified, return a structure of preferences for the preference group *group*.

If no arguments are specified, return a structure containing all groups of preferences and their values.

See also: [addpref], page 882, [setpref], page 882, [ispref], page 883, [rmpref], page 883.

```
setpref ("group", "pref", val)
setpref ("group", {"pref1", "pref2", ...}, {val1, val2, ...})
```

Set the preference *pref* to the given *val* in the named preference group *group*.

The named preference group must be a string.

The preference *pref* may be a string or a cell array of strings.

The corresponding value *val* may be any Octave value, .e.g., double, struct, cell array, object, etc. Or, if *pref* is a cell array of strings then *val* must be a cell array of values with the same size as *pref*.

If the named preference or group does not exist, it is added.

See also: [addpref], page 882, [getpref], page 881, [ispref], page 883, [rmpref], page 883.

```
addpref ("group", "pref", val)
addpref ("group", {"pref1", "pref2", ...}, {val1, val2, ...})
```

Add the preference *pref* and associated value *val* to the named preference group *group*.

The named preference group must be a string.

The preference *pref* may be a string or a cell array of strings. An error will be issued if the preference already exists.

The corresponding value *val* may be any Octave value, .e.g., double, struct, cell array, object, etc. Or, if *pref* is a cell array of strings then *val* must be a cell array of values with the same size as *pref*.

See also: [setpref], page 882, [getpref], page 881, [ispref], page 883, [rmpref], page 883.

```
rmpref ("group", "pref")
rmpref ("group", {"pref1", "pref2", ...})
rmpref ("group")
```

Remove the named preference *pref* from the preference group *group*.

The named preference group must be a string.

The preference *pref* may be a string or cell array of strings.

If *pref* is not specified, remove the preference group *group*.

It is an error to remove a nonexistent preference or group.

See also: [addpref], page 882, [ispref], page 883, [setpref], page 882, [getpref], page 881.

```
ispref ("group", "pref")
ispref ("group", {"pref1", "pref2", ...})
ispref ("group")
```

Return true if the named preference *pref* exists in the preference group *group*.

The named preference group must be a string.

The preference *pref* may be a string or a cell array of strings.

If *pref* is not specified, return true if the preference group *group* exists.

See also: [getpref], page 881, [addpref], page 882, [setpref], page 882, [rmpref], page 883.

```
prefdir
prefdir (1)
dir = prefdir
```

Return the directory that holds the preferences for Octave.

Examples:

Display the preferences directory

```
prefdir
```

Change to the preferences folder

```
cd (prefdir)
```

If called with an argument, the preferences directory is created if it doesn't already exist.

See also: [getpref], page 881, [setpref], page 882, [addpref], page 882, [rmpref], page 883, [ispref], page 883.

```
preferences
```

Display the GUI preferences dialog window for Octave.

36 System Utilities

This chapter describes the functions that are available to allow you to get information about what is happening outside of Octave, while it is still running, and use this information in your program. For example, you can get information about environment variables, the current time, and even start other programs from the Octave prompt.

36.1 Timing Utilities

Octave's core set of functions for manipulating time values are patterned after the corresponding functions from the standard C library. Several of these functions use a data structure for time that includes the following elements:

<code>usec</code>	Microseconds after the second (0-999999).
<code>sec</code>	Seconds after the minute (0-60). This number can be 60 to account for leap seconds.
<code>min</code>	Minutes after the hour (0-59).
<code>hour</code>	Hours since midnight (0-23).
<code>mday</code>	Day of the month (1-31).
<code>mon</code>	Months since January (0-11).
<code>year</code>	Years since 1900.
<code>wday</code>	Days since Sunday (0-6).
<code>yday</code>	Days since January 1 (0-365).
<code>isdst</code>	Daylight saving time flag.
<code>gmtoff</code>	Seconds offset from UTC.
<code>zone</code>	Time zone.

In the descriptions of the following functions, this structure is referred to as a *tm_struct*.

`seconds = time ()`

Return the current time as the number of seconds since the epoch.

The epoch is referenced to 00:00:00 UTC (Coordinated Universal Time) 1 Jan 1970. For example, on Monday February 17, 1997 at 07:15:06 UTC, the value returned by `time` was 856163706.

See also: [strftime], page 887, [strptime], page 889, [localtime], page 886, [gmtime], page 886, [mktime], page 887, [now], page 885, [date], page 890, [clock], page 889, [datenum], page 892, [datestr], page 893, [datevec], page 895, [calendar], page 895, [weekday], page 896.

`t = now ()`

Return the current local date/time as a serial day number (see `datenum`).

The integral part, `floor (now)` corresponds to the number of days between today and Jan 1, 0000.

The fractional part, `rem (now, 1)` corresponds to the current time.

See also: [clock], page 889, [date], page 890, [datenum], page 892.

`ctime (t)`

Convert a value returned from `time` (or any other non-negative integer), to the local time and return a string of the same form as `asctime`.

The function `ctime (time)` is equivalent to `asctime (localtime (time))`. For example:

```
ctime (time ())
⇒ "Mon Feb 17 01:15:06 1997\n"
```

See also: `[asctime]`, page 887, `[time]`, page 885, `[localtime]`, page 886.

`tm_struct = gmtime (t)`

Given a value returned from `time`, or any non-negative integer, return a time structure corresponding to UTC (Coordinated Universal Time).

For example:

```
gmtime (time ())
⇒ {
    usec = 0
    sec = 6
    min = 15
    hour = 7
    mday = 17
    mon = 1
    year = 97
    wday = 1
    yday = 47
    isdst = 0
    gmtoff = 0
    zone = GMT
}
```

See also: `[strftime]`, page 887, `[strptime]`, page 889, `[localtime]`, page 886, `[mktime]`, page 887, `[time]`, page 885, `[now]`, page 885, `[date]`, page 890, `[clock]`, page 889, `[datenum]`, page 892, `[datestr]`, page 893, `[datevec]`, page 895, `[calendar]`, page 895, `[weekday]`, page 896.

`tm_struct = localtime (t)`

Given a value returned from `time`, or any non-negative integer, return a time structure corresponding to the local time zone.

```

localtime (time ())
    ⇒ {
        usec = 0
        sec = 6
        min = 15
        hour = 1
        mday = 17
        mon = 1
        year = 97
        wday = 1
        yday = 47
        isdst = 0
        gmtoff = -21600
        zone = CST
    }

```

See also: [strftime], page 887, [strptime], page 889, [gmtime], page 886, [mktime], page 887, [time], page 885, [now], page 885, [date], page 890, [clock], page 889, [datenum], page 892, [datestr], page 893, [datevec], page 895, [calendar], page 895, [weekday], page 896.

seconds = mktime (tm_struct)

Convert a time structure corresponding to the local time to the number of seconds since the epoch.

For example:

```

mktime (localtime (time ()))
    ⇒ 856163706

```

See also: [strftime], page 887, [strptime], page 889, [localtime], page 886, [gmtime], page 886, [time], page 885, [now], page 885, [date], page 890, [clock], page 889, [datenum], page 892, [datestr], page 893, [datevec], page 895, [calendar], page 895, [weekday], page 896.

asctime (tm_struct)

Convert a time structure to a string using the following format: "ddd mmm mm HH:MM:SS yyyy\n".

For example:

```

asctime (localtime (time ()))
    ⇒ "Mon Feb 17 01:15:06 1997\n"

```

This is equivalent to `ctime (time ())`.

See also: [ctime], page 886, [localtime], page 886, [time], page 885.

strftime (fmt, tm_struct)

Format the time structure *tm_struct* in a flexible way using the format string *fmt* that contains '%' substitutions similar to those in `printf`.

Except where noted, substituted fields have a fixed size; numeric fields are padded if necessary. Padding is with zeros by default; for fields that display a single number,

padding can be changed or inhibited by following the ‘%’ with one of the modifiers described below. Unknown field specifiers are copied as normal characters. All other characters are copied to the output without change. For example:

```
strftime ("%r (%Z) %A %e %B %Y", localtime (time ()))
⇒ "01:15:06 AM (CST) Monday 17 February 1997"
```

Octave’s `strftime` function supports a superset of the ANSI C field specifiers.

Literal character fields:

<code>%%</code>	% character.
<code>%n</code>	Newline character.
<code>%t</code>	Tab character.

Numeric modifiers (a nonstandard extension):

- (dash) Do not pad the field.
- _ (underscore)
 - Pad the field with spaces.

Time fields:

<code>%H</code>	Hour (00-23).
<code>%I</code>	Hour (01-12).
<code>%k</code>	Hour (0-23).
<code>%l</code>	Hour (1-12).
<code>%M</code>	Minute (00-59).
<code>%p</code>	Locale’s AM or PM.
<code>%r</code>	Time, 12-hour (hh:mm:ss [AP]M).
<code>%R</code>	Time, 24-hour (hh:mm).
<code>%s</code>	Time in seconds since 00:00:00, Jan 1, 1970 (a nonstandard extension).
<code>%S</code>	Second (00-61).
<code>%T</code>	Time, 24-hour (hh:mm:ss).
<code>%X</code>	Locale’s time representation (%H:%M:%S).
<code>%z</code>	Offset from UTC (\pm hhmm), or nothing if no time zone is determinable.
<code>%Z</code>	Time zone (EDT), or nothing if no time zone is determinable.

Date fields:

<code>%a</code>	Locale’s abbreviated weekday name (Sun-Sat).
<code>%A</code>	Locale’s full weekday name, variable length (Sunday-Saturday).
<code>%b</code>	Locale’s abbreviated month name (Jan-Dec).
<code>%B</code>	Locale’s full month name, variable length (January-December).

%c	Locale's date and time (Sat Nov 04 12:02:33 EST 1989).
%C	Century (00-99).
%d	Day of month (01-31).
%e	Day of month (1-31).
%D	Date (mm/dd/yy).
%h	Same as %b.
%j	Day of year (001-366).
%m	Month (01-12).
%U	Week number of year with Sunday as first day of week (00-53).
%w	Day of week (0-6).
%W	Week number of year with Monday as first day of week (00-53).
%x	Locale's date representation (mm/dd/yy).
%y	Last two digits of year (00-99).
%Y	Year (1970-).

See also: [strptime], page 889, [localtime], page 886, [gmtime], page 886, [mktime], page 887, [time], page 885, [now], page 885, [date], page 890, [clock], page 889, [datenum], page 892, [datestr], page 893, [datevec], page 895, [calendar], page 895, [weekday], page 896.

[*tm_struct*, *nchars*] = strptime (*str*, *fmt*)

Convert the string *str* to the time structure *tm_struct* under the control of the format string *fmt*.

If *fmt* fails to match, *nchars* is 0; otherwise, it is set to the position of last matched character plus 1. Always check for this unless you're absolutely sure the date string will be parsed correctly.

See also: [strftime], page 887, [localtime], page 886, [gmtime], page 886, [mktime], page 887, [time], page 885, [now], page 885, [date], page 890, [clock], page 889, [datenum], page 892, [datestr], page 893, [datevec], page 895, [calendar], page 895, [weekday], page 896.

Most of the remaining functions described in this section are not patterned after the standard C library. Some are available for compatibility with MATLAB and others are provided because they are useful.

clock ()

Return the current local date and time as a date vector.

The date vector contains the following fields: current year, month (1-12), day (1-31), hour (0-23), minute (0-59), and second (0-61). The seconds field has a fractional part after the decimal point for extended accuracy.

For example:

```
fix (clock ())
⇒ 1993     8     20      4      56      1
```

`clock` is more accurate on systems that have the `gettimeofday` function.

See also: [\[now\]](#), page 885, [\[date\]](#), page 890, [\[datevec\]](#), page 895.

`date ()`

Return the current date as a character string in the form DD-MMM-YYYY.

For example:

```
date ()
⇒ 20-Aug-1993
```

See also: [\[now\]](#), page 885, [\[clock\]](#), page 889, [\[datestr\]](#), page 893, [\[localtime\]](#), page 886.

`etime (t2, t1)`

Return the difference in seconds between two time values returned from `clock (t2 - t1)`.

For example:

```
t0 = clock ();
# many computations later...
elapsed_time = etime (clock (), t0);
```

will set the variable `elapsed_time` to the number of seconds since the variable `t0` was set.

See also: [\[tic\]](#), page 891, [\[toc\]](#), page 891, [\[clock\]](#), page 889, [\[cputime\]](#), page 890, [\[addtodate\]](#), page 895.

`[total, user, system] = cputime ()`

Return the CPU time used by your Octave session.

The first output is the total time spent executing your process and is equal to the sum of second and third outputs, which are the number of CPU seconds spent executing in user mode and the number of CPU seconds spent executing in system mode, respectively.

If your system does not have a way to report CPU time usage, `cputime` returns 0 for each of its output values.

Note that because Octave used some CPU time to start, it is reasonable to check to see if `cputime` works by checking to see if the total CPU time used is nonzero.

See also: [\[tic\]](#), page 891, [\[toc\]](#), page 891.

`is_leap_year ()`

`is_leap_year (year)`

Return true if `year` is a leap year and false otherwise.

If no year is specified, `is_leap_year` uses the current year.

For example:

```
is_leap_year (2000)
⇒ 1
```

See also: [\[weekday\]](#), page 896, [\[eomday\]](#), page 896, [\[calendar\]](#), page 895.

```
tic ()
id = tic ()
    Initialize a wall-clock timer.
```

Calling `tic` without an output argument resets the internal timer. Subsequent calls to `toc` return the number of seconds since the timer was set.

If called with one output argument, `tic` creates a new timer instance and returns a timer identifier `id`. The `id` is a scalar of type `uint64` that may be passed to `toc` to check elapsed time on this timer, rather than the default internal timer.

Example 1 : benchmarking code with internal timer

```
tic;
# many computations later...
elapsed_time = toc;
```

Example 2 : mixed timer id and internal timer

```
tic;
pause (1);
toc
    ⇒ Elapsed time is 1.0089 seconds.
id = tic;
pause (2);
toc (id)
    ⇒ Elapsed time is 2.01142 seconds.
toc
    Elapsed time is 3.02308 seconds.
```

Calling `tic` and `toc` in this way allows nested timing calls.

If you are more interested in the CPU time that your process used, you should use the `cputime` function instead. The `tic` and `toc` functions report the actual wall clock time that elapsed between the calls. This may include time spent processing other jobs or doing nothing at all.

See also: [\[toc\]](#), page 891, [\[cputime\]](#), page 890.

```
toc ()
toc (id)
elapsed_time = toc (...)
```

Measure elapsed time on a wall-clock timer.

With no arguments, return the number of seconds elapsed on the internal timer since the last call to `tic`.

When given the identifier `id` of a specific timer, return the number of seconds elapsed since the timer `id` was initialized.

See [\[tic\]](#), page 891, for examples of the use of `tic/toc`.

See also: [\[tic\]](#), page 891, [\[cputime\]](#), page 890.

```
pause ()
pause (n)
old_state = pause ("on")
```

```
old_state = pause ("off")
old_state = pause ("query")
```

Suspend the execution of the program or change the state of the pause function.

If invoked without an input arguments then the program is suspended until a character is typed. If argument *n* is a positive real value, it indicates the number of seconds the program shall be suspended, for example:

```
tic; pause (0.05); toc
→ Elapsed time is 0.05039 seconds.
```

The following example prints a message and then waits 5 seconds before clearing the screen.

```
disp ("wait please...");
pause (5);
clc;
```

If invoked with a string argument "on", "off", or "query", the state of the pause function is changed or queried. When the state is "off", the pause function returns immediately. The optional return value contains the previous state of the pause function. In the following example pause is disabled locally:

```
old_state = pause ("off");
tic; pause (0.05); toc
→ Elapsed time is 3.00407e-05 seconds.
pause (old_state);
```

While the program is suspended Octave still handles figures painting and graphics callbacks execution.

See also: [\[kbhit\]](#), page 257.

```
days = datenum (datevec)
days = datenum (year, month, day)
days = datenum (year, month, day, hour)
days = datenum (year, month, day, hour, minute)
days = datenum (year, month, day, hour, minute, second)
days = datenum ("datestr")
days = datenum ("datestr", f)
days = datenum ("datestr", p)
[days, secs] = datenum (...)
```

Return the date/time input as a serial day number, with Jan 1, 0000 defined as day 1.

The integer part, `floor (days)` counts the number of complete days in the date input.

The fractional part, `rem (days, 1)` corresponds to the time on the given day.

The input may be a date vector (see `datevec`), `datestr` (see `datestr`), or directly specified as input.

When processing input datestrings, *f* is the format string used to interpret date strings (see `datestr`). If no format *f* is specified, then a relatively slow search is performed through various formats. It is always preferable to specify the format string *f* if it is known. Formats which do not specify a particular time component will have the

value set to zero. Formats which do not specify a date will default to January 1st of the current year.

p is the year at the start of the century to which two-digit years will be referenced. If not specified, it defaults to the current year minus 50.

The optional output *secs* holds the time on the specified day with greater precision than *days*.

Notes:

- Years can be negative and/or fractional.
- Months below 1 are considered to be January.
- Days of the month start at 1.
- Days beyond the end of the month go into subsequent months.
- Days before the beginning of the month go to the previous month.
- Days can be fractional.

Caution: datenums represent a specific time for the Earth as a whole. They do not take into account time zones (shifts in time based on location), nor seasonal changes due to Daylight Savings Time (shifts in time based on local regulation). Be aware that it is possible to create datenums that, when interpreted by a function which accounts for time zone and DST shifts such as `datestr`, are nonexistent or ambiguous.

Caution: this function does not attempt to handle Julian calendars so dates before October 15, 1582 are wrong by as much as eleven days. Also, be aware that only Roman Catholic countries adopted the calendar in 1582. It took until 1924 for it to be adopted everywhere. See the Wikipedia entry on the Gregorian calendar for more details.

Warning: leap seconds are ignored. A table of leap seconds is available on the Wikipedia entry for leap seconds.

See also: `[datestr]`, page 893, `[datevec]`, page 895, `[now]`, page 885, `[clock]`, page 889, `[date]`, page 890.

```
str = datestr (date)
str = datestr (date, f)
str = datestr (date, f, p)
```

Format the given date/time according to the format *f* and return the result in *str*.

date is a serial date number (see `datenum`), a date vector (see `datevec`), or a string or cell array of strings. In the latter case, it is passed to `datevec` to guess the input date format.

f can be an integer which corresponds to one of the codes in the table below, or a date format string.

p is the year at the start of the century in which two-digit years are to be interpreted in. If not specified, it defaults to the current year minus 50.

For example, the date 730736.65149 (2000-09-07 15:38:09.0934) would be formatted as follows:

Code	Format	Example
0	dd-mmm-yyyy HH:MM:SS	07-Sep-2000 15:38:09
1	dd-mmm-yyyy	07-Sep-2000
2	mm/dd/yy	09/07/00
3	mmm	Sep
4	m	S
5	mm	09
6	mm/dd	09/07
7	dd	07
8	ddd	Thu
9	d	T
10	yyyy	2000
11	yy	00
12	mmmyy	Sep00
13	HH:MM:SS	15:38:09
14	HH:MM:SS PM	3:38:09 PM
15	HH:MM	15:38
16	HH:MM PM	3:38 PM
17	QQ-YY	Q3-00
18	QQ	Q3
19	dd/mm	07/09
20	dd/mm/yy	07/09/00
21	mmm.dd,yyyy HH:MM:SS	Sep.07,2000 15:38:08
22	mmm.dd,yyyy	Sep.07,2000
23	mm/dd/yyyy	09/07/2000
24	dd/mm/yyyy	07/09/2000
25	yy/mm/dd	00/09/07
26	yyyy/mm/dd	2000/09/07
27	QQ-YYYY	Q3-2000
28	mmmyyyy	Sep2000
29	yyyy-mm-dd	2000-09-07
30	yyyymmddTHHMMSS	20000907T153808
31	yyyy-mm-dd HH:MM:SS	2000-09-07 15:38:08

If *f* is a format string, the following symbols are recognized:

Symbol	Meaning	Example
yyyy	Full year	2005
yy	Two-digit year	05
mmmm	Full month name	December
mmm	Abbreviated month name	Dec
mm	Numeric month number (padded with zeros)	01, 08, 12
m	First letter of month name (capitalized)	D
dddd	Full weekday name	Sunday
ddd	Abbreviated weekday name	Sun
dd	Numeric day of month (padded with zeros)	11
d	First letter of weekday name (capitalized)	S

HH	Hour of day, padded with zeros, or padded with spaces if PM is set	09:00 9:00 AM
MM	Minute of hour (padded with zeros)	10:05
SS	Second of minute (padded with zeros)	10:05:03
FFF	Milliseconds of second (padded with zeros)	10:05:03.012
AM	Use 12-hour time format	11:30 AM
PM	Use 12-hour time format	11:30 PM

If *f* is not specified or is -1, then use 0, 1 or 16, depending on whether the date portion or the time portion of *date* is empty.

If *p* is not specified, it defaults to the current year minus 50.

If a matrix or cell array of dates is given, a column vector of date strings is returned.

See also: [datenum], page 892, [datevec], page 895, [date], page 890, [now], page 885, [clock], page 889.

```
v = datevec (date)
v = datevec (date, f)
v = datevec (date, p)
v = datevec (date, f, p)
[y, m, d, h, mi, s] = datevec (...)
```

Convert a serial date number (see **datenum**) or date string (see **datestr**) into a date vector.

A date vector is a row vector with six members, representing the year, month, day, hour, minute, and seconds respectively.

f is the format string used to interpret date strings (see **datestr**). If *date* is a string, but no format is specified, then a relatively slow search is performed through various formats. It is always preferable to specify the format string *f* if it is known. Formats which do not specify a particular time component will have the value set to zero. Formats which do not specify a date will default to January 1st of the current year.

p is the year at the start of the century to which two-digit years will be referenced. If not specified, it defaults to the current year minus 50.

See also: [datenum], page 892, [datestr], page 893, [clock], page 889, [now], page 885, [date], page 890.

```
d = addtodate (d, q, f)
```

Add *q* amount of time (with units *f*) to the serial datenum, *d*.

f must be one of "year", "month", "day", "hour", "minute", "second", or "millisecond".

See also: [datenum], page 892, [datevec], page 895, [etime], page 890.

```
c = calendar ()
c = calendar (d)
c = calendar (y, m)
calendar (...)
```

Return the current monthly calendar in a 6x7 matrix.

If d is specified, return the calendar for the month containing the date d , which must be a serial date number or a date string.

If y and m are specified, return the calendar for year y and month m .

If no output arguments are specified, print the calendar on the screen instead of returning a matrix.

See also: [\[datenum\]](#), page 892, [\[datestr\]](#), page 893.

[n, s] = weekday (d)

[n, s] = weekday (d, format)

Return the day of the week as a number in n and as a string in s .

The days of the week are numbered 1–7 with the first day being Sunday.

d is a serial date number or a date string.

If the string $format$ is not present or is equal to "short" then s will contain the abbreviated name of the weekday. If $format$ is "long" then s will contain the full name.

Table of return values based on $format$:

<i>n</i>	"short"	"long"
1	Sun	Sunday
2	Mon	Monday
3	Tue	Tuesday
4	Wed	Wednesday
5	Thu	Thursday
6	Fri	Friday
7	Sat	Saturday

See also: [\[eomday\]](#), page 896, [\[is_leap_year\]](#), page 890, [\[calendar\]](#), page 895, [\[datenum\]](#), page 892, [\[datevec\]](#), page 895.

e = eomday (y, m)

Return the last day of the month m for the year y .

See also: [\[weekday\]](#), page 896, [\[datenum\]](#), page 892, [\[datevec\]](#), page 895, [\[is_leap_year\]](#), page 890, [\[calendar\]](#), page 895.

```
datetick ()
datetick (axis_str)
datetick (date_format)
datetick (axis_str, date_format)
datetick (..., "keeplimits")
datetick (..., "keerticks")
datetick (hax, ...)
```

Add date-formatted tick labels to an axis.

The axis to apply the ticks to is determined by $axis_str$ which can take the values "x", "y", or "z". The default value is "x".

The formatting of the labels is determined by the variable $date_format$, which can either be a string or positive integer that [datestr](#) accepts.

If the first argument *hax* is an axes handle, then plot into this axes, rather than the current axes returned by `gca`.

See also: `[datenum]`, page 892, `[datestr]`, page 893.

36.2 Filesystem Utilities

Octave includes many utility functions for copying, moving, renaming, and deleting files; for creating, reading, and deleting directories; for retrieving status information on files; and for manipulating file and path names.

```
movefile f1
movefile f1 f2
movefile f1 f2 f
movefile (f1)
movefile (f1, f2)
movefile (f1, f2, 'f')
[status, msg, msgid] = movefile (...)
```

Move the source file or directory *f1* to the destination *f2*.

The name *f1* may contain globbing patterns, or may be a cell array of strings. If *f1* expands to multiple filenames, *f2* must be a directory.

If no destination *f2* is specified then the destination is the present working directory. If *f2* is a filename then *f1* is renamed to *f2*.

When the force flag '*f*' is given any existing files will be overwritten without prompting.

If successful, *status* is 1, and *msg*, *msgid* are empty character strings (""). Otherwise, *status* is 0, *msg* contains a system-dependent error message, and *msgid* contains a unique message identifier. Note that the status code is exactly opposite that of the `system` command.

See also: `[rename]`, page 897, `[copyfile]`, page 897, `[unlink]`, page 898, `[delete]`, page 377, `[glob]`, page 903.

```
rename old new
[err, msg] = rename (old, new)
```

Change the name of file *old* to *new*.

If successful, *err* is 0 and *msg* is an empty string. Otherwise, *err* is nonzero and *msg* contains a system-dependent error message.

See also: `[movefile]`, page 897, `[copyfile]`, page 897, `[ls]`, page 923, `[dir]`, page 923.

```
copyfile f1 f2
copyfile f1 f2 f
copyfile (f1, f2)
copyfile (f1, f2, 'f')
[status, msg, msgid] = copyfile (...)
```

Copy the source file(s) or directory *f1* to the destination *f2*.

The name *f1* may contain globbing patterns, or may be a cell array of strings. If *f1* expands to multiple filenames, *f2* must be a directory.

When the force flag '`f`' is given any existing files will be overwritten without prompting.

If successful, `status` is 1, and `msg`, `msgid` are empty character strings (""). Otherwise, `status` is 0, `msg` contains a system-dependent error message, and `msgid` contains a unique message identifier. Note that the status code is exactly opposite that of the `system` command.

See also: [\[movefile\]](#), page 897, [\[rename\]](#), page 897, [\[unlink\]](#), page 898, [\[delete\]](#), page 377, [\[glob\]](#), page 903.

[err, msg] = unlink (file)

Delete the file named `file`.

If successful, `err` is 0 and `msg` is an empty string. Otherwise, `err` is nonzero and `msg` contains a system-dependent error message.

See also: [\[delete\]](#), page 377, [\[rmdir\]](#), page 899.

link old new

[err, msg] = link (old, new)

Create a new link (also known as a hard link) to an existing file.

If successful, `err` is 0 and `msg` is an empty string. Otherwise, `err` is nonzero and `msg` contains a system-dependent error message.

See also: [\[symlink\]](#), page 898, [\[unlink\]](#), page 898, [\[readlink\]](#), page 898, [\[lstat\]](#), page 900.

symlink old new

[err, msg] = symlink (old, new)

Create a symbolic link `new` which contains the string `old`.

If successful, `err` is 0 and `msg` is an empty string. Otherwise, `err` is nonzero and `msg` contains a system-dependent error message.

See also: [\[link\]](#), page 898, [\[unlink\]](#), page 898, [\[readlink\]](#), page 898, [\[lstat\]](#), page 900.

readlink symlink

[result, err, msg] = readlink (symlink)

Read the value of the symbolic link `symlink`.

If successful, `result` contains the contents of the symbolic link `symlink`, `err` is 0, and `msg` is an empty string. Otherwise, `err` is nonzero and `msg` contains a system-dependent error message.

See also: [\[lstat\]](#), page 900, [\[symlink\]](#), page 898, [\[link\]](#), page 898, [\[unlink\]](#), page 898, [\[delete\]](#), page 377.

mkdir dirname

mkdir parent dirname

mkdir (dirname)

mkdir (parent, dirname)

[status, msg, msgid] = mkdir (...)

Create a directory named `dirname` in the directory `parent`, creating any intermediate directories if necessary.

If *dirname* is a relative path, and no parent directory is specified, then the present working directory is used.

If successful, *status* is 1, and *msg* and *msgid* are empty strings (""). Otherwise, *status* is 0, *msg* contains a system-dependent error message, and *msgid* contains a unique message identifier.

When creating a directory permissions will be set to 0777 - UMASK.

See also: [\[rmdir\]](#), page 899, [\[pwd\]](#), page 924, [\[cd\]](#), page 922, [\[umask\]](#), page 899.

```
rmdir dir
rmdir (dir, "s")
[status, msg, msgid] = rmdir (...)
```

Remove the directory named *dir*.

If the optional second parameter is supplied with value "s", recursively remove all subdirectories as well.

If successful, *status* is 1, and *msg*, *msgid* are empty character strings (""). Otherwise, *status* is 0, *msg* contains a system-dependent error message, and *msgid* contains a unique message identifier.

See also: [\[mkdir\]](#), page 898, [\[confirm_recursive_rmdir\]](#), page 899, [\[pwd\]](#), page 924.

```
val = confirm_recursive_rmdir ()
old_val = confirm_recursive_rmdir (new_val)
confirm_recursive_rmdir (new_val, "local")
```

Query or set the internal variable that controls whether Octave will ask for confirmation before recursively removing a directory tree.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[rmdir\]](#), page 899.

```
err = mkfifo (name, mode)
[err, msg] = mkfifo (name, mode)
```

Create a FIFO special file named *name* with file mode *mode*.

mode is interpreted as an octal number and is subject to umask processing. The final calculated mode is *mode* - *umask*.

If successful, *err* is 0 and *msg* is an empty string. Otherwise, *err* is nonzero and *msg* contains a system-dependent error message.

See also: [\[pipe\]](#), page 917, [\[umask\]](#), page 899.

```
umask (mask)
```

Set the permission mask for file creation.

The parameter *mask* is an integer, interpreted as an octal number.

If successful, returns the previous value of the mask (as an integer to be interpreted as an octal number); otherwise an error message is printed.

The permission mask is a UNIX concept used when creating new objects on a file system such as files, directories, or named FIFOs. The object to be created has base

permissions in an octal number *mode* which are modified according to the octal value of *mask*. The final permissions for the new object are *mode - mask*.

See also: [fopen], page 273, [mkdir], page 898, [mkfifo], page 899.

```
[info, err, msg] = stat (file)
[info, err, msg] = stat (fid)
[info, err, msg] = lstat (file)
[info, err, msg] = lstat (fid)
```

Return a structure *info* containing the following information about *file* or file identifier *fid*.

dev	ID of device containing a directory entry for this file.
ino	File number of the file.
mode	File mode, as an integer. Use the functions S_ISREG, S_ISDIR, S_ISCHR, S_ISBLK, S_ISFIFO, S_ISLNK, or S_ISSOCK to extract information from this value.
modestr	File mode, as a string of ten letters or dashes as would be returned by <i>ls -l</i> .
nlink	Number of links.
uid	User ID of file's owner.
gid	Group ID of file's group.
rdev	ID of device for block or character special files.
size	Size in bytes.
atime	Time of last access in the same form as time values returned from time . See Section 36.1 [Timing Utilities], page 885.
mtime	Time of last modification in the same form as time values returned from time . See Section 36.1 [Timing Utilities], page 885.
ctime	Time of last file status change in the same form as time values returned from time . See Section 36.1 [Timing Utilities], page 885.
blksize	Size of blocks in the file.
blocks	Number of blocks allocated for file.

If the call is successful *err* is 0 and *msg* is an empty string. If the file does not exist, or some other error occurs, *info* is an empty matrix, *err* is -1, and *msg* contains the corresponding system error message.

If *file* is a symbolic link, **stat** will return information about the actual file that is referenced by the link. Use **lstat** if you want information about the symbolic link itself.

For example:

```
[info, err, msg] = stat ("/vmlinuz")
⇒ info =
```

```
{  
    atime = 855399756  
    rdev = 0  
    ctime = 847219094  
    uid = 0  
    size = 389218  
    blksize = 4096  
    mtime = 847219094  
    gid = 6  
    nlink = 1  
    blocks = 768  
    mode = -rw-r--r--  
    modestr = -rw-r--r--  
    ino = 9316  
    dev = 2049  
}  
⇒ err = 0  
⇒ msg =
```

See also: [\[lstat\]](#), page 900, [\[ls\]](#), page 923, [\[dir\]](#), page 923, [\[isfile\]](#), page 903, [\[isfolder\]](#), page 903.

S_ISBLK (*mode*)

Return true if *mode* corresponds to a block device.

The value of *mode* is assumed to be returned from a call to **stat**.

See also: [\[stat\]](#), page 900, [\[lstat\]](#), page 900.

S_ISCHR (*mode*)

Return true if *mode* corresponds to a character device.

The value of *mode* is assumed to be returned from a call to **stat**.

See also: [\[stat\]](#), page 900, [\[lstat\]](#), page 900.

S_ISDIR (*mode*)

Return true if *mode* corresponds to a directory.

The value of *mode* is assumed to be returned from a call to **stat**.

See also: [\[stat\]](#), page 900, [\[lstat\]](#), page 900.

S_ISFIFO (*mode*)

Return true if *mode* corresponds to a fifo.

The value of *mode* is assumed to be returned from a call to **stat**.

See also: [\[stat\]](#), page 900, [\[lstat\]](#), page 900.

S_ISLNK (*mode*)

Return true if *mode* corresponds to a symbolic link.

The value of *mode* is assumed to be returned from a call to **stat**.

See also: [\[stat\]](#), page 900, [\[lstat\]](#), page 900.

S_ISREG (*mode*)

Return true if *mode* corresponds to a regular file.

The value of *mode* is assumed to be returned from a call to **stat**.

See also: [\[stat\]](#), page 900, [\[lstat\]](#), page 900.

S_ISSOCK (*mode*)

Return true if *mode* corresponds to a socket.

The value of *mode* is assumed to be returned from a call to **stat**.

See also: [\[stat\]](#), page 900, [\[lstat\]](#), page 900.

fileattrib (*file*)**fileattrib ()****[*status*, *msg*, *msgid*] = fileattrib (...)**

Return information about *file*.

If successful, *status* is 1 and *msg* is a structure with the following fields:

Name Full name of *file*.

archive True if *file* is an archive (Windows).

system True if *file* is a system file (Windows).

hidden True if *file* is a hidden file (Windows).

directory

True if *file* is a directory.

UserRead

GroupRead

OtherRead

True if the user (group; other users) has read permission for *file*.

UserWrite

GroupWrite

OtherWrite

True if the user (group; other users) has write permission for *file*.

UserExecute

GroupExecute

OtherExecute

True if the user (group; other users) has execute permission for *file*.

If an attribute does not apply (i.e., archive on a Unix system) then the field is set to NaN.

If **attrib** fails, *msg* is a non-empty string containing an error message and *msg-id* is the non-empty string "fileattrib".

With no input arguments, return information about the current directory.

If *file* contains globbing characters, return information about all the matching files.

See also: [\[glob\]](#), page 903.

`tf = isfile (f)`

Return true if *f* is a regular file and false otherwise.

If *f* is a cell array of strings, *tf* is a logical array of the same size.

See also: [isfolder], page 903, [exist], page 134, [stat], page 900, [is_absolute_filename], page 906, [is_rooted_relative_filename], page 906.

`isdir (f)`

This function is not recommended. Use `isfolder` or `file_in_loadpath` instead.

Return true if *f* is a directory and false otherwise.

Compatibility Note: The MATLAB function of the same name will also search for *f* in the load path directories. To emulate this behavior use

`tf = ! isempty (file_in_loadpath (f))`

See also: [isfolder], page 903, [file_in_loadpath], page 198, [exist], page 134, [stat], page 900, [is_absolute_filename], page 906, [is_rooted_relative_filename], page 906.

`isfolder (f)`

`tf = isfolder (f)`

Return true if *f* is a directory and false otherwise.

If *f* is a cell array of strings, *tf* is a logical array of the same size.

See also: [isfile], page 903, [exist], page 134, [stat], page 900, [is_absolute_filename], page 906, [is_rooted_relative_filename], page 906.

`files = readdir (dir)`

`[files, err, msg] = readdir (dir)`

Return the names of files in the directory *dir* as a cell array of strings.

If an error occurs, return an empty cell array in *files*. If successful, *err* is 0 and *msg* is an empty string. Otherwise, *err* is nonzero and *msg* contains a system-dependent error message.

See also: [ls], page 923, [dir], page 923, [glob], page 903, [what], page 138.

`glob (pattern)`

Given an array of pattern strings (as a char array or a cell array) in *pattern*, return a cell array of filenames that match any of them, or an empty cell array if no patterns match.

The pattern strings are interpreted as filename globbing patterns (as they are used by Unix shells).

Within a pattern

- * matches any string, including the null string,
- ? matches any single character, and
- [...] matches any of the enclosed characters.

Tilde expansion is performed on each of the patterns before looking for matching filenames. For example:

`ls`

```

⇒
    file1  file2  file3  myfile1 myfile1b
glob ("*file1")
⇒
{
    [1,1] = file1
    [2,1] = myfile1
}
glob ("myfile?")
⇒
{
    [1,1] = myfile1
}
glob ("file[12]")
⇒
{
    [1,1] = file1
    [2,1] = file2
}

```

See also: [\[ls\]](#), page 923, [\[dir\]](#), page 923, [\[readdir\]](#), page 903, [\[what\]](#), page 138.

file_in_path (*path*, *file*)
file_in_path (*path*, *file*, "all")

Return the absolute name of *file* if it can be found in *path*.

The value of *path* should be a colon-separated list of directories in the format described for *path*. If no file is found, return an empty character string. For example:

```

file_in_path (EXEC_PATH, "sh")
⇒ "/bin/sh"

```

If the second argument is a cell array of strings, search each directory of the path for element of the cell array and return the first that matches.

If the third optional argument "all" is supplied, return a cell array containing the list of all files that have the same name in the path. If no files are found, return an empty cell array.

See also: [\[file_in_loadpath\]](#), page 198, [\[dir_in_loadpath\]](#), page 199, [\[path\]](#), page 197.

filesep ()
filesep ("all")

Return the system-dependent character used to separate directory names.

If "all" is given, the function returns all valid file separators in the form of a string. The list of file separators is system-dependent. It is '/' (forward slash) under UNIX or Mac OS X, '/' and '\' (forward and backward slashes) under Windows.

See also: [\[pathsep\]](#), page 198.

[dir, name, ext] = fileparts (filename)

Return the directory, name, and extension components of *filename*.

The input *filename* is a string which is parsed. There is no attempt to check whether the filename or directory specified actually exists.

See also: [\[fullfile\]](#), page 905, [\[filesep\]](#), page 904.

```
filename = fullfile (dir1, dir2, ..., file)
filenames = fullfile (... , files)
```

Build complete filename from separate parts.

Joins any number of path components intelligently. The return value is the concatenation of each component with exactly one file separator between each non empty part and at most one leading and/or trailing file separator.

If the last component part is a cell array, returns a cell array of filepaths, one for each element in the last component, e.g.:

```
fullfile ("/home/username", "data", {"f1.csv", "f2.csv", "f3.csv"})
⇒
{
    [1,1] = /home/username/data/f1.csv
    [1,2] = /home/username/data/f2.csv
    [1,3] = /home/username/data/f3.csv
}
```

On Windows systems, while forward slash file separators do work, they are replaced by backslashes; in addition drive letters are stripped of leading file separators to obtain a valid file path.

Note: `fullfile` does not perform any validation of the resulting full filename.

See also: [\[fileparts\]](#), page 904, [\[filesep\]](#), page 904.

```
tilde_expand (string)
tilde_expand (cellstr)
```

Perform tilde expansion on *string*.

If *string* begins with a tilde character, ('~'), all of the characters preceding the first slash (or all characters, if there is no slash) are treated as a possible user name, and the tilde and the following characters up to the slash are replaced by the home directory of the named user. If the tilde is followed immediately by a slash, the tilde is replaced by the home directory of the user running Octave.

If the input is a cell array of strings *cellstr* then tilde expansion is performed on each string element.

Examples:

```
tilde_expand ("~joeuser/bin")
⇒ "/home/joeuser/bin"
tilde_expand ("~/bin")
⇒ "/home/jwe/bin"
```

```
[cname, status, msg] = canonicalize_file_name (fname)
```

Return the canonical name of file *fname*.

If the file does not exist the empty string ("") is returned.

See also: [\[make_absolute_filename\]](#), page 906, [\[is_absolute_filename\]](#), page 906, [\[is_rooted_relative_filename\]](#), page 906.

make_absolute_filename (file)

Return the full name of *file* beginning from the root of the file system.

No check is done for the existence of *file*.

See also: [canonicalize_file_name], page 905, [is_absolute_filename], page 906, [is_rooted_relative_filename], page 906, [isfolder], page 903.

is_absolute_filename (file)

Return true if *file* is an absolute filename.

See also: [is_rooted_relative_filename], page 906, [make_absolute_filename], page 906, [isfolder], page 903.

is_rooted_relative_filename (file)

Return true if *file* is a rooted-relative filename.

See also: [is_absolute_filename], page 906, [make_absolute_filename], page 906, [isfolder], page 903.

val = recycle ()

old_val = recycle (new_val)

Query or set the preference for recycling deleted files.

When recycling is enabled, commands which would permanently erase files instead move them to a temporary location (such as the directory labeled Trash).

Programming Note: This function is provided for MATLAB compatibility, but recycling is not implemented in Octave. To help avoid accidental data loss an error will be raised if an attempt is made to enable file recycling.

See also: [delete], page 377, [rmdir], page 899.

36.3 File Archiving Utilities

bunzip2 (bzfile)

bunzip2 (bzfile, dir)

filelist = bunzip2 (...)

Unpack the bzip2 archive *bzfile*.

If *dir* is specified the files are unpacked in this directory rather than the one where *bzfile* is located.

The optional output *filelist* is a list of the uncompressed files.

See also: [bzip2], page 909, [unpack], page 908, [gunzip], page 907, [unzip], page 908, [untar], page 907.

filelist = gzip (files)

filelist = gzip (files, dir)

Compress the list of files and directories specified in *files*.

files is a character array or cell array of strings. Shell wildcards in the filename such as '*' or '?' are accepted and expanded. Each file is compressed separately and a new file with a ".gz" extension is created. The original files are not modified, but

existing compressed files will be silently overwritten. If a directory is specified then `gzip` recursively compresses all files in the directory.

If `dir` is defined the compressed files are placed in this directory, rather than the original directory where the uncompressed file resides. Note that this does not replicate a directory tree in `dir` which may lead to files overwriting each other if there are multiple files with the same name.

If `dir` does not exist it is created.

The optional output `filelist` is a list of the compressed files.

See also: [\[gunzip\]](#), page 907, [\[unpack\]](#), page 908, [\[bzip2\]](#), page 909, [\[zip\]](#), page 908, [\[tar\]](#), page 907.

```
gunzip (gzfile)
gunzip (gzfile, dir)
filelist = gunzip (...)
```

Unpack the gzip archive `gzfile`.

If `gzfile` is a directory, all gzfiles in the directory will be recursively unpacked.

If `dir` is specified the files are unpacked in this directory rather than the one where `gzfile` is located.

The optional output `filelist` is a list of the uncompressed files.

See also: [\[gzip\]](#), page 906, [\[unpack\]](#), page 908, [\[bunzip2\]](#), page 906, [\[unzip\]](#), page 908, [\[untar\]](#), page 907.

```
filelist = tar (tarfile, files)
filelist = tar (tarfile, files, rootdir)
```

Pack the list of files and directories specified in `files` into the TAR archive `tarfile`.

`files` is a character array or cell array of strings. Shell wildcards in the filename such as '*' or '?' are accepted and expanded. Directories are recursively traversed and all files are added to the archive.

If `rootdir` is defined then any files without absolute pathnames are located relative to `rootdir` rather than the current directory.

The optional output `filelist` is a list of the files that were included in the archive.

See also: [\[untar\]](#), page 907, [\[unpack\]](#), page 908, [\[bzip2\]](#), page 909, [\[gzip\]](#), page 906, [\[zip\]](#), page 908.

```
untar (tarfile)
untar (tarfile, dir)
filelist = untar (...)
```

Unpack the TAR archive `tarfile`.

If `dir` is specified the files are unpacked in this directory rather than the one where `tarfile` is located.

The optional output `filelist` is a list of the uncompressed files.

See also: [\[tar\]](#), page 907, [\[unpack\]](#), page 908, [\[bunzip2\]](#), page 906, [\[gunzip\]](#), page 907, [\[unzip\]](#), page 908.

```
filelist = zip (zipfile, files)
filelist = zip (zipfile, files, rootdir)
```

Compress the list of files and directories specified in *files* into the ZIP archive *zipfile*.

files is a character array or cell array of strings. Shell wildcards in the filename such as '*' or '?' are accepted and expanded. Directories are recursively traversed and all files are compressed and added to the archive.

If *rootdir* is defined then any files without absolute pathnames are located relative to *rootdir* rather than the current directory.

The optional output *filelist* is a list of the files that were included in the archive.

See also: [\[unzip\]](#), page 908, [\[unpack\]](#), page 908, [\[bzip2\]](#), page 909, [\[gzip\]](#), page 906, [\[tar\]](#), page 907.

```
unzip (zipfile)
unzip (zipfile, dir)
filelist = unzip (...)
```

Unpack the ZIP archive *zipfile*.

If *dir* is specified the files are unpacked in this directory rather than the one where *zipfile* is located.

The optional output *filelist* is a list of the uncompressed files.

See also: [\[zip\]](#), page 908, [\[unpack\]](#), page 908, [\[bunzip2\]](#), page 906, [\[gunzip\]](#), page 907, [\[untar\]](#), page 907.

```
files = unpack (file)
files = unpack (file, dir)
files = unpack (file, dir, filetype)
```

Unpack the archive *file* based on its extension to the directory *dir*.

If *file* is a list of strings, then each file is unpacked individually. Shell wildcards in the filename such as '*' or '?' are accepted and expanded.

If *dir* is not specified or is empty ([]), it defaults to the current directory. If a directory is in the file list, then *filetype* must also be specified.

The specific archive filetype is inferred from the extension of the file. The *filetype* may also be specified directly using a string which corresponds to a known extension.

Valid filetype extensions:

bz	
bz2	bzip archive
gz	gzip archive
tar	tar archive
tarbz	
tarbz2	
tbz	
tbz2	tar + bzip archive
targz	
tgz	tar + gzip archive

z compress archive
zip zip archive

The optional return value is a list of *files* unpacked.

See also: [bunzip2], page 906, [gunzip], page 907, [unzip], page 908, [untar], page 907, [bzip2], page 909, [gzip], page 906, [zip], page 908, [tar], page 907.

```
filelist = bzip2 (files)
filelist = bzip2 (files, dir)
```

Compress the list of files specified in *files*.

files is a character array or cell array of strings. Shell wildcards in the filename such as '*' or '?' are accepted and expanded. Each file is compressed separately and a new file with a ".bz2" extension is created. The original files are not modified, but existing compressed files will be silently overwritten.

If *dir* is defined the compressed files are placed in this directory, rather than the original directory where the uncompressed file resides. Note that this does not replicate a directory tree in *dir* which may lead to files overwriting each other if there are multiple files with the same name.

If *dir* does not exist it is created.

The optional output *filelist* is a list of the compressed files.

See also: [bunzip2], page 906, [unpack], page 908, [gzip], page 906, [zip], page 908, [tar], page 907.

36.4 Networking Utilities

gethostname ()

Return the hostname of the system where Octave is running.

36.4.1 FTP Objects

Octave supports the FTP protocol through an object-oriented interface. Use the function **ftp** to create an FTP object which represents the connection. All FTP functions take an FTP object as the first argument.

```
f = ftp (host)
```

```
f = ftp (host, username, password)
```

Connect to the FTP server *host* with *username* and *password*.

If *username* and *password* are not specified, user "anonymous" with no password is used. The returned FTP object *f* represents the established FTP connection.

The list of actions for an FTP object are shown below. All functions require an FTP object as the first argument.

Method	Description
ascii	Set transfer type to ascii
binary	Set transfer type to binary
cd	Change remote working directory

close	Close FTP connection
delete	Delete remote file
dir	List remote directory contents
mget	Download remote files
mkdir	Create remote directory
mput	Upload local files
rename	Rename remote file or directory
rmdir	Remove remote directory

close (f)
 Close the FTP connection represented by the FTP object *f*.
f is an FTP object returned by the **ftp** function.

mget (f, file)
mget (f, dir)
mget (f, remote_name, target)
 Download a remote file *file* or directory *dir* to the local directory on the FTP connection *f*.
f is an FTP object returned by the **ftp** function.
 The arguments *file* and *dir* can include wildcards and any files or directories on the remote server that match will be downloaded.
 If a third string argument *target* is given, then it must indicate the path to the local destination directory. *target* may be a relative or absolute path.

mput (f, file)
 Upload the local file *file* into the current remote directory on the FTP connection *f*.
f is an FTP object returned by the **ftp** function.
 The argument *file* is passed through the **glob** function and any files that match the wildcards in *file* will be uploaded.

cd (f)
cd (f, path)
 Get or set the remote directory on the FTP connection *f*.
f is an FTP object returned by the **ftp** function.
 If *path* is not specified, return the remote current working directory. Otherwise, set the remote directory to *path* and return the new remote working directory.
 If the directory does not exist, an error message is printed and the working directory is not changed.

lst = dir (f)
 List the current directory in verbose form for the FTP connection *f*.
f is an FTP object returned by the **ftp** function.

ascii (f)
 Set the FTP connection *f* to use ASCII mode for transfers.
 ASCII mode is only appropriate for text files as it will convert the remote host's newline representation to the local host's newline representation.
f is an FTP object returned by the **ftp** function.

binary (f)

Set the FTP connection *f* to use binary mode for transfers.

In binary mode there is no conversion of newlines from the remote representation to the local representation.

f is an FTP object returned by the `ftp` function.

delete (f, file)

Delete the remote file *file* over the FTP connection *f*.

f is an FTP object returned by the `ftp` function.

rename (f, oldname, newname)

Rename or move the remote file or directory *oldname* to *newname*, over the FTP connection *f*.

f is an FTP object returned by the `ftp` function.

mkdir (f, path)

Create the remote directory *path*, over the FTP connection *f*.

f is an FTP object returned by the `ftp` function.

rmdir (f, path)

Remove the remote directory *path*, over the FTP connection *f*.

f is an FTP object returned by the `ftp` function.

36.4.2 URL Manipulation

```
s = urlread (url)
[s, success] = urlread (url)
[s, success, message] = urlread (url)
[...] = urlread (url, method, param)
```

Download a remote file specified by its *url* and return its content in string *s*.

For example:

```
s = urlread ("ftp://ftp.octave.org/pub/README");
```

The variable *success* is 1 if the download was successful, otherwise it is 0 in which case *message* contains an error message.

If no output argument is specified and an error occurs, then the error is signaled through Octave's error handling mechanism.

This function uses libcurl. Curl supports, among others, the HTTP, FTP, and FILE protocols. Username and password may be specified in the URL. For example:

```
s = urlread ("http://user:password@example.com/file.txt");
```

GET and POST requests can be specified by *method* and *param*. The parameter *method* is either 'get' or 'post' and *param* is a cell array of parameter and value pairs. For example:

```
s = urlread ("http://www.google.com/search", "get",
             {"query", "octave"});
```

See also: [\[urlwrite\]](#), page 912.

```
urlwrite (url, localfile)
f = urlwrite (url, localfile)
[f, success] = urlwrite (url, localfile)
[f, success, message] = urlwrite (url, localfile)
Download a remote file specified by its url and save it as localfile.
```

For example:

```
urlwrite ("ftp://ftp.octave.org/pub/README",
          "README.txt");
```

The full path of the downloaded file is returned in *f*.

The variable *success* is 1 if the download was successful, otherwise it is 0 in which case *message* contains an error message.

If no output argument is specified and an error occurs, then the error is signaled through Octave's error handling mechanism.

This function uses libcurl. Curl supports, among others, the HTTP, FTP, and FILE protocols. Username and password may be specified in the URL, for example:

```
urlwrite ("http://username:password@example.com/file.txt",
          "file.txt");
```

GET and POST requests can be specified by *method* and *param*. The parameter *method* is either 'get' or 'post' and *param* is a cell array of parameter and value pairs. For example:

```
urlwrite ("http://www.google.com/search", "search.html",
          "get", {"query", "octave"});
```

See also: [\[urlread\]](#), page 911.

36.4.3 Base64 and Binary Data Transmission

Some transmission channels can not accept binary data. It is customary to encode binary data in Base64 for transmission and to decode the data upon reception.

```
s = base64_encode (x)
```

Encode a double matrix or array *x* into the base64 format string *s*.

See also: [\[base64_decode\]](#), page 912.

```
x = base64_decode (s)
```

```
x = base64_decode (s, dims)
```

Decode the double matrix or array *x* from the base64 encoded string *s*.

The optional input parameter *dims* should be a vector containing the dimensions of the decoded array.

See also: [\[base64_encode\]](#), page 912.

36.5 Controlling Subprocesses

Octave includes some high-level commands like `system` and `popen` for starting subprocesses. If you want to run another program to perform some task and then look at its output, you will probably want to use these functions.

Octave also provides several very low-level Unix-like functions which can also be used for starting subprocesses, but you should probably only use them if you can't find any way to do what you need with the higher-level functions.

```
system ("string")
system ("string", return_output)
system ("string", return_output, type)
[status, output] = system (...)
```

Execute a shell command specified by *string*.

If the optional argument *type* is "async", the process is started in the background and the process ID of the child process is returned immediately. Otherwise, the child process is started and Octave waits until it exits. If the *type* argument is omitted, it defaults to the value "sync".

If `system` is called with one or more output arguments, or if the optional argument `return_output` is true and the subprocess is started synchronously, then the output from the command is returned as a variable. Otherwise, if the subprocess is executed synchronously, its output is sent to the standard output. To send the output of a command executed with `system` through the pager, use a command like

```
[~, text] = system ("cmd");
more on;
disp (text);
```

or

```
more on;
printf ("%s\n", nthargout (2, "system", "cmd"));
```

The `system` function can return two values. The first is the exit status of the command and the second is any output from the command that was written to the standard output stream. For example,

```
[status, output] = system ("echo foo & exit 2");
```

will set the variable *output* to the string 'foo', and the variable *status* to the integer '2'.

For commands run asynchronously, *status* is the process id of the command shell that is started to run the command.

The shell used for executing commands varies with operating system and is typically `/bin/sh` for UNIX systems and `cmd.exe` for Windows systems.

See also: [\[unix\]](#), page 913, [\[dos\]](#), page 914.

```
unix ("command")
status = unix ("command")
[status, text] = unix ("command")
```

`[...] = unix ("command", "-echo")`

Execute a system command if running under a Unix-like operating system, otherwise do nothing.

Octave waits for the external command to finish before returning the exit status of the program in *status* and any output in *text*.

When called with no output argument, or the `"-echo"` argument is given, then *text* is also sent to standard output.

See also: [\[dos\]](#), page 914, [\[system\]](#), page 913, [\[isunix\]](#), page 927, [\[ismac\]](#), page 927, [\[ispc\]](#), page 927.

`dos ("command")`

`status = dos ("command")`

`[status, text] = dos ("command")`

`[...] = dos ("command", "-echo")`

Execute a system command if running under a Windows-like operating system, otherwise do nothing.

Octave waits for the external command to finish before returning the exit status of the program in *status* and any output in *text*.

When called with no output argument, or the `"-echo"` argument is given, then *text* is also sent to standard output.

See also: [\[unix\]](#), page 913, [\[system\]](#), page 913, [\[isunix\]](#), page 927, [\[ismac\]](#), page 927, [\[ispc\]](#), page 927.

`open file`

`output = open (file)`

Open the file *file* in Octave or in an external application based on the file type as determined by the filename extension.

By default, recognized file types are

`.m` Open file in the editor. No *output* value is returned.

`.mat`

`octave-workspace`

Open the data file with `load`. If no return value *output* is requested, variables are loaded in the base workspace. Otherwise *output* will be a structure containing loaded data. See [\[load function\]](#), page 261.

`.ofig` Open the figure with `hgload`. See [\[hgload function\]](#), page 390.

`.fig, .ofig`

Load the figure

`.exe` Execute the program (on Windows systems only). No *output* value is returned.

Custom file extensions may also be handled if a function `openxxx`, where *xxx* is the extension, is found in the load path. The function must accept the file name as input.

For example, in order to load ".dat" data files in the base workspace, as is done by default for ".mat" files, one may define "opendat.m" with the following contents:

```
function retval = opendat (fname)
    evalin ("base", sprintf ("load ('%s');", fname));
endfunction
```

Other file types are opened in the appropriate external application.

```
output = perl (scriptfile)
output = perl (scriptfile, argument1, argument2, ...)
[output, status] = perl (...)
```

Invoke Perl script *scriptfile*, possibly with a list of command line arguments.

Return output in *output* and optional status in *status*. If *scriptfile* is not an absolute filename it is searched for in the current directory and then in the Octave loadpath.

See also: [system], page 913, [python], page 915.

```
output = python (scriptfile)
output = python (scriptfile, argument1, argument2, ...)
[output, status] = python (...)
```

Invoke Python script *scriptfile*, possibly with a list of command line arguments.

Return output in *output* and optional status in *status*. If *scriptfile* is not an absolute filename it is searched for in the current directory and then in the Octave loadpath.

See also: [system], page 913, [perl], page 915.

```
fid = popen (command, mode)
Start a process and create a pipe.
```

The name of the command to run is given by *command*. The argument *mode* may be

- "r" The pipe will be connected to the standard output of the process, and open for reading.
- "w" The pipe will be connected to the standard input of the process, and open for writing.

The file identifier corresponding to the input or output stream of the process is returned in *fid*.

For example:

```
fid = popen ("ls -ltr / | tail -3", "r");
while (ischar (s = fgets (fid)))
    fputs (stdout, s);
endwhile

→ drwxr-xr-x 33 root root 3072 Feb 15 13:28 etc
→ drwxr-xr-x 3 root root 1024 Feb 15 13:28 lib
→ drwxrwxrwt 15 root root 2048 Feb 17 14:53 tmp
```

See also: [popen2], page 916.

pclose (fid)

Close a file identifier that was opened by `popen`.

The function `fclose` may also be used for the same purpose.

See also: [\[fclose\]](#), page 275, [\[popen\]](#), page 915.

[in, out, pid] = popen2 (command, args)

Start a subprocess with two-way communication.

The name of the process is given by `command`, and `args` is an array or cell array of strings containing options for the command.

The file identifiers for the input and output streams of the subprocess are returned in `in` and `out`. If execution of the command is successful, `pid` contains the process ID of the subprocess. Otherwise, `pid` is `-1`.

For example:

```
[in, out, pid] = popen2 ("sort", "-r");
fputs (in, "these\nare\nsome\nstrings\n");
fclose (in);
EAGAIN = errno ("EAGAIN");
done = false;
do
  s = fgets (out);
  if (ischar (s))
    fputs (stdout, s);
  elseif (errno () == EAGAIN)
    pause (0.1);
    fclear (out);
  else
    done = true;
  endif
until (done)
fclose (out);
waitpid (pid);

  -> these
  -> strings
  -> some
  -> are
```

Note that `popen2`, unlike `popen`, will not "reap" the child process. If you don't use `waitpid` to check the child's exit status, it will linger until Octave exits.

See also: [\[popen\]](#), page 915, [\[waitpid\]](#), page 918.

```
val = EXEC_PATH ()
old_val = EXEC_PATH (new_val)
EXEC_PATH (new_val, "local")
```

Query or set the internal variable that specifies a colon separated list of directories to append to the shell PATH when executing external programs.

The initial value of `is` is taken from the environment variable `OCTAVE_EXEC_PATH`, but that value can be overridden by the command line argument `--exec-path PATH`.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [IMAGE_PATH], page 807, [OCTAVE_HOME], page 927, [OCTAVE_EXEC_HOME], page 927.

In most cases, the following functions simply decode their arguments and make the corresponding Unix system calls. For a complete example of how they can be used, look at the definition of the function `popen2`.

`[pid, msg] = fork ()`

Create a copy of the current process.

Fork can return one of the following values:

- > 0 You are in the parent process. The value returned from `fork` is the process id of the child process. You should probably arrange to wait for any child processes to exit.
- 0 You are in the child process. You can call `exec` to start another process. If that fails, you should probably call `exit`.
- < 0 The call to `fork` failed for some reason. You must take evasive action. A system dependent error message will be waiting in `msg`.

`[err, msg] = exec (file, args)`

Replace current process with a new process.

Calling `exec` without first calling `fork` will terminate your current Octave process and replace it with the program named by `file`. For example,

```
exec ("ls", "-l")
```

will run `ls` and return you to your shell prompt.

If successful, `exec` does not return. If `exec` does return, `err` will be nonzero, and `msg` will contain a system-dependent error message.

`[read_fd, write_fd, err, msg] = pipe ()`

Create a pipe and return the reading and writing ends of the pipe into `read_fd` and `write_fd` respectively.

If successful, `err` is 0 and `msg` is an empty string. Otherwise, `err` is nonzero and `msg` contains a system-dependent error message.

See also: [mkfifo], page 899.

`[fid, msg] = dup2 (old, new)`

Duplicate a file descriptor.

If successful, `fid` is greater than zero and contains the new file ID. Otherwise, `fid` is negative and `msg` contains a system-dependent error message.

See also: [fopen], page 273, [fclose], page 275, [fcntl], page 920.

`[pid, status, msg] = waitpid (pid, options)`

Wait for process *pid* to terminate.

The *pid* argument can be:

- 1 Wait for any child process.
- 0 Wait for any child process whose process group ID is equal to that of the Octave interpreter process.
- > 0 Wait for termination of the child process with ID *pid*.

The *options* argument can be a bitwise OR of zero or more of the following constants:

- 0 Wait until signal is received or a child process exits (this is the default if the *options* argument is missing).
- WNOHANG** Do not hang if status is not immediately available.

WUNTRACED

Report the status of any child processes that are stopped, and whose status has not yet been reported since they stopped.

WCONTINUE

Return if a stopped child has been resumed by delivery of **SIGCONT**. This value may not be meaningful on all systems.

If the returned value of *pid* is greater than 0, it is the process ID of the child process that exited. If an error occurs, *pid* will be less than zero and *msg* will contain a system-dependent error message. The value of *status* contains additional system-dependent information about the subprocess that exited.

See also: [\[WCONTINUE\]](#), page 918, [\[WCOREDUMP\]](#), page 918, [\[WEXITSTATUS\]](#), page 919, [\[WIFCONTINUED\]](#), page 919, [\[WIFSIGNALED\]](#), page 919, [\[WIFSTOPPED\]](#), page 919, [\[WNOHANG\]](#), page 919, [\[WSTOPSIG\]](#), page 919, [\[WTERMSIG\]](#), page 920, [\[WUNTRACED\]](#), page 920.

WCONTINUE ()

Return the numerical value of the **WCONTINUE** macro.

WCONTINUE is the option argument that may be passed to **waitpid** to indicate that it should also return if a stopped child has been resumed by delivery of a **SIGCONT** signal.

See also: [\[waitpid\]](#), page 918, [\[WNOHANG\]](#), page 919, [\[WUNTRACED\]](#), page 920.

WCOREDUMP (status)

Given *status* from a call to **waitpid**, return true if the child produced a core dump.

This function should only be employed if **WIFSIGNALED** returned true. The macro used to implement this function is not specified in POSIX.1-2001 and is not available on some Unix implementations (e.g., AIX, SunOS).

See also: [\[waitpid\]](#), page 918, [\[WIFEXITED\]](#), page 919, [\[WEXITSTATUS\]](#), page 919, [\[WIFSIGNALED\]](#), page 919, [\[WTERMSIG\]](#), page 920, [\[WIFSTOPPED\]](#), page 919, [\[WSTOPSIG\]](#), page 919, [\[WIFCONTINUED\]](#), page 919.

WEXITSTATUS (*status*)

Given *status* from a call to `waitpid`, return the exit status of the child.

This function should only be employed if `WIFEXITED` returned true.

See also: [\[waitpid\]](#), page 918, [\[WIFEXITED\]](#), page 919, [\[WIFSIGNALED\]](#), page 919, [\[WTERMSIG\]](#), page 920, [\[WCOREDUMP\]](#), page 918, [\[WIFSTOPPED\]](#), page 919, [\[WSTOPSIG\]](#), page 919, [\[WIFCONTINUED\]](#), page 919.

WIFCONTINUED (*status*)

Given *status* from a call to `waitpid`, return true if the child process was resumed by delivery of `SIGCONT`.

See also: [\[waitpid\]](#), page 918, [\[WIFEXITED\]](#), page 919, [\[WEXITSTATUS\]](#), page 919, [\[WIFSIGNALED\]](#), page 919, [\[WTERMSIG\]](#), page 920, [\[WCOREDUMP\]](#), page 918, [\[WIFSTOPPED\]](#), page 919, [\[WSTOPSIG\]](#), page 919.

WIFSIGNALED (*status*)

Given *status* from a call to `waitpid`, return true if the child process was terminated by a signal.

See also: [\[waitpid\]](#), page 918, [\[WIFEXITED\]](#), page 919, [\[WEXITSTATUS\]](#), page 919, [\[WTERMSIG\]](#), page 920, [\[WCOREDUMP\]](#), page 918, [\[WIFSTOPPED\]](#), page 919, [\[WSTOPSIG\]](#), page 919, [\[WIFCONTINUED\]](#), page 919.

WIFSTOPPED (*status*)

Given *status* from a call to `waitpid`, return true if the child process was stopped by delivery of a signal.

This is only possible if the call was done using `WUNTRACED` or when the child is being traced (see `ptrace(2)`).

See also: [\[waitpid\]](#), page 918, [\[WIFEXITED\]](#), page 919, [\[WEXITSTATUS\]](#), page 919, [\[WIFSIGNALED\]](#), page 919, [\[WTERMSIG\]](#), page 920, [\[WCOREDUMP\]](#), page 918, [\[WSTOPSIG\]](#), page 919, [\[WIFCONTINUED\]](#), page 919.

WIFEXITED (*status*)

Given *status* from a call to `waitpid`, return true if the child terminated normally.

See also: [\[waitpid\]](#), page 918, [\[WEXITSTATUS\]](#), page 919, [\[WIFSIGNALED\]](#), page 919, [\[WTERMSIG\]](#), page 920, [\[WCOREDUMP\]](#), page 918, [\[WIFSTOPPED\]](#), page 919, [\[WSTOPSIG\]](#), page 919, [\[WIFCONTINUED\]](#), page 919.

WNOHANG ()

Return the numerical value of the `WNOHANG` macro.

`WNOHANG` is the option argument that may be passed to `waitpid` to indicate that it should return its status immediately instead of waiting for a process to exit.

See also: [\[waitpid\]](#), page 918, [\[WUNTRACED\]](#), page 920, [\[WCONTINUE\]](#), page 918.

WSTOPSIG (*status*)

Given *status* from a call to `waitpid`, return the number of the signal which caused the child to stop.

This function should only be employed if `WIFSTOPPED` returned true.

See also: [\[waitpid\]](#), page 918, [\[WIFEXITED\]](#), page 919, [\[WEXITSTATUS\]](#), page 919, [\[WIFSIGNALED\]](#), page 919, [\[WTERMSIG\]](#), page 920, [\[WCOREDUMP\]](#), page 918, [\[WIFSTOPPED\]](#), page 919, [\[WIFCONTINUED\]](#), page 919.

WTERMSIG (*status*)

Given *status* from a call to `waitpid`, return the number of the signal that caused the child process to terminate.

This function should only be employed if [\[WIFSIGNALED\]](#) returned true.

See also: [\[waitpid\]](#), page 918, [\[WIFEXITED\]](#), page 919, [\[WEXITSTATUS\]](#), page 919, [\[WIFSIGNALED\]](#), page 919, [\[WCOREDUMP\]](#), page 918, [\[WIFSTOPPED\]](#), page 919, [\[WSTOPSIG\]](#), page 919, [\[WIFCONTINUED\]](#), page 919.

WUNTRACED ()

Return the numerical value of the `WUNTRACED` macro.

`WUNTRACED` is the option argument that may be passed to `waitpid` to indicate that it should also return if the child process has stopped but is not traced via the `ptrace` system call

See also: [\[waitpid\]](#), page 918, [\[WNOHANG\]](#), page 919, [\[WCONTINUE\]](#), page 918.

[*err*, *msg*] = fcntl (*fid*, *request*, *arg*)

Change the properties of the open file *fid*.

The following values may be passed as *request*:

`F_DUPFD` Return a duplicate file descriptor.

`F_GETFD` Return the file descriptor flags for *fid*.

`F_SETFD` Set the file descriptor flags for *fid*.

`F_GETFL` Return the file status flags for *fid*. The following codes may be returned (some of the flags may be undefined on some systems).

`O_RDONLY` Open for reading only.

`O_WRONLY` Open for writing only.

`O_RDWR` Open for reading and writing.

`O_APPEND` Append on each write.

`O_CREAT` Create the file if it does not exist.

`O_NONBLOCK`

Non-blocking mode.

`O_SYNC` Wait for writes to complete.

`O_ASYNC` Asynchronous I/O.

`F_SETFL` Set the file status flags for *fid* to the value specified by *arg*. The only flags that can be changed are `O_APPEND` and `O_NONBLOCK`.

If successful, *err* is 0 and *msg* is an empty string. Otherwise, *err* is nonzero and *msg* contains a system-dependent error message.

See also: [\[fopen\]](#), page 273, [\[dup2\]](#), page 917.

```
[err, msg] = kill (pid, sig)
    Send signal sig to process pid.
    If pid is positive, then signal sig is sent to pid.
    If pid is 0, then signal sig is sent to every process in the process group of the current process.
    If pid is -1, then signal sig is sent to every process except process 1.
    If pid is less than -1, then signal sig is sent to every process in the process group -pid.
    If sig is 0, then no signal is sent, but error checking is still performed.
    Return 0 if successful, otherwise return -1.
```

SIG ()

Return a structure containing Unix signal names and their defined values.

36.6 Process, Group, and User IDs

```
pgid = getpgrp ()
    Return the process group id of the current process.
```

```
pid = getpid ()
    Return the process id of the current process.
```

See also: [\[getppid\]](#), page 921.

```
pid = getppid ()
    Return the process id of the parent process.
```

See also: [\[getpid\]](#), page 921.

```
euid = geteuid ()
    Return the effective user id of the current process.
```

See also: [\[getuid\]](#), page 921.

```
uid = getuid ()
    Return the real user id of the current process.
```

See also: [\[geteuid\]](#), page 921.

```
egid = getegid ()
    Return the effective group id of the current process.
```

See also: [\[getgid\]](#), page 921.

```
gid = getgid ()
    Return the real group id of the current process.
```

See also: [\[getegid\]](#), page 921.

36.7 Environment Variables

getenv (var)

Return the value of the environment variable *var*.

For example,

```
getenv ("PATH")
```

returns a string containing the value of your path.

See also: [setenv], page 922, [unsetenv], page 922.

setenv (var, value)

setenv (var)

putenv (...)

Set the value of the environment variable *var* to *value*.

If no *value* is specified then the variable will be assigned the null string.

See also: [unsetenv], page 922, [getenv], page 922.

status = unsetenv (var)

Delete the environment variable *var*.

Return 0 if the variable was deleted, or did not exist, and -1 if an error occurred.

See also: [setenv], page 922, [getenv], page 922.

homedir = get_home_directory ()

Return the current home directory.

On most systems, this is equivalent to `getenv ("HOME")`. On Windows systems, if the environment variable `HOME` is not set then it is equivalent to `fullfile (getenv ("HOMEDRIVE"), getenv ("HOMEPATH"))`

See also: [getenv], page 922.

36.8 Current Working Directory

cd dir

cd

old_dir = cd (dir)

chdir ...

Change the current working directory to *dir*.

If *dir* is omitted, the current directory is changed to the user's home directory ("~").

For example,

```
cd ~/octave
```

changes the current working directory to `~/octave`. If the directory does not exist, an error message is printed and the working directory is not changed.

`chdir` is an alias for `cd` and can be used in all of the same calling formats.

Compatibility Note: When called with no arguments, MATLAB prints the present working directory rather than changing to the user's home directory.

See also: [pwd], page 924, [mkdir], page 898, [rmdir], page 899, [dir], page 923, [ls], page 923.

```
ls
ls filenames
ls options
ls options filenames
list = ls (...)
```

List directory contents.

The `ls` command is implemented by calling the native operating system's directory listing command—available *options* will vary from system to system.

Filenames are subject to shell expansion if they contain any wildcard characters '*', '?', '[]'. To find a literal example of a wildcard character the wildcard must be escaped using the backslash operator '\'.

If the optional output *list* is requested then `ls` returns a character array with one row for each file/directory name.

Example usage on a UNIX-like system:

```
ls -l
  - total 12
  -rw-r--r--  1 jwe  users  4488 Aug 19 04:02 foo.m
  -rw-r--r--  1 jwe  users  1315 Aug 17 23:14 bar.m
```

See also: [\[dir\]](#), page 923, [\[readdir\]](#), page 903, [\[glob\]](#), page 903, [\[what\]](#), page 138, [\[stat\]](#), page 900, [\[filesep\]](#), page 904, [\[ls_command\]](#), page 923.

```
val = ls_command ()
old_val = ls_command (new_val)
```

Query or set the shell command used by Octave's `ls` command.

See also: [\[ls\]](#), page 923.

```
dir
dir (directory)
[list] = dir (directory)
```

Display file listing for directory *directory*.

If *directory* is not specified then list the present working directory.

If a return value is requested, return a structure array with the fields

<code>name</code>	File or directory name.
<code>folder</code>	Location of file or directory
<code>date</code>	Timestamp of file modification (string value).
<code>bytes</code>	File size in bytes.
<code>isdir</code>	True if name is a directory.
<code>datenum</code>	Timestamp of file modification as serial date number (double).
<code>stathash</code>	Information structure returned from <code>stat</code> .

If *directory* is a filename, rather than a directory, then return information about the named file. *directory* may also be a list rather than a single directory or file.

`directory` is subject to shell expansion if it contains any wildcard characters ‘*’, ‘?’, ‘[]’. To find a literal example of a wildcard character the wildcard must be escaped using the backslash operator ‘\’.

Note that for symbolic links, `dir` returns information about the file that the symbolic link points to rather than the link itself. However, if the link points to a nonexistent file, `dir` returns information about the link.

See also: [\[ls\]](#), page 923, [\[readdir\]](#), page 903, [\[glob\]](#), page 903, [\[what\]](#), page 138, [\[stat\]](#), page 900, [\[lstat\]](#), page 900.

```
pwd ()  
dir = pwd ()  
Return the current working directory.
```

See also: [\[cd\]](#), page 922, [\[dir\]](#), page 923, [\[ls\]](#), page 923, [\[mkdir\]](#), page 898, [\[rmdir\]](#), page 899.

36.9 Password Database Functions

Octave’s password database functions return information in a structure with the following fields.

<code>name</code>	The user name.
<code>passwd</code>	The encrypted password, if available.
<code>uid</code>	The numeric user id.
<code>gid</code>	The numeric group id.
<code>gecos</code>	The GECOS field.
<code>dir</code>	The home directory.
<code>shell</code>	The initial shell.

In the descriptions of the following functions, this data structure is referred to as a `pw_struct`.

```
pw_struct = getpwent ()
```

Return a structure containing an entry from the password database, opening it if necessary.

Once the end of the data has been reached, `getpwent` returns 0.

See also: [\[setpwent\]](#), page 925, [\[endpwent\]](#), page 925.

```
pw_struct = getpwuid (uid).
```

Return a structure containing the first entry from the password database with the user ID `uid`.

If the user ID does not exist in the database, `getpwuid` returns 0.

See also: [\[getpwnam\]](#), page 925.

`pw_struct = getpwnam (name)`

Return a structure containing the first entry from the password database with the user name *name*.

If the user name does not exist in the database, `getpwnam` returns 0.

See also: [\[getpwuid\]](#), page 924.

`setpwent ()`

Return the internal pointer to the beginning of the password database.

See also: [\[getpwent\]](#), page 924, [\[endpwent\]](#), page 925.

`endpwent ()`

Close the password database.

See also: [\[getpwent\]](#), page 924, [\[setpwent\]](#), page 925.

36.10 Group Database Functions

Octave's group database functions return information in a structure with the following fields.

`name` The user name.

`passwd` The encrypted password, if available.

`gid` The numeric group id.

`mem` The members of the group.

In the descriptions of the following functions, this data structure is referred to as a *grp_struct*.

`grp_struct = getgrent ()`

Return an entry from the group database, opening it if necessary.

Once the end of data has been reached, `getgrent` returns 0.

See also: [\[setgrent\]](#), page 925, [\[endgrent\]](#), page 926.

`grp_struct = getgrgid (gid)`

Return the first entry from the group database with the group ID *gid*.

If the group ID does not exist in the database, `getgrgid` returns 0.

See also: [\[getgrnam\]](#), page 925.

`grp_struct = getgrnam (name)`

Return the first entry from the group database with the group name *name*.

If the group name does not exist in the database, `getgrnam` returns 0.

See also: [\[getgrgid\]](#), page 925.

`setgrent ()`

Return the internal pointer to the beginning of the group database.

See also: [\[getgrent\]](#), page 925, [\[endgrent\]](#), page 926.

endgrent ()

Close the group database.

See also: [getgrent], page 925, [setgrent], page 925.

36.11 System Information

```
computer ()
c = computer ()
[c, maxsize] = computer ()
[c, maxsize, endian] = computer ()
arch = computer ("arch")
```

Print or return a string of the form *cpu-vendor-os* that identifies the type of computer that Octave is running on.

If invoked with an output argument, the value is returned instead of printed. For example:

```
computer ()
⇒ x86_64-pc-linux-gnu
```

```
mycomp = computer ()
⇒ mycomp = x86_64-pc-linux-gnu
```

If two output arguments are requested, also return the maximum number of elements for an array. This will depend on whether Octave has been compiled with 32-bit or 64-bit index vectors.

If three output arguments are requested, also return the byte order of the current system as a character ("B" for big-endian or "L" for little-endian).

If the argument "arch" is specified, return a string indicating the architecture of the computer on which Octave is running.

See also: [isunix], page 927, [ismac], page 927, [ispc], page 927.

```
[uts, err, msg] = uname ()
```

Return system information in the structure.

For example:

```
uname ()
⇒ {
    sysname = x86_64
    nodename = segfault
    release = 2.6.15-1-amd64-k8-smp
    version = Linux
    machine = #2 SMP Thu Feb 23 04:57:49 UTC 2006
}
```

If successful, *err* is 0 and *msg* is an empty string. Otherwise, *err* is nonzero and *msg* contains a system-dependent error message.

```
nproc ()
```

```
nproc (query)
```

Return the current number of available processors.

If called with the optional argument *query*, modify how processors are counted as follows:

all total number of processors.

current processors available to the current process.

overridable

same as **current**, but overridable through the `OMP_NUM_THREADS` environment variable.

ispc ()

Return true if Octave is running on a Windows system and false otherwise.

See also: [\[isunix\]](#), page 927, [\[ismac\]](#), page 927.

isunix ()

Return true if Octave is running on a Unix-like system and false otherwise.

See also: [\[ismac\]](#), page 927, [\[ispc\]](#), page 927.

ismac ()

Return true if Octave is running on a Mac OS X system and false otherwise.

See also: [\[isunix\]](#), page 927, [\[ispc\]](#), page 927.

isieee ()

Return true if your computer *claims* to conform to the IEEE standard for floating point calculations.

No actual tests are performed.

isdeployed ()

Return true if the current program has been compiled and is running separately from the Octave interpreter and false if it is running in the Octave interpreter.

Currently, this function always returns false in Octave.

isstudent ()

Return true if running in the student edition of MATLAB.

isstudent always returns false in Octave.

See also: [\[false\]](#), page 61.

OCTAVE_HOME ()

Return the name of the top-level Octave installation directory. `OCTAVE_HOME` corresponds to the configuration variable `prefix`.

See also: [\[EXEC_PATH\]](#), page 916, [\[IMAGE_PATH\]](#), page 807, [\[OCTAVE_EXEC_HOME\]](#), page 927.

OCTAVE_EXEC_HOME ()

Return the name of the top-level Octave installation directory for architecture-dependent files. If not specified separately, the value is the same as `OCTAVE_HOME`. `OCTAVE_EXEC_HOME` corresponds to the configuration variable `exec_prefix`.

See also: [\[EXEC_PATH\]](#), page 916, [\[IMAGE_PATH\]](#), page 807, [\[OCTAVE_HOME\]](#), page 927.

matlabroot ()

Return the name of the top-level Octave installation directory.

This is an alias for the function `OCTAVE_HOME` provided for compatibility.

See also: `[OCTAVE_HOME]`, page 927.

OCTAVE_VERSION ()

Return the version number of Octave as a string.

See also: `[ver]`, page 928, `[version]`, page 928.

```
v = version ()
[v, d] = version ()
v = version (feature)
```

Get version information for Octave.

If called without input argument, the first return value *v* gives the version number of Octave as a string. The second return value *d* holds the release date as a string.

The following options can be passed for *feature*:

`"-date"` for the release date of the running build,

`"-description"` for a description of the release (always an empty string),

`"-release"` for the name of the running build (always an empty string),

`"-java"` for version information of the Java VM,

`"-fftw"` for version information for the linked FFTW,

`"-blas"` for version information for the linked BLAS,

`"-lapack"` for version information for the linked LAPACK.

`"-hgid"` the mercurial ID of the sources used to build Octave.

The variant with no input and output argument is an alias for the function `OCTAVE_VERSION` provided for compatibility.

See also: `[OCTAVE_VERSION]`, page 928, `[ver]`, page 928.

```
ver
ver Octave
ver package
v = ver (...)
```

Display a header containing the current Octave version number, license string, and operating system. The header is followed by a list of installed packages, versions, and installation directories.

Use the package name *package* or Octave to limit the listing to a desired component.

When called with an output argument, return a vector of structures describing Octave and each installed package. The structure includes the following fields.

Name	Package name.
-------------	---------------

Version Version of the package.

Release Release of the package.

Date Date of the version/release.

See also: [version], page 928, [usejava], page 995, [pkg], page 933.

compare_versions (v1, v2, operator)

Compare two version strings using the given *operator*.

This function assumes that versions *v1* and *v2* are arbitrarily long strings made of numeric and period characters possibly followed by an arbitrary string (e.g., "1.2.3", "0.3", "0.1.2+", or "1.2.3.4-test1").

The version is first split into numeric and character portions and then the parts are padded to be the same length (i.e., "1.1" would be padded to be "1.1.0" when being compared with "1.1.1", and separately, the character parts of the strings are padded with nulls).

The operator can be any logical operator from the set

- "==" equal
- "<" less than
- "<=" less than or equal to
- ">" greater than
- ">=" greater than or equal to
- "!=" not equal
- "^=" not equal

Note that version "1.1-test2" will compare as greater than "1.1-test10". Also, since the numeric part is compared first, "a" compares less than "1a" because the second string starts with a numeric part even though double ("a") is greater than double ("1").

```
license
license inuse
license inuse feature
license ("inuse")
retval = license ("inuse")
retval = license ("test", feature)
retval = license ("checkout", feature)
[retval, errmsg] = license ("checkout", feature)
```

Get license information for Octave and Octave packages.

GNU Octave is free software distributed under the GNU General Public License (GPL), and a license manager makes no sense. This function is provided only for MATLAB compatibility.

When called with no extra input arguments, it returns the Octave license, otherwise the first input defines the operation mode and must be one of the following strings:

inuse, **test**, and **checkout**. The optional *feature* argument can either be "**octave**" (core), or an Octave package.

- "inuse"** Returns a list of loaded features, i.e., octave and the list of loaded packages. If an output is requested, it returns a struct array with the fields "**feature**", and "**user**".
- "test"** Return true if the specified *feature* is installed, false otherwise.
An optional third argument "**enable**" or "**disable**" is accepted but ignored.
- "checkout"** Return true if the specified *feature* is installed, false otherwise. An optional second output will have an error message if a package is not installed.

See also: [\[pkg\]](#), page 933, [\[ver\]](#), page 928, [\[version\]](#), page 928.

getrusage ()

Return a structure containing a number of statistics about the current Octave process. Not all fields are available on all systems. If it is not possible to get CPU time statistics, the CPU time slots are set to zero. Other missing data are replaced by NaN. The list of possible fields is:

- idrss** Unshared data size.
- inblock** Number of block input operations.
- isrss** Unshared stack size.
- ixrss** Shared memory size.
- majflt** Number of major page faults.
- maxrss** Maximum data size.
- minflt** Number of minor page faults.
- msgrcv** Number of messages received.
- msgsnd** Number of messages sent.
- nivcsw** Number of involuntary context switches.
- nsignals** Number of signals received.
- nswap** Number of swaps.
- nvcs** Number of voluntary context switches.
- oublock** Number of block output operations.
- stime** A structure containing the system CPU time used. The structure has the elements **sec** (seconds) **usec** (microseconds).
- utime** A structure containing the user CPU time used. The structure has the elements **sec** (seconds) **usec** (microseconds).

```
value = winqueryreg (rootkey, subkey, valuename)
value = winqueryreg (rootkey, subkey)
names = winqueryreg ("name", rootkey, subkey)
```

Query names or value from the Windows registry.

On Windows, return the value of the registry key *subkey* from the root key *rootkey*. You can specify the name of the queried registry value with the optional argument *valuename*. Otherwise, if called with only two arguments or *valuename* is empty, then the default value of *subkey* is returned. If the registry value is of type "REG_DWORD" then *value* is of class int32. If the value is of the type "REG_SZ" or "REG_EXPAND_SZ" a string is returned.

If the first argument is "name", a cell array of strings with the names of the values at that key is returned.

The variable *rootkey* must be a string with a valid root key identifier:

```
HKCR
HKEY_CLASSES_ROOT
HKEY_CURRENT_CONFIG
HKCU
HKEY_CURRENT_USER
HKLM
HKEY_LOCAL_MACHINE
HKU
HKEY_USERS
HKEY_PERFORMANCE_DATA
```

Examples:

Get a list of value names at the key 'HKCU\Environment':

```
valuenames = winqueryreg ("name", "HKEY_CURRENT_USER", ...
                           "Environment");
```

For each *valuenames*, display the value:

```
for k = 1:numel (valuenames)
    val = winqueryreg ("HKEY_CURRENT_USER", "Environment", ...
                           valuenames{k});
    str = sprintf ("%s = %s", valuenames{k}, num2str (val));
    disp (str);
endfor
```

On non-Windows platforms this function fails with an error.

36.12 Hashing Functions

It is often necessary to find if two strings or files are identical. This might be done by comparing them character by character and looking for differences. However, this can be slow, and so comparing a hash of the string or file can be a rapid way of finding if the files differ.

Another use of the hashing function is to check for file integrity. The user can check the hash of the file against a known value and find if the file they have is the same as the one that the original hash was produced with.

Octave supplies the **hash** function to calculate hash values of strings and files, the latter in combination with the **fileread** function. The **hash** function supports the most common used cryptographic hash functions, namely MD5 and SHA-1.

hash (hfun, str)

Calculate the hash value of the string *str* using the hash function *hfun*.

The available hash functions are given in the table below.

'MD2'	Message-Digest Algorithm 2 (RFC 1319).
'MD4'	Message-Digest Algorithm 4 (RFC 1320).
'MD5'	Message-Digest Algorithm 5 (RFC 1321).
'SHA1'	Secure Hash Algorithm 1 (RFC 3174)
'SHA224'	Secure Hash Algorithm 2 (224 Bits, RFC 3874)
'SHA256'	Secure Hash Algorithm 2 (256 Bits, RFC 6234)
'SHA384'	Secure Hash Algorithm 2 (384 Bits, RFC 6234)
'SHA512'	Secure Hash Algorithm 2 (512 Bits, RFC 6234)

To calculate for example the MD5 hash value of the string "abc" the **hash** function is called as follows:

```
hash ("md5", "abc")
→ ans = 900150983cd24fb0d6963f7d28e17f72
```

For the same string, the SHA-1 hash value is calculated with:

```
hash ("sha1", "abc")
→ ans = a9993e364706816aba3e25717850c26c9cd0d89d
```

And to compute the hash value of a file, e.g., *file* = "file.txt", call **hash** in combination with the **fileread**:

```
hash ("md5", fileread (file));
```

37 Packages

Since Octave is Free Software users are encouraged to share their programs with others. To aid this sharing Octave supports the installation of extra packages. The ‘Octave Forge’ project is a community-maintained set of packages that can be downloaded and installed in Octave. At the time of writing the ‘Octave Forge’ project can be found online at <https://octave.sourceforge.io>, but since the Internet is an ever-changing place this may not be true at the time of reading. Therefore it is recommended to see the Octave website for an updated reference.

37.1 Installing and Removing Packages

Assuming a package is available in the file `image-1.0.0.tar.gz` it can be installed from the Octave prompt with the command

```
pkg install image-1.0.0.tar.gz
```

If the package is installed successfully nothing will be printed on the prompt, but if an error occurred during installation it will be reported. It is possible to install several packages at once by writing several package files after the `pkg install` command. If a different version of the package is already installed it will be removed prior to installing the new package. This makes it easy to upgrade and downgrade the version of a package, but makes it impossible to have several versions of the same package installed at once.

To see which packages are installed type

```
pkg list
+ Package Name | Version | Installation directory
+-----+-----+
+       image *| 1.0.0 | /home/jwe/octave/image-1.0.0
```

In this case only version 1.0.0 of the `image` package is installed. The '*' character next to the package name shows that the `image` package is loaded and ready for use.

It is possible to remove a package from the system using the `pkg uninstall` command like this

```
pkg uninstall image
```

If the package is removed successfully nothing will be printed in the prompt, but if an error occurred it will be reported. It should be noted that the package file used for installation is not needed for removal, and that only the package name as reported by `pkg list` should be used when removing a package. It is possible to remove several packages at once by writing several package names after the `pkg uninstall` command.

To minimize the amount of code duplication between packages it is possible that one package depends on another one. If a package depends on another, it will check if that package is installed during installation. If it is not, an error will be reported and the package will not be installed. This behavior can be disabled by passing the `-nodeps` flag to the `pkg install` command

```
pkg install -nodeps my_package_with_dependencies.tar.gz
```

Since the installed package expects its dependencies to be installed it may not function correctly. Because of this it is not recommended to disable dependency checking.

```
pkg command pkg_name
pkg command option pkg_name
[out1, ...] = pkg (command, ...)
```

Manage or query packages (groups of add-on functions) for Octave.

Different actions are available depending on the value of *command* and on return arguments.

Available commands:

'install' Install named packages. For example,

```
pkg install image-1.0.0.tar.gz
```

installs the package found in the file `image-1.0.0.tar.gz`. The file containing the package can be an url, e.g.

```
pkg install 'http://somewebsite.org/image-1.0.0.tar.gz'
```

installs the package found in the given url. This requires an internet connection and the cURL library.

Security risk: no verification of the package is performed before the installation. It has the same security issues as manually downloading the package from the given url and installing it.

No support: the GNU Octave community is not responsible for packages installed from foreign sites. For support or for reporting bugs you need to contact the maintainers of the installed package directly (see the `DESCRIPTION` file of the package)

The *option* variable can contain options that affect the manner in which a package is installed. These options can be one or more of

-nodeps The package manager will disable dependency checking. With this option it is possible to install a package even when it depends on another package which is not installed on the system. **Use this option with care.**

-local A local installation (package available only to current user) is forced, even if the user has system privileges.

-global A global installation (package available to all users) is forced, even if the user doesn't normally have system privileges.

-forge Install a package directly from the Octave Forge repository. This requires an internet connection and the cURL library.

Security risk: no verification of the package is performed before the installation. There are no signature for packages, or checksums to confirm the correct file was downloaded. It has the same security issues as manually downloading the package from the Octave Forge repository and installing it.

-verbose The package manager will print the output of all commands as they are performed.

'update' Check installed Octave Forge packages against repository and update any outdated items. This requires an internet connection and the cURL library. Usage:

```
pkg update
```

To update a single package use `pkg install -forge`

'uninstall'

Uninstall named packages. For example,

```
pkg uninstall image
```

removes the `image` package from the system. If another installed package depends on the `image` package an error will be issued. The package can be uninstalled anyway by using the `-nodeps` option.

'load'

Add named packages to the path. After loading a package it is possible to use the functions provided by the package. For example,

```
pkg load image
```

adds the `image` package to the path.

'unload'

Remove named packages from the path. After unloading a package it is no longer possible to use the functions provided by the package.

'list'

Show the list of currently installed packages. For example,

```
pkg list
```

will produce a short report with the package name, version, and installation directory for each installed package. Supply a package name to limit reporting to a particular package. For example:

```
pkg list image
```

If a single return argument is requested then `pkg` returns a cell array where each element is a structure with information on a single package.

```
installed_packages = pkg ("list")
```

If two output arguments are requested `pkg` splits the list of installed packages into those which were installed by the current user, and those which were installed by the system administrator.

```
[user_packages, system_packages] = pkg ("list")
```

The `"-forge"` option lists packages available at the Octave Forge repository. This requires an internet connection and the cURL library. For example:

```
oct_forge_pkgs = pkg ("list", "-forge")
```

'describe'

Show a short description of installed packages. With the option `"-verbose"` also list functions provided by the package. For example,

```
pkg describe -verbose
```

will describe all installed packages and the functions they provide. Display can be limited to a set of packages:

```
## describe control and signal packages
```

```
pkg describe control signal
```

If one output is requested a cell of structure containing the description and list of functions of each package is returned as output rather than printed on screen:

```
desc = pkg ("describe", "secs1d", "image")
```

If any of the requested packages is not installed, `pkg` returns an error, unless a second output is requested:

```
[desc, flag] = pkg ("describe", "secs1d", "image")
```

`flag` will take one of the values "Not installed", "Loaded", or "Not loaded" for each of the named packages.

'prefix' Set the installation prefix directory. For example,

```
pkg prefix ~/my_octave_packages
```

sets the installation prefix to `~/my_octave_packages`. Packages will be installed in this directory.

It is possible to get the current installation prefix by requesting an output argument. For example:

```
pfx = pkg ("prefix")
```

The location in which to install the architecture dependent files can be independently specified with an addition argument. For example:

```
pkg prefix ~/my_octave_packages ~/my_arch_dep_pkgs
```

'local_list'

Set the file in which to look for information on locally installed packages. Locally installed packages are those that are available only to the current user. For example:

```
pkg local_list ~/.octave_packages
```

It is possible to get the current value of `local_list` with the following

```
pkg local_list
```

'global_list'

Set the file in which to look for information on globally installed packages. Globally installed packages are those that are available to all users. For example:

```
pkg global_list /usr/share/octave/octave_packages
```

It is possible to get the current value of `global_list` with the following

```
pkg global_list
```

'build'

Build a binary form of a package or packages. The binary file produced will itself be an Octave package that can be installed normally with `pkg`. The form of the command to build a binary package is

```
pkg build builddir image-1.0.0.tar.gz ...
```

where `builddir` is the name of a directory where the temporary installation will be produced and the binary packages will be found. The options `-verbose` and `-nodeps` are respected, while all other options are ignored.

'rebuild' Rebuild the package database from the installed directories. This can be used in cases where the package database has been corrupted.

See also: [ver], page 928, [news], page 21.

37.2 Using Packages

By default installed packages are not available from the Octave prompt, but it is possible to control this using the `pkg load` and `pkg unload` commands. The functions from a package can be added to the Octave path by typing

```
pkg load package_name
```

where `package_name` is the name of the package to be added to the path.

In much the same way a package can be removed from the Octave path by typing

```
pkg unload package_name
```

37.3 Administrating Packages

On UNIX-like systems it is possible to make both per-user and system-wide installations of a package. If the user performing the installation is `root` the packages will be installed in a system-wide directory that defaults to `OCTAVE_HOME/share/octave/packages/`. If the user is not `root` the default installation directory is `~/octave/`. Packages will be installed in a subdirectory of the installation directory that will be named after the package. It is possible to change the installation directory by using the `pkg prefix` command

```
pkg prefix new_installation_directory
```

The current installation directory can be retrieved by typing

```
current_installation_directory = pkg ("prefix")
```

To function properly the package manager needs to keep some information about the installed packages. For per-user packages this information is by default stored in the file `~/.octave_packages` and for system-wide installations it is stored in `OCTAVE_HOME/share/octave/octave_packages`. The path to the per-user file can be changed with the `pkg local_list` command

```
pkg local_list /path/to/new_file
```

For system-wide installations this can be changed in the same way using the `pkg global_list` command. If these commands are called without a new path, the current path will be returned.

37.4 Creating Packages

Internally a package is simply a gzipped tar file that contains a top level directory of any given name. This directory will in the following be referred to as `package` and may contain the following files:

`package/CITATION`

This is am optional file describing instructions on how to cite the package for publication. It will be displayed verbatim by the function `citation`.

package/COPYING

This is a required file containing the license of the package. No restrictions is made on the license in general. If however the package contains dynamically linked functions the license must be compatible with the GNU General Public License.

package/DESCRIPTION

This is a required file containing information about the package. See [Section 37.4.1 \[The DESCRIPTION File\], page 939](#), for details on this file.

package/ChangeLog

This is an optional file describing all the changes made to the package source files.

package/INDEX

This is an optional file describing the functions provided by the package. If this file is not given then one will be created automatically from the functions in the package and the `Categories` keyword in the `DESCRIPTION` file. See [Section 37.4.2 \[The INDEX File\], page 941](#), for details on this file.

package/NEWS

This is an optional file describing all user-visible changes worth mentioning. As this file increases in size, old entries can be moved into `package/ONEWS`.

package/ONEWS

This is an optional file describing old entries from the `NEWS` file.

package/PKG_ADD

An optional file that includes commands that are run when the package is added to the users path. Note that `PKG_ADD` directives in the source code of the package will also be added to this file by the Octave package manager. Note that symbolic links are to be avoided in packages, as symbolic links do not exist on some file systems, and so a typical use for this file is the replacement of the symbolic link

```
ln -s foo.oct bar.oct  
with an autoload directive like  
autoload ('bar', which ('foo'));
```

See [Section 37.4.3 \[PKG_ADD and PKG_DEL Directives\], page 942](#), for details on `PKG_ADD` directives.

package/PKG_DEL

An optional file that includes commands that are run when the package is removed from the users path. Note that `PKG_DEL` directives in the source code of the package will also be added to this file by the Octave package manager. See [Section 37.4.3 \[PKG_ADD and PKG_DEL Directives\], page 942](#), for details on `PKG_DEL` directives.

package/pre_install.m

This is an optional function that is run prior to the installation of a package. This function is called with a single argument, a struct with fields names after

the data in the `DESCRIPTION`, and the paths where the package functions will be installed.

`package/post_install.m`

This is an optional function that is run after the installation of a package. This function is called with a single argument, a struct with fields names after the data in the `DESCRIPTION`, and the paths where the package functions were installed.

`package/on_uninstall.m`

This is an optional function that is run prior to the removal of a package. This function is called with a single argument, a struct with fields names after the data in the `DESCRIPTION`, the paths where the package functions are installed, and whether the package is currently loaded.

Besides the above mentioned files, a package can also contain one or more of the following directories:

`package/inst`

An optional directory containing any files that are directly installed by the package. Typically this will include any `m`-files.

`package/src`

An optional directory containing code that must be built prior to the packages installation. The Octave package manager will execute `./configure` in this directory if this script exists, and will then call `make` if a file `Makefile` exists in this directory. `make install` will however not be called. The environment variables `MKOCTFILE`, `OCTAVE_CONFIG`, and `OCTAVE` will be set to the full paths of the programs `mkoctfile`, `octave-config`, and `octave`, respectively, of the correct version when `configure` and `make` are called. If a file called `FILES` exists all files listed there will be copied to the `inst` directory, so they also will be installed. If the `FILES` file doesn't exist, `src/*.m` and `src/*.oct` will be copied to the `inst` directory.

`package/doc`

An optional directory containing documentation for the package. The files in this directory will be directly installed in a sub-directory of the installed package for future reference.

`package/bin`

An optional directory containing files that will be added to the Octave `EXEC_PATH` when the package is loaded. This might contain external scripts, etc., called by functions within the package.

37.4.1 The `DESCRIPTION` File

The `DESCRIPTION` file contains various information about the package, such as its name, author, and version. This file has a very simple format

- Lines starting with '#' are comments.
- Lines starting with a blank character are continuations from the previous line.
- Everything else is of the form `NameOfOption: ValueOfOption`.

The following is a simple example of a DESCRIPTION file

```
Name: The name of my package
Version: 1.0.0
Date: 2007-18-04
Author: The name (and possibly email) of the package author.
Maintainer: The name (and possibly email) of the current
            package maintainer.
Title: The title of the package
Description: A short description of the package. If this
            description gets too long for one line it can continue
            on the next by adding a space to the beginning of the
            following lines.
License: GPLv3+
```

The package manager currently recognizes the following keywords

Name	Name of the package.
Version	Version of the package. A package version must be 3 numbers separated by dots.
Date	Date of last update.
Author	Original author of the package.
Maintainer	Maintainer of the package.
Title	A one line description of the package.
Description	A one paragraph description of the package.
Categories	Optional keyword describing the package (if no INDEX file is given this is mandatory).
Problems	Optional list of known problems.
Url	Optional list of homepages related to the package.
Depends	A list of other Octave packages that this package depends on. This can include dependencies on particular versions, with a format Depends: package (>= 1.0.0) Possible operators are <, <=, ==, >= or >. If the part of the dependency in () is missing, any version of the package is acceptable. Multiple dependencies can be defined as a comma separated list.
License	An optional short description of the used license (e.g., GPL version 3 or newer). This is optional since the file COPYING is mandatory.
SystemRequirements	These are the external install dependencies of the package and are not checked by the package manager. This is here as a hint to the distribution packager. They follow the same conventions as the Depends keyword.

BuildRequires

These are the external build dependencies of the package and are not checked by the package manager. This is here as a hint to the distribution packager. They follow the same conventions as the `Depends` keyword. Note that in general, packaging systems such as `rpm` or `deb` autoprobe the install dependencies from the build dependencies, and therefore a `BuildRequires` dependency usually removes the need for a `SystemRequirements` dependency.

The developer is free to add additional arguments to the `DESCRIPTION` file for their own purposes. One further detail to aid the packager is that the `SystemRequirements` and `BuildRequires` keywords can have a distribution dependent section, and the automatic build process will use these. An example of the format of this is

```
BuildRequires: libtermcap-devel [Mandriva] libtermcap2-devel
```

where the first package name will be used as a default and if the RPMs are built on a Mandriva distribution, then the second package name will be used instead.

37.4.2 The INDEX File

The optional `INDEX` file provides a categorical view of the functions in the package. This file has a very simple format

- Lines beginning with '#' are comments.
- The first non-comment line should look like this


```
toolbox >> Toolbox name
```
- Lines beginning with an alphabetical character indicates a new category of functions.
- Lines starting with a white space character indicate that the function names on the line belong to the last mentioned category.

The format can be summarized with the following example:

```
# A comment
toolbox >> Toolbox name
Category Name 1
    function1 function2 function3
    function4
Category Name 2
    function2 function5
```

If you wish to refer to a function that users might expect to find in your package but is not there, providing a work around or pointing out that the function is available elsewhere, you can use:

```
fn = workaround description
```

This workaround description will not appear when listing functions in the package with `pkg describe` but they will be published in the HTML documentation online. Workaround descriptions can use any HTML markup, but keep in mind that it will be enclosed in a bold-italic environment. For the special case of:

```
fn = use <code>alternate expression</code>
```

the bold-italic is automatically suppressed. You will need to use `<code>` even in references:

```
fn = use <a href="someothersite.html"><code>fn</code></a>
```

Sometimes functions are only partially compatible, in which case you can list the non-compatible cases separately. To refer to another function in the package, use `<f>fn</f>`. For example:

```
eig (a, b) = use <f>qz</f>
```

Since sites may have many missing functions, you can define a macro rather than typing the same link over and again.

```
$id = expansion
```

defines the macro `$id` anywhere in the description and it will be expanded. For example:

```
$TSA = see <a href="link_to_spctools">SPC Tools</a>
arcov = $TSA <code>armcv</code>
```

`id` is any string of letters, numbers and `_`.

37.4.3 PKG_ADD and PKG_DEL Directives

If the package contains files called `PKG_ADD` or `PKG_DEL` the commands in these files will be executed when the package is added or removed from the users path. In some situations such files are a bit cumbersome to maintain, so the package manager supports automatic creation of such files. If a source file in the package contains a `PKG_ADD` or `PKG_DEL` directive they will be added to either the `PKG_ADD` or `PKG_DEL` files.

In `m`-files a `PKG_ADD` directive looks like this

```
## PKG_ADD: some_octave_command
```

Such lines should be added before the `function` keyword. In `C++` files a `PKG_ADD` directive looks like this

```
// PKG_ADD: some_octave_command
```

In both cases `some_octave_command` should be replaced by the command that should be placed in the `PKG_ADD` file. `PKG_DEL` directives work in the same way, except the `PKG_ADD` keyword is replaced with `PKG_DEL` and the commands get added to the `PKG_DEL` file.

37.4.4 Missing Components

If a package relies on a component, such as another Octave package, that may not be present it may be useful to install a function which informs users what to do when a particular component is missing. The function must be written by the package maintainer and registered with Octave using `missing_component_hook`.

```
val = missing_component_hook ()
old_val = missing_component_hook (new_val)
missing_component_hook (new_val, "local")
```

Query or set the internal variable that specifies the function to call when a component of Octave is missing.

This can be useful for packagers that may split the Octave installation into multiple sub-packages, for example, to provide a hint to users for how to install the missing components.

When called from inside a function with the `"local"` option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

The hook function is expected to be of the form

fcn (*component*)

Octave will call *fcn* with the name of the function that requires the component and a string describing the missing component. The hook function should return an error message to be displayed.

See also: [missing_function_hook], page 1036.

Appendix A External Code Interface

"The sum of human wisdom is not contained in any one language"

— Ezra Pound

Octave is a fantastic language for solving many problems in science and engineering. However, it is not the only computer language and there are times when you may want to use code written in other languages. Good reasons for doing so include: 1) not re-inventing the wheel; existing function libraries which have been thoroughly tested and debugged or large scale simulation codebases are a good example, 2) accessing unique capabilities of a different language; for example the well-known regular expression functions of Perl (but don't do that because `regexp` already exists in Octave).

Performance should generally **not** be a reason for using compiled extensions. Although compiled extensions can run faster, particularly if they replace a loop in Octave code, this is almost never the best path to take. First, there are many techniques to speed up Octave performance while remaining within the language. Second, Octave is a high-level language that makes it easy to perform common mathematical tasks. Giving that up means shifting the focus from solving the real problem to solving a computer programming problem. It means returning to low-level constructs such as pointers, memory management, mathematical overflow/underflow, etc. Because of the low level nature, and the fact that the compiled code is executed outside of Octave, there is the very real possibility of crashing the interpreter and losing work.

Before going further, you should first determine if you really need to bother writing code outside of Octave.

- Can I get the same functionality using the Octave scripting language alone?

Even when a function already exists outside the language, it may be better to simply reproduce the behavior in an m-file rather than attempt to interface to the outside code.

- Is the code thoroughly optimized for Octave?

If performance is an issue you should always start with the in-language techniques for getting better performance. Chief among these is vectorization (see [Chapter 19 \[Vectorization and Faster Code Execution\], page 583](#)) which not only makes the code concise and more understandable but improves performance (10X-100X). If loops must be used, make sure that the allocation of space for variables takes place outside the loops using an assignment to a matrix of the right size, or zeros.

- Does the code make as much use as possible of existing built-in library routines?

These routines are highly optimized and many do not carry the overhead of being interpreted.

- Does writing a dynamically linked function represent a useful investment of your time, relative to staying in Octave?

It will take time to learn Octave's interface for external code and there will inevitably be issues with tools such as compilers.

With that said, Octave offers a versatile interface for including chunks of compiled code as dynamically linked extensions. These dynamically linked functions can be called from the interpreter in the same manner as any ordinary function. The interface is bi-directional and

external code can call Octave functions (like `plot`) which otherwise might be very difficult to develop.

The interface is centered around supporting the languages C++, C, and Fortran. Octave itself is written in C++ and can call external C++/C code through its native oct-file interface. The C language is also supported through the mex-file interface for compatibility with MATLAB. Fortran code is easiest to reach through the oct-file interface.

Because many other languages provide C or C++ APIs it is relatively simple to build bridges between Octave and other languages. This is also a way to bridge to hardware resources which often have device drivers written in C.

A.1 Oct-Files

A.1.1 Getting Started with Oct-Files

Oct-files are pieces of C++ code that have been compiled with the Octave API into a dynamically loadable object. They take their name from the file which contains the object which has the extension `.oct`.

Finding a C++ compiler, using the correct switches, adding the right include paths for header files, etc. is a difficult task. Octave automates this by providing the `mkoctfile` command with which to build oct-files. The command is available from within Octave or at the shell command line.

```
mkoctfile [-options] file ...
[output, status] = mkoctfile (...)
```

The `mkoctfile` function compiles source code written in C, C++, or Fortran. Depending on the options used with `mkoctfile`, the compiled code can be called within Octave or can be used as a stand-alone application.

`mkoctfile` can be called from the shell prompt or from the Octave prompt. Calling it from the Octave prompt simply delegates the call to the shell prompt. The output is stored in the `output` variable and the exit status in the `status` variable.

`mkoctfile` accepts the following options, all of which are optional except for the filename of the code you wish to compile:

- ‘-I DIR’ Add the include directory DIR to compile commands.
- ‘-D DEF’ Add the definition DEF to the compiler call.
- ‘-l LIB’ Add the library LIB to the link command.
- ‘-L DIR’ Add the library directory DIR to the link command.
- ‘-M’
- ‘--depend’ Generate dependency files (.d) for C and C++ source files.
- ‘-R DIR’ Add the run-time path to the link command.
- ‘-Wl,...’ Pass options to the linker like “-Wl,-rpath=...”. The quotes are needed since commas are interpreted as command separators.
- ‘-W...’ Pass options to the assembler like “-Wa,OPTION”.

‘-c’ Compile but do not link.
 ‘-g’ Enable debugging options for compilers.
 ‘-o FILE’
 ‘--output FILE’
 Output filename. Default extension is .oct (or .mex if ‘--mex’ is specified) unless linking a stand-alone executable.

‘-p VAR’
 ‘--print VAR’
 Print configuration variable VAR. There are three categories of variables:
 Octave configuration variables that users may override with environment variables. These are used in commands that `mkoctfile` executes.

ALL_CFLAGS	LAPACK_LIBS
ALL_CXXFLAGS	LDFLAGS
ALL_FFLAGS	LD_CXX
ALL_LDFLAGS	LD_STATIC_FLAG
BLAS_LIBS	LFLAGS
CC	LIBDIR
CFLAGS	LIBOCTAVE
CPICFLAG	LIBOCTINTERP
CPPFLAGS	OCTAVE_LINK_OPTS
CXX	OCTINCLUDEDIR
CXXFLAGS	OCTAVE_LIBS
CXXPICFLAG	OCTAVE_LINK_DEPS
DL_LD	OCTLIBDIR
DL_LDFLAGS	OCT_LINK_DEPS
F77	OCT_LINK_OPTS
F77_INTEGER8_FLAG	RDYNAMIC_FLAG
FFLAGS	SPECIAL_MATH_LIB
FPICFLAG	XTRA_CFLAGS
INCFLAGS	XTRA_CXXFLAGS
INCLUDEDIR	

Octave configuration variables as above, but currently unused by `mkoctfile`.

AR
DEPEND_EXTRA_SED_PATTERN
DEPEND_FLAGS
FFTW3F_LDFLAGS
FFTW3F_LIBS
FFTW3_LDFLAGS
FFTW3_LIBS
FFTW_LIBS
FLIBS
LIBS
RANLIB
READLINE_LIBS

Octave configuration variables that are provided for informational purposes only. Except for ‘OCTAVE_HOME’ and ‘OCTAVE_EXEC_HOME’, users may not override these variables.

If OCTAVE_HOME or OCTAVE_EXEC_HOME are set in the environment, then other variables are adjusted accordingly with OCTAVE_HOME or OCTAVE_EXEC_HOME substituted for the original value of the directory specified by the --prefix or --exec-prefix options that were used when Octave was configured.

API_VERSION	LOCALFCNFILEDIR
ARCHLIBDIR	LOCALOCTFILEDIR
BINDIR	LOCALSTARTUPFILEDIR
CANONICAL_HOST_TYPE	LOCALVERARCHLIBDIR
DATADIR	LOCALVERFCNFILEDIR
DATAROOTDIR	LOCALVEROCTFILEDIR
DEFAULT_PAGER	MAN1DIR
EXEC_PREFIX	MAN1EXT
EXEEXT	MANDIR
FCNFILEDIR	OCTAVE_EXEC_HOME
IMAGEDIR	OCTAVE_HOME
INFODIR	OCTAVE_VERSION
INFOFILE	OCTDATADIR
LIBEXECDIR	OCTDOCDIR
LOCALAPIARCHLIBDIR	OCTFILEDIR
LOCALAPIFCNFILEDIR	OCTFONTSDIR
LOCALAPOCTFILEDIR	STARTUPFILEDIR
LOCALARCHLIBDIR	

‘--link-stand-alone’

Link a stand-alone executable file.

‘--mex’ Assume we are creating a MEX file. Set the default output extension to “.mex”.

‘-s’

‘--strip’ Strip the output file.

‘-v’

‘--verbose’

Echo commands as they are executed.

‘file’

The file to compile or link. Recognized file types are:

```

.c      C source
.cc     C++ source
.cp     C++ source
.cpp    C++ source
.CPP    C++ source
.cxx    C++ source
.c++    C++ source
.C      C++ source
.f      Fortran source (fixed form)
.F      Fortran source (fixed form)
.f90   Fortran source (free form)
.F90   Fortran source (free form)
.o     object file
.a     library file

```

Consider the following short example which introduces the basics of writing a C++ function that can be linked to Octave.

```

#include <octave/oct.h>

DEFUN_DLD (helloworld, args, nargout,
           "Hello World Help String")
{
  octave_stdout << "Hello World has "
                << args.length () << " input arguments and "
                << nargout << " output arguments.\n";

  // Return empty matrices for any outputs
  octave_value_list retval (nargout);
  for (int i = 0; i < nargout; i++)
    retval(i) = octave_value (Matrix ());

  return retval;
}

```

The first critical line is `#include <octave/oct.h>` which makes available most of the definitions necessary for a C++ oct-file. Note that `octave/oct.h` is a C++ header and cannot be directly `#include`'ed in a C source file, nor any other language.

Included by `oct.h` is a definition for the macro `DEFUN_DLD` which creates a dynamically loaded function. This macro takes four arguments:

1. The function name as it will be seen in Octave,
2. The list of arguments to the function of type `octave_value_list`,
3. The number of output arguments, which can be—and often is—omitted if not used, and
4. The string to use for the help text of the function.

The return type of functions defined with `DEFUN_DLD` is always `octave_value_list`.

There are a couple of important considerations in the choice of function name. First, it must be a valid Octave function name and so must be a sequence of letters, digits, and underscores not starting with a digit. Second, as Octave uses the function name to define the filename it attempts to find the function in, the function name in the `DEFUN_DLD` macro must match the filename of the oct-file. Therefore, the above function should be in a file `helloworld.cc`, and would be compiled to an oct-file using the command

```
mkoctfile helloworld.cc
```

This will create a file called `helloworld.oct` that is the compiled version of the function. It should be noted that it is perfectly acceptable to have more than one `DEFUN_DLD` function in a source file. However, there must either be a symbolic link to the oct-file for each of the functions defined in the source code with the `DEFUN_DLD` macro or the `autoload` (Section 11.9 [Function Files], page 193) function should be used.

The rest of the function shows how to find the number of input arguments, how to print through the Octave pager, and how to return from the function. After compiling this function as above, an example of its use is

```
helloworld (1, 2, 3)
→ Hello World has 3 input arguments and 0 output arguments.
```

Subsequent sections show how to use specific classes from Octave's core internals. Base classes like `dMatrix` (a matrix of double values) are found in the directory `liboctave/array`. The definitive reference for how to use a particular class is the header file itself. However, it is often enough simply to study the examples in the manual in order to be able to use a class.

A.1.2 Matrices and Arrays in Oct-Files

Octave supports a number of different array and matrix classes, the majority of which are based on the `Array` class. The exception are the sparse matrix types discussed separately below. There are three basic matrix types:

`Matrix` A double precision matrix class defined in `dMatrix.h`

`ComplexMatrix`
A complex matrix class defined in `CMatrix.h`

`BoolMatrix`
A boolean matrix class defined in `boolMatrix.h`

These are the basic two-dimensional matrix types of Octave. In addition there are a number of multi-dimensional array types including

`NDArray` A double precision array class defined in `dNDArray.h`

`ComplexNDarray`
A complex array class defined in `CNDArray.h`

`boolNDArray`
A boolean array class defined in `boolNDArray.h`

```

int8NDArray
int16NDArray
int32NDArray
int64NDArray
    8, 16, 32, and 64-bit signed array classes defined in int8NDArray.h,
    int16NDArray.h, etc.

uint8NDArray
uint16NDArray
uint32NDArray
uint64NDArray
    8, 16, 32, and 64-bit unsigned array classes defined in uint8NDArray.h,
    uint16NDArray.h, etc.

```

There are several basic ways of constructing matrices or multi-dimensional arrays. Using the class `Matrix` as an example one can

- Create an empty matrix or array with the empty constructor. For example:

```
Matrix a;
```

This can be used for all matrix and array types.

- Define the dimensions of the matrix or array with a `dim_vector` which has the same characteristics as the vector returned from `size`. For example:

```
dim_vector dv (2, 3); // 2 rows, 3 columns
Matrix a (dv);
```

This can be used for all matrix and array types.

- Define the number of rows and columns in the matrix. For example:

```
Matrix a (2, 2)
```

This constructor can **only** be used with matrix types.

These types all share a number of basic methods and operators. Many bear a resemblance to functions that exist in the interpreter. A selection of useful methods include

<code>T& operator () (octave_idx_type)</code>	[Method]
<code>T& elem (octave_idx_type)</code>	[Method]

The `()` operator or `elem` method allow the values of the matrix or array to be read or set. These methods take a single argument, which is of type `octave_idx_type`, that is the index into the matrix or array. Additionally, the matrix type allows two argument versions of the `()` operator and `elem` method, giving the row and column index of the value to get or set.

Note that these functions do significant error checking and so in some circumstances the user might prefer to access the data of the array or matrix directly through the `fortran_vec` method discussed below.

<code>octave_idx_type numel (void) const</code>	[Method]
The total number of elements in the matrix or array.	

<code>size_t byte_size (void) const</code>	[Method]
The number of bytes used to store the matrix or array.	

`dim_vector dims (void) const` [Method]

The dimensions of the matrix or array in value of type `dim_vector`.

`int ndims (void) const` [Method]

The number of dimensions of the matrix or array. Matrices are always 2-D, but arrays can be N-dimensional.

`void resize (const dim_vector&)` [Method]

`void resize (nrows, ncols)` [Method]

A method taking either an argument of type `dim_vector`, or, in the case of a matrix, two arguments of type `octave_idx_type` defining the number of rows and columns in the matrix.

`T* fortran_vec (void)` [Method]

This method returns a pointer to the underlying data of the matrix or array so that it can be manipulated directly, either within Octave or by an external library.

Operators such as `+`, `-`, or `*` can be used on the majority of the matrix and array types. In addition there are a number of methods that are of interest only for matrices such as `transpose`, `hermitian`, `solve`, etc.

The typical way to extract a matrix or array from the input arguments of `DEFUN_DLD` function is as follows

```
#include <octave/oct.h>

DEFUN_DLD (addtwomatrices, args, , "Add A to B")
{
  if (args.length () != 2)
    print_usage ();

  NDArray A = args(0).array_value ();
  NDArray B = args(1).array_value ();

  return octave_value (A + B);
}
```

To avoid segmentation faults causing Octave to abort, this function explicitly checks that there are sufficient arguments available before accessing these arguments. It then obtains two multi-dimensional arrays of type `NDArray` and adds these together. Note that the `array_value` method is called without using the `is_matrix_type` method. If an error occurs when attempting to extract the value, Octave will print a message and throw an exception. The reason to prefer this coding structure is that the arguments might be a type which is not an `NDArray`, but for which it would make sense to convert them to one. The `array_value` method allows this conversion to be performed transparently when possible. If you need to catch errors like this, and perform some kind of cleanup or other operation, you can catch the `octave_execution_error` exception.

`A + B`, operating on two `NDArray` objects returns an `NDArray`, which is cast to an `octave_value` on the return from the function. An example of the use of this demonstration function is

```
addtwomatrices (ones (2, 2), eye (2, 2))
⇒ 2 1
    1 2
```

A list of the basic `Matrix` and `Array` types, the methods to extract these from an `octave_value`, and the associated header file is listed below.

Type	Function	Source Code
RowVector	row_vector_value	dRowVector.h
ComplexRowVector	complex_row_vector_value	CRowVector.h
ColumnVector	column_vector_value	dColVector.h
ComplexColumnVector	complex_column_vector_value	CColVector.h
Matrix	matrix_value	dMatrix.h
ComplexMatrix	complex_matrix_value	CMatrix.h
boolMatrix	bool_matrix_value	boolMatrix.h
charMatrix	char_matrix_value	chMatrix.h
NDArray	array_value	dNDArray.h
ComplexNDArray	complex_array_value	CNDArray.h
boolNDArray	bool_array_value	boolNDArray.h
charNDArray	char_array_value	charNDArray.h
int8NDArray	int8_array_value	int8NDArray.h
int16NDArray	int16_array_value	int16NDArray.h
int32NDArray	int32_array_value	int32NDArray.h
int64NDArray	int64_array_value	int64NDArray.h
uint8NDArray	uint8_array_value	uint8NDArray.h
uint16NDArray	uint16_array_value	uint16NDArray.h
uint32NDArray	uint32_array_value	uint32NDArray.h
uint64NDArray	uint64_array_value	uint64NDArray.h

A.1.3 Character Strings in Oct-Files

A character string in Octave is just a special `Array` class. Consider the example:

```
#include <octave/oct.h>

DEFUN_DLD (stringdemo, args, , "String Demo")
{
  if (args.length () != 1)
    print_usage ();

  octave_value_list retval;

  charMatrix ch = args(0).char_matrix_value ();

  retval(1) = octave_value (ch, '\'''); // Single Quote String

  octave_idx_type nr = ch.rows ();

  for (octave_idx_type i = 0; i < nr / 2; i++)
```

```

{
  std::string tmp = ch.row_as_string (i);

  ch.insert (ch.row_as_string (nr-i-1).c_str (), i, 0);
  ch.insert (tmp.c_str (), nr-i-1, 0);
}

retval(0) = octave_value (ch, "'"); // Double Quote String

return retval;
}

```

An example of the use of this function is

```

s0 = ["First String"; "Second String"];
[s1,s2] = stringdemo (s0)
⇒ s1 = Second String
      First String

⇒ s2 = First String
      Second String

typeinfo (s2)
⇒ sq_string
typeinfo (s1)
⇒ string

```

One additional complication of strings in Octave is the difference between single quoted and double quoted strings. To find out if an `octave_value` contains a single or double quoted string use one of the predicate tests shown below.

```

if (args(0).is_sq_string ())
  octave_stdout << "First argument is a single quoted string\n";
else if (args(0).is_dq_string ())
  octave_stdout << "First argument is a double quoted string\n";

```

Note, however, that both types of strings are represented by the `charNDArray` type, and so when assigning to an `octave_value`, the type of string should be specified. For example:

```

octave_value_list retval;
charNDArray ch;
...
// Create single quoted string
retval(1) = octave_value (ch); // default constructor is sq_string
OR
retval(1) = octave_value (ch, '\''); // explicitly create sq_string

// Create a double quoted string
retval(0) = octave_value (ch, "'");

```

A.1.4 Cell Arrays in Oct-Files

Octave's cell type is also available from within oct-files. A cell array is just an `Array` of `octave_values`, and thus each element of the cell array can be treated like any other `octave_value`. A simple example is

```
#include <octave/oct.h>
#include <octave/Cell.h>

DEFUN_DLD (celldemo, args, , "Cell Demo")
{
  if (args.length () != 1)
    print_usage ();

  Cell c = args(0).cell_value ();

  octave_value_list retval;
  retval.resize (c.numel ());      // faster code by pre-declaring size

  for (octave_idx_type i = 0; i < c.numel (); i++)
  {
    retval(i) = c(i);           // using operator syntax
    //retval(i) = c.elem (i);   // using method syntax
  }

  return retval;
}
```

Note that cell arrays are used less often in standard oct-files and so the `Cell.h` header file must be explicitly included. The rest of the example extracts the `octave_values` one by one from the cell array and returns them as individual output arguments. For example:

```
[b1, b2, b3] = celldemo ({1, [1, 2], "test"})
⇒
b1 = 1
b2 =
1 2
b3 = test
```

A.1.5 Structures in Oct-Files

A structure in Octave is a map between a number of fields represented and their values. The Standard Template Library `map` class is used, with the pair consisting of a `std::string` and an Octave `Cell` variable.

A simple example demonstrating the use of structures within oct-files is

```
#include <octave/oct.h>
#include <octave/ov-struct.h>
```

```

DEFUN_DLD (structdemo, args, , "Struct Demo")
{
  if (args.length () != 2)
    print_usage ();

  if (! args(0).isstruct ())
    error ("structdemo: ARG1 must be a struct");

  octave_scalar_map arg0 = args(0).scalar_map_value ();
//octave_map arg0 = args(0).map_value ();

  if (! args(1).is_string ())
    error ("structdemo: ARG2 must be a character string");

  std::string arg1 = args(1).string_value ();

  octave_value tmp = arg0.contents (arg1);
//octave_value tmp = arg0.contents (arg1)(0);

  if (! tmp.is_defined ())
    error ("structdemo: struct does not have a field named '%s'\n",
           arg1.c_str ());

  octave_scalar_map st;

  st.assign ("selected", tmp);

  return octave_value (st);
}

```

An example of its use is

```

x.a = 1; x.b = "test"; x.c = [1, 2];
structdemo (x, "b")
⇒ selected = test

```

The example above specifically uses the `octave_scalar_map` class which is for representing a single struct. For structure arrays, the `octave_map` class is used instead. The commented code shows how the demo could be modified to handle a structure array. In that case, the `contents` method returns a `Cell` which may have more than one element. Therefore, to obtain the underlying `octave_value` in the single struct example we would write

```
octave_value tmp = arg0.contents (arg1)(0);
```

where the trailing `(0)` is the `()` operator on the `Cell` object. If this were a true structure array with multiple elements we could iterate over the elements using the `()` operator.

Structures are a relatively complex data container and there are more functions available in `oct-map.h` which make coding with them easier than relying on just `contents`.

A.1.6 Sparse Matrices in Oct-Files

There are three classes of sparse objects that are of interest to the user.

`SparseMatrix`

A double precision sparse matrix class

`SparseComplexMatrix`

A complex sparse matrix class

`SparseBoolMatrix`

A boolean sparse matrix class

All of these classes inherit from the `Sparse<T>` template class, and so all have similar capabilities and usage. The `Sparse<T>` class was based on Octave's `Array<T>` class and users familiar with Octave's `Array` classes will be comfortable with the use of the sparse classes.

The sparse classes will not be entirely described in this section, due to their similarity with the existing `Array` classes. However, there are a few differences due the nature of sparse objects, and these will be described. First, although it is fundamentally possible to have N-dimensional sparse objects, the Octave sparse classes do not allow them at this time; All instances of the sparse classes **must** be 2-dimensional. This means that `SparseMatrix` is actually more similar to Octave's `Matrix` class than it is to the `NDArray` class.

A.1.6.1 Array and Sparse Class Differences

The number of elements in a sparse matrix is considered to be the number of nonzero elements, rather than the product of the dimensions. Therefore,

```
SparseMatrix sm;
...
int nnz = sm.nelem();
```

returns the number of nonzero elements (like the interpreter function `nnz`). If the user really requires the number of elements in the matrix, including the nonzero elements, they should use `numel` rather than `nelem`. Note that for very large matrices, where the product of the two dimensions is larger than the representation of an unsigned int, `numel` can overflow. An example is `speye (1e6)` which will create a matrix with a million rows and columns, but only a million nonzero elements. In this case, the number of rows multiplied by the number of columns is more than two hundred times the maximum value that can be represented by an unsigned 32-bit int. The use of `numel` should, therefore, be avoided unless it is known that it will not overflow.

Extreme care is also required when using the `elem` method or the `()` operator which perform essentially the same function. The reason is that if a sparse object is non-const, then Octave will assume that a request for a zero element in a sparse matrix is in fact a request to create this element so it can be filled. Therefore, a piece of code like

```
SparseMatrix sm;
...
for (int j = 0; j < nc; j++)
    for (int i = 0; i < nr; i++)
        std::cerr << "(" << i << "," << j << "): " << sm(i,j) << "\n";
```

is a great way of turning a sparse matrix into a dense one, and a very slow way at that since it reallocates the sparse object for each zero element in the matrix.

A simple way of preventing the above from happening is to create a temporary constant version of the sparse matrix. Note that only the container for the sparse matrix will be copied, while the actual representation of the data will be shared between the two versions of the sparse matrix; This is not a costly operation. The example above, re-written to prevent sparse-to-dense conversion, is

```
SparseMatrix sm;
...
const SparseMatrix tmp (sm);
for (int j = 0; j < nc; j++)
    for (int i = 0; i < nr; i++)
        std::cerr << "(" << i << "," << j << "): " << tmp(i,j) << "\n";
```

Finally, because the sparse types aren't represented by a contiguous block of memory, the `fortran_vec` method of `Array<T>` is not available. It is, however, replaced by three separate methods `ridx`, `cidx`, and `data`, that access the raw compressed column format that Octave sparse matrices are stored in. These methods can be used in a manner similar to `elem` to allow the matrix to be accessed or filled. However, it is up to the user to respect the sparse matrix compressed column format or the matrix will become corrupted.

A.1.6.2 Creating Sparse Matrices in Oct-Files

There are two useful strategies for creating a sparse matrix. The first is to create three vectors representing the row index, column index, and data values, and from these create the matrix. The second alternative is to create a sparse matrix with the appropriate amount of space, and then fill in the values. Both techniques have their advantages and disadvantages.

Below is an example of creating a small sparse matrix using the first technique

```
int nz, nr, nc;
nz = 4, nr = 3, nc = 4;

ColumnVector ridx (nz);
ColumnVector cidx (nz);
ColumnVector data (nz);

ridx(0) = 1; cidx(0) = 1; data(0) = 1;
ridx(1) = 2; cidx(1) = 2; data(1) = 2;
ridx(2) = 2; cidx(2) = 4; data(2) = 3;
ridx(3) = 3; cidx(3) = 4; data(3) = 4;
SparseMatrix sm (data, ridx, cidx, nr, nc);
```

which creates the matrix given in section [Section 22.1.1 \[Storage of Sparse Matrices\]](#), [page 617](#). Note that the compressed matrix format is not used at the time of the creation of the matrix itself, but is used internally.

As discussed in the chapter on Sparse Matrices, the values of the sparse matrix are stored in increasing column-major ordering. Although the data passed by the user need not respect this requirement, pre-sorting the data will significantly speed up creation of the sparse matrix.

The disadvantage of this technique for creating a sparse matrix is that there is a brief time when two copies of the data exist. For extremely memory constrained problems this may not be the best technique for creating a sparse matrix.

The alternative is to first create a sparse matrix with the desired number of nonzero elements and then later fill those elements in. Sample code:

```
int nz, nr, nc;
nz = 4, nr = 3, nc = 4;
SparseMatrix sm (nr, nc, nz);
sm(0,0) = 1; sm(0,1) = 2; sm(1,3) = 3; sm(2,3) = 4;
```

This creates the same matrix as previously. Again, although not strictly necessary, it is significantly faster if the sparse matrix is created and the elements are added in column-major ordering. The reason for this is that when elements are inserted at the end of the current list of known elements then no element in the matrix needs to be moved to allow the new element to be inserted; Only the column indices need to be updated.

There are a few further points to note about this method of creating a sparse matrix. First, it is possible to create a sparse matrix with fewer elements than are actually inserted in the matrix. Therefore,

```
int nr, nc;
nr = 3, nc = 4;
SparseMatrix sm (nr, nc, 0);
sm(0,0) = 1; sm(0,1) = 2; sm(1,3) = 3; sm(2,3) = 4;
```

is perfectly valid. However, it is a very bad idea because as each new element is added to the sparse matrix the matrix needs to request more space and reallocate memory. This is an expensive operation that will significantly slow this means of creating a sparse matrix. It is possible to create a sparse matrix with excess storage, so having *nz* greater than 4 in this example is also valid. The disadvantage is that the matrix occupies more memory than strictly needed.

Of course, it is not always possible to know the number of nonzero elements prior to filling a matrix. For this reason the additional unused storage of a sparse matrix can be removed after its creation with the `maybe_compress` function. In addition to deallocating unused storage, `maybe_compress` can also remove zero elements from the matrix. The removal of zero elements from the matrix is controlled by setting the argument of the `maybe_compress` function to be `true`. However, the cost of removing the zeros is high because it implies re-sorting the elements. If possible, it is better for the user to avoid adding the unnecessary zeros in the first place. An example of the use of `maybe_compress` is

```

int nz, nr, nc;
nz = 6, nr = 3, nc = 4;

SparseMatrix sm1 (nr, nc, nz);
sm1(0,0) = 1; sm1(0,1) = 2; sm1(1,3) = 3; sm1(2,3) = 4;
sm1.maybe_compress (); // No zero elements were added

SparseMatrix sm2 (nr, nc, nz);
sm2(0,0) = 1; sm2(0,1) = 2; sm2(0,2) = 0; sm2(1,2) = 0;
sm2(1,3) = 3; sm2(2,3) = 4;
sm2.maybe_compress (true); // Zero elements were added

```

The use of the `maybe_compress` function should be avoided if possible as it will slow the creation of the matrix.

A third means of creating a sparse matrix is to work directly with the data in compressed row format. An example of this advanced technique might be

```

octave_value arg;
...
int nz, nr, nc;
nz = 6, nr = 3, nc = 4; // Assume we know the max # nz
SparseMatrix sm (nr, nc, nz);
Matrix m = arg.matrix_value ();

int ii = 0;
sm.cidx (0) = 0;
for (int j = 1; j < nc; j++)
{
    for (int i = 0; i < nr; i++)
    {
        double tmp = m(i,j);
        if (tmp != 0.)
        {
            sm.data(ii) = tmp;
            sm.ridx(ii) = i;
            ii++;
        }
    }
    sm.cidx(j+1) = ii;
}
sm.maybe_compress (); // If don't know a priori the final # of nz.

```

which is probably the most efficient means of creating a sparse matrix.

Finally, it may sometimes arise that the amount of storage initially created is insufficient to completely store the sparse matrix. Therefore, the method `change_capacity` exists to reallocate the sparse memory. The above example would then be modified as

```

octave_value arg;
...

```

```

int nz, nr, nc;
nz = 6, nr = 3, nc = 4; // Guess the number of nz elements
SparseMatrix sm (nr, nc, nz);
Matrix m = arg.matrix_value ();

int ii = 0;
sm.cidx (0) = 0;
for (int j = 1; j < nc; j++)
{
    for (int i = 0; i < nr; i++)
    {
        double tmp = m(i,j);
        if (tmp != 0.)
        {
            if (ii == nz)
            {
                nz += 2; // Add 2 more elements
                sm.change_capacity (nz);
            }
            sm.data(ii) = tmp;
            sm.ridx(ii) = i;
            ii++;
        }
    }
    sm.cidx(j+1) = ii;
}
sm.maybe_compress (); // If don't know a priori the final # of nz.

```

Note that both increasing and decreasing the number of nonzero elements in a sparse matrix is expensive as it involves memory reallocation. Also because parts of the matrix, though not its entirety, exist as old and new copies at the same time, additional memory is needed. Therefore, if possible avoid changing capacity.

A.1.6.3 Using Sparse Matrices in Oct-Files

Most of the same operators and functions for sparse matrices that are available from the Octave interpreter are also available within oct-files. The basic means of extracting a sparse matrix from an `octave_value`, and returning it as an `octave_value`, can be seen in the following example.

```

octave_value_list retval;

SparseMatrix sm = args(0).sparse_matrix_value ();
SparseComplexMatrix scm = args(1).sparse_complex_matrix_value ();
SparseBoolMatrix sbm = args(2).sparse_bool_matrix_value ();
...
retval(2) = sbm;
retval(1) = scm;
retval(0) = sm;

```

The conversion to an `octave_value` is handled by the sparse `octave_value` constructors, and so no special care is needed.

A.1.7 Accessing Global Variables in Oct-Files

Global variables allow variables in the global scope to be accessed. Global variables can be accessed within oct-files by using the support functions `get_global_value` and `set_global_value`. `get_global_value` takes two arguments, the first is a string representing the variable name to obtain. The second argument is a boolean argument specifying what to do if no global variable of the desired name is found. An example of the use of these two functions is

```
#include <octave/oct.h>

DEFUN_DLD (globaldemo, args, , "Global Demo")
{
  if (args.length () != 1)
    print_usage ();

  octave_value retval;

  std::string s = args(0).string_value ();

  octave_value tmp = get_global_value (s, true);

  if (tmp.is_defined ())
    retval = tmp;
  else
    retval = "Global variable not found";

  set_global_value ("a", 42.0);

  return retval;
}
```

An example of its use is

```
global a b
b = 10;
globaldemo ("b")
⇒ 10
globaldemo ("c")
⇒ "Global variable not found"
num2str (a)
⇒ 42
```

A.1.8 Calling Octave Functions from Oct-Files

There is often a need to be able to call another Octave function from within an oct-file, and there are many examples of such within Octave itself. For example, the `quad` function is an oct-file that calculates the definite integral by quadrature over a user-supplied function.

There are also many ways in which a function could be given as input. It might be passed as one of

1. Function Handle
2. Anonymous Function Handle
3. Inline Function
4. String

The code below demonstrates all four methods of passing a function to an oct-file.

```
#include <octave/oct.h>
#include <octave/parse.h>

DEFUN_DLD (funcdemo, args, nargin, "Function Demo")
{
  int nargin = args.length ();

  if (nargin < 2)
    print_usage ();

  octave_value_list newargs;

  for (octave_idx_type i = nargin - 1; i > 0; i--)
    newargs(i-1) = args(i);

  octave_value_list retval;

  if (args(0).is_function_handle () || args(0).is_inline_function ())
  {
    octave_function *fcn = args(0).function_value ();

    retval = feval (fcn, newargs, nargin);
  }
  else if (args(0).is_string ())
  {
    std::string fcn = args(0).string_value ();

    retval = feval (fcn, newargs, nargin);
  }
  else
    error ("funcdemo: INPUT must be string, inline, or function handle");

  return retval;
}
```

The first input to the demonstration code is a user-supplied function and the remaining arguments are all passed to the function.

```

funcdemo (@sin, 1)
⇒ 0.84147
funcdemo (@(x) sin (x), 1)
⇒ 0.84147
funcdemo (inline ("sin (x)", 1)
⇒ 0.84147
funcdemo ("sin", 1)
⇒ 0.84147
funcdemo (@atan2, 1, 1)
⇒ 0.78540

```

When the user function is passed as a string the treatment of the function is different. In some cases it is necessary to have the user supplied function as an `octave_function` object. In that case the string argument can be used to create a temporary function as demonstrated below.

```

std::octave fcn_name = unique_symbol_name ("__fcn__");
std::string fcode = "function y = ";
fcode.append (fcn_name);
fcode.append ("(x) y = ");
fcn = extract_function (args(0), "funcdemo", fcn_name,
                        fcode, "; endfunction");

...
if (fcn_name.length ())
    clear_function (fcn_name);

```

There are two important things to know in this case. First, the number of input arguments to the user function is fixed, and in the above example is a single argument. Second, to avoid leaving the temporary function in the Octave symbol table it should be cleared after use. Also, by convention all internal function names begin and end with the character sequence ‘`__`’.

A.1.9 Calling External Code from Oct-Files

Linking external C code to Octave is relatively simple, as the C functions can easily be called directly from C++. One possible issue is that the declarations of the external C functions may need to be explicitly defined as C functions to the compiler. If the declarations of the external C functions are in the header `foo.h`, then the tactic to ensure that the C++ compiler treats these declarations as C code is

```

#ifndef __cplusplus
extern "C"
{
#endif
#include "foo.h"
#ifndef __cplusplus
} /* end extern "C" */
#endif

```

Calling Fortran code, however, can pose more difficulties. This is due to differences in the manner in which compilers treat the linking of Fortran code with C or C++ code. Octave supplies several macros that allow consistent behavior across a number of compilers.

The underlying Fortran code should use the `XSTOPX` function to replace the Fortran `STOP` function. `XSTOPX` uses the Octave exception handler to treat failing cases in the Fortran code explicitly. Note that Octave supplies its own replacement BLAS `XERBLA` function, which uses `XSTOPX`.

If the code calls `XSTOPX`, then the `F77_XFCN` macro should be used to call the underlying Fortran function. The Fortran exception state can then be checked with the global variable `f77_exception_encountered`. If `XSTOPX` will not be called, then the `F77_FCN` macro should be used instead to call the Fortran code.

There is no great harm in using `F77_XFCN` in all cases, except that for Fortran code that is short running and executes a large number of times, there is potentially an overhead in doing so. However, if `F77_FCN` is used with code that calls `XSTOP`, Octave can generate a segmentation fault.

An example of the inclusion of a Fortran function in an oct-file is given in the following example, where the C++ wrapper is

```
#include <octave/oct.h>
#include <octave/f77-fcn.h>

extern "C"
{
    F77_RET_T
    F77_FUNC (fortransub, FORTSUB)
        (const F77_INT&, F77_DBLE*, F77_CHAR_ARG_DECL F77_CHAR_ARG_LEN_DECL);
}

DEFUN_DLD (fortrandemo, args, , "Fortran Demo")
{
    if (args.length () != 1)
        print_usage ();

    NDArray a = args(0).array_value ();

    double *av = a.fortran_vec ();
    octave_idx_type na = a.numel ();

    OCTAVE_LOCAL_BUFFER (char, ctmp, 128);

    F77_XFCN (fortransub, FORTSUB,
               (na, av, ctmp F77_CHAR_ARG_LEN (128)));

    return ovl (a, std::string (ctmp));
}
```

and the Fortran function is

```
subroutine fortransub (n, a, s)
implicit none
character(*) s
```

```

real*8 a(*)
integer*4 i, n, ioerr
do i = 1, n
  if (a(i) .eq. 0d0) then
    call xstopx ('fortransub: divide by zero')
  else
    a(i) = 1d0 / a(i)
  endif
enddo
write (unit = s, fmt = '(a,i3,a,a)', iostat = ioerr)
$      'There are ', n,
$      ' values in the input vector', char(0)
if (ioerr .ne. 0) then
  call xstopx ('fortransub: error writing string')
endif
return
end

```

This example demonstrates most of the features needed to link to an external Fortran function, including passing arrays and strings, as well as exception handling. Both the Fortran and C++ files need to be compiled in order for the example to work.

```

mkoctfile fortrandemo.cc fortransub.f
[b, s] = fortrandemo (1:3)
⇒
b = 1.00000  0.50000  0.33333
s = There are  3 values in the input vector
[b, s] = fortrandemo (0:3)
error: fortrandemo: fortransub: divide by zero

```

A.1.10 Allocating Local Memory in Oct-Files

Allocating memory within an oct-file might seem easy, as the C++ new/delete operators can be used. However, in that case great care must be taken to avoid memory leaks. The preferred manner in which to allocate memory for use locally is to use the `OCTAVE_LOCAL_BUFFER` macro. An example of its use is

```
OCTAVE_LOCAL_BUFFER (double, tmp, len)
```

that returns a pointer `tmp` of type `double *` of length `len`.

In this case, Octave itself will worry about reference counting and variable scope and will properly free memory without programmer intervention.

A.1.11 Input Parameter Checking in Oct-Files

Because oct-files are compiled functions they open up the possibility of crashing Octave through careless function calls or memory faults. It is quite important that each and every function have a sufficient level of parameter checking to ensure that Octave behaves well.

The minimum requirement, as previously discussed, is to check the number of input arguments before using them to avoid referencing a nonexistent argument. However, in some cases this might not be sufficient as the underlying code imposes further constraints.

For example, an external function call might be undefined if the input arguments are not integers, or if one of the arguments is zero, or if the input is complex and a real value was expected. Therefore, oct-files often need additional input parameter checking.

There are several functions within Octave that can be useful for the purposes of parameter checking. These include the methods of the `octave_value` class like `is_real_matrix`, `is_numeric_type`, etc. (see `ov.h`). Often, with a knowledge of the Octave m-file language, you can guess at what the corresponding C++ routine will. In addition there are some more specialized input validation functions of which a few are demonstrated below.

```
#include <octave/oct.h>

DEFUN_DLD (paramdemo, args, nargout, "Parameter Check Demo")
{
  if (args.length () != 1)
    print_usage ();

  NDArray m = args(0).array_value ();

  double min_val = -10.0;
  double max_val = 10.0;

  octave_stdout << "Properties of input array:\n";

  if (m.any_element_is_negative ())
    octave_stdout << " includes negative values\n";

  if (m.any_element_is_inf_or_nan ())
    octave_stdout << " includes Inf or NaN values\n";

  if (m.any_element_not_one_or_zero ())
    octave_stdout << " includes other values than 1 and 0\n";

  if (m.all_elements_are_int_or_inf_or_nan ())
    octave_stdout << " includes only int, Inf or NaN values\n";

  if (m.all_integers (min_val, max_val))
    octave_stdout << " includes only integers in [-10,10]\n";

  return octave_value_list ();
}
```

An example of its use is:

```
paramdemo ([1, 2, NaN, Inf])
⇒ Properties of input array:
  includes Inf or NaN values
  includes other values than 1 and 0
  includes only int, Inf or NaN values
```

A.1.12 Exception and Error Handling in Oct-Files

Another important feature of Octave is its ability to react to the user typing **Control-C** during extended calculations. This ability is based on the C++ exception handler, where memory allocated by the C++ new/delete methods is automatically released when the exception is treated. When writing an oct-file which may run for a long time the programmer must periodically use the macro `OCTAVE_QUIT`, in order to allow Octave to check and possibly respond to a user typing **Control-C**. For example:

```
for (octave_idx_type i = 0; i < a.nelem (); i++)
{
    OCTAVE_QUIT;
    b.elem (i) = 2. * a.elem (i);
}
```

The presence of the `OCTAVE_QUIT` macro in the inner loop allows Octave to detect and acknowledge a **Control-C** key sequence. Without this macro, the user must either wait for the oct-file function to return before the interrupt is processed, or the user must press **Control-C** three times which will force Octave to exit completely.

The `OCTAVE_QUIT` macro does impose a very small performance penalty; For loops that are known to be small it may not make sense to include `OCTAVE_QUIT`.

When creating an oct-file that uses an external library, the function might spend a significant portion of its time in the external library. It is not generally possible to use the `OCTAVE_QUIT` macro in this case. The alternative code in this case is

```
BEGIN_INTERRUPT_IMMEDIATELY_IN_FOREIGN_CODE;
...
some code that calls a "foreign" function ...
END_INTERRUPT_IMMEDIATELY_IN_FOREIGN_CODE;
```

The disadvantage of this is that if the foreign code allocates any memory internally, then this memory might be lost during an interrupt, without being deallocated. Therefore, ideally Octave itself should allocate any memory that is needed by the foreign code, with either the `fortran_vec` method or the `OCTAVE_LOCAL_BUFFER` macro.

The Octave `unwind_protect` mechanism ([Section 10.8 \[The `unwind_protect` Statement\]](#), [page 174](#)) can also be used in oct-files. In conjunction with the exception handling of Octave, it ensures that certain recovery code is always run even if an exception occurs. An example of the use of this mechanism is

```
#include <octave/oct.h>
#include <octave/unwind-prot.h>

void
my_err_handler (const char *fmt, ...)
{
    // Do nothing!!
}

void
my_err_with_id_handler (const char *id, const char *fmt, ...)
{
    // Do nothing!!
```

```

}

DEFUN_DLD (unwinddemo, args, narginout, "Unwind Demo")
{
  if (args.length () < 2)
    print_usage ();

  NDArray a = args(0).array_value ();
  NDArray b = args(1).array_value ();

  // Declare unwind_protect frame which lasts as long as
  // the variable frame has scope.
  octave::unwind_protect frame;
  frame.add_fcn (set_liboctave_warning_handler,
                 current_liboctave_warning_handler);

  frame.add_fcn (set_liboctave_warning_with_id_handler,
                 current_liboctave_warning_with_id_handler);

  set_liboctave_warning_handler (my_err_handler);
  set_liboctave_warning_with_id_handler (my_err_with_id_handler);

  return octave_value (quotient (a, b));
}

```

As can be seen in the example:

```

unwinddemo (1, 0)
⇒ Inf
1 / 0
⇒ warning: division by zero
Inf

```

The warning for division by zero (and in fact all warnings) are disabled in the `unwinddemo` function.

A.1.13 Documentation and Testing of Oct-Files

The documentation for an oct-file is contained in the fourth string parameter of the DEFUN_DLD macro. This string can be formatted in the same manner as the help strings for user functions, however there are some issues that are particular to the formatting of help strings within oct-files.

The major issue is that the help string will typically be longer than a single line of text, and so the formatting of long multi-line help strings needs to be taken into account. There are several possible solutions, but the most common is illustrated in the following example,

```

DEFUN_DLD (do_what_i_want, args, nargout,
  "-**- texinfo -*-\n\
@deftypefn {} {} do_what_i_say (@var{n})\n\
A function that does what the user actually wants rather\n\
than what they requested.\n\
@end deftypefn")
{
...
}

```

where each line of text is terminated by `\n\` which is an embedded newline in the string together with a C++ string continuation character. Note that the final `\` must be the last character on the line.

Octave also includes the ability to embed test and demonstration code for a function within the code itself (see [Appendix B \[Test and Demo Functions\]](#), page 997). This can be used from within oct-files (or in fact any file) with certain provisos. First, the test and demo functions of Octave look for `%!` as the first two characters of a line to identify test and demonstration code. This is a requirement for oct-files as well. In addition, the test and demonstration code must be wrapped in a comment block to avoid it being interpreted by the compiler. Finally, the Octave test and demonstration code must have access to the original source code of the oct-file—not just the compiled code—as the tests are stripped from the compiled code. An example in an oct-file might be

```

/*
%!assert (sin ([1,2]), [sin(1),sin(2)])
%!error (sin ())
%!error (sin (1,1))
*/

```

A.2 Mex-Files

Octave includes an interface to allow legacy mex-files to be compiled and used with Octave. This interface can also be used to share compiled code between Octave and MATLAB users. However, as mex-files expose MATLAB's internal API, and the internal structure of Octave is different, a mex-file can never have the same performance in Octave as the equivalent oct-file. In particular, to support the manner in which variables are passed to mex functions there are a significant number of additional copies of memory blocks when invoking or returning from a mex-file function. For this reason, it is recommended that any new code be written with the oct-file interface previously discussed.

A.2.1 Getting Started with Mex-Files

The basic command to build a mex-file is either `mkoctfile --mex` or `mex`. The first command can be used either from within Octave or from the command line. To avoid issues with MATLAB's own `mex` command, the use of the command `mex` is limited to within Octave. Compiled mex-files have the extension `.mex`.

`mex [options] file ...`

Compile source code written in C, C++, or Fortran, to a MEX file.

This is equivalent to `mkoctfile --mex [options] file`.

See also: [mkostfile], page 946, [mexext], page 971.

mexext ()

Return the filename extension used for MEX files.

See also: [mex], page 970.

Consider the following short example:

```
#include "mex.h"

void
mexFunction (int nlhs, mxArray *plhs[],
             int nrhs, const mxArray *prhs[])
{
    mexPrintf ("Hello, World!\n");

    mexPrintf ("I have %d inputs and %d outputs\n", nrhs, nlhs);

    /* Return empty matrices for any outputs */
    int i;
    for (i = 0; i < nlhs; i++)
        plhs[i] = mxCreateDoubleMatrix (0, 0, mxREAL);
}
```

The first line `#include "mex.h"` makes available all of the definitions necessary for a mex-file. One important difference between Octave and MATLAB is that the header file "`matrix.h`" is implicitly included through the inclusion of "`mex.h`". This is necessary to avoid a conflict with the Octave file "`Matrix.h`" for operating systems and compilers that don't distinguish between filenames in upper and lower case.

The entry point into the mex-file is defined by `mexFunction`. The function takes four arguments:

1. The number of return arguments (# of left-hand side args).
2. An array of pointers to return arguments.
3. The number of input arguments (# of right-hand side args).
4. An array of pointers to input arguments.

Note that the function name definition is not explicitly included in `mexFunction` and so there can only be a single `mexFunction` entry point per file. Instead, the name of the function as seen in Octave is determined by the name of the mex-file itself minus the extension. If the above function is in the file `myhello.c`, it can be compiled with

```
mkostfile --mex myhello.c
```

which creates a file `myhello.mex`. The function can then be run from Octave as

```
myhello (1,2,3)
⇒ Hello, World!
⇒ I have 3 inputs and 0 outputs
```

It should be noted that the mex-file contains no help string. To document mex-files, there should exist an m-file in the same directory as the mex-file itself. Taking the above as an example, there would need to be a file `myhello.m` which might contain the text

```
%MYHELLO Simple test of the functionality of a mex-file.
```

In this case, the function that will be executed within Octave will be given by the mex-file, while the help string will come from the m-file. This can also be useful to allow a sample implementation of the mex-file within the Octave language itself for testing purposes.

Although there cannot be multiple entry points in a single mex-file, one can use the `mexFunctionName` function to determine what name the mex-file was called with. This can be used to alter the behavior of the mex-file based on the function name. For example, if

```
#include "mex.h"

void
mexFunction (int nlhs, mxArray *plhs[],
             int nrhs, const mxArray *prhs[])
{
    const char *nm;

    nm = mexFunctionName ();
    mexPrintf ("You called function: %s\n", nm);
    if (strcmp (nm, "myfunc") == 0)
        mexPrintf ("This is the principal function\n", nm);

    return;
}
```

is in the file `myfunc.c`, and is compiled with

```
mkostfile --mex myfunc.c
ln -s myfunc.mex myfunc2.mex
```

then as can be seen by

```
myfunc ()
⇒ You called function: myfunc
      This is the principal function
myfunc2 ()
⇒ You called function: myfunc2
```

the behavior of the mex-file can be altered depending on the function's name.

Although the user should only include `mex.h` in their code, Octave declares additional functions, `typedefs`, etc., available to the user to write mex-files in the headers `mexproto.h` and `mxarray.h`.

A.2.2 Working with Matrices and Arrays in Mex-Files

The basic mex type of all variables is `mxArray`. Any object, such as a matrix, cell array, or structure, is stored in this basic type. `mxArray` serves essentially the same purpose as the `octave_value` class in oct-files in that it acts as a container for all the more specialized types.

The `mxArray` structure contains at a minimum, the name of the variable it represents, its dimensions, its type, and whether the variable is real or complex. It can also contain a number of additional fields depending on the type of the `mxArray`. There are a number of func-

tions to create `mxArray` structures, including `mxCreateDoubleMatrix`, `mxCreateCellArray`, `mxCreateSparse`, and the generic `mxCreateNumericArray`.

The basic function to access the data in an array is `mxGetPr`. Because the mex interface assumes that real and imaginary parts of a complex array are stored separately, there is an equivalent function `mxGetPi` that gets the imaginary part. Both of these functions are only for use with double precision matrices. The generic functions `mxGetData` and `mxGetImagData` perform the same operation for all matrix types. For example:

```
 mxArray *m;
mwSize *dims;
UINT32_T *pr;

dims = (mwSize *) mxMalloc (2 * sizeof (mwSize));
dims[0] = 2; dims[1] = 2;
m = mxCreateNumericArray (2, dims, mxUINT32_CLASS, mxREAL);
pr = (UINT32_T *) mxGetData (m);
```

There are also the functions `mxSetPr`, etc., that perform the inverse, and set the data of an array to use the block of memory pointed to by the argument of `mxSetPr`.

Note the type `mwSize` used above, and also `mwIndex`, are defined as the native precision of the indexing in Octave on the platform on which the mex-file is built. This allows both 32- and 64-bit platforms to support mex-files. `mwSize` is used to define array dimensions and the maximum number of elements, while `mwIndex` is used to define indexing into arrays.

An example that demonstrates how to work with arbitrary real or complex double precision arrays is given by the file `mypow2.c` shown below.

```
#include "mex.h"

void
mexFunction (int nlhs, mxArray *plhs[],
             int nrhs, const mxArray *prhs[])
{
    mwSize n;
    mwIndex i;
    double *vri, *vro;

    if (nrhs != 1 || ! mxIsDouble (prhs[0]))
        mexErrMsgTxt ("ARG1 must be a double matrix");

    n = mxGetNumberOfElements (prhs[0]);
    plhs[0] = mxCreateNumericArray (mxGetNumberOfDimensions (prhs[0]),
                                   mxGetDimensions (prhs[0]),
                                   mxGetClassID (prhs[0]),
                                   mxIsComplex (prhs[0]));

    vri = mxGetPr (prhs[0]);
    vro = mxGetPr (plhs[0]);

    if (mxIsComplex (prhs[0]))
```

```

{
  double *vii, *vio;
  vii = mxGetPi (prhs[0]);
  vio = mxGetPi (plhs[0]);

  for (i = 0; i < n; i++)
  {
    vro[i] = vri[i] * vri[i] - vii[i] * vii[i];
    vio[i] = 2 * vri[i] * vii[i];
  }
}
else
{
  for (i = 0; i < n; i++)
    vro[i] = vri[i] * vri[i];
}
}

```

An example of its use is

```

b = randn (4,1) + 1i * randn (4,1);
all (b.^2 == mypow2 (b))
⇒ 1

```

The example above uses the functions `mxGetDimensions`, `mxGetNumberOfElements`, and `mxGetNumberOfDimensions` to work with the dimensions of multi-dimensional arrays. The functions `mxGetM`, and `mxGetN` are also available to find the number of rows and columns in a 2-D matrix (MxN matrix).

A.2.3 Character Strings in Mex-Files

As mex-files do not make the distinction between single and double quoted strings that Octave does, there is perhaps less complexity in the use of strings and character matrices. An example of their use that parallels the demo in `stringdemo.cc` is given in the file `mystring.c`, as shown below.

```

#include <string.h>
#include "mex.h"

void
mexFunction (int nlhs, mxArray *plhs[],
             int nrhs, const mxArray *prhs[])
{
  mwSize m, n;
  mwIndex i, j;
  mxChar *pi, *po;

  if (nrhs != 1 || ! mxIsChar (prhs[0])
      || mxGetNumberOfDimensions (prhs[0]) > 2)
    mexErrMsgTxt ("ARG1 must be a char matrix");

```

```

m = mxGetM (prhs[0]);
n = mxGetN (prhs[0]);
pi = mxGetChars (prhs[0]);
plhs[0] = mxCreateNumericMatrix (m, n, mxCHAR_CLASS, mxREAL);
po = mxGetChars (plhs[0]);

for (j = 0; j < n; j++)
    for (i = 0; i < m; i++)
        po[j*m + m - 1 - i] = pi[j*m + i];
}

```

An example of its expected output is

```

mystring (["First String"; "Second String"])
⇒ Second String
    First String

```

Other functions in the mex interface for handling character strings are `mxCreateString`, `mxArrayToString`, and `mxCreateCharMatrixFromStrings`. In a mex-file, a character string is considered to be a vector rather than a matrix. This is perhaps an arbitrary distinction as the data in the `mxArray` for the matrix is consecutive in any case.

A.2.4 Cell Arrays with Mex-Files

One can perform exactly the same operations on Cell arrays in mex-files as in oct-files. An example that duplicates the function of the `celldemo.cc` oct-file in a mex-file is given by `mycell.c` as shown below.

```

#include "mex.h"

void
mexFunction (int nlhs, mxArray *plhs[],
              int nrhs, const mxArray *prhs[])
{
    mwSize n;
    mwIndex i;

    if (nrhs != 1 || ! mxIsCell (prhs[0]))
        mexErrMsgTxt ("ARG1 must be a cell");

    n = mxGetNumberOfElements (prhs[0]);
    n = (n > nlhs ? nlhs : n);

    for (i = 0; i < n; i++)
        plhs[i] = mxDuplicateArray (mxGetCell (prhs[0], i));
}

```

The output is identical to the oct-file version as well.

```
[b1, b2, b3] = mycell ({1, [1, 2], "test"})
⇒
b1 = 1
b2 =
      2

b3 = test
```

Note in the example the use of the `mxDuplicateArray` function. This is needed as the `mxArray` pointer returned by `mxGetCell` might be deallocated. The inverse function to `mxGetCell`, used for setting Cell values, is `mxSetCell` and is defined as

```
void mxSetCell (mxArray *ptr, int idx, mxArray *val);
```

Finally, to create a cell array or matrix, the appropriate functions are

```
mxArray *mxCreateCellArray (int ndims, const int *dims);
mxArray *mxCreateCellMatrix (int m, int n);
```

A.2.5 Structures with Mex-Files

The basic function to create a structure in a mex-file is `mxCreateStructMatrix` which creates a structure array with a two dimensional matrix, or `mxCreateStructArray`.

```
mxArray *mxCreateStructArray (int ndims, int *dims,
                             int num_keys,
                             const char **keys);
mxArray *mxCreateStructMatrix (int rows, int cols,
                             int num_keys,
                             const char **keys);
```

Accessing the fields of the structure can then be performed with `mxGetField` and `mxSetField` or alternatively with the `mxGetFieldByNumber` and `mxSetFieldByNumber` functions.

```
mxArray *mxGetField (const mxArray *ptr, mwIndex index,
                     const char *key);
mxArray *mxGetFieldByNumber (const mxArray *ptr,
                           mwIndex index, int key_num);
void mxSetField (mxArray *ptr, mwIndex index,
                 const char *key, mxArray *val);
void mxSetFieldByNumber (mxArray *ptr, mwIndex index,
                        int key_num, mxArray *val);
```

A difference between the oct-file interface to structures and the mex-file version is that the functions to operate on structures in mex-files directly include an `index` over the elements of the arrays of elements per `field`; Whereas, the oct-file structure includes a Cell Array per field of the structure.

An example that demonstrates the use of structures in a mex-file can be found in the file `mystruct.c` shown below.

```
#include "mex.h"
```

```

void
mexFunction (int nlhs, mxArray *plhs[],
              int nrhs, const mxArray *prhs[])
{
    int i;
    mwIndex j;
    mxArray *v;
    const char *keys[] = { "this", "that" };

    if (nrhs != 1 || ! mxIsStruct (prhs[0]))
        mexErrMsgTxt ("ARG1 must be a struct");

    for (i = 0; i < mxGetNumberOfFields (prhs[0]); i++)
        for (j = 0; j < mxGetNumberOfElements (prhs[0]); j++)
    {
        mexPrintf ("field %s(%d) = ", mxGetFieldNameByNumber (prhs[0], i), j);■
        v = mxGetFieldByNumber (prhs[0], j, i);
        mexCallMATLAB (0, NULL, 1, &v, "disp");
    }

    v = mxCreateStructMatrix (2, 2, 2, keys);

    mxSetFieldByNumber (v, 0, 0, mxCreateString ("this1"));
    mxSetFieldByNumber (v, 0, 1, mxCreateString ("that1"));
    mxSetFieldByNumber (v, 1, 0, mxCreateString ("this2"));
    mxSetFieldByNumber (v, 1, 1, mxCreateString ("that2"));
    mxSetFieldByNumber (v, 2, 0, mxCreateString ("this3"));
    mxSetFieldByNumber (v, 2, 1, mxCreateString ("that3"));
    mxSetFieldByNumber (v, 3, 0, mxCreateString ("this4"));
    mxSetFieldByNumber (v, 3, 1, mxCreateString ("that4"));

    if (nlhs)
        plhs[0] = v;
}

```

An example of the behavior of this function within Octave is then

```

a(1).f1 = "f11"; a(1).f2 = "f12";
a(2).f1 = "f21"; a(2).f2 = "f22";
b = mystruct (a);
⇒ field f1(0) = f11
    field f1(1) = f21
    field f2(0) = f12
    field f2(1) = f22
b
⇒ 2x2 struct array containing the fields:

    this
    that

b(3)
⇒ scalar structure containing the fields:

    this = this3
    that = that3

```

A.2.6 Sparse Matrices with Mex-Files

The Octave format for sparse matrices is identical to the mex format in that it is a compressed column sparse format. Also, in both implementations sparse matrices are required to be two-dimensional. The only difference of importance to the programmer is that the real and imaginary parts of the matrix are stored separately.

The mex-file interface, in addition to using `mxGetM`, `mxGetN`, `mxSetM`, `mxSetN`, `mxGetPr`, `mxGetPi`, `mxSetPr`, and `mxSetPi`, also supplies the following functions.

```

mwIndex *mxGetIr (const mxArray *ptr);
mwIndex *mxGetJc (const mxArray *ptr);
mwSize mxGetNzmax (const mxArray *ptr);

void mxSetIr (mxArray *ptr, mwIndex *ir);
void mxSetJc (mxArray *ptr, mwIndex *jc);
void mxSetNzmax (mxArray *ptr, mwSize nzmax);

```

`mxGetNzmax` gets the maximum number of elements that can be stored in the sparse matrix. This is not necessarily the number of nonzero elements in the sparse matrix. `mxGetJc` returns an array with one additional value than the number of columns in the sparse matrix. The difference between consecutive values of the array returned by `mxGetJc` define the number of nonzero elements in each column of the sparse matrix. Therefore,

```

mwSize nz, n;
mwIndex *Jc;
mxArray *m;
...
n = mxGetN (m);
Jc = mxGetJc (m);
nz = Jc[n];

```

returns the actual number of nonzero elements stored in the matrix in `nz`. As the arrays returned by `mxGetPr` and `mxGetPi` only contain the nonzero values of the matrix, we also need a pointer to the rows of the nonzero elements, and this is given by `mxGetIr`. A complete example of the use of sparse matrices in mex-files is given by the file `mysparse.c` shown below.

```
#include "mex.h"

void
mexFunction (int nlhs, mxArray *plhs[],
             int nrhs, const mxArray *prhs[])
{
    mwSize m, n, nz;
    mxArray *v;
    mwIndex i;
    double *pr, *pi;
    double *pr2, *pi2;
    mwIndex *ir, *jc;
    mwIndex *ir2, *jc2;

    if (nrhs != 1 || ! mxIsSparse (prhs[0]))
        mexErrMsgTxt ("ARG1 must be a sparse matrix");

    m = mxGetM (prhs[0]);
    n = mxGetN (prhs[0]);
    nz = mxGetNzmax (prhs[0]);

    if (mxIsComplex (prhs[0]))
    {
        mexPrintf ("Matrix is %d-by-%d complex sparse matrix", m, n);
        mexPrintf (" with %d elements\n", nz);

        pr = mxGetPr (prhs[0]);
        pi = mxGetPi (prhs[0]);
        ir = mxGetIr (prhs[0]);
        jc = mxGetJc (prhs[0]);

        i = n;
        while (jc[i] == jc[i-1] && i != 0) i--;

        mexPrintf ("last nonzero element (%d, %d) = (%g, %g)\n",
                  ir[nz-1]+ 1, i, pr[nz-1], pi[nz-1]);

        v = mxCreateSparse (m, n, nz, mxCOMPLEX);
        pr2 = mxGetPr (v);
        pi2 = mxGetPi (v);
        ir2 = mxGetIr (v);
    }
}
```

```
jc2 = mxGetJc (v);

for (i = 0; i < nz; i++)
{
    pr2[i] = 2 * pr[i];
    pi2[i] = 2 * pi[i];
    ir2[i] = ir[i];
}
for (i = 0; i < n + 1; i++)
    jc2[i] = jc[i];

if (nlhs > 0)
    plhs[0] = v;
}
else if (mxIsLogical (prhs[0]))
{
    mxLogical *pbr, *pbr2;
    mexPrintf ("Matrix is %d-by-%d logical sparse matrix", m, n);
    mexPrintf (" with %d elements\n", nz);

    pbr = mxGetLogicals (prhs[0]);
    ir = mxGetIr (prhs[0]);
    jc = mxGetJc (prhs[0]);

    i = n;
    while (jc[i] == jc[i-1] && i != 0) i--;
    mexPrintf ("last nonzero element (%d, %d) = %d\n",
               ir[nz-1]+ 1, i, pbr[nz-1]);

    v = mxCreateSparseLogicalMatrix (m, n, nz);
    pbr2 = mxGetLogicals (v);
    ir2 = mxGetIr (v);
    jc2 = mxGetJc (v);

    for (i = 0; i < nz; i++)
    {
        pbr2[i] = pbr[i];
        ir2[i] = ir[i];
    }
    for (i = 0; i < n + 1; i++)
        jc2[i] = jc[i];

    if (nlhs > 0)
        plhs[0] = v;
}
else
{
```

```

mexPrintf ("Matrix is %d-by-%d real sparse matrix", m, n);
mexPrintf (" with %d elements\n", nz);

pr = mxGetPr (prhs[0]);
ir = mxGetIr (prhs[0]);
jc = mxGetJc (prhs[0]);

i = n;
while (jc[i] == jc[i-1] && i != 0) i--;
mexPrintf ("last nonzero element (%d, %d) = %g\n",
           ir[nz-1]+ 1, i, pr[nz-1]);

v = mxCreateSparse (m, n, nz, mxREAL);
pr2 = mxGetPr (v);
ir2 = mxGetIr (v);
jc2 = mxGetJc (v);

for (i = 0; i < nz; i++)
{
    pr2[i] = 2 * pr[i];
    ir2[i] = ir[i];
}
for (i = 0; i < n + 1; i++)
    jc2[i] = jc[i];

if (nlhs > 0)
    plhs[0] = v;
}
}

A sample usage of mysparse is
sm = sparse ([1, 0; 0, pi]);
mysparse (sm)
⇒
Matrix is 2-by-2 real sparse matrix with 2 elements
last nonzero element (2, 2) = 3.14159

```

A.2.7 Calling Other Functions in Mex-Files

It is possible to call other Octave functions from within a mex-file using `mexCallMATLAB`. An example of the use of `mexCallMATLAB` can be see in the example below.

```

#include "mex.h"

void
mexFunction (int nlhs, mxArray *plhs[],
             int nrhs, const mxArray *prhs[])
{
    char *str;

```

```

mexPrintf ("Starting file myfeval.mex\n");

mexPrintf ("I have %d inputs and %d outputs\n", nrhs, nlhs);

if (nrhs < 1 || ! mxIsChar (prhs[0]))
    mexErrMsgTxt ("ARG1 must be a function name");

str = mxArrayToString (prhs[0]);

mexPrintf ("I'm going to call the function %s\n", str);

if (nlhs == 0)
    nlhs = 1; // Octave's automatic 'ans' variable

/* Cast prhs just to get rid of 'const' qualifier and stop compile warning */
mexCallMATLAB (nlhs, plhs, nrhs-1, (mxArray**)prhs+1, str);

mxFree (str);
}

```

If this code is in the file `myfeval.c`, and is compiled to `myfeval.mex`, then an example of its use is

```

a = myfeval ("sin", 1)
⇒ Starting file myfeval.mex
I have 2 inputs and 1 outputs
I'm going to call the interpreter function sin
a = 0.84147

```

Note that it is not possible to use function handles or inline functions within a mex-file.

A.3 Standalone Programs

The libraries Octave uses itself can be utilized in standalone applications. These applications then have access, for example, to the array and matrix classes, as well as to all of the Octave algorithms. The following C++ program, uses class Matrix from `liboctave.a` or `liboctave.so`.

```

#include <iostream>
#include <octave/oct.h>

int
main (void)
{
    std::cout << "Hello Octave world!\n";

    int n = 2;
    Matrix a_matrix = Matrix (n, n);

```

```

    for (octave_idx_type i = 0; i < n; i++)
        for (octave_idx_type j = 0; j < n; j++)
            a_matrix(i,j) = (i + 1) * 10 + (j + 1);

    std::cout << a_matrix;

    return 0;
}

```

`mkoctfile` can be used to build a standalone application with a command like

```

$ mkoctfile --link-stand-alone standalone.cc -o standalone
$ ./standalone
Hello Octave world!
 11 12
 21 22
$
```

Note that the application `standalone` will be dynamically linked against the Octave libraries and any Octave support libraries. The above allows the Octave math libraries to be used by an application. It does not, however, allow the script files, oct-files, or built-in functions of Octave to be used by the application. To do that, the Octave interpreter needs to be initialized first. An example of how to do this can then be seen in the code

```

#include <iostream>
#include <octave/oct.h>
#include <octave/octave.h>
#include <octave/parse.h>
#include <octave/interpreter.h>

int
main (void)
{
    // Create interpreter.

    octave::interpreter interpreter;

    try
    {
        // Inhibit reading history file by calling
        //
        //   interpreter.initialize_history (false);

        // Set custom load path here if you wish by calling
        //
        //   interpreter.initialize_load_path (false);

        // Perform final initialization of interpreter, including
        // executing commands from startup files by calling
        //

```

```
//    int status interpreter.initialize ();
//
//    if (! interpreter.initialized ())
//    {
//        std::cerr << "Octave interpreter initialization failed!"
//                    << std::endl;
//        exit (status);
//    }
//
// You may skip this step if you don't need to do anything
// between reading the startup files and telling the interpreter
// that you are ready to execute commands.

// Tell the interpreter that we're ready to execute commands:

int status = interpreter.execute ();

if (status != 0)
{
    std::cerr << "creating embedded Octave interpreter failed!"
              << std::endl;
    return status;
}

octave_idx_type n = 2;
octave_value_list in;

for (octave_idx_type i = 0; i < n; i++)
    in(i) = octave_value (5 * (i + 2));

octave_value_list out = octave::feval ("gcd", in, 1);

if (out.length () > 0)
    std::cout << "GCD of [
                  << in(0).int_value ()
                  << ", "
                  << in(1).int_value ()
                  << "] is " << out(0).int_value ()
                  << std::endl;
else
    std::cout << "invalid\n";
}
catch (const octave::exit_exception& ex)
{
    std::cerr << "Octave interpreter exited with status = "
              << ex.exit_status () << std::endl;
}
```

```

    catch (const octave::execution_exception&)
    {
        std::cerr << "error encountered in Octave evaluator!" << std::endl;
    }

    return 0;
}

```

which, as before, is compiled and run as a standalone application with

```

$ mkoctfile --link-stand-alone embedded.cc -o embedded
$ ./embedded
GCD of [10, 15] is 5
$ 

```

It is worth re-iterating that, if only built-in functions are to be called from a C++ stand-alone program then it does not need to initialize the interpreter. The general rule is that for a built-in function named `function_name` in the interpreter, there will be a C++ function named `Ffunction_name` (note the prepended capital F) accessible in the C++ API. The declarations for all built-in functions are collected in the header file `builtin-defun-decls.h`. This feature should be used with care as the list of built-in functions can change. No guarantees can be made that a function that is currently a built-in won't be implemented as a .m file or as a dynamically linked function in the future. An example of how to call built-in functions from C++ can be seen in the code

```

#include <iostream>
#include <octave/oct.h>
#include <octave/builtin-defun-decls.h>

int
main (void)
{
    int n = 2;
    Matrix a_matrix = Matrix (n, n);

    for (octave_idx_type i = 0; i < n; i++)
        for (octave_idx_type j = 0; j < n; j++)
            a_matrix(i,j) = (i + 1) * 10 + (j + 1);

    std::cout << "This is a matrix:" << std::endl
           << a_matrix               << std::endl;

    octave_value_list in;
    in(0) = a_matrix;

    octave_value_list out = Fnorm (in, 1);
    double norm_of_the_matrix = out(0).double_value ();

    std::cout << "This is the norm of the matrix:" << std::endl
           << norm_of_the_matrix           << std::endl;
}

```

```
    return 0;
}
```

which is compiled and run as a standalone application with

```
$ mkoctfile --link-stand-alone standalonebuiltin.cc -o standalonebuiltin
$ ./standalonebuiltin
This is a matrix:
 11 12
 21 22

This is the norm of the matrix:
34.4952
$
```

A.4 Java Interface

The Java Interface is designed for calling Java functions from within Octave. If you want to do the reverse, and call Octave from within Java, try a library like `jopas` (<http://jopas.sourceforge.net>).

A.4.1 Making Java Classes Available

Java finds classes by searching a *classpath* which is a list of Java archive files and/or directories containing class files. In Octave the *classpath* is composed of two parts:

- the *static classpath* is initialized once at startup of the JVM, and
- the *dynamic classpath* which can be modified at runtime.

Octave searches the *static classpath* first, and then the *dynamic classpath*. Classes appearing in the *static classpath*, as well as in the *dynamic classpath*, will therefore be found in the *static classpath* and loaded from this location. Classes which will be used frequently, or must be available to all users, should be added to the *static classpath*. The *static classpath* is populated once from the contents of a plain text file named `javaclasspath.txt` (or `classpath.txt` historically) when the Java Virtual Machine starts. This file contains one line for each individual classpath to be added to the *static classpath*. These lines can identify directories containing class files, or Java archives with complete class file hierarchies. Comment lines starting with a '#' or a '%' character are ignored.

The search rules for the file `javaclasspath.txt` (or `classpath.txt`) are:

- First, Octave tries to locate it in the current directory (where Octave was started from). If such a file is found, it is read and defines the initial *static classpath*. Thus, it is possible to define a static classpath on a 'per Octave invocation' basis.
- Next, Octave searches in the user's home directory. If a file `javaclasspath.txt` exists here, its contents are appended to the static classpath (if any). Thus, it is possible to build an initial static classpath on a 'per user' basis.
- Finally, Octave looks for a `javaclasspath.txt` in the m-file directory where Octave Java functions live. This is where the function `javaclasspath.m` resides, usually something like `OCTAVE_HOME/share/octave/OCTAVE_VERSION/m/java/`. You can find this directory by executing the command

```
which javaclasspath
```

If this file exists here, its contents are also appended to the *static classpath*. Note that the archives and class directories defined in this last step will affect all users.

Classes which are used only by a specific script should be placed in the *dynamic classpath*. This portion of the classpath can be modified at runtime using the `javaaddpath` and `javarmpath` functions.

Example:

```
octave> base_path = "C:/Octave/java_files";

octave> # add two JAR archives to the dynamic classpath
octave> javaaddpath ([base_path, "/someclasses.jar"]);
octave> javaaddpath ([base_path, "/moreclasses.jar"]);

octave> # check the dynamic classpath
octave> p = javaclasspath;
octave> disp (p{1});
C:/Octave/java_files/someclasses.jar
octave> disp (p{2});
C:/Octave/java_files/moreclasses.jar

octave> # remove the first element from the classpath
octave> javarmpath ([base_path, "/someclasses.jar"]);
octave> p = javaclasspath;
octave> disp (p{1});
C:/Octave/java_files/moreclasses.jar

octave> # provoke an error
octave> disp (p{2});
error: A(I): Index exceeds matrix dimension.
```

Another way to add files to the *dynamic classpath* exclusively for your user account is to use the file `.octaverc` which is stored in your home directory. All Octave commands in this file are executed each time you start a new instance of Octave. The following example adds the directory `octave` to Octave's search path and the archive `myclasses.jar` in this directory to the Java search path.

```
# contents of .octaverc:
addpath ("~/octave");
javaaddpath ("~/octave/myclasses.jar");
```

A.4.2 How to use Java from within Octave

The function [\[javaObject\]](#), page 990, creates Java objects. In fact it invokes the public constructor of the class with the given name and with the given parameters.

The following example shows how to invoke the constructors `BigDecimal(double)` and `BigDecimal(String)` of the builtin Java class `java.math.BigDecimal`.

```
javaObject ("java.math.BigDecimal", 1.001 );
javaObject ("java.math.BigDecimal", "1.001");
```

Note that parameters of the Octave type `double` are implicitly converted into the Java type `double` and the Octave type (array of) `char` is converted into the java type `String`. A Java object created by [javaObject], page 990, is never automatically converted into an Octave type but remains a Java object. It can be assigned to an Octave variable.

```
a = 1.001;
b = javaObject ("java.math.BigDecimal", a);
```

Using [isjava], page 991, it is possible to check whether a variable is a Java object and its class can be determined as well. In addition to the previous example:

```
isjava (a)
⇒ ans = 0
class (a)
⇒ ans = double
isjava (b)
⇒ ans = 1
class (b)
⇒ ans = java.math.BigDecimal
```

The example above can be carried out using only Java objects:

```
a = javaObject ("java.lang.Double", 1.001);
b = javaObject ("java.math.BigDecimal", a);
```

```
isjava (a)
⇒ ans = 1
class (a)
⇒ ans = java.lang.Double
isjava (b)
⇒ ans = 1
class (b)
⇒ ans = java.math.BigDecimal
```

One can see, that even a `java.lang.Double` is not converted to an Octave `double`, when created by [javaObject], page 990. But ambiguities might arise, if the Java classes `java.lang.Double` or `double` are parameters of a method (or a constructor). In this case they can be converted into one another, depending on the context.

Via [javaObject], page 990, one may create all kinds of Java objects but arrays. The latter are created through [javaArray], page 991.

It is possible to invoke public member methods on Java objects in Java syntax:

```
a.toString
⇒ ans = 1.001
b.toString
⇒ ans = 1.00099999999999889865...
```

The second result may be surprising, but simply comes from the fact, that 1.001 cannot exactly be represented as `double`, due to rounding. Note that unlike in Java, in Octave methods without arguments can be invoked with and without parentheses () .

Currently it is not possible to invoke static methods with a Java like syntax from within Octave. Instead, one has to use the function [javaMethod], page 993, as in the following example:

```
java.math.BigDecimal.valueOf(1.001);           # does not work
javaMethod ("valueOf", "java.math.BigDecimal", 1.001); # workaround
```

As mentioned before, method and constructor parameters are converted automatically between Octave and Java types, if appropriate. For functions this is also true with return values, whereas for constructors this is not.

It is also possible to access public fields of Java objects from within Octave using Java syntax, with the limitation of static fields:

```
java.math.BigDecimal.ONE;           # does not work
java_get ("java.math.BigDecimal", "ONE"); # workaround
```

Accordingly, with [java_set], page 992, the value of a field can be set. Note that only public Java fields are accessible from within Octave.

The following example indicates that in Octave empty brackets [] represent Java's `null` value and how Java exceptions are represented.

```
javaObject ("java.math.BigDecimal", []);
⇒ error: [java] java.lang.NullPointerException
```

It is not recommended to represent Java's `null` value by empty brackets [], because `null` has no type whereas [] has type `double`.

In Octave it is possible to provide limited Java reflection by listing the public fields and methods of a Java object, both static or not.

```
fieldnames (<Java object>)
methods (<Java object>)
```

Finally, an examples is shown how to access the stack trace from within Octave, where the function [debug_java], page 996, is used to set and to get the current debug state. In debug mode, the Java error and the stack trace are displayed.

```
debug_java (true) # use "false" to omit display of stack trace
debug_java ()
⇒ ans = 1
javaObject ("java.math.BigDecimal", "1") ...
    .divide (javaObject ("java.math.BigDecimal", "0"))
```

A.4.3 Set up the JVM

In order to execute Java code Octave creates a Java Virtual Machine (JVM). By default the version of the JVM is used that was detected during configuration on Unix-like systems or that is pointed to from the registry keys at `HKEY_LOCAL_MACHINE\SOFTWARE\JavaSoft\JRE` or `HKEY_LOCAL_MACHINE\SOFTWARE\JavaSoft\Java Runtime Environment` on Windows. The default path to the JVM can be overridden by setting the environment variable `JAVA_HOME` to the path where the JVM is installed. On Windows that might be, for example, `C:\Program Files\Java\jre-10.0.2`. Make sure that you select a directory that contains the JVM with a bit-ness that matches Octave's.

The JVM is only loaded once per Octave session. Thus, to change the used version of the JVM, you might have to re-start Octave. To check which version of the JVM is currently being used, run `version -java`.

The JVM allocates a fixed amount of initial memory and may expand this pool up to a fixed maximum memory limit. The default values depend on the Java version (see

[javamem], page 995). The memory pool is shared by all Java objects running in the JVM. This strict memory limit is intended mainly to avoid runaway applications inside web browsers or in enterprise servers which can consume all memory and crash the system. When the maximum memory limit is hit, Java code will throw exceptions so that applications will fail or behave unexpectedly.

You can specify options for the creation of the JVM inside a file named `java.opts`. This is a text file where you enter lines containing `-X` and `-D` options that are then passed to the JVM during initialization.

The directory where the Java options file is located is specified by the environment variable `OCTAVE_JAVA_DIR`. If unset the directory where `javaclasspath.m` resides is used instead (typically `OCTAVE_HOME/share/octave/OCTAVE_VERSION/m/java/`). You can find this directory by executing

```
which javaclasspath
```

The `-X` options allow you to increase the maximum amount of memory available to the JVM. The following example allows up to 256 Megabytes to be used by adding the following line to the `java.opts` file:

```
-Xmx256m
```

The maximum possible amount of memory depends on your system. On a Windows system with 2 Gigabytes main memory you should be able to set this maximum to about 1 Gigabyte.

If your application requires a large amount of memory from the beginning, you can also specify the initial amount of memory allocated to the JVM. Adding the following line to the `java.opts` file starts the JVM with 64 Megabytes of initial memory:

```
-Xms64m
```

For more details on the available `-X` options of your Java Virtual Machine issue the command ‘`java -X`’ at the operating system command prompt and consult the Java documentation.

The `-D` options can be used to define system properties which can then be used by Java classes inside Octave. System properties can be retrieved by using the `getProperty()` methods of the `java.lang.System` class. The following example line defines the property `MyProperty` and assigns it the string `12.34`.

```
-DMyProperty=12.34
```

The value of this property can then be retrieved as a string by a Java object or in Octave:

```
octave> javaMethod ("getProperty", "java.lang.System", "MyProperty");
ans = 12.34
```

See also: `javamem`.

A.4.4 Java Interface Functions

The following functions are the core of the Java Interface. They provide a way to create a Java object, get and set its data fields, and call Java methods which return results to Octave.

```
jobj = javaObject (classname)
jobj = javaObject (classname, arg1, ...)
```

Create a Java object of class *classname*, by calling the class constructor with the arguments *arg1*,

The first example below creates an uninitialized object, while the second example supplies an initial argument to the constructor.

```
x = javaObject ("java.lang.StringBuffer")
x = javaObject ("java.lang.StringBuffer", "Initial string")
```

See also: [\[javaMethod\]](#), page 993, [\[javaArray\]](#), page 991.

```
jary = javaArray (classname, sz)
jary = javaArray (classname, m, n, ...)
```

Create a Java array of size *sz* with elements of class *classname*.

classname may be a Java object representing a class or a string containing the fully qualified class name. The size of the object may also be specified with individual integer arguments *m*, *n*, etc.

The generated array is uninitialized. All elements are set to null if *classname* is a reference type, or to a default value (usually 0) if *classname* is a primitive type.

Sample code:

```
jary = javaArray ("java.lang.String", 2, 2);
jary(1,1) = "Hello";
```

See also: [\[javaObject\]](#), page 990.

There are many different variable types in Octave, but only ones created through `javaObject` can use Java functions. Before using Java with an unknown object the type can be checked with `isjava`.

```
isjava (x)
```

Return true if *x* is a Java object.

See also: [\[class\]](#), page 39, [\[typeinfo\]](#), page 39, [\[isa\]](#), page 39, [\[javaObject\]](#), page 990.

Once an object has been created it is natural to find out what fields the object has, and to read (get) and write (set) them.

In Octave the `fieldnames` function for structures has been overloaded to return the fields of a Java object. For example:

```

dobj = javaObject ("java.lang.Double", pi);
fieldnames (dobj)
⇒
{
  [1,1] = public static final double java.lang.Double.POSITIVE_INFINITY
  [1,2] = public static final double java.lang.Double.NEGATIVE_INFINITY
  [1,3] = public static final double java.lang.Double.NaN
  [1,4] = public static final double java.lang.Double.MAX_VALUE
  [1,5] = public static final double java.lang.Double.MIN_NORMAL
  [1,6] = public static final double java.lang.Double.MIN_VALUE
  [1,7] = public static final int java.lang.Double.MAX_EXPONENT
  [1,8] = public static final int java.lang.Double.MIN_EXPONENT
  [1,9] = public static final int java.lang.Double.SIZE
  [1,10] = public static final java.lang.Class java.lang.Double.TYPE
}

```

The analogy of objects with structures is carried over into reading and writing object fields. To read a field the object is indexed with the ‘.’ operator from structures. This is the preferred method for reading fields, but Octave also provides a function interface to read fields with `java_get`. An example of both styles is shown below.

```

dobj = javaObject ("java.lang.Double", pi);
dobj.MAX_VALUE
⇒ 1.7977e+308
java_get ("java.lang.Float", "MAX_VALUE")
⇒ 3.4028e+38

val = java_get (obj, name)
Get the value of the field name of the Java object obj.

```

For static fields, *obj* can be a string representing the fully qualified name of the corresponding class.

When *obj* is a regular Java object, structure-like indexing can be used as a shortcut syntax. For instance, the following two statements are equivalent

```

java_get (x, "field1")
x.field1

```

See also: [\[java_set\]](#), page 992, [\[javaMethod\]](#), page 993, [\[javaObject\]](#), page 990.

```

obj = java_set (obj, name, val)

```

Set the value of the field *name* of the Java object *obj* to *val*.

For static fields, *obj* can be a string representing the fully qualified name of the corresponding Java class.

When *obj* is a regular Java object, structure-like indexing can be used as a shortcut syntax. For instance, the following two statements are equivalent

```

java_set (x, "field1", val)
x.field1 = val

```

See also: [\[java_get\]](#), page 992, [\[javaMethod\]](#), page 993, [\[javaObject\]](#), page 990.

To see what functions can be called with an object use `methods`. For example, using the previously created `dobj`:

```
methods (dobj)
⇒
Methods for class java.lang.Double:
boolean equals(java.lang.Object)
java.lang.String toString(double)
java.lang.String toString()
...
...
```

To call a method of an object the same structure indexing operator ‘.’ is used. Octave also provides a functional interface to calling the methods of an object through `javaMethod`. An example showing both styles is shown below.

```
dobj = javaObject ("java.lang.Double", pi);
dobj.equals (3)
⇒ 0
javaMethod ("equals", dobj, pi)
⇒ 1

ret = javaMethod (methodname, obj)
ret = javaMethod (methodname, obj, arg1, ...)
    Invoke the method methodname on the Java object obj with the arguments arg1,
    ...
For static methods, obj can be a string representing the fully qualified name of the corresponding class.
```

When *obj* is a regular Java object, structure-like indexing can be used as a shortcut syntax. For instance, the two following statements are equivalent

```
ret = javaMethod ("method1", x, 1.0, "a string")
ret = x.method1 (1.0, "a string")
```

`javaMethod` returns the result of the method invocation.

See also: [\[methods\]](#), page 837, [\[javaObject\]](#), page 990.

The following three functions are used to display and modify the class path used by the Java Virtual Machine. This is entirely separate from Octave’s PATH variable and is used by the JVM to find the correct code to execute.

```
javaclasspath ()
dpath = javaclasspath ()
[dpath, spath] = javaclasspath ()
clspath = javaclasspath (what)
```

Return the class path of the Java virtual machine in the form of a cell array of strings.

If called with no inputs:

- If no output is requested, the dynamic and static classpaths are printed to the standard output.
- If one output value *dpath* is requested, the result is the dynamic classpath.

- If two output values *dpath* and *spath* are requested, the first variable will contain the dynamic classpath and the second will contain the static classpath.

If called with a single input parameter *what*:

"dynamic"

Return the dynamic classpath.

"static"

Return the static classpath.

"all"

Return both the static and dynamic classpath in a single cellstr.

See also: [\[javaaddpath\]](#), page 994, [\[javarmpath\]](#), page 994.

javaaddpath (*clspath*)

javaaddpath (*clspath1*, ...)

Add *clspath* to the dynamic class path of the Java virtual machine.

clspath may either be a directory where `.class` files are found, or a `.jar` file containing Java classes. Multiple paths may be added at once by specifying additional arguments.

See also: [\[javarmpath\]](#), page 994, [\[javaclasspath\]](#), page 993.

javarmpath (*clspath*)

javarmpath (*clspath1*, ...)

Remove *clspath* from the dynamic class path of the Java virtual machine.

clspath may either be a directory where `.class` files are found, or a `.jar` file containing Java classes. Multiple paths may be removed at once by specifying additional arguments.

See also: [\[javaaddpath\]](#), page 994, [\[javaclasspath\]](#), page 993.

The following functions provide information and control over the interface between Octave and the Java Virtual Machine.

javachk (*feature*)

javachk (*feature*, *caller*)

msg = javachk (...)

Check for the presence of the Java *feature* in the current session. Return an error structure if *feature* is not available, not enabled, or not recognized.

Possible recognized features are:

"awt" Abstract Window Toolkit for GUIs.

"desktop"

Interactive desktop is running.

"jvm"

Java Virtual Machine.

"swing"

Swing components for lightweight GUIs.

If *feature* is not supported, a scalar struct with field "**message**" and "**identifier**" is returned. The field "**message**" contains an error message mentioning *feature* as well as the optional user-specified *caller*. This structure is suitable for passing in to the **error** function.

If *feature* is supported and available, an empty struct array is returned with fields "**message**" and "**identifier**".

javachk determines if specific Java features are available in an Octave session. This function is provided for scripts which may alter their behavior based on the availability of Java or specific Java runtime features.

Compatibility Note: The feature "**desktop**" is never available since Octave has no Java-based desktop.

See also: [\[usejava\]](#), page 995, [\[error\]](#), page 221.

usejava (*feature*)

Return true if the Java element *feature* is available.

Possible features are:

"**awt**" Abstract Window Toolkit for GUIs.

"**desktop**" Interactive desktop is running.

"**jvm**" Java Virtual Machine.

"**swing**" Swing components for lightweight GUIs.

usejava determines if specific Java features are available in an Octave session. This function is provided for scripts which may alter their behavior based on the availability of Java. The feature "**desktop**" always returns **false** as Octave has no Java-based desktop. Other features may be available if Octave was compiled with the Java Interface and Java is installed.

See also: [\[javachk\]](#), page 994.

javamem ()

***jmem* = javamem ()**

Show the current memory usage of the Java virtual machine (JVM) and run the garbage collector.

When no return argument is given the info is printed to the screen. Otherwise, the output cell array *jmem* contains Maximum, Total, and Free memory (in bytes).

All Java-based routines are run in the JVM's shared memory pool, a dedicated and separate part of memory claimed by the JVM from your computer's total memory (which comprises physical RAM and virtual memory / swap space on hard disk).

The maximum allowable memory usage can be configured using the file **java.opts**. The directory where this file resides is determined by the environment variable **OCTAVE_JAVA_DIR**. If unset, the directory where **javaaddpath.m** resides is used instead (typically **OCTAVE_HOME/share/octave/OCTAVE_VERSION/m/java/**).

java.opts is a plain text file with one option per line. The default initial memory size and default maximum memory size (which are both system dependent) can be overridden like so:

```
-Xms64m  
-Xmx512m
```

(in megabytes in this example). You can adapt these values to your own requirements if your system has limited available physical memory or if you get Java memory errors.

"**Total memory**" is what the operating system has currently assigned to the JVM and depends on actual and active memory usage. "**Free memory**" is self-explanatory. During operation of Java-based Octave functions the amount of Total and Free memory will vary, due to Java's own cleaning up and your operating system's memory management.

```
val = java_matrix_autoconversion ()  
old_val = java_matrix_autoconversion (new_val)  
java_matrix_autoconversion (new_val, "local")
```

Query or set the internal variable that controls whether Java arrays are automatically converted to Octave matrices.

The default value is false.

When called from inside a function with the "**local**" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[java_unsigned_autoconversion\]](#), page 996, [\[debug_java\]](#), page 996.

```
val = java_unsigned_autoconversion ()  
old_val = java_unsigned_autoconversion (new_val)  
java_unsigned_autoconversion (new_val, "local")
```

Query or set the internal variable that controls how integer classes are converted when `java_matrix_autoconversion` is enabled.

When enabled, Java arrays of class Byte or Integer are converted to matrices of class uint8 or uint32 respectively. The default value is true.

When called from inside a function with the "**local**" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[java_matrix_autoconversion\]](#), page 996, [\[debug_java\]](#), page 996.

```
val = debug_java ()  
old_val = debug_java (new_val)  
debug_java (new_val, "local")
```

Query or set the internal variable that determines whether extra debugging information regarding the initialization of the JVM and any Java exceptions is printed.

When called from inside a function with the "**local**" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[java_matrix_autoconversion\]](#), page 996, [\[java_unsigned_autoconversion\]](#), page 996.

Appendix B Test and Demo Functions

Octave includes a number of functions to allow the integration of testing and demonstration code in the source code of the functions themselves.

B.1 Test Functions

```
test name
test name quiet|normal|verbose
test ("name", "quiet|normal|verbose", fid)
test ("name", "quiet|normal|verbose", fname)
success = test (...)

[n, nmax, nxfail, nbug, nskip, nrtskip, nregression] = test (...)

[code, idx] = test ("name", "grabdemo")
test ([] , "explain", fid)
test ([] , "explain", fname)
```

Perform built-in self-tests from the first file in the loadpath matching *name*.

test can be called in either command or functional form. The exact operation of **test** is determined by a combination of mode (interactive or batch), reporting level ("quiet", "normal", "verbose"), and whether a logfile or summary output variable is used.

The default mode when **test** is called from the command line is interactive. In this mode, tests will be run until the first error is encountered, or all tests complete successfully. In batch mode, all tests are run regardless of any failures, and the results are collected for reporting. Tests which require user interaction, i.e., demo blocks, are never run in batch mode.

Batch mode is enabled by either 1) specifying a logfile using the third argument *fname* or *fid*, or 2) requesting an output argument such as *success*, *n*, etc.

The optional second argument determines the amount of output to generate and which types of tests to run. The default value is "normal". Requesting an output argument will suppress printing the final summary message and any intermediate warnings, unless verbose reporting is enabled.

"quiet" Print a summary message when all tests pass, or print an error with the results of the first bad test when a failure occurs. Don't run tests which require user interaction.

"normal" Display warning messages about skipped tests or failing xtests during test execution. Print a summary message when all tests pass, or print an error with the results of the first bad test when a failure occurs. Don't run tests which require user interaction.

"verbose" Display tests before execution. Print all warning messages. In interactive mode, run all tests including those which require user interaction.

The optional third input argument specifies a logfile where results of the tests should be written. The logfile may be a character string (*fname*) or an open file descriptor

ID (*fid*). To enable batch processing, but still print the results to the screen, use `stdout` for *fid*.

When called with just a single output argument `success`, `test` returns true if all of the tests were successful. If called with more than one output argument then the number of successful tests (*n*), the total number of tests in the file (*nmax*), the number of xtest failures (*nxfail*), the number of tests failed due known bugs (*nbug*), the number of tests skipped due to missing features (*nskip*), the number of tests skipped due to run-time conditions (*nrtskip*), and the number of regressions (*nregression*) are returned.

Example

```
test sind
⇒
PASSES 5 out of 5 tests

[n, nmax] = test ("sind")
⇒
n = 5
nmax = 5
```

Additional Calling Syntaxes

If the second argument is the string "`grabdemo`", the contents of any built-in demo blocks are extracted but not executed. The text for all code blocks is concatenated and returned as `code` with `idx` being a vector of positions of the ends of each demo block. For an easier way to extract demo blocks from files, See [example], page 1006.

If the second argument is "`explain`" then `name` is ignored and an explanation of the line markers used in `test` output reports is written to the file specified by `fname` or *fid*.

See also: [assert], page 1003, [fail], page 1004, [demo], page 1005, [example], page 1006, [error], page 221.

`test` scans the named script file looking for lines which start with the identifier '`%!`'. The prefix is stripped off and the rest of the line is processed through the Octave interpreter. If the code generates an error, then the test is said to fail.

Since `eval()` will stop at the first error it encounters, you must divide your tests up into blocks, with anything in a separate block evaluated separately. Blocks are introduced by valid keywords like `test`, `function`, or `assert` immediately following '`%!`'. A block is defined by indentation as in Python. Lines beginning with '`%!<whitespace>`' are part of the preceeding block.

For example:

```
%!test error ("this test fails!")
%!test "test doesn't fail. it doesn't generate an error"
```

When a test fails, you will see something like:

```
***** test error ("this test fails!")
!!!! test failed
this test fails!
```

Generally, to test if something works, you want to assert that it produces a correct value. A real test might look something like

```
%!test
%! a = [1, 2, 3; 4, 5, 6]; B = [1; 2];
%! expect = [ a ; 2*a ];
%! get = kron (b, a);
%! if (any (size (expect) != size (get)))
%!   error ("wrong size: expected %d,%d but got %d,%d",
%!         size (expect), size (get));
%! elseif (any (any (expect != get)))
%!   error ("didn't get what was expected.");
%! endif
```

To make the process easier, use the `assert` function. For example, with `assert` the previous test is reduced to:

```
%!test
%! a = [1, 2, 3; 4, 5, 6]; b = [1; 2];
%! assert (kron (b, a), [ a; 2*a ]);
```

`assert` can accept a tolerance so that you can compare results absolutely or relatively. For example, the following all succeed:

```
%!test assert (1+eps, 1, 2*eps)           # absolute error
%!test assert (100+100*eps, 100, -2*eps)  # relative error
```

You can also do the comparison yourself, but still have `assert` generate the error:

```
%!test assert (isempty ([]))
%!test assert ([1, 2; 3, 4] > 0)
```

Because `assert` is so frequently used alone in a test block, there is a shorthand form:

```
%!assert (...)
```

which is equivalent to:

```
%!test assert (...)
```

Occasionally a block of tests will depend on having optional functionality in Octave. Before testing such blocks the availability of the required functionality must be checked. A `%!testif HAVE_XXX` block will only be run if Octave was compiled with functionality ‘HAVE_XXX’. For example, the sparse single value decomposition, `svds()`, depends on having the ARPACK library. All of the tests for `svds` begin with

```
%!testif HAVE_ARPACK
```

Review `config.h` or `__octave_config_info__ ("build_features")` to see some of the possible values to check.

Sometimes during development there is a test that should work but is known to fail. You still want to leave the test in because when the final code is ready the test should pass, but you may not be able to fix it immediately. To avoid unnecessary bug reports for these known failures, mark the block with `xtest` rather than `test`:

```
%!xtest assert (1==0)
%!xtest fail ("success=1", "error")
```

In this case, the test will run and any failure will be reported. However, testing is not aborted and subsequent test blocks will be processed normally. Another use of `xtest` is for statistical tests which should pass most of the time but are known to fail occasionally.

Each block is evaluated in its own function environment, which means that variables defined in one block are not automatically shared with other blocks. If you do want to share variables, then you must declare them as `shared` before you use them. For example, the following declares the variable `a`, gives it an initial value (default is empty), and then uses it in several subsequent tests.

```
%!shared a
%! a = [1, 2, 3; 4, 5, 6];
%!assert (kron ([1; 2], a), [ a; 2*a ])
%!assert (kron ([1, 2], a), [ a, 2*a ])
%!assert (kron ([1,2; 3,4], a), [ a,2*a; 3*a,4*a ])
```

You can share several variables at the same time:

```
%!shared a, b
```

You can also share test functions:

```
%!function a = fn (b)
%! a = 2*b;
%!endfunction
%!assert (fn(2), 4)
```

Note that all previous variables and values are lost when a new shared block is declared.

Remember that `%!function` begins a new block and that `%!endfunction` ends this block. Be aware that until a new block is started, lines starting with '`%!<space>`' will be discarded as comments. The following is nearly identical to the example above, but does nothing.

```
%!function a = fn (b)
%! a = 2*b;
%!endfunction
%! assert (fn(2), 4)
```

Because there is a space after '`%!`' the `assert` statement does not begin a new block and this line is treated as a comment.

Error and warning blocks are like test blocks, but they only succeed if the code generates an error. You can check the text of the error is correct using an optional regular expression `<pattern>`. For example:

```
%!error <passes!> error ("this test passes!")
```

If the code doesn't generate an error, the test fails. For example:

```
%!error "this is an error because it succeeds."
```

produces

```
***** error "this is an error because it succeeds."
!!!! test failed: no error
```

It is important to automate the tests as much as possible, however some tests require user interaction. These can be isolated into demo blocks, which if you are in batch mode,

are only run when called with `demo` or the `verbose` option to `test`. The code is displayed before it is executed. For example,

```
%!demo
%! t = [0:0.01:2*pi]; x = sin (t);
%! plot (t, x);
%! # you should now see a sine wave in your figure window
```

produces

```
funcname example 1:
t = [0:0.01:2*pi]; x = sin (t);
plot (t, x);
# you should now see a sine wave in your figure window
```

Press <enter> to continue:

Note that demo blocks cannot use any shared variables. This is so that they can be executed by themselves, ignoring all other tests.

If you want to temporarily disable a test block, put `#` in place of the block type. This creates a comment block which is echoed in the log file but not executed. For example:

```
%!#demo
%! t = [0:0.01:2*pi]; x = sin (t);
%! plot (t, x);
%! # you should now see a sine wave in your figure window
```

The following trivial code snippet provides examples for the use of `fail`, `assert`, `error`, and `xtest`:

```
function output = must_be_zero (input)
if (input != 0)
    error ("Nonzero input!")
endif
output = input;
endfunction

%!fail ("must_be_zero (1)")
%!assert (must_be_zero (0), 0)
%!error <Nonzero> must_be_zero (1)
%!xtest error ("This code generates an error")
```

When putting this in a file `must_be_zero.m`, and running the test, we see

```
test must_be_zero verbose

⇒
>>> /path/to/must_be_zero.m
***** fail ("must_be_zero (1)")
***** assert (must_be_zero (0), 0)
***** error <Nonzero> must_be_zero (1)
***** xtest error ("This code generates an error")
!!!! known failure
This code generates an error
PASSES 3 out of 4 tests (1 expected failure)
```

Block type summary:

`%!test`

`%!test <MESSAGE>`

Check that entire block is correct. If `<MESSAGE>` is present, the test block is interpreted as for `xtest`.

`%!testif HAVE_XXX`

`%!testif HAVE_XXX, HAVE_YYY, ...`

`%!testif HAVE_XXX, HAVE_YYY ...; RUNTIME_COND`

`%!testif ... <MESSAGE>`

Check block only if Octave was compiled with feature `HAVE_XXX`. `RUNTIME_COND` is an optional expression to evaluate to check whether some condition is met when the test is executed. If `RUNTIME_COND` is false, the test is skipped. If `<MESSAGE>` is present, the test block is interpreted as for `xtest`.

`%!xtest`

`%!xtest <MESSAGE>`

Check block, report a test failure but do not abort testing. If `<MESSAGE>` is present, then the text of the message is displayed if the test fails, like this:

!!!! Known bug: `MESSAGE`

If the message is an integer, it is interpreted as a bug ID for the Octave bug tracker and reported as

!!!! Known bug: <https://octave.org/testfailure/?BUG-ID>

in which `BUG-ID` is the integer bug number. The intent is to allow clearer documentation of known problems.

`%!error`

`%!error <MESSAGE>`

`%!warning`

`%!warning <MESSAGE>`

Check for correct error or warning message. If `<MESSAGE>` is supplied it is interpreted as a regular expression pattern that is expected to match the error or warning message.

`%!demo` Demo only executes in interactive mode.

`%!#` Comment. Ignore everything within the block

```
%!shared x,y,z
    Declare variables for use in multiple tests.

%!function
    Define a function for use in multiple tests.

%!endfunction
    Close a function definition.

%!assert (x, y, tol)
%!assert <MESSAGE> (x, y, tol)
%!fail (CODE, PATTERN)
%!fail <MESSAGE> (CODE, PATTERN)
    Shorthand for %!test assert (x, y, tol) or %!test fail (CODE, PATTERN).
    If <MESSAGE> is present, the test block is interpreted as for xtest.
```

When coding tests the Octave convention is that lines that begin with a block type do not have a semicolon at the end. Any code that is within a block, however, is normal Octave code and usually will have a trailing semicolon. For example,

```
## bare block instantiation
%!assert (sin (0), 0)
```

but

```
## test block with normal Octave code
%!test
  %! assert (sin (0), 0);
```

You can also create test scripts for built-in functions and your own C++ functions. To do so, put a file with the bare function name (no .m extension) in a directory in the load path and it will be discovered by the `test` function. Alternatively, you can embed tests directly in your C++ code:

```
/*
%!test disp ("this is a test")
*/
```

or

```
#if 0
  %!test disp ("this is a test")
#endif
```

However, in this case the raw source code will need to be on the load path and the user will have to remember to type `test ("funcname.cc")`.

```
assert (cond)
assert (cond, errmsg)
assert (cond, errmsg, ...)
assert (cond, msg_id, errmsg, ...)
assert (observed, expected)
assert (observed, expected, tol)
```

Produce an error if the specified condition is not met.

`assert` can be called in three different ways.

```
assert (cond)
assert (cond, errmsg)
assert (cond, errmsg, ...)
assert (cond, msg_id, errmsg, ...)
```

Called with a single argument *cond*, `assert` produces an error if *cond* is false (numeric zero).

Any additional arguments are passed to the `error` function for processing.

```
assert (observed, expected)
```

Produce an error if *observed* is not the same as *expected*.

Note that *observed* and *expected* can be scalars, vectors, matrices, strings, cell arrays, or structures.

```
assert (observed, expected, tol)
```

Produce an error if *observed* is not the same as *expected* but equality comparison for numeric data uses a tolerance *tol*.

If *tol* is positive then it is an absolute tolerance which will produce an error if `abs (observed - expected) > abs (tol)`.

If *tol* is negative then it is a relative tolerance which will produce an error if `abs (observed - expected) > abs (tol * expected)`.

If *expected* is zero *tol* will always be interpreted as an absolute tolerance.

If *tol* is not scalar its dimensions must agree with those of *observed* and *expected* and tests are performed on an element-by-element basis.

See also: [\[fail\]](#), page 1004, [\[test\]](#), page 997, [\[error\]](#), page 221, [\[isequal\]](#), page 153.

```
fail (code)
fail (code, pattern)
fail (code, "warning")
fail (code, "warning", pattern)
```

Return true if *code* fails with an error message matching *pattern*, otherwise produce an error.

code must be in the form of a string that is passed to the Octave interpreter via the `evalin` function, i.e., a (quoted) string constant or a string variable.

Note that if *code* runs successfully, rather than failing, the error printed is:

```
expected error <..> but got none
```

If called with two arguments, the return value will be true only if *code* fails with an error message containing *pattern* (case sensitive). If the code fails with a different error than the one specified in *pattern* then the message produced is:

```
expected <pattern>
but got <text of actual error>
```

The angle brackets are not part of the output.

When called with the "warning" option `fail` will produce an error if executing the code produces no warning.

See also: [\[assert\]](#), page 1003, [\[error\]](#), page 221.

B.2 Demonstration Functions

```
demo name
demo name n
demo ("name")
demo ("name", n)
```

Run example code block *n* associated with the function *name*.

If *n* is not specified, all examples are run.

The preferred location for example code blocks is embedded within the script m-file immediately following the code that it exercises. Alternatively, the examples may be stored in a file with the same name but no extension located on Octave's load path. To separate examples from regular script code all lines are prefixed by `%!`. Each example must also be introduced by the keyword "demo" flush left to the prefix with no intervening spaces. The remainder of the example can contain arbitrary Octave code. For example:

```
%!demo
%! t = 0:0.01:2*pi;
%! x = sin (t);
%! plot (t, x);
%! title ("one cycle of a sine wave");
%! #-----
%! # the figure window shows one cycle of a sine wave
```

Note that the code is displayed before it is executed so that a simple comment at the end suffices for labeling what is being shown. For plots, labeling can also be done with `title` or `text`. It is generally **not** necessary to use `disp` or `printf` within the demo.

Demos are run in a stand-alone function environment with no access to external variables. This means that every demo must have separate initialization code. Alternatively, all demos can be combined into a single large demo with the code

```
%! input ("Press <enter> to continue: ", "s");
```

between the sections, but this usage is discouraged. Other techniques to avoid multiple initialization blocks include using multiple plots with a new `figure` command between each plot, or using `subplot` to put multiple plots in the same window.

Finally, because `demo` evaluates within a function context it is not possible to define new functions within the code. Anonymous functions make a good substitute in most instances. If function blocks **must** be used then the code `eval (example ("function", n))` will allow Octave to see them. This has its own problems, however, as `eval` only evaluates one line or statement at a time. In this case the function declaration must be wrapped with "`if 1 <demo stuff> endif`" where "if" is on the same line as "demo". For example:

```
%!demo if 1
%! function y = f(x)
%!     y = x;
%! endfunction
%! f(3)
%! endif
```

See also: [rundemos], page 1006, [example], page 1006, [test], page 997.

```
example name
example name n
example ("name")
example ("name", n)
[s, idx] = example (...)
```

Display the code for example *n* associated with the function *name*, but do not run it.

If *n* is not specified, all examples are displayed.

When called with output arguments, the examples are returned in the form of a string *s*, with *idx* indicating the ending position of the various examples.

See **demo** for a complete explanation.

See also: [demo], page 1005, [test], page 997.

```
rundemos ()
rundemos (directory)
```

Execute built-in demos for all m-files in the specified *directory*.

Demo blocks in any C++ source files (*.cc) will also be executed for use with dynamically linked oct-file functions.

If no directory is specified, operate on all directories in Octave's search path for functions.

See also: [demo], page 1005, [runtests], page 1006, [path], page 197.

```
runtests ()
runtests (directory)
```

Execute built-in tests for all m-files in the specified *directory*.

Test blocks in any C++ source files (*.cc) will also be executed for use with dynamically linked oct-file functions.

If no directory is specified, operate on all directories in Octave's search path for functions.

See also: [rundemos], page 1006, [test], page 997, [path], page 197.

```
speed (f, init, max_n, f2, tol)
[order, n, T_f, T_f2] = speed (...)
```

Determine the execution time of an expression (*f*) for various input values (*n*).

The *n* are log-spaced from 1 to *max_n*. For each *n*, an initialization expression (*init*) is computed to create any data needed for the test. If a second expression (*f2*) is given then the execution times of the two expressions are compared. When called without output arguments the results are printed to stdout and displayed graphically.

f The code expression to evaluate.

max_n The maximum test length to run. The default value is 100. Alternatively, use [*min_n*, *max_n*] or specify the *n* exactly with [*n1*, *n2*, ..., *nk*].

init	Initialization expression for function argument values. Use <i>k</i> for the test number and <i>n</i> for the size of the test. This should compute values for all variables used by <i>f</i> . Note that <i>init</i> will be evaluated first for <i>k</i> = 0, so things which are constant throughout the test series can be computed once. The default value is <code>x = randn (n, 1)</code> .
f2	An alternative expression to evaluate, so that the speed of two expressions can be directly compared. The default is <code>[]</code> .
tol	Tolerance used to compare the results of expression <i>f</i> and expression <i>f2</i> . If <i>tol</i> is positive, the tolerance is an absolute one. If <i>tol</i> is negative, the tolerance is a relative one. The default is <code>eps</code> . If <i>tol</i> is <code>Inf</code> , then no comparison will be made.
order	The time complexity of the expression $O(a * n^p)$. This is a structure with fields <i>a</i> and <i>p</i> .
n	The values <i>n</i> for which the expression was calculated AND the execution time was greater than zero.
T_f	The nonzero execution times recorded for the expression <i>f</i> in seconds.
T_f2	The nonzero execution times recorded for the expression <i>f2</i> in seconds. If required, the mean time ratio is simply <code>mean (T_f ./ T_f2)</code> .

The slope of the execution time graph shows the approximate power of the asymptotic running time $O(n^p)$. This power is plotted for the region over which it is approximated (the latter half of the graph). The estimated power is not very accurate, but should be sufficient to determine the general order of an algorithm. It should indicate if, for example, the implementation is unexpectedly $O(n^2)$ rather than $O(n)$ because it extends a vector each time through the loop rather than pre-allocating storage. In the current version of Octave, the following is not the expected $O(n)$.

```
speed ("for i = 1:n, y{i} = x(i); endfor", "", [1000, 10000])
```

But it is if you preallocate the cell array *y*:

```
speed ("for i = 1:n, y{i} = x(i); endfor", ...
"x = rand (n, 1); y = cell (size (x));", [1000, 10000])
```

An attempt is made to approximate the cost of individual operations, but it is wildly inaccurate. You can improve the stability somewhat by doing more work for each *n*. For example:

```
speed ("airy(x)", "x = rand (n, 10)", [10000, 100000])
```

When comparing two different expressions (*f*, *f2*), the slope of the line on the speedup ratio graph should be larger than 1 if the new expression is faster. Better algorithms have a shallow slope. Generally, vectorizing an algorithm will not change the slope of the execution time graph, but will shift it relative to the original. For example:

```
speed ("sum (x)", "", [10000, 100000], ...
"v = 0; for i = 1:length (x), v += x(i); endfor")
```

The following is a more complex example. If there was an original version of `xcorr` using for loops and a second version using an FFT, then one could compare the run

speed for various lags as follows, or for a fixed lag with varying vector lengths as follows:

```
speed ("xcorr (x, n)", "x = rand (128, 1);", 100,
      "xcorr_orig (x, n)", -100*eps)
speed ("xcorr (x, 15)", "x = rand (20+n, 1);", 100,
      "xcorr_orig (x, n)", -100*eps)
```

Assuming one of the two versions is in `xcorr_orig`, this would compare their speed and their output values. Note that the FFT version is not exact, so one must specify an acceptable tolerance on the comparison `100*eps`. In this case, the comparison should be computed relatively, as `abs ((x - y) ./ y)` rather than absolutely as `abs (x - y)`. Type `example ("speed")` to see some real examples or `demo ("speed")` to run them.

Appendix C Obsolete Functions

After being marked as deprecated for two major releases, the following functions have been removed from Octave. The third column of the table shows the version of Octave in which the function was removed. Prior to removal, each function in the list was marked as deprecated for at least two major releases. All deprecated functions issue warnings explaining that they will be removed in a future version of Octave, and which function should be used instead.

Replacement functions do not always accept precisely the same arguments as the obsolete function, but should provide equivalent functionality.

Obsolete Function	Replacement	Version
<code>beta_cdf</code>	<code>betacdf</code> in Octave Forge statistics pkg	3.4.0
<code>beta_inv</code>	<code>betainv</code> in Octave Forge statistics pkg	3.4.0
<code>beta_pdf</code>	<code>betapdf</code> in Octave Forge statistics pkg	3.4.0
<code>beta_rnd</code>	<code>betarnd</code> in Octave Forge statistics pkg	3.4.0
<code>binomial_cdf</code>	<code>binocdf</code> in Octave Forge statistics pkg	3.4.0
<code>binomial_inv</code>	<code>binoinv</code> in Octave Forge statistics pkg	3.4.0
<code>binomial_pdf</code>	<code>binopdf</code> in Octave Forge statistics pkg	3.4.0
<code>binomial_rnd</code>	<code>binornd</code> in Octave Forge statistics pkg	3.4.0
<code>chisquare_cdf</code>	<code>chi2cdf</code> in Octave Forge statistics pkg	3.4.0
<code>chisquare_inv</code>	<code>chi2inv</code> in Octave Forge statistics pkg	3.4.0
<code>chisquare_pdf</code>	<code>chi2pdf</code> in Octave Forge statistics pkg	3.4.0
<code>chisquare_rnd</code>	<code>chi2rnd</code> in Octave Forge statistics pkg	3.4.0
<code>clearplot</code>	<code>clf</code>	3.4.0
<code>com2str</code>	<code>num2str</code>	3.4.0
<code>exponential_cdf</code>	<code>expcdf</code> in Octave Forge statistics pkg	3.4.0
<code>exponential_inv</code>	<code>expinv</code> in Octave Forge statistics pkg	3.4.0
<code>exponential_pdf</code>	<code>exppdf</code> in Octave Forge statistics pkg	3.4.0

exponential_rnd	exprnd in Octave Forge statistics pkg	3.4.0
f_cdf	fcdf in Octave Forge statistics pkg	3.4.0
f_inv	finv in Octave Forge statistics pkg	3.4.0
f_pdf	fpdf in Octave Forge statistics pkg	3.4.0
f_rnd	frnd in Octave Forge statistics pkg	3.4.0
gamma_cdf	gamcdf in Octave Forge statistics pkg	3.4.0
gamma_inv	gaminv in Octave Forge statistics pkg	3.4.0
gamma_pdf	gampdf in Octave Forge statistics pkg	3.4.0
gamma_rnd	gamrnd in Octave Forge statistics pkg	3.4.0
geometric_cdf	geocdf in Octave Forge statistics pkg	3.4.0
geometric_inv	geoinv in Octave Forge statistics pkg	3.4.0
geometric_pdf	geopdf in Octave Forge statistics pkg	3.4.0
geometric_rnd	geornd in Octave Forge statistics pkg	3.4.0
hypergeometric_cdf	hygecdf in Octave Forge statistics pkg	3.4.0
hypergeometric_inv	hygeinv in Octave Forge statistics pkg	3.4.0
hypergeometric_pdf	hygepdf in Octave Forge statistics pkg	3.4.0
hypergeometric_rnd	hygernd in Octave Forge statistics pkg	3.4.0
intersection	intersect	3.4.0
is_bool	isbool	3.4.0
is_complex	iscomplex	3.4.0
is_list	None	3.4.0
is_matrix	ismatrix	3.4.0
is_scalar	isscalar	3.4.0
is_square	issquare	3.4.0
is_stream	None	3.4.0
is_struct	isstruct	3.4.0
is_symmetric	issymmetric	3.4.0
is_vector	isvector	3.4.0

lognormal_cdf	logncdf in Octave Forge statistics pkg	3.4.0
lognormal_inv	logninv in Octave Forge statistics pkg	3.4.0
lognormal_pdf	lognpdf in Octave Forge statistics pkg	3.4.0
lognormal_rnd	lognrnd in Octave Forge statistics pkg	3.4.0
meshdom	meshgrid	3.4.0
normal_cdf	normcdf in Octave Forge statistics pkg	3.4.0
normal_inv	norminv in Octave Forge statistics pkg	3.4.0
normal_pdf	normpdf in Octave Forge statistics pkg	3.4.0
normal_rnd	normrnd in Octave Forge statistics pkg	3.4.0
pascal_cdf	nbincdf in Octave Forge statistics pkg	3.4.0
pascal_inv	nbininv in Octave Forge statistics pkg	3.4.0
pascal_pdf	nbinpdf in Octave Forge statistics pkg	3.4.0
pascal_rnd	nbinrnd in Octave Forge statistics pkg	3.4.0
poisson_cdf	poisscdf in Octave Forge statistics pkg	3.4.0
poisson_inv	poissinv in Octave Forge statistics pkg	3.4.0
poisson_pdf	poisspdf in Octave Forge statistics pkg	3.4.0
poisson_rnd	poissrnd in Octave Forge statistics pkg	3.4.0
polyinteg	polyint	3.4.0
struct_contains	isfield	3.4.0
struct_elements	fieldnames	3.4.0
t_cdf	tcdf in Octave Forge statistics pkg	3.4.0
t_inv	tinv in Octave Forge statistics pkg	3.4.0
t_pdf	tpdf in Octave Forge statistics pkg	3.4.0
t_rnd	trnd in Octave Forge statistics pkg	3.4.0

uniform_cdf	unifcdf in Octave Forge statistics pkg	3.4.0
uniform_inv	unifinv in Octave Forge statistics pkg	3.4.0
uniform_pdf	unifpdf in Octave Forge statistics pkg	3.4.0
uniform_rnd	unifrnd in Octave Forge statistics pkg	3.4.0
weibull_cdf	wblcdf in Octave Forge statistics pkg	3.4.0
weibull_inv	wblinv in Octave Forge statistics pkg	3.4.0
weibull_pdf	wblpdf in Octave Forge statistics pkg	3.4.0
weibull_rnd	wblrnd in Octave Forge statistics pkg	3.4.0
wiener_rnd	wienrnd in Octave Forge statistics pkg	3.4.0
create_set	unique	3.6.0
dmult	diag (A) * B	3.6.0
iscommand	None	3.6.0
israwcommand	None	3.6.0
lchol	chol (... , "lower")	3.6.0
loadimage	load or imread	3.6.0
mark_as_command	None	3.6.0
mark_as_rawcommand	None	3.6.0
spatan2	atan2	3.6.0
spchol	chol	3.6.0
spchol2inv	chol2inv	3.6.0
spcholinv	cholinv	3.6.0
spcumprod	cumprod	3.6.0
spcumsum	cumsum	3.6.0
spdet	det	3.6.0
spdiag	sparse (diag (...))	3.6.0
spfind	find	3.6.0
sphcat	horzcat	3.6.0
spinv	inv	3.6.0
spkron	kron	3.6.0
splchol	chol (... , "lower")	3.6.0
split	char (strsplit (s, t))	3.6.0
splu	lu	3.6.0
spmax	max	3.6.0
spmin	min	3.6.0
spprod	prod	3.6.0
spqr	qr	3.6.0

spsum	sum	3.6.0
spsumsq	sumsq	3.6.0
spvcat	vertcat	3.6.0
str2mat	char	3.6.0
unmark_command	None	3.6.0
unmark_rawcommand	None	3.6.0
autocor	xcorr in Octave Forge signal pkg	3.8.0
autocov	xcov in Octave Forge signal pkg	3.8.0
betai	betainc	3.8.0
cellidx	ismember	3.8.0
cquad	quadcc	3.8.0
dispatch	None	3.8.0
fstat	stat	3.8.0
gammai	gammainc	3.8.0
glpkmex	glpk	3.8.0
is_duplicate_entry	unique	3.8.0
is_global	isglobal	3.8.0
krylovb	[Uret, ~, Ucols] = krylov (...)	3.8.0
perror	None	3.8.0
replot	refresh	3.8.0
saveimage	imwrite	3.8.0
setstr	char	3.8.0
strerror	None	3.8.0
values	unique	3.8.0
cut	histc	4.0.0
cor	corr	4.0.0
corrcoef	corr	4.0.0
_error_text__	lasterr	4.0.0
error_text	lasterr	4.0.0
polyderiv	polyder	4.0.0
shell_cmd	system	4.0.0
studentize	zscore	4.0.0
sylvester_matrix	hadamard (2^k)	4.0.0
default_save_options	save_default_options	4.2.0
gen_doc_cache	doc_cache_create	4.2.0
interp1q	interp1	4.2.0
isequalwithequalnans	isequaln	4.2.0
java_convert_matrix	java_matrix_ autoconversion	4.2.0
java_debug	debug_java	4.2.0
java_invoke	javaMethod	4.2.0
java_new	javaObject	4.2.0

java_unsigned_ conversion	java_unsigned_ autoconversion	4.2.0
javafields	fieldnames	4.2.0
javamethods	methods	4.2.0
re_read_readline_init_ file	readline_re_read_ init_file	4.2.0
read_readline_init_ file	readline_read_init_ file	4.2.0
saving_history	history_save	4.2.0
allow_noninteger_ range_as_index	None	4.4.0
bicubic	interp2	4.4.0
delaunay3	delaunay	4.4.0
do_braindead_ shortcircuit_ evaluation	None	4.4.0
dump_prefs	None 4.4.0	
find_dir_in_path	dir_in_loadpath	4.4.0
finite	isfinite	4.4.0
fmod	rem	4.4.0
fnmatch	glob or regexp	4.4.0
gmap40	None	4.4.0
loadaudio	audioread	4.4.0
luinc	ichol or ilu	4.4.0
mouse_wheel_zoom	mousehweelzoom property	4.4.0
nfields	numfields	4.4.0
octave_tmp_file_name	tempname	4.4.0
playaudio	audioplayer	4.4.0
saveaudio	audiowrite	4.4.0
setaudio	None	4.4.0
syl	sylvester	4.4.0
usage	print_usage	4.4.0
bitmax	flintmax	5.1.0
mahalanobis	mahal in Octave Forge statistics pkg	5.1.0
md5sum	hash	5.1.0
octave_config_info	__octave_config_ info__	5.1.0
onenormest	normest1	5.1.0
sleep	pause	5.1.0
usleep	pause	5.1.0
wavread	audioread	5.1.0
wavwrite	audiowrite	5.1.0

Appendix D Known Causes of Trouble

This section describes known problems that affect users of Octave. Most of these are not Octave bugs per se—if they were, we would fix them. But the result for a user may be like the result of a bug.

Some of these problems are due to bugs in other software, some are missing features that are too much work to add, and some are places where people’s opinions differ as to what is best.

D.1 Actual Bugs We Haven’t Fixed Yet

- Output that comes directly from Fortran functions is not sent through the pager and may appear out of sequence with other output that is sent through the pager. One way to avoid this is to force pending output to be flushed before calling a function that will produce output from within Fortran functions. To do this, use the command

```
fflush (stdout)
```

Another possible workaround is to use the command

```
page_screen_output (false);
```

to turn the pager off.

A list of ideas for future enhancements is distributed with Octave. See the file `PROJECTS` in the top level directory in the source distribution.

D.2 Reporting Bugs

Your bug reports play an essential role in making Octave reliable.

When you encounter a problem, the first thing to do is to see if it is already known. See [Appendix D \[Trouble\], page 1015](#). If it isn’t known, then you should report the problem.

Reporting a bug may help you by bringing a solution to your problem, or it may not. In any case, the principal function of a bug report is to help the entire community by making the next version of Octave work better. Bug reports are your contribution to the maintenance of Octave.

In order for a bug report to serve its purpose, you must include the information that makes it possible to fix the bug.

D.2.1 Have You Found a Bug?

If you are not sure whether you have found a bug, here are some guidelines:

- If Octave gets a fatal signal, for any input whatever, that is a bug. Reliable interpreters never crash.
- If Octave produces incorrect results, for any input whatever, that is a bug.
- Some output may appear to be incorrect when it is in fact due to a program whose behavior is undefined, which happened by chance to give the desired results on another system. For example, the range operator may produce different results because of differences in the way floating point arithmetic is handled on various systems.
- If Octave produces an error message for valid input, that is a bug.

- If Octave does not produce an error message for invalid input, that is a bug. However, you should note that your idea of “invalid input” might be my idea of “an extension” or “support for traditional practice”.
- If you are an experienced user of programs like Octave, your suggestions for improvement are welcome in any case.

D.2.2 Where to Report Bugs

To report a bug in Octave, submit a bug report to the Octave bug tracker <https://bugs.octave.org>.

Do not send bug reports to ‘help-octave’. Most users of Octave do not want to receive bug reports.

D.2.3 How to Report Bugs

Submit bug reports for Octave to the Octave bug tracker <https://bugs.octave.org>.

The fundamental principle of reporting bugs usefully is this: **report all the facts**. If you are not sure whether to state a fact or leave it out, state it!

Often people omit facts because they think they know what causes the problem and they conclude that some details don’t matter. Thus, you might assume that the name of the variable you use in an example does not matter. Well, probably it doesn’t, but one cannot be sure. Perhaps the bug is a stray memory reference which happens to fetch from the location where that name is stored in memory; perhaps, if the name were different, the contents of that location would fool the interpreter into doing the right thing despite the bug. Play it safe and give a specific, complete example.

Keep in mind that the purpose of a bug report is to enable someone to fix the bug if it is not known. Always write your bug reports on the assumption that the bug is not known.

Sometimes people give a few sketchy facts and ask, “Does this ring a bell?” This cannot help us fix a bug. It is better to send a complete bug report to begin with.

Try to make your bug report self-contained. If we have to ask you for more information, it is best if you include all the previous information in your response, as well as the information that was missing.

To enable someone to investigate the bug, you should include all these things:

- The version of Octave. You can get this by noting the version number that is printed when Octave starts, or running it with the ‘-v’ option.
- A complete input file that will reproduce the bug.

A single statement may not be enough of an example—the bug might depend on other details that are missing from the single statement where the error finally occurs.

- The command arguments you gave Octave to execute that example and observe the bug. To guarantee you won’t omit something important, list all the options.

If we were to try to guess the arguments, we would probably guess wrong and then we would not encounter the bug.

- The type of machine you are using, and the operating system name and version number.
- The command-line arguments you gave to the `configure` command when you installed the interpreter.

- A complete list of any modifications you have made to the interpreter source.
Be precise about these changes—show a context diff for them.
 - Details of any other deviations from the standard procedure for installing Octave.
 - A description of what behavior you observe that you believe is incorrect. For example, "The interpreter gets a fatal signal," or, "The output produced at line 208 is incorrect." Of course, if the bug is that the interpreter gets a fatal signal, then one can't miss it. But if the bug is incorrect output, we might not notice unless it is glaringly wrong. Even if the problem you experience is a fatal signal, you should still say so explicitly. Suppose something strange is going on, such as, your copy of the interpreter is out of sync, or you have encountered a bug in the C library on your system. Your copy might crash and the copy here would not. If you said to expect a crash, then when the interpreter here fails to crash, we would know that the bug was not happening. If you don't say to expect a crash, then we would not know whether the bug was happening. We would not be able to draw any conclusion from our observations.
- Often the observed symptom is incorrect output when your program is run. Unfortunately, this is not enough information unless the program is short and simple. It is very helpful if you can include an explanation of the expected output, and why the actual output is incorrect.
- If you wish to suggest changes to the Octave source, send them as context diffs. If you even discuss something in the Octave source, refer to it by context, not by line number, because the line numbers in the development sources probably won't match those in your sources.

Here are some things that are not necessary:

- A description of the envelope of the bug.

Often people who encounter a bug spend a lot of time investigating which changes to the input file will make the bug go away and which changes will not affect it. Such information is usually not necessary to enable us to fix bugs in Octave, but if you can find a simpler example to report *instead* of the original one, that is a convenience. Errors in the output will be easier to spot, running under the debugger will take less time, etc. Most Octave bugs involve just one function, so the most straightforward way to simplify an example is to delete all the function definitions except the one in which the bug occurs.

However, simplification is not vital; if you don't want to do this, report the bug anyway and send the entire test case you used.

- A patch for the bug. Patches can be helpful, but if you find a bug, you should report it, even if you cannot send a fix for the problem.

D.2.4 Sending Patches for Octave

If you would like to write bug fixes or improvements for Octave, that is very helpful. When you send your changes, please follow these guidelines to avoid causing extra work for us in studying the patches.

If you don't follow these guidelines, your information might still be useful, but using it will take extra work. Maintaining Octave is a lot of work in the best of circumstances, and we can't keep up unless you do your best to help.

- Send an explanation with your changes of what problem they fix or what improvement they bring about. For a bug fix, just include a copy of the bug report, and explain why the change fixes the bug.
- Always include a proper bug report for the problem you think you have fixed. We need to convince ourselves that the change is right before installing it. Even if it is right, we might have trouble judging it if we don't have a way to reproduce the problem.
- Include all the comments that are appropriate to help people reading the source in the future understand why this change was needed.
- Don't mix together changes made for different reasons. Send them *individually*. If you make two changes for separate reasons, then we might not want to install them both. We might want to install just one.
- Use ‘`diff -c`’ to make your diffs. Diffs without context are hard for us to install reliably. More than that, they make it hard for us to study the diffs to decide whether we want to install them. Unified diff format is better than contextless diffs, but not as easy to read as ‘`-c`’ format.

If you have GNU diff, use ‘`diff -cp`’, which shows the name of the function that each change occurs in.

- Write the change log entries for your changes.

Read the `ChangeLog` file to see what sorts of information to put in, and to learn the style that we use. The purpose of the change log is to show people where to find what was changed. So you need to be specific about what functions you changed; in large functions, it's often helpful to indicate where within the function the change was made.

On the other hand, once you have shown people where to find the change, you need not explain its purpose. Thus, if you add a new function, all you need to say about it is that it is new. If you feel that the purpose needs explaining, it probably does—but the explanation will be much more useful if you put it in comments in the code.

If you would like your name to appear in the header line for who made the change, send us the header line.

D.3 How To Get Help with Octave

The mailing list help@octave.org exists for the discussion of matters related to using and installing Octave. If you would like to join the discussion, please send a short note to help-request@octave.org.

Please do not send requests to be added or removed from the mailing list, or other administrative trivia to the list itself.

If you think you have found a bug in Octave or in the installation procedure, however, you should submit a complete bug report to the Octave bug tracker at <https://bugs.octave.org>. But before you submit a bug report, please read <https://www.octave.org/bugs.html> to learn how to submit a useful bug report.

D.4 How to Distinguish Between Octave and Matlab

Octave and MATLAB are very similar, but handle Java slightly different. Therefore it may be necessary to detect the environment and use the appropriate functions. The following

function can be used to detect the environment. Due to the persistent variable it can be called repeatedly without a heavy performance hit.

Example:

```
%%
%% Return: true if the environment is Octave.
%%
function retval = isOctave
    persistent cacheval; % speeds up repeated calls

    if isempty (cacheval)
        cacheval = (exist ("OCTAVE_VERSION", "builtin") > 0);
    end

    retval = cacheval;
end
```


Appendix E Installing Octave

The procedure for installing Octave from source on a Unix-like system is described next. Building on other platforms will follow similar steps. Note that this description applies to Octave releases. Building the development sources from the Mercurial archive requires additional steps as described in the development source itself.

E.1 Build Dependencies

Octave is a fairly large program with many build dependencies. You may be able to find pre-packaged versions of the dependencies distributed as part of your system, or you may have to build some or all of them yourself.

E.1.1 Obtaining the Dependencies Automatically

On some systems you can obtain many of Octave's build dependencies automatically. The commands for doing this vary by system. Similarly, the names of pre-compiled packages vary by system and do not always match exactly the names listed in [Section E.1.2 \[Build Tools\], page 1021](#), and [Section E.1.3 \[External Packages\], page 1022](#).

You will usually need the development version of an external dependency so that you get the libraries and header files for building software, not just for running already compiled programs. These packages typically have names that end with the suffix `-dev` or `-devel`.

On systems with `apt-get` (Debian, Ubuntu, etc.), you may be able to install most of the tools and external packages using a command similar to

```
apt-get build-dep octave
```

The specific package name may be `octave3.2` or `octave3.4`. The set of required tools and external dependencies does not change frequently, so it is not important that the version match exactly, but you should use the most recent one available.

On systems with `yum` (Fedora, Red Hat, etc.), you may be able to install most of the tools and external packages using a command similar to

```
yum-builddep octave
```

The `yum-builddep` utility is part of the `yum-utils` package.

For either type of system, the package name may include a version number. The set of required tools and external dependencies does not change frequently, so it is not important that the version exactly match the version you are installing, but you should use the most recent one available.

E.1.2 Build Tools

The following tools are required:

C++, C, and Fortran compilers

The Octave sources are primarily written in C++, but some portions are also written in C and Fortran. The Octave sources are intended to be portable. Recent versions of the GNU compiler collection (GCC) should work (<https://gcc.gnu.org>). If you use GCC, you should avoid mixing versions. For example, be sure that you are not using the obsolete `g77` Fortran compiler with modern versions of `gcc` and `g++`.

GNU Make

Tool for building software (<https://www.gnu.org/software/make>). Octave's build system requires GNU Make. Other versions of Make will not work. Fortunately, GNU Make is highly portable and easy to install.

AWK, sed, and other Unix utilities

Basic Unix system utilities are required for building Octave. All will be available with any modern Unix system and also on Windows with either Cygwin or MinGW and MSYS.

Additionally, the following tools may be needed:

Bison	Parser generator (https://www.gnu.org/software/bison). You will need Bison if you modify the <code>oct-parse.yy</code> source file or if you delete the files that are generated from it.
Flex	Lexer analyzer (https://www.gnu.org/software/flex). You will need Flex if you modify the <code>lex.ll</code> source file or if you delete the files that are generated from it.
Autoconf	Package for software configuration (https://www.gnu.org/software/autoconf). Autoconf is required if you modify Octave's <code>configure.ac</code> file or other files that it requires.
Automake	Package for Makefile generation (https://www.gnu.org/software/automake). Automake is required if you modify Octave's <code>Makefile.am</code> files or other files that they depend on.
Libtool	Package for building software libraries (https://www.gnu.org/software/libtool). Libtool is required by Automake.
gperf	Perfect hash function generator (https://www.gnu.org/software/gperf). You will need gperf if you modify the <code>octave.gperf</code> file or if you delete the file that is generated from it.
Texinfo	Package for generating online and print documentation (https://www.gnu.org/software/texinfo). You will need Texinfo to build Octave's documentation or if you modify the documentation source files or the docstring of any Octave function.

E.1.3 External Packages

The following external packages are required:

BLAS	Basic Linear Algebra Subroutine library. Accelerated BLAS libraries such as OpenBLAS (https://www.openblas.net/) or ATLAS (http://math-atlas.sourceforge.net) are recommended for best performance. The reference implementation (http://www.netlib.org/blas) is slow, unmaintained, and suffers from certain bugs in corner case inputs.
LAPACK	Linear Algebra Package (http://www.netlib.org/lapack).
PCRE	The Perl Compatible Regular Expression library (https://www.pcre.org).

The following external package is optional but strongly recommended:

GNU Readline

Command-line editing library (<https://www.gnu.org/s/readline>).

If you wish to build Octave without GNU readline installed, you must use the `--disable-readline` option when running the configure script.

The following external software packages are optional but recommended:

ARPACK	Library for the solution of large-scale eigenvalue problems (https://forge.scilab.org/index.php/p/arpac-ng). ARPACK is required to provide the functions <code>eigs</code> and <code>svds</code> .
cURL	Library for transferring data with URL syntax (https://curl.haxx.se). cURL is required to provide the <code>urlread</code> and <code>urlwrite</code> functions and the <code>ftp</code> class.
FFTW3	Library for computing discrete Fourier transforms (http://www.fftw.org). FFTW3 is used to provide better performance for functions that compute discrete Fourier transforms (<code>fft</code> , <code>ifft</code> , <code>fft2</code> , etc.)
FLTK	Portable GUI toolkit (http://www.fltk.org). FLTK is currently used to provide windows for Octave's OpenGL-based graphics functions.
fontconfig	Library for configuring and customizing font access (https://www.freedesktop.org/wiki/Software/fontconfig). Fontconfig is used to manage fonts for Octave's OpenGL-based graphics functions.
FreeType	Portable font engine (https://www.freetype.org). FreeType is used to perform font rendering for Octave's OpenGL-based graphics functions.
GLPK	GNU Linear Programming Kit (https://www.gnu.org/software/glpk). GLPK is required for the function <code>glpk</code> .
gl2ps	OpenGL to PostScript printing library (https://www.geuz.org/gl2ps/). gl2ps is required for printing when using OpenGL-based graphics toolkits (currently either FLTK or Qt).
gnuplot	Interactive graphics program (http://www.gnuplot.info). gnuplot is currently the default graphics renderer for Octave.
GraphicsMagick++	Image processing library (http://www.graphicsmagick.org). GraphicsMagick++ is used to provide the <code>imread</code> and <code>imwrite</code> functions.
HDF5	Library for manipulating portable data files (https://www.hdfgroup.org/HDF5). HDF5 is required for Octave's <code>load</code> and <code>save</code> commands to read and write HDF data files.
Java Development Kit	Java programming language compiler and libraries. The OpenJDK free software implementation is recommended (http://openjdk.java.net/), although other JDK implementations may work. Java is required to be able to call Java functions from within Octave.

LLVM	Compiler framework, (https://www.llvm.org). LLVM is required for Octave's experimental just-in-time (JIT) compilation for speeding up the interpreter.
OpenGL	API for portable 2-D and 3-D graphics (https://www.opengl.org). An OpenGL implementation is required to provide Octave's OpenGL-based graphics functions. Octave's OpenGL-based graphics functions usually outperform the gnuplot-based graphics functions because plot data can be rendered directly instead of sending data and commands to gnuplot for interpretation and rendering.
PortAudio	PortAudio (http://www.portaudio.com/) provides a very simple API for recording and/or playing sound using a simple callback function or a blocking read/write interface. It is required for the audio processing functions <code>audioplayer</code> , <code>audiorecorder</code> , and <code>audiodevinfo</code> .
Qhull	Computational geometry library (http://www.qhull.org). Qhull is required to provide the functions <code>convhull</code> , <code>convhulln</code> , <code>delaunay</code> , <code>delaunayn</code> , <code>voronoi</code> , and <code>voronoin</code> .
QRUPDATE	QR factorization updating library (https://sourceforge.net/projects/qrupdate). QRUPDATE is used to provide improved performance for the functions <code>qrdelete</code> , <code>qrinsert</code> , <code>qrshift</code> , and <code>qrupdate</code> .
QScintilla	Source code highlighter and manipulator; a Qt port of Scintilla (http://www.riverbankcomputing.co.uk/software/qscintilla). QScintilla is used for syntax highlighting and code completion in the GUI.
Qt	GUI and utility libraries (https://www.qt.io). Qt is required for building the GUI. It is a large framework, but the only components required are the GUI, core, and network modules.
SuiteSparse	Sparse matrix factorization library (http://faculty.cse.tamu.edu/davis/suitesparse.html). SuiteSparse is required to provide sparse matrix factorizations and solution of linear equations for sparse systems.
SUNDIALS	The SUite of Nonlinear and DIfferential/ALgebraic Equation Solvers (https://computation.llnl.gov/projects/sundials) is required for the Ordinary Differential Equations (ODE) solvers <code>ode15i</code> and <code>ode15s</code> .
zlib	Data compression library (https://zlib.net). The zlib library is required for Octave's <code>load</code> and <code>save</code> commands to handle compressed data, including MATLAB v5 MAT files.

E.2 Running Configure and Make

- Run the shell script `configure`. This will determine the features your system has (or doesn't have) and create a file named `Makefile` from each of the files named `Makefile.in`.

Here is a summary of the configure options that are most frequently used when building Octave:

- help Print a summary of the options recognized by the configure script.
- prefix=*prefix* Install Octave in subdirectories below *prefix*. The default value of *prefix* is `/usr/local`.
- srcdir=*dir* Look for Octave sources in the directory *dir*.
- disable-64 Disable using 64-bit integers for indexing arrays and use 32-bit integers instead. On systems with 32-bit pointers, this option is always disabled. If the configure script determines that your BLAS library uses 32-bit integers, then operations using the following libraries are limited to arrays with dimensions that are smaller than 2^{31} elements:
 - BLAS
 - LAPACK
 - QRUPDATE
 - SuiteSparse
 - ARPACKAdditionally, the following libraries use `int` internally, so maximum problem sizes are always limited:
 - GLPK
 - QhullSee [Section E.3 \[Compiling Octave with 64-bit Indexing\]](#), page 1029, for more details about building Octave with more complete support for large arrays.
- enable-address-sanitizer-flags Enable compiler options `-fsanitize=address` and `-fomit-frame-pointer` for memory access checking. This option is primarily used for debugging Octave. Building Octave with this option has a negative impact on performance and is not recommended for general use. It may also interfere with proper functioning of the GUI.
- disable-docs Disable building all forms of the documentation (Info, PDF, HTML). The default is to build documentation, but your system will need functioning Texinfo and TeX installs for this to succeed.
- enable-float-truncate This option allows for truncation of intermediate floating point results in calculations. It is only necessary for certain platforms.
- enable-readline Use the readline library to provide for editing of the command line in terminal environments. This option is on by default.

--enable-shared

Create shared libraries (this is the default). If you are planning to use the dynamic loading features, you will probably want to use this option. It will make your .oct files much smaller and on some systems it may be necessary to build shared libraries in order to use dynamically linked functions.

You may also want to build a shared version of `libstdc++`, if your system doesn't already have one.

--with-blas=<lib>

By default, configure looks for the best BLAS matrix libraries on your system, including optimized implementations such as the free ATLAS 3.0, as well as vendor-tuned libraries. (The use of an optimized BLAS will generally result in several-times faster matrix operations.) Use this option to specify a particular BLAS library that Octave should use.

--with-lapack=<lib>

By default, configure looks for the best LAPACK matrix libraries on your system, including optimized implementations such as the free ATLAS 3.0, as well as vendor-tuned libraries. (The use of an optimized LAPACK will generally result in several-times faster matrix operations.) Use this option to specify a particular LAPACK library that Octave should use.

--with-magick=<lib>

Select the library to use for image I/O. The two possible values are "`GraphicsMagick`" (default) or "`ImageMagick`".

--with-sepchar=<char>

Use <char> as the path separation character. This option can help when running Octave on non-Unix systems.

--without-amd

Don't use AMD, disable some sparse matrix functionality.

--without-camd

Don't use CAMD, disable some sparse matrix functionality.

--without-colamd

Don't use COLAMD, disable some sparse matrix functionality.

--without-ccolamd

Don't use CCOLAMD, disable some sparse matrix functionality.

--without-cholmod

Don't use CHOLMOD, disable some sparse matrix functionality.

--without-curl

Don't use the cURL library, disable the `ftp` objects, `urlread` and `urlwrite` functions.

--without-cxsparse

Don't use CXSPARSE, disable some sparse matrix functionality.

--without-fftw3
 Use the included FFTPACK library for computing Fast Fourier Transforms instead of the FFTW3 library.

--without-fftw3f
 Use the included FFTPACK library for computing Fast Fourier Transforms instead of the FFTW3 library when operating on single precision (float) values.

--without-glpk
 Don't use the GLPK library for linear programming.

--without-hdf5
 Don't use the HDF5 library, disable reading and writing of HDF5 files.

--without-opengl
 Don't use OpenGL, disable native graphics toolkit for plotting. You will need gnuplot installed in order to make plots.

--without-qhull
 Don't use Qhull, disable delaunay, convhull, and related functions.

--without-qrupdate
 Don't use QRUPDATE, disable QR and Cholesky update functions.

--without-umfpack
 Don't use UMFPACK, disable some sparse matrix functionality.

--without-zlib
 Don't use the zlib library, disable data file compression and support for recent MAT file formats.

--without-framework-carbon
 Don't use framework Carbon headers, libraries, or specific source code even if the configure test succeeds (the default is to use Carbon framework if available). This is a platform specific configure option for Mac systems.

--without-framework-opengl
 Don't use framework OpenGL headers, libraries, or specific source code even if the configure test succeeds. If this option is given then OpenGL headers and libraries in standard system locations are tested (the default value is **--with-framework-opengl**). This is a platform specific configure option for Mac systems.

See the file **INSTALL** for more general information about the command line options used by **configure**. That file also contains instructions for compiling in a directory other than the one where the source is located.

- Run **make**.

You will need a recent version of GNU Make as Octave relies on certain features not generally available in all versions of make. Modifying Octave's makefiles to work with other make programs is probably not worth your time; instead, we simply recommend installing GNU Make.

There are currently three options for plotting in Octave: the external program gnuplot, the internal graphics engine using OpenGL coupled with either FLTK or Qt widgets. Gnuplot is a command-driven interactive function plotting program.

To compile Octave, you will need a recent version of `g++` or other ANSI C++ compiler. In addition, you will need a Fortran 77 compiler or `f2c`. If you use `f2c`, you will need a script like `fort77` that works like a normal Fortran compiler by combining `f2c` with your C compiler in a single script.

If you plan to modify the parser you will also need GNU `bison` and `flex`. If you modify the documentation, you will need GNU Texinfo.

GNU Make, `gcc` (and `libstdc++`), `gnuplot`, `bison`, `flex`, and Texinfo are all available from many anonymous ftp archives. The primary site is ftp.gnu.org, but it is often very busy. A list of sites that mirror the software on ftp.gnu.org is available by anonymous ftp from [ftp://ftp.gnu.org/pub/gnu/GNUinfo/FTP](http://ftp.gnu.org/pub/gnu/GNUinfo/FTP).

Octave requires approximately 1.4 GB of disk storage to unpack and compile from source (significantly less, 400 MB, if you don't compile with debugging symbols). To compile without debugging symbols try the command

```
make CFLAGS=-O CXXFLAGS=-O LDFLAGS=
```

instead of just `make`.

- If you encounter errors while compiling Octave, first see [Section E.4 \[Installation Problems\], page 1031](#), for a list of known problems and if there is a workaround or solution for your problem. If not, see [Appendix D \[Trouble\], page 1015](#), for information about how to report bugs.
- Once you have successfully compiled Octave, run `make install`.

This will install a copy of Octave, its libraries, and its documentation in the destination directory. As distributed, Octave is installed in the following directories. In the table below, `prefix` defaults to `/usr/local`, `version` stands for the current version number of the interpreter, and `arch` is the type of computer on which Octave is installed (for example, '`i586-unknown-gnu`').

`prefix/bin`

Octave and other binaries that people will want to run directly.

`prefix/lib/octave-version`

Libraries like `liboctave.a` and `liboctinterp.a`.

`prefix/octave-version/include/octave`

Include files distributed with Octave.

`prefix/share`

Architecture-independent data files.

`prefix/share/man/man1`

Unix-style man pages describing Octave.

`prefix/share/info`

Info files describing Octave.

`prefix/share/octave/version/m`

Function files distributed with Octave. This includes the Octave version, so that multiple versions of Octave may be installed at the same time.

prefix/libexec/octave/version/exec/arch

Executables to be run by Octave rather than the user.

prefix/lib/octave/version/oct/arch

Object files that will be dynamically loaded.

prefix/share/octave/version/imagelib

Image files that are distributed with Octave.

E.3 Compiling Octave with 64-bit Indexing

Note: the following only applies to systems that have 64-bit pointers. Configuring Octave with `--enable-64` cannot magically make a 32-bit system have a 64-bit address space.

On 64-bit systems, Octave uses 64-bit integers for indexing arrays by default. If the configure script determines that your BLAS library uses 32-bit integers, then operations using the following libraries are limited to arrays with dimensions that are smaller than 2^{31} elements:

- BLAS
- LAPACK
- QRUPDATE
- SuiteSparse
- ARPACK

Additionally, the following libraries use `int` internally, so maximum problem sizes are always limited:

- GLPK
- Qhull

Except for GLPK and Qhull, these libraries may also be configured to use 64-bit integers, but most systems do not provide packages built this way. If you wish to experiment with large arrays, the following information may be helpful.

The following instructions were tested with the development version of Octave and GCC 4.3.4 on an x86_64 Debian system and may be out of date now. Please report any problems or corrections on the Octave bug tracker.

The versions listed below are the versions used for testing. If newer versions of these packages are available, you should try to use them, although there may be some differences.

All libraries and header files will be installed in subdirectories of `$prefix64` (you must choose the location of this directory).

- BLAS and LAPACK (<http://www.netlib.org/lapack>)

Reference versions for both libraries are included in the reference LAPACK 3.2.1 distribution from netlib.org.

- Copy the file `make.inc.example` and name it `make.inc`. The options `-fdefault-integer-8` and `-fPIC` (on 64-bit CPU) have to be added to the variable `OPTS` and `NOOPT`.
- Once you have compiled this library make sure that you use it for compiling Suite Sparse and Octave. In the following we assume that you installed the LAPACK library as `$prefix64/lib/liblapack.a`.

- QRUPDATE (<https://sourceforge.net/projects/qrupdate>)

In the `Makeconf` file:

- Add `-fdefault-integer=8` to `FFLAGS`.
- Adjust the BLAS and LAPACK variables as needed if your 64-bit aware BLAS and LAPACK libraries are in a non-standard location.
- Set `PREFIX` to the top-level directory of your install tree.
- Run `make solib` to make a shared library.
- Run `make install` to install the library.

- SuiteSparse (<http://faculty.cse.tamu.edu/davis/suitesparse.html>)

Pass the following options to `make` to enable 64-bit integers for BLAS library calls. On 64-bit Windows systems, use `-DLONGBLAS="long long"` instead.

```
CFLAGS=' -DLONGBLAS=long'
CXXFLAGS=' -DLONGBLAS=long'
```

The SuiteSparse makefiles don't generate shared libraries. On some systems, you can generate them by doing something as simple as

```
top=$(pwd)
for f in *.a; do
    mkdir tmp
    cd tmp
    ar vx ../$f
    gcc -shared -o ../../${f%%.a}.so *.*
    cd $top
    rm -rf tmp
done
```

Other systems may require a different solution.

- ARPACK (<https://forge.scilab.org/index.php/p/arpack-ng/>)
 - Add `-fdefault-integer=8` to `FFLAGS` when running `configure`.
 - Run `make` to build the library.
 - Run `make install` to install the library.
- ATLAS instead of reference BLAS and LAPACK

Suggestions on how to compile ATLAS would be most welcome.
- GLPK
- Qhull (<http://www.qhull.org>)

Both GLPK and Qhull use `int` internally so maximum problem sizes may be limited.
- Octave

Octave's 64-bit index support is activated with the `configure` option `--enable-64`.

```
./configure \
LD_LIBRARY_PATH="$prefix64/lib" \
CPPFLAGS="-I$prefix64/include" LDFLAGS="-L$prefix64/lib" \
--enable-64
```

You must ensure that all Fortran sources except those in the `liboctave/external/ranlib` directory are compiled such that INTEGERS are 8-bytes wide. If you are using

gfortran, the configure script should automatically set the Makefile variable `F77_INTEGER_8_FLAG` to `-fdefault-integer-8`. If you are using another compiler, you must set this variable yourself. You should NOT set this flag in `FFLAGS`, otherwise the files in `liboctave/external/ranlib` will be miscompiled.

- Other dependencies

Probably nothing special needs to be done for the following dependencies. If you discover that something does need to be done, please submit a bug report.

- pcre
- zlib
- hdf5
- fftw3
- cURL
- GraphicsMagick++
- OpenGL
- freetype
- fontconfig
- fltk

E.4 Installation Problems

This section contains a list of problems (and some apparent problems that don't really mean anything is wrong) that may show up during installation of Octave.

- On some SCO systems, `info` fails to compile if `HAVE_TERMIOS_H` is defined in `config.h`. Simply removing the definition from `info/config.h` should allow it to compile.
- If `configure` finds `dlopen`, `dlsym`, `dlclose`, and `dlerror`, but not the header file `dlfcn.h`, you need to find the source for the header file and install it in the directory `usr/include`. This is reportedly a problem with Slackware 3.1. For Linux/GNU systems, the source for `dlfcn.h` is in the `ldso` package.
- Building `.oct` files doesn't work.

You should probably have a shared version of `libstdc++`. A patch is needed to build shared versions of version 2.7.2 of `libstdc++` on the HP-PA architecture. You can find the patch at <ftp://ftp.cygnus.com/pub/g++/libg++-2.7.2-hppa-gcc-fix>.

- On some DEC alpha systems there may be a problem with the `libdxm1` library, resulting in floating point errors and/or segmentation faults in the linear algebra routines called by Octave. If you encounter such problems, then you should modify the `configure` script so that `SPECIAL_MATH_LIB` is not set to `-ldxm1`.
- On FreeBSD systems Octave may hang while initializing some internal constants. The fix appears to be to use

```
options      GPL_MATH_EMULATE
```

rather than

```
options      MATH_EMULATE
```

in the kernel configuration files (typically found in the directory `/sys/i386/conf`). After making this change, you'll need to rebuild the kernel, install it, and reboot.

- If you encounter errors like

```
passing 'void (*)()' as argument 2 of
  'octave_set_signal_handler(int, void (*)(int))'
```

or

```
warning: ANSI C++ prohibits conversion from '(int)'
          to '(...)'
```

while compiling `sighandlers.cc`, you may need to edit some files in the `gcc` include subdirectory to add proper prototypes for functions there. For example, Ultrix 4.2 needs proper declarations for the `signal` function and the `SIG_IGN` macro in the file `signal.h`.

On some systems the `SIG_IGN` macro is defined to be something like this:

```
#define SIG_IGN (void ())1
```

when it should really be something like:

```
#define SIG_IGN (void (*)(int))1
```

to match the prototype declaration for the `signal` function. This change should also be made for the `SIG_DFL` and `SIG_ERR` symbols. It may be necessary to change the definitions in `sys signal.h` as well.

The `gcc fixincludes` and `fixproto` scripts should probably fix these problems when `gcc` installs its modified set of header files, but I don't think that's been done yet.

You should not change the files in /usr/include. You can find the `gcc` include directory tree by running the command

```
gcc -print-libgcc-file-name
```

The directory of `gcc` include files normally begins in the same directory that contains the file `libgcc.a`.

- Some of the Fortran subroutines may fail to compile with older versions of the Sun Fortran compiler. If you get errors like

```
zgemm.f:
      zgemm:
warning: unexpected parent of complex expression subtree
zgemm.f, line 245: warning: unexpected parent of complex
      expression subtree
warning: unexpected parent of complex expression subtree
zgemm.f, line 304: warning: unexpected parent of complex
      expression subtree
warning: unexpected parent of complex expression subtree
zgemm.f, line 327: warning: unexpected parent of complex
      expression subtree
pcc_binval: missing IR_CONV in complex op
make[2]: *** [zgemm.o] Error 1
```

when compiling the Fortran subroutines in the `liboctave/external` subdirectory, you should either upgrade your compiler or try compiling with optimization turned off.

- On NeXT systems, if you get errors like this:

```
/usr/tmp/cc007458.s:undefined local
    symbol LBB7656
/usr/tmp/cc007458.s:undefined local
    symbol LBE7656
```

when compiling `Array.cc` and `Matrix.cc`, try recompiling these files without `-g`.

- Some people have reported that calls to `system()` and the pager do not work on SunOS systems. This is apparently due to having `G_HAVE_SYS_WAIT` defined to be 0 instead of 1 when compiling `libg++`.
- On systems where the reference BLAS library is used the following matrix-by-vector multiplication incorrectly handles NaN values of the form `NaN * 0`.

```
[NaN, 1; 0, 0] * [0; 1]
⇒
[ 1
  0 ]
```

```
correct result ⇒
[ NaN
  0 ]
```

Install a different BLAS library such as OpenBLAS or ATLAS to correct this issue.

- On NeXT systems, linking to `libsyst_s.a` may fail to resolve the following functions

```
_tcgetattr
_tcsetattr
_tcflow
```

which are part of `libposix.a`. Unfortunately, linking Octave with `-posix` results in the following undefined symbols.

```
.destructors_used
.constructors_used
_objc_msgSend
_NXGetDefaultValue
_NXRegisterDefaults
_OBJC_CLASS_NAME_NXStringTable
_OBJC_CLASS_NAME_NXBundle
```

One kluge around this problem is to extract `termios.o` from `libposix.a`, put it in Octave's `src` directory, and add it to the list of files to link together in the makefile. Suggestions for better ways to solve this problem are welcome!

- If Octave crashes immediately with a floating point exception, it is likely that it is failing to initialize the IEEE floating point values for infinity and NaN.

If your system actually does support IEEE arithmetic, you should be able to fix this problem by modifying the function `octave_ieee_init` in the file `lo-ieee.cc` to correctly initialize Octave's internal infinity and NaN variables.

If your system does not support IEEE arithmetic but Octave's configure script incorrectly determined that it does, you can work around the problem by editing the file `config.h` to not define `HAVE_ISINF`, `HAVEFINITE`, and `HAVE_ISNAN`.

In any case, please report this as a bug since it might be possible to modify Octave's configuration script to automatically determine the proper thing to do.

- If Octave is unable to find a header file because it is installed in a location that is not normally searched by the compiler, you can add the directory to the include search path by specifying (for example) `CPPFLAGS=-I/some/nonstandard/directory` as an argument to `configure`. Other variables that can be specified this way are `CFLAGS`, `CXXFLAGS`, `FFLAGS`, and `LDFLAGS`. Passing them as options to the `configure` script also records them in the `config.status` file. By default, `CPPFLAGS` and `LDFLAGS` are empty, `CFLAGS` and `CXXFLAGS` are set to "`-g -O2`" and `FFLAGS` is set to "`-O`".

Appendix F Grammar and Parser

This appendix will eventually contain a semi-formal description of Octave's language.

F.1 Keywords

The following identifiers are keywords, and may not be used as variable or function names:

<code>--FILE--</code>	<code>--LINE--</code>	<code>break</code>
<code>case</code>	<code>catch</code>	<code>classdef</code>
<code>continue</code>	<code>do</code>	<code>else</code>
<code>elseif</code>	<code>end</code>	<code>end_try_catch</code>
<code>end_unwind_protect</code>	<code>endclassdef</code>	<code>endenumeration</code>
<code>endevents</code>	<code>endfor</code>	<code>endfunction</code>
<code>endif</code>	<code>endmethods</code>	<code>endparfor</code>
<code>endproperties</code>	<code>endswitch</code>	<code>endwhile</code>
<code>enumeration</code>	<code>events</code>	<code>for</code>
<code>function</code>	<code>global</code>	<code>if</code>
<code>methods</code>	<code>otherwise</code>	<code>parfor</code>
<code>persistent</code>	<code>properties</code>	<code>return</code>
<code>switch</code>	<code>try</code>	<code>until</code>
<code>unwind_protect</code>	<code>unwind_protect_cleanup</code>	<code>while</code>

The function `iskeyword` can be used to quickly check whether an identifier is reserved by Octave.

```
iskeyword ()
iskeyword (name)
```

Return true if `name` is an Octave keyword.

If `name` is omitted, return a list of keywords.

See also: [\[isvarname\]](#), page 127, [\[exist\]](#), page 134.

F.2 Parser

The parser has a number of variables that affect its internal operation. These variables are generally documented in the manual alongside the code that they affect.

In addition, there are three non-specific parser customization functions. `add_input_event_hook` can be used to schedule a user function for periodic evaluation. `remove_input_event_hook` will stop a user function from being evaluated periodically.

```
id = add_input_event_hook (fcn)
id = add_input_event_hook (fcn, data)
```

Add the named function or function handle `fcn` to the list of functions to call periodically when Octave is waiting for input.

The function should have the form

```
fcn (data)
```

If `data` is omitted, Octave calls the function without any arguments.

The returned identifier may be used to remove the function handle from the list of input hook functions.

See also: [\[remove_input_event_hook\]](#), page 1036.

```
remove_input_event_hook (name)
remove_input_event_hook (fcn_id)
```

Remove the named function or function handle with the given identifier from the list of functions to call periodically when Octave is waiting for input.

See also: [\[add_input_event_hook\]](#), page 1035.

Finally, when the parser cannot identify an input token it calls a particular function to handle this. By default, this is the internal function "`__unimplemented__`" which makes suggestions about possible Octave substitutes for MATLAB functions.

```
val = missing_function_hook ()
old_val = missing_function_hook (new_val)
missing_function_hook (new_val, "local")
```

Query or set the internal variable that specifies the function to call when an unknown identifier is requested.

When called from inside a function with the "local" option, the variable is changed locally for the function and any subroutines it calls. The original variable value is restored when exiting the function.

See also: [\[missing_component_hook\]](#), page 942.

Appendix G GNU GENERAL PUBLIC LICENSE

Version 3, 29 June 2007

Copyright © 2007 Free Software Foundation, Inc. <https://fsf.org/>

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Preamble

The GNU General Public License is a free, copyleft license for software and other kinds of works.

The licenses for most software and other practical works are designed to take away your freedom to share and change the works. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change all versions of a program—to make sure it remains free software for all its users. We, the Free Software Foundation, use the GNU General Public License for most of our software; it applies also to any other work released this way by its authors. You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for them if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs, and that you know you can do these things.

To protect your rights, we need to prevent others from denying you these rights or asking you to surrender the rights. Therefore, you have certain responsibilities if you distribute copies of the software, or if you modify it: responsibilities to respect the freedom of others.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must pass on to the recipients the same freedoms that you received. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

Developers that use the GNU GPL protect your rights with two steps: (1) assert copyright on the software, and (2) offer you this License giving you legal permission to copy, distribute and/or modify it.

For the developers' and authors' protection, the GPL clearly explains that there is no warranty for this free software. For both users' and authors' sake, the GPL requires that modified versions be marked as changed, so that their problems will not be attributed erroneously to authors of previous versions.

Some devices are designed to deny users access to install or run modified versions of the software inside them, although the manufacturer can do so. This is fundamentally incompatible with the aim of protecting users' freedom to change the software. The systematic pattern of such abuse occurs in the area of products for individuals to use, which is precisely where it is most unacceptable. Therefore, we have designed this version of the GPL to prohibit the practice for those products. If such problems arise substantially in other domains, we stand ready to extend this provision to those domains in future versions of the GPL, as needed to protect the freedom of users.

Finally, every program is threatened constantly by software patents. States should not allow patents to restrict development and use of software on general-purpose computers, but in those that do, we wish to avoid the special danger that patents applied to a free program could make it effectively proprietary. To prevent this, the GPL assures that patents cannot be used to render the program non-free.

The precise terms and conditions for copying, distribution and modification follow.

TERMS AND CONDITIONS

0. Definitions.

“This License” refers to version 3 of the GNU General Public License.

“Copyright” also means copyright-like laws that apply to other kinds of works, such as semiconductor masks.

“The Program” refers to any copyrightable work licensed under this License. Each licensee is addressed as “you”. “Licensees” and “recipients” may be individuals or organizations.

To “modify” a work means to copy from or adapt all or part of the work in a fashion requiring copyright permission, other than the making of an exact copy. The resulting work is called a “modified version” of the earlier work or a work “based on” the earlier work.

A “covered work” means either the unmodified Program or a work based on the Program.

To “propagate” a work means to do anything with it that, without permission, would make you directly or secondarily liable for infringement under applicable copyright law, except executing it on a computer or modifying a private copy. Propagation includes copying, distribution (with or without modification), making available to the public, and in some countries other activities as well.

To “convey” a work means any kind of propagation that enables other parties to make or receive copies. Mere interaction with a user through a computer network, with no transfer of a copy, is not conveying.

An interactive user interface displays “Appropriate Legal Notices” to the extent that it includes a convenient and prominently visible feature that (1) displays an appropriate copyright notice, and (2) tells the user that there is no warranty for the work (except to the extent that warranties are provided), that licensees may convey the work under this License, and how to view a copy of this License. If the interface presents a list of user commands or options, such as a menu, a prominent item in the list meets this criterion.

1. Source Code.

The “source code” for a work means the preferred form of the work for making modifications to it. “Object code” means any non-source form of a work.

A “Standard Interface” means an interface that either is an official standard defined by a recognized standards body, or, in the case of interfaces specified for a particular programming language, one that is widely used among developers working in that language.

The “System Libraries” of an executable work include anything, other than the work as a whole, that (a) is included in the normal form of packaging a Major Component, but which is not part of that Major Component, and (b) serves only to enable use of the work with that Major Component, or to implement a Standard Interface for which an implementation is available to the public in source code form. A “Major Component”, in this context, means a major essential component (kernel, window system, and so on) of the specific operating system (if any) on which the executable work runs, or a compiler used to produce the work, or an object code interpreter used to run it.

The “Corresponding Source” for a work in object code form means all the source code needed to generate, install, and (for an executable work) run the object code and to modify the work, including scripts to control those activities. However, it does not include the work’s System Libraries, or general-purpose tools or generally available free programs which are used unmodified in performing those activities but which are not part of the work. For example, Corresponding Source includes interface definition files associated with source files for the work, and the source code for shared libraries and dynamically linked subprograms that the work is specifically designed to require, such as by intimate data communication or control flow between those subprograms and other parts of the work.

The Corresponding Source need not include anything that users can regenerate automatically from other parts of the Corresponding Source.

The Corresponding Source for a work in source code form is that same work.

2. Basic Permissions.

All rights granted under this License are granted for the term of copyright on the Program, and are irrevocable provided the stated conditions are met. This License explicitly affirms your unlimited permission to run the unmodified Program. The output from running a covered work is covered by this License only if the output, given its content, constitutes a covered work. This License acknowledges your rights of fair use or other equivalent, as provided by copyright law.

You may make, run and propagate covered works that you do not convey, without conditions so long as your license otherwise remains in force. You may convey covered works to others for the sole purpose of having them make modifications exclusively for you, or provide you with facilities for running those works, provided that you comply with the terms of this License in conveying all material for which you do not control copyright. Those thus making or running the covered works for you must do so exclusively on your behalf, under your direction and control, on terms that prohibit them from making any copies of your copyrighted material outside their relationship with you.

Conveying under any other circumstances is permitted solely under the conditions stated below. Sublicensing is not allowed; section 10 makes it unnecessary.

3. Protecting Users’ Legal Rights From Anti-Circumvention Law.

No covered work shall be deemed part of an effective technological measure under any applicable law fulfilling obligations under article 11 of the WIPO copyright treaty adopted on 20 December 1996, or similar laws prohibiting or restricting circumvention of such measures.

When you convey a covered work, you waive any legal power to forbid circumvention of technological measures to the extent such circumvention is effected by exercising rights under this License with respect to the covered work, and you disclaim any intention to limit operation or modification of the work as a means of enforcing, against the work's users, your or third parties' legal rights to forbid circumvention of technological measures.

4. Conveying Verbatim Copies.

You may convey verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice; keep intact all notices stating that this License and any non-permissive terms added in accord with section 7 apply to the code; keep intact all notices of the absence of any warranty; and give all recipients a copy of this License along with the Program.

You may charge any price or no price for each copy that you convey, and you may offer support or warranty protection for a fee.

5. Conveying Modified Source Versions.

You may convey a work based on the Program, or the modifications to produce it from the Program, in the form of source code under the terms of section 4, provided that you also meet all of these conditions:

- a. The work must carry prominent notices stating that you modified it, and giving a relevant date.
- b. The work must carry prominent notices stating that it is released under this License and any conditions added under section 7. This requirement modifies the requirement in section 4 to "keep intact all notices".
- c. You must license the entire work, as a whole, under this License to anyone who comes into possession of a copy. This License will therefore apply, along with any applicable section 7 additional terms, to the whole of the work, and all its parts, regardless of how they are packaged. This License gives no permission to license the work in any other way, but it does not invalidate such permission if you have separately received it.
- d. If the work has interactive user interfaces, each must display Appropriate Legal Notices; however, if the Program has interactive interfaces that do not display Appropriate Legal Notices, your work need not make them do so.

A compilation of a covered work with other separate and independent works, which are not by their nature extensions of the covered work, and which are not combined with it such as to form a larger program, in or on a volume of a storage or distribution medium, is called an "aggregate" if the compilation and its resulting copyright are not used to limit the access or legal rights of the compilation's users beyond what the individual works permit. Inclusion of a covered work in an aggregate does not cause this License to apply to the other parts of the aggregate.

6. Conveying Non-Source Forms.

You may convey a covered work in object code form under the terms of sections 4 and 5, provided that you also convey the machine-readable Corresponding Source under the terms of this License, in one of these ways:

- a. Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by the Corresponding Source fixed on a durable physical medium customarily used for software interchange.
- b. Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by a written offer, valid for at least three years and valid for as long as you offer spare parts or customer support for that product model, to give anyone who possesses the object code either (1) a copy of the Corresponding Source for all the software in the product that is covered by this License, on a durable physical medium customarily used for software interchange, for a price no more than your reasonable cost of physically performing this conveying of source, or (2) access to copy the Corresponding Source from a network server at no charge.
- c. Convey individual copies of the object code with a copy of the written offer to provide the Corresponding Source. This alternative is allowed only occasionally and noncommercially, and only if you received the object code with such an offer, in accord with subsection 6b.
- d. Convey the object code by offering access from a designated place (gratis or for a charge), and offer equivalent access to the Corresponding Source in the same way through the same place at no further charge. You need not require recipients to copy the Corresponding Source along with the object code. If the place to copy the object code is a network server, the Corresponding Source may be on a different server (operated by you or a third party) that supports equivalent copying facilities, provided you maintain clear directions next to the object code saying where to find the Corresponding Source. Regardless of what server hosts the Corresponding Source, you remain obligated to ensure that it is available for as long as needed to satisfy these requirements.
- e. Convey the object code using peer-to-peer transmission, provided you inform other peers where the object code and Corresponding Source of the work are being offered to the general public at no charge under subsection 6d.

A separable portion of the object code, whose source code is excluded from the Corresponding Source as a System Library, need not be included in conveying the object code work.

A “User Product” is either (1) a “consumer product”, which means any tangible personal property which is normally used for personal, family, or household purposes, or (2) anything designed or sold for incorporation into a dwelling. In determining whether a product is a consumer product, doubtful cases shall be resolved in favor of coverage. For a particular product received by a particular user, “normally used” refers to a typical or common use of that class of product, regardless of the status of the particular user or of the way in which the particular user actually uses, or expects or is expected to use, the product. A product is a consumer product regardless of whether the product has substantial commercial, industrial or non-consumer uses, unless such uses represent the only significant mode of use of the product.

“Installation Information” for a User Product means any methods, procedures, authorization keys, or other information required to install and execute modified versions of a covered work in that User Product from a modified version of its Corresponding Source.

The information must suffice to ensure that the continued functioning of the modified object code is in no case prevented or interfered with solely because modification has been made.

If you convey an object code work under this section in, or with, or specifically for use in, a User Product, and the conveying occurs as part of a transaction in which the right of possession and use of the User Product is transferred to the recipient in perpetuity or for a fixed term (regardless of how the transaction is characterized), the Corresponding Source conveyed under this section must be accompanied by the Installation Information. But this requirement does not apply if neither you nor any third party retains the ability to install modified object code on the User Product (for example, the work has been installed in ROM).

The requirement to provide Installation Information does not include a requirement to continue to provide support service, warranty, or updates for a work that has been modified or installed by the recipient, or for the User Product in which it has been modified or installed. Access to a network may be denied when the modification itself materially and adversely affects the operation of the network or violates the rules and protocols for communication across the network.

Corresponding Source conveyed, and Installation Information provided, in accord with this section must be in a format that is publicly documented (and with an implementation available to the public in source code form), and must require no special password or key for unpacking, reading or copying.

7. Additional Terms.

“Additional permissions” are terms that supplement the terms of this License by making exceptions from one or more of its conditions. Additional permissions that are applicable to the entire Program shall be treated as though they were included in this License, to the extent that they are valid under applicable law. If additional permissions apply only to part of the Program, that part may be used separately under those permissions, but the entire Program remains governed by this License without regard to the additional permissions.

When you convey a copy of a covered work, you may at your option remove any additional permissions from that copy, or from any part of it. (Additional permissions may be written to require their own removal in certain cases when you modify the work.) You may place additional permissions on material, added by you to a covered work, for which you have or can give appropriate copyright permission.

Notwithstanding any other provision of this License, for material you add to a covered work, you may (if authorized by the copyright holders of that material) supplement the terms of this License with terms:

- a. Disclaiming warranty or limiting liability differently from the terms of sections 15 and 16 of this License; or
- b. Requiring preservation of specified reasonable legal notices or author attributions in that material or in the Appropriate Legal Notices displayed by works containing it; or
- c. Prohibiting misrepresentation of the origin of that material, or requiring that modified versions of such material be marked in reasonable ways as different from the original version; or

- d. Limiting the use for publicity purposes of names of licensors or authors of the material; or
- e. Declining to grant rights under trademark law for use of some trade names, trademarks, or service marks; or
- f. Requiring indemnification of licensors and authors of that material by anyone who conveys the material (or modified versions of it) with contractual assumptions of liability to the recipient, for any liability that these contractual assumptions directly impose on those licensors and authors.

All other non-permissive additional terms are considered “further restrictions” within the meaning of section 10. If the Program as you received it, or any part of it, contains a notice stating that it is governed by this License along with a term that is a further restriction, you may remove that term. If a license document contains a further restriction but permits relicensing or conveying under this License, you may add to a covered work material governed by the terms of that license document, provided that the further restriction does not survive such relicensing or conveying.

If you add terms to a covered work in accord with this section, you must place, in the relevant source files, a statement of the additional terms that apply to those files, or a notice indicating where to find the applicable terms.

Additional terms, permissive or non-permissive, may be stated in the form of a separately written license, or stated as exceptions; the above requirements apply either way.

8. Termination.

You may not propagate or modify a covered work except as expressly provided under this License. Any attempt otherwise to propagate or modify it is void, and will automatically terminate your rights under this License (including any patent licenses granted under the third paragraph of section 11).

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, you do not qualify to receive new licenses for the same material under section 10.

9. Acceptance Not Required for Having Copies.

You are not required to accept this License in order to receive or run a copy of the Program. Ancillary propagation of a covered work occurring solely as a consequence of using peer-to-peer transmission to receive a copy likewise does not require acceptance.

However, nothing other than this License grants you permission to propagate or modify any covered work. These actions infringe copyright if you do not accept this License. Therefore, by modifying or propagating a covered work, you indicate your acceptance of this License to do so.

10. Automatic Licensing of Downstream Recipients.

Each time you convey a covered work, the recipient automatically receives a license from the original licensors, to run, modify and propagate that work, subject to this License. You are not responsible for enforcing compliance by third parties with this License.

An “entity transaction” is a transaction transferring control of an organization, or substantially all assets of one, or subdividing an organization, or merging organizations. If propagation of a covered work results from an entity transaction, each party to that transaction who receives a copy of the work also receives whatever licenses to the work the party’s predecessor in interest had or could give under the previous paragraph, plus a right to possession of the Corresponding Source of the work from the predecessor in interest, if the predecessor has it or can get it with reasonable efforts.

You may not impose any further restrictions on the exercise of the rights granted or affirmed under this License. For example, you may not impose a license fee, royalty, or other charge for exercise of rights granted under this License, and you may not initiate litigation (including a cross-claim or counterclaim in a lawsuit) alleging that any patent claim is infringed by making, using, selling, offering for sale, or importing the Program or any portion of it.

11. Patents.

A “contributor” is a copyright holder who authorizes use under this License of the Program or a work on which the Program is based. The work thus licensed is called the contributor’s “contributor version”.

A contributor’s “essential patent claims” are all patent claims owned or controlled by the contributor, whether already acquired or hereafter acquired, that would be infringed by some manner, permitted by this License, of making, using, or selling its contributor version, but do not include claims that would be infringed only as a consequence of further modification of the contributor version. For purposes of this definition, “control” includes the right to grant patent sublicenses in a manner consistent with the requirements of this License.

Each contributor grants you a non-exclusive, worldwide, royalty-free patent license under the contributor’s essential patent claims, to make, use, sell, offer for sale, import and otherwise run, modify and propagate the contents of its contributor version.

In the following three paragraphs, a “patent license” is any express agreement or commitment, however denominated, not to enforce a patent (such as an express permission to practice a patent or covenant not to sue for patent infringement). To “grant” such a patent license to a party means to make such an agreement or commitment not to enforce a patent against the party.

If you convey a covered work, knowingly relying on a patent license, and the Corresponding Source of the work is not available for anyone to copy, free of charge and under the terms of this License, through a publicly available network server or other readily accessible means, then you must either (1) cause the Corresponding Source to be so

available, or (2) arrange to deprive yourself of the benefit of the patent license for this particular work, or (3) arrange, in a manner consistent with the requirements of this License, to extend the patent license to downstream recipients. “Knowingly relying” means you have actual knowledge that, but for the patent license, your conveying the covered work in a country, or your recipient’s use of the covered work in a country, would infringe one or more identifiable patents in that country that you have reason to believe are valid.

If, pursuant to or in connection with a single transaction or arrangement, you convey, or propagate by procuring conveyance of, a covered work, and grant a patent license to some of the parties receiving the covered work authorizing them to use, propagate, modify or convey a specific copy of the covered work, then the patent license you grant is automatically extended to all recipients of the covered work and works based on it.

A patent license is “discriminatory” if it does not include within the scope of its coverage, prohibits the exercise of, or is conditioned on the non-exercise of one or more of the rights that are specifically granted under this License. You may not convey a covered work if you are a party to an arrangement with a third party that is in the business of distributing software, under which you make payment to the third party based on the extent of your activity of conveying the work, and under which the third party grants, to any of the parties who would receive the covered work from you, a discriminatory patent license (a) in connection with copies of the covered work conveyed by you (or copies made from those copies), or (b) primarily for and in connection with specific products or compilations that contain the covered work, unless you entered into that arrangement, or that patent license was granted, prior to 28 March 2007.

Nothing in this License shall be construed as excluding or limiting any implied license or other defenses to infringement that may otherwise be available to you under applicable patent law.

12. No Surrender of Others’ Freedom.

If conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot convey a covered work so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not convey it at all. For example, if you agree to terms that obligate you to collect a royalty for further conveying from those to whom you convey the Program, the only way you could satisfy both those terms and this License would be to refrain entirely from conveying the Program.

13. Use with the GNU Affero General Public License.

Notwithstanding any other provision of this License, you have permission to link or combine any covered work with a work licensed under version 3 of the GNU Affero General Public License into a single combined work, and to convey the resulting work. The terms of this License will continue to apply to the part which is the covered work, but the special requirements of the GNU Affero General Public License, section 13, concerning interaction through a network will apply to the combination as such.

14. Revised Versions of this License.

The Free Software Foundation may publish revised and/or new versions of the GNU General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies that a certain numbered version of the GNU General Public License “or any later version” applies to it, you have the option of following the terms and conditions either of that numbered version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of the GNU General Public License, you may choose any version ever published by the Free Software Foundation.

If the Program specifies that a proxy can decide which future versions of the GNU General Public License can be used, that proxy’s public statement of acceptance of a version permanently authorizes you to choose that version for the Program.

Later license versions may give you additional or different permissions. However, no additional obligations are imposed on any author or copyright holder as a result of your choosing to follow a later version.

15. Disclaimer of Warranty.

THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM “AS IS” WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

16. Limitation of Liability.

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MODIFIES AND/OR CONVEYS THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

17. Interpretation of Sections 15 and 16.

If the disclaimer of warranty and limitation of liability provided above cannot be given local legal effect according to their terms, reviewing courts shall apply local law that most closely approximates an absolute waiver of all civil liability in connection with the Program, unless a warranty or assumption of liability accompanies a copy of the Program in return for a fee.

END OF TERMS AND CONDITIONS

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively state the exclusion of warranty; and each file should have at least the “copyright” line and a pointer to where the full notice is found.

one line to give the program's name and a brief idea of what it does.

Copyright (C) year name of author

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see <https://www.gnu.org/licenses/>.

Also add information on how to contact you by electronic and paper mail.

If the program does terminal interaction, make it output a short notice like this when it starts in an interactive mode:

program Copyright (C) year name of author

This program comes with ABSOLUTELY NO WARRANTY; for details type ‘show w’.

This is free software, and you are welcome to redistribute it under certain conditions; type ‘show c’ for details.

The hypothetical commands ‘show w’ and ‘show c’ should show the appropriate parts of the General Public License. Of course, your program’s commands might be different; for a GUI interface, you would use an “about box”.

You should also get your employer (if you work as a programmer) or school, if any, to sign a “copyright disclaimer” for the program, if necessary. For more information on this, and how to apply and follow the GNU GPL, see <https://www.gnu.org/licenses/>.

The GNU General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Lesser General Public License instead of this License. But first, please read <https://www.gnu.org/licenses/why-not-lgpl.html>.

Concept Index

A

acknowledgements.....	1
addition.....	149, 846
and operator.....	153, 846
anonymous functions.....	214
ans	127
answers, incorrect.....	1015, 1017
application-defined data.....	459
apply.....	588
area series.....	465
arguments in function call.....	145
arithmetic operators.....	149, 846
array, creating a Java array.....	991
assignment expressions.....	156
assignment operators.....	156
axes graphics object	395

B

bar series.....	466
batch processing.....	36
block comments.....	37
body of a loop.....	169
boolean expressions.....	153, 846
boolean operators.....	153, 846
break statement.....	172
broadcast.....	585
broadcasting.....	585
BSX.....	585
bug criteria.....	1015
bug tracker.....	1016
bugs.....	1015
bugs, investigating.....	1017
bugs, known.....	1015
bugs, reporting.....	1016
built-in data types	39

C

callbacks.....	457
calling Java from Octave.....	986
calling Octave from Java.....	986
case statement.....	167
catch	174
cell arrays.....	44, 114
character strings.....	43, 67
Cholesky factorization	555
Citations.....	5
Citing Octave.....	5
classes, making available to Octave.....	986
classpath, adding new path	994
classpath, difference between static and dynamic	986
classpath, displaying	993

classpath, dynamic	993, 994
classpath, removing path	994
classpath, setting	986
classpath, static	993
classpath.txt	986
clearing the screen	26
code profiling	244
colors, graphics	456
comma separated lists	124
command and output logs	34
command completion	27
command descriptions	13
command echoing	34
command history	28
command options	15
command-line editing	25
comments	37
comparison expressions	152, 846
complex-conjugate transpose	149, 846
containers	101
continuation lines	175
continue statement	173
contour series	467
contributing to Octave	6
contributors	1
conversion specifications (printf)	278
conversion specifications (scanf)	284
copy-on-write	596
copyright	1037
core dump	1015
COW	596
creating graphics objects	395
cs-lists	124
customizing readline	32
customizing the prompt	32

D

DAE	677
data sources in object groups	465
data structures	43, 101
data types	39
data types, built-in	39
data types, user-defined	44
decrement operator	159
default arguments	193
default graphics properties	455
defining functions	177
Degree Symbol	381
deprecated functions	1009
description format	12
diagonal and permutation matrices	609
diagonal matrix expressions	612

dialog, displaying a dialog for
 selecting directories 861
 dialog, displaying a dialog for selecting files 861
 dialog, displaying a dialog for storing files 862
 dialog, displaying a help dialog 863
 dialog, displaying a list dialog 864
 dialog, displaying a message dialog 865
 dialog, displaying a modal dialog 867
 dialog, displaying a question dialog 866
 dialog, displaying a warning dialog 867
 dialog, displaying an error dialog 863
 dialog, displaying an input dialog 864
 diary of commands and output 34
 differential equations 677
 diffs, submitting 1017
 distribution of Octave 6
 division 149, 846
do-until statement 170
 documentation fonts 11
 documentation notation 11
 documenting functions 38
 documenting Octave programs 37
 documenting user scripts 38
 Dulmage-Mendelsohn decomposition 636
 dynamic classpath 986, 993
 dynamic classpath, adding new path 994
 dynamic naming 106
 dynamic-linking 945
 Dynamically Linked Functions 945

E

echoing executing commands 34
 editing the command line 25
 element-by-element evaluation 153
else statement 165
elseif statement 165
end statement 165
end, indexing 141
end: and **:end** 141
end_try_catch 174
end_unwind_protect 174
endfor statement 170
endfunction statement 177
endif statement 165
endswitch statement 167
endwhile statement 169
 equality operator 152, 846
 equality, tests for 152, 846
 equations, nonlinear 599
 erroneous messages 1015
 erroneous results 1015, 1017
 error bar series 468
 error ids 226
 error message notation 12
 error messages 35
 error messages, incorrect 1015
 escape sequence notation 67

evaluation notation 11
 executable scripts 36
 execution speed 596
 exiting octave 7, 19
 exponentiation 149, 846
 expression, range 52
 expressions 139
 expressions, assignment 156
 expressions, boolean 153, 846
 expressions, comparison 152, 846
 expressions, logical 153, 846

F

factorial function 147
 fatal signal 1015
 field, returning value of Java object field 992
 field, setting value of Java object field 992
 fields, displaying available fields
 of a Java object 991
 figure deletefcn 148
 figure graphics object 395
 finding minimums 603
 finish.m 19
 flag character (**printf**) 279
 flag character (**scanf**) 284
for statement 170
 Frobenius norm 552
 function application 588
 function descriptions 12
 function file 193
 function handle 148
function statement 177
 functions, deprecated 1009
 functions, obsolete 1009
 functions, user-defined 177
 funding Octave development 6

G

general p-norm 552
global statement 129
 global variables 129
 grammar rules 1035
 graphics 295
 graphics colors 456
 graphics data structures 393
 graphics line styles 456
 graphics marker styles 457
 graphics object properties 405
 graphics object, axes 395
 graphics object, figure 395
 graphics object, image 395
 graphics object, light 395
 graphics object, line 395
 graphics object, patch 395
 graphics object, root figure 395
 graphics object, surface 395

graphics object, text	395
graphics objects	395
graphics objects, saving	404
graphics properties, default	455
graphics toolkits	473
greater than operator	152, 846
group objects	469, 470, 472

H

handle functions	399
handle, function handles	214
hash table	114
help, online	20
help, user-defined functions	38
help, where to find	1018
Hermitian operator	149, 846
Hessenberg decomposition	557
history	1
history of commands	28

I

if statement	165
image graphics object	395
improving Octave	1016, 1017
incorrect error messages	1015
incorrect output	1015, 1017
incorrect results	1015, 1017
increment operator	159
indirect function call	148
infinity norm	552
initialization	18
inline, inline functions	214
input conversions, for <code>scanf</code>	285
input history	28
<code>~/.inputrc</code>	32
installation trouble	1015
installing Octave	1021
instance, creating a Java instance	990
introduction	7
introduction to graphics structures	393
invalid input	1015

J

Java, calling from Octave	986
Java, using with Octave	986
<code>javaclasspath.txt</code>	986

K

Kendall's Tau	740
key/value store	114
keywords	1035
known causes of trouble	1015

L

language definition	1035
less than operator	152, 846
light graphics object	395
line graphics object	395
line series	468
line styles, graphics	456
linear algebra	545
linear algebra, techniques	545
loading data	258
local minimum	603
logging commands and output	34
logical expressions	153, 846
logical operators	153, 846
loop	169
looping over structure elements	171
LP	697
LU decomposition	557
lvalue	157

M

map	588
Map	114
marker styles, graphics	457
matching failure, in <code>scanf</code>	284
matrices	48
matrices, diagonal and permutation	609
matrix factorizations	555
matrix functions, basic	545
matrix multiplication	149, 846
matrix, functions of	568
matrix, permutation functions	614
matrix, specialized solvers	570
matrix, zero elements	615
maximum field width (<code>scanf</code>)	285
memory management	596
memory, displaying Java memory status	995
memory, limitations on JVM	989
messages, error	35
method, invoking a method of a Java object	993
methods, displaying available methods	
of a Java object	993
mex	970
mex-files	970
minimum field width (<code>printf</code>)	279
missing data	43
<code>mkOctfile</code>	946
multi-line comments	37
multiplication	149, 846

N

negation	149, 846
NLP	697
nonlinear equations	599
nonlinear programming	697
not operator	153, 846
numeric constant	42, 47
numeric value	42, 47

O

object groups	460
object, creating a Java object	990
obsolete functions	1009
oct	946
oct-files	946
Octave and MATLAB, how to	
distinguish between	1018
Octave API	945
Octave command options	15
Octave, calling from Java	986
.octaverc	16, 19
~/.octaverc	16, 19
ODE	677
online help	20
opengl rendering slow windows	474
operator precedence	159
operators, arithmetic	149, 846
operators, assignment	156
operators, boolean	153, 846
operators, decrement	159
operators, increment	159
operators, logical	153, 846
operators, relational	152, 846
optimization	596, 697
options, Octave command	15
--braindead	17
--built-in-docstrings-file filename	15
--debug	15
--debug-jit	15
--doc-cache-file filename	15
--echo-commands	15
--eval code	15
--exec-path path	15
--gui	16
--help	16
--image-path path	16
--info-file filename	16
--info-program program	16
--interactive	16
--jit-compiler	16
--line-editing	16
--no-gui	16
--no-history	16
--no-init-file	16
--no-init-path	16
--no-line-editing	16
--no-site-file	16

--no-window-system	16
--norc	17
--path path	17
--persist	17
--quiet	17
--silent	17
--texi-macros-file filename	17
--traditional	17
--verbose	17
--version	18
-d	15
-f	17
-h	16
-H	16
-i	16
-p path	17
-q	17
-v	18
-V	17
-W	16
-x	15
or operator	153, 846
oregonator	679
otherwise statement	167
output conversions, for printf	280

P

parser	1035
patch graphics object	395
patches, submitting	1017
path, adding to classpath	994
path, removing from classpath	994
permutation matrix functions	614
persistent statement	130
persistent variables	130
personal startup file	19
PKG_ADD	938
PKG_DEL	938
plotting	295
plotting, high-level	295
plotting, multiple plot windows	372
plotting, multiple plots per figure	371
plotting, object manipulation	372
plotting, saving and printing plots	381
plotting, three-dimensional	331
plotting, two-dimensional functions	327
plotting, window manipulation	374
precision (printf)	280
printing notation	11
printing plots	381
profiler	244
program, self contained	36
Progress Bar	868
project startup file	19
prompt customization	32
pseudoinverse	553, 612

Q

- QP 697
 QR factorization 559
 quadratic programming 697
 quitting octave 7, 19
 quiver group 469
 quotient 149, 846

R

- range expressions 52
readline customization 32
 recycling 585
 relational operators 152, 846
 reporting bugs 1015, 1016
 results, incorrect 1015, 1017
 root figure graphics object 395

S

- saving data 258
 saving graphics objects 404
 saving plots 381
 scatter group 470
 Schur decomposition 564
 script files 177
 scripts 36
 select JVM version 989
 self contained programs 36
 series objects 465, 466, 467, 468, 471
 short-circuit evaluation 155
 side effect 157
 SIMD 585
 singular value decomposition 565
 site exiting file 19
 site startup file 16, 18
 Spearman's Rho 740
 speedups 596
 stair group 470
 startup 18
 startup files 18
startup.m 19
 statements 165
 static classpath 986, 993
 stem series 471
 strings 43, 67
 structural rank 642
 structure elements, looping over 171
 structures 43, 101
 submitting diffs 1017
 submitting patches 1017

- subtraction 149, 846
 suggestions 1016
 surface graphics object 395
 surface group 472
switch statement 167

T

- test functions 997
 tests for equality 152, 846
 text graphics object 395
 toolkit customization 473
 toolkits, graphics 473
 transform groups 472
 transpose 149, 846
 transpose, complex-conjugate 149, 846
 troubleshooting 1015
try statement 174

U

- uitable properties 446
 unary minus 149, 846
 undefined behavior 1015
 undefined function value 1015
unwind_protect statement 174
unwind_protect_cleanup 174
 use of comments 37
 user-defined data types 44
 user-defined functions 177
 user-defined variables 127
 using Octave with Java 986

V

- varargin** 188
varargout 191
 variable-length argument lists 188
 variable-length return lists 191
 variables, global 129
 variables, persistent 130
 variables, user-defined 127
 vectorization 583
 vectorize 583
 version startup file 19

W

- warning ids 230
 warranty 1037
while statement 169
 wrong answers 1015, 1017

Function Index

A

abs	511
accumarray	593
accumdim	594
acos	513
acosd	516
acosh	514
acot	514
acotd	516
acoth	515
acscl	514
acscl	516
acsch	515
add_input_event_hook	1035
addlistener	462
addpath	196
addpref	882
addproperty	461
addtodate	895
airy	526
all	475
allchild	404
amd	633
ancestor	403
and	154
angle	511
annotation	369
any	475
arch_fit	796
arch_rnd	797
arch_test	797
area	321, 322
arg	511
argnames	217
argv	18
arma_rnd	797
arrayfun	588
ascii	910
asctime	887
asec	513
asecd	516
asech	514
asin	513
asind	516
asinh	514
assert	1003
assignin	164
atan	513
atan2	515
atan2d	516
atand	516
atanh	514
atexit	20
audiodevinfo	828

audioformats	828
audioinfo	827
audioplayer	829
audioread	827
audiorecorder	830
audiowrite	828
autoload	203
autoreg_matrix	798
autumn	818
available_graphics_toolkits	473
axes	396
axis	323, 324

B

balance	545
bandwidth	546
bar	301
barh	302
bartlett	798
base2dec	94
base64_decode	912
base64_encode	912
beep	223
beep_on_error	223
besselh	529
besselj	528
besselk	527
bessely	527
beta	530
betainc	530
betaincinv	530
betaln	531
bicg	570
bicgstab	572
bin2dec	92
binary	911
bincoeff	531
bitand	59
bitcmp	59
bitget	58
bitor	59
bitpack	41
bitset	58
bitshift	60
bitunpack	42
bitxor	59
blackman	798
blanks	70
blkdiag	489
blkmm	569
bone	818
bounds	713
box	367

brighten	822	clock	889
bsxfun	586	close	377, 910
built_in_docstrings_file	23	closereq	378
builtin	202	cmpermute	824
bunzip2	906	cmunique	823
byte_size	951	colamd	634
bzip2	909	colloc	668
C			
calendar	895	colon	844
camlight	346, 347	colorbar	367, 368
camlookat	350	colorcube	818
camorbit	351	colormap	816
campos	350	colperm	635
camroll	352	colstyle	457
camtarget	352	columns	44
camup	353	comet	323
camva	354	comet3	323
camzoom	354	command_line_path	199
canonicalize_file_name	905	common_size	477
cart2pol	539	commutation_matrix	532
cart2sph	539	compan	748
cast	40	compare_versions	929
cat	481	compass	320
caxis	325	completion_append_char	27
cbrt	510	completion_matches	27
ccolamd	633	complex	48
cd	910, 922	computer	926
ceil	518	cond	546
cell	116	condeig	547
cell2mat	124	condest	640
cell2struct	124	confirm_recursive_rmdir	899
celldisp	115	conj	511
cellfun	590	containers.Map	114
cellindexmat	122	contour	310
cellslices	120	contour3	312
cellstr	123	contourc	311
center	736	contourf	311
cgs	574	contrast	822
char	71	conv	749
chdir	922	conv2	750
chol	555	convhull	786
chol2inv	555	convhulln	787
choldelete	556	convn	749
cholinsert	556	cool	818
cholinv	555	copper	819
cholshift	556	copyfile	897
cholupdate	556	copyobj	404
circshift	483	corr	739
citation	5	corrcoef	739
cla	377	cos	513
clabel	366	cosd	515
class	39	cosh	514
clc	26	cosint	532
clear	135	cot	513
clearvars	136, 137	cotd	516
clf	376	coth	514
		cov	739
		cplxpairs	511
		cputime	890

crash.dumps_octave_core.....	271	dec2hex	93
cross.....	522	decic.....	684
csc.....	513	deconv.....	750
cscd.....	515	deg2rad.....	512
csch.....	514	del2.....	523
cstrcat.....	73	delaunay.....	775
csvread.....	265	delaunayn.....	776
csvwrite.....	264	delete.....	377, 911
csymamd.....	635	dellistener.....	462
ctime.....	886	demo.....	1005
ctranspose.....	150	det.....	547
cubehelix.....	819	detrend.....	798
cummax.....	520	diag.....	489
cummin.....	521	dialog.....	867
cumprod.....	518	diary.....	14, 34
cumsum.....	517	diff.....	476
cumtrapz.....	668	diffpara.....	798
curl.....	523	diffuse.....	345
cylinder.....	362	dims.....	952
D			
daspect.....	358	dir.....	910, 923
daspk.....	688	dir_in_loadpath.....	199
daspk_options.....	689	disable_diagonal_matrix.....	609
dasrt.....	694	disable_permutation_matrix.....	609
dasrt_options.....	695	disable_range.....	52
dassl.....	691	discrete_cdf.....	741
dassl_options.....	693	discrete_inv.....	741
date.....	890	discrete_pdf.....	741
datenum.....	892	discrete_rnd.....	742
datestr.....	893	disp.....	251
datetick.....	896	display.....	850
datevec.....	895	divergence.....	522
dawson.....	532	dlmread.....	264
dbclear.....	240, 241	dlmwrit.....	263
dbcont.....	238	dmperm.....	636
dbdown.....	244	do_string_escapes.....	97
dblist.....	242	doc.....	21
dblquad.....	669	doc_cache_create.....	24
dbnext.....	243	doc_cache_file.....	23
dbquit.....	238	dos.....	914
dbstack.....	243	dot.....	522
dbstatus.....	240	double.....	47
dbstep.....	243	drawnow.....	374
dbstop.....	238	dsearch.....	781
dbtype.....	242	dsearchn.....	782
dbup.....	244	dup2.....	917
dbwhere.....	242	duplication_matrix.....	532
deal.....	191	durbinlevinson.....	799
deblank.....	77		
debug_java.....	996		
debug_jit.....	596		
debug_on_error.....	237		
debug_on_interrupt.....	237		
debug_on_warning.....	237		
dec2base.....	93		
dec2bin.....	92		

E

e.....	540
echo.....	34
edit.....	194
edit_history.....	29
EDITOR.....	31
eig.....	547
eigs.....	644
elem.....	951
ellipj.....	532
ellipke.....	533
ellipsoid.....	363
empirical_cdf.....	741
empirical_inv.....	741
empirical_pdf.....	741
empirical_rnd.....	742
endgrent.....	926
endpwent.....	925
eomday.....	896
eps.....	542
eq.....	153
erase.....	87
erf.....	533
erfc.....	534
erfcinv.....	534
erfcx.....	534
erfi.....	534
erfinv.....	534
errno.....	227
errno_list.....	227
error.....	221
errorbar.....	313
errordlg.....	863
etime.....	890
etree.....	626
etreeplot.....	627
eval.....	161
evalc.....	161
evalin.....	164
example.....	1006
exec.....	917
EXEC_PATH.....	916
exist.....	134
exit.....	19
exp.....	509
expint.....	534
expm.....	568
expm1.....	509
eye.....	489
ezcontour.....	329
ezcontourf.....	330
ezmesh.....	359
ezmeshc.....	360
ezplot.....	328
ezplot3.....	359
ezpolar.....	330
ezsurf.....	361
ezsurfc.....	361

F

factor.....	524
factorial.....	524
fail.....	1004
false.....	61
fclear.....	291
fclose.....	275
fcntl.....	920
fdisp.....	263
feather.....	320, 321
feof.....	291
ferror.....	291
feval.....	162
fflush.....	256
fft.....	791
fft2.....	791
fftconv.....	794
fftfilt.....	794
fftn.....	792
fftshift.....	799
fftw.....	793
fgetl.....	276
fgets.....	276
fieldnames.....	109
figure.....	372
file_in_loadpath.....	198
file_in_path.....	904
fileattrib.....	902
fileparts.....	904
fileread.....	263
filesep.....	904
fill.....	322
filter.....	794
filter2.....	795
find.....	477
findall.....	454, 455
findfigs.....	404
findobj.....	453, 454
findstr.....	78
fix.....	518
fixed_point_format.....	51
flag.....	819
flintmax.....	56
flip.....	479
fliplr.....	479
flipud.....	479
floor.....	518
fminbnd.....	603
fminsearch.....	605
fminunc.....	604
foo.....	13
fopen.....	274
fork.....	917
format.....	252
formula.....	217
fortran_vec.....	952
fplot.....	327, 328
fprintf.....	277

fputs.....	276	getpwent.....	924
fractdiff.....	799	getpwnam.....	925
frame2im.....	816	getpwuid.....	924
fread.....	286	getrusage.....	930
freport.....	291	getuid.....	921
freqz.....	795	ginput.....	392
freqz_plot.....	796	givens.....	548
frewind.....	292	glob.....	903
fscanf.....	283	glpk.....	697
fseek.....	292	gls.....	708
fskip1.....	277	gmres.....	576
fsolve.....	599	gmtime.....	886
ftell.....	292	gnuplot_binary.....	473
ftp.....	909	gplot.....	627
full.....	622	grabcode.....	209
fullfile.....	905	gradient.....	521
func2str.....	215	graphics_toolkit.....	473
functions.....	214	gray.....	819
fwrite.....	289	gray2ind.....	814
fzero.....	602	grid.....	367

G

gallery	500, 501, 502, 503, 504, 505
gamma.....	535
gammainc.....	535
gammairninv.....	536
gammaln.....	537
gca.....	401
gcbf.....	459
gcbo.....	459
gcd.....	524
gcf.....	400
gco.....	401
ge.....	153
genpath.....	197
get.....	402, 830, 832
get_first_help_sentence.....	25
get_help_text.....	24
get_help_text_from_file.....	24
get_home_directory.....	922
getappdata.....	459
getaudiodata.....	831
getegid.....	921
getenv.....	922
geteuid.....	921
getfield.....	111
getframe.....	815
getgid.....	921
getgrent.....	925
getgrgid.....	925
getgrnam.....	925
gethostname.....	909
getpggrp.....	921
getpid.....	921
getplayer.....	831
getppid.....	921
getpref	881, 882

H

hadamard.....	505
hamming.....	799
hankel.....	505
hanning.....	799
hash.....	932
have_window_system.....	879
hdl2struct.....	404
help.....	20
helpdlg.....	863
hess.....	557
hex2dec.....	93
hex2num.....	94
hgroup.....	460
hgload.....	390
hgsave.....	390
hgtransform.....	472
hidden.....	335
hilb.....	505
hist.....	303
histc.....	736
history.....	28
history_control.....	30
history_file.....	31
history_save.....	30

history_size.....	31	inputParser.StructExpand.....	187
history_timestamp_format_string.....	31	inputParser.Unmatched.....	186
hold.....	375, 376	inputParser.UsingDefaults.....	186
home.....	26	int16.....	55
horzcat.....	481	int2str.....	75
hot.....	819	int32.....	55
housh.....	567	int64.....	55
hsv.....	820	int8.....	54
hsv2rgb.....	824	integral.....	665
humps.....	607	integral2.....	672
hurst.....	800	integral3.....	673
hypot.....	521	interp1.....	765
		interp2.....	769
I		interp3.....	770
i.....	541	interpft.....	767
ichol.....	652	interpn.....	771
idivide.....	57	intersect.....	743
ifelse.....	156	intmax.....	55
ifft.....	791	intmin.....	56
ifft2.....	792	inv.....	549
ifftn.....	792	invhilb.....	506
ifftshift.....	799	ipermute.....	482
ignore_function_time_stamp.....	196	iqr.....	713
ilu.....	653, 654	is_absolute_filename.....	906
im2double.....	813	is_dq_string.....	68
im2frame.....	816	is_function_handle.....	214
imag.....	512	is_leap_year.....	890
image.....	812	is_rooted_relative_filename.....	906
IMAGE_PATH.....	807	is_sq_string.....	68
imagesc.....	812	is_valid_file_id.....	275
imfinfo.....	808	isa.....	39
imformats.....	810	isalnum.....	98
importdata.....	270	isalpha.....	98
imread.....	805	isappdata.....	460
imshow.....	811	isargout.....	191
imwrite.....	806	isascii.....	99
ind2gray.....	814	isaxes.....	400
ind2rgb.....	814	isbanded.....	65
ind2sub.....	144	isbool.....	63
index.....	79	iscell.....	116
Inf.....	541	iscellstr.....	123
inf.....	541	ischar.....	68
inferiorito.....	847	iscntrl.....	99
info.....	22	iscolormap.....	817
info_file.....	22	iscolumn.....	64
info_program.....	22	iscomplex.....	63
inline.....	216	isdebugmode.....	243
inpolygon.....	785	isdefinite.....	65
input.....	256	isdeployed.....	927
inputdlg.....	864	isdiag.....	65
inputname.....	180	isdigit.....	99
inputParser.....	186	isdir.....	903
inputParser.CaseSensitive.....	186	isempty.....	46
inputParser.FunctionName.....	187	isequal.....	153
inputParser.KeepUnmatched.....	187	isequaln.....	153
inputParser.Parameters.....	186	isfield.....	110
inputParser.Results	186	isfigure.....	400
		isfile.....	903

isfinite	477	isvector	64
isfloat	63	isxdigit	99
isfolder	903	I	541
isglobal	130		
isgraph	99		
isgraphics	399		
isguirunning	879	J	
ishandle	399	j	541
ishermitian	64	java_get	992
ishghandle	399	java_matrix_autoconversion	996
ishold	376	java_set	992
isieee	927	java_unsigned_autoconversion	996
isindex	145	javaaddpath	994
isinf	476	javaArray	991
isinteger	54	javachk	994
isjava	991	javaclasspath	993
iskeyword	1035	javamem	995
isletter	98	javaMethod	993
islogical	63	javaObject	991
islower	98	javarmpath	994
ismac	927	jet	820
ismatrix	63	jit_enable	595
ismember	744	jit_failcnt	596
ismethod	837	jit_startcnt	595
isna	43	J	541
isnan	476		
isnull	46		
isnumeric	63	K	
isobject	836	kbhit	257
isocaps	340	kendall	740
isocolors	341, 342	keyboard	241
isonormals	340	kill	921
isosurface	338	kron	569
ispc	927	krylov	567
isplaying	830	kurtosis	716
ispref	883		
isprime	66	L	
isprint	99	lasterr	225
isprop	395	lasterror	224
ispunct	99	lastwarn	230
isreal	63	lcm	524
isrecording	831	ldivide	150
isrow	64	le	153
isscalar	64	legend	364
issorted	486	legendre	536
isspace	99	length	45
issparse	624	lgamma	537
issquare	64	license	929
isstring	68	light	398
isstrprop	100	lighting	345
isstruct	109	lin2mu	832
isstudent	927	line	396
issymmetric	64	lines	820
istril	65	link	898
istriu	66	linkaxes	464
isunix	927	linkprop	463
isupper	99	linsolve	550
isvarname	127	linspace	493

list_in_columns.....	251	mexext	971
list_primes.....	526	mfilename	195
listdlg	864	mget	910
load.....	261	mgorth	553
loaded_graphics_toolkits.....	473	min	520
loadobj	840	minus	150
localfunctions	199	mislocked	205
localtime	886	missing_component_hook	942
log.....	509	missing_function_hook	1036
log10.....	509	mkdir	898, 911
log1p.....	509	mkfifo	899
log2.....	509	mkoctfile	946
logical	61	mkpp	761
loglog	300	mkstemp	289
loglogerr	316	mktime	887
logm.....	568	mldivide	150
logspace	494	mlock	204
lookfor	21	mod	525
lookup	478	mode	712
lower	96	moment	716
ls	923	more	255
ls_command	923	movefile	897
lscov.....	709	movegui	879
lsode.....	677	movfun	719
lsode_options	678	movie	815
lsqnonneg.....	708	movmad	721
lstat.....	900	movmax	723
lt	153	movmean	724
lu	557	movmedian	726
luupdate	558	movmin	727
		movprod	729
		movslice	721
		movstd	731
		movsum	732
		movvar	734
		mpoles	749
		mpower	150
		mput	910
		mrdivide	151
		msgbox	865
		mtimes	151
		mu2lin	832
		munlock	204

M

mad.....	713
magic.....	506
make_absolute_filename.....	906
makeinfo_program.....	22
mat2cell	118
mat2str	74
material	346
matlab.lang.makeUniqueStrings	128
matlab.lang.makeValidName	127
matlabroot	928
matrix_type	550
max	519
max_recursion_depth	147
max_stack_depth	148
mean	711
meansq	714
median	712
menu	257
merge	156
mesh	333
meshc	334
meshgrid	347
meshz	334
methods	837
mex	970

N

namelengthmax	128
nan	542
NaN	542
nargin	179
narginchk	183
nargout	182
nargoutchk	183
native_float_format	263
native2unicode	97
NA	43
nchoosek	737
ndgrid	348

ndims..... 44, 952
 ne..... 153
 newplot..... 375
 news..... 21
 nextpow2..... 510
 nnz..... 624
 nonzeros..... 624
 norm..... 551
 normest..... 639
 normest1..... 639, 640
 not..... 154
 now..... 885
 nproc..... 926
 nth_element..... 486
 nthargout..... 181
 nthroot..... 510
 null..... 552
 num2cell..... 117
 num2hex..... 94
 num2str..... 74
 numel..... 44, 951
 numfields..... 109
 nzmax..... 624

O

ocean..... 820
 octave_core_file_limit..... 272
 octave_core_file_name..... 272
 octave_core_file_options..... 272
 OCTAVE_EXEC_HOME..... 927
 OCTAVE_HOME..... 927
 OCTAVE_VERSION..... 928
 ode15i..... 683
 ode15s..... 682
 ode23..... 681
 ode45..... 680
 odeget..... 687
 odeplot..... 687
 odeset..... 685
 ols..... 707
 onCleanup..... 227
 ones..... 490
 open..... 914
 openfig..... 391
 openvar..... 880
 operator..... 951
 optimget..... 710
 optimize_subsasgn_calls..... 842
 optimset..... 709
 or..... 155
 ordeig..... 565
 orderfields..... 111
 ordschur..... 565
 orient..... 389
 orth..... 552
 ostrsplit..... 83
 output_precision..... 50

P

P_tmpdir..... 290
 pack..... 137
 padecoeff..... 760
 page_output_immediately..... 256
 page_screen_output..... 255
 PAGER..... 255
 PAGER_FLAGS..... 255
 pan..... 372, 373
 pareto..... 309
 parseparams..... 189
 pascal..... 506
 patch..... 397
 path..... 197
 pathdef..... 198
 pathsep..... 198
 pause..... 830, 831, 891, 892
 pbaspect..... 358
 pcg..... 647
 pchip..... 800
 pclose..... 916
 pcOLOR..... 321
 pcr..... 650
 peaks..... 393
 periodogram..... 800
 perl..... 915
 perms..... 737
 permute..... 482
 pi..... 540
 pie..... 317, 318
 pie3..... 318
 pink..... 820
 pinv..... 553
 pipe..... 917
 pkg..... 933, 934
 planerot..... 549
 play..... 829, 831
 playblocking..... 829
 plot..... 296
 plot3..... 348
 plotmatrix..... 308
 plotyy..... 298
 plus..... 151
 pol2cart..... 539
 polar..... 317
 poly..... 763
 polyaffine..... 753
 polyarea..... 784
 polyder..... 752
 polyeig..... 748
 polyfit..... 753
 polygcd..... 750
 polyint..... 753
 polyout..... 763
 polyreduce..... 763
 polyval..... 747
 polyvalm..... 747
 popen..... 915

popen2 916
 postpad 489
 pow2 510
 power 151
 powerset 745
 ppder 762
 ppint 762
 ppjumps 762
 ppval 762
 pqpnonneg 704
 prctile 718
 prefdir 883
 preferences 883
 prepad 488
 primes 525
 print 383
 print_empty_dimensions 51
 print_struct_array_contents 103
 print_usage 223
 printd 305
 printf 277
 prism 821
 prod 517
 profexplore 246
 profexport 246
 profile 244
 profshow 245
 program_invocation_name 18
 program_name 18
 PS1 33
 PS2 33
 PS4 33
 psi 537
 publish 207
 putenv 922
 puts 276
 pwd 924
 python 915

Q

qmr 579
 qp 703
 qr 559
 qrdelete 562
 qrinsert 561
 qrshift 562
 qrupdate 561
 quad 660
 quad_options 661
 quad2d 670
 quadcc 664
 quadgk 663
 quadl 662
 quadv 661
 quantile 717
 questdlg 866
 quit 19

quiver 318, 319
 quiver3 319
 qz 562
 qzhess 563

R

rad2deg 512
 rainbow 821
 rand 494
 rande 496
 randg 498
 randi 495
 randn 496
 randp 497
 randperm 499
 range 713
 rank 553
 ranks 738
 rat 538
 rats 538
 rcond 554
 rdivide 151
 readdir 903
 readline_re_read_init_file 32
 readline_read_init_file 32
 readlink 898
 real 512
 reallog 509
 realmax 543
 realmin 543
 realpow 510
 realsqrt 510
 record 831, 833
 recordblocking 831
 rectangle 331
 rectint 784
 recycle 906
 reducepatch 343
 reducevolume 343
 refresh 374
 refreshdata 465
 regexp 88
 regexpi 90
 regexprep 91
 regextranslate 91
 register_graphics_toolkit 473
 rehash 198
 rem 525
 remove_input_event_hook 1036
 rename 897, 911
 repelem 491
 repelems 491
 repmat 491
 reset 456
 reshape 482
 residue 751
 resize 483, 952

restoredefaultpath	199	sech	514
resume	830, 831	SEEK_CUR	292
rethrow	226	SEEK_END	292
return	193	SEEK_SET	292
rgb2gray	825	semilogx	299
rgb2 hsv	824	semilogxerr	315
rgb2ind	814	semilogy	299
rgbplot	817	semilogyerr	316
ribbon	355, 356	set	402, 830, 832
rindex	79	setappdata	459
rmappdata	460	setdiff	744
rmdir	899, 911	setenv	922
rmfield	111	setfield	110
rmpath	197	setgrent	925
rmpref	883	setpref	882
roots	748	setpwent	925
rose	309	setxor	744
rosser	506	shading	356
rot90	480	shg	377
rotate	373	shift	484
rotate3d	373	shiftdim	484
rotdim	480	shrinkfaces	344
round	519	sighup_dumps_octave_core	271
roundb	519	sign	526
rows	44	signbit	526
rref	554	sigquit_dumps_octave_core	271
rsf2csf	564	sigterm_dumps_octave_core	272
rticks	326	SIG	921
run	163	silent_functions	180
run_count	738	sin	513
run_history	29	sinc	796
rundemos	1006	sind	515
runlength	738	sinetone	801
runtests	1006	sinewave	801
S			
S_ISBLK	901	single	53
S_ISCHR	901	sinh	514
S_ISDIR	901	sinint	538
S_ISFIFO	901	size	45
S_ISLNK	901	size_equal	46
S_ISREG	902	sizemax	49
S_ISSOCK	902	sizeof	46
save	258	skewness	715
save_default_options	260	slice	355
save_header_format_string	261	smooth3	342
save_precision	260	sombrero	393
saveas	389	sort	485
savefig	391	sortrows	486
saveobj	840	sound	833
savepath	197	soundsc	833
scanf	283	source	207
scatter	307	spalloc	622
scatter3	356	sparse	622
schur	564	sparse_auto_mutate	629
sec	513	spaugment	643
secd	515	spconvert	623
		spdiags	619
		spearman	740
		spectral_adf	801

spectral_xdf	802	strtok	81
specular	345	strtrim	78
speed	1006	strtrunc	78
spencer	802	struct	108
speye	620	struct_levels_to_print	103
spfun	590	struct2cell	113
sph2cart	540	struct2hdl	404
sphere	362	structfun	592
spinmap	822	strvcat	72
spline	769	sub2ind	143
splinefit	754	subplot	371
split_long_rows	50	subsasgn	842
spones	620	subsindex	843
spparms	641	subspace	565
sprand	620	subsref	840
sprandn	620	substr	87
sprandsym	621	substruct	112
sprank	642	sum	516
spring	821	summer	821
sprintf	278	sumsq	518
spstats	624	superioro	847
spy	626	suppress_verbose_help_message	24
sqp	705	surf	335
sqrt	510	surface	398
sqrtm	569	surfc	336
squeeze	46	surfl	336
sscanf	284	surfnorm	337
stairs	305	svd	565
stat	900	svd_driver	567
statistics	718	svds	646
std	714	swapbytes	41
stderr	273	sylvester	569
stdin	273	symamd	636
stdout	273	symbfact	642
stem	305	symlink	898
stem3	307	symrcm	637
stemleaf	304	symvar	217
stft	802	synthesis	802
stop	830, 831	system	913
str2double	95		
str2func	215		
str2num	96		
strcat	72	T	
strchr	78	tan	513
strcmp	76	tand	515
strcmpi	77	tanh	514
strfind	79	tar	907
strftime	887	tempdir	290
string_fill_char	69	tempname	290
strjoin	80	terminal_size	252
strjust	96	test	997
strmatch	80	tetramesh	778
strncmp	76	texi_macros_file	23
strncmpi	77	text	366
strptime	889	textread	265
strrread	84	textscan	266
strrep	86	tfqmr	580
strsplit	81	thetaticks	326
		tic	891

tilde_expand	905	unpack	908
time	885	unsetenv	922
times	151	untabify	91
title	364	untar	907
tmpfile	289	unwrap	796
toc	891	unzip	908
toeplitz	507	uplus	152
tolower	96	upper	97
toupper	97	urlread	911
trace	554	urlwrite	912
transpose	152	usejava	995
trapz	667		
treelayout	627		
treeplot	627		
tril	487		
trimesh	777		
triplequad	670		
triplot	777		
trisurf	778		
triu	487		
true	61		
tsearch	781		
tsearchn	781		
type	137		
typecast	40		
typeinfo	39		

U

uibuttongroup	870
uicontextmenu	876
uicontrol	870
uigetdir	861
uigetfile	861
uimenu	875
uint16	55
uint32	55
uint64	55
uint8	55
uipanel	869
uipushtool	877
uiputfile	862
uiresume	881
uitable	872
uitoggletool	878
uitoolbar	877
uiwait	880
umask	899
uminus	152
uname	926
undo_string_escapes	97
unicode_idx	92
unicode2native	97
union	744
unique	743
unix	913
unlink	898
unmkpp	761

V

validateattributes	184
validatestring	184
vander	507
var	715
vec	488
vech	488
vecnorm	554
vectorize	584
ver	928
version	928
vertcat	481
view	349
viridis	821
voronoi	782
voronoin	783

W

waitbar	868, 869
waitfor	881
waitForbuttonpress	392
waitpid	918
warndlg	867
warning	228
warranty	22
waterfall	357
WCONTINUE	918
WCOREDUMP	918
weekday	896
WEXITSTATUS	919
what	138
which	138
white	821
whitebg	822, 823
who	132
whos	132
whos_line_format	133
WIFCONTINUED	919
WIFEXITED	919
WIFSIGNALLED	919
WIFSTOPPED	919
wilkinson	507
winqueryreg	931
winter	822

WNOHANG	919
WSTOPSIG	919
WTERMSIG	920
WUNTRACED	920

X

xlabel	366
xlim	326
xor	476
xticklabels	327
xticks	326

Y

yes_or_no	257
ylabel	366
ylim	325
yticklabels	327
yticks	326
yulewalker	803

Z

zeros	490
zip	908
zlabel	366
zlim	325
zoom	373
zscore	736
zticklabels	327
zticks	326

Operator Index

!	
!.....	154, 846
!=.....	152, 153, 846
"	
".....	43, 67
#	
# comment marker	37
#! self-contained script	36
#{ block comment marker	37
%	
% comment marker	37
%{ block comment marker	37
&	
&.....	154, 846
&&.....	155
,	
'.....	43, 67, 150, 846
(
(.....	139
)	
).....	139
*	
*.....	149, 151, 846
**.....	149, 150
*=.....	158
+	
+.....	149, 150, 151, 152, 846
++.....	159
+=.....	158
,	
,.....	48
-	
-.....	149, 150, 152, 846
--.....	159
-=.....	158
.	
.	101
.!	150, 152, 846
.*	149, 151, 846
.**	149, 151
.+	149
. . continuation marker	175
./	149, 151, 846
.^	149, 151, 846
.\	149, 150, 846
/	
/.....	149, 151, 846
/=.....	158
:	
:.....	139, 846
:: indexing expressions	140
:: range expressions	52
;	
;;.....	48
<	
<.....	152, 153, 846
<=.....	152, 153, 846
=	
=.....	156
==.....	152, 846
>	
>.....	152, 153, 846
>=.....	152, 153, 846
@	
@ class methods	835
@ function handle	148

[48	{	114
]	48	}	
^			
^	149, 150, 846		154, 155, 846
\	149, 150, 846	~	154
\ continuation marker	175	=	152, 153, 846

Graphics Properties Index

Axes Properties	412
Figure Properties	407
Image Properties	425
Light Properties	434
Line Properties	419
Patch Properties	427
Root Figure Properties	405
Surface Properties	430
Text Properties	422
Uibuttongroup Properties	437
Uicontextmenu Properties	440
Uicontrol Properties	444
Uimenu Properties	436
Uipanel Properties	441
Uipushtool Properties	450
Uitoggletool Properties	452
UIToolbar Properties	449
axes fontsize	414
axes fontsmoothing	414
axes fontunits	414
axes fontweight	414
axes gridalpha	414
axes gridalphemode	414
axes gridcolor	414
axes gridcolormode	414
axes gridlinestyle	415
axes handlevisibility	415
axes hittest	415
axes interruptible	415
axes labelfontsizemultiplier	415
axes layer	415
axes linestyleorder	415
axes linestyleorderindex	415
axes linewidth	415
axes minorgridalpha	415
axes minorgridalphemode	415
axes minorgridcolor	415
axes minorgridcolormode	415
axes minorgridlinestyle	415
axes mousewheelzoom	415
axes nextplot	415
axes outerposition	416
axes parent	416
axes pickableparts	416
axes plotboxaspectratio	416
axes plotboxaspectratiomode	416
axes position	416
axes projection	416
axes selected	416
axes selectionhighlight	416
axes sortmethod	416
axes tag	416
axes tickdir	416
axes tickdirmode	416
axes ticklabelinterpreter	417
axes ticklength	417
axes tightinset	417
axes title	417
axes titlefontsizemultiplier	417
axes titlefontweight	417
axes type	417
axes uicontextmenu	417
axes units	417
axes userdata	417
axes view	417
axes visible	417
axes xaxislocation	417
axes xcolor	417
axes xcolormode	417
axes xdir	417
axes xgrid	417

axes xlabel	418	figure currentpoint	407
axes xlim	418	figure deletefcn	408
axes xlimmode	418	figure dockcontrols	408
axes xminorgrid	418	figure filename	408
axes xminortick	418	figure graphicssmoothing	408
axes xscale	418	figure handlevisibility	408
axes xtick	418	figure hittest	408
axes xticklabel	418	figure integerhandle	408
axes xticklabelmode	418	figure interruptible	408
axes xticklabelrotation	418	figure inverthardcopy	408
axes xtickmode	418	figure keypressfcn	408
axes yaxislocation	418	figure keyreleasefcn	409
axes ycolor	418	figure menubar	409
axes ycolormode	418	figure name	409
axes ydir	418	figure nextplot	409
axes ygrid	418	figure number	409
axes ylabel	418	figure numbertitle	409
axes ylim	418	figure outerposition	409
axes ylimmode	418	figure paperorientation	409
axes yminorgrid	418	figure paperposition	409
axes yminortick	418	figure paperpositionmode	410
axesyscale	419	figure papersize	410
axes ytick	419	figure papertype	410
axes yticklabel	419	figure paperunits	410
axes yticklabelmode	419	figure parent	410
axes yticklabelrotation	419	figure pickableparts	410
axes ytickmode	419	figure pointer	410
axes zcolor	419	figure pointershapecdata	410
axes zcolormode	419	figure pointershapehotspot	410
axes zdir	419	figure position	410
axes zgrid	419	figure renderer	410
axes zlabel	419	figure renderermode	411
axes zlim	419	figure resize	411
axes zlimmode	419	figure resizefcn	411
axes zminorgrid	419	figure selected	411
axes zminortick	419	figure selectionhighlight	411
axes zscale	419	figure selectiontype	411
axes ztick	419	figure sizechangedfcn	411
axes zticklabel	419	figure tag	411
axes zticklabelmode	419	figure toolbar	411
axes zticklabelrotation	419	figure type	411
axes ztickmode	419	figure uicontextmenu	411
		figure units	411
		figure userdata	411
		figure visible	411
		figure windowbuttondownfcn	411
		figure windowbuttonmotionfcn	411
		figure windowbuttonupfcn	411
		figure windowkeypressfcn	412
		figure windowkeyreleasefcn	412
		figure windowscrollwheelfcn	412
		figure windowstyle	412

F

figure alphamap	407
figure beingdeleted	407
figure busyaction	407
figure buttondownfcn	407
figure children	407
figure clipping	407
figure closerequestfcn	407
figure color	407
figure colormap	407
figure createfcn	407
figure currentaxes	407
figure currentcharacter	407
figure currentobject	407

I

image alphadata	425
image alphadatamapping	425
image beingdeleted	425
image busyaction	425
image buttondownfcn	425
image cdata	425
image cdatamapping	425
image children	425
image clipping	425
image createfcn	425
image deletefcn	425
image handlevisibility	425
image hittest	425
image interruptible	426
image parent	426
image pickableparts	426
image selected	426
image selectionhighlight	426
image tag	426
image type	426
image uicontextmenu	426
image userdata	426
image visible	426
image xdata	426
image ydata	426

line deletefcn	420
line displayname	420
line handlevisibility	420
line hitest	420
line interruptible	420
line linejoin	421
line linestyle	421
line linewidth	421
line marker	421
line markeredgecolor	421
line markerfacecolor	421
line markersize	421
line parent	421
line pickableparts	421
line selected	421
line selectionhighlight	421
line tag	421
line type	421
line uicontextmenu	421
line userdata	421
line visible	422
line xdata	422
line xdatasource	422
line ydata	422
line ydatasource	422
line zdata	422
line zdatasource	422

L

light beingdeleted	434
light busyaction	434
light buttondownfcn	434
light children	434
light clipping	434
light color	434
light createfcn	434
light deletefcn	435
light handlevisibility	435
light hitest	435
light interruptible	435
light parent	435
light pickableparts	435
light position	435
light selected	435
light selectionhighlight	435
light style	435
light tag	435
light type	435
light uicontextmenu	435
light userdata	436
light visible	436
line beingdeleted	420
line busyaction	420
line buttondownfcn	420
line children	420
line clipping	420
line color	420
line createfcn	420

P

patch alphadatamapping	427
patch ambientstrength	427
patch backfacelighting	427
patch beingdeleted	427
patch busyaction	427
patch buttondownfcn	427
patch cdata	427
patch cdatamapping	427
patch children	427
patch clipping	427
patch createfcn	427
patch deletefcn	428
patch diffusestrength	428
patch displayname	428
patch edgealpha	428
patch edgecolor	428
patch edgelighting	428
patch facealpha	428
patch facecolor	428
patch facelighting	428
patch facenormals	428
patch facenormalsmode	428
patch faces	428
patch facevertexalphadata	429
patch facevertexcdata	429
patch handlevisibility	429
patch hitest	429
patch interruptible	429

patch linestyle	429
patch linewidth	429
patch marker	429
patch markeredgecolor	429
patch markerfacecolor	429
patch markersize	429
patch parent	429
patch pickableparts	429
patch selected	429
patch selectionhighlight	430
patch specularcolorreflectance	430
patch specularexponent	430
patch specularstrength	430
patch tag	430
patch type	430
patch uicontextmenu	430
patch userdata	430
patch vertexnormals	430
patch vertexnormalsmode	430
patch vertices	430
patch visible	430
patch xdata	430
patch ydata	430
patch zdata	430

R

root beingdeleted	405
root busyaction	405
root buttondownfcn	405
root callbackobject	405
root children	405
root clipping	405
root commandwindowsize	405
root createfcn	405
root currentfigure	405
root deletefcn	405
root fixedwidthfontname	405
root handlevisibility	405
root hittest	405
root interruptible	406
root monitorpositions	406
root parent	406
root pickableparts	406
root pointerlocation	406
root pointerwindow	406
root screendepth	406
root screenpixelsperinch	406
root screensize	406
root selected	406
root selectionhighlight	406
root showhiddenhandles	406
root tag	406
root type	406
root uicontextmenu	406
root units	406
root userdata	406
root visible	406

S

surface alphadata	430
surface alphadatamapping	430
surface ambientstrength	430
surface backfacelighting	431
surface beingdeleted	431
surface busyaction	431
surface buttondownfcn	431
surface cdata	431
surface cdatamapping	431
surface cdatasource	431
surface children	431
surface clipping	431
surface createfcn	431
surface deletefcn	431
surface diffusestrength	431
surface displayname	431
surface edgealpha	431
surface edgecolor	431
surface edgelighting	432
surface facealpha	432
surface facecolor	432
surface facelighting	432
surface facenormals	432
surface facenormalsmode	432
surface handlevisibility	432
surface hittest	432
surface interruptible	432
surface linestyle	432
surface linewidth	433
surface marker	433
surface markeredgecolor	433
surface markerfacecolor	433
surface markersize	433
surface meshstyle	433
surface parent	433
surface pickableparts	433
surface selected	433
surface selectionhighlight	433
surface specularcolorreflectance	433
surface specularexponent	433
surface specularstrength	433
surface tag	433
surface type	433
surface uicontextmenu	433
surface userdata	433
surface vertexnormals	434
surface vertexnormalsmode	434
surface visible	434
surface xdata	434
surface xdatasource	434
surface ydata	434
surface ydatasource	434
surface zdata	434
surface zdatasource	434

T

text backgroundcolor	422
text beingdeleted	422
text busyaction	422
text buttondownfcn	422
text children	422
text clipping	422
text color	422
text createfcn	422
text deletefcn	423
text edgecolor	423
text editing	423
text extent	423
text fontangle	423
text fontname	423
text fontsize	423
text fontunits	423
text fontweight	423
text handlevisibility	423
text hittest	423
text horizontalalignment	423
text interpreter	423
text interruptible	423
text linestyle	424
text linewidth	424
text margin	424
text parent	424
text pickableparts	424
text position	424
text rotation	424
text selected	424
text selectionhighlight	424
text string	424
text tag	424
text type	424
text uicontextmenu	424
text units	424
text userdata	424
text verticalalignment	424
text visible	425

U

uibuttongroup backgroundcolor	438
uibuttongroup beingdeleted	438
uibuttongroup bordertype	438
uibuttongroup borderwidth	438
uibuttongroup busyaction	438
uibuttongroup buttondownfcn	438
uibuttongroup children	438
uibuttongroup clipping	438
uibuttongroup createfcn	438
uibuttongroup deletefcn	438
uibuttongroup fontangle	438
uibuttongroup fontname	438
uibuttongroup fontsize	438
uibuttongroup fontunits	438
uibuttongroup fontweight	438

uibuttongroup foregroundcolor	438
uibuttongroup handlevisibility	439
uibuttongroup highlightcolor	439
uibuttongroup hittest	439
uibuttongroup interruptible	439
uibuttongroup parent	439
uibuttongroup pickableparts	439
uibuttongroup position	439
uibuttongroup resizefcn	439
uibuttongroup selected	439
uibuttongroup selectedobject	439
uibuttongroup selectionchangedfcn	439
uibuttongroup selectionhighlight	439
uibuttongroup shadowcolor	439
uibuttongroup sizechangedfcn	439
uibuttongroup tag	439
uibuttongroup title	439
uibuttongroup titleposition	440
uibuttongroup type	440
uibuttongroup uicontextmenu	440
uibuttongroup units	440
uibuttongroup userdata	440
uibuttongroup visible	440
uiicontextmenu beingdeleted	440
uiicontextmenu busyaction	440
uiicontextmenu buttondownfcn	440
uiicontextmenu callback	440
uiicontextmenu children	440
uiicontextmenu clipping	440
uiicontextmenu createfcn	440
uiicontextmenu deletefcn	440
uiicontextmenu handlevisibility	440
uiicontextmenu hittest	441
uiicontextmenu interruptible	441
uiicontextmenu parent	441
uiicontextmenu pickableparts	441
uiicontextmenu position	441
uiicontextmenu selected	441
uiicontextmenu selectionhighlight	441
uiicontextmenu tag	441
uiicontextmenu type	441
uiicontextmenu uicontextmenu	441
uiicontextmenu userdata	441
uiicontextmenu visible	441
uiicontrol backgroundcolor	444
uiicontrol beingdeleted	444
uiicontrol busyaction	444
uiicontrol buttondownfcn	444
uiicontrol callback	444
uiicontrol cdata	444
uiicontrol children	444
uiicontrol clipping	444
uiicontrol createfcn	444
uiicontrol deletefcn	444
uiicontrol enable	444
uiicontrol extent	445
uiicontrol fontangle	445
uiicontrol fontname	445

uicontrol fontsize	445	uipanel bordertype.....	442
uicontrol fontunits	445	uipanel borderwidth	442
uicontrol fontweight	445	uipanel busyaction	442
uicontrol foregroundcolor	445	uipanel buttondownfcn	442
uicontrol handlevisibility	445	uipanel children	442
uicontrol hittest	445	uipanel clipping	442
uicontrol horizontalalignment	445	uipanel createfcn	442
uicontrol interruptible	445	uipanel deletefcn	442
uicontrol keypressfcn	445	uipanel fontangle	442
uicontrol listboxtop	445	uipanel fontname	442
uicontrol max	445	uipanel fontsize	442
uicontrol min	445	uipanel fontunits	442
uicontrol parent	445	uipanel fontweight	442
uicontrol pickableparts	445	uipanel foregroundcolor	442
uicontrol position	446	uipanel handlevisibility	443
uicontrol selected	446	uipanel highlightcolor	443
uicontrol selectionhighlight	446	uipanel hittest	443
uicontrol sliderstep	446	uipanel interruptible	443
uicontrol string	446	uipanel parent	443
uicontrol style	446	uipanel pickableparts	443
uicontrol tag	446	uipanel position	443
uicontrol tooltipstring	446	uipanel resizefcn	443
uicontrol type	446	uipanel selected	443
uicontrol uicontextmenu	446	uipanel selectionhighlight	443
uicontrol units	446	uipanel shadowcolor	443
uicontrol userdata	446	uipanel sizechangedfcn	443
uicontrol value	446	uipanel tag	443
uicontrol verticalalignment	446	uipanel title	443
uicontrol visible	446	uipanel titleposition	443
uimenu accelerator	436	uipanel type	443
uimenu beingdeleted	436	uipanel uicontextmenu	443
uimenu busyaction	436	uipanel units	444
uimenu buttondownfcn	436	uipanel userdata	444
uimenu callback	436	uipanel visible	444
uimenu checked	436	uipushtool beingdeleted	450
uimenu children	436	uipushtool busyaction	450
uimenu clipping	436	uipushtool buttondownfcn	450
uimenu createfcn	436	uipushtool cdata	450
uimenu deletefcn	436	uipushtool children	450
uimenu enable	436	uipushtool clickedcallback	450
uimenu foregroundcolor	436	uipushtool clipping	450
uimenu handlevisibility	436	uipushtool createfcn	450
uimenu hittest	436	uipushtool deletefcn	451
uimenu interruptible	437	uipushtool enable	451
uimenu label	437	uipushtool handlevisibility	451
uimenu parent	437	uipushtool hittest	451
uimenu pickableparts	437	uipushtool interruptible	451
uimenu position	437	uipushtool parent	451
uimenu selected	437	uipushtool pickableparts	451
uimenu selectionhighlight	437	uipushtool selected	451
uimenu separator	437	uipushtool selectionhighlight	451
uimenu tag	437	uipushtool separator	451
uimenu type	437	uipushtool tag	451
uimenu uicontextmenu	437	uipushtool tooltipstring	451
uimenu userdata	437	uipushtool type	452
uimenu visible	437	uipushtool uicontextmenu	452
uipanel backgroundcolor	442	uipushtool userdata	452
uipanel beingdeleted	442	uipushtool visible	452

uitable backgroundcolor	446	uitoggletool buttondownfcn	452
uitable beingdeleted	446	uitoggletool cdata	452
uitable busyaction	446	uitoggletool children	452
uitable buttondownfcn	446	uitoggletool clickedcallback	452
uitable celleditcallback	446	uitoggletool clipping	452
uitable cellselectioncallback	447	uitoggletool createfcn	452
uitable children	447	uitoggletool deletefcn	452
uitable clipping	447	uitoggletool enable	452
uitable columneditable	447	uitoggletool handlevisibility	452
uitable columnformat	447	uitoggletool hittest	452
uitable columnname	447	uitoggletool interruptible	453
uitable columnwidth	447	uitoggletool offcallback	453
uitable createfcn	447	uitoggletool oncallback	453
uitable data	447	uitoggletool parent	453
uitable deletefcn	447	uitoggletool pickableparts	453
uitable enable	447	uitoggletool selected	453
uitable extent	447	uitoggletool selectionhighlight	453
uitable fontangle	447	uitoggletool separator	453
uitable fontname	447	uitoggletool state	453
uitable fontsize	447	uitoggletool tag	453
uitable fontunits	447	uitoggletool tooltipstring	453
uitable fontweight	447	uitoggletool type	453
uitable foregroundcolor	447	uitoggletool uicontextmenu	453
uitable handlevisibility	447	uitoggletool userdata	453
uitable hittest	447	uitoggletool visible	453
uitable interruptible	448	uitoolbar beingdeleted	449
uitable keypressfcn	448	uitoolbar busyaction	449
uitable keyreleasefcn	448	uitoolbar buttondownfcn	449
uitable parent	448	uitoolbar children	449
uitable pickableparts	448	uitoolbar clipping	449
uitable position	448	uitoolbar createfcn	449
uitable rearrangeablecolumns	448	uitoolbar deletefcn	449
uitable rowname	448	uitoolbar handlevisibility	449
uitable rowstriping	448	uitoolbar hittest	449
uitable selected	448	uitoolbar interruptible	449
uitable selectionhighlight	448	uitoolbar parent	449
uitable tag	448	uitoolbar pickableparts	450
uitable tooltipstring	448	uitoolbar selected	450
uitable type	448	uitoolbar selectionhighlight	450
uitable uicontextmenu	448	uitoolbar tag	450
uitable units	448	uitoolbar type	450
uitable userdata	448	uitoolbar uicontextmenu	450
uitable visible	448	uitoolbar userdata	450
uitoggletool beingdeleted	452	uitoolbar visible	450
uitoggletool busyaction	452		

